arXiv:2610.17768v1 [cs.IR] 13 Oct 2026

More Reader Data or a Better Teacher? Updating a Model of Crowd Highlights in Late 2026

Kazuki Nakayashiki Keisuke Watanabe
Glasp Inc.
{kazuki,kei}@glasp.co

Abstract

A model trained on readers' highlights can show a newcomer where earlier readers converged on a page, even on pages with few readers. A year of stronger language models suggests two upgrades: use them directly, or use them as teachers. We compare both with a third, more reader data, on one social highlighting platform, with the comparisons, the selection rules and a sealed set of fresh pages fixed in advance. On 215 development pages, seven language models released by October 2026 rank sentences zero-shot at an average precision of at most 0.3744 against the sentences half the readers marked most, and 0.3931 when fused with reading position, while a 3-billion-parameter model trained on reader highlights alone reaches 0.4346. Used as the teacher that breaks ties among readers' counts, the stronger fusion changes a trained model by +0.0013 [−0.0012, +0.0040], whereas going from half to all of its training pages adds +0.0242 [+0.0179, +0.0301]. A 27-billion-parameter backbone on the reader-only recipe adds +0.0225 on the development pages; on 131 fresh pages, read once and too few for a confirmatory comparison, its difference from the 3-billion model is unresolved, both score above the fusion, and in an exploratory split their lead over the previous model sits on pages from domains that their fresh training pages also cover. In this recipe, more reader data improves a trained model where a stronger tie-breaking teacher yields no material gain, and frontier models used as rankers still trail a small model trained on readers.

1 Introduction

Social highlighting platforms let readers see which passages other readers marked. On a page with few readers there is little to show, and a model of where readers converge (crowd salience) can fill the gap. Our earlier work found that such a model is learnable from the page alone [9], that language models agree with one another far more than with readers [11], and that fusing several frontier models with reading position recovers more of the predictable structure than any single model [12]. The best model we had in August 2026 was a 3-billion-parameter language model distilled from such a fusion and then trained on readers' highlights: it ranked sentences well above every model it learned from.

Two months later the frontier had moved. Seven new models were available, and the natural plan is to ask whether they close the gap zero-shot, and if not, whether they make a better teacher for the trained model. A third option needs no new model: more reader data. This paper compares the three on the highlights readers leave in ordinary use, with the comparisons fixed in advance.

This paper makes one claim: for predicting where readers converge, reader data still beats frontier models. As rankers, alone or fused, they trail a small model

trained on reader highlights; and in the recipe we tested, more reader pages improve that model while a stronger frontier teacher, used to break ties among readers' counts, yields no material gain. No single new model clears the earlier reference by the margin we set in advance; the seven fused with reading position do, and still trail a 3-billion-parameter model trained on reader highlights alone by 0.042 average precision. That stronger fusion, used as the teacher that breaks ties, moves the trained model by +0.0013, while going from half to all of the reader pages moves it by +0.0242. A larger backbone scores higher on the pages it was selected on; on fresh pages its advantage could not be confirmed, and in an exploratory split the gains of the new models over the previous one sit on familiar domains.

Related work. Distilling a large model's rankings into a small one is established in retrieval, where a 440-million-parameter student of ChatGPT's rankings outperformed a 3-billion-parameter supervised ranker [8], in the tradition of knowledge distillation [4], and a language model has outperformed crowd workers on several annotation tasks [3]. Readers' highlights carry information about their comprehension and interests [13], and in our earlier work, asking a model which sentences a reader would highlight agreed with readers less than asking it which sentences to keep [10]. What we add is a measured comparison, on naturally accumulated reader behaviour, of the two levers a platform can pull: a stronger teacher, or more of its own readers' data.

2 Setting

Pages, readers and labels. We use public highlights on one social highlighting platform. A page's text is split into sentences and each reader's highlights are mapped onto them; a page enters an evaluation set with at least 6 readers and 30 sentences. Following our earlier protocol [12], the readers of a page are split in half 40 times; on each split the sentences half B marked most (up to round$(0.15N)$ of a page's N sentences, only sentences half B marked, ties broken at random) are the positives, every model's ranking of the page's sentences is scored by average precision (AP) against them, and half A's counts give a split-half reference (how well half the readers predict the other half). A page's AP is the mean over its splits, and a set's AP the mean over its pages. Differences are paired by page with 95% intervals from a bootstrap that resamples domains (4,000 draws) [2]; they reflect the sampling of pages, not variation between training runs. Families of comparisons are Holm-adjusted [6].

Sets. The *development* set is our earlier 217-page evaluation set, re-extracted under the rules below; 215 pages over 84 domains pass the gates. The *training* pages are the 819 pages of the August model with their readers re-extracted under the same rules (801 keep at least one reader). *Fresh* pages are pages that never entered any earlier training or evaluation set, by page identity, URL and text near-duplication. The fresh pages with 3–5 readers form a second training set (3,160 pages); those with at least 6 readers and fetchable public text form the *test* set (131 pages over 42 domains). Near-duplicates of test pages are kept out of the second training set. The test set was sealed when it was written and read once, after the candidates were frozen. Only highlights public under the platform's own rule, recorded before a fixed cutoff, were read.

What was fixed in advance. The comparisons, the margins that decide each of them, the rules that choose the candidates for the test set and the three comparisons on it were registered before any of them was run. At the test set's reader depth the smallest difference it could detect (0.0247) exceeds half the effect it was meant to confirm (0.0179), a condition we fixed in advance, so every number on the test set is reported as a description, not a confirmatory test.

first sample, against `gemini-3.1-pro-preview`	difference [95%]	Holm p
`gpt-6-astra` (low)	+0.0159 [+0.0050, +0.0261]	0.065
`gpt-6-astra` (high)	+0.0138 [+0.0045, +0.0229]	0.050
`gpt-6.1-sol`	+0.0131 [+0.0029, +0.0227]	0.094
`claude-opus-5-5`	+0.0120 [−0.0019, +0.0240]	0.336
`claude-fable-5-1`	+0.0073 [−0.0060, +0.0193]	0.582
`gemini-3.8-flash`	+0.0015 [−0.0082, +0.0097]	0.741
`claude-haiku-5-5`	−0.0180 [−0.0320, −0.0056]	0.050
`gpt-6-luna`	−0.0248 [−0.0370, −0.0109]	0.015
the seven fused with position	+0.0345 [+0.0228, +0.0450]	0.001
against the earlier five-model fusion	+0.0154 [+0.0072, +0.0238]	

Table 1: Zero-shot ranking on the development set (215 pages, one scoring run): each model's first sample against the cached reference. Holm p is the Holm-adjusted permutation p of the difference from zero, over the nine rows. The registered rule: a model or the new fusion passes when its point estimate reaches the margin (+0.02 over the reference; for the new fusion also +0.015 over the earlier fusion) and its 95% interval lies above zero, with a Holm p below 0.05 for the comparison with the reference. The margin bounds the point estimate, not the interval.

3 Frontier models zero-shot

Seven models released by October 2026 (`claude-opus-5-5`, `claude-fable-5-1` and `claude-haiku-5-5` at low reasoning effort, `gpt-6.1-sol` and `gpt-6-astra` at low effort, `gpt-6-luna` without reasoning, `gemini-3.8-flash` at its default) each ranked every sentence of a page, given the numbered sentences and asked for the most important first; we drew two samples per page. A ranking becomes scores by Borda count [1], and a fusion sums the members' normalised Borda scores with reading position, as in [12]. The references are the cached July 2026 draw of `gemini-3.1-pro-preview`, the best single model then [12], and the earlier fusion of five models with position [12].

No single model clears the margin of +0.02 we set in advance (Table 1); the nearest, `gpt-6-astra`, gains +0.0159, and high reasoning effort scores no higher. Two samples of the same model agree at a median per-page Spearman correlation between 0.70 and 0.97, and the mean of two samples adds at most 0.006 AP. Fusion does what it did before [12]: the seven with position beat the reference by +0.0345 and the earlier fusion by +0.0154, both resolved. The fusion reaches 0.3909 AP in this run.

In an exploratory check added after the test set was read, on the development set, a prompt aimed at readers does not help. Asked instead for the sentences a typical reader would highlight (the instruction of [10], adapted to rank every sentence; Appendix A), `gpt-6-astra` changes by −0.0412 [−0.0585, −0.0282] and `claude-opus-5-5` by −0.0363 [−0.0582, −0.0187] (first samples, in a scoring run of their own), as `gpt-5.4` did in [10].

4 Training on readers

Model. A pretrained language model reads the numbered sentences of a page (up to 8,192 tokens) and a two-layer head scores each sentence from the hidden state at its last token; the base is quantised to 4 bits and adapted with LoRA [7, 5]. The target ranks sentences by how

model	training	AP
27B, readers only	Qwen3.8-27B from base; fresh pages, then the 819; position breaks ties	0.4572
27B, teacher breaks ties	Qwen3.8-27B from base; the 819; the five-model fusion breaks ties	0.4465
August model	SmolLM3-3B, distilled start; the 819 (August readers)	0.4349
August recipe, re-run	the same recipe trained again here, on the same GPU as every arm	0.4345
new teacher	distilled start, 100%, the seven-model fusion breaking ties	0.4350
3B, readers only	SmolLM3-3B from base; fresh pages, then the 819	0.4346
distilled start, 100%	SmolLM3-3B, distilled start; the 819, refreshed readers	0.4337
3B, teacher breaks ties	SmolLM3-3B from base; the 819; the five-model fusion breaks ties	0.4240
3B, readers only, 819	SmolLM3-3B from base; the 819; position breaks ties	0.4101
distilled start, 50% / 25%	as 100% on half / a quarter of the pages	0.4095 / 0.3807
the seven fused with position	zero-shot	0.3931
the earlier five-model fusion	zero-shot	0.3765
best single model	`gpt-6-astra`, mean of two samples	0.3744
S3	Qwen3-8B distilled from the five-model fusion [12]	0.3641
reading position	first sentences first	0.2509
split-half reference	half the readers predicting the other half	0.5136

Table 2: Development set, 215 pages, all models in one scoring run; the differences in the text are computed in this run from unrounded per-page values. "Fresh pages" are the 3,160 fresh pages with 3–5 readers.

many readers marked them, and ties, which dominate on pages with few readers, are broken either by a teacher's ranking or by position (earlier first). The loss is an equal mix of squared error on the standardised target and a listwise cross-entropy at temperature 0.1. The *August model* is SmolLM3-3B first distilled from the earlier five-model fusion and then trained on the 819 training pages with that fusion breaking ties; it comes from an unpublished comparison of backbones and is reported here for the first time. Every model trained for this paper is trained twice with different seeds and reported as the mean of the two; the August model and the other backbones of its comparison are single earlier runs.

More pages help; refreshed labels and a stronger tie-breaking teacher add little. Table 2 gives every arm, and every comparison on the development set below is computed in that one scoring run (the registered decisions and the freeze of Section 5 were taken in earlier runs and hold in this one). On nested quarters, halves and all of the 819 pages, the distilled start gains +0.0242 [+0.0179, +0.0301] from the half to the whole (Figure 1). Re-extracting the readers of the same pages changes the trained model by −0.0008 [−0.0030, +0.0015] against the August recipe re-run, a change that includes losing the 18 pages left without a reader. The seven-model fusion, which ranks pages +0.0166 better than the earlier fusion in this run, was collected for

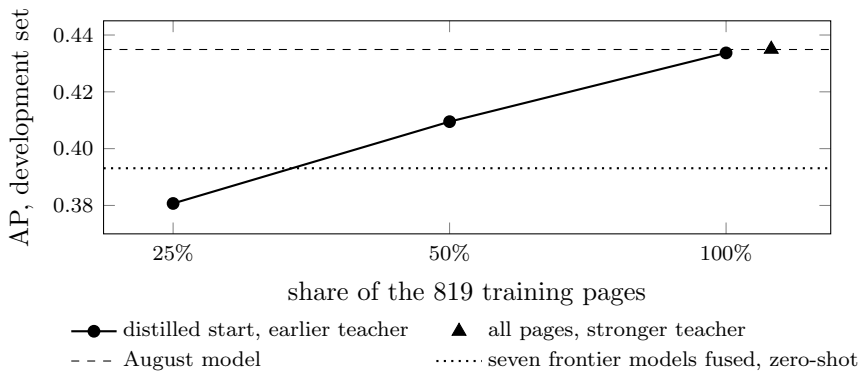

Figure 1: More pages against a stronger tie-breaking teacher on the development set (the run of Table 2): the distilled start trained on a quarter, half and all of the 819 training pages with the earlier teacher breaking ties, and on all of them with the stronger teacher (the triangle, drawn right of 100% for legibility).

all 819 training pages (each model's first sample, fused with position) and used to break ties in an otherwise identical run: the model moves by +0.0013 [−0.0012, +0.0040], below the +0.005 we required, and by +0.0013 [−0.0017, +0.0044] and +0.0013 [−0.0012, +0.0039] with either seed alone. The swap changes the target's order for 8% of the sentence pairs on these pages: the target orders sentences by readers' counts, so a teacher orders only pairs with equal counts (54% of them), and the two teachers order 85% of those alike; 86% of the pairs that change are of two sentences no reader marked. A teacher is not useless: from the base model on the 819 pages, breaking ties with the five-model fusion instead of position adds +0.0139 [+0.0037, +0.0235]. But once a teacher breaks the ties, a stronger one adds little next to more reader pages.

A model trained on readers alone comes close to the distilled one. Trained from the base model on reader highlights with position breaking ties, SmolLM3-3B reaches 0.4101 on the 819 pages and 0.4346 when the 3,160 fresh pages come first (+0.0245 [+0.0100, +0.0472]), a point estimate close to the August model's (−0.0003 [−0.0219, +0.0287]; the difference is unresolved). No output of a frontier model enters it. An exploratory control on the development set, added after the test set was read, shows that the gain comes from the pages, not from the 6,320 steps that come with them: trained as long on the 819 pages alone, the model changes by −0.0201 [−0.0273, −0.0134] against the shorter run, and the fresh pages lead that control by +0.0446 [+0.0299, +0.0699] (a scoring run of their own).

A larger backbone scores higher here. The same recipes on Qwen3.8-27B, of another family and a newer generation than SmolLM3-3B, score +0.0225 [+0.0153, +0.0308] above SmolLM3-3B with the five-model fusion breaking ties, and +0.0225 [+0.0155, +0.0295] above it with readers only; the 27B reader-only model reaches 0.4572, against the August model's 0.4349 (+0.0222 [−0.0003, +0.0493]). These are the pages on which the candidates were selected.

Smaller backbones (exploratory). After the test set was read, we trained the reader-only recipe on Qwen3-0.6B and Qwen3-1.7B and scored them in the same run, next to the other backbones of the comparison the August model came from (Table 3). At both sizes the reader-only model's point estimate is close to that of the distilled model of the same size (+0.0020 [−0.0180, +0.0307] at 0.6B, +0.0034 [−0.0165, +0.0323] at 1.7B; unresolved), as at 3B. The score rises with size: the 0.6B reader-only model trails the 3B one (−0.0189 [−0.0252, −0.0107]), yet its

backbone	August recipe	readers only
ModernBERT-large (0.4B, encoder)	0.4043	—
Qwen3-0.6B	0.4138	0.4158
Qwen3-1.7B	0.4234	0.4268
SmolLM3-3B	0.4349	0.4346
Qwen3-4B	0.4318	—
Qwen3-8B	0.4371	—
Qwen3.8-27B	—	0.4572
the seven fused with position (zero-shot)	0.3931	

Table 3: Backbone size on the development set (exploratory; AP in the run of Table 2). *August recipe*: distilled from the five-model fusion, then trained on the 819 pages with that fusion breaking ties; one run each, from the comparison the August model (the SmolLM3-3B row) came from. *Readers only*: the recipe of the 3B reader-only model, mean of two seeds; the 0.6B and 1.7B models were trained after the test set was read and were not scored on it.

model	AP	comparison	difference [98.3%]
3B, readers only	0.3663	C1: August model − S3	+0.0420 [+0.0190, +0.0671]
27B, readers only	0.3632	C2: 27B − August model	+0.0377 [−0.0036, +0.1120]
August model	0.3255	C3: 27B − the fusion	+0.0603 [+0.0164, +0.1276]
the seven fused with position	0.3029		
S3	0.2834		
reading position	0.2086		
random (expected)	0.1702		

Table 4: The sealed test set: 131 fresh pages over 42 domains, one scoring run, intervals Bonferroni-adjusted over the three registered comparisons. Descriptive: the set is smaller than a confirmatory comparison needed.

point estimate still exceeds the seven-model fusion's (+0.0227 [−0.0010, +0.0512]), as does that of the 0.4B encoder on the August recipe (+0.0112 [−0.0080, +0.0313]). Neither new model was a candidate, and neither was scored on the test set.

5 Fresh pages

We froze two candidates by the registered rules (the best 3B reader-only model and the best model overall, the 27B reader-only model, which is also the larger size) and the best zero-shot configuration (the seven-model fusion, collected on the test pages with one sample per model), and scored them once on the sealed test set with the August model and S3. The checks of the scoring held: a random ranking sits at its expectation (+0.0042 [−0.00003, +0.0115]), a permuted ranking at random (−0.0006 [−0.0137, +0.0078]), and S3 above position (+0.0748 [+0.0550, +0.0927]).

The August model keeps its lead over S3 on fresh pages (C1 in Table 4), on pages with 6–9 readers (+0.0422 [+0.0194, +0.0707], 107 pages) and with 10–19 (+0.0429 [+0.0021, +0.0862], 22 pages). The 27B model leads the seven-model fusion (C3), and so does the 3B (+0.0634 [+0.0095, +0.1632], an unregistered comparison). The 27B's lead over the August model (C2) does not exclude zero, and its difference from the 3B is unresolved (−0.0031 [−0.0372, +0.0160],

an unregistered comparison). The split-half reference (0.3275) sits below the two reader-only candidates because most test pages have 6–7 readers, so three readers predict three or four.

Exploratory, after the look. 120 of the 131 test pages come from domains that also hold a fresh training page. On those, the 3B and 27B reader-only models lead the August model by +0.0496 [+0.0020, +0.1377] and +0.0445 [+0.0066, +0.1058]; on the other 11 pages, from 11 domains (10 of them absent from the 819 training pages too), they trail it by −0.0545 [−0.1095, −0.0148] and −0.0362 [−0.0804, −0.0028] (95% intervals). On the development set the 3B's and the 27B's differences from the August model are +0.0103 [−0.0241, +0.0653] and +0.0337 [−0.0010, +0.0823] on the 125 pages of such domains, and −0.0150 [−0.0342, −0.0002] and +0.0062 [−0.0162, +0.0236] on the other 90. Eleven pages cannot settle it, but the gain of the new models may be coverage of the domains a platform's readers keep reading rather than a better general ranker.

6 Discussion

Frontier models rank sentences by what a model finds important, and readers converge on something else [11]. A newer generation narrows that gap only a little: the best single model of October 2026 stays 0.060 AP below a 3B model trained on readers, and fusing seven of them still leaves 0.042. The trained model's advantage comes from the readers, and the parts of its recipe that rely on a frontier model matter less than they seemed. A teacher that breaks ties helps a model trained on few reader pages, but a stronger one added little once a teacher was there, and with 3,160 more reader pages a model trained from its base model on readers alone reaches the distilled model's point estimate. For a platform this is the useful direction: the improvement comes from data it already collects, needs no model provider's outputs in training, and stays with a small model (Table 3).

The larger backbone is the open question. On the development set it adds as much as doubling the pages did; on fresh pages its advantage could not be confirmed, and the exploratory split points at domain coverage as a common source of the new models' gains there. A test on pages from domains that no training page covers would separate a better ranker from a familiar one.

Limits. One platform, public pages whose text could be fetched, and Latin-script pages only in the fresh training and test sets. The test set is small, was read as a description, and tests the frozen candidates, not the development-set comparisons behind them. The stronger teacher replaced only the tie-breaker: the distilled start kept the earlier teacher, so distilling anew from the stronger one, or giving the fresh pages to the distilled model, was not tested. The zero-shot models are specific versions at specific settings, which date quickly; a reader-prediction instruction lowered the two we tried, but other prompts, examples or fine-tuning might narrow the gap. The trained models were selected on the development set, so their numbers there favour them.

7 Conclusion

Seven language models released by October 2026 do not close the gap to a small model trained on reader highlights, alone or fused, and the stronger fusion, as the teacher that breaks ties, did not improve it by the margin we required. More reader pages did, and a model trained on reader highlights alone now reaches a point estimate close to one distilled from frontier models. On this

platform and in the recipe we tested, the data that improves a model of where readers converge is the readers'.

Ethics and data

This is a retrospective, non-interventional secondary analysis of service records collected under the platform's Terms of Service and Privacy Policy, which name research among the purposes of processing. The study read production records read-only, kept only highlight documents public under the platform's own rule (public is the default), and wrote records without account identifiers or reader keys. Page text came from the platform's stored copies or, for fresh pages, from the public web without credentials; model providers received public page text only, and the training service received page text and highlight positions keyed by ordinals. Record-level files are not released, and every reported result is an aggregate over pages with at least six readers. The operating company is not a research institution and has no institutional review board; no institutional review board reviewed this study, and no formal ethics approval was sought. Both authors are affiliated with the platform's operator.

AI assistance

The study design, the extraction, training and analysis code, the statistics, the tables and a first draft of this paper were produced with an AI coding agent (Anthropic's Claude) under the authors' direction. OpenAI's Codex and a separate Claude instance audited the design, the code and the write-ups. The authors take full responsibility for the content.

Data and code availability

Record-level data cannot be released. The ancillary files contain the pre-registration with its amendments, the results write-up, every aggregate artifact, and the code that builds the prompts, scores the rankings and builds the training targets. Every number in this paper regenerates from those artifacts (`make-numbers.py`), and `check-claims.py` checks the paper's listed qualitative comparisons against them; retraining a model needs the record-level data, so the training and scoring code is there as specification.

References

[1] J. A. Aslam and M. Montague. Models for metasearch. In *Proc. SIGIR*, pages 276–284, 2001.

[2] A. C. Cameron, J. B. Gelbach, and D. L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008.

[3] F. Gilardi, M. Alizadeh, and M. Kubli. ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 120(30):e2305016120, 2023.

[4] G. Hinton, O. Vinyals, and J. Dean. Distilling the knowledge in a neural network. arXiv:1503.02531, 2015.

[5] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems*, 2023.

[6] S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

[7] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. LoRA: Low-rank adaptation of large language models. In *Proc. ICLR*, 2022.

[8] W. Sun, L. Yan, X. Ma, S. Wang, P. Ren, Z. Chen, D. Yin, and Z. Ren. Is ChatGPT good at search? Investigating large language models as re-ranking agents. In *Proc. EMNLP*, 2023.

[9] K. Nakayashiki and K. Watanabe. The long tail, not the front page: Cold-start prediction of crowd highlight salience. arXiv:2606.11654, 2026.

[10] K. Nakayashiki and K. Watanabe. Measuring alignment with reader highlights net of position and length. arXiv:2607.27739, 2026.

[11] K. Nakayashiki and K. Watanabe. Language models agree with each other, not with readers. arXiv:2607.29274, 2026.

[12] K. Nakayashiki and K. Watanabe. Floor, ceiling, and the fusion gap: How much of crowd reading attention can machines predict? arXiv:2608.01704, 2026.

[13] A. Winchell, A. Lan, and M. Mozer. Highlights as an early predictor of student comprehension and interests. *Cognitive Science*, 44(11):e12901, 2020.

A Reproduction details

Zero-shot prompt. Every model received the system message "You compress documents by selecting the most important sentences. Output only what is asked for." and a user message that lists the page's sentences, one per line as "[i] sentence" after "Sentences from one article, in order:", and ends "You must compress this document by keeping only the most important sentences. Rank ALL N sentence numbers from most important to least important. Return ONLY a JSON array of every sentence number, most important first." It is the prompt of [12], kept byte-identical so that the new models compare with its cached draws. The reader-prediction instruction (Section 3) takes the system message of [10], "You select the sentences a reader would most likely highlight or underline while reading. Readers highlight memorable, surprising, quotable, or personally striking passages — not merely topic sentences or definitions. Output only JSON.", and after the same list of sentences asks "Rank ALL N sentence numbers from the one a typical reader would most likely highlight to the one least likely. Return ONLY a JSON array of every sentence number, most likely first."

Models and settings (collected on 2026-10-08 and 2026-10-09). `gpt-6.1-sol` and `gpt-6-astra` through OpenAI's API at reasoning effort low (`gpt-6-astra` also at high, one sample); `gpt-6-luna` at effort none and temperature 0; `anthropic/claude-opus-5.5`, `anthropic/claude-fable-5.1` and `anthropic/claude-haiku-5.5` at effort low, pinned to Anthropic, and `google/gemini-3.8-flash` at its default, through OpenRouter. No other temperature was sent; the output cap was 32,000 tokens (64,000 at high effort). Batch and synchronous calls carried identical requests. A ranking that leaves sentences out keeps the ranked ones above the rest; a refusal, or a filtered, truncated or unparseable answer, is scored as a tie of all sentences (7 of the 1,505 first samples on the development set).

Training. The base model is quantised to 4-bit NF4 and adapted with LoRA of rank 16 ($\alpha = 32$, dropout 0.05) on the attention and MLP projections [7, 5]; the head is two linear layers with 256 hidden units and GELU. Each sentence is written as "i: sentence" and a page is cut

at 8,192 tokens; sentences past the cut score below every scored one, in document order. The target of sentence i is the rank fraction, within the page, of $c_i + t_i$, where c_i is its reader count and $t_i \in [0, 1)$ the rank fraction of the teacher's score or of position, so the tie-breaker orders sentences only within the same count. The loss is 0.5 times the squared error on the standardised target plus 0.5 times the Kullback–Leibler divergence between the softmax of the target and that of the prediction at temperature 0.1. AdamW at 2×10^{-5} when continuing the distilled start and 10^{-4} for a fresh adapter; gradients are accumulated over 8 pages. A run takes a fixed budget of 2,457 page steps on the training pages (306 updates, 304 on the quarter of the pages; August's three passes over its 819 pages), after 6,320 steps on the fresh pages for the reader-only recipe (two passes, 790 updates). The page order is shuffled once per seed, the model is the last step (validation is recorded, not used to choose), and every arm trained here has two seeds, on one NVIDIA H100. Qwen3-0.6B and Qwen3-1.7B (Table 3) use the settings of the 3B reader-only model. The equal-steps control (Section 4) is the 3B reader-only model's command with the first stage reading the 819 training pages instead of the fresh ones. The August-recipe rows of Table 3 are the stored sentence scores of the earlier comparison's runs, one each, made with the August model's two stages; the 8B's distilled start is S3, and ModernBERT-large, an encoder, was fully fine-tuned, from a distilled start trained for eight passes at 10^{-4}.

Pages and readers. A reader's highlights are mapped onto the sentences they cover, and a sentence's count is the number of readers who marked it. Of the 819 training pages, 18 kept no reader whose highlights still anchor to the stored text after re-extraction, so the runs on the re-extracted readers use the other 801 (a subset, the part of them it holds), while the August recipe re-run uses the August readers of all 819; the teachers were collected for all 819.

Registration. The pre-registration was fixed on 2026-10-08, before any record was read; its dated amendments, the candidates frozen on 2026-10-09 before the test set was opened, and the results write-up are in the ancillary files, with the code that builds the prompts, scores the rankings and regenerates every number.