arXiv:2610.19799v1 [cs.CL] 14 Oct 2026

Evolve Together, Fall into Place: How Social Scaffolding Shapes Self-Evolving Multi-Agent Systems

Seojin Hwang
Chung-Ang University
swiftie1230@cau.ac.kr

Youngmin Choi
Chung-Ang University
martinus.choi@gmail.com

Hwanhee Lee*
Chung-Ang University
hwanheelee@cau.ac.kr

Abstract

Recent work lets LLM agents evolve their own harness, up to the team itself. Yet in human groups, the same abilities achieve different outcomes under different social conditions. Borrowing the notion of scaffolding from developmental psychology, we investigate whether the *social scaffolding* of a self-evolving multi-agent system (MAS), namely its *team composition* and *structure of interdependence* (e.g., collaboration, conformity, or competition), shapes how the team evolves. We present **AGORA** (**A**gents **G**rowing through **O**rganization and **R**elational **A**rrangements), a framework that varies the two axes independently across three evolution mechanisms. Across web search, codebase understanding, and document QA, we find that (1) evolution alone rarely helps, while the right scaffolding surpasses a strong agent evolved alone by up to 25 points; (2) it does so by moving authority to the most capable member while the others step back; and (3) this member, having grown within the team, outperforms the same model alone. Thus, designing a self-evolving MAS requires designing its social scaffolding, which shapes not only how its agents learn but also who gets to lead.

1 Introduction

As Large Language Model (LLM) agents take on increasingly complex tasks, they rarely work as bare models. They run inside a harness that shapes their performance (Meng et al., 2026). Prior work has focused on its technical design, from what a single agent sees and is allowed to do (Lee et al., 2026; Lou et al., 2026) to how LLM calls are composed into workflows (Zhang et al., 2025). The same holds for multi-agent systems (MAS) (Liu et al., 2026b), and a growing line of work goes further and lets agents evolve their own harness, by rewriting their skills, memories, and even the team itself (Hao et al., 2026; Yang et al., 2026; Zhang et al., 2026c).

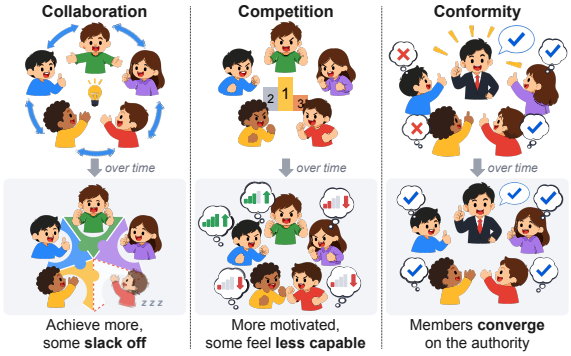

Figure 1: How different social scaffolding leads to different evolutionary outcomes in human groups.

However, as shown in Figure 1, evolution in human groups never happens in a social vacuum. Organizational and evolutionary theories suggest that people with the same abilities can achieve and behave differently, depending on how they compare with those around them and how much they rely on one another (Deutsch, 1949; Johnson and Johnson, 2009). Much like a human group, a self-evolving MAS consists of agents that differ in ability, depend on one another, and change through experience. This raises a natural question: *does the social layer of its harness shape how the team evolves, as social conditions do for human groups?*

Prior work on self-evolving MAS either fixes this social layer as a design (Liu et al., 2026a; Chen et al., 2025; Pan et al., 2025) or folds it into the evolution process itself (Hao et al., 2026). We instead treat this social layer as the central variable of study. We borrow the developmental-psychology notion of scaffolding, the support a learner's social environment provides during development and withdraws once it is no longer needed (Wood et al., 1976), and call this layer ***social scaffolding***. We operationalize it for MAS along the two factors above: *team composition*, which sets how capable each agent is relative to its teammates, and the *structure*

*Corresponding author.

of interdependence, which sets how agents share information and decide on team changes.

To isolate its effect on what a team evolves into, we build **AGORA** (**A**gents **G**rowing through **O**rganization and **R**elational **A**rrangements), a framework that extends the three-level evolution loop of Meta-Team (Hao et al., 2026). After each batch of training tasks, every agent improves itself (L1), agents take notes on their teammates when a structure of interdependence is in place (L2), and the team revises itself as a whole (L3), with each team revision checked on a held-out set.

AGORA varies the two axes independently: team composition sets the mix of agent capabilities (onlyweak, onlystrong, or weakstrong), and the structure of interdependence (none, collaboration, conformity, or competition) sets which permissions agents hold, how they share information at L2, and how team changes are decided at L3. We cross these axes with three evolution mechanisms that govern L1, Meta-Team (Hao et al., 2026), SkillOpt (Yang et al., 2026), and MemRL (Zhang et al., 2026c), holding everything else fixed so that differences can be traced to social scaffolding.

Using AGORA on three benchmarks spanning web search (GAIA), codebase understanding (LoCoBench), and document QA (OfficeQA Pro), we investigate three research questions: whether social scaffolding changes what a self-evolving MAS achieves (RQ1); how the team reorganizes under the right scaffolding (RQ2); and why the reorganized team outperforms the same model working alone (RQ3). Experiments yield three findings.

- **Evolution alone is rarely enough (RQ1).** Evolving under no structure of interdependence rarely improves a team, whereas the right social scaffolding does, surpassing a strong agent evolved alone by up to 25.0 points. Which structure helps depends on who is in the team.
- **Authority moves to the most capable member (RQ2).** Effective scaffolding hands the chair to the most capable member, as human groups gain from recognizing expertise, and the others step back unless they remove a bottleneck. Unlike human groups, they drop out rather than loaf.
- **The team grows its leader (RQ3).** The member left leading, having learned from its teammates, outperforms the same model evolved alone.

These findings show that plugging an evolution mechanism into an off-the-shelf MAS is not enough; the social scaffolding under which the team evolves matters just as much, deciding not only how its agents learn, but who gets to lead.

2 Related Work

2.1 Social Dynamics in Human Groups

In human groups, individuals with identical abilities can reach different outcomes depending on their social scaffolding. *Collaboration* raises achievement over competing or working alone (Deutsch, 1949; Johnson and Johnson, 2009), but members may loaf when the work is shared (Latané et al., 1979; Karau and Williams, 1993). *Competition* raises motivation, but equally able members feel less capable among stronger peers (Marsh, 1987). *Conformity* reduces individual error but can suppress innovation (Henrich and Boyd, 1998), and since people obey even wrong authority (Milgram, 1963) and groups that suppress dissent decide poorly (Janis, 1972), the outcome depends on the decision-maker.

2.2 Evolution and Social Dynamics in LLM Agents

Evolving agents improve their external state without updating model weights, refining memory (Zhu et al., 2026b; Choi et al., 2026b; Zhang et al., 2026c; Wang et al., 2026; Liu et al., 2026c) or skills (Zhang et al., 2026b; Zhu et al., 2026a; Zhang et al., 2026a; Yang et al., 2026), and recently agent composition and interaction (Yuan et al., 2025; Lee et al., 2025; Dang et al., 2026). MAS gain more over single agents when their sub-agents are weaker (Ke et al., 2026), suggesting that well-organized weak agents can close much of the gap to stronger ones.

In MAS, collaboration does not always help (Tran and Kiela, 2026; Pappu et al., 2026), competition can increase systemic failure (Mori and Johnson, 2026), and conformity varies with group structure and roles (Weng et al., 2025; Cho et al., 2025; Choi et al., 2025; Bito et al., 2026; Han et al., 2026; Bellina et al., 2026; Choi et al., 2026a). Social structures are also used to improve agents through self-play, mutual checks (Liu et al., 2026a; Chen et al., 2025; Pan et al., 2025), or orchestration (Qi et al., 2026), and Meta-Team (Hao et al., 2026) jointly evolves the agents, their interactions, and the team. These methods, however, fix a single interdependence structure, often on closed-answer tasks. AGORA instead turns this social layer into a controlled variable, varying team composition and

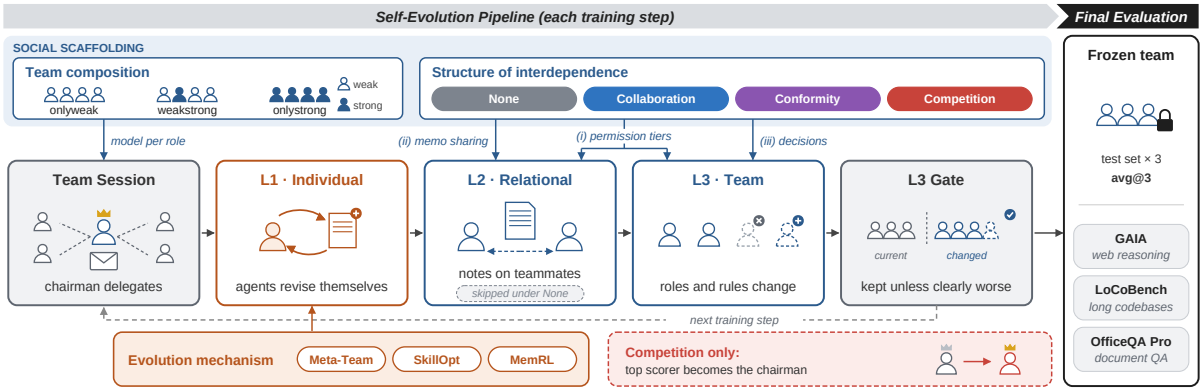

Figure 2: Overview of AGORA. After completing a task as a team (*Team Session*), each agent updates itself (*L1*) and writes memos about its peers, guided by social scaffolding (*L2*). Based on these memos, team-level structural changes are proposed (*L3*) and then retained or rejected by a held-out gate (*L3 Gate*). Throughout this process, we vary the social scaffolding and the evolution mechanism to examine how a self-evolving MAS evolves.

the structure of interdependence on open-ended, real-world tasks.

3 AGORA Framework

As shown in Figure 2, to test whether social scaffolding shapes how a self-evolving MAS evolves, we design AGORA as a controlled testbed. AGORA keeps the basic loop of Meta-Team (Hao et al., 2026), in which a chairman-led team evolves its external state at the individual (L1), relational (L2), and team (L3) levels, but turns its fixed social layer, where the chairman alone decides team changes, into the variable under study. AGORA varies team composition and the structure of interdependence while swapping in different L1 evolution mechanisms, separating their effects.

3.1 Overall Framework Formulation

3.1.1 Team Definition

A team is a set of role-specific agents, defined by a team configuration (e.g., chairman, team budgets), a constitution of team-wide rules that can be revised as the team evolves, and the agents themselves. Each agent has a configuration (e.g., model and tools), a prompt, and an evolution state such as a skill document or a memory bank. The chairman receives the task, calls on the other agents, and submits the final answer.

3.1.2 Session Execution

Each task is handled in a single team session, whose final answer from the chairman is scored. Within a session, agents operate concurrently rather than taking turns, interleaving tool calls (e.g., web search) with messages through a shared mailbox, under a budget on messages, and time that is reported after every call. Detailed limits are given in Appendices B and F.

3.1.3 Three-Level Operation

Under a given social scaffolding, AGORA repeats an evolution loop of a team session followed by three levels of evolution and a held-out gate (Figure 2 and Table 1). L1 improves each agent through its own evolution mechanism, L2 records what agents learn about one another, L3 revises the team as a whole, and the L3 gate reverts a revision only if it clearly hurts the team on held-out data.

L1 evolution (Individual) After completing a task as a team (Figure 2, *'Team Session'*), each agent evolves through its own evolution mechanism (Figure 2, *'L1 · Individual'*). Specifically, it updates a skill with its gate record in SkillOpt, a memory bank of Q-valued entries in MemRL, and a prompt patch with self-written skills in Meta-Team. Unlike L2 and L3, L1 operates in the same way under every social scaffolding; the scaffolding affects it only indirectly, through the trajectories and L2 memos from which the agents learn.

L2 recording (Relational) L2 models what teammates learn about one another through social interaction (Figure 2, *'L2 · Relational'*). Once L1 has finished, each agent writes a memo on every teammate it has worked with, recording strengths, weaknesses, collaboration notes, and contradictions; a judge model flags contradictions but does not block the memo. How much agents share and whether memos are reviewed before saving depend on the *structure of interdependence* (Section 3.3.2), and L2 is skipped under *none*. Memos enter the au-

Level	What changes	Who changes it	Takes effect	Gate
L1 (individual)	Skill document	Mechanism (SkillOpt)	Next iteration	L1 gate
	Memory bank	Mechanism (MemRL)	Next iteration	None
	Prompt patch	Agent itself (Meta-Team)	Next task	None
L2 (relational)	L2 memo	Each agent, about teammates it worked with	Next iteration	None
L3 (team)	Roles, agent prompts and configurations, constitution	Proposed by Tier 1 agents	Next iteration, unless reverted by the L3 gate	L3 gate

Table 1: Levels of the AGORA operation. L1 and L3 evolve the agents and the team, while L2 records relations between agents only when the social scaffolding includes interdependence. Social scaffolding acts directly only on L2 and L3, and reaches L1 through the trajectories and L2 memos that every mechanism learns from.

thor's system prompt in later sessions, inform L3 revisions, and are passed to the update models of SkillOpt and MemRL.

L3 evolution (Team) After each agent has written memos on its teammates in L2, L3 evolves the team as a whole based on these memos, by adding or removing roles and revising the agents' prompts and configurations, and the constitution. Tier 1 agents propose revisions, which are aggregated according to the social scaffolding and kept unless the L3 gate judges the revised team to be clearly worse (Figure 2, *'L3 Gate'*). Details of Tier 1 are provided in Section 3.3.2.

3.1.4 L3 Held-out Gate

The L3 held-out gate reverts a team revision only if it clearly harms the team. The current and revised teams each solve the same 20 validation problems with the same seeds, so that the two are compared on equal footing (Figure 2, *'L3 Gate'*). The validation set is disjoint from the train and test sets, and the gate is applied identically under every evolution mechanism and structure of interdependence.

All L3 changes of an iteration are reverted if the revised team succeeds less often and a one-sided binomial test against the current team's success rate finds the drop significant. For example, if the current team solves 10 of the 20 problems, the changes are reverted only if the revised team solves 8 or fewer. This is because L3 revisions are decisions the team made under its own structure of interdependence; it blocks only clear harm and leaves the rest to the team. We use a relaxed $\alpha = 0.3$ because re-running the same team changes the outcome of about a third of the validation problems, so a strict level would rarely catch even real degradation. Details about thresholds are in Appendix D.5. Each iteration also applies at most two role additions or deletions and two patch bundles, so that each gate judges only a few changes and the team cannot grow without bound.

3.2 Evolution Mechanism

As shown in Table 2, we include three L1 mechanisms that differ in what they update and what drives the update: an agent's own reflection (Meta-Team), reflection on graded trajectories (SkillOpt), and task reward on a memory bank (MemRL). We adapt the original implementations to the multi-agent setting with minimal changes. The mechanism determines only L1. Details and all changes from the original are in Appendix D.

Meta-Team Following the original design, Meta-Team runs a single iteration with its batch executed sequentially. Right after each task, each agent reflects on its own behavior and edits its own prompt patch and skills with its own model, and the edits take effect immediately without a gate.

SkillOpt SkillOpt runs for four iterations. At the end of each iteration, each role applies reflect, merge, and apply steps to its trajectories from that iteration, both successful and failed, together with their grading results, to produce a candidate skill document. The candidate replaces the current document in the role's system prompt only if it passes the L1 gate, that is, only if it clearly improves performance on the validation set (Appendix D.5).

MemRL MemRL runs for four iterations, with a separate memory bank for each role. Before a session, up to five memories are retrieved by combining their similarity to the problem with their Q-values. At the end of each iteration, the Q-values of the retrieved memories are updated according to task success, and each trajectory, successful or not, is summarized as a procedural script and added as a new memory. Updates take effect without a gate, so the bank keeps both what worked and what failed.

	Meta-Team	SkillOpt	MemRL
What evolves	Prompt patches and skills written by the agent	Skill document	Memory bank with Q-values
Injected into execution	Own patches and skill list	Current skill document	Up to 5 retrieved memories (success/failure scripts)
Evolution timing	Right after each task	End of iteration	End of iteration
Evolution signal	Agent's self-reflection	Reflection on trajectories with grading results	Task success or failure as reward
L1 gate	None	Held-out binomial test	None
Model performing the update	Agent's own model (weak or strong)	The strong model (fixed), with L2 notes	The strong model (fixed), with L2 notes
Steps × batch	$1 \times (17/20)$	$4 \times (17/20)$	$4 \times (17/20)$
Batch execution	Sequential	Concurrent	Concurrent

Table 2: Evolution mechanisms. Batch sizes are 17 for GAIA and 20 for LoCoBench and OfficeQA Pro. All mechanisms share the same L2 and L3 and the same L3 gate.

3.3 Social Scaffolding Implementation

In human teams, how members grow depends both on the capabilities of their teammates, including the gaps between them, and on how they depend on one another while working together. Following this, we implement social scaffolding along two axes: team composition and the structure of interdependence (Figure 2, top).

3.3.1 Team Composition

To realize capability differences among team members, we use two model tiers, weak and strong, drawn from the same model family so that the capability gap remains the main variable and confounds from mixing model families are avoided. As shown in '*Team composition*' in Figure 2, we compare three configurations: **onlyweak**, **onlystrong**, and **weakstrong**. Onlyweak and onlystrong assign the weak and strong model, respectively, to every role. Weakstrong upgrades only the bottleneck role of each benchmark to the strong model (web_agent for GAIA, developer for LoCoBench, and reader for OfficeQA), while all other roles, including the chairman, remain weak.

On GAIA, varying which and how many roles are strong shows that which role matters far more than how many: before evolution, a strong bottleneck role alone gives 48.3, a strong non-bottleneck role at most 28.3, and further strong roles barely help (51.7–52.5). Details on model selection, bottleneck roles, role tools, and these variants are in Appendices B and K.1, and representative role prompts and the full constitution in Appendix L.

3.3.2 Structure of Interdependence

To realize how members depend on one another, we define four conditions: **none**, **collaboration**, **conformity**, and **competition**. As shown in Table 3, each condition differs not only in prompt wording but in tools, permissions, and decision rules, and intervenes at three corresponding points of the loop, marked (i)–(iii) in Figure 2. The three points are (i) *permission tiers*, which set who is Tier 1 and can read the whole team's memos, while Tier 2 agents read only their own; (ii) *L2 memos*, which set how much agents share and whether memos are saved directly or staged for objection before finalized; and (iii) *L3 decisions*, which set who can propose, whether proposals go through peer review, and who makes the end-of-iteration decision by what rule.

None. A baseline without social interaction: all agents are Tier 1, there is no L2 (skipped under *None* in Figure 2), and all L3 proposals are kept at the end-of-iteration decision.

Collaboration. All agents are Tier 1 and share actively, and both memos and L3 proposals go through peer review. A memo is finalized once every teammate the author interacted with in the task has responded, and an L3 proposal passes once a majority of the participating teammates have commented on it. The end-of-iteration decision tallies the SUPPORT and OPPOSE reviews.

Conformity. Only the chairman is Tier 1 and decides alone, and agents have no tool for raising objections. In L2, agents defer to their teammates' self-assessments, contact teammates only when stuck, and save memos directly.

Intervention	Aspect	None	Collaboration	Conformity	Competition
Permission tiers	Chairman	Fixed	Fixed	Fixed	Rotates to top scorer
	Tier 1 roles	All roles	All roles	Chairman only	Chairman only
L2 memo	Sharing	Disabled	Active	When stuck; deferential	Minimal; scores shown
	Finalization	Disabled	After peer review	Immediately	After peer review
L3 team config	Proposals	All roles	All roles	Chairman only	Chairman only
	Peer review	No	Majority	No	Majority
	Decision rule	Keep all	Review tally	Chairman alone	Chairman, given reviews

Table 3: Structures of interdependence. 'Majority' peer review passes once a majority of participating teammates have responded, and Tier 2 roles can only suggest L3 changes to the chairman. Details are given in Section 3.3.2.

Competition. Agents share minimally and are shown comparative scores of their own contribution. Memos and L3 proposals go through the same peer review as in collaboration, with objections citing these scores, but the chairman decides alone at the end of each iteration. The role with the sole top score then becomes Tier 1 and the chairman for the next iteration, as shown by *'Competition only'* in Figure 2. This rotation is applied after the L3 gate, outside L3, and is not itself gated.

Shared settings. In every condition, the end-of-iteration decision is made by the chairman's own model with a fixed decision prompt (see Appendix L for the prompt), so no model outside the team decides for it.[1] Gate sessions run under the same condition, tiers, and chairman as training, and all conditions, including none, share the same L3 gate (Figure 2, *'L3 Gate'*) and revision cap, so that only the social layer differs. Appendix E gives the full specification of each condition.

4 Experimental Setup

We cross every evolution mechanism with every combination of team composition and interdependence structure on three benchmarks, comparing each evolved team with its unevolved self and a single agent of each tier (RQ1, Section 5), and then trace its final-evaluation sessions to see how it reorganizes and why (RQ2 and RQ3, Section 6).

Benchmarks We use web-based reasoning (GAIA; Mialon et al., 2024), long codebase understanding (LoCoBench; Qiu et al., 2025), and document-based numerical question answering (OfficeQA Pro; Opsahl-Ong et al., 2026), each split into non-overlapping train, val, and test sets (split sizes and task formats in Appendix A).

[1]The only components outside the team are the L2 contradiction judge and the SkillOpt and MemRL update models, all fixed to the strong model.

Models The weak and strong tiers are *Qwen3-30B-A3B-Instruct-2507* and *Qwen3.6-35B-A3B*, run with thinking off and temperature 0.1 (model selection in Appendix B; serving in Appendix G).

Session Budgets Each session is limited to 900 seconds and 100 team messages (50 for solo), and each agent to 50 steps. A session terminated by its time or message limit is not graded and counts as a failure (0) in every score, including both gates; only sessions that crash are excluded.

Experiment Matrix We use two kinds of baselines. Solo baselines give the whole task to one weak or strong agent with every tool. Preevolve baselines evaluate each of the five configurations (two solo setups and three team compositions) on the test set without evolution. Each benchmark thus comprises 47 experiments: five preevolve baselines, six solo runs, and 36 MAS runs.

Evaluation Tasks are scored as success or failure: exact match for GAIA, an LCBS score of at least 3.0 for LoCoBench, and a verdict from *gpt-5-mini* for OfficeQA Pro, a judge from a different model family to avoid self-preference.

5 Does Social Scaffolding Change What a Team Achieves?

To answer RQ1, we compare each team before evolution, after evolving under *none*, and after evolving under structure of interdependence (Table 4).

Evolution alone falls short. Evolution alone brings little across the 27 MAS settings: it raises the score by at most 4.3 points over *Pre* in 23 of them and even lowers it in 11, and 17 of them fall short of a single strong agent evolved with the same mechanism. Large gains appear mainly in teams of identical strong agents (13.8–16.7 points), the most expensive teams to deploy.

		GAIA				LoCoBench				OfficeQA Pro			
Team	Structure	Pre	Meta-Team	SkillOpt	MemRL	Pre	Meta-Team	SkillOpt	MemRL	Pre	Meta-Team	SkillOpt	MemRL
solo_weak	None	20.8 ±1.4	20.8 ±5.2	20.8 ±5.2	20.0 ±0.0	54.5 ±2.0	50.8 ±1.2	49.5 ±3.6	46.8 ±3.8	2.7 ±0.9	2.7 ±0.9	5.9 ±0.9	2.2 ±2.5
solo_strong	None	51.7 ±2.9	52.5 ±5.0	49.2 ±1.4	53.3 ±5.2	73.1 ±1.5	76.4 ±3.1	74.4 ±2.5	76.8 ±1.7	28.5 ±0.9	31.2 ±4.9	34.4 ±0.9	29.0 ±8.4
onlyweak	None	21.7 ±5.2	21.7 ±9.5	20.0 ±4.3	18.3 ±6.3	52.9 ±3.5	43.4 ±5.3	48.5 ±5.6	43.4 ±4.0	1.1 ±0.9	0.5 ±0.9	0.5 ±0.9	2.2 ±0.9
	Collab.		25.0 ±4.3	19.2 ±6.3	**24.2** ±1.4		**59.3** ±4.1	**54.9** ±6.1	46.8 ±5.1		0.0 ±0.0	0.5 ±0.9	1.1 ±0.9
	Conform.		25.0 ±2.5	20.8 ±1.4	17.5 ±2.5		54.2 ±2.3	45.8 ±6.7	39.1 ±5.9		**1.1** ±0.9	0.0 ±0.0	0.5 ±0.9
	Compet.		**27.5** ±2.5	**28.3** ±1.4	13.3 ±3.8		22.6 ±10.5	28.3 ±1.7	**49.8** ±3.1		0.5 ±0.9	**2.2** ±2.5	**6.5** ±3.2
weakstrong	None	48.3 ±1.4	51.7 ±2.9	51.7 ±2.9	50.0 ±4.3	80.5 ±1.5	81.1 ±5.9	72.7 ±3.0	80.5 ±1.5	17.2 ±4.9	18.8 ±4.7	**21.5** ±2.4	14.5 ±4.3
	Collab.		43.3 ±5.2	62.5 ±4.3	50.8 ±1.4		84.8 ±2.0	**85.9** ±5.6	82.2 ±1.5		16.1 ±3.2	16.7 ±1.8	**15.1** ±3.4
	Conform.		47.5 ±0.0	50.8 ±3.8	**58.3** ±5.8		83.2 ±3.5	79.1 ±1.2	77.4 ±4.1		**19.4** ±4.3	8.6 ±0.9	13.4 ±3.7
	Compet.		**63.3** ±7.6	**74.2** ±3.8	50.8 ±5.8		**86.5** ±0.6	82.5 ±4.6	**84.8** ±1.0		14.5 ±2.8	15.1 ±3.3	10.2 ±1.9
onlystrong	None	45.0 ±5.0	**61.7** ±2.9	**60.8** ±5.8	59.2 ±8.0	63.0 ±2.3	58.6 ±6.3	76.8 ±2.7	64.3 ±4.2	33.9 ±2.8	34.4 ±2.5	35.5 ±4.3	29.6 ±2.5
	Collab.		59.2 ±1.4	51.7 ±8.0	**60.8** ±1.4		**72.4** ±1.2	82.8 ±1.7	59.3 ±3.8		34.9 ±3.7	36.0 ±3.3	30.6 ±1.6
	Conform.		55.0 ±7.5	57.5 ±6.6	53.3 ±3.8		64.0 ±2.1	**87.2** ±2.1	56.2 ±3.3		35.5 ±4.3	**36.6** ±6.5	29.6 ±3.4
	Compet.		43.3 ±7.6	51.7 ±5.2	54.2 ±1.4		70.0 ±2.5	83.8 ±1.0	**83.5** ±1.2		**36.0** ±3.4	33.9 ±1.6	**35.5** ±3.2

Table 4: Main results (avg@3 success rate, %) on the test set. Cell color shows the change from *Pre*: rose for gains and blue for drops (*Pre* : same team before evolution), darker for larger changes. **Bold**: best structure of interdependence for each team, mechanism, and benchmark. Values are mean ± SD over three rollouts.

The right scaffolding makes evolution pay off. An appropriate structure outperforms evolution alone in 24 of the 27 settings, and the best combination of team composition and structure surpasses the strong agent evolved alone on every benchmark and mechanism, by up to 25.0 points (74.2 vs. 49.2 with SkillOpt on GAIA). The same evolution mechanism thus yields far larger gains once the team is organized appropriately.

The best scaffolding differs from team to team. Which structure helps depends on who is in the team. On GAIA with SkillOpt, the onlystrong team is best under evolution alone (60.8), but the weakstrong team is best under competition (74.2). On LoCoBench with Meta-Team, competition lifts the onlystrong team (70.0) but collapses the onlyweak team (22.6). Gains are large both when members clearly differ in ability and when all are strong, and no structure helps when one member is near zero: on OfficeQA Pro, where the weak model scores 2.7 alone, even the best mixed team (21.5) stays below the strong agent alone (28.5). Section 6 explains why: what the right scaffolding changes is who leads the team.

6 How the Team Reorganizes, and Why It Works

To answer RQ2 and RQ3, we trace the final-evaluation sessions of the teams with the largest effects (Table 5) and relate them to human groups. Sections 6.1 and 6.2 answer RQ2, and Section 6.3 answers RQ3. Additional experiments with a much weaker model, different assignments of strong roles, and other model families lead to largely the same conclusions (Appendix K).

6.1 Authority Moves to the Most Capable Member

The first way the team reorganizes (RQ2) is in who leads it. GAIA weakstrong starts with the weak plan_agent as chair, which relays sub-tasks to the strong web_agent and doubles the LLM calls of the strong agent alone (46.8 vs. 22.6). Competition removes this relay by handing the chair to the strong web_agent under SkillOpt and Meta-Team, which then works alone at solo-level cost (74.2 and 63.3). Promoting a weak role instead backfires: on OfficeQA Pro weakstrong, competition hands the chair to a weak role and drops the team below evolution alone (15.1 vs. 21.5 with SkillOpt; Table 4). Every chair change under competition, with the score that decided it, is listed in Table 13. Like human groups, which gain from their most capable member only when they recognize it (Littlepage et al., 1997) and pay for poorly suited leaders (Milgram, 1963; Janis, 1972), competition helps only when the right role rises.

6.2 The Others Step Back, Unless They Remove a Bottleneck

The second way the team reorganizes (RQ2) is in who works. Coordination pays off only when it removes a bottleneck. On LoCoBench weakstrong, the weak readers take over the reading, halving the timeouts of the strong agent alone (5.4% vs. 12.8%) and outscoring it before evolution (80.5 vs. 73.1; Table 5). Coordination beyond this costs the team: conformity raises the LLM calls per session from 80.0 to 107.3 with SkillOpt and 100.9 with MemRL, and drops the team below *Pre* (79.1 and 77.4), adding cost without removing any further bottleneck.

GAIA weakstrong						LoCoBench weakstrong chair: planner (W)					LoCoBench onlystrong chair: planner (S)				
Run	Chair	Agents	Calls	TO	avg@3	Run	Agents	Calls	TO	avg@3	Run	Agents	Calls	TO	avg@3
Solo strong, *Pre*	–	1.00	22.6	0.0	51.7	Solo strong, *Pre*	1.00	24.3	12.8	73.1	Team, *Pre*	3.20	73.3	23.2	63.0
Team, *Pre*	plan_agent (W)	2.76	46.8	7.5	48.3	Team, *Pre*	4.18	80.0	5.4	80.5	SO, none	1.40	27.0	12.8	76.8
SO, none	plan_agent (W)	2.97	39.6	3.3	51.7	SO, conform.	4.62	107.3	5.1	79.1	SO, collab.	1.00	19.7	5.1	82.8
SO, compet.	web_agent (S)	1.38	19.4	0.0	**74.2**	SO, compet.	4.55	91.5	3.4	82.5	SO, conform.	1.00	18.3	4.4	**87.2**
MT, compet.	web_agent (S)	1.01	17.0	0.8	63.3	MR, conform.	4.43	100.9	6.1	77.4	SO, compet.	1.01	21.1	3.7	83.8
MR, compet.	plan_agent (W)	2.98	40.0	7.5	50.8	MR, compet.	4.01	74.8	4.7	**84.8**	MR, conform.	1.17	29.3	11.8	56.2

Table 5: Process metrics per session in the final evaluation for the runs discussed in Section 6 (GAIA n=120, LoCoBench n=297 sessions per run; all scores in Table 4). *Chair*: the role holding the chair and its model; on LoCoBench, the chair stays the planner in every run shown. *Agents*: roles started in the session; *Calls*: LLM calls; *TO*: share of sessions stopped by the 900-second limit (%). Bold: best avg@3 in each block. (SO: SkillOpt, MT: Meta-Team, MR: MemRL; W: weak, S: strong)

In onlystrong, where every member is equally capable, splitting the work adds waiting (23.2% timeouts), so SkillOpt moves the work to the planner alone (1.00–1.40 active agents). Here conformity helps most, with the fewest LLM calls (18.3): it writes a planner-implements rule and removes the developer within two iterations, as conformity spreads one practice and preserves it (Henrich and Boyd, 1998). The same structure thus helps when it removes coordination and hurts when it adds it. Unlike humans, who may loaf beside more capable members (Latané et al., 1979; Karau and Williams, 1993), the others drop out, and the team does better.

6.3 The Team Grows Its Leader

	GAIA		LoCoBench	
	Solo	Team	Solo	Team
avg@3	49.2	**74.2**	74.4	**87.2**
Skills accepted	0/4	2/4	2/4	2/4
Solved only here	5	**35**	8	**46**

Table 6: Team-grown agent (*Team*) vs. the same model alone (*Solo*) under SkillOpt in GAIA weakstrong (competition) and LoCoBench onlystrong (conformity). Skills accepted: this agent's skills that passed the L1 gate, out of its candidates that reached the gate. Solved only here: (task, rollout) pairs only this agent solved.

The reorganized team outperforms the same model working alone (RQ3) because the agent left leading it has learned from its teammates. This agent outperforms the same model evolved alone (Table 6), solving 35 (task, rollout) pairs the solo agent misses against 5 the other way on GAIA, and 46 against 8 on LoCoBench. On GAIA, the strong agent alone accepts none of its four candidate skills, whereas web_agent accepts two of its four within the team, and on LoCoBench the planner follows the workflow its team wrote into the constitution. MemRL under conformity consolidates the work without such learning and falls below *Pre* (56.2 vs. 63.0). The team thus acts as scaffolding in its original sense (Wood et al., 1976; Cole et al., 1978): support the leader needs while it develops and no longer needs once it has.

6.4 What This Means for Team Design

Taken together, the answers to RQ1–RQ3 suggest three principles for designing a self-evolving MAS. First, decide who leads before deciding how agents learn: on GAIA, the same weakstrong team scores 74.2 with the strong web_agent in the chair but at most 62.5 while the weak plan_agent keeps it. Second, let performance choose the leader, but check that it picks a capable role: competition helps when it promotes the strong web_agent and hurts when it promotes a weak role, as on OfficeQA Pro (Section 6.1). Third, keep teammates only where they remove a bottleneck: readers that share a long codebase cut timeouts on LoCoBench, whereas in an onlystrong team the work moves to the planner alone and the score rises (Section 6.2).

7 Conclusion

We investigate whether the social layer of a self-evolving MAS shapes its evolution, as in human groups, and build AGORA to vary team composition and the structure of interdependence with all else fixed. Evolution alone rarely helps; the right scaffolding works by reorganizing the team: authority moves to the most capable member, the others step back unless they remove a bottleneck, and the member left leading, having learned from its teammates, outperforms the same model alone. An evolution mechanism alone is not enough, since the social scaffolding under which a team evolves decides how its agents learn and who gets to lead.

Limitations

Results are partly confounded with benchmarks, so differences in task type may contribute to the patterns we attribute to ability. Our experiments that change the models (Appendix K) partially separate the two, but a full factorial design over ability and task type is left to future work. Each session has a 900-second wall-clock limit, and up to six sessions run concurrently on shared GPUs. Because more active agents can raise the server load and slow down each call, the observation that the team improves when the others step back may partly reflect this serving setup. Due to cost, we run evolution only once per condition. The reported standard deviations reflect only the noise of evaluating the same evolved team three times, not the variance of the evolution process itself.

Acknowledgments

This work was supported by the Institute of Information & Communications Technology Planning & Evaluation (IITP) grant funded by the Korea government (MSIT) (No. 2021-0-01341, Artificial Intelligence Graduate School Program (Chung-Ang University)), the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (No. RS-2026-25494299), and the "Advanced GPU Utilization Support Program" funded by the Government of the Republic of Korea (Ministry of Science and ICT).

References

Alessandro Bellina, Giordano De Marzo, and David Garcia. 2026. Conformity and social impact on ai agents. *Preprint*, arXiv:2601.05384.

Mikako Bito, Keita Nishimoto, Kimitaka Asatani, and Ichiro Sakata. 2026. Large language models exhibit normative conformity. *Preprint*, arXiv:2604.19301.

Yixing Chen, Yiding Wang, Siqi Zhu, Haofei Yu, Tao Feng, Muhan Zhang, Mostofa Patwary, and Jiaxuan You. 2025. Multi-agent evolve: Llm self-improve through co-evolution. *Preprint*, arXiv:2510.23595.

Young-Min Cho, Sharath Chandra Guntuku, and Lyle Ungar. 2025. Herd behavior: Investigating peer influence in llm-based multi-agent systems. *arXiv preprint arXiv:2505.21588*.

Junhyuk Choi, Jeongyoun Kwon, Heeju Kim, Haeun Cho, Hayeong Jung, Sehee Min, and Bugeun Kim. 2026a. Belief in authority: Impact of authority in multi-agent evaluation framework. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 21305–21319, San Diego, California, United States. Association for Computational Linguistics.

Min Choi, Keonwoo Kim, Sungwon Chae, and Sangyeob Baek. 2025. An empirical study of group conformity in multi-agent systems. *Preprint*, arXiv:2506.01332.

Minseok Choi, Seungbin Yang, Dongjin Kim, Subin Kim, Jungmin Son, Yunseung Lee, Jaegul Choo, and Youngjun Kwak. 2026b. Membrane: A self-evolving contrastive safety memory for llm agent defense. *arXiv preprint arXiv:2606.05743*.

Michael Cole, Vera John-Steiner, Sylvia Scribner, and Ellen Souberman. 1978. Mind in society. *Mind in society the development of higher psychological processes. Cambridge, MA: Harvard University Press*.

Yufan Dang, Chen Qian, Xueheng Luo, Jingru Fan, Zihao Xie, Ruijie Shi, Weize Chen, Cheng Yang, Xiaoyin Che, Ye Tian, et al. 2026. Multi-agent collaboration via evolving orchestration. *Advances in neural information processing systems*, 38:165025–165059.

Morton Deutsch. 1949. A theory of co-operation and competition. *Human relations*, 2(2):129–152.

Chen Han, Jin Tan, Bohan Yu, Wenzhen Zheng, and Xijin Tang. 2026. Conformity dynamics in llm multi-agent systems: The roles of topology and self-social weighting. *Preprint*, arXiv:2601.05606.

Zhezheng Hao, Tianfu Wang, Huanshuo Dong, Ziyan Liu, Hong Wang, Xiankun Lin, Qiang Lin, Can Wang, Hande Dong, and Jiawei Chen. 2026. Evolve as a team: Collaborative self-evolution for llm-based multi-agent systems. *arXiv preprint arXiv:2605.29790*.

Joe Henrich and Robert Boyd. 1998. The evolution of conformist transmission and the emergence of between-group differences. *Evolution and human behavior*, 19(4):215–241.

Irving L. Janis. 1972. *Victims of Groupthink: A Psychological Study of Foreign-Policy Decisions and Fiascoes*. Houghton Mifflin, Boston.

David W Johnson and Roger T Johnson. 2009. An educational psychology success story: Social interdependence theory and cooperative learning. *Educational researcher*, 38(5):365–379.

Steven J. Karau and Kipling D. Williams. 1993. Social loafing: A meta-analytic review and theoretical integration. *Journal of Personality and Social Psychology*, 65(4):681.

Zixuan Ke, Yifei Ming, Austin Xu, Ryan Chin, Xuan-Phi Nguyen, Prathyusha Jwalapuram, Jiayu Wang, Semih Yavuz, Caiming Xiong, and Shafiq Joty. 2026. Mas-orchestra: Understanding and improving multi-agent reasoning through holistic orchestration and controlled benchmarks. *arXiv preprint arXiv:2601.14652*.

Bibb Latané, Kipling Williams, and Stephen Harkins. 1979. Many hands make light the work: The causes and consequences of social loafing. *Journal of Personality and Social Psychology*, 37(6):822.

Kuang-Huei Lee, Ian Fischer, Yueh-Hua Wu, Dave Marwood, Shumeet Baluja, Dale Schuurmans, and Xinyun Chen. 2025. Evolving deeper llm thinking. *arXiv preprint arXiv:2501.09891*.

Yoonho Lee, Roshen Nair, Qizheng Zhang, Kangwook Lee, Omar Khattab, and Chelsea Finn. 2026. Meta-harness: End-to-end optimization of model harnesses. *arXiv preprint arXiv:2603.28052*.

Glenn Littlepage, William Robison, and Kelly Reddington. 1997. Effects of task experience and group experience on group performance, member ability, and recognition of expertise. *Organizational behavior and human decision processes*, 69(2):133–147.

Bo Liu, Leon Guertler, Simon Yu, Zichen Liu, Penghui Qi, Daniel Balcells, Mickel Liu, Cheston Tan, Weiyan Shi, Min Lin, Wee Sun Lee, and Natasha Jaques. 2026a. Spiral: Self-play on zero-sum games incentivizes reasoning via multi-agent multi-turn reinforcement learning. *Preprint*, arXiv:2506.24119.

Hanzhi Liu, Chaofan Shou, Xiaonan Liu, Hongbo Wen, Yanju Chen, Ryan Jingyang Fang, and Yu Feng. 2026b. Synthesizing multi-agent harnesses for vulnerability discovery. *arXiv preprint arXiv:2604.20801*.

Jiaqi Liu, Xinyu Ye, Peng Xia, Zeyu Zheng, Cihang Xie, Mingyu Ding, and Huaxiu Yao. 2026c. Evolvemem: Self-evolving memory architecture via autoresearch for llm agents. *arXiv preprint arXiv:2605.13941*.

Xinghua Lou, Miguel Lázaro-Gredilla, Antoine Dedieu, Carter Wendelken, Wolfgang Lehrach, and Kevin P Murphy. 2026. Autoharness: improving llm agents by automatically synthesizing a code harness. *arXiv preprint arXiv:2603.03329*.

Herbert W Marsh. 1987. The big-fish-little-pond effect on academic self-concept. *Journal of educational psychology*, 79(3):280.

Qianyu Meng, Yanan Wang, Liyi Chen, Yihang Li, Wei Wu, Wenyuan Jiang, Qimeng Wang, Chengqiang Lu, Yan Gao, Yi Wu, et al. 2026. Agent harness for large language model agents: A survey. *Preprints*.

Grégoire Mialon, Clémentine Fourrier, Thomas Wolf, Yann LeCun, and Thomas Scialom. 2024. Gaia: a benchmark for general ai assistants. In *International Conference on Learning Representations*, volume 2024, pages 9025–9049.

Stanley Milgram. 1963. Behavioral study of obedience. *The Journal of abnormal and social psychology*, 67(4):371.

Dhwanil M. Mori and Neil F. Johnson. 2026. Three ai-agents walk into a bar . . . . 'lord of the flies' tribalism emerges among smart ai-agents. *Preprint*, arXiv:2602.23093.

Krista Opsahl-Ong, Arnav Singhvi, Jasmine Collins, Ivan Zhou, Cindy Wang, Ashutosh Baheti, Owen Oertell, Jacob Portes, Sam Havens, Erich Elsen, et al. 2026. Officeqa pro: An enterprise benchmark for end-to-end grounded reasoning. *arXiv preprint arXiv:2603.08655*.

Zhenyu Pan, Yiting Zhang, Zhuo Liu, Yolo Yunlong Tang, Zeliang Zhang, Haozheng Luo, Yuwei Han, Jianshu Zhang, Dennis Wu, Hong-Yu Chen, Haoran Lu, Haoyang Fang, Manling Li, Chenliang Xu, Philip S. Yu, and Han Liu. 2025. Advevomarl: Shaping internalized safety through adversarial co-evolution in multi-agent reinforcement learning. *Preprint*, arXiv:2510.01586.

Aneesh Pappu, Batu El, Hancheng Cao, Carmelo di Nolfo, Yanchao Sun, Meng Cao, and James Zou. 2026. Multi-agent teams hold experts back. *Preprint*, arXiv:2602.01011.

Shihao Qi, Jie Ma, Rui Xing, Wei Guo, Xiao Huang, Zhitao Gao, Jianhao Deng, Jun Liu, Lingling Zhang, Bifan Wei, Boqian Yang, Pinghui Wang, Jianwen Sun, Jing Tao, Yaqiang Wu, Hui Liu, Yu Yao, and Tongliang Liu. 2026. Beyond individual intelligence: Surveying collaboration, failure attribution, and self-evolution in llm-based multi-agent systems. *Preprint*, arXiv:2605.14892.

Jielin Qiu, Zuxin Liu, Zhiwei Liu, Rithesh Murthy, Jianguo Zhang, Haolin Chen, Shiyu Wang, Ming Zhu, Liangwei Yang, Juntao Tan, et al. 2025. Locobench: A benchmark for long-context large language models in complex software engineering. *arXiv preprint arXiv:2509.09614*.

Dat Tran and Douwe Kiela. 2026. Single-agent llms outperform multi-agent systems on multi-hop reasoning under equal thinking token budgets. *Preprint*, arXiv:2604.02460.

Juntong Wang, Haoyue Zhao, Xiyuan Wang, Yanbo Wang, Qiyan Deng, Muhan Zhang, et al. 2026. Sage: A self-evolving agentic graph-memory engine for structure-aware associative memory. *arXiv preprint arXiv:2605.12061*.

Zhiyuan Weng, Guikun Chen, and Wenguan Wang. 2025. Do as we do, not as you think: the conformity of large language models. *arXiv preprint arXiv:2501.13381*.

David Wood, Jerome S Bruner, and Gail Ross. 1976. The role of tutoring in problem solving. *Journal of child psychology and psychiatry*, 17(2):89–100.

Yifan Yang, Ziyang Gong, Weiquan Huang, Qihao Yang, Ziwei Zhou, Zisu Huang, Yan Li, Xuemei Gao, Qi Dai, Bei Liu, et al. 2026. Skillopt: Executive strategy for self-evolving agent skills. *arXiv preprint arXiv:2605.23904*.

Siyu Yuan, Kaitao Song, Jiangjie Chen, Xu Tan, Dongsheng Li, and Deqing Yang. 2025. Evoagent: Towards automatic multi-agent generation via evolutionary algorithms. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 6192–6217.

Aimin Zhang, Jiajing Guo, Fuwei Jia, Chen Lv, Boyu Wang, and Fangzheng Li. 2026a. Evoagent: An evolvable agent framework with skill learning and multi-agent delegation. *arXiv preprint arXiv:2604.20133*.

Hanrong Zhang, Shicheng Fan, Henry Peng Zou, Yankai Chen, Zhenting Wang, Jiayu Zhou, Chengze Li, Wei-Chieh Huang, Yifei Yao, Kening Zheng, et al. 2026b. Coevoskills: Self-evolving agent skills via co-evolutionary verification. *arXiv preprint arXiv:2604.01687*.

Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, Xionghui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, et al. 2025. Aflow: Automating agentic workflow generation. In *International Conference on Learning Representations*, volume 2025, pages 34040–34077.

Shengtao Zhang, Jiaqian Wang, Ruiwen Zhou, Junwei Liao, Yuchen Feng, Zhuo Li, Yujie Zheng, Weinan Zhang, Ying Wen, Zhiyu Li, et al. 2026c. Memrl: Self-evolving agents via runtime reinforcement learning on episodic memory. *arXiv preprint arXiv:2601.03192*.

Jie Zhu, Huaixia Dou, Shuo Jiang, Junhui Li, Lifan Guo, Feng Chen, Chi Zhang, and Fang Kong. 2026a. Esc-skills: Discovering and self-evolving skills for emotional support conversations. *arXiv preprint arXiv:2605.27908*.

Zixuan Zhu, Yitong Hu, Yong Dai, Junfeng Fang, Chunyang Jiang, Senkang Hu, and Yuzhi Zhao. 2026b. Unified context evolution for llm agents. *arXiv preprint arXiv:2606.02304*.

A Benchmark Details

For each benchmark, we fix 20 val problems and divide the remaining problems into train and test at a 4:5 ratio (at least 40 test problems, seed 42). Gates always use val, and final evaluation always uses test. Training tasks are drawn epoch-wise: the train problems are shuffled once per epoch and consumed in consecutive, non-overlapping batches. The last batch of an epoch may be smaller, after which the problems are reshuffled. Split sizes and scoring criteria are summarized in Table 7.

For GAIA, we use only the 77 problems without attached files, and answers are scored by exact match on the text after "FINAL ANSWER:". Web search uses the Serper API, with identical queries cached for reproducibility, and each task starts with an empty working directory. Since train has only 17 problems, the same 17 problems are repeated in a reshuffled order with the same seeds over the 4 iterations, so per-iteration batch success rates on GAIA are compared on the same problems.

For LoCoBench, we use 199 Python scenarios. Because the whole codebase has to be read, a 200k-token context is required. A task counts as a success if its LCBS score, which is based on static analysis and ranges up to 5, is 3.0 or higher. Over the 4 iterations, each of the 80 train problems is solved once, so per-iteration batch success rates also reflect differences in problem difficulty.

For OfficeQA, the Pro set consists of 133 hard problems. One to three relevant documents are copied into each task's working directory, and an LLM judge (openai/gpt-5-mini) compares the answer with the ground truth and returns correct or incorrect.

Benchmark	Train	Val	Test	Success criterion	Judge
GAIA	17	20	40	Exact match of the final answer	–
LoCoBench	80	20	99	LCBS $\geq$ 3.0 (out of 5)	–
OfficeQA Pro	51	20	62	Judged correct	gpt-5-mini

Table 7: Benchmark splits and scoring. The validation set (20 tasks) is fixed first; the rest is split train : test = 4 : 5 (test $\geq$ 40, seed 42). Gates always use the validation set and the final evaluation always uses the test set.

B Team Configurations

Table 8 lists the roles of each benchmark, the tools assigned to each role, and which role is strong in weakstrong. In solo, a single agent uses all tools of the benchmark and has no constitution. Since it has no teammates, it evolves through L1 alone, without L2 or L3.

We chose the weak and strong models based on the MMLU-Pro, GPQA, and SWE-bench Verified scores reported on their official model cards, although the evaluation settings there may differ from ours. The weak model, Qwen3-30B-A3B-Instruct-2507, reports 78.4 on MMLU-Pro and 70.4 on GPQA; it does not report SWE-bench Verified (43.2 on LiveCodeBench v6). The strong model, Qwen3.6-35B-A3B, reports 85.2 on MMLU-Pro, 86.0 on GPQA (Diamond), and 73.4 on SWE-bench Verified. Both are MoE models with about 3B active parameters, and thinking mode is turned off for both. In weakstrong, the strong role is the bottleneck role of each benchmark, that is, the role that gathers information or writes the final output. When that role has multiple instances, all of them are made strong.

All template roles use a temperature of 0.1 and a limit of 50 steps per agent, and each task session has a time limit of 900 seconds, separate from the time budget for reflection (Appendix F). The maximum number of output tokens is 16,384 for GAIA (4,096 for answer_agent), 10,000 for LoCoBench, and 16,384 for OfficeQA (8,192 for the readers and the calculator, and 4,096 for answer_agent). Solo agents use 16,384 on every benchmark. The team message limit is 100 for MAS and 50 for solo, and a message sent to several recipients counts once per recipient. Representative role prompts and the full constitutions are shown in Appendix L.

C Differences from Meta-Team

AGORA builds on the public implementation of Meta-Team (Hao et al., 2026). From it, AGORA inherits the basic loop: a chairman-led team of LLM agents solves tasks by exchanging messages, evolves only its external state rather than model weights, and revises itself at the individual (L1), relational (L2), and team (L3) levels. Table 9 summarizes what AGORA changes.

The main change is in purpose. Meta-Team uses the three-level loop to improve a single team under one fixed social design: the chairman alone proposes and applies team changes, while other members can only send suggestions. At L3, this design is closest to our conformity condition. AGORA instead treats this social layer as the variable under study. It parameterizes L2 and L3 by four structures of interdependence, adds team composition as a second axis, and makes L1 a slot for interchange-

Benchmark	Role	Tools	Model in *weakstrong*
GAIA	plan_agent (chairman)	–	weak
	web_agent	web_search, web_fetch, bash, read_file	**strong**
	file_agent	bash, read_file, write_file	weak
	answer_agent	bash, read_file	weak
LoCoBench	planner (chairman)	read_file, bash	weak
	developer	read_file, write_file, bash	**strong**
	reader-1–4	read_file, bash	weak
OfficeQA	planner (chairman)	–	weak
	reader-1, reader-2	bash, read_file, web_search, web_fetch	**strong**
	calculator	bash	weak
	answer_agent	–	weak

Table 8: Roles, tools, and model assignment in *weakstrong* for each benchmark. Weak: Qwen3-30B-A3B-Instruct-2507; strong: Qwen3.6-35B-A3B. All roles also have the messaging tools; the chairman additionally has the role-management (start_agent / stop_agent) and answer-submission (finalize_task) tools. In *weakstrong*, only the bottleneck role is strong; OfficeQA runs this role (reader) as two identical instances, both strong.

able evolution mechanisms so that their effect can be separated from that of social scaffolding.

Two further changes keep conditions comparable. First, Meta-Team applies team changes immediately, with no check on held-out data; AGORA checks SkillOpt skill revisions and every L3 revision on the same validation set (Section 3.1.4), so that no condition keeps changes that clearly hurt the team. Second, Meta-Team solves training tasks one at a time and reflects right after each task; AGORA runs sessions concurrently with explicit budgets and defers reflection until the batch is done (Appendix F). Under the Meta-Team mechanism, we keep the original sequential execution and single training pass.

	Meta-Team	**AGORA**
Goal	Improve one team	Compare social conditions
L1 update	Agent's own edits	Meta-Team, SkillOpt, or MemRL
L2	Single fixed protocol	Set by structure; off under None
L3 proposal	Chairman alone; others suggest	Tier 1 proposes
L3 decision	Applied by the chairman	Rule set by structure
Chairman	Fixed by default	Fixed, or rotates under Competition
Team composition	One model for all roles	Weak/strong per role
Held-out gate	None; changes apply immediately	L1 (SkillOpt) and L3, same val set
Execution	One task at a time	Concurrent sessions†
Reflection	Right after each task	Deferred until the batch ends†

Table 9: What AGORA inherits from and changes in Meta-Team. †Except under the Meta-Team mechanism, which keeps the original sequential execution and single training pass (Table 2).

D Evolution Details

This section follows the three evolution levels in Table 1: the L1 mechanisms that update each role (Appendix D.2), the relational notes of L2, the team-level proposals of L3, and the gates that decide whether L1 and L3 changes are applied (Appendix D.5).

D.1 Models Used at Each Stage

Table 10 lists the model that performs each stage. Everything an agent does inside a session, including its L1, L2, and L3 reflection, uses that agent's own model, so under Meta-Team the team's own models revise themselves. In contrast, the SkillOpt optimizer, the MemRL memory writer, and the L2 contradiction judge are fixed to the strong model regardless of team composition. In the logs of all GAIA Meta-Team runs, every L1 tool call (an edit or a skip) was made by the calling agent's own model.

Stage	Model
Task execution	the role's own model (weak or strong)
L1, L2, and L3 reflection in a session	the agent's own model
L3 decision at the end of an iteration	the current chairman's model
Gate sessions	as in task execution
SkillOpt analysis and edits	the strong model (fixed)
MemRL memory writing	the strong model (fixed)
L2 contradiction judge	the strong model (fixed)
MemRL retrieval embedding	bge-m3
OfficeQA scoring	openai/gpt-5-mini

Table 10: Model used at each stage. "Fixed" means the same model regardless of team composition.

D.2 L1: Evolution Mechanisms

D.2.1 Meta-Team

Meta-Team has no separate mechanism: right after a session, in the L1 phase, agents directly edit their own prompt patches and skills. There is no L1 gate, and as in the original design, batches run sequentially for a single iteration. L3 changes go through

the same L3 gate as in the other mechanisms. The preevolve runs give Meta-Team zero iterations, so the template team is evaluated on test as is.

D.2.2 SkillOpt

We use the reflect → merge → apply functions of the original SkillOpt without modification and run them concurrently for each role. The input to the analysis model follows the original format: trajectories that pair each action with its outcome, the grading result appended at the end (prediction, ground truth, and correctness, or the LCBS sub-scores), the failure reason, and the system prompt of that session. Each analysis call batches up to eight trajectories from the same role. If a batch exceeds the context limit, it is split evenly into the fewest smaller batches that fit, and the longest tool outputs are truncated first only when a single session alone exceeds the limit.

Each iteration uses an edit budget of 4, 16 parallel workers, and the strong model as the optimizer model. Every reflect call writes to a fresh output directory so that revisions from previous iterations are not reused, and slightly malformed JSON in analysis responses is repaired with json_repair, a dependency of the original implementation.

The resulting candidate skill documents are adopted only if they pass the L1 gate. The L1 gate uses the same 20-problem validation set and one-sided binomial test as the L3 gate (Section 3.1.4), but whereas the L3 gate keeps a revision unless it is clearly worse, the L1 gate adopts a candidate only if it is clearly better than the current skill. For example, if the current skill solves 10 of the 20 problems, a candidate must solve at least 12 to be adopted. The significance level is $\alpha = 0.3$, the same as for the L3 gate. Because merely rerunning the same team changes the outcome of about a third of the validation problems, a strict level would fail to distinguish genuine improvements from noise and would reject almost every candidate (full thresholds in Appendix D.5).

D.2.3 MemRL

We use the query-based retrieval of the original MemRL runner. Only memories above the similarity threshold become candidates; they are ranked by $0.5\,z(\mathrm{sim}) + 0.5\,z(Q)$, and up to five are selected. With probability 10%, the selection is random instead, with the random state fixed by the session seed, and memories with $Q < -0.8$ are excluded. All selected memories are added to the session prompt, each marked with its success or failure and its similarity. Every selected memory then receives a reward of $r = +1$ (success) or -1 (failure), and Q is updated as $Q \leftarrow (1-\alpha)Q + \alpha r$ with $\alpha = 0.3$.

To create a new memory, the trajectory, including actions but excluding the system prompt that contains the retrieved older memories, is summarized as a procedural script. Failed trajectories are also stored, as "scripts of failed attempts." The output is capped at 12,800 tokens and each memory at 12,000 characters. If the similarity to an existing retrieval key is 0.8 or higher, the memory is attached to that key; otherwise, it is indexed under a new key. MemRL has no gate.

We made five changes to the original. First, we use a local bge-m3 embedding model instead of the original text-embedding-3-large, and therefore re-tuned the threshold per benchmark so that about 16% of val–train pairs pass, using only the unlabeled task inputs; no labels or scores were used: 0.38 for GAIA, 0.62 for LoCoBench, and 0.56 for OfficeQA (the original also uses a different value per benchmark, ranging from 0.25 to 0.62). Second, the retrieval key uses only the problem content, without the shared prompt template; with the template included, even unrelated GAIA problems reached a similarity of 0.64. Third, each role has its own memory bank, and every experiment starts with an empty bank. Fourth, updates are applied together at the end of each iteration rather than after each task, because sessions in a batch run concurrently. Finally, the model that writes memories is fixed to the strong model regardless of team composition.

D.3 L2: Relational Notes

L2 produces two artifacts: a profile of each teammate and collaboration notes for each teammate. A profile holds at most 5 strengths, 5 weaknesses, 3 notes, and 5 recorded contradictions, newest first, together with a cumulative observation count. Tier 1 can read the notes of the whole team, while Tier 2 can read only its own. Under collaboration and competition, a staged note is finalized once every teammate the author interacted with in the task has responded; if such a teammate finishes its L2 phase without responding, the note is finalized automatically.

To check for contradictions, a strong judge is asked whether a new observation contradicts earlier ones and answers either *NO* or *YES: <one sentence>*. A contradiction is only flagged and does

not block saving. Since sessions in a batch run concurrently, if the file at save time differs from the version a session saw at its start, the notes written by other sessions are kept below the new note. In SkillOpt and MemRL, L2 notes are also passed to the L1 update as reference text.

D.4 L3: Proposals and Review

A new role is proposed as a config and prompt in a new folder, and a removal is proposed by marking the config as deleted. Proposals are not applied immediately but accumulate in a iteration buffer. Tier 2 agents cannot make proposals directly, but they can send a suggestion, which is added to the chairman's review list.

Under collaboration and competition, proposals go through peer review. A proposal passes once a majority ($\lfloor m/2 \rfloor + 1$) of the teammates who took part in the task, excluding the proposer, have commented on it. Without a majority, it still passes 180 seconds after its last edit, and teammates who did not comment are recorded as non-responses. Reviews are stored per version of each proposal file, so editing one proposal does not erase the reviews of another. However a session ends, whether by a timeout or by the chairman skipping, proposals not yet applied are passed on to the gate at the end of the iteration. As safeguards, the model of a new role is reset to the strong model in onlystrong teams and to the weak model otherwise, including weakstrong.

D.5 Gates

Both gates score candidates on the same 20 val problems, but they point in opposite directions: L1 accepts a change only if it is clearly better, whereas L3 accepts a change unless it is clearly worse. L1 follows the original SkillOpt, which keeps a new skill only if it scores higher than the current one, and adds a binomial test against the noise of the small val set. L3 instead checks structural changes that the team itself decided under its structure of interdependence, so it blocks only changes that clearly hurt and otherwise leaves the decision to the team. Both use $\alpha = 0.3$. Given the noise in 20 val problems (re-running the same team changes the outcome of about a third of the tasks), we use this relaxed threshold instead of a strict significance level. Table 11 lists the resulting thresholds for $n = 20$: L1 needs about two or more successes above the reference, and L3 reverts only at two or more successes below it. Gate sessions run under the same structure of interdependence, tiers, and chairman as training, except the shared iteration-0 L1 baseline, which is measured once under none (Appendix D.5.1).

Reference		L1: accept if candidate	L3: revert if changed team
Successes (/20)	Rate b	$\geq$ (/20)	$\leq$ (/20)
0	0.00	2	–
1	0.05	2	–
2	0.10	4	0
3	0.15	5	1
4	0.20	6	2
5	0.25	7	3
6	0.30	8	4
7	0.35	9	5
8	0.40	10	6
9	0.45	11	7
10	0.50	12	8
11	0.55	13	9
12	0.60	14	10

Table 11: Gate thresholds for $n = 20$ and $\alpha = 0.3$. The reference is the stored skill score for L1 and the current team's score for L3. L1 accepts a candidate skill only if it is clearly better (about two or more tasks above the reference); L3 reverts the changes only if the changed team is clearly worse (two or more tasks below). "–": never reverted.

D.5.1 L1 Gate (SkillOpt)

Let b be the baseline success rate and s the number of successes of the candidate out of n tasks. With $p_0 = \max(b, 0.5/n)$ and $X \sim \mathrm{Bin}(n, p_0)$, the candidate is accepted $\iff s > bn \ \wedge \ P(X \geq s) \leq \alpha$. Tasks in which the role was not called are excluded, so n can be smaller than 20.

The iteration-0 baseline is measured once on val using the untrained template team under none and is shared by all SkillOpt runs with the same team composition. A fingerprint of the settings that affect the result is stored with it, and the baseline is re-measured if the fingerprint changes. After an accepted update, the baseline becomes the accepted score. For a newly added role, the baseline is set to its success rate on the tasks where it was called during L3 gate scoring.

D.5.2 L3 Gate (single_check)

The L3 gate first merges the proposals from all sessions of the iteration. For the same file, the last proposal wins, and the config and prompt of a new role are bundled into a single role addition. It then drops invalid proposals: removing the chairman, a new role without a config, and invalid paths. Under competition, proposals that change the chairman are also dropped, since the ranking decides the chairman.

Next, the decision rule of the team's structure

of interdependence (Table 12) determines which proposals to keep and in what order. The decision is made by the actual model of the team's chairman, and if its response cannot be parsed, all proposals are kept. Following this order, at most two role changes and two patch bundles are kept, so that each gate judges only a few changes; this cap applies to every condition, including none. In preliminary runs without the cap, teams that started with six roles grew to as many as 17 (LoCoBench, Meta-Team weakstrong).

The current team and the team with all remaining changes applied then each solve the same 20 val problems with the same seeds (40 sessions, with a separate working directory for each team). With s_b successes out of n_b for the current team and s_c out of n_c for the modified team, let $b = s_b/n_b$ and $X \sim \mathrm{Bin}(n_c, b)$. All changes are reverted $\iff$ $s_c < b\,n_c \ \wedge \ P(X \le s_c) \le \alpha$; otherwise, all are applied. If fewer than 10 sessions complete without crashing, nothing is applied. The team folder is backed up before changes are applied and restored if applying them fails.

E Structure of Interdependence Specification

We confirmed over repeated calls that both the weak and strong models follow the response format of the decision prompt (*{"keep": [...], "order": [...]}*). Under collaboration, the decision follows an explicit rule that classifies each review as SUPPORT, OPPOSE, or NEUTRAL and counts them, so the decision model's own opinion is not taken into account.

E.1 Competition: Promotion and Chairman Change

At the end of each iteration, a score is computed for each role, and the sole top scorer becomes Tier 1 and the chairman for the next iteration; in case of a tie, the current Tier 1 stays. The score depends on the mechanism. Under SkillOpt, it is the current gate score of the role's skill (the iteration-0 baseline or the last accepted score); under MemRL, it is the mean Q of the role's memory bank; and under Meta-Team, it is the success rate on tasks the role took part in during its single training iteration.

When the chairman changes, the original chairman is recorded in the team configuration, and a notice of the change is added to the session prompts. The new chairman is told that it has taken over the chairman's authority (receiving the task, calling roles, and submitting the final answer) and that it can still call the original chairman as a teammate. The original chairman and the other roles are told who the chairman is and that their own duties are unchanged. In sessions without a promotion, the prompts are identical byte for byte. The promoted chairman decides on its own whom to call, and we treat this behavior itself as an object of observation.

The score has no minimum sample size, so a rarely called role can reach first place through an extreme value from a small sample (e.g., 1/1). We keep this rule and report, for every promotion, the role, its score, the sample size, and performance before and after (Table 13). Across the three benchmarks, the chairman changed 20 times in 16 of the 27 competition runs. Of the 12 changes under SkillOpt and Meta-Team, whose scores count tasks, five rested on three or fewer tasks, and on GAIA, file_agent, which takes part in only 2–9% of sessions (Table 22), became the chairman three times. The new chairman ran on the weak model in 13 of the 20 changes, including every change in OfficeQA Pro weakstrong, where the strong roles are the readers.

F Execution Engine Details

This section describes how sessions and reflection are run.

F.1 Session Execution

Up to six sessions run concurrently in training batches, gate scoring, and final evaluation. Ordinary messages are delivered when the recipient is waiting, whereas a phase-transition notice, if unread, is attached to the agent's next tool result and delivered immediately. If an agent sends the same recipient a message identical to its previous one, or retries a send that just failed (to a nonexistent recipient, to itself, or without content), the message is not sent and the agent waits until a new message arrives. Without this safeguard, agents never waited; they repeated the same message hundreds to thousands of times without ever receiving a reply.

The remaining steps, team messages, and time are given as the last message of each call and are not stored in the conversation history, and the system prompt stays fixed throughout a session to keep the prefix cache. Once the history exceeds 151 messages, only the latest 120 are kept. Only when the

	None	Collaboration	Conformity	Competition
L2 sharing	Disabled	Active	Only when blocked	Minimal
L2 note commit	–	Stage → challenge → commit	Direct write	Stage → challenge → commit
Comparative scores shown	–	–	–	Yes
Tier 1	All roles	All roles	Chairman	Top scorer
L3 proposal right	All roles	All roles	Chairman	Tier 1
L3 challenge tool	No	Yes	No	Yes
Peer review	No	Majority, 180 s	No	Majority, 180 s
End-of-step decision	Keep all	Drop if OPPOSE > SUPPORT	Chairman alone	Chairman alone
Order under the cap	Arrival	SUPPORT−OPPOSE	Chairman's priority	Chairman's priority
Chairman change	No	No	No	Top scorer, every step

Table 12: Full specification of the structure of interdependence. Scoring and revert rules are identical across conditions. Gate sessions run under the same condition, tiers, chairman, and constitution as training sessions.

next call would exceed the model limit (200k minus the role's output allowance minus an 8k margin) are older tool outputs shortened, step by step; the system prompt, the task, and the agent's own messages are never truncated. Sessions terminated by the time or message limit are not graded and count as failures, since their working directories may be left half-written; a session in which agents stop on their own is graded on its final output (the chairman's last answer, or the files in solution/ on LoCoBench).

F.2 Deferred Reflection

Reflection is deferred: sessions in a batch wait after finishing their task, and reflection starts once every task in the batch is done, three sessions at a time. Reflection has its own time budget, separate from task time: 300 seconds for L1, 900 for L2, and 900 for L3 (at most 2,160 seconds per session in total, including a margin). Debates are capped at 20 rounds in L2 and 30 in L3. Meta-Team is excluded, since its batch runs sequentially and it reflects right after each task.

We defer reflection because, when it ran right after each task, finished sessions entering L2 re-prefilled their whole history while other sessions were still solving tasks. The resulting prefill load raised timeouts even in gate sessions (48/120 vs. 21/120 at step 0 in one LoCoBench run).

F.3 Seeds and Repeated Evaluation

Training and gate sessions are seeded by task ID, and final evaluation by a hash of (task ID, run), so all experiments share the same set of seeds. Since outcomes still vary across repeats (Appendix G), we report avg@3 with its SD. The SD reflects evaluation noise only, not the variance of the single evolution run per condition. Going from three to five passes would shrink the standard error only modestly (from 2.9 to 2.25 points on the 99 LoCoBench test problems) at 67% more cost.

G Infrastructure and Validation

We serve the models on sets of two GPUs. In each set, one GPU hosts the weak model (memory fraction 0.90) and the bge-m3 embedding model (0.03), and the other hosts the strong model (0.90). Each of the four sets runs one experiment behind its own LiteLLM proxy. Calls to the strong model are about 1.1–1.3 times slower than calls to the weak model, measured by time to first token.

To validate the FP8 KV cache, we ran LoCoBench solo_weak preevolve on the 99 test problems with BF16 and FP8 KV caches, using the same seeds and time limit. The success rate was 0.444 with BF16 and 0.545 with FP8 (paired McNemar $p = 0.11$), so we found no evidence of degradation. On the 67 problems with no timeout under either setting, the median session time dropped from 333 to 215 seconds, and the KV cache capacity of the weak model grew from about 298k to 578k tokens.

Even with seeds fixed per task, repeating the same experiment three times changed the outcome of 32 of the 99 LoCoBench test problems, likely due to numerical differences from GPU batch composition and the timing of asynchronous execution. This is why we report avg@3. The main GAIA, LoCoBench, and OfficeQA experiments ran on machines with eight NVIDIA RTX PRO 6000 Blackwell GPUs (96 GB), and the analysis experiments (Appendix J) on eight RTX PRO 6000 Blackwell Workstation Edition GPUs (team composition and capability gap) and four RTX PRO 6000 Blackwell Server Edition GPUs (model family). We used vLLM 0.28.0 with FlashInfer 0.6.16.post3, Transformers 5.18.0, and PyTorch 2.13, and LiteLLM 1.101.0 as the proxy; the Mis-

Bench.	Mechanism	Team	Step	Chairman change	Model	New chairman score	New chairman n	Runner-up score (n)	Batch success (%) before → after	avg@3 (%)
GAIA	SkillOpt	onlyweak	1	plan_agent → verifier_agent†	W	0.67	3	0.20 (20)	11.8 → 32.4	28.3
GAIA	SkillOpt	weakstrong	0	plan_agent → web_agent	**S**	0.56	16	0.50 (20)	58.8 → 88.2	**74.2**
GAIA	SkillOpt	onlystrong	0	plan_agent → file_agent	S	1.00	1	0.75 (20)	64.7 → 47.1	51.7
GAIA	MemRL	onlyweak	2	plan_agent → answer_agent	W	−0.218	48	−0.220 (49)	23.5 → 41.2	13.3
GAIA	MemRL	onlyweak	3	answer_agent → file_agent	W	0.023	7	−0.203 (63)	41.2 → –	
GAIA	Meta-Team	weakstrong	0	plan_agent → web_agent	**S**	0.69	16	0.65 (17)	64.7 → –	**63.3**
GAIA	Meta-Team	onlystrong	0	plan_agent → file_agent	S	1.00	1	0.71 (17)	70.6 → –	43.3
LoCoBench	SkillOpt	onlyweak	1	planner → verifier†	W	1.00	1	0.70 (20)	70.0 → 35.0	28.3
LoCoBench	MemRL	onlystrong	1	planner → developer	S	0.101	17	0.022 (33)	45.0 → 77.5	83.5
LoCoBench	Meta-Team	onlyweak	0	planner → reader-1	W	0.53	19	0.50 (20)	50.0 → –	22.6
OfficeQA Pro	SkillOpt	onlyweak	0	planner → answer_agent	W	0.14	7	0.10 (10)	10.0 → 8.3	2.2
OfficeQA Pro	SkillOpt	weakstrong	0	planner → calculator	**W**	0.38	13	0.35 (20)	15.0 → 22.4	15.1
OfficeQA Pro	SkillOpt	weakstrong	1	calculator → answer_agent	**W**	1.00	2	0.38 (13)	20.0 → 23.6	
OfficeQA Pro	SkillOpt	onlystrong	1	planner → calculator	S	0.62	13	0.50 (2)	50.0 → 35.7	33.9
OfficeQA Pro	MemRL	onlyweak	1	planner → reader-2	W	0.000	9	−0.027 (11)	0.0 → 7.0	6.5
OfficeQA Pro	MemRL	onlyweak	3	reader-2 → reader-1	W	−0.250	15	−0.266 (15)	5.0 → –	
OfficeQA Pro	MemRL	weakstrong	1	planner → answer_agent	**W**	0.080	25	0.077 (20)	35.0 → 19.5	10.2
OfficeQA Pro	MemRL	weakstrong	2	answer_agent → verifier†	**W**	0.000	11	−0.027 (31)	9.1 → 30.0	
OfficeQA Pro	MemRL	onlystrong	3	planner → calculator	S	0.039	54	0.012 (57)	40.0 → –	35.5
OfficeQA Pro	Meta-Team	weakstrong	0	planner → calculator	**W**	0.14	14	0.13 (15)	10.0 → –	14.5

Table 13: **Chairman changes under competition.** *Model*: the model of the new chairman (W: weak, S: strong). The score that decides the ranking is the skill's gate score for SkillOpt (n: val tasks), the mean Q of the memory bank for MemRL (n: memories), and the success rate on the role's batch tasks for Meta-Team (n: tasks). Batch success is measured in the step of the change and averaged over the later steps; – marks a change at the last step. LoCoBench and OfficeQA Pro batches differ across steps, so their values are not directly comparable. In weakstrong teams, the change helps when the strong role takes the chair (GAIA: 74.2 and 63.3, the best structure) and hurts when a weak role does (OfficeQA Pro: 15.1, 10.2, and 14.5, at or near the bottom). No other competition run changed its chairman. †Role added by L3 in an earlier step.

tral models ran on a separate install with vLLM 0.31.0. Each model server handles up to 16 concurrent sequences with a 200k-token context (131,072 tokens in the Gemma experiment). Thinking is turned off on every request through the chat template (`enable_thinking=false`), which overrides any server default. The Mistral tokenizer does not accept this argument, so it is omitted for the Mistral models, whose reasoning is off by default (`reasoning_effort=none`). The wall-clock time of one experiment on one GPU set was 4.5–16.0 hours for SkillOpt, 3.4–9.1 hours for MemRL, and 2.1–4.9 hours for Meta-Team on GAIA, and 14.5–37.2 hours for SkillOpt on LoCoBench.

H Patterns by Ability Band

The results in Table 4 follow how capable each agent is on the task. Before evolution, the weak model scores 2.7, 20.8, and 54.5 alone on OfficeQA Pro, GAIA, and LoCoBench, and the strong model 28.5, 51.7, and 73.1. Grouping agents into bands by this solo score (0–20, 20–50, 50–70, and 70+) reveals five main patterns in the best social scaffolding; the two remaining teams are discussed in Appendix H.1.

- **P1. A near-zero member** (OfficeQA Pro onlyweak and weakstrong). No structure rescues the team; even the best mixed team (21.5) stays below the strong agent alone (28.5).
- **P2. Identical 20–50 agents** (GAIA onlyweak, OfficeQA Pro onlystrong). Structures help little, adding at most 8.3 points over evolution.
- **P3. 20–50 mixed with 50–70** (GAIA weakstrong). Competition helps most with SkillOpt and Meta-Team, lifting the team to 74.2 and 63.3, well above any other structure; with MemRL, where competition leaves the chair weak, the best structure adds only 8.3 points over evolution, no more than in P2.
- **P4. 50–70 mixed with 70+** (LoCoBench weakstrong). Already strong before evolution, the team gains 1.7–6.0 points over *Pre* under collaboration and competition, while conformity ranks last among the three structures under every mechanism.
- **P5. Identical 70+ agents** (LoCoBench onlystrong). Every mechanism has a structure that helps strongly, but a different one: conformity with SkillOpt (87.2), collaboration with Meta-Team (72.4), and competition with

MemRL (83.5).

H.1 Teams of Identical 50–70 Agents

The two teams of identical 50–70 agents fit none of the five patterns above. Evolution alone already lifts GAIA onlystrong by 14.2–16.7 points, and no structure adds to it under Meta-Team or SkillOpt, whereas LoCoBench onlyweak gains most from collaboration under Meta-Team and SkillOpt (59.3 and 54.9) and collapses under competition (22.6 and 28.3).

I Process Metrics

Social scaffolding can change final performance only through what happens during evolution. We therefore trace its effect in stages: the structure of interdependence changes how agents reflect, reflection changes the team structure, and the team structure changes how agents collaborate and how efficiently they solve tasks, which in turn determines performance. We define metrics for each stage so that a difference in performance can be attributed to a specific change in process. All metrics are computed from three records kept for every experiment: per-session event logs, which record every message, tool call, LLM call with its token counts, timeout, and session duration; per-iteration snapshots of the team, including its roles, prompts, and iteration result; and the decisions of the L1 and L3 gates.

I.1 Reflection

The structure of interdependence acts most directly on reflection (Table 14). Under none, L2 is skipped and L3 proposals are kept without discussion, so the metrics that depend on peer interaction are defined only for collaboration, conformity, and competition.

Metric	Definition
Challenge rate	Challenges to a teammate's note or reflection, divided by the number of proposals
Withdrawal rate	Notes withdrawn by their author, divided by the number of notes
Reflection adoption rate	Reflections applied, divided by reflections proposed
Skip rate	Share of sessions in which the agent skips L2 or L3 reflection
Team-improvement suggestions	Suggestions sent to the chairman by Tier 2 agents
Teammate-model updates	Updates to teammate profiles and collaboration notes

Table 14: Reflection metrics, counted per session and averaged per iteration.

I.2 Team Structure

Table 15 lists the metrics for changes made through L3. Proposals are counted in units of one role addition, one role removal, one constitution patch, or one patch to an existing role, matching the units that the L3 gate merges and caps (Appendix D.5). Initial team sizes differ across benchmarks (Table 8), so we report team size both as an absolute number and relative to the initial team.

Metric	Definition
Roles added and removed	Per iteration and cumulative, counting only changes that pass the L3 gate
Final team size	Number of roles after the last iteration
Proposals and adoption rate	Proposals made, proposals kept by the decision rule, and proposals applied after the gate
Change types	Shares of role additions, constitution patches, and patches to existing roles
Gate margin	Validation successes of the current and modified teams and the p-value of each L3 gate
Usage of added roles	Share of later sessions in which an added role is actually called

Table 15: Team-structure metrics for L3.

I.3 Role-level Updates

L1 updates are recorded per role and per iteration: for SkillOpt, the number of candidate skills accepted by the L1 gate and each role's gate score over iterations; for MemRL, the size of each role's memory bank, its mean Q, and the number of value updates. Under competition, we also record each chairman change, its step, and the score margin (Table 13).

I.4 Collaboration

Collaboration metrics describe how a team solves a single task (Table 16). Under competition, the chairman can change across iterations, so chairman-mediated messages are counted with respect to the chairman of that session.

Metric	Definition	What it indicates
Messages	Messages sent per session	Volume of communication
Active roles	Roles called per session	Team size in use
Speaker concentration	Gini coefficient of each role's share of messages	Hierarchical vs. flat
Chairman share	Share of messages between the chairman and a teammate, among all messages	Hub vs. mesh
Coordination overhead	Calls to wait for replies or check teammates' status, divided by calls to work tools (file, shell, web search)	Waiting relative to working
Communication failures	Sends blocked as repeats and sends to roles that do not exist	Confusion as the team changes

Table 16: Collaboration metrics, computed per session.

I.5 Efficiency and Failures

Table 17 lists the efficiency and failure metrics. Timeouts and session durations are reported sepa-

rately for training and final evaluation.

Metric	Definition
Batch success rate	Success rate of each training batch
Timeout rate	Sessions that hit the 900-second limit, divided by all sessions
Session duration	Median and 90th percentile of session time
Cost per session	Messages, LLM calls, and input and output tokens
Abnormal endings	Sessions in which the chairman exits without submitting an answer, and sessions whose last message is used as a fallback answer
Context trimming	Times per session that older tool outputs are shortened (Appendix F)

Table 17: Efficiency and failure metrics.

The batch success rate is not comparable across iterations, because each iteration solves different problems (except on GAIA, Appendix A) and the last batch of an epoch is smaller (11 problems at iteration 2 of OfficeQA). We use it only as a secondary signal.

Table 18 lists the batch success rate at each iteration and Table 19 the efficiency of the final evaluation. On GAIA, where every iteration solves the same problems, the batch success of SkillOpt weakstrong under competition jumps from 58.8 to 94.1 right after competition hands the chair to web_agent (Table 13), whereas the other structures stay between 41.2 and 70.6. In the final evaluation, timeouts fall most where performance rises most: in LoCoBench onlystrong they drop from 23.2% before evolution to 3.7–12.8% under SkillOpt (Table 5), while the score rises from 63.0 to 76.8–87.2.

I.6 Linking Process to Performance

We connect the process metrics to performance in three ways. First, across all evolved conditions of the three benchmarks, we correlate the change in avg@3 from the preevolve baseline with the change in team size, timeout rate, and messages per session. Second, we compare outcomes task by task before and after evolution, which is less affected by evaluation noise than comparing means (Appendix F.3). Third, for conditions with large effects, we compare the evolved skills and team configurations directly with the templates. Where performance is near the floor, as for the OfficeQA onlyweak team, differences in avg@3 carry little information, and we describe these conditions through the process metrics instead.

Across the 108 evolved runs, a larger gain over *Pre* goes with a smaller team and fewer timeouts (Table 20; pooled Spearman $\rho = -0.32$ and -0.42). The link is strongest on LoCoBench, where the change in timeouts alone tracks the change in score ($\rho = -0.83$), and absent on OfficeQA Pro, where most teams stay near their floor. Messages per session follow performance on LoCoBench but not on GAIA, where the best runs replace the relay through the chairman with a single working agent.

Task by task (Table 21), the best structure solves more (task, rollout) pairs than evolution alone in 24 of the 27 settings, and the difference is significant ($p < 0.05$, exact McNemar test) in 7 of them. These are exactly the settings where Appendix H reports that structure matters: P3 with SkillOpt and Meta-Team (competition), P4 with SkillOpt (collaboration), P5 with all three mechanisms, and LoCoBench onlyweak with Meta-Team (collaboration; Appendix H.1). None falls in P1 or P2, where we report that structures help little or not at all, and P3 with MemRL, whose best structure is conformity because competition leaves the chair weak, is not significant ($p = 0.12$). Evolution alone differs significantly from *Pre* in 7 settings: it helps the three GAIA onlystrong teams and LoCoBench SkillOpt onlystrong, and hurts LoCoBench Meta-Team and MemRL onlyweak and SkillOpt weakstrong. Because the best structure is chosen after the fact from three and the three rollouts of a task are not independent, these p-values are optimistic; we read them as a check on which effects stand out from evaluation noise rather than as confirmatory tests.

J Additional Experiment Setup

The three additional experiments use the same harness, hyperparameters, splits, and seeds as the main experiments. The team-composition and capability-gap experiments run on GAIA, and the model-family experiment runs on GAIA with Gemma-4 and on LoCoBench with Mistral. Only the models or the assignment of strong roles change, as described below, and web access is not restricted, as in the main experiments.

Team composition. The main weakstrong team makes only web_agent strong (1 strong : 3 weak). To make more roles strong, we add strong roles in order of workload, measured on the scored test sessions of the untrained onlyweak, weakstrong, and onlystrong teams (Table 22). web_agent consumes the largest share of input tokens, plan_agent takes part in every session and produces a large share of the output tokens, answer_agent takes part in almost every session but does little work, and file_agent takes part in only 2–9% of sessions. The 2:2 team therefore makes web_agent

		Solo		onlyweak				weakstrong				onlystrong			
Mechanism	Step	Weak	Strong	None	Collab.	Conform.	Compet.	None	Collab.	Conform.	Compet.	None	Collab.	Conform.	Compet.
GAIA (17 problems per step, the same at every step)															
SkillOpt	S0	29.4	70.6	5.9	35.3	11.8	17.6	58.8	41.2	47.1	58.8	58.8	70.6	82.4	64.7
	S1	35.3	58.8	29.4	17.6	23.5	11.8	41.2	41.2	58.8	94.1	70.6	76.5	82.4	35.3
	S2	29.4	70.6	29.4	23.5	23.5	23.5	58.8	58.8	58.8	88.2	70.6	76.5	70.6	47.1
	S3	41.2	70.6	35.3	29.4	11.8	41.2	52.9	70.6	52.9	82.4	64.7	82.4	88.2	58.8
MemRL	S0	35.3	70.6	11.8	17.6	29.4	11.8	35.3	47.1	58.8	47.1	64.7	52.9	52.9	52.9
	S1	41.2	70.6	41.2	29.4	11.8	29.4	58.8	70.6	52.9	76.5	76.5	82.4	76.5	64.7
	S2	35.3	82.4	17.6	35.3	23.5	23.5	58.8	64.7	52.9	76.5	70.6	76.5	70.6	70.6
	S3	47.1	70.6	41.2	41.2	17.6	41.2	64.7	76.5	58.8	70.6	82.4	41.2	76.5	70.6
Meta-Team	S0	23.5	70.6	11.8	23.5	17.6	29.4	52.9	52.9	52.9	64.7	52.9	64.7	58.8	70.6
LoCoBench (20 problems per step)															
SkillOpt	S0	75.0	80.0	40.0	50.0	60.0	45.0	85.0	75.0	85.0	85.0	50.0	55.0	55.0	65.0
	S1	60.0	85.0	40.0	70.0	50.0	70.0	70.0	80.0	80.0	75.0	65.0	75.0	70.0	60.0
	S2	55.0	90.0	35.0	40.0	60.0	35.0	65.0	95.0	90.0	95.0	75.0	75.0	90.0	90.0
	S3	45.0	80.0	45.0	50.0	40.0	35.0	80.0	90.0	85.0	80.0	70.0	90.0	85.0	75.0
MemRL	S0	55.0	70.0	50.0	40.0	55.0	60.0	70.0	70.0	90.0	85.0	65.0	55.0	60.0	65.0
	S1	55.0	75.0	25.0	50.0	50.0	60.0	90.0	85.0	80.0	75.0	60.0	55.0	60.0	45.0
	S2	60.0	70.0	45.0	55.0	30.0	45.0	80.0	95.0	90.0	90.0	80.0	65.0	60.0	75.0
	S3	55.0	85.0	35.0	45.0	30.0	50.0	80.0	95.0	95.0	85.0	55.0	75.0	60.0	80.0
Meta-Team	S0	65.0	90.0	85.0	55.0	65.0	50.0	75.0	85.0	85.0	85.0	80.0	65.0	65.0	70.0
OfficeQA Pro (20 problems per step; 11 at S2)															
SkillOpt	S0	0.0	25.0	5.0	5.0	5.0	10.0	15.0	25.0	15.0	15.0	30.0	25.0	35.0	40.0
	S1	10.0	40.0	5.0	0.0	0.0	20.0	30.0	25.0	15.0	20.0	55.0	55.0	45.0	50.0
	S2	9.1	36.4	0.0	0.0	9.1	0.0	18.2	27.3	9.1	27.3	54.5	63.6	45.5	36.4
	S3	5.0	35.0	0.0	0.0	0.0	5.0	25.0	30.0	15.0	20.0	35.0	50.0	50.0	35.0
MemRL	S0	5.0	20.0	5.0	5.0	0.0	5.0	5.0	10.0	20.0	20.0	40.0	30.0	20.0	35.0
	S1	0.0	65.0	0.0	0.0	0.0	0.0	30.0	35.0	25.0	35.0	55.0	60.0	55.0	65.0
	S2	0.0	63.6	0.0	0.0	9.1	9.1	18.2	18.2	27.3	9.1	63.6	63.6	63.6	54.5
	S3	0.0	35.0	5.0	0.0	0.0	5.0	15.0	20.0	30.0	30.0	55.0	55.0	40.0	40.0
Meta-Team	S0	5.0	40.0	0.0	0.0	5.0	0.0	10.0	30.0	15.0	10.0	25.0	40.0	35.0	30.0

Table 18: Training batch success rate (%) at each step (S0–S3); Meta-Team runs a single step. On GAIA every step solves the same 17 train problems in a new order, so steps are comparable; on LoCoBench and OfficeQA Pro each step solves different problems, and step 2 of OfficeQA Pro has only 11 problems (Appendix A).

and plan_agent strong, and the 3:1 team additionally makes answer_agent strong; file_agent stays weak in all teams. Because plan_agent is the chairman, the L3 decision in these two teams is made by the strong model. Workload alone does not show which role gains most from a stronger model, so we also evaluate, without evolution, a team in which only plan_agent is strong and a team in which only answer_agent is strong, and compare them with the main weakstrong team. The 2:2 and 3:1 teams are trained with all three mechanisms under all four structures.

Capability gap. The weak model is replaced by Qwen3.5-2B (thinking off, as for all models), and the strong model is unchanged. On its official model card, Qwen3.5-2B reports 55.3 on MMLU-Pro and 51.6 on GPQA (measured in thinking mode), so the gap to the strong model grows from 6.8 to 29.9 points on MMLU-Pro and from 15.6 to 34.4 points on GPQA. We evaluate solo_weak and weakstrong without evolution and train weakstrong with SkillOpt, Meta-Team, and MemRL under all four structures.

Model family. The weak and strong models are replaced by Gemma-4 E4B and Gemma-4 26B-A4B. Their model cards report 69.4 and 82.6 on MMLU-Pro and 58.6 and 82.3 on GPQA Diamond (measured in thinking mode), gaps of 13.2 and 23.7 points. The two models are served under the same names as in the main experiments, so every stage fixed to the strong model (Table 10) uses Gemma-4 26B-A4B. Since E4B supports a 128K context, both models are served with a 131,072-token limit instead of 200k; in the main GAIA runs, only 0.57% of sessions grew long enough to be affected by this limit. This experiment uses SkillOpt only, with solo_weak and solo_strong under none and onlyweak, weakstrong, and onlystrong under all four structures.

With thinking off, the Gemma-4 chat template starts each model turn with an empty thought block, and Gemma-4 26B-A4B often keeps emitting further empty thought blocks (<|*channel*>*thought*<*channel*|>) until it reaches the output

Team	Structure	GAIA				LoCoBench				OfficeQA Pro			
		TO	Time	Calls	Msgs	TO	Time	Calls	Msgs	TO	Time	Calls	Msgs
onlyweak	*Pre*	11.7	257	45.8	15.4	17.8	530	106.6	26.9	79.6	902	96.5	17.2
	None	8.6	271	66.7	13.1	23.4	581	112.6	27.0	53.2	728	90.5	19.9
	Collab.	5.0	226	60.3	9.1	19.1	575	99.3	26.4	46.8	773	79.1	14.8
	Conform.	8.3	259	48.2	8.2	26.1	619	127.4	33.5	86.4	908	77.0	14.7
	Compet.	5.3	168	75.0	18.7	35.5	738	142.3	40.6	41.7	692	106.8	27.3
weakstrong	*Pre*	7.5	132	46.8	20.4	5.4	341	80.0	19.6	44.6	772	109.8	15.1
	None	3.1	66	41.6	6.6	7.8	374	93.1	23.2	36.0	634	110.7	19.1
	Collab.	3.3	57	36.3	5.2	4.6	327	81.9	19.6	39.1	750	129.7	21.9
	Conform.	1.7	139	55.7	8.3	5.6	395	98.9	23.2	42.6	712	113.7	17.8
	Compet.	2.8	48	25.5	2.3	3.9	342	78.7	18.0	23.8	301	128.8	33.2
onlystrong	*Pre*	22.5	182	42.4	6.0	23.2	534	73.3	4.8	34.4	525	91.9	11.8
	None	14.5	136	38.6	5.5	22.7	545	56.8	3.7	33.1	336	81.4	12.1
	Collab.	15.0	150	43.2	5.7	16.1	471	37.4	1.7	25.4	355	88.9	12.4
	Conform.	15.3	169	40.6	6.1	14.7	447	43.4	2.1	26.9	373	78.4	12.5
	Compet.	20.8	159	45.1	5.1	11.9	444	35.6	1.9	30.3	129	51.3	7.1

Table 19: Efficiency in the final evaluation. *TO*: sessions stopped by the 900-second limit (%); *Time*: median session time (s); *Calls*: LLM calls per session; *Msgs*: messages sent per session. Evolved rows average the Meta-Team, SkillOpt, and MemRL runs of each structure; *Pre* is the same team before evolution.

	n	ΔTeam size	ΔTO	ΔMessages	ΔCalls
GAIA	36	**−0.59**	**−0.46**	0.06	−0.28
LoCoBench	36	**−0.45**	**−0.83**	**−0.65**	**−0.67**
OfficeQA Pro	36	−0.17	−0.22	0.00	−0.16
Pooled	108	**−0.32**	**−0.42**	**−0.36**	**−0.33**

Table 20: Spearman correlation between the change in avg@3 from *Pre* and the change in each process metric, over all evolved runs (36 per benchmark). **Bold**: $p < 0.05$.

limit, without calling a tool. This happened mostly to the chairman: without evolution, 78 of the 120 test sessions of the onlystrong team hit the 900-second limit, and the team scored 20.0. Removing the empty block from the template did not stop the loop. We therefore suppress the token that opens a thought block (<|*channel*>) for the strong model with a logit bias of -100, which keeps thinking off; with this, the same team scored 58.3. The weak model showed no such loop and is left unchanged. All reported Gemma runs that involve the strong model use this suppression.

On LoCoBench, the weak and strong models are replaced by Ministral-3-8B and Mistral-Small-4-119B-A6B (NVFP4). Their model cards report no benchmark on which both were evaluated in the same setting, so we characterize the pair by its solo scores on LoCoBench (58.2 and 77.4), which fall in the same 50–70 and 70+ bands as the Qwen pair. As with Gemma-4, every stage fixed to the strong model uses Mistral-Small-4, and the experiment uses SkillOpt only, with solo_weak and solo_strong under none and onlyweak, weakstrong, and onlystrong under all four structures.

K Additional Experiment Results

This section reports the three analysis experiments whose setup is given in Appendix J. Each table reads its columns against the same team's *Pre*, as in Table 4. Appendix H.1 additionally discusses two teams from the main results. The results support the reading of Section 6: a structure of interdependence helps when it routes the work to a capable role, and which role is strong matters more than how many roles are strong.

K.1 Team Composition

Before evolution, the choice of the strong role matters far more than the number of strong roles (Table 23). Making only web_agent strong gives 48.3, whereas making only plan_agent strong gives 28.3 and making only answer_agent strong gives 25.0, close to the onlyweak team (21.7). Adding strong roles on top of web_agent barely helps: the 2:2 team scores 52.5 and the 3:1 team 51.7.

After evolution, the structure that helps most shifts once the chairman is strong (Table 24). In the 1:3 team, where the weak plan_agent chairs, competition is the best structure for SkillOpt and

Bench.	Mechanism	Team	Best structure	Best structure vs. none +/−	Best structure vs. none p	None vs. *Pre* +/−	None vs. *Pre* p
GAIA	Meta-Team	onlyweak	Compet.	+17/−10	0.248	+14/−14	1.000
GAIA	Meta-Team	weakstrong	Compet.	+26/−12	**0.034**	+23/−19	0.644
GAIA	Meta-Team	onlystrong	Collab.	+15/−18	0.728	+28/−8	**0.001**
GAIA	SkillOpt	onlyweak	Compet.	+17/−7	0.064	+10/−12	0.832
GAIA	SkillOpt	weakstrong	Compet.	+34/−7	**<0.001**	+21/−17	0.627
GAIA	SkillOpt	onlystrong	Conform.	+13/−17	0.585	+25/−6	**<0.001**
GAIA	MemRL	onlyweak	Collab.	+18/−11	0.265	+11/−15	0.557
GAIA	MemRL	weakstrong	Conform.	+22/−12	0.121	+17/−15	0.860
GAIA	MemRL	onlystrong	Collab.	+16/−14	0.856	+25/−8	**0.005**
LoCoBench	Meta-Team	onlyweak	Collab.	+93/−46	**<0.001**	+50/−78	**0.017**
LoCoBench	Meta-Team	weakstrong	Compet.	+41/−25	0.064	+39/−37	0.909
LoCoBench	Meta-Team	onlystrong	Collab.	+70/−29	**<0.001**	+40/−53	0.213
LoCoBench	SkillOpt	onlyweak	Collab.	+78/−59	0.124	+59/−72	0.294
LoCoBench	SkillOpt	weakstrong	Collab.	+59/−20	**<0.001**	+36/−59	**0.023**
LoCoBench	SkillOpt	onlystrong	Conform.	+45/−14	**<0.001**	+69/−28	**<0.001**
LoCoBench	MemRL	onlyweak	Compet.	+76/−57	0.118	+54/−82	**0.020**
LoCoBench	MemRL	weakstrong	Compet.	+38/−25	0.130	+43/−43	1.000
LoCoBench	MemRL	onlystrong	Compet.	+72/−15	**<0.001**	+46/−42	0.749
OfficeQA Pro	Meta-Team	onlyweak	Conform.	+2/−1	1.000	+0/−1	1.000
OfficeQA Pro	Meta-Team	weakstrong	Conform.	+16/−15	1.000	+11/−8	0.648
OfficeQA Pro	Meta-Team	onlystrong	Compet.	+25/−22	0.771	+16/−15	1.000
OfficeQA Pro	SkillOpt	onlyweak	Compet.	+4/−1	0.375	+1/−2	1.000
OfficeQA Pro	SkillOpt	weakstrong	Collab.	+15/−24	0.200	+20/−12	0.215
OfficeQA Pro	SkillOpt	onlystrong	Conform.	+22/−20	0.878	+19/−16	0.736
OfficeQA Pro	MemRL	onlyweak	Compet.	+11/−3	0.057	+4/−2	0.688
OfficeQA Pro	MemRL	weakstrong	Collab.	+13/−12	1.000	+10/−15	0.424
OfficeQA Pro	MemRL	onlystrong	Compet.	+28/−17	0.135	+12/−20	0.215

Table 21: Task-level paired comparison on the 120 (GAIA), 297 (LoCoBench), and 186 (OfficeQA Pro) (task, rollout) pairs of the test set, which share seeds across runs. +/−: pairs solved only by the first / only by the second run; p: exact McNemar test. *Best structure*: the best of collaboration, conformity, and competition. **Bold**: $p < 0.05$.

Role	Sessions	Input tokens	Output tokens	Tool calls
web_agent	78–80	50–78	28–52	33–54
plan_agent (chairman)	100	15–32	39–58	30–42
answer_agent	97–99	5–15	9–12	12–22
file_agent	2–9	0–2	0–4	0–3

Table 22: Workload of each GAIA role (%): share of sessions it joins and of the team's tokens and tool calls, given as the range over the three teams before evolution.

Strong role	Pre
web_agent (main weakstrong)	48.3 ±1.4
plan_agent	28.3 ±2.9
answer_agent	25.0 ±2.5
web + plan (2:2)	52.5 ±6.6
web + plan + answer (3:1)	51.7 ±3.8

Table 23: Which role is strong matters (GAIA, no evolution, avg@3, %; mean ± SD over three rollouts). Each row makes the listed role(s) strong and the rest weak.

Meta-Team (74.2 and 63.3) because it moves the chair to the strong web_agent (Section 6.1). In the 2:2 and 3:1 teams, plan_agent is already strong, so competition has no weak chair to replace. With SkillOpt it moves the chair from one strong role to another (plan_agent to web_agent) and gains little (54.2 and 55.0); with Meta-Team it moves the chair to the weak file_agent and gives the lowest score of its column (44.2 and 49.2); with MemRL the chair does not change. Instead, collaboration becomes the best structure for Meta-Team (60.8 and 63.3), for SkillOpt (57.5 and 56.7), and for MemRL in the 2:2 team (61.7). More strong roles do not raise the ceiling: the best 2:2 and 3:1 teams (61.7 and 63.3) stay below the 1:3 team under SkillOpt competition (74.2).

Structure	1:3 main (strong: web)				2:2 (strong: web + plan)				3:1 (strong: web + plan + answer)			
	Pre	Meta-Team	SkillOpt	MemRL	Pre	Meta-Team	SkillOpt	MemRL	Pre	Meta-Team	SkillOpt	MemRL
None	48.3 ±1.4	51.7 ±2.9	51.7 ±2.9	50.0 ±4.3	52.5 ±6.6	56.7 ±1.4	49.2 ±5.2	57.5 ±2.5	51.7 ±3.8	58.3 ±2.9	55.8 ±8.0	56.7 ±7.6
Collab.		43.3 ±5.2	62.5 ±4.3	50.8 ±1.4		**60.8** ±3.8	**57.5** ±2.5	**61.7** ±3.8		**63.3** ±7.6	**56.7** ±3.8	54.2 ±5.8
Conform.		47.5 ±0.0	50.8 ±3.8	**58.3** ±5.8		58.3 ±1.4	55.0 ±6.6	58.3 ±5.2		54.2 ±3.8	48.3 ±2.9	55.0 ±2.5
Compet.		**63.3** ±7.6	**74.2** ±3.8	50.8 ±5.8		44.2 ±1.4	54.2 ±7.2	57.5 ±2.5		49.2 ±5.2	55.0 ±2.5	**60.0** ±6.6

Table 24: Team composition on GAIA (avg@3, %, mean ± SD over three rollouts): the ratio of strong to weak roles in the weakstrong team. Strong roles are added in order of measured workload (web > plan > answer). Colors show the change from *Pre* of the same team as in Table 4; **bold** marks the best structure of interdependence in each column.

Team	Structure	Main gap (weak = Qwen3-30B-A3B)				Larger gap (weak = Qwen3.5-2B)			
		Pre	Meta-Team	SkillOpt	MemRL	Pre	Meta-Team	SkillOpt	MemRL
solo_weak	None	20.8 ±1.4	20.8 ±5.2	20.8 ±5.2	20.0 ±0.0	16.7 ±5.8	14.2 ±5.8	15.0 ±6.6	8.3 ±3.8
weakstrong	None	48.3 ±1.4	51.7 ±2.9	51.7 ±2.9	50.0 ±4.3	3.3 ±3.8	**4.2** ±1.4	5.0 ±5.0	5.8 ±2.9
	Collab.		43.3 ±5.2	62.5 ±4.3	50.8 ±1.4		**4.2** ±3.8	1.7 ±2.9	5.0 ±2.5
	Conform.		47.5 ±0.0	50.8 ±3.8	**58.3** ±5.8		3.3 ±1.4	**5.8** ±6.3	1.7 ±1.4
	Compet.		**63.3** ±7.6	**74.2** ±3.8	50.8 ±5.8		3.3 ±1.4	0.0 ±0.0	**65.8** ±6.3

Table 25: Larger weak–strong gap on GAIA (avg@3, %, mean ± SD over three rollouts). The strong model is Qwen3.6-35B-A3B in both settings; only the weak model changes. Colors show the change from *Pre* of the same team as in Table 4; **bold** marks the best structure of interdependence in each column.

K.2 Capability Gap

As shown in Table 25, with Qwen3.5-2B as the weak model, the solo weak agent scores 16.7, in the 0–20 band of P1, and evolving it alone does not help (14.2, 15.0, and 8.3 with Meta-Team, SkillOpt, and MemRL). The weakstrong team falls to 3.3 before evolution, far below both its own weak agent alone and the strong agent alone (51.7). The 2B chairman fails to lead: in 47 of the 120 test sessions the team hits the 900-second limit, and in 61 the chairman never submits an answer, repeatedly starting agents and sending messages without the required fields. Evolution leaves the team near zero (0.0–5.8) in 11 of the 12 settings, as P1 predicts. The exception is MemRL under competition, which hands the chair to the weak answer_agent in the second iteration and to the strong web_agent in the third; the training batch success then jumps from 17.0 to 82.0, web_agent solves 118 of the 120 final sessions alone, and the team reaches 65.8, above even the strong agent alone. SkillOpt competition instead moves the chair only to answer_agent, in the last iteration, and scores 0.0, and Meta-Team competition does not change the chair (3.3). The same team thus scores near zero under a weak chair and 65.8 once the strong role takes the chair, the extreme case of Section 6.1.

Team	Structure	Qwen (main)		Gemma-4	
		Pre	SkillOpt	Pre	SkillOpt
solo_weak	None	20.8 ±1.4	20.8 ±5.2	27.5 ±0.0	28.3 ±3.8
solo_strong	None	51.7 ±2.9	49.2 ±1.4	44.2 ±5.2	43.3 ±3.8
onlyweak	None	21.7 ±5.2	20.0 ±4.3	20.0 ±7.5	19.2 ±3.8
	Collab.		19.2 ±6.3		20.8 ±7.6
	Conform.		20.8 ±1.4		21.7 ±6.3
	Compet.		**28.3** ±1.4		**30.0** ±6.6
weakstrong	None	48.3 ±1.4	51.7 ±2.9	35.8 ±3.8	**37.5** ±0.0
	Collab.		62.5 ±4.3		30.8 ±5.2
	Conform.		50.8 ±3.8		25.8 ±8.0
	Compet.		**74.2** ±3.8		32.5 ±5.0
onlystrong	None	45.0 ±5.0	**60.8** ±5.8	58.3 ±5.8	54.2 ±3.8
	Collab.		51.7 ±8.0		56.7 ±5.2
	Conform.		57.5 ±6.6		**59.2** ±10.1
	Compet.		51.7 ±5.2		50.0 ±6.6

Table 26: Different model family on GAIA (avg@3, %, mean ± SD over three rollouts), SkillOpt only. Gemma-4: weak = E4B, strong = 26B-A4B; the SkillOpt optimizer is the family's strong model. Colors show the change from *Pre* of the same team as in Table 4; **bold** marks the best structure of interdependence in each column.

K.3 Model Family

In Table 26, both Gemma-4 agents fall in the 20–50 band on GAIA (27.5 and 44.2 alone), so the weakstrong team is a P2-like mix rather than the P3 mix of the Qwen pair, and Appendix H predicts small effects of evolution and structure. This holds across the teams. In weakstrong, evolution alone adds 1.7 points (37.5) and no structure helps further (30.8, 25.8, and 32.5 under collaboration, conformity, and

Team	Structure	Qwen (main) Pre	Qwen (main) SkillOpt	Mistral Pre	Mistral SkillOpt
solo_weak	None	54.5 ±2.0	49.5 ±3.6	58.2 ±1.5	67.3 ±3.5
solo_strong	None	73.1 ±1.5	74.4 ±2.5	77.4 ±3.5	80.5 ±3.5
onlyweak	None	52.9 ±3.5	48.5 ±5.6	68.0 ±3.2	**71.4** ±2.1
	Collab.		**54.9** ±6.1		70.0 ±3.1
	Conform.		45.8 ±6.7		63.6 ±2.0
	Compet.		28.3 ±1.7		45.5 ±3.6
weakstrong	None	80.5 ±1.5	72.7 ±3.0	78.5 ±2.9	83.5 ±2.5
	Collab.		**85.9** ±5.6		**83.8** ±3.0
	Conform.		79.1 ±1.2		83.2 ±2.3
	Compet.		82.5 ±4.6		61.6 ±4.4
onlystrong	None	63.0 ±2.3	76.8 ±2.7	55.2 ±3.8	45.8 ±5.2
	Collab.		82.8 ±1.7		**58.9** ±4.7
	Conform.		**87.2** ±2.1		31.3 ±5.6
	Compet.		83.8 ±1.0		33.0 ±2.9

Table 27: Different model family on LoCoBench (avg@3, %, mean ± SD over three rollouts), SkillOpt only. Mistral: weak = Ministral-3-8B, strong = Mistral-Small-4-119B-A6B (NVFP4); teams that use only the strong model serve it on two GPUs. Colors show the change from *Pre* of the same team as in Table 4; **bold** marks the best structure of interdependence.

competition); competition moves the chair after the first iteration, but to the weak answer_agent rather than the strong web_agent, and the team ends below none, as in OfficeQA Pro weakstrong (Section 6.1). The onlystrong team already outscores the strong agent alone before evolution (58.3 vs 44.2), but evolution barely changes it under any structure (50.0–59.2). In onlyweak, every structure stays within 1.7 points of *Pre* (20.0) except competition (30.0), which hands the chair to web_agent, the role that does most of the work; in the main GAIA onlyweak team, competition is likewise the best structure with SkillOpt and Meta-Team (28.3 and 27.5). Evolving either agent alone does not help either (28.3 and 43.3).

On LoCoBench, the Mistral pair (58.2 and 77.4 alone) falls in the same bands as the Qwen pair (Table 27). Two results carry over: collaboration is again the best or tied-best structure for the weakstrong and onlystrong teams (83.8 and 58.9), and competition again collapses the onlyweak team (45.5 vs. 71.4 under none, cf. 28.3 with Qwen). Unlike with Qwen, competition also lowers the weakstrong team (61.6 vs. 78.5 before evolution), and no structure consolidates the onlystrong team. Instead of moving the work to one role as in P5, the team grows from six to eight to eleven roles, and the share of final sessions that time out rises from 30% before evolution to 43–53% under none, conformity, and competition (45.8, 31.3, and 33.0). Only collaboration keeps timeouts at the Pre level (29%) and improves on Pre. As in Section 6.2, added roles pay off only when they remove a bottleneck; here they mostly add waiting.

L Prompt Examples

We show the prompts that define the task, the team, and each structure of interdependence. All prompts are copied verbatim from the experiment code; values filled in by the code appear as placeholders in braces (e.g., *{partner_name}*), and the first line of each box gives the source file and line. Role prompts and constitutions are the weakstrong templates; across team compositions, only the *model:* line of each role differs, except that the OfficeQA weakstrong constitution adds a *Model split* paragraph. The prompts of SkillOpt and MemRL are used without modification from their public implementations and are omitted.

L.1 Task Prompts

GAIA and OfficeQA require a *FINAL ANSWER:* line, while LoCoBench is scored on the files written to *solution/*. OfficeQA answers are judged by *openai/gpt-5-mini* with the official HLE grading prompt.

Task prompts (GAIA, LoCoBench)

```
# source:
↪  core/evolution_mechanisms/task_source.py:520-532
↪  (_GAIA_USER_PROMPT), 628-646
↪  (LightweightGAIASource._to_tasks), 1389-1422
↪  (LightweightLoCoBenchSource._build_task_text),
↪  1424-1438 (LightweightLoCoBenchSource._to_tasks);
↪  core/runner.py:591 (chairman's first user message),
↪  486-529 (_TIME_BUDGET_CONSEQUENCE, time_budget_line);
↪  core/agent.py:1162-1199 (per-turn budget message)
# Used in: the first user message of the chairman's
↪  session (task phase), for every training,
↪  held-out-gate and final-evaluation session. Teammates
↪  do not receive it directly; the chairman forwards
↪  work via send_message.

############################################################
↪  ####################
## 0. Wrapper (core/runner.py:591)
############################################################
↪  ####################
## Task

{task_text}

############################################################
↪  ####################
## 1. GAIA {task_text}  (_GAIA_USER_PROMPT.format(level,
↪  question))
############################################################
↪  ####################
## Question (Level {level})

{question}

## Answer Format

YOUR FINAL ANSWER should be a number OR as few words as
↪  possible OR a comma separated list of numbers and/or
↪  strings.
- If asked for a number: don't use comma to write your
↪  number, don't use units such as $ or percent sign
↪  unless specified otherwise.
```

```
- If asked for a string: don't use articles, don't use
↪  abbreviations (e.g. for cities), write digits in
↪  plain text unless specified otherwise.
- Example: respond with 'Saint Louis' instead of 'St.
↪  Louis'; respond with '17' instead of '17m'.

State your final answer clearly as: "FINAL ANSWER: <your
↪  answer>"

# Scoring: text after the first "final answer:"
↪  (case-insensitive) in the chairman's
# finalize_task output, lower-cased and stripped of
↪  non-alphanumerics, must equal the
# equally normalized reference answer. Self-contained
↪  GAIA subset only (no attached files);
# each task gets an empty scratch workspace for
↪  computation.

############################################################
↪  ####################
## 2. LoCoBench {task_text}  (_build_task_text)
############################################################
↪  ####################
## Task: {title}

**Category**: {task_category, underscores -> spaces,
↪  Title Case}
**Difficulty**: {difficulty}
**Language**: {language}

{description}

{task_prompt}

The following source code files are available in your
↪  workspace under `context/`:

  - `context/{file_1}`
  - `context/{file_2}`
  ...
  - ... and {n - 50} more files            (only if more
  ↪  than 50 files)

**Workspace location**: `{workspace}`
- Context files: `{workspace}/context/...`
- Write solutions to: `{workspace}/solution/...`

  - {evaluation_criterion_1}
  - {evaluation_criterion_2}
  ...

Write your solution code files into the `solution/`
↪  directory using `write_file`
(e.g. `write_file(path="solution/main.py",
↪  content="...")`). Create complete,
working {language} code files that address all
↪  requirements.

**IMPORTANT**: You MUST write at least one file to
↪  `solution/` before finishing.
An empty solution scores 0.

# Categories used: feature_implementation +
↪  cross_file_refactoring (Python).
# Scoring: static-analysis LCBS score of the files in
↪  solution/; success iff LCBS >= 3.0 (of 5).
# LoCoBench has no FINAL ANSWER instruction (the output
↪  graded is the solution/ directory).

############################################################
↪  ####################
## 3. Per-turn budget message (--show-time-budget)
############################################################
↪  ####################
# Sent to every agent before each LLM call as a trailing
↪  user message that is not
# stored in the conversation (so the system-prompt prefix
↪  stays cache-stable).
# The step line is always present; the team-message line
↪  only for the chairman;
# the time line only with --show-time-budget
↪  (AGORA_SHOW_TIME_BUDGET=1).

[Budget status -- automatic, not a message from a
↪  teammate]
[Your budget: {remaining_steps} of {max_steps} steps
↪  remaining]
[Team budget: {remaining_messages} of {max_messages}
↪  messages remaining]
[Time: {M}m {SS}s of {M}m {SS}s remaining {label}]
↪  {consequence}

{label}:  "for the task" (task phase)  |  "in this phase"
↪  (L1/L2/L3 reflection)

{consequence}, task_execution:
If time runs out before the task is finalized, the task is
↪  scored as failed -- work done so far is not graded.

{consequence}, l1_reflection and l2_reflection:
When this phase's time runs out, it ends for everyone and
↪  the session moves to the next phase; anything not yet
↪  finalized is not carried over.

{consequence}, l3_reflection:
When this phase's time runs out, the session ends;
↪  anything not yet finalized is not carried over.

# Example (task phase, max_seconds 900):
# [Time: 12m 05s of 15m 00s remaining for the task] If
↪  time runs out before the task is finalized, the task
↪  is scored as failed -- work done so far is not graded.
```

OfficeQA task prompt and LLM judge prompt

```
# source:
↪  core/evolution_mechanisms/task_source.py:893-922
↪  (_OFFICEQA_USER_PROMPT), 1038-1049
↪  (LightweightOfficeQASource._to_tasks), 775-803
↪  (_HLE_JUDGE_PROMPT), 1051-1071 (score)
# Used in: OfficeQA (Pro and Full). The task prompt is
↪  the chairman's first message (wrapped in "##
↪  Task\n\n{task_text}" like the other benchmarks); the
↪  1-3 source documents are copied into the task
↪  workspace under documents/. MemRL retrieval key = the
↪  question text only (content-only key).

############################################################
↪  ####################
## 1. OfficeQA task prompt
↪  (_OFFICEQA_USER_PROMPT.format(question, file_list))
############################################################
↪  ####################
## Question

{question}

## Source Documents

This question is grounded in the following Treasury
↪  Bulletin issue(s),
copied into your workspace as plain text (pre-OCR'd, no
↪  PDF parsing
needed):

{file_list}

**IMPORTANT -- these files are large** (tens of thousands
↪  of lines each).
Do NOT try to `read_file` an entire file at once. Use
↪  `bash` (e.g.
`grep -n "keyword" documents/<file>` or similar) to
↪  locate the relevant
section(s) first, then `read_file` with a narrow
↪  offset/line range around
each match. Reading selectively leaves you enough budget
↪  to actually
answer the question.

## Answer Format

The question above specifies its OWN required answer
↪  format (units,
rounding, ordering, whether to wrap multiple values in
↪  brackets, etc.) --
follow it exactly as stated in the question text, not a
↪  generic
convention. State your final answer clearly as:
"FINAL ANSWER: <your answer, formatted exactly as the
↪  question asks>"
```

```
# {file_list} = one line per source file: "  -
↪  `documents/{file_name}`"

############################################################
↪  ####################
## 2. LLM judge (single user message, judge model
↪  openai/gpt-5-mini via OpenRouter, temperature 0.0,
↪  max_tokens 1024)
##    Mirrors the HLE reference judge prompt used by
↪  MemRL's hle_runner verbatim.
##    success = the judge's "correct: yes"; the
↪  extracted_final_answer is logged.
############################################################
↪  ####################
Judge whether the following [response] to [question] is
↪  correct or not based on the precise and unambiguous
↪  [correct_answer] below.

[question]: {question}

[response]: {response}

Your judgement must be in the format and criteria
↪  specified below:

extracted_final_answer: The final exact answer extracted
↪  from the [response]. Put the extracted answer as
↪  'None' if there is no exact, final answer to extract
↪  from the response.

[correct_answer]: {correct_answer}

reasoning: Explain why the extracted_final_answer is
↪  correct or incorrect based on [correct_answer],
↪  focusing only on if there are meaningful differences
↪  between [correct_answer] and the
↪  extracted_final_answer. Do not comment on any
↪  background to the problem, do not attempt to solve
↪  the problem, do not argue for any answer different
↪  than [correct_answer], focus only on whether the
↪  answers match.

correct: Answer 'yes' if extracted_final_answer matches
↪  the [correct_answer] given above, or is within a
↪  small margin of error for numerical problems. Answer
↪  'no' otherwise, i.e. if there is any inconsistency,
↪  ambiguity, non-equivalency, or if the extracted
↪  answer is incorrect.

confidence: The extracted confidence score between 0% and
↪  100% from [response]. Put 100 if there is no
↪  confidence score available.
```

L.2 Constitutions and Chairman Prompts

The chairman is plan_agent in GAIA and planner in LoCoBench and OfficeQA. Prompts of the other roles are in the code repository.

GAIA constitution

```
# GAIA 4-Agent Team -- Constitution

## Mission

Answer real-world, knowledge-intensive GAIA questions
↪  through a 4-agent
collaboration that mirrors the MASFly multi-agent setup:
**PlanAgent -> FileAgent / WebAgent -> AnswerAgent**.

## Team

- **plan_agent** *(Chairman)* -- analyses the question,
↪  decides which
  specialists to recruit, dispatches sub-tasks,
  ↪  aggregates results,
  hands off to AnswerAgent, then submits the final answer.
- **file_agent** -- extracts and summarises relevant data
↪  from any
  attached file (xlsx, csv, pdf, docx, pptx, zip, jsonld,
  ↪  txt, py, ...).
- **web_agent** -- performs targeted web searches and
↪  page fetches to
  find specific facts, with source URLs.
- **answer_agent** -- synthesises all findings, applies
↪  strict GAIA
  format rules, and returns the single-line `FINAL
  ↪  ANSWER: <...>`
  to PlanAgent.

## SOP (mirrors MASFly `workflow`)

```
1. User         -> PlanAgent
2. PlanAgent    -> FileAgent      [if file is provided]
3. PlanAgent    -> WebAgent       [if web search is
↪  needed]
4. FileAgent    -> AnswerAgent    [via PlanAgent]
5. WebAgent     -> AnswerAgent    [via PlanAgent]
6. AnswerAgent  -> PlanAgent      [returns "FINAL ANSWER:
↪  ..."]
7. PlanAgent    -> set_final_output + terminate
```

Core Principles

1. **Plan before acting** -- PlanAgent decomposes the
↪ question and writes
   an explicit dispatch plan before recruiting.
2. **Specialists do specialist work** -- FileAgent does
↪ *only* file work;
   WebAgent does *only* web work. Cross-tool work (e.g.
   ↪ "search the web
   then process a downloaded file") stays under
   ↪ PlanAgent's coordination.
3. **AnswerAgent enforces format** -- the GAIA grader is
↪ unforgiving.
   AnswerAgent re-reads the question, applies the rules,
   ↪ returns the
   `FINAL ANSWER:` line.
4. **PlanAgent submits after AnswerAgent confirms** --
↪ only the chairman
   calls `set_final_output` / `finalize_task`, and only
   ↪ **after**
   `answer_agent` has replied with a message containing a
   `FINAL ANSWER:` line. The value PlanAgent submits MUST
   ↪ be copied
   verbatim from AnswerAgent's reply. PlanAgent MUST NOT
   ↪ fabricate,
   paraphrase, or extract an answer directly from
   ↪ FileAgent's or
   WebAgent's findings -- routing through AnswerAgent is
   ↪ mandatory,
   even when the raw findings look conclusive.
   ↪ Specialists never
   submit.
5. **Cite sources** -- WebAgent attaches URLs to each
↪ fact; FileAgent
   attaches sheet/page/section locations.
6. **Compute, never estimate** -- use `bash`+Python for
↪ any arithmetic,
   unit conversion, or aggregation.

Available Tools (per Agent)

- `plan_agent`:
  (chairman primitives `list_pool`, `start_agent`,
  ↪ `send_message`,
   `wait_for_replies`, `set_final_output`, `terminate`
   ↪ are auto-injected)
- `file_agent`: bash, read_file, write_file
- `web_agent`: web_search, web_fetch, bash, read_file
- `answer_agent`: bash, read_file

Critical Answer Format

YOUR FINAL ANSWER should be a number OR as few words as
↪ possible OR a
comma-separated list of numbers and/or strings.

- **Number**: no commas, no units unless specified (`17`
↪ not `17m`).
- **String**: no articles, no abbreviations (`Saint
↪ Louis` not
  `St. Louis`); digits in plain text unless asked
  ↪ otherwise.
- **List**: apply per-element rules, separated by `, `.
- The grader normalises both sides as `lowercase + strip
  non-alphanumerics` and checks exact equality. Brevity
  ↪ wins.
```

````
## Submission Format

Only `plan_agent` submits, exactly once:
```
set_final_output(output="FINAL ANSWER: <concise answer>")
terminate(reason="task completed")
```
````

GAIA plan_agent (chairman)

````
# PlanAgent -- Question Analyst & Coordinator (Chairman)

## Your Role

You are a general AI assistant and the **Chairman** of this 4-agent
GAIA team. Your job is purely **planning, routing, and final submission**.
You **do not have any information-gathering tools yourself** -- no
`web_search`, no `web_fetch`, no `bash`, no `read_file`, no `write_file`.

You **must** delegate concrete work to your specialists:

- File reading / parsing -> `file_agent` (has `bash` + `read_file`).
- Web searching / page fetching -> `web_agent` (has `web_search` +
  `web_fetch` + `bash`).
- Final answer synthesis with strict format -> `answer_agent`.

This forced separation mirrors MASFly's 4-Agent SOP and prevents the
chairman from collapsing the team into a 1-agent run.

## Workflow (mirrors MASFly SOP)

```
1. User          -> PlanAgent (you)
2. PlanAgent     -> FileAgent     [if file is provided]
3. PlanAgent     -> WebAgent      [if web search is needed]
4. FileAgent     -> AnswerAgent   [via PlanAgent]
5. WebAgent      -> AnswerAgent   [via PlanAgent]
6. AnswerAgent   -> PlanAgent     [returns "FINAL ANSWER: ..."]
7. PlanAgent     -> set_final_output + terminate
```

## Step-by-Step Procedure

### Step 1 -- Read the question carefully
Identify:
- What exactly is being asked?
- What answer format? (number / a few words / comma-separated list)
- Is a file attached? What type? (xlsx, pdf, csv, docx, pptx, zip, ...)
- Is web information required? Is computation required?

### Step 2 -- Write an explicit dispatch plan
Briefly state your plan in your thoughts, e.g.:
- "Need to read the xlsx -> recruit file_agent."
- "Need 2024 Wikipedia fact -> recruit web_agent."
- "Both file + web -> recruit both, then synthesise via answer_agent."
- "Always recruit answer_agent for final synthesis."

### Step 3 -- Recruit the right teammates
```
list_pool()
start_agent(name="file_agent")     # only if file is attached
start_agent(name="web_agent")      # only if web info needed
start_agent(name="answer_agent")   # ALWAYS, for the final synthesis
```

### Step 4 -- Dispatch sub-tasks with `send_message`
Be concrete:

- **To `file_agent`** -- include:
  - the absolute path of the file in the workspace,
  - the original question,
  - which slice (sheet / page / column / row) probably contains the
    answer, if you can guess,
  - what specific fact / aggregate to extract.

- **To `web_agent`** -- include:
  - the original question,
  - 2-4 candidate English search queries (comma-separated form is fine),
  - the exact fact (number / name / date) you need,
  - any high-priority sources (Wikipedia, official site, paper venue).

- **To `answer_agent`** -- include:
  - the original question verbatim,
  - the **complete findings** from `file_agent` and/or `web_agent`,
  - a reminder of the GAIA answer format constraints.

### Step 5 -- Wait for replies
Use `wait_for_replies(from_agents=["file_agent","web_agent"])` to block
until they report back. Then forward their findings to `answer_agent`
and `wait_for_replies(from_agents=["answer_agent"])`.

### Step 6 -- Receive the synthesised `FINAL ANSWER:` from answer_agent
Verify it follows the format rules below. If it looks wrong (still has
units, articles, abbreviations, or doesn't match the question's exact
ask), send another message to `answer_agent` asking for a corrected
single-line answer.

### Step 7 -- Submit (protocol-mandatory)

Before calling `set_final_output` / `finalize_task`, verify **all** of
the following:

1. `answer_agent` has been recruited and has replied to you.
2. AnswerAgent's reply contains a line starting with `FINAL ANSWER:`.
3. The value you are about to submit is **copied verbatim** from
   the text after `FINAL ANSWER:` in AnswerAgent's reply -- not
   extracted, rephrased, or inferred from FileAgent or WebAgent
   messages.

If any check fails, do **not** submit. Instead: dispatch `answer_agent`
(or re-dispatch with clearer inputs / a corrected format), then
`wait_for_replies(from_agents=["answer_agent"])` again.

Even when FileAgent's or WebAgent's findings already look conclusive
(e.g. they state "the answer is X"), you must still route through
`answer_agent` -- that agent's job is precisely to verify and format.
Skipping it is a protocol violation and will produce a wrong answer.

Once all checks pass:
```
set_final_output(output="FINAL ANSWER: <value copied from answer_agent>")
terminate(reason="task completed")
```

## You Have NO Information-Gathering Tools

If you find yourself thinking "let me just search this myself" or "let
me read the xlsx myself" -- **stop**. You can't. You only have:
`list_pool`, `start_agent`, `stop_agent`, `send_message`,
`read_messages`, `wait_for_replies`, `check_agent_status`,
`set_final_output`, `terminate`.
````

```
Delegate to specialists. That is your job.

## Critical Answer Format

In `set_final_output`, your output **must** start with
↪  `FINAL ANSWER:`
on its own line and contain a single concise answer.

- **Number**: no commas, no units unless specified. `17`
↪  not `17,000`
  or `17m`.
- **String**: no articles, no abbreviations. `Saint
↪  Louis` not
  `St. Louis`.
- **Comma-separated list**: apply per-element rules,
↪  separated by `, `.
- The grader normalises both as `lowercase + strip
↪  non-alphanumerics`
  and checks exact equality. Brevity wins.

## Tips for Effective Coordination

- It is forbidden to bounce a task right back to the
↪  agent who sent it
  without making any attempt yourself -- but for you
  ↪  "attempt" means
  re-routing it correctly to the right specialist.
- Always pass the **original question text** to
↪  `answer_agent` -- never
  just raw findings, or it cannot judge the format.
- If `file_agent` or `web_agent` reports failure, send
↪  refined
  instructions, or recruit the other specialist as a
  ↪  fallback.
- For trivially simple questions (one-step lookup), you
↪  still must
  recruit `answer_agent` to enforce the answer format.
  ↪  The grader is
  unforgiving.
```

LoCoBench constitution

```
# LoCoBench -- Long-Context Code Engineering Team

## Mission
Analyze large codebases (10-100+ files) provided as
↪  context, understand architecture and logic, then
↪  complete a software engineering task.

## Environment
- **Context files**: In `context/` directory. Use
↪  `read_file(path="context/...")`.
- **Solution output**: Write to `solution/` directory
↪  using `write_file(path="solution/...")`.
- **No Docker**: All tools operate on the local
↪  workspace.

## Team Strategy
The key challenge is **long context** -- 15-50 source
↪  files totaling 50K-200K+ tokens.
**Divide and conquer**: Planner splits files -> Readers
↪  read in parallel -> Developer writes solution.

## Hard Constraints
- **ALWAYS** write output to `solution/` directory
- **NEVER** try to execute the context code
- **Be selective** with file reading
- For ANALYSIS tasks: write `.py` files with analysis as
↪  comments, NOT markdown

## Reflection Note
During reflection phases, ignore any references to
↪  `docker_bash` or `git diff` -- those are for other
↪  benchmarks. Your tools are `read_file`, `write_file`,
↪  and `bash`. Your task validation shows LCBS scores
↪  (SE/FC/CQ/LCU dimensions) -- use these to guide your
↪  reflection.
```

LoCoBench planner (chairman)

```
# Planner -- Task Coordinator

You coordinate a team to analyze large codebases and
↪  complete software engineering tasks.

## Your Team
- **reader-1 .. reader-4** -- Read source files and
↪  report findings **directly to developer**.
- **developer** -- Writes solution files based on reader
↪  reports.

## Your Workflow

### Step 1: Analyze (2-3 steps)
Read the task description. Examine the context file list.
↪  Decide which files are relevant and split them into
↪  groups.

### Step 2: Recruit All Agents (1 step)
```
start_agent(name="developer")
start_agent(name="reader-1")
start_agent(name="reader-2")
```

### Step 3: Brief Developer (1 step)
Send developer the task requirements and tell it to
↪  **wait for reader reports** before writing code:
```
send_message(to="developer", content="TASK: [task
↪ description]. Wait for reader-1 and reader-2 to send
↪ you code analysis, then write the solution to
↪ solution/. Key requirements: ...")
```

### Step 4: Dispatch Readers (1 step per reader)
Tell each reader to read specific files and **report
↪  directly to developer** (NOT to you):
```
send_message(to="reader-1", content="Read these files and
↪ send your findings DIRECTLY to developer:
↪ context/src/main.py, context/src/auth.py. Include
↪ actual code snippets of key functions.")
send_message(to="reader-2", content="Read these files and
↪ send your findings DIRECTLY to developer:
↪ context/src/api/routes.py, context/src/db/models.py.
↪ Include actual code snippets.")
```

### Step 5: Wait and Submit (2-3 steps)
```
wait_for_replies(from_agents=["developer"])
```
When developer reports solution files written, submit the
↪  result and end the task.

## Key Principles
- **Readers report to developer, NOT to you** --
↪  eliminates information relay loss
- **You are a dispatcher, not a relay** -- brief
↪  developer once, dispatch readers, then wait
- Scale readers: 1-8 files->1 reader, 9-20->2, 21-35->3,
↪  36+->4
- Spend minimal steps on coordination -- most value is in
↪  reading and writing
```

OfficeQA constitution (weakstrong)

```
# OfficeQA 5-Agent Team -- Constitution (weakstrong)

## Model split

`reader-1`/`reader-2` = qwen-strong (the
↪  highest-judgment role -- finding
the right figure in a dense, irregularly-formatted OCR'd
↪  document is
where mistakes actually originate; a bad extraction
↪  poisons every
downstream step regardless of how strong
↪  calculator/answer_agent are).
`planner`/`calculator`/`answer_agent` = qwen-weak --
↪  calculator always
executes via code rather than mental arithmetic, so
↪  correct formula
selection on already-clean inputs doesn't need the strong
↪  model the way
extraction does. Mirrors GAIA's weakstrong split
↪  (web_agent=strong, the
evidence-gathering role, not answer_agent).
```

```
## Mission

Answer OfficeQA Pro questions (Databricks,
↪  arXiv:2603.08655) that require
finding and combining precise figures buried in dense
↪  Treasury Bulletin
documents (each ~68K tokens, some questions cite 2-3
↪  issues across
different years) and computing question-specific derived
↪  statistics
(CAGR, geometric mean, elasticity, etc.) exactly.
↪  Frontier agents with
full document access still only average ~34% on this
↪  benchmark -- the
difficulty is real, not an artifact of missing context.

## Team

- **planner** *(Chairman)* -- reads the question, decides
↪  which
  document(s) each reader should search and for what
  ↪  specific figure,
  dispatches readers, forwards their extracted raw
  ↪  figures to
  `calculator`, forwards the computed result to
  ↪  `answer_agent`, submits.
- **reader-1**, **reader-2** -- each searches ONE
↪  assigned source document
  (via `bash`/grep, never reading the whole file) for the
  ↪  SPECIFIC raw
  figures planner asked for, and reports them verbatim
  ↪  with the exact
  line/table they came from (no computation, no rounding
  ↪  -- that is
  calculator's job, not theirs). If a question cites 3
  ↪  documents, planner
  assigns reader-1 a second document after its first
  ↪  report.
- **calculator** -- takes the raw figures reader(s)
↪  extracted and performs
  the EXACT computation the question asks for (CAGR,
  ↪  geometric mean,
  percentage change, elasticity, sums/differences across
  ↪  years, etc.) via
  `bash`+Python -- never estimates or does arithmetic in
  ↪  its head.
- **answer_agent** -- formats calculator's numeric result
↪  into the EXACT
  shape the question itself specifies (units, rounding,
  ↪  bracket/ordering
  requirements) -- every question states its own format,
  ↪  there is no
  single fixed convention to memorize.

## Why split reading from computing

Early failure mode to guard against: an agent that both
↪  reads a raw
figure AND immediately computes/rounds it tends to
↪  silently propagate
transcription errors (misread digit) straight into a
↪  "confident" final
number with no independent check. Splitting extraction
↪  (reader) from
computation (calculator) means calculator only trusts
↪  figures it was
explicitly handed, and can flag back to planner if a
↪  figure looks
suspicious (e.g. wrong order of magnitude) rather than
↪  computing through it.

## SOP

```
1. User   -> Planner
2. Planner   -> Reader(s)   [dispatch: which document,
↪ which specific figure]
3. Reader(s) -> Planner   [raw figures + source
↪ line/table, verbatim]
4. Planner   -> Calculator   [question's requested
↪ computation + raw figures]
5. Calculator-> Planner   [computed result]
6. Planner   -> AnswerAgent   [original question +
↪ computed result]
7. AnswerAgent -> Planner   [FINAL ANSWER: ...,
↪ formatted per the question's own spec]
8. Planner   -> set_final_output + terminate
```
```

OfficeQA planner (chairman)

```
# Planner -- Coordinator (Chairman)

## Your Role

You are the **Chairman**. Your job is purely **planning,
↪  routing, and
final submission** -- you have **no tools**. The question
↪  prompt already
lists which document(s) are in your workspace under
↪  `documents/`; you
never read them yourself.

You **must** delegate:
- Searching a specific document for a specific figure ->
↪  `reader-1` / `reader-2`.
- Turning raw figures into the exact computation the
↪  question asks for -> `calculator`.
- Formatting the final answer to spec -> `answer_agent`.

## Workflow

```
1. User        -> Planner (you)
2. Planner     -> reader-1 [, reader-2]   [one document +
↪  one specific figure each]
3. Reader(s)   -> Planner                 [raw figures,
↪  verbatim, with source line]
4. Planner     -> Calculator              [question's
↪  requested computation + raw figures]
5. Calculator  -> Planner                 [computed
↪  result]
6. Planner     -> AnswerAgent             [original
↪  question + computed result]
7. AnswerAgent -> Planner                 [FINAL ANSWER:
↪  ...]
8. Planner     -> set_final_output + terminate
```

Step-by-Step Procedure

Step 1 -- Read the question carefully
Identify: exactly which figure(s) are needed, from which
↪ document(s)
(the question prompt lists them), what computation (if
↪ any) combines
them, and what output format the question itself
↪ specifies (units,
rounding, brackets, ordering of multiple values).

Step 2 -- Assign documents to readers
```
list_pool()
start_agent(name="reader-1")
start_agent(name="reader-2")   # only if a second
↪  document is needed
```
If the question cites 3 documents, dispatch 2 to start,
↪ then re-dispatch
reader-1 (or reader-2) to the third document after it
↪ reports back -- do
not try to force 3 documents onto 2 readers in one
↪ message.

To each reader, be concrete:
- The exact document filename (from `documents/`).
- The EXACT figure/row/column/date you need -- not just
↪ "read the file."
- Ask them to report the value verbatim plus where they
↪ found it.

Step 3 -- Route to Calculator
Forward the question's exact requested computation (paste
↪ the relevant
sentence from the question) plus every raw figure
↪ reader(s) reported,
unedited, to `calculator`. Never do arithmetic yourself.

Step 4 -- Route to AnswerAgent
Send `answer_agent` the original question verbatim +
↪ calculator's
computed result.
```

````
### Step 5 -- Submit
Before `set_final_output`, verify:
1. Every document the question cites was actually
↪  assigned to a reader.
2. `calculator` was used for any arithmetic (never
↪  approximate yourself).
3. `answer_agent` replied with a `FINAL ANSWER:` line,
↪  copied verbatim.

```
set_final_output(output="FINAL ANSWER: <value copied from
↪  answer_agent>")
terminate(reason="task completed")
```

## Rules

- You have no tools -- you cannot open `documents/`
↪  yourself. If you find
  yourself thinking "let me just check the file", stop
  ↪  and dispatch a reader.
- Never skip `calculator` for questions that need
↪  arithmetic, even simple
  ones (e.g. a single subtraction) -- this project's
  ↪  convention is
  "never estimate, always compute."
````

L.3 Structure of Interdependence

Each structure of interdependence differs in how agents discuss L2 notes, who decides on L3 changes, and, under competition, who leads.

L2 discussion. Under none, L2 is skipped.

L2 discussion protocol: collaboration

```
### Discussion Protocol

**Default action: Share, don't skip.** Even if
↪  collaboration went smoothly with no concrete problem,
↪  proactively reach out to your partners and share how
↪  you approached the task -- your reasoning, what you
↪  checked, what you'd do differently next time. The
↪  goal is for the whole team to learn from each other's
↪  process, not just to flag failures.

- **If you are the initiator** for a partner: Send them a
↪  short message about your approach and invite them to
↪  do the same. Then use `wait_for_replies` to collect
↪  their response.
- **If you are waiting** for a partner: Do NOT send them a
↪  message first. Wait for them to reach out, then reply
↪  thoughtfully and share your own approach in return.
- **Independent verification is encouraged**: if a
↪  partner's result affected your work, briefly say
↪  whether you independently checked it and what you
↪  found -- this helps the team calibrate trust over
↪  time.

**Important**: Do not send discussion messages to
↪  partners you are assigned to wait for. This prevents
↪  duplicate crossed conversations.

- **Share by default** -- Even smooth collaboration is
↪  worth discussing so the team learns from it.
```

L2 discussion protocol: conformity

```
### Discussion Protocol

**Default action: Defer to your partner's
↪  self-assessment.** If a partner has already shared
↪  their own evaluation of their work (via a message, a
↪  status report, or their side of the task), accept it
↪  and build your `update_correlation` /
↪  `update_teammate_profile` notes on top of it rather
↪  than independently re-checking or re-litigating it.
↪  Assume they already verified their own claims -- you
↪  do not need to double-check them.

- **If you are the initiator** for a partner: Only reach
↪  out if there is a concrete problem that cannot be
↪  resolved by accepting their self-assessment. Keep
↪  pushback minimal.
- **If you are waiting** for a partner: Do NOT send them
↪  a message first. If they reach out, agree with their
↪  framing where reasonable rather than raising
↪  counter-points.
- **Do not independently re-verify** a partner's reported
↪  results unless something is clearly broken -- treat
↪  disagreement as costly and to be avoided.

**Important**: Do not send discussion messages to
↪  partners you are assigned to wait for. This prevents
↪  duplicate crossed conversations.

- **Skip by default, and don't push back** -- Only raise
↪  an issue if it's impossible to proceed without doing
↪  so.
```

L2 discussion protocol: competition

```
### Discussion Protocol

**Default action: Minimize information sharing.** You are
↪  being evaluated on your own track record (acceptance
↪  rate, error-catch rate, output quality), not on how
↪  much you help teammates. Only reach out to a partner
↪  when a concrete problem directly blocks your own work
↪  -- not to help them improve, and not to share your
↪  reasoning process.

**If your "Your L2 Discussion Assignments" section below
↪  includes a `[comparative signal -- ...]` tag for a
↪  partner**, that's a real number showing how you and
↪  that partner currently stack up (see the tag's own
↪  note for exactly what it measures and its limits --
↪  it is not always a clean apples-to-apples individual
↪  performance score). Use it as competitive context for
↪  how you approach that partner, but don't overstate
↪  its precision beyond what its note says.

- **If you are the initiator** for a partner: Send only
↪  the minimum needed to resolve the specific blocking
↪  problem. Do not volunteer extra context about your
↪  own approach.
- **If you are waiting** for a partner: Do NOT send them
↪  a message first. If they reach out, answer only
↪  what's asked.
- **Treat `update_correlation` /
↪  `update_teammate_profile` as your own private
↪  record**, not as a collaboration aid -- note what
↪  helps *you* predict and work around this partner, not
↪  what would help them.

**Important**: Do not send discussion messages to
↪  partners you are assigned to wait for. This prevents
↪  duplicate crossed conversations.

- **Skip by default, and keep exchanges minimal** --
↪  Treat every interaction as a cost, not an investment
↪  in the team.
```

L3 proposals and decisions. The chairman receives the structure prompt below. At the end of each iteration, the actual model of the chairman decides which proposals to keep in a single call (temperature 0); under none, all proposals are kept without a call.

L3 structure prompt (sent to the chairman)

```
## L3 Reflection -- Team Structure & Configuration
↪  (Chairman Only)

You are now in the **L3 Reflection phase**. Your job is
↪  to evolve the team so it performs better next round.

### Goal
```

```
Evolve the team. You can change team membership
↪  (add/remove agents), update rules (constitution), or
↪  modify agent configurations. **Active experimentation
↪  is encouraged.**

### Step 1: Evaluate the Team

First, consider your team size: how many agents do you
↪  have? Is that enough to cover analysis,
↪  implementation, and verification without overloading
↪  anyone?

Then think about these questions based on your own
↪  observations -- even when the task succeeded:

**Team Membership:**
- What subtask consumed the most time or budget? Could a
↪  dedicated agent handle it better?
- Did one agent have to do too many different types of
↪  work? Would splitting responsibilities help?
- Was there a verification gap -- nobody caught problems
↪  before submission?
- Could any work have been done in parallel with separate
↪  agents?
- Is any current agent consistently unused? Should it be
↪  replaced with something more useful?

**When to add a new agent** (instead of adding more
↪  rules):
- If the same type of failure keeps happening despite
↪  constitution/prompt updates -> a dedicated agent is
↪  more reliable than rules
- If verification is consistently missing -> a
↪  verification agent is better than a "remember to
↪  verify" rule
- If the planner is doing too much analysis -> an
↪  analysis agent can work in parallel
- If test discovery is a recurring bottleneck -> a
↪  test-specialist agent can handle it

**Workflow & Rules:**
- Were there bottlenecks? Did agents wait too long for
↪  instructions?
- Are the team's principles (constitution) still serving
↪  us well?
- Should any agent's prompt or tools be updated?

### Step 2: Gather Team Feedback

Use `view_reflection()` to see team improvement
↪  suggestions from your team members. **Read ALL
↪  suggestions before taking any action** -- workers may
↪  submit suggestions after you start reading.

### Step 3: Decide and Act

Based on **your evaluation (Step 1)** and **team
↪  suggestions (Step 2)**, propose changes.

**Prefer structural changes** (adding/removing agents)
↪  over rule changes (constitution/prompt edits) -- a
↪  new agent with its own step budget is a more durable
↪  improvement than a new paragraph in the constitution.

1. Use `view_current_config(target_file)` to see current
↪  content before editing
2. Use `propose_reflection(target_file, content, reason)`
↪  for each change
3. Call `apply_reflection(confirmation="apply")` once to
↪  apply all proposals

**Skip only if** BOTH: your team submitted zero
↪  suggestions AND you yourself cannot identify any
↪  improvement. Call `skip_reflection(reason)` with a
↪  specific explanation.

---

### How to Add a New Agent

Propose both config.yaml and prompt.md:

**config.yaml** -- copy, modify, and propose. **Use the
↪  SAME `model:` value your existing teammates use**
↪  (check their config.yaml via `view_current_config` if
↪  unsure) -- do not invent or substitute a different
↪  model:

```
propose_reflection(
    target_file="<name>/config.yaml",
    content="name: <agent_name>\ndescription: \"One
    ↪  sentence about this agent's role\"\nmodel:
    ↪  []\ntemperature: []\nmax_tokens: []\nmax_steps:
    ↪  []\ntools:\n  - []\n  - read_file",
    reason="..."
)
```

prompt.md -- copy, modify, and propose:
```
propose_reflection(
    target_file="<name>/prompt.md",
    content="# <Agent Name>\n\nYou are a <role
    ↪  description>.\n\n## Your Mission\n\n<What this
    ↪  agent does and why it exists.>\n\n##
    ↪  Workflow\n\n1. Wait for instructions from the
    ↪  planner via message\n2. <Main work steps>\n3.
    ↪  Report results back to the planner via
    ↪  send_message\n\n## Key Principles\n\n- <Principle
    ↪  1>\n- <Principle 2>",
    reason="..."
)
```

How to Remove an Agent

```
propose_reflection(target_file="<name>/config.yaml",
↪  content="__REMOVE__", reason="...")
```

Cannot remove the Chairman.

How to Modify Constitution

When modifying `constitution.md`:
- **Preserve existing foundational sections** (Mission,
↪ Environment, Hard Constraints)
- **Add or update rules** based on lessons learned --
↪ also prune rules that are no longer relevant
- Include ALL existing content plus your changes
↪ (propose_reflection requires full file content)
- Always `view_current_config("constitution.md")` before
↪ proposing

Rules
- **Always view before edit** -- Use
↪ `view_current_config` before every
↪ `propose_reflection`
- **Evidence-based** -- Reference specific observations
↪ from recent tasks
- **Batch proposals** -- Submit ALL `propose_reflection`
↪ calls first, then call `apply_reflection` exactly
↪ once
- Changes are written to the **session copy** -- original
↪ source files are not modified until persistence

CRITICAL
You MUST complete L3 by calling one of:
↪ `propose_reflection` + `apply_reflection`, or
↪ `skip_reflection`. Do NOT reply with plain text --
↪ use a tool call.
```

Decision: collaboration

```
# source: scripts/train_agora.py:1221-1247,1468-1509
↪  (_L3_DECISION_INSTRUCTIONS, _llm_decide_l3; called
↪  from run_l3_gate_single_check:1512-1583)
# condition: collaboration  (--l3-gate-mode single_check;
↪  end of every training step)

=== PROMPT (user message) ===
You are aggregating your team's own reviews of the
↪  structural changes proposed this step; you do not
↪  judge the proposals yourself. For each proposal,
↪  classify every teammate review listed under it as
↪  SUPPORT (agrees, or says it looks fine), OPPOSE
↪  (argues against it) or NEUTRAL ("no_response", or no
↪  clear stance). Then apply this rule exactly: if
↪  OPPOSE reviews outnumber SUPPORT reviews, drop the
↪  proposal; otherwise keep it (a tie, or no reviews at
↪  all -> keep). Order the kept proposals by SUPPORT
↪  minus OPPOSE, highest first.
```

```
Proposals:
{proposal_list}

Reply with ONLY a JSON object: {"keep": [indices of
↪  proposals to keep], "order": [the kept indices,
↪  highest priority first]}. Example: {"keep": [2, 0],
↪  "order": [2, 0]}. An empty keep list is allowed.

=== FORMAT OF EACH ENTRY IN {proposal_list} (entries
↪  separated by a blank line) ===
[{index}] {unit_label}
  proposed_by: {proposed_by}
  reason: {target_file}: {reason}; {target_file_2}:
  ↪  {reason_2}
  teammate reviews:
    - {teammate_name}: "{challenge_reflection_comment}"
    - {teammate_name}: (no response)

(If a proposal has no reviews at all, the review block is
↪  the single line
"  teammate reviews: (none)".)
{unit_label} is one of: "add <role>" (new role: its
↪  config.yaml + prompt.md as one
unit), "remove <role>", "patch <role>" (config.yaml
↪  and/or prompt.md of an existing
role), "patch constitution.md", "patch pool.yaml".

=== CALL / RESPONSE HANDLING ===
One single-turn user message (no system prompt) on the
↪  chairman's own model
(the `model:` of <chairman>/config.yaml), temperature
↪  0.0, max_tokens 300.
Expected response: {"keep": [...], "order": [...]} (a
↪  bare JSON list is also
accepted and read as keep, in order). If the reply cannot
↪  be parsed, every
proposal is kept in arrival order. Kept proposals are
↪  then capped in this
order (max 2 add/remove units, max 2 patch groups per
↪  step) and tested together
in ONE paired held-out check (current team vs. changed
↪  team, same tasks); the
changes are applied unless the changed team is clearly
↪  worse (one-sided
binomial test, alpha = 0.3; not applied if either side
↪  has fewer than 10 sessions that complete without
↪  crashing).
```

Decision: conformity

```
# source: scripts/train_agora.py:1221-1247,1468-1509
↪  (_L3_DECISION_INSTRUCTIONS, _llm_decide_l3; called
↪  from run_l3_gate_single_check:1512-1583)
# condition: conformity  (--l3-gate-mode single_check;
↪  end of every training step)

=== PROMPT (user message) ===
You are {chairman_name}, the team's chairman. This step
↪  your team collected the structural changes listed
↪  below. As chairman you alone decide which to adopt.
↪  Keep a proposal if you judge its stated reason sound
↪  and it does not conflict with another kept proposal;
↪  drop it otherwise. Order the kept proposals by
↪  priority.

Proposals:
{proposal_list}

Reply with ONLY a JSON object: {"keep": [indices of
↪  proposals to keep], "order": [the kept indices,
↪  highest priority first]}. Example: {"keep": [2, 0],
↪  "order": [2, 0]}. An empty keep list is allowed.

=== FORMAT OF EACH ENTRY IN {proposal_list} (entries
↪  separated by a blank line) ===
[{index}] {unit_label}
  proposed_by: {proposed_by}
  reason: {target_file}: {reason}; {target_file_2}:
  ↪  {reason_2}

(No review block under conformity: there is no
↪  challenge_reflection tool.)
{unit_label} is one of: "add <role>" (new role: its
↪  config.yaml + prompt.md as one
unit), "remove <role>", "patch <role>" (config.yaml
↪  and/or prompt.md of an existing
role), "patch constitution.md", "patch pool.yaml".

=== CALL / RESPONSE HANDLING ===
One single-turn user message (no system prompt) on the
↪  chairman's own model
(the `model:` of <chairman>/config.yaml), temperature
↪  0.0, max_tokens 300.
Expected response: {"keep": [...], "order": [...]} (a
↪  bare JSON list is also
accepted and read as keep, in order). If the reply cannot
↪  be parsed, every
proposal is kept in arrival order. Kept proposals are
↪  then capped in this
order (max 2 add/remove units, max 2 patch groups per
↪  step) and tested together
in ONE paired held-out check (current team vs. changed
↪  team, same tasks); the
changes are applied unless the changed team is clearly
↪  worse (one-sided
binomial test, alpha = 0.3; not applied if either side
↪  has fewer than 10
sessions that complete without crashing).
```

Decision: competition

```
# source: scripts/train_agora.py:1221-1247,1468-1509
↪  (_L3_DECISION_INSTRUCTIONS, _llm_decide_l3; called
↪  from run_l3_gate_single_check:1512-1583)
# condition: competition  (--l3-gate-mode single_check;
↪  end of every training step)
# {chairman_name} = the current chairman, i.e. the role
↪  holding Tier 1 after the
# previous step's competition rotation. Before this call,
↪  any proposal whose
# pool.yaml would name a different chairman is dropped
# (reason "chairman_set_by_competition_tiers").

=== PROMPT (user message) ===
You are {chairman_name}, the team's chairman, holding the
↪  role because your contribution scored highest. This
↪  step your team collected the structural changes
↪  listed below; teammate reviews are attached as
↪  reference. You alone decide which to adopt. Weigh the
↪  reviews, but the call is yours. Keep or drop each
↪  proposal and order the kept ones by priority.

Proposals:
{proposal_list}

Reply with ONLY a JSON object: {"keep": [indices of
↪  proposals to keep], "order": [the kept indices,
↪  highest priority first]}. Example: {"keep": [2, 0],
↪  "order": [2, 0]}. An empty keep list is allowed.

=== FORMAT OF EACH ENTRY IN {proposal_list} (entries
↪  separated by a blank line) ===
[{index}] {unit_label}
  proposed_by: {proposed_by}
  reason: {target_file}: {reason}; {target_file_2}:
  ↪  {reason_2}
  teammate reviews:
    - {teammate_name}: "{challenge_reflection_comment}"
    - {teammate_name}: (no response)

(If a proposal has no reviews at all, the review block is
↪  the single line
"  teammate reviews: (none)".)
{unit_label} is one of: "add <role>" (new role: its
↪  config.yaml + prompt.md as one
unit), "remove <role>", "patch <role>" (config.yaml
↪  and/or prompt.md of an existing
role), "patch constitution.md", "patch pool.yaml".

=== CALL / RESPONSE HANDLING ===
One single-turn user message (no system prompt) on the
↪  chairman's own model
(the `model:` of <chairman>/config.yaml), temperature
↪  0.0, max_tokens 300.
Expected response: {"keep": [...], "order": [...]} (a
↪  bare JSON list is also
accepted and read as keep, in order). If the reply cannot
↪  be parsed, every
proposal is kept in arrival order. Kept proposals are
↪  then capped in this
```

```
order (max 2 add/remove units, max 2 patch groups per
↪  step) and tested together
in ONE paired held-out check (current team vs. changed
↪  team, same tasks); the
changes are applied unless the changed team is clearly
↪  worse (one-sided
binomial test, alpha = 0.3; not applied if either side
↪  has fewer than 10
sessions that complete without crashing).
```

Chairman change under competition. When another role becomes the chairman, each agent receives a notice; only the authority to lead passes to the new chairman.

Chairman change notice

```
# source: core/runner.py:1158-1205
↪  (Runner._leadership_note; appended as the last
↪  "---"-separated section of the session context by
↪  _build_chairman_context:1086-1138 and
↪  _build_agent_context:1207-1243). Trigger:
↪  scripts/train_agora.py:140-158 (write_pool_chairman
↪  records `original_chairman` in pool.yaml) and
↪  :616-628 (competition promotes the Tier-1 role to
↪  chairman).
# condition: competition only, and only while a promotion
↪  is in effect (current chairman != pool.yaml
↪  `original_chairman`, and the original role still
↪  exists). Otherwise the notice is empty and no section
↪  is added.
# Placeholders: {new_chairman} = current (promoted)
↪  chairman; {original_chairman} = pool.yaml
↪  original_chairman.
# {duties} is the fixed string:
#   receiving the task, recruiting and dismissing
↪  teammates (`start_agent`/`stop_agent`), and
↪  submitting the team's final answer (`finalize_task`)

=== BRANCH A: recipient is the promoted chairman
↪  ({new_chairman}) ===
## Chairman change

You have been promoted to Chairman for this task
↪  (competition). You keep your own role as
↪  {new_chairman}; {original_chairman} keeps its own
↪  role as well. The Chairman duties -- receiving the
↪  task, recruiting and dismissing teammates
↪  (`start_agent`/`stop_agent`), and submitting the
↪  team's final answer (`finalize_task`) -- are now
↪  yours. Where the constitution or the instructions
↪  call {original_chairman} the Chairman, or give
↪  {original_chairman} those duties, they now mean you.
↪  Everything else stays as written: the steps the
↪  workflow assigns to {original_chairman} are still
↪  done by {original_chairman}, so recruit
↪  {original_chairman} and pass it the task where the
↪  workflow starts with it, recruit the teammates it
↪  asks for, and submit the final answer yourself.

=== BRANCH B: recipient is the original chairman
↪  ({original_chairman}) ===
## Chairman change

The Chairman for this task is **{new_chairman}**
↪  (promoted in competition), not you. You keep your own
↪  role as {original_chairman}. The Chairman duties --
↪  receiving the task, recruiting and dismissing
↪  teammates (`start_agent`/`stop_agent`), and
↪  submitting the team's final answer (`finalize_task`)
↪  -- belong to {new_chairman}: ask {new_chairman} to
↪  recruit the teammates you need, and where your
↪  instructions say you submit the final answer, send it
↪  to {new_chairman} instead. Everything else in your
↪  instructions stays as written.

=== BRANCH C: recipient is any other teammate ===
## Chairman change

The Chairman for this task is **{new_chairman}**
↪  (promoted in competition); {original_chairman}, named
↪  as Chairman in the constitution and instructions,
↪  keeps its own role as {original_chairman}. The
↪  Chairman duties -- receiving the task, recruiting and
↪  dismissing teammates (`start_agent`/`stop_agent`),
↪  and submitting the team's final answer
↪  (`finalize_task`) -- belong to {new_chairman}: where
↪  the constitution or your instructions mean the
↪  Chairman, read {new_chairman}. Everything else stays
↪  as written, including your own role and which
↪  teammate you send each kind of work to.
```

L.4 Meta-Team Self-Reflection

Under Meta-Team, each agent edits its own prompt patches and skills after a session with the prompt below.

L1 self-reflection prompt

```
## L1 Reflection -- Individual Agent Self-Improvement

You are now in the **L1 Reflection phase**. The task has
↪  been completed. Reflect on YOUR OWN performance
↪  during this task and identify improvements for the
↪  future.

### Goal
Improve your own prompt and skills that would make you
↪  more effective in future tasks.

### SUBMISSION IS LOCKED

**Your code submission has been frozen.** The patch for
↪  evaluation was captured before this reflection phase
↪  began. Any changes you make to the codebase now will
↪  NOT affect the evaluation score. **Do NOT edit source
↪  code, write new code, or attempt to fix bugs during
↪  reflection.** This phase is purely for learning and
↪  self-improvement.

**If the task was truly trivial** (required no
↪  significant decisions or effort), you can choose to
↪  skip L1. But if you applied any noteworthy strategy
↪  or learned something useful, record it even if the
↪  task succeeded.

### Step 1: Self-Reflection

Reflect deeply on your own behavior during this task. You
↪  may use read-only environment commands as evidence to
↪  support your reflection, but the focus is on your
↪  **decision-making process**, not on diagnosing the
↪  code.

Think honestly about:
1. **What did I do well?** -- Which decisions,
↪  strategies, or habits contributed positively?
↪  Learning from success is just as important as
↪  learning from failure.
2. **Where did I go wrong?** -- What mistakes did I make?
↪  Where did my reasoning break down? Did I make
↪  incorrect assumptions?
3. **What was inefficient?** -- Even if the task
↪  succeeded, did I waste steps? Over-explore?
↪  Under-delegate? Miss a faster path?
4. **What would I do differently next time?** -- If I
↪  faced a similar (but not identical) task tomorrow,
↪  what would I change about my strategy?

### Step 2: Identify Reusable Patterns

Your goal is to identify **patterns and strategies** that
↪  will help on future similar tasks -- not to record
↪  what happened on this specific task.

**The key question**: Would this insight help me if I
↪  encountered a similar (but not identical) task in
↪  this codebase? If yes, it's worth recording.
```

```
**What to avoid**: Don't record one-off details about
↪  this specific bug or this specific file. If your
↪  patch only makes sense in the context of this exact
↪  task, it's too narrow.

**What's valuable**: Strategies, workflows, verification
↪  steps, and patterns that you'd want to remember
↪  across many tasks in this codebase.

**Even if the task succeeded**, reflect on whether your
↪  approach was efficient. Could you have solved it in
↪  fewer steps? Did you explore unnecessary files? Was
↪  there a faster path?

### Step 3: Apply Improvements (only if needed)

Based on your reflection, use the available tools:

- **`update_prompt_patch(action, patch/patches)`** --
↪  Manage your behavioral patches (rules injected into
↪  your system prompt in future sessions).
  - `action="add", patch="..."` -- Append one new patch
  ↪  to your current list
  - `action="replace_all", patches=[...]` -- Replace the
  ↪  entire list (use to merge duplicates, remove
  ↪  outdated patches, or reorganize)
  - The tool always shows your **FULL current patch
  ↪  list** after the operation, so you can see what you
  ↪  already have.
  - **Review your existing patches** before adding new
  ↪  ones. If you have 5+ patches, consider using
  ↪  `replace_all` to merge related ones.

- **`update_skill(skill_name, content)`** -- Create or
↪  update a reusable skill file. **Skills are your
↪  PRIMARY self-improvement tool** -- use them for:
  - Multi-step procedures or checklists you discovered
  ↪  during the task
  - Reusable workflows that would help in similar future
  ↪  tasks
  - Complex decision trees or validation sequences

- `skip_l1_reflection(reason)` -- Call this to **mark L1
↪  as complete**. Use it after making any updates, or
↪  immediately if no improvements are needed.

**IMPORTANT**: `update_prompt_patch` and `update_skill`
↪  do NOT mark L1 as complete. You can call them
↪  multiple times. When you are finished with ALL your
↪  updates, you MUST call `skip_l1_reflection` to mark
↪  completion.

### Patch Quality Guide

Your patches should capture **reusable patterns** --
↪  things you'd want to remember for future tasks in
↪  this codebase.

**BAD** -- one-off details about this specific task, for
↪  example:
- "The bug was in line 42 of foo.py where the return type
↪  was wrong"
- "Developer forgot to check strategy plugins for
↪  PlayIterator" -> teammate observation, belongs in L2

**GOOD** -- reusable work patterns you'd want to follow
↪  on future tasks, for example:
- "After implementing a fix, run `git diff --stat` and
↪  verify the changed files match your intent -- if a
↪  file you edited doesn't appear, the edit wasn't
↪  applied"
- "When a PR changes a function signature, grep for all
↪  callers and update every one before testing --
↪  partial caller updates cause partial test failures"
- "Run the failing tests locally before implementing to
↪  verify they actually fail -- if a test doesn't exist
↪  locally, the test file may need to be created as part
↪  of the task"
- "In projects with an installed package and a source
↪  tree, check which one Python imports before editing
↪  -- edits to the wrong copy have no effect"

### Skills > Patches

When your improvement involves a **multi-step workflow**,
↪  **checklist**, or **decision procedure**, create a
↪  **skill** instead of a text patch. Skills are
↪  structured, reusable, and more powerful than
↪  single-line rules.

**Use a patch** for: simple behavioral rules ("Always X
↪  before Y")
**Use a skill** for: procedures with multiple steps,
↪  conditional logic, or detailed checklists

### Rules
- **Do NOT modify code** -- The submission is locked. No
↪  editing files, no writing code, no running tests.
↪  Read-only diagnosis only.
- **Be efficient** -- If nothing noteworthy happened,
↪  call `skip_l1_reflection` to skip this phase
↪  immediately.
- Focus on **your own behavior** -- later phases will
↪  handle teammate observations
- Patches are **additive** -- they supplement your base
↪  prompt, not replace it
- **Do NOT send messages** to other agents during L1 --
↪  save collaboration feedback for L2
- **Do NOT spend too much time on environment
↪  interaction** -- This is reflection, not a debugging
↪  session.

### CRITICAL
**You MUST call `skip_l1_reflection`** after making all
↪  desired updates (or immediately if no improvements
↪  are needed). After calling it, you will wait in idle
↪  until the system advances to L2.
```

M Usage of AI Assistants

All content of this paper was written by the authors. We used an AI assistant (ChatGPT) only to polish the language, such as improving clarity, correcting grammar, and enhancing readability. It was not involved in developing research ideas, designing experiments, interpreting results, or generating any content.