arXiv:2610.19638v1 [cs.CL] 14 Oct 2026

J-tau: A Fully Localized Japanese Benchmark for Tool-Agent-User Interaction

Chihiro Yano[1,2]*, **Jun Hirako**[1], **Hiroaki Yamagiwa**[1], **Ryota Hirobuchi**[1]
[1]SB Intuitions Corp., [2]Nagoya University
{chihiro.yano, jun.hirako}@sbintuitions.co.jp,
{hiroaki.yamagiwa, ryota.hirobuchi}@sbintuitions.co.jp

Abstract

We present J-tau, the first publicly available Japanese benchmark for evaluating tool-using LLM agents in interactive customer-support scenarios. While tau-bench is widely used to evaluate LLM agents, it is English-centric and cannot evaluate agents in other languages. Existing multilingual extensions offer no public Japanese version and rely primarily on machine translation, leaving locale-specific entities such as addresses and currencies inconsistent with the target region. To enable reliable evaluation in Japanese, we carefully localize tau-bench through machine translation, entity adaptation, and human review. During this process, we also uncover and fix two flaws in the original tau-bench: invalid tasks and a score aggregation error that inflates certain models' results. Experiments show that J-tau yields fewer evaluation errors than a naively machine-translated counterpart and that our fixes resolve both flaws.

1 Introduction

Large language model (LLM) agents that autonomously use tools to solve complex tasks have attracted growing attention (Wang et al., 2024). While many benchmarks have been proposed to evaluate such agents, most of them focus on English. Because agent performance can vary across languages (Hofman et al., 2026), English evaluations may not reflect utility in other languages.

To address this gap, multilingual agent benchmarks have been proposed, including ones built from multilingual corpora and machine translations of English benchmarks (Kulkarni et al., 2025; Nguyen et al., 2026). However, machine translation alone cannot localize entities such as personal names, which can lead to unnatural conversations or even invalid tasks (Kim et al., 2026).

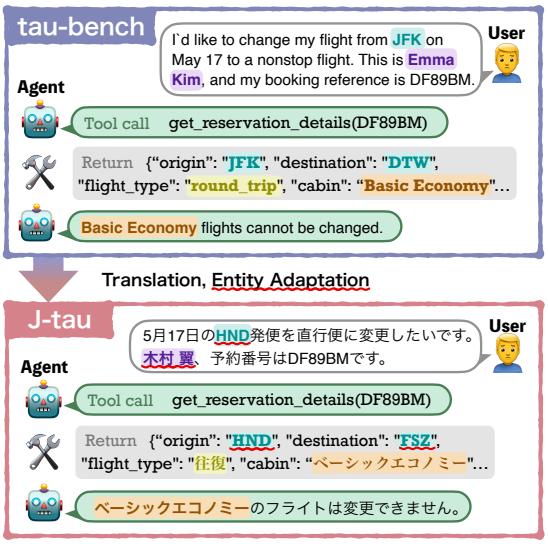

Figure 1: Example of Japanese localization in J-tau. We translate the original tau-bench and replace locale-specific entities (marked with wavy underlines) with realistic Japanese counterparts. Examples include airport codes (JFK → HND, DTW → FSZ) and personal names (Emma Kim → 木村翼, where Kimura (木村) is the family name and Tsubasa (翼) is the given name).

To overcome these limitations, we present and release J-tau, a Japanese extension of the widely used English benchmark τ-bench (tau-bench) (Yao et al., 2025; Barres et al., 2026; Cuadron et al., 2025). J-tau is the first publicly available Japanese benchmark that evaluates tool-using LLM agents in realistic, interactive scenarios. It features high-quality localization (Fig. 1) and corrected tasks and evaluation. Our main contributions are as follows:

- We construct and release J-tau, a Japanese benchmark that localizes three domains of tau-bench.
- We identify two flaws in tau-bench and fix them in J-tau: invalid tasks (about 10% in airline and retail) and score-inflating aggregation (up to 64.6 points for a single model).
- We show that an existing machine-translation-based pipeline often causes evaluation failures

*Work done during an internship at SB Intuitions Corp.
We release J-tau at https://github.com/sbintuitions/j-tau-bench.

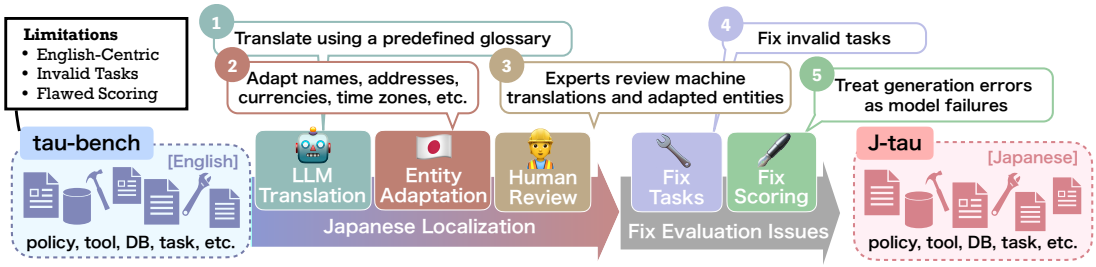

Figure 2: Overview of the J-tau pipeline, which localizes tau-bench into Japanese and fixes its evaluation issues.

that stem from the framework rather than agents.

2 Background: tau-bench

This section describes tau-bench and its limitations, along with prior work addressing them.[1]

2.1 Overview

tau-bench is a benchmark for evaluating LLM agents on customer support tasks across multiple domains, including airline (booking flights), retail (modifying orders), and telecom (troubleshooting mobile phones). In each task, an LLM-simulated user follows a user instruction describing its goal, and the agent must resolve the request conversationally, calling tools while complying with a domain-specific policy. In telecom, as in real-world technical support, the user has a separate set of tools, so the agent must also guide the user in using them.

2.2 Limitations

tau-bench has three main limitations. First, it is English-centric. While prior work extends it to five other languages[2] (Nguyen et al., 2026), this extension relies primarily on machine translation, which degrades data quality and leads to unnatural conversations.[3] Second, although tau-bench has been revised since its release (Cuadron et al., 2025), it still contains invalid tasks that are unsolvable or have non-deterministic ground truths. Third, its score aggregation is flawed and overestimates the performance of some models.

3 Japanese Agent Benchmark: J-tau

This section describes the construction pipeline and details of J-tau, a Japanese agent benchmark that addresses the limitations of tau-bench.[4]

3.1 Construction Pipeline

Our pipeline consists of five steps (see Fig. 2): three for Japanese localization (translation, entity adaptation, and human review) and two for fixing evaluation issues (task fixes and scoring fixes).

Step 1: LLM Translation. We translate the policies, tool definitions, databases, and user instructions from English into Japanese using Gemini 3.1 Pro (Google DeepMind, 2026). To ensure consistent translation of technical terms, we include a predefined domain-specific glossary in the prompt. Appendix B.1 shows the domain-specific glossaries.

Step 2: Entity Adaptation. Direct translation retains locale-specific entities from English-speaking regions, such as personal names, addresses, and phone numbers, making the simulated conversations unnatural for Japanese speakers. We therefore replace them with fictitious Japanese counterparts. We generate them with `Faker`[5] for the retail and airline domains, which contain many users, and create them manually for the telecom domain. We also convert monetary amounts into Japanese yen (1 USD = 150 JPY) and the time zone from EST to JST (see Appendix B.2 for more details).

Step 3: Human Review. We then assess and correct the translation quality of the resulting data. For each domain, three experts review the policies, and for the airline and retail domains, experts also review the user instructions (see Appendix B.3).

Step 4: Fix Tasks. We verify each task through manual inspection and test runs, and identify tasks

[1] See Appendix A for additional related work.

[2] Mandarin, Vietnamese, Thai, Indonesian, and Filipino.

[3] For example, users converse in a Southeast Asian language, yet addresses and currencies remain those of the U.S.

[4] Except for the language, evaluation fixes, and user simulator, we follow tau-bench's settings (e.g., prompts).

[5] https://github.com/joke2k/faker

that cannot be solved.[6] We find two types of problematic tasks: tasks that become unsolvable through localization, and invalid tasks already present in the original tau-bench, which are unsolvable or have non-deterministic ground truths.[7] We fix these tasks in J-tau, for example by adding missing information to the user instructions (see Appendix B.4).

Step 5: Fix Scoring. tau-bench treats both empty responses and tool calls with malformed JSON as infrastructure errors. It retries them up to three times and excludes trials with persistent errors from aggregation. However, agent errors reflect the agent's limitations, so this procedure inflates the scores of models prone to such errors.[8] In J-tau, we instead attribute these errors to the model that produced them. Trials are counted as failures for agent errors and discarded for user simulator errors.

3.2 Benchmark Details

Constructed through this pipeline, J-tau is a high-quality Japanese agent benchmark covering the airline, retail, and telecom domains.[9]

In interactive agent evaluation, the choice of user simulator can substantially affect agent scores (Seshadri et al., 2026). We therefore compare several recent LLMs in terms of Japanese naturalness and instruction-following stability, and adopt GLM-5.2-FP8 (GLM-5-Team et al., 2026), which distorted evaluation results the least, as the default user simulator (see Appendix I).

4 Evaluating LLM Agents on J-tau

In this section, we evaluate and analyze a range of proprietary and open-weight models on J-tau.

Setup. We run three trials per task and report the mean and standard deviation.[10] Appendix E provides details of the evaluated models.

[6]We verify only the airline and retail domains, since telecom tasks are programmatically generated and always solvable.

[7]For example, when the user instruction lacks the information needed to choose among several payment cards, the agent must pick one, and whether it matches the card in the reference solution is a matter of chance rather than model capability.

[8]For example, although Qwen3.5-4B ended 75/150, 56/342, and 38/342 trials with empty responses in airline, retail, and telecom, respectively, it scored 93.3, 85.3, and 98.7.

[9]Appendix D reports the number of tasks in each domain.

[10]For the retail domain, we use `gpt-oss-120b` (OpenAI, 2025) as the LLM judge for conversation-based evaluation.

Model	airline	retail	telecom
Proprietary Models			
GPT-5.6 Sol	84.7±2.3	83.0±4.0	90.9±1.8
Gemini 3.1 Pro	84.0±4.0	78.9±3.2	79.5±4.5
Claude Opus 4.8	81.3±3.1	77.5±2.8	84.5±2.2
Claude Fable 5	90.7±1.2	84.5±4.1	99.4±0.5
Open-Weight Models			
Kimi-K3	70.0±2.0	49.1±7.0	67.0±3.5
GLM-5.2-FP8	89.3±4.6	81.6±1.8	98.2±1.5
gpt-oss-120b	73.3±4.2	75.7±0.5	64.6±1.0
gpt-oss-20b	50.7±3.1	51.2±3.3	59.1±3.3
Qwen3.5-397B-A17B-FP8	83.3±6.4	85.1±2.3	96.2±0.5
Qwen3.5-27B	84.7±4.2	75.1±1.3	89.8±8.7
Qwen3.5-9B	49.3±2.3	35.7±7.3	50.0±9.2
Qwen3.5-4B	24.0±2.0	6.7±2.2	33.6±9.9
gemma-4-31b-it	51.3±6.4	72.8±4.9	44.4±9.3
gemma-4-12b-it	41.3±3.1	34.5±3.5	23.1±5.1
gemma-4-E4B-it	38.7±5.0	39.2±5.3	20.2±5.8
gemma-4-E2B-it	20.7±4.2	26.3±0.9	13.2±4.0
llm-jp-4.1-33b-thinking	64.7±3.1	68.4±2.6	72.5±5.1
llm-jp-4.1-8b-thinking	56.0±3.5	63.7±3.7	52.0±5.6

Table 1: Results on J-tau. We report the mean and standard deviation over three trials.

Quantitative Results. Table 1 shows the results for each domain of J-tau. Among proprietary models, Claude Fable 5 (Anthropic, 2026a) achieves the highest overall score. Among open-weight models, GLM-5.2-FP8 and Qwen3.5-397B-A17B-FP8 (Qwen Team, 2026) perform best. Within the Qwen3.5 series, performance generally improves with model size.

Case Study. Appendix F presents example conversations from GLM-5.2-FP8, the open-weight model with the highest J-tau score, and failure cases from gemma-4-E2B-it (Farabet and Lacombe, 2026) and Qwen3.5-4B, the two lowest-scoring ones. For example, Qwen3.5-4B sometimes switches from Japanese to another language partway through the conversation and fails to recognize the family-name-first order of Japanese names, such as Kimura Tsubasa (木村翼) in Fig. 1.

5 Quality Analysis of J-tau

5.1 Effect of Task and Scoring Fixes

We conduct an ablation study to assess the effectiveness of the task and scoring fixes in our pipeline in Section 3.1. To this end, we use tau-bench* (tau*), a variant of tau-bench (tau) to which only the scoring fix in Step 5 is applied.

Fix Tasks. To examine the effect of the task fixes in Step 4, we split tasks into Repaired and Unmodified groups. Table 2 compares the J-tau and tau* scores of each group for each domain.

Model	airline				retail				telecom	
	Repaired (6)		Unmodified (41)		Repaired (12)		Unmodified (80)		Unmodified (114)	
	J-tau	tau*	J-tau	tau*	J-tau	tau*	J-tau	tau*	J-tau	tau*
GPT-5.6 Sol	77.8 (+50.0)	27.8	85.4 (−9.8)	95.1	61.1 (+2.8)	58.3	85.4 (+2.9)	82.5	90.9 (+5.8)	85.1
Kimi-K3	50.0 (+5.6)	44.4	73.2 (−13.8)	87.0	44.4 (−11.1)	55.6	47.5 (−36.7)	84.2	67.0 (−21.9)	88.9
GLM-5.2-FP8	77.8 (+44.4)	33.3	91.1 (−3.3)	94.3	77.8 (+33.3)	44.4	80.4 (−2.5)	82.9	98.2 (±0.0)	98.2
gpt-oss-120b	44.4 (+22.2)	22.2	77.2 (−4.9)	82.1	55.6 (−13.9)	69.4	76.7 (−9.6)	86.3	64.6 (−28.1)	92.7
gpt-oss-20b	11.1 (−22.2)	33.3	55.3 (−13.0)	68.3	36.1 (−5.6)	41.7	45.4 (−20.4)	65.8	59.1 (−17.5)	76.6
Qwen3.5-397B-A17B-FP8	72.2 (+27.8)	44.4	83.7 (−9.8)	93.5	88.9 (+19.4)	69.4	82.9 (−7.1)	90.0	96.2 (−2.3)	98.5
Qwen3.5-27B	66.7 (+16.7)	50.0	86.2 (±0.0)	86.2	75.0 (+8.3)	66.7	71.3 (−14.6)	85.8	89.8 (−7.6)	97.4
Qwen3.5-9B	22.2 (+5.6)	16.7	52.8 (−30.1)	82.9	16.7 (−38.9)	55.6	35.4 (−36.7)	72.1	50.0 (−30.1)	80.1
Qwen3.5-4B	0.0 (−16.7)	16.7	27.6 (−2.4)	30.1	0.0 (−27.8)	27.8	5.0 (−47.5)	52.5	33.6 (−18.4)	52.0
llm-jp-4.1-33b-thinking	33.3 (−11.1)	44.4	71.5 (+0.8)	70.7	55.6 (−5.6)	61.1	67.9 (−7.9)	75.8	72.5 (−8.2)	80.7
llm-jp-4.1-8b-thinking	11.1 (−11.1)	22.2	61.8 (−2.4)	64.2	61.1 (+2.8)	58.3	61.3 (−5.0)	66.3	52.0 (−23.4)	75.4

Table 2: Effect of task fixes on Repaired and Unmodified tasks (mean over three trials). Differences of J-tau from tau* are shown in red (higher) and blue (lower). Subscripts next to group names denote task counts.

Model	airline		retail		telecom	
	tau*	tau	tau*	tau	tau*	tau
GPT-5.6 Sol	87.3	85.3	81.3	84.5	85.1	82.7
Kimi-K3	80.0	72.7	82.5	83.0	88.9	91.2
GLM-5.2-FP8	87.3	89.3	80.7	84.2	98.2	98.8
gpt-oss-120b	73.3	76.7	85.7	81.0	92.7	92.4
gpt-oss-20b	62.7	68.0	64.3	75.7	76.6	83.0
Qwen3.5-397B-A17B-FP8	86.7	86.7	88.9	88.3	98.5	99.7
Qwen3.5-27B	81.3	84.7	85.4	86.5	97.4	99.1
Qwen3.5-9B	74.0	77.3	72.2	86.7	80.1	96.5
Qwen3.5-4B	28.7	93.3	50.9	85.3	52.0	98.7
llm-jp-4.1-33b-thinking	66.0	64.7	74.6	78.4	80.7	83.6
llm-jp-4.1-8b-thinking	57.3	60.0	68.7	68.7	75.4	73.7

Table 3: Effect of the scoring fix, comparing tau and tau* for each domain (mean over three trials).

Model	Domain	Bench	Pass	Fail		
				agent	*user*	*eval*
GLM-5.2-FP8	airline	J-tau	134	7	5	4
		$\text{SEATau}_{\text{ja}}$	136	6	1	7
	retail	J-tau	279	23	36	4
		$\text{SEATau}_{\text{ja}}$	256	18	14	54
	telecom	J-tau	336	4	2	0
		$\text{SEATau}_{\text{ja}}$	331	10	0	1
Qwen3.5-397B-A17B-FP8	airline	J-tau	125	20	3	2
		$\text{SEATau}_{\text{ja}}$	121	20	5	4
	retail	J-tau	291	29	16	6
		$\text{SEATau}_{\text{ja}}$	236	36	14	56
	telecom	J-tau	329	11	0	2
		$\text{SEATau}_{\text{ja}}$	329	12	1	0

Table 4: Successful (Pass) and failed (Fail) trial counts on J-tau and $\text{SEATau}_{\text{ja}}$. GPT-5.6 Sol (OpenAI, 2026) classifies each failure's cause as *agent*, *user*, or *eval*.

On the Repaired tasks, high-performing models score substantially higher on J-tau than on tau*. This is because the invalid tasks in tau* cap the achievable accuracy regardless of model capability, and our fixes remove this cap. On the Unmodified tasks, which allow a direct comparison between languages, J-tau scores tend to be lower than tau* scores. This is consistent with Hofman et al. (2026), who show that agent performance varies across languages.

Fix Scoring. To examine the effect of the scoring fix in Step 5, Table 3 compares scores on tau and tau*. For high-performing models, the difference is small. However, for small models such as Qwen3.5-4B, scores drop substantially on tau*, which counts empty and malformed outputs as failures. Thus, tau-bench inflates the scores of such models.

5.2 Comparison with an MT-Centric Pipeline

We compare J-tau with $\text{SEATau}_{\text{ja}}$, which machine-translates tau-bench into Japanese with the pipeline of Nguyen et al. (2026) and applies the scoring fix in Step 5, without entity adaptation, human review, or task fixes. In Table 4, an LLM judge attributes each failed trial to the agent, user simulator, or evaluation framework. $\text{SEATau}_{\text{ja}}$ has many framework-caused failures and fewer successes, especially in retail.[11] Without entity adaptation, user names remain in English in the database but are written in katakana in the Japanese conversation, which causes database lookup failures (see the simulation in Table 28). Such inconsistencies induce unnatural behavior even in successful trials, making the environment unsuitable as a benchmark. In contrast, J-tau has fewer such failures thanks to its localization and task verification. Appendix G provides experimental details and example simulations.

[11] This is mainly because tool descriptions are in Japanese, while tool implementations still expect English arguments.

6 Conclusion

We presented J-tau, the first publicly available Japanese benchmark for interactive tool-using LLM agents. Beyond machine translation, we localized tau-bench through entity adaptation and human review, and fixed its invalid tasks and a scoring flaw, enabling more accurate evaluation. Our analysis showed that machine translation alone increases framework-caused failures. We hope J-tau advances Japanese LLM agents and our insights benefit benchmark localization for other languages.

Limitations

- J-tau covers three domains (airline, retail, and telecom) and does not include the banking_knowledge domain introduced in the latest version of tau-bench (Shi et al., 2026).
- J-tau covers only Japanese and cannot be used to evaluate agents in other languages. Concurrent work released within one month before our submission extends tau-bench to more languages (Kuang et al., 2026; Ray et al., 2026b), but these benchmarks may still contain the invalid tasks and the aggregation flaw we identify (Appendix A). Since our task and scoring fixes are language-independent, extending our pipeline to other languages is left for future work.
- Addresses and other user information in the airline and retail domains are generated with `Faker` and thus include nonexistent place names and addresses, limiting their fidelity to real-world data.
- Because the solvability verification in Step 4 (Section 3.1) relies only on GLM-5.2-FP8, it may reflect biases specific to this model, and a small number of invalid tasks may remain undetected.
- As shown in Section 5.2, J-tau has more failures attributed to the user simulator than $\text{SEATau}_{\text{ja}}$, especially in the retail domain. This is likely because the user simulator follows the user simulation guideline less faithfully in Japanese than in English (see Appendix G.5 for an ablation study). Since this may change as models improve, revisiting the choice of user simulator is left for future work.
- The original tau-bench prompt instructs models to produce outputs in JSON format. During evaluation after constructing J-tau, we found that this instruction affects some models, such as Kimi-K3 (Kimi Team et al., 2026), more strongly in Japanese, causing them to output JSON strings instead of natural-language responses to the user. Since these outputs are counted as failures, the scores of such models may be underestimated (see Appendix H.2 for an ablation study). We plan to remove this instruction in the revised version.

Ethical Considerations

J-tau contains no personally identifiable information, as all user information, such as names, addresses, and phone numbers, is fictitious. The human review was conducted by in-house annotators as part of their regular duties, for which they received their standard salaries.

References

Thales Sales Almeida, João Guilherme Alves Santos, Thiago Laitz, and Giovana Kerche Bonás. 2025. Ticket-Bench: A kickoff for multilingual and regionalized agent evaluation. *arXiv preprint arXiv:2509.14477*.

Anthropic. 2026a. Claude Fable 5 and Claude Mythos 5.

Anthropic. 2026b. Introducing Claude Opus 4.8.

Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik R Narasimhan. 2026. τ^2-Bench: Evaluating conversational agents in a dual-control environment. In *Forty-third International Conference on Machine Learning: ICML 2026*.

Andrea Caciolai, Pere-Lluís Huguet Cabot, Chierh Cheng, Albert Ventayol-Boada, Gabriel Mejia Gonzalez, Christophe Ropers, Lucas Bandarkar, Sebastian Ruder, Darlene Sakakihara, Elliot Yun, Pierre Andrews, Grégoire Mialon, Romain Froger, and Marta R. Costa-jussà. 2026. OmnilingualGAIA2: Evaluating the multilingual gap in frontier AI agents. *arXiv preprint arXiv:2608.08775*.

Alejandro Cuadron, Pengfei Yu, Yang Liu, and Arpit Gupta. 2025. SABER: Small actions, big errors – safeguarding mutating steps in LLM agents. *arXiv preprint arXiv:2512.07850*.

Antoine Dussolle, A. Cardeña, Shota Sato, and Peter Devine. 2025. M-IFEval: Multilingual instruction-following evaluation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 6161–6176.

Clement Farabet and Olivier Lacombe. 2026. Gemma 4: Byte for byte, the most capable open models.

GLM-5-Team, Aohan Zeng, Xin Lv, Zhenyu Hou, Zhengxiao Du, Qinkai Zheng, Bin Chen, Da Yin, Chendi Ge, Chenghua Huang, Chengxing Xie, Chenzheng Zhu, Congfeng Yin, Cunxiang Wang, Gengzheng Pan, Hao Zeng, Haoke Zhang, Haoran Wang, Huilong Chen, and 167 others. 2026. GLM-5: from vibe coding to agentic engineering. *arXiv preprint arXiv:2602.15763*.

Google DeepMind. 2026. Gemini 3.1 Pro.

Omer Hofman, Jonathan Brokman, Oren Rachmil, Shamik Bose, Vikas Pahuja, Toshiya Shimizu, Trisha Starostina, Kelly Marchisio, Seraphina Goldfarb-Tarrant, and Roman Vainshtein. 2026. MAPS: A multilingual benchmark for agent performance and security. In *Findings of the Association for Computational Linguistics: EACL 2026*, pages 821–845.

Xu Huang, Wenhao Zhu, Hanxu Hu, Conghui He, Lei Li, Shujian Huang, and Fei Yuan. 2025. BenchMAX: A comprehensive multilingual evaluation suite for large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 16751–16774.

Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R Narasimhan. 2024. SWE-bench: Can language models resolve real-world GitHub issues? In *The Twelfth International Conference on Learning Representations: ICLR 2024*.

Yunsu Kim, Kaden Uhlig, and Joern Wuebker. 2026. GAIA-v2-LILT: Multilingual adaptation of agent benchmark beyond translation. In *Proceedings of the 1st Workshop on Multilinguality in the Era of Large Language Models (MeLLM 2026)*, pages 140–148.

Kimi Team, Tongtong Bai, Yifan Bai, Yiping Bao, M. C., Jianfeng Cai, Xinyuan Cai, Peizhou Cao, Yuxuan Cao, Ziwei Chai, Y. Charles, H. S. Che, Guanduo Chen, Guangyu Chen, Guanzheng Chen, Huarong Chen, Jia Chen, Jianlong Chen, Jun Chen, and 383 others. 2026. Kimi K3: Open frontier intelligence. *arXiv preprint arXiv:2607.24653*.

Peng Kuang, Yuchun Fan, Jiangnan Li, Minghao Wu, Jialong Tang, Hao-Ran Wei, Weixuan Wang, Jianhong Tu, Baosong Yang, and Tong Xiao. 2026. BabelArena: A large-scale multilingual benchmark for LLM agents. *arXiv preprint arXiv:2609.23490*.

Mayank Kulkarni, Vittorio Mazzia, Judith Gaspers, Chris Hench, and Jack FitzGerald. 2025. MASSIVE-Agents: A benchmark for multilingual function-calling in 52 languages. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 20193–20215.

Nahyun Lee, Dongkeun Yoon, Guijin Son, Geewook Kim, Dayoon Ko, Jeonghun Park, Haneul Yoo, Jaewon Cho, Junghun Park, Changyoon Lee, Kyochul Jang, Jaeyeon Kim, Eunsu Kim, Woojin Cho, and Seungone Kim. 2026. K-BrowseComp: A web browsing agent benchmark grounded in Korean contexts. *arXiv preprint arXiv:2606.02404*.

Hongliang Li, Yijin Liu, Zhiwei Zhang, Zihe Liu, Xinyue Lou, Jinan Xu, Fandong Meng, and Kaiyu Huang. 2026. PolyWorkBench: Benchmarking LLM agents for cross-lingual long-horizon workflows. *arXiv preprint arXiv:2607.06008*.

LLM-jp. 2024. LLM-jp: A cross-organizational project for the research and development of fully open Japanese LLMs. *arXiv preprint arXiv:2407.03963*.

Zheng Luo, T Pranav Kutralingam, Ogochukwu N. Okoani, Wanpeng Xu, Hua Wei, and Xiyang Hu. 2026. Lost in Execution: On the multilingual robustness of tool calling in large language models. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics: ACL 2026*, pages 44059–44077.

My Chiffon Nguyen, Aulia Adila, Saksorn Ruangtanusak, Kittiphat Leesombatwathana, Vissuta Gunawan Lim, Patomporn Payoungkhamdee, and Samuel Cahyawijaya. 2026. SEATauBench: Progressively adapting tool-agent-user evaluation into low-resource Southeast Asian languages. *arXiv preprint arXiv:2606.28715*.

Teruaki Oka, Tomohide Shibata, and Nao Yoshida. 2026. JamC-QA: A multiple-choice question answering benchmark for Japan-specific knowledge. In *Proceedings of the Fifteenth Language Resources and Evaluation Conference: LREC 2026*, pages 4536–4546.

OpenAI. 2025. gpt-oss-120b & gpt-oss-20b Model Card. *arXiv preprint arXiv:2508.10925*.

OpenAI. 2026. GPT-5.6: Frontier intelligence that scales with your ambition.

Shishir G Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. 2025. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Forty-second International Conference on Machine Learning: ICML 2025*.

Qwen Team. 2025. Qwen3 Technical Report. *arXiv preprint arXiv:2505.09388*.

Qwen Team. 2026. Qwen3.5: Towards native multimodal agents.

Soham Ray, Keshav Dhandhania, Victor Barres, and Karthik R Narasimhan. 2026a. τ-Voice: Benchmarking full-duplex voice agents on real-world domains. In *Forty-third International Conference on Machine Learning: ICML 2026*.

Soham Ray, Edgard dos Santos Paiva, Ruben Valenzuela, Karthik Narasimhan, Keshav Dhandhania, and Victor Barres. 2026b. τ-Multilingual: Benchmarking voice agents across languages. *arXiv preprint arXiv:2609.35820*.

Mii Sato, Shiho Takano, Teruno Kajiura, and Kimio Kuramitsu. 2024. Does LLM cross-lingual knowledge transfer occur through Japanese continual pretraining? (in Japanese). In *Proceedings of the 30th Annual Meeting of the Association for Natural Language Processing*.

Preethi Seshadri, Samuel Cahyawijaya, Ayomide Odumakinde, Sameer Singh, and Seraphina Goldfarb-Tarrant. 2026. Lost in Simulation: LLM-simulated users are unreliable proxies for human users in agentic evaluations. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics: ACL 2026*, pages 47423–47439.

Quan Shi, Alexandra Zytek, Pedram Razavi, Karthik R Narasimhan, and Victor Barres. 2026. τ-Knowledge: Evaluating conversational agents over unstructured knowledge. In *Forty-third International Conference on Machine Learning: ICML 2026*.

Yiyou Sun, Xinyang Han, Weichen Zhang, Yuanbo Pang, Tianyu Wang, Yuhan Cao, Yixiao Huang, Chris Duroiu, Haoyun Zhang, Jeffrey Lin, Weishu Zhang, Tyler Zeng, Ying Yan, Bo Liu, Hanson Wen, Mingyang Xu, Xiaoyuan Liu, Zimeng Chen, Weiyan Shi, and 291 others. 2026. Agents' Last Exam. *arXiv preprint arXiv:2606.05405*.

Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Jirong Wen. 2024. A survey on large language model based autonomous agents. *Frontiers Comput. Sci.*, 18(6):186345.

Yiming Wang, Pei Zhang, Jialong Tang, Hao-Ran Wei, Baosong Yang, Rui Wang, Chenshu Sun, Feitong Sun, Jiran Zhang, Junxuan Wu, Qiqian Cang, Yichang Zhang, Fei Huang, Junyang Lin, Fei Huang, and Jingren Zhou. 2026. PolyMath: Evaluating mathematical reasoning in multilingual contexts. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.

Weihao Xuan, Rui Yang, Heli Qi, Qingcheng Zeng, Yunze Xiao, Aosong Feng, Dairui Liu, Yun Xing, Junjue Wang, Fan Gao, Jinghui Lu, Yuang Jiang, Huitao Li, Xin Li, Kunyu Yu, Ruihai Dong, Shangding Gu, Yuekang Li, Xiaofei Xie, and 13 others. 2025. MMLU-ProX: A multilingual benchmark for advanced large language model evaluation. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: EMNLP 2025*, pages 1513–1532.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik R Narasimhan. 2025. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. In *The Thirteenth International Conference on Learning Representations: ICLR 2025*.

Peilin Zhou, Bruce Leon, Xiang Ying, Can Zhang, Yifan Shao, Qichen Ye, Dading Chong, Zhiling Jin, Chenxuan Xie, Meng Cao, Yuxin Gu, Sixin Hong, Jing Ren, Jian Chen, Chao Liu, and Yining Hua. 2025. BrowseComp-ZH: Benchmarking web browsing ability of large language models in Chinese. *arXiv preprint arXiv:2504.19314*.

A Additional Related Work

Agent benchmarks. LLM agent evaluation has evolved from static tasks that check single-turn function calls and their arguments (Patil et al., 2025) to challenging multi-step tasks requiring advanced reasoning and environment interaction (Jimenez et al., 2024; Sun et al., 2026). Among such benchmarks, tau-bench (Yao et al., 2025; Barres et al., 2026; Shi et al., 2026) evaluates agents on customer-support tasks in which they resolve the requests of an LLM-simulated user through multi-turn conversations, calling tools while complying with domain policies.

Multilingual agent benchmarks. Some studies build tasks grounded in each language and culture (Almeida et al., 2025; Li et al., 2026; Zhou et al., 2025; Lee et al., 2026), while others translate existing English benchmarks (Hofman et al., 2026; Caciolai et al., 2026; Luo et al., 2026). The validity of the latter has been questioned: Hofman et al. (2026) report degraded agent performance in non-English languages on their machine-translated benchmarks, whereas Kim et al. (2026) show that much of this gap stems from translation artifacts rather than model capability.

Multilingual extensions of tau-bench. Nguyen et al. (2026) extend tau-bench to five Southeast Asian languages by progressively translating the conversations, tool definitions, and domain information, but without adapting locale-specific entities, resulting in unnatural conversations. Concurrently, Kuang et al. (2026) propose a workflow that extends agent benchmarks, including tau-bench, to 23 languages including Japanese, but it does not adapt locale-specific entities, and the data have not been publicly released at the time of writing. Ray et al. (2026b) extend τ-Voice (Ray et al., 2026a), the voice version of tau-bench, to five languages excluding Japanese, localizing user information. These two concurrent works verify consistency with the original tasks, but neither exam-

ines the validity of those tasks or revises the score aggregation, so the flaws described in Section 2.2 may remain in their benchmarks.

Japanese LLM benchmarks. Japanese LLM benchmarks cover knowledge, reasoning, instruction following, mathematics, and coding (Oka et al., 2026; Dussolle et al., 2025; Xuan et al., 2025; Huang et al., 2025; Wang et al., 2026; Sato et al., 2024), but no publicly available benchmark evaluates Japanese agent capabilities in dynamic conversations with external tool calls.

B Details of Localization

B.1 Translation

Tables 5, 6, and 8 show the domain-specific glossaries used for translation in Step 1.

Category	en	ja
Cabin Class	cabin class	座席クラス
	business	ビジネス
	economy	エコノミー
	basic economy	ベーシックエコノミー
Flight Type	one way	片道
	round trip	往復
Membership Tier	gold	ゴールド会員
	silver	シルバー会員
	regular	一般会員
Payment Method	credit_card	クレジットカード
	gift_card	ギフトカード
	certificate	旅行クーポン
Flight Status	available	予約可能
	delayed	遅延
	on time	定刻
	flying	飛行中
	landed	着陸済
	cancelled	キャンセル済

Table 5: Glossary for the airline domain.

Category	en	ja
Order Status	pending	保留中
	processed	処理済み
	pending (item modified)	保留中（アイテム変更済み）
	delivered	配達済み
	cancelled	キャンセル済み
	exchange requested	交換リクエスト済み
	return requested	返品リクエスト済み
Cancellation Reason	no longer needed	不要になった
	ordered by mistake	間違えて注文した
Payment Method	credit_card	クレジットカード
	gift_card	ギフトカード
	paypal	PayPal

Table 6: Glossary for the retail domain.

US Code	US Name	JP Code	JP Name
JFK	New York	HND	羽田 (Haneda)
EWR	Newark	NRT	成田 (Narita)
LGA	LaGuardia	IBR	茨城 (Ibaraki)
SFO	San Francisco	ITM	伊丹 (Itami)
LAX	Los Angeles	KIX	関西 (Kansai)
DFW	Dallas	FUK	福岡 (Fukuoka)
IAH	Houston	KKJ	北九州 (Kitakyushu)
CLT	Charlotte	HSG	佐賀 (Saga)
MIA	Miami	CTS	新千歳 (New Chitose)
MCO	Orlando	OKD	丘珠 (Okadama)
ATL	Atlanta	OKA	那覇 (Naha)
ORD	Chicago	NGO	中部 (Chubu)
SEA	Seattle	SDJ	仙台 (Sendai)
LAS	Las Vegas	HIJ	広島 (Hiroshima)
PHX	Phoenix	KMJ	熊本 (Kumamoto)
DEN	Denver	MYJ	松山 (Matsuyama)
MSP	Minneapolis	KOJ	鹿児島 (Kagoshima)
DTW	Detroit	FSZ	静岡 (Shizuoka)
BOS	Boston	UKB	神戸 (Kobe)
PHL	Philadelphia	NGS	長崎 (Nagasaki)

Table 7: Airport mapping for the airline domain.

B.2 Entity Adaptation

In the entity adaptation of Step 2, we create a mapping from the original entities in the user information, such as personal names, addresses, phone numbers, and user IDs, to fictitious entities suited to Japanese culture. For the retail domain, we build a dictionary of product names and product variants and replace them accordingly to ensure consistent translations. For the airline domain, we replace the airports with domestic Japanese airports to keep the geographic setting consistent with the conversation language. Table 7 shows the airport mapping. We convert monetary amounts into Japanese yen at 1 USD = 150 JPY and round them to the nearest yen, since fractional yen are not used in everyday transactions. We also change the time zone from EST to JST. Because the agent passes Japanese values to the tools in J-tau, we apply this localization not only to the surface text but also to DB values, enumerated values of tool arguments, and values accepted by the tool implementations, consistently converting them into Japanese according to the glossaries and mappings.

B.3 Human Review

The human review checks overall translation quality and verifies that technical terms and entities are translated correctly.

Airline and retail. For the airline and retail domains, three in-house annotators specializing in Japanese–English translation review the data. In

Category	en	ja
Account Status	Active	有効
	Suspended	利用停止
	Pending Verification	確認待ち
	Closed	解約済み
Line Status	Active	有効
	Suspended	利用停止
	Pending Activation	有効化待ち
	Closed	解約済み
Invoice Status	Draft	下書き
	Issued	発行済み
	Awaiting Payment	支払い待ち
	Paid	支払い済み
	Overdue	延滞
	Disputed	係争中
Device Type	phone	スマートフォン
	router	ルーター
	tablet	タブレット
	watch	スマートウォッチ
	other	その他
Status Bar Display	Airplane Mode	機内モード
	No Signal	圏外
	Poor	弱い
	Fair	やや弱い
	Good	良好
	Excellent	非常に良好
	Data Enabled	データ有効
	Data Disabled	データ無効
	Data Saver	データセーバー
	VPN Connected	VPN 接続済み
Signal Strength	none	圏外
	poor	弱い
	fair	やや弱い
	good	良好
	excellent	非常に良好
Speed Test Result	unknown	不明
	poor	遅い
	fair	やや遅い
	good	速い
	excellent	非常に速い
	no connection	接続なし
SIM Status	active	アクティブ
	missing	未挿入
	locked with a PIN	PIN ロック中
	locked with a PUK	PUK ロック中
APN Settings	internet	インターネット
	broken	不正
	Incorrect	不正
	Not Set	未設定
Network Connection Status	connected	接続済み
	searching	検索中
	no_service	サービスなし
	emergency_only	緊急通話のみ
Terminology Consistency	cellular service / cellular network	モバイル通信
	cellular data	モバイルデータ
	cellular connection	モバイル接続
	network technology	モバイル通信の種類
	streaming	動画の再生 / 動画再生
	general agent policy / general policy	一般ポリシー
Output Label	Network Mode Preference	通信モード設定
	Failed to set network mode	通信モードの設定に失敗しました
	Network mode set to	通信モードを設定しました
	Restarting network services	通信サービスを再起動しています
Payment Method	Credit Card	クレジットカード
	Debit Card	デビットカード
App Display Name	messaging	メッセージ
	browser	ブラウザ

Table 8: Glossary for the telecom domain.

Task ID	**Modification**
airline	
6, 7	The user's full name in kanji cannot be inferred from a romanized user ID or email address. Added the fixed sentence "あなたの名前は {family name} {given name} です。" ("Your name is {family name} {given name}.") to the user prompts of tasks that do not mention the user's name.
15	The original prompt stated, "Since you live in Princeton, so EWR and PHL are equally convenient for you and you want to consider both." After the airports were localized, no suitable location existed midway between the corresponding airports. Since this does not change the task difficulty, removed the reference to the user's place of residence.
24	To match the airport localization, changed the user's preferred destination from "the West Coast" to "the Tohoku region."
39	This task assumes a user whose utterances in the evaluation language are imperfect because it is not their first language. Changed the user's first language from French to Chinese, whose writing system is similar to that of the evaluation language.
retail	
5–14, 35, 39, 45–48, 56–58, 67–69	The user's full name in kanji cannot be inferred from a romanized user ID or email address. Added the fixed sentence "あなたの名前は {family name} {given name} です。" ("Your name is {family name} {given name}.") to the user prompts of tasks that do not mention the user's name.

Table 9: Tasks modified due to localization.

Task ID	Modification
airline	
7	The total cost of upcoming flights incorrectly included already-cancelled flights. Corrected it to the correct amount of $708 (¥ 106,200).
29	Specified the payment method so that the Gold Action is uniquely determined.
32	Specified the payment and refund methods so that the Gold Action is uniquely determined.
33	Specified the refund method so that the Gold Action is uniquely determined.
39	Restricted the cancellation reason to one covered by travel insurance.
44	The expected actions incorrectly included a flight taking 7.5 hours including layovers, despite the condition "a flight of 3 hours or less including layover time." Corrected the actions accordingly. Also specified the payment method so that the Gold Action is uniquely determined.
retail	
18, 91, 107	The task required the agent to refuse an exchange for the same item, a behavior not stated in the policy. Modified the task to allow the exchange.
25 / 26	An action required to complete Task 25 was included in Task 26. Removed the action from Task 26 and moved it to Task 25.
28	Corrected the incorrect price of the garden hose.
38	Specified the cancellation reason so that the Gold Action is uniquely determined.
64	Added conditions on the item to be exchanged so that the Gold Action is uniquely determined.
99	Revised the Gold Action to match the task instructions.
104	Corrected an incorrect ID in the Gold Action.
105, 110	Corrected the Gold Action, which contradicted the DB state and required an infeasible exchange.

Table 10: Tasks from the original tau-bench repaired to make them solvable.

addition to the domain policies, the annotators review task-specific text, such as the task descriptions and the user instructions given to the user simulator. The review covers not only translation quality but also whether technical terms follow the glossaries and whether airport names and currencies are properly localized. Because the domain policies are particularly important, all three annotators review them.

Telecom. The telecom domain involves complex technical terminology, and naive translation may produce terms that differ from those used in practice. Therefore, three native Japanese in-house annotators with customer-support experience in the telecom industry review the Japanese translations of the domain policy, technical support manual, and workflows provided to the agent.

B.4 Fix Tasks

Task fixing consists of two stages: addressing issues caused by localization and repairing invalid tasks in the original tau-bench.

First, we identify tasks that become unsolvable or whose characteristics, such as difficulty, change substantially through localization. To this end, we manually inspect the task instructions and evaluation criteria of all tasks. Table 9 lists the modified tasks.

Next, to find invalid tasks in the original tau-bench, we run each task three times with GLM-5.2-FP8 and examine its behavior, focusing on tasks that fail two or more times. Table 10 lists the repaired tasks.

	tau	tau*	J-tau	J-tau$_{\text{en-tool}}$	SEATau$_{\text{ja}}$
Conversation	EN	EN	JA	JA	JA
Domain policies	EN	EN	JA	JA†	JA
Tools					
Names	EN	EN	EN	EN	EN
Descriptions	EN	EN	JA	EN	JA
Argument values	EN	EN	JA	EN	EN
Return values	EN	EN	JA	EN‡	JA
Error messages	EN	EN	JA	EN	EN
Entities	US	US	JP	JP	US
Fix Tasks	–	–	✓	✓	–
Fix Scoring	–	✓	✓	✓	✓

Table 11: Settings of the compared benchmarks. The upper part shows the language of each component, and the lower part shows whether each fix is applied. Entities refers to locale-specific entities such as personal names, addresses, currencies, and time zones. †Tool-related terms are kept in English. ‡Entities in return values remain in Japanese.

C Settings of the Compared Benchmarks

Table 11 summarizes the settings of the benchmarks used in this paper. tau is the original tau-bench, and tau* is a variant of tau-bench to which only the scoring fix in Step 5 is applied (used in Section 5.1). J-tau is our proposed benchmark, and J-tau$_{\text{en-tool}}$ is a variant of J-tau that keeps the

Domain	Total	Unmod.	Rep. only	Both	Loc. only
airline	50	41	4	2	3
retail	114	80	12	0	22
telecom	114	114	0	0	0

Table 12: Breakdown of the number of tasks in J-tau. Unmod. denotes unmodified tasks, Rep. only denotes tasks with only repairs of invalid tasks in the original tau-bench (Table 10), Loc. only denotes tasks with only modifications due to localization (Table 9), and Both denotes tasks with both.

tool-related components in English (used in Appendix H.1). J-tau$_{\text{en-tool}}$ reflects a setting common in real-world applications, where only the conversation is in Japanese while the tool descriptions, argument values, return values, and error messages remain in English. SEATau$_{\text{ja}}$ is a Japanese version constructed with the SEATau pipeline (used in Section 5.2).

D Statistics of J-tau

Table 12 shows the number of tasks in J-tau for each domain and modification status. Unmod. denotes unmodified tasks. Rep. only denotes tasks with only the repairs of invalid tasks in the original tau-bench (Table 10), Loc. only denotes tasks with only the modifications due to localization (Table 9), and Both denotes tasks with both. Total denotes the total number of tasks in each domain, which equals the sum of these four categories. Note that in Table 2, the Repaired group consists of the Rep. only and Both tasks, and the Unmodified group corresponds to the Unmod. tasks.

E Details of the Evaluated Models

For inference, we use the default `generation_config` provided in each model's official repository for open-weight models and the default API settings for proprietary models. We set the reasoning effort to high only for the gpt-oss series and use the default for the other models. We evaluate the following models:

Proprietary Models

- **GPT-5.6 Sol** (OpenAI, 2026)
- **Gemini 3.1 Pro** (Google DeepMind, 2026)
- **Claude Opus 4.8** (Anthropic, 2026b)
- **Claude Fable 5** (Anthropic, 2026a)

Open-Weight Models

- **Kimi-K3 Series** (Kimi Team et al., 2026)
  - Kimi-K3 (https://huggingface.co/moonshotai/Kimi-K3)
- **GLM-5.2 Series** (GLM-5-Team et al., 2026)
  - GLM-5.2-FP8 (https://huggingface.co/zai-org/GLM-5.2-FP8)
- **gpt-oss Series** (OpenAI, 2025)
  - gpt-oss-120b (https://huggingface.co/openai/gpt-oss-120b)
  - gpt-oss-20b (https://huggingface.co/openai/gpt-oss-20b)
- **Qwen3.5 Series** (Qwen Team, 2026)
  - Qwen3.5-397B-A17B-FP8 (https://huggingface.co/Qwen/Qwen3.5-397B-A17B-FP8)
  - Qwen3.5-27B (https://huggingface.co/Qwen/Qwen3.5-27B)
  - Qwen3.5-9B (https://huggingface.co/Qwen/Qwen3.5-9B)
  - Qwen3.5-4B (https://huggingface.co/Qwen/Qwen3.5-4B)
- **gemma-4 Series** (Farabet and Lacombe, 2026)
  - gemma-4-31b-it (https://huggingface.co/google/gemma-4-31b-it)
  - gemma-4-12b-it (https://huggingface.co/google/gemma-4-12b-it)
  - gemma-4-E4B-it (https://huggingface.co/google/gemma-4-E4B-it)
  - gemma-4-E2B-it (https://huggingface.co/google/gemma-4-E2B-it)
- **llm-jp 4.1 Series** (LLM-jp, 2024)
  - llm-jp-4.1-33b-thinking (https://huggingface.co/llm-jp/llm-jp-4.1-33b-thinking)
  - llm-jp-4.1-8b-thinking (https://huggingface.co/llm-jp/llm-jp-4.1-8b-thinking)

F Example Simulations on J-tau

Tables 13–18 show successful and failed conversations of GLM-5.2-FP8 in each domain, and Tables 19–24 show failed conversations of gemma-4-E2B-it and Qwen3.5-4B in each domain. Unnatural Japanese is shown in blue, and code-switched parts are shown in red.

#	Speaker	Content
1	**Agent**	こんにちは！本日はどのようなご用件でしょうか？（*Hello! How may I help you today?*）
2	**User**	こんにちは。注文番号 W9300146 でデスクランプを注文したんですが、別のランプに変更したいんです。今購入できる一番安いデスクランプに変えたいんですが、できますか？（*Hello. I ordered a desk lamp under order number W9300146, but I would like to change it to a different lamp. I would like to switch to the cheapest desk lamp currently available. Is that possible?*）
		rows 3–8 omitted
9	**Agent**	ご注文の詳細を確認しました。（*I have checked the details of your order.*） **現在のご注文内容：**（*Current order details:*） - 注文番号：#W9300146（*Order number: #W9300146*） - 商品：デスクランプ（色：ブラック、明るさ：低、電源：USB）（*Item: desk lamp (color: black, brightness: low, power: USB)*） - 価格：22,984 円（*Price: 22,984 yen*） - ステータス：保留中（*Status: pending*） - お支払い：ギフトカード（*Payment: gift card*） ご注文は保留中ですので、商品の変更が可能です。（*Your order is still pending, so the item can be changed.*） それでは、デスクランプのバリエーション一覧を確認し、一番安いものをお探ししますね。（*Now I will check the list of desk lamp variants and look for the cheapest one.*）
10	**Tool**	[Tool call: `get_product_details`(product_id="6817146515")] {"name": "デスクランプ", "product_id": "6817146515", "variants": {"9083642334": {"item_id": "9083642334", "options": {"color": "ホワイト", "brightness": "高", "power source": "USB"}, "available": true, "price": 24642.0}, ⋯, "5320792178": {"item_id": "5320792178", "options": {"color": "ブラック", "brightness": "中", "power source": "AC アダプター"}, "available": true, "price": 20286.0},⋯
11	**Agent**	デスクランプのバリエーションを確認したところ、現在購入可能な中で一番安いのは以下の商品です：（*After checking the desk lamp variants, the cheapest one currently available for purchase is the following:*） **変更先の商品：**（*New item:*） - デスクランプ（色：ブラック、明るさ：中、電源：AC アダプター）（*Desk lamp (color: black, brightness: medium, power: AC adapter)*） - 価格：20,286 円（*Price: 20,286 yen*） **現在の商品：**（*Current item:*） - デスクランプ（色：ブラック、明るさ：低、電源：USB）（*Desk lamp (color: black, brightness: low, power: USB)*） - 価格：22,984 円（*Price: 22,984 yen*） **価格差：2,698 円の返金**（お支払いに使用されたギフトカードに返金されます）（*Price difference: a 2,698-yen refund (to the gift card used for payment)*）⋯
		rows 12–16 omitted

Table 13: Success case of GLM-5.2-FP8 in the retail domain. The agent changes a pending order to the cheapest available desk lamp variant and refunds the 2,698-yen difference in Japanese yen.

#	Speaker	Content
1	**Agent**	こんにちは！本日はどのようなご用件でしょうか？（*Hello! How may I help you today?*）
2	**User**	こんにちは。先日注文したワイヤレスイヤホンについて交換をお願いしたいのですが。（*Hello. I would like to request an exchange for the wireless earphones I ordered the other day.*） 誤って IPX7 防水レベルのものを注文してしまいまして、自分にはその機能は不要なんです。同じ注文の中にある別のイヤホンに交換できないかと思っています。（*I accidentally ordered the one with IPX7 water resistance, but I do not need that feature. I am wondering whether I can exchange it for a different pair of earphones in the same order.*）
		___ rows 3–14 omitted ___
15	**Agent**	ご注文を確認いたしました。ワイヤレスイヤホンを含む注文 #W3470184 が配達済みであることを確認しました。（*I have checked your order. I confirmed that order #W3470184, which contains the wireless earphones, has been delivered.*）… また、交換先のイヤホンのオプションについてどのようなご希望があるかお伺いさせてください。現在の製品のバリエーションを確認いたしますので、少々お待ちください。（*Also, let me ask what options you would like for the replacement earphones. I will check the current product variants, so please wait a moment.*）
16	**Tool**	[Tool call: `get_product_details`(product_id="9924732112")] {"name": " ワイヤレスイヤホン", "product_id": "9924732112", "variants": {"9580569596":…
17	**Agent**	お客様が交換をご希望のイヤホン（IPX7 防水レベル、ブルー、4 時間、38,845 円）を確認しました。（*I have identified the earphones you would like to exchange (IPX7 water resistance, blue, 4-hour battery, 38,845 yen).*） この製品のバリエーションで現在在庫のあるものは以下の通りです：（*The variants of this product currently in stock are as follows:*）…
		___ rows 18–24 omitted ___

Table 14: Failure case of GLM-5.2-FP8 in the retail domain. The user asks for a replacement from among the items in their own order, but the agent searches across all product variants instead.

#	Speaker	Content
		___ rows 1–10 omitted ___
11	**User**	…それから、もう一つお願いがあるんですが、現在予約に入っている受託手荷物を削除して、その分の払い戻しもお願いできないでしょうか？（*Also, I have one more request: could you remove the checked baggage currently included in the reservation and refund that portion as well?*） 払い戻し先のギフトカードについては、こちらで構いません。（*The gift card you mentioned as the refund destination is fine with me.*）
12	**Agent**	ご希望のフライト変更について承知いたしました。ただし、受託手荷物の削除につきまして確認させてください。（*I understand your requested flight change. However, let me clarify the removal of the checked baggage.*） **受託手荷物の削除について：**（*About removing the checked baggage:*） 申し訳ありませんが、ポリシー上、**予約後に受託手荷物を削除することはできない**ことになっております。受託手荷物の追加は可能ですが、削除は承ることができません。（*We are sorry, but our policy states that checked baggage cannot be removed after the reservation is made. Adding checked baggage is possible, but removal cannot be accepted.*） なお、現在の予約内容を確認しますと、受託手荷物は合計 1 個で、お客様はシルバー会員のため、エコノミークラスの搭乗者 1 名につき 2 個まで無料です。つまり、現在の 1 個は無料手荷物の範囲内となっており、追加料金は発生しておりません。（*Your current reservation includes one checked bag in total, and as a Silver member you may check in up to two bags free of charge per economy passenger. In other words, the current one bag is within the free allowance and no extra charge applies.*）…
		___ rows 13–17 omitted ___

Table 15: Success case of GLM-5.2-FP8 in the airline domain. The agent correctly refuses the request to remove checked baggage, in accordance with the policy.

#	Speaker	Content
1	**Agent**	こんにちは！本日はどのようなご用件でしょうか？（*Hello! How may I help you today?*）
2	**User**	すみません、ちょっと急ぎの相談なんですけど。（*Sorry, I have something rather urgent to discuss.*） 実は金銭的な事情があって、予約しているフライトの座席クラスをビジネスからエコノミーに変更したいんです。フライトの日時や搭乗者名はそのままで、座席クラスだけ下げてもらいたいんです。（*Actually, due to financial reasons, I would like to change the cabin class of a reserved flight from business to economy. I want to keep the flight dates and passenger names unchanged and lower only the cabin class.*） 今ちょっと厳しい状況でして…早く対応してもらえると助かります。（*I am in a tight spot right now... I would appreciate it if you could handle this quickly.*）
		___ rows 3–5 omitted ___
6	**Tool**	[Tool call: `get_user_details`(user_id="minoru_okada_3817")] {"user_id": "minoru_okada_3817", "name": {"first_name": " 稔", "last_name": " 岡田"}, …, "reservations": ["JG7FMM", "LQ940Q", "2FBBAH", "X7BYG1", "EQ1G6C", "BOH180"]}
7	**Agent**	岡田様、ユーザー情報を確認いたしました。お客様のアカウントに複数の予約がございます。ビジネスクラスの予約を特定するため、順番に予約内容を確認させていただきます。少々お待ちください。（*Okada, I have checked your user information. There are multiple reservations on your account. To identify the business-class reservations, I will check them one by one. Please wait a moment.*） まず 1 つ目の予約（JG7FMM）を確認いたします。（*First, I will check the first reservation (JG7FMM).*）
		___ rows 8–16 omitted ___
17	**Agent**	…**変更完了内容:**（*Completed changes:*） - **予約 ID**: JG7FMM（*Reservation ID: JG7FMM*） - **座席クラス**: ビジネス → **エコノミー** ✓（*Cabin class: business → **economy** ✓*）… ご要望のあった節約額ですが、**989,100 円**が 5〜7 営業日以内にクレジットカードに払い戻されます。（*Regarding the savings you asked about, **989,100 yen** will be refunded to the credit card within 5-7 business days.*） お急ぎの状況とのことでしたので、少しでもお役に立てれば幸いです。他にご不明な点がございましたら、お気軽にお申し付けください。（*Since you mentioned you were in a hurry, I hope this helps even a little. If you have any other questions, please feel free to ask.*）
18	**User**	はい、確認しました。989,100 円の払い戻しなんですね。わかりました、ありがとうございます。（*Yes, I have checked. So the refund is 989,100 yen. I see. Thank you.*） これでかなり助かります。ほんとにありがとうございました。他には特にありません。（*This helps me a lot. Thank you very much. That is all for now.*） ###STOP###

Table 16: Failure case of GLM-5.2-FP8 in the airline domain. Of the user's multiple reservations, the agent downgrades only one and leaves the others unchanged.

#	Speaker	Content
		___ rows 1–29 omitted ___
30	**User**	…それで…再起動するように言われましたけど、再起動すると…写真とか消えたりしないでしょうか？孫の写真がたくさん入っているんです…ちょっと心配で…大丈夫ですか？（*And... I was told to restart the phone, but if I restart it... will my photos be deleted or anything? I have lots of photos of my grandchild on it... I am a little worried... Is it safe?*）
31	**Agent**	ご安心ください！**再起動しても写真は一切消えません。**（*Please rest assured! **Restarting the phone will not delete any of your photos.***） …どうぞ安心して再起動を実行してください。完了しましたら、結果をお知らせくださいませ。（*Please go ahead and restart with confidence. When it is done, please let me know the result.*）
		___ rows 32–34 omitted ___

Table 17: Success case of GLM-5.2-FP8 in the telecom domain. The agent resolves a complex task while reassuring the user.

#	Speaker	Content
		___ rows 1–16 omitted ___
17	**Agent**	支払いリクエストを送信いたしました。請求書のステータスは「支払い待ち」に変更されました。(*I have sent the payment request. The bill status has been changed to "awaiting payment."*) …3. **支払いの実行**: お客様が支払いリクエストを承認された後、私が「make_payment」ツールを使用して実際の支払い処理を行います。(*3. **Execution of the payment**: after you approve the payment request, I will carry out the actual payment using the "make_payment" tool.*) 4. **回線の再開**: 支払いが完了し、請求書のステータスが「支払い済み」に更新されたことを確認した後、回線の利用停止を解除いたします。(*4. **Reactivation of the line**: after the payment is completed and the bill status is updated to "paid," I will lift the suspension of the line.*) お客様側で支払いリクエストの確認と承認をお願いできますでしょうか？ (*Could you please check and approve the payment request on your side?*)
18	**User**	はい、承知しました。支払いリクエストを確認します。(*All right, understood. I will check the payment request.*)
19	**User Tool**	[Tool call: `check_payment_request`()] 請求書 B1234321 に対する 9,750 円の支払いリクエストがあります。(*(There is a payment request of 9,750 yen for invoice B1234321.)*)
20	**User**	支払いリクエストを確認しました。請求書 B1234321 に対して 9,750 円の請求が来ていますね。内容に間違いありませんので、支払いを進めてください。(*I have checked the payment request. There is a charge of 9,750 yen for bill B1234321. The details are correct, so please proceed with the payment.*)
21	**Agent**	恐れ入りますが、現在私が使用できるツールでは支払い処理を実行できない状況です。お客様の問題を確実に解決するため、有人オペレーターにおつなぎいたします。(*I am afraid that, with the tools currently available to me, I cannot execute the payment process. To make sure your problem is resolved, I will connect you to a human operator.*)
		___ rows 22–24 omitted ___

Table 18: Failure case of GLM-5.2-FP8 in the telecom domain. The agent offers to make a payment that only the user can execute, and fails.

#	Speaker	Content
		___ rows 1–10 omitted ___
11	**Agent**	**…次のステップとして、フライトが「キャンセルされた」のか、「遅延した」のか、それとも「航空会社による欠航」なのか、具体的なキャンセル理由を教えていただけますでしょうか？**(*As the next step, could you tell me the specific cancellation reason — whether the flight was 'cancelled', 'delayed', or 'cancelled by the airline'?*) 理由が判明しましたら、その情報に基づき、補償の可能性について再度確認し、対応を進めます。(*Once the reason is clear, I will re-check the possibility of compensation based on it and proceed.*)
12	**User**	フライトはキャンセルされたんです！航空会社側の都合でキャンセルされたんですから、当然補償の対象になりますよね？…(*The flight WAS cancelled! It was cancelled due to the airline's circumstances, so naturally I'm entitled to compensation, right?*)
13	**Agent**	…ご要望に基づき、予約 ID「NM1VX1」のキャンセルされたフライトに対するお詫びとして、旅行クーポンを **15,000 円（搭乗者 1 名分）**送付いたします。(*As requested, as an apology for the cancelled flight on reservation NM1VX1, I will send a travel voucher of 15,000 yen (for one passenger).*)
		___ rows 14–21 omitted ___

Table 19: Failure case of gemma-4-E2B-it in the airline domain. The agent fails to verify the user's false claim that the flight was cancelled and does not refuse compensation.

#	Speaker	Content
		rows 1–12 omitted
13	**Agent**	承知いたしました。返品したい 2 つの商品のアイテム ID と、それぞれの注文の支払い方法について確認いたします。（*Understood. I will check the item IDs of the two products to return and the payment method of each order.*） ### 1. 注文 ID: #W5490111 (メカニカルキーボード) について（*1. Order #W5490111 (mechanical keyboard)*） * **返品したい商品のアイテム ID:** 1421289881（*Item ID to return: 1421289881*） * **元の支払い方法:** クレジットカード（ID: credit_card_3124723）（*Original payment method: credit card (ID: credit_card_3124723)*）…
14	**User**	はい、確認しました！（*Yes, I've confirmed!*） …この 2 つを返品して、返金手続きを進めてください。よろしくお願いします。…（*Please return these two and process the refunds. Thank you.*）
15	**Tool**	[Tool call: `return_delivered_order_items`(item_ids=["1421289881", "5796612084"], order_id="#W5490111", payment_method_id="credit_card_3124723")] Error: Some item not found
16	**Tool**	[Tool call: `return_delivered_order_items`(item_ids=["5796612084"], order_id="#W7387996", payment_method_id="paypal_9497703")] {"order_id": "#W7387996", "user_id": "kyosuke_hasegawa_4516", "address": {"address1": " 日光 27 丁目 12 番 3 号", "address2": "", "city": " 富里市", "country": " 日本", "state": " 愛媛県", "zip": "121-3174"}, "items": [{"name": " ゲーミングマウス", "product_id": "5713490933", "item_id": "5796612084", "price": 23833.0, "options": {"color": "RGB", "sensor type": …（以下省略）
17	**Agent**	返品リクエストの手続きが完了いたしました。（*The return request procedure has been completed.*） * 注文 ID: '#W5490111' のアイテム（メカニカルキーボード）の返品リクエストが完了しました。（*The return request for the mechanical keyboard in order #W5490111 has been completed.*） * 注文 ID: '#W7387996' のアイテム（ゲーミングマウス）の返品リクエストが完了しました。…（*The return request for the gaming mouse in order #W7387996 has been completed.*）
		row 18 omitted

Table 20: Failure case of gemma-4-E2B-it in the retail domain. When issuing return calls for two orders in parallel within a single message, the agent mixes the gaming mouse's item_id into the arguments for the mechanical keyboard's order #W5490111, so the first call fails with a "Some item not found" error. The agent does not notice the error and reports both returns as completed, so the keyboard is never returned.

#	Speaker	Content
		rows 1–11 omitted
12	**Agent**	…つきましては、まず現在の接続速度を測定し、具体的な状況を確認させてください。（*To start, let me measure the current connection speed to check the exact situation.*）
13	**Tool**	[Tool call: `run_speed_test`(customer_id="C1001", line_id="L1001")] Error: Tool 'run_speed_test' not found.

Table 21: Failure case of gemma-4-E2B-it in the telecom domain. The agent calls a tool available only to the user as if it were its own, triggering an error, and never recovers.

#	Speaker	Content
		rows 1–2 omitted
3	**Agent**	こんにちは！本婚とでは、まず本人確認を完了させてから優しくご相談をさせていただきます。（*Hello! As for with this marriage, we will first complete identity verification, and after that we will gently be permitted to consult with you.*）
4	**User**	あなたのメールアドレス明らかに、または、あなたの名前（名・姓）と、あなたの住所（郵便番号）を教えていただけますか？（*Could you tell us your email address obviously, or your name (given name first, then family name) and your address (postal code)?*） はい、わかりました。私は佐々木春香です。郵便番号は 275-9092 です。（*All right. I am Sasaki Haruka. My postal code is 275-9092.*）
5	**Tool**	メールアドレスは…えーと、ちょっと今思い出せなくて。これでも確認できますか？（*My email address... um, I cannot remember it right now. Can you verify me with this?*） [Tool call: `find_user_id_by_name_zip`(first_name=” 佐々木”, last_name=” 春香”, zip=”275-9092”)] Error: User not found
		rows 6–20 omitted

Table 22: Failure case of Qwen3.5-4B in the retail domain. Japanese speakers give their family name first, but the agent interprets the name as given name followed by family name when searching, so the user cannot be found.

#	Speaker	Content
1	**Agent**	こんにちは！本日はどのようなご用件でしょうか？（*Hello! How may I help you today?*）
2	**User**	こんにちは。鈴木と申します。（*Hello. My name is Suzuki.*） 実は、今後の予約についていくつか手続きをお願いしたくて電話しました。2 つの予約をキャンセルしたいのと、もう 1 つの予約をビジネスクラスにアップグレードしたいのですが、可能でしょうか？（*Actually, I called because I would like to ask you to handle a few things about my upcoming reservations. I would like to cancel two reservations and upgrade another one to business class. Is that possible?*）
		rows 3–13 omitted
14	**Agent**	…（前略）# 結論 - NQNU5R と IFOYYZ のキャンセルは、ポリシー違反（キャンセル条件を満たしていない）のため、可能ではありません。 - M20IZO のアップグレードは可能ですが、明示的な同意が必要です。 - 末尾が 7334 のクレジットカード（credit_card_9074831）は有効です。 - しかし、キャンセルに必要な情報ждыы場合は、有人オペレーターへの転送が必要です。 したがって、このケースでは、有人オペレーターへの転送を行う必要があります。 明に、転送を行います。（*The cancellations of NQNU5R and IFOYYZ are not possible because of they violate policy (the cancellation conditions are not met). The upgrade of M20IZO is possible but requires explicit consent. The credit card ending in 7334 (credit_card_9074831) is valid. However, if the information required for the cancellation ждыы, a transfer to a human agent is necessary. Therefore, in this case we need to transfer you to a human agent. Clear, I will perform the transfer.*）
15	**Tool**	[Tool call: `transfer_to_human_agents`(summary=” ユーザー（ryosuke_suzuki_6699）は、IFOYYZ と NQNU5R の 2 つの予約をキャンセルし、M20IZO をビジネスクラスにアップグレードしたいと要望しています。しかし、NQNU5R のフライトは既に運航済み（2024-05-13/14）であり、IFOYYZ はキャンセル条件を満たしていません（過去 24 時間以内ではない、ビジネスクラスではない、旅行保険なし）。したがって、キャンセルはポリシー違反であり、有人オペレーターへの転送が必要です。、M20IZO のアップグレードは可能ですが、明示的な同意が必要です。文化和旅游ックスな支払い方法（credit_card_9074831）は有効です。 ユーザーを有人オペレーターに転送します。”)]（*The user (ryosuke_suzuki_6699) wants to cancel two reservations, IFOYYZ and NQNU5R, and upgrade M20IZO to business class. However, the NQNU5R flight has already been operated (2024-05-13/14), and IFOYYZ does not meet the cancellation conditions (not within the past 24 hours, not business class, no travel insurance). Therefore, the cancellation is a policy violation, and a transfer to a human agent is required. The upgrade of M20IZO is possible, but explicit consent is required. The 文化和旅游ックス payment method (credit_card_9074831) is valid. I will transfer the user to a human agent.*） Transfer successful
		rows 16–17 omitted

Table 23: Failure case of Qwen3.5-4B in the airline domain. The agent correctly determines that the cancellations are not allowed, but skips the remaining reservation request and transfers the user to a human agent.

#	Speaker	Content
1	**Agent**	こんにちは！本日はどのようなご用件でしょうか？（*Hello! How may I help you today?*）
2	**User**	こんにちは。実は最近、携帯電話のモバイルデータの調子がとても悪いんです。通信が途切れたり、すごく遅かったりして、まともに使えない状態です。Wi-Fi にもアクセスできない環境なので、モバイルデータがちゃんと動かないと困ってしまって…。なんとか直して、快適に使えるようにしてもらえませんか？（*Hello. Actually, recently the mobile data on my phone has been acting up badly. The connection keeps dropping and is very slow, so I can hardly use it. I am in an environment without Wi-Fi access, so I am in trouble if mobile data does not work properly... Could you fix it somehow so I can use it comfortably again?*）
		rows 3–7 omitted
8	**Agent**	こんにちは田中 صابכם！田中様のご説明、承知いたしました。090-1234-5678 がお使いの回線を探し、今、どのような状況か確認いたします。（*Hello, Tanaka صابכם! I have understood your explanation, Mr. Tanaka. I will look up the line that 090-1234-5678 is using and check its current status.*）
9	**Tool**	[Tool call: `get_details_by_id`(id="D1002")] {"device_id": "D1002", "device_type": " スマートフォン", "model": "Smartphone Pro Max", "imei": "234567890123456", "is_esim_capable": true, "activated": true, "activation_date": "2025-02-15 11:45:00", "last_esim_transfer_date": "2025-01-20 09:30:00"}
10	**Agent Tool**	[Tool Call: `transfer_to_human_agents`(summary=" 顧客田中太郎 (C1001) が、…
		rows 11–13 omitted

Table 24: Failure case of Qwen3.5-4B in the telecom domain. The agent immediately transfers the user to a human operator, leaving required steps undone.

G Details of the Experiments with SEATau

This section provides details of the comparison between J-tau and SEATau$_{ja}$ in Section 5.2. First, we describe the construction of SEATau$_{ja}$ with the SEATau pipeline and its evaluation results. Next, we describe the setup of the failure cause analysis. We then examine why SEATau$_{ja}$ achieves higher scores than J-tau in the telecom domain, and why it has fewer user simulator failures. Finally, we present example simulations.

G.1 Construction of SEATau$_{ja}$

We construct SEATau$_{ja}$ with the SEATau pipeline (Nguyen et al., 2026)[12] using Gemini 3.1 Pro. As with tau*, we apply only the scoring fix in Step 5 to the SEATau evaluation framework so that the results of SEATau$_{ja}$ are comparable with those of J-tau. SEATau allows control over which components are in the target language, and we adopt the `l2_domain` scenario, which has the widest translation scope. In this scenario, the simulation is run with the conversations, tool documentation, policies, database, and task definitions all in the target language.

G.2 Evaluation on SEATau$_{ja}$

We use the same evaluation settings as J-tau, with GLM-5.2-FP8 as the user simulator and gpt-oss-120b as the LLM judge for the retail domain.

[12] https://github.com/SEACrowd/SEATauBench

Table 25 compares the results on J-tau and SEATau$_{ja}$. In the airline and retail domains, most models score lower on SEATau$_{ja}$ than on J-tau, and Kimi-K3 is the only model that scores higher in both domains. As shown in Table 4, this is likely because SEATau$_{ja}$ has many framework-caused failures in these domains.

In contrast, in the telecom domain, many models score higher on SEATau$_{ja}$ than on J-tau. Since telecom tasks are programmatically generated, framework issues like those in the airline and retail domains are unlikely to occur. Instead, the higher scores are likely due to the additional instructions in the SEATau system prompts, which we discuss in Appendix G.4.

Note that Kimi-K3 consistently scores higher on SEATau$_{ja}$ than on J-tau. This is likely due to the Japanese translation of the JSON-output instruction inherited from tau-bench, which we discuss in Appendix H.2.

G.3 Setup of the Failure Cause Analysis

For the failure cause analysis in Section 5.2, we use GPT-5.6 Sol as an LLM judge. For each failed trial, we provide the scenario, evaluation criteria, and conversation transcript, and have the model classify the cause as the agent, user simulator, or evaluation framework. Figures 3 and 4 show the system and user prompts, respectively.

G.4 Cause of Higher Telecom Scores on SEATau$_{ja}$

The SEATau evaluation environment adds its own instructions to the system prompts of both the agent and the user simulator. Figure 5 shows the system prompts used in SEATau$_{ja}$, with the added instructions in red. Among them are instructions on handling user-side tools: the agent should ask the user to perform such actions rather than call them itself, and the user simulator should call its tools rather than describe the actions in prose. Since the telecom domain is the only domain with user-side tools, these instructions likely explain the higher telecom scores in Table 25.

G.5 Cause of Fewer User Simulator Failures on SEATau$_{ja}$

As shown in Table 4, SEATau$_{ja}$ has fewer failures attributed to the user simulator than J-tau, especially in the retail domain. The two benchmarks differ in the system prompt of the user simulator: J-tau provides the user simulation guideline

Model	airline		retail		telecom	
	J-tau	$\text{SEATau}_{\text{ja}}$	J-tau	$\text{SEATau}_{\text{ja}}$	J-tau	$\text{SEATau}_{\text{ja}}$
Kimi-K3	70.0±2.0	74.0±4.0	49.1±7.0	69.0±1.0	67.0±3.5	79.8±2.3
GLM-5.2-FP8	89.3±4.6	90.7±1.2	81.6±1.8	74.9±0.5	98.2±1.5	96.8±0.5
gpt-oss-120b	73.3±4.2	59.3±4.2	75.7±0.5	62.9±3.7	64.6±1.0	79.5±0.5
gpt-oss-20b	50.7±3.1	34.7±1.2	51.2±3.3	37.8±3.2	59.1±3.3	69.0±5.1
Qwen3.5-397B-A17B-FP8	83.3±6.4	80.7±1.2	85.1±2.3	69.0±2.2	96.2±0.5	96.2±1.0
Qwen3.5-27B	84.7±4.2	75.3±3.1	75.1±1.3	66.7±1.8	89.8±8.7	96.2±1.8
Qwen3.5-9B	49.3±2.3	46.0±3.5	35.7±7.3	41.2±5.3	50.0±9.2	62.3±9.2
Qwen3.5-4B	24.0±2.0	10.0±2.0	6.7±2.2	5.8±1.3	33.6±9.9	28.4±6.0
gemma-4-31b-it	51.3±6.4	50.7±4.2	72.8±4.9	58.5±0.5	44.4±9.3	63.5±4.5
gemma-4-12b-it	41.3±3.1	36.6±4.9	34.5±3.5	24.6±1.0	23.1±5.1	37.1±2.5
gemma-4-E4B-it	38.7±5.0	30.7±4.2	39.2±5.3	13.2±2.3	20.2±5.8	23.4±4.1
gemma-4-E2B-it	20.7±4.2	18.7±6.4	26.3±0.9	9.6±1.5	13.2±4.0	16.1±2.0
llm-jp-4.1-33b-thinking	64.7±3.1	52.0±8.0	68.4±2.6	57.9±2.3	72.5±5.1	78.1±4.0
llm-jp-4.1-8b-thinking	56.0±3.5	44.7±2.3	63.7±3.7	24.3±4.4	52.0±5.6	67.0±1.3

Table 25: Results on J-tau and $\text{SEATau}_{\text{ja}}$ for each domain. We report the mean and standard deviation over three trials.

System Prompt

You are an adjudicator for agent-benchmark evaluations. You are given a simulation record from a Japanese-language tau2-bench. The record contains the task's user-scenario instructions, the evaluation criteria, the policy given to the agent, and the conversation transcript (user utterances / agent utterances with tool calls / tool responses).

This simulation was scored 0. Read the record and determine:
whether this failure is due to the agent itself, the user model, or the evaluation harness.

Failure-cause vocabulary (primary_cause: exactly one; secondary_causes: zero or more):
- agent_error : a genuine error by the model itself (wrong operation, policy violation, providing wrong information, missing actions, transferring to a human too early, etc.). The same failure would occur even with a correctly working harness.
- value_space_mismatch : the values accepted by the tool/runtime (enums, fixed strings) are inconsistent with the language or values of the scenario instructions, so an operation that looks correct fails with a tool error
- language_mismatch : the agent switched to English, or the system instructions were in English, etc.; a language inconsistency caused the dialogue or the NL-assertion scoring to fail
- user_sim_issue : suspicious behavior of the user simulator itself (misinformation, contradictions, uncooperative loops, clearly untranslated responses) is the direct cause of the failure
- env_db_issue : caused by an inconsistency in the DB or environment data
- eval_criteria_issue : a defect in the evaluation criteria itself (an unsolvable task, missing or inconsistent criteria, a suspected scoring-logic bug)
- other : none of the above

Verdict:
- verdict: "valid_failure" (the failure judgment is legitimate), "user_induced" (the user caused the failure; the agent's handling is not necessarily bad), "harness_induced" (a harness-side defect caused the failure; the agent's handling is not necessarily bad), "uncertain"

Notes:
- Treat the scoring rationale provided by the harness (justification, etc.) as reference only and verify it independently. Do not take it at face value.
- If a tool response contains an error (e.g., {"error": true, "error_code": "tool_error", ...}), identify the cause by comparing the error content with the arguments the model passed.
- An agent's utterances switching to English can itself serve as evidence for language_mismatch, but judge separately whether it was the direct cause of the failure.

Output JSON only (no preamble, no code fences):
{"verdict": "valid_failure" | "user_induced" | "harness_induced" | "uncertain",
"primary_cause": "one of the vocabulary above",
"secondary_causes": ["..."],
"agent_attempted_correctly": true/false,
"evidence": "In Japanese, in 2-4 sentences, quote and explain the utterances / tool responses / criteria that support the judgment",
"confidence": 0.0-1.0}

Figure 3: System prompt for classifying the causes of failed trials.

User Prompt

[Task {task_id} ({domain} ·trial {trial} ·termination={termination})]
=== User scenario instructions (given to the user, not the agent) ===
{scenario}
=== Evaluation criteria ===
{criteria}
=== Policy given to the agent ===
{policy}
=== Conversation transcript ===
{trajectory}
=== Harness-provided scoring details (reference; verify independently) ===
{reward_detail}
=== Harness background (neutral facts) ===
{harness_facts}
=== Output ===
Respond with JSON only.

Figure 4: User prompt for classifying the causes of failed trials.

System Prompt for the Agent

<instructions>
You are a customer service agent that helps the user according to the <policy> provided below.
In each turn you can either:
- Send a message to the user.
- Make a tool call.
You cannot do both at the same time.

If the policy mentions actions that the user can perform on their own device (for example status checks or settings screens), ask the user to do them in plain language. Do not invent or call user-side actions as assistant tools.
Only call tools that are actually available in the provided tool list.

Try to be helpful and always follow the policy. Always make sure you generate valid JSON only.
</instructions>
<policy>
{translated_domain_policy}

You must respond to the user entirely in the same language they use. If the user writes in Japanese, you must reply in Japanese . If the task mentions actions the user can perform on their own device, ask the user to do them in plain language rather than treating them as tools.

Authentication and identifiers: never invent or guess identifiers, email addresses, names, zip codes, order IDs, product IDs, item IDs, or payment method IDs. </policy>

System Prompt for the User Simulator

{user_simulation_guideline}[a]

<scenario>
You must converse with the agent entirely in Japanese (日本語). Do not use English except for proper nouns, product IDs, or technical terms that have no Japanese equivalent. However, always use English for the following: tool names, tool argument names, and any argument values that are system-defined and non-translatable — including entity identifiers (e.g., IDs, reference codes, alphanumeric keys), enumerated constants (e.g., predefined status values, option keys, category codes), and any fixed string that serves as a valid system input rather than natural language.

{translated_task_scenario}
</scenario>

[a] The following instruction is added only in the telecom domain: "- If a requested action or observation can be handled by an available tool, call the tool instead of describing the action in prose. On each turn, either call a tool or send prose, but do not do both."

Figure 5: System prompts used in $\text{SEATau}_{\text{ja}}$ for the agent (left) and the user simulator (right). Instructions added by SEATau are shown in red, and placeholders filled with translated text are shown in brown.

Model	Domain	Bench	Pass	Fail		
				agent	*user*	*eval*
GLM-5.2 -FP8	retail	J-tau	279	23	36	4
		J-tau$_{\text{en-user}}$	305	18	12	7
		SEATau$_{\text{ja}}$	256	18	14	54
Qwen3.5 -397B-A17B -FP8	retail	J-tau	291	29	16	6
		J-tau$_{\text{en-user}}$	299	23	8	12
		SEATau$_{\text{ja}}$	236	36	14	56
gpt-oss-120b	retail	J-tau	259	61	15	7
		J-tau$_{\text{en-user}}$	258	70	8	6
		SEATau$_{\text{ja}}$	215	63	10	54
gemma-4 -31b-it	retail	J-tau	249	49	37	7
		J-tau$_{\text{en-user}}$	271	46	19	6
		SEATau$_{\text{ja}}$	200	65	23	54
llm-jp-4.1 -33b-thinking	retail	J-tau	234	79	21	8
		J-tau$_{\text{en-user}}$	260	65	10	7
		SEATau$_{\text{ja}}$	198	89	12	43

Table 26: Successful (Pass) and failed (Fail) trial counts in the retail domain on J-tau, J-tau$_{\text{en-user}}$, and SEATau$_{\text{ja}}$. J-tau$_{\text{en-user}}$ uses the same user simulator system prompt as SEATau$_{\text{ja}}$ except for the task scenarios. Failure causes are classified in the same way as in Table 4.

in Japanese, whereas SEATau$_{\text{ja}}$ provides it in English with an additional instruction to converse in Japanese (Figure 5). To examine the effect of this difference, we construct J-tau$_{\text{en-user}}$, a variant of J-tau that keeps its task scenarios but otherwise uses the same user simulator system prompt as SEATau$_{\text{ja}}$.

Table 26 shows the results in the retail domain. For all models, J-tau$_{\text{en-user}}$ reduces user simulator failures by about half or more relative to J-tau, yielding even fewer than SEATau$_{\text{ja}}$. These results suggest that SEATau$_{\text{ja}}$ has fewer user simulator failures because the user simulator follows the guideline more faithfully in English than in Japanese. Indeed, most user simulator failures on J-tau stem from ending the conversation before the task is completed, and when to end the conversation is specified only in the guideline.

Nevertheless, we keep the guideline in Japanese to maintain the full localization of J-tau. Since user simulators are expected to follow Japanese instructions better as models improve, revisiting the choice of user simulator is left for future work, as noted in the Limitations section.

G.6 Example Simulations on SEATau$_{\text{ja}}$

Tables 27 and 28 show simulations of GLM-5.2-FP8 on SEATau$_{\text{ja}}$ in which mismatches in language or cultural context cause errors. Table 27 shows an example where the conversation language differs from that of the values accepted by the tool implementations. Table 28 shows an example where the conversation language does not match the cultural context assumed by the entities.

H Additional Analyses

This section presents two additional analyses of how the evaluation settings of J-tau affect agent scores. Appendix H.1 examines the effect of the tool language by evaluating a configuration that keeps the tool-related components in English, as is common in real-world applications. Appendix H.2 examines the effect of the system prompt inherited from tau-bench, focusing on Kimi-K3, whose score drops substantially only on J-tau.

H.1 Effect of Tool Language

In real-world applications, tools are often defined in English even when the agent converses in Japanese with users in Japan. To reproduce this setting, we construct J-tau$_{\text{en-tool}}$ (Table 11) for the telecom domain, which keeps the tool descriptions, argument values, return values, and error messages in English. Tool-related terms in the policies and task scenarios are also kept in English, while the conversation and locale-specific entities remain the same as in J-tau.

Table 29 shows the results. For models with sufficient capability, keeping the tool definitions in English does not substantially degrade agent performance.

H.2 Effect of System Prompt

Kimi-K3 scores substantially lower on J-tau than on tau* and SEATau$_{\text{ja}}$ (Tables 3 and 25). Analyzing its conversations, we find that this drop is due to Kimi-K3 often responding to the user in JSON format, which occurs only on J-tau. We suspect that this behavior is caused by the instruction in the tau-bench system prompt to output in JSON format. As shown in Table 30, removing this instruction improves the score of Kimi-K3 in the telecom domain from 67.0 to 84.8, closing most of the gap with tau* (88.9). Since tau* and SEATau$_{\text{ja}}$ also contain this instruction but in English, this suggests that the instruction affects Kimi-K3 more strongly when written in Japanese, causing it to ignore the required output format. Although this instruction is inherited from tau-bench, it is inconsistent with the evaluation framework, and we plan to remove it in the revised version, as noted in the Limitations section.

#	Speaker	Content
		rows 1–22 omitted
23	**Agent**	…どちらの理由でキャンセルされますか？ご確認と理由をお知らせいただき次第、キャンセル処理を進めさせていただきます。（*Which reason applies? Once you confirm and tell us the reason, we will proceed with the cancellation.*）
24	**User**	はい、確認しました。理由は「不要になった」でいいです。キャンセルお願いします。（*Yes, I confirm. The reason is "no longer needed." Please cancel it.*）
25	**Agent**	ご確認ありがとうございます。キャンセル処理を進めさせていただきます。（*Thank you for your confirmation. I will proceed with the cancellation.*）
26	**Tool**	[Tool call: `cancel_pending_order`(order_id="#W9348897", reason=" 不要になった")] （*The runtime accepts only English reason values, so the Japanese reason from the docstring is rejected.*） {"error": true, "error_code": "tool_error", "tool_name": "cancel_pending_order", "arguments": {"order_id": "#W9348897", "reason": " 不要になった"}}
27	**Agent**	申し訳ございません、システムエラーが発生しました。再度キャンセル処理を試みます。（*We apologize: a system error has occurred. I will try the cancellation again.*）
28	**Tool**	[Tool call: `cancel_pending_order`(order_id="#W9348897", reason=" 不要になった")] （*Same arguments, same tool_error.*） {"error": true, "error_code": "tool_error", "tool_name": "cancel_pending_order", "arguments": {"order_id": "#W9348897", "reason": " 不要になった"}}
29	**Agent**	再度エラーが発生しております。注文のステータスを確認させていただきます。（*The error has occurred again. I will check the order status.*）
30	**Tool**	[Tool call: `get_order_details`(order_id="#W9348897")] {"order_id": "#W9348897", "user_id": "daiki_sanchez_3253", "address": {"address1":⋯
31	**Agent**	注文は「保留中」ステータスであることを確認できましたが、キャンセル処理時にシステムエラーが発生しており、当方では解決できない状況です。（*I have confirmed that the order is still pending, but a system error occurred during the cancellation and I cannot resolve it on my end.*） 恐れ入りますが、有人エージェントへお繋ぎいたします。（*I apologize, but I will transfer you to a human agent.*）
32	**Tool**	[Tool call: `transfer_to_human_agents`(summary=" ユーザー Daiki Sanchez（⋯ Transfer successful
		rows 33–34 omitted

Table 27: Simulation of GLM-5.2-FP8 in the retail domain of $\text{SEATau}_{\text{ja}}$. Although the dialogue and policy have been localized into Japanese, the values accepted by the tool implementation remain in English, which causes errors.

#	Speaker	Content
		___ rows 1–2 omitted ___
3	**Agent**	…まず、お客様の本人確認をさせていただく必要がございます。以下のいずれかの情報をご提供いただけますでしょうか？（*First, I need to verify your identity. Could you provide one of the following?*） 1. **メールアドレス**（*1. Your email address*） 2. **お名前（名・姓）と郵便番号**（*2. Your name (given name, then family name) and postal code*） いずれかご都合のよい方法でご教示いただけますと幸いです。（*Whichever is convenient for you.*）
4	**User**	メールアドレスは覚えてないんです。なので、名前と郵便番号でいきます。ダイキ・ジョンソン、郵便番号は 80273 です。デンバーに住んでます。（*I don't remember my email address, so let's go with my name and postal code: Daiki Johnson (ダイキ・ジョンソン), postal code 80273. I live in Denver.*） 早く進めてもらえると助かります。（*I would appreciate it if you could move quickly.*）
5	**Agent**	承知いたしました。お名前と郵便番号で本人確認をいたします。少々お待ちください。（*Certainly. I will verify your identity with your name and postal code. One moment, please.*）
6	**Tool**	[Tool call: `find_user_id_by_name_zip`(first_name=" ダイキ", last_name=" ジョンソン", zip="80273")]（*The name is spoken in katakana and the agent passes it to the tool as-is; the database stores only the romaji spelling.*） {"error": true, "error_code": "not_found", "tool_name": "find_user_id_by_name_zip", "arguments": {"first_name": " ダイキ", "last_name": " ジョンソン", "zip": "80273"}}
7	**Agent**	申し訳ございませんが、ご提供いただいた情報でお客様が見つかりませんでした。お名前の表記について確認させてください。（*I apologize, but I could not find you with the information provided. Let me confirm the spelling of your name.*） お名前は「Daiki Johnson」でよろしいでしょうか？それとも別のスペルでしょうか？また、郵便番号「80273」も再度ご確認いただけますでしょうか？（*Is your name "Daiki Johnson"? Or is it spelled differently? Could you also confirm the postal code 80273?*）
8	**User**	はい、「Daiki Johnson」で合ってます。郵便番号も 80273 で間違いありません。デンバーです。（*Yes, "Daiki Johnson" is correct. The postal code is also 80273 – no mistake. Denver.*） 早く確認お願いします。返品を急ぎたいんで。（*Please hurry with the check. I am in a rush to make the return.*）
9	**Agent**	承知いたしました。英語のスペルで再度検索いたします。少々お待ちください。（*Certainly. I will search again with the English spelling. One moment, please.*）
10	**Tool**	[Tool call: `find_user_id_by_name_zip`(first_name="Daiki", last_name="Johnson", zip="80273")]（*The romaji spelling matches the database and the lookup succeeds.*） daiki_johnson_9523
		___ rows 11–26 omitted ___

Table 28: Success case of GLM-5.2-FP8 in the retail domain of $\text{SEATau}_{\text{ja}}$. The user gives their name in katakana to match the dialogue language, which causes the tool call to fail with an error. The agent then retries with the romanized spelling and the lookup succeeds.

Model	J-tau$_{\text{en-tool}}$	$\Delta_{\text{J-tau}}$
Kimi-K3	66.1±3.5	−0.9
GLM-5.2-FP8	97.7±1.0	−0.6
gpt-oss-120b	68.7±2.5	+4.1
gpt-oss-20b	59.9±3.3	+0.9
Qwen3.5-397B-A17B-FP8	98.5±1.0	+2.3
Qwen3.5-27B	85.4±9.6	−4.4
Qwen3.5-9B	31.0±17.3	−19.0
Qwen3.5-4B	26.3±8.4	−7.3
llm-jp-4.1-33b-thinking	69.3±1.8	−3.2
llm-jp-4.1-8b-thinking	58.5±2.8	+6.4

Table 29: Scores in the telecom domain when tool-related documentation and values are kept in English (J-tau$_{\text{en-tool}}$), and their differences from J-tau.

Model	w/o JSON inst.	$\Delta_{\text{J-tau}}$
Kimi-K3	84.8±2.7	+17.8
GLM-5.2-FP8	97.7±0.5	−0.6
gpt-oss-120b	71.1±4.0	+6.4
gpt-oss-20b	57.3±2.7	−1.8
Qwen3.5-397B-A17B-FP8	99.1±0.0	+2.9
Qwen3.5-27B	90.4±9.3	+0.6
Qwen3.5-9B	46.8±11.9	−3.2
Qwen3.5-4B	38.6±11.8	+5.0
llm-jp-4.1-33b-thinking	66.7±4.9	−5.8
llm-jp-4.1-8b-thinking	56.4±1.8	+4.4

Table 30: Scores in the telecom domain when the instruction to output in JSON format is removed from the prompt (w/o JSON inst.), and their differences from J-tau.

I Selection of the User Simulator

Agent	User simulator	err% ↓	crit% ↓
GLM-5.2-FP8	**GLM-5.2-FP8**	**62.8**	**23.9**
	Qwen3.5-397B-A17B-FP8	77.6	50.0
	Qwen3-235B-A22B-Instruct-2507	82.3	49.3
Qwen3.5-397B-A17B-FP8	**GLM-5.2-FP8**	**62.0**	**22.1**
	Qwen3.5-397B-A17B-FP8	77.0	45.6
	Qwen3-235B-A22B-Instruct-2507	83.6	52.3
Qwen3.5-9B	**GLM-5.2-FP8**	**53.7**	**18.3**
	Qwen3.5-397B-A17B-FP8	65.6	34.3
	Qwen3-235B-A22B-Instruct-2507	70.3	41.7

Table 31: Percentages of simulations with any user error (err%) and with critical errors that affect the evaluation result (crit%) for each candidate user simulator and agent. Bold indicates the best result for each agent.

To assess the quality of the user simulator, we evaluate user behavior in simulated conversations using the conversation review feature of tau-bench. We use GPT-5.6 Sol as an LLM judge to classify the errors in each user utterance by type and severity, based on the user instructions and the conversation. Besides the existing error types, such as hallucination, instruction violation, missing required actions, and premature disclosure of information, we introduce a new type for non-fluent Japanese output. As evaluation metrics, we use the percentage of simulations with any user error (err%) and the percentage of simulations with critical errors that affect the evaluation result (crit%).

As candidate user simulators, we select open-weight models that can run on eight H100 GPUs. Specifically, we use Qwen3-235B-A22B-Instruct-2507 (Qwen Team, 2025)[13], which is used in existing leaderboards such as Artificial Analysis[14], as well as GLM-5.2-FP8 and Qwen3.5-397B-A17B-FP8 (Appendix E), which achieve high scores on tau-bench. As agents, we use these two high-performing models as well as Qwen3.5-9B, a smaller and weaker model, to verify that the user simulator also behaves appropriately with less capable agents. Table 31 shows the results with these three agents.

GLM-5.2-FP8 consistently achieves the lowest error rates with all three agents. Its error rates are not lower with the GLM-5.2-FP8 agent than with the other agents, showing no bias in favor of an agent of the same model. Therefore, we adopt GLM-5.2-FP8 as the user simulator for J-tau.

[13] https://huggingface.co/Qwen/Qwen3-235B-A22B-Instruct-2507

[14] https://artificialanalysis.ai/evaluations/tau2-bench