arXiv:2610.19848v1 [cs.CL] 14 Oct 2026

OpenTWBench (Open Taiwan Benchmark): Taiwan-Domain Knowledge Evaluation for Large Language Models, and How the Measurement Apparatus Shapes the Score

Liang-Hsun Huang[1,2,3] Min-Yi Chen[1,2,3] Chao-Chun Chuang[4,5]
Yu-Wei Liu[6] Che-Yu Wu[5] Po-Feng Lee[7]
[1]Twinkle AI [2]APMIC [3]Open Taiwan Benchmark*
[4]National Center for High-performance Computing (NCHC), Hsinchu, Taiwan
[5]MonoLab [6]Google Developer Expert, Cloud AI
[7]Open Culture Foundation (OCF)

Abstract

We present OpenTWBench, an open evaluation suite for Taiwan-domain knowledge in large language models, built around two complementary benchmarks. TW-EXAM-BENCH contains 255,854 four-option questions across 148 academic subjects, parsed from fifteen years of Taiwan's national examinations published as open data by the Ministry of Examination; answer keys agree across 31,312 duplicated question groups at 99.91%, and every one of the 27 disagreements is accounted for. A reviewed second pass withdrew a further 2,180 items whose stems depend on a figure the text does not carry, so every score we report is computed on the remaining 253,674. TW-FORMOSA-BENCH contains 1,696 difficulty-curated questions about everyday and contemporary Taiwan — the Taiwan that licensing examinations never ask about — generated from a community-curated corpus under a construction discipline in which no model ever writes a ground truth, and curated against a frozen seven-model pool so that the pool scored at the passing margin (46.9–74.6%, median ≈54%), with a leave-one-out experiment showing the selection preserves model ranking at near-zero transfer bias. Within six weeks the frontier had cleared that band: of the 108 models now on the public board, 22 score above its ceiling and the best scores 91.6%.

Alongside the datasets, we document how much of a benchmark score is contributed by the measurement apparatus rather than the model. On identical questions and models, switching answer extraction from pattern matching to explicit `\boxed{}` moves accuracy from 20.8% to 93.9%; a single placeholder character in a system prompt moved one model from last place (58.2%) to first (83.9%) on the same benchmark; doubling the thinking budget from 8k to 16k tokens bought 0.2% while shuffling which questions failed; and a three-judge validation panel proved to be one judge in disguise, with two members approving 92–99% of everything shown to them. We release all datasets, per-item difficulty data, curation manifests, and a public leaderboard at `https://opentwbench.ai`.

1 Introduction

Traditional Chinese as written in Taiwan is not a font variant of Chinese; it is a linguistic and institutional context of its own, with its own calendar (the Minguo era), legal system, medical

Open Taiwan Benchmark is a non-profit organization that develops and maintains open, reproducible evaluation benchmarks for Taiwan-domain language models, including the OpenTWBench suite presented in this paper.

licensing regime, place names, vocabulary, and everyday culture. Language models are routinely deployed for Taiwanese users, yet the evaluations behind those deployments are mostly translations of English benchmarks or Simplified-Chinese suites, and the few Taiwan-native benchmarks that exist are small, static, and increasingly saturated. This paper describes OpenTWBench, an open evaluation suite for Taiwan-domain knowledge, and — with equal emphasis — what we learned about how much of a reported benchmark number is produced by the measurement apparatus rather than by the model being measured.

Two complementary benchmarks. TW-EXAM-BENCH asks what national examinations ask: 255,854 four-option questions across 148 academic subjects, parsed from fifteen years of open-data papers published by Taiwan's Ministry of Examination, with official answer keys whose internal consistency we can measure (99.91% agreement across 31,312 duplicated question groups, every one of the 27 disagreements manually attributed; Section 3.2). TW-FORMOSA-BENCH asks what examinations never ask: 1,696 closed-book questions about everyday and contemporary Taiwan — food, transit rules, public figures, local institutions — generated from a community-curated corpus under a construction discipline in which no model ever writes a ground truth, then *difficulty-curated* against a frozen pool of seven contemporary open-weight models, so that the pool scored at the passing margin (46.9–74.6%, median ≈54%). A leave-one-out experiment shows the curation preserves model ranking with near-zero transfer bias for mid-field models (Section 4.2). The band held for about six weeks; Section 6.3 reports what the frontier did to it, and what that says about how long any difficulty-targeted benchmark can expect to hold.

The apparatus is part of the result. Building the suite forced us to measure things benchmark papers usually leave implicit, and the sizes are uncomfortable. On identical questions and models, switching answer extraction from pattern matching over free-form output to an explicit `\boxed{}` protocol moves accuracy from 20.8% to 93.9% — a 73% swing contributed entirely by the harness. A single placeholder character in a system prompt (甲 in the format example) moved one 675B model from last place (58.2%) to first place (83.9%) on the same benchmark, because the model copied the example's answer alphabet; the same prompt cost a different model 17 questions out of 60,389. Doubling a reasoning model's thinking budget from 8k to 16k tokens bought 0.21% of accuracy at 1.72× the cost, while silently exchanging *which* questions failed. And a three-judge "consensus" validation panel turned out to be one judge in disguise: two of its members approved 92–99% of everything shown to them (Section 5). None of these effects are visible in a leaderboard number; all of them are larger than typical claimed model differences.

Contributions.

1. TW-EXAM-BENCH: the largest Taiwan-native examination benchmark to date (255,854 items, 148 subjects, 15 years), with a measured key-consistency floor and a fully attributed conflict list.

2. TW-FORMOSA-BENCH: a closed-book everyday-Taiwan benchmark with a zero-model-authored-ground-truth construction pipeline and a difficulty-curation method validated by leave-one-out transfer.

3. A quantified account of measurement-apparatus effects — extraction protocol, prompt placeholders, thinking budget, judge-panel degeneracy — each reproducible from released data.

4. A public leaderboard with per-item difficulty sidecars and frozen curation manifests, so every selection decision is auditable.

All datasets, manifests, and evaluation code are released; the leaderboard lives at `https://opentwbench.ai`.

2 Related Work

Taiwan-native benchmarks. TMMLU+ [8] is the closest prior work: multiple-choice questions in Traditional Chinese drawn largely from the same Ministry of Examination source we use, organized into 66 subjects. OpenTWBench is best understood as a scaled-up successor on the examination axis — 255,854 items against TMMLU+'s ∼22k, 148 subjects against 66, with fifteen years of coverage and a measured key-consistency floor — plus an everyday-knowledge axis (TW-FORMOSA-BENCH) that examination sources cannot provide at all. Because the two suites share a source, the same duplicate audit we run on our corpus (Section 3.2) applies to TMMLU+, and our build drops 2,948 items whose content fingerprints match TMMLU+ so that the suites stay disjoint. Auditing TMMLU+ internally by order-insensitive content matching finds 821 duplicate groups (1,720 of 23,072 rows); compared by answer *text* rather than letter — option order differs across appearances, and letter comparison manufactures 321 false conflicts — exactly one group truly disagrees, a very low label-error rate. More consequentially, 431 of those groups span different splits, and 136 of the 20,160 test items appear verbatim in a non-test split (train, val, or the few-shot dev pool) — a contamination channel for anyone who trains or few-shots on those splits and reports on test (Appendix B). We release the audit script and matching lists.

TC-Eval [3] and the TAIDE bench cover Traditional Chinese understanding at smaller scale; MT-Bench-TW and derivatives translate English benchmarks, inheriting their topic distribution rather than Taiwan's.

Chinese-language suites. C-Eval [4] and CMMLU [5] are Simplified-Chinese examination benchmarks for the PRC context; their subject matter (PRC law, PRC accounting standards) and orthography make them poor proxies for Taiwan deployment, and we treat Simplified-character and mainland-vocabulary screening as a first-class filter (Section 4.1).

Examination-derived English benchmarks. MMLU [2] established examination questions as an evaluation substrate and also established their failure mode: label errors and saturation [1]. Our response to the error problem is to measure it at the source — duplicate answer keys across fifteen years give an empirical noise floor ($\sim 10^{-3}$) rather than an assumption — and our response to saturation is difficulty curation with transfer validation rather than manual re-authoring.

Benchmark-measurement effects. A growing literature documents that scores depend on the harness: prompt-format sensitivity [7], choice-ordering artifacts [9], and LLM-judge biases [10]. Our contribution here is not the phenomenon but its magnitude and specificity in a CJK setting — a single CJK placeholder character re-ranking a leaderboard, extraction protocol dominating model choice by an order of magnitude — with paired same-model, same-item data released for each effect.

3 tw-exam-bench: 255,854 Questions from Fifteen Years of National Examinations

3.1 Source and extraction

Taiwan's Ministry of Examination (考選部) publishes every national examination paper — civil service, professional licensing, and technician certification — as open-data PDF files, together with the official answer keys. We parse 15 years of these papers into a single benchmark following the Twinkle Eval multiple-choice contract (`question`, `A`–`D`, `answer`), organized by *academic subject* rather than by profession: papers are grouped by what they examine (pharmacology, constitutional law, railway engineering), not by which professional category sat them. The suite covers 148 subjects across 19 fields, from medicine and nursing through maritime navigation to indigenous affairs.

Extraction is deliberately conservative; a paper that cannot be parsed with confidence is rejected rather than approximated. Three failure classes consumed most of the engineering effort and are worth recording because each would have silently corrupted the dataset:

Private-use glyphs. The Ministry's PDFs encode circled option markers as End-User Defined Characters in the Unicode Private Use Area (U+E18C–U+E190); text extraction yields tofu unless these are mapped back to A–E. Normalization must use NFC rather than NFKC, which would destroy other distinctions the papers rely on.

Answer-grid alignment. Official answer keys are laid out as visual grids whose reading order does not match question order, and the first cell of a row is frequently glued to its row label. Keys are therefore paired to questions by page coordinates, with column tolerance derived from the grid's own pitch. A one-cell shift here assigns every answer to its neighbor — the kind of error that is invisible in aggregate statistics and fatal to a benchmark; it was caught by manually verifying individual questions against domain knowledge.

Phantom questions. Bare numbers inside question stems (statute articles, dosages) can be mistaken for question numbers. Question numbering is enforced to be monotonic, and any paper whose stated question count (單選題數, 題數) disagrees with the parsed count is rejected whole.

3.2 Deduplication and key consistency

Questions recur across years and examination categories. An order-insensitive fingerprint (stem plus sorted option texts) identifies 31,312 duplicate groups. Crucially, duplicates are compared by *answer text*, not answer letter: option order varies across appearances, and comparing letters manufactures false conflicts. (We initially made this mistake, observing 322 apparent conflicts in an external dataset that reduced to exactly one real conflict once compared by text; see Section 2.)

Across all duplicate groups, official answer keys agree at **99.9138%**. All 27 flagged groups were manually attributed (Appendix A): 14 depend on figures absent from the extracted text, 9 are detector artifacts (identical keys differing only in whitespace or punctuation), one option set was legitimately rewritten between years, one stem points at a case vignette outside the item, and only **two** are true conflicts — self-contained stems whose official keys genuinely disagree. We quote the conservative detector-level 99.9138% as the noise floor a benchmark built from these papers inherits; the attributed floor is an order of magnitude lower.

3.3 Filtering

Four filters remove items that cannot be answered from text alone, each calibrated against real false positives rather than assumed correct: figure-dependent items (detected by deictic reference only — a loose match on 圖 deletes library-science questions about 圖書分類); items continuing a previous item (承上題); English-language items, judged on the stem only after two failed iterations that deleted 18,634 valid medical questions containing English option terms; and items whose correct option is pathologically longer than its distractors.

Table 1 gives the exact pipeline accounting, aggregated from the per-subject build statistics carried on each dataset's card, so every row can be re-derived from what is published.

Pipeline stage	Count
Exam papers seen / parsed	13,812 / 11,621
rejected: incomplete parse / no questions / download failed	2,002 / 188 / 1
Raw items extracted	629,343
Folded as re-printings of an earlier year's item	265,736
English-language items	70,386
Not four-choice	14,779
Figure-dependent	6,113
Officially voided (送分)	5,430
Multi-select	3,513
Owned by an upstream benchmark (TMMLU+ overlap)	2,948
Continues a previous item (承上題)	2,876
Bloated correct option	1,198
Duplicate options within an item	326
Stem too short	132
Cross-year answer-key conflict	27
No question text	25
Items published	**255,854**
Figure-dependent, second pass (Section 3.4)	2,180
Items scored	**253,674**

Table 1: TW-EXAM-BENCH build accounting across all 148 subjects. Rows below the first rule are filtered items; folding of duplicate groups (31,312 groups across the 148 published subjects) is counted separately since those items are re-printings, not defects. The published set is what the dataset ships; the scored set is what every number in this paper is computed on, after the second-pass withdrawal described in Section 3.4.

3.4 A second pass the first one needed

The deictic filter of Section 3.3 removed 6,113 figure-dependent items, and we reported that as the figure problem solved. It was not. Reviewing low-scoring subjects after the first release, we found items whose stems depend on an absent figure without ever using a deictic marker the filter could match — 圖中標示星號（*）之肌肉是, where the reference is a typographic symbol, or stems that describe a waveform or an ultrasound image in prose and then ask what it shows. A reviewed pass over all 148 subjects withdrew a further **2,180** items across 97 subjects: 85.8% figure references, 7.0% items continuing a question group, and the remainder referencing an absent table, passage or attachment. The withdrawals concentrate exactly where images carry the content — dentistry

(277), electronics (166), diagnostic radiology (142) — which is why aggregate scores hid them.

Two things follow. Every score in this paper is computed on the remaining 253,674 items, and the withdrawal list ships with the dataset, so the exclusion is auditable rather than a private correction.

The second is about how a filter gets validated. Section 3.3 measures what the figure rule wrongly deleted, because a wrongly deleted item is visible: the library-science questions were there to be counted. What the rule let through is not visible in the same way — nothing marks an item a filter failed to catch — so the only way to count those was to read all 148 subjects again. A filter checked on its false-positive rate has been checked on the half of its errors that never reach the published set; the second pass is the measurement the other half requires. We would expect any corpus assembled this way to need one.

4 tw-formosa-bench: Difficulty-Curated Everyday Taiwan

National examinations measure professional knowledge. They never ask who designed the Weiwuying arts center, why a famous fish-head restaurant switched to turkey during the foot-and-mouth outbreak, or what the MRT rules say about drinking water. TW-FORMOSA-BENCH covers this everyday, contemporary Taiwan: closed-book, four-option questions derived from Taiwan.md, a community-curated CC BY-SA corpus of roughly 4.4M characters across 924 articles.

4.1 Construction: no model writes a ground truth

Items are drafted by a generator model (gpt-oss-120b), one per paragraph-aligned window of the corpus (2,972 windows, median 1,119 characters, boundaries never mid-sentence). The prompt asks for facts checkable outside the corpus and forbids document-local trivia. Quality then comes from filters, not from trusting the generator:

- **Grounding.** Every keyed answer is verified against the window it was written from by a judge that had previously caught 100% of planted key defects in a calibration study (Section 5.4); 256 of 2,939 well-formed items failed and were dropped.
- **Self-containment.** Stems that open on a pronoun (她取得美國國籍的年份是？) name nobody once the source is gone and are rejected; pronouns *inside* a stem resolve locally and are legal. Stems referencing the passage (根據本文) are stripped where the opening is pure scaffolding and rejected where the reference is load-bearing.
- **Anti-shortcut.** The generator placed the correct answer in option A 84% of the time; options are permuted by a hash of each item's source window, after which "always answer A" scores 24.9% and "pick the longest option" 26.6% (chance = 25%).
- **Orthography and vocabulary.** Zero Simplified-only characters and zero mainland vocabulary written in Traditional characters (網絡, 軟件), enforced by a lexicon whose strict terms fire zero times across 4.4M characters of Taiwanese prose and 255k examination questions — membership in the strict list requires exactly that.

An earlier design filtered instead for items *unanswerable without the source*. It kept registration fees and edition numbers — corpus memorization, not Taiwan knowledge — and was discarded; the episode is a useful caution that a filter optimizes exactly what it tests.

4.2 Difficulty curation with transfer validation

From 2,665 grounded items we publish the 1,696 that at least one model of a frozen seven-model pool answers incorrectly (band pass $\leq 6/7$; pool, band and date frozen in a public manifest). On the pool this lands scores at the passing margin: 46.9–74.6%, median ≈54%. That is a statement about that pool on that date, not about the benchmark in perpetuity; Section 6.3 reports how long the band survived contact with models released after it was frozen.

Pool-relative selection invites a fair objection: is the resulting set adversarial to the pool, with scores that do not transfer to new models? We answer it empirically with a leave-one-out experiment: select the band using six models, then measure the held-out seventh on the selection.

Held-out model	n	Full set	Held-out	In-pool mean	Gap
Mistral-Large-675B	1,679	83.9%	75.4%	54.7%	+20.7
gemma-4-31B	1,633	73.9%	61.3%	56.2%	+5.1
Gemma-3-27B	1,599	70.1%	56.3%	56.3%	+0.0
Nemotron-Super-120B	1,643	70.9%	56.1%	57.2%	−1.2
gpt-oss-120b	1,622	69.3%	54.2%	57.1%	−2.9
Llama-3.3-70B	1,505	66.2%	52.9%	54.7%	−1.8

Table 2: Leave-one-out difficulty transfer. Each row selects the band pass $\leq 5/6$ using the other six models and measures the held-out model on the resulting n items. "Full set" is the held-out model's accuracy on all 2,665 grounded items. Selection compresses scores by 8–16%; mid-field models transfer with near-zero bias; stronger models rise above the pool mean, as ability — not selection artifact — predicts.

Mid-field held-out models land within ±3% of the in-pool mean; the two strongest held-out models exceed it by exactly the margin their full-set superiority predicts, and the held-out ranking matches the full-set ranking everywhere. As an additional validation with a model entirely outside the pool and outside the generator's family, we evaluate Qwen3.8-27B, which never participated in selection, under the identical protocol. It scores 72.0% on the full 2,665 items — mid-field, just under the pool's full-set mean of 73.2% — and **59.0%** [56.6–61.3] on the published 1,696: a 13.0-point compression squarely inside the leave-one-out range, landing +1.0 points above the pool's curated-set mean of 57.9%, within the ±3-point transfer bias the leave-one-out experiment predicts for mid-field models. Its nearest full-set neighbor (gemma-4-31B, 73.9% full) scores 59.0% on the curated set — a dead tie with Qwen's 59.0%, i.e. the adjacent ranking survives curation exactly.

4.3 A note on fairness

The generator (gpt-oss-120b) also serves in the difficulty pool; its own score on this benchmark carries a self-advantage and is annotated as such on the leaderboard. The curation manifest and the per-item difficulty sidecar are published so that any selection choice we made can be audited or redone.

5 The Measurement Apparatus Shapes the Score

Every result in this section holds the questions and the model constant and varies only the apparatus. The effects are ordered by size.

5.1 Extraction protocol: 20.8% vs. 93.9%

Twinkle Eval's `pattern` strategy sends the bare question with no system prompt and pattern-matches an option letter out of free-form output; its `box` strategy instructs the model to end with `\boxed{X}` and counts an unparseable response as wrong. On the same civics subject and the same model, `pattern` reports 20.8% and `box` reports 93.9%. The pattern strategy was not merely noisy: with no stop discipline the model rambled past the answer, and the matcher once extracted the letter "W" — not an option — from surrounding prose. The 73% gap is larger than the gap between the best and worst model we evaluate anywhere in this paper. Every score we publish uses `box`, and we publish parse rates next to accuracies so that "the model got it wrong" and "the harness could not read the answer" remain separable.

What the leaderboard did instead, and why it was wrong. Stating a protocol is not the same as running one. Until 14 October 2026 the public board scored each response with the `box` answer where the model emitted one and a pattern match where it did not — a fallback that reads a bare leading letter, or failing that any option letter in the text. The intent was charitable and the effect was regressive: the *less* a model complied with the answer format, the more of its score came from pattern matching, so a model that never wrote `\boxed{}` was scored entirely leniently while one that always wrote it was scored entirely strictly. The effect was not hypothetical: a model that wrote `\boxed{}` on 6% of examination items sat two places above one that wrote it on 95% and was held to it, and a model that wrote it on 32% sat six places above that same model. 185 of the 215 published cells moved when we corrected this, 122 of them by more than a point and 49 by more than ten; the median cell moved -2.4 points. The board now publishes the `box` score, with the lenient figure beside it, so that what a model knows and what format it writes it in are two columns rather than one. We record the error rather than quietly repairing it: a suite whose stated protocol and published numbers disagreed for six weeks is evidence for this section's argument, not an exception to it.

5.2 One placeholder character re-ranks the leaderboard: 58.2% → 83.9%

Our first system prompt illustrated the format as 例如 $\boxed{\text{甲}}$ — using 甲, the first of the traditional Chinese ordinals, as the placeholder. Mistral-Large-675B copied the example's alphabet and answered with 甲乙丙丁 throughout: 827 of 828 "unparsed" responses on one subject were in fact real answers in the wrong alphabet, and the model scored 58.2% — last place. Mapping the ordinals back to A–D predicted 83.4%; rewriting the prompt to name the alphabet ("one of A, B, C, D") without an example letter yielded a measured 83.9% with a 100% parse rate — first place among the pool at the time. The same defective prompt cost Qwen3.8-27B only 17 responses out of 60,389.

The lesson is sharper than "prompts matter": the defect was *model-dependent*, so it did not shift the leaderboard — it *re-ordered* it. A constant apparatus bias would cancel in comparisons; a model-conditional one silently masquerades as capability difference. We now treat any subject×model cell with an outlier parse rate as an apparatus alarm, not a model result.

5.3 Thinking budget: doubling buys 0.21%

On accounting-audit-law, raising the generation cap from 8,192 to 16,384 tokens changed accuracy by +0.21% (CI ±2.6%) at 1.72× the wall-clock cost, and unparsed responses fell only from 15.9% to 14.0%. Item-level comparison shows why the aggregate is so stable: the larger budget rescued 130 items and newly lost 103 — the unparsed set *churns* rather than shrinks. Of responses that hit the

cap, 97% ran to the limit but only 8% were degenerate repetition loops; the dominant behavior is genuine indecision, circling between statute readings without committing. Budget, in other words, is not the binding constraint on hard legal items — decisiveness is.

5.4 A judge panel that was one judge: 99% / 92% / 36%

An early Formosa validation pipeline used a three-judge panel with a unanimity rule, which sounds rigorous. Measured per judge, two members approved 99% and 92% of everything shown to them; the third approved 36%. "Three judges agree" was operationally "one judge decides." We replaced the panel after a calibration study: ten candidate judges were shown items with deliberately planted defects of three construction classes (wrong key, ungrounded fact, malformed options). Every competent candidate caught 100% of planted defects; all disagreement between judges lay in unanchored subjective dimensions ("is this interesting?"). The grounded-verification judge used in Section 4.1 is the one that passed this calibration — and the general rule we draw is that a judge is only as meaningful as the planted-defect study behind it.

5.5 Answer-key consistency as a measured noise floor

Section 3.2 treats the 99.9138% cross-year key agreement not as a quality boast but as an apparatus parameter: it bounds from below the error any evaluation on this source inherits. Reporting model differences finer than the source's own $\sim 10^{-3}$ key noise, on subsets small enough for that noise to bite, is measuring the apparatus again.

6 Results

6.1 tw-exam-bench

Full-corpus anchor. Qwen3.8-27B is our full-corpus anchor: it answered all 148 subjects at a 16,384-token budget, scoring **82.19%** [82.04–82.34] over the 253,674-item cleaned canon — the 255,854 published items less the 2,180 later withdrawn as unanswerable without their missing figure — at a 94.34% parse rate.

The 54% subject spread. The aggregate hides the structure. The model scores 97–98% on clinical medicine subjects and 43–50% on the law cluster (accounting-audit law, tax law, civil procedure) — a 54% spread within one model on one benchmark. Low parse rates (64–75%) co-occur with the law cluster: the same indecision documented in Section 5.3 burns the budget without committing to an answer. Subject-level scores are the unit at which this benchmark is informative; we publish all 148.

Multi-model stratified sample. Running 255,854 items per model is impractical for a leaderboard, so the multi-model exam board uses a deterministic stratified sample: per subject, $\min(150, n)$ items selected by content hash (22,200 rows drawn, 22,021 after the withdrawals above), with the full-corpus Qwen run as the anchor that calibrates sample-to-population error.

Items are counted as rows, not as distinct question texts. Taiwan's examinations reuse a stem verbatim across years with a fresh set of options — 下圖為何種通氣模式？ is six different questions in mechanical ventilation between 2014 and 2019, with three different answer keys — so 119 stems in the sample stand for 247 separate items. Keying results on the stem would silently discard 128 of them; it is also why the withdrawal list removes 179 rows where it names 174 stems.

The anchor calibrates that sample. Within a subject, items are selected by content hash, which is independent of whether any model answers them correctly; a sample drawn that way can differ from the population only through composition. Weighting each subject's full-corpus accuracy by the number of items the sample draws from it gives **80.1%** against the 82.19% population score — a −2.04-point shift produced by the sample's deliberate equal weighting of subjects, which moves weight away from the large, high-scoring medical corpus (subjects with ≥2,000 items average 85.5%; those under 500 average 77.4%) toward the smaller legal and administrative subjects where the model is weakest. Sample scores are therefore comparable across models but sit systematically below what a full-corpus run would report.

Model	Accuracy	`\boxed{}`
Mistral-Large-675B	83.4%	99.9%
gemma-4-31B	82.1%	100.0%
Qwen3.8-27B (anchor, composition-weighted)	80.1%	—
gpt-oss-120b	77.1%	99.9%
Nemotron-Super-120B	76.0%	89.2%
Gemma-3-27B	66.7%	100.0%
Llama-3.3-70B	66.7%	94.0%
gemma-3-4B-T1-it	46.5%	95.4%
Llama-3.2-3B-F1 (Instruct / Reasoning)	45.4% / 45.3%	98.5% / 96.7%
Llama-3.1-TAIDE-LX-8B	16.5%	32.2%

Table 3: TW-EXAM-BENCH stratified sample (22,021 rows), selected rows; the full board of 107 models is published at `https://opentwbench.ai`. The second column is the share of responses that contained a `\boxed{}` at all, and the accuracy is computed on those alone — which is why Llama-3.1-TAIDE-LX-8B, whose knowledge is comparable to the 3–4B models above it, sits last: it wrote the requested format on under a third of items. The anchor row places Qwen's full-corpus run on the sample's scale by weighting its per-subject accuracies by the sample's per-subject counts; it is an estimate on that scale, not a measurement taken on it (its full-corpus figure is 82.2%). Gemma-3-27B and Llama-3.3-70B differ by 0.06 points — a statistical tie. Scoring is row-wise over the 22,021 rows; 102 of the 106 sampled models answered all of them and four runs fell 9–36 rows short.

6.2 tw-formosa-bench

Scores of the frozen pool on the published (curated) benchmark:

Model	Accuracy
Mistral-Large-675B	74.6%
gemma-4-31B	59.0%
Nemotron-Super-120B	54.3%
Gemma-3-27B	53.1%
gpt-oss-120b	51.8%[†]
Llama-3.3-70B	46.9%

Table 4: Frozen-pool scores on TW-FORMOSA-BENCH (1,696 items). Scores sit at the passing margin by construction (Section 4.2). [†]gpt-oss-120b generated the draft items and its score carries a self-advantage annotation.

Qwen3.8-27B — outside the pool, outside the generator family, never seen during selection — scores 59.0% [56.6–61.3] on the published set (72.0% on the ungated 2,665), the fresh-model transfer validation detailed in Section 4.2. Across the frozen pool and this held-out model — the eight models the band was calibrated against — scores span 46.9–74.6% with a median of 56.6%: the passing margin the curation targeted.

Four Taiwan-tuned models land well above where their parameter count predicts on the everyday-Taiwan axis (all served locally in bf16 — there is no quantization anywhere in the suite). Twinkle AI's gemma-3-4B-T1-it scores **46.2%** [43.9–48.6] on TW-FORMOSA-BENCH, within a point of Llama-3.3-70B's 46.9% at 1/17 the parameters; the Llama-3.2-3B-F1 pair lands at 40.9% (Instruct) and 39.0% (Reasoning); TAIDE's Llama-3.1-TAIDE-LX-8B-Chat scores **38.8%** [36.5–41.1]. All four are well clear of the 24.9% shuffled-chance floor, and all four sit above where 3–8B parameters would put them.

The TAIDE figure needs its second column to be read at all. That model wrote `\boxed{}` on 67.0% of TW-FORMOSA-BENCH items and 32.2% of examination items, so both of its scores are bounded by its compliance rather than its knowledge; held to the same pattern-matching extraction the rest of this paragraph's models barely need, it scores 55.4% and 47.9%. The comparison we can make without that caveat is among the compliant models, where `\boxed{}` rates run 83–99%: a 4B Taiwan-tuned model is level with a 70B general model on everyday Taiwan while scoring 20 points below it on professional examinations.

On TW-FORMOSA-BENCH the reasoning-tuned F1 variant scores *below* its instruct sibling (39.0% vs. 40.9%) at an indistinguishable examination score (45.3% vs. 45.4%) — closed-book recall is not a reasoning problem, and thinking longer cannot retrieve what the weights do not store (the same mechanism as the budget result of Section 5.3).

6.3 Benchmark half-life: the band held six weeks

The difficulty band of Section 4.2 was calibrated against a pool frozen on 24 August 2026. We report what happened next, because it bounds how a difficulty-targeted benchmark should be read.

The public board carries 108 models on TW-FORMOSA-BENCH as of 14 October 2026, including frontier systems released after the freeze. Their scores span 0.0–91.6% with a median of 45.4%. Of those 108, only 29 fall inside the 46.9–74.6% band the curation targeted; 22 score above its ceiling, 8 of them above 85%, and the best scores 91.6%. On TW-EXAM-BENCH the best model scores 96.4%. Seven weeks after freezing, in other words, the band no longer describes the top of the field — it describes the middle of it.

We do not read this as the curation having failed. The leave-one-out experiment of Section 4.2 asked whether the selection transfers to models *outside the pool*, and it does: ranking is preserved and mid-field transfer bias stays within ±3 points, which is why the band still orders the 2026-vintage field correctly. What the band cannot do is stay centred on a moving frontier, and no pool-relative selection can. The honest form of the claim is therefore dated: these items sat at the passing margin for the models available in August 2026, and the difficulty sidecar is published per item so the band can be re-cut against a newer pool without regenerating the benchmark.

The practical reading is that a difficulty-targeted benchmark has two distinct lifetimes. Its *discriminative* lifetime — how long it separates models — is long: at the time of writing the published set still spreads the field over 91 points, 66 of them between models that all write the answer format and are separated by knowledge alone. Its *calibration* lifetime — how long a stated score band remains true — was about six weeks. Papers that report the second as though it were the first will be wrong within a release cycle, and we would have been had we published this one in August without dating the claim.

6.4 What the two benchmarks disagree about

Over the 107 models carrying both scores, the two axes agree far more than they disagree: Spearman's $\rho = 0.97$. That figure is inflated by range — a board spanning 0% to 96% will correlate on almost any two measures of capability — but restricting it does not rescue the difference: ρ is 0.89 over the top twenty models by examination score and 0.94 over the 27 scoring above 80%. On the evidence, a second Taiwan benchmark buys very little *general* ranking information over the first.

What it buys is a specific kind of disagreement, and the size of it depends on which extraction the question is asked under. Ranking each model on both axes and taking the difference, the seven Taiwan-tuned models on the board — the TAIDE, Twinkle AI, Breeze and Llama-3-Taiwan families — gain a mean of **6.9** places on Formosa relative to their examination rank (median +3), while the other 100 models move by a mean of −0.5 places (median 0). On the published `\boxed{}` scores that is the whole effect, and it is not a robust one: three of those seven models write the answer format on under 7% of items and sit near zero on both axes, carrying no signal in either direction, and the mean is dominated by one model whose compliance differs by a factor of two between the two benchmarks.

Holding extraction constant isolates the part of the effect that is about knowledge. On the lenient figure published beside every score — the same responses, read by a pattern matcher instead of only inside `\boxed{}` — the seven gain a mean of **15.6** places (median +17) against −1.2 for the other 94, and the median model scores 7.1 points *lower* on Formosa than on the examination while Llama-3.1-TAIDE-LX-8B-Chat scores 7.5 points *higher* and gemma-3-4B-T1-it 3.7 points higher. We report both because they answer different questions. What a leaderboard owes its readers is one stated protocol applied to everyone, which is the `\boxed{}` column. What a claim about Taiwan-specific tuning needs is extraction held fixed, because otherwise a model is being measured on a format convention that has nothing to do with Taiwan.

Either way this is a sharper justification for two axes than "the rankings reorder," and a narrower one. The examination axis measures professional knowledge and is well predicted by general capability; the Formosa axis measures everyday Taiwan, where general capability predicts most of the variance and Taiwan-specific tuning accounts for a systematic remainder — one that is clearly measurable with extraction fixed and partly obscured by format compliance without it. A suite that carried only the examination would rank the field almost identically, and would miss precisely the models built for the thing it claims to measure.

6.5 Taiwan-specific failure modes

Reading failures, not just counting them, surfaces a taxonomy specific to the Taiwan setting (drawn from structured-output evaluations during suite development; the underlying benchmark is second-wave, but the failure classes are observable in free-text answers as well):

- **Minguo-calendar arithmetic.** Given a bare Minguo year (115年 without the 民國 prefix), Llama-3.3-70B converts with a fixed off-by-one (115 → 2027); Mistral-Large hallucinates an unrelated date entirely, yet converts correctly whenever the 民國 prefix is present — a one-token feature flips the behavior.
- **CJK glyph corruption.** Llama-3.3-70B invents nonexistent characters (左骟 for 左營) — an orthography violation, not a knowledge error, and one our filters treat as a specification failure rather than an acceptable variant.
- **Romanization drift.** TAIDE emits pinyin (`banqiao`, `zuoying`) where Chinese station names are required.

- **Error detectability decides survival.** A corrupted station name fails loudly and gets retried into correctness; a wrong-but-legal date fails silently and survives to the final answer. Downstream validation is a safety net only where errors are loud.

7 Discussion

Quality cannot be requested, only filtered. A recurring pattern across the suite's construction: asking a generator model for a property does not produce it. Explicitly prompting for corpus-independent questions *lowered* the corpus-independence rate (40% → 32% in our generator trials); asking three judges for consensus produced one judge's opinion. Every quality property we actually obtained came from a filter with a measurable false-positive story — grounding verification calibrated on planted defects, shortcut baselines measured after permutation, screening lexicons validated for zero false fires on 4.4M characters of native prose. We suggest this as a design stance for benchmark construction generally: generation proposes, filtering disposes.

Report the apparatus with the score. The effects in Section 5 range from 0.2% to 73%. Model cards and leaderboards report the score; almost none report the extraction protocol, the prompt's placeholder characters, the parse rate, or the judge approval distribution — parameters that our data show can dominate the ranking. Our leaderboard publishes parse rate beside accuracy for every cell, and we would argue any CJK benchmark should: the placeholder incident is invisible without it.

Difficulty curation is a lens, not a distortion — if validated. Selecting items a frozen pool finds hard invites the charge of building an adversarial set. The leave-one-out experiment is the answer we owe: ranking preserved, mid-field transfer bias within ±3%, and stronger models rising above the pool mean exactly as ability predicts. What the experiment does not buy is permanence: a pool-relative band tracks the pool, and Section 6.3 measures how quickly a stated band stops describing the top of the field. Curation choices should be published with the date and pool attached, as ours are. The alternative interventions we measured were far weaker: pruning all-correct items moved top scores by only 0.6–3.0% (the all-correct fraction is small, and numerator and denominator shrink together). Scores are lowered by writing harder items or selecting harder subsets — not by deleting easy ones.

Limitations. (1) The difficulty pool is seven open-weight models served through one provider, frozen in August 2026. The board has since grown to 108 models and the frontier has cleared the band (Section 6.3), which validates the selection's transfer but also dates it: the published band reflects one pool at one moment, and re-cutting it against a current pool is left to the per-item difficulty sidecar rather than done here. (2) TW-FORMOSA-BENCH's draft items come from one generator; the zero-ground-truth discipline bounds the damage a generator can do to *correctness*, but topic distribution still reflects its choices and the source corpus's coverage. (3) TW-EXAM-BENCH inherits the Ministry's own $\sim 10^{-3}$ key noise; we can measure it but not remove it. (4) Both benchmarks are four-option multiple choice; free-form and agentic Taiwan-domain evaluation (tool use with Minguo dates and CJK arguments, long-context, instruction following) is built and measured but deliberately held for a second release, where its own apparatus story — including a cross-leaderboard calibration showing our first multi-turn design ran 2.4–7.6× easier than BFCL v3 [6] — deserves full treatment.

Knowledge first, then agency. Most benchmark effort at the time of writing goes to coding and agentic ability, and that is a reasonable place for the field's attention — but it is not a place a Taiwan-specific suite can start. An agent booking a high-speed rail seat, reading a statute or filling in a government form fails on a Minguo date or a corrupted station name long before it fails on planning (Section 6.5): those are knowledge and orthography failures, and no coding benchmark will ever surface them. TW-EXAM-BENCH and TW-FORMOSA-BENCH are therefore the first two instruments of an open Taiwan evaluation suite rather than the whole of it. The agentic layer — tool use with Minguo arguments and CJK parameters, long context, instruction following in Traditional Chinese — is built and measured already, and follows in the second release. Open Taiwan Benchmark's remit is that whole sequence, published on the same terms throughout: open data, frozen manifests, dated pools, and the apparatus reported beside every score.

Conclusion. OpenTWBench contributes the largest Taiwan-native examination benchmark, a validated difficulty-curated everyday-knowledge benchmark, and — we would argue most durably — a worked demonstration that the measurement apparatus is a first-class experimental variable. The datasets will age — we measured how fast, and it was faster than we expected — but the discipline of publishing parse rates, curation manifests, dated pools, withdrawal lists and planted-defect judge calibrations alongside scores should not.

Acknowledgements

Four kinds of help made this suite possible, and no two of them came from the same place.

The **National Center for High-performance Computing** provided the B200 compute and the API access behind the leaderboard, the latter through **Taiwan AI RAP** (`https://rap.genai.nchc.org.tw/`), which served part of the board. Both halves mattered and they are not interchangeable: the hosted endpoints reach the frontier models that cannot be run locally at all, and the B200 capacity serves the large open-weight models that no single workstation can hold. The full-corpus anchor — 253,674 items across 148 subjects — and the 108-model board are both downstream of that support.

The **Open Culture Foundation** sponsored the Apple Silicon Mac Studio that serves every model our CUDA machine cannot. The limit there was not only size: a 24 GB consumer GPU capped at CUDA 12.5 cannot start a runtime new enough for several recent architectures, whatever their parameter count, and 256 GB of unified memory removes both constraints at once. The smaller open-weight models on the board, down to the sub-billion ones, are evaluated there under the identical protocol.

The **Google Developer Expert Program** provided the Gemini credits behind the four Gemini models on the board. They are not a marginal addition to it: they hold the first three places on both benchmarks, and the 91.6% TW-FORMOSA-BENCH ceiling and 96.4% TW-EXAM-BENCH ceiling that Section 6.3 measures the difficulty band against are theirs. A benchmark that cannot reach the frontier cannot report where the frontier is, and the half-life result is the part of this paper that depended most directly on being able to. The programme that provided them recognises one of this paper's authors; we note that rather than leave it implicit, in the same spirit as the Taiwan.md and generator annotations.

TW-FORMOSA-BENCH exists because **Taiwan.md** exists: an openly licensed, community-curated corpus of everyday and contemporary Taiwan, created and maintained by Che-Yu Wu, broad enough that a benchmark about what the examinations never ask could be built from it at all. We note the overlap rather than leave it implicit — the corpus one of this paper's authors maintains is the

corpus TW-FORMOSA-BENCH is drawn from — in the same spirit as the generator annotation of Section 4.3. A benchmark can only measure what somebody first took the trouble to write down, and that corpus was years of work before there was an evaluation suite to use it.

References

[1] Aryo Pradipta Gema, Joshua Ong Jun Leang, Giwon Hong, Alessio Devoto, Alberto Carlo Maria Mancino, Rohit Saxena, Xuanli He, Yu Zhao, Xiaotang Du, Mohammad Reza Ghasemi Madani, Claire Barale, Robert McHardy, Joshua Harris, Jean Kaddour, Emile van Krieken, and Pasquale Minervini. Are we done with mmlu? In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5069–5096, 2025.

[2] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021.

[3] Chan-Jan Hsu, Chang-Le Liu, Feng-Ting Liao, Po-Chun Hsu, Yi-Chang Chen, and Da-shan Shiu. Advancing the evaluation of traditional chinese language models: Towards a comprehensive benchmark suite. *arXiv preprint arXiv:2309.08448*, 2023.

[4] Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu, Chuancheng Lv, Yikai Zhang, Jiayi Lei, Yao Fu, Maosong Sun, and Junxian He. C-eval: A multi-level multi-discipline chinese evaluation suite for foundation models. *Advances in Neural Information Processing Systems*, 36, 2023.

[5] Haonan Li, Yixuan Zhang, Fajri Koto, Yifei Yang, Hai Zhao, Yeyun Gong, Nan Duan, and Timothy Baldwin. Cmmlu: Measuring massive multitask language understanding in chinese. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 11260–11285, 2024.

[6] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The berkeley function calling leaderboard (bfcl): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392, 2025.

[7] Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models' sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. In *International Conference on Learning Representations*, 2024.

[8] Zhi-Rui Tam, Ya-Ting Pai, Yen-Wei Lee, Jun-Da Chen, Wei-Min Chu, Sega Cheng, and Hong-Han Shuai. An improved traditional chinese evaluation suite for foundation models. *arXiv preprint arXiv:2403.01858*, 2024.

[9] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations*, 2024.

[10] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica.

Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36, 2023.

A Answer-Key Conflicts

Table 5 lists every one of the 27 duplicate groups (of 31,312) our detector flagged as answer-key disagreements across printings, each manually attributed from its full printings. The attribution matters more than the count: **14** depend on figures absent from the extracted text (two different flow-cytometry plots, ECG traces, or ventilator waveforms reusing the same stem — the keys "disagree" because the pictures differ); **9** are detector artifacts where every printing keys the same letter and the answer texts differ only in whitespace or punctuation (0℃為100Ω vs. 0℃為100 Ω); **1** had its option set legitimately rewritten between years so the key followed; **1** has a deictic stem (該脊髓損傷患者) pointing at a case vignette outside the item; and only **2** are true conflicts — self-contained, text-only stems whose official keys genuinely disagree across years. The figure-dependent and vignette groups are exactly what the filters of Section 3.3 remove; the two true conflicts put the irreducible label noise of the source at $2/31{,}312 \approx 6 \times 10^{-5}$ per duplicate group, an order of magnitude below the conservative detector-level figure quoted in the body.

Subject	Question (truncated)	Official keys by year	Class
clinical-hematology	下圖是一位健康成年人的骨髓細胞以流式細胞儀（flow cytometry）分...	104: D (lymphocyte); 105: D (monocyte)	figure
clinical-physiology-pathology	這張心電圖最恰當的診斷為何？...	105: A (第一度房室間傳導阻斷); 107: C (第三度房室間傳導阻斷)	figure
critical-care-respiratory	先天性心臟病的兒童經過開心手術後，常有許多術後心臟功能異常，下列心血管壓力監...	103: A (左心室功能異常：右心房壓力正); 103: A (左心室功能異常：右心房壓力：)	typo.
dental-technology	圖中標示星號（*）之肌肉是：...	101: A (顳肌（temporal mu); 101: D (咬肌（masseteric)	figure
dental-technology	圖中，上顎齒列依殘留齒與缺牙區的位置關係為甘迺迪（Kennedy）分類法中之...	102: C (I 類-3 級); 102: D (I 類-4 級)	figure
dentistry	圖中黑色箭頭所指的解剖位置為何？...	111: D (left mandible); 112: A (left maxilla)	figure
freight-operations	依據國際海運危險品準則（IMDG CODE）規定，此標籤屬於何種類別之物質？...	101: B (放射性物質); 102: C (腐蝕性物質)	figure
freight-operations	依據配艙法則，為滿足船舶具有適當穩度，對具有二層艙之船舶來說，其下大艙的貨重...	101: B (1/3); 103: D (2/3)	**true**
legal-knowledge	關於婦女保障名額之規定，下列何者正確？...	106: B (我國各選舉區選出之直轄市議員)	typo.
marine-electrical-control	於直流電路中有兩只同型電容器，其規格為500 V、0.1μF。今若予串聯，試...	102: C (1000 V、0.05μF); 103: C (1000 V，0.05μF)	typo.

Subject	Question (truncated)	Official keys by year	Class
marine-electrical-control	使用磁脈衝感測器，其磁頭輸出的信號近似：...	102: C (正弦波); 103: A (方波)	**true**
marine-engine-maintenance	PT100之温度量測元件，在檢驗好壞時，可使用歐姆計量測其電阻，其基準電阻為...	101: B (0℃為100Ω); 104: B (0℃為100 Ω)	typo.
maritime-communications	下圖為救援單位指引遇險船舶人員登岸之信號，其意義為：...	101: B (此處登岸極危險); 101: A (這是最佳登岸地點) ...	figure
maritime-communications	下圖為搜救船舶回答搜救飛機之信號，其意義為：...	101: C (收到信號，無法配合); 101: B (收到信號，可配合)	figure
mechanical-ventilation	下列圖形顯示病人使用的通氣模式為何？...	102: B (壓力目標控制型通氣（pres); 103: A (容積目標控制型通氣（volu) ...	figure
mechanical-ventilation	病人使用呼吸器顯示下列圖形時，何種處理方式最正確？...	102: D (尋找漏氣的原因); 102: D (氣道分泌物的抽吸) ...	figure
mechanical-ventilation	病人使用容積控制型通氣模式，恆定吸氣流速，病人生理參數如下，下列敘述何者正確...	105: B (僅②③⑥); 106: C (僅②④⑤)	revised
mechanical-ventilation	下圖為何種通氣模式？...	106: B (容積控制型通氣（volume); 106: A (壓力控制型通氣（pressu)	figure
molecular-diagnostics	下圖為何種儀器？...	102: A (尿液比重計（urinomet); 104: B (折射儀（refractome)	figure
neuro-physical-therapy	該脊髓損傷患者符合那種臨床症候群？...	112: C (布朗－斯夸症候群); 114: B (中髓症候群)	vignette
occupational-therapy	人類職能模式（ Model of Human Occupation） 認為人...	101: B (職能參與（ occupati); 105: B (職能參與（occupatio)	typo.
occupational-therapy-techniques	有關射飛鏢活動之分析，下列敘述何者錯誤？...	101: C (Biceps 是射出飛鏢的重); 105: C (Biceps是射出飛鏢的重要)	typo.
pharmaceutical-analysis	請問下圖為何種儀器的裝置示意圖？...	104: B (螢光光譜儀); 105: C (紫外光－可見光光譜儀)	figure
physical-dysfunction-ot	下列何項感覺測試，不適合用來評估脊髓損傷的受傷層級？...	101: A (觸辨覺（ stereogno); 106: A (觸辨覺（stereognos)	typo.
radiographic-equipment	加馬攝影機中，下列有關其輸出信號的脈衝振幅的敘述，何者正確？...	101: C (與撞擊NaI(Tl)晶體的加); 104: C (與撞擊NaI（Tl）晶體的加)	typo.
respiratory-diseases	下列ECG是那一種心律不整？...	101: D (Paroxysmal sup); 103: C (Atrial fibrill)	figure
veterinary-pharmacology	何者為 enalapril 擴張血管的主要作用機制？...	101: D (抑制 angiotensin); 104: D (抑制angiotensin-)	typo.

Subject	Question (truncated)	Official keys by year	Class

Table 5: All 27 flagged answer-key conflict groups with manual attribution. *figure*: stem names a figure the text does not carry; *typo.*: printings differ only in whitespace or punctuation (detector artifact, not a conflict); *revised*: option set legitimately rewritten between years; *vignette*: deictic stem pointing at case context outside the item; *true*: self-contained stem whose official keys genuinely disagree.

B TMMLU+ Overlap Lists

We ran the same duplicate audit on TMMLU+ that Section 3.2 runs on our own source, grouping byte-identical items (NFKC-normalised stem plus the *sorted* option multiset, so that a re-printing that shuffles the letters is the same item and a letter change alone is not a conflict). Across its 23,072 rows: **821** duplicate groups covering 1,720 items; exactly **one** true answer-key conflict — a polymer-counting chemistry item keyed once as "5" and once as "7"; **431** groups recur across different splits (899 items); and **136** of the 20,160 test items appear verbatim in a non-test split (train, val, or the few-shot dev pool), including 22 in train itself. The label accuracy is excellent — one conflict in 821 groups (0.12%), the same order of magnitude as our source's 27 in 31,312 (0.09%) — but the cross-split recurrence means few-shot evaluation on TMMLU+ can show a model its own test items as exemplars. The audit script and machine-readable lists are released alongside the datasets (`scripts/audit_tmmluplus.py`, `data/build-stats/tmmluplus_overlap.json`).

C Prompts

Evaluation system prompt. Both versions differ only in the sentence that names the answer format; the defective version's example letter 甲 is what Section 5.2 traces the Mistral collapse to.

Defective (with placeholder example):
你是一位專業的測驗作答助理。請仔細閱讀題目，以繁體中文（臺灣用語）思考， 並在最後將唯一正確選項的英文字母放入 `\boxed{}` 中， 例如 `\boxed{甲}` 的位置應填入選項字母。 除了 `\boxed{}` 之外不要使用其他格式標示答案。

Final:
你是一位專業的測驗作答助理。請仔細閱讀題目，以繁體中文（臺灣用語）思考， 並在最後輸出 `\boxed{}`，大括號內只填入唯一正確選項的 英文字母，也就是 A、B、C、D 其中一個。除了 `\boxed{}` 之外不要使用其他格式標示答案。

Formosa generation prompt (verbatim).

你是臺灣的測驗命題委員，正在編寫一份「認識臺灣」的知識評測集。 受測者只會看到題目，看不到任何參考資料，必須靠自己對臺灣的認識作答。

短文只是你的取材來源，不是題目的一部分。請從短文中挑出**值得知道的 臺灣知識**來出題。

必須遵守：

1. 題幹要自己說得清楚。人名、地名、機構名、作品名一律寫出來。 絕對不可以用「他」「她」「它」「該」「此」「這位」來指涉短文裡的人事物。 反例（禁止）：「她取得美國國籍的年份是？」 正例：「作家陳若曦取得美國國籍是在哪一年？」

2. 絕對不可以出現「短文」「本文」「文中」「上文」「文章」 「根據以上」等字眼，也不可以用「製作說明中」「文件裡」 這種暗指來源的說法。

3. 只出在其他地方也查得到的臺灣知識 —— 歷史事件、人物事蹟、地理、制度法規、 文化習俗、著名作品、產業發展、自然生態。

4. 不要出只有這份文件才知道的細節：報名費與價格、某一屆活動的數字、 網站或專案的內部統計、功能上線或下架的日期、單一公司的內部流程、 某個專案有幾個圖層。自問：「一個很懂臺灣但沒讀過這份文件的人， 有沒有可能知道答案？」若不可能，換一題。

5. 三個誘答選項要合理但明確錯誤，長度與正確答案相近。

6. 一律使用臺灣用語與繁體中文（「網路」而非「網絡」、「軟體」而非「軟件」）。

7. 題幹不可直接出現答案文字。

Grounding-judge prompt (verbatim).

你是臺灣測驗題目的審題委員。你會看到一段來源文字和一道依據它出的選擇題。

請判斷兩件事：

— supported：標示為正確的那個選項，是否真的能從來源文字中找到依據？ 改寫、摘要、合併來源中的多個事實都算有依據。若該選項的內容在來源裡找不到， 或來源其實支持別的選項，就是 false。

— single_answer：是否只有一個選項正確？若有第二個選項也說得通， 就是 false。

只輸出 JSON：{"supported": true, "single_answer": true, "reason": "一句話"}

D Reproducibility

All evaluations use Twinkle Eval's `box` strategy with temperature 0, a 16,384-token generation budget for the full-corpus exam run, and published parse rates per cell. The Formosa curation manifest freezes the pool membership, band, and date; per-item difficulty sidecars allow any alternative band to be recomputed. Table 6 lists every evaluation run behind a number in this paper.

Run	Models	Serving	Dates (2026)
TW-EXAM-BENCH full corpus	Qwen3.8-27B	self-hosted vLLM	Aug 22–24
TW-EXAM-BENCH budget study	Qwen3.8-27B	self-hosted vLLM	Aug 22
Formosa pool (7 models)	NCHC catalog (see caption)	NCHC API	Aug 23
Formosa selection frozen	—	manifest v2.1	Aug 24
Formosa held-out	Qwen3.8-27B	self-hosted vLLM	Aug 24
Formosa generation	gpt-oss-120b	NCHC API	Aug 23
Judge calibration	10 candidates	NCHC API	Aug 23

Table 6: Runs, serving paths, and dates. The frozen pool is Mistral-Large-3-675B-Instruct-2512, Llama-3.1-405B-Instruct-FP8, gemma-4-31B-it, NVIDIA-Nemotron-3-Super-120B-A12B, Gemma-3-27B, gpt-oss-120b, and Llama-3.3-70B-Instruct, all served by Taiwan's National Center for High-performance Computing (NCHC). All runs use Twinkle Eval's `box` strategy, temperature 0, `top_p` 0.9, `shuffle_options` on; generation caps are 16,384 tokens for the full-corpus exam run and 4,096 for all Formosa runs. Datasets, per-item difficulty sidecars, the curation manifest, per-cell parse rates, and the build manifests are published with the benchmark repositories.