arXiv:2610.22039v1 [cs.SE] 15 Oct 2026

WEAVE: Weaving Real-World Code Evolution into Dependency-Structured Tasks for Benchmarking and Training Code Agents

Weihao Li[1]*, Zihang Xu[2]*, Zhongqiu Wang[1]*
Songwei Yu[4], Wenqiang Zhu[1], Yuechen Luo[3]
Jingxuan Xu[4], Yanan Wu[1], Haoyang Huang[1], Ken Deng[1]†
[1]Independent Researchers, [2]Shanghai Artificial Intelligence Laboratory
[3]Tsinghua University, [4]National University of Singapore
12532572@mail.sustech.edu.cn

Abstract

Real-world software engineering rarely consists of isolated changes. Features, migrations, and systematic fixes often unfold across multiple pull requests (PRs), requiring coordinated modifications across a repository. Yet most executable repository-level benchmarks retain a single issue or PR as the task unit, while existing multi-PR settings rely on coarse aggregation and structure-agnostic evaluation. This mismatch obscures dependencies among changes and limits progress measurement and learning signals for long-horizon code agents. We introduce WEAVE, a framework that transforms verified PR histories into executable, dependency-structured tasks. WEAVE leverages repository-calibrated typed change loci to organize related PRs into bounded DAGs and instantiate each as a unified executable task. We further introduce Graph Credit, a deterministic, test-grounded metric that measures dependency-consistent partial progress and supports dense, verifiable rewards for reinforcement learning. Experiments on task construction, agent evaluation, supervised fine-tuning, and reinforcement learning show that WEAVE yields coherent executable tasks, enables more faithful progress measurement than flat scoring, and provides effective supervision and verifiable rewards for training code agents.

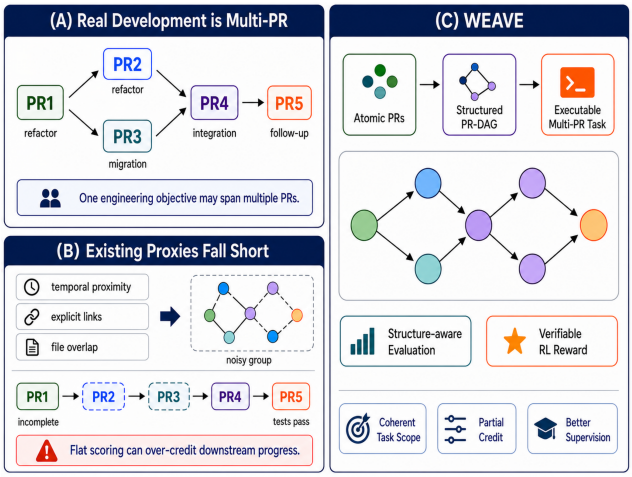

Figure 1: Real-world objectives span dependent PRs, but existing proxies obscure task structure and credit assignment. WEAVE constructs executable dependency-structured tasks for verifiable evaluation and rewards.

1 Introduction

Executable repository-level benchmarks have increasingly become a central testbed for software engineering agents. SWE-bench and its variants pair a natural-language issue with a repository state and an executable test oracle, enabling standardized, verifiable evaluation of generated patches (Jimenez et al. 2023; Zan et al. 2025; Deng et al. 2025). Most benchmarks, however, retain a single issue or pull request (PR) as the basic task unit. This abstraction captures isolated bug fixes and localized features, but only a restricted portion of real software development. In practice, a single feature, migration, dependency upgrade, or systematic fix often evolves across multiple PRs. These changes may be sequential, branch in parallel, or converge through later integration. The missing abstraction is therefore not merely longer context or a larger patch, but the *dependency structure* connecting multiple changes into a coherent engineering objective. Figure 1 summarizes this mismatch between dependency-structured development, proxy-based task construction, and structure-agnostic credit assignment.

Recent work has moved beyond isolated PRs by grouping changes through release windows, temporal spans, explicit cross-references, or object-level dependencies (Thai et al. 2025; Chen et al. 2026; Jiang et al. 2026; Zhou et al. 2026). These efforts advance long-horizon software engineering, yet two coupled challenges remain. First, these aggregation signals are imperfect proxies for task coherence: temporal proximity can merge unrelated changes, explicit references are sparse, and coarse overlap is often dominated by frequently modified utilities, configuration files, or dependency manifests. Moreover, a related change set does not itself define a reproducible task with a valid base state, task specification, and executable oracle. Second, current evaluation poorly captures structured partial progress. Binary resolution ignores progress on unsolved tasks, while flat test- or PR-level aggregation overlooks dependencies and may over-credit downstream behavior built on incomplete upstream changes. The same issue weakens reinforcement learning: binary rewards are sparse, whereas patch-similarity and learned rewards may

*These authors contributed equally.
†Corresponding author.

Table 1: Comparison of executable software-engineering benchmarks **Struct.** denotes task structure retained for evaluation; **Exec.** indicates an executable oracle; and **Train** indicates available executable training data. ∼ denotes partial support.

Benchmark	Unit	Construction	Struct.	Exec.	Credit	Train	Lang.
SWE-bench Verified	Issue–PR	Issue linked	Atomic	✓	Binary	–	Py
Multi-SWE-bench	Issue–PR	Issue linked	Atomic	✓	Binary	✓	Multi
SWE-Gym	Issue–PR	Real issues	Atomic	✓	Binary	✓	Py
SWE-EVO	Release delta	Release based	Flat	✓	Flat F2P	–	Py
SWE-CI	Commit span	Evolution span	Sequence	✓	EvoScore	–	Py
SWE-Milestone	Milestone	Commit + LLM	DAG	✓	Milestone	–	Multi
daVinci-Agency	PR chain	Explicit references	Chain	∼	–	✓	Multi
WEAVE (Ours)	**Coupled PRs**	**Typed-locus coupling**	**DAG**	✓	**Graph Credit**	✓	**Multi**

not align with executable task semantics (Tao et al. 2025; Wei et al. 2025; Wang et al. 2026; Shum et al. 2025). Addressing these limitations requires a unified, code-grounded representation for executable task construction, dependency-aware evaluation, and verifiable learning signals.

We introduce WEAVE, a framework for weaving real-world code evolution into executable, dependency-structured tasks for benchmarking and training code agents. As illustrated in Figure 1, WEAVE starts from verified atomic PRs, identifies code-grounded relations through typed change loci calibrated by repository-local frequencies, and organizes related changes into bounded dependency DAGs. Each DAG is instantiated as a reproducible repository-level task with a unified specification and executable oracle, while remaining evaluator-side metadata unavailable to the agent. Beyond task construction, the DAG provides a common interface for measuring and learning from structured progress. We introduce Graph Credit, a deterministic and test-grounded metric that complements binary resolution by rewarding dependency-consistent partial completion. The same structure supports dense, verifiable rewards for reinforcement learning, while the constructed corpus provides long-horizon trajectories for supervised fine-tuning. WEAVE thereby unifies executable task construction, structure-aware evaluation, and agent training over realistic multi-PR software evolution.

To assess both the benchmarking and training utility of WEAVE, we conduct a four-part empirical study covering task construction, agent evaluation, supervised fine-tuning, and reinforcement learning. WEAVE yields 3200 training tasks and 300 held-out benchmark tasks spanning six programming languages. Its code-grounded aggregation produces more coherent task groupings and a higher executable-task yield than temporal and coarse-overlap baselines. Evaluations of eight software engineering agents further show that Graph Credit exposes structurally inconsistent credit obscured by flat scoring. Under matched supervised-token budgets, WEAVE-based supervision improves downstream agent performance, while graph-structured rewards outperform binary and flat-reward baselines. Together, these results establish dependency-structured tasks as a principled foundation for both evaluating and improving code agents.

Contributions. Our contributions are threefold:

- **Executable dependency-structured task construction.** We introduce WEAVE, a framework that constructs bounded dependency DAGs from verified PR histories using repository-calibrated typed change loci, and instantiates each graph as a unified executable task.
- **Structure-aware evaluation and verifiable credit.** We introduce Graph Credit, a deterministic, test-grounded metric for evaluating dependency-consistent partial progress. By conditioning credit on upstream resolution, it complements binary and flat scoring while providing a dense, verifiable signal for agent optimization.
- **Empirical validation for benchmarking and training.** Extensive experiments show that WEAVE yields coherent executable tasks, captures progress missed by structure-agnostic evaluation, and provides effective supervision and verifiable rewards for training code agents.

2 Related Work

Table 1 compares representative executable software-engineering benchmarks by task unit, construction, retained structure, evaluation granularity, and training support.

2.1 Single-PR Software Engineering Benchmarks

SWE-bench established executable repository-level evaluation by pairing an issue and a pre-change repository state with F2P/P2P tests (Jimenez et al. 2023). Subsequent benchmarks improve instance quality or expand the same atomic issue–PR formulation across visual software, programming languages, recent repository states, and more challenging enterprise-scale changes (Jimenez et al. 2023; Yang et al. 2024; Zan et al. 2025; Zhang et al. 2025a; Deng et al. 2025). Executable workbenches such as SWE-Gym, R2E-Gym, SWE-Smith, and SWE-rebench V2 further scale this primitive for trajectory generation and model training (Pan et al. 2024; Jain et al. 2025; Yang et al. 2025; Badertdinov et al. 2025). Together with recent contamination analyses, these works establish the foundations of verifiable SWE-agent evaluation and training, while generally retaining an individual issue, commit, or PR as the basic task unit.

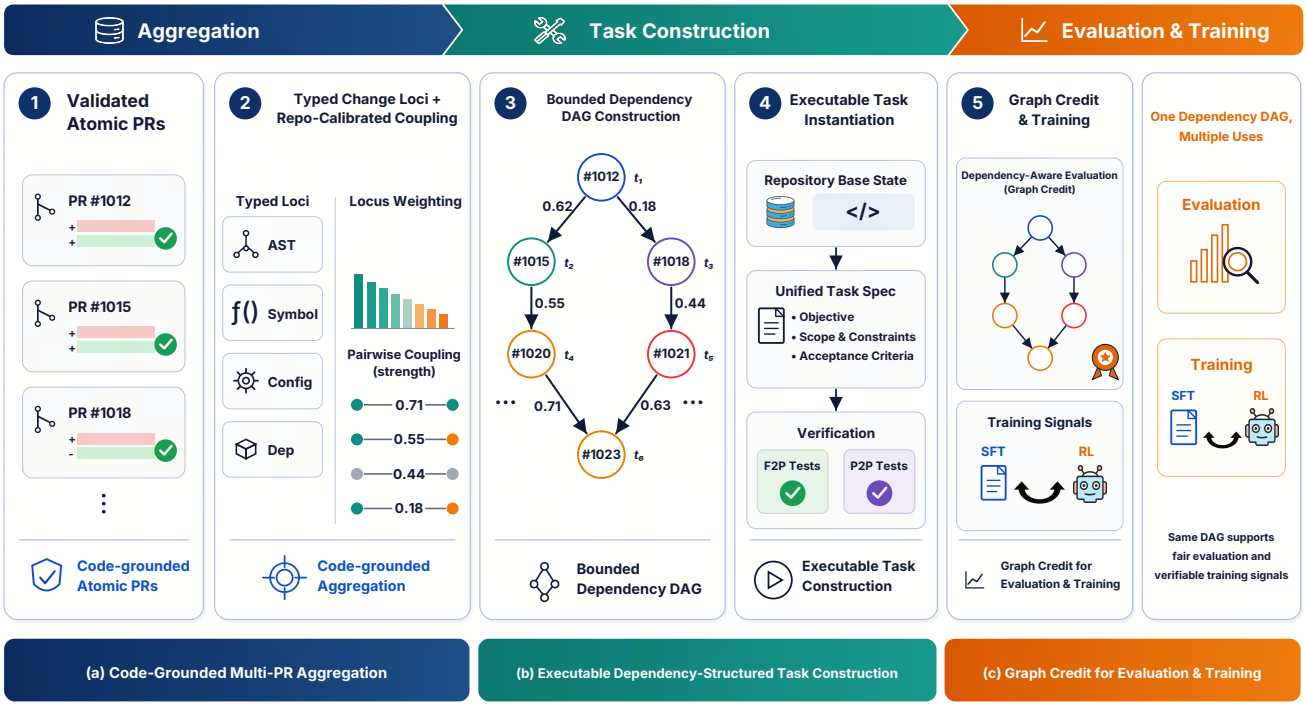

Figure 2: Overview of WEAVE. Validated PRs are represented by typed change loci, coupled into bounded dependency DAGs, and instantiated as executable tasks with unified specifications and node-attributed test oracles. The resulting graphs support structure-aware evaluation and reinforcement learning.

2.2 Multi-PR Software Engineering Benchmarks

Recent work extends software-engineering evaluation to broader units by aggregating changes across releases, features, or development roadmaps (Thai et al. 2025; Zhou et al. 2026; Xu et al. 2026), or by preserving repository evolution through temporal PR ranges, commit spans, release chains, and chronologically linked fixes (Shastry et al. 2026; Chen et al. 2026; Lam et al. 2026; Jin et al. 2026). Other approaches reconstruct milestone DAGs or mine explicitly referenced PR chains (Deng et al. 2026; Jiang et al. 2026). These benchmarks commonly measure progress through flat test completion, stage- or component-level scores, or milestone resolution. Related work has also explored denser signals based on patch similarity, weighted tests, learned reward models, process feedback, enhanced test oracles, and multidimensional evaluation (Wei et al. 2025; Wang et al. 2026; Shum et al. 2025; Dai et al. 2024; Yu et al. 2025; Tao et al. 2025; He et al. 2025; Miserendino et al. 2025), while graph structures and cross-language code representations have been used to organize agent workflows or support structural program analysis (Chen et al. 2024; Song et al. 2024; Zhang et al. 2025b; Vogel et al. 2026). WEAVE complements these directions by constructing executable PR-DAG tasks from repository-calibrated code coupling and deriving graph-conditioned partial credit from node-attributed tests, which can also serve as a dense, verifiable outcome signal during optimization.

3 Dependency-Structured Task Construction

3.1 Overview and Task Definition

WEAVE constructs executable dependency-structured tasks from validated atomic PR instances. As illustrated in Figure 2, the construction contains two stages. First, WEAVE represents each PR by the program entities it modifies and aggregates strongly coupled PRs into bounded dependency DAGs. Second, each candidate graph is instantiated as a single-environment task with a unified specification, executable tests, and node-level test attribution. The resulting graph is retained as evaluator-side metadata and is subsequently used for the graph-conditioned credit and reinforcement-learning methods in Section 4.

A final WEAVE instance is defined as

$$\mathcal{I} = (S_{\text{base}}, q, G, \mathcal{T}_{\text{F2P}}, \mathcal{T}_{\text{P2P}}, A)\,, \tag{1}$$

where S_{base} is the repository state presented to the agent, q is a unified task specification, $G = (V, E)$ is the dependency DAG, $\mathcal{T}_{\text{F2P}}$ and $\mathcal{T}_{\text{P2P}}$ are the functional and regression test sets, and A denotes test-to-node attribution. The agent observes only S_{base} and q; individual PR descriptions, identifiers, historical decomposition, and the graph itself are not exposed.

3.2 Repository-Calibrated DAG Construction

Temporal proximity and file-level overlap are insufficient for reliable dependency recovery: nearby PRs may be unrelated,

while frequently modified files may connect otherwise independent changes. WEAVE therefore represents PRs at the level of typed change loci and calibrates their shared evidence against repository-specific modification patterns.

Typed change loci. For each PR p, we extract a set of typed change loci

$$\mathcal{L}(p) = \mathcal{L}_{\text{ast}}(p) \cup \mathcal{L}_{\text{sym}}(p) \cup \mathcal{L}_{\text{cfg}}(p) \cup \mathcal{L}_{\text{decl}}(p) \cup \mathcal{L}_{\text{dep}}(p). \tag{2}$$

These loci capture AST declarations, lightweight symbols, structured configuration or schema entries, and dependency keys. Thus, a PR is represented by the concrete program entities and structured entries it changes, rather than by the files it touches. Type-specific extraction prevents configuration files, manifests, and lockfiles from being treated as single shared locations. Test files, generated or vendored code, and binary diffs are excluded from aggregation. Detailed extraction rules are provided in Appendix.

Repository-calibrated coupling. The evidential value of a locus depends on how frequently it is modified within its repository. Let $\mathcal{P}_R$ be the candidate PR set in repository R, $N_R = |\mathcal{P}_R|$, and

$\text{df}_R(\ell) = |\{p \in \mathcal{P}_R : \ell \in \mathcal{L}(p)\}|$ be the number of PRs touching locus ℓ. We assign ℓ the repository-local inverse-frequency weight

$$w_R(\ell) = \log\left(\frac{N_R + 1}{\text{df}_R(\ell) + 1}\right) + 1. \tag{3}$$

The coupling score between PRs p_i and p_j is then defined as

$$s_R(p_i, p_j) = \frac{\sum_{\ell \in \mathcal{L}(p_i) \cap \mathcal{L}(p_j)} w_R(\ell)}{\sum_{\ell \in \mathcal{L}(p_i) \cup \mathcal{L}(p_j)} w_R(\ell)}. \tag{4}$$

This weighted overlap suppresses repository-wide hubs and emphasizes specific, infrequently modified loci that provide stronger evidence of continuous engineering work.

Edge filtering and bounded DAG assembly. Candidate relations are constructed only within the same repository:

$$E_R^{\text{cand}} = \{(p_i, p_j) \mid t_i < t_j,\ t_j - t_i \le \Delta,\ s_R(p_i, p_j) \ge \theta\}, \tag{5}$$

where Δ is the search radius and θ the minimum coupling threshold. Merge time only constrains search and orients edges; temporal proximity does not indicate coupling.

We remove relations induced by repeated or reversed changes, including reverts, backports, release replays, and near-duplicate patches. The remaining edges define repository-level connected components. Components of at most K nodes are retained; larger ones are reduced to high-scoring connected subgraphs. Preserving all surviving edges among selected nodes yields bounded dependency DAGs with linear, fork, join, and mixed structures. Merge-time linearization is used only as a baseline. Appendix details further filtering and graph assembly.

3.3 Executable Task Instantiation

A dependency DAG defines the scope of a candidate task but does not by itself provide an executable benchmark instance. WEAVE therefore constructs a unified repository state, task specification, and test oracle for each retained graph.

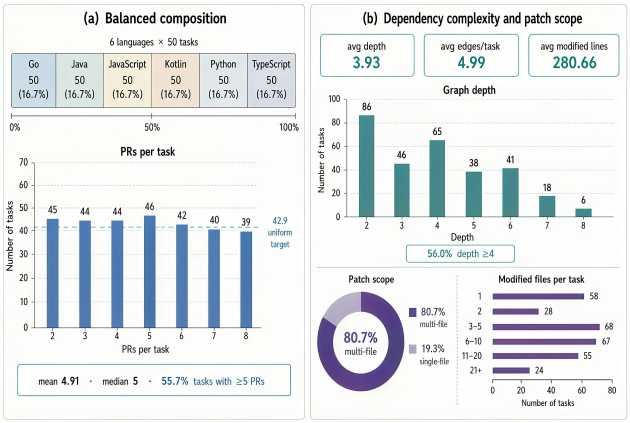

Figure 3: Composition and structural characteristics of WEAVE-BENCH. **(a)** Language balance and PR scale. **(b)** Dependency-DAG complexity and patch scope.

Executable task instantiation. For each candidate graph G, we start from a repository state S_{gold} containing all constituent PRs and revert their changes to construct the counterfactual base state S_{base}, while preserving changes outside G. We then consolidate the PR objectives into a single task specification q, removing PR identifiers, commit metadata, and historical implementation details. The agent solves the task in a single repository state and returns one patch, without reproducing the original PR sequence.

The combined executable oracle must satisfy

$$\mathcal{T}_{\text{F2P}}(S_{\text{base}}) = \text{fail}, \qquad \mathcal{T}_{\text{F2P}}(S_{\text{gold}}) = \text{pass}, \tag{6}$$

while preserving the required $\mathcal{T}_{\text{P2P}}$ behavior. Candidates are discarded when repository reconstruction, gold-patch application, or executable validation fails.

Node attribution and corpus construction. To support graph-conditioned credit, each valid F2P test is assigned to a PR node when high-confidence attribution is possible:

$$A : \mathcal{T}_{\text{F2P}} \to V \cup \{\text{ambiguous}, \text{unattributed}\}. \tag{7}$$

Attributed tests enable node-level scoring, while all valid tests determine full-task correctness. We retain only executable instances with adequate attribution coverage and no recoverable target commits or repository history. Appendix details reconstruction and attribution.

Validated instances are partitioned into WEAVE-TRAIN and the held-out WEAVE-BENCH. We enforce temporal separation and eliminate overlapping target PRs, duplicate or near-duplicate patches, and other direct solution leakage across splits. As shown in Figure 3, WEAVE-BENCH balances language coverage while spanning diverse task scales, dependency structures, and patch scopes, including non-linear PR dependencies and coordinated multi-file changes.

4 Graph Credit

Binary resolution provides a definitive measure of full-task correctness but cannot distinguish progress among unsolved tasks. Flat test scoring offers finer granularity but ignores

dependencies: a downstream node may pass its local tests while an upstream prerequisite remains unresolved. We call the resulting structurally unsupported score *phantom credit*. Graph Credit addresses this limitation by conditioning node-level completion on the task DAG.

4.1 Dependency-Conditioned Scoring

For each evaluable node v, let $\mathcal{T}_v$ denote its confidently attributed F2P tests. Its local completion is

$$\ell_v = \frac{1}{|\mathcal{T}_v|} \sum_{t \in \mathcal{T}_v} \mathbb{I}[t \text{ passes}]. \tag{8}$$

Let $V_A = \{v \in V : |\mathcal{T}_v| > 0\}$ and let $G_A = (V_A, E_A)$ be the induced scoring graph. Ambiguous and unattributed tests remain part of the full-task oracle but are excluded from node-level scoring.

To enforce dependency consistency, we recursively define the effective credit of node v as

$$c_v = \begin{cases} \ell_v, & \mathrm{pa}_{G_A}(v) = \varnothing, \\ \min\left\{ \ell_v, \min\limits_{u \in \mathrm{pa}_{G_A}(v)} c_u \right\}, & \text{otherwise.} \end{cases} \tag{9}$$

Because G_A is acyclic, this recursion is well defined. Graph Credit is the mean effective credit over evaluable nodes:

$$\mathrm{GraphCredit}(G_A) = \frac{1}{|V_A|} \sum_{v \in V_A} c_v. \tag{10}$$

The recursive minimum caps each node's credit by its own completion and that of its upstream prerequisites. Thus, downstream test passes cannot receive more credit than the least-complete prerequisite on their dependency paths. Graph Credit relies solely on executable outcomes and evaluator-side graph metadata; it neither imposes the historical PR order nor uses patch similarity or learned judges.

4.2 Evaluation and Agent Optimization

Evaluation Protocol We report Graph Credit alongside two complementary metrics. Resolved measures full-task correctness and remains the primary benchmark metric, while Flat-F2P measures the fraction of functional tests passed. Graph Credit instead captures dependency-consistent partial progress, providing a structural diagnostic for unsolved tasks. Alternative graph operators, linearized scoring, and randomized-graph controls are examined in RQ2.

Reinforcement-Learning Reward Graph Credit also provides a structural signal for reinforcement learning. We combine flat functional progress, dependency-conditioned credit, full-task resolution, and regression preservation:

$$\begin{aligned} R_{\text{WEAVE}} = \lambda_{\text{f2p}}\ \mathrm{FlatF2P} + \lambda_{\text{graph}}\ \mathrm{GraphCredit}(G_A) \\ + \lambda_{\text{solve}}\ \mathrm{Resolved} - \lambda_{\text{reg}}\ \mathrm{Reg}_{\text{P2P}}\,. \end{aligned} \tag{11}$$

Here, $\mathrm{Reg}_{\text{P2P}}$ penalizes regressions on previously passing tests, and all coefficients are nonnegative. Graph Credit is therefore a structural component of the reward rather than a standalone objective. Since every term is computed from executable test outcomes on the generated patch, R_{WEAVE} provides dense, verifiable feedback. Training details and reward comparisons are presented in RQ4.

5 Experiments

We evaluate WEAVE along the complete pipeline from task construction to agent evaluation and improvement. Specifically, we consider the following research questions:

- **RQ1** Does WEAVE construct coherent and executable dependency DAGs?
- **RQ2** How do current software engineering agents perform on WEAVE-Bench, and does Graph Credit provide diagnostic information beyond flat evaluation?
- **RQ3** Do WEAVE tasks provide more effective and token-efficient supervision for trajectory SFT?
- **RQ4** Does graph-conditioned credit provide additional gains when used as a reinforcement-learning reward?

Common settings. We use WEAVE-Train for supervised fine-tuning and reinforcement learning, a held-out calibration subset for threshold selection and hyperparameter tuning, and WEAVE-Bench exclusively for final evaluation. We enforce temporal separation across splits and remove overlapping target PRs, identical or near-duplicate patches, and recoverable target-commit metadata. Unless otherwise specified, agent runs use the same tool interface, interaction budget, context limit, and execution timeout. We report task-level macro averages and obtain confidence intervals through task-level bootstrap resampling. Complete implementation details and hyperparameters are provided in the appendix.

5.1 RQ1: Construction Validity

To validate WEAVE's construction pipeline, we evaluate whether typed-locus coupling recovers true dependency edges more reliably than single-signal alternatives. The evaluation covers 30 repositories containing 1,198 validated atomic PRs. To establish ground truth, three expert annotators independently assess candidate PR pairs and label each as *dependency*, *same-task but independent*, *loosely related*, or *unrelated*; a senior annotator adjudicates disagreements, yielding 219 gold dependency edges (Fleiss' $\kappa = 0.791$). Each method produces a scored ranking of candidate edges; we report macro-averaged AP over repositories, Coverage—the fraction of gold edges recalled at full emission—and pooled Precision at rank 50 (P@50). All methods feed into the identical downstream executable construction pipeline (§3.3).

We compare WEAVE against four single-signal baselines: Temporal Window (merge-time adjacency within $\Delta = 30$ days), Traceability Links (explicit cross-PR references such as `closes` and `depends on`), File Overlap (weighted Jaccard over changed-file paths), and Symbol Overlap (shared function and class declarations via lightweight static analysis). Score-based methods retain at most $3N_R$ top-scoring edges per repository (where N_R is the PR count); threshold sensitivity is reported in Appendix.

Results. WEAVE achieves the highest AP (37.7%), outperforming the next-best baseline (File Overlap, 26.6%) by 11.1 points (paired ΔAP 95% CI: [+2.7, +20.1]), with significance against all four baselines after Holm–Bonferroni correction ($p \leq 0.009$; Table 3). Crucially, WEAVE achieves the highest recall: 84.5% Coverage exceeds all baselines,

Table 2: Agent performance on WEAVE-BENCH (8 models × 2 scaffolds). **Resolved**: all F2P pass, no P2P regressions. **Flat-F2P**: attributed F2P pass rate. **Graph Credit**: dependency-conditioned partial score (§4.1). **OverCredit**: phantom credit = Flat-F2P − Graph Credit. **Turns/Tok**: per-task averages. Models grouped by scaffold and ordered by model family (MiniMax, DeepSeek, Kimi, GLM, Claude), with earlier versions listed first within each family; best per-column **bolded** within each scaffold.

Scaffold	Model	Resolved↑	Flat-F2P↑	GraphCr.↑	OverCr.↓	Turns↓	Tok (M)↓
OpenCode	MiniMax-2.7	12.3	24.2	13.4	10.8	90.9	5.98
	DeepSeek-V4-Flash	21.0	21.8	14.6	**7.2**	78.7	6.03
	Kimi-K2.5	14.0	12.4	5.1	7.3	88.6	6.65
	Kimi-K2.6	29.3	30.1	21.4	8.7	96.3	8.16
	GLM-5.1	28.7	24.1	14.5	9.6	114.8	7.88
	GLM-5.2	34.0	32.5	23.8	8.7	145.1	12.17
	Claude Opus 4.6	30.3	24.9	15.8	9.1	90.8	4.34
	Claude Opus 4.7	**41.0**	**45.6**	**37.2**	8.4	**68.6**	**4.40**
Cline	MiniMax-2.7	9.7	19.1	10.5	8.6	114.7	5.71
	DeepSeek-V4-Flash	34.3	58.1	41.3	16.8	143.9	7.80
	Kimi-K2.5	18.7	27.5	16.8	10.7	103.7	5.30
	Kimi-K2.6	37.3	**69.1**	57.2	11.9	108.7	5.28
	GLM-5.1	21.0	18.8	11.6	**7.2**	191.3	10.42
	GLM-5.2	43.0	59.5	44.1	15.4	138.2	7.52
	Claude Opus 4.6	36.3	61.3	47.6	13.7	**44.5**	**1.65**
	Claude Opus 4.7	**48.7**	64.4	**50.8**	13.6	46.8	2.11

Table 3: Dependency recovery quality (30 repos, 219 gold edges, Fleiss' $\kappa = 0.791$). AP: per-repo macro-average. Coverage: gold-edge recall at emission. P@50: pooled top-50 precision. Paired bootstrap + Holm correction; ** $p<0.01$.

Method	AP↑	Coverage↑	P@50↑
Temporal Window	5.1	73.1	2.0
Traceability Links	17.2	12.8	28.0
File Overlap	26.6	73.5	10.0
Symbol Overlap	24.0	69.0	30.0
WEAVE (Ours)	**37.7****	**84.5**	**38.0**

including those designed for breadth (Temporal Window 73.1%, File Overlap 73.5%). No baseline matches WEAVE on both axes—typed-locus coupling is the only method that escapes the precision–recall trade-off characterizing single-signal alternatives. The advantage stems from two mechanisms: repository-calibrated IDF weighting suppresses hub-file false edges (P@50: File Overlap 10% vs. WEAVE 38%), while multi-channel evidence recovers dependencies invisible to any single signal—e.g., Traceability Links covers only the 12.8% of true edges carrying explicit cross-references. Additional P@K curves, per-repository AP distributions, and locus-type ablations are provided in Appendix.

5.2 RQ2: Benchmark Performance

We evaluate eight frontier models from five providers—Claude Opus 4.7, Claude Opus 4.6, GLM-5.2, GLM-5.1, DeepSeek-V4-Flash, Kimi-K2.6, Kimi-K2.5, and MiniMax-2.7—under two agent scaffolds (OpenCode, Cline) on the full WEAVE-BENCH suite (300 tasks, 16 configurations). All configurations share the same task specification, repository state, and tool permissions. Table 2 reports three scoring regimes of increasing structural awareness—RESOLVED (binary full-task correctness), Flat-F2P (fraction of attributed F2P tests passed (Thai et al. 2025)), and Graph Credit (§4.1)—together with OverCredit = Flat-F2P − Graph Credit and per-task efficiency (interaction turns, total tokens).

Results. The strongest configuration (Claude Opus 4.7 + Cline, 48.7% Resolved) outperforms the weakest (MiniMax-2.7 + Cline, 9.7%) by $>5\times$, confirming that dependency-structured multi-PR tasks are qualitatively harder than isolated single-PR benchmarks where frontier models exceed 50% Resolved (Jimenez et al. 2023). Across all 16 configurations, Flat-F2P exceeds Graph Credit by **10.5 pp** on average (range 7.2–16.8 pp), exposing substantial phantom credit. The gap is scaffold-sensitive: Cline configurations exhibit larger OverCredit (mean 12.2 pp) than OpenCode (mean 8.6 pp), indicating that Cline's more exploratory strategy produces more downstream test passes that do not survive dependency conditioning. Notably, OverCredit does not correlate monotonically with Flat-F2P: Kimi-K2.6 + Cline achieves the highest Flat-F2P (69.1%) yet exhibits only moderate OverCredit (11.9 pp), while DeepSeek-V4-Flash + Cline reaches 58.1% Flat-F2P with the largest OverCredit (16.8 pp)—suggesting that the former resolves upstream dependencies more faithfully despite extensive downstream progress. Replacing the true DAG with node-permuted random graphs reduces the gap to < 2 pp (Appendix), verifying that the signal arises from real dependency structure. The same graph that enables this diagnostic next improves train-

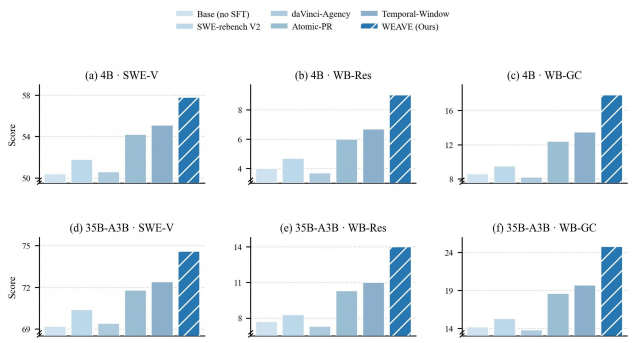

Figure 4: SFT comparison across different data construction strategies.

ing (§5.3).

5.3 RQ3: SFT Utility and Token Efficiency

We examine whether dependency-structured tasks provide more effective supervision for trajectory-level SFT. We fine-tune two open-weight backbones, Qwen3.5-4B and Qwen3.5-35B-A3B, using the Cline scaffold for both trajectory generation and evaluation.

Data sources. We compare three variants constructed from the same repository pool: **Atomic-PR**, which treats each PR independently; **Temporal-Window**, which groups PRs merged within a 30-day window; and **WEAVE**, which aggregates PRs according to typed-locus dependencies. These variants share the same trajectory generator, quality filter, and training configuration, differing only in task construction. We additionally include SWE-rebench V2 (Badertdinov et al. 2025), using 5K sampled trajectories, and daVinci-Agency (Jiang et al. 2026), using its full corpus, as external references. Full training details are provided in the Appendix.

Results. Under matched training conditions, WEAVE consistently outperforms both controlled alternatives across model scales and benchmarks (Figure 4). On WEAVE-Bench, WEAVE improves Resolved over Atomic-PR by 3.0/3.7 pp at 4B/35B, with larger Graph Credit gains (5.4/6.2 pp), indicating stronger dependency-consistent progress even without full resolution. Improvements transfer to single-PR benchmarks: SWE-bench Verified improves by 3.6/2.8 pp over Atomic-PR, with consistent gains on SWE-bench Pro and Multilingual. Temporal-Window provides only modest improvements, confirming that temporal proximity alone introduces weaker supervision than verified dependency structure.

External corpora yield substantially smaller gains: SWE-rebench V2 improves SWE-bench Verified by only 1.2–1.4 pp over base models, while daVinci-Agency shows no consistent improvement—suggesting that training utility depends on how changes are organized into coherent tasks, not merely data volume. Token-scaling analysis (Appendix) further shows that WEAVE maintains stronger performance across training budgets, indicating more information-dense supervision.

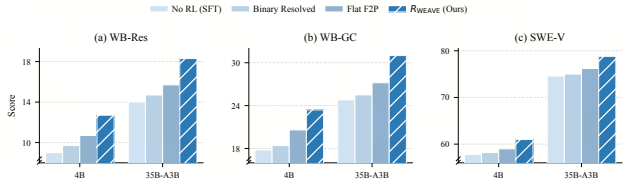

Figure 5: RL reward ablation on two model scales. All Runs start from the WEAVE-SFT checkpoint with matched rollout budgets.

5.4 RQ4: Graph-Conditioned Reinforcement Learning

We assess whether Graph Credit provides optimization value beyond supervised training. Starting from the WEAVE-SFT checkpoint (Figure 4), we apply GRPO (Shao et al. 2024) on the same WEAVE-Train tasks with identical rollout budgets (8 rollouts per task), varying only the reward formulation: (1) no further RL (SFT-only baseline); (2) binary Resolved reward; (3) flat F2P-based reward (r = FlatF2P); and (4) R_{WEAVE} (§4.2), which places Graph Credit as the dominant component ($\lambda_{\text{graph}} = 0.5$, $\lambda_{\text{solve}} = 0.25$, $\lambda_{\text{f2p}} = 0.15$, $\lambda_{\text{reg}} = 0.1$), with a three-phase DAG-complexity curriculum ($\sim$1.0 pp over flat scheduling; Appendix).

Results. R_{WEAVE} consistently outperforms all alternatives (Figure 5). On Qwen3.5-4B, it improves Resolved by 3.7 pp over SFT and 2.0 pp over flat F2P; on Qwen3.5-35B-A3B, gains are 4.3 and 2.6 pp respectively, showing that graph-conditioned reward scales with model capacity. Graph Credit separation is larger (+5.6/+6.2 pp over SFT at 4B/35B), indicating disproportionate improvement in structured partial progress. The gains transfer to SWE-bench Verified (+3.2/+4.2 pp over SFT), confirming out-of-distribution generalization.

Binary Resolved reward yields marginal gains ($\leq$0.7 pp), consistent with sparse-signal limitations on long-horizon tasks (Wei et al. 2025). Flat F2P improves over binary but under-performs R_{WEAVE} on all metrics, confirming that dependency conditioning eliminates phantom-credit shortcuts. All settings maintain P2P regression below 2.3%; component ablations in Appendix.

6 Conclusion

This paper addresses a fundamental limitation of current software engineering agent evaluation: the lack of a task abstraction that captures causal dependencies among multiple repository changes. We introduce WEAVE, which recovers dependency DAGs from verified PR histories through repository-calibrated typed-locus coupling. The same graph representation unifies task construction, dependency-conditioned partial credit (Graph Credit), and structured reward design, enabling consistent evaluation and optimization of long-horizon code agents. Experiments across construction quality, agent benchmarking, supervised training, and reinforcement learning demonstrate the effectiveness of WEAVE as both a challenging evaluation workbench and an effective

substrate for improving code agents.

References

Badertdinov, I.; Golubev, A.; Nekrashevich, M.; Shevtsov, A.; Karasik, S.; Andriushchenko, A.; Trofimova, M.; Litvintseva, D.; and Yangel, B. 2025. SWE-rebench: An Automated Pipeline for Task Collection and Decontaminated Evaluation of Software Engineering Agents. arXiv:2602.23866.

Chen, D.; Lin, S.; Zeng, M.; Zan, D.; Wang, J.-G.; Cheshkov, A.; Sun, J.; Yu, H.; Dong, G.; Aliev, A.; et al. 2024. CodeR: Issue Resolving with Multi-Agent and Task Graphs. arXiv:2406.01304.

Chen, J.; Xu, X.; Wei, H.; Chen, C.; and Zhao, B. 2026. SWE-CI: Evaluating Agent Capabilities in Maintaining Codebases via Continuous Integration. arXiv:2603.03823.

Dai, N.; Wu, Z.; Zheng, R.; Wei, Z.; Shi, W.; Jin, X.; Liu, G.; Dun, C.; Huang, L.; and Yan, L. 2024. Process Supervision-Guided Policy Optimization for Code Generation. arXiv:2410.17621.

Deng, G.; Chen, Z.; Yu, Z.; Fan, H.; Liu, Y.; Yang, Y.; Parikh, D.; Kannan, R.; Cong, L.; Wang, M.; et al. 2026. SWE-Milestone: Evaluating AI Agents on Continuous Software Evolution. In *International Conference on Machine Learning (ICML)*. Formerly EvoClaw.

Deng, X.; Da, J.; Pan, E.; He, Y.; Ide, C.; Garg, K.; Lauffer, N.; Park, A.; Pasari, N.; Rane, C.; et al. 2025. SWE-Bench Pro: Can AI Agents Solve Long-Horizon Software Engineering Tasks? arXiv:2509.16941.

He, X.; Liu, Q.; Du, M.; Yan, L.; Fan, Z.; Huang, Y.; Yuan, Z.; and Ma, Z. 2025. SWE-Perf: Can Language Models Optimize Code Performance on Real-World Repositories? arXiv:2507.12415.

Jain, N.; Singh, J.; Shetty, M.; Zheng, L.; Sen, K.; and Stoica, I. 2025. R2E-Gym: Procedural Environments and Hybrid Verifiers for Scaling Open-Weights SWE Agents. arXiv:2504.07164.

Jiang, M.; Fu, D.; Shi, J.; Zeng, J.; Si, W.; Li, K.; Li, X.; Xiao, Y.; Li, W.; Wang, D.; et al. 2026. daVinci-Agency: Unlocking Long-Horizon Agency Data-Efficiently. arXiv:2602.02619.

Jimenez, C. E.; Yang, J.; Wettig, A.; Yao, S.; Pei, K.; Press, O.; and Narasimhan, K. 2023. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? In *International Conference on Learning Representations*.

Jin, Q.; Tung, L.; Li, K.; Shi, Q.; She, Y.; Jia, H.; Zhao, H.; Xia, K.; Du, Z.; Zhang, Y.; et al. 2026. ChainSWE: Benchmarking Coding Agents on Multi-Bug Software Maintenance. arXiv:2607.02606.

Lam, M. H.; Wang, C.; Liu, H.; Xiao, J.; Li, H.-s.; Huang, J.-t.; Zhuo, T. Y.; and Lyu, M. R. 2026. SWE-Chain: Benchmarking Coding Agents on Chained Release-Level Package Upgrades. arXiv:2605.14415.

Miserendino, S.; Wang, M.; Patwardhan, T.; and Heidecke, J. 2025. SWE-Lancer: Can Frontier LLMs Earn $1 Million from Real-World Freelance Software Engineering? OpenAI blog; arXiv preprint.

Pan, J.; Wang, X.; Neubig, G.; Jaitly, N.; Ji, H.; Suhr, A.; and Zhang, Y. 2024. Training Software Engineering Agents and Verifiers with SWE-Gym. arXiv:2412.21139.

Shao, Z.; Wang, P.; Zhu, Q.; Xu, R.; Song, J.; Zhang, M.; Li, Y.; Wu, Y.; and Guo, D. 2024. DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models. In *arXiv.org*. Introduces GRPO (Group Relative Policy Optimization).

Shastry, K. A.; Senrayan, G.; Satapara, S.; Panda, P.; and Devaguptapu, C. 2026. Beyond Isolated Tasks: A Framework for Evaluating Coding Agents on Sequential Software Evolution. SWE-STEPS, arXiv:2604.03035.

Shum, K.; Hui, B.; Chen, J.; Zhang, L.; x., W.; Yang, J.; Huang, Y.; Lin, J.; and He, J. 2025. SWE-RM: Execution-free Feedback For Software Engineering Agents. arXiv:2512.21919.

Song, Y.; Lothritz, C.; Tang, D.; Bissyandé, T. F.; and Klein, J. 2024. Revisiting code similarity evaluation with abstract syntax tree edit distance. In *Annual Meeting of the Association for Computational Linguistics*, 38–46. Association for Computational Linguistics.

Tao, L.; Kulikov, I.; Saha, S.; Wang, T.; Xu, J.; Li, Y.; Weston, J.; and Yu, P. 2025. Hybrid Reinforcement: When Reward Is Sparse, It's Better to Be Dense. arXiv:2510.07242.

Thai, M.-T.; Le, T.; Mạnh, D. N.; Nhat, H. P.; and Bui, N. D. Q. 2025. SWE-EVO: Benchmarking Coding Agents in Long-Horizon Software Evolution Scenarios. arXiv:2512.18470.

Vogel, M.; Meyer-Eschenbach, F.; Kohler, S.; Grünewald, E.; and Balzer, F. 2026. Codebase-Memory: Tree-Sitter-Based Knowledge Graphs for LLM Code Exploration via MCP. arXiv:2603.27277.

Wang, L.; Liu, Y.; Wu, X.; Hu, X.; Fan, Y.; Yu, K.; Weng, Q.; Xi, W.; and Li, X. 2026. Beyond Binary: Turning Partial Success into Dense Verifiable Rewards for Reinforcement Learning in Code Generation. arXiv:2601.03525.

Wei, Y.; Duchenne, O.; Copet, J.; Carbonneaux, Q.; Zhang, L.; Fried, D.; Synnaeve, G.; Singh, R.; and Wang, S. 2025. SWE-RL: Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution. In *Neural Information Processing Systems*.

Xu, X.; Yang, R.; Shen, H.; Xu, W.; Gao, B.; Wu, R.; Shi, K.; Xie, W.; Chen, X.; Wu, M.; et al. 2026. RoadmapBench: Evaluating Long-Horizon Agentic Software Development Across Version Upgrades. arXiv:2605.15846.

Yang, J.; Jimenez, C. E.; Zhang, A. L.; Lieret, K. A.; Yang, J.; Wu, X.; Press, O.; Muennighoff, N.; Synnaeve, G.; Narasimhan, K. R.; et al. 2024. SWE-bench Multimodal: Do AI Systems Generalize to Visual Software Domains? In *arXiv.org*.

Yang, J.; Lieret, K.; Jimenez, C. E.; Wettig, A.; Khandpur, K.; Zhang, Y.; Hui, B.; Press, O.; Schmidt, L.; and Yang, D. 2025. SWE-smith: Scaling Data for Software Engineering Agents. arXiv:2504.21798.

Yu, B.; Zhu, Y.; He, P.; and Kang, D. 2025. UTBoost: Rigorous Evaluation of Coding Agents on SWE-Bench. In *Annual*

Meeting of the Association for Computational Linguistics, 3762–3774. Association for Computational Linguistics.

Zan, D.; Huang, Z.; Liu, W.; Chen, H.; Zhang, L.; Xin, S.; Chen, L.; Liu, Q.; Zhong, X.; Li, A.; et al. 2025. Multi-SWE-bench: A Multilingual Benchmark for Issue Resolving. In *arXiv.org*.

Zhang, L.; He, S.; Zhang, C.; Kang, Y.; Li, B.; Xie, C.; Wang, J.; Wang, M.; Huang, Y.; Fu, S.; et al. 2025a. SWE-bench Goes Live! In *Neural Information Processing Systems*.

Zhang, Y.; Zhao, X.; Wang, Z.; Yang, C.; Wei, J.; and Wu, T. 2025b. cAST: Enhancing Code Retrieval-Augmented Generation with Structural Chunking via Abstract Syntax Tree. arXiv:2506.15655.

Zhou, Q.; Zhang, J.; Wang, H.; Hao, R.; Wang, J.; Han, M.; Yang, Y.; Wu, S.; Pan, F.; Fan, L.; et al. 2026. FeatureBench: Benchmarking Agentic Coding for Complex Feature Development. In *arXiv.org*.

A Construction Pipeline Details

This appendix details the three stages of the WEAVE construction pipeline introduced in the main paper: typed change loci extraction (§A.1), edge filtering and bounded DAG assembly (§A.2), and rollback construction with test attribution (§A.3).

A.1 Typed Change Loci Extraction

Recovering dependency structure from PR histories requires a representation granular enough to distinguish genuine engineering dependencies from coincidental co-modification. File-level overlap—the most common proxy in prior work—conflates distinct changes within shared files and over-connects PRs that modify frequently touched configuration manifests, barrel imports, or dependency lockfiles. WEAVE addresses this by representing each PR as a set of *typed change loci*: fine-grained program entities extracted from its code diff, each annotated with a structural category that determines its evidential role in coupling computation. We define five locus types, each designed to capture a distinct channel of cross-PR evidence:

1. **AST declarations** ($\mathcal{L}_{\text{ast}}$): Function, method, and class definitions identified via language-specific AST parsers. We extract fully qualified names (e.g., `module.Class.method`) to disambiguate identically named entities across files.
2. **Lightweight symbols** ($\mathcal{L}_{\text{sym}}$): Exported symbols, type aliases, constants, and enum members extracted from declaration-level analysis without full semantic resolution. This captures entities that do not correspond to function bodies but serve as cross-module interfaces.
3. **Configuration and schema entries** ($\mathcal{L}_{\text{cfg}}$): Structured keys in YAML, TOML, JSON, and INI configuration files; database migration column definitions; API route declarations; and protobuf/GraphQL field paths. Each entry is represented as a dot-separated key path (e.g., `server.http.timeout`).
4. **Declarations and signatures** ($\mathcal{L}_{\text{decl}}$): Type signatures, interface definitions, abstract method declarations, and trait/protocol requirements. These capture contract-level changes that downstream implementations depend on.
5. **Dependency keys** ($\mathcal{L}_{\text{dep}}$): Package manager entries (e.g., `requirements.txt`, `package.json`, `go.mod`), import statements, and build-system target declarations. These capture supply-chain relationships between PRs that modify shared external dependencies.

Language support. Table 4 summarizes the extraction capabilities for the six languages in WEAVE-Bench and WEAVE-Train. AST parsing is implemented via Tree-sitter grammars, which provide incremental, error-tolerant parsing for all six languages. Symbol and declaration extraction build on language-specific Tree-sitter queries; for JavaScript, which lacks a native type system, $\mathcal{L}_{\text{decl}}$ relies on heuristic extraction from JSDoc annotations and runtime patterns. Configuration and dependency loci are extracted via format-specific parsers (JSON, YAML, TOML, XML, Gradle) rather than language grammars; the partial marks for Go, Java, and Kotlin reflect lower ecosystem reliance on structured configuration files relative to Python and TypeScript/JavaScript projects.

Table 4: Language support matrix for typed change loci extraction. ✓: full support; ∼: partial (heuristic-based).

Language	$\mathcal{L}_{\text{ast}}$	$\mathcal{L}_{\text{sym}}$	$\mathcal{L}_{\text{cfg}}$	$\mathcal{L}_{\text{decl}}$	$\mathcal{L}_{\text{dep}}$
Python	✓	✓	✓	✓	✓
TypeScript	✓	✓	✓	✓	✓
JavaScript	✓	✓	✓	∼	✓
Go	✓	✓	∼	✓	✓
Java	✓	✓	∼	✓	✓
Kotlin	✓	✓	∼	✓	✓

Exclusion rules. Certain file categories are excluded from loci extraction because they introduce coupling signals unrelated to engineering dependencies. **Test files** (matching patterns such as `test_*`, `*_test.*`, and language-specific test directories) are excluded because test co-modification frequently reflects shared test infrastructure rather than production-code dependency. **Generated code**, vendored dependencies (e.g., `vendor/`, `node_modules/`), and binary files are excluded as they represent mechanical artifacts. **Formatting-only changes**—PRs where all hunks contain only whitespace, comment, or import-order modifications (detected via AST-level diff equivalence)—are excluded to prevent automated formatters from producing spurious coupling edges.

Deduplication and normalization. Loci are canonicalized before comparison: fully qualified names are lowercased for case-insensitive languages (Go exported symbols excepted), trailing version suffixes in dependency keys are stripped, and configuration paths are normalized to dot notation regardless of source format.

Illustrative example. We illustrate typed change loci extraction on a real WEAVE-Bench instance from a TypeScript monorepo. PR_1 adds federation v2 link handling to the `import` package, modifying both source code and the dependency manifest:

```
# PR1 diff (abridged)
--- a/.../import/package.json
+++ b/.../import/package.json
+  "@theguild/federation-composition":
+    "^0.18.3",
--- a/.../import/src/index.ts
+++ b/.../import/src/index.ts
+import { extractLinkImplementations }
+  from '@theguild/federation-composition';
+function importFederatedSchemaLinks(
+    def, defsByName) {
+  const { matchesImplementation,
+          resolveImportName }
+    = extractLinkImplementations({...});
+  const fedUrl = '.../federation';
+  if (matchesImplementation(
+      fedUrl, 'v2.0'))
+    addDef(resolveImportName(
+      fedUrl, '@key'));
```

+}

Extracted typed change loci for PR_1:

- $\mathcal{L}_{\text{ast}}$: `importFederatedSchemaLinks` (new function)
- $\mathcal{L}_{\text{sym}}$: `extractLinkImplementations`, `matchesImplementation`, `resolveImportName`
- $\mathcal{L}_{\text{dep}}$: `@theguild/federation-composition @^0.18.3`

A subsequent PR_2 extends federation directive support in the `merge` package. Its diff adds the same dependency (`@theguild/federation-composition @^0.16.0`) to a different `package.json` and calls the identical symbol sequence (`extractLinkImplementations`, `matchesImplementation`, `resolveImportName`) in `merge-typedefs.ts`. The shared $\mathcal{L}_{\text{dep}}$ locus (same dependency package) and $\mathcal{L}_{\text{sym}}$ loci (same function symbols) produce a coupling score well above the emission threshold, despite the two PRs modifying different packages with no file-level overlap. A file-overlap baseline would miss this edge entirely, since the only shared file is the monorepo-wide `yarn.lock`—a hub file whose repository-calibrated weight is near zero.

A.2 Edge Filtering and DAG Assembly

Mechanical-change filtering. Not all PR pairs with high coupling scores reflect genuine engineering dependencies. Certain repository maintenance patterns produce mechanically related patches that inflate coupling metrics without meaningful task structure. We identify and remove four categories before DAG assembly:

- **Reverts**: PR p_j is identified as a revert of p_i via two complementary signals: (i) title-level patterns (`Revert "..."`) and (ii) content-level detection where the diff of p_j is the near-exact ($\geq$95% line overlap) inverse of p_i's diff.
- **Backports and cherry-picks**: PRs explicitly labeled as cherry-picks or backports (via commit message markers or branch metadata) to release branches represent the same logical change applied across code states.
- **Release-branch references**: PRs whose title or body references release, stable, maintenance, or hotfix branches (detected via text patterns such as `release/`, `stable/`, `maint/`, `hotfix/`) are removed, as they typically represent integration merges rather than incremental development.
- **Near-duplicate patches**: PR pairs with high file-set overlap (Jaccard ≥ 0.8) and high containment (≥ 0.9) indicate template-generated or copy-paste changes rather than incremental development.

Across our construction corpus, backports and cherry-picks dominate the filtered set (>95%), as multi-branch workflows frequently produce mechanically identical patches across release channels. Reverts, release-branch references, and near-duplicate patches each contribute a small fraction of the remainder.

Component partitioning. After edge filtering, the remaining edges within each repository induce connected components. We process these components as follows:

- Components with $|V| \leq K$ (where $K = 8$) are retained as complete candidate DAGs.
- Larger components are reduced via greedy subgraph extraction: we iteratively select the highest-weight connected subgraph of size $\leq K$ by starting from the densest node (highest sum of incident edge weights) and expanding via best-first search. Extracted subgraphs are removed from the component, and the process repeats until no subgraph of size ≥ 2 with mean edge weight $\geq \theta$ remains.

Node budget. The budget K controls the maximum number of PR nodes in a single task. We set K=8 based on two considerations.

First, *rollback feasibility*: constructing executable instances requires reverting all constituent PRs from the repository's integrated state (§A.3). In our corpus, the median merge-time span (first PR to last PR) is 51 days for K=4 tasks and 78 days for K=8 tasks. Although larger K does not monotonically increase span—K=8 chains are often dense feature or migration bursts with short inter-PR intervals ($\sim$10–15 days), while smaller chains can span longer periods with sparser commits—the absolute span at K=8 (median $\sim$2.5 months, range 17–332 days) means the repository accumulates substantial intervening changes from outside the task scope, raising the probability of non-trivial merge conflicts during rollback construction. Indeed, approximately 35% of candidate DAGs are discarded during construction (§A.3).

Second, *ecological validity*: coherent engineering objectives in open-source repositories—feature additions, systematic refactors, dependency upgrades—rarely produce more than eight tightly coupled PRs before coupling evidence between peripheral changes falls below the detection threshold. In our candidate pool, connected components exceeding 8 nodes predominantly arise from hub-file co-modification (shared configuration, CI pipelines, or utility modules) rather than focused development objectives. This observation is consistent with prior work on software change coupling, where meaningful change sets typically span 2–6 related commits.

The resulting benchmark reflects this design: WEAVEBENCH tasks range from 2 to 8 PR nodes (mean 4.91, median 5), with a broad distribution across this range, spanning the full complexity spectrum from simple linear chains to deep branching dependency structures.

Algorithm 1 summarizes the complete pipeline from typed change loci extraction through bounded DAG assembly.

DAG orientation. All edges are oriented by merge timestamp: if p_i was merged before p_j and the coupling score exceeds θ, the edge is directed $p_i \to p_j$. This temporal orientation reflects the incremental nature of software development without requiring explicit dependency annotations from developers. Cycles (which would arise from conflicting timestamps in rebased histories) are broken by removing the lowest-weight back-edge in each strongly connected component.

Algorithm 1: Dependency DAG Assembly

Require: Repository R, merged PRs $\{p_1, \ldots, p_n\}$ (by timestamp), threshold θ, budget K, search radius Δ
Ensure: Set of candidate DAGs $\mathcal{G}$

// Phase 1: Typed Change Loci Extraction
for each $p_i \in \{p_1, \ldots, p_n\}$ **do**
  $L(p_i) \leftarrow \mathcal{L}_{\text{ast}}(p_i) \cup \mathcal{L}_{\text{sym}}(p_i) \cup \mathcal{L}_{\text{cfg}}(p_i) \cup \mathcal{L}_{\text{decl}}(p_i) \cup \mathcal{L}_{\text{dep}}(p_i)$
  Apply exclusion rules (test files, generated code, vendored)
end for
// Phase 2: Candidate Edge Generation
$E^{\text{cand}} \leftarrow \varnothing$
for each pair (p_i, p_j) with $|t(p_j) - t(p_i)| \leq \Delta$ **do**
  Compute $s_R(p_i, p_j) = \frac{\sum_{\ell \in L(p_i) \cap L(p_j)} w_R(\ell)}{\sum_{\ell \in L(p_i) \cup L(p_j)} w_R(\ell)}$
  if $s_R(p_i, p_j) \geq \theta$ **then**
    $E^{\text{cand}} \leftarrow E^{\text{cand}} \cup \{(p_i, p_j)\}$
  end if
end for
// Phase 3: Mechanical-Change Filtering
Remove reverts, backports, release-branch refs, near-duplicates from E^{cand}
// Phase 4: Component Partitioning
Compute connected components $\mathcal{C}$ from (V, E^{cand})
$\mathcal{G} \leftarrow \varnothing$
for each component $C_i \in \mathcal{C}$ **do**
  if $|C_i| \leq K$ **then**
    Orient edges by timestamp; break cycles
    $\mathcal{G} \leftarrow \mathcal{G} \cup \{\text{DAG}(C_i)\}$
  else
    while $\exists$ subgraph $S \subseteq C_i$ with $|S| \geq 2$, mean weight $\geq \theta$ **do**
      Extract densest-first subgraph S with $|S| \leq K$
      Orient edges; break cycles; $\mathcal{G} \leftarrow \mathcal{G} \cup \{\text{DAG}(S)\}$
      $C_i \leftarrow C_i \setminus S$
    end while
  end if
end for
// Phase 5: Degenerate Filtering
Remove star graphs, trivially serial chains, disconnected DAGs from $\mathcal{G}$
return $\mathcal{G}$

Degenerate structure filtering. A subset of assembled DAGs are structurally uninformative for evaluation despite passing coupling and size thresholds. We apply three topological filters:

- **Star graphs**, where a single hub node accounts for >80% of all incident edges with no inter-leaf connections, typically indicate a utility-file artifact rather than a genuine multi-party dependency chain.
- **Trivially serial chains** with uniformly low edge weights (below the 25th percentile of the repository's weight distribution) and no branching structure, which provide limited evaluation signal beyond what single-PR benchmarks already capture.
- **Post-filtering disconnection**: DAGs that become disconnected after mechanical-change removal are repartitioned into their connected components and re-evaluated independently.

A.3 Rollback Construction and Test Attribution

Base state construction. Given a candidate DAG $G = (V, E)$ and the repository state S_{gold} where all constituent PRs have been merged, we construct the base state S_{base} by removing the changes introduced by all nodes in V. The primary mechanism is *reverse-topological revert*: we compute a topological ordering σ of V and process nodes in reverse order σ^{-1} (leaves first), applying `git revert -no-commit` for each node's merge commit. PRs external to G that were merged between the earliest and latest constituent PRs are preserved in S_{base}, ensuring the base state reflects realistic repository evolution outside the task scope.

When git revert produces non-trivial merge conflicts—typically caused by overlapping modifications from intervening commits or cross-cutting refactors—we employ an *agent-assisted carve* as a fallback. In this path, an LLM-based agent receives the PR metadata (title, description, and code diff) and removes the target changes from the current repository state while preserving surrounding code. The carve operates semantically rather than mechanically: the agent identifies which hunks correspond to the target PR's contributions and removes them while maintaining syntactic and build validity.

After either path, the resulting S_{base} is validated: it must compile successfully and all P2P (pass-to-pass) tests must pass. Candidates where neither revert nor carve produces a valid base state are discarded.

Construction success rates. Approximately 35% of candidate DAGs are discarded during construction because neither revert nor carve produces a valid, test-passing base state. Among the successfully constructed instances (which form the final WEAVE-BENCH and WEAVE-TRAIN corpora), 73% are built via clean git revert and 27% require agent-assisted carve. The higher carve rate for Java and Kotlin tasks reflects stricter type-system constraints that propagate through revert sequences, while Python and JavaScript tasks exhibit higher clean-revert rates due to fewer cross-module compilation dependencies.

Test attribution. For graph-conditioned scoring, each F2P test must be associated with one or more PR nodes in the task DAG. The primary attribution mechanism is *isolated revert*: for each node $v \in V$, we construct $S_v = S_{\text{gold}} \setminus \{v\}$ by reverting only v's changes, then execute all F2P tests against S_v. If exactly one node v causes test t to fail, t is *attributed* to v; if multiple nodes independently cause failure, t is marked *ambiguous*; if t fails only under joint reversion, it is *unattributed*. In cases where isolated revert is infeasible for individual nodes (e.g., due to compilation dependencies between constituent PRs), an agent-based attribution step supplements the automated procedure by analyzing test semantics against the PR change descriptions.

Algorithm 2 formalizes the rollback construction and test attribution procedures.

Algorithm 2: Rollback Construction and Test Attribution

Require: DAG $G = (V, E)$, gold repository state S_{gold}
Ensure: Base state S_{base}, test attribution map $A : T_{\text{F2P}} \to V \cup \{\text{amb}, \text{unattr}\}$

// Phase 1: Rollback Construction
Compute topological order σ of V
$S \leftarrow S_{\text{gold}}$
for each v in σ^{-1} (leaves first) **do**
  Apply `git revert -no-commit` of v's merge commit to S
  if non-trivial conflict **then**
    Attempt agent-assisted carve for v
    if carve fails **then**
      discard this DAG instance;
      return $\varnothing$
    end if
  end if
end for
$S_{\text{base}} \leftarrow S$
Validate: compile S_{base}; run P2P tests
// Phase 2: F2P Test Discovery
$T_{\text{F2P}} \leftarrow \{t : t \text{ fails in } S_{\text{base}} \wedge t \text{ passes in } S_{\text{gold}}\}$
// Phase 3: Test Attribution
for each $t \in T_{\text{F2P}}$ **do**
  $\text{candidates}(t) \leftarrow \varnothing$
  for each $v \in V$ **do**
    $S_v \leftarrow S_{\text{gold}} \setminus \{v\}$ {revert only v}
    if t fails in S_v **then**
      $\text{candidates}(t) \leftarrow \text{candidates}(t) \cup \{v\}$
    end if
  end for
  if $|\text{candidates}(t)| = 1$ **then**
    $A(t) \leftarrow$ the unique $v \in \text{candidates}(t)$
  else if $|\text{candidates}(t)| > 1$ **then**
    $A(t) \leftarrow \text{ambiguous}$
  else
    $A(t) \leftarrow \text{unattributed}$
  end if
end for
return S_{base}, A

B Graph Credit Analysis

This appendix supplements the Graph Credit analysis presented in RQ2. We examine how the dependency graph—constructed via typed change loci and repository-calibrated coupling—interacts with different scoring operators and task topologies.

B.1 Alternative Graph Operators

The main paper defines Graph Credit using the recursive minimum (min-bottleneck) operator. We evaluate three alternative graph operators to validate the design choice:

Table 5: Comparison of graph operators. Spearman ρ with binary Resolved (higher = better diagnostic); *residual* OverCredit is the operator score minus the min-bottleneck score, i.e., the phantom credit the operator fails to remove (lower = better). The linearized row therefore recovers the corpus-level OverCredit of 10.5 pp reported in the main paper.

Operator	Spearman $\rho\uparrow$	Resid. OverCredit$\downarrow$
Linearized (no graph)	0.69	10.5 pp
Depth-weighted	0.73	7.1 pp
Mean-parent	0.77	3.2 pp
Min-bottleneck (ours)	**0.82**	**0.0 pp**

Mean-parent operator. Instead of taking the minimum over parent credits, the mean-parent variant averages them:

$$c_v^{\text{mean}} = \begin{cases} \ell_v, & \text{pa}(v) = \varnothing, \\ \min\left(\ell_v,\ \frac{1}{|\text{pa}(v)|}\sum_{u \in \text{pa}(v)} c_u^{\text{mean}}\right), & \text{otherwise.} \end{cases} \tag{12}$$

This operator is less conservative: it allows partial credit to flow through even when one parent is unresolved, provided other parents are complete.

Depth-weighted operator. This variant discounts credit by the depth of the node in the DAG:

$$c_v^{\text{depth}} = \ell_v \cdot \gamma^{d(v)}, \tag{13}$$

where $d(v)$ is the longest-path depth from any root and $\gamma = 0.95$ is a decay factor. Deeper nodes receive inherently less credit regardless of upstream resolution.

Linearized scoring. Nodes are scored independently of the DAG, using only test attribution:

$$c_v^{\text{linear}} = \ell_v. \tag{14}$$

This reduces to Flat-F2P restricted to attributed tests (no dependency conditioning), providing a direct comparison to the dependency-free baseline.

Empirical comparison. Table 5 reports the correlation between each operator's score and actual task resolution, measured across all WEAVE-BENCH tasks evaluated by 8 models $\times$ 2 scaffolds (4,800 agent runs).

The min-bottleneck operator achieves the highest correlation with Resolved and, by construction, eliminates all phantom credit. The mean-parent operator permits residual phantom credit (3.2 pp) because it averages rather than gates upstream deficiencies; the residual is modest only because WEAVE-BENCH is chain-dominated (median anti-chain width 1), and the two operators coincide exactly at single-parent nodes. The depth-weighted operator penalizes deep nodes regardless of actual upstream completion, reducing its diagnostic fidelity. These results justify the conservative min-bottleneck design.

Connections to credit assignment in RL. Graph Credit's recursive conditioning bears structural similarity to hierarchical credit assignment in multi-agent reinforcement learning and process reward models. However, Graph Credit operates on a *known* DAG derived from real code dependencies

rather than a learned latent structure, making it fully deterministic and interpretable. The min-bottleneck operator is analogous to the conjunction of preconditions in classical planning—a task node receives earned credit only when all its causal prerequisites are satisfied.

B.2 Randomized-Graph Controls

To confirm that Graph Credit's diagnostic value arises from the *true* dependency structure rather than arbitrary graph constraints, we evaluate node-permuted random graphs as controls.

Experimental setup. For each WEAVE-BENCH task, we generate 10 random permutations of the node labels while preserving the DAG topology. Graph Credit is recomputed under each permuted graph, and results are averaged across permutations.

Results. Under the true DAG, the mean OverCredit (Flat-F2P − Graph Credit) is 10.5 pp across all agent configurations. Under randomized graphs:

- Mean OverCredit drops to 1.6 pp (<2 pp as reported in the main text).
- The standard deviation across permutations is 0.7 pp, indicating stable behavior.
- For 89% of tasks, the randomized OverCredit is within 2 pp of zero.

This demonstrates that the gap between Flat-F2P and Graph Credit under the true DAG reflects genuine phantom credit arising from dependency violations, not an artifact of imposing arbitrary graph structure.

B.3 Chain-Depth Interpretation

Graph depth requires an explicit counting convention. The benchmark summary in Figure 3 and Table 19 counts nodes on the longest path of the full task DAG G. For scoring and curriculum thresholds, depth counts edges; a path containing k nodes has $k - 1$ edges. A single-node graph therefore has edge depth zero. The scoring graph G_A is induced by evaluable nodes and can have a shorter longest path than the full task graph.

The recursive minimum in Graph Credit caps downstream completion by upstream prerequisites at every edge. Longer chains consequently provide more opportunities for locally passing tests to be capped by an unresolved ancestor. This is a structural property of the scoring rule; it does not by itself establish an empirical relationship between chain depth and OverCredit. The overall OverCredit results are reported in Section 5.2.

C Experimental Setup

This appendix provides implementation details necessary for reproducing the experiments in the main paper: computing environment (§C.1), agent evaluation configuration (§C.2), and sensitivity analysis of the two key construction hyperparameters (§C.3).

Table 6: Scaffold evaluation configurations.

Parameter	Value
Max interaction turns	200
Total task timeout	2 hours
Test suite timeout	30 minutes
Tool interface	Shell + File Edit
Context window	Model-provider default
Temperature	Model-provider default

C.1 Computing Infrastructure

All experiments—including construction pipeline execution, agent evaluation, and model training—are conducted on a private cloud cluster equipped with NVIDIA datacenter GPUs.

Software environment. Agent evaluation runs execute in isolated Docker containers. The primary software stack uses Python 3.12 with PyTorch 2.x. Model training employs a distributed training framework supporting ZeRO-3 parallelism, sequence parallelism for long-context batches, and mixed-precision (BF16) optimization. For agent evaluation (RQ2), all eight evaluated models are accessed through their official provider endpoints, so that measured behavior reflects production inference settings. Trajectory generation for WEAVE-TRAIN (§E.1) instead uses a self-hosted deployment of the open-weight generator, which makes the corpus reproducible without dependence on a hosted endpoint.

Resource allocation. Each agent evaluation run executes in an isolated container with a 2-hour task-level timeout. SFT and RL training use multi-GPU distributed configurations. We report the hyperparameters that determine the optimization trajectory (§E.2, §E.3) rather than device counts or sharding topology: the former fully specify the training run, while the latter are deployment choices that vary with available hardware and are omitted per organizational policy.

C.2 Evaluation Configuration

Table 6 summarizes the shared evaluation parameters for both scaffolds in RQ2.

Both scaffolds share the configuration in Table 6. Notably, no per-turn timeout, file-read limit, or patch-size restriction is imposed—agents may produce arbitrarily large edits within the task-level budget. Context window sizes and sampling temperatures follow the default settings of each model provider's official API; we do not override these parameters to ensure results reflect production-equivalent inference behavior.

Scaffold differences. OpenCode and Cline differ primarily in their system prompts and tool invocation patterns. OpenCode uses a structured XML-based action format with explicit planning steps, while Cline employs a conversational format with implicit planning. Both expose identical underlying capabilities (shell execution, file reading/writing, patch submission). The choice of two scaffolds enables assessment of scaffold-invariant model capabilities versus scaffold-specific behaviors.

Table 7: Sensitivity to coupling threshold θ (dev set, Δ=90 days). **Bold**: selected configuration.

θ	AP↑	Coverage↑	Exec. Rate↑
0.05	31.4	92.0	62.4
0.10	35.6	88.0	71.2
0.15	37.6	84.0	76.8
0.20	37.2	78.0	82.4
0.25	34.0	68.0	80.6
0.30	30.2	58.0	78.8

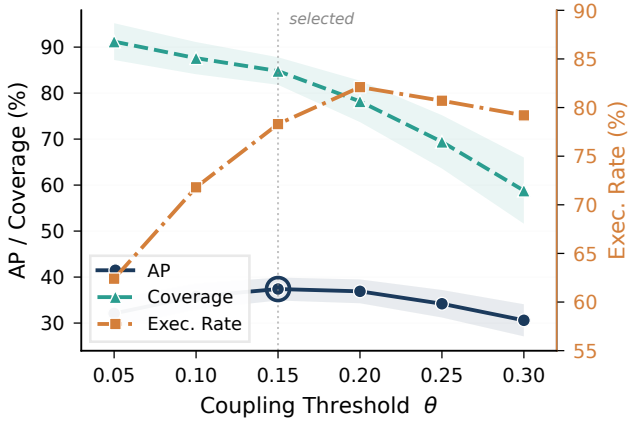

Figure 6: Sensitivity of construction quality to coupling threshold θ. Left axis: AP (solid) and Coverage (dashed); right axis: executable construction rate. The selected $\theta = 0.20$ (vertical dotted line) balances precision and recall while maintaining high executability. Shaded region indicates ± 1 standard deviation across development-set repositories.

C.3 Threshold and Hyperparameter Sensitivity

Coupling threshold θ. Table 7 reports construction quality metrics as a function of the coupling threshold θ on the development set (10 repositories, 50 gold edges).

We select $\theta = 0.20$ as it achieves the best tradeoff between edge precision and executable construction success: AP remains within 0.4 pp of the peak ($\theta = 0.15$) while the executable construction rate reaches its maximum. Lower thresholds introduce spurious edges that increase merge conflicts during rollback construction, while higher thresholds sacrifice recall of genuine dependency relationships.

Search radius Δ. Table 8 reports sensitivity to the temporal search radius.

We select $\Delta = 90$ days. The median inter-PR interval in our corpus is ∼2–3 weeks, well within this window; 90 days serves as a conservative upper bound that captures long-lived feature arcs without omitting genuine dependencies. Coverage saturates by 60 days and gains only 0.6 pp from 60 to 90, confirming that 90 days is a safe ceiling rather than a tight fit.

Table 8: Sensitivity to search radius Δ (dev set, θ=0.20). **Bold**: selected configuration.

Δ (days)	AP↑	Coverage↑	Candidates
7	38.4	48.0	326
14	38.0	60.0	708
30	37.6	72.0	1,482
60	37.4	77.4	3,716
90	37.2	78.0	6,218

Table 9: Precision at various rank cutoffs. Candidate edges from all 30 repositories are merged into one score-ranked list; P@K is the fraction of gold dependency edges among its top K entries.

Method	P@10	P@20	P@50	P@100
Temporal Window	10.0	5.0	2.0	2.0
Traceability Links	60.0	45.0	28.0	17.0
File Overlap	20.0	15.0	10.0	8.0
Symbol Overlap	50.0	40.0	30.0	22.0
WEAVE	**70.0**	**50.0**	**38.0**	**29.0**

D Additional Experimental Results

This section provides extended results for all four research questions. The dependency graph—as the central computational object—connects these results: RQ1 validates graph construction quality, RQ2 evaluates agent performance under graph-conditioned scoring, and RQ3–RQ4 demonstrate that the same graph structure improves both supervised and reinforcement-based training.

D.1 RQ1 Extended Results

Precision at K curves. Table 9 reports precision at multiple rank cutoffs for all methods, providing a more complete view of ranking quality beyond P@50. Following the main-text protocol, candidate edges from all 30 repositories are merged into a single score-ranked list, and P@K is the fraction of gold dependency edges among the top K entries of that pooled ranking. Because the pooled list is scored globally, P@K is quantized to multiples of $100/K$ percentage points.

WEAVE attains the highest precision at every cutoff, and no single baseline is consistently second: Traceability Links is the strongest alternative at $K \leq 20$, where its explicit cross-PR references are almost always genuine, but it exhausts its recoverable edges early (Coverage = 12.8%) and falls below Symbol Overlap by $K = 50$. This crossover is precisely why neither signal suffices alone.

Pooled P@K and per-repository AP (reported in the RQ1 table of the main text) measure different properties and need not agree. AP is computed within each repository and then macro-averaged, so it is insensitive to whether scores are comparable *across* repositories; pooled P@K merges all repositories into one ranking and therefore additionally rewards cross-repository score calibration. File Overlap illustrates the gap: its raw Jaccard scores rank reasonably within a repository (AP 26.6%) but are dominated globally

Table 10: Locus-type ablation. "Single": using only that type. "−Type": removing that type from the full set.

Configuration	AP↑	Coverage↑
Single type		
$\mathcal{L}_{\text{ast}}$ only	28.4	62.1
$\mathcal{L}_{\text{sym}}$ only	22.6	55.8
$\mathcal{L}_{\text{cfg}}$ only	14.2	38.4
$\mathcal{L}_{\text{decl}}$ only	18.8	44.2
$\mathcal{L}_{\text{dep}}$ only	12.4	32.6
Leave-one-out		
$-\mathcal{L}_{\text{ast}}$	31.2	76.8
$-\mathcal{L}_{\text{sym}}$	34.6	80.2
$-\mathcal{L}_{\text{cfg}}$	36.8	83.4
$-\mathcal{L}_{\text{decl}}$	35.4	82.1
$-\mathcal{L}_{\text{dep}}$	37.0	84.0
All types (WEAVE)	**37.7**	**84.5**

by large repositories with many shared files (P@50 10.0%). WEAVE's repository-local inverse-frequency weighting normalizes locus evidence against each repository's own modification statistics, which keeps its scores comparable across repositories and yields close agreement between the two views (AP 37.7% vs. P@50 38.0%).

Per-repository AP distributions. Across the 30 evaluation repositories, WEAVE achieves:

- Median AP: 36.2% (IQR: 28.4–47.1%).
- AP $>$ 50% in 6 repositories.
- AP $<$ 20% in 3 repositories (all with $<$5 gold edges, limiting statistical power).

The next-best method (File Overlap) achieves median AP of 24.8% (IQR: 18.2–35.6%). The macro-average (37.7%) exceeds the median, reflecting a mild right skew: a small number of repositories with dense, well-separated dependency structure achieve high AP, while the bulk of the distribution is concentrated between the quartiles. The interquartile ranges of the two methods are disjoint at the median, indicating that WEAVE's advantage is distributional rather than driven by a few favorable repositories.

Locus-type ablation. Table 10 evaluates the contribution of each locus type by training with individual types and with leave-one-out configurations.

$\mathcal{L}_{\text{ast}}$ is the most informative single type (AP 28.4%), and its removal causes the largest AP drop (−6.5 pp). However, no single type matches the full combination, confirming that the five locus types capture complementary evidence channels. Configuration and dependency loci contribute less individually but improve coverage by capturing relationships invisible to code-structural analysis alone.

Two comparisons against the main-text baselines are worth separating. $\mathcal{L}_{\text{ast}}$ alone (AP 28.4%) already exceeds the strongest baseline, File Overlap (26.6%), showing that moving from file granularity to declaration granularity is itself the dominant source of improvement. The Symbol Overlap baseline (24.0%) is *not* the counterpart of $\mathcal{L}_{\text{sym}}$ alone (22.6%): that baseline matches shared function and class declarations, a scope spanning both $\mathcal{L}_{\text{ast}}$ and $\mathcal{L}_{\text{sym}}$, whereas $\mathcal{L}_{\text{sym}}$ in our decomposition covers only exported symbols, type aliases, constants, and enum members. The like-for-like comparison is therefore $\mathcal{L}_{\text{ast}}$ (28.4%) against Symbol Overlap (24.0%), which isolates the contribution of repository-calibrated weighting.

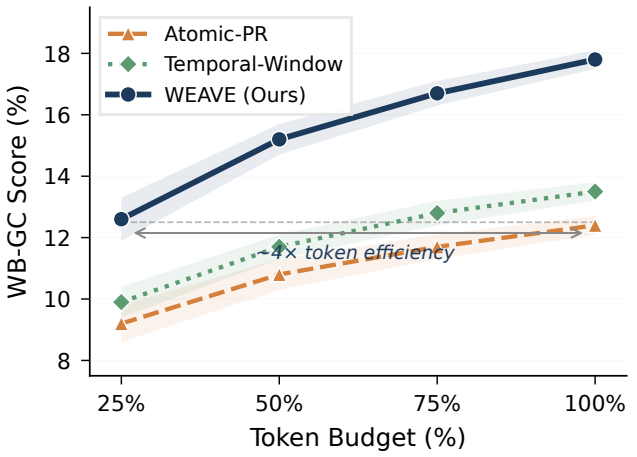

Figure 7: Token-scaling curves for SFT training (Qwen3.5-4B, measured by WEAVE-Bench Graph Credit). WEAVE (solid blue) achieves at 25% of the token budget what Atomic-PR (dashed orange) achieves at 100%, indicating a token-efficiency advantage of at least ∼4×. Temporal-Window (dotted green) falls between the two. The three curves differ in slope rather than offset, indicating that dependency-structured supervision is denser per token rather than merely better initialized. Error bars show ±1 std across 3 random seeds.

D.2 RQ3 Full Numerical Results and Token Scaling

Full SFT comparison table. Table 11 provides the complete numerical results underlying the SFT comparison figure of the main text (RQ3).

Token-scaling analysis. We evaluate how each data construction strategy scales with training token budget by training Qwen3.5-4B at 25%, 50%, 75%, and 100% of the full budget (∼158 M tokens; §E.1). Table 12 reports the resulting WB-GC scores and Figure 7 plots the corresponding curves.

All three strategies scale approximately log-linearly in token budget, gaining a roughly constant amount per doubling (∼1.6 pp for Atomic-PR, ∼1.8 pp for Temporal-Window, ∼2.6 pp for WEAVE), with no strategy saturating within the budgets tested. The curves are separated by slope rather than by offset, which is the signature of a difference in supervision density rather than a fixed head start. WEAVE reaches at 25% budget (12.6%) what Atomic-PR reaches only at 100% (12.4%). Since we do not evaluate below 25%, this establishes a token-efficiency advantage of *at least* ∼4× at this operating point rather than an exact ratio.

Table 11: Full SFT comparison across data construction strategies (matched token budget). All values are percentages (%). **Bold**: best per model scale; underline: best overall.

Model	Training Data	SWE-V	SWE-Pro	SWE-ML	WB-Res	WB-GC
Qwen3.5-4B						
	Base (no SFT)	50.4	24.2	35.3	4.0	8.6
	External Reference					
	SWE-rebench V2	51.8	25.4	36.6	4.7	9.5
	daVinci-Agency	50.6	24.0	35.7	3.7	8.2
	Controlled (same repos, matched tokens)					
	Atomic-PR	54.2	27.8	38.5	6.0	12.4
	Temporal-Window	55.1	28.6	39.2	6.7	13.5
	WEAVE	**57.8**	**31.0**	**41.7**	**9.0**	**17.8**
Qwen3.5-35B-A3B						
	Base (no SFT)	69.2	44.6	60.3	7.7	14.2
	External Reference					
	SWE-rebench V2	70.4	45.8	61.4	8.3	15.3
	daVinci-Agency	69.4	44.2	60.5	7.3	13.8
	Controlled (same repos, matched tokens)					
	Atomic-PR	71.8	47.6	62.7	10.3	18.6
	Temporal-Window	72.4	48.3	63.2	11.0	19.7
	WEAVE	**74.6**	**50.8**	**65.9**	**14.0**	**24.8**

Table 12: Token-scaling analysis (Qwen3.5-4B, WB-GC metric).

Budget	25%	50%	75%	100%
Atomic-PR	9.2	10.8	11.7	12.4
Temporal-Window	9.9	11.7	12.8	13.5
WEAVE	**12.6**	**15.2**	**16.7**	**17.8**

D.3 RQ4 Component Ablation and Curriculum

Reward component ablation. Table 13 evaluates the contribution of each component in R_{WEAVE} by progressive removal.

Graph Credit is the most impactful individual component: its removal reduces WB-GC by 4.4 pp and WB-Res by 1.7 pp. The Resolved bonus is the second most important component, providing an explicit signal for full-task completion. The regression penalty (Reg_{P2P}) has minimal impact on performance metrics but is critical for maintaining code quality: among the four reward variants compared in RQ4 the P2P regression rate stays at 1.8%, whereas removing the penalty from R_{WEAVE} raises it to 4.7%. This ablation-only configuration is the sole setting that exceeds the 2.3% regression ceiling reported in the main text.

The two single-signal rewards bound the ablation from below. Binary RESOLVED improves over the SFT checkpoint by at most 0.7 pp on any metric: on WEAVE-BENCH a rollout earns reward only when every constituent PR is solved, so most rollouts return identical zero reward and GRPO receives almost no advantage signal to differentiate them. Flat F2P restores a dense signal and recovers roughly half the gap to R_{WEAVE}, but because it credits downstream test passes irrespective of upstream state, it optimizes precisely the phantom credit that Graph Credit is designed to remove. Removing Graph Credit from the full reward (retaining Resolved, FlatF2P, and the regression penalty) remains better than Flat F2P alone, confirming that the residual components contribute independently.

Table 13: Reward component ablation (Qwen3.5-4B, after RL). All runs start from the same WEAVE-SFT checkpoint with matched rollout budgets. Δ is the WB-Res difference relative to full R_{WEAVE}. After each removal the remaining coefficients are renormalized to sum to one, so differences reflect the missing signal rather than a change in reward scale.

Reward Config	WB-Res	WB-GC	SWE-V	Δ
SFT only (no RL)	*9.0*	*17.8*	*57.8*	*−3.7*
R_{WEAVE} (full)	12.7	23.4	61.0	–
− GraphCredit	11.0	19.0	59.4	−1.7
− Resolved	11.6	22.3	60.1	−1.1
− FlatF2P	12.1	22.7	60.5	−0.6
− Reg_{P2P}	12.4	23.1	60.7	−0.3
Binary only (Resolved)	9.7	18.4	58.3	−3.0
Flat F2P only	10.7	18.8	59.0	−2.0

DAG-complexity curriculum. Table 14 compares the three-phase curriculum against flat (uniform random) task scheduling.

The three-phase curriculum provides consistent improvements of $\sim$1.0 pp across metrics. We observe that this benefit arises primarily in the first 30% of training: the curriculum prevents early divergence on overly complex tasks before the policy has developed basic multi-file editing capabilities.

Table 14: Curriculum scheduling comparison (Qwen3.5-4B). Phases defined by DAG complexity: P1 (depth $\leq$2, $|V| \leq$3), P2 (depth $\leq$3, $|V| \leq$5), P3 (all tasks).

Scheduling	WB-Res	WB-GC	SWE-V
Flat (uniform random)	11.7	22.2	60.2
3-phase curriculum	**12.7**	**23.4**	**61.0**
Δ	+1.0 pp	+1.2 pp	+0.8 pp

Curriculum phase details. The three phases are defined by DAG complexity thresholds and are allocated equal thirds of the update budget. Task counts and the schedule itself are given in §E.3; here we report the difficulty gradient the curriculum induces:

- **Phase 1** (0–33% of updates): depth $\leq$2 and $|V| \leq 3$. Mean $|V|$: 2.39. Mean gold patch: 166 lines.
- **Phase 2** (33–66%): depth $\leq$3 and $|V| \leq 5$. Mean $|V|$: 3.03. Mean gold patch: 211 lines.
- **Phase 3** (66–100%): all tasks, sampled uniformly. Mean $|V|$: 4.16. Mean gold patch: 289 lines.

Mean patch size scales near-linearly with node count at roughly 70 lines per PR node, so the curriculum increases edit scope by $\sim$1.7$\times$ from the first phase to the last.

E Training Details

The same dependency graph that defines evaluation tasks and Graph Credit scoring also structures the training pipeline. This section details how WEAVE leverages the PR dependency DAG to generate high-quality SFT trajectories and graph-structured RL rewards, completing the unified framework from construction through optimization.

E.1 Trajectory Generation and Quality Filtering

Trajectory generation setup. Training trajectories are generated using the Cline scaffold with the following configuration:

- **Generator model**: GLM-5.2.
- **Max turns**: 200.
- **Context window**: 128K tokens.
- **Temperature**: 0.6 (for diversity).
- **Tasks attempted**: 3,200 (all of WEAVE-Train).
- **Attempts per task**: 3 (best-of-3 by Graph Credit).

Choice of generator. We generate with GLM-5.2 rather than the strongest evaluated configuration (Claude Opus 4.7 + Cline, per the RQ2 results in the main text) for three reasons. First, GLM-5.2 has openly released weights, so the resulting corpus can be published without the derived-data restrictions that proprietary model outputs carry, and the generation stage is self-hostable and therefore independently reproducible. Second, GLM-5.2 is the second-strongest configuration on WEAVE-Bench (43.0% versus 48.7% Resolved); the best-of-3 protocol absorbs most of this 5.7 pp gap at the corpus level. Third, generator and student share a broadly similar tool-invocation format, which reduces format adaptation during distillation. The cost is throughput: GLM-5.2 averages 138.2 turns and 7.52 M cumulative tokens per benchmark task, against 46.8 turns and 2.11 M for Claude Opus 4.7. This 3.6$\times$ generation overhead is why we cap attempts at three per task rather than sampling more widely.

Quality filtering. Generated trajectories pass through five filters. These are deliberately strict: a trajectory that satisfies downstream tests while leaving upstream prerequisites unimplemented exhibits precisely the phantom credit that WEAVE-Bench is designed to expose, and distilling such behavior into the student would be self-defeating. We therefore gate on Graph Credit rather than on flat test pass rate.

1. **Dependency-consistent success filter**: Retain trajectories satisfying both $\mathrm{GraphCredit}(G_A) \geq 0.85$ and $\min_{v \in V_A} c_v \geq 0.50$. The first condition requires near-complete resolution of the task graph. The second is applied per node and forbids any single node from remaining grossly unimplemented; because c_v is already bottlenecked by its ancestors under the recursive definition of Graph Credit in the main text, it constrains every dependency path simultaneously. A trajectory with high Flat-F2P but an unresolved upstream node fails this filter by construction.
2. **Length filter**: Discard trajectories shorter than 8 turns, and any trajectory that terminated by exhausting the 200-turn or context budget rather than by submitting a patch, since truncated rollouts are incomplete by construction. The resulting effective upper bound is 150 turns, 75% of the interaction cap.
3. **Regression filter**: Discard trajectories exhibiting any P2P regression. Failing P2P tests are re-executed three times and counted as regressions only when they fail consistently, so that flaky tests do not discard otherwise valid trajectories.
4. **Degenerate pattern filter**: Remove trajectories with three or more consecutive identical tool calls (same tool, same arguments), or with edit churn above 0.5—the fraction of edit operations subsequently overwritten or reverted within the same trajectory—which indicates thrashing rather than incremental progress.
5. **Leakage filter**: Discard trajectories in which the agent recovered the original change through repository history or network access, detected by scanning tool-call arguments for `git log`, `git show`, `.git` traversal, and outbound requests. Task instantiation already removes recoverable target commits and repository history; this filter enforces the same guarantee at the trajectory level, where an agent could otherwise reach them at run time.

Surviving rollouts are deduplicated per task: we keep the attempt with the highest Graph Credit, break ties by shorter trajectory length, and treat remaining attempts whose final patches exceed 0.9 normalized edit-distance similarity to the retained one as duplicates.

Corpus statistics after filtering. Table 15 reports the yield at each stage.

Table 15: Trajectory generation and filtering yield. All rows count rollouts except the last, which counts one retained trajectory per task.

Stage	Retained	Pass rate
Rollouts generated ($3{,}200 \times 3$)	9,600	–
Dependency-consistent success	4,320	45.0%
Length	4,058	93.9%
Regression	3,857	95.0%
Degenerate pattern	3,699	95.9%
Leakage	3,644	98.5%
Per-task deduplication	**2,304**	**72.0%**[†]

[†]Fraction of the 3,200 tasks retaining at least one trajectory.

Table 16: SFT training configuration. Search ranges are given for parameters tuned on the held-out calibration set; "–" marks values fixed a priori.

Parameter	4B	35B-A3B	Search range
Learning rate	5e-6	3e-6	{2,3,5,10}e-6
LR schedule	Cosine, 10% floor		–
Warmup ratio	0.03	0.03	$\{0.01, 0.03, 0.05\}$
Global batch (seq.)	16	16	–
Training epochs	3	3	$\{2, 3, 4\}$
Max sequence length	131,072		–
Weight decay	0.01	0.01	–
Adam β_1/β_2	0.9 / 0.95		–
Gradient clipping	1.0	1.0	–
Precision	BF16		–
Loss masking	Assistant tokens only		–

The final corpus contains 2,304 trajectories covering 72.0% of WEAVE-TRAIN tasks, with mean length 61.4 turns and mean size 68.4K tokens, totaling $\sim$158 M training tokens. Task coverage remains approximately uniform across dependency depth—74.1% for tasks of depth ≤ 2 versus 69.4% for depth ≥ 4—so the strict success filter does not bias the corpus toward shallow structures, which would otherwise undercut the supervision signal that motivates dependency-structured construction.

E.2 SFT Configuration

Table 16 details the SFT hyperparameters for both model scales. We report only parameters that affect the optimization objective or its trajectory; sharding strategy, device placement, and accumulation factors are implementation choices that follow from the global batch size and are omitted.

Three choices warrant comment. **Loss masking**: gradients are computed only on assistant-generated tokens; tool results, file contents, and test output are supplied as context but excluded from the objective, since these constitute environment observations rather than target behavior. **Sequence length**: the 131,072-token window matches the generator's context budget, so no retained trajectory is truncated during training; training at this length uses sequence parallelism. **Learning rate**: with only 2,304 trajectories, the corpus is small enough that the higher end of the search range overfits within one epoch. Qwen3.5-35B-A3B activates roughly 3B parameters per token, so its optimal rate is close to—though slightly below—that of the 4B dense model rather than an order of magnitude lower.

Both scales use identical batch size and epoch count so that the controlled comparison in RQ3 varies only task construction. The resulting schedule is 144 optimizer steps per epoch and 432 steps overall, of which the first 13 are warmup. All runs use a single fixed random seed; reported results are single-run unless stated otherwise.

Data mixing strategy. For controlled comparisons, we ensure matched token budgets across all data sources:

- Each training arm uses $\sim$158 M tokens.
- For external references (SWE-rebench V2, daVinci-Agency), we sample trajectories uniformly to match the token budget.
- For controlled variants (Atomic-PR, Temporal-Window, WEAVE), trajectories are generated from the same repository pool using the same generator and quality filter, differing only in task construction.
- All variants use identical SFT configuration (Table 16).

Checkpoint selection. We select the final checkpoint by Graph Credit on a held-out calibration set of 50 tasks disjoint from WEAVE-BENCH, evaluating every 36 steps (one quarter epoch) for 12 evaluation points across the 432-step schedule.

E.3 RL Configuration

Table 17 details the GRPO configuration for RQ4. Unlike SFT, RL requires no gold trajectory—only an executable environment—so the prompt pool is the full WEAVE-TRAIN set of 3,200 tasks rather than the 2,304 tasks for which a demonstration survived filtering.

We specify the interaction budget in terms of the quantities that define it—pool size, group size, tasks per update, and passes over the pool—and report the induced update and episode counts as derived values, so that the budget can be reconstructed exactly.

Both scales share the same interaction budget so that the reward comparison in RQ4 is not confounded by differing amounts of experience. The rollout temperature of 1.0 is deliberate: GRPO derives its advantage signal from variation *within* each group of G rollouts, so insufficient sampling diversity collapses the group to a common reward and contributes no gradient.

Reward coefficients. The reward R_{WEAVE} uses the following coefficients (selected via grid search on calibration set):

- $\lambda_{\mathrm{graph}} = 0.5$ (dominant component)
- $\lambda_{\mathrm{solve}} = 0.25$ (full-task bonus)
- $\lambda_{\mathrm{f2p}} = 0.15$ (flat progress)
- $\lambda_{\mathrm{reg}} = 0.1$ (regression penalty)

Table 17: RL (GRPO) training configuration. Rows marked $*$ are derived from the rows above them. Search ranges are given for parameters tuned on the calibration set.

Parameter	4B	35B-A3B	Search range
RL prompt pool	3,200 tasks		–
Rollouts per task (G)	8	8	–
Tasks per update	16	16	–
Passes over pool	2	2	–
$*$ Policy updates	400	400	–
$*$ Total episodes	51,200		–
Learning rate	1e-6	5e-7	{5e-7, 1e-6, 2e-6}
LR schedule	Cosine decay		–
KL coefficient (β)	0.01	0.01	{0.005, 0.01, 0.04}
Clip ratio (ϵ)	0.2	0.2	–
Rollout temperature	1.0	1.0	{0.8, 1.0}
Max turns per episode	200	200	–
Reference policy	SFT checkpoint		–
Advantage	Group-relative		–
Precision	BF16		–

Advantage normalization. GRPO is critic-free: for each task we sample $G = 8$ rollouts, score each with R_{WEAVE}, and take the advantage of a rollout to be its reward centered and scaled by the mean and standard deviation *within its own group*. No value network, and hence no generalized advantage estimation, is used. Group-relative normalization makes the gradient invariant to task difficulty, so that easy and hard tasks contribute comparably.

Group-relative normalization also explains why reward density matters here beyond the usual sparse-reward argument: a group whose G rollouts all receive the same reward contributes no gradient at all, so a sparse reward wastes rollout budget outright rather than merely learning slowly. Because Graph Credit is continuous and dependency-conditioned, its groups almost always exhibit within-group spread. The empirical consequence is quantified in §D.3.

Three-phase curriculum. RL training follows a DAG-complexity curriculum over the 3,200-task pool. Phases are cumulative—each admits all tasks from the previous phase—and are allocated equal thirds of the 400 policy updates:

- **Phase 1** (updates 1–132): DAGs with depth $\leq$2 and $|V| \leq 3$. Eligible tasks: 1,048.
- **Phase 2** (updates 133–264): DAGs with depth $\leq$3 and $|V| \leq 5$. Eligible tasks: 1,864.
- **Phase 3** (updates 265–400): all tasks, sampled uniformly. Eligible tasks: 3,200.

Phase membership follows directly from the WEAVE-Train node and depth distributions (Table 20); depth is measured as the number of edges on the longest path, consistent with §B.3.

F Dataset Statistics and Case Studies

This section provides comprehensive statistics on WEAVE-Bench and WEAVE-Train, followed by qualitative case studies demonstrating how Graph Credit operates on concrete dependency structures produced by the construction pipeline.

Table 18: Language distribution across dataset splits.

Language	WEAVE-Bench	WEAVE-Train
Python	50 (16.7%)	720 (22.5%)
TypeScript	50 (16.7%)	680 (21.3%)
JavaScript	50 (16.7%)	640 (20.0%)
Go	50 (16.7%)	480 (15.0%)
Kotlin	50 (16.7%)	360 (11.3%)
Java	50 (16.7%)	320 (10.0%)
Total	**300**	**3,200**

Table 19: DAG structural statistics across splits. Depth counts nodes on the longest path, matching the benchmark distribution in Figure 3. Training depth is converted from edge count by adding one; curriculum thresholds use edge count.

Metric	WEAVE-Bench		WEAVE-Train	
	Mean	Median	Mean	Median
Nodes ($\lvert V\rvert$)	4.91	5	4.16	4
Edges ($\lvert E\rvert$)	4.99	5	4.28	4
Depth (path nodes)	3.93	4	3.71	4
Width (max anti-chain)	1.50	1	1.31	1

F.1 WEAVE-Bench and WEAVE-Train Statistics

Language distribution. Table 18 reports the language composition of both WEAVE-Bench and WEAVE-Train.

DAG structure statistics. Table 19 characterizes the dependency graph structures.

Table 20 gives the full node-count distribution for WEAVE-Train, which determines the eligible task pool at each phase of the RL curriculum (§E.3).

Topology distribution.

- **Linear chains** (all nodes have in/out-degree $\leq$1): 44.3% of WEAVE-Bench tasks.
- **Diamond structures** (fork followed by join): 15.7%.
- **General DAGs** (mixed fork/join patterns): 15.0%.
- **Fan structures** (one root, multiple leaves): 13.3%.
- **Cross structures** (multiple interleaved dependencies): 11.7%.

Patch and test statistics.

Repository coverage. WEAVE-Bench draws from 52 unique repositories. The most-represented repository contributes 23 tasks (7.7%); the top-12 repositories account for $\sim$60% of tasks, with 12 repositories contributing only a single task each. WEAVE-Train covers $\sim$120 repositories.

F.2 Qualitative Examples

We present two case studies illustrating how Graph Credit reveals phantom credit in specific tasks.

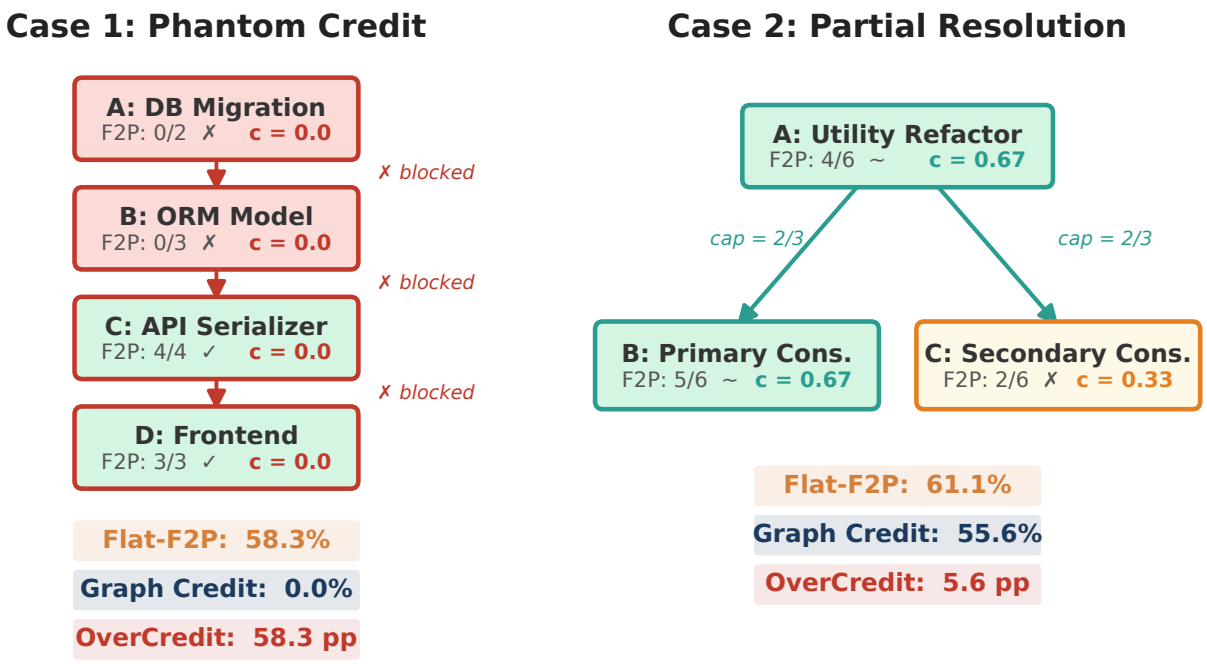

Figure 8: Dependency graphs for Case Studies 1 and 2. **Left**: 4-node DAG (Case 1). Node A (DB migration) → Node B (ORM model) → Node C (API serializer) → Node D (frontend). Green fill indicates agent-completed nodes; red fill indicates unresolved. **Right**: 3-node DAG (Case 2). Node A (utility refactor) is partially resolved (4/6 tests) and branches to Node B (primary consumer, 5/6 tests) and Node C (secondary consumer, 2/6 tests). Effective credits are $c_A = c_B = 2/3$ and $c_C = 1/3$; Node B is capped by its upstream prerequisite. Right-panel fill intensity encodes c_v.

Table 20: Node-count distribution of WEAVE-Train ($|V|$ mean 4.16, median 4).

$\|V\|$	2	3	4	5	6	7	8
Tasks	640	704	640	480	352	224	160
Share	20.0	22.0	20.0	15.0	11.0	7.0	5.0

Table 21: Patch and test statistics for WEAVE-Bench.

Metric	Mean	Median	P25	P75
Gold patch (lines)	378	274	139	528
Files modified	8.4	6	3	12
F2P tests	18.6	12	6	24
P2P tests	142	87	34	198
Attributed tests (%)	72.8	76.4	61.2	86.3

Case 1: Database migration with downstream API changes. **Task structure:** A 4-node DAG in a Python web application. Node A introduces a new database column; Node B modifies the ORM model to expose the column; Node C updates the API serializer; Node D adds frontend validation. The dependency graph captures that C and D cannot function correctly without B's ORM changes, which in turn require A's schema migration.

Agent behavior (Claude Opus 4.7 / Cline): The agent successfully implements Nodes C and D (API serializer and frontend validation) by pattern-matching existing code, but fails to modify the database migration (Node A) or ORM model (Node B). The downstream tests for Nodes C and D pass because the test fixtures mock the database layer.

Key F2P tests and attribution:

- Node A: `test_migration_adds_column`, `test_rollback_migration` (2 tests, both fail)
- Node B: `test_model_field_exists`, `test_model_serialization`, `test_model_validation` (3 tests, all fail)
- Node C: `test_api_serialize_new_field`, `test_api_response_schema`, `test_api_pagination_with_field`, `test_api_filter_by_field` (4 tests, all pass via mock)
- Node D: `test_frontend_validation_rule`, `test_form_field_rendering`, `test_error_message_display` (3 tests, all pass via mock)

Scoring comparison:

- **Flat-F2P**: 58.3% (7/12 attributed tests pass).
- **Graph Credit**: 0.0% (Nodes C and D depend on Node B, which depends on Node A; since $c_A = 0$, all downstream

credits are capped at 0).

- **OverCredit**: 58.3 pp.

This case demonstrates the most extreme form of phantom credit: the agent receives majority credit under flat scoring despite implementing none of the foundational changes. Graph Credit correctly identifies that no genuine progress has been made toward a coherent solution.

Case 2: Partial resolution with mixed upstream completion. Task structure: A 3-node DAG in a TypeScript project. Node A refactors a shared utility module; Node B (depends on A) updates the primary consumer; Node C (depends on A) updates the secondary consumer. This fork structure tests whether Graph Credit correctly attributes partial completion across parallel downstream nodes that share a common dependency.

Agent behavior (GLM-5.2 / OpenCode): The agent refactors Node A but misses two behaviors of the original utility (edge-case handling and error propagation), implements Node B (primary consumer update) against the new interface, and produces an incorrect implementation of Node C that passes only 2 of 6 attributed tests.

Key F2P tests and attribution:

- Node A: `test_util_new_api`, `test_util_backward_compat`, `test_util_edge_cases`, `test_util_error_handling`, `test_util_types`, `test_util_async` (6 tests, 4 pass; `edge_cases` and `error_handling` fail)
- Node B: `test_primary_consumer_integration`, `test_primary_consumer_error_path`, `test_primary_consumer_caching`, `test_primary_consumer_retry`, `test_primary_consumer_events`, `test_primary_consumer_metrics` (6 tests, 5 pass)
- Node C: `test_secondary_consumer_basic`, `test_secondary_consumer_batch`, `test_secondary_consumer_timeout`, `test_secondary_consumer_fallback`, `test_secondary_consumer_logging`, `test_secondary_consumer_config` (6 tests, 2 pass)

Scoring comparison:

- **Flat-F2P**: 61.1% (11/18 attributed tests pass: 4/6 for A, 5/6 for B, 2/6 for C).
- **Graph Credit**: 55.6%. Node credits: $c_A = 4/6 = 0.67$ (no parents), $c_B = \min(5/6, c_A) = 0.67$, $c_C = \min(2/6, c_A) = 0.33$. Mean $= (2/3 + 2/3 + 1/3)/3 = 5/9 \approx 0.556$.
- **OverCredit**: 5.6 pp.

This fork exercises both branches of the recursive minimum. Node B's local score (5/6) exceeds its parent's completion ($c_A = 2/3$) and is therefore capped: flat scoring would report the primary consumer as nearly complete even though two behaviors of the utility module it consumes are still missing. Node C, by contrast, already scores below its parent, so its own local completion remains the binding constraint and Graph Credit leaves it unchanged. The entire 5.6 pp OverCredit is thus attributable to Node B—precisely the structurally unsupported portion of its score.

Aggregate patterns. Across all 4,800 agent runs on WEAVE-BENCH:

- 68% of runs exhibit >0 pp OverCredit.
- 23% exhibit >15 pp OverCredit (substantial phantom credit).
- 8% exhibit >30 pp OverCredit (severe misrepresentation under flat scoring).
- The correlation between Flat-F2P and actual Resolved is $\rho = 0.69$, while Graph Credit achieves $\rho = 0.82$—a 19% improvement in diagnostic fidelity.