arXiv:2610.21755v1 [cs.AI] 15 Oct 2026

DeciCAD: Harness-Grounded Learning of Action Policies for CAD Agents

Yufeng Peng[1], Yixin Zhang[1], Ziyi Liu[1], Zhikun Xu[1], Jiarong Yang[1], Sijie Wang[2], Wenjie Feng[1], Qiyu Kang[1†]

[1]University of Science and Technology of China
[2]Nanyang Technological University

github.com/PhysAGI/PhysCAD

Abstract

Large Language Model agents operate through closed-loop action policies whose decision space is defined by an agent harness. We investigate how this operating space shapes agent behavior and policy learning in computer-aided design (CAD), a professional domain characterized by persistent state, geometric constraints, and long-horizon procedures. We instantiate this perspective in DeciCAD, a domain-native harness that exposes semantic CAD actions, decision-oriented observations, environment feedback, and persistent state transitions. Building on this harness, we develop a learning pipeline that extracts causal decision supervision from interaction trajectories and improves the policy through supervised fine-tuning (SFT) and geometry-grounded online reinforcement learning (RL). Across four CAD benchmarks, DeciCAD achieves the best F1 and intersection over union (IoU) while reducing model-output tokens by 45–88% relative to a general-purpose coding harness. Under matched supervision, its SFT policy improves F1 from 0.6883 to 0.7437 and outperforms the coding policy; subsequent online RL further raises geometry availability from 66.5% to 73.5% and F1 to 0.7880. These results demonstrate that the decision space induced by a harness materially affects both agent execution and policy learning in CAD.

1 Introduction

Large Language Models (LLMs) are increasingly being used as agents to operate professional software and complete long-horizon tasks (Achiam et al. 2023; Naveed et al. 2025; Guo et al. 2024). Computer-aided design (CAD) is a particularly demanding setting for such agents (Willis et al. 2020; Wu, Xiao, and Zheng 2021). Unlike one-shot content generation, CAD modeling requires a sequence of state-dependent operations over a persistent design document. An agent must create and modify sketches, constraints, features, and geometric objects; inspect intermediate results; recover from execution failures; and maintain references across long interaction trajectories. In this process, the model outputs are interpreted as actions: each action changes the current CAD state, produces execution feedback, and conditions the next decision (Yao et al. 2023; Schick et al. 2023). The effectiveness of a CAD agent therefore depends not only on the underlying model and its reasoning strategy, but also on how

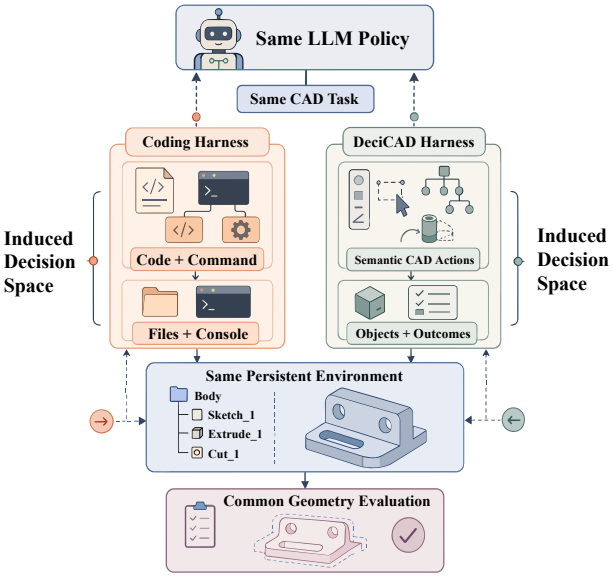

Figure 1: Harness-induced CAD decision spaces under controlled model and environment conditions. Given the same LLM policy, CAD task, persistent design state, and geometry evaluator, the coding harness represents decisions through code and shell commands and returns files and console output. DeciCAD instead exposes semantic CAD actions and returns object-level state and execution outcomes. The two interfaces therefore induce different action–observation histories over the same underlying modeling process.

CAD actions, states, and execution outcomes are represented to the model.

Such a model–environment boundary is defined by an *agent harness* (Yang et al. 2024). In CAD, a harness determines which design operations the model can express, which parts of the evolving design state it can inspect, and how execution outcomes are returned for subsequent decisions. A general-purpose coding harness may expose the CAD system through programs, libraries, and runtime feedback, whereas a domain-native harness may expose structured CAD actions and object-level observations. Consequently, even when the model, task, and CAD backend remain fixed, different harnesses can present the model with substantially different de-

†Correspondence to: Qiyu Kang <qiyukang@ustc.edu.cn>.

cision spaces, as illustrated in Figure 1. Existing work has studied CAD code generation, tool-using agents, execution repair, and interactive modeling (Mallis et al. 2024; Guan et al. 2025; Doris et al. 2025; Wang et al. 2025; Barkley, Loghmani, and Barati Farimani 2026; Gong et al. 2026), but the harness itself is typically treated as implementation infrastructure rather than as a variable that shapes both agent execution and policy learning. This raises a central question: *How should a CAD agent harness represent actions and states so that it preserves the decisions required for complex modeling while avoiding representation burden that is irrelevant to the task?*

This trade-off is important in CAD. A fine-grained coding interface gives the model broad control over the underlying software, but the same CAD operation can often be expressed through many programs, command sequences, variable names, temporary objects, and API combinations. These implementations may produce the same geometric state transition, but they appear as different generation and supervision targets. The model must then spend capacity on implementation choices that do not affect the intended state transition. A semantic interface can remove many of these choices by representing the intended operation directly. However, excessive compression can hide state differences that require different actions or remove operations needed for correct modeling and error recovery. A useful harness must therefore reduce irrelevant representation choices while preserving the actions and observations needed for correct decisions.

To make this trade-off explicit, we formalize the model-facing interface as a *Harness-Induced CAD Decision Space*: a harness maps model outputs to executable operations, maps internal CAD states to model-visible observations, and returns execution results as feedback, thereby determining which operations the agent can perform, which state differences and failures it can identify, and which policies it can execute and learn. We then introduce the *Semantic Decision Bottleneck* as a design principle for this space: it compresses behaviorally equivalent implementations and removes state information that does not affect the next decision, while preserving the actions and observations needed to reach the target geometry and recover from errors.

Following this principle, we develop DeciCAD, a domain-native harness for closed-loop CAD agents. DeciCAD represents task-relevant design transitions as semantic CAD actions, grounds these actions in a persistent CAD document, and returns structured execution feedback together with decision-oriented observations. We further develop a harness-grounded learning pipeline that reconstructs causally conditioned action supervision from successful and recovered interaction trajectories, initializes CAD action policies through supervised fine-tuning (SFT), and supports online reinforcement learning (RL) using policy-generated trajectories and geometry-grounded rewards. Figure 2 summarizes this execution and learning loop.

Across four CAD benchmarks, DeciCAD consistently produces more accurate final geometry than competitive CAD agents and a general-purpose coding harness, while reducing model-output tokens by 45–88% relative to the coding harness on text-to-CAD tasks. Observation ablations show that showing more CAD state does not always help: unfiltered state inflates the input context, and coarse feedback strips out what the agent needs in order to recover. With matched training settings, open-weight models trained in the DeciCAD decision space learn stronger long-horizon CAD action policies than the same models trained through code-level interaction.

Our contributions are summarized as follows:

- We define the *Harness-Induced CAD Decision Space* and introduce the *Semantic Decision Bottleneck* as a design principle for CAD agent interfaces. We design DeciCAD, which combines semantic CAD actions, compact observations, structured feedback, and persistent object state.
- We develop a harness-grounded policy-learning pipeline that reconstructs causally conditioned supervision from interaction trajectories, initializes policies through SFT, and supports geometry-grounded online RL.
- Experiments across four benchmarks verify DeciCAD improves geometric accuracy, token efficiency, error recovery, and long-horizon policy learnability significantly.

2 Related Work

Our work connects agent interaction interfaces, decision-preserving abstraction, and policy learning in interactive CAD environments. We review these directions through the common question of how an agent's operating representation affects execution and learning.

2.1 Agent Interfaces and Decision Abstraction

Tool-augmented language models connect generation to executable operations (Parisi, Zhao, and Fiedel 2022; Karpas et al. 2022; Schick et al. 2023), while API-Bank, Gorilla, and ToolLLM extend tool learning to executable dialogues and large API ecosystems (Li et al. 2023; Patil et al. 2023; Qin et al. 2024). ReAct interleaves reasoning, actions, and observations (Yao et al. 2023); SWE-agent, WebArena, and OSWorld further show that interaction design affects behavior in stateful coding, web, and computer environments (Yang et al. 2024; Zhou et al. 2024; Xie et al. 2024). They establish the empirical importance of interfaces, but generally treat them as tool provision or runtime infrastructure. We instead study the harness as the joint action–state representation that defines the policy space presented to the model.

This view is related to abstraction in sequential decision-making. MDP homomorphisms and state-abstraction theory characterize when compressed representations preserve transition, reward, value, or policy structure (Ravindran and Barto 2003; Li, Walsh, and Littman 2006; Abel, Hershkowitz, and Littman 2016; Abel et al. 2020); the Information Bottleneck similarly retains target-relevant information while suppressing irrelevant input dependence (Tishby, Pereira, and Bialek 2000). Classical abstractions operate within an already specified decision process. An LLM harness instead jointly defines model actions, visible observations, execution bindings, and feedback semantics. Our Harness-Induced Decision Space describes this transformation, and the Semantic Decision Bottleneck constrains its compression by decision preservation.

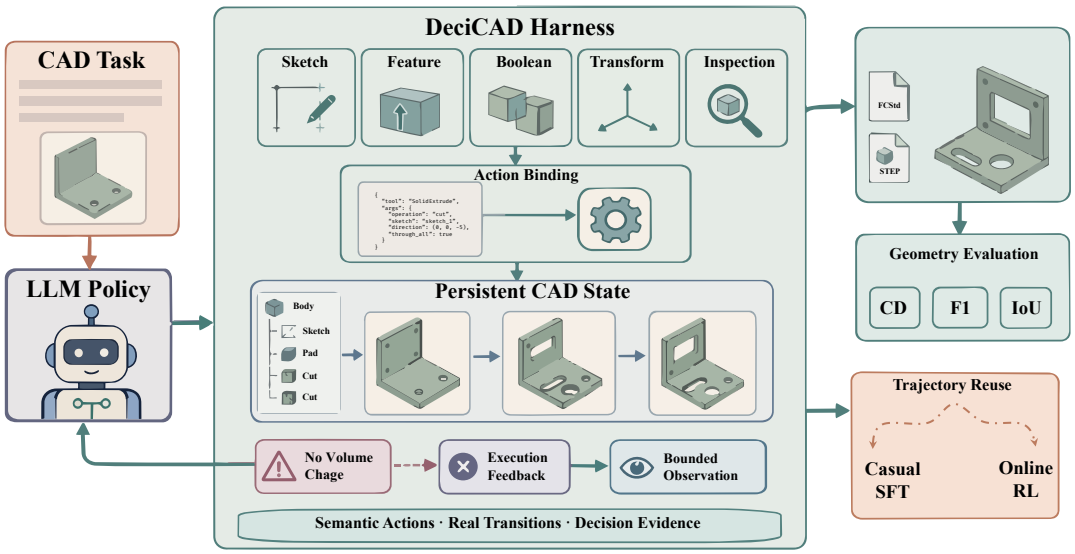

Figure 2: Overview of DeciCAD. Semantic CAD actions are bound to real transitions in persistent CAD state, with bounded observations and execution feedback; the resulting trajectories support causal SFT and geometry-grounded online RL.

2.2 CAD Generation and Interactive Agents

CAD generation has been studied through construction sequences, executable programs, and interactive actions. DeepCAD and Fusion 360 Gallery represent structured modeling histories (Wu, Xiao, and Zheng 2021; Willis et al. 2020), while Text2CAD and CADmium learn parametric command sequences from language (Khan et al. 2024; Govindarajan et al. 2026). CAD-Coder systems generate executable CAD code (Guan et al. 2025; Doris et al. 2025), and CADFusion, CADSmith, and CADDesigner incorporate rendering, execution repair, or multi-agent feedback (Wang et al. 2025; Barkley, Loghmani, and Barati Farimani 2026; Fan et al. 2025). Code is expressive, but also makes program organization and API realization part of the model output.

Interactive systems instead place decisions in evolving design state. CAD-Assistant combines a Python interpreter with FreeCAD operations (Mallis et al. 2024); ToolCAD trains tool-using CAD agents from demonstrations and online curriculum learning (Gong et al. 2026); CADMate learns sequential assembly actions using geometric and physical rewards (Chen et al. 2026). Our contribution is complementary: under fixed models, tasks, and terminal geometry, we compare the decision spaces induced by domain-semantic actions and a general-purpose coding harness.

2.3 Policy Learning from Interaction

Tool policies commonly acquire initial capability from generated or curated demonstrations (Schick et al. 2023; Li et al. 2023; Qin et al. 2024). In a stateful environment, however, action validity depends on the preceding trajectory and feedback; imitation therefore encounters distribution shift when student-induced states depart from teacher trajectories (Ross, Gordon, and Bagnell 2011). Online methods address this gap using execution outcomes on states visited by the learned policy (Feng et al. 2025; Qian et al. 2025), including curriculum reinforcement learning in CAD (Gong et al. 2026). We align causal supervision, student rollouts, and environment feedback within the same Harness-Induced Decision Space, enabling a controlled study of how representation affects both imitation and online learning.

3 Method

The harness determines how the closed CAD loop is presented to the language model: whether the model writes general-purpose code or selects a semantic CAD action, which part of the current design is shown, and how execution failures are reported. We formalize this interaction, formulate the Semantic Decision Bottleneck, instantiate it in DeciCAD, and use the same representation for SFT and online RL.

3.1 From a CAD Backend to a Model-Facing Decision Process

CAD tasks and backend state. We use g to denote one concrete CAD task, which may contain a natural language instruction, optional initial state or reference image, and a final design criterion. At step t, its internal state $s_t \in \mathcal{S}$ includes the current document, feature structure, geometry, parameters, and dependencies; note that this is the CAD system's internal configuration, rather than text shown to the LLM. An element $u_t \in \mathcal{U}$ is a backend operation, and $s_{t+1} \sim P(\cdot \mid s_t, u_t)$ describes how the document changes after its execution. Task g specifies an initial-state distribution μ_g, a maximum number of interaction steps H, and reward $r_g(s_t, u_t, s_{t+1})$. This process describes what the CAD system can execute. It does not yet describe what the language model is allowed to say or see.

What the harness changes. A language model does not directly issue an internal backend operation u_t, nor does

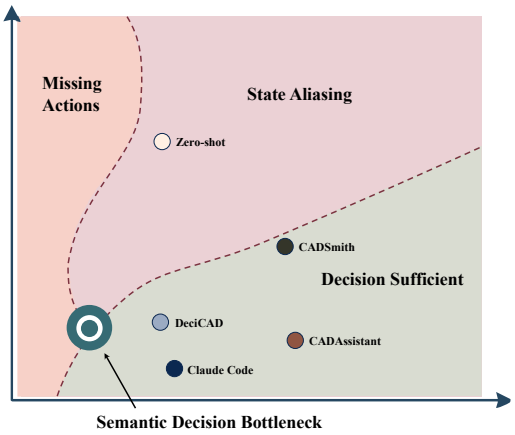

Figure 3: Semantic CAD Decision Bottleneck. Missing actions and state aliasing violate action and observation sufficiency, respectively; DeciCAD is designed to approach the decision-sufficient boundary. Method locations are conceptual and not to scale.

it receive the full backend state s_t. A harness provides the model-facing interface between the two. We represent it as

$$\mathcal{H} = (\mathcal{A}_{\mathcal{H}}, \beta_{\mathcal{H}}, O_{\mathcal{H}}), \tag{1}$$

whose components have direct operational meanings:

- $\mathcal{A}_{\mathcal{H}}$ is the set of admissible model-level actions exposed by the harness. In a coding harness, an action may be a Python program. In DeciCAD, it is a typed operation, such as `Extrude(sketch=S3, distance=10mm)`.
- $\beta_{\mathcal{H}}$ interprets the model output in the current CAD document and binds it to an executable backend operation.
- $O_{\mathcal{H}}$ determines what the model receives after execution, including compact descriptions of success, failure, and affected design objects.

The harness thus translates between model outputs, backend operations, and model-visible execution evidence without replacing the CAD kernel.

One closed-loop interaction step. One interaction step is

$$\begin{aligned} a_t &\sim \pi_\theta(\cdot \mid h_t; \mathcal{H}), \\ u_t &= \beta_{\mathcal{H}}(a_t, s_t), \\ s_{t+1} &\sim P(\cdot \mid s_t, u_t), \\ y_{t+1} &\sim O_{\mathcal{H}}(\cdot \mid s_t, a_t, s_{t+1}), \end{aligned} \tag{2}$$

Here, a_t is what the model emits, u_t is what the CAD backend executes, and y_{t+1} is what the model sees afterward. The model chooses its next action from the visible history

$$h_t = (g, y_0, a_0, y_1, \ldots, a_{t-1}, y_t).$$

This contains the task specification and previous model actions and harness responses, but not the hidden backend state s_t. Read-only inspection is also represented as an action, keeping information gathering and geometry-changing operations in the same closed loop.

Combining backend execution and harness response gives the induced transition kernel

$$T_{\mathcal{H}}(s', y \mid s, a) = P(s' \mid s, \beta_{\mathcal{H}}(a, s))\, O_{\mathcal{H}}(y \mid s, a, s'). \tag{3}$$

We call the model-facing interaction defined by Eqs. (2)–(3) the *Harness-Induced CAD Decision Space*, comprising the actions, responses, and interaction histories exposed by the harness.

The distinction is important. A coding harness and DeciCAD may call the same CAD kernel yet induce different sequential prediction problems. The former requires choices about program structure, API calls, and error handling; the latter requires choices about CAD operations and their geometric parameters, with different evidence after each step.

3.2 Semantic Decision Bottleneck

The preceding formulation exposes two sources of unnecessary difficulty. On the output side, many syntactically different programs can realize the same CAD change. On the input side, the CAD backend contains much more information than each decision requires. A useful harness should compress both directions without removing a necessary operation or hiding a distinction that changes the correct next action. Figure 3 illustrates this trade-off.

Preserving necessary actions. Let $V^*_{g,t}(s)$ be the best return achievable from backend state s at step t, and let $Q^*_{g,t}(s, u)$ be the return obtained by first executing backend operation u and then acting optimally. A harness is *action-sufficient* if

$$\max_{a \in \mathcal{A}_{\mathcal{H}}} Q^*_{g,t}(s, \beta_{\mathcal{H}}(a, s)) = V^*_{g,t}(s) \tag{4}$$

for every task-reachable state. The harness need not expose every low-level CAD command under Eq. (4), only an action realizing an optimal next operation when needed.

Preserving necessary state distinctions. A harness response need not reveal the backend state, so one visible history may be consistent with several backend states, as in a partially observable decision process (Kaelbling, Littman, and Cassandra 1998). Define

$$\mathcal{B}_{\mathcal{H}}(h_t) := \{ s \in \mathcal{S} : \Pr(s_t = s \mid h_t; \mathcal{H}) > 0 \} \tag{5}$$

as the backend states consistent with h_t. We call h_t *reachable* when this set is nonempty. For a backend state s, let

$$\mathcal{A}^*_{\mathcal{H},g,t}(s) := \arg\max_{a \in \mathcal{A}_{\mathcal{H}}} Q^*_{g,t}(s, \beta_{\mathcal{H}}(a, s)). \tag{6}$$

These are the model actions optimal at s. The harness is *observation-sufficient* if the states in Eq. (5) share an action from the sets in Eq. (6):

$$\bigcap_{s \in \mathcal{B}_{\mathcal{H}}(h_t)} \mathcal{A}^*_{\mathcal{H},g,t}(s) \neq \varnothing, \quad \forall \text{ reachable } h_t,\ t < H. \tag{7}$$

The harness may map two backend states to the same visible history when the same next action is optimal in both; it must distinguish them when they require incompatible next actions. A generic "operation failed" message is therefore insufficient if different failures require different repairs.

A harness that is both action-sufficient and observation-sufficient is *decision-sufficient*. Any selection from the intersection in Eq. (7) depends only on h_t and, by action sufficiency, is optimal at the true state; backward induction gives $\mathbb{E}_{s_0\sim\mu_g}[V^*_{g,0}(s_0)]$ for the resulting history-dependent policy. An inexpressible operation therefore violates action sufficiency, whereas a response merging states that require incompatible repairs violates observation sufficiency.

Compressing implementation and context. Action and observation sufficiency specify what a harness must preserve, but do not uniquely determine its model-facing representation. Among decision-sufficient interfaces, we prefer those that reduce unnecessary burden on both model output and input. We refer to this principle as the *Semantic Decision Bottleneck*: compressing action representations and visible context while retaining the distinctions required for correct action selection and error recovery. DeciCAD instantiates this principle through domain-specific action schemas, concise execution feedback, and targeted inspection operations that retrieve additional state only when needed. A formal preference relation over decision-sufficient harnesses is provided in the appendix.

3.3 DeciCAD: A Domain-Native CAD Harness

DeciCAD implements the preceding design principle through semantic actions, state-grounded execution, and decision-oriented responses. These three components correspond directly to $\mathcal{A}_\mathcal{H}$, $\beta_\mathcal{H}$, and $O_\mathcal{H}$ in Eq. (1).

Semantic CAD Actions DeciCAD represents each model action as

$$a_t = (\omega_t, \mathbf{x}_t),$$

where ω_t identifies a typed CAD operation and $\mathbf{x}_t$ contains its structured arguments. These actions express design intent such as sketch and constraint editing, parametric features, Boolean operations, transformations, inspection, and artifact persistence (Willis et al. 2020; Wu, Xiao, and Zheng 2021; Mallis et al. 2024). Using these transitions as the policy-action unit removes distinctions caused solely by program organization, variable naming, or API expansion, while typed parameters and state preconditions retain the information needed for execution.

State-Grounded Execution The binding map $\beta_\mathcal{H}$ resolves each domain action against the CAD document that exists at that moment and invokes the corresponding backend transition. All accepted actions modify the same persistent document, and stable model-visible identifiers let later actions refer to previously created entities. If parsing, reference resolution, precondition checking, or backend execution fails, the harness keeps the document recoverable and reports the failure. Action semantics are thus grounded in actual CAD transitions rather than textual well-formedness alone.

Decision-Oriented Responses After executing a_t, DeciCAD returns $y_{t+1} = (e_t, o_{t+1})$, where e_t summarizes the execution consequence and o_{t+1} contains additional state information needed for subsequent decisions. The response indicates whether the action executed, which objects were modified, which postconditions hold, and where a failure occurred. When more information is needed, inspection actions retrieve object hierarchy, constraints, parameters, dependencies, local measurements, or terminal artifacts.

A complete state dump repeatedly injects irrelevant context; overly generic feedback can instead map states requiring different recovery actions to the same observation. DeciCAD uses compact consequence-related evidence by default and exposes targeted inspection actions so that the policy acquires additional evidence only when the next decision requires it. Detailed action schemas, reference mechanisms, backend adaptation, and runtime isolation are provided in the appendix.

The resulting DeciCAD interaction loop is

$$h_t \xrightarrow{\pi_\theta} a_t \xrightarrow{\beta_\mathcal{H}, P} s_{t+1} \xrightarrow{O_\mathcal{H}} y_{t+1},$$

which concretely instantiates Eq. (2).

3.4 Harness-Grounded Policy Learning

Execution and learning use the same CAD decision space. Teacher supervision, student rollouts, and environment optimization retain the same action semantics, persistent references, and feedback protocol. For each accepted teacher trajectory under harness $\mathcal{H}$, we reconstruct the causal history $h_{n,t}$ visible immediately before action $a_{n,t}$. The resulting pairs form $\mathcal{D}^\mathcal{H}_{\mathrm{SFT}} = \{(h_{n,t}, a_{n,t})\}_{n,t}$, and we optimize

$$\mathcal{L}_{\mathrm{SFT}}(\theta) = -\mathbb{E}_{(h,a)\sim\mathcal{D}^\mathcal{H}_{\mathrm{SFT}}}[\log \pi_\theta(a \mid h; \mathcal{H})].$$

Later feedback, actions, and terminal evaluation never enter the input, preserving causal conditioning in sequential imitation (Ross, Gordon, and Bagnell 2011; Qin et al. 2024). Recovered failures are retained because their error state, feedback, and corrective action provide genuine CAD recovery supervision.

Uniform decision sampling would overweight long trajectories, repeated inspection, and frequent operations. We therefore stratify causal decisions by task difficulty, design stage, preceding execution outcome, and target action type, preserving modeling, inspection, and recovery decisions under their original histories.

SFT covers teacher-visited CAD states, whereas closed-loop execution reaches states induced by student actions. We therefore sample current-policy rollouts by the same harness and optimize groups containing evaluable terminal geometry. Geometry quality supplies the primary group-relative signal; execution and protocol evidence describe the trajectory but cannot make an all-failure group trainable. Reference-regularized Group Relative Policy Optimization (GRPO) with curriculum sampling instantiates this update (Bengio et al. 2009; Shao et al. 2024; Qian et al. 2025; Gong et al. 2026); reward coefficients, sampling schedules, and optimizer settings are deferred to the appendix.

4 Experiments

Our evaluation follows the methodological chain developed above. We ask whether a semantic CAD decision

Method	CADPrompt				P3D-Text-to-3D				CADInstruct				P3D-Assembly-3D			
	CD↓	F1↑	IoU↑	Avail.↑	CD↓	F1↑	IoU↑	Avail.↑	CD↓	F1↑	IoU↑	Avail.↑	CD↓	F1↑	IoU↑	Avail.↑
Zero-shot	0.32	61.84	74.10	93.50	0.35	37.11	48.15	90.00	0.93	19.28	23.24	**100.00**	0.52	25.40	22.35	71.92
CADSmith	**0.26**	63.70	72.08	**100.00**	0.33	40.74	52.73	**100.00**	0.85	20.11	23.14	**100.00**	0.64	26.27	21.19	96.06
CAD-Assistant	0.41	56.15	65.77	**100.00**	0.45	35.83	53.04	**100.00**	0.98	18.07	14.17	**100.00**	0.64	25.43	21.12	99.51
Claude Code	0.29	65.44	73.86	**100.00**	0.34	40.44	51.20	**100.00**	0.92	20.10	22.62	**100.00**	**0.32**	31.95	23.22	**100.00**
DeciCAD	0.28	**65.83**	**75.12**	**100.00**	**0.25**	**43.63**	**54.22**	**100.00**	**0.84**	**23.49**	**23.35**	**100.00**	0.49	**32.11**	**27.07**	**100.00**
Tokens per task (thousands, input / output)																
Zero-shot	0.36 / 0.22				0.90 / 1.08				0.82 / 0.80				0.76 / 1.66			
CADSmith	8.15 / 1.40				30.15 / 6.91				22.73 / 5.81				59.23 / 19.06			
CAD-Assistant	26.39 / 2.54				36.20 / 5.47				42.40 / 6.41				51.95 / 8.37			
Claude Code	509.66 / 9.37				704.34 / 85.57				702.14 / 110.64				1399.69 / 23.02			
DeciCAD	293.61 / 4.27				390.54 / 10.54				351.13 / 61.38				655.06 / 10.31			

Table 1: Overall performance across four benchmarks, ordered by increasing calibrated difficulty. CD, F1, and IoU are scaled by 100; Avail. is the percentage of tasks producing evaluable geometry. Best results are in **bold** and second-best results are underlined. Token counts in the lower block are averaged per task, are descriptive, and are not ranked.

space preserves the capabilities required for complex modeling, whether its benefits are associated with lower task-irrelevant representation burden, and whether different harness-induced policy spaces exhibit different finite-data learnability and online improvement.

4.1 Experimental Setup

Benchmark. We evaluate harness-level capability, in increasing calibrated difficulty, on CADPrompt (Alrashedy et al. 2025), a hard subset of CADInstruct (Lv and Bao 2025), and two P3D-derived benchmarks: P3D-Text-to-3D from Text2CAD and the multimodal P3D-Assembly-3D from the Fusion 360 Gallery (Khan et al. 2024; Willis et al. 2020).

Difficulty. We assign each task a calibrated difficulty that combines semantic complexity, reference-mesh complexity, and description ambiguity, and divide each benchmark into seven equal-frequency quantiles $Q1$–$Q7$. Quantiles therefore express relative difficulty within a benchmark rather than the same absolute complexity across benchmarks. The full definition and the distribution of each benchmark are given in Appendix D.1 and Figure 10.

Baseline. Using identical task inputs and Claude Opus 4.7 (Anthropic 2026b), DeciCAD is compared with zero-shot CadQuery generation, CADSmith, CAD-Assistant, and Claude Code (Anthropic 2026a). We study policy learning on the standard train–test split of Text2CAD/DeepCAD (Khan et al. 2024; Wu, Xiao, and Zheng 2021), fixing Qwen3.5-9B (Qwen Team 2026) as the student and Claude Sonnet 4.6 (Anthropic 2026c) as the teacher in both the Coding and DeciCAD decision spaces. Generated and reference meshes are normalized, center-aligned, and rigidly registered.

Metric. We report Chamfer Distance (CD), surface F1, volumetric Intersection over Union (IoU), and Geometry Availability. F1, IoU, and CD are computed over valid assets; Geometry Availability measures the percentage of tasks entering geometric evaluation. Harness-level and ablation tables scale CD, F1, and IoU by 100, whereas policy-learning tables report their original values. These point-set and occupancy-based measures follow established 3D reconstruction evaluation practice (Fan, Su, and Guibas 2017; Mescheder et al. 2019). To measure representation cost, we report model-visible input and generated output tokens averaged per task; provider-specific cache statistics are excluded. Across tables, best results are bold and second best are underlined. Policy learning for both DeciCAD and the coding harness used eight NVIDIA H800 GPUs; complete configurations are provided in the appendix.

4.2 Decision Sufficiency and CAD Capability

Table 1 compares the same frontier model under different CAD representations and interaction mechanisms. DeciCAD matches or exceeds the strongest coding harness in geometric fidelity while producing evaluable assets across all four benchmarks.

To check that this advantage is not driven only by easy tasks, we divide each benchmark into seven calibrated difficulty quantiles and report end-to-end F1, which assigns zero to missing or invalid outputs (Figure 4). DeciCAD and Claude Code have similar slopes on CADPrompt and both decline slowly on P3D-Assembly-3D; DeciCAD is flatter on P3D-Text-to-3D, indicating a more stable advantage under complex textual constraints. Across the individual cells, DeciCAD leads 14 of the 28 dataset–quantile combinations and exceeds Claude Code at the hardest quantile of every benchmark, where missing operations and unresolved state ambiguity would be most visible. Semantic CAD actions therefore preserve the transitions needed for complex modeling while supporting more reliable final geometry. The full quantile heatmaps and difficulty distributions are provided in the appendix.

4.3 Validating the Semantic Decision Bottleneck

We test both sides of the Semantic Decision Bottleneck: output scaling characterizes action representation, while observation ablations test whether compression preserves decision-relevant state.

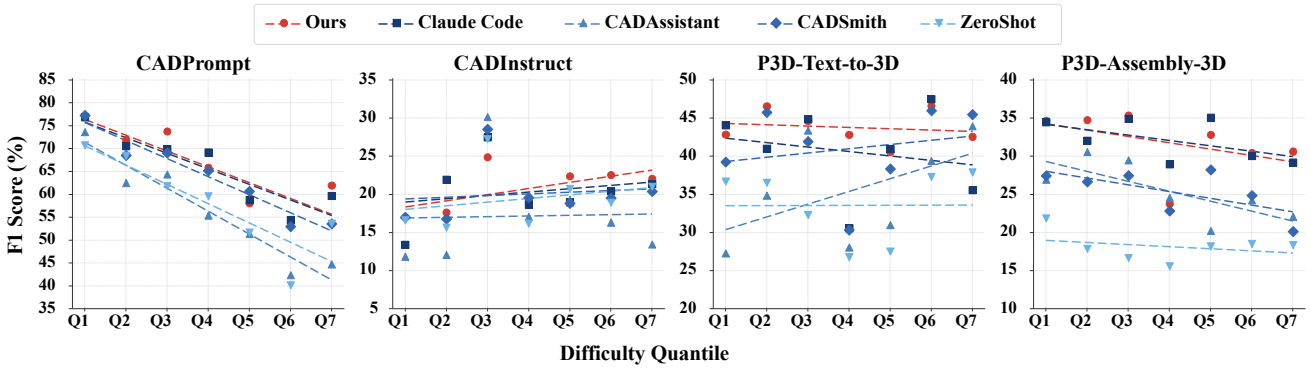

Figure 4: Linear summaries of end-to-end F1 across difficulty quantiles. Dashed lines are least-squares fits over the seven quantiles.

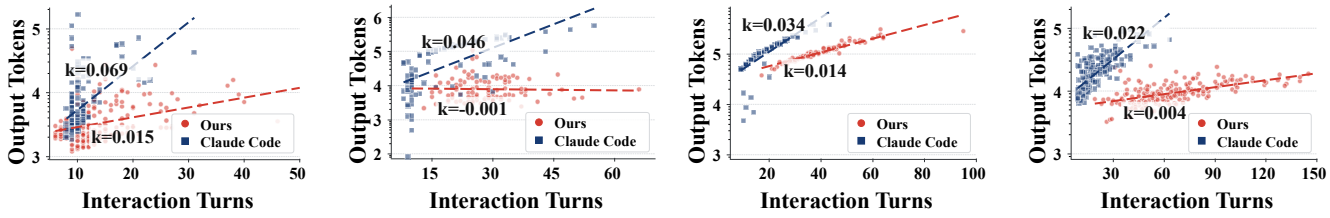

(a) CADPrompt (b) P3D-Text-to-3D (c) CADInstruct (d) P3D-Assembly-3D

Figure 5: Interaction length versus cumulative output tokens. Each point is a complete task trajectory; dashed lines are log-linear fits.

Action-side representation. Token totals in Table 1 depend on trajectory length, feedback history, and caching, and therefore do not by themselves identify the source of the output reduction. Figure 5 compares interaction turns with cumulative output at the task level. Claude Code exhibits consistently positive log-linear slopes of approximately 0.022–0.069, whereas DeciCAD's slopes range from approximately −0.001 to 0.015: its output remains comparatively stable as trajectories lengthen. The reduction does not arise from early termination. DeciCAD trajectories may still be long, but each prediction contains an action and its arguments rather than an expanded implementation, which identifies action representation, rather than shortened reasoning, as the source of the efficiency gain.

Observation-side representation. Table 2 isolates observation sufficiency on P3D-Text-to-3D. More raw state degrades geometry despite increasing context, whereas removing active inspection or coarsening failure semantics hides distinctions needed by subsequent actions. Effective compression therefore retains decision-identifying evidence rather than simply minimizing or maximizing observation volume. Together, both sides support a decision-sufficient space with lower representation burden.

4.4 Harness-Grounded CAD Policy Learning

Frontier models can operate both harnesses, but the corresponding policies need not be equally learnable by a smaller student. We reconstruct causal decision supervision from harness-matched teacher trajectories and compare Qwen3.5-9B policies on the complete Text2CAD test set.

Metric	SDO	SDO-CFF	RSO	EFO
Avail. (%)↑	**100.00**	95.00	55.00	87.00
CD↓	**0.25**	0.44	0.99	0.61
F1↑	**43.63**	32.26	21.62	31.75
IoU↑	**54.22**	43.78	33.06	40.71
Input tokens↓	**390.54**	3,050.00	3,450.00	3,200.00

Table 2: Ablation of observation and feedback design on P3D-Text-to-3D. Variants are Structured Decision-Oriented Observation (SDO; full), SDO with Coarse Failure Feedback (SDO-CFF), Raw State Observation (RSO), and Execution-Feedback-Only Observation (EFO). CD, F1, and IoU are scaled by 100; input tokens are in thousands per task.

Policy acquisition from interaction trajectories. Table 3 compares Base and SFT policies within each induced decision space.

SFT substantially improves the coding policy, showing that the student can acquire program construction, execution, and repair behavior. Under matched supervision, however, the DeciCAD policy produces more evaluable assets and better geometry across all reported metrics. Coding supervision must resolve geometric intent together with program struc-

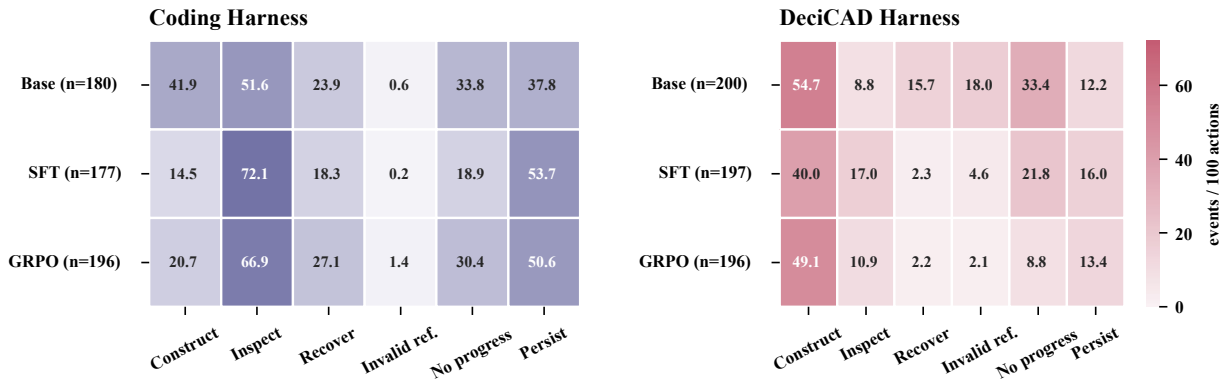

Figure 6: Observable policy behavior across Base, SFT, and GRPO. The left and right panels show the Coding and DeciCAD harnesses. Columns report events per 100 decisions under a shared behavioral mapping; n is the number of trajectories containing parseable action records.

Harness	Policy	Avail. (%)	F1	IoU	CD
Coding	Base	12.5	0.5902	0.5871	0.00972
	SFT	60.5	0.6460	0.6286	0.01130
DeciCAD	Base	61.5	0.6883	0.6487	0.01023
	SFT	**66.5**	**0.7437**	**0.6631**	**0.00782**

Table 3: Closed-loop policy acquisition on the Text2CAD/ DeepCAD test set. Availability is computed over all tasks; F1, IoU, and CD are computed over valid assets.

ture, API choice, and implementation form, whereas DeciCAD binds supervision directly to state-conditioned domain actions. The result supports harness-dependent learnability under finite supervision.

Online improvement from environment feedback. Starting from each SFT policy, we sample student rollouts in its own induced decision space, and apply reference-regularized GRPO with geometry-grounded rewards. The student model, task exposure, samples per task, and evaluator are matched. Improvement is judged by the frozen unseen-task evaluation rather than training reward, trajectory length, or token consumption.

Harness	Policy	Avail. (%)	F1	IoU	CD
Coding	SFT	60.5	0.6460	0.6286	0.01130
	SFT+GRPO	65.5	0.6253	0.5987	0.02253
DeciCAD	SFT	66.5	0.7437	0.6631	0.00782
	SFT+GRPO	**73.5**	**0.7880**	**0.7036**	**0.00570**

Table 4: Online policy improvement under matched training and evaluation. F1, IoU, and CD are computed over valid assets.

Table 4 shows that online feedback improves DeciCAD across availability and all three geometric metrics, whereas the coding policy does not realize the same gains under the matched inference budget, supporting harness-dependent online learnability.

What the policies learn. To compare the learning paths, we map the external effects of both harnesses into six shared categories: Construct, Inspect, Recover, Invalid Reference, No Progress, and Persist, and report events per 100 decisions (Figure 6). Coding SFT increases Inspect from 51.6 to 72.1 events per 100 decisions, while Construct and Recover both decline. GRPO restores part of the construction and recovery behavior, but No Progress rises from 18.9 to 30.4, showing that more repair attempts do not consistently produce geometric progress. In DeciCAD, SFT reduces Invalid Reference and No Progress from 18.0 and 33.4 to 4.6 and 21.8, and GRPO further reduces them to 2.1 and 8.8 while retaining dense construction behavior. Coding supervision therefore mainly changes execution, inspection, and repair behavior, whereas DeciCAD learning more directly changes CAD state transitions and suppresses interactions that fail to advance the model.

5 Conclusion

We formalized Harness-Induced CAD Decision Spaces and instantiated the Semantic Decision Bottleneck in DeciCAD through semantic actions, persistent state, and decision-oriented feedback. Across execution, SFT, and online RL, DeciCAD preserves complex modeling capability while reducing irrelevant burden and enabling more learnable and optimizable policies. Together, these findings show that a harness determines both the CAD operations a model can execute and the distinctions its policy must learn, while environment feedback improves that policy within the same decision representation. This establishes harness design as central to CAD agents. Our study has limitations. The harness comparison uses one CAD backend and one frontier model, policy learning uses one student model on one dataset, and mesh-based metrics do not capture editability, constraint satisfaction, or feature-tree validity. Extending the comparison to further backends, models, and CAD-native correctness measures is left to future work.

References

Abel, D.; Hershkowitz, D. E.; and Littman, M. L. 2016. Near Optimal Behavior via Approximate State Abstraction. In *Proceedings of the 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, 2915–2923. PMLR.

Abel, D.; Umbanhowar, N.; Khetarpal, K.; Arumugam, D.; Precup, D.; and Littman, M. L. 2020. Value Preserving State-Action Abstractions. In *Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics*, volume 108 of *Proceedings of Machine Learning Research*, 1639–1650. PMLR.

Achiam, J.; Adler, S.; Agarwal, S.; Ahmad, L.; Akkaya, I.; Aleman, F. L.; Almeida, D.; Altenschmidt, J.; Altman, S.; Anadkat, S.; et al. 2023. GPT-4 Technical Report. arXiv:2303.08774.

Alrashedy, K.; Tambwekar, P.; Zaidi, Z. H.; Langwasser, M.; Xu, W.; and Gombolay, M. C. 2025. Generating CAD Code with Vision-Language Models for 3D Designs. In *International Conference on Learning Representations*.

Anthropic. 2026a. Claude Code: An Agentic Coding System. Accessed July 29, 2026.

Anthropic. 2026b. Introducing Claude Opus 4.7. Accessed July 29, 2026.

Anthropic. 2026c. Introducing Claude Sonnet 4.6. Accessed July 29, 2026.

Barkley, J.; Loghmani, R.; and Barati Farimani, A. 2026. CADSmith: Multi-Agent CAD Generation with Programmatic Geometric Validation. arXiv:2603.26512.

Bengio, Y.; Louradour, J.; Collobert, R.; and Weston, J. 2009. Curriculum Learning. In *Proceedings of the 26th Annual International Conference on Machine Learning*, 41–48. ACM.

Chen, J.; Fu, D.; Hei, X.; Liu, Y.; Chen, Y.; Xie, J.; Fan, W.; and Cai, Y. 2026. CADMate: Generating CAD Assembly Plan with Geometric Chain-of-Thought and Spatial Physical Rewards. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 18324–18348. San Diego, California, United States: Association for Computational Linguistics.

Doris, A. C.; Alam, M. F.; Nobari, A. H.; and Ahmed, F. 2025. CAD-Coder: An Open-Source Vision-Language Model for Computer-Aided Design Code Generation. arXiv:2505.14646.

Fan, F.; Ni, J.; Yin, X.; Wang, S.; Lu, X.; Zou, Q.; Tong, R.; Tang, M.; and Du, P. 2025. CADDesigner: Conceptual CAD Model Generation with a General-Purpose Agent. arXiv:2508.01031.

Fan, H.; Su, H.; and Guibas, L. J. 2017. A Point Set Generation Network for 3D Object Reconstruction from a Single Image. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2463–2471.

Feng, J.; Huang, S.; Qu, X.; Zhang, G.; Qin, Y.; Zhong, B.; Jiang, C.; Chi, J.; and Zhong, W. 2025. ReTool: Reinforcement Learning for Strategic Tool Use in LLMs. arXiv:2504.11536.

Gong, Y.; Wu, X.; Liu, W.; and Tu, K. 2026. TOOL-CAD: Exploring Tool-Using Large Language Models in Text-to-CAD Generation with Reinforcement Learning. arXiv:2604.07960.

Govindarajan, P.; Baldelli, D.; Pathak, J.; Fournier, Q.; and Chandar, S. 2026. CADmium: Fine-Tuning Code Language Models for Text-Driven Sequential CAD Design. *Transactions on Machine Learning Research*.

Guan, Y.; Wang, X.; Ming, X.; Zhang, J.; Xu, D.; and Yu, Q. 2025. CAD-Coder: Text-to-CAD Generation with Chain-of-Thought and Geometric Reward. arXiv:2505.19713.

Guo, T.; Chen, X.; Wang, Y.; Chang, R.; Pei, S.; Chawla, N. V.; Wiest, O.; and Zhang, X. 2024. Large Language Model Based Multi-Agents: A Survey of Progress and Challenges. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*, 8048–8057.

Kaelbling, L. P.; Littman, M. L.; and Cassandra, A. R. 1998. Planning and Acting in Partially Observable Stochastic Domains. *Artificial Intelligence*, 101(1–2): 99–134.

Karpas, E.; Abend, O.; Belinkov, Y.; Lenz, B.; Lieber, O.; Ratner, N.; Shoham, Y.; Bata, H.; Levine, Y.; Leyton-Brown, K.; Muhlgay, D.; Rozen, N.; Schwartz, E.; Shachaf, G.; Shalev-Shwartz, S.; Shashua, A.; and Tennenholtz, M. 2022. MRKL Systems: A Modular, Neuro-Symbolic Architecture That Combines Large Language Models, External Knowledge Sources and Discrete Reasoning. arXiv:2205.00445.

Khan, M. S.; Sinha, S.; Sheikh, T. U.; Stricker, D.; Ali, S. A.; and Afzal, M. Z. 2024. Text2CAD: Generating Sequential CAD Designs from Beginner-to-Expert Level Text Prompts. In *Advances in Neural Information Processing Systems*, volume 37.

Li, L.; Walsh, T. J.; and Littman, M. L. 2006. Towards a Unified Theory of State Abstraction for MDPs. In *Proceedings of the 9th International Symposium on Artificial Intelligence and Mathematics*.

Li, M.; Zhao, Y.; Yu, B.; Song, F.; Li, H.; Yu, H.; Li, Z.; Huang, F.; and Li, Y. 2023. API-Bank: A Comprehensive Benchmark for Tool-Augmented LLMs. arXiv:2304.08244.

Lv, C.; and Bao, J. 2025. CADInstruct: A Multimodal Dataset for Natural Language-Guided CAD Program Synthesis. *Computer-Aided Design*, 188: 103926.

Mallis, D.; Karadeniz, A. S.; Cavada, S.; Rukhovich, D.; Foteinopoulou, N.; Cherenkova, K.; Kacem, A.; and Aouada, D. 2024. CAD-Assistant: Tool-Augmented VLLMs as Generic CAD Task Solvers? arXiv:2412.13810.

Mescheder, L.; Oechsle, M.; Niemeyer, M.; Nowozin, S.; and Geiger, A. 2019. Occupancy Networks: Learning 3D Reconstruction in Function Space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4460–4470.

Naveed, H.; Khan, A. U.; Qiu, S.; Saqib, M.; Anwar, S.; Usman, M.; Akhtar, N.; Barnes, N.; and Mian, A. 2025. A Comprehensive Overview of Large Language Models. *ACM Transactions on Intelligent Systems and Technology*, 16(5): 106:1–106:72.

Parisi, A.; Zhao, Y.; and Fiedel, N. 2022. TALM: Tool Augmented Language Models. arXiv:2205.12255.

Patil, S. G.; Zhang, T.; Wang, X.; and Gonzalez, J. E. 2023. Gorilla: Large Language Model Connected with Massive APIs. arXiv:2305.15334.

Qian, C.; Acikgoz, E. C.; He, Q.; Wang, H.; Chen, X.; Hakkani-Tür, D.; Tur, G.; and Ji, H. 2025. ToolRL: Reward is All Tool Learning Needs. In *Advances in Neural Information Processing Systems*, volume 38.

Qin, Y.; Liang, S.; Ye, Y.; Zhu, K.; Yan, L.; Lu, Y.; Lin, Y.; Cong, X.; Tang, X.; Qian, B.; Zhao, S.; Hong, L.; Tian, R.; Xie, R.; Zhou, J.; Gerstein, M.; Li, D.; Liu, Z.; and Sun, M. 2024. ToolLLM: Facilitating Large Language Models to Master 16000+ Real-World APIs. In *International Conference on Learning Representations*.

Qwen Team. 2026. Qwen3.5-9B Model Card. Hugging Face. Accessed July 29, 2026.

Ravindran, B.; and Barto, A. G. 2003. SMDP Homomorphisms: An Algebraic Approach to Abstraction in Semi-Markov Decision Processes. In *Proceedings of the 18th International Joint Conference on Artificial Intelligence*, 1011–1018. Morgan Kaufmann.

Ross, S.; Gordon, G. J.; and Bagnell, J. A. 2011. A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, volume 15 of *Proceedings of Machine Learning Research*, 627–635. PMLR.

Schick, T.; Dwivedi-Yu, J.; Dessì, R.; Raileanu, R.; Lomeli, M.; Zettlemoyer, L.; Cancedda, N.; and Scialom, T. 2023. Toolformer: Language Models Can Teach Themselves to Use Tools. In *Advances in Neural Information Processing Systems*, volume 36.

Shao, Z.; Wang, P.; Zhu, Q.; Xu, R.; Song, J.; Bi, X.; Zhang, H.; Zhang, M.; Li, Y. K.; Wu, Y.; and Guo, D. 2024. DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models. arXiv:2402.03300.

Sutton, R. S.; and Barto, A. G. 2018. *Reinforcement Learning: An Introduction*. Cambridge, MA: MIT Press, 2 edition.

Tishby, N.; Pereira, F. C.; and Bialek, W. 2000. The Information Bottleneck Method. arXiv:physics/0004057.

Wang, R.; Yuan, Y.; Sun, S.; and Bian, J. 2025. Text-to-CAD Generation Through Infusing Visual Feedback in Large Language Models. In *Proceedings of the 42nd International Conference on Machine Learning*.

Willis, K. D. D.; Pu, Y.; Luo, J.; Chu, H.; Du, T.; Lambourne, J. G.; Solar-Lezama, A.; and Matusik, W. 2020. Fusion 360 Gallery: A Dataset and Environment for Programmatic CAD Construction from Human Design Sequences. arXiv:2010.02392.

Wu, R.; Xiao, C.; and Zheng, C. 2021. DeepCAD: A Deep Generative Network for Computer-Aided Design Models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 6772–6782.

Xie, T.; Zhang, D.; Chen, J.; Li, X.; Zhao, S.; Cao, R.; Hua, T. J.; Cheng, Z.; Shin, D.; Lei, F.; Liu, Y.; Xu, Y.; Zhou, S.; Savarese, S.; Xiong, C.; Zhong, V.; and Yu, T. 2024. OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments. In *Advances in Neural Information Processing Systems*, volume 37.

Yang, J.; Jimenez, C. E.; Wettig, A.; Lieret, K.; Yao, S.; Narasimhan, K.; and Press, O. 2024. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. In *Advances in Neural Information Processing Systems*, volume 37.

Yao, S.; Zhao, J.; Yu, D.; Du, N.; Shafran, I.; Narasimhan, K.; and Cao, Y. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations*.

Zhou, S.; Xu, F. F.; Zhu, H.; Zhou, X.; Lo, R.; Sridhar, A.; Cheng, X.; Ou, T.; Bisk, Y.; Fried, D.; Alon, U.; and Neubig, G. 2024. WebArena: A Realistic Web Environment for Building Autonomous Agents. In *International Conference on Learning Representations*.

A Overview of the Appendix

This appendix complements the main text with the theoretical, implementation, experimental, and system details needed to connect the Harness-Induced CAD Decision Space to its realization in DeciCAD. The remainder of this document is organized as follows:

1. Appendix B completes the theoretical analysis of decision sufficiency and the Semantic Decision Bottleneck, including the conditions under which action and observation compression preserve optimal CAD decisions.
2. Appendix C explains how DeciCAD instantiates these conditions through semantic CAD actions, state grounding, compact execution feedback, and active hierarchical inspection.
3. Appendix D provides the complete evaluation protocol and additional evidence for harness comparison, causal SFT, and geometry-grounded online RL.
4. Appendix E presents DeciCAD as an interactive Agentic CAD system, covering its user workflow, multi-agent control plane, backend SDK, the official FreeCAD integration, and the supported backends.

Together, these appendices connect the formal decision-space argument to its software realization and to the empirical evidence reported in the main text.

B Theoretical Foundations

The main text models a harness as the decision interface between a language model and a CAD backend. It characterizes decision sufficiency through the actions that the interface can express and the state distinctions retained by its observations, connecting partial observability with value-preserving state abstraction (Kaelbling, Littman, and Cassandra 1998; Ravindran and Barto 2003; Li, Walsh, and Littman 2006; Abel, Hershkowitz, and Littman 2016). This section completes the corresponding theoretical argument and then formalizes the compression preference underlying the Semantic Decision Bottleneck. Figure 7 summarizes the resulting criterion: interface compression is desirable only within the decision-sufficient region.

B.1 Decision Sufficiency

We first restate the objects needed for the analysis. A finite-horizon CAD task g specifies an initial-state distribution μ_g, a horizon H, and a reward $r_g(s_t, u_t, s_{t+1})$. The CAD backend has state $s_t \in \mathcal{S}$, admits backend operations $u_t \in \mathcal{U}$, and evolves according to the standard controlled-transition formalism (Sutton and Barto 2018):

$$s_{t+1} \sim P(\cdot \mid s_t, u_t). \tag{8}$$

A state is *task-reachable* if it can be reached from $s_0 \sim \mu_g$ within H steps by operations admissible for the task.

The model does not act on this backend process directly. A harness

$$\mathcal{H} = (\mathcal{A}_{\mathcal{H}}, \beta_{\mathcal{H}}, O_{\mathcal{H}}) \tag{9}$$

exposes model actions $a_t \in \mathcal{A}_{\mathcal{H}}$, binds them to backend operations through $\beta_{\mathcal{H}}$, and returns model-visible responses through $O_{\mathcal{H}}$. One closed-loop step is

$$\begin{aligned} a_t &\sim \pi_\theta(\cdot \mid h_t; \mathcal{H}), \\ u_t &= \beta_{\mathcal{H}}(a_t, s_t), \\ s_{t+1} &\sim P(\cdot \mid s_t, u_t), \\ y_{t+1} &\sim O_{\mathcal{H}}(\cdot \mid s_t, a_t, s_{t+1}), \end{aligned} \tag{10}$$

where the policy observes only

$$h_t = (g, y_0, a_0, y_1, \ldots, a_{t-1}, y_t). \tag{11}$$

Let $V^*_{g,t}(s)$ be the optimal return available from backend state s at step t, and $Q^*_{g,t}(s,u)$ the return from first applying backend operation u and then acting optimally. The harness is *action-sufficient* when

$$\max_{a \in \mathcal{A}_{\mathcal{H}}} Q^*_{g,t}(s, \beta_{\mathcal{H}}(a,s)) = V^*_{g,t}(s) \tag{12}$$

at every task-reachable state. This condition requires the model-facing action space to retain an executable realization of an optimal backend operation, not every operation or every surface form.

Because s_t is hidden from the policy, define the states consistent with a visible history as

$$\mathcal{B}_{\mathcal{H}}(h_t) := \{s \in \mathcal{S} : \Pr(s_t = s \mid h_t; \mathcal{H}) > 0\}. \tag{13}$$

A history is *reachable* when this set is nonempty. The model actions that are optimal at a particular backend state are

$$\mathcal{A}^*_{\mathcal{H},g,t}(s) := \arg\max_{a \in \mathcal{A}_{\mathcal{H}}} Q^*_{g,t}(s, \beta_{\mathcal{H}}(a,s)). \tag{14}$$

We assume that this set is nonempty at every task-reachable state. The harness is *observation-sufficient* when

$$\bigcap_{s \in \mathcal{B}_{\mathcal{H}}(h_t)} \mathcal{A}^*_{\mathcal{H},g,t}(s) \neq \varnothing \tag{15}$$

for every reachable h_t and $t < H$. Thus, hidden states may share a visible representation precisely when they retain a common optimal next action. A harness satisfying Eqs. (12) and (15) is *decision-sufficient*.

Theorem 1 (Decision sufficiency)**.** *If a harness $\mathcal{H}$ is action-sufficient at every task-reachable state and observation-sufficient at every reachable history, then there exists a policy $\pi^*_{\mathcal{H}}$ depending only on the model-visible history h_t that attains $V^*_{g,t}(s_t)$ at every reachable decision point. Its expected return from the initial-state distribution is therefore*

$$\mathbb{E}_{s_0 \sim \mu_g}\left[V^*_{g,0}(s_0)\right].$$

Thus, the harness preserves the optimal return available through the CAD backend.

Proof. For any reachable history h_t, Eq. (15) permits selecting an action $a^*(h_t)$ from the common intersection. The true backend state s_t is necessarily an element of $\mathcal{B}_{\mathcal{H}}(h_t)$; hence $a^*(h_t)$ is optimal for every backend state consistent with what the model has observed. Action sufficiency in Eq. (12) then ensures that the bound operation $\beta_{\mathcal{H}}(a^*(h_t), s_t)$ attains $V^*_{g,t}(s_t)$. Backward induction from the terminal step shows that the history-dependent policy defined by $a^*(h_t)$ attains the optimal continuation value at every reachable decision point. Taking the expectation over $s_0 \sim \mu_g$ gives the expected value stated above. □

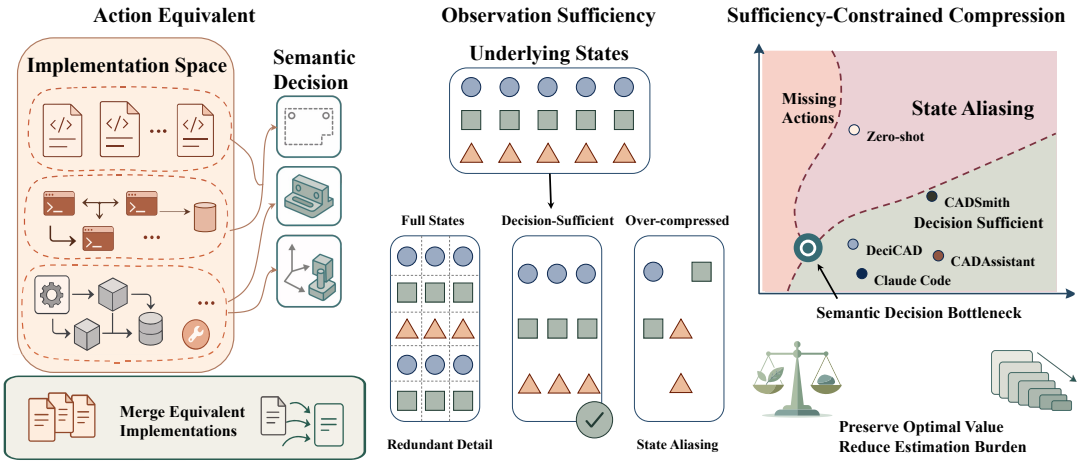

Figure 7: The Semantic Decision Bottleneck compares harnesses by their decision capabilities while preserving optimal CAD performance.

The theorem clarifies what a model-facing CAD interface must retain. The model does not require direct access to the complete backend state or to every low-level command. It requires an executable representation of an optimal operation whenever one is needed, together with observations that distinguish backend states whenever those states demand incompatible next actions.

The two conditions also describe distinct failure boundaries. If action sufficiency fails at a reachable state, no action exposed by the harness can attain $V^*_{g,t}(s)$; increasing model capacity or adding policy data cannot recover an operation absent from the interface. If observation sufficiency fails at a reachable history, the states consistent with that history share no optimal model action. Any policy conditioned only on the visible history must therefore be suboptimal for at least one of them. The resulting expected loss depends on how often the aliased states occur and on the value gap between the chosen and optimal operations.

Missing actions and state aliasing consequently constrain different sides of the harness-induced decision space. The former limits what the policy can execute; the latter limits what it can distinguish. These boundaries determine which action and observation details may be compressed while preserving the optimal CAD policy.

B.2 Semantic Decision Bottleneck

Decision sufficiency specifies what a harness must preserve, but does not uniquely determine how those capabilities are represented to the model. Two harnesses may connect to the same CAD backend and support the same task-level operations while exposing different numbers of implementation choices, different state granularity, and different forms of execution feedback. The Semantic Decision Bottleneck compares these decision-sufficient representations, adapting the compression principle of the information bottleneck to a model-facing action–observation interface (Tishby, Pereira, and Bialek 2000).

Decision-preserving compression. We call a transformation from harness $\mathcal{H}_2$ to harness $\mathcal{H}_1$ a *decision-preserving compression* when it satisfies two conditions. On the action side, model actions may be merged only when, at every task-reachable state where they are admissible, their bound backend operations induce the same distribution over next states and the same task reward for every transition of nonzero probability. The compressed harness must also retain an executable representation of every consequence needed to attain Eq. (12). This removes differences in syntax or implementation without changing a decision-relevant CAD consequence. On the observation side, visible histories may be identified only when all backend states consistent with any of those histories still share an optimal action in the compressed harness. This is precisely the criterion in Eq. (15), evaluated after the histories are identified. The action condition retains an optimal executable consequence at every task-reachable state, while the observation condition retains a common optimal action for every compressed history. They therefore preserve action and observation sufficiency, respectively.

For decision-sufficient harnesses, we write

$$\mathcal{H}_1 \preceq_{\mathrm{SDB}} \mathcal{H}_2 \tag{16}$$

if $\mathcal{H}_1$ can be obtained from $\mathcal{H}_2$ through a decision-preserving compression. The relation expresses a representation preference under fixed decision capability: $\mathcal{H}_1$ does not require the model to distinguish action implementations or visible states that are irrelevant to optimal CAD decisions in $\mathcal{H}_2$.

Decision-preserving compressions can be applied successively: each step retains the action consequences and state distinctions required by Eqs. (12) and (15). The relation therefore supports comparing harnesses by removable semantic structure without selecting a particular action syntax or observation format as canonical.

Compression frontier. The Semantic Decision Bottleneck is the decision-sufficient frontier of this relation. A harness lies on the frontier when every further semantic compression violates at least one sufficiency condition. Further action

compression crosses this frontier when it removes the only representation of an operation required by Eq. (12). Further observation compression crosses it when the common-action intersection becomes empty, aliasing states that demand incompatible next actions. The relation need not select a unique frontier harness: different action vocabularies or observation organizations may be incomparable while preserving the same optimal CAD capability. The bottleneck is therefore neither a unique interface nor the shortest one in characters or tokens; it is the frontier of semantic compression subject to preserving the optimal decision policy.

Equation (16) concerns semantic redundancy rather than token count itself. In a language-model interface, redundant action implementations and state descriptions often increase model input or output, so token use provides an observable measure of representation burden. Nevertheless, semantic compression and token compression are not equivalent. We therefore evaluate a harness jointly by task performance, observation ablations, and token consumption. Nor does the relation by itself guarantee easier optimization for every finite model; the matched SFT and online RL experiments in the main text test whether the induced representation improves policy learning in practice.

Multi-objective harness optimization. A natural direction for future work is to treat harness construction as an explicit multi-objective optimization problem. Let L_{in} and L_{out} denote the numbers of model-facing input and output tokens, respectively, and let N_{inspect} denote the number of targeted inspection operations. One may seek a Pareto-efficient harness according to

$$\min_{\mathcal{H}\in\mathfrak{H}_{\text{suff}}} \mathbf{C}(\mathcal{H}) = (\mathbb{E}[L_{\text{in}}], \mathbb{E}[L_{\text{out}}], \mathbb{E}[N_{\text{inspect}}]) \,, \qquad (17)$$

where $\mathfrak{H}_{\text{suff}}$ is the set of harnesses satisfying both action and observation sufficiency, and the expectations are taken over the task distribution and policy-induced trajectories. Different deployment settings may scalarize these objectives using different weights: for example, interactive systems may penalize inspection latency more heavily, whereas resource-constrained deployments may prioritize context and generation length. More broadly, the harness and policy could be optimized jointly, subject to semantic sufficiency constraints. We leave automated exploration of this harness–policy Pareto frontier to future work.

C Implementation Details

The main text characterizes a decision-sufficient CAD harness through action and observation sufficiency. This section follows an actual modeling decision through DeciCAD to show how these conditions are instantiated. The model first expresses geometric intent in a semantic action space, the harness grounds that intent in the current CAD state, and the environment returns the execution evidence needed for the next decision. Together, these mechanisms determine the CAD policies that the model can learn and execute.

C.1 Semantic CAD Action Space

A modeling policy first requires an action vocabulary capable of expressing design intent. DeciCAD represents each model output as a CAD action with geometric arguments. The vocabulary covers documents and artifacts, sketches and constraints, solid features, Booleans and patterns, edge treatments, object placement, assembly, and state inspection. Each action exposes the variables that determine the CAD state transition. An extrusion, for example, specifies its profile, material mode, extent, and direction, while object resolution and CAD API invocation remain within the harness. This feature-oriented organization follows structured construction histories used in parametric CAD corpora while exposing them as executable interaction steps (Wu, Xiao, and Zheng 2021; Willis et al. 2020; Khan et al. 2024).

The following abstract interface summarizes the model-visible action space. `Ref<T>` denotes a stable environment-generated object reference that subsequent actions can reuse; the expression after `->` describes the geometric postcondition returned after execution.

```
CADAction ::=
    Document {Create, Open, Save, Export}
  | Sketch   {Create, AddGeometry, Constrain,
              Pattern, Trim, Project}
  | Solid    {CreateBody, Extrude, Revolve, Sweep,
              Loft, Boolean, Repeat, Fillet,
              Chamfer, Shell}
  | Entity   {SetPlacement}
  | Assembly {Create, AddComponent, Ground, AddMate,
              RemoveMate, Solve}
  | Inspect  {List, Get}

Extrude(
    profile:   Ref<Sketch>,
    mode:      Add | Cut,
    extent:    Distance(d) | ThroughAll,
    direction: SketchNormal | Vector3D
) -> {
    feature:     Ref<SolidFeature>,
    body:        Ref<Body>,
    solid_count: Int,
    volume_mm3:  Float
}
```

This representation identifies behaviorally equivalent choices of variable names, control flow, and API syntax as the same action, while retaining the parameters that determine geometric outcome. Stable references compose these actions: a sketch can produce a solid feature; that solid can participate in Booleans, transformations, and edge treatments; and parts can enter an assembly and receive constraints. Each action describes a local geometric transition, while an action sequence incrementally constructs the full design. For the task distribution considered in this work, this compositional closure allows long-horizon modeling policies to remain within one semantic action space, instantiating action sufficiency in the main text.

C.2 Grounding Semantic Actions in CAD State

Semantic actions specify what the model can select; state grounding determines what a selection means in the current design. A point, direction, or rotation has a unique geometric meaning only relative to an explicit coordinate frame. DeciCAD therefore interprets all action arguments under a shared spatial contract. Linear quantities use millimeters and angular quantities use degrees. Documents use a right-handed world frame with (+X) right, (+Y) back, and (+Z) up, and

rotations follow the right-hand rule. Two-dimensional geometry lies in the right-handed local frame of its sketch, while solid placement explicitly selects world or parent coordinates. System instructions state these rules, typed arguments mark the space of each value, and the execution layer reports authoritative frames read from the actual CAD placement.

```
SpatialContract {
    length_unit = millimeter
    angle_unit  = degree
    world_frame = right_handed(+X right, +Y back, +Z
      up)
    rotation    = right_hand_rule(axis, angle)

    Point2D   := coordinate in Ref<Sketch>.
     local_frame
    Point3D   := coordinate in world_frame
    Placement(space = World | Parent)
    Frame     := {origin, unit_x, unit_y, unit_z}
}

CreateSketch(support, plane)
    -> {sketch: Ref<Sketch>, local_frame: Frame}

SetPlacement(target, placement, space)
    -> {target, effective_placement, space}
```

The spatial contract fixes the meaning of arguments, while object references fix their targets. Documents, entities, and reusable topological elements have stable identifiers returned by creation or inspection and reused unchanged in later actions. The harness resolves these references in the current document, executes the native CAD operation, and verifies action-specific postconditions. A valid result is committed to the persistent document so that later actions operate on the resulting state; a failed result preserves the previous state or marks the execution outcome as unknown.

```
Execute(action, state s_t):
    args       <- validate(action, SpatialContract)
    objects    <- resolve(action.references, s_t)
    transaction.begin()
    native_out <- apply(action.name, args, objects)
    evidence   <- verify_postcondition(native_out)

    if evidence is valid:
        transaction.commit()
        return Success(evidence), s_{t+1}
    else:
        transaction.rollback()
        return Failure(recovery_evidence), s_t
```

Action grounding is therefore more than a lookup from a tool name to a backend function. The shared spatial contract determines parameter semantics, stable references determine the operated objects, and transactions with postcondition checks determine whether an action actually took effect. Every successful action in the model history corresponds to a verified transition in the persistent CAD document. Together, these constraints implement the binding map defined in the main text.

C.3 Decision-Oriented Environment Observations

After an action executes, the model must continue from the resulting environment state. Expanding the complete CAD document repeatedly introduces objects and properties irrelevant to the current decision, whereas reporting only call-level success cannot support geometric judgment or recovery. DeCiCAD separates these needs: passive execution feedback describes the state transition that just occurred, while active hierarchical inspection retrieves local state needed by the next decision.

Passive Execution Feedback The harness returns success or failure after every action. A success response contains affected objects, action-specific postconditions, and the document generation. A failure response contains an error category, the relevant argument or object, and a recovery direction. Invalid input, missing references, geometric failure, conflict, timeout, and unknown outcome remain distinct because they call for different next actions.

```
ActionResult :=
    Success {
        affected:       [StableRef],
        postcondition:  ActionSpecificEvidence,
        document_state: Generation
    }
  | Failure {
        category: InvalidInput | MissingObject |
                  InvalidState | GeometryFailure |
                  Conflict | Timeout |
                  BackendUnavailable |
     OutcomeUnknown,
        location: Optional<ArgumentPath | StableRef
     >,
        message:  ShortExplanation,
        recovery: RetryWithArguments | InspectState
     |
                  ChooseAnotherAction
    }
```

For example, a successful extrusion returns its generated feature, owning body, and solid count; invalid arguments identify the responsible field; and an unknown outcome directs the model to inspect the document before acting again. These distinctions determine whether the policy should continue modeling, revise an argument, or first verify state. Passive feedback therefore retains causal evidence for local continuation and recovery while keeping each step compact.

Active Hierarchical Inspection Passive feedback focuses on the action just completed, but later decisions may depend on a broader portion of the design. An edge treatment may require locating topological edges with particular geometry, while an assembly mate may require checking component placement and connectivity. For these cases, `InspectList` first discovers relevant state. The model selects a bounded collection—such as entities, children, topology, sketch geometry, constraints, dependencies, or assembly interference—and receives stable references with compact summaries through filtering and pagination. `InspectGet` then reads the canonical detail of one precise reference, including identity, placement or local frame, bounding box, shape or constraint summary, relations, and validity.

```
InspectList(
    collection: Documents | Entities | Children |
                Topology | SketchGeometry |
                SketchConstraints | Dependencies |
                Interferences | ArtifactComponents,
    scope:  Optional<Ref<Document | Entity> |
     ArtifactPath>,
    filter: Optional<{name_pattern, entity_type}>,
    offset: Int,
```

```
    limit:  Int
) -> {
    items: [{ref, kind, compact_summary}],
    generation: Generation,
    next_offset: Optional<Int>
}

InspectGet(
    target: Ref<Document | Entity | Face | Edge |
     Vertex>
            | ArtifactPath
) -> {
    identity, placement_or_frame, bounding_box,
    shape_or_constraint_summary, relation_summary,
     validity
}
```

The two inspection levels answer, respectively, which objects are relevant to the current decision and what the authoritative state of one target is. A policy can first list the topological edges of a solid, select candidates from their summaries, retrieve exact geometric properties, and reuse the same reference in a fillet action. Separating discovery from retrieval bounds each observation while allowing the model to recover state distinctions that change the next action. Passive feedback and active inspection thereby instantiate observation sufficiency in the main text.

C.4 Implementing the Semantic Decision Bottleneck

Across this loop, DeciCAD coordinates compression of model-facing actions and observations. Semantic actions merge behaviorally equivalent implementation forms, while the spatial contract, stable references, and execution verification ensure that the compressed actions still denote well-defined CAD transitions. Passive feedback omits the complete internal state, and active inspection restores local distinctions when a decision requires them. Compression therefore remains constrained by action and observation sufficiency.

Together, these mechanisms instantiate the Semantic Decision Bottleneck from the main text. The model generates CAD decisions rather than their expanded implementations and observes state evidence supporting the next decision rather than the full internal representation of the CAD system. Coordinating actions, grounding, and observations in the same closed loop reduces representation burden without removing the decision paths required for task completion and error recovery.

In the concrete code organization, these mechanisms are separated into three connected components. The harness runtime organizes the decision loop across model invocation, action execution, and feedback. An interface-contract layer defines semantic actions, geometric representations, and execution results. A backend execution layer binds those contracts to native operations in a specific CAD system. This organization allows the interaction policy, the model-visible interface, and CAD execution logic to evolve independently. A new CAD backend implements the same action and feedback contracts while preserving the model-facing decision space.

D Supplementary Experiments

This section completes the evaluation protocol used in the main text and further analyzes how the harness affects CAD capability, interaction representation, and policy learning. The main text reports aggregate geometry results, difficulty trends, output scaling, observation ablations, policy performance, and learned behavior. Here we provide the protocols and evidence not shown in the main text: difficulty calibration and per-quantile geometry, the geometry–context trade-off, qualitative assets, supervision construction, online-RL design, and a per-task trajectory example. All analyses retain the frozen tasks, models, and evaluator from the main experiments.

D.1 Complete Experimental Setup

We first specify the common benchmarks, baselines, difficulty calibration, geometry evaluator, and policy-learning protocol underlying the subsequent analyses.

Benchmarks and baselines. The harness-level evaluation covers four CAD settings. CADPrompt uses its complete evaluation set. We manually reviewed CADInstruct and P3D-Text-to-3D and fixed their most challenging subsets before running any method. P3D-Assembly-3D comprises multimodal assembly tasks that provide both textual requirements and target images. Together, the benchmarks cover parametric part modeling, complex text-conditioned geometry, and multi-component assembly (Alrashedy et al. 2025; Lv and Bao 2025; Khan et al. 2024; Willis et al. 2020).

Benchmark	Samples	Input	Task type
CADPrompt	200	Text	Part modeling
P3D-Text-to-3D	100	Text	Part modeling
CADInstruct	100	Text	Part modeling
P3D-Assembly-3D	203	Text + image	Assembly modeling

Table 5: Statistics of the benchmarks used in our experiments.

Table 5 summarizes the input modality and task type of each benchmark. Together they span text-conditioned part modeling and multimodal assembly with progressively richer geometric and relational requirements.

We compare zero-shot CadQuery, CADSmith, CAD-Assistant, Claude Code, and DeciCAD. Zero-shot emits a complete CadQuery program in one response. CADSmith coordinates planning, program generation, and verification across multiple agents. CAD-Assistant iteratively revises a design through a Python interpreter and external CAD operations. Claude Code provides a general coding harness for CadQuery generation, execution, and repair. DeciCAD instead exposes semantic CAD actions and object-level environment feedback (Mallis et al. 2024; Barkley, Loghmani, and Barati Farimani 2026; Anthropic 2026a). Harness-level experiments use Claude Opus 4.7, identical task inputs and geometry evaluation, and only the system instructions required by each interface (Anthropic 2026b).

Difficulty calibration. Aggregate means alone do not locate a method's boundary on complex tasks. For task x,

(a) CADPrompt

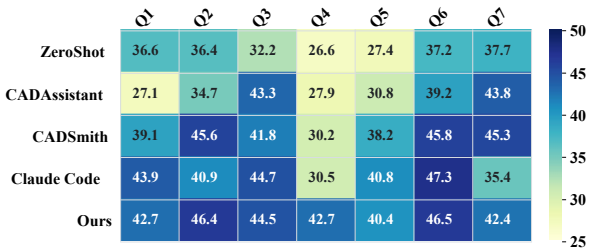

(b) P3D-Text-to-3D

(c) CADInstruct

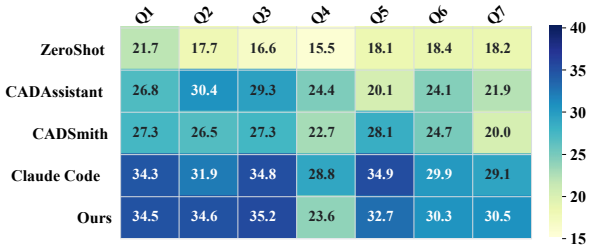

(d) P3D-Assembly-3D

Figure 8: Mean end-to-end F1 across calibrated within-dataset difficulty quantiles.

let $s_{\text{sem}}(x)$ denote semantic complexity, $s_{\text{mesh}}(x)$ reference-mesh complexity, and $a(x)$ description ambiguity. We first define geometric complexity and then combine it with ambiguity:

$$g_{\text{geo}}(x) = \alpha s_{\text{sem}}(x) + \beta s_{\text{mesh}}(x), \tag{18}$$

$$d(x) = \lambda_1 g_{\text{geo}}(x) + \lambda_2 a(x) + \lambda_3 g_{\text{geo}}(x) a(x). \tag{19}$$

The coefficients are fixed across benchmarks. The interaction term assigns higher difficulty when complex geometry is also described ambiguously. We compute $d(x)$ by the same rule for every sample and form seven equal-frequency quantiles $Q1$–$Q7$ within each dataset. Quantiles thus represent relative difficulty within a benchmark rather than identical absolute complexity across benchmarks.

Component	Setting
Student / teacher	Qwen3.5-9B / Claude Sonnet 4.6
SFT LoRA	rank 32, alpha 64, dropout 0.05
SFT learning rate	1×10^{-5}, cosine
SFT context	65,536 tokens
Checkpoint interval	25 optimizer steps
GRPO LoRA	rank 32, alpha 64
Rollout group	16 tasks $\times$ 8 samples
Rollout context	131,072 tokens
Sampled-sequence cap	65,536 tokens
Hardware	Eight NVIDIA H800 GPUs

Table 6: Shared policy-learning configuration.

Geometry and execution evaluation. Generated and reference assets are converted to a common mesh representation, scale-normalized, center-aligned, and rigidly aligned by ICP. Chamfer Distance (CD) and surface F1 use 10,000 uniformly sampled surface points per asset; the F1 nearest-neighbor threshold and IoU voxel resolution are both 0.01. Geometry Availability is the fraction of tasks producing a nonempty, parseable asset that enters geometry evaluation. F1, IoU, and CD are computed over those valid assets. For the difficulty-stratified Figures 8 and 4, we additionally report end-to-end F1: missing, unparsable, or invalid outputs are assigned F1=0. Availability and conditional geometry quality are kept separate rather than collapsed into a single success score. The surface-distance and occupancy measures follow established point-set and implicit-geometry evaluation practice (Fan, Su, and Guibas 2017; Mescheder et al. 2019).

Table 6 reports the shared policy-learning configuration. The two harness tracks use the same student and teacher models, LoRA capacity, context limits, checkpoint schedule, and hardware.

Policy-learning protocol. Policy learning uses the Text2CAD/DeepCAD 800/200 train–test split. Claude Sonnet 4.6 generates teacher trajectories independently in the coding harness and DeciCAD; Qwen3.5-9B is the shared student. Both tracks use the same task descriptions, reference geometry, CAD environment, and evaluator. The final test set is excluded from trajectory selection, balancing, hyperparameter adjustment, reward design, and checkpoint selection.

SFT and online RL use BF16 LoRA. The SFT tracks share learning rate, batch configuration, context length, and LoRA settings. GRPO begins from the selected SFT checkpoint for each harness and samples student rollouts through the same interface used during supervision.

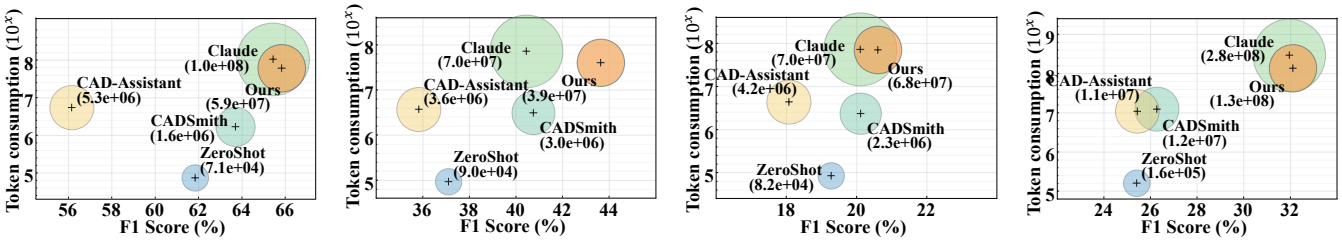

(a) CADPrompt (b) P3D-Text-to-3D (c) CADInstruct (d) P3D-Assembly-3D

Figure 9: Geometry–context trade-off across benchmarks. Each point reports mean F1 and cumulative model-visible input tokens for one method.

D.2 Extended Harness Comparison

We extend the aggregate comparison in the main text along task difficulty, model-visible context, and representative output geometry.

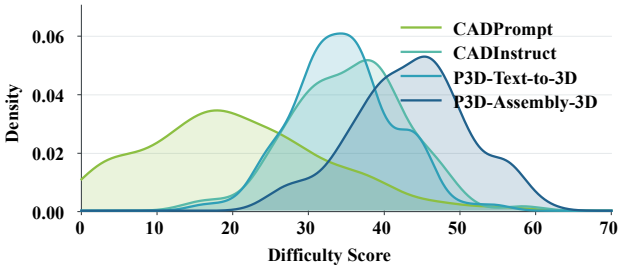

Figure 10: Calibrated difficulty distributions of the four benchmarks.

Difficulty distribution and stratified performance. Figure 10 shows the calibrated difficulty distributions. CADPrompt concentrates in the lower range, CADInstruct and P3D-Text-to-3D shift toward greater complexity, and P3D-Assembly-3D occupies the highest range. This ordering reflects the additional need to preserve component identity, relative placement, and cross-object relations in assembly.

Figure 8 expands the aggregate F1 results across $Q1$–$Q7$. DeciCAD leads 14 of the 28 dataset–quantile cells, more than any baseline, and exceeds Claude Code at $Q7$ on every benchmark. It ranks first in five of seven P3D-Assembly-3D quantiles and leads from $Q5$ through $Q7$ on CADInstruct. On P3D-Text-to-3D, its F1 remains ahead of Claude Code at both $Q4$ (42.7 versus 30.5) and $Q7$ (42.4 versus 35.4).

The aggregate advantage is therefore not driven only by easy tasks. If the semantic action space omitted operations required by complex designs, performance would deteriorate systematically at high difficulty. Its competitive $Q7$ behavior instead indicates that compressing implementation-level choices retains the necessary CAD transitions. CADInstruct remains non-monotonic, so difficulty serves as a stratification variable rather than a deterministic predictor of failure.

Figure 4 in the main text summarizes these quantile trends with linear fits. Because CADInstruct is non-monotonic, its slope is used only as a compact summary of the heatmap.

Geometry–context trade-off. Figure 9 jointly compares geometric quality and model-visible input tokens. On the three text-to-CAD benchmarks, DeciCAD occupies a stronger performance–cost region than Claude Code: it attains higher F1 while consuming substantially less input context. The geometric advantage therefore cannot be attributed to exposing the model to a larger interaction history. P3D-Assembly-3D is more demanding for both harnesses because multi-object references and repeated state updates enlarge the context, yet DeciCAD retains a competitive trade-off. Together with the output-side analysis below, these results separate representation efficiency from trajectory truncation or reduced task coverage.

Output scaling with interaction length. Figure 5 in the main text compares interaction turns with cumulative output at the task level and shows that the output reduction does not arise from early termination.

Decision allocation. We group decisions as Planning, Modeling, or Verification according to their external effect. Figure 11 shows that DeciCAD consistently devotes a larger share of observable decisions to Modeling across datasets and difficulty levels, whereas Claude Code devotes a larger combined share to Planning and Verification.

Qualitative comparison. Figure 12 shows representative assets from P3D-Text-to-3D and CADPrompt. Zero-shot and lightweight specialized agents more often omit local features or produce unevaluable artifacts. Claude Code usually recovers the global shape but may retain geometric-detail errors despite executable code. DeciCAD more consistently preserves the requested profiles, openings, and repeated features, matching the aggregate trends in F1, IoU, and Availability.

D.3 Harness-Grounded SFT

We next describe how closed-loop teacher trajectories are converted into causal supervision and how the resulting policies are selected and evaluated.

Causal decision supervision and data composition. Policy supervision must preserve the decision conditions under which each action was produced. For every teacher action a_t, we reconstruct the model-visible history h_t immediately before that action, including the system instruction, task, preceding model actions, tool results, and available tool schema.

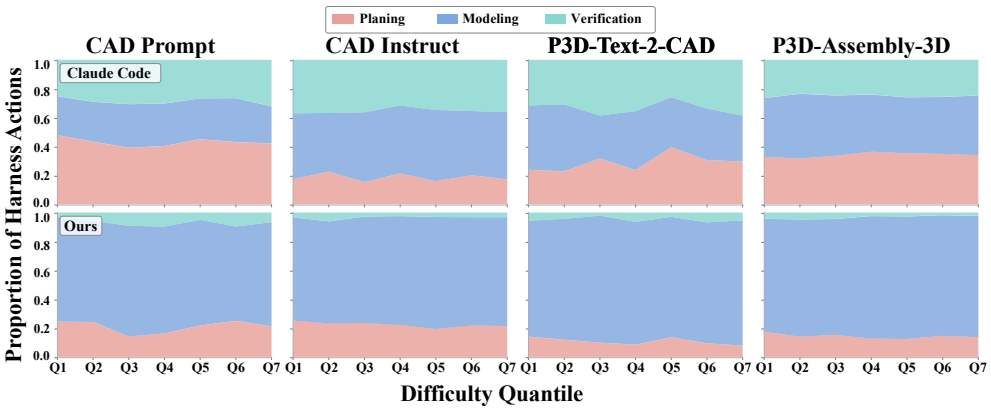

Figure 11: Allocation of observable decisions to Planning, Modeling, and Verification across datasets and difficulty quantiles. The upper and lower rows show Claude Code and DeciCAD, respectively.

The target is the teacher's next coding response or semantic CAD action. When a trajectory contains a local failure, the failure result remains in the history and supervision is applied to the subsequent recovery decision. This causal construction follows the principle that interactive policies must be trained on histories induced by preceding actions rather than on future trajectory information (Ross, Gordon, and Bagnell 2011; Qin et al. 2024).

Both datasets originate from the official Text2CAD training split. Claude Sonnet 4.6 executes the same tasks independently through the Coding and DeciCAD harnesses, and the resulting assets are evaluated against the same target geometry. Geometry- and protocol-audited trajectories are then expanded into causal decision examples (Khan et al. 2024; Anthropic 2026c; Qwen Team 2026). The DeciCAD dataset contains 642 audited tasks, divided by task identity into 542 training and 100 development tasks. Each task retains at most 11 regular and three corrective decisions covering initialization, core modeling, inspection, recovery, and artifact persistence. The Coding dataset contains 499 geometry-valid tasks, of which 471 are used for training and 28 for development, yielding 3,651 training targets and 207 development targets; 336 training targets follow an observed execution error.

Harness	Audited tasks	Train / dev.	Target cap	Recovery targets
Coding	499	471 / 28	12	Up to 4 observed
DeciCAD	642	542 / 100	11 + 3	Up to 3 observed or controlled

Table 7: SFT supervision obtained through the two harnesses. Both tracks use the same student model and optimization protocol.

The two tracks share the source tasks, teacher and student models, causal truncation, context limit, and optimization protocol, while retaining the natural granularity of their respective decision spaces. DeciCAD examples are stratified by task difficulty, CAD stage, accumulated document state, preceding action and outcome, history pattern, and target action class. Generic query and computation actions are capped, whereas geometry-changing, inspection, and recovery actions retain their causal histories. Coding data follow the same task-level filtering principle and prioritize genuine recovery responses.

Table 7 records the audited task coverage and target construction used for the two SFT tracks. These counts reflect the executable supervision available in each harness-induced decision space.

Checkpoint selection and learned policies. Both tracks save a checkpoint every 25 optimizer steps. Training and token-level validation loss monitor optimization, while the final checkpoint is selected on frozen development tasks using closed-loop Availability, F1, IoU, and CD.

SFT equips both Qwen3.5-9B policies with closed-loop CAD capability, but the learned outcomes differ. Coding SFT raises Availability from 12.5% to 60.5%, F1 from 0.5902 to 0.6460, and IoU from 0.5871 to 0.6286, while CD changes from 0.00972 to 0.01130. The student therefore acquires program construction, execution, and repair, although local geometric error remains after the large coverage expansion. DeciCAD SFT raises Availability from 61.5% to 66.5%, F1 from 0.6883 to 0.7437, and IoU from 0.6487 to 0.6631, while reducing CD from 0.01023 to 0.00782. Semantic CAD supervision thus improves both asset coverage and the geometry of valid outputs. The aggregate stage-wise geometry results and the resulting policy behavior are reported in the main text; below we detail the online optimization that produces them.

D.4 Geometry-Grounded Online RL

Starting from the selected SFT policies, we detail the curriculum, reward, and update constraints used for online optimization before analyzing the learned behavior.

Two-stage online optimization. Online RL continues from states induced by the current student. Training tasks

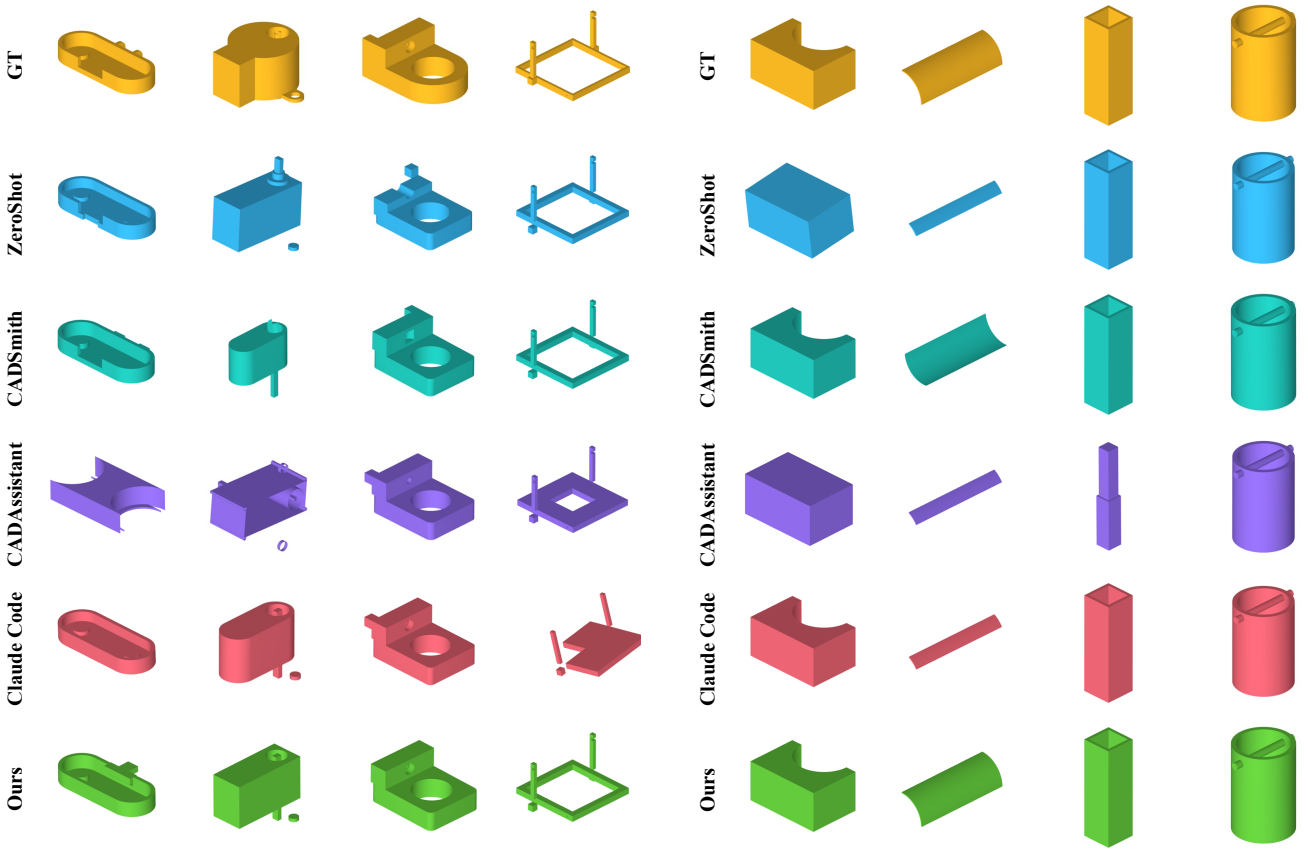

Figure 12: Qualitative comparison on P3D-Text-to-3D (left) and CADPrompt (right).

serve four roles: `completion_frontier` tasks frequently fail to form a valid asset; `recovery` tasks expose locally recoverable failures; `quality_frontier` tasks complete but admit geometric improvement; and `anchor` tasks are solved reliably and constrain forgetting. The first stage emphasizes completion and recovery to expand valid-asset coverage. Its policy becomes the frozen parent of a second, geometry-oriented stage. The final second-stage iteration samples 16 tasks—12 quality-frontier, two completion-frontier, and two anchor tasks—with eight rollouts per task. This ordering combines curriculum learning with group-relative online optimization used in tool-learning policies (Bengio et al. 2009; Shao et al. 2024; Qian et al. 2025; Gong et al. 2026).

Each group contains three conservative, three standard, and two exploratory rollouts with $(T, p) = (0.65, 0.90)$, $(0.90, 0.95)$, and $(1.10, 0.98)$, respectively. Interaction budgets increase with task stage from 32, 40, and 48 to 64 decisions. Both stages retain the same harness, tool contracts, and terminal evaluator; the second stage changes policy preference within the existing CAD decision space.

Geometry and process reward. For a rollout with valid terminal geometry, the normalized geometry reward is

$$R_{\text{geo}} = w_F F1^{p_F} + w_I IoU^{p_I} + w_C \exp\left(-\frac{CD}{\tau_C}\right), \quad (20)$$

where $w_F = 0.45$, $w_I = 0.35$, $w_C = 0.20$, $p_F = p_I = 2$, and $\tau_C = 0.002$. Squaring retains discrimination in the high-quality region, and the exponential maps unbounded CD to $[0, 1]$. Missing or unparsable assets, and assets with $R_{\text{geo}} < 0.15$, receive zero geometry reward.

Process evidence describes observable CAD progress. At decision t, let v_t, c_t, r_t, and b_t indicate a valid call, a new state change, recovery from the preceding failure, and repeated failure or no-progress regression. Then

$$R_{\text{proc},t} = \text{clip}\left(0.20v_t + 0.55c_t + 0.25r_t - 0.45b_t, -1, 1\right), \quad (21)$$

and trajectory process reward is its decision-wise mean. Protocol reward is the valid-call fraction. The second-stage reward is

$$R = \begin{cases} 0, & R_{\text{geo}} = 0, \\ \eta R_{\text{geo}} + \rho R_{\text{proc}} + \xi R_{\text{proto}}, & R_{\text{geo}} > 0, \end{cases} \quad (22)$$

with $(\eta, \rho, \xi) = (0.94, 0.05, 0.01)$. Zero-geometry gating places the terminal asset first, while the smaller process and protocol terms distinguish no-progress behavior, invalid calls, and executable modeling sequences.

Group-relative advantage and update gating. Eight rollouts of one task form a GRPO group. A group enters optimization only if at least one rollout has nonzero geometry

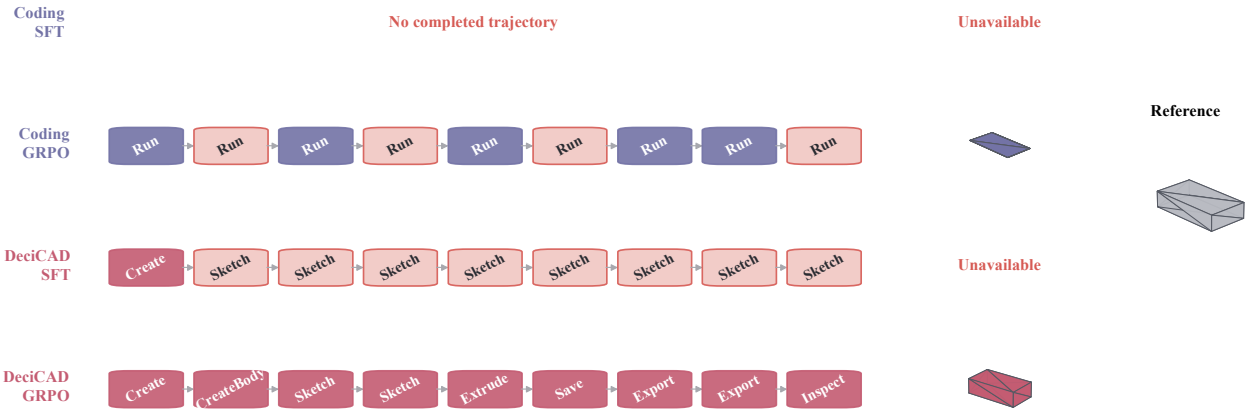

Figure 13: Cross-harness policy trajectories and terminal geometry on the same CAD task. Rows show Coding SFT, Coding GRPO, DeciCAD SFT, and DeciCAD GRPO. Nodes summarize external calls and their execution status; the right side shows generated assets and the reference geometry. Unavailable indicates that no evaluable asset was produced.

reward and its reward standard deviation is at least 0.02. This excludes all-failure groups whose trajectories differ only in procedural score.

The second stage combines completion and quality advantage at $0.15 : 0.85$. Quality-frontier and anchor candidates must satisfy an absolute quality floor and remain within the tolerance of the frozen parent anchor. Each group retains at most one positive and one negative trajectory, balancing task influence. The policy objective uses learning rate 3×10^{-9}, two-sided clip 0.03, and token-level KL coefficient $\beta = 0.10$ to the frozen first-stage parent. A stop action receives positive credit only when geometry exceeds both an absolute threshold and a high within-group quantile; premature termination receives negative credit.

The learned second-stage LoRA increment is finally applied to the parent at strength 0.375. We preserve the adapter weights and change LoRA alpha from 64 to 24, scaling the learned offset by $24/64 = 0.375$. This continuous control retains the parent's coverage and recovery behavior while adding the quality preference learned in the second stage.

Online optimization results. DeciCAD GRPO raises Availability from 66.5% to 73.5%, F1 from 0.7437 to 0.7880, and IoU from 0.6631 to 0.7036, while reducing CD from 0.00782 to 0.00570. Coding GRPO raises Availability from 60.5% to 65.5%, but F1 falls from 0.6460 to 0.6253, IoU from 0.6286 to 0.5987, and CD increases from 0.01130 to 0.02253. Thus, identical terminal geometry feedback expands executable coverage in both harnesses, while stable geometry improvement depends on how credit propagates through the harness-induced decision space.

Observable capabilities in rollouts. Figure 6 in the main text reports policy behavior under six shared categories. Construct creates or modifies geometric implementation; Inspect executes code, reads state, or checks a result; Recover changes action after an explicit failure and restores progress; Invalid Reference and No Progress record invalid targets and repeated non-progress; and Persist saves or exports the terminal asset. Events are counted per 100 external decisions so that trajectories of different lengths are comparable.

Figure 13 grounds the aggregate difference in one shared task. Coding SFT produces no completed trajectory. Coding GRPO emits an asset after repeated runs and repairs, but the output collapses to a thin plate unlike the reference structure. DeciCAD SFT enters the domain modeling process but repeated failed Sketch actions leave no evaluable asset. DeciCAD GRPO instead completes a chain of creation, sketching, extrusion, persistence, export, and inspection, producing terminal geometry consistent with the reference. The example clarifies that recovery count alone does not imply useful recovery: subsequent actions must continue to advance the actual CAD state.

Aggregate geometry, behavior statistics, and trajectory-level evidence jointly define the scope of online improvement. GRPO increases valid-asset coverage and reduces invalid references and repeated non-progress, while small features in long-horizon tasks remain limited by terminal reward resolution. F1 and IoU measure surface and volumetric agreement, and CD adds local distance, but feature count, topological correspondence, repeated structures, and constraint relations can still differ under similar terminal scores. We therefore use closed-loop geometry as the primary evidence and the observed behavior and assets to show how improvement forms within each harness decision space.

E DeciCAD: An Interactive Agentic CAD System

DeciCAD is distributed as a terminal-native Agentic CAD application. It organizes model providers, specialized agents, CAD backends, and project resources in one workspace, allowing a user to create, modify, inspect, assemble, and export native CAD documents through natural-language interaction. This section presents the resulting product workflow, its sup-

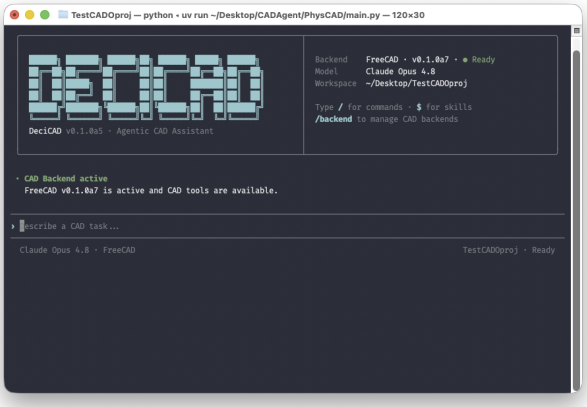

(a) Project workspace and active CAD runtime.

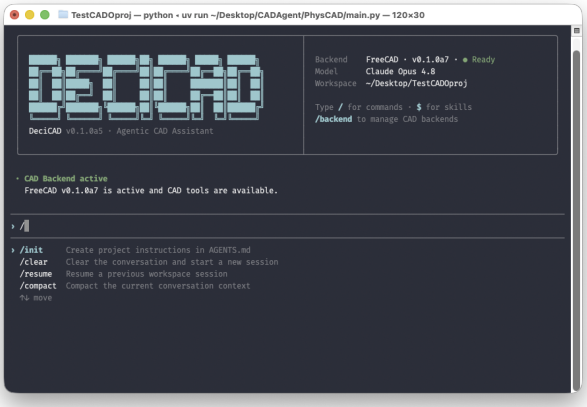

(b) Session commands available from the composer.

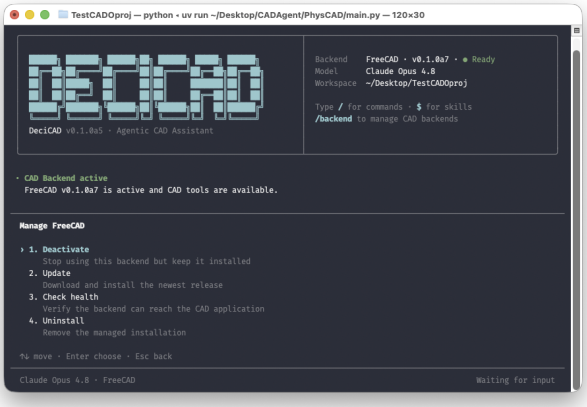

(c) Lifecycle and health controls for the FreeCAD backend.

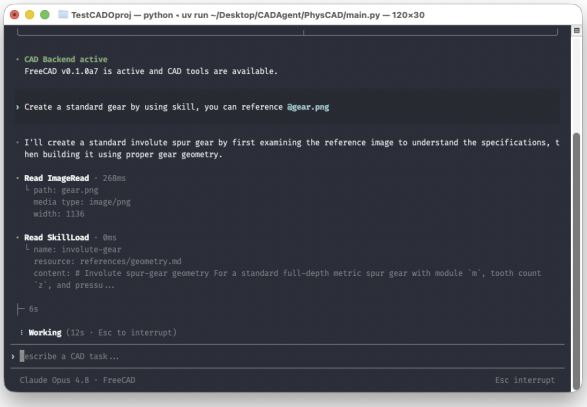

(d) A live task using an image reference and a CAD skill.

Figure 14: The project-level DeciCAD interface. The workspace keeps natural-language interaction, project resources, CAD skills, backend state, and streamed execution events in one persistent session. Users can inspect and control the runtime while semantic CAD actions continue to operate on the active document.

port for long-horizon multi-agent projects, the backend SDK available to developers, and our official FreeCAD integration. Its streamed action–observation loop follows the interaction pattern established by language agents while grounding each domain action in persistent CAD state (Yao et al. 2023; Yang et al. 2024).

E.1 From Natural Language to Editable CAD Documents

DeciCAD provides native release packages for macOS, Linux, and Windows and uses a managed installer for application updates. Running `decicad` inside a project opens that directory as the active workspace; `decicad -project /path/to/project` selects it explicitly. During initial setup, the application configures a model provider and requests trust before accessing project files or CAD resources.

The composer accepts part modeling, assembly, geometric revision, inspection, and export requests. An `@` reference attaches a project resource such as an image or specification, while `$` loads an inline CAD skill that supplies task-specific procedures. Slash commands expose runtime controls in the same interface: `/model` selects the provider, model, and reasoning effort; `/backend` installs, activates, updates, and checks a CAD backend; and `/status` summarizes the current runtime and token usage. Figure 14 shows these operations in the terminal interface.

Model actions and environment results then appear in temporal order. File and image reads identify the evidence used by the model, skill loads expose the selected domain procedure, and CAD calls report their execution state and compact result. Successful calls immediately update the persistent document; failed calls return information that supports argument correction, targeted inspection, or selection of another action. The user can interrupt execution or add a follow-up request at any point, and subsequent actions continue from

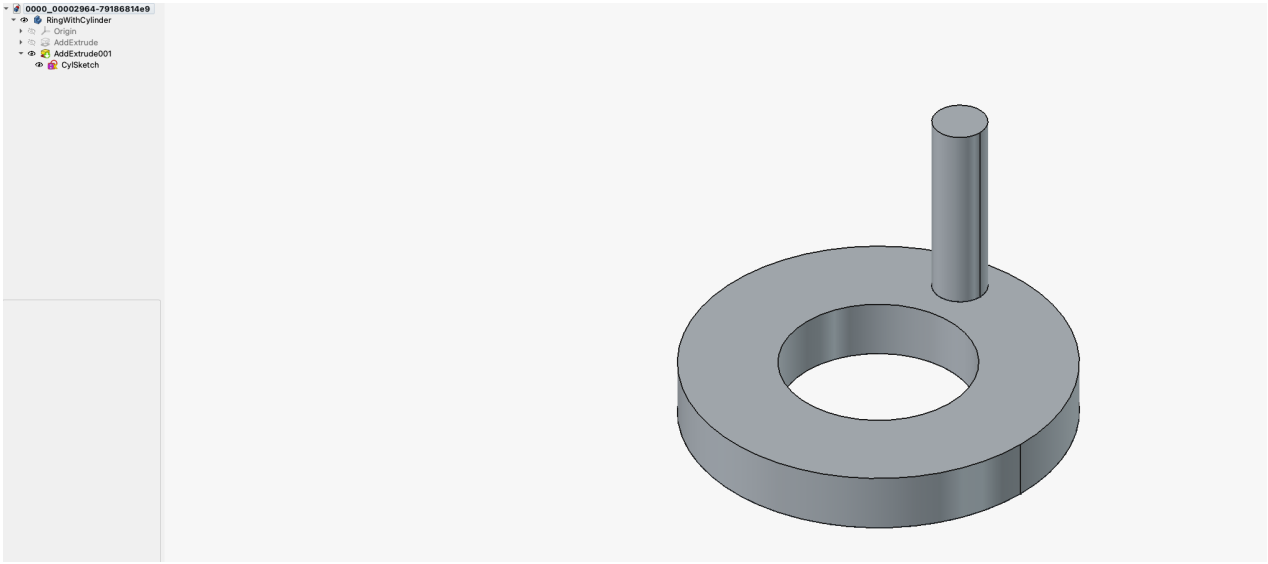

Figure 15: An editable CAD document produced through the FreeCAD backend. The viewport shows the generated solid, while the document tree retains the body, additive features, and underlying sketch created during the interaction. This persistent structure allows later inspection and modification through either DeciCAD or the native FreeCAD interface.

the current CAD state rather than regenerating the artifact from scratch.

The resulting artifact is also available in the native CAD application. Users can inspect the three-dimensional view and access the created Body, Sketch, Feature, Assembly, and Component objects in the document tree. The document can be saved in its native editable format or exported as a standard geometry asset, preserving a path from agent generation to manual inspection and continued engineering work.

E.2 Control Plane for Multi-Agent CAD Work

The interactive loop is supported by an application-level control plane that coordinates agents, tools, resources, permissions, and persistent project state. CAD backends execute domain operations, but they do not decide which agent may act, which calls may run concurrently, or whether an action was based on a current document revision. Keeping these responsibilities in the control plane allows the same orchestration policy to operate across compatible CAD systems.

Agent orchestration. For larger projects, a leader agent retains the user conversation and overall goal while delegating bounded subtasks to four built-in roles. The Designer turns functional and geometric intent into a feature and interface plan; the Modeler owns one complete part or a dependency-ordered modeling stage; the Assembler inserts components and manages placements, grounding, mates, and solve state; and the Verifier performs an independent, predominantly read-only evidence pass over the final document and exported assets. Each child receives the original request together with a bounded assignment, separate context, role-specific instructions, tool allowlist, and interaction budget. A child cannot recursively create more agents or modify the leader's task board. Its final response is returned as evidence together with the resources changed during execution, leaving acceptance and follow-up decisions with the leader.

Task and dependency state. Multi-stage work is represented by a compact project task graph. Tasks record their status, description, dependencies, blocking relations, and completion evidence. This graph separates conversational planning from execution state: the leader can recover which components are ready, which assembly step is blocked, and which verification remains after context compaction or session resumption. Dependency updates are serialized by the control plane, preventing two agents from concurrently producing inconsistent task-board transitions.

Resource resolution and authorization. Every tool specification declares static resource behavior, including the arguments that identify reads, writes, observations, creations, or filesystem materialization. Before execution, the resource manager resolves those fields to concrete document, entity, topology, assembly, or file keys. Each agent also maintains the revisions it has actually observed. A write is authorized only when its dependent resources have been observed at their current revisions; if another agent has modified one of them, the stale write is rejected with a recoverable conflict and the agent must inspect the updated state. File edits follow the same rule and require a complete prior read. This observation-based authorization prevents a plausible but outdated model context from silently overwriting newer CAD or project state.

Concurrency and commit. Independent child assignments can be emitted together and run concurrently. Tool specifications additionally assign calls to execution lanes and mark them as parallel or exclusive. After resource resolution, write scopes are sorted and locked in a deterministic order;

calls with disjoint scopes proceed in parallel, while overlapping CAD writes are serialized. Authorization is revalidated immediately before execution to close the interval between scheduling and dispatch. A successful backend result reports concrete `RuntimeResourceEffects`; the control plane commits these effects, advances resource generations, records the responsible agent, and invalidates stale observations held elsewhere. A failed call does not claim a state transition. If transport fails after a potentially mutating call, the affected scope is conservatively invalidated so that later decisions re-inspect it.

This protocol supports useful concurrency without treating the CAD document as an ordinary shared text buffer. Different Modelers may construct independent parts simultaneously, after which an Assembler consumes their committed outputs. Calls that touch the same body, assembly, or output file are ordered, and each downstream agent receives only the document revisions known to be valid at delegation time.

Permissions and project boundaries. Workspace trust and permission checks are evaluated before protected files or CAD resources are exposed to an agent. The user remains the authority for access that crosses the active project boundary or carries elevated effects. Because permissions are evaluated in the application rather than reimplemented by every backend, adding a CAD integration does not create a second, backend-specific authorization path.

Sessions, memory, and observability. The control plane projects model messages, tool calls, execution outcomes, permissions, resource commits, task changes, and compaction events onto an append-only rollout. The same event stream drives the terminal interface and provides the state required for replay and session restoration. `/resume` reconstructs an earlier project session, while `/compact` replaces older interaction with a bounded summary and keeps recent decisions available to the model. Durable memory is extracted from completed root rollouts with links to its source turns, separating persistent design facts from transient dialogue. Skills add reusable CAD procedures, hooks connect project workflows to agent and tool lifecycle events, and a project-level `AGENTS.md` carries coordinate conventions, naming rules, and delivery requirements into subsequent sessions. Together, these services allow long-running CAD work to remain recoverable, inspectable, and governed by one consistent project state.

E.3 Adapting a CAD Backend with the DeciCAD SDK

The DeciCAD SDK is the contract boundary shared by the application and backend implementations. It defines canonical tool names, typed input and output models, geometry types, opaque object references, permission requirements, runtime resource effects, structured success values, and model-facing errors. The application remains responsible for model interaction, agent orchestration, permissions, sessions, and resource coordination. A backend developer maps the canonical state transitions to the native API of a target CAD system.

This separation makes the SDK a semantic contract rather than a convenience wrapper around one vendor API. Canonical tools are grouped into document, inspection, entity, sketch, solid, and assembly domains. Their Pydantic models normalize units, frames, geometric values, operation modes, and result shapes, while `DocumentRef` and `EntityRef` preserve backend-owned object identity as opaque values. A backend imports these specifications instead of copying or redefining them. Consequently, field descriptions, model-facing tool schemas, resource declarations, and output validation remain identical across integrations, while implementation instructions can still describe behavior specific to the target CAD application.

A backend first imports a canonical `ToolSpec` and binds it to an async implementation with `@tool`. Inputs have already passed the spec's Pydantic validation before execution. The implementation returns the canonical output model together with accurate `RuntimeResourceEffects`; these effects let the application update resource versions, permissions, and locks. Recoverable CAD failures use a compact `ToolError`, keeping operational failures distinct from backend defects and transport errors. The following example summarizes the binding used for native document creation.

Listing 1: Minimal binding of a canonical SDK tool to a native CAD operation.

```
CREATE = DOCUMENT_TOOLS["DocumentCreate"]

@tool(spec=CREATE,
      instructions="Create one native document.")
async def create_document(arguments):
    result = await execute_cad_write(
        operation="new_document",
        arguments=arguments,
    )
    if result.ok:
        record_document_create_effect(result)
    return result
```

The backend then declares a manifest containing its backend, SDK, protocol, and tool-contract versions. A `BackendServer` registers the complete canonical tool roster and exposes a health check for the target CAD runtime, as illustrated in Listing 2. During activation, a `BackendClient` starts the isolated process and verifies the manifest, protocol version, tool registry, contract digest, and health response. Only after this initialization does it bind remote executors as ordinary SDK tools inside DeciCAD.

Listing 2: Registering tool domains and exposing backend health through the SDK server.

```
server = BackendServer(
    manifest=MANIFEST,
    health_check=health_check,
)
for domain in (document, entity, inspection,
               sketch, solid, assembly):
    server.register(domain.TOOLS)
await server.serve()
```

The server reserves standard output for newline-delimited JSON-RPC frames and sends operational diagnostics to standard error, keeping transport parseable under subprocess supervision. It refuses to initialize when a canonical tool is

Domain	Canonical operations	Supported behavior
Document	`DocumentCreate`, `DocumentOpen`, `DocumentSave`, `DocumentExport`, `DocumentClose`, `DocumentDeleteEntity`	Create and open native FreeCAD documents; save editable documents; export selected targets to STEP or STL; close documents; and delete unused entities without downstream dependents.
Inspection	`InspectList`, `InspectGet`	Page through bounded document, entity, geometry, topology, and assembly collections, then retrieve canonical detail for one stable reference.
Entity	`EntitySetPlacement`	Apply rigid transforms in world or parent space to placeable CAD entities and component instances.
Sketch	`SketchCreate`, `SketchAddLines`, `SketchAddCircle`, `SketchAddArc`, `SketchAddPoint`, `SketchConstrain`, `SketchProject`, `SketchPattern`, `SketchTrim`	Create plane-bound sketches; add local lines, circles, arcs, and points; apply geometric and dimensional constraints; project external geometry; create linear or circular patterns; and trim at intersections.
Solid	`SolidCreateBody`, `SolidExtrude`, `SolidRevolve`, `SolidSweep`, `SolidLoft`, `SolidBoolean`, `SolidRepeat`, `SolidFillet`, `SolidChamfer`, `SolidShell`	Create feature bodies; add or cut material by extrusion, revolution, sweep, or loft; perform body booleans; repeat or mirror features and bodies; and apply fillets, chamfers, and inward or outward shells.
Assembly	`AssemblyCreate`, `AssemblyAddComponent`, `AssemblySetGrounded`, `AssemblyMateAdd`, `AssemblyMateRemove`, `AssemblySolve`	Create assemblies; insert placed component instances; ground or release components; add and remove geometry-referenced mates; and solve the resulting assembly constraints.

Table 8: Canonical CAD operations implemented by the current FreeCAD backend. The 34 tools constitute the complete SDK backend registry used by DeciCAD.

missing, a tool name is duplicated, an imported spec differs from the SDK, or the manifest targets incompatible SDK, protocol, or tool-contract versions. Successful results are validated against the exact output model before they cross the process boundary; operational failure results contain a `ToolError` and cannot claim resource effects.

These constraints reduce the work required from a backend developer. The SDK provides model schemas, validation, RPC framing, process lifecycle, contract negotiation, error serialization, and the client-side `Tool` bindings. The developer concentrates on object resolution, native CAD calls, readiness, and accurate resource effects. A real `BackendClient`/`BackendServer` initialization test exercises the same registry and wire contract used by the product, catching integration drift before an agent can invoke the backend. The application and backend may then be released independently whenever their declared contracts remain compatible; an incompatible release fails explicitly at activation instead of producing ambiguous behavior midway through a modeling task.

E.4 FreeCAD as the Reference Backend

Our official reference backend instantiates this SDK boundary for FreeCAD, and all experiments in this paper use this backend. The backend manager shown in Figure 14 controls its installation, activation, health checking, update, and removal through an isolated managed environment. In development mode, the setup command also installs a FreeCAD Addon into the versioned user directory discovered from the actual FreeCAD runtime, supporting system, portable, and custom installations.

The integration uses two local communication layers. The SDK `BackendClient` communicates with an independent backend process; that process connects through an authenticated loopback bridge to FreeCAD's embedded Python runtime. The bridge executes CAD operations on the FreeCAD GUI thread, wraps document writes in native transactions, and encodes the result under the canonical output model. The bridge panel exposes connection state, dynamic endpoint, protocol version, age of the last authenticated request, and start, stop, restart, and automatic-start controls, while the application sees only the stable SDK contract. Keeping the bridge protocol separate is necessary because FreeCAD owns the embedded interpreter and GUI event loop; it also keeps FreeCAD object lookup, argument decoding, result encoding, transaction handling, and native diagnostics out of the application process.

Table 8 summarizes the complete canonical operation set currently implemented by the reference backend. Stable opaque references connect documents, entities, geometry, topology, features, and assembly objects across calls. Successful tools report concrete resource effects, and failures are mapped into shared error categories, so the general agent and resource runtime do not depend on the FreeCAD Python API.

Inspection is hierarchical rather than a full document dump. Paginated list operations expose document structure and stable references, after which get operations retrieve detailed properties, placement, geometry, topology, or assembly state only for selected objects. Modeling tools return compact native results and resource effects immediately after execution. This supports the same active and passive observation pattern used by the DeciCAD decision loop while retaining direct access to FreeCAD's editable feature history.

Figure 15 shows a document state produced through this path. The three-dimensional solid is accompanied by its body, additive features, and sketch in the FreeCAD tree, rather than being exposed only as a flattened mesh or image.

For batch evaluation and data collection, multiple workers run isolated FreeCAD and backend processes with separate bridge endpoints, authentication tokens, instance identifiers, and result caches. A worker resets and verifies all open doc-

uments before accepting another episode. This isolation prevents CAD state from leaking across tasks while retaining the same backend contract used by the interactive product.

Parallelism is therefore process-isolated rather than implemented by issuing concurrent CAD operations inside one FreeCAD instance. Each worker owns one FreeCAD process, one Addon bridge, and one matching backend process. Distinct endpoint descriptor files lead to independent dynamic loopback ports and authentication tokens. This design supports parallel interactive services, trajectory collection, and benchmark evaluation without sharing mutable FreeCAD state across tasks.

Capability	FreeCAD	SolidWorks	AutoCAD
Documents and objects	✓	✓	✓
Inspection and query	✓	✓	✓
Parametric modeling	✓	✓	
Assemblies and mates	✓	✓	
2D drafting			✓
Exchange formats	STEP, STL	STEP, STL	DXF

Table 9: Capabilities supported by the three CAD backends. All backends implement the same SDK contract; all experiments in this paper use the FreeCAD backend.

E.5 Additional Backends

The same SDK contract also supports two commercial CAD systems, SolidWorks and AutoCAD. Each backend declares the capabilities it implements, and the application loads the corresponding tools when it connects; agent orchestration, permissions, resource coordination, and sessions are shared across all backends. Table 9 summarizes the capabilities supported by the three backends.