arXiv:2610.21496v1 [cs.CL] 15 Oct 2026

OPRD: On-Policy Reminder Distillation for Long-Horizon Coding Agents

Yuechen Luo[1,*], **Zhiyi Lai**[2,*], **Miao Ziqi**[3,*], **Zihang Xu**[3], **Jingxuan Xu**[4], **Songwei Yu**[4], **Wenqiang Zhu**[2], **Weihao Li**[2], **Junhao Fan**[1], **Haoyang Huang**[2], **Yanan Wu**[2], **Ken Deng**[2,†]

[1]Tsinghua University, [2]Independent Researcher
[3]Shanghai Artificial Intelligence Laboratory, [4]National University of Singapore
luo-yc24@mails.tsinghua.edu.cn

Abstract

Reliable repository-level coding requires agents not only to generate correct code but also to preserve constraints and respond to tool feedback across long interactions. However, our diagnostic analysis shows that agents often fail to act on relevant information already available in context, while final Pass/Fail outcomes provide only sparse process supervision. We introduce **On-Policy Reminder Distillation (OPRD)**, which combines Multi-Scope Reminder supervision with outcome-based routing on reminder-free student rollouts. Successful rollouts provide autoregressive self-imitation targets, whereas failed rollouts receive Reminder-conditioned reverse-KL supervision from a teacher that re-evaluates the same student-generated prefixes. A Global Reminder converts final-test diagnostics into trajectory-level guidance, while Turn-local Reminders apply only to the next decision following the relevant environment observation. The teacher and student share a two-stage initialization from Coding SFT and Instruction-Following (IF) SFT. At inference, the student operates without Reminders. On SWE-bench Verified, SWE-bench Pro, and SWE-bench Multilingual, OPRD improves the benchmark-averaged task resolution rate by 11.3 points over the open-source Qwen3.5-4B base model. Under the same diagnostic protocol, it reduces the aggregate error-event count from 446 to 222, a 50.2% reduction.

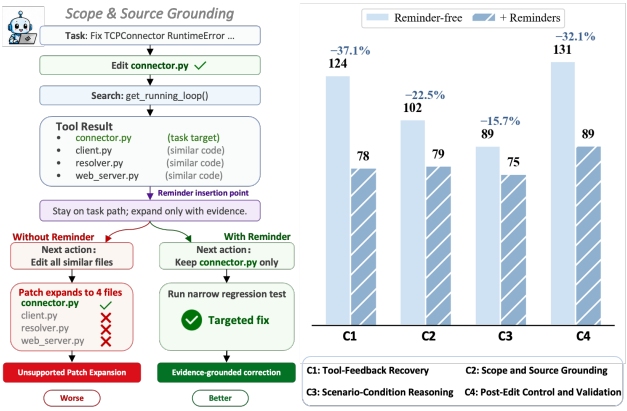

Figure 1: Diagnostic analysis of Qwen3.5-4B on 200 tasks from the SWE-rebench v2 validation set. Left: a *Scope and Source Grounding* example in which a Turn-local Reminder prevents unsupported patch expansion. Right: aggregate error-event counts across four categories under Reminder-free and Reminder-augmented settings; percentages indicate relative reductions.

1 Introduction

Large language models have advanced in reasoning (Zhan et al. 2025a; Deng et al. 2025a; Xu et al. 2026b; Miao et al. 2026) and repository-level software engineering (Jimenez et al. 2024; Yang et al. 2024; Zhang et al. 2024; Ouyang et al. 2025; Xia et al. 2025). Reliable long-horizon coding nevertheless requires agents to preserve task constraints and adapt to tool and execution feedback (Yao et al. 2023; Chen et al. 2024). We examine this challenge on 200 SWE-rebench V2 validation tasks (Badertdinov et al. 2026), using two Qwen3.5-4B rollouts per task under identical inference settings (Qwen Team 2026). In the reminder-free rollout, deterministic rules identify explicit tool and execution failures, while GPT-5.4 classifies remaining context-dependent events from the task and complete trajectory (OpenAI 2026). Each event belongs to one of four categories in Table 6: *Tool-Feedback Recovery*, *Scope and Source Grounding*, *Scenario-Condition Reasoning*, and *Post-Edit Control and Validation*. In the second rollout, the agent restarts from the task under the same settings, and the Reminder mechanism examines each environment observation to decide whether to insert a Turn-local Reminder before the next action. Applying the same annotation protocol reveals fewer error events across all four categories (Fig. 1). Post-hoc annotation audits completed trajectories, while runtime Reminder selection uses only the currently visible interaction prefix. This finding suggests that reliable long-horizon coding requires models to preserve task-level constraints across a trajectory and use local interaction feedback at appropriate decision points. A natural direction is therefore to provide process-level corrective information with explicit temporal scope.

However, coding-agent training relies on supervised imitation of successful trajectories or executable-outcome optimization (Pan et al. 2025; Jain et al. 2025; Wei et al. 2025). These signals improve task completion but do not indicate which constraints should persist or which observation should guide the next decision. At inference time, agents can adapt to feedback (Yao et al. 2023; Chen et al. 2024) or recover missed information through reflection and self-critique (Shinn et al.

*Equal contribution.
†Corresponding author.

2023; Madaan et al. 2023; Gou et al. 2024). Runtime reminders require deployment-time construction and injection, consume context, and tie behavior to an agent scaffold. On-policy distillation trains students on self-generated outputs to reduce train–inference mismatch (Agarwal et al. 2024; Gu et al. 2024). Related work uses context- and feedback-conditioned distillation to internalize behavior induced by auxiliary instructions or environment feedback (Ye et al. 2026; Hübotter et al. 2026). Applying these ideas to long-horizon coding leaves two challenges. First, training-time corrective information can encode persistent constraints and local corrections from observations relevant to one decision. Without explicit scopes, local corrections may be diluted in long contexts or propagated to unrelated decisions. Second, models without instruction-following adaptation may ignore constraints and feedback, so teacher reminders may not produce reliable distillation targets.

This raises a practical question: *Can we internalize scope-specific corrective information during training so that a coding agent can operate without reminders at inference time?*

We address this with **On-Policy Reminder Distillation (OPRD)**, which combines Multi-Scope Reminder supervision with outcome-based routing on rollouts from a reminder-free student. For failed trajectories, we construct reminders from task, trajectory, and execution/validation records. A Global Reminder converts final-test diagnostics into trajectory-level guidance. A Turn-local Reminder is anchored to an observation and applies only to the next decision. This combines trajectory-level feedback with local correction without propagating local feedback to unrelated decisions. For reliable supervision, our two-stage cold start uses Coding SFT to establish coding capabilities and Instruction-Following (IF) SFT to improve responsiveness to constraints, feedback, and reminders. Both models are initialized from the IF checkpoint. Based on executable verification, OPRD routes successful rollouts to autoregressive self-imitation and failed rollouts to reminder-conditioned distillation without reinforcing erroneous decisions. For failures, the teacher re-evaluates student-generated prefixes under reminders, while the student matches the teacher distribution through Top-100 reverse KL. OPRD internalizes scope-specific corrections, enabling the student to act on corresponding constraints and feedback without reminders at inference time. We evaluate OPRD on SWE-bench Verified, SWE-bench Pro, and SWE-bench Multilingual (OpenAI 2024; Deng et al. 2025b; Yang et al. 2025). OPRD improves the benchmark-averaged task resolution rate by 11.3 points over the open-source Qwen3.5-4B base model. Under the same diagnostic protocol, the trained OPRD model reduces the aggregate error-event count across the four categories from 446 to 222, a 50.2% reduction relative to the base model.

Our main contributions are threefold:

- **Multi-Scope Reminder supervision for long-horizon coding.** We combine a Global Reminder that provides guidance across the trajectory with Turn-local Reminders that target only the next decision, enabling process-level supervision at complementary temporal scopes.
- **On-Policy Reminder Distillation with Outcome-Based Routing.** We propose OPRD, which learns from verified successful rollouts through autoregressive self-imitation and from failed rollouts through reminder-conditioned reverse KL distillation.
- **A two-stage supervised cold start for reliable reminder distillation.** Coding SFT followed by IF SFT provides a shared teacher–student initialization that improves responsiveness to constraints, execution feedback, and reminders, thereby providing a stronger starting point for subsequent reminder-conditioned distillation.

2 Related Work

Repository-Level Software Engineering Agents. SWE-bench established executable evaluation on repository issues (Jimenez et al. 2024). SWE-agent, AutoCodeRover, and Agentless support issue resolution through tool interfaces, program-aware search, and structured repair pipelines (Yang et al. 2024; Zhang et al. 2024; Xia et al. 2025). Training resources include agent trajectories and verifiers (Pan et al. 2025; Jain et al. 2025), generated or freshly collected tasks (Yang et al. 2025; Badertdinov et al. 2025), and reinforcement learning from executable outcomes (Wei et al. 2025). SWE-Compass and KAT-Coder further study broad capability evaluation and coding-agent training (Xu et al. 2025; Zhan et al. 2025b).

Feedback-Guided Reasoning and Reflection. ReAct conditions actions on observations, while Reflexion and self-refinement use feedback to revise subsequent attempts (Yao et al. 2023; Shinn et al. 2023; Madaan et al. 2023). FTTT formulates feedback use as test-time optimization (Li, Lyu, and Wang 2025). Tool-grounded critique, search, and code-specific execution or runtime verification further guide correction (Gou et al. 2024; Zhou et al. 2024; Chen et al. 2024; Zhong, Wang, and Shang 2024). FLARE extends process feedback through a generative reward model for inference and post-training (Xu et al. 2026a). These approaches motivate learning to use corrective signals without requiring runtime Reminder injection.

On-Policy and Privileged-Context Distillation. On-policy distillation obtains teacher supervision on student-generated sequences to reduce train–inference mismatch (Agarwal et al. 2024). MiniLLM and DistiLLM study reverse or skew KL objectives (Gu et al. 2024; Ko et al. 2024), with extensions using entropy-aware or black-box feedback (Jin et al. 2026; Ye et al. 2025). Context distillation internalizes behavior induced by auxiliary instructions, demonstrations, or scratchpads (Snell, Klein, and Zhong 2022); privileged-context methods additionally condition teachers on verified solutions or other training-only information (Zhao et al. 2026; Penaloza et al. 2026). OPCD matches student trajectories to a context-conditioned teacher through reverse KL (Ye et al. 2026). OPRD combines explicitly scoped Reminders with outcome-based routing for long-horizon coding.

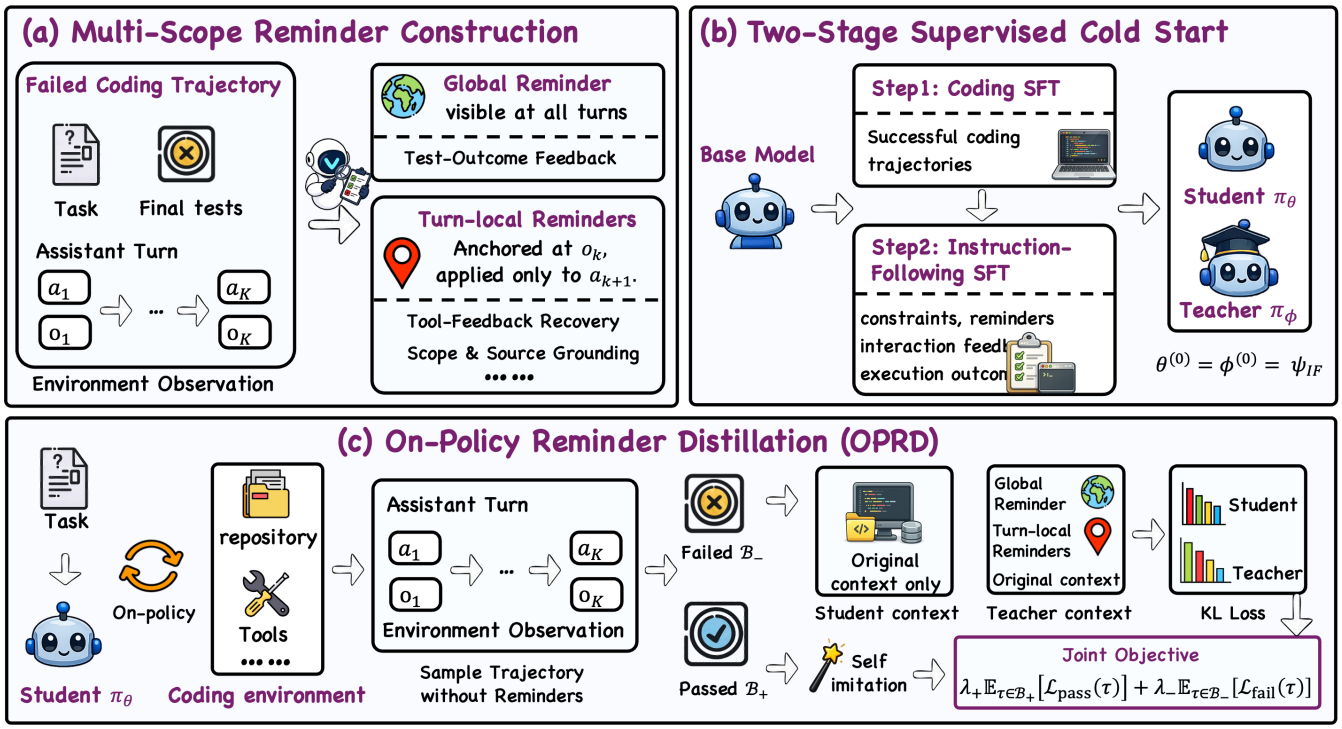

Figure 2: **Overview of OPRD.** OPRD constructs Global Reminders and observation-anchored Turn-local Reminders, and initializes the teacher and student through Coding SFT followed by IF SFT. During on-policy training, successful rollouts are optimized through autoregressive self-imitation, while failed rollouts receive reverse-KL supervision from a Reminder-conditioned teacher.

3 Method

Long-horizon coding agents must not only generate correct code, but also consistently follow user constraints throughout multi-turn interactions and adapt their behavior to tool responses and execution feedback. However, final verification outcomes provide limited process-level supervision for constraint adherence and feedback use at different scopes. To address this limitation, we propose **O**n-**P**olicy **R**eminder **D**istillation **(OPRD)**, which distills the behavior of a teacher conditioned on Multi-Scope Reminders into a student that operates without reminders at inference time. As illustrated in Fig. 2, we formulate the long-horizon coding task (Sec. 3.1), then introduce Multi-Scope Reminders (Sec. 3.2) and a two-stage supervised cold start (Sec. 3.3), and present Multi-Scope Reminder-Conditioned Distillation (Sec. 3.4).

3.1 Problem Formulation

Given a repository-level coding task x, an agent interacts with a code execution environment over multiple turns. We denote the k-th assistant turn by a_k and the resulting environment observation by o_k. The complete interaction trajectory is

$$\tau = (x, a_1, o_1, \ldots, a_K, o_K). \tag{1}$$

An executable verifier evaluates the final repository state and returns $c(\tau) \in \{0, 1\}$, where $c(\tau) = 1$ indicates that the trajectory passes verification and $c(\tau) = 0$ indicates failure. We denote the trainable student by π_θ and the teacher by π_ϕ. The token sequence of the k-th assistant turn is $a_k = (y_{k,1}, \ldots, y_{k,L_k})$, where L_k is the number of tokens in that turn. The interaction history preceding the k-th assistant turn and the student context for generating its j-th token are defined as

$$H_k = (x, a_1, o_1, \ldots, a_{k-1}, o_{k-1}), \qquad h^S_{k,j} = (H_k, y_{k,<j}), \tag{2}$$

where $y_{k,<j}$ denotes the token prefix preceding $y_{k,j}$ within the current turn. System messages, user messages, and tool responses provide generation context but are not prediction targets. Training is restricted to valid assistant tokens, whose index set is

$$\Omega(\tau) \subseteq \{(k, j) \mid 1 \le k \le K,\ 1 \le j \le L_k\}. \tag{3}$$

3.2 Multi-Scope Reminders

The constraints and feedback received by a long-horizon coding agent operate over different temporal scopes. Task-level constraints, such as permitted modification scope, interface requirements, and output format, must hold throughout the trajectory. In contrast, test errors, tool responses, and file inspection results typically pertain to a specific decision point. Nevertheless, the final verifier compresses the entire trajectory into a binary outcome and cannot represent these distinct scopes. Treating all corrective information as a single trajectory-wide prompt may cause local errors to be diluted by the long context, whereas propagating local feedback to

all subsequent turns may introduce irrelevant or conflicting supervision as the interaction state evolves. We therefore organize corrective information from a failed trajectory into a Multi-Scope Reminder. A Global Reminder provides persistent task-level guidance across turns, while a Turn-local Reminder is aligned with a specific decision point and highlights environmental feedback that the model ignored or misinterpreted at that point. For a failed trajectory τ, its Multi-Scope Reminder is represented as

$$R(\tau) = \left(r^G, \{(k, r_k^L)\}_{k\in\mathcal{I}_\tau}\right), \qquad \mathcal{I}_\tau \subseteq \{1, \ldots, K-1\}, \tag{4}$$

where r^G denotes the Global Reminder, r_k^L denotes the Turn-local Reminder anchored to environment observation o_k, and $\mathcal{I}_\tau$ is the set of observation indices associated with local corrective signals. The Global Reminder remains visible throughout the teacher's evaluation of assistant turns. In contrast, r_k^L is visible only when the teacher evaluates the assistant turn a_{k+1} following o_k, and does not propagate to later turns. Instruction-Following (IF) SFT first strengthens the model's ability to understand and respond to explicit constraints and feedback. During OPRD, Multi-Scope Reminders serve as training-time privileged information available only to the teacher, and the behavior they induce is distilled into a student that does not observe reminders.

Reminder Construction. When a trajectory contains a correctable failure, we construct reminders post hoc from the task, trajectory, interaction feedback, and final test results. We determine the scope of each corrective signal and then assign its functional category. The Global Reminder contains one category, *Test-Outcome Feedback*, which converts final-test diagnostics into trajectory-level guidance. Turn-local Reminders target correctable decisions identified from each environment observation and assistant action. They comprise four categories: *Tool-Feedback Recovery* for responding to explicit tool feedback; *Scope and Source Grounding* for evidence-grounded localization and modification; *Scenario-Condition Reasoning* for task-specific execution conditions; and *Post-Edit Control and Validation* for controlled post-edit verification. We retain grounded and actionable candidates, discarding ambiguous, redundant, unverifiable, reference-patch-revealing, or unsupported content. Detailed category definitions, generation prompts, trigger rules, and filtering procedures are provided in Appendix C.

3.3 Two-Stage Supervised Cold Start

The base model first acquires capabilities for repository understanding, tool interaction, and code editing through Coding SFT. However, coding competence alone does not ensure that the model will follow constraints and feedback that arise during an interaction. A standard coding model may still ignore user requirements, tool responses, test errors, or harness rules that are already present in the context. Given a base model ψ_0, we first train on successful coding trajectories $\mathcal{D}_{\text{code}}$ to establish core agentic coding capabilities. We then use IF data $\mathcal{D}_{\text{IF}}$ to improve the model's responsiveness to user constraints, interaction feedback, execution outcomes, and reminders, yielding an IF checkpoint ψ_{IF} that combines coding and instruction-following capabilities:

$$\psi_{\text{code}} = \text{SFT}(\psi_0, \mathcal{D}_{\text{code}}), \qquad \psi_{\text{IF}} = \text{SFT}(\psi_{\text{code}}, \mathcal{D}_{\text{IF}}). \tag{5}$$

3.4 Multi-Scope Reminder-Conditioned Distillation

OPRD uses shared initialization. The teacher π_ϕ and student π_θ use the same model architecture and are initialized with the parameters of the same IF checkpoint ψ_{IF}, such that $\theta^{(0)} = \phi^{(0)} = \psi_{\text{IF}}$. During training, the teacher is periodically updated toward the student, while the student is optimized at every step. The teacher additionally receives reminders as training-time privileged information. For a task $x \sim \mathcal{D}_{\text{task}}$, the current student interacts with the coding environment without access to reminders to sample $\tau \sim \pi_\theta(\cdot \mid x)$. As the student is updated, subsequent rollouts expose different states and errors, keeping reminder supervision aligned with the distribution visited by the current policy.

Outcome-Based Routing. After each rollout, an executable verifier determines whether it succeeds but does not provide fine-grained credit assignment. A successful rollout provides executable evidence that the student's current behavior can solve the task and can therefore serve directly as an autoregressive self-imitation signal. Applying the same objective to a failed rollout would instead increase the likelihood of erroneous decisions. OPRD thus routes successful rollouts to a self-imitation branch and failed rollouts to a reminder-conditioned distillation branch. For the current rollout batch $\mathcal{B}$, we define $\mathcal{B}_+ = \{\tau \in \mathcal{B} \mid c(\tau) = 1\}$ and $\mathcal{B}_- = \{\tau \in \mathcal{B} \mid c(\tau) = 0\}$.

Autoregressive Self-Imitation from Successful Rollouts. For each $\tau \in \mathcal{B}_+$, we compute the autoregressive loss over all valid assistant tokens to increase the likelihood that the student reproduces successful behavior discovered through its own on-policy sampling:

$$\mathcal{L}_{\text{pass}}(\tau) = -\frac{1}{|\Omega(\tau)|} \sum_{(k,j)\in\Omega(\tau)} \log \pi_\theta(y_{k,j} \mid h_{k,j}^S). \tag{6}$$

Reminder-Conditioned Distillation from Failed Rollouts. For each $\tau \in \mathcal{B}_-$, reminders are constructed after the rollout and provided only to the teacher. The student always uses the original context without reminders. The teacher does not resample the trajectory. Instead, it evaluates the same student-generated token prefix $y_{k,<j}$, ensuring that the distillation targets remain aligned with the states visited by the current policy. For the k-th assistant turn, the teacher history is defined as

$$\widetilde{H}_k^T = \begin{cases} \text{Inject}(H_k, r^G, r_{k-1}^L), & \substack{k\geq 2,\\ k-1\in\mathcal{I}_\tau,} \\ \text{Inject}(H_k, r^G), & \text{otherwise.} \end{cases} \tag{7}$$

The corresponding teacher token context is $h_{k,j}^T = (\widetilde{H}_k^T, y_{k,<j})$. The teacher augments the student's original context with the reminders active at the current scope. The Global Reminder is visible at every assistant turn, whereas

r^L_{k-1} is injected only when the teacher evaluates a_k and is removed from subsequent turns. Reminder scope is determined by its visibility in the teacher context, while the loss for a failed trajectory still covers all valid assistant tokens in $\Omega(\tau)$. Within this unified distillation signal, the Global Reminder provides task-level correction across turns, while the Turn-local Reminder provides fine-grained guidance at the corresponding decision point. Let $\mathcal{V}$ denote the full vocabulary, and let z_θ and z_ϕ denote the student and teacher logits, respectively. For $v \in \mathcal{V}$, their token distributions at temperature $\mathcal{T}$ are

$$p^{(\mathcal{T})}_{\theta,k,j}(v) = \mathrm{softmax}\left(\frac{z_\theta(h^S_{k,j})}{\mathcal{T}}\right)_v, \qquad p^{(\mathcal{T})}_{\phi,k,j}(v) = \mathrm{softmax}\left(\frac{z_\phi(h^T_{k,j})}{\mathcal{T}}\right)_v. \tag{8}$$

We then compute the reverse Kullback–Leibler (KL) divergence over the Top-K-vocabulary distributions:

$$\mathcal{L}_{\mathrm{fail}}(\tau) = \frac{\mathcal{T}^2}{|\Omega(\tau)|} \sum_{(k,j)\in\Omega(\tau)} D_{\mathrm{KL}}\left[p^{(\mathcal{T})}_{\theta,k,j} \,\middle\|\, \mathrm{sg}\left[p^{(\mathcal{T})}_{\phi,k,j}\right]\right], \tag{9}$$

where sg stops gradients through the teacher branch and $\mathcal{T}^2$ compensates for gradient-scale changes from temperature scaling. Reverse KL weights discrepancies by the student distribution, prioritizing corrections to tokens to which the student assigns high probability in visited states.

Joint Objective and Inference. After training, only the student is deployed. The teacher, the reminder-construction module, and all reminders are removed at inference time. Let λ_+ and λ_- control the relative strengths of self-imitation on successful rollouts and reminder distillation on failed rollouts, respectively. The OPRD objective is

$$\mathcal{L}_{\mathrm{OPRD}} = \lambda_+ \mathbb{E}_{\tau\in\mathcal{B}_+}[\mathcal{L}_{\mathrm{pass}}(\tau)] + \lambda_- \mathbb{E}_{\tau\in\mathcal{B}_-}[\mathcal{L}_{\mathrm{fail}}(\tau)]. \tag{10}$$

4 Experiments

4.1 Experimental Setup

Models and Training Data. We use Qwen3.5-4B and Qwen3.5-9B as base models (Qwen Team 2026). Training consists of three stages. Coding SFT uses 8,000 verified successful trajectories, including 7,262 from SWE-REBENCH V2 (Badertdinov et al. 2026) and 738 from Multi-SWE-RL (Zan et al. 2025). IF SFT then uses 2,000 reminder-following examples from SWE-REBENCH V2. OPRD operates on a pool of 3,630 SWE-REBENCH V2 tasks and 370 Multi-SWE-RL tasks, from which the student generates reminder-free rollouts online. In the OPRD variant without IF SFT, both teacher and student are initialized from the Coding SFT checkpoint while all other settings remain unchanged.

Table 1 summarizes the three stage-specific data configurations. Coding SFT and IF SFT operate on trajectories or reminder-following examples, whereas OPRD uses a task pool from which the current student generates new rollouts online. The OPRD counts therefore describe available training tasks rather than a fixed set of trajectories.

Stage	Unit	SWE-REBENCH V2	Multi-SWE-RL
Coding SFT	Trajectory	7,262	738
IF SFT	Example	2,000	0
OPRD	Task	3,630	370

Table 1: Training data used at each stage.

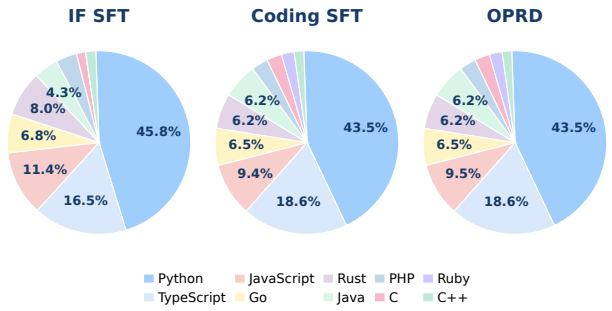

Figure 3: Programming-language distributions for IF SFT, Coding SFT, and the OPRD task pool. Percentages are computed from 2,000 examples, 8,000 examples, and 4,000 examples, respectively. The same language-color mapping is used in all three pies.

Figure 3 reports the programming-language composition of the three stages. Python is the largest source in every stage (43.5–45.8%), followed by TypeScript and JavaScript. Coding SFT and the OPRD task pool use nearly identical proportions by construction, while the IF set is slightly more concentrated in Python, JavaScript, and Rust and contains no Ruby example. The distribution nevertheless covers nine languages in IF SFT and all ten languages in Coding SFT and OPRD.

Implementation Details. Training implementation and hyperparameters are detailed in Appendix B.

Evaluation Setup. We evaluate task resolution rate on SWE-BENCH VERIFIED, SWE-BENCH PRO, and SWE-BENCH MULTILINGUAL (OpenAI 2024; Deng et al. 2025b; Yang et al. 2025). All models use the same agent harness and inference configuration. Except for the controlled Reminder interventions in the scope and functional-category analyses, all main results and training ablations use reminder-free inference.

All models are evaluated with a shared Claude Code agent scaffold and the executable verifier associated with each benchmark. Every subtask is allocated 16 CPU cores and 32,768 MiB of memory, with a maximum of 200 agent turns. Decoding uses temperature 1.0 and top-p 0.95. The vLLM deployment permits at most 128 concurrent sequences, a maximum model length of 262,144 tokens, and 8,192 batched tokens.

Training Configuration	SWE-bench Verified (%)	SWE-bench Pro (%)	SWE-bench Multilingual (%)	Average (%)
Qwen3.5-4B	50.4	24.2	35.3	36.6
+ Coding SFT	53.6	25.6	38.7	39.3
+ Coding SFT → IF SFT	55.8	26.7	42.3	41.6
+ Coding SFT → OPRD (w/o IF SFT)	<u>58.4</u>	<u>28.3</u>	<u>46.3</u>	<u>44.3</u>
+ Coding SFT → IF SFT → OPRD (Ours)	**61.2**	**30.5**	**52.0**	**47.9**
Qwen3.5-9B	59.6	32.3	41.3	44.4
+ Coding SFT	61.4	34.7	47.7	47.9
+ Coding SFT → IF SFT	64.0	36.7	50.7	50.5
+ Coding SFT → OPRD (w/o IF SFT)	<u>67.2</u>	<u>38.9</u>	<u>53.7</u>	<u>53.3</u>
+ Coding SFT → IF SFT → OPRD (Ours)	**70.0**	**40.9**	**56.7**	**55.9**

Table 2: Main results with Qwen3.5-4B and Qwen3.5-9B on SWE-bench Verified, SWE-bench Pro, and SWE-bench Multilingual. We report task resolution rates (%) and their unweighted average. Within each backbone, the best and second-best results in each column are shown in bold and underlined, respectively.

4.2 Main Results

OPRD provides consistent gains across software engineering benchmarks. Table 2 compares the base and supervised checkpoints with the OPRD variants to assess whether Reminder-conditioned process supervision transfers to reminder-free inference. The complete training recipe reaches an average resolution rate of 47.9, improving over the Qwen3.5-4B base model by 11.3 points and over the IF SFT checkpoint by 6.3 points. Gains are consistent across all three benchmarks, with the largest improvement on SWE-bench Multilingual. Removing IF SFT before OPRD reduces the average by 3.6 points. This gap suggests that coding adaptation alone is insufficient for fully exploiting Reminder-conditioned supervision, whereas IF SFT provides a more suitable shared initialization for teacher and student. Because all entries in Table 2 use reminder-free inference, the improvement does not depend on test-time Reminder injection.

OPRD gains persist with the 9B backbone. As shown in Table 2, the complete Qwen3.5-9B training recipe reaches an average resolution rate of 55.9, improving over the open-source checkpoint by 11.5 points and over the IF SFT checkpoint by 5.4 points. The improvement over the open-source model is consistent on SWE-bench Verified, SWE-bench Pro, and SWE-bench Multilingual, with gains of 10.4, 8.6, and 15.4 points, respectively. Removing IF SFT before OPRD lowers the average by 2.6 points, showing that the supervised instruction-following preparation remains useful at the larger model scale.

4.3 Ablation Study

We examine the distillation design, loss balance, teacher refresh, supervised preparation, and Reminder scope and functional categories, followed by a process-level error analysis. Training ablations use the same task pool and on-policy rollout protocol; Reminder interventions are compared under shared inference settings.

The OPRD variant without IF SFT initializes both the teacher and student from the Coding SFT checkpoint and otherwise retains the complete OPRD configuration. In the distillation-formulation ablation, All-Rollout OPSD applies

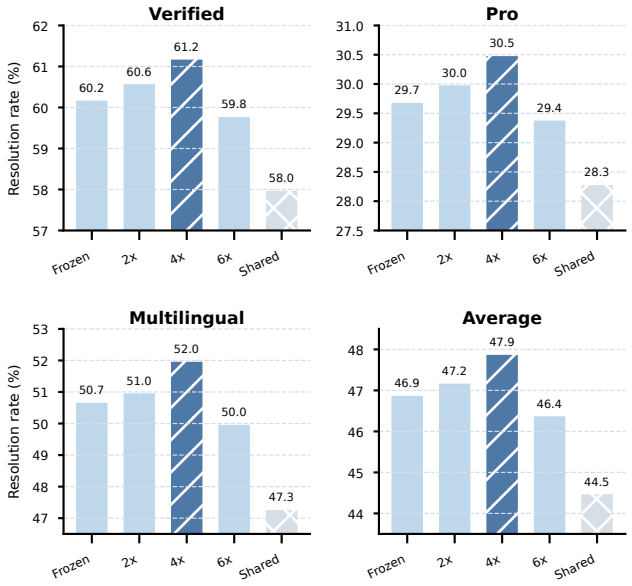

Figure 4: Effect of the teacher refresh schedule on Qwen3.5-4B. We compare Frozen Teacher, Periodic Refresh ($2\times$, $4\times$, and $6\times$), and Shared Parameters on three benchmarks and their unweighted average.

the selected KL objective to both successful and failed on-policy rollouts. Outcome-Routed OPRD instead applies autoregressive self-imitation to verified successes and Multi-Scope Reminder-conditioned KL to failures. The compared rows use the same training tasks, on-policy rollout budget, Top-100 support, and distillation temperature.

The teacher-refresh ablation compares a frozen teacher, periodic refresh schedules, and shared teacher–student parameters. The frozen teacher remains at its initial checkpoint. Our four-refresh setting copies the current student parameters to the teacher at 20%, 40%, 60%, and 80% of training. Under shared parameters, the teacher context is evaluated with the current student parameters at every update while gradients remain stopped through the teacher branch.

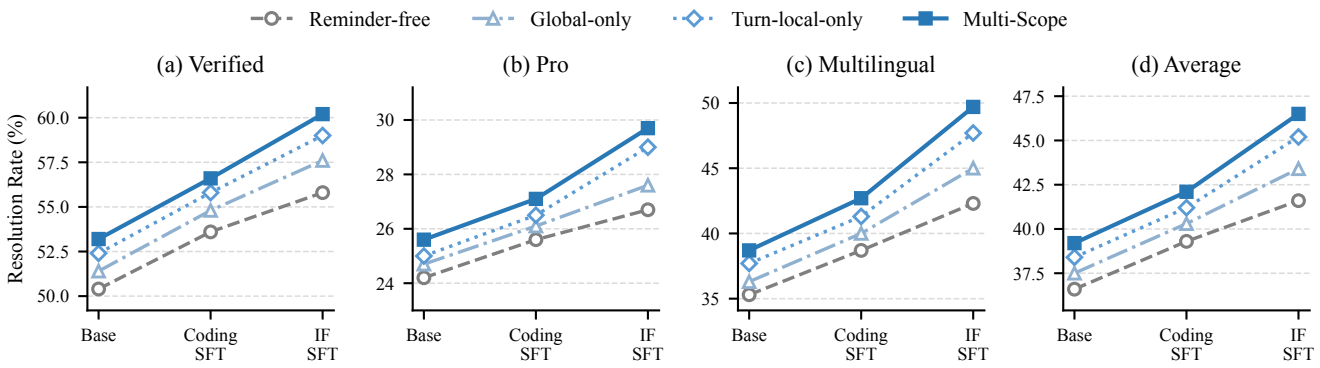

Figure 5: Effect of Reminder scope and supervised preparation on Qwen3.5-4B. We compare reminder-free, Global-only, Turn-local-only, and Multi-Scope Reminder conditioning for the Base, Coding SFT, and IF SFT checkpoints.

Setting	SWE-bench Verified (%)	SWE-bench Pro (%)	SWE-bench Multilingual (%)	Average (%)
KL Support and Distillation Temperature				
Full Vocabulary, $\mathcal{T} = 0.6$	59.8	29.4	50.3	46.5
Top-100, $\mathcal{T} = 0.6$	60.2	29.7	50.7	46.9
Top-100, $\mathcal{T} = 1.0$ (Ours)	**61.2**	**30.5**	**52.0**	**47.9**
Top-100, $\mathcal{T} = 1.5$	<u>60.4</u>	<u>29.8</u>	<u>51.0</u>	<u>47.1</u>
Outcome-Based Routing and KL Direction				
All-Rollout OPSD, Forward KL	58.8	28.6	48.3	45.2
All-Rollout OPSD, Reverse KL	59.6	29.1	49.3	46.0
Outcome-Routed OPRD, Forward KL	<u>60.2</u>	<u>29.7</u>	<u>50.3</u>	<u>46.7</u>
Outcome-Routed OPRD, Reverse KL (Ours)	**61.2**	**30.5**	**52.0**	**47.9**

Table 3: Ablations of KL support, distillation temperature, outcome-based routing, and KL direction with Qwen3.5-4B. Scores are task resolution rates (%), and *Average* is the unweighted mean across the three benchmarks. The best and second-best results within each group are shown in bold and underlined, respectively.

Effect of KL support and distillation temperature. The KL-support ablation compares full-vocabulary and Top-100 supports and varies the distillation temperature to identify an effective distributional target for OPRD. At the shared temperature of $\mathcal{T} = 0.6$, Top-100 and full-vocabulary KL perform similarly, with Top-100 obtaining a slightly higher average score. This indicates that a sparse high-probability support is sufficient for effective distillation in this setting. Among the Top-100 variants, $\mathcal{T} = 1.0$ performs best with an average score of 47.9, while both lower and higher temperatures perform worse, suggesting that a moderate temperature balances teacher preferences with useful alternatives.

Effect of outcome-based routing and KL direction. A paired comparison of All-Rollout OPSD and Outcome-Routed OPRD under forward and reverse KL isolates the effects of outcome-based routing and KL direction. Reverse KL outperforms forward KL under both training strategies. At a fixed KL direction, OPRD consistently improves over OPSD, with a 1.9-point gain under reverse KL. These results support assigning verified successes to autoregressive self-imitation and failures to Reminder-conditioned distillation, rather than applying a single KL objective uniformly to all rollouts. Reverse KL further focuses the update on candidates that the student assigns high probability but the Reminder-conditioned teacher considers inappropriate.

$\lambda_+ : \lambda_-$	Ver.	Pro	Multi.	Avg.
0.25:1	59.6	28.7	49.7	46.0
0.5:1	60.4	29.5	51.0	47.0
1:1 (Ours)	**61.2**	**30.5**	**52.0**	**47.9**
2:1	60.8	29.8	51.3	47.3
4:1	59.8	28.9	50.3	46.3

Table 4: Ablation of the relative weights assigned to successful-trajectory autoregressive self-imitation (λ_+) and failed-trajectory Reminder-conditioned distillation (λ_-). Columns report SWE-bench Verified, Pro, Multilingual, and their average.

Balancing successful and failed trajectories. The balanced 1:1 setting performs best on all three benchmarks and obtains a 47.9 average. Reducing λ_+ over-concentrates training on correcting failed states and weakens retention of successful behavior already discovered by the current policy. Increasing λ_+ gradually shifts the objective toward successful-

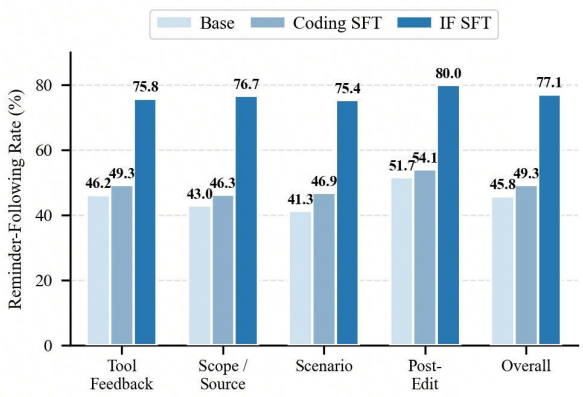

Figure 6: Reminder-following rates on 200 SWE-REBENCH V2 validation tasks. After a category-specific error is detected, the corresponding Reminder is inserted at the relevant interaction point, and the subsequent behavior is judged by whether it corrects or avoids that error. Each task–category pair contributes at most one successful follow event. *Overall* aggregates the four functional categories.

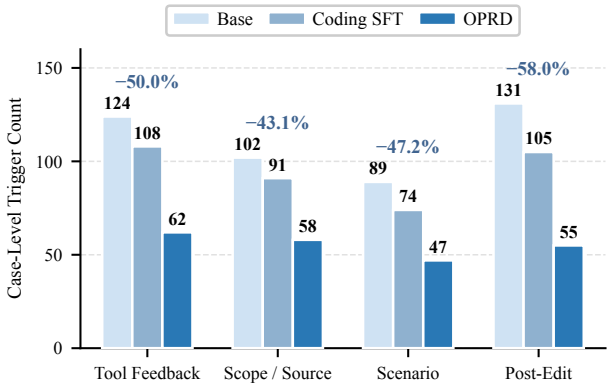

Figure 7: Diagnostic error reduction on 200 SWE-REBENCH V2 validation tasks. Bars show case-level trigger counts for Qwen3.5-4B (Base), Coding SFT, and OPRD. Percentages report relative reductions from Qwen3.5-4B to OPRD.

trajectory self-imitation and underuses the process supervision supplied by Reminders on failed states. SWE-BENCH PRO and SWE-BENCH MULTILINGUAL are more sensitive to this imbalance because they contain more tasks that require correction from tool feedback and cross-file context.

Effect of the teacher refresh schedule. Varying the teacher refresh schedule examines the trade-off between adapting to the evolving student and maintaining a stable distillation target. As shown in Fig. 4, periodic refreshes exhibit a non-monotonic trend, with four evenly spaced refreshes performing best on every benchmark and reaching an average resolution rate of 47.9. This pattern suggests that a frozen teacher may become misaligned with states visited by the evolving student, whereas overly frequent refreshes may reduce target stability. The Shared Parameters setting performs worst, further supporting the importance of a relatively stable teacher target for Reminder-conditioned distillation.

Effect of Reminder scope. A controlled comparison of four Reminder conditions across the Base, Coding SFT, and IF SFT checkpoints measures Reminder utility and model responsiveness. Both Global-only and Turn-local-only conditioning improve every checkpoint on all benchmarks, while combining them consistently performs best. Turn-local Reminders generally provide larger gains, while the additional improvement from Global Reminders shows that persistent cross-turn guidance complements feedback aligned with a specific environment observation. The average benefit of Multi-Scope Reminders grows across preparation stages and reaches 4.9 points after IF SFT, indicating that IF SFT improves the model's ability to use corrective information at different scopes.

Effect of IF SFT on Reminder responsiveness. As shown in Fig. 6, the aggregate Reminder-following rate increases from 45.8 for the open-source base model to 49.3 after Coding SFT, and then to 77.1 after IF SFT. Relative to Coding SFT, IF SFT improves the rate by 27.8 points overall and by 25.9–30.4 points across the four functional categories. Its category-level rates are consistently high, ranging from 75.4 for Scenario-Condition Reasoning to 80.0 for Post-Edit Control and Validation. The largest gain over Coding SFT occurs for Scope and Source Grounding, which increases from 46.3 to 76.7. These results isolate behavioral adherence to an inserted Reminder rather than end-to-end task resolution, and support using the IF checkpoint as the shared teacher–student initialization for OPRD.

Effect of functional Reminder categories. Grouping Reminders by content rather than temporal scope evaluates the contribution of each functional category. All four categories in Fig. 8 outperform reminder-free evaluation, with Post-Edit Control and Validation providing the strongest individual result. Enabling all categories raises the average resolution rate from 36.6 to 39.2, suggesting that they provide complementary corrective signals and broader coverage when combined.

Process-level error reduction. Applying the same four-category taxonomy as Fig. 1 tests whether OPRD's task-level gains coincide with fewer process errors. Each checkpoint generates one reminder-free trajectory for the same 200 SWE-REBENCH V2 validation tasks under identical inference settings, and we record case-level triggers from the resulting trajectories. As shown in Fig. 7, the total number of triggers falls from 446 for Qwen3.5-4B to 378 after Coding SFT and 222 after OPRD, a 50.2% reduction relative to Qwen3.5-4B. Trigger counts decline by 43.1–58.0% across categories, with the largest decrease in Post-Edit Control and Validation. These results show that OPRD improves both final task resolution and intermediate process behavior.

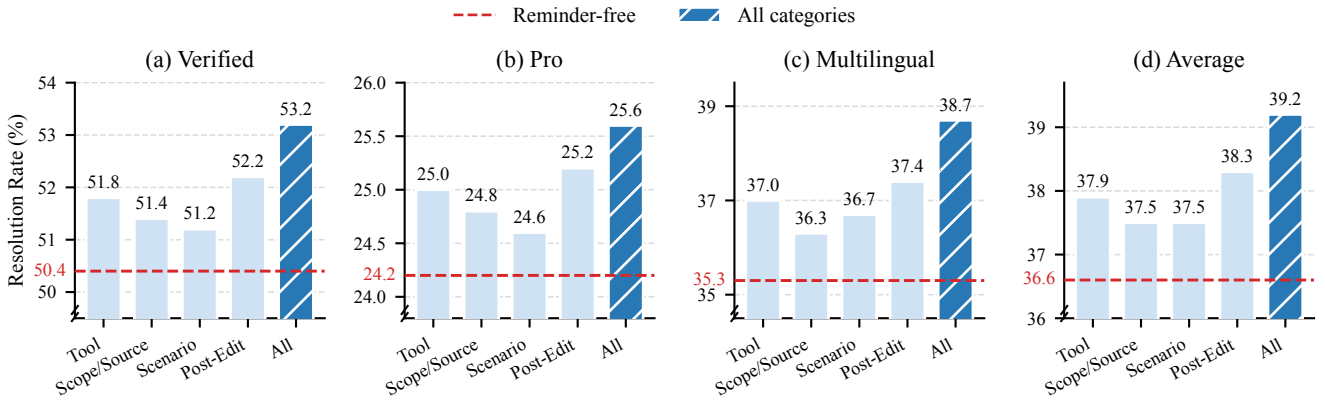

Figure 8: Effect of functional Reminder categories on the Qwen3.5-4B base checkpoint. We evaluate Tool-Feedback Recovery, Scope and Source Grounding, Scenario-Condition Reasoning, and Post-Edit Control and Validation, together with their combination. The dashed line denotes reminder-free performance.

5 Conclusion

Long-horizon coding agents must preserve task constraints and respond to interaction feedback while producing correct code. Our diagnostic analysis highlights process-level failures throughout repository-level problem solving. We introduce OPRD, which uses Multi-Scope Reminders as training-time privileged information and applies outcome-based routing to on-policy rollouts. A Global Reminder converts final-test diagnostics into trajectory-level guidance, while Turn-local Reminders anchor observation-derived corrections to the next assistant decision. OPRD routes successful rollouts to autoregressive self-imitation and failed rollouts to Reminder-conditioned reverse-KL distillation. Its two-stage supervised cold start, consisting of Coding SFT followed by Instruction-Following SFT, improves responsiveness to these corrective signals. Experiments show that OPRD improves task resolution and reduces adherence failures. By internalizing trajectory-level feedback and local corrections into model parameters, OPRD enables long-horizon coding agents to operate without Reminders at inference time.

References

Agarwal, R.; Vieillard, N.; Zhou, Y.; Stanczyk, P.; Ramos Garea, S.; Geist, M.; and Bachem, O. 2024. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. In *The Twelfth International Conference on Learning Representations*.

Badertdinov, I.; Golubev, A.; Nekrashevich, M.; Shevtsov, A.; Karasik, S.; Andriushchenko, A.; Trofimova, M.; Litvintseva, D.; and Yangel, B. 2025. SWE-rebench: An Automated Pipeline for Task Collection and Decontaminated Evaluation of Software Engineering Agents. In *Advances in Neural Information Processing Systems*, volume 38.

Badertdinov, I.; Nekrashevich, M.; Shevtsov, A.; and Golubev, A. 2026. SWE-rebench V2: Language-Agnostic SWE Task Collection at Scale. arXiv:2602.23866.

Chen, X.; Lin, M.; Schärli, N.; and Zhou, D. 2024. Teaching Large Language Models to Self-Debug. In *The Twelfth International Conference on Learning Representations*.

Dao, T. 2023. FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning. arXiv:2307.08691.

Deng, K.; Zhan, Z.; Xiang, W.; Zhu, W.; Peng, T.; Lei, X.; et al. 2025a. HiPO: Hybrid Policy Optimization for Dynamic Reasoning in LLMs. *arXiv preprint arXiv:2509.23967*.

Deng, X.; Da, J.; Pan, E.; He, Y. Y.; Ide, C.; Garg, K.; Lauffer, N.; Park, A.; Pasari, N.; Rane, C.; Sampath, K.; Krishnan, M.; Kundurthy, S.; Hendryx, S.; Wang, Z.; Bharadwaj, V.; Holm, J.; Aluri, R.; Zhang, C. B. C.; Jacobson, N.; Liu, B.; and Kenstler, B. 2025b. SWE-Bench Pro: Can AI Agents Solve Long-Horizon Software Engineering Tasks? arXiv:2509.16941.

Gou, Z.; Shao, Z.; Gong, Y.; Shen, Y.; Yang, Y.; Duan, N.; and Chen, W. 2024. CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing. In *The Twelfth International Conference on Learning Representations*.

Gu, Y.; Dong, L.; Wei, F.; and Huang, M. 2024. MiniLLM: Knowledge Distillation of Large Language Models. In *The Twelfth International Conference on Learning Representations*.

Hübotter, J.; Lübeck, F.; Behric, L.; Baumann, A.; Bagatella, M.; Marta, D.; Hakimi, I.; Shenfeld, I.; Kleine Buening, T.; Guestrin, C.; and Krause, A. 2026. Reinforcement Learning via Self-Distillation. arXiv:2601.20802.

Jacobs, S. A.; Tanaka, M.; Zhang, C.; Zhang, M.; Song, S. L.; Rajbhandari, S.; and He, Y. 2023. DeepSpeed Ulysses: System Optimizations for Enabling Training of Extreme Long Sequence Transformer Models. arXiv:2309.14509.

Jain, N.; Singh, J.; Shetty, M.; Zhang, T.; Zheng, L.; Sen, K.; and Stoica, I. 2025. R2E-Gym: Procedural Environment Generation and Hybrid Verifiers for Scaling Open-Weights SWE Agents. In *Second Conference on Language Modeling*.

Jimenez, C. E.; Yang, J.; Wettig, A.; Yao, S.; Pei, K.; Press, O.; and Narasimhan, K. R. 2024. SWE-bench: Can Language

Models Resolve Real-World GitHub Issues? In *The Twelfth International Conference on Learning Representations*.

Jin, W.; Min, T.; Yang, Y.; Wei, D.; Zhou, Y.; Kadhe, S. R.; Baracaldo, N.; and Lee, K. 2026. Entropy-Aware On-Policy Distillation of Language Models. arXiv:2603.07079.

Ko, J.; Kim, S.; Chen, T.; and Yun, S.-Y. 2024. DistiLLM: Towards Streamlined Distillation for Large Language Models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 24872–24895. PMLR.

Li, Y.; Lyu, M. R.; and Wang, L. 2025. Learning to Reason from Feedback at Test-Time. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 5241–5253. Vienna, Austria: Association for Computational Linguistics.

Ma, Q.; Zheng, Y.; Shi, Z.; Zhao, Z.; Jia, B.; Huang, Z.; Lin, Z.; Li, Y.; Yang, J.; Peng, Y.; Zhang, Z.; and Liu, X. 2025. VeOmni: Scaling Any Modality Model Training with Model-Centric Distributed Recipe Zoo. arXiv:2508.02317.

Madaan, A.; Tandon, N.; Gupta, P.; Hallinan, S.; Gao, L.; Wiegreffe, S.; Alon, U.; Dziri, N.; Prabhumoye, S.; Yang, Y.; Gupta, S.; Majumder, B. P.; Hermann, K.; Welleck, S.; Yazdanbakhsh, A.; and Clark, P. 2023. Self-Refine: Iterative Refinement with Self-Feedback. In *Advances in Neural Information Processing Systems*, volume 36, 46534–46594.

Miao, Z.; Jia, H.; Li, L.; Qian, C.; Xiong, Y.; Yan, W.; and Shao, J. 2026. Seeing with You: Perception-Reasoning Coevolution for Multimodal Reasoning. *arXiv preprint arXiv:2603.28618*.

OpenAI. 2024. Introducing SWE-bench Verified.

OpenAI. 2026. Introducing GPT-5.4.

Ouyang, S.; Yu, W.; Ma, K.; Xiao, Z.; Zhang, Z.; Jia, M.; Han, J.; Zhang, H.; and Yu, D. 2025. RepoGraph: Enhancing AI Software Engineering with Repository-level Code Graph. In *International Conference on Learning Representations*, volume 2025, 30098–30121.

Pan, J.; Wang, X.; Neubig, G.; Jaitly, N.; Ji, H.; Suhr, A.; and Zhang, Y. 2025. Training Software Engineering Agents and Verifiers with SWE-Gym. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, 47717–47737. PMLR.

Penaloza, E.; Vattikonda, D.; Gontier, N.; Lacoste, A.; Charlin, L.; and Caccia, M. 2026. Privileged Information Distillation for Language Models. arXiv:2602.04942.

Qwen Team. 2026. Qwen3.5: Towards Native Multimodal Agents.

Shinn, N.; Cassano, F.; Gopinath, A.; Narasimhan, K.; and Yao, S. 2023. Reflexion: Language Agents with Verbal Reinforcement Learning. In *Advances in Neural Information Processing Systems*, volume 36, 8634–8652.

Snell, C.; Klein, D.; and Zhong, R. 2022. Learning by Distilling Context. arXiv:2209.15189.

Wei, Y.; Duchenne, O.; Copet, J.; Carbonneaux, Q.; Zhang, L.; Fried, D.; Synnaeve, G.; Singh, R.; and Wang, S. I. 2025. SWE-RL: Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution. In *Advances in Neural Information Processing Systems*, volume 38, 78500–78525.

Xia, C. S.; Deng, Y.; Dunn, S.; and Zhang, L. 2025. Demystifying LLM-Based Software Engineering Agents. *Proceedings of the ACM on Software Engineering*, 2(FSE): 801–824.

Xu, J.; Deng, K.; Li, W.; Yu, S.; Tang, H.; Huang, H.; et al. 2025. SWE-Compass: Towards Unified Evaluation of Agentic Coding Abilities for Large Language Models. *arXiv preprint arXiv:2511.05459*.

Xu, J.; Wu, G.; Wu, Y.; Mou, Y.; Yu, S.; He, T.; et al. 2026a. FLARE: A Full-Lifecycle Dense Supervision Paradigm for Long-Horizon Coding Agents via Generative Reward Model. *arXiv preprint arXiv:2609.23808*.

Xu, Z.; Xie, H.; Miao, Z.; Gong, W.; Qian, C.; and Li, L. 2026b. Stable Adaptive Thinking via Advantage Shaping and Length-Aware Gradient Regulation. *arXiv preprint arXiv:2602.22556*.

Yang, J.; Jimenez, C. E.; Wettig, A.; Lieret, K.; Yao, S.; Narasimhan, K. R.; and Press, O. 2024. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. In *Advances in Neural Information Processing Systems*, volume 37, 50528–50652.

Yang, J.; Lieret, K.; Jimenez, C. E.; Wettig, A.; Khandpur, K.; Zhang, Y.; Hui, B.; Press, O.; Schmidt, L.; and Yang, D. 2025. SWE-smith: Scaling Data for Software Engineering Agents. In *Advances in Neural Information Processing Systems*, volume 38.

Yao, S.; Zhao, J.; Yu, D.; Du, N.; Shafran, I.; Narasimhan, K.; and Cao, Y. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *The Eleventh International Conference on Learning Representations*.

Ye, T.; Dong, L.; Chi, Z.; Wu, X.; Huang, S.; and Wei, F. 2025. Black-Box On-Policy Distillation of Large Language Models. arXiv:2511.10643.

Ye, T.; Dong, L.; Wu, X.; Huang, S.; and Wei, F. 2026. On-Policy Context Distillation for Language Models. arXiv:2602.12275.

Zan, D.; Huang, Z.; Liu, W.; Chen, H.; Zhang, L.; Xin, S.; Chen, L.; Liu, Q.; Zhong, X.; Li, A.; Liu, S.; Xiao, Y.; Chen, L.; Zhang, Y.; Su, J.; Liu, T.; Long, R.; Shen, K.; and Xiang, L. 2025. Multi-SWE-bench: A Multilingual Benchmark for Issue Resolving. arXiv:2504.02605.

Zhan, Z.; Deng, K.; Tang, H.; Xiang, W.; Wu, K.; Li, W.; et al. 2025a. KAT-V1: Kwai-AutoThink Technical Report. *arXiv preprint arXiv:2507.08297*.

Zhan, Z.; Deng, K.; Zhang, X.; Wang, J.; Tang, H.; Lai, Z.; et al. 2025b. KAT-Coder Technical Report. *arXiv preprint arXiv:2510.18779*.

Zhang, Y.; Ruan, H.; Fan, Z.; and Roychoudhury, A. 2024. AutoCodeRover: Autonomous Program Improvement. In *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*, 1592–1604. ACM.

Zhao, S.; Xie, Z.; Liu, M.; Huang, J.; Pang, G.; Chen, F.; and Grover, A. 2026. Self-Distilled Reasoner: On-Policy Self-Distillation for Large Language Models. arXiv:2601.18734.

Zhao, Y.; Gu, A.; Varma, R.; Luo, L.; Huang, C.-C.; Xu, M.; Wright, L.; Shojanazeri, H.; Ott, M.; Shleifer, S.; Desmaison, A.; Balioglu, C.; Damania, P.; Nguyen, B.; Chauhan, G.; Hao, Y.; Mathews, A.; and Li, S. 2023. PyTorch FSDP: Experiences on Scaling Fully Sharded Data Parallel. arXiv:2304.11277.

Zhong, L.; Wang, Z.; and Shang, J. 2024. Debug like a Human: A Large Language Model Debugger via Verifying Runtime Execution Step by Step. In *Findings of the Association for Computational Linguistics: ACL 2024*, 851–870. Bangkok, Thailand: Association for Computational Linguistics.

Zhou, A.; Yan, K.; Shlapentokh-Rothman, M.; Wang, H.; and Wang, Y.-X. 2024. Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 62138–62160. PMLR.

A Appendix

In this appendix, we provide training implementation and hyperparameters (Section B), together with supplementary materials on Reminder construction and diagnostic annotation. **Section C** details the construction procedures, functional taxonomy, Global Reminder evidence processing, and complete runtime category-judge prompt. **Section D** describes the deterministic and model-based annotation protocol used to analyze coding-agent errors. Training data, evaluation settings, results, and ablations are presented in Section 4.

B Training Implementation and Hyperparameters

We implement OPRD using VeOmni (Ma et al. 2025) with FlashAttention 2 (Dao 2023), fused MoE operators, FSDP2 (Zhao et al. 2023), and Ulysses sequence parallelism (Jacobs et al. 2023). Distributed checkpoints are stored with DCP. Table 5 lists the supplied Qwen3.5-4B configuration.

The main Qwen3.5-4B configuration uses a maximum sequence length of 131,072, a global batch size of 32, 3 training epochs, and a peak learning rate of 1×10^{-5} with cosine decay. OPRD uses Top-100 reverse KL at temperature $\mathcal{T} = 1.0$.

Teacher refresh. The teacher is not optimized by gradient descent. At 20%, 40%, 60%, and 80% of OPRD training, its parameters are replaced by the current student parameters and then kept fixed until the next refresh. This schedule yields four teacher refreshes over a complete run.

Top-k KL support. The main configuration evaluates reverse KL on one shared Top-100 support at each valid assistant-token position. After temperature scaling, we compute a joint importance score $s(v) = \max\{\log p_\theta(v), \log p_\phi(v)\}$ for every vocabulary item v, where p_θ and p_ϕ denote the student and stop-gradient teacher distributions. We select the 100 tokens with largest $s(v)$, then renormalize both distributions on exactly this common index set before evaluating reverse KL. Selecting by the pointwise maximum retains tokens considered plausible by either model while keeping the support size fixed.

The main loss uses $\lambda_+ = 1$ for autoregressive self-imitation on verified successful trajectories and $\lambda_- = 1$ for Reminder-conditioned distillation on failed trajectories.

C Reminder Construction Details

C.1 Construction Procedure and Functional Taxonomy

Reminder construction and Reminder categorization describe two distinct parts of our pipeline. Construction specifies how a candidate is obtained, filtered, inserted, and verified, whereas the functional taxonomy specifies which coding-agent failure the retained Reminder addresses. These two attributes are independent of the temporal scope used by OPRD: a Global Reminder provides persistent final-test guidance, while the runtime construction procedure below is used to construct and evaluate observation-anchored Turn-local Reminders.

We construct runtime Reminder candidates from SWE-REBENCH V2 tasks and naturally occurring coding-agent trajectories. For each trajectory, we align the task requirement, assistant action, tool call, tool result, next assistant action, and subsequent validation outcome. A candidate is retained only when its trigger is observable without hidden correctness labels, its desired next action is local and behaviorally verifiable, and its text neither reveals a reference patch nor invents a new task requirement. Correctness is determined separately by the executable ReBench evaluator; a trajectory alone is used only as behavioral evidence.

Passive and active construction modes. A *passive* Reminder is appended immediately after a tool result exposes a concrete failure, ambiguity, scope drift, or validation signal. It is one-shot: the next assistant decision observes it, but it does not automatically persist to later turns. An *active* Reminder is evaluated at a controlled pre-action boundary. We first obtain a probe action, use it only to determine whether the trigger fires, discard it when an intervention is required, insert the Reminder into the same visible state, and reroll the action. The active verifier then evaluates the rerolled tool choice, arguments, target path, edit scope, and immediate follow-up behavior. Thus passive construction mines recovery from observed feedback, whereas active construction provides a matched intervention for measuring causal behavior change.

Runtime selection and insertion. At each legal insertion boundary, deterministic hard rules are matched against the latest tool result before a model judge considers non-hard contextual strategies. Non-hard candidates must reach confidence 0.90 and satisfy their phase gate. We then apply a per-family cap, a per-case budget, repeated-text removal, and hard-error-signature deduplication. Post-edit candidates require a visible edit, and validation-wrap-up candidates require concrete test, build, compiler, diff, formatting, or completion evidence. If all candidates fail these gates, no Reminder is inserted; the system does not use a fallback category. Hard-rule text is copied from the rubric, and every inserted Reminder is serialized after, rather than inside, the corresponding tool response.

We separately record whether trajectory capture completed, a Reminder was inserted, the required local behavior was followed, and the final source-only patch passed executable verification. These quantities use different denominators and are not substituted for one another.

C.2 Global Reminder Construction and Evidence Judge

Global Reminders are constructed only after a complete coding attempt has been rejected by executable verification. Unlike a Turn-local Reminder, which is anchored to one observation and governs the next decision, a Global Reminder summarizes the previous patch and its verifier artifacts and is placed at the input of the next full retry. A candidate is eligible only if it contains a non-empty patch, receives zero benchmark reward, is not marked non-retryable, is successfully applied in the verification environment, and is not accepted

Table 5: Training configuration for the Qwen3.5-4B OPRD experiment.

Parameter	Value	Parameter	Value
Model	Qwen3.5-4B	Maximum sequence length	131,072
Training framework	VeOmni	LR start	0
Epochs	3	GPU	16 NVIDIA H20
Global batch size	32	Micro batch size	1
Peak learning rate	1×10^{-5}	Minimum learning rate	1×10^{-6}
LR schedule	Cosine	Warmup ratio	0.005
LR decay ratio	1	Weight decay	0
Distillation temperature	1.0	KL direction	Reverse
KL vocabulary support	Top-100	Teacher refresh	20/40/60/80%
Self-imitation weight λ_+	1.0	Distillation weight λ_-	1.0
Expert parallel size	1	Ulysses size	4
Data-parallel shard size	2	Attention implementation	FlashAttention 2

Table 6: The four reporting classes used for functional Reminder analysis. The mapping is a fixed partition of the 32 original rubric categories.

ID	Functional Class	Original Rubric Source	Count
R1	Tool-Feedback Recovery	Deterministic hard rules	10
R2	Scope and Source Grounding	Scope/patch and source/path slots	10
R3	Scenario-Condition Reasoning	Scenario/condition slot	5
R4	Post-Edit Control and Validation	Post-edit drift and validation/wrap-up slots	7
	Total	10 hard rules and 22 non-hard strategies	32

by the benchmark. These gates exclude empty or unapplied patches, successful attempts, and infrastructure-only records.

The evidence judge parses the benchmark result JSON and raw test log, records available pass/pending/fail counts and touched files, deduplicates failures with a shared root cause, and selects at most five signals. Root-cause-bearing interface, argument, syntax, reference, symbol, path, and signature failures are ranked before assertion, command, and generic failures. The resulting `verification_feedback` states what failed and where the evidence came from, but does not reveal a reference patch or prescribe an unsupported implementation. It is visible to the teacher throughout the retry trajectory.

Example. Suppose a patch applies cleanly but leaves the targets `no_system_header_includes` and `test_wrap_static_fns` failing, with three neighboring targets passing. The generated Global Reminder reports the successful patch application, the 3/2 pass–fail summary, the touched files, and the two unresolved targets, then asks the agent to re-read the task and reconstruct a consistent fix from the implicated function signatures and command flow. This evidence distinguishes an implementation failure from a malformed patch without exposing the hidden solution.

C.3 Four-Class Functional Taxonomy

The finalized rubric contains 32 fine-grained Reminder categories: 10 deterministic hard rules and 22 model-judged contextual strategies organized in five family slots. We map them into four mutually exclusive reporting classes without introducing new fine-grained categories. After the most specific original category is selected for an error event, its coarse class is fixed by Table 6. A trajectory may contain multiple independent events, but a single event receives exactly one most-specific category; multiple Reminders are allowed only when the visible evidence supports distinct risks and distinct next actions.

R1: Tool-Feedback Recovery. This class covers commands, searches, edits, compilation, paths, validation conditions, and toolchain failures that can be identified deterministically from the latest tool response.

R2: Scope and Source Grounding. This class combines constraints on the initial patch surface with evidence about the true production path, caller, source value, and data flow.

R3: Scenario-Condition Reasoning. This class covers branch, default, fallback, cache, classification, and one-condition repairs. It separates the failing scenario from sibling scenarios that must retain their previous behavior.

R4: Post-Edit Control and Validation. This class applies after a production edit. It prevents unsupported expansion of the patch and uses concrete assertions, tests, compiler output, diffs, and formatting signals to complete a local repair.

C.4 Complete Runtime Judge Prompt

The deployed API judge receives only the task, recent visible trajectory, latest tool result, and previously inserted Reminders. Figures 9–13 present the complete prompt instantiated from the runtime YAML, including all categories, gates, caps, and the JSON output contract.

D Diagnostic Error Annotation

For each task, deterministic rules first identify failures exposed by explicit tool and execution signals. GPT-5.4 then receives the original problem and the complete agent trajectory to classify the remaining context-dependent failures. The combined annotations use Tool-Feedback Recovery, Scope and Source Grounding, Scenario-Condition Reasoning, and Post-Edit Control and Validation. For case-level analyses, each fine-grained category is counted at most once per trajectory. This post-hoc analysis is distinct from the online judge in Fig. 9: the online judge operates only on the current visible prefix and decides whether to intervene, whereas diagnostic annotation audits a completed trajectory for aggregate error statistics.

Runtime Reminder category-judge prompt

```
You are the runtime error-category classifier and Reminder judge for a
repository-level coding-agent trajectory. Use only the visible task, recent
history, tool calls, tool results, edits, build feedback, tests, diffs, and
previous Reminders supplied in the input.

OBJECTIVES
1. Decide whether the latest observation and recent trajectory justify one or
   more existing Reminder categories from the rubric below.
2. For each independently justified event, select exactly one most-specific
   original category.
3. Assign its coarse_class from the fixed R1--R4 mapping.
4. Produce one concise, imperative, evidence-grounded Reminder.
5. Return no decision when visible evidence is insufficient.

Never create or rename a category. Do not use hidden chain of thought, final
oracle labels, reference patches, or dataset-specific answer rules. GRM is
disabled. There is no fallback insertion.

RUNTIME CONFIGURATION
- provider: vllm; judge temperature: 0.6
- multi_strategy: true; strategy_selection_mode: family_cap
- max_nonhard_per_case: 5; min_confidence: 0.90
- hard_error_text_dedupe: true
- dedupe_text_per_case: false; dedupe_category_per_case: false
- eligibility_mode: all; disable_grm: true; disable_gpt55_judge: false

INPUT CONTRACT
{
  "task_context": "visible repository-level request",
  "recent_history": ["visible messages, calls, results, edits, and tests"],
  "latest_tool_result": {
    "index": 0,
    "tool_name": "Read|Grep|Glob|Edit|MultiEdit|Write|Bash|other",
    "call_id": "visible identifier",
    "occurrence": 1,
    "output_text": "complete visible output"
  },
  "previous_reminders": ["previous category and text records"]
}
Treat absent information as unknown. Never invent paths, symbols, callers,
conditions, literals, defaults, thresholds, tests, or validation results.

FIXED COARSE-CLASS MAP
R1_TOOL_FEEDBACK_RECOVERY:
  hard_go_fmt_cross_dir, hard_ambiguous_edit_old_string,
  hard_go_signature_callers, hard_semantic_search_auth,
  hard_malformed_search, hard_mockey_gcflags, hard_go_toolchain_env,
  hard_noisy_validation_output_narrow, hard_go_compile_error,
  hard_missing_file_path_narrow.

R2_SCOPE_SOURCE_GROUNDING:
  selective_scope_contract, scope_contract_gate,
  polarity_neutral_scope_contract, pre_edit_patch_surface_contract,
  production_before_test_guard, scenario_source_scope_guard,
  named_path_source_guard, named_path_over_neighbor,
  prove_value_source_before_shared_change, read_failing_test_first.

R3_SCENARIO_CONDITION_REASONING:
```


Figure 9: Complete runtime Reminder category-judge prompt (Part 1 of 5).

```
  scenario_matrix_guard, sibling_scenario_invariant,
  static_branch_pair_read_mature, one_condition_repair,
  strict_one_condition_literal.

R4_POST_EDIT_VALIDATION_CONTROL:
  change_surface_budget_guard, edit_scope_evidence_guard,
  safe_default_or_data_model_path_guard, assertion_diff_local_repair,
  test_failure_localization_guard, compile_mechanical_only,
  diff_wrapup_format_test.

Once a category is selected, copy its coarse class from this map. A single
event has one category. Multiple decisions require distinct evidence and
distinct immediate actions.

HARD RULES -- EVALUATE FIRST
A hard rule requires an explicit signal in the latest tool result, uses
confidence 1.0, and copies the corresponding action below.

hard_go_fmt_cross_dir
Trigger: go fmt names files from multiple directories.
Action: use gofmt -w on named files or run go fmt separately per directory.

hard_ambiguous_edit_old_string
Trigger: Edit fails because old_string matches multiple locations.
Action: retry with a longer old_string anchored by surrounding code; avoid broad
replace_all unless every occurrence must change.

hard_go_signature_callers
Trigger: Go compiler output shows caller/interface mismatch after signature edits.
Action: search and update affected implementations and call sites before narrow validation.

hard_semantic_search_auth
Trigger: a discovery helper fails with authentication, identity, or permission errors.
Action: stop retrying it; use reliable paths, symbols, local rg/find, and targeted reads.

hard_malformed_search
Trigger: rg/Grep fails because of malformed pattern, escaping, type, or arguments.
Action: retry with a simpler quoted or fixed-string search, or use find, before editing.

hard_mockey_gcflags
Trigger: a Mockey-related Go test fails without visible -gcflags="all=-N -l" evidence.
Action: rerun the narrow test with those gcflags before treating it as a code bug.

hard_go_toolchain_env
Trigger: Go validation is blocked by a concrete toolchain/version mismatch.
Action: treat it as an environment blocker; use the repository-required toolchain
and do not keep editing product code from that signal.

hard_noisy_validation_output_narrow
Trigger: validation output is truncated or dominated by noise.
Action: use the first concrete FAIL, panic, assertion, or compiler diagnostic;
otherwise switch to narrower validation without changing product code.

hard_go_compile_error
Trigger: Go build/test output contains concrete compiler diagnostics.
Action: fix only reported file/line issues first, then rerun narrow validation.

hard_missing_file_path_narrow
```


Figure 10: Complete runtime Reminder category-judge prompt (Part 2 of 5).

```
Trigger: Read/Edit/stat/open fails because a concrete path does not exist.
Action: correct only the path using visible parent/import/filename evidence and
keep the task target fixed.

Do not repeat the same hard category for the same normalized error signature.

NON-HARD FAMILY 1 -- scope_patch_slot -- maximum one per case
selective_scope_contract
Trigger: before risky production editing, multiple paths/layers/shared surfaces
or source/sibling uncertainty are visible.
Action: choose one evidence-backed production path and stable sibling/source;
keep the first patch there unless read evidence requires expansion.

scope_contract_gate
Trigger: before the first substantial edit, concrete path/source/sibling risk
exists but no single-path scope contract has been established.
Action: bind the fix to one evidence-backed path, name the stable condition, and
keep the first patch local.

polarity_neutral_scope_contract
Trigger: a scope contract is needed but read evidence does not prove boolean,
default, version, or threshold polarity.
Action: bind path and stable condition without deciding polarity from task wording.

pre_edit_patch_surface_contract
Trigger: before the first edit, the plan spans files, packages, layers, tests,
go.mod, generated files, shared models, common context, or broad helpers.
Action: budget one production file/function first and expand only with source proof.

production_before_test_guard
Trigger: tests, fixtures, or go.mod are edited/planned before the production
contract is fixed, or test edits dominate production edits.
Action: return to production first; edit tests only for test-only tasks or when an
already-read production contract requires fixture synchronization.

NON-HARD FAMILY 2 -- scenario_condition_slot -- maximum one per case
scenario_matrix_guard
Trigger: a branch/classification/fallback/default/cache/error edit is likely and
sibling scenarios are visible.
Action: build a small matrix with one changing scenario and one protected sibling;
use a narrow condition.

sibling_scenario_invariant
Trigger: a local fix risks breaking same-function or same-test sibling behavior.
Action: identify and preserve the sibling invariant before editing.

static_branch_pair_read_mature
Trigger: read/search evidence exposes failing and sibling scenarios but does not
separate branch/default/version/threshold polarity.
Action: separate them using read evidence; if only one side is evidenced, guard
narrowly for it and preserve the other.

one_condition_repair
Trigger: an existing condition/value separates the failing scenario from a sibling.
Action: name both scenarios and the exact separator; change only that condition.

strict_one_condition_literal
Trigger: a visible literal/enum/field/fallback/assertion uniquely justifies a repair.
```


Figure 11: Complete runtime Reminder category-judge prompt (Part 3 of 5).

```
Action: use only the visible separator; invent no threshold, default, version, or polarity.

NON-HARD FAMILY 3 -- source_path_slot -- maximum one per case
scenario_source_scope_guard
Trigger: a scenario edit is likely but the selecting source value/request field/
condition/config/caller is not pinned.
Action: identify failing scenario, source selector, and stable sibling; patch only
that production branch.

named_path_source_guard
Trigger: task/test/search/read evidence names a likely path, but the trajectory
drifts to a neighbor/shared helper without source proof.
Action: prefer the named path; prove the neighbor through caller/source/data flow
before editing it.

named_path_over_neighbor
Trigger: evidence points to a task-specific implementation while the next action
follows a similarly named generic/neighbor implementation.
Action: refocus on the task-specific path and use neighbors only as reference.

prove_value_source_before_shared_change
Trigger: a shared model/context/helper/data shape may be edited before proving
which value, request field, or caller drives the failure.
Action: prove the value source/caller/request field before changing shared surfaces.

read_failing_test_first
Trigger: a failing test/subtest/fixture is named or discoverable, but production
editing is about to begin without reading it.
Action: read the failing test or closest fixture before the first production edit.

NON-HARD FAMILY 4 -- post_edit_drift_slot -- maximum one per case
This family requires a visible Edit, MultiEdit, or Write.

change_surface_budget_guard
Trigger: after an edit, more files/layers/helpers are planned without new evidence.
Action: stop expanding; keep the patch on the minimum already-edited production path.

edit_scope_evidence_guard
Trigger: after an edit, the patch drifts from the named task path, data model,
caller, or failure site.
Action: re-anchor to visible task/test/caller/read evidence and avoid unsupported edits.

safe_default_or_data_model_path_guard
Trigger: after an edit, the right area is touched but a required default, fallback,
data source, or caller-specific path is likely changed.
Action: preserve the required default/source and change only the evidenced field/path.

NON-HARD FAMILY 5 -- validation_wrapup_slot -- maximum one per case
This family requires concrete validation, compiler, build, diff, formatting, or
wrap-up evidence.

assertion_diff_local_repair
Trigger: feedback contains Expected/Actual, assertion, panic, named subtest, or
concrete build feedback.
Action: treat it as a local invariant, inspect the named test and producing branch,
and make the smallest production change for that scenario.

test_failure_localization_guard
```


Figure 12: Complete runtime Reminder category-judge prompt (Part 4 of 5).

```
Trigger: feedback names a failing test/file/function and the next action risks
broad or unrelated edits.
Action: localize the next edit to the named target and rerun narrow validation.

compile_mechanical_only
Trigger: compile diagnostics are visible and the model risks changing design/tests
or adding shared types to fix them.
Action: fix only the diagnostic in the edited file/package and rerun narrow validation.

diff_wrapup_format_test
Trigger: code was edited and the model appears ready to finish after diff,
formatting, validation, or near-completion evidence.
Action: run formatting and the narrowest relevant build/test before concluding.

SELECTION AND GATING
- Hard rules precede non-hard judgment.
- At most one category from each non-hard family and at most five non-hard
  Reminders per case.
- Non-hard confidence must be at least 0.90.
- Enforce both phase gates exactly.
- Prefer exact literals over generic scenario advice, named paths over generic
  scope, and concrete assertions/diagnostics over generic validation advice.
- Do not emit semantically duplicate decisions for one event.
- Reminder text must be concise, imperative, immediately actionable, and use
  only entities already visible in the input.
- If all candidates fail evidence, confidence, phase, family, budget, or
  duplicate gates, return no Reminder.

OUTPUT CONTRACT
Return JSON only:
{
  "decisions": [
    {
      "should_insert": true,
      "coarse_class": "R1_TOOL_FEEDBACK_RECOVERY|R2_SCOPE_SOURCE_GROUNDING|R3_SCENARIO_COND
        ITION_REASONING|R4_POST_EDIT_VALIDATION_CONTROL",
      "category": "one exact original category identifier",
      "decision_source": "hard_rule|model_judge",
      "confidence": 0.93,
      "evidence": "concise quotation or faithful visible-evidence summary",
      "reason": "why this exact category is independently justified",
      "reminder_text": "one concise imperative Reminder",
      "scope": "the immediate governed action/path/branch/validation step"
    }
  ],
  "no_decision_reason": "empty when decisions exist"
}

If nothing is justified, return exactly:
{
  "decisions": [],
  "no_decision_reason": "No hard rule or confidence-qualified non-hard strategy is
    justified by the visible trajectory."
}

Before returning, verify exact category spelling, fixed coarse mapping, visible
evidence, valid phase, available family slot and budget, and non-duplication.
```


Figure 13: Complete runtime Reminder category-judge prompt (Part 5 of 5).