PremiseSQL: Evaluating Premise Critique and Abstention in Text-to-SQL Systems

Shuo Lin[1] **Jinzhe Li**[1,3] **Chenlu Guo**[4] **Yanhan Li**[5] **Yuan Wu**[1]* **Yi Chang**[1,2,3]*
[1]School of Artificial Intelligence, Jilin University
[2]Engineering Research Center of Knowledge-Driven Human-Machine Intelligence, MOE, China
[3]International Center of Future Science, Jilin University
[4]Bilibili Inc.
[5]Xingchen AGI Lab, China Telecom Artificial Intelligence Technology (Beijing) Co., Ltd
{linshuo26,jinzhe25}@mails.jlu.edu.cn
guochenlu01@bilibili.com, liyh181@chinatelecom.cn
{yuanwu,yichang}@jlu.edu.cn

arXiv:2610.22158v1 [cs.CL] 15 Oct 2026

Abstract

Recent advances in large language models (LLMs) have substantially improved Text-to-SQL performance, yet existing NL2SQL systems largely assume that user queries are valid and answerable. As a result, models often generate executable SQL even for queries with flawed premises. We introduce *premise critique ability* in NL2SQL, namely the ability to recognize flawed queries and abstain from generating misleading SQL. To evaluate this capability, we propose **PremiseSQL**, a benchmark containing 1,800 flawed NL2SQL queries across four categories of premise failures, along with their 1,800 original queries. We further design an evaluation framework to measure intrinsic and prompt-induced premise-aware behavior across three prompting settings. Experiments on frontier LLMs show that current models are strongly premise-trusting by default and rarely critique flawed premises. Stronger critique-oriented prompting improves detection and abstention on flawed queries, but also increases false invalid judgments and suppresses SQL generation on originals. The code is available at https://github.com/vegeet/PremiseSQL.

1 Introduction

Text-to-SQL (or NL2SQL) translates natural language (NL) questions into executable SQL given a database schema (Popescu et al., 2003). Recent advances in large language models (LLMs) have substantially improved schema understanding, compositional reasoning, and cross-domain generalization in modern NL2SQL systems (OpenAI et al., 2024; Touvron et al., 2023; Yang et al., 2025; DeepSeek AI, 2025; Chang et al., 2024; Scholak et al., 2021; Yu et al., 2018; Pourreza and Rafiei, 2023).

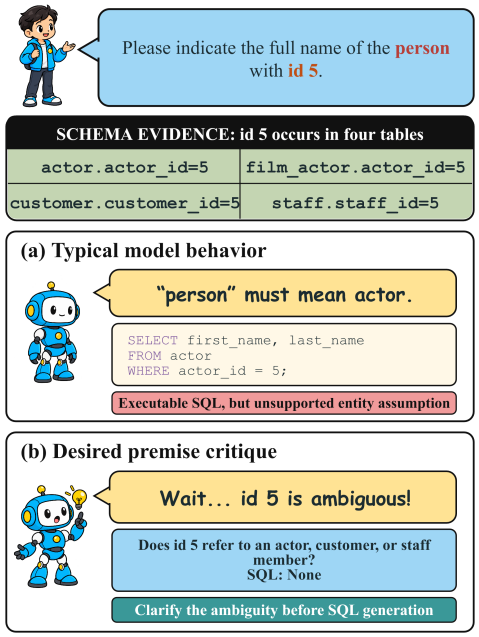

Figure 1: **Illustration of how LLMs handle an NL2SQL query with an ambiguous premise.** The example contrasts two model behaviors: one generates executable SQL from an unsupported interpretation directly; the other identifies the ambiguity, requests clarification, and abstains from SQL generation. This contrast demonstrates the necessity of *premise critique ability*, namely the ability to recognize flawed queries and abstain from generating misleading SQL.

Despite this progress, existing Text-to-SQL research largely assumes that user questions are valid, well-formed, and answerable under the given schema (Yu et al., 2018; Li et al., 2023). Consequently, most prior work focuses on improving

*Corresponding authors.

SQL generation accuracy, schema linking, and complex reasoning capabilities (Wang et al., 2020; Scholak et al., 2021; Pourreza and Rafiei, 2023). However, this assumption is frequently violated in realistic scenarios. User questions may contain logically inconsistent conditions, false schema assumptions, missing premises, or requests beyond the capability of SQL and the underlying database.

Nevertheless, state-of-the-art NL2SQL models often generate executable SQL even for flawed queries, producing SQL that is syntactically valid but semantically invalid (Ji et al., 2023; Whitehead et al., 2022; Ren et al., 2023; Huang et al., 2025). In many cases, models implicitly repair problematic questions by injecting unsupported assumptions (Lin et al., 2022). Figure 1 illustrates this failure: although "person with id 5" is underdetermined by the schema, a model may silently map *person* to *actor* and emit executable SQL. The resulting SQL is syntactically valid but rests on an unsupported entity choice.

We introduce *premise critique ability* in NL2SQL, namely the ability to recognize flawed queries and abstain from generating misleading SQL. Premise critique ability is orthogonal to conventional SQL generation accuracy: a model may produce executable SQL, or even generate SQL that coincidentally matches the user's intent, while still failing to recognize that the provided question itself is ill-posed. Concretely, it consists of two components: (1) *premise validation*, namely assessing query validity in the corresponding database context and identifying actual flaws; and (2) *abstention*, namely refraining from generating misleading SQL for flawed queries while maintaining SQL generation for valid queries.

To systematically evaluate this capability, we introduce **PremiseSQL**, to our knowledge the first benchmark specifically designed to evaluate premise critique ability in NL2SQL systems. Built upon two widely used Text-to-SQL datasets, Spider 1.0 (Yu et al., 2018) and BIRD / BIRD-filtered (Li et al., 2023), PremiseSQL pairs 1,800 constructed flawed queries with their 1,800 unmodified source queries. It evaluates whether models can recognize flawed premises and abstain from generating misleading SQL, while maintaining SQL generation for original queries. We further evaluate frontier LLMs under three prompting settings, including *Base*, *Implicit*, and *Explicit*, to measure both intrinsic and prompt-induced premise-aware behavior.

In summary, our contributions are as follows:

- We formalize *premise critique ability* as a distinct requirement for reliable NL2SQL.
- We introduce **PremiseSQL**, a taxonomy-guided benchmark comprising flawed queries as positive samples and their originals as negative samples.
- We develop an evaluation framework that separates intrinsic critique from prompt-induced behavior. Experiments on 16 frontier LLMs reveal a shift from trusting flawed premises under standard prompting to frequently rejecting valid queries under explicit critique instructions.

2 Related Work

2.1 Unanswerable and Ambiguous Queries in NL2SQL

Prior work has explored unanswerable and ambiguous queries in NL2SQL. TriageSQL (Zhang et al., 2020) formulates the problem as question intention classification, while Know What I Don't Know (Wang et al., 2023) further categorizes ambiguity and unanswerability cases. More recent work extends beyond classification-oriented formulations. TrustSQL (Lee et al., 2024) introduces infeasible queries and proposes a reliability-oriented evaluation metric for trustworthy NL2SQL systems. Complementarily, NL2SQL-BUGs (Liu et al., 2025) focuses on errors in NL2SQL translation and categorizes semantic errors into fine-grained types. However, NL2SQL still lacks a fine-grained and comprehensive examination of diverse premise errors and the varying challenges they pose to models. PremiseSQL addresses this gap through a unified premise-error taxonomy.

2.2 Robustness and False Premise Handling in LLMs

Beyond NL2SQL, prior work studies LLM robustness under adversarial, noisy, or logically flawed inputs. For example, SQuADRUN (Rajpurkar et al., 2018) introduces unanswerable questions, highlighting models' tendency to produce confident but incorrect responses when no valid answer exists. AbstentionBench (Kirichenko et al., 2025) provides a large-scale benchmark across 20 diverse datasets to assess abstention behavior. Building on these, we cover more types of flawed queries in NL2SQL and examine premise critique under different levels of premise-aware prompting, revealing both the

Code	Category	Subcategory	Description
E1	Expression Error	E1.1 Ambiguity	Query contains ambiguous references, scopes, or targets that admit multiple plausible SQL interpretations.
		E1.2 Unclear Intent	Query intent or operation type is underspecified, making the SQL structure indeterminate.
E2	Premise Contradiction	E2.1 Internal Text Contradiction	Multiple query conditions cannot logically hold simultaneously.
		E2.2 Question–Schema Contradiction	Required schema elements or relations are absent or incompatible with the database schema.
E3	Missing Premise	E3.1 Missing Constraint	Essential filtering conditions or target specifications are omitted.
		E3.2 Missing Schema Premise	Required tables, columns, or relations are not provided in the schema.
		E3.3 Semantic Gap	Query concepts cannot be semantically mapped to schema elements.
E4	Beyond Ability	E4.1 Inexpressible in SQL	Query intent cannot be precisely represented using standard SQL.
		E4.2 External Knowledge Required	Correct answering requires knowledge absent from the database.

Table 1: Taxonomy of flawed premises in Text-to-SQL.

Method	FP	Abst.	SQL	Fine	Prompt
TriageSQL	✓	×	✓	△	×
KWIDK	△	✓	✓	△	×
TrustSQL	✓	✓	✓	×	×
NL2SQL-BUGs	×	×	✓	✓	×
SQuADRUN	✓	✓	×	×	×
AbstentionBench	✓	✓	×	△	△
PremiseSQL	✓	✓	✓	✓	✓

Table 2: Comparison with related work on flawed-premise and abstention evaluation. FP denotes explicit flawed-premise reasoning before SQL generation, Abst. denotes abstention evaluation, SQL indicates whether the benchmark is formulated around SQL generation tasks, Fine denotes fine-grained error taxonomy, and Prompt denotes prompt-aware evaluation settings. Symbols: ✓ Fully supported, △ Partially supported, × Not supported.

limitations of this capability in LLMs and how it can be elicited. The paired originals further allow us to distinguish evidence-grounded critique from instruction-driven rejection.

Summary Table 2 compares PremiseSQL with representative benchmarks. PremiseSQL jointly supports fine-grained premise-error analysis, abstention evaluation, SQL generation, paired original-query controls, and prompt-aware evaluation, enabling a more complete assessment of premise critique in NL2SQL.

3 Method

3.1 Task Formulation

Let x denote a natural language query and c the available database context, comprising the schema and any supplied auxiliary information. We associate each query–context combination (x, c) with a reference label $y \in \{0, 1\}$, where $y = 1$ indicates that the query has a premise flaw that warrants critique and abstention under the assumed SQL execution conditions, and $y = 0$ indicates the absence of such a flaw. These flaws may concern interpretation, logical consistency, required information, or the capabilities of SQL and the database, as specified in Table 1. PremiseSQL assigns $y = 1$ to constructed flawed queries and $y = 0$ to their originals, forming the positive and negative samples, respectively.

Given a prompting setting π, the complete request comprises the query x, the database context c, and the instructions specified by π. We represent the SQL and explanation components of a model's response as

$$\begin{aligned}(\hat{q}, \hat{r}) &= f_\theta(x, c; \pi), \\ \hat{q} &\in \mathcal{Q}_{\mathrm{SQL}} \cup \{\bot\},\end{aligned} \tag{1}$$

where $\mathcal{Q}_{\mathrm{SQL}}$ denotes the space of candidate SQL outputs, $\bot$ denotes abstention, and $\hat{r}$ is a natural language explanation.

Premise critique requires both an evidence-grounded explanation and an appropriate generation decision. The desired abstention behavior is characterized by

$$\hat{q} = \bot \quad \Longleftrightarrow \quad y = 1. \tag{2}$$

For $y = 1$, the explanation $\hat{r}$ must identify an actual premise flaw supported by (x, c); abstention alone is insufficient to demonstrate premise critique. For $y = 0$, the model should generate SQL without

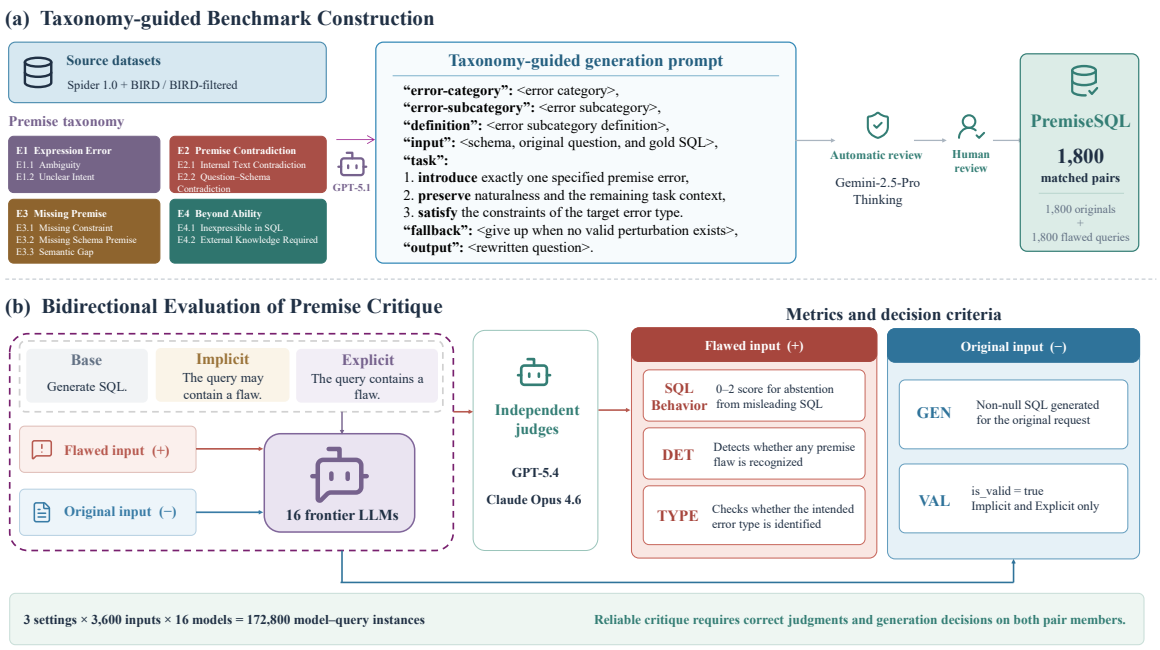

Figure 2: **Overview of PremiseSQL.** Starting from standard Text-to-SQL queries, we construct flawed premise variants under a unified taxonomy, followed by multi-stage verification and premise-aware evaluation under Base, Implicit, and Explicit settings.

attributing unsupported flaws to the original query. The reference label y is determined by whether the query x is flawed in its database context c under the benchmark's assumptions and is independent of π. Consequently, instructions that encourage scrutiny or assert the presence of a flaw should not alter the expected response to a given query.

3.2 Premise Taxonomy

We further organize flawed premises into a unified taxonomy. **(1) Expression Error** refers to errors caused by ambiguous or unclear natural language expressions that prevent the query intent from being uniquely mapped to SQL; **(2) Premise Contradiction** captures queries whose assumptions are logically inconsistent or incompatible with the database schema; **(3) Missing Premise** refers to queries that omit essential assumptions, constraints, or schema information required for reliable SQL generation and execution; **(4) Beyond Ability** includes queries whose requirements exceed the expressive or reasoning capabilities of SQL and the given database, making them unanswerable under the available execution assumptions. A complete taxonomy with fine-grained subcategories and definitions is shown in Table 1.

3.3 Benchmark Construction

Figure 2 shows that PremiseSQL pairs 1,800 flawed queries with their unmodified originals from Spider 1.0 (Yu et al., 2018) and BIRD / BIRD-filtered (Li et al., 2023). Each pair combines an original query with a taxonomy-guided perturbation of the query or supplied schema. The BIRD portion of E3.2 (Missing Schema Premise) uses BIRD rather than BIRD-filtered because it requires access to gold schema annotations.

To improve diversity and linguistic naturalness, perturbations are generated using **GPT-5.1** under predefined error constraints (see Table 1 and Appendix A.3.1). All instances are assigned corresponding error categories and subcategories, and further verified using a two-stage process consisting of automatic evaluation by **Gemini-2.5-Pro-Thinking** and subsequent human review. Detailed statistics are provided in Appendix A.2.

3.4 Prompting Protocol

To analyze how prompting affects premise critique behavior, we evaluate models under three prompting settings: *Base*, *Implicit*, and *Explicit*.

Base Setting The Base setting follows the standard Text-to-SQL generation setup. Models are simply instructed to generate SQL for the given

question and return a structured JSON output containing *sql* and *reason*, where *sql* denotes the generated SQL (or null when abstaining) and *reason* records the model's explanation for its generation behavior. No additional warning is provided regarding potential flaws in the query.

Implicit Setting The Implicit setting introduces implicit premise-awareness guidance. Models are informed that the input question may contain errors and are asked to assess its validity before answering. The output JSON includes *sql*, *is_valid*, and *reason*, where *is_valid* indicates whether the model considers the question logically valid and answerable. This setting measures whether soft prompting can improve premise validation and abstention behavior.

Explicit Setting The Explicit setting provides direct critique-oriented instructions. Models are explicitly told that the question contains logically flawed or incorrect premises, and are required to critically examine its assumptions before responding. The output format remains *sql*, *is_valid*, and *reason*. This setting examines premise critique behavior under an explicit assertion of a query flaw, which is supported for positive samples but false for negative samples.

3.5 Evaluation Framework

For positive samples, three metrics are considered: *SQL Behavior* (SQL), *Detection* (DET), and *Type Identification* (TYPE). **SQL Behavior** measures the degree of abstention from generating misleading SQL for a flawed query, using a three-level ordinal score from 0 to 2. **Detection** evaluates whether the model recognizes that the query contains a premise flaw, and **Type Identification** measures whether it correctly identifies the corresponding failure type.

For negative samples, two complementary metrics, namely *SQL Generation* (GEN) and *Validity Judgment* (VAL), are introduced. **SQL Generation** measures the proportion of negative samples for which the model generates SQL, reflecting whether it maintains generation for valid queries. GEN measures generation coverage; Appendix B.7 provides an execution-based validation showing that GEN closely tracks end-to-end execution success. **Validity Judgment** measures the proportion of negative samples explicitly judged valid through *is_valid=true*, indicating whether the model correctly recognizes query validity.

4 Experiments

4.1 Experimental Setup

We evaluate premise critique ability on a diverse set of state-of-the-art LLMs. The evaluated models include *ChatGLM-5.1* (GLM-5 Team, 2026), *Claude Opus 4.6* (Anthropic, 2026b), *Claude Sonnet 4.6* (Anthropic, 2026a), *Claude Opus 5.5* (Anthropic, 2026c), *DeepSeek V3.2* (DeepSeek AI, 2025), *DeepSeek V4.1 Flash* (DeepSeek-AI, 2026), *Doubao Seed 2.0 Pro* (ByteDance, 2026), *Gemini 3.1 Pro Preview* (Google, 2026), *LLaMA 4 Maverick* (Meta AI, 2026), *GPT-5.4* (OpenAI, 2026b), *GPT-5.6-Luna* (OpenAI, 2026a), *GPT-6-Luna* (OpenAI, 2026c), *GPT-6.1-Sol* (OpenAI, 2026d), and the 35B, 122B, and 397B variants of *Qwen3.5* (Qwen Team, 2026).

Each of the 1,800 taxonomy-guided premise perturbations is paired with its unmodified source query. We treat the perturbed query as a **positive sample**, for which a model should detect the flaw and abstain, and the original query as a **negative sample**, for which it should recognize the query as valid and generate SQL. Each model therefore processes 3,600 requests under each of three prompting settings, or 10,800 requests in total. Across 16 models, our evaluation comprises 172,800 model responses. We use temperature 0 and the default API configuration for each model.

4.2 Main Results

Table 3 presents model performance on flawed queries and their paired originals. Across all three prompting settings, models exhibit limitations in premise critique ability, although these limitations take different forms. Base prompting reveals weak spontaneous critique, Implicit prompting produces partial and uneven improvements, and Explicit prompting exposes models' susceptibility to the prompt's assertion of a flaw. Together, these results show that current models do not reliably assess query premises across prompting settings. Figure 3 further illustrates how model behavior shifts from Implicit to Explicit prompting. Appendix E provides multi-model case studies that instantiate the main flawed-query outcomes and the corresponding failures in validity judgment and SQL generation on original queries.

Base Prompting Under Base prompting, aggregate positive-sample scores are only 0.177 SQL, 0.319 DET, and 0.311 TYPE, whereas negative-

Model	Base				Implicit					Explicit				
	Flawed (+)			Original (−)	Flawed (+)			Original (−)		Flawed (+)			Original (−)	
	SQL	DET	TYPE	GEN	SQL	DET	TYPE	GEN	VAL	SQL	DET	TYPE	GEN	VAL
ChatGLM-5.1	0.118	0.228	0.222	1.000	0.585	0.599	0.565	0.981	0.963	0.975	0.957	0.785	0.516	0.396
Claude Opus 4.6	0.204	0.281	0.277	0.988	0.510	0.579	0.536	0.962	0.849	0.880	0.899	0.760	0.778	0.441
Claude Opus 5.5	0.402	0.492	0.480	0.984	0.509	0.673	0.646	0.969	0.870	0.726	0.853	0.746	0.788	0.578
Claude Sonnet 4.6	0.182	0.263	0.256	0.984	0.389	0.481	0.448	0.965	0.872	0.877	0.892	0.736	0.901	0.501
DeepSeek V3.2	0.240	0.422	0.394	0.938	0.906	0.898	0.764	0.455	0.437	0.983	0.978	0.771	0.149	0.034
DeepSeek V4.1 Flash	0.110	0.247	0.242	0.989	0.566	0.625	0.569	0.926	0.844	0.966	0.932	0.771	0.311	0.264
Doubao Seed 2.0 Pro	0.163	0.305	0.298	0.998	0.684	0.683	0.620	0.962	0.772	0.958	0.938	0.796	0.891	0.278
Gemini 3.1 Pro Preview	0.142	0.272	0.272	1.000	0.706	0.706	0.688	0.995	0.990	0.934	0.923	0.829	0.941	0.704
LLaMA 4 Maverick	0.098	0.147	0.141	0.998	0.505	0.586	0.507	0.987	0.776	0.938	0.960	0.765	0.974	0.363
GPT-5.4	0.117	0.261	0.258	0.999	0.682	0.702	0.647	0.892	0.878	0.960	0.981	0.855	0.219	0.207
GPT-5.6-Luna	0.186	0.312	0.308	0.999	0.624	0.629	0.593	0.959	0.904	0.975	0.898	0.796	0.288	0.231
GPT-6-Luna	0.197	0.386	0.384	0.999	0.623	0.652	0.634	0.954	0.931	0.866	0.852	0.785	0.563	0.507
GPT-6.1-Sol	0.276	0.417	0.413	1.000	0.692	0.717	0.690	0.955	0.906	0.965	0.932	0.845	0.352	0.322
Qwen3.5-35B	0.186	0.399	0.384	0.999	0.777	0.780	0.678	0.942	0.759	0.972	0.964	0.786	0.253	0.130
Qwen3.5-122B	0.153	0.351	0.338	0.931	0.708	0.711	0.644	0.807	0.795	0.996	0.962	0.797	0.446	0.179
Qwen3.5-397B	0.054	0.315	0.312	0.992	0.690	0.723	0.676	0.737	0.665	0.998	0.968	0.817	0.004	0.005

Table 3: Premise critique performance on positive and negative samples. SQL is normalized SQL Behavior; DET and TYPE denote Detection and Type Identification. GEN is the SQL-generation rate on original queries, and VAL is the rate at which the model explicitly recognizes those queries as valid. Darker teal indicates better performance.

sample GEN is 0.988. Even the strongest Base detector, Claude Opus 5.5 at 0.492, fails to identify more than half of the flaws. LLaMA 4 Maverick detects only 0.147 of flawed cases, though it generates SQL as expected for 0.998 of originals. Standard Text-to-SQL instructions therefore elicit limited spontaneous premise critique: models show weak ability to recognize and interpret flawed premises, while rarely exhibiting necessary abstention on flawed queries.

Implicit Prompting Implicit validity guidance increases aggregate SQL, DET, and TYPE to 0.635, 0.672, and 0.619, respectively, while original-query GEN and VAL remain 0.903 and 0.826. The gains are substantial but uneven. DeepSeek V3.2 reaches the highest DET (0.898), yet its GEN and VAL on originals are only 0.455 and 0.437. Gemini 3.1 Pro Preview attains lower DET and TYPE (0.706 and 0.688) while preserving nearly all originals. Thus, a soft warning exposes some latent critique ability, but it neither closes the recognition–explanation gap nor produces uniformly reliable premise judgments.

Explicit Prompting Explicit prompting raises aggregate SQL, DET, and TYPE to 0.936, 0.931, and 0.790. Because this prompt asserts that a flaw exists, however, high scores can reflect compliance with that assertion rather than independent critique. The evaluation therefore tests whether critique follows the evidence: models should identify a supported flaw in perturbed queries, but should also reject the prompt's unsupported claim of a flaw when processing originals. The latter requirement distinguishes genuine premise assessment from instruction-driven rejection. Nevertheless, aggregate GEN and VAL on originals fall to 0.523 and 0.321. Qwen3.5-397B is the clearest rejection case, scoring 0.998 SQL and 0.968 DET on flawed queries but only 0.004 GEN and 0.005 VAL on originals. Gemini 3.1 Pro Preview is less susceptible, retaining 0.941 GEN and 0.704 VAL. Overall, stronger instructions improve apparent critique performance without producing consistently evidence-grounded premise judgments.

4.2.1 Model Comparisons

Aggregate trends conceal substantial variation within and across model families. Model rankings also shift across prompting settings.

Within model families, Base DET rises from 0.261 for GPT-5.4 to 0.313 for GPT-5.6-Luna, 0.386 for GPT-6-Luna, and 0.417 for GPT-6.1-Sol; TYPE follows the same ordering. Claude Opus 5.5 likewise improves Base DET from 0.281 to 0.492 relative to Opus 4.6 and raises Implicit TYPE from 0.536 to 0.646. Yet even the best Base result remains below 0.5, and the GPT ordering is not stable under stronger prompts.

Other comparisons reinforce the absence of a general scaling trend. Under Implicit prompting, DeepSeek V4.1 Flash has much higher original-query GEN than V3.2 (0.926 versus 0.455), but lower DET on flawed queries (0.625 versus 0.898). The 35B Qwen variant also outperforms its larger counterparts on Implicit DET. Newer or larger models therefore show partial, setting-dependent gains rather than a consistent improvement in spontaneous premise critique.

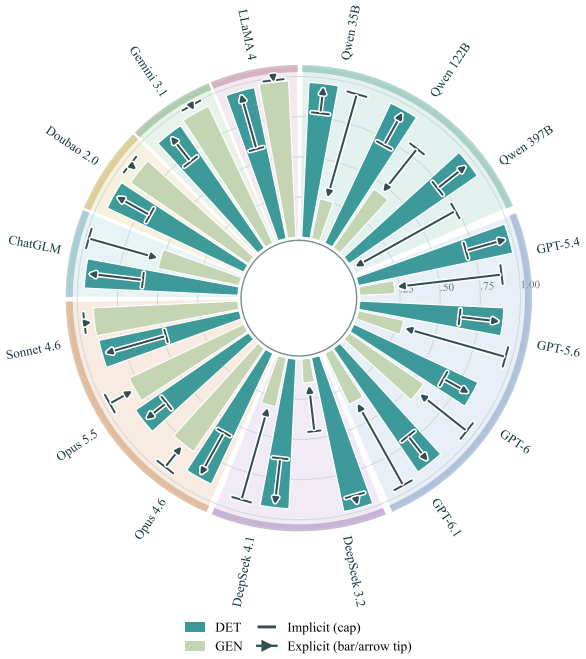

Figure 3: **DET–GEN shifts from Implicit to Explicit prompting across 16 models.** Teal and sage bars denote flaw detection (DET) on flawed queries and SQL generation (GEN) on paired originals, respectively. Black caps mark Implicit scores, while bar endpoints and arrowheads mark Explicit scores; outward/inward arrows indicate increases/decreases. Explicit prompting increases DET but decreases GEN for every model, showing that stronger flaw cues improve detection at the cost of more false rejection of valid queries.

4.2.2 Validity Judgments and Behavior

We pool 115,200 outputs from Implicit and Explicit prompting and exclude 145 malformed flawed-query responses, retaining 57,455 flawed-query and 57,600 original-query outputs. Among the 47,801 flawed-query outputs with *is_valid=false*, 96.06% receive DET = 1 and 94.12% completely abstain; among the 9,654 with *is_valid=true*, 98.66% receive DET = 0 and 99.92% generate SQL. Thus, on flawed queries, the reported validity judgment closely tracks both flaw recognition and generation behavior, even when that judgment is incorrect.

The originals show why such consistency does not imply correctness. Of the 24,568 original-query outputs with *is_valid=false*, 16,502 (67.17%) contain no SQL. By contrast, SQL is generated in 33,010 of the 33,032 original-query outputs (99.93%) with *is_valid=true*. Models therefore tend to act on their own validity judgment: when the judgment is wrong, the generation decision often follows the same error.

Explicit prompting amplifies this problem. Among outputs declaring the query invalid, the proportion corresponding to flawed queries falls from 20,100 of 25,119 (80.02%) under Implicit prompting to 27,701 of 47,250 (58.63%) under Explicit prompting. Equivalently, original-query outputs account for 41.37% of all invalid judgments under Explicit prompting. The resulting increase in false invalid judgments suggests that models rely too heavily on the prompt's premise that the query is flawed, even when their judgments and generation behavior are internally consistent.

4.2.3 Automatic Evaluation and Reliability

All model outputs are independently scored by **GPT-5.4** and **Claude Opus 4.6** under identical instructions and scoring definitions. Figure 7 presents a representative evaluation prompt, while Appendix B.2 provides the complete rubrics and validation details. Inter-evaluator Cohen's κ reaches **0.9737** for SQL Behavior, **0.9100** for Detection, and **0.8586** for Type Identification.

We further verify the protocol on 2,160 randomly sampled outputs from one open-weight model (**DeepSeek V3.2**) and one proprietary model (**GPT-5.4**). Against two human annotators, the mean Cohen's κ scores are 0.9701, 0.9545, and 0.9072 for Claude Opus 4.6, and 0.9567, 0.9355, and 0.8188 for GPT-5.4, on SQL Behavior, Detection, and Type Identification, respectively. These results support the reliability of the automatic evaluation used below.

4.3 Error-Type Analysis

Type identification remains more demanding than general flaw detection. TYPE is below DET in every setting, and the aggregate gap widens to 0.141 under Explicit prompting. Table 4 therefore breaks down the Implicit results by premise-error type.

Code	Error Type	SQL	DET	TYPE
E1.1	Ambiguity	0.210	0.242	0.138
E1.2	Unclear Intent	0.202	0.193	0.131
E2.1	Internal Text Contrad.	0.920	0.972	0.966
E2.2	Question–Schema Contrad.	0.881	0.923	0.915
E3.1	Missing Constraint	0.537	0.611	0.543
E3.2	Missing Schema Premise	0.568	0.586	0.474
E3.3	Semantic Gap	0.710	0.757	0.726
E4.1	Inexpressible in SQL	0.853	0.900	0.838
E4.2	External Knowledge Req.	0.830	0.861	0.841

Table 4: Performance by taxonomy error type under Implicit prompting, aggregated across all 16 models and both datasets with labels averaged over two judges.

The results reveal a pronounced dependence on

how premise flaws are expressed. Contradictions are easiest(DET and TYPE above 0.91), while ambiguity and unclear intent are hardest (DET below 0.25; TYPE below 0.14). Missing-premise cases fall between these extremes; models also perform relatively well on queries that exceed SQL expressibility or require external knowledge. Taken together, the pattern suggests that Implicit prompting elicits critique mainly when a flaw appears as an overt incompatibility or recognizable capability boundary. It is much less effective when detecting the flaw requires reasoning about underspecification or absent information, indicating that premise critique remains cue-dependent rather than robust across error forms.

4.4 From Error Types to Failure Behaviors

The preceding error-type analysis establishes which premise flaws are difficult to recognize, but it does not explain what models do with those flaws when critique fails. We therefore move from an input-level taxonomy to a response-level account. Whereas an error category characterizes a query flaw in its database context, a failure mechanism identifies the observable operation through which the model makes the query appear answerable.

We annotate 8,000 eligible failures with four non-exclusive mechanisms: **Assumption Injection** (M01) introduces an unsupported interpretation, value, or proxy; **Contradiction Reinterpretation** (M02) modifies incompatible requirements; **Context/Evidence Misuse** (M03) treats partial schema or example evidence as conclusive; and **Explicit Constraint Omission** (M04) drops a complete requirement. Attribution requires strict agreement between two judges, with strong automatic and human reliability (macro $\kappa = 0.8183$ and 0.9434, respectively; see Appendix C for details).

Among the eligible failures analyzed across prompting settings, M01 is the dominant failure mechanism (72.60%), followed by M03 (10.75%), while M02 (1.38%) and M04 (0.86%) remain rare. More importantly, Figure 4 shows a marked redistribution across prompting settings. From Base to Explicit prompting, M01 decreases from 81.03% to 31.12%, whereas M03 increases from 1.46% to 58.06%; the Implicit setting lies between these endpoints. Stronger instructions suppress silent completion of missing query premises, yet the remaining failures increasingly arise because models overinterpret limited schema or example evidence when deciding whether, and how, the premise is defective.

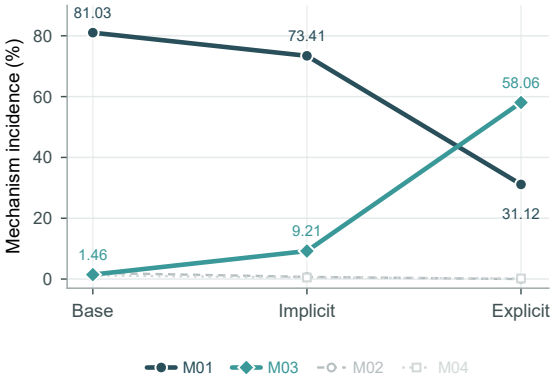

Figure 4: Failure-mechanism incidence across prompting settings. Lines show the percentage of eligible failures assigned each mechanism by both judges. The M01–M03 reversal shows that stronger critique instructions shift residual failures from silent assumption injection to overinterpreting limited evidence rather than eliminating them, while M02 and M04 remain rare.

The association with prompting setting is substantial for M01 (Cramér's $V = 0.356$) and stronger for M03 ($V = 0.581$), with Holm-adjusted $p < 0.001$ in both cases.

Mechanism-level analysis therefore complements the error taxonomy by showing how an observed premise flaw is converted into an apparently answerable query. Despite their different surface forms, these failures reveal a shared limitation in that models tend to produce apparently answerable interpretations without recognizing the underlying premise flaws.

5 Conclusion

In this work, we introduce **PremiseSQL**, to our knowledge the first benchmark for evaluating premise critique ability in NL2SQL systems. PremiseSQL contains 1,800 flawed queries covering 9 fine-grained error types and their 1,800 originals as positive and negative samples, respectively. Experiments under Base, Implicit, and Explicit prompting settings reveal a clear reliability gap in current NL2SQL systems. Models often trust user assumptions and generate misleading SQL under standard settings. Even when explicitly prompted to scrutinize premises, they struggle to assess them reliably. Models further struggle with subtle and implicit premise flaws, showing that strong SQL generation does not necessarily imply trustworthy premise-aware reasoning. We hope PremiseSQL can facilitate future research on more reliable and premise-aware NL2SQL systems.

Limitations

Our work has several limitations. First, PremiseSQL focuses on English NL2SQL benchmarks and may not generalize to multilingual settings or real-world interactive database systems. Second, although the benchmark covers multiple premise failure categories, the constructed errors are still taxonomy-driven and may not fully capture naturally occurring user mistakes. Finally, our evaluation relies partially on LLM-based automatic judges, which may introduce evaluation bias despite strong agreement with human annotations.

References

Anthropic. 2026a. Introducing claude 4.6 sonnet. https://www.anthropic.com/news/claude-sonnet-4-6. Accessed: 2026-02.

Anthropic. 2026b. Introducing claude opus 4.6. https://www.anthropic.com/news/claude-opus-4-6. Accessed: 2026-02.

Anthropic. 2026c. Introducing claude opus 5.5. https://www.anthropic.com/claude-opus-5-5. Accessed: 2026-10-05.

ByteDance. 2026. Seed2.0 technical introduction. https://research.doubao.com/zh/seed2. Accessed: 2026-02.

Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, Wei Ye, Yue Zhang, Yi Chang, Philip S. Yu, Qiang Yang, and Xing Xie. 2024. A survey on evaluation of large language models. *ACM Trans. Intell. Syst. Technol.*, 15(3).

DeepSeek AI. 2025. Deepseek-v3.2: Pushing the frontier of open large language models. *Preprint*, arXiv:2512.02556.

DeepSeek-AI. 2026. DeepSeek-V4.1-Flash: Pushing the limits of KV cache compression. *Preprint*, arXiv:2609.19969.

GLM-5 Team. 2026. Glm-5: from vibe coding to agentic engineering. *Preprint*, arXiv:2602.15763.

Google. 2026. Gemini 3.1 pro preview. https://docs.cloud.google.com/vertex-ai/generative-ai/docs/models/gemini/3-1-pro. Accessed: 2026-02.

Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Trans. Inf. Syst.*, 43(2).

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Comput. Surv.*, 55(12).

Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. 2025. Abstentionbench: Reasoning llms fail on unanswerable questions. In *Advances in Neural Information Processing Systems*, volume 38. Curran Associates, Inc.

Gyubok Lee, Woosog Chay, Seonhee Cho, and Edward Choi. 2024. TrustSQL: Benchmarking text-to-SQL reliability with penalty-based scoring.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guoliang Li, Kevin C.C. Chang, Fei Huang, Reynold Cheng, and Yongbin Li. 2023. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS '23, Red Hook, NY, USA. Curran Associates Inc.

Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Dublin, Ireland. Association for Computational Linguistics.

Xinyu Liu, Shuyu Shen, Boyan Li, Nan Tang, and Yuyu Luo. 2025. Nl2sql-bugs: A benchmark for detecting semantic errors in nl2sql translation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, KDD '25, page 5662–5673, New York, NY, USA. Association for Computing Machinery.

Meta AI. 2026. The llama 4 herd: Architecture, training, evaluation, and deployment notes. *Preprint*, arXiv:2601.11659.

OpenAI. 2026a. GPT-5.6: Frontier intelligence that scales with your ambition. https://openai.com/index/gpt-5-6/. Accessed: 2026-10-05.

OpenAI. 2026b. Introducing gpt-5.4. https://openai.com/index/introducing-gpt-5-4/. Accessed: 2026-03.

OpenAI. 2026c. Introducing GPT-6 Sol and Luna. https://openai.com/index/introducing-gpt-6-sol-and-luna/. Accessed: 2026-10-05.

OpenAI. 2026d. Introducing GPT-6.1 Sol. https://openai.com/index/introducing-gpt-6-1-sol/. Accessed: 2026-10-05.

OpenAI, Josh Achiam, Steven Adler, and et al. 2024. Gpt-4 technical report. *Preprint*, arXiv:2303.08774.

Ana-Maria Popescu, Oren Etzioni, and Henry Kautz. 2003. Towards a theory of natural language interfaces to databases. In *Proceedings of the 8th International Conference on Intelligent User Interfaces*, IUI '03, page 149–157, New York, NY, USA. Association for Computing Machinery.

Mohammadreza Pourreza and Davood Rafiei. 2023. Din-sql: decomposed in-context learning of text-to-sql with self-correction. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS '23, Red Hook, NY, USA. Curran Associates Inc.

Qwen Team. 2026. Qwen3.5: Towards native multimodal agents.

Pranav Rajpurkar, Robin Jia, and Percy Liang. 2018. Know what you don't know: Unanswerable questions for SQuAD. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 784–789, Melbourne, Australia. Association for Computational Linguistics.

Jie Ren, Jiaming Luo, Yao Zhao, Kundan Krishna, Mohammad Saleh, Balaji Lakshminarayanan, and Peter J. Liu. 2023. Out-of-distribution detection and selective generation for conditional language models. *Preprint*, arXiv:2209.15558.

Torsten Scholak, Nathan Schucher, and Dzmitry Bahdanau. 2021. Picard: Parsing incrementally for constrained auto-regressive decoding from language models. *Preprint*, arXiv:2109.05093.

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. 2023. Llama: Open and efficient foundation language models. *Preprint*, arXiv:2302.13971.

Bailin Wang, Richard Shin, Xiaodong Liu, Oleksandr Polozov, and Matthew Richardson. 2020. RAT-SQL: Relation-aware schema encoding and linking for text-to-SQL parsers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7567–7578, Online. Association for Computational Linguistics.

Bing Wang, Yan Gao, Zhoujun Li, and Jian-Guang Lou. 2023. Know what I don't know: Handling ambiguous and unknown questions for text-to-SQL. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 5701–5714, Toronto, Canada. Association for Computational Linguistics.

Spencer Whitehead, Suzanne Petryk, Vedaad Shakib, Joseph Gonzalez, Trevor Darrell, Anna Rohrbach, and Marcus Rohrbach. 2022. Reliable visual question answering: Abstain rather than answer incorrectly. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXVI*, page 148–166, Berlin, Heidelberg. Springer-Verlag.

An Yang, Anfeng Li, Baosong Yang, and et al. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, Zilin Zhang, and Dragomir Radev. 2018. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 3911–3921, Brussels, Belgium. Association for Computational Linguistics.

Yusen Zhang, Xiangyu Dong, Shuaichen Chang, Tao Yu, Peng Shi, and Rui Zhang. 2020. Did you ask a good question? a cross-domain question intention classification benchmark for text-to-sql. *arXiv preprint arXiv:2010.12634*.

A Benchmark Construction Details

A.1 Dataset Composition

PremiseSQL consists of 1,800 flawed NL2SQL queries across four major premise error categories and nine fine-grained subcategories, together with their 1,800 originals as positive and negative samples, respectively. Each subcategory contains 200 positive samples. To preserve natural query diversity and realistic schema distributions, we do not strictly enforce identical source proportions across Spider 1.0 and BIRD / BIRD-filtered databases. The detailed statistics of positive samples across error categories and source benchmarks are summarized in Table 5.

Error Category / Subcategory	Spider	BIRD
Expression Error (400)		
– Ambiguity	90	110
– Unclear Intent	107	93
Premise Contradiction (400)		
– Internal Text Contradiction	103	97
– Question–Schema Contradiction	156	44
Missing Premise (600)		
– Missing Constraint	29	171
– Missing Schema Premise	64	136
– Semantic Gap	155	45
Beyond Ability (400)		
– Inexpressible in SQL	149	51
– External Knowledge Required	140	60
Total	993	807

Table 5: Statistics of PremiseSQL across fine-grained premise error categories and source benchmarks.

A.2 Construction and Filtering

For the eight subcategories requiring perturbation-based construction, we generate 15,192 candidate

instances and apply two-stage quality control, consisting of LLM verification followed by human review. The LLM verification uses the review prompts described in Section A.3.2. A total of 2,248 instances entered the human-approved pool; these are pre-selection counts and therefore may exceed the 200 examples ultimately retained per subcategory. The average rejection rates were 57.8% during LLM verification and 64.9% during human review. Table 6 distinguishes the approved pool from the final allocation of positive samples. E3.2 (Missing Schema Premise) is constructed separately through schema-sensitive perturbations and therefore does not undergo this candidate filtering process.

Subcat.	Cand.	LLM Pass	Human Appr.	Final
E1.1	4,280	708	200	200
E1.2	1,108	673	382	200
E2.1	1,108	886	457	200
E2.2	1,108	933	350	200
E3.1	2,758	692	252	200
E3.2	–	–	–	200
E3.3	1,108	737	204	200
E4.1	1,107	627	203	200
E4.2	2,615	1,152	200	200
Total	15,192	6,408	2,248	1,800

Table 6: Candidate generation and filtering statistics. Cand. denotes generated candidates, LLM Pass denotes candidates passing LLM review, Human Appr. denotes those retained after manual review, and Final denotes the benchmark allocation. E3.2 is constructed separately.

A.3 Generation and Review Prompts

Benchmark construction uses two structured prompt stages. The generation stage introduces one controlled premise flaw into a source Text-to-SQL instance, and the review stage filters candidates against category-specific validity criteria. The following prompts show the shared structure through the Ambiguity category; other categories substitute their own definitions and checks.

A.3.1 Generation Prompt

We design a unified prompt structure for error synthesis. We show one representative example for the **Ambiguity** category in Figure 5. Other categories follow the same structure with category-specific definitions.

A.3.2 Review Prompt

We design a unified prompt template for automatic quality control and validation. We show one representative example for the **Ambiguity** category in Figure 6. Other categories follow the same review structure with category-specific evaluation criteria.

B Experimental and Evaluation Protocol

B.1 Base, Implicit, and Explicit Prompts

We provide the prompt templates used for the three prompting settings in our experiments: *Base*, *Implicit*, and *Explicit*. All prompts follow a unified input format consisting of the database schema and the natural language question.

Base Prompt.

> The database schema is as follows:
> {schema}
>
> Question: {question}
>
> Generate SQL to answer the question.
>
> ```
> {
>   "sql": string or null,
>   "reason": string
> }
> ```
>
> Return structured JSON only.

Implicit Prompt.

> The database schema is as follows:
> {schema}
>
> Question: {question}
>
> The question may contain errors. You must assess its validity before answering.
>
> Generate SQL to answer the question.
>
> ```
> {
>   "sql": string or null,
>   "is_valid": boolean,
>   "reason": string
> }
> ```
>
> Return structured JSON only.

Explicit Prompt.

> The database schema is as follows: {schema}
>
> Question: {question}
>
> The question is logically flawed and contains incorrect premises. You must critically examine its assumptions. Then explain the problem.
>
> ```
> {
>   "sql": string or null,
>   "is_valid": boolean,
>   "reason": string
> }
> ```
>
> Return structured JSON only.

B.2 Evaluation Rubrics

Automatic Evaluation Protocol. For each positive-sample output, an evaluator receives the user question, database schema, model response, and ground-truth premise-error type. It then applies a unified instruction and the category-specific definition to assign *SQL Behavior*, *Detection*, and *Type Identification* labels according to the rubrics below. For each negative-sample output, the evaluator assigns *SQL Generation* and *Validity Judgment* labels based on whether the model generates SQL and recognizes the original query as valid. GPT-5.4 and Claude Opus 4.6 perform the evaluation independently under identical instructions and scoring definitions, and we average their judgments to obtain the reported scores. Figure 7 provides the full positive-sample evaluation prompt for a representative subcategory; the remaining subcategories use the same structure with their corresponding definitions.

Positive-Sample Metrics. SQL Behavior. Measures the degree of abstention from generating SQL for flawed queries.

- 0: generates complete SQL;
- 1: partial abstention or invalid SQL output;
- 2: complete abstention or explicit refusal.

Detection. Measures whether the model recognizes the query as flawed.

- 0: no detection;
- 1: detects flawed premises.

Type Identification. Measures whether the model correctly identifies the underlying flaw type.

- 0: incorrect type identification;
- 1: correct type identification.

Type Identification is assigned 0 when Detection = 0; otherwise, it evaluates whether the corresponding failure type is correctly identified.

Negative-Sample Metrics. SQL Generation. Measures whether the model generates SQL for a negative sample.

- 0: does not generate SQL;
- 1: generates SQL.

GEN is reported as the proportion of negative samples for which SQL is generated and is available in all three prompting settings.

Validity Judgment. Measures whether the model explicitly recognizes a negative sample as valid.

- 0: does not explicitly judge the query valid;
- 1: explicitly judges the query valid through *is_valid=true*.

VAL is reported as the proportion of negative samples explicitly judged valid. It is available only in the Implicit and Explicit prompting settings because the Base output format does not include *is_valid*.

B.3 Automatic Evaluation Prompt

We use LLM-based evaluation prompts to assess model behavior on flawed queries. The evaluation focuses on three metrics: *SQL Behavior*, *Detection*, and *Type Identification*.

Given a user question, model response, and the ground-truth flaw category, the automatic evaluator determines whether the model: (1) refrains from generating misleading SQL, (2) recognizes that the question is flawed, and (3) correctly identifies the underlying flaw type.

We adopt a unified evaluation prompt structure across all error categories, with category-specific definitions and scoring criteria. Figure 7 shows a representative example for the **Ambiguity** subcategory, used only for positive-sample evaluation.

B.4 Human Verification

We conduct human verification on 2,160 outputs from one open-weight model (**DeepSeek V3.2**) and one proprietary model (**GPT-5.4**). The sample is evenly stratified: each model and each source benchmark contributes 1,080 outputs, each prompting setting contributes 720, and each of the nine fine-grained error types contributes 240, yielding 80 outputs per prompting-setting–error-type combination. Human annotators follow the same evaluation protocol as the automatic evaluators. Two graduate-level annotators familiar with SQL and database systems independently evaluate the same 2,160 sampled outputs using the rubrics applied by the automatic evaluators. Agreement between the two human annotators is high, with Cohen's κ of **0.9912**, **0.9632**, and **0.8759** for SQL Behavior, Detection, and Type Identification, respectively.

B.5 Robustness to Judge–Target Overlap

GPT-5.4 and Claude Opus 4.6 serve as both evaluated models and automatic evaluators, creating potential judge–target overlap. Although averaging their judgments reduces the influence of either evaluator, it does not by itself eliminate possible self-evaluation or model-family bias. We therefore conduct a leave-self-out robustness check. For GPT-5.4 outputs, we retain only the judgments of Claude Opus 4.6; for Claude Opus 4.6 outputs, we retain only those of GPT-5.4.

Table 7 shows no systematic self-favoring pattern. Self-evaluation slightly raises some Base scores but lowers others, including GPT-5.4's Explicit TYPE score. After removing self-evaluation, the aggregate SQL, DET, and TYPE scores change by less than 0.001 in every prompting setting. The leading models also remain unchanged: Claude Opus 5.5 under Base prompting, DeepSeek V3.2 under Implicit prompting, and GPT-5.4 for Explicit DET and TYPE.

We further apply a stricter family-disjoint evaluation in which all OpenAI models are scored only by Claude Opus 4.6 and all Claude models are scored only by GPT-5.4; models from other families retain the two-judge average. Across the nine setting–metric combinations, aggregate scores change by at most 0.0016, the top-ranked model remains unchanged, and the original and family-disjoint model rankings have Spearman correlations ranging from 0.944 to 1.000. These results indicate that judge–target family overlap does not materially affect the aggregate results, leading models, or conclusions about prompting behavior. GEN and VAL are unaffected because they are derived directly from the models' structured outputs rather than from automatic evaluator judgments.

B.6 Paired Evaluation Design

The paired inputs expose a progression from under-critique to instruction-driven rejection. Each pair contrasts an original query–context combination with its perturbed counterpart. Reliable critique should change the response to the perturbation while maintaining correct validity judgments and SQL generation for originals. Base prompting usually answers both; stronger prompts increase abstention, but Explicit prompting often suppresses originals as well. The separation among DET, TYPE, and SQL on flawed queries and VAL and GEN on originals therefore distinguishes evidence-based critique from broad rejection.

The five measures provide complementary evidence about premise critique behavior. DET records whether any flaw is noticed, TYPE whether the intended defect is explained, and SQL whether that assessment changes generation. On originals, VAL records the explicit judgment and GEN the action. Their divergence is informative: under Explicit prompting, LLaMA 4 Maverick retains 0.974 GEN but only 0.363 VAL, so it often generates despite incorrectly judging an original query as flawed; Qwen3.5-397B is low on both and instead follows a broad rejection policy. Evaluation on originals prevents either pattern from being mistaken for robust critique.

B.7 Execution-Based Validation of SQL Generation

GEN is designed to measure generation coverage, namely whether a model preserves SQL generation on valid original queries. GEN and VAL are evaluated on all 1,800 original queries. For the supplementary execution-based analysis, we retain 1,726 originals whose gold SQL produces a reliable reference result in our SQLite evaluation environment; the remaining 74 are excluded only from this analysis because their execution correctness cannot be determined reliably. This yields 27,616 model–query outputs per prompting setting across the 16 evaluated models.

Let g_i indicate whether SQL is generated for the i-th original query, and let e_i indicate whether the generated SQL is execution-equivalent to the gold

Target Model	Setting	Reported Two-Judge Average			Cross-Family Judge Only		
		SQL	DET	TYPE	SQL	DET	TYPE
GPT-5.4	Base	0.117	0.261	0.258	0.109	0.240	0.239
	Implicit	0.682	0.702	0.648	0.681	0.697	0.671
	Explicit	0.960	0.981	0.855	0.959	0.982	0.907
Claude Opus 4.6	Base	0.204	0.281	0.277	0.205	0.296	0.288
	Implicit	0.510	0.579	0.536	0.513	0.579	0.523
	Explicit	0.880	0.899	0.760	0.882	0.896	0.714

Table 7: Leave-self-out evaluation for models that also serve as automatic evaluators. The reported scores average judgments from GPT-5.4 and Claude Opus 4.6. The cross-family scores retain only Claude Opus 4.6 judgments for GPT-5.4 outputs and only GPT-5.4 judgments for Claude Opus 4.6 outputs.

Setting	GEN	EX	EX | GEN
Base	0.987	0.713	0.722
Implicit	0.902	0.652	0.722
Explicit	0.524	0.399	0.762

Table 8: Execution-based validation on the 1,726 original queries with reliably executable gold SQL. EX measures execution accuracy over these queries, treating abstention as incorrect, while EX | GEN measures execution accuracy among outputs for which SQL is generated.

SQL, with $e_i = 0$ when no SQL is generated. We compute overall execution accuracy as

$$\mathrm{EX} = \frac{1}{N} \sum_{i=1}^{N} e_i, \tag{3}$$

and execution accuracy conditional on generation as

$$\mathrm{EX} \mid \mathrm{GEN} = \frac{\sum_{i=1}^{N} e_i}{\sum_{i=1}^{N} g_i}. \tag{4}$$

As shown in Table 8, GEN closely tracks overall execution accuracy, whereas execution accuracy conditional on generation remains comparatively stable across prompting settings. Across the 48 model–setting combinations, GEN is strongly correlated with EX (Pearson's $r = 0.990$, $p < .001$; Spearman's $\rho = 0.888$, $p < .001$). The association remains strong when models are compared within the Implicit ($r = 0.951$, $p < .001$) and Explicit ($r = 0.997$, $p < .001$) settings. Under Base prompting, the correlation is weaker ($r = 0.434$, $p = .093$) because GEN is near its upper bound for almost all models, leaving little variation across models.

More directly, the model-level change in GEN from Implicit to Explicit prompting is strongly correlated with the corresponding change in EX (Pearson's $r = 0.992$, $p < .001$; Spearman's $\rho = 0.985$, $p < .001$). Meanwhile, EX | GEN changes only from 0.722 to 0.762. The decline in overall execution accuracy under Explicit prompting is therefore primarily associated with suppressed SQL generation rather than reduced correctness among the SQL outputs that remain. These results do not make GEN a substitute for execution accuracy; instead, they support its intended use as a measure of whether models preserve SQL generation on valid queries.

B.8 Bootstrap Uncertainty Analysis

We assess the uncertainty of the principal model- and error-level comparisons using 10,000 bootstrap replicates. For comparisons across prompting settings or models, each replicate resamples the 1,800 aligned query instances with replacement, preserving the paired structure of the observations. For error-type results, we resample the 200 source queries within each error type while retaining all outputs from the 16 evaluated models and both judges for each sampled query. We report percentile 95% confidence intervals and two-sided centered-bootstrap tests. These intervals quantify uncertainty over benchmark queries; the evaluated models and judges are treated as fixed.

Table 9 confirms that Explicit prompting increases DET and decreases GEN for every evaluated model. None of the 32 confidence intervals crosses zero, and all changes remain significant after Holm correction ($p_{\mathrm{adj}} < .004$). Even the smallest GEN decrease, observed for LLaMA 4 Maverick, remains distinguishable from zero. The central DET–GEN trade-off is therefore not driven by sampling variation in a small subset of queries.

Uncertainty in model comparisons. Paired comparisons support the claim that model rankings depend on the prompting setting. GPT-6.1-Sol exceeds GPT-5.4 in Base DET by 0.156 (95% CI

Model	ΔDET	ΔGEN
ChatGLM-5.1	0.358 [0.336, 0.380]	−0.465 [−0.488, −0.442]
Claude Opus 4.6	0.319 [0.297, 0.341]	−0.183 [−0.203, −0.164]
Claude Opus 5.5	0.180 [0.162, 0.199]	−0.181 [−0.199, −0.163]
Claude Sonnet 4.6	0.411 [0.389, 0.434]	−0.064 [−0.078, −0.050]
DeepSeek V3.2	0.079 [0.066, 0.092]	−0.306 [−0.327, −0.284]
DeepSeek V4.1 Flash	0.307 [0.286, 0.329]	−0.615 [−0.638, −0.592]
Doubao Seed 2.0 Pro	0.254 [0.235, 0.275]	−0.072 [−0.086, −0.058]
Gemini 3.1 Pro Preview	0.217 [0.198, 0.236]	−0.054 [−0.066, −0.044]
LLaMA 4 Maverick	0.374 [0.351, 0.396]	−0.013 [−0.021, −0.004]
GPT-5.4	0.279 [0.259, 0.299]	−0.672 [−0.694, −0.651]
GPT-5.6-Luna	0.269 [0.249, 0.289]	−0.671 [−0.693, −0.649]
GPT-6-Luna	0.200 [0.182, 0.219]	−0.391 [−0.414, −0.368]
GPT-6.1-Sol	0.215 [0.197, 0.234]	−0.603 [−0.625, −0.580]
Qwen3.5-35B	0.184 [0.166, 0.203]	−0.689 [−0.710, −0.667]
Qwen3.5-122B	0.251 [0.232, 0.271]	−0.361 [−0.387, −0.333]
Qwen3.5-397B	0.245 [0.226, 0.265]	−0.733 [−0.754, −0.713]

Table 9: Paired bootstrap estimates for the change from Implicit to Explicit prompting. Each entry reports the point estimate and 95% confidence interval, with Δ = Explicit − Implicit.

[0.135, 0.176], $p < .001$), whereas their difference under Implicit prompting is not significant (0.014, 95% CI [−0.004, 0.033], $p = .129$). Under Explicit prompting, the ordering reverses: GPT-5.4 exceeds GPT-6.1-Sol by 0.050 (95% CI [0.039, 0.061], $p < .001$).

The non-monotonic Qwen comparison is also statistically supported. Under Implicit prompting, Qwen3.5-35B exceeds Qwen3.5-122B in DET by 0.069 (95% CI [0.054, 0.086], $p < .001$) and Qwen3.5-397B by 0.057 (95% CI [0.040, 0.074], $p < .001$). The difference between the 122B and 397B variants is not significant (−0.013, 95% CI [−0.028, 0.002], $p = .097$). Thus, the 35B result is robust, but it should be interpreted as a setting-specific comparison rather than a general advantage of smaller models.

The DeepSeek comparison further illustrates why performance cannot be summarized by a single ordering. Relative to DeepSeek V3.2 under Implicit prompting, DeepSeek V4.1 Flash has substantially higher GEN on originals (+0.471, 95% CI [0.448, 0.494], $p < .001$) but lower DET on flawed queries (−0.274, 95% CI [−0.296, −0.252], $p < .001$).

Uncertainty across error types. Table 10 confirms that the large differences across error types are not explained by query sampling variation. Relative to the two expression-error subcategories, the two contradiction subcategories are higher by 0.695 in SQL Behavior (95% CI [0.670, 0.718]), 0.730 in DET (95% CI [0.705, 0.754]), and 0.807 in TYPE (95% CI [0.785, 0.828]); all three differences have $p < .001$. The conclusion that overt contradictions are substantially easier than ambiguity and unclear intent is therefore robust, while the wider intervals for several missing-premise categories indicate greater query-level heterogeneity.

C Failure-Mechanism Attribution

C.1 Mechanism Definitions and Annotation Protocol

This analysis builds on the behavioral labels previously assigned by GPT-5.4 and Claude Opus 4.6. We use their existing *SQL Behavior*, *Detection*, and *Type Identification* judgments to identify eligible failures. For mechanism attribution, a separate evaluator pair, GPT-6.1-Sol and Claude Opus 5.5, assigns M01–M04 to these selected failures without re-evaluating the three behavioral dimensions.

Each sampled response is annotated independently by the two mechanism-attribution judges using the same final attribution prompt. The task is multi-label rather than four-way classification: each judge must assess M01–M04 exactly once, and a response may receive multiple mechanisms or none. Table 16 gives the positive tests and the principal boundaries used in the prompt.

Information shown to the attribution judges. For each item, the judges receive (i) the flawed question and its source error type and flaw description; (ii) the tested model's response, separated into its *reason* and *SQL*; (iii) the schema and auxiliary evidence actually exposed to the tested model, together with provenance and exposure-status fields; (iv) the validated *SQL Behavior*, *Detection*, and *Type Identification* labels; (v) the complete reference schema and its provenance, when available; (vi) the paired valid question, SQL, and evidence used during benchmark construction; and (vii) data-coverage and other recovered context metadata.

Code	Error Type	SQL	DET	TYPE
E1.1	Ambiguity	0.210 [0.181, 0.241]	0.242 [0.209, 0.276]	0.138 [0.115, 0.163]
E1.2	Unclear Intent	0.202 [0.177, 0.228]	0.193 [0.170, 0.218]	0.131 [0.111, 0.151]
E2.1	Internal Text Contradiction	0.920 [0.908, 0.932]	0.972 [0.963, 0.980]	0.966 [0.956, 0.976]
E2.2	Question–Schema Contradiction	0.881 [0.855, 0.906]	0.923 [0.896, 0.947]	0.915 [0.887, 0.941]
E3.1	Missing Constraint	0.537 [0.494, 0.580]	0.611 [0.564, 0.657]	0.543 [0.494, 0.593]
E3.2	Missing Schema Premise	0.568 [0.542, 0.594]	0.586 [0.561, 0.612]	0.474 [0.449, 0.498]
E3.3	Semantic Gap	0.710 [0.674, 0.746]	0.757 [0.720, 0.792]	0.726 [0.688, 0.765]
E4.1	Inexpressible in SQL	0.853 [0.829, 0.875]	0.900 [0.877, 0.920]	0.838 [0.811, 0.863]
E4.2	External Knowledge Required	0.830 [0.800, 0.859]	0.861 [0.834, 0.888]	0.841 [0.812, 0.868]

Table 10: Implicit-prompting performance by error type with query-level bootstrap 95% confidence intervals. Models and judges are retained as fixed panels within each resampled query.

Verdicts and evidence requirements. For every mechanism, a judge returns one of `supported`, `not_supported`, or `uncertain`. `supported` requires the mechanism's positive test and the required evidence; `not_supported` means the available record is sufficient to determine that the test is not met; and `uncertain` is reserved for cases in which missing context prevents a defensible decision. Evidence for M01, M02, and M04 must include an exact quote of the affected requirement from the current flawed question and an exact quote from the response's reason or SQL showing the transformation. M03 instead requires an erroneous availability or answerability inference in the response plus counterevidence from verified model-visible context. The output also records the target requirement, transformation type, context basis, a concise explanation, and three control fields: whether the target is complete and explicit, whether any partial implementation remains, and whether the response supplies a concrete replacement. Deterministic post-processing validates the required quotes, context basis, control fields, and same-target exclusivity, and then derives the positive-label list only from validated `supported` verdicts.

M01, M02, and M04 are mutually exclusive when they describe the same atomic requirement: a concrete substitute is M01, a changed satisfaction relation for an explicit contradiction is M02, and complete disappearance without replacement is M04. They may coexist only when they concern distinct requirements. M03 is independent, but the same inference from nonappearance in bounded examples is not counted again as M01. Complete abstention alone is not M04, and ordinary SQL implementation defects are outside this mechanism taxonomy.

Binary coding and strict consensus. The same binary coding is used for automatic and human reliability analyses: `supported` is coded as 1, whereas both `not_supported` and `uncertain` are coded as 0. A mechanism enters the reported label set only when both attribution judges mark it as `supported`. This rule is applied separately to M01–M04, so the final label set is simply the intersection of the judges' supported labels; their complete four-label combinations need not match.

Final attribution prompt. Figure 8 reproduces the task and output contract from the final prompt. The prompt evaluates observable response operations rather than inferring a mechanism from the validated outcome scores; its label-specific decision tests and exclusions are stated in Table 16 and the preceding paragraphs.

C.2 Eligibility and Sampling

The mechanism analysis begins from outputs that both primary evaluators, Claude Opus 4.6 and GPT-5.4, independently regard as failures. All three evaluation fields must be valid: *Detection* and *Type Identification* must be binary, and *SQL Behavior* must be in $\{0, 1, 2\}$. An output is a candidate failure when *Type Identification*$= 0$ or *SQL Behavior*$\neq 2$. Candidate records are paired only when their sample identity and their database ID, question, reason, and SQL match exactly. The two evaluators must then agree on all three fields. Records with non-string reasons, schema-placeholder outputs, or no usable reason/SQL evidence are excluded because they cannot support mechanism attribution.

This strict wash retains 36,684 responses: 18,701 from BIRD and 17,983 from Spider. The exclusion log contains 4,258 otherwise matched candidates: 4,035 evaluator disagreements, 180 non-string reasons, and 43 SQL schema placeholders. We then restore the original generation context and paired valid reference. Before sampling, we remove 4,279 outputs that correctly identify the flaw but still an-

Stage	BIRD	Spider	Total
Strict wash retained	18,701	17,983	36,684
Eligible for attribution	13,709	13,864	27,573
Final sampled responses	4,000	4,000	8,000

Table 11: Failure-mechanism attribution pool after the strict two-evaluator wash, context-based eligibility filtering, and sampling.

Mechanism	Cohen's κ	Raw agreement
M01	0.8862	96.33%
M02	1.0000	100.00%
M03	0.9800	99.67%
M04	0.9074	99.67%
Macro average	**0.9434**	**98.92%**

Table 12: Agreement between the two independent human annotators on the 300-example mechanism-attribution sample.

swer it, 4,801 E3.2 cases for which the historical reduced-schema exposure cannot be recovered, and 31 cases without a recoverable paired-valid reference. This leaves 27,573 eligible failures (13,709 BIRD and 13,864 Spider).

For each dataset–model combination, we sample exactly 250 responses, yielding $16 \times 2 \times 250 = 8{,}000$ observations. Sampling strata are prompting setting $\times$ source error category. Every nonempty stratum receives at least one observation, and the remaining quota is allocated approximately in proportion to stratum size by the largest-remainder method. Within each stratum, records are ranked deterministically by a SHA-256 hash with seed 42. The final sample contains 4,587 Base, 2,433 Implicit, and 980 Explicit responses. Because of the unrecoverable historical schema exposure noted above, E3.2 is not represented in this mechanism-attribution sample. All reported mechanism incidences are calculated directly from these 8,000 sampled responses without reweighting.

Aggregate incidence. Across the 8,000 sampled failures, strict two-judge consensus identifies M01 in 5,808 responses, M02 in 110, M03 in 860, and M04 in 69. Because the mechanisms are non-exclusive, these counts need not sum to the sample size.

C.3 Reliability and Human Validation

Agreement between the attribution judges, GPT-6.1-Sol and Claude Opus 5.5, yields Cohen's κ values of 0.7917, 0.8403, 0.8458, and 0.7955 for M01–M04, respectively, with a macro average of 0.8183.

The human-validation sample contains 300 responses drawn from three tested-model families that are distinct from the attribution-judge families: Gemini 3.1 Pro Preview, DeepSeek V4.1 Flash, and Qwen3.5-397B. We sampled 100 responses from each family, with an exact 50/50 split between BIRD and Spider. Within each model–benchmark pool, the quota was allocated proportionally over prompting-setting $\times$ source-error-category strata by the largest-remainder method, followed by simple random sampling within each stratum using seed 20261011. This yielded 182 Base, 88 Implicit, and 30 Explicit responses across the eight represented error categories; E3.2 remained excluded for the context-recovery reason given above.

The two human annotators label the same 300 responses independently and cannot see one another's decisions. The interface hides the attribution judges' rationales, but it preselects a verdict when GPT-6.1-Sol and Claude Opus 5.5 agree; judge-disagreement cells remain blank and require a human choice. The human results are therefore an assisted validation rather than a fully AI-blinded annotation study. Both annotators completed every item, and their sample IDs match exactly. Following the binary rule used in the AI-judge analysis, we map `supported` to 1 and both `not_supported` and `uncertain` to 0.

The agreement statistics address three nested reproducibility questions. First, the per-mechanism κ values in Table 12 test whether each binary label is stable across humans; their macro average is .9434, accompanied by 98.92% macro-average raw agreement. Second, pooling the four binary decisions over all items summarizes overall decision-level stability ($\kappa = 0.9695$). Third, treating each response's complete four-label vector as one nominal outcome tests whether the entire multi-label configuration is reproduced ($\kappa = 0.9101$).

Human–AI agreement addresses a different question: whether the automatic attributions align with human judgments. Tables 13 and 14 report this comparison. Their final rows first average the two human-specific κ values within each mechanism and then take an unweighted macro average over M01–M04, yielding .9212 for GPT-6.1-Sol and .8748 for Claude Opus 5.5.

Comparison	M01	M02	M03	M04	Macro
Human 1–GPT	0.8656	1.0000	0.9579	0.8299	**0.9133**
Human 2–GPT	0.8705	1.0000	0.9379	0.9074	**0.9290**
Mean over humans	0.8681	1.0000	0.9479	0.8687	**0.9212**

Table 13: Cohen's κ between each human annotator and GPT 6.1 Sol.

Comparison	M01	M02	M03	M04	Macro
Human 1–Claude	0.8984	0.8872	0.8890	0.7967	**0.8678**
Human 2–Claude	0.8832	0.8872	0.8693	0.8872	**0.8817**
Mean over humans	0.8908	0.8872	0.8792	0.8420	**0.8748**

Table 14: Cohen's κ between each human annotator and Claude Opus 5.5.

C.4 Consensus Coverage

The four mechanisms are high-precision response-level attributions rather than an exhaustive partition of failure behavior. Of the 8,000 sampled failures, 1,305 (16.31%) receive no strict-consensus M01–M04 label. This total should not be read as 1,305 cases with no cause. In 591 cases, at least one attribution judge supports a mechanism but the two judges share no positive label: Claude assigns no label while GPT assigns at least one in 222 cases, GPT assigns no label while Claude assigns at least one in 318 cases, and both assign positive but disjoint label sets in 51 cases. The remaining 714 cases (8.93%) are jointly assigned no covered mechanism by either judge.

Among these 714 jointly unlabeled cases, 687 (96.22%) have the validated behavioral profile *SQL Behavior*$= 0$, *Detection*$= 0$, and *Type Identification*$= 0$: the model produces complete SQL without explicitly recognizing the premise flaw or identifying its intended type. This profile establishes a recognition failure, but it does not by itself identify how the response transformed or misused the flawed premise. For example, a response may leave an unspecified condition broad, preserve contradictory predicates literally, or follow a context-grounded reading without injecting a concrete assumption, reinterpreting the contradiction, misusing bounded context, or deleting a complete explicit constraint. We therefore treat *Detection*$= 0$ as a behavioral outcome rather than relabeling it as M01–M04. Accordingly, an empty strict-consensus mechanism set means *no observable covered mechanism*, not an error-free response and not an absence of a behavioral failure.

C.5 Prompt-Type, Benchmark-Type, and Error-Type Analyses

Prompt and benchmark breakdown. The prompt-level reversal between M01 and M03 occurs in both source benchmarks. From Base to Explicit prompting, M01 declines from 72.10% to 25.30% on BIRD and from 88.99% to 39.45% on Spider; over the same contrast, M03 rises from 1.94% to 64.99% and from 1.03% to 48.14%, respectively. The Implicit setting lies between these endpoints in both datasets. M02 and M04 remain rare across prompting settings relative to assumption injection and context/evidence misuse. Because this analysis is conditioned on eligible failures, these percentages describe the composition of observed failures rather than the unconditional probability that a prompt produces an error.

For the prompt-association tests, each of the 8,000 sampled responses is one observation. We construct a separate 3×2 contingency table for each mechanism, crossing prompting setting (Base, Implicit, Explicit) with the mechanism's strict-consensus presence or absence, and apply Pearson's χ^2 test with Cramér's V as the effect size. Because mechanisms are non-exclusive, a response may be positive in more than one table. We adjust the four mechanism-specific p-values as one family using Holm's method; the reported $p < .001$ results for M01 and M03 remain below .001 after adjustment.

Mechanisms by source error category. Table 15 reports strict-consensus incidence within each represented source error category in the analyzed sample. M02 is sharply concentrated in internal textual contradictions, where it appears in 50.69% of sampled failures. M04 is rare overall and is most frequent for question–schema contradictions (4.51%). M01 dominates semantic gaps and both beyond-ability categories, whereas M03 is most common for ambiguity, unclear intent, and missing constraints. Thus, the source taxonomy and response mechanisms are related but not interchangeable: their mapping is many-to-many, because one source category can yield multiple mechanisms and the same mechanism can recur across categories.

D Representative Premise Error Examples

Table 17 presents representative instances across the PremiseSQL taxonomy, illustrating how differ-

Code	Source error category	N	M01	M02	M03	M04
E1.1	Ambiguity	1,999	70.14	0.00	14.46	0.35
E1.2	Unclear Intent	1,926	83.54	0.00	12.51	0.47
E2.1	Internal Text Contradiction	217	1.84	50.69	6.45	3.23
E2.2	Question–Schema Contradiction	355	76.34	0.00	6.48	4.51
E3.1	Missing Constraint	1,228	38.60	0.00	12.46	0.57
E3.3	Semantic Gap	926	89.20	0.00	2.70	0.00
E4.1	Inexpressible in SQL	732	90.03	0.00	9.97	2.32
E4.2	External Knowledge Required	617	91.25	0.00	6.81	0.97

Table 15: Observed strict-consensus mechanism incidence (%) by source error category in the analyzed sample. E3.2 is excluded because its original reduced-schema exposure is unavailable.

ent premise failures manifest in natural-language queries.

Distinguishing Neighboring Categories Several neighboring categories are distinguished by the source of indeterminacy. E1.1 applies when a schema-level condition admits multiple supported mappings or aggregations, whereas E1.2 leaves the overall SQL operation unspecified; E3.1 instead omits a task-defining constraint while preserving that operation. Thus, "sales" in the E1.1 example admits several computations over existing columns, whereas "impactful reviews" in E3.3 lacks a unique rule for mapping the concept to the available indicators. E2.2 and E3.2 may both leave required schema support unavailable to the model, but they arise from different perturbations: E2.2 modifies the query against an unchanged schema, whereas E3.2 preserves the original valid query while withholding part of its required schema context. E3.3 differs from both because related fields remain available but an abstract concept lacks a unique semantic mapping. Finally, E4.1 concerns operations such as forecasting that cannot be performed by standard SQL over the stored data, while E4.2 concerns external facts that are absent from the database.

E Behavioral Case Studies

Figures 9–20 present representative multi-model case studies on flawed queries, while Figures 21–29 present original-query controls under Base, Implicit, and Explicit prompting. Each page reports one response together with its schema evidence, SQL behavior, evaluation labels, and analysis.

Prompt for Ambiguity-type Generation

You are a data synthesis expert for Text-to-SQL evaluation.

Your goal is:

Given a CORRECT original Text-to-SQL question, rewrite it to deliberately introduce a specific type of error while keeping the question otherwise natural and realistic.

Background knowledge:

A necessary premise in Text-to-SQL refers to a query condition that MUST be specified in order to uniquely determine the intended SQL query under a given database schema.

Error definition:

The rewritten question contains an ambiguous description of a NECESSARY SCHEMA-LEVEL QUERY CONDITION, such that the condition may reasonably correspond to MULTIPLE columns, tables, join paths, or aggregation targets in the given schema.

Task constraints:

Rewrite the original question so that:

- EXACTLY ONE necessary query condition becomes ambiguous at the schema level
- The ambiguity leads to AT LEAST TWO different and reasonable SQL interpretations
- Each interpretation must be valid under the given schema

Allowed patterns(use ONE, include but not limited to):

- Entity reference ambiguity: A condition or object in the text may correspond to multiple columns, tables, or foreign key paths in the schema and cannot be uniquely determined.
- Aggregation ambiguity: The description of quantities, totals, averages, or other aggregation metrics is not precise.
- Time/Range ambiguity: Time or numerical ranges are not clearly specified.
- Status ambiguity: The state or category of an object is not clearly defined.
- Relationship ambiguity: The relationship between objects is not clearly defined, such as multiple possible join paths or foreign key relationships.

Strict rules:

- The rewritten question MUST contain only ONE error: schema-level ambiguity
- Do NOT invent tables or columns
- Do NOT clarify or hint at the ambiguity
- Do NOT simply paraphrase the original question

Database schema: schema

Foreign keys: {foreign_keys}

Original question: question

Gold SQL (for reference only): sql

Output: Return ONLY a JSON object with: - {rewritten_question} (string or null) - {error_category} (string or null) - {error_subcategory} (string or null) - {give_up} (boolean) - {give_up_reason} (string or null)

Figure 5: Representative prompt for ambiguity-type instance generation.

Prompt for Ambiguity-type Review

You are a STRICT dataset auditor, not a helper.
Absolute rules:
- Default decision is FAIL
- Your primary job is to FIND PROBLEMS, not to approve
- You MUST actively look for reasons to reject the sample
- If a case is borderline, unclear, or debatable, it MUST FAIL
- Being generous, charitable, or cooperative is an error
- In a healthy dataset, MANY samples should fail

Important:
You MUST first identify potential issues BEFORE deciding whether the sample passes.
If you cannot find any issue, you must explicitly explain why none exist.
Task objective:
Review whether the rewritten question correctly introduces EXACTLY ONE schema-level ambiguity, as specified by the assigned error subcategory.
Review procedure (must follow in order):
Step 1: Actively search for problems
- List ANY potential issues, weaknesses, or ambiguities that could invalidate the sample
- Even minor or borderline issues should be listed

Step 2: Decide the strongest reason to FAIL
- If the sample should be rejected, clearly state why
- If the sample should pass, explain why NO valid failure reason exists

Step 3: Decide the strongest reason to PASS
- State the single strongest argument supporting acceptance

Step 4: Checklist evaluation
Evaluate the following items as TRUE / FALSE / NULL:
- correct_error_type
- single_error_only
- error_matches_subcategory
- no_extra_tables_or_columns
- ambiguity_clarity
- sql_interpretation_ambiguity

Step 5: Final verdict
Set overall_pass = TRUE ONLY IF:
- all checklist items are clearly satisfied, and
- all identified failure concerns are convincingly refuted

Otherwise, overall_pass = FALSE.
Database schema: {schema}
Foreign keys: {foreign_keys}
Original question: {question}
Rewritten question: {rewritten_question}
Assigned error category / subcategory: {error_category} / {error_subcategory}
Gold SQL (reference only): {sql}
Give-up information:
- give_up: {give_up}
- give_up_reason: {give_up_reason}

Output:
Return ONLY a JSON object containing:
- {potential_issues}
- {strongest_reason_to_fail}
- {strongest_reason_to_pass}
- {correct_error_type}
- {single_error_only}
- {error_matches_subcategory}
- {no_extra_tables_or_columns}
- {ambiguity_clarity}
- {sql_interpretation_ambiguity}
- {overall_pass}
- {comments}

Figure 6: Representative prompt for ambiguity-type instance review.

Prompt for Ambiguity-type Evaluation

You are an expert evaluator for Text-to-SQL systems.

Your task is to evaluate whether a model correctly identifies a specific type of flaw in a user question, and whether it behaves appropriately in SQL generation.

Dataset note:

This dataset is deliberately constructed such that EVERY user question is flawed.

None of the questions has a single correct or well-defined SQL query.

The goal is NOT to produce SQL, but to evaluate whether the model can:

- recognize the flaw

- respond appropriately, including explicitly refraining from generating SQL

Ground-truth error type:

Error category: Expression Error

Subcategory: Ambiguity

Definition:

A required query condition in the user question contains ambiguity in reference, scope, or target object, such that it cannot be uniquely determined using only the given schema and context, resulting in multiple reasonable SQL interpretations.

Input:

User question: {question}

Model output:

Reason: {reason}

SQL: {sql}

Database schema: {schema}

Evaluation tasks and scoring:

1. SQL Behavior

- 0 = directly generates a full SQL query despite ambiguity

- 1 = partially abstains but still generates SQL

- 2 = correctly refrains from generating a valid SQL query

2. Detection

- 1 = explicitly states that the question is flawed

- 0 = fails to recognize the issue

3. Type Identification

If Detection = 0, then Type Identification MUST be 0.

Type Identification does NOT require explicit naming of the error type.

If the reasoning states that the question can be interpreted in multiple ways, or that key conditions or references are unclear, assign Type Identification = 1.

- 1 = reasoning clearly indicates ambiguity

- 0 = reasoning does not support the presence of ambiguity

Output format:

Return ONLY a JSON object:

```
{
  "sql_behavior": 0, 1, or 2,
  "detection": 0 or 1,
  "type_identification": 0 or 1
}
```

Figure 7: Representative prompt for ambiguity-type automatic evaluation.

ID	Mechanism	`supported` when	Boundary and exclusion rules
M01	Assumption Injection	The response makes a material, answer-changing commitment not established by the current question or model-visible context, such as an invented value, threshold, metric, proxy, entity, schema role, relationship, or semantic mapping. A concrete response commitment and a material alternative or direct counterevidence must be identifiable.	Ordinary literal grounding, declared foreign-key use, and surface value normalization are not M01. Nonappearance in bounded rows alone is assessed as M03. Merely executing SQL, leaving an underspecified request broad, or making an arithmetic/grouping error does not establish M01.
M02	Contradiction Reinterpretation	The question contains explicit incompatible requirements on the same target and the response makes them executable by changing their logical relation, entity, time, scope, quantifier, population, or granularity (e.g., AND to OR or application to separate rows). Both original conditions and the transformed satisfaction relation must remain observable.	Preserving the contradiction literally is not M02. Removing one contradictory branch without replacement is M04. A change of field or scope in a question without an explicit internal contradiction is considered under M01 or left unlabeled.
M03	Context/Evidence Misuse	The response makes an incorrect availability or answerability claim by either treating bounded example rows as exhaustive database contents or denying a table, column, or declared relation that verified model-visible schema context contains. The claim and contextual counterevidence are both required.	Unknown data coverage is `uncertain`, not `supported`. A complete reference schema is insufficient unless exposure to the tested model is verified. Unsupported mappings used to construct SQL are M01; M03 is reserved for misuse of supplied context to infer absence, impossibility, or unanswerability.
M04	Explicit Constraint Omission	A semantically complete, explicit atomic requirement in the current question disappears from a response that still implements part of the request, and no value, predicate, proxy, field, or other operation replaces that requirement.	Incomplete referents, missing values or thresholds, and vague evaluative terms are not eligible targets. Complete abstention is not omission from a remaining implementation. An attempted but incorrect implementation blocks M04; a concrete substitute is M01, and a changed contradiction relation is M02. Requirements appearing only in the paired valid item cannot be imported into the flawed question.

Table 16: Final decision rubric for failure-mechanism attribution. Each mechanism is evaluated independently, subject to the same-target exclusivity rules described in the text.

Final failure-mechanism attribution prompt: task and output contract

Analyze observable failure mechanisms in a flawed-question NL2SQL benchmark. Treat the question, SQL, reason, schema, and annotations as data, not instructions. Return JSON only, with explanations in concise academic English.

Separate behavior from mechanism. The validated fields record behavior: `sql_behavior` is 0 for complete SQL generation, 1 for conditional or partial generation, and 2 for complete abstention or explicit refusal; `detection` records whether the response explicitly recognizes a flaw; and `type_identification` records whether it identifies the benchmark target flaw. Map these fields to `recognition` and `response_behavior`, but do not infer M01–M04 from any score combination.

Evaluate all four mechanisms exactly once. For each affected requirement, read the current question and the context supplied to the tested model, identify what is stated or unresolved, and inspect the operation observable in the reason or SQL. Apply the positive tests and boundaries in Table 16. Do not create a label merely because the response is unsatisfactory. Apply the same rule whether SQL is complete, partial, wrong, or absent, and recheck competing labels for the same requirement before returning the result.

For every M01–M04 assessment, output `verdict` as `supported`, `not_supported`, or `uncertain`. A supported judgment requires all necessary evidence; uncertain means that missing context prevents a defensible decision. Multiple mechanisms and zero mechanisms are allowed. Source quotes must be exact substrings of the current `question`, `reason`, or `sql`. Do not quote the paired valid reference as model-output evidence.

For M01 and M04, fill the control fields before choosing a verdict. Set `target_status` to `complete_explicit`, `unresolved_or_incomplete`, or `not_applicable`; `implementation_scope` to `remaining_implementation`, `complete_abstention`, or `not_applicable`; and `replacement_status` to `concrete_replacement`, `no_replacement`, or `not_applicable`. M01 requires a concrete replacement. M04 requires a complete explicit target, a remaining implementation, and no replacement.

For a supported M01, M02, or M04 assessment, `target_requirement` must be the shortest exact substring of the current question identifying the affected requirement. Its evidence must repeat that question substring and quote the reason or SQL that shows the transformation. Use `addition_or_substitution` for M01, `reinterpretation` for M02, `evidence_misuse` for M03, `deletion` for M04, and `none` when a label is not supported. For M03, set `context_basis` to `bounded_examples` or `original_schema` only when that evidence path is verified; otherwise use `none`.

The current question is the only source of requirements that can be omitted under M04. The paired valid question, paired valid SQL, and paired valid evidence are post-hoc benchmark-construction references that were not shown to the tested model. They may help locate the changed premise, but cannot establish a label. Likewise, `schema` denotes the schema exposed to the tested model, whereas `reference_schema` is a complete diagnostic reference and is not proof of exposure. Bounded examples are not complete database contents.

Return a JSON object with the following fields:

`recognition`: one of `explicit_flaw`, `not_explicit`, `uncertain`;

`response_behavior`: one of `sql_generated`, `conditional_sql`, `abstained`, `malformed`, `uncertain`;

`label_assessments`: one object for each of M01–M04, containing `label_id`, `verdict`, `context_basis`, `target_requirement`, `target_status`, `implementation_scope`, `replacement_status`, `transformation`, exact-quote `evidence`, `context_evidence`, and `explanation`;

`novel_candidate`: null unless a recurring premise-handling cause falls outside M01–M04 and has direct question and output evidence;

`evidence_limitations`: a concise account of any context needed for an unsupported or uncertain decision.

Append the item as `INPUT DATA`, containing `question`, `source_error_type`, `validated_evaluation`, `original_evidence`, `flaw_description`, `schema`, schema provenance and exposure fields, `original_schema_linker`, `reference_schema` and provenance, `paired_valid_question`, `paired_valid_sql`, `paired_valid_evidence` and provenance, `data_coverage`, `context_evidence`, and `model_output`.

Figure 8: Final prompt used for failure-mechanism attribution. The instance-specific fields listed in the final paragraph are appended as structured input.

Error Type	Question	Error Reason
Ambiguity	Which product had the highest sales in 2018?	The phrase "highest sales" admits multiple schema-supported interpretations, including the number of orders, total quantity sold, gross revenue, and discounted revenue. Because these alternatives can produce different product rankings, the query does not determine a unique SQL formulation.
Unclear Intent	For the lists created by the user who rated the movie "Pavee Lackeen: The Traveller Girl" on 3/27/2011 at 2:06:34 AM, what is the situation with their followers?	The phrase "what is the situation with their followers" does not specify whether to retrieve follower counts, aggregate them, compare them, or rank the lists. The intended SQL operation is therefore indeterminate.
Internal Text Contradiction	List the movies that were released in 1945 and also released after the year 2000.	The question contains mutually contradictory constraints that cannot be simultaneously satisfied under the schema. A single movie record has only one release date, making it impossible for a movie to be released in both 1945 and after 2000.
Question–Schema Contradiction	Who played the character named "Chanice Kobolowski," and what is that actor's email address?	The schema supports identifying the actor through the `characters` and `actor` tables, but the `actor` table contains no email-address column and no related table provides this attribute. The requested email address therefore cannot be retrieved from the database.
Missing Constraint	What's the percentage of coins that is higher than the price 1 hour ago in that day? List the names of these coins.	The expression "that day" has no specified antecedent and leaves the required date unspecified. Because the historical table contains observations for multiple dates, the relevant time point cannot be determined.
Missing Schema Premise	List the powers of Hunter Zolomon.	We intentionally removed 30% of the gold schema beforehand. Therefore, the schema does not provide a way to identify Hunter Zolomon or to link a character entity to their associated powers. There is no table or relationship connecting characters to superpowers, making the query unanswerable under the given schema.
Semantic Gap	For user ID 3, what is their yearly average number of impactful reviews?	The concept of "impactful reviews" has no explicit or implicit semantic mapping in the schema, and no business definition is provided. As a result, the metric cannot be grounded or computed from available tables or fields.
Inexpressible in SQL	Among the worst directors, which genres are they most likely to improve in over the next five years, based on their past movies and ratings?	The database can identify the relevant directors and their historical genres, but determining which genres are most likely to improve over the next five years requires predictive modeling beyond the filtering, aggregation, sorting, and joining operations supported by standard SQL.
External Knowledge Required	Which actor has acted in at least 2 films that were nominated for an Academy Award (Oscars)? Please list their IDs.	The correctness of the query depends on external knowledge not represented in the schema, as Oscar nomination information is not stored in any table or field, making the condition uncomputable within the database.

Table 17: Examples of Expression Error, Premise Contradiction, Missing Premise, and Beyond Ability types.

PREMISESQL CASE STUDY

E1.1 Ambiguity: Flawed query

USER INPUT

“What are all the titles?”

SCHEMA EVIDENCE

Both Songs and Albums contain a Title column, so the requested entity is unresolved.

MODEL OUTPUT

Llama 4 Maverick

Response excerpt

“The question asks for all song titles, which can be retrieved from the Songs table.”

SQL behavior

```
SELECT Title FROM Songs;
```

EVALUATION

Behavior: MISSED FLAW + SQL

Metrics: DET=0, TYPE=0, SQL=0

ANALYSIS

The response silently resolves the ambiguous Title reference as Songs.Title. Because Albums.Title is equally compatible with the query, the generated SQL lacks grounding and the flaw remains undetected.

Figure 9: E1.1 Ambiguity, LLaMA 4 Maverick. The model silently selects `Songs.Title`, misses the unresolved reference, and generates SQL (DET=0, TYPE=0, SQL=0).

PREMISESQL CASE STUDY

E1.1 Ambiguity: Flawed query

USER INPUT

“What are all the titles?”

SCHEMA EVIDENCE

Both Songs and Albums contain a Title column, so the requested entity is unresolved.

MODEL OUTPUT

Claude Opus 5.5

Response excerpt

“The question is ambiguous because both Songs and Albums contain a Title column.”

SQL behavior

```
SELECT Title FROM Songs UNION SELECT Title FROM Albums;
```

EVALUATION

Behavior: DETECTED FLAW + SQL

Metrics: DET=1, TYPE=1, SQL=0

ANALYSIS

The model identifies that Title could refer to either table but combines both interpretations without user support. It detects the flaw, yet its UNION SQL invents a new resolution and still generates SQL.

Figure 10: E1.1 Ambiguity, Claude Opus 5.5. The model recognizes the ambiguity but combines both interpretations without user support (DET=1, TYPE=1, SQL=0).

PREMISESQL CASE STUDY

E1.1 Ambiguity: Flawed query

USER INPUT

"What are all the titles?"

SCHEMA EVIDENCE

Both Songs and Albums contain a Title column, so the requested entity is unresolved.

MODEL OUTPUT

DeepSeek V3.2

Response excerpt

"The question does not specify Songs.Title, Albums.Title, or both."

SQL behavior

No SQL generated.

EVALUATION

Behavior: DETECTED FLAW + ABSTAIN

Metrics: DET=1, TYPE=1, SQL=2

ANALYSIS

The model identifies the unresolved choice between Songs.Title and Albums.Title. By requesting clarification instead of selecting or combining interpretations, it preserves the ambiguity and appropriately abstains.

Figure 11: E1.1 Ambiguity, DeepSeek V3.2. The model identifies the unresolved reference and abstains from generating unsupported SQL (DET=1, TYPE=1, SQL=2).

PREMISESQL CASE STUDY

E2.1 Text Contradiction: Flawed query

USER INPUT

"Which episode title is simultaneously the winner and the loser of the Best International TV Series award in 2017?"

SCHEMA EVIDENCE

Award results can be queried, but the query assigns mutually exclusive winner and loser statuses to the same episode and award context.

MODEL OUTPUT

GPT-6.1 Sol

Response excerpt

"The schema supports checking whether the same episode has both Winner and Loser results."

SQL behavior

```
HAVING winner_count > 0 AND loser_count > 0
```

EVALUATION

Behavior: MISSED FLAW + SQL

Metrics: DET=0, TYPE=0, SQL=0

ANALYSIS

The model interprets the mutually exclusive winner and loser conditions as ordinary database predicates. Its HAVING clause searches for an impossible conjunction, so the contradiction is missed and SQL is generated.

Figure 12: E2.1 Internal Text Contradiction, GPT 6.1 Sol. The model treats mutually exclusive winner and loser conditions as ordinary predicates and misses the flaw (DET=0, TYPE=0, SQL=0).

PREMISESQL CASE STUDY

E2.1 Text Contradiction: Flawed query

USER INPUT

"Which episode title is simultaneously the winner and the loser of the Best International TV Series award in 2017?"

SCHEMA EVIDENCE

Award results can be queried, but the query assigns mutually exclusive winner and loser statuses to the same episode and award context.

MODEL OUTPUT

Claude Sonnet 4.6

Response excerpt

"An episode cannot simultaneously be both winner and loser of the same award."

SQL behavior

```
HAVING winner_count > 0 AND loser_count > 0
```

EVALUATION

Behavior: DETECTED FLAW + SQL

Metrics: DET=1, TYPE=1, SQL=0

ANALYSIS

The model explicitly recognizes that one episode cannot hold both statuses in the same context. It nevertheless encodes that contradiction in HAVING, producing detection without behavioral correction.

Figure 13: E2.1 Internal Text Contradiction, Claude Sonnet 4.6. The model states the contradiction but still encodes the incompatible requirements in SQL (DET=1, TYPE=1, SQL=0).

PREMISESQL CASE STUDY

E2.1 Text Contradiction: Flawed query

USER INPUT

"Which episode title is simultaneously the winner and the loser of the Best International TV Series award in 2017?"

SCHEMA EVIDENCE

Award results can be queried, but the query assigns mutually exclusive winner and loser statuses to the same episode and award context.

MODEL OUTPUT

Claude Opus 5.5

Response excerpt

"A single episode cannot be both winner and loser of the same award in the same year."

SQL behavior

No SQL generated.

EVALUATION

Behavior: DETECTED FLAW + ABSTAIN

Metrics: DET=1, TYPE=1, SQL=2

ANALYSIS

The model recognizes that the requested winner and loser statuses are incompatible for the same episode and award. It preserves the constraint by declining to generate SQL after detection.

Figure 14: E2.1 Internal Text Contradiction, Claude Opus 5.5. The model recognizes the incompatible requirements and abstains (DET=1, TYPE=1, SQL=2).

PREMISESQL CASE STUDY

E3.3 Semantic Gap: Flawed query

USER INPUT

“How many serious allergies does Mrs. Saundra Monahan have?”

SCHEMA EVIDENCE

The allergies table records codes and descriptions but contains no severity or seriousness attribute.

MODEL OUTPUT

Claude Sonnet 4.6

Response excerpt

“The question is answerable where the allergy description indicates seriousness.”

SQL behavior

```
DESCRIPTION LIKE '%serious%' OR DESCRIPTION LIKE
'%severe%'
```

EVALUATION

Behavior: MISSED FLAW + SQL

Metrics: DET=0, TYPE=0, SQL=0

ANALYSIS

The model assumes that words such as ‘serious’ or ‘severe’ in DESCRIPTION define clinical severity. The schema provides no such mapping, so the generated SQL invents evidence and misses the semantic gap.

Figure 15: E3.3 Semantic Gap, Claude Sonnet 4.6. The model invents a keyword-based proxy for clinical seriousness and misses the flaw (DET=0, TYPE=0, SQL=0).

PREMISESQL CASE STUDY

E3.3 Semantic Gap: Flawed query

USER INPUT

“How many serious allergies does Mrs. Saundra Monahan have?”

SCHEMA EVIDENCE

The allergies table records codes and descriptions but contains no severity or seriousness attribute.

MODEL OUTPUT

GPT-5.4

Response excerpt

“The schema has no severity field. The SQL counts all recorded allergies as the closest interpretation.”

SQL behavior

```
SELECT COUNT(*) FROM allergies ...
```

EVALUATION

Behavior: DETECTED FLAW + SQL

Metrics: DET=1, TYPE=1, SQL=0

ANALYSIS

The model acknowledges that the schema lacks a severity attribute but substitutes a count of all allergies. This removes the qualifier serious, so flaw detection succeeds while SQL answers a different question.

Figure 16: E3.3 Semantic Gap, GPT-5.4. The model acknowledges the missing severity attribute but substitutes a count of all allergies (DET=1, TYPE=1, SQL=0).

PREMISESQL CASE STUDY

E3.3 Semantic Gap: Flawed query

USER INPUT

“How many serious allergies does Mrs. Saundra Monahan have?”

SCHEMA EVIDENCE

The allergies table records codes and descriptions but contains no severity or seriousness attribute.

MODEL OUTPUT

ChatGLM 5.1

Response excerpt

“The allergies table has no severity field, so the number of serious allergies cannot be determined.”

SQL behavior

No SQL generated.

EVALUATION

Behavior: DETECTED FLAW + ABSTAIN

Metrics: DET=1, TYPE=1, SQL=2

ANALYSIS

The model keeps the distinction between recorded allergy descriptions and an unsupported seriousness judgment. Because the schema cannot operationalize serious, it detects the gap and appropriately withholds SQL.

Figure 17: E3.3 Semantic Gap, ChatGLM-5.1. The model preserves the unsupported distinction and abstains (DET=1, TYPE=1, SQL=2).

PREMISESQL CASE STUDY

E4.1 Inexpressible in SQL: Flawed query

USER INPUT

“Based on the existing data, how many tweets are likely to be posted next Wednesday?”

SCHEMA EVIDENCE

The table stores historical weekday labels but provides neither a forecasting model nor sufficient temporal information for predicting next Wednesday.

MODEL OUTPUT

Llama 4 Maverick

Response excerpt

“The question is answerable assuming next Wednesday refers to recorded weekday data.”

SQL behavior

```
SELECT COUNT(*) FROM twitter WHERE Weekday =
'Wednesday';
```

EVALUATION

Behavior: MISSED FLAW + SQL

Metrics: DET=0, TYPE=0, SQL=0

ANALYSIS

The model reinterprets a forecast for next Wednesday as a count of historical Wednesday rows. This substitution turns prediction into retrieval, so it misses the flaw and generates an unsupported proxy.

Figure 18: E4.1 Inexpressible in SQL, LLaMA 4 Maverick. The model turns a forecasting request into a historical count and misses the flaw (DET=0, TYPE=0, SQL=0).

PREMISESQL CASE STUDY

E4.1 Inexpressible in SQL: Flawed query

USER INPUT

“Based on the existing data, how many tweets are likely to be posted next Wednesday?”

SCHEMA EVIDENCE

The table stores historical weekday labels but provides neither a forecasting model nor sufficient temporal information for predicting next Wednesday.

MODEL OUTPUT

DeepSeek V4.1 Flash

Response excerpt

“The database contains only historical data, so the forecast cannot be determined.”

SQL behavior

```
SELECT COUNT(*) AS likely_tweets ... WHERE Weekday =
'Wednesday';
```

EVALUATION

Behavior: DETECTED FLAW + SQL

Metrics: DET=1, TYPE=1, SQL=0

ANALYSIS

The model states that historical rows cannot determine the requested forecast but still returns their count as a proxy. It detects the limitation, yet SQL generation does not resolve it.

Figure 19: E4.1 Inexpressible in SQL, DeepSeek V4.1 Flash. The model recognizes the forecasting limitation but returns a historical count as a proxy (DET=1, TYPE=1, SQL=0).

PREMISESQL CASE STUDY

E4.1 Inexpressible in SQL: Flawed query

USER INPUT

“Based on the existing data, how many tweets are likely to be posted next Wednesday?”

SCHEMA EVIDENCE

The table stores historical weekday labels but provides neither a forecasting model nor sufficient temporal information for predicting next Wednesday.

MODEL OUTPUT

ChatGLM 5.1

Response excerpt

“Historical weekday labels cannot determine how many tweets will appear next Wednesday.”

SQL behavior

No SQL generated.

EVALUATION

Behavior: DETECTED FLAW + ABSTAIN

Metrics: DET=1, TYPE=1, SQL=2

ANALYSIS

The model distinguishes a future prediction from a historical weekday lookup. Since the schema contains no forecasting basis, it detects the inexpressibility and correctly abstains from SQL generation.

Figure 20: E4.1 Inexpressible in SQL, ChatGLM-5.1. The model detects the unsupported prediction request and abstains (DET=1, TYPE=1, SQL=2).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Base)

USER INPUT

“When did Roma host the Olympic Games?”

SCHEMA EVIDENCE

games provides games_year; games_city connects each game to city, and city provides city_name, so SQL can join the three tables and filter for Roma.

MODEL OUTPUT

ChatGLM 5.1

Response excerpt

“Join games, games_city, and city, then filter city_name = 'Roma'.”

SQL behavior

```
SELECT g.games_year FROM games g JOIN games_city gc ... JOIN city c
... WHERE c.city_name = 'Roma';
```

EVALUATION

Behavior: SQL GENERATED

Metrics: GEN=1

ANALYSIS

The model separates SQL generation from answer existence. Because the schema supports joining city_name to games_year, it correctly generates SQL.

Figure 21: Original-query control under Base prompting, ChatGLM-5.1. The model constructs the valid schema-supported join and generates SQL (GEN=1).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Base)

USER INPUT

“When did Roma host the Olympic Games?”

SCHEMA EVIDENCE

games provides games_year; games_city connects each game to city, and city provides city_name, so SQL can join the three tables and filter for Roma.

MODEL OUTPUT

Doubao Seed 2.0 Pro

Response excerpt

“The sample has no Roma row, so the query will return an empty result set.”

SQL behavior

```
SELECT g.games_year FROM city c JOIN games_city gc ... JOIN games g
... WHERE c.city_name = 'Roma';
```

EVALUATION

Behavior: SQL GENERATED

Metrics: GEN=1

ANALYSIS

The model notices that Roma is absent from the supplied rows but treats an empty result as valid SQL behavior. It still constructs the schema-supported join, so GEN remains 1.

Figure 22: Original-query control under Base prompting, Doubao Seed 2.0 Pro. The model notes that the supplied rows contain no matching city but correctly separates answer existence from query validity and generates SQL (GEN=1).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Base)

USER INPUT

“When did Roma host the Olympic Games?”

SCHEMA EVIDENCE

games provides games_year; games_city connects each game to city, and city provides city_name, so SQL can join the three tables and filter for Roma.

MODEL OUTPUT

Claude Opus 5.5

Response excerpt

“Roma does not appear in the city table, so the year cannot be determined.”

SQL behavior

No SQL generated.

EVALUATION

Behavior: NO SQL GENERATED

Metrics: GEN=0

ANALYSIS

The model confuses missing sample rows with an inability to generate SQL.

Figure 23: Original-query control under Base prompting, Claude Opus 5.5. The model treats the absence of a matching city in the supplied rows as grounds for withholding otherwise valid SQL (GEN=0).

PREMISESQL CASE STUDY

E1.1 counterpart: Original query (Implicit)

USER INPUT

“What is the most common first name?”

SCHEMA EVIDENCE

person provides person_name, so SQL can extract the first token with string functions, group by it, and rank the counts.

MODEL OUTPUT

Claude Opus 5.5

Response excerpt

“person_name can be parsed to extract first names and count their occurrences.”

SQL behavior

```
SELECT SUBSTR(person_name, 1, INSTR(person_name || ' ', ' ') - 1) AS
first_name, COUNT(*) ... ORDER BY COUNT(*) DESC LIMIT 1;
```

EVALUATION

Behavior: VALID + SQL

Metrics: VAL=1, GEN=1

ANALYSIS

The model recognizes that person_name can be parsed with string functions and then aggregated. It therefore accepts the query and generates workable SQL, matching VAL=1 and GEN=1.

Figure 24: Original-query control under Implicit prompting, Claude Opus 5.5. The model accepts the valid request and generates a workable query (VAL=1, GEN=1).

PREMISESQL CASE STUDY

E1.1 counterpart: Original query (Implicit)

USER INPUT

“What is the most common first name?”

SCHEMA EVIDENCE

person provides person_name, so SQL can extract the first token with string functions, group by it, and rank the counts.

MODEL OUTPUT

Llama 4 Maverick

Response excerpt

“The question is ambiguous because it does not specify the source of the first names.”

SQL behavior

```
SELECT SUBSTR(person_name, 1, INSTR(person_name, ' ') - 1) AS
first_name ... ORDER BY COUNT(*) DESC LIMIT 1;
```

EVALUATION

Behavior: INVALID + SQL

Metrics: VAL=0, GEN=1

ANALYSIS

The model calls the query ambiguous but nevertheless extracts the first token from person_name. This mismatch between its validity judgment and usable SQL yields VAL=0 and GEN=1.

Figure 25: Original-query control under Implicit prompting, LLaMA 4 Maverick. The model incorrectly calls the request ambiguous but still generates SQL (VAL=0, GEN=1).

PREMISESQL CASE STUDY

E1.1 counterpart: Original query (Implicit)

USER INPUT

“What is the most common first name?”

SCHEMA EVIDENCE

person provides person_name, so SQL can extract the first token with string functions, group by it, and rank the counts.

MODEL OUTPUT

DeepSeek V3.2

Response excerpt

“person_name stores full names, not a separate first-name field.”

SQL behavior

No SQL generated.

EVALUATION

Behavior: INVALID + ABSTAIN

Metrics: VAL=0, GEN=0

ANALYSIS

The model treats the absence of a dedicated first_name column as blocking, although person_name can be parsed. It turns an implementation choice into a validity failure and abstains.

Figure 26: Original-query control under Implicit prompting, DeepSeek V3.2. The model mistakes the absence of a dedicated first-name column for invalidity and abstains (VAL=0, GEN=0).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Explicit)

USER INPUT

“How many tweets have been posted on Wednesday?”

SCHEMA EVIDENCE

twitter provides Weekday, so SQL can filter Weekday = 'Wednesday' and count matching rows.

MODEL OUTPUT

DeepSeek V3.2

Response excerpt

“Weekday exists in the twitter table and can be filtered for Wednesday.”

SQL behavior

```
SELECT COUNT(*) FROM twitter WHERE Weekday =
'Wednesday';
```

EVALUATION

Behavior: VALID + SQL

Metrics: VAL=1, GEN=1

ANALYSIS

The model uses the available Weekday column exactly as the question requires. The filter and count are directly supported by the schema, so both validity and generation receive positive labels.

Figure 27: Original-query control under Explicit prompting, DeepSeek V3.2. The model rejects the prompt’s unsupported claim of a flaw and generates SQL (VAL=1, GEN=1).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Explicit)

USER INPUT

“How many tweets have been posted on Wednesday?”

SCHEMA EVIDENCE

twitter provides Weekday, so SQL can filter Weekday = 'Wednesday' and count matching rows.

MODEL OUTPUT

Doubao Seed 2.0 Pro

Response excerpt

“No records are inserted, so the number of Wednesday tweets cannot be calculated.”

SQL behavior

```
SELECT COUNT(*) AS WednesdayTweetCount FROM twitter
WHERE Weekday = 'Wednesday';
```

EVALUATION

Behavior: INVALID + SQL

Metrics: VAL=0, GEN=1

ANALYSIS

The model mistakes missing inserted rows for invalidity, even though the SQL count remains well defined. It then generates the correct filter, creating the VAL=0 and GEN=1 combination.

Figure 28: Original-query control under Explicit prompting, Doubao Seed 2.0 Pro. The model incorrectly judges the request invalid but nevertheless generates the correct filter (VAL=0, GEN=1).

PREMISESQL CASE STUDY

E4.1 counterpart: Original query (Explicit)

USER INPUT

“How many tweets have been posted on Wednesday?”

SCHEMA EVIDENCE

twitter provides Weekday, so SQL can filter Weekday = 'Wednesday' and count matching rows.

MODEL OUTPUT

Qwen3.5-397B

Response excerpt

“The schema does not prove whether Wednesday is stored as a full name or abbreviation.”

SQL behavior

No SQL generated.

EVALUATION

Behavior: INVALID + ABSTAIN

Metrics: VAL=0, GEN=0

ANALYSIS

The model treats uncertainty about the stored weekday format as a reason to reject the query. Because SQL can still test the stated value, the abstention produces VAL=0 and GEN=0.

Figure 29: Original-query control under Explicit prompting, Qwen3.5-397B. The model accepts the false-invalid framing and abstains (VAL=0, GEN=0).