arXiv:2610.26406v1 [cs.CL] 19 Oct 2026

Expert L2 English Speech Evaluation on a Gaming PC:
Accuracy, Fluency, and Prosody with XLS-R, XEUS, and PhoneticXeus

Richard Rose | HUFS | richard.rose@hufs.ac.kr

Abstract

oollim is a locally executable system for evaluating L2 English pronunciation accuracy, fluency, and prosody on a consumer computer. The system combines pretrained XLS-R, XEUS, and PhoneticXeus representations with reference-text features and ridge regressions trained on SpeechOcean762, the Korean Learners' L2 Speech Corpus (English), J-AESOP, and their combinations. On 305 independently held-out recordings assessed by three expert raters, several ridge configurations substantially outperform archived Azure API scores obtained in January 2026 in both Spearman correlation and quadratic weighted kappa across all three constructs. These results support expert speech evaluation on consumer hardware using pretrained representations and construct-specific regressions, without neural-model fine-tuning or calibration on the private evaluation data.

1 Introduction

Pronunciation evaluation requires distinct judgments of accuracy, fluency, and prosody. Accuracy concerns the realization of expected consonants and vowels [1, 2]; fluency concerns speech continuity, including speed, pauses, and repairs [3, 4]; prosody concerns rhythm, stress, intonation, and phrasing [5, 6]. These related dimensions inform pronunciation judgments but are not interchangeable [2, 7]. Segmental accuracy concerns distinctions between speech sounds, whereas intelligibility concerns what listeners understand and comprehensibility their perceived processing difficulty; a noticeable accent need not prevent understanding [8]. An accuracy score does not directly measure intelligibility. Fluency research distinguishes underlying processing efficiency, observable timing, and perceived smoothness [9]. Speech rate, pause placement and duration, and repairs contribute differently to listener judgments, so fluency is not simply speaking quickly [3, 4]. Prosody organizes speech through lexical stress, prominence, phrasing, and intonation, with pitch, duration, and intensity providing interacting cues [5, 6]. Appropriate stress can improve listeners' recall and evaluations of a speaker [10]; contextual appropriateness therefore matters alongside acoustic variation. Although their acoustic correlates overlap, the three constructs require separate targets and useful measurements for educators; the scores need not come from the same representation combination or scoring function. XLS-R, XEUS, and PhoneticXeus provide complementary sources of speech information. This study examines whether these existing models, combined with modest supervised scoring functions, can support expert L2 English evaluation on an ordinary gaming PC. With computationally expensive pretraining supplied by the original model developers, the contribution is the selection and combination of their representations, reference-text evidence, and corpus-specific assessment supervision. The central evidence is an external evaluation on independently collected, privately held recordings, while training-corpus development results serve model selection and reproducibility and are provided in the repository. The external comparison includes Azure pronunciation assessment [11], seven ridge-training families, and two human agreement baselines, with both rank association and numerical agreement measured. This design separates three questions: whether scores rank recordings similarly to human judgments, whether they agree on the numerical scale, and how that agreement compares with agreement among the human raters [12].

2 Speech models and related work

Automated assessment has long combined acoustic evidence with supervised mappings to human judgments. SpeechRater derives pronunciation, fluency, and other proficiency features from recognized speech and uses learned scoring models [13]. Neural assessment extends this approach by learning representations and their relationship to rated performance. GOPT uses phonetic posterior features and canonical phone identities in a transformer trained jointly on phone-, word-, and utterance-level human annotations [14]. On SpeechOcean762, the 2,500-utterance training and test partitions provide the same assessment targets on different recordings [15, 14]. The scoring model is trained on human ratings and evaluated against ratings of held-out recordings. Kim et al. adapt wav2vec 2.0

and HuBERT with CTC, then combine contextual representations and reference text in a bidirectional LSTM scoring model, evaluating against expert scores on separate partitions of Korean learner speech and SpeechOcean762 [16]. Cai et al. connect hierarchical neural assessment to applied-linguistics constructs, adapting scoring toward human judgments of segmental, suprasegmental, and temporal performance [2]. Saito et al. use regression over phonological, prosodic, and temporal features to predict comprehensibility, testing generalization on additional recordings and an interview task [17]. These studies train scoring models against human ratings and assess multiple dimensions or transfer to new data, although differences in targets, rating panels, and metrics prevent direct numerical comparison with *oollim*. XLS-R extends cross-lingual self-supervised speech representation learning [18]; XEUS learns multilingual representations across thousands of languages [19]; PhoneticXeus adds phone recognition and supplies hidden representations and phone evidence [20]. In *oollim* their pretrained weights remain unchanged, while ridge regression learns regularized linear mappings to assessment labels [21]. *oollim* combines these representations with reference-phone diagnostics and construct-specific regressions trained on one or more corpora, then evaluates them on an independently collected passage-reading cohort to test transfer beyond holding out recordings within a training corpus. Earlier neural-head, temporal, and acoustic-feature experiments motivated the ridge baseline; the evaluation does not cover every neural assessment architecture. The comparison therefore focuses on the external transfer of compact scoring models.

3 Method

3.1 Audio, reference text, and representations

Whole-recording embeddings and phone diagnostics derived from required reference text provide the scoring inputs. FFmpeg [22] decodes audio-bearing containers; a common preprocessing path averages channels to mono and resamples to 16 kHz. The full recording is processed as one utterance, without averaging separately scored sentences. Each encoder adapter returns frame embeddings and valid-frame lengths or masks, with XLS-R 300M, XEUS, and PhoneticXeus each providing 1,024-dimensional frame representations in the configurations used here [18, 19, 20]. Statistics pooling [23], using masked temporal means and population standard deviations, with 10^{-5} added to the variance before the square root, yields 2,048-dimensional utterance vectors. Model revisions and implementation checksums are pinned in the repository. A pinned CMU pronunciation dictionary supplies ARPAbet reference pronunciations [24]. ARPAbet–IPA correspondences are documented in CMU dictionary tools [25]; the versioned mapping in *oollim* converts dictionary phones to the PhoneticXeus token inventory, including explicit diphthong splitting and unstressed-vowel handling. PhoneticXeus provides frame-level phone probabilities and a greedy connectionist temporal classification (CTC) sequence [20, 26]. Seventeen diagnostics summarize phone probabilities and entropy [27], blank behavior, predicted and expected phone rates, reference coverage, relative sequence length, edit operations [28], and CTC negative log-likelihoods normalized by frame and reference-token counts [26]. The CTC calculation sums valid paths for the expected phone sequence rather than assigning a single forced alignment [26]. These features combine pronunciation, timing, and uncertainty evidence; they are not error-free phonetic diagnoses. Out-of-vocabulary references are rejected explicitly. The evaluated accuracy configuration uses XEUS + PhoneticXeus + IPA diagnostics, whereas fluency and prosody use XLS-R + XEUS. The scoring pipeline can select a representation combination and training family separately for each construct.

3.2 Training corpora and label mappings

Three research corpora supply assessment labels mapped to a common 0–10 range. SpeechOcean762 supplies utterance accuracy, fluency, and prosodic ratings on the 0–10 scale. The official SpeechOcean762 training partition is divided by speaker into 2,000 training and 500 development utterances; the official test partition is excluded. The Korean Learners' L2 Speech Corpus (English) [29] is distributed as AIHub dataset 71463 [30]. Eligible Training PRN recordings yield 2,699 training and 636 development examples. The corpus rubric rates pronunciation accuracy (segment clarity and errors) and prosodic fluency (stress, rhythm, intonation, rate, and pauses) on 1–5 scales [29]. The supplied articulation score trains accuracy, while the prosodic-fluency score trains both fluency and prosody after mapping with $2.5(y - 1)$. J-AESOP [31] supplies 180 Japanese L2 speakers reading the North Wind and Sun passage (task 6_01). Deterministic speaker partitions allocate 108 recordings to training, 36 to development, and 36 to a reserved partition; native controls are excluded. Ratings from 16 raters across three sections are averaged per construct. Segment ratings supply accuracy, alongside fluency and prosody; $10(y - 1)/9$ maps the 1–10 source scale to 0–10. Partitions are speaker-disjoint within each corpus to separate development from fitting [32]. In the Korean Learners' L2 Speech Corpus (English), recorder identities linked by identical decoded audio are

grouped before splitting; duration and dictionary-coverage checks determine eligibility. Exact sampling rules, audio audits, mappings, and development results are retained in the reproduction package. Equal numerical ranges do not establish score equivalence [33]; the shared prosody/fluency target does not provide an independent fluency annotation. The mappings permit joint fitting while retaining these differences in supervision.

3.3 Ridge fitting and family selection

Seven training families compare single-corpus and mixed-corpus supervision. The single-corpus families are `speechocean`, `korean`, and `aesop`; `korean` abbreviates the Korean Learners' L2 Speech Corpus (English) in model identifiers. The pairs are `koreanocean`, `aesopocean`, and `koreanaesop`, while the three-corpus family is `aesopkoreanocean`. In the matched family procedure, a recording from corpus c receives weight $w_i = 1/(Kn_c)$, where K is the number of included corpora and n_c their respective training counts. Training-only weighted means and standard deviations standardize features [34]; a block of width d_b receives gain $\sqrt{2048/d_b}$. Coordinates with standard deviation below 10^{-8} use scale 1. Each construct fits a weighted ridge objective [21, 34]

$$\min_{\beta,b} \quad \sum_i w_i(y_i - z_i^\top \beta - b)^2 + \alpha\|\beta\|_2^2, \tag{1}$$

with an unpenalized intercept. Penalties $\{0.001, 0.01, 0.1, 1, 10, 100\}$ are compared using mean development Spearman correlation over the included corpora; ties prefer lower mean MAE and then larger regularization, keeping selection on development data [32]. Multi-output array storage does not couple the separate squared-loss objectives. All seven ridge families were scored for the present study. The deployed SpeechOcean ridge used an earlier tuning procedure, including development accuracy MAE, and its scores were carried forward from the first evaluation round rather than replaced with matched-refit predictions. The retained deployment is version `speechocean-stella-v2`; the other six families use the matched `seven-families-v1` heads. Request metadata are retained where available; the original SpeechOcean scores were imported rather than regenerated. Public-development selection retained koreanocean for accuracy and aesopkoreanocean for fluency and prosody; these ranking-oriented research selections are distinct from the column maxima in the external table and are not claimed to maximize external QWK. The representation sensitivity analysis is separate from the fixed-combination external evaluation. A sensitivity check of all 315 family–construct–representation configurations on public development data retained the overall construct-specific selections under the 0.01 Spearman improvement threshold and corpus-level degradation and uncertainty checks recorded before fitting in the October 18, 2026 ablation protocol; the full protocol and results are provided in the reproduction materials. All coefficient estimation, preprocessing fitting, and score calibration exclude private evaluation ratings.

Table 1: Agreement with human consensus ($N = 305$).

	Accuracy		Fluency		Prosody	
Method	ρ	QWK	ρ	QWK	ρ	QWK
Human consensus	**0.809**	**0.704**	**0.807**	**0.739**	**0.807**	**0.742**
Azure	0.761	0.161	0.633	0.078	0.737	0.370
speechocean	0.844	0.546	0.842	**0.671**	0.847	**0.570**
korean	0.856	0.702	0.860	0.495	0.874	0.489
aesop	0.852	0.241	0.850	0.276	0.850	0.261
koreanocean	0.853	**0.821**	0.840	0.585	0.850	0.546
aesopocean	**0.872**	0.598	0.874	0.613	0.881	0.541
koreanaesop	0.867	0.477	**0.876**	0.462	**0.883**	0.431
aesopkoreanocean	0.870	0.585	0.874	0.527	0.881	0.480

Note. Agreement is averaged over two-rater references (other raters for humans). One negative AESOP accuracy score was capped at zero for QWK only. Bold marks the human baseline and ridge-family values within .001 of each column maximum (reported precision). Full comparisons and 95% CIs are in the appendices.

3.4 Local execution and capacity

Local operation was verified on consumer hardware using CPU inference. On the Windows gaming PC (Intel Core i7-11700K, 32 GB physical RAM, Ubuntu under WSL2 with approximately 24 GiB available RAM and 8 GiB

swap), an archived 120-second capacity test completed in approximately 169 seconds, reaching 22.52 GiB peak process resident memory and 5.51 GiB peak system-wide swap use. The installed RTX 3080 was not required for that inference test. A separate Apple Silicon installation test processed an 8.21-second public recording in 42.13 seconds with four CPU threads, including decoding, all three encoders, phone features, and scoring. Software versions for the same WSL installation are recorded in the appendices. These single-run checks establish operation rather than representative throughput; the released duration limit is 120 seconds, and long recordings can require substantial RAM and swap.

4 External evaluation

4.1 Independent recordings and ratings

The collection comprised university EFL learners from China ($n = 98$), Korea ($n = 101$), and Japan ($n = 107$), who recorded the same "Please Call Stella" passage on personal smartphones. One recording was excluded because of its length, leaving 305 samples. The three university professors had approximately 20–30 years of teaching experience, including pronunciation instruction, Cambridge English speaking assessment, and high-stakes speech rating. Based in those countries, the raters calibrated their judgments on separate recordings spanning proficiency levels, discussing discrepancies to establish shared criteria for accuracy, fluency, and prosody. Each professor then rated every recording independently in randomized order, blinded to participant identity and automated scores, using a 0–100 scale in five-point increments. The private evaluation recordings and ratings were held out from model training, hyperparameter tuning, and score calibration. The original human ratings and Azure API scores are archival; the seven ridge families were evaluated for the present study against those existing ratings. Azure scores were obtained in January 2026 using the `en-US` locale, with the Stella reference text supplied and each complete recording assessed as a whole. The returned accuracy, fluency, and prosody scores were retained. The service revision and remaining request options were not retained in the supplied evaluation archive; the comparison therefore concerns those recorded API outputs rather than a newly standardized Azure benchmark.

4.2 Metrics and human baselines

Evaluation distinguishes rank consistency from numerical agreement [35, 12]. Spearman's ρ measures rank association [36]. Quadratic weighted kappa (QWK) measures chance-corrected agreement with squared category-distance penalties [37]; this study uses fixed integer categories 0–100 and half-up rounding. Model outputs are multiplied by ten; no calibration is fitted. Predictions are clipped to 0–100 for QWK, including bootstrap replicates, whereas Spearman retains raw predictions. A reference mean requires all three human scores; partial means are not substituted. Human–Human pairwise agreement averages the three pairwise statistics, whereas Human–Human consensus agreement averages each rater's agreement with the mean of the other two, never including that rater in their own reference. For each system, the three-rater row compares the system's scores with the mean of all humans, while the two-rater row averages the three separate agreement statistics against the three possible two-human means. Statistics are averaged, not observations or reference means pooled. The two-rater system row and Human–Human consensus row therefore provide the appropriate reference-size comparison when recomputed on the same recordings. Percentile 95% intervals resample speaker clusters jointly, recomputing all component statistics within each replicate [38, 39]. The supplementary marginal intervals retain the supplied evaluation export (2,000 replicates, seed 42). Additional paired comparisons use 200,000 bootstrap replicates (seed 42), subtracting statistics within each replicate on the same sampled recordings. An independent 200,000-replicate run (seed 2026) produced the same interval-exclusion decisions for all contrasts. Each unique recording URL is treated as one speaker. Both percentile 95% intervals and Bonferroni-adjusted intervals across all 90 reported differences are supplied: eight systems versus human consensus and seven ridges versus Azure, for three constructs and two metrics. These post-evaluation comparisons condition on the observed raters and fitted models; adjustment covers the reported contrasts, not the preceding model search. The appendices report all reference comparisons and intervals for the same 305 recordings. Matched reference sizes and cohorts make the human and model estimates comparable.

5 Results

Table 1 shows that every ridge family exceeds Azure in rank association across all three constructs, with mixed-corpus families generally strongest. Model correlations also exceed the human consensus baseline, but numerical agreement presents a different pattern: the leading accuracy model exceeds the human QWK baseline,

whereas fluency and prosody remain below the human QWK baseline. Several families outperform Azure on both metrics, although the advantage is not universal for QWK. No family leads both metrics across all constructs, supporting construct-specific scoring choices. Across all 90 contrasts, Bonferroni-adjusted paired intervals remain above zero for every ridge–Azure Spearman difference and for koreanocean accuracy QWK versus human consensus. `aesopocean`, `koreanaesop`, and `aesopkoreanocean` also exceed human-consensus rank agreement across all three constructs under that adjustment; `korean` does so for prosody. No ridge has a positive adjusted QWK difference from human consensus for fluency or prosody. Full differences and score-distribution summaries appear in the appendices; external family rankings remain descriptive.

6 Discussion

The results favor selecting scorers by construct and intended use. Public-development selection prioritizes rank correlation; users seeking numerical feedback should also consider QWK, for which the external table favors koreanocean accuracy and speechocean fluency/prosody. Those external preferences remain descriptive rather than independently validated deployment choices. Several mixed-corpus families have higher correlations but lower QWK values than single-corpus families, so scorer selection depends on whether the intended use requires ranking, numerical agreement, or both; no single family needs to generate all three constructs. Azure scores average 23.63, 27.88, and 10.57 points above the three-human means for accuracy, fluency, and prosody, respectively; these offsets help explain low QWK despite positive rank associations. No correction was fitted to remove those offsets. XEUS alone approaches the retained combinations on public-development macro correlation (gaps of 0.0017 for accuracy, 0.0016 for fluency, and 0.0041 for prosody in the selected families), so the sweep does not establish a substantial benefit from every additional encoder. A lighter XEUS-only system warrants separate external evaluation; the present private results apply to the evaluated combinations. Korean-only fluency and prosody share features and targets, so their predictions do not constitute independent learning of two constructs. Human baselines also answer different questions, with agreement against a two-rater consensus generally higher than agreement with one other human because averaging reduces some rater-specific variation. On the same reference pairs and recordings, leading model point estimates exceed the Human–Human consensus baseline for rank association across constructs and for accuracy QWK, whereas fluency and prosody QWK remain below that baseline. The advantage over human agreement therefore depends on the construct and metric, with paired uncertainty reported in the appendices. External family rankings are descriptive rather than prospectively validated selections [32]. Agreement with human ratings supports a particular score interpretation, rather than every educational use of the scores [40]. The software, trained scoring weights, model configurations, and reproduction materials are available at `https://github.com/cburst/oollim`, subject to the licenses documented in the repository. Practical scorer selection should therefore consider the target construct, agreement metric, and teaching context.

7 Limitations

The evaluation supports conclusions about one passage and three university EFL settings. Recording-level scores permit paired analysis, but URL-based grouping assumes one recording per speaker. The cohort excludes one long recording, and the 120-second operational limit constrains input coverage. Score-scale mappings do not establish equivalent constructs: the Korean Learners' L2 Speech Corpus (English) supplies shared prosody/fluency supervision, so differences among training families cannot be attributed solely to additional data. Speaker-disjoint fitting also cannot resolve possible upstream exposure of PhoneticXeus to learner corpora. The Azure service revision and remaining request options are unavailable, limiting exact replication of that comparator. The public-development ablation used a recorded selection rule, but neither bootstrap intervals nor contrast adjustment account for the entire model-search process. Broader claims require additional passages, populations, and intended uses [40].

8 Conclusion

XLS-R, XEUS, and PhoneticXeus, combined with construct-specific ridge scoring, support local evaluation of L2 English accuracy, fluency, and prosody on the tested consumer configuration. The external passage evaluation shows strong rank association with human ratings, with every ridge family exceeding correlations from the archived January 2026 Azure API outputs across all three constructs. Numerical agreement varies by construct and training family and identifies different leading scorers from rank association, supporting the reporting of both metrics for each construct. The additional public-development ablation retained the overall construct-specific configurations under

the 0.01 improvement rule. Additional passages and populations are needed to establish the scope of comparative human-agreement claims.

9 Generative AI Use Disclosure

OpenAI Codex assisted with software development, experiment execution, and manuscript drafting and revision; the author is responsible for the submitted content and its verification.

References

[1] D. Deterding. Segmentals. In M. Reed and J. M. Levis, eds., *The Handbook of English Pronunciation*, pp. 67–84. Wiley, 2015.

[2] D. Cai, B. Naismith, M. Kostromitina, Z. Teng, K. P. Yancey, and G. T. LaFlair. Developing an Automatic Pronunciation Scorer: Aligning Speech Evaluation Models and Applied Linguistics Constructs. *Language Learning*, 75(S1):170–203, 2025.

[3] J. Kormos and M. Dénes. Exploring measures and perceptions of fluency in the speech of second language learners. *System*, 32(2):145–164, 2004.

[4] H. R. Bosker, A.-F. Pinget, H. Quené, T. Sanders, and N. H. de Jong. What makes speech sound fluent? The contributions of pauses, speed and repairs. *Language Testing*, 30(2):159–175, 2013.

[5] J. Cole. Prosody in context: A review. *Language, Cognition and Neuroscience*, 30(1–2):1–31, 2015.

[6] A. Cutler. Lexical Stress in English Pronunciation. In M. Reed and J. M. Levis, eds., *The Handbook of English Pronunciation*, pp. 106–124. Wiley, 2015.

[7] T. Isaacs and R. I. Thomson. Rater Experience, Rating Scale Length, and Judgments of L2 Pronunciation: Revisiting Research Conventions. *Language Assessment Quarterly*, 10(2):135–159, 2013.

[8] M. J. Munro and T. M. Derwing. Foreign Accent, Comprehensibility, and Intelligibility in the Speech of Second Language Learners. *Language Learning*, 45(1):73–97, 1995.

[9] N. Segalowitz. *Cognitive Bases of Second Language Fluency*. Routledge, 2010.

[10] L. D. Hahn. Primary Stress and Intelligibility: Research to Motivate the Teaching of Suprasegmentals. *TESOL Quarterly*, 38(2):201–223, 2004.

[11] Microsoft. Use pronunciation assessment. *Microsoft Learn*. `https://learn.microsoft.com/azure/ai-services/speech-service/how-to-pronunciation-assessment`. Accessed October 18, 2026.

[12] D. M. Williamson, X. Xi, and F. J. Breyer. A Framework for Evaluation and Use of Automated Scoring. *Educational Measurement: Issues and Practice*, 31(1):2–13, 2012.

[13] L. Chen et al. Automated Scoring of Nonnative Speech Using the SpeechRater v. 5.0 Engine. *ETS Research Report Series*, 2018(1):1–31, 2018.

[14] Y. Gong, Z. Chen, I.-H. Chu, P. Chang, and J. Glass. Transformer-Based Multi-Aspect Multi-Granularity Non-Native English Speaker Pronunciation Assessment. *ICASSP*, 2022. arXiv:2205.03432.

[15] J. Zhang et al. speechocean762: An Open-Source Non-Native English Speech Corpus for Pronunciation Assessment. Interspeech, 2021. `https://www.openslr.org/101/`.

[16] E. Kim, J.-J. Jeon, H. Seo, and H. Kim. Automatic Pronunciation Assessment using Self-Supervised Speech Representation Learning. In *Proc. Interspeech*, pp. 1411–1415, 2022.

[17] K. Saito, K. Macmillan, M. Kachlicka, T. Kunihara, and N. Minematsu. Automated assessment of second language comprehensibility: Review, training, validation, and generalization studies. *Studies in Second Language Acquisition*, 45(1):234–263, 2023.

[18] A. Babu et al. XLS-R: Self-supervised Cross-lingual Speech Representation Learning at Scale. In *Proc. Interspeech*, pp. 2278–2282, 2022.

[19] W. Chen et al. Towards Robust Speech Representation Learning for Thousands of Languages. 2024. https://arxiv.org/abs/2407.00837.

[20] S. Bharadwaj et al. An Empirical Recipe for Universal Phone Recognition. 2026. https://arxiv.org/abs/2603.29042.

[21] A. E. Hoerl and R. W. Kennard. Ridge Regression: Biased Estimation for Nonorthogonal Problems. *Technometrics*, 12(1):55–67, 1970.

[22] FFmpeg Developers. *FFmpeg Documentation*. https://ffmpeg.org/documentation.html. Accessed October 18, 2026.

[23] D. Snyder, D. Garcia-Romero, D. Povey, and S. Khudanpur. Deep Neural Network Embeddings for Text-Independent Speaker Verification. In *Proc. Interspeech*, pp. 999–1003, 2017.

[24] CMU Sphinx. *CMU Pronouncing Dictionary*. https://github.com/cmusphinx/cmudict. Accessed October 18, 2026.

[25] CMU Sphinx. *CMU Pronunciation Dictionary Tools: Phone Tables and Phoneme Sets*. https://github.com/cmusphinx/cmudict-tools. Accessed October 18, 2026.

[26] A. Graves, S. Fernández, F. Gomez, and J. Schmidhuber. Connectionist Temporal Classification: Labelling Unsegmented Sequence Data with Recurrent Neural Networks. In *Proc. ICML*, pp. 369–376, 2006.

[27] C. E. Shannon. A Mathematical Theory of Communication. *The Bell System Technical Journal*, 27(3):379–423, 1948.

[28] R. A. Wagner and M. J. Fischer. The String-to-String Correction Problem. *Journal of the ACM*, 21(1):168–173, 1974.

[29] S. Han, M. Chung, and S. Kim. Constructing Korean Learners' L2 Speech Corpus of Seven Languages for Automatic Pronunciation Assessment. In *Proc. LREC-COLING*, pp. 3772–3781, 2024.

[30] National Information Society Agency. Korean Learners' L2 Speech Corpus (English). AIHub dataset 71463. https://aihub.or.kr/aihubdata/data/view.do?dataSetSn=71463.

[31] K. Yazawa, T. Konishi, and M. Kondo. The J-AESOP corpus: Design, application, and future directions of a Japanese-English bilingual speech corpus. Acoustical Science and Technology, 46(5):553–563, 2025. https://doi.org/10.1250/ast.e25.06.

[32] G. C. Cawley and N. L. C. Talbot. On Over-fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation. *Journal of Machine Learning Research*, 11:2079–2107, 2010.

[33] N. J. Dorans. Scores, Scales, and Score Linking. In P. Irwing, T. Booth, and D. J. Hughes, eds., *The Wiley Handbook of Psychometric Testing*. Wiley, 2018.

[34] T. Hastie, R. Tibshirani, and J. Friedman. *The Elements of Statistical Learning*, 2nd ed. Springer, 2009, secs. 3.4 and 7.

[35] S. E. Stemler. A Comparison of Consensus, Consistency, and Measurement Approaches to Estimating Interrater Reliability. *Practical Assessment, Research & Evaluation*, 9(4), 2004.

[36] C. Spearman. The Proof and Measurement of Association between Two Things. *The American Journal of Psychology*, 15(1):72–101, 1904.

[37] J. Cohen. Weighted kappa: Nominal scale agreement with provision for scaled disagreement or partial credit. *Psychological Bulletin*, 70(4):213–220, 1968.

[38] B. Efron. Bootstrap Methods: Another Look at the Jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

[39] C. A. Field and A. H. Welsh. Bootstrapping clustered data. *Journal of the Royal Statistical Society: Series B*, 69(3):369–390, 2007.

[40] M. T. Kane. Validating the Interpretations and Uses of Test Scores. *Journal of Educational Measurement*, 50(1):1–73, 2013.

A Full external agreement comparisons

Table 2: Supplementary accuracy comparisons on the matched cohort.

Method	Agreement comparison	N	ρ	95% CI	QWK	95% CI
Human raters	**Human–Human pairwise agreement**	**305**	**0.759**	**[0.710, 0.799]**	**0.642**	**[0.597, 0.685]**
Human raters	**Human–Human consensus agreement (other two raters)**	**305**	**0.809**	**[0.769, 0.842]**	**0.704**	**[0.663, 0.743]**
Azure	Model–Human consensus agreement (two raters)	305	0.761	[0.709, 0.804]	0.161	[0.129, 0.196]
Azure	Model–Human consensus agreement (three raters)	305	0.776	[0.723, 0.820]	0.158	[0.126, 0.193]
`speechocean`	Model–Human consensus agreement (two raters)	305	0.844	[0.807, 0.872]	0.546	[0.494, 0.592]
`speechocean`	Model–Human consensus agreement (three raters)	305	0.862	[0.825, 0.890]	0.555	[0.500, 0.605]
`korean`	Model–Human consensus agreement (two raters)	305	0.856	[0.822, 0.882]	0.702	[0.662, 0.739]
`korean`	Model–Human consensus agreement (three raters)	305	0.872	[0.839, 0.898]	0.726	[0.686, 0.763]
`aesop`	Model–Human consensus agreement (two raters)	305	0.852	[0.818, 0.878]	0.241	[0.215, 0.265]
`aesop`	Model–Human consensus agreement (three raters)	305	0.869	[0.836, 0.894]	0.244	[0.217, 0.269]
`koreanocean`	Model–Human consensus agreement (two raters)	305	0.853	[0.816, 0.882]	0.821	[0.789, 0.847]
`koreanocean`	Model–Human consensus agreement (three raters)	305	0.870	[0.833, 0.898]	0.860	[0.832, 0.884]
`aesopocean`	Model–Human consensus agreement (two raters)	305	0.872	[0.842, 0.894]	0.598	[0.563, 0.631]
`aesopocean`	Model–Human consensus agreement (three raters)	305	0.889	[0.859, 0.911]	0.615	[0.579, 0.649]
`koreanaesop`	Model–Human consensus agreement (two raters)	305	0.867	[0.835, 0.891]	0.477	[0.439, 0.512]
`koreanaesop`	Model–Human consensus agreement (three raters)	305	0.884	[0.853, 0.908]	0.486	[0.448, 0.522]
`aesopkoreanocean`	Model–Human consensus agreement (two raters)	305	0.870	[0.839, 0.894]	0.585	[0.548, 0.620]
`aesopkoreanocean`	Model–Human consensus agreement (three raters)	305	0.888	[0.857, 0.911]	0.601	[0.563, 0.636]

Note. $N = 305$ for every row. Human–Human pairwise agreement averages the three rater-pair statistics. Human–Human consensus agreement compares each rater with the other two raters' mean and averages the three statistics. Model–Human consensus agreement (two raters) averages the three statistics against each possible two-rater mean; the three-rater comparison uses the mean of all three humans. Two-rater model rows precede three-rater rows to align with the human consensus baseline. Intervals are percentile speaker-cluster bootstrap intervals for each statistic, not paired differences. QWK uses fixed 0–100 categories and half-up rounding without fitted calibration. Bold identifies human baselines only; model rankings are discussed in the text. One AESOP accuracy prediction was capped at zero for both QWK comparisons and their bootstrap intervals; Spearman retained it unchanged.

Table 3: Supplementary fluency comparisons on the matched cohort.

Method	Agreement comparison	N	ρ	95% CI	QWK	95% CI
Human raters	**Human–Human pairwise agreement**	**305**	**0.747**	**[0.697, 0.791]**	**0.680**	**[0.631, 0.724]**
Human raters	**Human–Human consensus agreement (other two raters)**	**305**	**0.807**	**[0.768, 0.840]**	**0.739**	**[0.695, 0.777]**
Azure	Model–Human consensus agreement (two raters)	305	0.633	[0.556, 0.698]	0.078	[0.056, 0.107]
Azure	Model–Human consensus agreement (three raters)	305	0.649	[0.571, 0.715]	0.077	[0.055, 0.105]
`speechocean`	Model–Human consensus agreement (two raters)	305	0.842	[0.801, 0.874]	0.671	[0.627, 0.707]
`speechocean`	Model–Human consensus agreement (three raters)	305	0.859	[0.819, 0.891]	0.704	[0.660, 0.740]
`korean`	Model–Human consensus agreement (two raters)	305	0.860	[0.825, 0.888]	0.495	[0.452, 0.535]
`korean`	Model–Human consensus agreement (three raters)	305	0.879	[0.844, 0.905]	0.510	[0.466, 0.550]
`aesop`	Model–Human consensus agreement (two raters)	305	0.850	[0.813, 0.879]	0.276	[0.247, 0.305]
`aesop`	Model–Human consensus agreement (three raters)	305	0.868	[0.833, 0.896]	0.283	[0.253, 0.312]
`koreanocean`	Model–Human consensus agreement (two raters)	305	0.840	[0.799, 0.873]	0.585	[0.537, 0.628]
`koreanocean`	Model–Human consensus agreement (three raters)	305	0.857	[0.817, 0.890]	0.606	[0.557, 0.650]
`aesopocean`	Model–Human consensus agreement (two raters)	305	0.874	[0.842, 0.899]	0.613	[0.575, 0.647]
`aesopocean`	Model–Human consensus agreement (three raters)	305	0.893	[0.863, 0.917]	0.633	[0.595, 0.667]
`koreanaesop`	Model–Human consensus agreement (two raters)	305	0.876	[0.846, 0.900]	0.462	[0.424, 0.496]
`koreanaesop`	Model–Human consensus agreement (three raters)	305	0.895	[0.866, 0.917]	0.474	[0.436, 0.509]
`aesopkoreanocean`	Model–Human consensus agreement (two raters)	305	0.874	[0.842, 0.898]	0.527	[0.486, 0.564]
`aesopkoreanocean`	Model–Human consensus agreement (three raters)	305	0.892	[0.861, 0.916]	0.543	[0.501, 0.580]

Note. $N = 305$ for every row. Human–Human pairwise agreement averages the three rater-pair statistics. Human–Human consensus agreement compares each rater with the other two raters' mean and averages the three statistics. Model–Human consensus agreement (two raters) averages the three statistics against each possible two-rater mean; the three-rater comparison uses the mean of all three humans. Two-rater model rows precede three-rater rows to align with the human consensus baseline. Intervals are percentile speaker-cluster bootstrap intervals for each statistic, not paired differences. QWK uses fixed 0–100 categories and half-up rounding without fitted calibration. Bold identifies human baselines only; model rankings are discussed in the text. No prediction required QWK clipping for this construct.

Table 4: Supplementary prosody comparisons on the matched cohort.

Method	Agreement comparison	N	ρ	95% CI	QWK	95% CI
Human raters	**Human–Human pairwise agreement**	**305**	**0.751**	**[0.702, 0.795]**	**0.681**	**[0.634, 0.722]**
Human raters	**Human–Human consensus agreement (other two raters)**	**305**	**0.807**	**[0.768, 0.842]**	**0.742**	**[0.701, 0.778]**
Azure	Model–Human consensus agreement (two raters)	305	0.737	[0.672, 0.792]	0.370	[0.318, 0.421]
Azure	Model–Human consensus agreement (three raters)	305	0.750	[0.684, 0.806]	0.369	[0.318, 0.422]
`speechocean`	Model–Human consensus agreement (two raters)	305	0.847	[0.810, 0.878]	0.570	[0.525, 0.610]
`speechocean`	Model–Human consensus agreement (three raters)	305	0.865	[0.829, 0.895]	0.593	[0.547, 0.634]
`korean`	Model–Human consensus agreement (two raters)	305	0.874	[0.842, 0.899]	0.489	[0.447, 0.526]
`korean`	Model–Human consensus agreement (three raters)	305	0.893	[0.862, 0.917]	0.502	[0.460, 0.540]
`aesop`	Model–Human consensus agreement (two raters)	305	0.850	[0.814, 0.878]	0.261	[0.232, 0.288]
`aesop`	Model–Human consensus agreement (three raters)	305	0.867	[0.833, 0.895]	0.266	[0.237, 0.295]
`koreanocean`	Model–Human consensus agreement (two raters)	305	0.850	[0.812, 0.881]	0.546	[0.501, 0.586]
`koreanocean`	Model–Human consensus agreement (three raters)	305	0.869	[0.832, 0.900]	0.563	[0.517, 0.604]
`aesopocean`	Model–Human consensus agreement (two raters)	305	0.881	[0.850, 0.905]	0.541	[0.502, 0.576]
`aesopocean`	Model–Human consensus agreement (three raters)	305	0.899	[0.870, 0.922]	0.557	[0.517, 0.592]
`koreanaesop`	Model–Human consensus agreement (two raters)	305	0.883	[0.852, 0.906]	0.431	[0.394, 0.465]
`koreanaesop`	Model–Human consensus agreement (three raters)	305	0.902	[0.873, 0.924]	0.442	[0.404, 0.477]
`aesopkoreanocean`	Model–Human consensus agreement (two raters)	305	0.881	[0.850, 0.906]	0.480	[0.440, 0.516]
`aesopkoreanocean`	Model–Human consensus agreement (three raters)	305	0.900	[0.870, 0.924]	0.493	[0.452, 0.530]

Note. $N = 305$ for every row. Human–Human pairwise agreement averages the three rater-pair statistics. Human–Human consensus agreement compares each rater with the other two raters' mean and averages the three statistics. Model–Human consensus agreement (two raters) averages the three statistics against each possible two-rater mean; the three-rater comparison uses the mean of all three humans. Two-rater model rows precede three-rater rows to align with the human consensus baseline. Intervals are percentile speaker-cluster bootstrap intervals for each statistic, not paired differences. QWK uses fixed 0–100 categories and half-up rounding without fitted calibration. Bold identifies human baselines only; model rankings are discussed in the text. No prediction required QWK clipping for this construct.

B Paired differences and score distributions

The same 305 complete recordings were resampled jointly 200,000 times (seed 42), treating each distinct recording URL as one speaker. Human consensus agreement averages each rater's agreement with the other two; system agreement averages comparisons against those same three reference pairs. Differences are recomputed within each replicate. Intervals are percentile approximations; adjusted endpoints use probabilities $0.025/90$ and $1 - 0.025/90$, covering 90 reported contrasts. A second 200,000-replicate run (seed 2026) reproduced all 90 interval-exclusion decisions; adjusted endpoints differed by at most 0.00367. Each adjusted tail now contains approximately 56 draws, although percentile intervals remain approximations. The supplementary marginal intervals above retain the original 2,000-replicate export. These post-evaluation comparisons condition on the panel and fitted models, without correcting for prior model exploration. All 48 system point estimates match the supplied export to numerical precision. Private scores were used only for evaluation; no coefficients or calibration were fitted.

B.1 Accuracy differences

System	Reference	Metric	Difference	95% CI	Adjusted CI
azure	Human	ρ	-0.049	[-0.096, -0.005]	[-0.137, 0.027]
azure	Human	QWK	-0.543	[-0.580, -0.504]	[-0.606, -0.471]
speechocean	Human	ρ	+0.035	[0.004, 0.065]	[-0.021, 0.091]
speechocean	Human	QWK	-0.158	[-0.206, -0.110]	[-0.243, -0.073]
speechocean	Azure	ρ	+0.083	[0.041, 0.128]	[0.009, 0.167]
speechocean	Azure	QWK	+0.385	[0.341, 0.426]	[0.306, 0.455]
korean	Human	ρ	+0.046	[0.018, 0.075]	[-0.004, 0.101]
korean	Human	QWK	-0.002	[-0.041, 0.037]	[-0.071, 0.066]
korean	Azure	ρ	+0.095	[0.055, 0.139]	[0.025, 0.177]
korean	Azure	QWK	+0.541	[0.500, 0.580]	[0.470, 0.609]
aesop	Human	ρ	+0.043	[0.012, 0.074]	[-0.013, 0.101]
aesop	Human	QWK	-0.463	[-0.493, -0.431]	[-0.515, -0.405]
aesop	Azure	ρ	+0.092	[0.054, 0.132]	[0.028, 0.169]
aesop	Azure	QWK	+0.080	[0.053, 0.105]	[0.030, 0.122]
koreanocean	Human	ρ	+0.044	[0.013, 0.074]	[-0.012, 0.099]
koreanocean	Human	QWK	+0.116	[0.082, 0.152]	[0.055, 0.181]
koreanocean	Azure	ρ	+0.093	[0.051, 0.137]	[0.020, 0.176]
koreanocean	Azure	QWK	+0.660	[0.626, 0.690]	[0.600, 0.712]
aesopocean	Human	ρ	+0.062	[0.033, 0.093]	[0.010, 0.119]
aesopocean	Human	QWK	-0.106	[-0.140, -0.071]	[-0.165, -0.045]
aesopocean	Azure	ρ	+0.111	[0.071, 0.155]	[0.043, 0.195]
aesopocean	Azure	QWK	+0.437	[0.402, 0.470]	[0.375, 0.495]
koreanaesop	Human	ρ	+0.057	[0.028, 0.088]	[0.005, 0.114]
koreanaesop	Human	QWK	-0.227	[-0.261, -0.193]	[-0.287, -0.167]
koreanaesop	Azure	ρ	+0.106	[0.066, 0.150]	[0.037, 0.191]
koreanaesop	Azure	QWK	+0.316	[0.281, 0.350]	[0.254, 0.376]
aesopkoreanocean	Human	ρ	+0.061	[0.031, 0.091]	[0.009, 0.118]
aesopkoreanocean	Human	QWK	-0.119	[-0.153, -0.084]	[-0.179, -0.058]
aesopkoreanocean	Azure	ρ	+0.110	[0.070, 0.154]	[0.041, 0.194]
aesopkoreanocean	Azure	QWK	+0.425	[0.389, 0.459]	[0.360, 0.486]

B.2 Fluency differences

System	Reference	Metric	Difference	95% CI	Adjusted CI
azure	Human	ρ	-0.174	[-0.244, -0.110]	[-0.304, -0.067]
azure	Human	QWK	-0.661	[-0.697, -0.621]	[-0.723, -0.588]
speechocean	Human	ρ	+0.035	[0.000, 0.069]	[-0.029, 0.095]

System	Reference	Metric	Difference	95% CI	Adjusted CI
speechocean	Human	QWK	-0.068	[-0.113, -0.024]	[-0.145, 0.010]
speechocean	Azure	ρ	+0.209	[0.145, 0.279]	[0.099, 0.338]
speechocean	Azure	QWK	+0.593	[0.554, 0.627]	[0.523, 0.653]
korean	Human	ρ	+0.054	[0.022, 0.085]	[-0.002, 0.111]
korean	Human	QWK	-0.244	[-0.281, -0.208]	[-0.308, -0.179]
korean	Azure	ρ	+0.228	[0.164, 0.297]	[0.121, 0.355]
korean	Azure	QWK	+0.417	[0.379, 0.452]	[0.351, 0.478]
aesop	Human	ρ	+0.044	[0.009, 0.077]	[-0.017, 0.104]
aesop	Human	QWK	-0.463	[-0.494, -0.430]	[-0.517, -0.403]
aesop	Azure	ρ	+0.218	[0.157, 0.284]	[0.116, 0.342]
aesop	Azure	QWK	+0.198	[0.174, 0.221]	[0.155, 0.238]
koreanocean	Human	ρ	+0.033	[-0.004, 0.068]	[-0.034, 0.096]
koreanocean	Human	QWK	-0.154	[-0.195, -0.114]	[-0.227, -0.084]
koreanocean	Azure	ρ	+0.207	[0.140, 0.279]	[0.094, 0.339]
koreanocean	Azure	QWK	+0.507	[0.464, 0.546]	[0.428, 0.575]
aesopocean	Human	ρ	+0.068	[0.038, 0.098]	[0.015, 0.123]
aesopocean	Human	QWK	-0.126	[-0.159, -0.093]	[-0.184, -0.066]
aesopocean	Azure	ρ	+0.242	[0.180, 0.310]	[0.139, 0.369]
aesopocean	Azure	QWK	+0.535	[0.501, 0.566]	[0.473, 0.590]
koreanaesop	Human	ρ	+0.069	[0.040, 0.099]	[0.019, 0.124]
koreanaesop	Human	QWK	-0.277	[-0.311, -0.243]	[-0.335, -0.216]
koreanaesop	Azure	ρ	+0.243	[0.182, 0.311]	[0.142, 0.369]
koreanaesop	Azure	QWK	+0.384	[0.350, 0.416]	[0.325, 0.439]
aesopkoreanocean	Human	ρ	+0.067	[0.037, 0.098]	[0.015, 0.123]
aesopkoreanocean	Human	QWK	-0.212	[-0.247, -0.179]	[-0.273, -0.152]
aesopkoreanocean	Azure	ρ	+0.241	[0.177, 0.311]	[0.136, 0.370]
aesopkoreanocean	Azure	QWK	+0.449	[0.412, 0.482]	[0.382, 0.508]

B.3 Prosody differences

System	Reference	Metric	Difference	95% CI	Adjusted CI
azure	Human	ρ	-0.070	[-0.124, -0.023]	[-0.172, 0.010]
azure	Human	QWK	-0.372	[-0.419, -0.325]	[-0.455, -0.289]
speechocean	Human	ρ	+0.040	[0.009, 0.069]	[-0.017, 0.092]
speechocean	Human	QWK	-0.172	[-0.214, -0.128]	[-0.247, -0.094]
speechocean	Azure	ρ	+0.110	[0.063, 0.163]	[0.032, 0.211]
speechocean	Azure	QWK	+0.201	[0.145, 0.256]	[0.100, 0.298]
korean	Human	ρ	+0.067	[0.039, 0.096]	[0.018, 0.120]
korean	Human	QWK	-0.253	[-0.287, -0.220]	[-0.312, -0.193]
korean	Azure	ρ	+0.137	[0.098, 0.183]	[0.073, 0.225]
korean	Azure	QWK	+0.119	[0.077, 0.160]	[0.043, 0.189]
aesop	Human	ρ	+0.043	[0.012, 0.073]	[-0.012, 0.097]
aesop	Human	QWK	-0.481	[-0.510, -0.451]	[-0.530, -0.427]
aesop	Azure	ρ	+0.113	[0.073, 0.159]	[0.048, 0.200]
aesop	Azure	QWK	-0.109	[-0.147, -0.072]	[-0.176, -0.044]
koreanocean	Human	ρ	+0.042	[0.010, 0.074]	[-0.017, 0.098]
koreanocean	Human	QWK	-0.196	[-0.232, -0.161]	[-0.259, -0.135]
koreanocean	Azure	ρ	+0.113	[0.069, 0.163]	[0.040, 0.208]
koreanocean	Azure	QWK	+0.176	[0.127, 0.223]	[0.088, 0.258]
aesopocean	Human	ρ	+0.073	[0.048, 0.100]	[0.029, 0.123]
aesopocean	Human	QWK	-0.201	[-0.232, -0.170]	[-0.255, -0.146]
aesopocean	Azure	ρ	+0.144	[0.099, 0.195]	[0.072, 0.243]
aesopocean	Azure	QWK	+0.171	[0.127, 0.213]	[0.092, 0.244]

System	Reference	Metric	Difference	95% CI	Adjusted CI
koreanaesop	Human	ρ	+0.076	[0.049, 0.104]	[0.029, 0.128]
koreanaesop	Human	QWK	-0.311	[-0.342, -0.279]	[-0.365, -0.254]
koreanaesop	Azure	ρ	+0.146	[0.105, 0.194]	[0.080, 0.238]
koreanaesop	Azure	QWK	+0.062	[0.022, 0.100]	[-0.009, 0.128]
aesopkoreanocean	Human	ρ	+0.074	[0.047, 0.101]	[0.027, 0.124]
aesopkoreanocean	Human	QWK	-0.262	[-0.293, -0.231]	[-0.315, -0.208]
aesopkoreanocean	Azure	ρ	+0.144	[0.101, 0.194]	[0.074, 0.240]
aesopkoreanocean	Azure	QWK	+0.111	[0.068, 0.151]	[0.034, 0.181]

B.4 Prediction-scale diagnostics

Means, sample standard deviations, mean signed errors (bias), and mean absolute errors (MAE) use unrounded predictions on the 0–100 scale. Errors compare predictions with the three-human mean; these descriptive summaries do not recalibrate scores.

Construct	System	Mean	SD	Bias	MAE
Accuracy	azure	89.29	7.50	+23.63	23.63
Accuracy	speechocean	72.54	6.47	+6.87	8.22
Accuracy	korean	58.57	12.39	-7.10	7.66
Accuracy	aesop	38.49	13.31	-27.17	27.17
Accuracy	koreanocean	64.85	9.77	-0.81	4.38
Accuracy	aesopocean	55.27	11.66	-10.40	10.51
Accuracy	koreanaesop	50.09	14.23	-15.58	15.59
Accuracy	aesopkoreanocean	54.37	12.52	-11.30	11.38
Fluency	azure	94.18	5.87	+27.88	27.88
Fluency	speechocean	63.26	6.34	-3.04	5.36
Fluency	korean	53.31	10.75	-12.99	12.99
Fluency	aesop	44.43	10.97	-21.87	21.87
Fluency	koreanocean	57.68	8.45	-8.62	8.74
Fluency	aesopocean	56.63	10.99	-9.67	9.79
Fluency	koreanaesop	51.09	12.44	-15.21	15.21
Fluency	aesopkoreanocean	54.18	10.50	-12.12	12.13
Prosody	azure	77.13	6.15	+10.57	11.04
Prosody	speechocean	60.56	6.00	-5.99	6.69
Prosody	korean	53.31	10.75	-13.25	13.25
Prosody	aesop	44.14	10.62	-22.42	22.42
Prosody	koreanocean	56.89	8.32	-9.66	9.72
Prosody	aesopocean	55.25	9.89	-11.31	11.33
Prosody	koreanaesop	50.26	12.36	-16.29	16.29
Prosody	aesopkoreanocean	53.05	10.42	-13.50	13.50

C Public-development comparisons

The matched seven-family experiment uses fixed representation combinations and penalty selection over included-corpus development sets. All entries are development Spearman correlations, not final-test results. SpeechOcean rows here are fresh matched refits, not the original deployed SpeechOcean checkpoint in the external comparison. Korean fluency and prosody share a source target. J-AESOP has only 36 development recordings, so large correlations should not be read as precise generalization estimates.

Construct	Family	SpeechOcean (500)	Korean (636)	J-AESOP (36)
Accuracy	speechocean	0.816	0.699	0.967
Accuracy	korean	0.735	0.759	0.969

Construct	Family	SpeechOcean (500)	Korean (636)	J-AESOP (36)
Accuracy	aesop	0.659	0.673	0.987
Accuracy	koreanocean	0.808	0.755	0.964
Accuracy	aesopocean	0.807	0.715	0.970
Accuracy	koreanaesop	0.683	0.755	0.976
Accuracy	aesopkoreanocean	0.804	0.748	0.971
Fluency	speechocean	0.886	0.668	0.856
Fluency	korean	0.837	0.778	0.919
Fluency	aesop	0.682	0.586	0.966
Fluency	koreanocean	0.876	0.770	0.924
Fluency	aesopocean	0.881	0.728	0.952
Fluency	koreanaesop	0.834	0.779	0.959
Fluency	aesopkoreanocean	0.875	0.769	0.953
Prosody	speechocean	0.890	0.653	0.929
Prosody	korean	0.829	0.778	0.938
Prosody	aesop	0.685	0.544	0.975
Prosody	koreanocean	0.876	0.771	0.949
Prosody	aesopocean	0.878	0.747	0.971
Prosody	koreanaesop	0.830	0.777	0.973
Prosody	aesopkoreanocean	0.872	0.770	0.974

D Korean supervision and representation sensitivity

D.1 Korean scoring and target mapping

Han, Chung, and Kim [29] define two 1–5 ratings: pronunciation accuracy and prosodic fluency. Accuracy addresses segmental clarity and errors; prosodic fluency combines stress, rhythm, intonation, speed, and pausing. The stored fields are `articulationScore` and `prosodyScore`. The second field supplies both fluency and prosody targets in the present implementation. Mapping $2.5(y-1)$ produces 0, 2.5, 5, 7.5, and 10 from the five integer levels; mapping does not equate the rubric with SpeechOcean or create an independent fluency annotation. Korean-only heads use the same features and targets for both constructs, yielding identical predictions in the evaluated configuration.

Source	Mapped	Accuracy (paraphrased)	Prosodic fluency (paraphrased)
1	0	Extensive segmental difficulty; understanding is hard.	Strongly disrupted prosody, very slow delivery, excessive pauses.
2	2.5	Frequent errors substantially burden understanding.	Awkward prosody and slow delivery with prolonged, poorly placed pauses.
3	5	Recurring errors interfere with understanding.	Uneven prosody and rate, with frequent or misplaced breaks.
4	7.5	Minor errors with little intelligibility impact.	Mostly steady delivery with mild prosodic or pausing irregularities.
5	10	Clear segments without noticeable errors.	Natural prosody, rate, and phrase boundaries.

D.2 Ablation protocol

The protocol was recorded before fitting on October 18, 2026 (repository: `results/ablation-mac-v1/PROTOCOL.md`). The completed public-development sweep crossed seven training families, three constructs, and all 15 nonempty subsets of four feature blocks, giving 315 configurations and 945 corpus-level comparisons. X denotes XEUS, R XLS-R, P PhoneticXeus embeddings, and I the 17 phone diagnostics; I still requires PhoneticXeus inference, so omitting P does not necessarily avoid that encoder. Existing speaker-disjoint partitions, corpus-balanced weights, training-only standardization, and block gains were preserved. Each configuration selected a penalty from $\{0.001, 0.01, 0.1, 1, 10, 100\}$ using equal-weight mean Spearman across all three development corpora, breaking ties by

lower macro MAE and then larger penalty. This differs from the earlier included-corpus tuning; no corpus in these tables is an untouched cross-corpus test. Private recordings and ratings were excluded from fitting and selection.

Each candidate was compared with the fixed-feature baseline refitted under the same sweep protocol (X+P+I for accuracy, X+R otherwise). Candidates within 0.01 of the numerical best were shortlisted. Promotion required macro Spearman gain at least 0.01, no corpus correlation loss exceeding 0.02, no corpus MAE increase exceeding 0.10 on the 0–10 scale, and a positive lower 95% paired speaker-bootstrap bound (2,000 draws, seed 42). Fewer blocks were preferred where benefits were inconclusive. These screening gates do not constitute correction for all search comparisons. Deployed heads were not automatically replaced.

The decision table distinguishes numerical maxima from candidates passing the complete gate. Two AESOP-only alternatives passed: X+P for fluency and X+P+I for prosody. Their macro gains over the refitted AESOP baselines were 0.0341 (95% interval [0.0118, 0.0619]) for fluency and 0.0389 ([0.0161, 0.0638]) for prosody. Both exceed the 0.01 within-family threshold. Neither became the overall construct selection, and neither was substituted into the external seven-family evaluation. The overall retained choices were koreanocean X+P+I accuracy and aesopkoreanocean X+R fluency/prosody. Thus retaining the overall configurations does not mean that no individual family improved. XEUS-only macro correlations for the selected families were 0.8407 for koreanocean accuracy and 0.8644/0.8680 for aesopkoreanocean fluency/prosody, compared with 0.8424 and 0.8660/0.8721 for the retained combinations. These small development differences do not establish that all additional blocks are necessary; no XEUS-only private evaluation or matched resource benchmark is reported. The retained choices reflect the recorded development rule, not independently established representation optimality. Matrices report all 315 macro correlations; the accompanying CSV preserves all 945 corpus-specific correlations and MAEs.

Family	Construct	Numerical best	Best ρ	Retained/gated	Promotion
speechocean	accuracy	X+P+I	0.8276	X+P+I	No
speechocean	fluency	X+P+I	0.8655	X+R	No
speechocean	prosody	X+I	0.8640	X+R	No
korean	accuracy	X+R+I	0.8322	X+P+I	No
korean	fluency	X+I	0.8544	X+R	No
korean	prosody	X	0.8560	X+R	No
aesop	accuracy	X+P+I	0.7813	X+P+I	No
aesop	fluency	X+P	0.7961	X+P	Yes
aesop	prosody	X+P+I	0.7893	X+P+I	Yes
koreanocean	accuracy	X+P+I	0.8424	X+P+I	No
koreanocean	fluency	X+P	0.8685	X+R	No
koreanocean	prosody	X+P+I	0.8691	X+R	No
aesopocean	accuracy	X+P+I	0.8309	X+P+I	No
aesopocean	fluency	X+R+P+I	0.8654	X+R	No
aesopocean	prosody	X+R+P+I	0.8653	X+R	No
koreanaesop	accuracy	X+R	0.8337	X+P+I	No
koreanaesop	fluency	X+R+P	0.8600	X+R	No
koreanaesop	prosody	X+R	0.8599	X+R	No
aesopkoreanocean	accuracy	X+R+P+I	0.8425	X+P+I	No
aesopkoreanocean	fluency	X+R+P+I	0.8710	X+R	No
aesopkoreanocean	prosody	X+R	0.8721	X+R	No

D.3 Accuracy representation matrix

Columns S, K, A, SK, SA, KA, and SKA denote the seven corpus families in that order. Values are macro development Spearman correlations.

Features	S	K	A	SK	SA	KA	SKA
X	0.8208	0.8321	0.7655	0.8407	0.8226	0.8318	0.8336
R	0.6715	0.7460	0.3186	0.7646	0.7465	0.7766	0.7972
P	0.7860	0.7366	0.7292	0.8228	0.8065	0.7600	0.8273
I	0.6080	0.6742	0.6201	0.6961	0.3225	0.5164	0.4372
X+R	0.7957	0.8297	0.7414	0.8322	0.8257	0.8337	0.8361
X+P	0.8190	0.8030	0.7732	0.8350	0.8231	0.8159	0.8354
X+I	0.8253	0.8321	0.7522	0.8411	0.8249	0.8149	0.8362
R+P	0.7505	0.7690	0.6650	0.8192	0.8008	0.7871	0.8287
R+I	0.7290	0.7875	0.5625	0.7968	0.7765	0.7830	0.8085
P+I	0.7991	0.7750	0.7582	0.8308	0.8155	0.7171	0.8356
X+R+P	0.7962	0.8104	0.7453	0.8280	0.8235	0.8225	0.8350
X+R+I	0.7984	0.8322	0.7223	0.8363	0.8266	0.8183	0.8382
X+P+I	0.8276	0.8212	0.7813	0.8424	0.8309	0.8048	0.8409
R+P+I	0.7723	0.7932	0.7065	0.8260	0.8104	0.7740	0.8381
X+R+P+I	0.8048	0.8218	0.7565	0.8363	0.8304	0.8132	0.8425

D.4 Fluency representation matrix

Columns S, K, A, SK, SA, KA, and SKA denote the seven corpus families in that order. Values are macro development Spearman correlations.

Features	S	K	A	SK	SA	KA	SKA
X	0.8551	0.8543	0.7782	0.8630	0.8588	0.8545	0.8644
R	0.7312	0.7734	0.4373	0.8213	0.8000	0.8100	0.8319
P	0.8428	0.7912	0.7640	0.8592	0.8495	0.8127	0.8599
I	0.6288	0.6757	0.6577	0.6876	0.4060	0.5606	0.5002
X+R	0.8364	0.8444	0.7620	0.8567	0.8599	0.8574	0.8660
X+P	0.8613	0.8465	0.7961	0.8685	0.8645	0.8578	0.8702
X+I	0.8569	0.8544	0.7671	0.8618	0.8576	0.8538	0.8662
R+P	0.8178	0.8162	0.7153	0.8555	0.8498	0.8415	0.8653
R+I	0.7833	0.8076	0.6342	0.8407	0.8291	0.8200	0.8458
P+I	0.8469	0.8099	0.7748	0.8570	0.8511	0.8004	0.8593
X+R+P	0.8481	0.8458	0.7730	0.8609	0.8654	0.8600	0.8696
X+R+I	0.8443	0.8516	0.7400	0.8593	0.8598	0.8533	0.8652
X+P+I	0.8655	0.8508	0.7866	0.8666	0.8644	0.8498	0.8677
R+P+I	0.8340	0.8283	0.7268	0.8576	0.8515	0.8333	0.8653
X+R+P+I	0.8509	0.8481	0.7635	0.8639	0.8654	0.8556	0.8710

D.5 Prosody representation matrix

Columns S, K, A, SK, SA, KA, and SKA denote the seven corpus families in that order. Values are macro development Spearman correlations.

Features	S	K	A	SK	SA	KA	SKA
X	0.8631	0.8560	0.7817	0.8688	0.8599	0.8573	0.8680
R	0.7489	0.7740	0.3583	0.8172	0.7891	0.8178	0.8377
P	0.8313	0.7768	0.7520	0.8503	0.8416	0.7970	0.8526
I	0.6228	0.6834	0.6322	0.6960	0.4179	0.5379	0.5035
X+R	0.8410	0.8483	0.7504	0.8657	0.8653	0.8599	0.8721
X+P	0.8620	0.8384	0.7872	0.8664	0.8615	0.8508	0.8681
X+I	0.8640	0.8524	0.7627	0.8659	0.8609	0.8499	0.8693
R+P	0.8077	0.8070	0.6886	0.8517	0.8409	0.8296	0.8586
R+I	0.7789	0.8076	0.5893	0.8412	0.8229	0.8208	0.8498
P+I	0.8361	0.7979	0.7792	0.8527	0.8449	0.7793	0.8539
X+R+P	0.8444	0.8408	0.7602	0.8662	0.8647	0.8540	0.8707
X+R+I	0.8485	0.8512	0.7281	0.8663	0.8639	0.8495	0.8687
X+P+I	0.8635	0.8459	0.7893	0.8691	0.8649	0.8420	0.8690
R+P+I	0.8235	0.8210	0.7086	0.8522	0.8437	0.8214	0.8608
X+R+P+I	0.8485	0.8455	0.7554	0.8668	0.8653	0.8477	0.8715

E Execution environment

The same Windows/WSL machine used for the reported capacity test was inspected on October 19, 2026. The installation reports Ubuntu 26.04.1 LTS, kernel `6.18.54.1-microsoft-standard-WSL2`, Python 3.11.17, and PyTorch `2.11.0+cu128` (CUDA build 12.8). Hardware comprises an Intel Core i7-11700K with 16 logical CPUs and an NVIDIA RTX 3080 with 10 GiB VRAM; the host has 32 GB RAM, with approximately 24 GiB allocated to WSL and 8 GiB swap. The reported inference capacity test used the CPU, not CUDA. These versions document the inspected installation; the archived timing and memory trace, rather than a new benchmark, supplies the reported performance measurements.

F External rater backgrounds and procedure

The original human-evaluation study reports three experienced university raters. The Korea-based rater had more than 20 years of teaching experience, including a decade of university pronunciation instruction and speaking assessment. The Japan-based rater had nearly 30 years of teaching experience and more than 10 years as a Cambridge English speaking assessor. The China-based rater had nearly 20 years of experience, including high-stakes speech rating and university speaking assessment in an international joint-degree program. All three held professorial positions, as confirmed by the study author.

Before evaluating the source recordings, the raters jointly assessed separate Stella recordings spanning proficiency levels and discussed discrepancies to establish shared criteria for segmental accuracy, fluency, and prosody. They then scored the 306 source recordings independently in randomized order, blinded to participant identity and automated scores. Ratings used a 0–100 scale in five-point increments, with communicative-effectiveness anchors labelled “High Intelligibility” above 85 and “Low Intelligibility” below 40. These are the documented general scale anchors, not a newly constructed analytic rubric for each dimension. The original study’s authors served as its raters; their ratings were reused as the external reference for the present model evaluation. One recording was subsequently excluded for excessive duration, leaving the matched 305-recording cohort. Human calibration used separate recordings and did not fit or calibrate any automated scorer.