arXiv:2610.25599v1 [cs.LG] 19 Oct 2026

Text Dominance in Paralinguistic Conflict: Tracing Suppression Through Audio LLM Circuits

Ankur Banerjee[*,1] **Vishnu Guntakala**[*,1] **Danity Pike**[†,1]
Advay Roongta[†,1] **Kevin Yang**[2] **Ruizhe Li**[3]

[1]Independent Researcher [2]Brown University [3]University of Birmingham
{ankurbaner09, vishnun1231, roongtaadvay}@gmail.com
pikedm27@episcopalacademy.org kevin_c_yang@brown.edu r.li.7@bham.ac.uk

Abstract

Audio LLMs deployed in settings such as crisis hotlines and customer support are only useful beyond speech-to-text pipelines if they use acoustic cues rather than the transcript alone. We test whether they prioritise transcript content over paralinguistic cues (emotion, speaker identity, sarcasm) and trace the mechanism when they do. Across six audio LLMs under LISTEN-style conflict conditions [Chen et al., 2025], the mean text-dominance rate on mismatched text+audio items is 0.78 (0.86 excluding sarcasm). In the most text-dominant screened model, probing shows speaker identity remains linearly decodable at ceiling at the language model's input, so the failure is not perceptual. Back-patching, reported to repair lexical text dominance, has no measurable effect here; but it also fails on its own source task in our replication (n=300), so we treat its null as uninformative rather than as evidence of non-transfer. Attention knockout partially restores audio-grounded output by blocking the claim from propagating through intermediate tokens during open-ended generation, whereas blocking the answer position's direct connection to the conflicting claim has no effect. The claim's influence therefore travels through an intermediate relay rather than a single severable edge. The effect appears only in free-form generation and not in closed-set multiple choice, suggesting that benchmarks relying on that format may be systematically blind to interventions of this kind. Code and data are publicly available at `https://github.com/AdvayRoongta/audio-llm-mech-interp`.

1 Introduction

Audio LLMs are increasingly used in conversational applications, including customer service, interview transcription, and voice-based assistants [Chu et al., 2024, Tang et al., 2024]. But they are largely trained and evaluated on transcription and other lexically-centred tasks, which encourages them to prioritise the words spoken over acoustic and non-verbal information such as prosody, emotion, and speaker identity. That tendency undercuts the central advantage of processing raw audio: access to information an automatic speech recognition pipeline would discard before a text-based LLM ever saw it.

The cost is concrete wherever voice carries meaning the transcript does not. A caller on a crisis line may say everything is fine while their prosody says otherwise; a speaker's voice may contradict their claimed identity, a failure mode central to fraud detection [Hansen and Hasan, 2015]; a customer's sarcasm may exist only in their tone. Depression and trauma screening depend on vocal affect for the

[*]Equal contribution.
[†]Equal contribution.

Presented at the InterpScience Workshop at the 40th Conference on Neural Information Processing Systems (NeurIPS 2026).

same reason [Cummins et al., 2015]. In each case the non-verbal signal modifies or contradicts the lexical content, and a text-dominant model discards it.

Text dominance in audio LLMs has been evaluated behaviourally, typically by pairing audio with transcript text asserting a conflicting attribute [Chen et al., 2025, Wang et al., 2025], and studied mechanistically for *lexical* conflicts, where back-patching has been proposed as a training-free repair [Cho et al., 2026]. What remains unexamined is the mechanism behind suppression of *paralinguistic* evidence: whether the circuits that produce lexical text dominance also govern conflicts carried by prosody, emotion, and speaker identity, and whether the interventions developed there transfer. We localise the responsible components, test those interventions, and identify one that works.

We approach this problem mechanistically rather than treating text dominance as a purely behavioural failure. Observing that a model follows text over audio does not tell us why: the model may have failed to encode the acoustic signal in the first place, or it may have encoded the signal but failed to use it when resolving the conflict. We therefore investigate the failure in four stages. First, linear probing tests whether paralinguistic information survives the audio encoder and projector and remains available to the language model. Second, attribution patching identifies where the model relies on the audio and textual signals when they conflict. Third, back-patching and direct amplification test whether the audio pathway can be strengthened enough to overcome the textual signal. Finally, attention knockout tests the converse hypothesis: that reducing the influence of the conflicting text can restore audio-grounded behaviour. We first establish whether the relevant information is present, then ask how it is used, and finally intervene on the pathway implicated by that analysis. To avoid attributing missing information to suppression, we apply this mechanistic procedure only to capabilities the model can recover from audio alone.

Our contributions are fourfold. First, we introduce a 3,594-item paralinguistic conflict benchmark spanning emotional prosody, speaker identity, and sarcasm/pacing, and use it to measure text dominance across six audio-language models and two encoder families. Second, we show that in the model selected for mechanistic analysis, paralinguistic information remains fully recoverable at the language model's input, reaching ceiling-level probe accuracy from encoder layer 13 onward; the observed text dominance is therefore not explained by a failure to perceive or encode the acoustic signal. Third, we systematically test whether the audio pathway can be strengthened through activation back-patching and direct amplification, and find no configuration that restores audio-grounded answers; replicating back-patching on the lexical task for which it was designed also yields no repair, so the intervention provides no working positive control and its null on paralinguistic conflict is uninformative about mechanism. Fourth, we localise the textual influence to an intermediate relay pathway rather than a direct claim-to-answer connection, and show that attenuating this pathway during free-form generation reduces text dominance substantially while preserving matched-item and text-only performance. This effect is not captured by the closed-set multiple-choice evaluation used in standard audio-language-model benchmarks.

2 Related Work

Text dominance in audio LLMs. Audio LLMs encode raw audio and map that representation into the language model's input space alongside the text instruction [Chu et al., 2024, Tang et al., 2024, Radford et al., 2023], giving the model access to acoustic detail that transcription would discard. In practice they follow textual information over contradictory audio, a behaviour known as text dominance. Cho et al. [2026] analysed this mechanistically and found that text and audio propagate through distinct pathways early in processing before converging in later layers, with the text branch inhibiting rather than replacing audio representations; they proposed back-patching as a training-free repair. Billa [2026] examined audio–text conflicts across languages with the ALME benchmark and found that audio–text conflicts produce far higher text-dominance rates than text–text conflicts. Both address conflicts in lexical or semantic content, whereas we ask whether the same mechanisms apply when the conflicting information is carried by prosody, emotion, or speaker identity.

Mechanistic interpretability methods. Mechanistic interpretability identifies the internal components involved in a model's behaviour rather than characterising that behaviour only at the level of inputs and outputs [Wang et al., 2023]. Activation patching is the standard causal tool: a component's activation is replaced with one from a clean or corrupted input and the change in output is measured [Vig et al., 2020, Meng et al., 2022]. Done exhaustively it requires one forward pass per component,

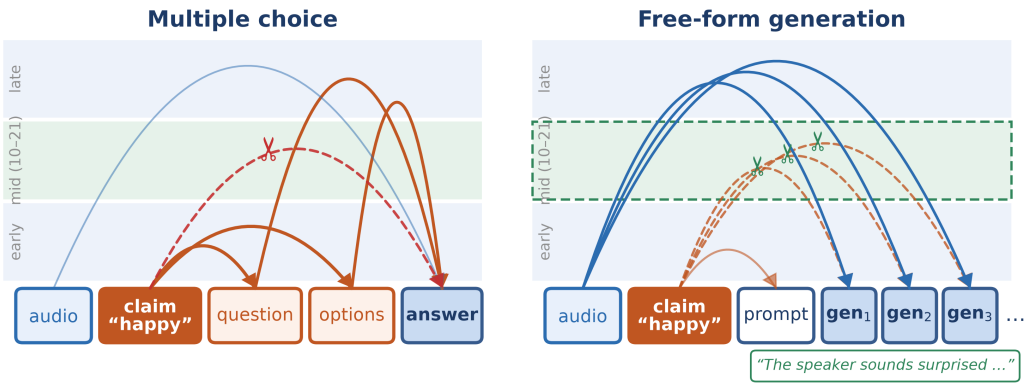

Figure 1: Attention knockout, the intervention at the centre of our analysis. Tokens run left to right; each edge is one position attending to another, drawn at the layer band where that attention occurs. Left: in closed-set multiple choice the claim is copied into the question and option tokens early and the answer position reads it from them, so blocking the direct claim→answer edge at all 32 layers leaves TDR unchanged (0.98 → 0.98). Right: in free-form generation we block every generated token from the claim span at mid layers (10–21), re-applied at each decoding step; with the relay cut, generation follows the audio (gen-TDR 1.00 → 0.38). The direct-edge and mid-layer blocks are the measured interventions; the placement of relay-building and audio-reading edges in the early and late bands is schematic.

so attribution patching substitutes a gradient-based approximation requiring far fewer [Syed et al., 2023]. That approximation can misrank components, so a circuit analysis must include a faithfulness check verifying that the selected components reproduce the full model's behaviour [Wang et al., 2023, Hanna et al., 2024]. Cho et al. [2026] applied this pipeline to lexical text dominance and proposed back-patching, in which late-layer activations are routed into early layers. Whether the same circuit or the same repair applies to paralinguistic conflict has not been tested.

Paralinguistic representation and benchmarks. Paralinguistic cues include tone, pitch, intensity, timing, emphasis, and speaker characteristics such as age or gender [Schuller et al., 2013], none of which are recoverable from a transcript alone. Chen et al. [2025] introduced the LISTEN benchmark, separating lexical from acoustic information through neutral, matched, mismatched, and paralinguistic conditions, and found consistent lexical dominance across six audio LLMs, with models often defaulting to neutral interpretations under conflict. MCR-Bench [Wang et al., 2025] measures adherence to adversarial textual descriptions on sound and vocal classification, and PARSA-Bench [Kalahroodi et al., 2026] reports that text-only baselines outperform audio counterparts in Persian, indicating the reliance is not confined to English. A parallel line of work steers attention at inference time to resolve conflicts between competing *textual* sources: PASTA [Zhang et al., 2024] reweights a small set of attention heads to emphasise user-specified spans, while other methods shift activations along probe-identified directions [Li et al., 2023] or use contrastive decoding that upweights context-conditioned predictions [Shi et al., 2024]; the broader problem of contextual versus parametric knowledge conflict has been studied extensively in text-only settings [Longpre et al., 2021, Xu et al., 2024]. These establish behavioural rates or operate entirely within a single textual stream; none localises the components responsible when the competing evidence is acoustic and must first pass through an audio encoder and projector. Appendix A.1 gives an extended treatment.

3 Method

3.1 Dataset Construction and Model Selection

We created a 3,594 item dataset in three categories from existing speech corpora with acoustic and prosodic annotations, using labels provided by MUStARD++ [Ray et al., 2022] and emotion information encoded in RAVDESS [Livingstone and Russo, 2018] file names. Using the LISTEN

benchmark [Chen et al., 2025], we created 3 types of conflicts: Emotional Prosody Conflict (e.g., cheerful text paired with an angry vocal tone), Speaker Identity Conflict (e.g., a transcript claiming a male speaker paired with a female voice), and Sarcasm Pacing Conflict (e.g., a sincere transcript paired with an overtly sarcastic tone).

We evaluated 6 audio language models using two families of audio encoders to determine whether text dominance is a phenomenon specific to the Whisper-based encoders [Radford et al., 2023] or common across different encoder architectures. Four models utilize Whisper-family encoders: Qwen2-Audio [Chu et al., 2024],

[Tang et al., 2024], DeSTA2.5 [Lu et al., 2025], and Audex [Labs, 2026], a 30-billion-parameter model included to test whether text dominance continues to occur at a larger scale. Two models use non-Whisper audio encoders: Phi-4-multimodal [Abouelenin et al., 2025], which uses a custom convolutional-conformer encoder, and Baichuan-Omni-1.5 [Li et al., 2025], which uses a residual-vector quantized audio tokenizer. Baichuan-Omni-1.5 shares the Qwen2.5-7B [Yang et al., 2024] language-model backbone with Qwen2-Audio, but we include it because our comparison concerns the audio encoder rather than the language model backbone. All inferences were run with 16-bit precision to prevent confoundation of accuracy and text-dominance measurements central to our analysis.

3.2 Mechanistic Interpretability Procedure

Figure 1 summarizes the procedure. To identify the mechanism behind this behavior in the model showing the strongest bias, we probed internal activations, applied attribution patching, tested back-patching, and used attention knockout. In the first set of experiments, we applied linear probing [Alain and Bengio, 2016, Belinkov, 2022] on the activation of the model layers without updating their weights to verify whether paralinguistic cues are encoded at each level of representation starting from the audio encoder up to the projector layer where the language model receives input. Second, attribution patching [Syed et al., 2023, Hanna et al., 2024] was performed as a gradient-based approximation approach to estimate the causal effect of each component on the model's decisions without the cost of testing them independently, and confirmed its faithfulness by verifying if the highly weighted components can replicate the output of the model alone. Third, back-patching [Cho et al., 2026] was performed as an intervention from previous research related to lexical text dominance, where activations from the later layers were replicated into earlier layers varying on multiple dimensions to determine whether it transfers to paralinguistic conflict. Fourth, we used attention knockout [Geva et al., 2023] to examine how the effect of the claim would propagate within the neural network by knocking out the attention of specific token positions to the false text claim. We evaluated each intervention under both multiple-choice and open-ended generation settings. Isolated controls determined if any effect seen was because of actual repair versus damage done to the output of the model.

4 Results

4.1 Statistical Significance of Text-Dominance Differences

Text-dominance rates differ significantly between models in all three categories (Cochran's Q [Cochran, 1950], the appropriate omnibus test for binary outcomes measured repeatedly on the same items: emotional prosody $Q = 201.04$, $df = 5$, $n = 304$; speaker identity $Q = 218.09$, $df = 5$, $n = 331$; sarcasm/pacing $Q = 345.00$, $df = 5$, $n = 369$; all $p < 10^{-40}$). Complete-case n varies because we drop items on which any model returned an unparseable response, which Baichuan-Omni-1.5 accounts for almost entirely. Averaged over the six models above on cells with modality-split scoring, the text-dominance rate on mismatched text+audio items is 0.78 (0.86 excluding sarcasm); SALMONN is omitted because its run predates modality-split scoring, and pooling in text-only items would inflate the rate (Figure 2).

Pairwise McNemar tests [McNemar, 1947] with Holm–Bonferroni correction [Holm, 1979] across all 15 comparisons are reported in Appendix A (Tables 15–17). We note one result here because it bounds what our behavioural data can support: on speaker identity, the five models other than SALMONN are statistically indistinguishable from one another despite an 8-point spread in raw text-dominance rate, so that spread should not be read as a model ranking at this sample size.

Our mechanistic findings come from a single model. Qwen2-Audio-7B-Instruct [Chu et al., 2024] showed the strongest text-dominance bias among models passing capability screening, and all probing, attribution patching, back-patching and attention-knockout results below are from it.

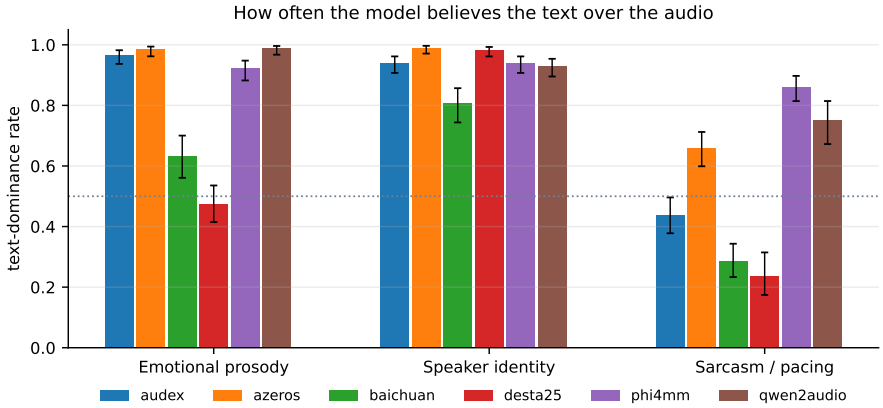

Figure 2: Text-dominance rate on mismatched `text_audio` items, by model and conflict category. Error bars are 95% Wilson intervals; the dotted line marks chance. SALMONN-13B is omitted because its behavioural run predates modality-split scoring; its text-dominance rates appear in Table 4. AZeroS was evaluated behaviourally and is shown for completeness; it is not among the six models in the statistical comparison.

4.2 Capability Screening

Attribution-based localization assumes that the model can perform the task on clean inputs: where audio-only accuracy is at chance, a high text-dominance rate reflects the absence of a usable acoustic representation rather than the suppression of one [Chen et al., 2025]. We therefore screen every model in every category on the `Paralinguistic` (audio-only) condition and keep only those exceeding chance for mechanistic analysis.

Five of six models pass the screen in both categories we analyze mechanistically. Baichuan-Omni-1.5 falls below chance on emotional prosody (0.099 against 0.125) and is excluded from that category. SALMONN-13B clears emotional prosody (0.150 against 0.125) but falls below chance on speaker identity (0.453 against 0.500) under the option log-probability scoring used by the mechanistic screen, and is excluded on that basis.

We note that SALMONN's behavioural accuracy on speaker identity (0.600) exceeds its screened accuracy, because the behavioural scorer excludes responses falling outside the option set from its denominator and SALMONN produces such responses on roughly a fifth of items. The log-probability screen cannot fail to parse and is the criterion we apply throughout. The screen therefore measures whether a model answers correctly from audio in our prompt format, and cannot separate a model that does not perceive prosody from one that perceives it but does not answer in the form we ask; we treat the exclusion as a statement about what we can measure rather than about what SALMONN-13B can hear.

All probing, attribution patching, back-patching and attention-knockout results in the main text are from Qwen2-Audio-7B-Instruct [Chu et al., 2024], which shows the strongest and most consistently measurable bias among screened models. It uses a Whisper-family encoder [Radford et al., 2023], so our account may be specific to that encoder family; the two non-Whisper models in our behavioural set would settle this, and we did not run them mechanistically.

4.3 Audio Information Retention and Attribution

Next, we evaluate whether audio information is retained in the model prior to the decoder and reasoning steps. Using a linear probe classifier [Alain and Bengio, 2016], we found that speaker gender was recovered at an accuracy of 1.000 in the encoder from layer 13 onward, and that this

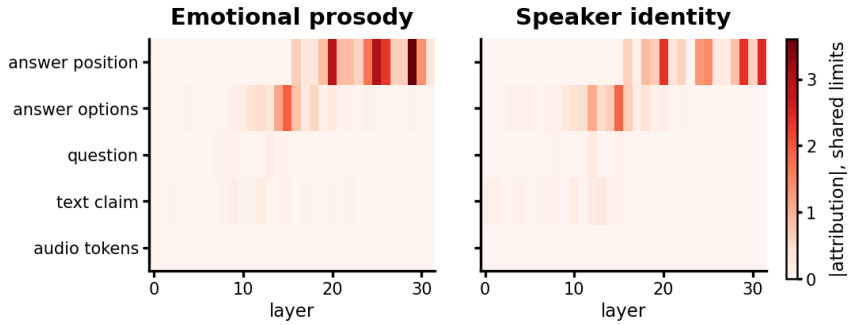

Figure 3: Qwen2-Audio-7B text circuit attribution by layer and prompt position

information was not lost in the projector. This is consistent with prior benchmark findings that speaker gender is near-ceiling for models in this family [Kalahroodi et al., 2026]. We paired this with attribution patching [Syed et al., 2023] to identify how much audio information was truly used by the model. Text was found to contribute 3.29 compared to audio's 0.18 contribution, a ratio of roughly 18:1. We applied a faithfulness test [Wang et al., 2023, Hanna et al., 2024] by comparing the output of the full model to the top 20% of components identified by attribution patching, which reproduced 100% of the original model's behavior.

4.4 Backpatching Produces No Measurable Repair

Due to the disproportionate signal in text, we tried backpatching which previous work has shown reduces TDR for lexical text dominance [Cho et al., 2026]. The hidden states of the audio token span in decoder layers 31/32, identified by their position within the sequence, were written into earlier layer 0, and the model was re-run from that point. We expected this to work because adding a later activation into an earlier one would give the model another opportunity to process the audio evidence, increasing its effective weight.

Attribution data from the previous experiment revealed the influence of audio peaks at layer 9 and falls to 0.5% by the later layers. The best source layer was layer 16, which gained 0.511 nats against a baseline of -7.32 nats, closing 7% of the gap; sources in later layers made things worse, with margins falling as low as -0.381 nats below baseline. Widening the destination window (widths 1, 3, 5, and 7) proved ineffective, as did amplifying the audio span (1.5x, 2.0x, 3.0x, 5.0x). At the highest amplification (5.0x), capability was reduced from 1.000 to 0.100, while the decision margin became undefined. The norm profile grew 4.9x between the early and late layers, but rescaling the injected vectors to compensate did not help either.

To establish whether this intervention functions as a positive control in our pipeline, we additionally rebuilt the lexical conflict task it was developed for from the ALME benchmark [Billa, 2026], resynthesizing audio with TTS, and ran each of the four English flip types at its own published layer configuration. Results are given in Table 1.

The phenomenon reproduces in kind — text dominance of 0.793 against an audio-only ceiling of 0.908 — but the repair does not: back-patching yields no measurable gain on the lexical task either, altering the outcome on 26 of 300 items and helping on 14 while hurting on 12. Because the published configurations are tuned per flip type on the original authors' audio and we resynthesized with a different TTS system, we read this as configuration sensitivity of a layer-pinned intervention rather than as a challenge to the original result. What these results jointly establish is narrower than transfer failure: back-patching provides no working positive control in our hands, in either setting, so we draw no conclusion from its null about the relationship between lexical and paralinguistic mechanisms. The audio-side evidence that does bear weight comes from probing and amplification, reported above and taken up in Section 5.

4.5 Attention Knockout

We next tried reducing the text's weight instead of increasing the audio's. We applied attention knockout [Geva et al., 2023], blocking the answer position from attending to the text claim tokens, to

Table 1: Replication of back-patching on the lexical conflict task it was developed for, rebuilt from ALME [Billa, 2026] with TTS resynthesis. Each flip type is run at its own published layer configuration, pinned rather than swept. Paired 95% CI and McNemar exact test on discordant pairs.

Flip type	Config (src→dst, window)	n	Baseline → Patched	Gain
Adjective swap	$15 \to 4$, w=3	75	0.213 → 0.227	+0.013
Number swap	$12 \to 4$, w=1	75	0.213 → 0.213	+0.000
Time swap	$13 \to 4$, w=3	75	0.320 → 0.360	+0.040
Negation add	$31 \to 6$, w=3	75	0.080 → 0.053	−0.027
Pooled	—	**300**	**0.207 → 0.213**	+0.007

Paired 95% CI $[-0.027, +0.040]$; 14 items fixed, 12 broken; McNemar exact $p = 0.845$.
Reported in prior work: +0.160 [Cho et al., 2026].

prevent the claim from carrying too much weight when audio evidence contradicted it. This did not reduce TDR because information from the counterfactual text claim was merged into the question and option tokens early: by the time the model answers, the claim is no longer confined to the claim tokens: it has already been copied into the intervening question and option tokens, and the answer position reads it from them.

Attention knockout remained effective, but only when applied throughout free-form generation rather than solely at the final answer position. In the multiple-choice setting, the model generates an answer from a small number of fixed positions, whereas free-form generation creates a sequence of new tokens, each of which can attend to information from earlier positions. We therefore tested which positions needed to be prevented from attending to the conflicting textual claim (Table 2). Blocking only the later positions where the answer would be produced had little effect: emotional-prosody accuracy remained unchanged (0.000 to 0.000), while speaker-identity accuracy increased modestly (0.067 to 0.083). In contrast, blocking every position after the claim from attending to it substantially reduced textual dominance, increasing audio-consistent accuracy from 0.000 to 0.533 for emotional prosody and from 0.067 to 0.717 for speaker identity. Under this intervention, TDR fell from 0.983 to 0.150 for prosody and from 0.933 to 0.283 for identity, at a capability cost of 1.000 to 0.367 and 1.000 to 0.800 respectively. These results indicate that the relevant distinction is not which layers are intervened on, but which token positions are blocked: in the multiple-choice setting, only the final answer position is affected, whereas free-form generation allows the intervention to block the entire sequence of tokens that could otherwise carry the textual claim forward. The results therefore suggest that the model's textual bias is mediated by an intermediate relay pathway, in which the claim is repeatedly re-encoded into later tokens before influencing the final response.

Table 2: Routing bisection: audio-consistent accuracy under closed-set scoring when different token positions are blocked from attending to the claim (all layers, γ=0). Blocking the answer position alone recovers none of the effect; blocking every position downstream of the claim recovers all of it.

	Emotional Prosody			Speaker Identity		
Positions blocked	Audio acc.	Δ margin	Capability	Audio acc.	Δ margin	Capability
None (baseline)	0.000	—	1.000	0.067	—	1.000
Answer position only	0.000	−0.00	1.000	0.083	+0.01	1.000
All positions after claim	**0.533**	+10.54	0.367	**0.717**	+5.80	0.800
Every position	0.533	+10.54	0.367	0.717	+5.80	0.800

4.6 Controls

We ran controls to confirm that the intervention was not simply damaging model performance in a way that mechanically inflated TDR, and that similar interventions inconsistent with our account do not work. On matched items, accuracy increased from 0.767 to 0.900 for prosody and from 0.933 to 1.000 for identity. The model demonstrated identical accuracy for prosody and speaker identity when given text alone with no audio. This effect appeared only when blocking mid-layers; blocking late layers did not affect TDR for either category. Finally, blocking at all layers suppressed TDR further but impaired the model's ability, dropping matched-case prosody accuracy to 0.667. This rules out

the possibility that the intervention simply deletes the text claim from the input and forces the model to rely on audio by default. Full per-condition results are given in Table 3.

Table 3: Free-form generation under attention knockout. gen-TDR is the fraction of descriptions naming the false claimed label (lower is better); "names true" is the fraction naming the audio-supported label. Capability is accuracy on matched items; cost is accuracy on text-only items, where no audio is present and the claim is the only evidence. Fluency is mean per-token log-probability under the unmodified model, scored in a claim-free context. 95% Wilson intervals.

Category	Condition	gen-TDR [95% CI]	Names true	Capability	Cost	Fluency
Prosody	Baseline	1.000 [0.86, 1.00]	0.000	0.767	0.567	-1.28
	Late γ=0	1.000 [0.86, 1.00]	0.000	0.900	0.467	-1.13
	Mid γ=0	**0.375** [0.21, 0.57]	**0.500**	0.900	0.567	-0.64
	All γ=0	0.182 [0.07, 0.39]	0.600	0.667	0.533	-0.46
Identity	Baseline	1.000 [0.89, 1.00]	0.000	0.933	1.000	-1.45
	Late γ=0	1.000 [0.87, 1.00]	0.033	1.000	1.000	-1.54
	Mid γ=0	**0.433** [0.27, 0.61]	**0.567**	1.000	1.000	-0.78
	All γ=0	0.300 [0.17, 0.48]	0.700	0.933	0.900	-0.69

5 Discussion

Our repair evidence divides into what back-patching can and cannot tell us. Because it produced no measurable effect even on the lexical task it was designed for (Table 1), its null on paralinguistic conflict carries no information about whether the two phenomena share a mechanism, and we draw no such conclusion. The audio-side evidence that does bear weight comes from probing and amplification. This is not an availability failure: probes recover speaker gender from the encoder at ceiling from layer 13 onward and the information remains linearly decodable at the projector output. The acoustic evidence reaches the language model and is ignored. Amplification points the same way, where scaling the audio span 5x left TDR unchanged while capability collapsed from 1.000 to 0.100, so interventions that raise audio salience are targeting the wrong variable.

The two attention-knockout conditions locate where the asymmetry lives. Blocking the answer position from attending to the claim leaves TDR unchanged, which indicates the claim has already been incorporated into the question and option representations in the early layers. Blocking every generated position across mid-layers during free-form generation moves it substantially. We conclude that the claim's influence is distributed across intermediate positions and reconstructed at each generation step, rather than passed through a single site that one patch can sever. Concurrent work applying causal tracing to audio–text fusion reports a compatible structure in DeSTA, Qwen and Voxtral: the final sequence position acts as an informational bottleneck, and intermediate positions perform an attention-like retrieval of task-relevant audio context [Chen et al., 2026]. That analysis is descriptive and does not involve conflicting inputs, so it does not test suppression; but it independently locates the two components our routing account depends on.

Two alternative explanations should be addressed. First, the mid-layer knockout might degrade processing generally rather than remove the claim specifically. Our controls rule this out by showing that matched-item accuracy under the intervention rises rather than falls in both categories. Second, the gap between multiple-choice and free-form results is larger than the gap between blocking patterns, so the evaluation format may itself be driving the effect. These are not independent, and our account predicts exactly this: option tokens multiply the positions carrying claim-informed content, so removing any one leaves others to propagate it. Evaluation format is therefore not incidental to measuring text dominance; it determines how much redundancy the mechanism has to work with.

6 Limitations

Our datasets use acted performances from RAVDESS [Livingstone and Russo, 2018] and MUStARD++ [Ray et al., 2022], so their paralinguistic signals and audio-text conflicts may be more categorical than in natural conversation. All conflicts are also in English, leaving open whether the observed mechanisms generalize across languages, particularly where prior work reports different

conflict rates [Billa, 2026]. We evaluate emotional prosody, speaker identity, and sarcasm/pacing, but the mechanistic analysis excludes sarcasm/pacing because the audio-only labels were insufficiently determined to support a suppression claim. Other forms of paralinguistic conflict, including non-speech events, accent/dialect, and vocal effort, remain untested.

Our mechanistic analysis and repair are primarily demonstrated in a single model and on two conflict categories. The second model establishes the prerequisite that paralinguistic information is recoverable, but does not independently validate the identified mechanism. The intervention also has important scope conditions: it reduces text dominance during free-form generation but not closed-set multiple-choice scoring, and its effect depends on attenuating the textual claim across the intermediate positions that can relay its information. Our lexical back-patching replication produced a null despite the procedure's reported effectiveness; because our task and TTS instantiation differ from the original study, we interpret this cautiously as evidence of configuration sensitivity rather than a contradiction.

Finally, the intervention is unconditional: it attenuates the claim whether or not the claim is correct. Although text-only performance is preserved, we do not test cases where ambiguous audio should appropriately be overridden by a reliable textual claim. The free-form intervention experiments also use a limited number of items, so individual configuration differences should not be over-interpreted. More broadly, the same sensitivity to paralinguistic information that enables audio-grounded reasoning can facilitate inference of sensitive attributes such as emotion or identity, raising privacy and consent concerns for deployment.

7 Conclusion

We investigated whether paralinguistic suppression in audio LLMs is mechanistically related to lexical text dominance and whether it can be repaired by the interventions developed for that setting. Across six models, text-dominance under conflict is widespread and statistically robust (Cochran's Q, all categories $p < 1e\text{-}40$). In the model with the strongest and most consistently measurable bias, probing showed that paralinguistic information survives to the language model's input in full (ceiling accuracy at the projector), so the failure is not one of perception. Direct audio amplification failed to restore audio-grounded answers at any scale tested, ruling out an availability-based account; back-patching [Cho et al., 2026] produced no measurable effect in either setting, including our replication on the lexical task it was designed for, and therefore constrains no hypothesis. A routing analysis located the mechanism: the claim's influence is not carried through a single connection to the answer position but is re-encoded into intervening tokens, which is why blocking that direct connection has no effect. Targeting the relay instead by attenuating the claim across mid-layers during free-form generation reduced text-dominance substantially in both categories tested, with matched-item accuracy and text-only performance unaffected. The effect appears only in open-ended generation but not in the closed-set multiple-choice format common in audio LLM evaluation, suggesting that benchmarks relying on that format may be systematically blind to interventions of this kind.

Reproducibility Statement

To facilitate the reproducibility of our mechanistic diagnosis pipeline, detailed prompt templates and data construction methodology for the 3,594-item paralinguistic conflict benchmark—derived from the MUSTARD++ and RAVDESS corpora—are provided in Appendix A.6. This includes comprehensive examples of prompt formatting across all conditions for emotional prosody, speaker identity, and sarcasm/pacing conflicts. The complete evaluation codebase, including all scripts required for linear probing, attribution patching, and attention knockout, along with the full dataset manifests and experimental hyperparameters, will be made publicly available in an open-source repository upon publication.

References

A. Abouelenin, A. Ashfaq, A. Atkinson, H. Awadalla, N. Bach, J. Bao, A. Benhaim, M. Cai, V. Chaudhary, C. Chen, et al. Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-LoRAs. arXiv preprint arXiv:2503.01743, 2025.

G. Alain and Y. Bengio. Understanding intermediate layers using linear classifier probes. arXiv preprint arXiv:1610.01644, 2016.

Y. Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.

Billa. When audio-LLMs don't listen: A cross-linguistic study of modality arbitration (ALME benchmark). 2026, 2026.

J. Chen, Z. Guo, J. Chun, P. Wang, A. Perrault, and M. Elsner. Do audio LLMs really LISTEN, or just transcribe? Measuring lexical vs. acoustic emotion cues reliance. arXiv preprint arXiv:2510.10444, 2025.

Wei-Chih Chen, Chien-yu Huang, and Hung-yi Lee. Causal Tracing of Audio-Text Fusion in Large Audio Language Models. arXiv preprint arXiv:2603.13768, 2026.

H. Cho, S. Yoo, J. Jang, C. Kim, and J. Chung. Who wins the conflict? Mechanistic interpretability of text bias in audio LLMs. arXiv preprint arXiv:2606.18924, 2026.

Y. Chu, J. Xu, Q. Yang, H. Wei, X. Wei, Z. Guo, Y. Leng, Y. Lv, J. He, J. Lin, et al. Qwen2-Audio technical report. arXiv preprint arXiv:2407.10759, 2024.

W. G. Cochran. The comparison of percentages in matched samples. *Biometrika*, 1950.

N. Cummins, S. Scherer, J. Krajewski, S. Schnieder, J. Epps, and T. F. Quatieri. A review of depression and suicide risk assessment using speech analysis. *Speech Communication*, 2015.

M. Geva, J. Bastings, K. Filippova, and A. Globerson. Dissecting recall of factual associations in auto-regressive language models. In *EMNLP*, 2023.

M. Hanna, S. Pezzelle, and Y. Belinkov. Have faith in faithfulness: Going beyond circuit overlap when finding model mechanisms. arXiv preprint arXiv:2403.17806, 2024.

J. H. L. Hansen and T. Hasan. Speaker recognition by machines and humans: A tutorial review. *IEEE Signal Processing Magazine*, 32(6):74–99, 2015.

S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

Mohammad Javad Ranjbar Kalahroodi, Mohammad Amini, Parmis Bathayan, Heshaam Faili, and Azadeh Shakery. PARSA-Bench: A Comprehensive Persian Audio-Language Model Benchmark. arXiv preprint arXiv:2603.14456, 2026.

NVIDIA Nemotron Labs. Unified audio intelligence without regressing on text intelligence. arXiv preprint arXiv:2607.05196, 2026.

K. Li, O. Patel, F. Viégas, H. Pfister, and M. Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *NeurIPS*, 2023.

Y. Li et al. Baichuan-Omni-1.5 technical report. arXiv preprint arXiv:2501.15368, 2025.

S. R. Livingstone and F. A. Russo. The Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS): A dynamic, multimodal set of facial and vocal expressions in North American English. *PLoS ONE*, 13(5):e0196391, 2018.

S. Longpre, K. Perisetla, A. Chen, N. Ramesh, C. DuBois, and S. Singh. Entity-based knowledge conflicts in question answering. In *EMNLP*, 2021.

K.-H. Lu, Z. Chen, S.-W. Fu, C.-H. H. Yang, J. Balam, B. Ginsburg, Y.-C. F. Wang, and H.-Y. Lee. DeSTA2.5-Audio: Toward general-purpose large audio language model with self-generated cross-modal alignment. arXiv preprint arXiv:2507.02768, 2025.

Q. McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.

K. Meng, D. Bau, A. Andonian, and Y. Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever. Robust speech recognition via large-scale weak supervision. In *International Conference on Machine Learning (ICML)*, 2023.

A. Ray, S. Mishra, A. Nunna, and P. Bhattacharyya. A multimodal corpus for emotion recognition in sarcasm. In *Language Resources and Evaluation Conference (LREC)*, 2022.

B. Schuller, S. Steidl, A. Batliner, F. Burkhardt, L. Devillers, C. Müller, and S. Narayanan. Paralinguistics in speech and language—state-of-the-art and the challenge. *Computer Speech & Language*, 27(1):4–39, 2013.

W. Shi, X. Han, M. Lewis, Y. Tsvetkov, L. Zettlemoyer, and S. W. Yih. Trusting your evidence: Hallucinate less with context-aware decoding. In *NAACL*, 2024.

A. Syed, C. Rager, and A. Conmy. Attribution patching outperforms automated circuit discovery. *arXiv preprint*, 2023.

C. Tang, W. Yu, G. Sun, X. Chen, T. Tan, W. Li, L. Lu, Z. Ma, and C. Zhang. SALMONN: Towards generic hearing abilities for large language models. In *International Conference on Learning Representations (ICLR)*, 2024.

J. Vig, S. Gehrmann, Y. Belinkov, S. Qian, D. Nevo, Y. Singer, and S. Shieber. Investigating gender bias in language models using causal mediation analysis. In *NeurIPS*, 2020.

C. Wang, G. Deng, X. Yang, H. Qiu, and T. Zhang. When audio and text disagree: Benchmarking text bias in large audio-language models under cross-modal inconsistencies. arXiv preprint arXiv:2508.15407, 2025.

K. Wang, A. Variengien, A. Conmy, B. Shlegeris, and J. Steinhardt. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small. In *ICLR*, 2023.

R. Xu, Z. Qi, Z. Guo, C. Wang, H. Wang, Y. Zhang, and W. Xu. Knowledge conflicts for LLMs: A survey. In *EMNLP*, 2024.

A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, et al. Qwen2.5 technical report. arXiv preprint arXiv:2412.15115, 2024.

Q. Zhang, C. Singh, L. Liu, X. Liu, B. Yu, J. Gao, and T. Zhao. Tell your model where to attend: Post-hoc attention steering for LLMs. In *ICLR*, 2024.

A Appendix

A.1 Extended Related Work

Inference-time attention steering. A line of work resolves conflicts between competing sources of information without retraining. PASTA [Zhang et al., 2024] identifies a small set of attention heads and reweights them to emphasise user-specified text spans, improving performance on tasks involving conflicting information within a text context. Other methods steer models toward contextual or parametric knowledge by shifting activations along directions identified from probing [Li et al., 2023] or by contrastive decoding that upweights context-conditioned predictions [Shi et al., 2024]. The broader problem of conflict between contextual and parametric knowledge has been studied extensively in text-only settings [Longpre et al., 2021, Xu et al., 2024]. Together these show that modifying attention at inference time can influence which source a language model prioritises, but all concern conflicts entirely within text: the competing information arrives through the same textual stream, whether from different spans, documents, or the model's own parameters.

Our setting differs in that one source is acoustic and must pass through an audio encoder and projector before entering the residual stream, where it competes with text. It is therefore not obvious that steering methods designed for text conflicts transfer. We find that the direct adaptation — blocking

the answer position from attending to the conflicting claim — is insufficient, because the claim is re-encoded into intermediate tokens that carry it forward. Reducing textual dominance requires intervening on this relay rather than at the claim's source or at the answer position.

Behavioural benchmarks for audio–text conflict. MCR-Bench [Wang et al., 2025] measures whether audio LLMs follow adversarial textual descriptions over their own perception on sound and vocal classification tasks. PARSA-Bench [Kalahroodi et al., 2026], a Persian-language audio benchmark, reports that text-only baselines consistently outperform their audio counterparts, indicating that reliance on transcript content over acoustic evidence is not confined to English. LISTEN [Chen et al., 2025] separates lexical and acoustic information through neutral, matched, mismatched, and paralinguistic conditions, finding consistent lexical dominance across six models with accuracy improving only slightly when cues align. All three establish behavioural rates without localising the components responsible.

A.2 Per-Model Text Dominance Rate

Table 4 reports the complete-case Text Dominance Rate (TDR) for all six models across the three conflict categories, supporting the TDR claims summarized in the main Results section. Complete-case TDR excludes trials with invalid or refused model outputs (see the main Results section for exclusion criteria); per-category dropped-trial counts are given in Table 5.

Table 4: Complete-case Text Dominance Rate (%) by model and conflict category.

Category	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Emotional Prosody	98.7	78.3	68.8	93.1	98.7	83.9
Speaker Identity	96.4	89.1	98.5	96.1	94.9	70.7
Sarcasm/Pacing	34.4	24.7	20.9	65.3	58.5	56.1

Table 5: Complete-case sample sizes and dropped trials (invalid/refused outputs) by category.

Category	Total Trials	Complete-case n	Dropped	Driven by
Emotional Prosody	399	304	95	Baichuan (86), DeSTA2.5 (11), Qwen2-Audio (2)
Speaker Identity	450	331	119	Baichuan (119)
Sarcasm/Pacing	399	369	30	Qwen2-Audio (15), Baichuan (12), DeSTA2.5 (4)

A.3 Per-Model Accuracy by Condition

Tables 6–8 report accuracy (%) for all six models across the four evaluation conditions (Neutral-Text, Category-Matched, Category-Mismatched, Paralinguistic-only), for each conflict category.

Table 6: Emotional Prosody: accuracy (%) by condition and model.

Condition	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Neutral-Text	75.9	44.2	55.5	52.3	49.8	16.5
Category-Matched	99.2	78.3	86.9	93.2	100.0	90.2
Category-Mismatched	34.1	43.8	52.1	32.3	33.0	29.3
Paralinguistic-only	60.9	9.9	50.4	13.5	54.1	15.0

A.4 Audio-Side Intervention Sweeps

Table 9 reports the source-layer, destination-window, and amplitude sweeps for the audio-side interventions described in the main Results section, on emotional prosody with untruncated audio.

Table 7: Speaker Identity: accuracy (%) by condition and model.

Condition	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Neutral-Text	92.7	47.1	99.2	57.3	98.6	58.0
Category-Matched	99.3	94.1	100.0	96.9	100.0	65.8
Category-Mismatched	37.3	54.7	34.7	37.3	38.4	41.6
Paralinguistic-only	92.0	56.1	99.3	51.3	98.7	44.7

Table 8: Sarcasm/Pacing: accuracy (%) by condition and model.

Condition	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Neutral-Text	20.3	20.0	21.4	23.3	16.9	13.9
Category-Matched	46.1	29.6	21.3	70.9	68.5	56.6
Category-Mismatched	14.0	14.2	12.4	9.5	9.1	7.3
Paralinguistic-only	19.5	13.8	21.1	13.5	16.9	13.5

Table 9: Audio-side intervention sweeps (emotional prosody, untruncated audio). Decision margin is $\log P(\text{audio label}) - \log P(\text{claimed label})$ in nats; baseline margin is -7.407. No configuration closes more than 12% of the gap.

(a) Source layer sweep (w=1, dst=0)								
Source layer	4	9	13	16	20	24	28	31
Δ margin (nats)	$+0.17$	$+0.42$	$+0.66$	$\mathbf{+0.87}$	$+0.47$	$+0.04$	-0.11	-0.19
(b) Destination window sweep (src=13)								
Window width	1		3		5		7	
Δ margin (nats)	$+0.66$		$+0.60$		$+0.58$		$+0.59$	
(c) Amplitude sweep (audio span scaled in place)								
Scale factor	1.5×		2.0×		3.0×		5.0×	
Capability	1.000		1.000		1.000		0.100	

Table 10: Layer-index convention check. The published configurations are given without a stated indexing base; re-running all four under the alternative reading (both endpoints shifted down one) leaves the null unchanged.

Flip type	0-indexed gain	1-indexed gain
Adjective swap	$+0.013$	$+0.013$
Number swap	$+0.000$	$+0.000$
Time swap	$+0.040$	$+0.013$
Negation add	-0.027	-0.013
Pooled (n=300)	$\mathbf{+0.007}$	$\mathbf{+0.003}$
Paired 95% CI	$[-0.027, +0.040]$	$[-0.028, +0.035]$
McNemar exact p	0.845	1.000

A.5 Layer-Index Convention Check

A.6 Dataset Construction Examples

Tables 11–13 show example prompts for each of the four conditions (Neutral-Text, Category-Matched, Category-Mismatched, Paralinguistic-only), for each of the three conflict categories used in the final dataset.

Table 11: Example prompts: Emotional Prosody Conflict (audio emotion = happy).

Condition	Example
Neutral-Text	Transcript is "The kids are sitting by the door." How would you describe the emotional tone of the speaker?
Category-Matched	The speaker's emotion is happy. Listen to the audio and classify the emotion. How would you describe the emotional tone of the speaker? Beware that the audio description and audio itself may not match.
Category-Mismatched	The speaker's emotion is angry. Listen to the audio and classify the emotion. How would you describe the emotional tone of the speaker? Beware that the audio description and audio itself may not match.
Paralinguistic-only	Listen to the audio and classify the emotion. What emotion is expressed in the speaker's voice?

Table 12: Example prompts: Speaker Identity Conflict (speaker identity = adult female).

Condition	Example
Neutral-Text	Transcript is "The kids are sitting by the door." What is the speaker's gender identity?
Category-Matched	The speaker is an adult female. Listen to the audio and classify the gender identity. What is the speaker's gender identity? Beware that the audio description and audio itself may not match.
Category-Mismatched	The speaker is an adult male. Listen to the audio and classify the gender identity. What is the speaker's gender identity? Beware that the audio description and audio itself may not match.
Paralinguistic-only	Listen to the audio and classify the gender identity. What is the speaker's gender identity?

Table 13: Example prompts: Sarcasm/Pacing Conflict (speaker is being sincere).

Condition	Example
Neutral-Text	Transcript is "The dog is playing outside." How sincere or satirical is the speaker being?
Category-Matched	The speaker is being sincere. Listen to the audio and classify how sincere or satirical the speaker is. Beware that the audio description and audio itself may not match.
Category-Mismatched	The speaker is being sarcastic. Listen to the audio and classify how sincere or satirical the speaker is. Beware that the audio description and audio itself may not match.
Paralinguistic-only	Listen to the audio and classify how sincere or sarcastic the speaker is.

A.7 Dataset and Model Licenses

RAVDESS is released under CC BY-NC-SA 4.0 [Livingstone and Russo, 2018]; MUStARD++ is released for academic research use with no formal license specified [Ray et al., 2022]. Of the evaluated models, Qwen2-Audio-7B-Instruct and SALMONN-13B are released under Apache 2.0, Phi-4-multimodal-instruct under MIT, Baichuan-Omni-1.5 under Apache 2.0 with an accompanying community license, and Audex under NVIDIA's one-way noncommercial license; DeSTA2.5-Audio's model authors do not specify a license. All models and datasets were used strictly for non-commercial academic research, consistent with these terms.

A.8 Statistical Tests

Table 14 reports the Cochran's Q omnibus test for each conflict category (testing whether TDR differs across the six models), followed by pairwise McNemar tests with Holm-Bonferroni correction (Tables 15–17).

Table 14: Cochran's Q omnibus test for between-model TDR differences.

Category	n (complete-case)	Q	df	p-value	Significant ($p < .05$)
Emotional Prosody	304	201.04	5	1.702×10^{-41}	Yes
Speaker Identity	331	218.09	5	3.815×10^{-45}	Yes
Sarcasm/Pacing	369	345.00	5	2.082×10^{-72}	Yes

Table 15: Emotional Prosody: pairwise McNemar test, Holm-Bonferroni corrected p-values.

	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Audex	—	7.79×10^{-16}	2.85×10^{-25}	0.0024	1.0000	1.97×10^{-10}
Baichuan	7.79×10^{-16}	—	0.0115	8.11×10^{-7}	1.82×10^{-14}	0.1931
DeSTA2.5	2.85×10^{-25}	0.0115	—	1.66×10^{-15}	3.22×10^{-24}	0.0002
Phi-4-mm	0.0024	8.11×10^{-7}	1.66×10^{-15}	—	0.0024	0.0014
Qwen2-Audio	1.0000	1.82×10^{-14}	3.22×10^{-24}	0.0024	—	6.35×10^{-10}
SALMONN	1.97×10^{-10}	0.1931	0.0002	0.0014	6.35×10^{-10}	—

Table 16: Speaker Identity: pairwise McNemar test, Holm-Bonferroni corrected p-values.

	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Audex	—	0.0035	0.4813	1.0000	1.0000	6.09×10^{-19}
Baichuan	0.0035	—	3.35×10^{-6}	0.0112	0.0768	1.03×10^{-6}
DeSTA2.5	0.4813	3.35×10^{-6}	—	0.4813	0.0768	9.68×10^{-23}
Phi-4-mm	1.0000	0.0112	0.4813	—	1.0000	7.62×10^{-20}
Qwen2-Audio	1.0000	0.0768	0.0768	1.0000	—	4.81×10^{-22}
SALMONN	6.09×10^{-19}	1.03×10^{-6}	9.68×10^{-23}	7.62×10^{-20}	4.81×10^{-22}	—

Table 17: Sarcasm/Pacing: pairwise McNemar test, Holm-Bonferroni corrected p-values.

	Audex	Baichuan	DeSTA2.5	Phi-4-mm	Qwen2-Audio	SALMONN
Audex	—	0.0074	2.53×10^{-6}	4.63×10^{-24}	2.42×10^{-14}	2.14×10^{-9}
Baichuan	0.0074	—	0.3956	3.87×10^{-30}	7.58×10^{-23}	2.05×10^{-17}
DeSTA2.5	2.53×10^{-6}	0.3956	—	2.33×10^{-37}	3.44×10^{-31}	3.78×10^{-20}
Phi-4-mm	4.63×10^{-24}	3.87×10^{-30}	2.33×10^{-37}	—	0.0106	0.0044
Qwen2-Audio	2.42×10^{-14}	7.58×10^{-23}	3.44×10^{-31}	0.0106	—	0.4558
SALMONN	2.14×10^{-9}	2.05×10^{-17}	3.78×10^{-20}	0.0044	0.4558	—

A.9 Supplementary Figures

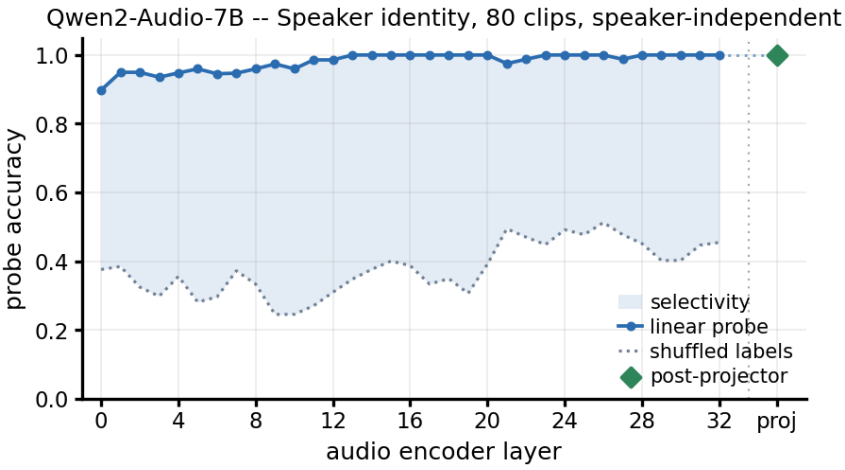

Figure 4: Linear-probe accuracy for speaker gender across encoder layers and at the projector output, with speaker-independent splits and a shuffled-label control.

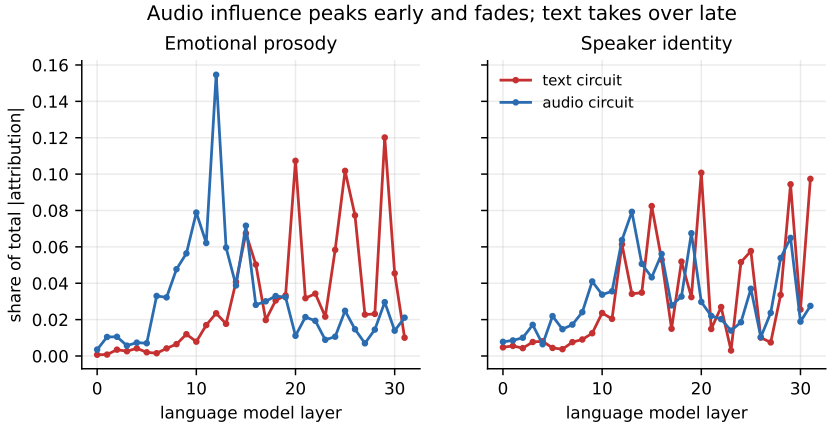

Figure 5: Full attribution map, including the audio position groups omitted from Figure 3. Text and audio panels share colour limits, so the near-uniform audio panels are the 18:1 text-to-audio ratio reported in Section 4.3 shown directly rather than summarised.

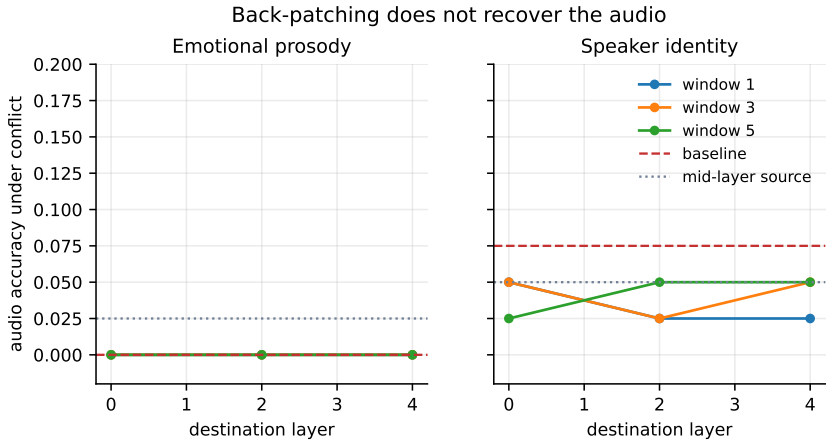

Figure 6: Back-patching sweep on paralinguistic conflicts across source layers, destination layers, and window widths.

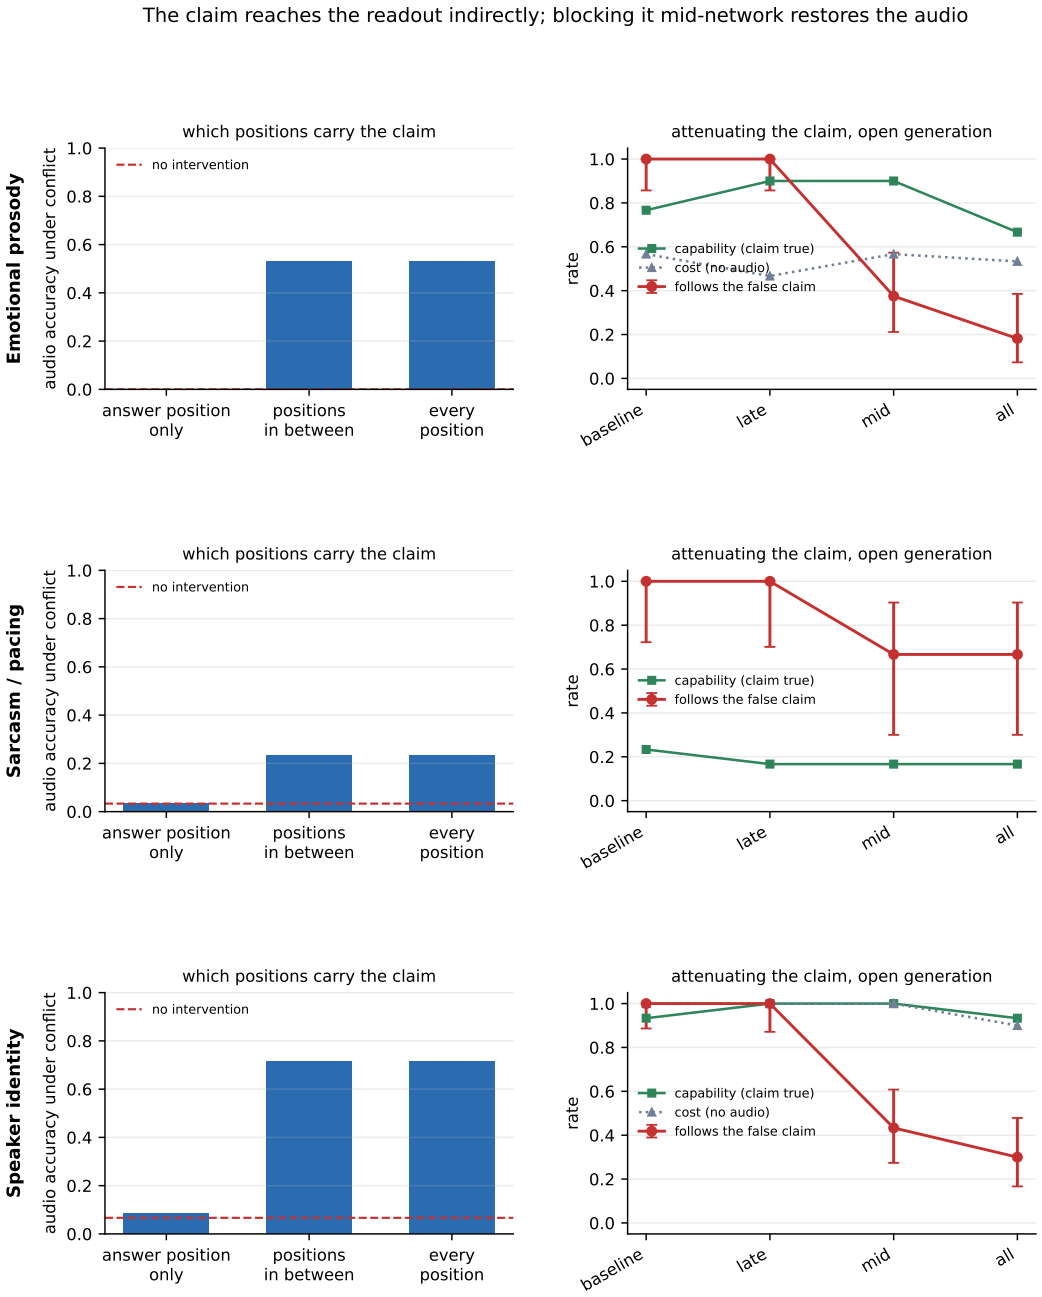

Figure 7: Left: routing bisection — blocking the answer position's direct view of the claim recovers none of the effect, while blocking the intervening positions recovers all of it. Right: the generative intervention with 95% Wilson intervals, alongside the capability and cost controls.

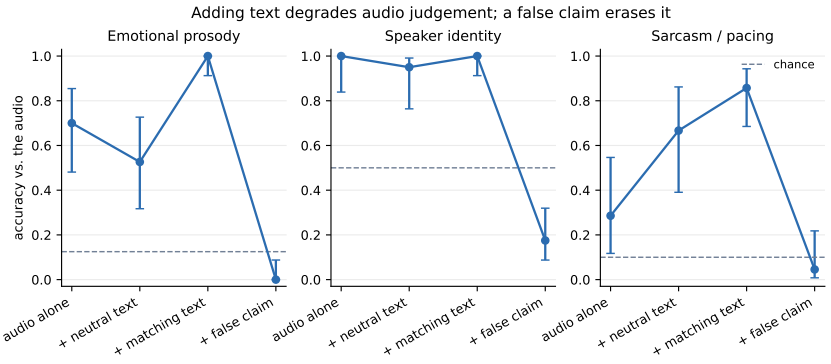

Figure 8: Accuracy across the four LISTEN conditions, by model and category.

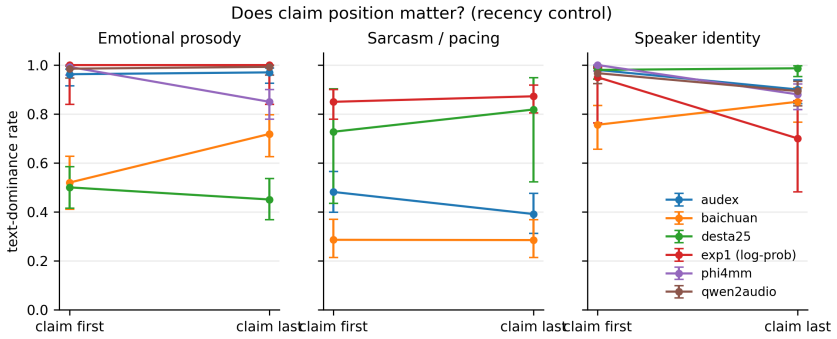

Figure 9: Text-dominance by claim order (claim-first vs. claim-last), as a control against recency bias. For emotional prosody and speaker identity, most models' text-dominance rate stays near ceiling regardless of claim position, indicating the effect is not simply an artifact of which claim appears more recently in the prompt. Sarcasm/pacing shows more movement between claim positions, but the direction differs by model, so claim order does not explain the overall text-dominance effect.