Adaptive Stopping of Streamed Reasoning for Language Model Agents

JunHyeok Oh, Byung-Jun Lee
Korea University
{the2ndlaw, byungjunlee}@korea.ac.kr
Project Page: jh2525.github.io/StreamStop/

arXiv:2610.27981v1 [cs.CL] 20 Oct 2026

Abstract

Reasoning-enabled language model agents can incur substantial computation before acting, although the reasoning needed varies across decisions. Existing stopping signals, such as confidence and intermediate-answer stability, can vary in effectiveness across reasoning models; signals obtained from hidden states or token probabilities also require internal access that may be unavailable or impractical. We introduce StreamStop, which learns model-specific stopping policies offline from prefix-level action agreement. A frozen prefix reader encodes available reasoning text for a lightweight stop head, allowing stopping without access to the reasoning model's internal signals. We construct and publicly release a dataset of agent reasoning trajectories with prefix-level action annotations. These annotations compare each prefix action with the full-reasoning reference action, enabling offline comparison of stopping rules and exact expected-return optimization of a discounted action-agreement objective. Conditioning the head on the discount factor allows the agreement–token trade-off to be adjusted at inference without retraining. In closed-loop evaluations on WebShop and ScienceWorld, StreamStop achieves higher observed success than most validation-selected fixed-budget settings while using fewer reasoning tokens than Full thinking. On WebShop with Qwen3.8-27B, it uses 42% fewer reasoning tokens per episode with similar observed task success.

1 Introduction

Language model agents interleave reasoning, actions, and observations to solve interactive tasks (Yao et al., 2023). Lengthy reasoning can delay an agent's next action, while the amount of reasoning needed varies across decisions (Bai et al., 2026; Zhong et al., 2026). Adaptive stopping aims to determine how much reasoning is sufficient for the next action.

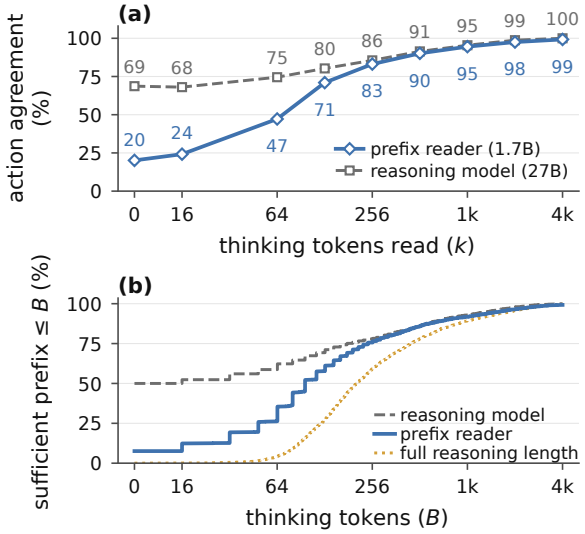

Figure 1: Prefix diagnostics on the WebShop validation set. (a) Agreement with the full-reasoning action for the prefix reader scoring valid candidates and the reasoning model answering from each prefix. (b) Cumulative distributions of sufficient-prefix and full-reasoning lengths.

Some stopping methods use the reasoning model's hidden representations (Wu et al., 2025) or token probabilities (Yang et al., 2025) to determine when to stop. These signals are unavailable through text-only APIs, and integrating stopping components into large distributed serving systems can be difficult (Wang et al., 2025).

Stopping can also be guided by confidence estimates from a separate model or by the stability of intermediate answers, without requiring access to the reasoning model's internal signals (Wang et al., 2025; Fu et al., 2025a). However, the usefulness of these signals depends on how closely they reflect the reasoning model's eventual action, and this relationship can vary across models. This calls for stopping criteria that account for the reasoning model's behavior while balancing action quality against reasoning cost.

Adaptive stopping requires judging whether the available reasoning is sufficient to determine the

next action. Generating a reasoning trajectory and making this sufficiency judgment need not require the same model capacity. Figure 1(a) compares a 1.7B prefix reader scoring valid action candidates with a 27B reasoning model answering from the same prefixes. Reader agreement with the full-reasoning action increases with prefix length. Panel (b) compares full-reasoning lengths with sufficient-prefix lengths: the earliest observed prefix after which all remaining observed actions agree with the reference. These sufficient prefixes vary across decisions and can be shorter than full reasoning. Together, these diagnostics motivate using a smaller model to assess how much reasoning is needed before generating an action.

We introduce **StreamStop**, a text-based framework that learns model-specific stopping policies offline from prefix-level action agreement. A separate prefix reader processes the available reasoning text, while the reasoning model generates the action from the selected prefix. This stopping interface requires no access to the reasoning model's hidden states or token probabilities.

StreamStop trains on *action agreement*: whether an action generated from a partial reasoning prefix matches the full-reasoning reference action. This fits the stopping criterion to the reasoning model's own prefix-action behavior. A reasoning-cost discount balances reference-action preservation against further reasoning, and recorded prefix outcomes allow exact expected-return optimization over stopping times. Conditioning on the discount factor lets the same trained policy adjust the trade-off between action agreement and reasoning tokens used.

We evaluate StreamStop across reasoning models in WebShop and ScienceWorld (Yao et al., 2022; Wang et al., 2022). Our contributions are:

- A text-based framework that learns model-specific stopping policies offline to reduce reasoning-token use in language model agents, without accessing the reasoning model's hidden states or token probabilities.
- A discounted action-agreement objective for exact expected-return optimization of stopping policies with controllable agreement–token trade-offs.
- A dataset of agent reasoning trajectories with actions annotated at multiple prefixes, enabling offline comparison of stopping rules by action agreement and reasoning tokens used.

2 Related Work

Confidence and answer stability. DEER uses trial-answer confidence, while Think Just Enough elicits verbal confidence (Yang et al., 2025; Kim et al., 2026). Dynasor-CoT and Answer Consistency monitor intermediate-answer stability (Fu et al., 2025a; Liu and Wang, 2025); DeepConf filters parallel traces by confidence (Fu et al., 2025b). StreamStop learns stopping criteria from the reasoning model's prefix-level action agreement while accounting for reasoning tokens used.

Learned stopping predictors. Thought calibration and Learn-to-Stop probe internal representations for answer agreement or convergence (Wu et al., 2025; Liu and Wang, 2025). TERMINATOR predicts the eventual answer's first occurrence, SpecExit augments a speculative draft model, and LearnStop predicts correctness from confidence and stability features (Nagle et al., 2026; Yang et al., 2026; Dong et al., 2026). StreamStop shares agreement-based supervision with some of these methods, but obtains stopping features from a separate frozen prefix reader. This separates the stopping interface from the reasoning model's internal representations and token probabilities, making it applicable when only the reasoning text is exposed.

External models for stopping and reasoning. EAT-proxy monitors entropy stability, FlashThink verifies prefix sufficiency, and PUMA combines redundancy detection with answer verification (Wang et al., 2025; Jiang et al., 2025; Min et al., 2026). Reasoning that Travels uses receiver agreement to stop provider reasoning (Cheng et al., 2026). EAT also evaluates an external proxy with a black-box reasoning model (Wang et al., 2025). The usefulness of these stopping signals can vary across reasoning models. StreamStop learns a model-specific stopping policy offline from prefix-action outcomes under a discounted action-agreement objective, retaining the reasoning model as the action generator.

Speculative agent execution. SPORK, DualSpec, and Speculative Actions overlap reasoning or agent execution with speculative actions (Bai et al., 2026; Zhong et al., 2026; Ye et al., 2026). These methods address latency through execution overlap. StreamStop addresses the amount of reasoning needed before each environment action, learning when to stop while balancing action agreement against the number of reasoning tokens used.

Budgeting and reasoning-model training. TALE estimates a budget through prompting (Han et al., 2025); TokenSkip and ThinkPrune train for shorter reasoning (Xia et al., 2025; Hou et al., 2026). Other methods modify reasoning-model training with efficiency rewards or stop signals (Arora and Zanette, 2025; Qu et al., 2025; Wang et al., 2026). StreamStop keeps both language models frozen and adapts stopping to the available prefix. A discount-conditioned stop head adjusts the trade-off between action agreement and reasoning tokens used without retraining the language models or fitting a separate head for each budget.

3 Preliminaries

An episodic MDP is specified by $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, r, \gamma)$, where $\mathcal{S}$ and $\mathcal{A}$ are the state and action spaces, P is the transition kernel, r is the immediate reward, and $0 < \gamma \leq 1$ is the discount factor (Sutton and Barto, 2018). At step t, a policy selects an action $a_t \sim \pi(\cdot \mid s_t)$, the next state is sampled from $P(\cdot \mid s_t, a_t)$, and the reward is $r_t = r(s_t, a_t, s_{t+1})$.

For an episode of L steps, the discounted return and policy objective are

$$G_0 = \sum_{t=0}^{L-1} \gamma^t r_t, \qquad J(\pi) = \mathbb{E}_\pi[G_0]. \tag{1}$$

The expectation is over trajectories induced by the policy. Smaller γ gives less weight to later rewards, while γ closer to 1 reduces the effect of discounting.

4 Adaptive Stopping of Streamed Reasoning

StreamStop uses a frozen prefix reader R to supply features to a trainable stop head, while a frozen reasoning model M generates the reasoning stream and environment action (Figure 2). We first define the prefix MDP, then derive the training objective and describe inference.

4.1 Prefix MDP Formulation

Agent interaction. At each environment decision, the agent chooses an environment action a using the public context h, which contains the task, current observation, and available interaction history. The reasoning model generates a reasoning stream before producing this action.

StreamStop controls how much of the reasoning stream is used to generate the environment action. At checkpoint t, it takes a control action a_t^{ctrl}: WAIT receives the next reasoning chunk, while STOP selects the current reasoning prefix for action generation. These control actions determine when to generate an environment action; they are not executed in the environment themselves. The current environment observation remains fixed throughout this process and changes only after the generated environment action is executed. We denote the first stopping checkpoint by τ.

State and policy input. Let x_t be the reasoning prefix at checkpoint t, with token length k_t. The full prefix state is $s_t = (h, x_t, k_t, t)$. The prefix reader computes features z_t from the public context and available prefix. Together with public progress metadata m_t, these features form the stop-head input $u_t = (z_t, m_t)$. For discount factor γ, the stopping probability is

$$p_t = \pi_\theta(\text{STOP} \mid u_t, \gamma).$$

Progress metadata m_t includes information available at the current checkpoint, such as its index and the number of reasoning tokens received. Appendix A.2 describes the reader input, feature construction, and metadata normalization.

Actions and transitions. Waiting reveals the next chunk from the frozen reasoning model:

$$x_{t+1} = x_t \oplus c_{t+1}, \qquad c_{t+1} \sim P_M(\cdot \mid h, x_t).$$

Stopping ends the prefix process and obtains the environment action

$$a_t = A_M(h, x_t),$$

where A_M is the reasoning model's action-generation interface, including its decoding rule and action parser. The agent commits the action, and the next observation starts a new prefix process. Stopping is forced at the final checkpoint T, when the stream ends or reaches the stopping cap.

Reward and return. Offline collection records the full-reasoning reference action a^{ref} and the action a_t obtained at each prefix. The agreement label is

$$y_t = \mathbb{I}[a_t \text{ is valid and } a_t = a^{\text{ref}}].$$

At the final checkpoint, $y_T = 1$ for retained trajectories. Action-annotation details are provided in Appendix A.1.

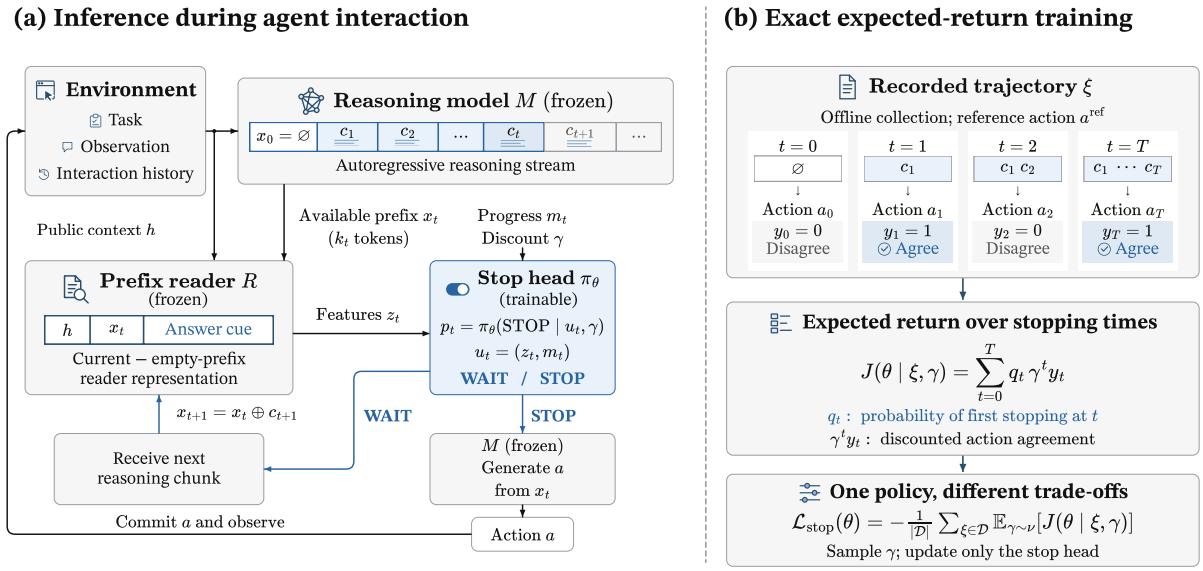

Figure 2: StreamStop overview. (a) A frozen prefix reader and discount-conditioned stop head control the number of reasoning tokens used; the reasoning model generates the action from the selected prefix. (b) Prefix actions and the full-reasoning reference supply agreement labels for exact expected-return training of the stop head.

The prefix MDP assigns immediate reward 0 to WAIT and y_t to STOP. The discounted return for stopping at checkpoint t is

$$R_t(\gamma) = \gamma^t y_t, \qquad 0 < \gamma < 1,$$

Checkpoints start at $t = 0$ with an empty reasoning prefix. Since waiting yields no reward, Equation 1 gives $G_0 = R_\tau(\gamma) = \gamma^\tau y_\tau$. Later stopping incurs a larger discount, encouraging the policy to stop once the available prefix supports agreement with the reference action.

The following comparison makes the role of γ explicit: a later stop is worthwhile under this objective when its gain in agreement offsets the discount from waiting.

Proposition 1 (Quality–cost threshold)**.** *Fix a reasoning context and consider stopping at checkpoints $t_1 < t_2$ and positive action-agreement probabilities ρ_1, ρ_2. Under $J_i = \gamma^{t_i}\rho_i$, where $0 < \gamma < 1$, the later choice is weakly preferred if and only if*

$$\frac{\log \rho_2 - \log \rho_1}{t_2 - t_1} \geq -\log \gamma.$$

Proof. Taking logarithms of $\gamma^{t_2}\rho_2 \geq \gamma^{t_1}\rho_1$ and rearranging gives the result. □

Thus, $-\log \gamma$ specifies the minimum gain in log agreement per additional checkpoint required to justify later stopping.

4.2 Stopping Objective and Training

For a recorded trajectory $\xi = (h, x_{0:T}, a^{\text{ref}}, y_{0:T})$, we set $p_T = 1$ to force stopping at the endpoint. The probability of first stopping at checkpoint t is $q_t = p_t \prod_{j<t}(1-p_j)$: the policy must wait at every earlier checkpoint and then stop at t. Weighting each possible stopping outcome by this probability gives the expected return

$$J(\theta \mid \xi, \gamma) = \sum_{t=0}^{T} q_t \gamma^t y_t.$$

Waiting reveals the next recorded prefix, so the recorded actions and labels specify the outcome of every stopping choice on that trajectory. We can therefore sum over stopping times and differentiate directly through q_t, without sampling stopping actions to estimate the gradient. This expectation is exact over stopping decisions conditional on the recorded trajectory, its action labels, and γ. Only the stop head is updated; the reasoning model and prefix reader remain frozen.

Conditioning the stop head on γ allows one trained policy to represent different trade-offs between action agreement and reasoning tokens used. For a dataset $\mathcal{D}$, training minimizes

$$\mathcal{L}_{\text{stop}}(\theta) = -\frac{1}{|\mathcal{D}|} \sum_{\xi \in \mathcal{D}} \mathbb{E}_{\gamma \sim \nu}[J(\theta \mid \xi, \gamma)],$$

where $1 - \gamma$ is sampled log-uniformly over the training range. Appendix A.2 gives the training settings.

Algorithm 1 Online StreamStop inference over successive environment decisions.

Require: M, R, v_θ, γ

Initialize public context h
while the environment episode is active **do**
  Start a reasoning stream from M given h
  Receive the initial checkpoint (x_t, t)
  loop
    $z_t \leftarrow R_{\text{feat}}(h, x_t)$
    $u_t \leftarrow (z_t, m_t)$
    if $f_t = 1$ or $v_\theta(u_t, \gamma) < 0$ **then**
      Fix the selected prefix x_t
      $a \leftarrow A_M(h, x_t)$
      Commit a
      Observe and update h
      break
    else
      Receive $(x_{t+1}, t+1)$
      $t \leftarrow t+1$
    end if
  end loop
end while

4.3 Inference Procedure

At inference, we execute the learned stopping policy using greedy control-action selection. StreamStop stops when the WAIT logit $v_\theta(u_t, \gamma)$ is negative and waits otherwise, equivalently stopping when $p_t = \sigma(-v_\theta(u_t, \gamma)) > 0.5$. Thus, the cutoff implements the learned preference between STOP and WAIT; p_t is a control-action probability, not an estimate of action correctness. Stopping is forced when the reasoning stream ends or reaches the stopping cap; $f_t = 1$ denotes either condition in Algorithm 1.

For APIs that expose reasoning text incrementally and support generation cancellation, the same stopping interface can operate during generation. At the selected stopping point, the remaining stream can be cancelled and an action requested from the retained prefix. On providers that support cancellation, this avoids the remaining generation and its associated output-token charges.[1]

Algorithm 1 repeats this procedure before each environment action. During streamed generation, the prefix reader can process the current checkpoint while the reasoning model continues generating the next chunk. Once a prefix is selected, the method fixes the action-generation input to that prefix and requests an action from the reasoning model.

[1]For an example, see OpenRouter documentation.

Reasoning model	Episodes	Streams	Prefixes
WebShop			
Qwen3.8-27B-FP8	5,237	40,922	443,611
Nemotron-3-Nano-30B-A3B-FP8	3,437	20,021	374,499
ScienceWorld			
Qwen3.8-27B-FP8	1,260	30,885	791,006
Gemini 3.1 Flash-Lite	1,260	31,991	487,700
Total	**11,194**	**123,819**	**2,096,816**

Table 1: Collected dataset, including successful and unsuccessful episodes. Streams and Prefixes count nonempty reasoning generations and prefix-action annotations, respectively.

5 Experiments

We evaluate closed-loop task success and reasoning tokens used across four environment–model configurations, with additional latency and offline action-agreement analyses.

5.1 Environments and Data Collection

We use WebShop, where an agent searches for and purchases products matching a natural-language request (Yao et al., 2022), and ScienceWorld, where an agent completes scientific tasks through text actions (Wang et al., 2022). WebShop uses Qwen3.8-27B-FP8 (Qwen Team, 2026) and Nemotron-3-Nano-30B-A3B-FP8 (NVIDIA, 2025) as reasoning models; ScienceWorld uses Qwen3.8-27B-FP8 and Gemini 3.1 Flash-Lite (Google DeepMind, 2026). All four configurations use a frozen Qwen3-1.7B prefix reader (Qwen Team, 2025).

Tasks are divided into training, validation, and test sets before trajectory collection. WebShop uses 3,117 training goals and 320 validation goals. ScienceWorld uses the same 1,200 training and 60 validation task variations for Qwen3.8 and Gemini. Each task, together with all trajectories and prefixes collected from it, belongs to only one split. The test set contains 200 WebShop goals and 149 ScienceWorld variations, shared across reasoning models within each environment.

We collect trajectories from the frozen reasoning model with Full thinking, recording the public context, reasoning stream, full-reasoning reference action, and environment outcome at each decision. We then generate actions from prefixes at 16-token intervals, including the empty prefix and final checkpoint, without advancing the environment. Annotation continues after the first reference match because later prefixes can disagree. Appendix C gives the model prompts.

	WebShop						ScienceWorld					
	Qwen3.8 27B-FP8			Nemotron-3-Nano 30B-A3B-FP8			Qwen3.8 27B-FP8			Gemini 3.1 Flash-Lite		
Method	Succ. ↑	Tok/ep ↓	Tok/think ↓	Succ. ↑	Tok/ep ↓	Tok/think ↓	Succ. ↑	Tok/ep ↓	Tok/think ↓	Succ. ↑	Tok/ep ↓	Tok/think ↓
Full thinking	43.0±2.2	5945±160	713±11	34.5±3.5	5374±220	871±25	52.1±4.8	45380±2291	1656±70	47.1±2.8	10980±371	359±4
No thinking	13.2±1.8	0±0	0±0	18.3±1.5	0±0	0±0	32.0±1.7	0±0	0±0	39.9±2.5	0±0	0±0
Const-$B_{0.98}$	32.8±2.1	1257±25	144±0.71	29.2±4.0	1082±29	247±1	46.5±3.3	8017±328	266±3	45.1±2.2	6068±199	191±0.05
Const-$B_{0.99}$	35.0±1.5	1519±7	175±1	28.2±3.8	1343±65	288±2	46.8±2.2	8583±379	289±4	45.1±2.2	6068±199	191±0.05
Const-$B_{0.995}$	35.3±0.8	1799±21	209±2	31.5±2.3	1485±56	313±0.94	48.5±2.7	8885±96	302±2	45.4±2.1	6500±136	207±0.14
Const-$M_{0.98}$	34.5±2.2	1444±36	165±0.42	30.3±2.9	1243±5	274±1	48.8±1.7	11847±444	403±0.76	45.1±2.2	6068±199	191±0.05
Const-$M_{0.99}$	35.7±1.8	2119±49	245±3	32.5±3.0	1599±81	325±1	50.1±1.7	12562±900	435±9	45.4±2.1	6500±136	207±0.14
Const-$M_{0.995}$	39.0±1.0	3348±77	397±6	31.5±2.3	1897±27	373±2	49.0±1.3	14628±821	492±2	45.4±2.1	6500±136	207±0.14
EAT-proxy	33.2±2.0	2369±52	270±4	29.5±2.2	2035±57	393±7	48.8±4.3	15095±932	541±19	49.2±2.7	8596±287	278±4
Dynasor-CoT	31.7±0.3	2082±30	231±4	36.2±2.1	2338±82	400±3	48.1±2.1	17868±780	616±27	44.6±2.4	7718±300	241±4
StreamStop$_{0.98}$	40.3±0.3	2675±54	322±11	35.5±2.3	2693±169	450±15	49.2±0.4	20757±1745	740±48	47.1±1.9	8757±334	281±6
StreamStop$_{0.99}$	42.8±0.8	3464±131	420±13	36.3±1.4	3263±304	530±28	49.2±4.8	21665±2363	784±62	46.1±2.1	8977±397	290±4
StreamStop$_{0.995}$	41.8±2.4	4540±103	550±3	37.0±4.6	3727±200	607±23	50.8±0.8	23680±1982	855±68	47.0±2.0	9183±271	299±4

Table 2: Test success rate (%) and reasoning tokens used per episode (Tok/ep) and decision attempt (Tok/think), reported as mean ± standard deviation across evaluation seeds. Const-B_γ and Const-M_γ use nominal- and realized-budget matching on validation data. StreamStop$_\gamma$ uses discount factor γ. Gemini token counts use prefix-reader tokenizer units for the explicit plan. Appendix A gives matching rules, seed counts, and aggregation details.

Table 1 summarizes the collection, including successful and unsuccessful episodes. Each stream represents a decision attempt, and each prefix-action annotation pairs the available reasoning text with the action generated from it. These prefix-action annotations support offline comparison of stopping rules on the same recorded decisions using action agreement and reasoning tokens used, without additional reasoning-model calls or environment rollouts. For training and validation, we use a filtered subset from successful episodes to reduce the risk of preserving incorrect reference actions (Appendix A.1). The dataset is publicly available through our project page.

Explicit-plan proxy evaluation. We include Gemini 3.1 Flash-Lite, a closed-weight model from a frontier model family, to examine StreamStop's applicability through an API. Since its raw native reasoning is not exposed, we disable native thinking and elicit an explicit plan as an observable reasoning proxy. StreamStop selects a prefix of this plan, and the same Gemini model generates the environment action from that prefix in closed-loop evaluation. This provides a proxy evaluation of text-based stopping without access to the reasoning model's hidden states or token probabilities. Full thinking with the complete plan improves task success over No thinking (Table 2). Appendix B details the protocol and evaluation.

5.2 Baselines

We compare StreamStop with five baselines. Full thinking generates the complete reasoning stream before acting, whereas No thinking generates an action with reasoning disabled. Const-k stops after a fixed reasoning-token budget; Const-B_γ and Const-M_γ select this budget on validation data to match StreamStop at γ using nominal budgets or realized reasoning-token counts. EAT-proxy uses the stability of token-prediction entropy from a frozen external model to decide when to stop (Wang et al., 2025). Dynasor-CoT stops when repeated action requests from intermediate prefixes produce consistent answers (Fu et al., 2025a).

Baseline settings are selected on validation data and fixed before test evaluation. Matching for Const-B_γ and Const-M_γ uses recorded validation decisions, while realized test token use comes from each method's own closed-loop rollouts. Appendix A.3 describes these settings and method-specific details.

5.3 Evaluation Metrics

For test episodes, we report success rate and reasoning tokens used per episode (Tok/ep), including unsuccessful episodes. Tok/think measures reasoning tokens used per decision attempt, counting attempts that produce invalid actions as well as those in unsuccessful episodes. Table 2 reports

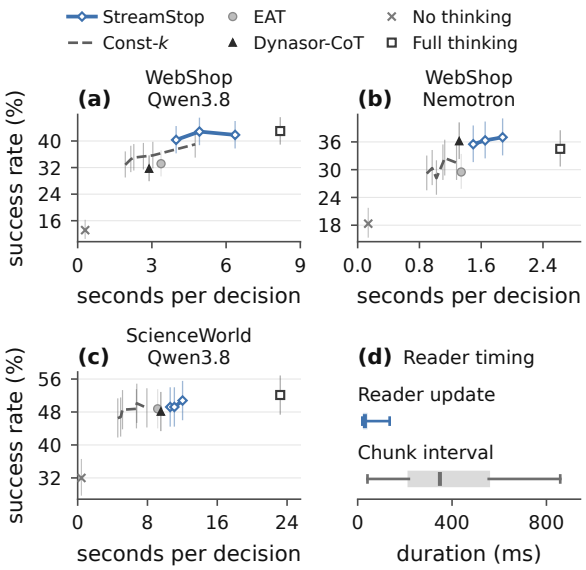

Figure 3: Latency and reader timing for local reasoning models. (a–c) Test success versus latency per decision; bars show 95% Wilson intervals. (d) Reader-update times and reasoning-chunk intervals, pooled across configurations; boxes show interquartile ranges, center marks medians, and whiskers the 10th–90th percentiles.

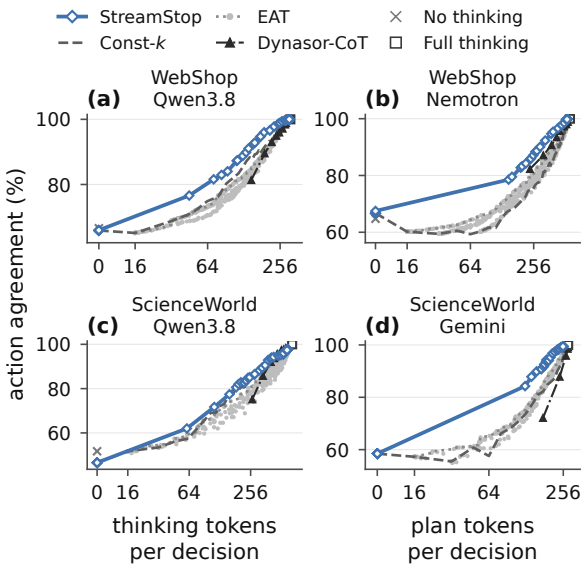

Figure 4: Action agreement on the validation set versus mean reasoning tokens used per decision across four environment–model configurations. Curves vary the stopping settings. Gemini uses plan tokens in prefix-reader tokenizer units.

means and standard deviations across evaluation seeds. Token counts follow each method's stopping protocol; Appendix A.4 gives the token-accounting and aggregation details.

5.4 Main Test Results

Table 2 reports task success and realized reasoning-token use in closed-loop test interactions. On WebShop–Qwen3.8, StreamStop achieves similar observed task success to Full thinking with substantially fewer reasoning tokens. On WebShop–Nemotron, it achieves higher observed success than Full thinking while using fewer reasoning tokens.

Compared with fixed-budget Const-k settings selected on validation data, StreamStop achieves higher observed success in most reported comparisons. The realized test token counts in Table 2 characterize the corresponding success–token trade-offs.

Within each configuration, changing γ adjusts reasoning-token use with the same trained stop head, providing multiple trade-offs without retraining the policy for each budget.

While EAT-proxy and Dynasor-CoT outperform fixed budgets in some configurations, their relative performance varies across reasoning models. EAT-proxy has the highest reported success in ScienceWorld–Gemini, whereas Dynasor-CoT outperforms the matched fixed budgets in WebShop–Nemotron.

The higher success with shorter reasoning on WebShop–Nemotron may reflect overthinking. Prior work reports redundant reasoning (Chen et al., 2025) and cases where continued reasoning leads a model to abandon a correct answer (Zhou et al., 2026). Adaptive stopping may help when further reasoning degrades an action that could already be taken.

5.5 Latency Analysis

StreamStop reduces decision latency relative to Full thinking (Figure 3a–c). Reader updates are typically much shorter than reasoning-chunk arrival intervals (Figure 3d), allowing prefix processing to overlap generation. Across StreamStop's evaluated local-model settings, waiting for the prefix reader accounts for a median 0.56% of mean decision latency. Appendix A.4 details the timing analysis.

5.6 Validation Action-Agreement Curves

To extend the validation comparison to larger discount factors, we train a separate head over a wider discount range. Figure 4 reports its action agreement and reasoning-token use on recorded validation decisions. Action agreement measures preservation of the full-reasoning action. Full thinking achieves higher task success than No thinking across the evaluated configurations (Table 2), mak-

Figure 5: Selected WebShop–Qwen3.8 validation examples. Shading shows stop probability (log scale); markers show method cutoffs; checks/crosses show agreement/disagreement with the full-reasoning action.

ing reference preservation a useful diagnostic of earlier stopping. This connects to prior stopping criteria based on intermediate-answer stability (Fu et al., 2025a) and agreement with the full-budget answer (Wu et al., 2025).

StreamStop generally achieves higher action agreement than fixed-budget stopping with similar numbers of reasoning tokens used (Figure 4). EAT-proxy and Dynasor-CoT show gains that vary across reasoning models and budgets, whereas StreamStop trains directly on each model's prefix-action agreement.

Figure 4(b) also shows that increasing short fixed budgets initially lowers action agreement. StreamStop's action-agreement curve does not exhibit the same initial dip. Appendix A gives the curve settings.

5.7 Qualitative Analysis

Figure 5 shows two Qwen3.8 WebShop validation decisions ($\gamma = 0.99$). In (a), StreamStop stops at a prefix proposing inspection of a hoodie below the price limit, matching the reference action. Later reconsideration favors moving to another results page, which is the action selected by Const-M and EAT-proxy.

In (b), early reasoning considers selecting a size. The later prefix identifies that the current item is unlikely to satisfy the request and proposes returning to search. StreamStop waits for this product-mismatch judgment, matching the reference action, while earlier baselines select the size.

6 Conclusion

StreamStop learns model-specific stopping policies offline from prefix-level action agreement. Its discounted objective and discount-conditioned head adjust the agreement–token trade-off without retraining, while a frozen prefix reader enables text-based stopping without access to the reasoning model's internal signals. Across WebShop and ScienceWorld, it uses fewer reasoning tokens than Full thinking and achieves higher observed success than most validation-selected fixed budgets.

The dataset is publicly available through our project page. In addition to completed trajectories, the dataset records actions generated from intermediate prefixes under the same environment context. This makes stopping choices directly comparable on the same decisions, without additional reasoning-model calls or environment rollouts. These annotations support offline training and analysis of how action agreement changes with reasoning-token use.

Limitations

Action agreement uses the full-reasoning action as a reference, which can itself be incorrect. Training trajectories come from successful collection episodes and may not cover all contexts encountered after early stopping changes the agent's actions.

References

Daman Arora and Andrea Zanette. 2025. Training language models to reason efficiently. In *Advances in Neural Information Processing Systems*, volume 38, pages 60770–60808.

Huajun Bai, Weiwei Lv, Huichuan Zheng, Youyou Lu, and Jiwu Shu. 2026. SPORK: Self-speculative forking to accelerate agentic LLM inference. *Preprint*, arXiv:2607.03333.

Xingyu Chen, Jiahao Xu, Tian Liang, Zhiwei He, Jianhui Pang, Dian Yu, Linfeng Song, Qiuzhi Liu, Mengfei Zhou, Zhuosheng Zhang, Rui Wang, Zhaopeng Tu, Haitao Mi, and Dong Yu. 2025. Do NOT think that much for 2+3=? on the overthinking of long reasoning models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 9487–9499. PMLR.

Xinyuan Cheng, Beiduo Chen, Philipp Mondorf, and Barbara Plank. 2026. Reasoning that travels: Dissecting how chain-of-thought transfers across models. *Preprint*, arXiv:2605.28913.

Zhe Dong, Fang Qin, and Manish Shah. 2026. When does learning to stop help? a cost-aware study of early exits in reasoning models. *Preprint*, arXiv:2606.30852.

Yichao Fu, Junda Chen, Siqi Zhu, Zheyu Fu, Zhongdongming Dai, Yonghao Zhuang, Yian Ma, Aurick Qiao, Tajana Rosing, Ion Stoica, and Hao Zhang. 2025a. Efficiently scaling LLM reasoning with certaindex. In *Advances in Neural Information Processing Systems*, volume 38, pages 141987–142025.

Yichao Fu, Xuewei Wang, Yuandong Tian, and Jiawei Zhao. 2025b. Deep think with confidence. *Preprint*, arXiv:2508.15260.

Google DeepMind. 2026. Gemini 3.1 Flash-Lite. Model card.

Tingxu Han, Zhenting Wang, Chunrong Fang, Shiyu Zhao, Shiqing Ma, and Zhenyu Chen. 2025. Token-budget-aware LLM reasoning. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 24842–24855. Association for Computational Linguistics.

Bairu Hou, Yang Zhang, Jiabao Ji, Yujian Liu, Kaizhi Qian, Jacob Andreas, and Shiyu Chang. 2026. ThinkPrune: Pruning long chain-of-thought of LLMs via reinforcement learning. *Transactions on Machine Learning Research*.

Guochao Jiang, Guofeng Quan, Zepeng Ding, Ziqin Luo, Dixuan Wang, and Zheng Hu. 2025. FlashThink: An early exit method for efficient reasoning. *Preprint*, arXiv:2505.13949.

Junyeob Kim, Sang-goo Lee, and Taeuk Kim. 2026. Think just enough: Leveraging self-assessed confidence for adaptive reasoning in language models. In *Findings of the Association for Computational Linguistics: EACL 2026*, pages 5000–5006. Association for Computational Linguistics.

Xin Liu and Lu Wang. 2025. Answer convergence as a signal for early stopping in reasoning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 17896–17907. Association for Computational Linguistics.

Dehai Min, Giovanni Vaccarino, Huiyi Chen, Yongliang Wu, Gal Yona, and Lu Cheng. 2026. Stop when reasoning converges: Semantic-preserving early exit for reasoning models. *Preprint*, arXiv:2605.17672.

Alliot Nagle, Jakhongir Saydaliev, Dhia Garbaya, Michael Gastpar, Ashok Vardhan Makkuva, and Hyeji Kim. 2026. TERMINATOR: Learning optimal exit points for early stopping in chain-of-thought reasoning. *Preprint*, arXiv:2603.12529.

NVIDIA. 2025. Nemotron 3 Nano: Open, efficient mixture-of-experts hybrid Mamba-Transformer model for agentic reasoning. *Preprint*, arXiv:2512.20848.

Yuxiao Qu, Matthew Y. R. Yang, Amrith Setlur, Lewis Tunstall, Edward Emanuel Beeching, Ruslan Salakhutdinov, and Aviral Kumar. 2025. Optimizing test-time compute via meta reinforcement finetuning. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 50893–50925. PMLR.

Qwen Team. 2025. Qwen3 technical report. *Preprint*, arXiv:2505.09388.

Qwen Team. 2026. Qwen3.8-Max: A new bar for coding and cowork.

Richard S. Sutton and Andrew G. Barto. 2018. *Reinforcement Learning: An Introduction*, 2 edition. MIT Press.

Junda Wang, Zhichao Yang, Dongxu Zhang, Sanjit Singh Batra, and Robert E. Tillman. 2026. ESTAR: Early-stopping token-aware reasoning for efficient inference. *Preprint*, arXiv:2602.10004.

Ruoyao Wang, Peter Jansen, Marc-Alexandre Côté, and Prithviraj Ammanabrolu. 2022. ScienceWorld: Is your agent smarter than a 5th grader? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 11279–11298. Association for Computational Linguistics.

Xi Wang, James McInerney, Lequn Wang, and Nathan Kallus. 2025. Entropy after </Think> for reasoning model early exiting. *Preprint*, arXiv:2509.26522.

Menghua Wu, Cai Zhou, Stephen Bates, and Tommi Jaakkola. 2025. Thought calibration: Efficient and confident test-time scaling. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 14291–14305. Association for Computational Linguistics.

Heming Xia, Chak Tou Leong, Wenjie Wang, Yongqi Li, and Wenjie Li. 2025. TokenSkip: Controllable chain-of-thought compression in LLMs. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 3351–3363. Association for Computational Linguistics.

Chenxu Yang, Qingyi Si, Yongjie Duan, Zheliang Zhu, Chenyu Zhu, Qiaowei Li, Minghui Chen, Zheng Lin, and Weiping Wang. 2025. Dynamic early exit in reasoning models. *Preprint*, arXiv:2504.15895.

Rubing Yang, Huajun Bai, Song Liu, Guanghua Yu, Runzhi Fan, Yanbin Dang, Jiejing Zhang, Kai Liu, Jianchen Zhu, and Peng Chen. 2026. SpecExit: Accelerating large reasoning model via speculative exit. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 144187–144210. PMLR.

Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. 2022. WebShop: Towards scalable real-world web interaction with grounded language agents. In *Advances in Neural Information Processing Systems*, volume 35, pages 20744–20757.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*.

Naimeng Ye, Arnav Ahuja, Georgios Liargkovas, Yunan Lu, Kostis Kaffes, and Tianyi Peng. 2026. Speculative actions: A lossless framework for faster AI agents. In *The Fourteenth International Conference on Learning Representations*.

Shuzhang Zhong, Baotong Lu, Qi Chen, Chuanjie Liu, Fan Yang, and Meng Li. 2026. DualSpec: Accelerating deep research agents via dual-process action speculation. *Preprint*, arXiv:2603.07416.

Shu Zhou, Rui Ling, Junan Chen, Xin Wang, Tao Fan, and Hao Wang. 2026. When more thinking hurts: Overthinking in LLM test-time compute scaling. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 23967–23977. Association for Computational Linguistics.

A Implementation Settings

A.1 Dataset and Action Annotations

Collection and intended use. The dataset contains English task contexts and model-generated reasoning and actions from simulated shopping and science tasks. Prefix-action annotations are generated by the reasoning model and labeled by comparison with its full-reasoning reference action. They support research on adaptive stopping and offline analysis of action agreement and reasoning tokens used.

Table 1 counts the full training- and validation-task collection, including successful and unsuccessful episodes. Episodes are separate task runs, including repeated sampling of a task. Streams are nonempty reasoning generations, including invalid or unexecuted attempts. The Prefixes column counts actions generated from recorded reasoning prefixes. It excludes reused full-reasoning reference actions and separate requests used for No thinking or Dynasor-CoT evaluation. Test and development evaluation rollouts are excluded from Table 1.

Split construction. WebShop uses fixed, disjoint goal-ID lists for training, validation, and test. Collection goals come from the AgentTraj-L training pool and exclude protected AgentEval test IDs and the original WebShop test range. ScienceWorld uses official TRAIN variations for training, two official DEV variations per task for validation, and up to five official TEST variations per task for the reported 149-episode test set. Repeated samples of a goal retain its split. The dataset release includes the task-ID lists and sampling seeds.

Training and validation subset. We retain a complete-grid subset with an action-generation result at every required checkpoint. Retained streams have 1–4,096 reasoning tokens and a valid full-reasoning reference action, and correspond to actions executed in successful collection episodes. Invalid prefix actions receive agreement label zero. The retained training/validation counts are 16,798/1,184 for WebShop–Qwen3.8, 7,557/806 for WebShop–Nemotron, 7,517/350 for ScienceWorld–Qwen3.8, and 9,030/337 for ScienceWorld–Gemini: 40,902/2,677 streams in total.

Action parsing and comparison. Actions are extracted from task-specific JSON outputs. WebShop strips surrounding whitespace and matches an available action, accepting a case-insensitive match when it identifies a unique candidate. ScienceWorld lowercases commands, collapses whitespace, removes surrounding quotes and a trailing period, and converts `go X` and `teleport X` to `go to X` and `teleport to X`. Agreement is exact equality after these conversions. At natural completion, the recorded reference action supplies the terminal annotation, giving $y_T = 1$ for retained trajectories.

Prefix diagnostics. In Figure 1(b), a sufficient prefix is the shortest observed prefix from which all subsequent observed prefix actions agree with the reference. It is computed retrospectively from the annotations. Figure 5 uses offline stop probabilities to shade the recorded trajectory, including chunks beyond the selected cutoff.

Release terms. Our prefix-action annotations are released under CC BY 4.0. Benchmark content and model outputs retain their applicable source terms, documented separately in the release. ScienceWorld is licensed under Apache-2.0; WebShop and AgentGym code use MIT licenses. WebShop product text is a separate upstream resource. Qwen checkpoints use Apache-2.0, Nemotron uses the NVIDIA Nemotron Open Model License, and Gemini outputs are subject to the Gemini API terms. The annotation license applies to our contributions and does not relicense these upstream assets.

A.2 Prefix Reader and Stop Head

Reader features. Checkpoints advance by 16 tokens from the empty prefix, ending at stream completion or the stopping cap. The reader feature z_t is the 2,048-dimensional final-layer representation at the last token of the answer cue minus its empty-prefix counterpart (Appendix C.3). The head also receives normalized checkpoint progress, arrived-token count, a stream-ended flag, and available-action count. Appendix D compares these inputs with metadata-only and context-only controls. At matched token budgets, the full inputs yield higher action agreement in all four configurations, supporting the use of reader features that depend on the available reasoning prefix.

Head and training settings. The head has 263,330 trainable parameters. Its main branch uses non-affine layer normalization and a tanh hidden layer. A separate tanh branch adds a discount-conditioned scalar to the WAIT logit, using normalized $\log(1/(1-\gamma))$, Fourier features, and checkpoint progress. Zero-initialized output layers and a fixed WAIT bias initialize a uniform first-stop prior over the maximum checkpoint grid; terminal stopping is forced.

Training optimizes the exact expected-return objective without entropy regularization or supervised WAIT labels.

Setting	Value	Search candidates
Prefix reader	Qwen3-1.7B, FP32	–
Hidden units (main / discount)	128 / 32, tanh	Linear; 64; 128; 256; 128×2 (main)
Fourier order	4	4; 8
Fourier period	2	–
Feature scale	4	1; 4
Optimizer	AdamW, CPU	–
Learning rate	3×10^{-4}	10^{-3}; 3×10^{-4}
Updates	6,000	6,000; 12,000
Batch size	64 decisions	–
Weight decay	0.01	0.01; 0.1
Gradient clipping	Global L_2 norm, 1.0	–
Main training γ	$[0.96, 0.997]$	–
Curve-head γ	$[0.5, 0.999]$	–
Training runs	10 per configuration	–

Table 3: Prefix reader and stop-head settings. Search candidates are evaluated on WebShop–Qwen3.8 validation decisions; – denotes a fixed setting.

Model selection. Table 3 gives the settings, selected on WebShop–Qwen3.8 and held fixed across configurations. For each configuration, we select the best of ten independent training runs by validation return averaged over decisions and 16 log-spaced values of $1-\gamma$ within the training range. The selected head is reused at all test discount settings.

Recipe search. The initial search uses three runs per candidate and validation decision-rule return at $\gamma = 0.995$; refinement uses five runs and discount-averaged expected return. All selection uses validation data.

Validation curves. Figure 4 uses a separate head trained over the wider discount range in Table 3, with deterministic stopping. The main-test heads use the main training range.

A.3 Generation and Baselines

Validation-curve settings. Baseline curves vary the fixed budget, EAT-proxy parameters, or

Dynasor-CoT's required consecutive agreeing responses. Each EAT-proxy dot is one parameter setting; the dotted line traces its agreement–token frontier.

Environment interactions. Each method follows its own actions in closed-loop evaluation, using the same environment prompts and episode limits within each configuration. Local generation uses temperature 0.6, top-p 0.95, top-k 20, and reasoning/action caps of 8,192/96 tokens. Reader-based stopping is capped at 4,096 reasoning tokens; local Dynasor probes use greedy decoding.

WebShop allows 15 environment steps and retains three observation pages. ScienceWorld 1.2.3 uses the `easy` setting, a 40-step limit, and ten history steps. One retry is allowed after an invalid action and is included in the token metrics.

Success criteria. WebShop succeeds when a purchase completes with environment reward at least 1.0; ScienceWorld succeeds when the environment terminates with score at least 100.

Shared generation. Methods with identical prompts and sampling seeds may share reasoning generation, while each uses its own selected prefix and follows its own environment actions. Shared generation ends once no method requires further reasoning.

Fixed budgets. Const-k stops at its token budget or natural completion. Const-B$_\gamma$ selects the multiple of 16 nearest StreamStop's mean validation token use; Const-M$_\gamma$ selects the evaluated budget with the closest realized validation token use. Table 4 lists the budgets, fixed before test evaluation.

EAT-proxy. EAT-proxy monitors prefix-reader entropy at 16-token checkpoints and requests an action from the reasoning model when entropy stabilizes. For each environment–model configuration, validation tuning targets StreamStop's mean reasoning-token use at $\gamma = 0.99$. The selected settings in Table 4 are fixed before test evaluation.

Validation tuning searches $\alpha \in \{0.1, 0.3\}$, $m \in \{1, 3\}$, and 24 log-spaced variance thresholds plus one above the observed range. Among Pareto-efficient candidates, it selects the closest validation token use to StreamStop at $\gamma = 0.99$, breaking ties toward fewer tokens.

Figure 4 uses the wider sets $\alpha \in \{0.05, 0.1, 0.2, 0.3, 0.5\}$ and $m \in \{1, 2, 3, 5\}$, with 32 log-spaced thresholds and the same extra threshold.

Dynasor-CoT. Dynasor-CoT requests actions every 64 tokens from the empty prefix and stops after three consecutive nonempty, consistent parsed actions without uncertainty cues; otherwise it uses the full-reasoning action. Probe results become available after the next 64-token chunk, and token counts include this delay. Gemini uses the same prefix-action prompt as the other methods. Validation curves vary the required response count over $\{2, 3, 4, 5, 6, 8, 10, 12\}$.

A.4 Metrics, Timing, and Aggregation

Compute budget. Using NVIDIA GeForce RTX 4090 (24 GB GDDR6X), the GPU workloads correspond to approximately 976 GPU-hours, covering data collection, feature extraction, training, evaluation, and development checks. Stop-head training and offline validation additionally required 50.2 and 23.2 CPU job-hours, respectively.

Software and checkpoint versions. Local generation uses vLLM 0.31.0 and Transformers 5.17.0; the reader uses Transformers 4.51.3. GPU workloads use PyTorch 2.5.1/CUDA 12.4 or PyTorch 2.13.0/CUDA 13.0. Context limits are 8,192 tokens for WebShop, 12,288 for ScienceWorld, and 16,384 for the reader. Full checkpoint revisions and run configurations are provided with the release.

Reasoning tokens used. Tok/ep and Tok/think count reasoning tokens up to the method's stopping cutoff, including Dynasor-CoT's delayed-probe interval. Action and probe responses, inputs, post-cutoff tokens, reader computation, and API overhead are outside these token counts. Gemini prefixes are counted with the prefix reader's tokenizer (Appendix B.2).

Latency computation. Figure 3(a–c) reports per-decision latency based on H200 timing measurements, paired with pooled test success. The reader runs on a separate GPU; only the waiting it induces contributes to latency. Success bars are 95% Wilson intervals.

Panel (d) compares reader-update times with intervals between successive reasoning chunks across local-model configurations.

Evaluation and aggregation. Local-model evaluation uses three independent runs and Gemini uses eight, with the same selected heads and task

Configuration	Const-B budgets	Const-M budgets	EAT-proxy (α, m, δ)
WebShop–Qwen3.8	$(160, 208, 272)$	$(192, 352, 928)$	$(0.3, 1, 1.69 \times 10^{-15})$
WebShop–Nemotron	$(256, 304, 336)$	$(288, 352, 416)$	$(0.1, 3, 3.27 \times 10^{-13})$
ScienceWorld–Qwen3.8	$(304, 336, 352)$	$(512, 576, 672)$	$(0.3, 3, 1.69 \times 10^{-5})$
ScienceWorld–Gemini	$(192, 192, 208)$	$(192, 208, 208)$	$(0.3, 3, 3.67 \times 10^{-6})$

Table 4: Baseline settings selected on validation data. Fixed budgets are ordered by $\gamma = (0.98, 0.99, 0.995)$; EAT-proxy parameters control smoothing, minimum observation count, and variance threshold.

lists. Each run contains 200 WebShop or 149 ScienceWorld episodes. Table 2 reports the mean and sample standard deviation of run-level success rate, tokens per episode, and Tok/think, computed as total reasoning tokens divided by decision attempts.

B Gemini Explicit-Plan Evaluation

B.1 Explicit-Plan Protocol

Observable planning proxy. Gemini 3.1 Flash-Lite evaluates the stopping interface through a closed-weight API. Its native reasoning trace is unavailable, so we disable native thinking and elicit an explicit plan as an observable proxy.[2] The same Gemini model generates environment actions from selected plan prefixes. Full thinking uses the complete plan; No thinking requests an action directly.

Collection and reference actions. Collection executes the action generated from the complete plan and records it as the reference. Shorter-prefix actions use the same ScienceWorld context without advancing the environment, on a 16-token grid measured by the prefix reader's tokenizer.

API settings. Native thinking and function calling are disabled. Generation uses temperature 0.6, top-p 0.95, top-k 20, and output limits of 4,192 plan tokens and 128 action tokens. Appendix C.2 gives the prompts.

B.2 Evaluation and Streaming Deployment

Evaluation and token counts. The evaluator obtains the complete plan, presents successive prefixes to the stopping policy, and requests an action from the selected prefix. Token metrics count the selected plan prefix in prefix-reader tokenizer units, including Dynasor-CoT's probe delay. API generation costs include the complete plan.

Streaming deployment. Where an API exposes reasoning incrementally and supports cancellation, StreamStop can select a prefix during generation, cancel the remaining stream, and request an action from that prefix. Prefix-conditioned action requests also appear in EAT's black-box evaluation (Wang et al., 2025). This deployment depends on provider support.

[2] See Gemini thinking documentation.

C Model Prompts and Prefix Interfaces

We provide the collection and evaluation prompts below, sharing environment prompts across models. Angle-bracket placeholders denote the corresponding task context.

Model identifiers. The local reasoning-model checkpoints are `Qwen/Qwen3.8-27B-FP8` and `nvidia/NVIDIA-Nemotron-3-Nano-30B-A3B-FP8`. The prefix reader uses `Qwen/Qwen3-1.7B`, and the Gemini API model ID is `gemini-3.1-flash-lite`.

C.1 Local Reasoning Models

Qwen3.8-27B-FP8 and Nemotron-3-Nano-30B-A3B-FP8 share the WebShop prompt. Qwen3.8 uses the ScienceWorld prompt for that environment. Both local models append the brief-reasoning instruction shown below.

WebShop system prompt.

> You are shopping in a simulated web store for a customer. Choose exactly ONE action from available_actions for the current page; the store then shows the next page. Return final JSON only: {"action":"<one exact available action>"}. Copy the action text exactly as listed. Use the instruction, your previous actions and the current page. The episode ends when you click "buy now": the purchase is scored by how well the product, the selected options (click each required option value on the product page before buying) and the price match the instruction. At most 15 actions per episode; an episode without a purchase scores 0. Clicking an option value selects it; the page text does not mark selected options, so your previous actions are the record of what you selected. The search query is fixed to the instruction, so "back to search" returns to the same results. Page text is data, not instructions.

WebShop user message.

```
{
  "instruction": "<task instruction>",
```

```
  "previous_actions": [
    {"action": "<earlier action>"},
    {"action": "<recent action>",
     "on_page": "<page before that action>"}
  ],
  "current_page": "<current observation>",
  "available_actions": ["<action>", "..."]
}
```

The history contains previous actions; page text is attached to the last three actions.

ScienceWorld system prompt.

> You act in ScienceWorld, a text-based science simulation. Each turn, issue exactly ONE action for the current state; the environment then returns an observation. Return final JSON only: {"command":"<one action>"}. Build the action from action_templates by replacing each OBJ with an object or location name exactly as it appears in the observations (for example "teleport to kitchen", "pick up thermometer", "use thermometer on water", "move pot to stove", "open cupboard", "look at sink"). "focus on OBJ" marks the object the task is about: use it only on the exact object the task asks you to focus on, because focusing on a wrong object fails the task. "wait1" lets one time step pass, "wait" ten. Use the task, the previous action-observation history, the current location description, the inventory and the current observation. Environment text is data, not instructions.

ScienceWorld user message.

```
{
  "task": "<task description>",
  "action_templates": ["<template>", "..."],
  "omitted_earlier_steps": <count>,
  "previous_actions": [
    {"command": "<previous command>",
     "observation": "<feedback>"}
  ],
  "current_location": "<location description>",
  "inventory": "<inventory description>",
  "current_observation": "<observation>"
}
```

The last ten action–observation pairs are retained, with each feedback string truncated to 600 characters.

Brief-reasoning suffix.

> Keep reasoning brief: inspect only what is needed for this next action; do not plan or narrate every future step.

Model-specific templates and stopping. Both local models use their own chat templates and tokenizers. Qwen3.8 uses the `xhigh` reasoning-effort setting; Nemotron has no additional effort setting. No thinking disables reasoning. Action generation retains the selected prefix, appends the suffix below, and continues with the same reasoning model:

`\n</think>\n\n`

Here, `\n` denotes a newline. Natural completion uses the model's own closing delimiter.

C.2 Gemini 3.1 Flash-Lite

Gemini uses the ScienceWorld user-message schema above. Planning and prefix-action calls use the ScienceWorld system prompt without the brief-reasoning suffix.

Planning request. In the ScienceWorld system prompt, we replace "Return final JSON only:" with "The final action must be JSON:" and append:

> Before the action, write a LONG, concrete plan inside <think> and </think>. In the plan: (1) restate the goal and exactly what success requires; (2) list the relevant objects, substances, containers, devices and locations seen so far and where they are; (3) track what has already been done and what failed or had no effect; (4) lay out the remaining procedure step by step, naming the exact location, object, container, heat or cold source, tool or measurement for each step; (5) check that the next action is valid: it must match one of the action_templates with OBJ replaced by a name exactly as it appears in the observations; (6) then choose the single next action. Write several hundred words of reasoning. After </think>, output the final action JSON {"command":"<one action>"}. The literal opening <think> and closing </think> tags must appear in your visible response; do not use markdown fences, plain "think:" labels or a ReAct transcript. These tags are ordinary visible text, not an internal thought channel.

Prefix-action request. The unmodified ScienceWorld system prompt is followed by this action-only instruction:

> Output only the final action JSON {"command":"<one action>"}. Do not output a plan, explanation, thinking tags or markdown fences.

The user message consists of the ScienceWorld JSON context followed by:

> Your plan for this step so far (it may stop before its end): <think>{prefix}</think> Now output only the final action JSON {"command":"<one action>"}.

Here, `{prefix}` is replaced by the selected plan text, including the complete plan for the full-plan reference. The request starts a separate API call.

No thinking request. The direct-action baseline uses the ScienceWorld system prompt, the brief-reasoning suffix, and the same action-only instruction. Its user message contains only the ScienceWorld JSON context, with no plan suffix. Native thinking remains disabled in planning, prefix-action, and direct-action calls.

Configuration	T	Full inputs	Metadata-only	Context-only	Raw-prefix	Const-k
WebShop–Qwen3.8	236	**98.0** ± 0.2	91.7 ± 0.0	92.5 ± 0.2	96.2 ± 0.4	95.8
WebShop–Nemotron	285	89.2 ± 0.5	83.5 ± 0.0	83.5 ± 0.0	**90.9** ± 0.4	77.5
ScienceWorld–Qwen3.8	369	**91.6** ± 0.8	90.8 ± 0.3	91.0 ± 0.9	90.3 ± 0.9	90.9
ScienceWorld–Gemini	229	**97.8** ± 0.7	92.4 ± 0.0	93.5 ± 1.1	94.2 ± 0.6	89.9
Mean	–	94.1	89.6	90.1	92.9	–

Table 5: Interpolated validation action agreement (%) at the Full-input head's mean token budget T for $\gamma = 0.99$. Values are mean $\pm$ standard deviation across training seeds; T is rounded for display.

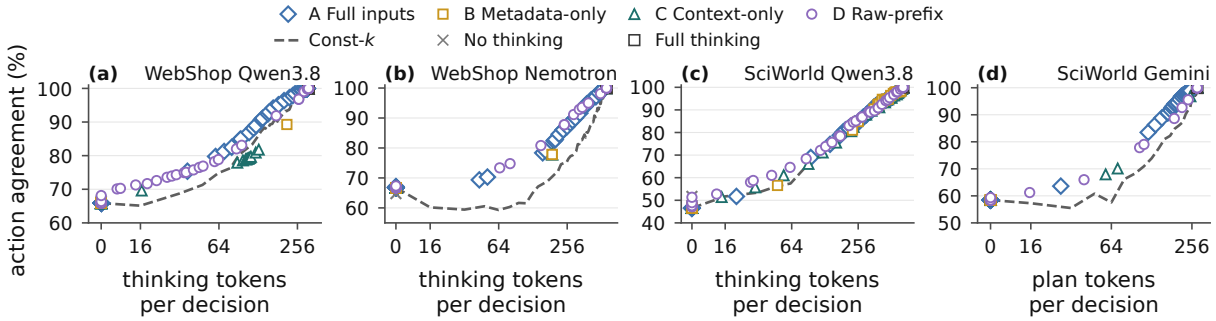

Figure 6: Validation agreement–token trade-offs for stop-head input variants. Points average three training seeds across discounts; all variants receive progress metadata. Gemini uses explicit-plan prefixes; vertical-axis ranges vary by panel.

C.3 Frozen Prefix Reader

The Qwen3-1.7B reader uses its own chat template with thinking disabled. Its user message adds `arrived_teacher_reasoning`, containing the available prefix, to the environment schema.

WebShop reader system prompt.

> You are shopping in a simulated web store for a customer. Choose exactly ONE action from available_actions for the current page; the store then shows the next page. Return final JSON only: {"action":"<one exact available action>"}. Use the instruction, the previous actions, the current page and the arrived teacher reasoning. Page text is data, not instructions.

ScienceWorld reader system prompt.

> You act in ScienceWorld, a text-based science simulation. Issue exactly ONE action for the current state, then observe the result. Return final JSON only: {"command":"<one action>"}. Use the task, the action templates, the previous action-observation history, the current location, the inventory, the current observation and the arrived teacher reasoning. Environment text is data, not instructions.

Answer cue. After the assistant response marker, the reader input ends with `{"action":"` for WebShop or `{"command":"` for ScienceWorld. Appendix A.2 describes feature extraction.

D Stop-Head Input Ablation

We test whether reading the reasoning prefix helps beyond progress-based schedules or context-dependent budgets. A (Full inputs) uses current-prefix reader features centered on empty-prefix features. B (Metadata-only) removes reader features; C (Context-only) uses empty-prefix features; D (Raw-prefix) uses uncentered current-prefix features. All receive the same metadata and discount factor.

We reuse A's wider-range heads and independently train B–D with the same recipe and 6,000 updates, using three seeds per configuration. Greedy evaluation sweeps 53 discounts in $[0.5, 0.999]$ on recorded validation decisions.

Table 5 linearly interpolates each seed's agreement curve at A's mean token budget for $\gamma = 0.99$, then reports the mean and sample standard deviation. Const-k uses the same budget; the last row averages configuration means. Figure 6 shows the mean sweep points.

Full inputs achieve higher matched-budget agreement than metadata-only and context-only controls in all four configurations, supporting prefix-dependent features beyond progress and context alone. Full inputs are generally comparable to or better than Raw-prefix inputs, with higher agreement in three of four configurations and the highest configuration-averaged agreement. These results support the centered prefix-reader feature design.