© 2026 IEEE. Personal use of this material is permitted. Permission from IEEE must be obtained for all other uses, in any current or future media, including reprinting/republishing this material for advertising or promotional purposes, creating new collective works, for resale or redistribution to servers or lists, or reuse of any copyrighted component of this work in other works.

A Systematic Comparison of Search Space Complexity and Efficiency for Time-Synchronous End-to-End ASR Models

Simon Berger
Human Language Technology and Pattern Recognition, Computer Science Department, RWTH Aachen University, Germany
AppTek GmbH, Germany
berger@hltpr.rwth-aachen.de

Larissa Kleppel
Human Language Technology and Pattern Recognition, Computer Science Department, RWTH Aachen University, Germany
lkleppel@hltpr.rwth-aachen.de

Ralf Schlüter
Human Language Technology and Pattern Recognition, Computer Science Department, RWTH Aachen University, Germany
AppTek GmbH, Germany
schlueter@hltpr.rwth-aachen.de

arXiv:2610.31259v1 [cs.CL] 22 Oct 2026

***Abstract*—Search is a central component of Automatic Speech Recognition (ASR) systems and has been deeply investigated for classical ASR modeling based on Hidden Markov Models (HMMs). However, in work on current sequence-to-sequence ASR architectures it is often simplified to a small fixed-size beam search without external Language Model (LM). In this work, we present a comparative analysis of search space complexity and efficiency for Connectionist Temporal Classification (CTC) and neural transducer models under different pruning strategies and LM configurations. We show that these end-to-end (E2E) models work with a small search space most of the time but exhibit large temporary peaks in hypothesis count during periods of uncertainty. Score-based pruning dynamically adapts the search to model uncertainty, leading to a significantly lower Real-Time Factor (RTF). Our results highlight the limitations of search with low fixed-size beam pruning and provide practical guidelines for efficient search in E2E ASR.**

***Index Terms*—ASR, search space, pruning strategies, CTC, neural transducer**

I. Introduction & Related Work

Search (or decoding) is the central process that brings together the components of an ASR system in order to turn acoustic evidence into a transcription by navigating a space of intermediate hypotheses over time. In the large vocabulary continuous speech recognition (LVCSR) era of the 1990s-2000s, search spaces often contained thousands or tens of thousands of hypotheses per time step, which drove the need for efficient search strategies and methods for search space reduction [1]–[4]. Modern neural ASR models such as CTC [5], neural transducers [6] and Attention Encoder-Decoder (AED) [7] models have changed this setting. They often exhibit sharp posterior distributions where most of the time the probability mass is concentrated on only one output class [8]. Additionally, these models are commonly defined on an orthographic (e.g. character or Byte-Pair Encoding (BPE) [9]) rather than a phonetic subword vocabulary, making search possible without a pronunciation lexicon or external LM. Due to this development, more sophisticated search methods have been largely disregarded in recent literature, and many works just report fixed beam-sizes of 1-8 without external LM integration. Dynamic score-threshold pruning and lexical prefix-tree based search organization are known techniques from classical ASR, but their practical role for modern time-synchronous E2E architectures has been underrepresented.

Most publicly available toolkits reflect this shift of focus away from search and only offer limited search capabilities for modern model architectures. Flashlight [10] offers lexicon-constrained and lexicon-free search as well as word- and subword-level LM integration, but separate pruning of within-word and word-end hypotheses is not possible and it does not provide support for transducer models. ESPnet [11] provides transducer search with neural LMs at the subword-level but does not provide lexicon- or word-level capabilities and only fixed-size beam pruning instead of dynamic score-based pruning. k2/icefall [12] also offers many options for CTC search, including a lexicon and n-gram LM but only implements neural LMs via rescoring and again lacks dynamic pruning based on a score threshold for transducer search. NeMo [13] and recent publications with it [14]–[16] aim to optimize the transducer search process fully on GPU in order to take advantage of parallelization and avoid inter-device communication. However, this approach again prohibits word-level processing and dynamic pruning.

In this paper, we challenge the assumption that small fixed-size beam search is sufficient for optimal performance with modern model architectures, in terms of both accuracy and speed. In our open-source toolkit RASR [17]–[19] [1] we implemented a generic search module which supports various search strategies for various E2E ASR models, allowing a comprehensive study of their behavior and search space under comparable conditions. We restrict the analysis in this work to time-synchronous CTC and monotonic transducer models because AED, Speech-Large Language Model (LLM) and related sequence-to-sequence models search label-synchronously which means that hypotheses at one search step may correspond to different input positions and finished hypotheses may have different lengths. This makes active-hypothesis-count and pruning statistics not directly comparable and introduces the length-normalization problem which is out of scope for this discussion.

The main contributions of our work are as follows:

- We systematically compare the performance of BPE and phoneme-based CTC and neural transducer models under various search conditions such as with and without lexicon-constraint as well as with and without subword- or word-level LMs.
- We analyze search space behavior across different models and task difficulty levels by computing search errors and hypothesis counts.
- We quantify the impact of dynamic vs. fixed-size beam pruning on the search accuracy (Word Error Rate (WER)) and efficiency (RTF) across datasets and search conditions.

[1] https://github.com/rwth-i6/rasr

Algorithm 1 One step of lexicon-free time-synchronous beam search. Scores are negative log probabilities. The subword-level LM scorer and the corresponding intermediate pruning step are optional.

Require: beam $\mathcal{H}$, output labels $\mathcal{A}$, ASR score function $\text{score}_{\text{ASR}}$, optional subword-LM score function $\text{score}_{\text{SubLM}}$
Require: score thresholds $\Delta_{\text{ASR}}, \Delta_{\text{SubLM}}$, intermediate max beam size B_{ASR}, final max beam size B_{SubLM}
$\mathcal{E} \leftarrow \emptyset$
for all $q \in \mathcal{H}$ and $a \in \mathcal{A}$ **do**
  $\text{score} \leftarrow \text{score}(q) + \text{score}_{\text{ASR}}(q, a)$
  $\mathcal{E} \leftarrow \mathcal{E} \cup \{(q, a, \text{score})\}$
end for
$\mathcal{E} \leftarrow \textsc{ScorePrune}(\mathcal{E}, \Delta_{\text{ASR}})$
if LM is present **then**
  $\mathcal{E} \leftarrow \textsc{BeamPrune}(\mathcal{E}, B_{\text{ASR}})$
  for all $e = (q, a, \text{score}) \in \mathcal{E}$ **do**
    if $a \neq \langle\text{blank}\rangle$ **then**
      $\text{score}(e) \leftarrow \text{score}(e) + \lambda \cdot \text{score}_{\text{SubLM}}(q, a)$
    end if
  end for
  $\mathcal{E} \leftarrow \textsc{ScorePrune}(\mathcal{E}, \Delta_{\text{SubLM}})$
end if
$\mathcal{H}' \leftarrow \{\textsc{Extend}(q, a, \text{score}) : (q, a, \text{score}) \in \mathcal{E}\}$
$\mathcal{H}' \leftarrow \textsc{Recombine}(\mathcal{H}')$
$\mathcal{H}' \leftarrow \textsc{BeamPrune}(\mathcal{H}', B_{\text{SubLM}})$
return $\mathcal{H}'$

- We give compact algorithmic descriptions and tuned decoding parameters for reproducibility.

Our setups are published on GitHub[2].

II. Background

In modern ASR, the search maps a sequence of encoder states h_1^T which are computed from acoustic features $x_1^{T'}$ to the output sequence with maximum probability. This can be conducted in a lexicon-free (open-vocabulary) or a lexicon-constrained (closed-vocabulary) manner. In the lexicon-free case, the maximization is carried out directly over subword sequences a_1^S:

$$h_1^T \rightarrow \arg\max_{a_1^S} \left\{ p_{\text{SubLM}}(a_1^S)^\lambda \cdot \max_{y_1^T : a_1^S} p_{\text{ASR}}(y_1^T \mid h_1^T) \right\}. \quad (1)$$

Here, $y_1^T : a_1^S$ denotes the set of alignments allowed by the label topology that map or collapse to the subword sequence a_1^S. For each subword sequence there is a score contribution from a subword LM p_{SubLM} scaled by a factor λ (which may be 0) and an ASR model p_{ASR} which is broken down to frame-wise alignments.

Lexicon-constrained search works in a similar way but has an additional outer maximization over word sequences w_1^N:

$$h_1^T \rightarrow \arg\max_{w_1^N} \Bigg\{ p_{\text{WordLM}}(w_1^N)^{\lambda_1} \cdot \max_{a_1^S : w_1^N} \Big\{ p_{\text{SubLM}}(a_1^S)^{\lambda_2} \cdot \max_{y_1^T : a_1^S} p_{\text{ASR}}(y_1^T \mid h_1^T) \Big\} \Bigg\}. \quad (2)$$

In the inner maximization, only those label sequences a_1^S that correspond to a valid word sequence according to the lexicon are considered. This type of search also allows the integration of a word-level LM p_{WordLM} in addition to p_{SubLM}. In practice, however, usually only a word-level LM is used during lexicon-constrained search.

[2]https://github.com/rwth-i6/returnn-experiments/tree/master/2026-search-space

Algorithm 2 One step of lexicon-constrained time-synchronous tree search. Scores are negative log probabilities. Subword- and word-LM scoring are optional; word ends update the word-LM history and return to a tree root.

Require: beam $\mathcal{H}$, lexical prefix tree $\mathcal{T}$, subword-level ASR score $\text{score}_{\text{ASR}}$, optional subword-LM score $\text{score}_{\text{SubLM}}$, optional word-LM score $\text{score}_{\text{WordLM}}$
Require: score thresholds $\Delta_{\text{ASR}}, \Delta_{\text{SubLM}}, \Delta_{\text{WordLM}}$, intermediate max beam size B_{ASR}, subword-level max beam size B_{SubLM}, word-level max beam size B_{WordLM}
$\mathcal{E}_{\text{Subword}} \leftarrow \emptyset$
for all $q \in \mathcal{H}$ and $(a, s_{\text{next}}) \in \textsc{Successors}(\mathcal{T}, \text{state}(q))$ **do**
  $\text{score} \leftarrow \text{score}(q) + \text{score}_{\text{ASR}}(q, a)$
  $\mathcal{E}_{\text{Subword}} \leftarrow \mathcal{E}_{\text{Subword}} \cup \{(q, a, s_{\text{next}}, \text{score})\}$
end for
$\mathcal{E}_{\text{Subword}} \leftarrow \textsc{ScorePrune}(\mathcal{E}_{\text{Subword}}, \Delta_{\text{ASR}})$
if SubLM is present **then**
  $\mathcal{E}_{\text{Subword}} \leftarrow \textsc{BeamPrune}(\mathcal{E}_{\text{Subword}}, B_{\text{ASR}})$
  for all $e = (q, a, s_{\text{next}}, \text{score}) \in \mathcal{E}_{\text{Subword}}$ **do**
    if $a \neq \langle\text{blank}\rangle$ **then**
      $\text{score}(e) \leftarrow \text{score}(e) + \lambda_2 \cdot \text{score}_{\text{SubLM}}(q, a)$
    end if
  end for
  $\mathcal{E}_{\text{Subword}} \leftarrow \textsc{ScorePrune}(\mathcal{E}_{\text{Subword}}, \Delta_{\text{SubLM}})$
end if
$\mathcal{H}_{\text{Subword}} \leftarrow \{\textsc{ExtendSubword}(e) : e \in \mathcal{E}_{\text{Subword}}\}$
$\mathcal{H}_{\text{Subword}} \leftarrow \textsc{RecombineSubword}(\mathcal{H}_{\text{Subword}})$
$\mathcal{H}_{\text{Subword}} \leftarrow \textsc{BeamPrune}(\mathcal{H}_{\text{Subword}}, B_{\text{SubLM}})$
$\mathcal{E}_{\text{word}} \leftarrow \emptyset$
for all $q \in \mathcal{H}_{\text{Subword}}$ and $(w, s_{\text{root}}) \in \textsc{WordEnds}(\mathcal{T}, \text{state}(q))$ **do**
  $\text{score} \leftarrow \text{score}(q)$
  if WordLM is present **then**
    $\text{score} \mathrel{+}= \lambda_1 \cdot \text{score}_{\text{WordLM}}(q, w)$
  end if
  $\mathcal{E}_{\text{word}} \leftarrow \mathcal{E}_{\text{word}} \cup \{(q, w, s_{\text{root}}, \text{score})\}$
end for
$\mathcal{E}_{\text{word}} \leftarrow \textsc{ScorePrune}(\mathcal{E}_{\text{word}}, \Delta_{\text{WordLM}})$
$\mathcal{H}_{\text{word}} \leftarrow \{\textsc{ExtendWordEnd}(e) : e \in \mathcal{E}_{\text{word}}\}$
$\mathcal{H}_{\text{word}} \leftarrow \textsc{RecombineWordEnd}(\mathcal{H}_{\text{word}})$
$\mathcal{H}_{\text{word}} \leftarrow \textsc{BeamPrune}(\mathcal{H}_{\text{word}}, B_{\text{WordLM}})$
return $\mathcal{H}_{\text{Subword}} \cup \mathcal{H}_{\text{word}}$

In this work, we focus on CTC and monotonic neural transducer [20] ASR models, since both of them work on a frame-by-frame basis and allow for a meaningful comparison. In the case of CTC, the alignment probability in search can be stated as follows, including a frame-wise prior probability $p_{\text{Prior}}(y_t)$ with scale β:

$$p_{\text{ASR}}(y_1^T \mid h_1^T) = \prod_{t=1}^{T} \frac{p_{\text{CTC}}(y_t \mid h_t)}{p_{\text{Prior}}(y_t)^\beta}. \quad (3)$$

For the transducer model, we use a limited context of one history label, which has previously shown to yield performance equivalent to full context [21]. We also correct by the Internal Language Model (ILM) probability $p_{\text{ILM}}(a_s \mid a_{s-1})$ which we estimate by the zero-encoder method [22]. For convenience of notation, we define $a(y_1^{t-1})$ as the last non-blank label in y_1^{t-1} and set $p_{\text{ILM}}(y \mid a) = 1$ when

y is the blank label since the ILM only contributes for non-blank labels:

$$p_{\text{ASR}}(y_1^T \mid h_1^T) = \prod_{t=1}^{T} \frac{p_{\text{Transducer}}(y_t \mid a(y_1^{t-1}), h_t)}{p_{\text{ILM}}(y_t \mid a(y_1^{t-1}))^{\beta}}. \quad (4)$$

Algorithms 1 and 2 implement the global objectives defined in equations (1) and (2) approximately in negative-log space by accumulating local ASR and LM scores. Here, $\text{score}_{\text{ASR}} = -\log p_{\text{ASR}}$ is instantiated by Equation (3) for CTC and Equation (4) for the transducer. Recombination keeps the best hypothesis among equivalent contexts. For lexicon-free search, this recombination context consists of prediction network state (for transducer) and the subword-LM history. For lexicon-constrained search, the recombination context additionally contains tree state and the word-LM history.

During the search process, we apply score-based and max-beam-size-based pruning strategies. In score-based pruning with a threshold Δ, the current best hypothesis is selected from the beam and all hypotheses whose score (i.e., negative log probability) is worse than the current best plus the threshold are pruned. Let $\mathcal{H}$ be the current set of hypotheses, then this is formalized as:

$$\textsc{ScorePrune}(\mathcal{H}, \Delta) = \left\{ q \in \mathcal{H} \mid \text{score}(q) \leq \min_{q' \in \mathcal{H}} \{\text{score}(q')\} + \Delta \right\} \quad (5)$$

On the other hand, with max-beam-size pruning, we simply keep a fixed maximum amount B of hypotheses while pruning away the rest, i.e.:

$$\textsc{BeamPrune}(\mathcal{H}, B) = \text{Top}_B(\mathcal{H}) \quad (6)$$

Pure fixed-beam search disables ScorePrune and only applies BeamPrune. In our search implementation, we apply score- and max-beam-size-pruning with the option to use different thresholds after each individual score contribution, i.e., first after $\text{score}_{\text{ASR}}$, then after $\text{score}_{\text{SubLM}}$ and finally after $\text{score}_{\text{WordLM}}$. We also prune within-word hypotheses separately from word-end hypotheses. We denote the ASR-stage, subword-LM and word-LM maximum beam sizes by B_{ASR}, B_{SubLM} and B_{WordLM}, and their score thresholds by Δ_{ASR}, Δ_{SubLM} and Δ_{WordLM}, respectively.

All decoding experiments use RASR as the implementation framework. It supports, among others, CTC and transducer scorers, subword- and word-level LMs, lexicon-free and lexicon-constrained decoding, and the staged pruning described above. Algorithms 1 and 2 summarize the corresponding search steps implementations, but omit some optimizations such as on-the-fly extension pruning. Notably, all results characterize the model and search procedure, not RASR-specific behavior.

III. Experimental Setup

Experiments are conducted on three corpora: LibriSpeech, Loquacious, and an in-house Spanish telephony corpus. LibriSpeech [23] uses the full 960h training set and is evaluated on dev-other and test-other. For Loquacious, we use the public medium portion with 2500h of 16kHz training data [24]. Loquacious evaluation is performed on dev.all and test.all, which aggregate VoxPopuli, LibriSpeech, YODAS, and Common Voice. The Spanish corpus contains around 12000h of 8kHz telephony training data and is evaluated on an in-house 4.5h test set, providing a harder-domain cross-lingual condition.

We train models on BPE and phoneme alphabets. The LibriSpeech BPE vocabulary contains 185 labels including blank while the end-of-word (EOW) augmented phoneme vocabulary contains 80 labels including blank. The Loquacious setup uses BPE/EOW-phoneme vocabularies of 184/80 labels including blank. On the Spanish task, we use a vocabulary of 63 EOW-phoneme labels.

For acoustic features, we extract 80-dimensional log-mel features from the 16kHz LibriSpeech and Loquacious data, and 60-dimensional log-mel features from the 8kHz Spanish data. The features are passed through a convolutional front-end and subsampled to a step-size of 40 ms. Afterwards, they are processed by a Conformer [25] encoder with 12 layers for LibriSpeech and Loquacious, and 16 layers for Spanish. The transducer models additionally use a 2-layer feed-forward prediction network for the limited history and a 1-layer feed-forward joiner network.

Each LibriSpeech model is trained for 100 epochs, each Loquacious model for 40 epochs, and the Spanish model for 60 epochs. All models use a One-Cycle Learning Rate (OCLR) [26] schedule and apply SpecAugment [27], speed perturbation [28], dropout, and weight decay for regularization.

For word-level LM experiments, we apply a 4-gram LM; for LibriSpeech, this is the official 4-gram LM, for Loquacious, this is the 4-gram LM setup described in [29], and for Spanish, it is trained on in-house data. On LibriSpeech, we additionally present results with a small 2-layer Long Short-Term Memory (LSTM) [30] LM with 55M parameters at the BPE level and larger 96-layer Transformer LMs at BPE and word level with around 300M and 430M parameters respectively [31].

TABLE I
WER and RTF results on LibriSpeech for different models given various search conditions and LM combinations.

Model	Lex.	LM		WER [%]		RTF $\cdot 10^3$
		Subword	Word	dev-o	test-o	
BPE CTC	No	-	-	7.0	7.1	0.62
		LSTM	-	5.1	5.1	22
		TraFo		4.6	4.7	3357
	Yes	-	-	6.5	6.6	0.93
		-	4-gram	5.2	5.6	1.6
		LSTM	-	5.2	5.3	9.0
		LSTM	4-gram	5.1	5.3	7.5
		-	TraFo	4.6	4.7	154
Phon. CTC	Yes	-	4-gram	5.5	5.7	1.5
		-	TraFo	4.5	4.8	290
BPE Transd.	No	-	-	6.0	5.9	7.6
		LSTM	-	4.5	4.8	27
		TraFo	-	4.1	4.3	3488
	Yes	-	-	5.9	5.9	8.1
		-	4-gram	4.8	5.1	13
		LSTM	-	4.7	5.1	21
		LSTM	4-gram	4.8	5.1	20
		-	TraFo	3.8	4.2	310

IV. Search Strategies

Table I shows WER and search RTF results for the LibriSpeech models under different search conditions and LM integrations. The search RTF only measures the computational speed of the search process, which includes hypothesis extension, predictor and joiner forwarding for transducer, LM forwarding, pruning and bookkeeping but excludes the encoder forwarding. It is measured with all neural components (i.e., neural LMs and transducer predictor and joiner) running on an A10 GPU while the search logic and 4-gram LM run on CPU (note: the encoder RTF on an A10 is about 3.0×10^{-3}). The presented results cover decoding with no external LM, subword-level neural LSTM or Transformer LM integration, and word-level Transformer-LM or 4-gram integration. Lexicon constraints improve CTC without an external LM (from 7.0/7.1% to 6.5/6.6% WER) and

TABLE II
WER AND RTF RESULTS ON LOQUACIOUS FOR DIFFERENT MODELS GIVEN VARIOUS SEARCH CONDITIONS AND LM COMBINATIONS.

Model	Lex.	Word LM	WER [%]		RTF $\cdot 10^3$
			dev.all	test.all	
BPE CTC	No	-	11.7	12.8	0.62
	Yes	-	11.0	12.0	0.73
		4-gram	9.5	10.4	0.89
BPE Transd.	No	-	9.0	10.2	7.7
	Yes	-	9.0	10.1	8.6
		4-gram	8.0	8.9	13
Phon. Transd.	Yes	4-gram	8.4	9.2	13

TABLE III
WER AND HYPOTHESIS COUNT STATISTICS FOR DIFFERENT PRUNING PARAMETERS WITH BPE CTC LEXICON-CONSTRAINED SEARCH + WORD-LEVEL TRANSFORMER LM ON LIBRISPEECH DEV-OTHER.

Score Thresh.	Beam Size	WER [%]	Search Err. [%]	#Hypotheses		
				p50	p90	p99
0.0	1	10.4	30.4	1	1	1
4.0	16	9.4	14.7	1	2	7
8.0	128	6.7	4.9	2	6	41
12.0	256	4.8	1.3	5	28	242
16.0	256	4.6	0.6	35	167	254

have little effect for the transducer. All external-LM combinations outperform the standalone ASR models by a good margin, with the best setup reaching 3.8/4.2% WER. Overall, the word-level Transformer LM gives the best WER, while the 4-gram and subword-level LSTM LMs provide decent accuracy at much higher speed. The subword-level Transformer LM is expensive in terms of execution time because it needs to score every (non-blank) subword extension, while not improving WER over the word-level Transformer LM which only needs to be forwarded at word ends. The subword-level LSTM LM slightly improves over the 4-gram LM, but is also slower; combining the LSTM and 4-gram LMs gives results similar to the LSTM alone. A natural direction for future work is to extend this comparison to larger ASR models trained on more data, where integration of external LMs may have lower impact.

The phoneme CTC model is slightly worse than BPE CTC with a 4-gram LM (5.5/5.7% vs. 5.2/5.6%) and on par with the Transformer LM. Search without a lexicon and word-LM is not applicable to phoneme models since phonetically equivalent words cannot be distinguished. The transducer outperforms the CTC models, especially when no lexicon or LM is present (6.0/5.9% vs. 7.0/7.1% WER) and remains better by between 0.3% to 0.7% absolute WER in all other conditions. Compared to CTC, transducer search is much slower without LM or with only a 4-gram LM, but the gap shrinks with a neural LM because LM forwarding makes up a large part of the search cost.

Table II reports the same core search conditions without neural LMs on Loquacious, providing evidence on a public heterogeneous dataset beyond LibriSpeech.

V. SEARCH SPACE COMPLEXITY

We first investigate the relation between pruning parameters, search space size, and recognition performance. To optimize the pruning parameters, we first set a very large maximum beam-size B while tuning the score thresholds and then decrease B until performance starts to degrade. Tables III and IV show LibriSpeech lexicon-constrained BPE CTC and BPE transducer search with a word-level Transformer LM. Search errors count segments that were transcribed incorrectly although the ground-truth transcription has a better score than the recognized one, i.e., the ground-truth was erroneously pruned away during search. The tables also include percentile statistics of the within-word hypothesis counts at each search step. The p50 value shows that the median active hypothesis set is small while the p90 and p99 statistics show that in some rare difficult regions the signal causes locally increased hypothesis counts. This effect is visualized in Figure 1 which shows the hypothesis count by timestep for a selected utterance with a CTC and transducer model. The search space stays small for most search steps and exhibits one large temporary peak between steps 300 and 400.

TABLE IV
WER AND HYPOTHESIS COUNT STATISTICS FOR DIFFERENT PRUNING PARAMETERS WITH BPE TRANSDUCER LEXICON-CONSTRAINED SEARCH + WORD-LEVEL TRANSFORMER LM ON LIBRISPEECH DEV-OTHER.

Score Thresh.	Beam Size	WER [%]	Search Err. [%]	#Hypotheses		
				p50	p90	p99
0.0	1	10.0	24.1	1	1	1
4.0	16	11.0	20.3	1	2	6
8.0	128	6.9	7.2	2	7	51
12.0	256	4.2	1.7	7	46	250
16.0	1024	3.8	0.4	55	433	1012

TABLE V
WER AND HYPOTHESIS COUNT STATISTICS FOR DIFFERENT PRUNING PARAMETERS WITH BPE TRANSDUCER LEXICON-CONSTRAINED SEARCH + WORD-LEVEL 4-GRAM LM ON LOQUACIOUS DEV.ALL.

Score Thresh.	Beam Size	WER [%]	Search Err. [%]	#Hypotheses		
				p50	p90	p99
0.0	1	11.9	16.8	1	1	1
2.0	16	10.9	10.9	1	2	4
4.0	32	8.6	2.6	1	3	11
6.0	32	8.0	0.7	1	5	30
8.0	32	7.9	0.5	2	11	31

Table V reports the corresponding analysis for Loquacious with BPE transducer decoding and a word 4-gram LM. Overall, the hypothesis counts are lower on Loquacious. This effect can be explained by the fact that the experiments on Loquacious only use a 4-gram LM and a much lower LM scale (see Table X). Since the 4-gram only relies on a three-word history, the recombination step can recombine many more hypotheses and with the low LM scale, score jumps from the LM score contributions are much smaller and need a lower score threshold to cover for them. Still, the same non-uniform behavior can be observed on this dataset: the median active set remains small, while p99 reflects rare high-complexity regions and is comparatively much higher.

To test more difficult acoustic conditions, we collected similar

TABLE VI
WER AND HYPOTHESIS COUNT STATISTICS FOR DIFFERENT PRUNING PARAMETERS WITH PHONEME CTC LEXICON-CONSTRAINED SEARCH + WORD-LEVEL 4-GRAM LM ON SPANISH TELEPHONY.

Score Thresh.	Beam Size	WER [%]	#Hypotheses			
			Avg.	p50	p90	p99
0.0	1	46.6	1.0	1	1	1
4.0	128	17.9	6.1	3	11	79
6.0	512	15.5	88.3	55	180	512
8.0	1024	15.0	251.0	148	675	1024
10.0	4096	14.9	932.2	451	2926	4096

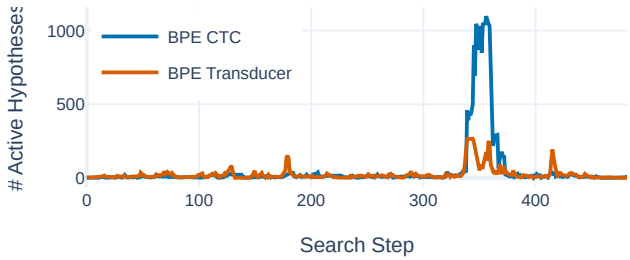

Fig. 1. Illustrative number of active within-word hypotheses at each search step for a selected LibriSpeech dev-other utterance with BPE CTC/transducer using lexicon-constrained search including a word-level 4-gram LM.

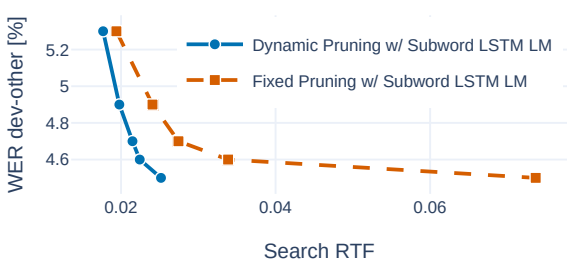

Fig. 2. Search RTF vs. WER for fixed and dynamic pruning on LibriSpeech dev-other with BPE transducer + subword-level LSTM LM.

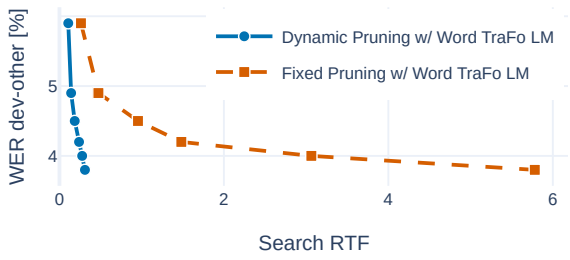

Fig. 3. Search RTF vs. WER for fixed and dynamic pruning on LibriSpeech dev-other with BPE transducer + word-level Transformer LM.

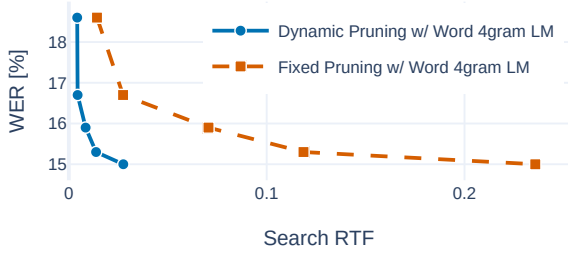

Fig. 4. Search RTF vs. WER for fixed and dynamic pruning on Spanish telephony with phoneme CTC + word-level 4-gram LM.

statistics for a phoneme CTC model on the in-house Spanish 8kHz telephony test set, as shown in Table VI. This task has a much higher WER and a low hypothesis count in the tens is not nearly enough to reach full performance in this case. Even at the median, several hundred hypotheses are required and the p90 and p99 values go into the thousands in order to reach optimal performance.

These results confirm that small fixed beams are insufficient when model uncertainty is higher and when external LM contribution makes up a large part of the score. Dynamic pruning adapts to the model uncertainty and LM score contributions and keeps the hypothesis counts comparatively low in the median.

VI. Search Efficiency

Search efficiency is tied to the size of the active search space. We compare fixed max-beam pruning with dynamic score-threshold pruning while preserving the WER. For each max-beam size B_{ASR}, the fixed baseline uses only max-beam pruning while the dynamic variant additionally performs score-threshold pruning with the tightest possible thresholds that do not degrade the WER. We report the relative search RTF reduction when switching from fixed to dynamic pruning in percent:

$$\text{RTF reduction}\,[\%] = 100 \cdot \left(1 - \frac{\text{RTF}_{\text{dynamic}}}{\text{RTF}_{\text{fixed}}}\right). \qquad (7)$$

Figures 2 and 3 contrast score-based pruning with pure max-beam pruning for LibriSpeech BPE transducer search with subword and word LMs. In both cases, dynamic pruning achieves the same WERs at much faster inference times, especially at higher beam sizes. With dynamic pruning, the RTF only increases marginally while fixed beam pruning quickly becomes very expensive. For the smaller LSTM LM the search times are faster than with the larger Transformer LM but the relative difference between fixed and dynamic pruning is still clear. Figure 4 shows the same effect on Spanish telephony with a 4-gram LM.

Tables VII and VIII provide detailed comparisons over the max beam size B_{ASR} for the subword-level LSTM LM and word-level Transformer LM setups on LibriSpeech, including the relative RTF reduction. In Table IX we summarize the reduction in average hypothesis count and search RTF for different ASR models and external LMs on all three datasets. When no external LM is used, greedy search already gives optimal results and dynamic pruning provides no advantage in this case. For all experiments with external LM, however, dynamic pruning provides significant efficiency im-

TABLE VII
WER and Search RTF breakdown for fixed and dynamic pruning with different max beam sizes B_{ASR} on LibriSpeech dev-other with BPE Transducer + subword-level LSTM LM.

B_{ASR}	WER [%]	Search RTF $\cdot 10^3$		RTF reduction [%]
		Fixed	Dynamic	
4	5.3	19	18	5.3
16	4.7	27	21	22.2
64	4.6	49	24	51.0
128	4.5	74	25	66.2

TABLE VIII
WER and Search RTF breakdown for fixed and dynamic pruning with max beam sizes B_{ASR} on LibriSpeech dev-other with BPE Transducer + word-level Transformer LM.

B_{ASR}	WER [%]	Search RTF $\cdot 10^3$		RTF reduction [%]
		Fixed	Dynamic	
4	10.4	72	60	16.7
16	5.9	261	107	59.0
64	4.5	955	186	80.5
256	4.0	3063	276	91.0
1024	3.8	5781	310	94.6

TABLE IX
RTF REDUCTION THROUGH DYNAMIC PRUNING FOR DIFFERENT SETUPS. MEASURED ON DEV-OTHER FOR LIBRISPEECH, DEV.ALL FOR LOQUACIOUS AND AN IN-HOUSE TEST-SET FOR SPANISH. AVG. HYPS COUNTS ACTIVE SUBWORD-LEVEL HYPOTHESES AFTER PRUNING.

Corpus	Model	LM	WER [%]	Fixed Pruning		Dynamic Pruning		RTF
				Avg. Hyps	RTF $\cdot 10^3$	Avg. Hyps	RTF $\cdot 10^3$	red. [%]
LibriSpeech	BPE CTC	-	7.0	1	0.62	1.0	0.62	0.0
		Word 4-gram	5.2	256	326	13.5	1.6	99.5
		Subword LSTM	5.1	128	140	27.7	22	84.3
		Word Trafo	4.6	256	1203	61.2	154	87.2
	Phoneme CTC	Word 4-gram	5.5	256	128	13.2	1.5	98.8
		Word Trafo	4.5	256	1320	60.0	292	77.9
	BPE Transducer	-	6.0	1	7.6	1.0	7.6	0.0
		Word 4-gram	4.8	128	97	12.7	13	86.6
		Subword LSTM	4.5	128	74	13.8	25	66.2
		Word Trafo	3.8	1024	5781	150.6	310	94.6
Loquacious	BPE CTC	-	11.7	1	0.62	1.0	0.62	0.0
		Word 4-gram	9.5	64	91	4.2	0.89	99.0
	BPE Transducer	-	9.0	1	7.7	1.0	7.7	0.0
		Word 4-gram	8.0	64	121	2.9	13	89.3
	Phoneme Transducer	Word 4-gram	8.4	64	70	4.6	13	81.4
Spanish In-House	Phoneme CTC	Word 4-gram	14.9	4096	1860	932.2	90	95.2

TABLE X
TUNED DECODING PARAMETERS FOR DYNAMIC PRUNING SETUPS.

Corpus	Model	LM	LM Scale	ILM / Prior Scale	B_{ASR}	B_{LM}	Δ_{ASR}	Δ_{LM}
LibriSpeech	BPE CTC	-	-	-	1	-	0.0	-
		Word 4-gram	0.6	0.2	256	16	12.0	6.0
		Subword LSTM	0.8	0.2	1024	128	14.0	12.0
		Word Trafo	0.8	0.2	256	16	16.0	8.0
	Phoneme CTC	Word 4-gram	0.6	0.2	256	8	12.0	6.0
		Word Trafo	0.8	0.2	256	16	16.0	8.0
	BPE Transducer	-	-	-	1	-	0.0	-
		Word 4-gram	0.6	0.2	128	4	12.0	4.8
		Subword LSTM	0.8	0.2	64	32	10.0	8.0
		Word Trafo	0.8	0.2	1024	64	16.0	12.8
Loquacious	BPE CTC	-	-	-	1	-	0.0	-
		Word 4-gram	0.3	0.2	64	16	6.0	3.0
	BPE Transducer	-	-	-	1	-	0.0	-
		Word 4-gram	0.3	0.2	64	16	6.0	3.0
	Phoneme Transducer	Word 4-gram	0.3	0.2	64	16	6.0	3.0
Spanish In-House	Phoneme CTC	Word 4-gram	0.5	0.3	4096	16	10.0	5.0

provements to reach the same WER. The average hypothesis count and search RTF are often reduced by around an order of magnitude. The relative RTF reduction is largest for CTC and 4-gram LM setups because their search cost is dominated by hypothesis expansion and CPU-side bookkeeping, while transducer prediction/joiner networks and neural LMs use batched GPU computations which do not scale linearly with the active set.

Table X lists tuned dynamic-pruning parameters. The LM scale weights external-LM scores; the ILM/prior scale controls transducer ILM or CTC prior correction. $B_{\text{ASR}}/B_{\text{LM}}$ are maximum beam sizes and $\Delta_{\text{ASR}}/\Delta_{\text{LM}}$ are score thresholds pruning at the ASR and LM stages respectively. These parameters define the pruning settings used in Table IX.

VII. CONCLUSIONS

We presented a systematic comparison of search strategies for time-synchronous BPE- and phoneme-based CTC and monotonic neural transducer models in E2E ASR. Comparisons of WER and RTF show that dynamic score-based pruning and external LM combination remain relevant for decoding with modern E2E ASR models. We have seen that typically, these models operate with a low active hypothesis count for the majority of search steps but spike to hundreds or thousands of hypotheses in rare difficult regions. This behavior is verified on the public LibriSpeech task as well as the more heterogeneous public Loquacious task and a more difficult in-house telephony dataset.

The fixed-vs-dynamic pruning experiments show that dynamic pruning can preserve WER while reducing search RTF by adapting the active hypothesis set to local uncertainty, especially with external LMs where low fixed-size beams are insufficient for optimal performance. We provide tuned decoding parameters and open source implementations for reproducibility.

This study is limited to time-synchronous CTC and transducer search; label-synchronous and streaming decoding require some fundamentally different analysis and remain future work.

Overall, our results suggest that external LM integration and dynamic pruning are underutilized in modern E2E ASR applications.

Acknowledgements

This work was partially supported by the project RESCALE within the program *AI Lighthouse Projects for the Environment, Climate, Nature and Resources* funded by the Federal Ministry for the Environment, Nature Conservation, Nuclear Safety and Consumer Protection (BMUV), funding ID: 67KI32006A.

REFERENCES

[1] Volker Steinbiss, Bach-Hiep Tran, and Hermann Ney, "Improvements in beam search," in *International Conference on Spoken Language Processing*, Yokohama, Japan, September 1994, pp. 2143–2146.

[2] Xavier L Aubert, "An overview of decoding techniques for large vocabulary continuous speech recognition," *Computer Speech & Language*, vol. 16, no. 1, pp. 89–114, 2002.

[3] Sascha Wendt, Gernot A Fink, and Franz Kummert, "Dynamic search-space pruning for time-constrained speech recognition," in *Proc. Interspeech*, Denver, CO, USA, September 2002, pp. 377–380.

[4] Janne Pylkkönen, "New pruning criteria for efficient decoding.," in *Proc. Interspeech*, Lisbon, Portugal, September 2005, pp. 581–584.

[5] Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber, "Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks," in *Proc. ICML*, Pittsburgh, PA, USA, June 2006, pp. 369–376.

[6] Alex Graves, "Sequence transduction with recurrent neural networks," *arXiv preprint arXiv:1211.3711*, 2012.

[7] Jan Chorowski, Dzmitry Bahdanau, Dmitriy Serdyuk, Kyunghyun Cho, and Yoshua Bengio, "Attention-based models for speech recognition," in *Proc. NIPS*, Montréal, Canada, December 2015, pp. 577–585.

[8] Albert Zeyer, Ralf Schlüter, and Hermann Ney, "Why does CTC result in peaky behavior?," *arXiv preprint arXiv:2105.14849*, 2021.

[9] Philip Gage, "A new algorithm for data compression," *The C Users Journal*, vol. 12, no. 2, pp. 23–38, February 1994.

[10] Jacob D Kahn, Vineel Pratap, Tatiana Likhomanenko, Qiantong Xu, et al., "Flashlight: Enabling innovation in tools for machine learning," in *Proc. ICML*, Baltimore, MD, USA, July 2022, pp. 10557–10574.

[11] Shinji Watanabe, Takaaki Hori, Shigeki Karita, Tomoki Hayashi, et al., "ESPnet: End-to-end speech processing toolkit," in *Proc. Interspeech*, Hyderabad, India, September 2018, pp. 2207–2211.

[12] Daniel Povey, Arnab Ghoshal, Gilles Boulianne, Lukáš Burget, et al., "The Kaldi speech recognition toolkit," in *Proc. ASRU*, Waikoloa, HI, USA, December 2011.

[13] Oleksii Kuchaiev, Jason Li, Huyen Nguyen, Oleksii Hrinchuk, et al., "NeMo: a toolkit for building AI applications using neural modules," *arXiv preprint arXiv:1909.09577*, 2019.

[14] Hainan Xu, Vladimir Bataev, Lilit Grigoryan, and Boris Ginsburg, "WIND: Accelerated RNN-T decoding with windowed inference for non-blank detection," in *Proc. Interspeech*, Rotterdam, The Netherlands, August 2025, pp. 639–643.

[15] Vladimir Bataev, Andrei Andrusenko, Lilit Grigoryan, Aleksandr Laptev, et al., "NGPU-LM: GPU-accelerated n-gram language model for context-biasing in greedy ASR decoding," in *Proc. Interspeech*, Rotterdam, The Netherlands, August 2025, pp. 644–648.

[16] Lilit Grigoryan, Vladimir Bataev, Andrei Andrusenko, Hainan Xu, et al., "Pushing the limits of beam search decoding for transducer-based ASR models," in *Proc. Interspeech*, Rotterdam, The Netherlands, August 2025, pp. 649–653.

[17] Simon Wiesler, Alexander Richard, Pavel Golik, Ralf Schlüter, and Hermann Ney, "RASR/NN: The RWTH neural network toolkit for speech recognition," in *Proc. ICASSP*, Florence, Italy, May 2014, pp. 3281–3285.

[18] David Rybach, Stefan Hahn, Patrick Lehnen, David Nolden, et al., "RASR - the RWTH Aachen University open source speech recognition toolkit," in *Proc. ASRU*, Waikoloa, HI, USA, December 2011, pp. 3281–3285.

[19] Wei Zhou, Eugen Beck, Simon Berger, Ralf Schlüter, and Hermann Ney, "RASR2: The RWTH ASR toolkit for generic sequence -to-sequence speech recognition," in *Proc. Interspeech 2023*, Dublin, Ireland, August 2023, pp. 4094–4098.

[20] Anshuman Tripathi, Han Lu, Hasim Sak, and Hagen Soltau, "Monotonic recurrent neural network transducer and decoding strategies," in *Proc. ASRU*, Singapore, Singapore, December 2019, pp. 944–948.

[21] Wei Zhou, Simon Berger, Ralf Schlüter, and Hermann Ney, "Phoneme based neural transducer for large vocabulary speech recognition," in *Proc. ICASSP*, Toronto, ON, Canada, June 2021, pp. 5644–5648.

[22] Ehsan Variani, David Rybach, Cyril Allauzen, and Michael Riley, "Hybrid autoregressive transducer (HAT)," in *Proc. ICASSP*, Barcelona, Spain, May 2020, pp. 6139–6143.

[23] Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur, "LibriSpeech: An ASR corpus based on public domain audio books," in *Proc. ICASSP*, South Brisbane, QLD, Australia, April 2015, pp. 5206–5210.

[24] Titouan Parcollet, Yuan Tseng, Shucong Zhang, and Rogier C. van Dalen, "Loquacious Set: 25,000 Hours of Transcribed and Diverse English Speech Recognition Data for Research and Commercial Use," in *Interspeech 2025*, 2025, pp. 4053–4057.

[25] Anmol Gulati, James Qin, Chung-Cheng Chiu, Niki Parmar, et al., "Conformer: Convolution-augmented transformer for speech recognition," in *Proc. Interspeech 2020*, Shanghai, China, October 2020, pp. 5036–5040.

[26] Leslie N Smith and Nicholay Topin, "Super-convergence: Very fast training of neural networks using large learning rates," in *Artificial intelligence and machine learning for multi-domain operations applications*, Baltimore, MD, USA, April 2019, pp. 369–386.

[27] Daniel S Park, William Chan, Yu Zhang, Chung-Cheng Chiu, et al., "SpecAugment: A simple data augmentation method for automatic speech recognition," in *Proc. Interspeech 2019*, Graz, Austria, September 2019, pp. 2613–2617.

[28] Tom Ko, Vijayaditya Peddinti, Daniel Povey, and Sanjeev Khudanpur, "Audio augmentation for speech recognition," in *Proc. Interspeech*, Dresden, Germany, September 2015, pp. 3586–3589.

[29] Nick Rossenbach, Robin Schmitt, Tina Raissi, Simon Berger, Larissa Kleppel, and Ralf Schlüter, "Supplementary resources and analysis for automatic speech recognition systems trained on the loquacious dataset," in *Proceedings of the Fifteenth Language Resources and Evaluation Conference (LREC 2026)*, Stelios Piperidis, Núria Bel, Henk van den Heuvel, Nancy Ide, Simon Krek, and Antonio Toral, Eds., Palma, Mallorca, Spain, May 2026, pp. 5839–5848, European Language Resources Association (ELRA).

[30] Sepp Hochreiter and Jürgen Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, November 1997.

[31] Kazuki Irie, Albert Zeyer, Ralf Schlüter, and Hermann Ney, "Language modeling with deep transformers," in *Proc. Interspeech*, Graz, Austria, September 2019, pp. 3905–3909.