How Well Do Decision Models Transfer to Vietnamese? An Empirical Study of Text Classification

Phat Huu Ha[1,2] **and Nguyen Xuan Phuc**[1,2] **and Nong Nhut Quy**[1,2] **and Trinh Tran Tran**[1,2] **and Nguyen Minh Triet**[1,2] **and Dang Van Thin**[1,2]
[1]University of Information Technology, Ho Chi Minh City, Vietnam
[2]Vietnam National University HoChiMinh City, Vietnam
Corresponding Author: thindv@uit.edu.vn

arXiv:2610.31422v1 [cs.CL] 22 Oct 2026

Abstract

Recent advances in specialized decision models have expanded the possibilities for zero-shot classification, yet their effectiveness on Vietnamese remains underexplored. This study evaluates four specialized decision models across 13 classification settings spanning six task groups, examining zero-shot performance, sensitivity to label-description language, and robustness to Vietnamese linguistic variations. The evaluation considers English, Vietnamese, and bilingual label descriptions, alongside common linguistic perturbations involving teencode, typographical errors, and diacritic removal. Under English label descriptions, JEV achieves the strongest overall performance, ranking first in 11 classification settings. The effects of label-description language vary across models and datasets, with bilingual descriptions providing no consistent improvement. Robustness experiments reveal that diacritic removal causes the most substantial performance degradation, reducing mean macro-F_1 by 9.51–15.47 points across model configurations, considerably more than teencode or typographical errors. These findings highlight the importance of label-description choices and Vietnamese orthographic variation when evaluating specialized decision models for zero-shot text classification.

1 Introduction

Text classification is a fundamental task in natural language processing, covering applications such as sentiment analysis, intent detection, and natural language inference. General-purpose large language models (LLMs) have demonstrated the ability to perform such tasks through zero-shot prompting. Recently, specialized decision models have emerged to provide structured predictions over predefined candidate options without free-form text generation. Models such as JEV supports classification through typed decision interfaces or user-defined schemas (TypeSafe AI, 2026).

Recent studies have evaluated specialized decision models across various tasks and domains. Deußer et al. (2026) evaluate JEV across 37 datasets, reporting strong performance on several tasks alongside considerable variation in results. Other studies examine decision models and compare their performance with general-purpose LLMs (Duan et al., 2026; Ibrahim and Zaki, 2026). Despite these efforts, their performance on Vietnamese text classification remains insufficiently characterized, particularly across tasks with different label inventories and linguistic characteristics.

Classification performance may also depend on how labels are specified and how input text is written. Previous research has shown that option labels can influence predictions in typed decision models even when label definitions are provided (Azizi et al., 2026). This motivates examining the effects of label-description language on Vietnamese classification tasks. Vietnamese user-generated text also contains abbreviations, non-standard spellings, and omitted diacritics (Nguyen et al., 2024, 2023). Previous evaluations of Vietnamese pretrained language models have reported substantial performance degradation under diacritic removal (Nguyen et al., 2023). These findings motivate further investigation of decision-model sensitivity to both label descriptions and Vietnamese linguistic variations.

This study evaluates four specialized decision models across 13 Vietnamese classification settings spanning six task groups, with additional comparisons against previously published LLM results. The evaluation examines zero-shot classification performance, the effects of label-description language, and robustness to Vietnamese linguistic perturbations. Under English label descriptions, JEV achieves the highest performance in 11 settings. The effects of label-description language vary across models and datasets, while diacritic removal causes the largest performance degradation

among the evaluated perturbations.

The main contributions of this study are:

- **Zero-shot evaluation of decision models for Vietnamese.** A systematic comparison across diverse classification tasks, identifying differences in model performance at both task and class levels.
- **Analysis of label-description sensitivity and linguistic robustness.** An investigation of the effects of label-description language and Vietnamese linguistic variations, revealing model-dependent behavior and substantial sensitivity to diacritic removal.

2 Related Work

2.1 Specialized Decision Models and Their Evaluation

Specialized decision models produce structured predictions over predefined candidate options without free-form text generation. JEV and LAYA provide typed decision interfaces, while GLiNER2.5-multi-Decide supports multilingual classification with user-specified labels (TypeSafe AI, 2026; Convai Innovations, 2026; Zaratiana et al., 2025). These systems support related classification tasks despite differences in their architectures and inference interfaces.

Recent evaluations show that decision-model performance varies across tasks and evaluation conditions. Broad benchmarks report competitive results on several classification and knowledge tasks, alongside limitations on more demanding tasks and settings (Deußer et al., 2026; Duan et al., 2026; Ibrahim and Zaki, 2026; Li et al., 2026). More focused studies reveal behavior not captured by aggregate scores. Experiments on legal inference identify sensitivity to request configuration, comparisons of JEV and LAYA show different performance and robustness profiles, and analyses of JEV report biases in the distribution of selected labels and ordinal scores (Zhang et al., 2026; Li, 2026; Gao et al., 2026). However, evidence across diverse Vietnamese classification tasks remains limited.

2.2 Label Specifications and Vietnamese Linguistic Robustness

The specification of candidate labels can influence zero-shot classification. Studies of generative LLMs show that label wording and definitions affect classification performance (Mu et al., 2024; Li et al., 2025). Related findings have been reported for typed decision models, where option names can exert greater influence than their accompanying definitions (Azizi et al., 2026). This evidence motivates examining the language of label descriptions as a separate factor in Vietnamese classification.

Robustness evaluation examines model behavior under controlled changes to input text. CheckList and TextFlint establish methods for testing linguistic transformations, while DecisionRobust extends perturbation-based evaluation to decision models and emphasizes the importance of validating that transformed inputs retain their original labels (Ribeiro et al., 2020; Wang et al., 2021; Lan et al., 2026). For Vietnamese, ViLexNorm documents nonstandard lexical forms in social media, and ViSoBERT reports performance degradation under diacritic removal (Nguyen et al., 2024, 2023). These findings motivate evaluating the effects of Vietnamese linguistic variations on specialized decision models.

3 Experimental Setup

3.1 Tasks and Datasets

The evaluation covers 13 classification settings from 12 Vietnamese datasets across six task groups: topic classification, emotion recognition, hate speech detection, sentiment analysis, intent detection, and natural language inference (NLI). The datasets include UIT-VSFC (Nguyen et al., 2018), NEU-ESC (Mai et al., 2025), UIT-VSMEC (Ho et al., 2019), UIT-ViOCD (Nhung et al., 2021), ViHSD (Luu et al., 2021), UIT-ViCTSD (Nguyen et al., 2021), PhoATIS (Dao et al., 2021), ViMQ (Huy et al., 2021), ViNLI (Huynh et al., 2022), ViANLI (Huynh et al., 2026), ViMedNLI (Phan et al., 2023), and vnNLI (Anh et al., 2022). UIT-VSFC is evaluated for both sentiment and topic classification.

3.2 Models

Four specialized decision models are evaluated. JEV (TypeSafe AI, 2026) is accessed through its hosted API using the typesafe/jev-1.13 model version. The remaining three models are evaluated locally using open-weight checkpoints: LAYA Multilingual (322M parameters)[1], GLiNER2.5-multi-

[1] https://huggingface.co/convaiinnovations/laya-multilingual

Table 1: Zero-shot classification performance (%) using English label descriptions and a supervised fine-tuned PhoBERT baseline. Acc, WF1, and MF1 denote accuracy, weighted-F_1, and macro-F_1. Bold and underlined MF1 indicate the highest and second-highest reported scores. Dashes denote unavailable values. Means cover the 12 settings shared by all four zero-shot models.

Dataset	LAYA			GLiNER			Quyet			JEV			Fine-tuning BERT		
	Acc	WF1	**MF1**	Acc	WF1	**MF1**	Acc	WF1	**MF1**	Acc	WF1	**MF1**	Acc	WF1	**MF1**
Topic classification															
NEU-ESC	26.49	30.02	14.79	62.30	58.80	27.10	47.19	49.19	23.08	51.79	55.83	37.31	78.65	–	**62.73**
UIT-VSFC	49.15	57.08	46.02	44.91	55.03	48.99	82.56	81.86	63.63	83.77	84.61	74.36	88.22	–	**78.08**
Emotion recognition															
UIT-VSMEC	34.34	35.80	31.89	18.76	6.13	5.10	56.28	56.59	51.76	55.70	54.58	49.99	64.65	63.80	**59.06**
Hate speech detection															
ViHSD	58.95	65.38	38.35	10.33	1.98	6.25	29.12	37.45	30.59	83.70	83.87	59.34	86.75	86.52	**64.76**
UIT-ViCTSD	87.80	85.61	58.08	87.30	83.88	50.90	89.50	88.84	69.89	89.50	88.73	69.32	90.78	90.27	**71.31**
UIT-ViOCD	51.55	38.55	39.00	63.39	58.17	58.41	87.07	86.99	86.98	92.90	92.86	92.85	94.35	94.35	**94.35**
Sentiment analysis															
UIT-VSFC	49.81	58.71	45.67	42.89	50.80	39.49	75.87	81.12	64.46	86.58	88.44	73.44	93.87	93.45	**82.83**
Intent detection															
PhoATIS	15.79	17.94	19.86	23.74	26.29	27.72	–	–	–	81.08	87.60	**79.26**	90.82	90.82	–
ViMQ	44.00	44.71	37.96	48.40	46.48	38.77	38.70	31.34	37.15	79.90	79.42	80.20	–	–	**89.50**
Natural language inference															
ViNLI	47.61	46.84	46.66	28.28	23.86	23.63	51.19	48.69	48.48	79.91	80.85	**81.00**	80.67	–	80.69
ViANLI	31.80	31.21	31.21	32.10	20.24	20.25	34.40	33.02	33.01	53.80	53.48	**53.49**	–	–	–
ViMedNLI	58.44	57.28	57.28	36.78	23.99	23.99	59.07	59.03	59.03	73.14	73.40	73.40	73.70	–	**73.70**
vnNLI	71.99	71.79	71.79	40.24	32.20	32.19	75.34	75.24	75.24	85.71	85.89	**85.90**	73.76	–	73.80
Mean (12 shared)	50.99	51.92	43.23	42.97	38.46	31.26	60.52	60.78	53.61	76.37	76.83	**69.22**	–	–	–

Decide (287M parameters)[2], and Quyet-1.0-Small (328M parameters)[3]. Quyet supports a maximum of ten candidate labels and is therefore excluded from PhoATIS, which contains 28 intent labels.

3.3 Zero-Shot Evaluation Protocol

All models are evaluated with frozen weights, without task-specific fine-tuning or labeled demonstrations. Each model receives the task input and its predefined candidate labels. To examine the effects of label-description language, three configurations are considered: English (EN), Vietnamese (VI), and bilingual English–Vietnamese (EN+VI). Label identities and their order remain fixed across configurations, while EN+VI combines the English and Vietnamese descriptions. English descriptions serve as the reference configuration for the main evaluation.

Macro-F_1 is the primary evaluation metric, with accuracy and weighted-F_1 reported as supplementary measures. Macro-F_1 is computed over classes represented in the gold test set. Robustness is measured by the change in macro-F_1 between original and perturbed inputs under the same model configuration. The perturbation procedures are described in Section 5.

[2] https://huggingface.co/fastino/GLiNER2.5-multi-Decide

[3] https://huggingface.co/chinhnc/Quyet-1.0-Small

4 Results and Analysis

4.1 Overall Performance Across Tasks

Table 1 presents the zero-shot classification results under English label descriptions. JEV achieves the highest macro-F_1 in 11 of the 13 settings, with an average of 69.22 across the 12 settings shared by all models, followed by Quyet (53.61), LAYA (43.23), and GLiNER (31.26). Its advantage is particularly pronounced in intent classification, achieving 79.26 macro-F_1 on PhoATIS and 80.20 on ViMQ, exceeding the next-best models by 51.54 and 41.43 points, respectively. Notably, PhoATIS contains 28 candidate intent classes, indicating strong performance even with a relatively large label inventory. Substantial variation is also observed within task families. In topic classification, JEV achieves 74.36 macro-F_1 on UIT-VSFC but only 37.31 on

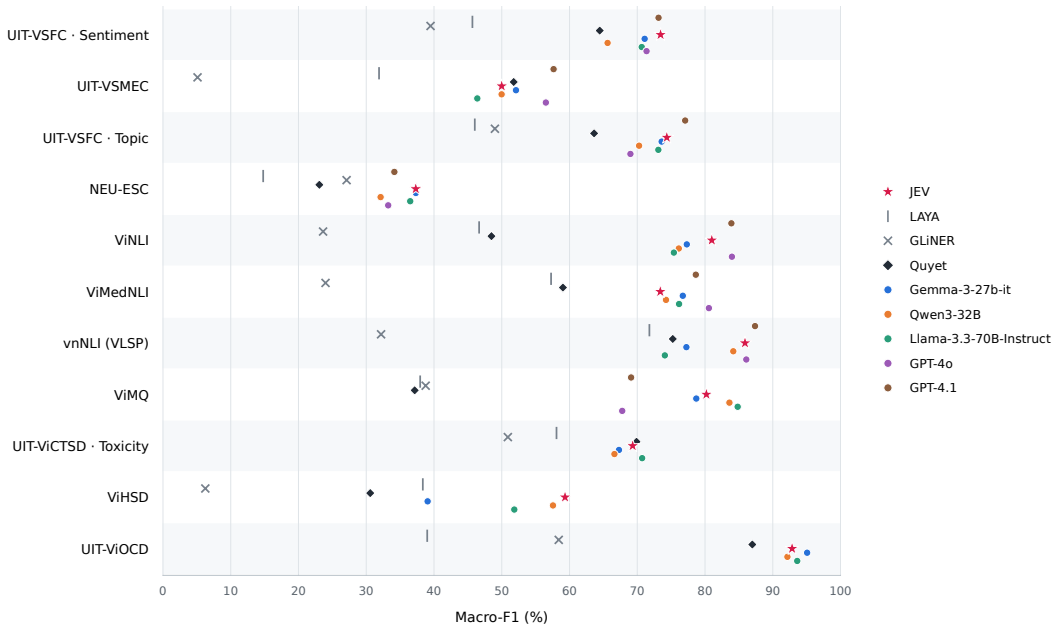

Figure 1: Zero-shot macro-F_1 (%) of specialized decision models and general-purpose LLMs across 11 Vietnamese classification settings.

NEU-ESC. Across the NLI benchmarks, all four models obtain their lowest scores on ViANLI, a dataset constructed using adversarial examples to challenge existing NLI systems. Nevertheless, JEV does not consistently outperform the other models, with Quyet achieving slightly higher scores on UIT-VSMEC and UIT-ViCTSD. Differences are also observed across evaluation metrics. On NEU-ESC, GLiNER achieves the highest accuracy (62.30%), whereas JEV obtains the highest macro-F_1 (37.31%), illustrating how model rankings can differ depending on the evaluation metric.

4.2 Comparison with LLMs

To contextualize the performance of specialized decision models, we compare their zero-shot classification results with those reported for five general-purpose LLMs in prior work (Thin et al., 2024). Figure 1 presents macro-F_1 scores on 11 overlapping classification settings, with detailed results provided in Table 10 in the appendix. The comparison uses English label descriptions for the decision models and previously reported zero-shot prompting results for the LLMs.

JEV demonstrates competitive performance against substantially larger general-purpose models, outperforming all three open-weight LLMs on five of the 11 settings: UIT-VSFC topic and sentiment classification, ViHSD, ViNLI, and vnNLI. It also achieves the highest reported macro-F_1 on ViHSD (59.34), exceeding Qwen3-32B (57.56), and on UIT-VSFC sentiment classification (73.44), slightly surpassing GPT-4.1 (73.14). On NEU-ESC, JEV performs comparably to the strongest LLM, Gemma-3-27b-it (37.31 versus 37.32). Among the remaining decision models, Quyet is particularly competitive on UIT-ViCTSD, reaching 69.89 macro-F_1 compared with 70.72 for Llama-3.3-70B-Instruct. Nevertheless, the relative performance varies across tasks. On UIT-VSMEC, GPT-4.1 exceeds JEV by 7.66 macro-F_1 points, while GPT-4o achieves a 7.17-point advantage on ViMedNLI. Llama-3.3-70B-Instruct also outperforms JEV on ViMQ by 4.62 points.

4.3 Error Analysis

On PhoATIS, JEV frequently predicts compound intents for queries annotated with a single intent. Of its 169 incorrect predictions, 140 include the gold intent together with an additional intent. The most frequent confusion is $\text{flight} \rightarrow \text{flight\#flight_time}$ (109 cases). For example, JEV assigns $\text{flight\#flight_time}$ to the query "cho tôi biết các chuyến bay từ thành phố hà nội đến thanh hóa khởi hành sau 5 giờ chiều thứ tư" ("show me flights from Hanoi to Thanh Hoa departing after 5 p.m. on Wednesday"), although its gold label is flight. The departure time specifies a search constraint, not a separate request for flight-time information. JEV correctly identifies 19 of the

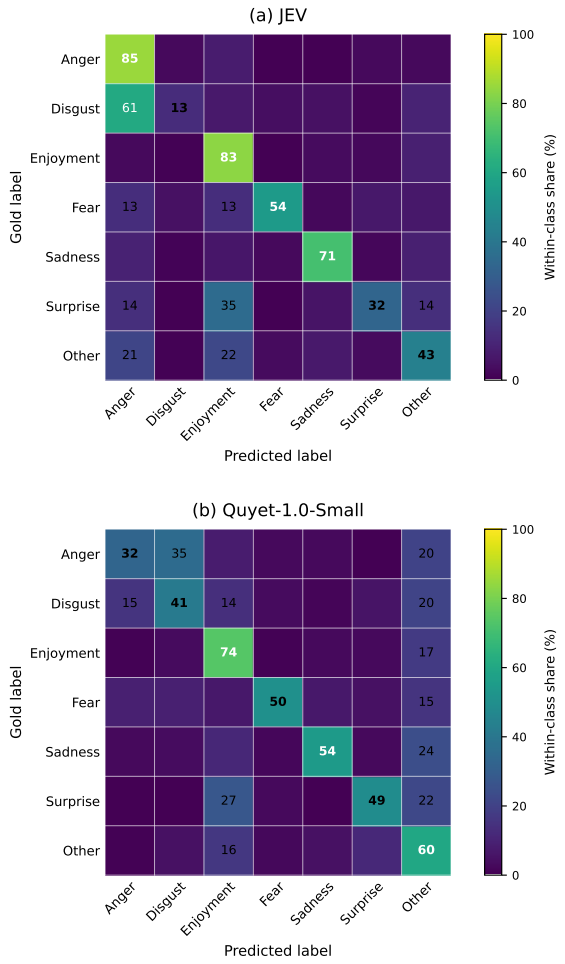

Figure 2: Comparison of JEV and Quyet confusion matrices on UIT-VSMEC.

20 examples annotated with compound intents, suggesting that distinguishing requested information from query constraints is a prominent difficulty in this setting. However, the small number of compound examples limits conclusions about compound-intent recognition more broadly.

On UIT-VSMEC, a prominent difference between JEV and Quyet concerns the distinction between *anger* and *disgust*. JEV misclassifies 80 of the 132 *disgust* examples as *anger*, obtaining an F_1 of 21.94 for *disgust*, compared with 49.09 for Quyet. Conversely, JEV achieves higher recall for *anger* (85.00% versus 32.50%). Figure 2 shows these contrasting confusion profiles, which are not reflected in the relatively close overall macro-F_1 scores.

Another recurring error pattern is the disproportionate selection of particular labels. On UIT-VSFC sentiment classification, all four models overpredict *neutral*, assigning it to 395–1,875 examples against 167 gold instances, with neutral precision ranging from 7.36% to 27.09%. For example, the positive comment "trình bày vấn đề một cách chính xác và dễ hiểu" ("presents the material accurately and clearly") is classified as *neutral* by JEV, LAYA, and GLiNER. On ViNLI, 551 of JEV's 601 errors involve predicting *neutral* for examples with other gold labels. More extreme prediction concentration occurs for GLiNER, which assigns *other* to 692 of 693 UIT-VSMEC examples and *hate* to 6,678 of 6,680 ViHSD examples. These distributions indicate that, in some evaluated configurations, errors are dominated by repeated selection of particular labels rather than confusions spread across multiple categories.

4.4 Effects of Label-Description Language

Table 2 shows distinct patterns of sensitivity to label-description language across models. JEV exhibits the most stable average performance, with macro-F_1 varying by only 0.35 points across the three configurations and reaching its highest value under EN+VI descriptions (69.44). Quyet and LAYA achieve their highest average scores with Vietnamese descriptions, although the improvements over English are modest and both models perform substantially worse with bilingual descriptions. GLiNER, in contrast, performs best with English descriptions and exhibits the largest average degradation under VI and EN+VI. These results suggest that multilingual label descriptions do not necessarily benefit models built on multilingual encoders.

The effects are also strongly task-dependent. Quyet, for example, gains 21.06 macro-F_1 points on ViNLI with Vietnamese descriptions, largely through improved discrimination between *other* and *neutral*. However, the same configuration reduces its performance on ViHSD by 14.23 points, accompanied by increased misclassification of *clean* examples as *offensive*. Additional instance-level prediction comparisons are provided in Appendix A.1. This contrast suggests that the effect of description language is closely related to the distinctions between task-specific labels, rather than a general preference for Vietnamese descriptions. The degradation observed for LAYA and GLiNER is particularly pronounced in some settings. LAYA loses 18.70 macro-F_1 points on ViMQ with bilingual descriptions, while GLiNER assigns *other* to every ViNLI example under both Vietnamese and bilingual descriptions. These findings indicate that

Table 2: Changes in macro-F_1 relative to English label descriptions (percentage points).

Dataset	LAYA		GLiNER		Quyet		JEV	
	VI	EN+VI	VI	EN+VI	VI	EN+VI	VI	EN+VI
Topic classification								
NEU-ESC	+6.06	−4.07	−1.37	+1.32	+3.63	−3.18	−0.67	−0.32
UIT-VSFC	−5.24	−15.08	−11.65	−13.68	+2.96	−1.42	+0.26	+0.84
Emotion recognition								
UIT-VSMEC	−2.36	−4.83	+0.57	+0.59	+3.38	−11.86	+1.42	+1.55
Hate speech detection								
ViHSD	+0.68	−4.02	−0.01	−0.01	−14.23	−18.94	−0.62	+0.34
UIT-ViCTSD	−1.70	−0.69	−3.31	−2.23	−1.15	−3.86	+0.14	+0.77
UIT-ViOCD	+13.33	−4.25	−9.64	−23.48	−1.93	−19.50	−0.19	+0.37
Sentiment analysis								
UIT-VSFC	+8.00	+5.28	−4.43	−23.44	−6.66	−11.18	−1.32	−0.22
Intent detection								
PhoATIS	+2.14	0.00	−4.68	−3.41	–	–	−4.82	−1.09
ViMQ	−2.37	−18.70	−22.46	−22.06	−0.86	−7.24	−0.48	+1.15
Natural language inference								
ViNLI	+4.56	−1.16	−13.85	−13.85	+21.06	+15.97	+1.50	+0.07
ViANLI	−0.55	−0.23	−2.61	−2.68	−1.65	+0.05	−0.12	−0.77
ViMedNLI	−2.67	−4.78	−6.45	−6.45	+1.03	−3.87	−0.67	−0.81
vnNLI	−0.92	−1.27	−12.74	−12.01	−0.29	−0.88	−0.73	−0.33
Mean (12 shared)	+1.40	−4.48	−7.33	−9.83	+0.44	−5.49	−0.12	+0.22

label-description language affects not only classification performance but also the distribution of predicted labels, with substantial differences across models and tasks.

5 Robustness to Vietnamese Linguistic Variation

5.1 Perturbation Design

We consider three types of Vietnamese linguistic variation. *Teencode* introduces informal lexical substitutions commonly found in online communication, such as "không" → "ko" ("not"). *Typographical noise* simulates typing errors through Telex-related substitutions, diacritic modifications, and character-level edits. *Diacritic removal* removes Vietnamese tone and vowel diacritics, including the conversion of *đ* to *d*. Each perturbation is applied independently to the original test inputs while retaining the original labels. For NLI, teencode and diacritic removal affect both the premise and hypothesis, whereas typographical noise is applied only to the hypothesis.

Table 3 summarizes perturbation coverage and character edit burden across the 13 evaluation settings, with dataset-level statistics provided in Appendix C. Coverage measures the proportion of modified inputs, while edit burden is the total character-level Levenshtein distance normalized by the original input length. Diacritic removal produces the highest coverage and edit burden, followed by typographical noise and teencode.

Table 3: Perturbation coverage and character edit burden across the 13 evaluation settings.

Perturbation	Changed (%)	Edit burden (%)
Teencode	61.70	2.14
Typo	90.18	4.93
No accents	98.10	18.28

5.2 Performance Under Linguistic Perturbations

Table 4 reports changes in macro-F_1 relative to unperturbed inputs. For each model, we use the label-description configuration that achieves the highest mean macro-F_1 on the original test data (EN+VI for JEV, EN for GLiNER, and VI for LAYA and Quyet) and keep it fixed across all perturbation conditions. JEV maintains the strongest overall performance across all three perturbation types. Its mean macro-F_1 across the 12 shared settings decreases from 69.44 on the original inputs to 59.93 under diacritic removal, 68.34 under teencode, and 66.37 under typographical noise, remaining above

Table 4: Changes in macro-F_1 under linguistic perturbations relative to original inputs (percentage points).

Dataset	LAYA			GLiNER			Quyet			JEV		
	NoAcc	Teen	Typo	NoAcc	Teen	Typo	NoAcc	Teen	Typo	NoAcc	Teen	Typo
Topic classification												
NEU-ESC	−9.93	−0.31	−3.42	−12.74	−0.10	−2.30	−9.68	+0.02	−1.38	−2.48	−0.38	−0.62
UIT-VSFC	−27.08	−0.59	−4.31	−46.60	+0.50	−7.86	−29.42	+0.08	−4.08	−13.22	+0.36	−3.36
Emotion recognition												
UIT-VSMEC	−15.85	−1.74	−6.85	−0.62	−0.62	+0.59	−26.38	−4.93	−9.80	−18.64	−5.13	−6.43
Hate speech detection												
ViHSD	−8.50	−0.45	−3.41	−0.02	−0.01	0.00	−5.93	−0.34	−1.38	−16.07	−2.39	−4.97
UIT-ViCTSD	−2.75	−0.39	−1.40	−3.81	+1.41	−1.95	−15.83	+0.29	−4.71	−8.23	−0.62	−2.60
UIT-ViOCD	+1.88	+2.67	+9.58	−25.44	−6.28	−8.97	−24.79	−1.89	−6.56	−11.49	−0.71	−4.84
Sentiment analysis												
UIT-VSFC	−13.36	−1.98	−3.23	−34.12	−4.12	−8.85	−14.61	−0.62	−3.26	−12.39	−1.02	−4.11
Intent detection												
PhoATIS	−19.27	+0.63	−3.53	−22.51	−0.46	−5.74	–	–	–	−12.04	+1.46	−3.67
ViMQ	−12.54	+1.63	−1.92	−21.98	+2.66	−4.34	−5.08	+1.78	−3.61	−4.94	−0.39	−2.07
Natural language inference												
ViNLI	−21.76	+1.68	−4.10	−11.62	−1.51	−4.65	−12.55	+1.71	−3.19	−6.07	−0.57	−1.61
ViANLI	−4.67	+1.68	−1.17	−2.51	−0.99	−2.01	−5.73	−0.24	−2.00	−7.95	+0.27	−0.66
ViMedNLI	−19.05	−5.34	−6.20	−4.67	−1.09	−1.53	−11.19	−0.79	−6.22	−9.15	−1.78	−4.01
vnNLI	−31.62	−3.13	−8.01	−12.91	−1.92	−5.43	−24.48	−0.76	−8.31	−3.46	−0.85	−1.50
Mean (12 shared)	−13.77	−0.52	−2.87	−14.75	−1.01	−3.94	−15.47	−0.47	−4.54	−9.51	−1.10	−3.07

the other models in every condition. JEV achieves the highest macro-F_1 in all 13 settings under diacritic removal and in 12 of 13 settings under each of the other two perturbations. Its lead over the next-best model, Quyet, also widens from 15.39 points on the original inputs to 21.35 points under diacritic removal.

Diacritic removal produces the largest average performance loss for every model, with decreases in 47 of the 48 model–setting comparisons shared by all four models. JEV exhibits the smallest mean reduction (9.51 points), followed by LAYA (13.77), GLiNER (14.75), and Quyet (15.47). The effects nevertheless vary substantially across tasks. On UIT-VSFC topic classification, GLiNER loses 46.60 points and predicts *others* for 3,164 of 3,166 instances; on the sentiment task, it predicts *neutral* for 3,112 instances, accompanied by a 34.12-point decline. Differences are also pronounced on vnNLI, where LAYA and Quyet lose 31.62 and 24.48 points, respectively, compared with only 3.46 for JEV. Prediction-level comparisons further show that diacritic removal changes 19.57% of JEV's decisions on average, compared with 40.60% for Quyet and 54.29% for LAYA. GLiNER exhibits a similarly low flip rate (19.83%), but its predictions remain concentrated on particular labels in several settings. Detailed prediction transitions are reported in Appendix A.2.

Teencode and typographical noise have smaller average effects, with mean reductions of 0.47–1.10 and 2.87–4.54 points, respectively. Although JEV retains the highest absolute performance, it does not consistently exhibit the smallest degradation: Quyet has the lowest mean loss under teencode, while LAYA is least affected by typographical noise. Notably, LAYA improves on UIT-ViOCD under all three perturbations, gaining 9.58 points under typographical noise. This improvement coincides with a reduction in its preference for *non_complaint*, from 480 of 549 predictions on the original inputs to 291 under typographical noise, closer to the 270 gold instances. Conversely, teencode causes notable declines on UIT-VSMEC for JEV (5.13 points) and Quyet (4.93 points). A targeted analysis of the Vietnamese negation marker *không* in Appendix A.3 further identifies differences in models' sensitivity to its informal written forms.

6 Conclusion

We presented a systematic evaluation of four specialized decision models for Vietnamese zero-shot text classification across 13 settings span-

ning six task families. Beyond overall classification performance, we examined the effects of label-description language, class-level prediction patterns, and robustness to Vietnamese linguistic variation. JEV achieved the highest macro-F_1 in 11 of the 13 settings under English label descriptions and maintained the strongest average performance across all three linguistic perturbations. However, model behavior varied considerably across tasks and evaluation conditions. Label-description language produced model- and task-dependent effects, with bilingual descriptions offering no consistent advantage. Diacritic removal caused the largest average performance degradation for all four models, although JEV exhibited the smallest mean reduction under this condition. These findings demonstrate that strong zero-shot performance does not necessarily imply consistent behavior across label configurations and linguistic variations, highlighting the importance of evaluating specialized decision models beyond aggregate benchmark scores.

References

Hoang Tuan Anh, Nguyen Thi Minh Huyen, Nguyen Lien, et al. 2022. VLSP 2021 – vnNLI challenge: Vietnamese and English-Vietnamese textual entailment. *VNU Journal of Science: Computer Science and Communication Engineering*, 38(2).

Seyedarmin Azizi, Erfan Baghaei Potraghloo, and Massoud Pedram. 2026. Labels override definitions in jev-style typed decision models. *Preprint*, arXiv:2610.02586.

Convai Innovations. 2026. Laya Multilingual: Model card. https://huggingface.co/convaiinnovations/laya-multilingual. Accessed 2026-10-11; exact evaluation revision must be recorded.

Mai Hoang Dao, Thinh Hung Truong, and Dat Quoc Nguyen. 2021. Intent detection and slot filling for Vietnamese. In *Proceedings of the 22nd Annual Conference of the International Speech Communication Association (INTERSPEECH)*.

Tobias Deußer, Lorenz Sparrenberg, and Rafet Sifa. 2026. Evaluating and benchmarking the system one model jev. *arXiv preprint arXiv:2609.37647*.

Feiyu Duan, Jiayu Lin, Jia Wang, Jun Xiang, Jialiang Wu, Xinnong Zhang, Hanqi Yan, Siyuan Wang, and Zhongyu Wei. 2026. General decision models: Benchmarking and insights beyond jev. *Preprint*, arXiv:2610.03935.

Tianxiang Gao, Jinzhe Li, Zhiyuan Li, Yi Chang, and Yuan Wu. 2026. More choices, fewer decisions: Ordinal-scale bias in JEV-like direct-decision models. *Preprint*, arXiv:2609.38827.

Vong Anh Ho, Duong Huynh-Cong Nguyen, Danh Hoang Nguyen, Linh Thi-Van Pham, Duc-Vu Nguyen, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2019. Emotion recognition for Vietnamese social media text. In *International Conference of the Pacific Association for Computational Linguistics (PACLING)*, pages 319–333. Springer.

Ta Duc Huy, Nguyen Anh Tu, Tran Hoang Vu, Nguyen Phuc Minh, Nguyen Phan, Trung H. Bui, and Steven Q. H. Truong. 2021. ViMQ: A Vietnamese medical question dataset for healthcare dialogue system development. In *International Conference on Neural Information Processing (ICONIP)*, pages 657–664. Springer.

Tin Van Huynh, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2022. ViNLI: A Vietnamese corpus for studies on open-domain natural language inference. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 3858–3872.

Tin Van Huynh, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2026. A new benchmark dataset and mixture-of-experts language models for adversarial natural language inference in vietnamese. *Expert Systems with Applications*, 306:130109.

Hazem Ibrahim and Yasir Zaki. 2026. Evaluating decision models for text annotation in computational social science. *Preprint*, arXiv:2609.24574.

Yifan Lan, Zhongye Liu, Yurui Chang, Hanyu Wang, Yongkang Du, Lu Lin, and Jinghui Chen. 2026. DecisionRobust: Benchmarking robustness for JEV style decision models across text and vision. *Preprint*, arXiv:2610.19118.

Jiawei Li. 2026. Fast models, slow evidence: A paired and self-audited evaluation of system-1 decision models for llm agent harnesses. *Preprint*, arXiv:2610.02267.

Xing Li, Qingcheng Chang, Jinzhong Ning, Changfeng Xu, Shenlong Zhang, Yijia Zhang, Ling Luo, and Hongfei Lin. 2026. Specialized decision models vs. general-purpose llms: Benchmarking jev across knowledge, reasoning, and multilingual tasks. *Preprint*, arXiv:2610.11978.

Yue Li, Zhixue Zhao, and Carolina Scarton. 2025. Label set optimization via activation distribution kurtosis for zero-shot classification with generative models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 31724–31741.

Son T. Luu, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2021. A large-scale dataset for hate speech detection on Vietnamese social media texts. In *Advances and Trends in Artificial Intelligence. Artificial Intelligence Practices (IEA/AIE)*, pages 415–426. Springer.

Phan Quoc Hung Mai, Quang Hung Nguyen, Phuong Giang Duong, Hong Hanh Nguyen, and Nguyen Tuan Long. 2025. NEU-ESC: A comprehensive Vietnamese dataset for educational sentiment analysis and topic classification toward multitask learning. *arXiv preprint arXiv:2506.23524*.

Yida Mu, Ben P. Wu, William Thorne, Ambrose Robinson, Nikolaos Aletras, Carolina Scarton, Kalina Bontcheva, and Xingyi Song. 2024. Navigating prompt complexity for zero-shot classification: A study of large language models in computational social science. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 12074–12086.

Kiet Van Nguyen, Vu Duc Nguyen, Phu X. V. Nguyen, Tham T. H. Truong, and Ngan Luu-Thuy Nguyen. 2018. UIT-VSFC: Vietnamese students' feedback corpus for sentiment analysis. In *2018 10th International Conference on Knowledge and Systems Engineering (KSE)*, pages 19–24. IEEE.

Luan Thanh Nguyen, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2021. Constructive and toxic speech detection for open-domain social media comments in Vietnamese. In *Advances and Trends in Artificial Intelligence. Artificial Intelligence Practices (IEA/AIE)*, pages 572–583. Springer.

Nam Nguyen, Thang Phan, Duc-Vu Nguyen, and Kiet Nguyen. 2023. ViSoBERT: A pre-trained language model for Vietnamese social media text processing. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5191–5207, Singapore. Association for Computational Linguistics.

Thanh-Nhi Nguyen, Thanh-Phong Le, and Kiet Nguyen. 2024. ViLexNorm: A lexical normalization corpus for Vietnamese social media text. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1421–1437.

Thi-Hong Nguyen Nhung, Phuong Phan-Dieu Ha, Luan Thanh Nguyen, Kiet Van Nguyen, and Ngan Luu-Thuy Nguyen. 2021. Vietnamese complaint detection on e-commerce websites. In *New Trends in Intelligent Software Methodologies, Tools and Techniques (SoMeT)*.

Long Phan, Tai Dang, Hieu Tran, Trieu H. Trinh, Vy Phan, Lam D. Chau, and Minh-Thang Luong. 2023. Enriching biomedical knowledge for low-resource language through large-scale translation. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 3131–3142.

Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. Beyond accuracy: Behavioral testing of NLP models with CheckList. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4902–4912.

Dang Van Thin, Duong Ngoc Hao, and Ngan Luu-Thuy Nguyen. 2024. Prompt engineering with large language models for Vietnamese sentiment classification. In *Proceedings of the 38th Pacific Asia Conference on Language, Information and Computation*, pages 181–192, Tokyo, Japan. Tokyo University of Foreign Studies.

TypeSafe AI. 2026. Jev: Introduction. https://docs.typesafe.ai/introduction. Accessed 2026-10-08.

Xiao Wang, Qin Liu, Tao Gui, Qi Zhang, Yicheng Zou, Xin Zhou, Jiacheng Ye, Yongxin Zhang, Rui Zheng, Zexiong Pang, Qinzhuo Wu, Zhengyan Li, Chong Zhang, Ruotian Ma, Zichu Fei, Ruijian Cai, Jun Zhao, Xingwu Hu, Zhiheng Yan, and 15 others. 2021. TextFlint: Unified multilingual robustness evaluation toolkit for natural language processing. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing: System Demonstrations*, pages 347–355.

Urchade Zaratiana, Gil Pasternak, Oliver Boyd, George Hurn-Maloney, and Ash Lewis. 2025. Gliner2: An efficient multi-task information extraction system with schema-driven interface. *Preprint*, arXiv:2507.18546.

Fan Zhang, Yankai Chen, Zhuohan Xie, Yixi Zhou, Sijia Peng, Lei Fan, Xinhua Ji, Cunyuan Zheng, Huangyong Shan, Philip S. Yu, Xue Liu, Yu Chen, Preslav Nakov, and Songwei He. 2026. Same scores, different decisions: Evaluating JEV and language models for legal document understanding. *arXiv preprint arXiv:2609.27678*.

A Prediction-Level Sensitivity

We examine instance-level prediction changes under two evaluation conditions: variations in label-description language and linguistic perturbations of the input text. A prediction flip occurs when the predicted label differs between paired configurations. W→C, C→W, and W→W denote transitions from incorrect to correct, correct to incorrect, and between two incorrect predictions, respectively.

A.1 Label-Description Language

Table 5 presents selected prediction changes underlying the macro-F_1 differences reported in Section 4.4. On ViNLI, switching from English to Vietnamese descriptions changes 123 predictions for JEV but 930 for Quyet. For Quyet, this change corrects 682 previously incorrect predictions while introducing 151 new errors, consistent with its substantial improvement in macro-F_1. The effect

is reversed on ViHSD, where the same change introduces 1,247 errors and corrects only 147. For LAYA, switching from Vietnamese to bilingual descriptions on ViMQ produces 393 newly incorrect predictions compared with 88 corrections. GLiNER also exhibits substantial prediction changes on ViNLI, with 987 of its 2,266 flips occurring between two incorrect labels.

A.2 Linguistic Perturbations

Table 6 summarizes prediction changes under diacritic removal, the perturbation producing the largest performance degradation in Section 5.2. Results are averaged across the 12 shared evaluation settings, with each setting weighted equally. JEV exhibits considerably fewer prediction changes under diacritic removal than Quyet and LAYA. In particular, 11.73% of its predictions change from correct to incorrect, compared with 23.33% for Quyet and 27.95% for LAYA. Despite a similarly low flip rate (19.83%), GLiNER achieves substantially lower classification performance, with its predictions dominated by a single label in several settings. Its apparent prediction stability therefore partly reflects persistent label concentration.

A.3 Negation-Cue Sensitivity under Teencode

We further examine the sensitivity of NLI predictions to informal variants of the Vietnamese negation marker *không*. We identify 194 instances across the four NLI benchmarks in which *không* occurs in the hypothesis but not the premise and is replaced by a teencode variant, such as "ko" or "k", without modifying the remaining input text. Fixed expressions in which *không* does not function as negation, such as "không khí", "không gian", and "hàng không", are excluded. Replacing *không* substantially reduces the proportion of *contradiction* predictions for LAYA and GLiNER, by 20.62 and 14.43 percentage points, respectively. The corresponding changes are considerably smaller for JEV (3.09 points) and Quyet (2.06 points). LAYA and GLiNER also experience larger accuracy reductions on these instances. These results indicate greater sensitivity to the surface realization of negation cues in LAYA and GLiNER, even when the remaining input text is unchanged.

B Prompt and Schema Templates

Each evaluation setting specifies a task instruction, input fields, and an ordered set of candidate labels with descriptions. Single-text tasks use a text field, whereas NLI uses separate premise and hypothesis fields. Gold labels and example identifiers are excluded from model inputs. Figure 3 provides a schematic illustration using sentiment classification. The task specification is adapted to each model's inference interface. JEV, LAYA, and Quyet use typed-choice questions with instructions and criteria, while GLiNER uses a classification schema with label descriptions and single-label prediction. We evaluate three label-description configurations: English (EN), Vietnamese (VI), and bilingual English-Vietnamese (EN+VI), with English preceding Vietnamese in the bilingual configuration.

State
text: {input}

Choice question
type: choice
instructions:
Classify the overall sentiment expressed in Vietnamese student feedback. Select one label based on the provided definitions.

Candidate labels
negative: The author expresses dissatisfaction, criticism, disapproval, or a negative evaluation.

neutral: The text is factual or descriptive without a clear positive or negative evaluation.

positive: The author expresses satisfaction, praise, approval, or a positive evaluation.

Figure 3: Illustrative typed-choice schema for sentiment classification. The instruction is translated into English for readability.

C Perturbation Implementation Details

C.1 Perturbation Generation

Teencode and typographical perturbations are generated using nom-vn v0.2.37, while diacritic removal uses the character mapping provided by pyvi v0.1.1. All inputs are normalized to NFC before transformation. Table 8 summarizes the stochastic generator parameters for teencode and typographical noise. OCR and word-segmentation perturbations are disabled in both conditions. Teencode uses dictionary-based slang substitutions, whereas typographical noise combines Telex-related errors, diacritic modifications, confusion-set substitutions, and character-level operations.

Character-level operations are evaluated sequentially using a shared random draw, so the thresholds in Table 8 do not represent independent mutation probabilities. The max_edit_ratio param-

Table 5: Selected instance-level prediction changes across label-description configurations.

Model	Dataset	Description change	Flips / n (%)	W→C	C→W	W→W	ΔMF1
LAYA	ViMQ	VI → EN+VI	760 / 1,000 (76.00)	88	393	279	−16.33
GLiNER	ViNLI	EN → VI	2,266 / 2,991 (75.76)	580	699	987	−13.85
Quyet	ViNLI	EN → VI	930 / 2,991 (31.09)	682	151	97	+21.06
Quyet	ViHSD	EN → VI	1,555 / 6,680 (23.28)	147	1,247	161	−14.23
JEV	ViNLI	EN → VI	123 / 2,991 (4.11)	84	28	11	+1.50

Table 6: Instance-level prediction transitions under diacritic removal (%).

Model	Flips	W→C	C→W	W→W
LAYA	54.29	12.15	27.95	14.19
GLiNER	19.83	3.33	12.29	4.21
Quyet	40.60	8.13	23.33	9.15
JEV	19.57	4.21	11.73	3.63

Table 7: Effects of rewriting *không* on NLI predictions (194 paired instances).

	Contradiction (%)		**Accuracy (%)**	
Model	Original	Teencode	Original	Teencode
LAYA	52.06	31.44	67.01	44.85
GLiNER	23.71	9.28	35.05	24.23
Quyet	58.76	56.70	74.23	70.10
JEV	50.00	46.91	82.99	78.35

eter limits the internal operation budget for each unprotected segment, with a minimum budget of one, rather than imposing a strict bound on Levenshtein distance. Both stochastic transformations protect URLs, mentions, hashtags, digit-containing expressions, and emoji or symbol sequences. Typographical noise additionally protects selected Vietnamese negation forms, including “không”, “chưa”, “chẳng”, “chả”, and “đừng”. Random generation uses a base seed of 42, with seeds derived from example identifiers, input fields, and segment positions.

C.2 Perturbation Statistics

Table 9 reports the realized perturbation coverage and character edit burden for each evaluation setting. Coverage denotes the proportion of inputs modified, while edit burden is calculated as the sum of field-level Levenshtein distances divided by the total original input length. Both measures are computed on NFC-normalized strings. For NLI, the edit denominator includes both the premise and hypothesis, including the unchanged premise under

Table 8: Stochastic perturbation generator parameters.

Operation / parameter	Teencode	Typo
Token-level operations		
Slang substitution	0.80	–
Telex-related errors	–	0.08
Full diacritic removal	–	0.05
Partial diacritic removal	–	0.03
Confusion-set substitution	–	0.02
Character-level thresholds		
Keyboard substitution	–	0.015
Adjacent-character swap	–	0.025
Character deletion	–	0.035
Character duplication	–	0.040
max_edit_ratio	0.50	0.20

typographical perturbations. The statistics cover 13 evaluation settings, comprising 33,350 setting–example pairs. UIT-VSFC is counted separately for topic and sentiment classification.

Table 9: Perturbation coverage and character edit burden by evaluation setting (%).

Dataset	n	Teencode		Typo		No accents	
		Changed	Edit	Changed	Edit	Changed	Edit
Topic classification							
NEU-ESC	6,613	68.30	3.57	90.20	7.42	99.17	20.42
UIT-VSFC	3,166	42.61	2.20	90.24	7.49	99.87	20.44
Emotion recognition							
UIT-VSMEC	693	71.72	5.05	87.73	7.73	99.86	21.62
Hate speech detection							
ViHSD	6,680	49.03	3.30	78.95	7.23	91.53	18.99
UIT-ViCTSD	1,000	73.90	2.99	95.00	7.55	99.90	21.24
UIT-ViOCD	549	75.59	3.83	91.07	7.55	99.82	21.49
Sentiment analysis							
UIT-VSFC	3,166	43.02	2.26	89.70	7.47	99.87	20.44
Intent detection							
PhoATIS	893	51.96	1.66	94.96	7.35	100.00	19.70
ViMQ	1,000	56.20	2.73	95.60	7.08	100.00	20.47
Natural language inference							
ViNLI	2,991	80.51	1.62	98.90	3.16	100.00	19.78
ViANLI	1,000	76.40	1.55	98.40	2.74	100.00	19.69
ViMedNLI	1,422	69.48	2.43	82.49	1.83	100.00	20.51
vnNLI	4,177	77.47	1.11	99.64	3.32	100.00	14.31

Table 10: Zero-shot macro-F_1 on 11 overlapping classification settings. Decision models use English label descriptions, while LLM results are taken from prior zero-shot prompting experiments. Bold and underlined values indicate the highest and second-highest reported scores per setting, respectively. Unreported results are denoted by –.

Dataset	Large Language Models					Decision Models			
	Gemma-3	Qwen3	Llama-3.3	GPT-4o	GPT-4.1	LAYA	GLiNER	Quyet	JEV
Topic classification									
NEU-ESC	**37.32**	32.12	36.49	33.24	34.15	14.79	27.10	23.08	<u>37.31</u>
UIT-VSFC	73.60	70.27	73.11	68.99	**77.07**	46.02	48.99	63.63	<u>74.36</u>
Emotion recognition									
UIT-VSMEC	52.10	49.98	46.39	<u>56.51</u>	**57.65**	31.89	5.10	51.76	49.99
Hate speech detection									
ViHSD	39.06	<u>57.56</u>	51.85	–	–	38.35	6.25	30.59	**59.34**
UIT-ViCTSD	67.30	66.63	**70.72**	–	–	58.08	50.90	<u>69.89</u>	69.32
UIT-ViOCD	**95.07**	92.16	<u>93.61</u>	–	–	39.00	58.41	86.98	92.85
Sentiment analysis									
UIT-VSFC	71.09	65.62	70.66	71.37	<u>73.14</u>	45.67	39.49	64.46	**73.44**
Intent detection									
ViMQ	78.71	<u>83.60</u>	**84.82**	67.78	69.09	37.96	38.77	37.15	80.20
Natural language inference									
ViNLI	77.31	76.13	75.40	**83.97**	<u>83.89</u>	46.66	23.63	48.48	81.00
ViMedNLI	76.72	74.24	76.17	**80.57**	<u>78.64</u>	57.28	23.99	59.03	73.40
vnNLI	77.25	84.17	74.07	<u>86.09</u>	**87.38**	71.79	32.19	75.24	85.90