arXiv:2610.30828v1 [cs.CL] 22 Oct 2026

TutorCaseBench: A Case-Grounded Benchmark for Evaluating LLM Tutors in Supporting Students with Challenging Behaviors

Yeon Su Park[♥] **Chanyoung Park**[♡] **Tae Soo Kim**[♥] **Seungju Kim**[♠]
Hyunwoo Kim[♥] **Daho Jung**[♥] **Jinsung Yoon**[♣] **Juho Kim**[♥,◇]

[♥]KAIST [♡]Seoul National University [♣]Google Cloud AI Research
[♠]Korea Institute for Curriculum and Evaluation [◇]SkillBench

{yeonsupark, taesoo.kim, khw0726, dhj1104, juhokim}@kaist.ac.kr
pcy1203@snu.ac.kr seun8ju@kice.re.kr jsyoon0823@gmail.com

Abstract

LLMs are increasingly used as AI tutors, yet existing benchmarks overlook the *challenging behaviors* common among students with learning and behavioral difficulties. We introduce TutorCaseBench, a benchmark that grounds AI-tutor evaluation in qualitative case studies from pedagogy, educational psychology, and special education. Our pipeline converts documented student categories into case-specific scenarios, tutor–student dialogues, and rubrics grounded in source evidence. The benchmark contains 923 dialogues across five challenge types, evaluated through open-ended response generation (1,560 items) and multiple-choice selection (2,537 items). Across 11 state-of-the-art LLMs, even the strongest model satisfies a rubric criterion in only 55.8% of open-ended cases, while models are better at selecting appropriate strategies than generating situated responses. These results reveal a gap between intervention recognition and context-specific adaptation, especially for open-weight models. All data are publicly available at https://github.com/kixlab/TutorCaseBench.

1 Introduction

Large language models (LLMs) are increasingly deployed as conversational AI tutors for scalable, individualized instruction (Jurenka et al., 2025). Yet, effective tutoring requires more than providing correct answers: a tutor who misreads emotion, fails to re-engage, or escalates refusal can demotivate the student, making affective and behavioral factors as consequential as cognitive ones (Baker et al., 2010; D'Mello and Graesser, 2013; Karumbaiah et al., 2023). This is especially important for students with learning and behavioral difficulties, whose struggles often appear as refusal, withdrawal, emotional distress, or social conflict (Sugai and Horner, 2006). We use *challenging behaviors* to refer to such disruptive or disengaged responses, which make tutoring difficult to sustain and require

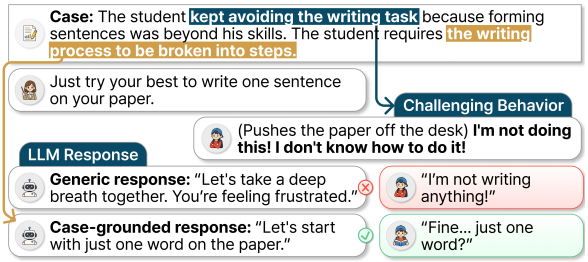

Figure 1: We introduce TutorCaseBench, a benchmark grounded in qualitative case studies for evaluating LLM tutors' ability to respond to challenging student behaviors. Example from Park and Pinkelman (2017).

support tailored to the learner and context (Pepe and Addimando, 2013; Patnaik and Subban, 2023).

Existing AI-tutor benchmarks (Macina et al., 2023, 2025; Sonkar et al., 2023) focus on cognitive aspects of tutoring, such as math-error detection and scaffolding. Their inputs often depict idealized student behavior rather than disengagement, resistance, or other challenging behaviors (Macina et al., 2023; Markel et al., 2023; Liu et al., 2024b), while their rubrics rely on common pedagogical criteria (Shridhar et al., 2022; Liu et al., 2024a). As a result, they offer limited evidence on whether AI tutors can provide case-sensitive support for students with challenging behaviors.

Our work grounds this evaluation in a resource the NLP community has so far overlooked: qualitative case literature from pedagogy, educational psychology, and special education. These studies document individual students, their behaviors in learning contexts, and interventions that practitioners observed to be effective or ineffective. Unlike hypothetical personas, such cases provide evidence about student backgrounds, precipitating events, and practitioner responses, allowing us to synthesize realistic pedagogical scenarios involving challenging behaviors and documented interventions.

We introduce an automated LLM pipeline that converts qualitative case-study papers into struc-

tured tutoring evaluation cases. For each documented student category, the pipeline extracts verbatim learner and instructional evidence and constructs scenarios with learning contexts, challenging behaviors, and intervention strategies. Each scenario is then converted into a grounded tutor–student dialogue ending with the student exhibiting challenging behavior, paired with a rubric derived from the documented interventions. The resulting items are evaluated in *immediate* and *long-term* settings through an LLM-as-a-judge checklist rubric (Zheng et al., 2023) and a multiple-choice task in the style of Hendrycks et al. (2021).

Applying this pipeline to qualitative case studies, we construct TutorCaseBench, a benchmark of 923 English dialogues across five challenge types: emotional distress, task resistance, social-interpersonal difficulty, attention/engagement failure, and academic struggle. It contains 1,560 open-ended and 2,537 multiple-choice evaluation items.

We evaluate 11 state-of-the-art LLMs on TutorCaseBench and find that they remain limited in providing case-specific support. Even the strongest model satisfies a documented intervention in only 55.8% of open-ended cases, while models perform better at selecting provided strategies than at generating tailored responses. Our benchmark provides systematic evidence that LLM tutors struggle to support students with challenging behaviors even in controlled settings, motivating further investigation in more interactive and situated tutoring contexts.

Contributions.

1. A case-grounded pipeline that converts pedagogical case studies into tutoring scenarios, dialogues, and evaluation criteria.

2. TutorCaseBench, a benchmark of 923 case-grounded dialogues covering five types of challenging student behavior, with both open-ended and multiple-choice evaluation settings.

3. An empirical analysis of 11 LLM tutors, revealing difficulty in providing case-specific support, especially in long-term interventions.

2 Related Work

2.1 Literature-Grounded Student Simulations

AI tutor evaluation requires realistic student–tutor interactions, but collecting them at scale is difficult (Macina et al., 2023). LLM-based student simulation offers a useful alternative, typically modeling learner knowledge and reasoning, such as misconceptions and knowledge components (Macina et al., 2023; Lu and Wang, 2024; Qi et al., 2026), or learner traits, such as personas and motivational profiles (Markel et al., 2023; Liu et al., 2024b; Jin et al., 2025; Marquez-Carpintero et al., 2026).

However, even fine-tuned models show limited alignment with real student behavior (Scarlatos et al., 2026). Simulated students also rarely express frustration or disengagement, remain unnaturally attentive and linguistically sophisticated, and accept tutor feedback without resistance (Martynova et al., 2025). Thus, existing frameworks mostly simulate idealized student behavior, overlooking disengaged, resistant, or off-task behaviors.

Source documents offer another way to ground dialogue synthesis. Prior work shows that document grounding improves coherence and factual consistency (Feng et al., 2020, 2021), and supports persona simulation when documents describe behavior and internal states (Wang et al., 2025). In education, textbook-grounded dialogue generation improves content accuracy (Wang et al., 2024a). Hence, we ground synthesis in pedagogy, educational psychology, and special education literature, enabling dialogues with realistic challenging behaviors beyond cooperative learner settings.

2.2 Evaluation of AI Tutors

AI tutor evaluation has shifted from general-purpose NLG metrics to pedagogy-grounded frameworks. Earlier approaches relied on BLEU (Papineni et al., 2002) and BERTScore (Zhang et al., 2020), comparing tutor responses against reference answers, but such metrics poorly capture pedagogical quality because multiple valid responses may exist (Tack and Piech, 2022; Jurenka et al., 2025).

To address this, prior work has proposed pedagogy-grounded criteria. Tack and Piech (2022) found that AI responses fell short of human teachers in teacher-likeness, student understanding, and helpfulness. Later frameworks introduced criteria such as coherence, correctness, equitable tutoring (Macina et al., 2023), as well as usefulness, care, and human-likeness (Wang et al., 2024b). However, these criteria remain fragmented and lack consensus (Jurenka et al., 2025). Recent benchmarks have begun to consolidate them, including MRBench's eight-dimensional taxonomy (Maurya et al., 2025) and MathTutorBench's focus on mathematical expertise, student understanding, and scaffolding quality (Macina et al., 2025).

Despite this progress, existing benchmarks largely use fixed criteria, although support for students exhibiting challenging behaviors is often context-dependent. Thus, we take a case-specific approach, deriving evaluation criteria from interventions documented in pedagogical case studies.

3 Data Construction

We construct TutorCaseBench from case-study papers documenting diverse learning and behavioral challenges, using an automated LLM pipeline to convert them into structured tutor–student dialogues in English for evaluation. The benchmark includes two tasks: open-ended response generation and multiple-choice question (MCQ) selection.

We use `gemini-3-flash-preview` (Google DeepMind, 2025) for extractive stages (Stage 1–2) and `gemini-3.1-pro-preview` (Google DeepMind, 2026a) for generative stages (Stage 3–5), with default thinking `high` and `temperature=0`. For consistency, we used the Gemini models across the pipeline and rerouted only a number of persistent failures to OpenAI models (see Appendix A).

3.1 Paper Collection

Stage 0: Paper Collection. We collect candidate case-study papers from the ERIC database[1] using the ERIC API for papers that include at least one student *and* one methodology descriptor. The descriptors are listed in Appendix B. We restrict publications to 2010–2026 to keep the corpus contemporary while ensuring sufficient volume. Since descriptor matching does not ensure sufficient detail on student behavior and teacher actions, we further filter papers by their abstracts (see Appendix C).

3.2 Open-Ended Response Task

TutorCaseBench's open-ended response task evaluates whether a model produces a next tutor response aligned with documented interventions for a tutor–student dialogue grounded in case studies.

3.2.1 Tutor–Student Dialogue Construction

Stages 1–4 convert the collected papers into tutor–student dialogues paired with a hidden answer key of effective and ineffective interventions. Figure 2 provides an overview of this pipeline.

Stage 1: Student Category Extraction. Starting from the initial set of papers, the LLM rates how concretely each paper documents three dimensions essential to our benchmark—*student behavior*, *learning context*, and *teacher actions*—on a 1–5 scale (prompt in Appendix T.2). Papers scoring below 3 on any dimension are filtered out. We set this cutoff after inspecting samples and finding them too general to ground a concrete dialogue.

For papers that pass screening, the model extracts *student categories*—distinct student profiles within the paper, each with a label, rationale, demographics, and supporting excerpts for the three dimensions above (prompt in Appendix T.3). We use each category as a unit for scenario construction, since a paper may document multiple students or groups with different behaviors, contexts, or instructional needs. Categories are then filtered by the same dimensions at a stricter threshold (3 or below), since downstream dialogue synthesis requires per-category evidence to be concrete on its own.

Stage 2: Evidence Extraction. Next, the LLM reads each paper from Stage 1 twice—once for learner evidence, once for instructional evidence—tagging each verbatim span with the categories it applies to. The first read extracts learner evidence along four dimensions—*observable behaviors*, *cognitive states*, *affective states*, and *difficulty triggers*—along with verbatim student quotes (prompt in Appendix T.4). These dimensions adapt Pintrich's self-regulated learning framework (Pintrich, 2000; Panadero, 2017), replacing its contextual axis with difficulty triggers to surface what specifically prompted the student's struggle.

The second read takes the paper and the prior learner spans as input, and extracts spans of instructional evidence, recording *when* the strategy applies, *how* it was enacted, and its documented *outcome* (prompt in Appendix T.5). Each instructional span is tagged as *effective* or *ineffective*, given a temporal scope of *immediate*, *long-term*, or *both*.

Stage 3: Scenario Construction. The LLM then composes one or more concrete *scenarios* per category from the extracted evidence (prompt in Appendix T.6). We construct multiple scenarios when evidence spans distinct tasks or classroom contexts, rather than collapsing them into a generic scenario.

Each scenario is organized as two parallel blocks: a *content* block (task, classroom context, triggering behavior, related behavior patterns) that frames what a tutor sees in the moment, paired with a *hidden strategies* block (immediate effective, long-term effective, and ineffective) that is later used

[1] https://eric.ed.gov/

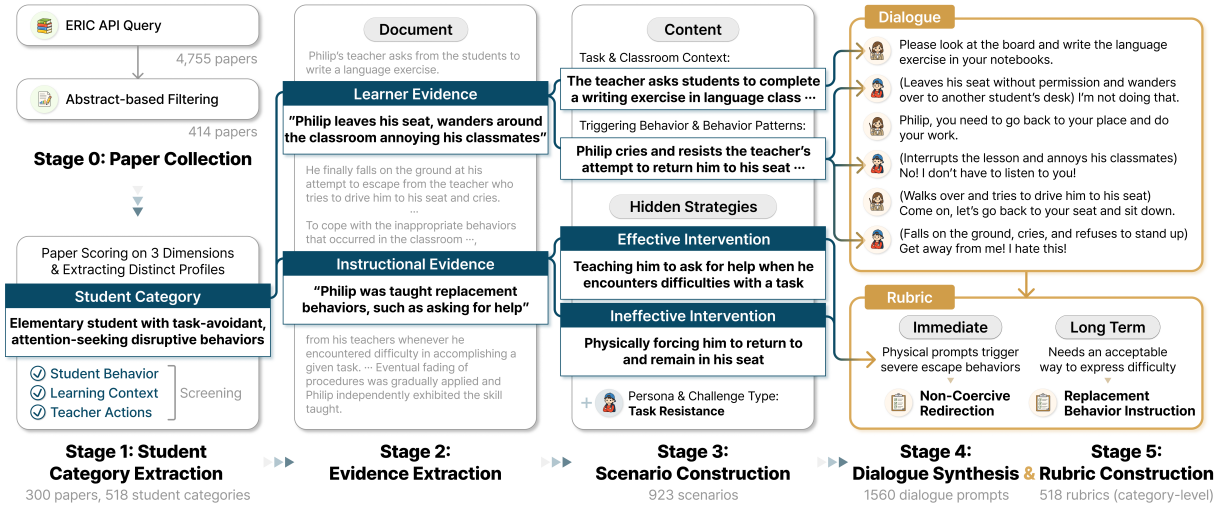

Figure 2: Overview of the TutorCaseBench construction pipeline. Based on extracted evidence, we compose case-grounded tutor–student dialogues. Examples are adapted from Theodoridou and Koutsoklenis (2013).

to derive the evaluation rubric. Each element is linked back to Stage 2 evidence spans. We retain scenarios only when all content fields are specified, and at least one effective strategy is documented.

At the category level, the model produces two additional artifacts. The *student persona* captures the learner's recurring patterns, triggers, and what tends to work or fail, ensuring consistent student turns across all scenarios in the category. Each category is also assigned a *primary challenge type* (task resistance, emotional distress, attention/engagement failure, social-interpersonal difficulty, or academic struggle) using a dialogue-based taxonomy of challenge types (see Appendix D for details), allowing us to examine model responses across qualitatively different learner difficulties.

Stage 4: Dialogue Synthesis. In the final stage, the LLM converts each scenario into a 6-turn tutor–student *dialogue* (i.e., three tutor and three student turns) that ends with the student exhibiting the challenging behavior (prompt in Appendix T.7). Given the student category, persona, and scenario content, the model generates a coherent dialogue in which every student turn is grounded in one or more excerpts. The hidden strategies block is excluded and tutor turns are constrained to be neutral.

The model also generates a one-sentence *minimal context* containing only the task and setting, leaving the student's needs to be inferred from the interaction. The full scenario content is preserved as *full context* for downstream grading.

We use six turns to provide sufficient interaction context around the challenging behavior, as tutor responses may depend on how the interaction has developed across multiple exchanges. Limiting dialogue length also keeps synthesized interactions grounded in case evidence, avoiding unsupported details and repetition. To assess this choice, we additionally test 4- and 8-turn variants and observe consistent results (see Appendix E for details).

3.2.2 Open-Ended Response Evaluation

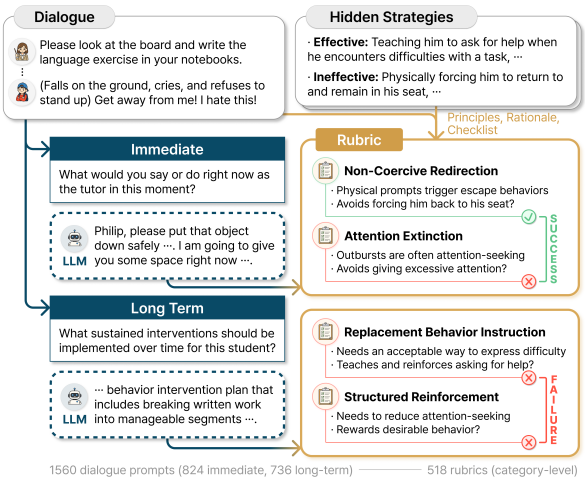

Figure 3: Rubric construction and evaluation process.

Stage 5: Rubric Construction. To evaluate tutor responses, we create case-specific rubrics by student category and response type (immediate or long-term) (prompt in Appendix T.8). Rather than scoring responses on fixed dimensions, each rubric is built from effective intervention strategies documented for that student. We define rubrics at the category level because scenarios within a category

share a learner profile and intervention evidence.

Given the student persona, challenge type, synthesized dialogues, and hidden strategies from the category's scenarios, the model produces a checklist of criteria, each capturing an evidence-supported response principle with a rationale and check condition. Ineffective strategies are included as negative anchors, and a brief student description is produced as auxiliary context. Figure 3 shows the rubric construction and evaluation process.

Each criterion is judged independently against its check condition and scored as satisfied or not (0–1). The criteria represent alternative evidence-supported approaches to the same student difficulty, rather than requirements to be combined in a single response. Requiring multiple criteria could favor broad and abstract responses over those that clearly apply a single approach. Thus, we count a response as successful if it satisfies at least one criterion.

3.3 MCQ Selection Task

To complement open-ended evaluation, we convert scenarios from §3.2.1 into 4-option MCQs. Each question pairs one effective strategy from the target scenario with three from other scenarios as distractors, testing whether models identify the strategy that best fits the student and dialogue. MCQ accuracy thus reflects the ability to select interventions supported by case-specific evidence, without implying that alternative strategies are ineffective.

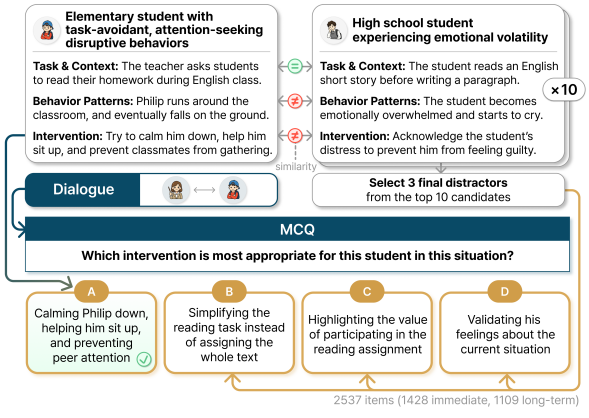

Figure 4: 4-option MCQ construction with cross-scenario distractors. Examples from Sander et al. (2010), Lichtinger and Kaplan (2015), Newman (2012).

Figure 4 illustrates the MCQ construction process. We retrieve candidate distractors by embedding similarity, select three with an LLM judge, and rewrite all options into a uniform surface form. We also construct a binary variant pairing one effective and one ineffective strategy from the same scenario as a sanity check. Full construction details for both variants are provided in Appendix F.

3.4 Data Statistics

Overall. Starting from 4,755 ERIC candidates, the pipeline produces 923 scenarios from 300 source papers and 518 student categories (M = 1.78 scenarios per category, SD = 0.77). The final benchmark includes 1,560 open-ended items: 824 immediate and 736 long-term. Category-level rubrics contain 2.90 immediate criteria (SD = 0.91) and 2.45 long-term criteria (SD = 0.87) on average. The MCQ set includes 2,537 items: 1,428 immediate and 1,109 long-term. Additional details on student demographics and corresponding model performance are provided in Appendix G.

Challenge-Type Distribution. Our taxonomy of student challenges (Appendix D) covers the dataset scenarios as follows: Task Resistance (28.7%), Attention/Engagement Failure (20.4%), Academic Struggle (18.4%), Emotional Distress (13.7%), and Social-Interpersonal Difficulty (18.9%).

3.5 Data Validation

We conducted a human validation study to assess whether outputs from our pipeline were grounded in the source papers. We randomly sampled 30 papers and recruited 12 students (three undergraduates, nine graduates) from South Korea through online university forums, all with prior experience in reading and analyzing academic papers. Each reviewer evaluated five papers, with every paper independently assessed by two reviewers so that the validation did not depend on a single evaluator.

Reviewers first read and summarized the source paper, then evaluated key outputs from each pipeline stage, providing a rationale for every judgment (see Appendix H for the study interface). On a three-point scale (No=1, Partly=2, Yes=3), mean ratings were 2.62 (SD = 0.56) for student categories (Stage 1), 2.42 (SD = 0.56) for evidence (Stage 2), 2.57 (SD = 0.57) for scenarios and 2.48 (SD = 0.54) for instructional strategies (Stage 3), 2.60 (SD = 0.60) for dialogues (Stage 4), and 2.69 (SD = 0.46) for rubrics (Stage 5). Challenge types were evaluated on a binary scale (No=0, Yes=1), yielding a mean rating of 0.76 (SD = 0.43). Further details on human validation results are provided in Appendix I. Overall, these results suggest that our pipeline outputs align with human judgments of what is supported by the source literature.

4 Experiments

We evaluate 11 state-of-the-art models on TutorCaseBench across three tasks: open-ended response generation, MCQ selection, and reference-based response generation. Experimental settings and model details are provided in Appendix J.

Open-Ended Response. Each model produces a free-form response to each open-ended item, either an immediate next-turn response or a long-term intervention plan (prompts in Appendix T.11). Immediate items ask what the tutor should say or do next, while long-term items ask what sustained interventions should be implemented over several weeks. In both cases, the model is shown only the scenario's minimal context and the dialogue up to the challenging student utterance, requiring it to infer the student's needs and respond appropriately.

We evaluate each response with `gemini-3.1-pro-preview` using the rubric from §3.2.2 (prompt in Appendix T.12). The evaluator receives the rubric, student description, dialogue, the scenario's full context, and the model response, and assigns each criterion a binary score (1 if satisfied, 0 otherwise). A response is marked successful if it satisfies at least one criterion. We also report the average criterion satisfaction rate in Appendix K. Since Gemini models are used for both construction and evaluation, we additionally test for model-family bias by varying the construction models (Appendix L) and evaluation judge (Appendix M), finding consistent results for both.

MCQ Selection. For each MCQ from §3.3, we provide the same minimal context and dialogue used in the open-ended setting, along with four answer options, and ask the model to select one (prompt in Appendix T.13). A response counts as correct if it matches the correct option, which is grounded in the scenario's effective strategy.

Reference-Based Response. We also create a reference-based variant to test whether models can produce better open-ended responses when relevant guidance is available. For each test dialogue, we retrieve the most similar scenarios from a separate pool and add their intervention strategies to the prompt as a reference (Figure 5). We evaluate the resulting responses with the same rubrics as the dialogue-only setting, enabling a paired comparison. Full construction details are in Appendix N.

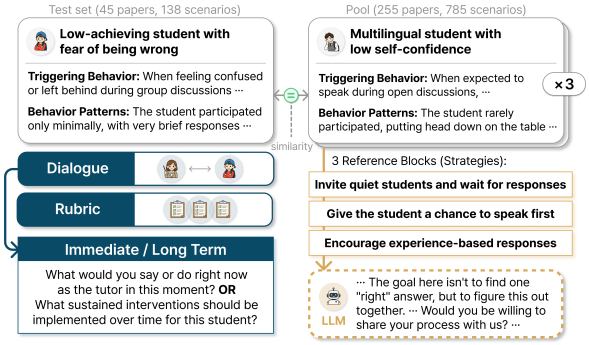

Figure 5: We use strategies from similar cases as references. Examples from Lack et al. (2014), Collett (2023), Alghamdi and Siddiqui (2016), Cassidy et al. (2018).

5 Evaluation Results

5.1 Open-Ended Response

Overall Performance. Table 1 reports the success rate of open-ended responses (i.e., the proportion satisfying at least one criterion in the rubric). Overall performance remains limited even for the strongest models. Proprietary models occupy the top ranks, with Claude Opus 4.7 performing best at 55.8% and GPT-5.5 following at 54.7%. Among open-weight models, DeepSeek V4 Pro performs best at 47.0%, while Llama 4 Maverick is lowest at 31.7%. These results indicate that generating pedagogically grounded responses for students with challenging behaviors remains difficult.

Table 1: Success rate of open-ended response (%)

Model	Immediate (N=824)	Long-Term (N=736)	Overall (N=1,560)
Claude Opus 4.7	67.0	43.3	55.8
Claude Haiku 4.5	59.0	40.1	50.1
Gemini 3.1 Pro	67.5	35.2	52.2
Gemini 3 Flash	58.7	32.2	46.2
GPT-5.5	67.2	40.6	54.7
GPT-5.4 Mini	49.2	33.3	41.7
Llama 4 Maverick	43.0	19.0	31.7
Gemma 4 31B	52.4	29.1	41.4
DeepSeek V4 Pro	55.6	37.4	47.0
GPT-OSS 120B	52.9	34.6	44.3
Qwen 3.5-397B	51.7	27.4	40.3

A consistent gap appears between immediate and long-term responses. All models perform better on immediate responses, with the largest gap for Gemini 3.1 Pro (67.5% vs. 35.2%) and the smallest for GPT-5.4 Mini (49.2% vs. 33.3%). This suggests that models are better at producing plausible in-the-moment responses, but struggle to formulate sustained interventions tailored to student needs.

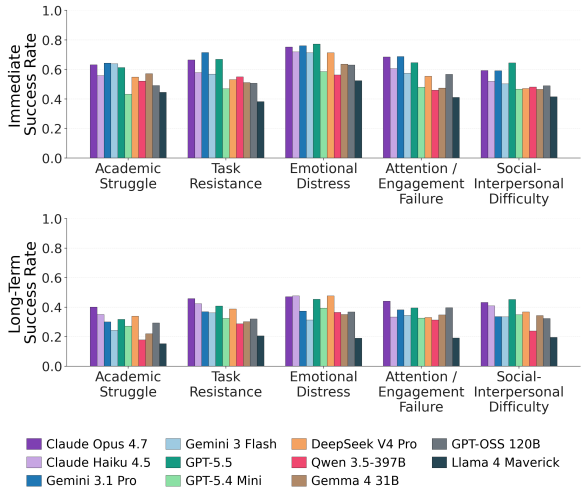

Figure 6: Success rate for open-ended response across models broken down by challenge type.

Figure 6 breaks down open-ended performance by challenge type. For immediate responses, social-interpersonal difficulty has the lowest performance and ranks last for 6 of 11 models. For long-term interventions, academic struggle has the lowest performance and ranks last for all 11 models.

Diversity and Case-Specificity. To better understand model-level differences, we additionally test whether stronger models produce less templated and more case-specific responses. We measure lexical diversity using the distinct-n ratio—the fraction of unique n-grams in each model's responses—and compute Spearman correlations with success rates.

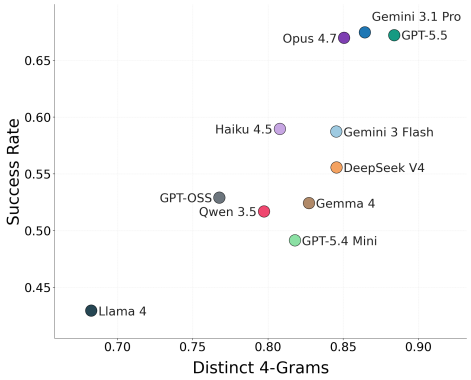

Figure 7: Correlation between distinct 4-grams ratio and immediate response success rate across models.

For immediate responses, 4-gram diversity shows a positive correlation with success rates ($\rho = 0.800$, $p < 0.01$), as shown in Figure 7. This suggests that higher-performing models tend to vary their responses more across cases rather than rely on repeated phrasing. For example, Llama 4 Maverick's most frequent 4-gram ("take a deep breath") appears in 50.4% of its immediate responses, whereas the most frequent 4-grams of Claude Opus 4.7, GPT-5.5, and Gemini 3.1 Pro appear in only 6–9%. The relationship is not significant for long-term interventions ($p > 0.05$). Results for 3- and 5-grams are in Appendix O.

We complement lexical diversity with case specificity. For each response, we use `text-embedding-3-large` (OpenAI, 2024) to retrieve the nearest response from another model and check whether it comes from the same dialogue case. We define the resulting case-recovery rate as a measure of case specificity and compute its Spearman correlation with success rates. Case-recovery rate is positively correlated with immediate-response success ($\rho = 0.691$, $p < 0.05$; see Appendix P), suggesting that stronger models produce more case-specific responses. The correlation is not significant for long-term interventions.

Failure Analysis. We qualitatively analyze a sample of 50 open-ended cases where all models failed (25 immediate, 25 long-term). We identify four recurring failure patterns, shown in Table 2.

Empathy/Praise. A student refuses to present because she believes her work is wrong and calls herself "stupid" (MacMath et al., 2010). The case-prescribed response is to verify that her work is accurate before presenting. Models instead reassure her that mistakes are acceptable, leaving her belief that the work is wrong unaddressed.

Generic Scaffolding. Another student with language-comprehension difficulties interprets "5 times more" additively and writes 10 (Bruno et al., 2024). The case-prescribed response is to use simpler terms such as "double" or "triple." Models instead correct the answer, explain multiplication, or ask why he added the numbers, addressing the error but not the underlying language difficulty.

Over-Intervention. A normally quiet student becomes nervous while co-teaching a science lesson (Adjapong and Emdin, 2015). The case-prescribed response is for the teacher to take a student role and intervene only for content issues. Models instead keep guiding him to breathe, speak louder, or begin with a suggested sentence, reducing his opportunity to lead independently.

Task Redesign Missed. A student refuses an assigned reading task and becomes disruptive during independent work (Guardino and Fullerton, 2010). The case-prescribed response is to replace the task

Table 2: Recurring failure modes in open-ended response evaluation missed by all models.

Failure Mode	Description
Empathy/Praise	Comforts the student when the response should challenge their belief or behavior.
Generic Scaffolding	Applies a standard support plan without identifying the student's central difficulty.
Over-Intervention	Steps in to correct or redirect when the tutor should step back instead.
Task Redesign Missed	Helps the student cope with a task that itself should be changed.

with a simple activity such as a crossword puzzle, coloring, or journaling. All models instead focus on calming the student and returning them to the original task, leaving the rejected task unchanged.

5.2 MCQ Selection

Table 3 reports accuracy on the MCQ task. Scores are higher and less variable than in the open-ended task, suggesting that models can recognize an appropriate intervention when provided as an option, even when they struggle to formulate one from scratch. Claude Opus 4.7 performs best at 63.0%, followed by Gemini 3.1 Pro at 62.6% and GPT-5.5 at 62.4%. The top open-weight model, Gemma 4 31B, reaches 57.3%, while GPT-OSS 120B is the lowest at 49.2%, still above the 25% chance level.

Table 3: Accuracy on 4-option MCQ cases (%).

Model	Immediate (N=1,428)	Long-Term (N=1,109)	Overall (N=2,537)
Claude Opus 4.7	63.2	62.7	63.0
Claude Haiku 4.5	53.9	55.9	54.8
Gemini 3.1 Pro	61.6	63.8	62.6
Gemini 3 Flash	55.5	56.9	56.1
GPT-5.5	62.7	62.0	62.4
GPT-5.4 Mini	51.9	54.6	53.1
Llama 4 Maverick	51.9	51.4	51.7
Gemma 4 31B	55.3	59.8	57.3
DeepSeek V4 Pro	52.4	54.7	53.4
GPT-OSS 120B	48.3	50.2	49.2
Qwen 3.5-397B	55.9	56.0	55.9

Challenge-type patterns partly mirror the open-ended setting. For immediate responses, social-interpersonal difficulty is the weakest MCQ category for all 11 models, reinforcing its difficulty in open-ended generation. Academic struggle shows the opposite pattern, performing strongly in MCQ despite being the weakest category for long-term responses (see Appendix Q). We also observe systematic convergence in MCQ errors (see Appendix R).

5.3 Reference-Based Response

We compare responses generated from the dialogue alone with those generated with strategies from similar prior cases as references. Figure 8 shows the effect on immediate responses. Most models benefit, with 10 of 11 improving and an average gain of 9.9 points overall. Llama 4 Maverick improves the most, from 41.0% to 64.8%, while Gemini 3.1 Pro does not improve (71.3% in both conditions).

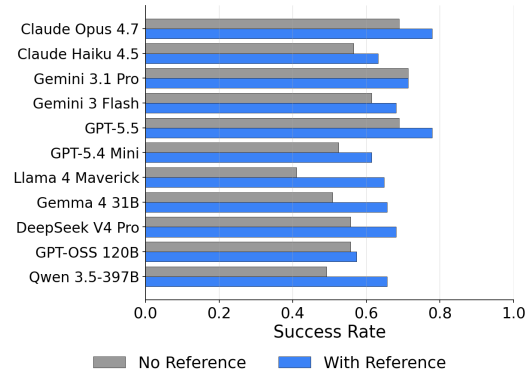

Figure 8: Success rate across models for immediate response, with and without references.

The reference effect does not extend to long-term plans (see Appendix S for details). This suggests that providing references is more effective when they directly guide the next tutor move.

6 Conclusion

We introduce TutorCaseBench, a case-grounded benchmark for evaluating whether LLM tutors can support students with challenging behaviors. TutorCaseBench converts qualitative case studies into tutor–student dialogues and case-specific evaluation criteria, avoiding idealized learner simulations and fixed pedagogical rubrics. Across 923 dialogues and 11 state-of-the-art LLMs, we find that current models remain limited in generating support aligned with documented interventions, even when they often recognize appropriate strategies in multiple-choice settings. These results highlight the need for AI tutors that move beyond plausible responses toward context-sensitive support grounded in individual student needs. Overall, TutorCaseBench provides a resource for developing tutors that address not only cognitive errors but also students' affective and behavioral difficulties.

Limitations

Our work has several limitations. First, although TutorCaseBench is grounded in qualitative case-study literature, its construction primarily relies on the Gemini model family, which may introduce model-specific interpretations into the generated scenarios and rubrics. While human validation (§3.5) and robustness analyses with alternative construction models (Appendix L) and evaluation judges (Appendix M) support the consistency of our findings, future work could explore combining different model families across pipeline stages.

Second, the benchmark has limited coverage of some student populations and educational contexts. Although we draw on accessible, contemporary ERIC case studies with detailed evidence about student behavior, learning context, and teacher actions, the resulting categories are concentrated in particular regions and age groups (Appendix G). Future work could broaden this coverage by incorporating cases from additional databases, disciplines, countries, age groups, or classroom settings.

Third, TutorCaseBench evaluates LLM tutors' responses to challenging behaviors in text-based dialogues, enabling controlled evaluation independently of behavior detection. Although the dialogues describe actions and nonverbal cues, text may not fully capture their complexity in real interactions. Future work could extend this approach to interactive, multimodal settings where tutors must recognize and respond to challenging behaviors.

Lastly, we construct student categories separately within each paper rather than grouping similar cases across the corpus. This preserves case-specific intervention evidence but limits analysis of recurring patterns in student challenges and support strategies. Future work could cluster related categories to examine these patterns across cases.

Ethical Considerations

The human validation study was reviewed and approved by the Institutional Review Board (IRB) of our institution prior to data collection. All participants provided informed consent before participation and were informed that they could withdraw at any time. As the study was conducted remotely, participants could contact the authors with any questions or concerns throughout the process. Validating each paper took approximately 90 minutes, and participants received 120,000 KRW (≈ 89.68 USD) in compensation for their time.

Source documents were identified through ERIC (Education Resources Information Center), a U.S. Department of Education database for education research literature. We used only publicly accessible full-text materials, including documents hosted by ERIC and papers reachable through DOI or publisher links listed in ERIC records. Copyright remains with the original authors and publishers, and the materials were used only for non-commercial academic research. Our released benchmark consists only of LLM-synthesized tutor–student dialogues, rubrics, and multiple-choice items. We do not redistribute the full source texts or release intermediate extraction artifacts. Source excerpts used as examples in the paper are explicitly cited. For provenance, we provide source-paper metadata for the documents used in dataset construction.

Since TutorCaseBench focuses on students with learning and behavioral difficulties, some cases may include descriptions of distress, refusal, aggression, or other challenging behaviors. These behaviors are included only when grounded in the source literature and are used to evaluate supportive tutoring responses, not to stereotype or label students. We frame such behaviors as context-dependent student needs rather than as fixed traits.

We also acknowledge the use of AI assistants in preparing this work. We used ChatGPT[2] and Claude[3] for writing assistance, as well as for generating icons used in figures. We also used Codex[4] and Claude Code[5] for coding assistance in implementing the benchmark construction pipeline.

[2] https://chatgpt.com/
[3] https://claude.ai
[4] https://openai.com/codex/
[5] https://code.claude.com/docs

References

Thomas M Achenbach. 1991. *Manual for the child behavior checklist/4-18 and 1991 profile*. University of Vermont, Department of Psychiatry.

Edmund S Adjapong and Christopher Emdin. 2015. Rethinking pedagogy in urban spaces: Implementing hip-hop pedagogy in the urban science classroom. *Journal of Urban Learning, Teaching, and Research*, 11:66–77.

Fatimah M A Alghamdi and Ozma Siddiqui. 2016. Supporting low-achieving efl learners: Expectations, procedure and significance of remedial sessions at a saudi university. *Journal of Education and Training Studies*, 4(12):204–212.

Anthropic. 2025. System Card: Claude Haiku 4.5. Accessed: 2026-10-12.

Anthropic. 2026a. System Card: Claude Opus 4.7. Accessed: 2026-10-12.

Anthropic. 2026b. System Card: Claude Sonnet 5.5. Accessed: 2026-10-12.

Ryan S.J.d. Baker, Sidney K. D'Mello, Ma.Mercedes T. Rodrigo, and Arthur C. Graesser. 2010. Better to be frustrated than bored: The incidence, persistence, and impact of learners' cognitive–affective states during interactions with three different computer-based learning environments. *International Journal of Human-Computer Studies*, 68(4):223–241.

Alicia Bruno, Irene Polo-Blanco, Steven Van Vaerenbergh, Raúl Fernández-Cobos, and María José González-López. 2024. Strategies for solving multiplicative problems using a conceptual model-based problem-solving approach. a case study with a student with autism spectrum disorder. *ZDM Mathematics Education*, 56:1239–1256.

Claire Cassidy, Helen Marwick, Lynn Deeney, and Gillian McLean. 2018. Philosophy with children, self-regulation and engaged participation for children with emotional-behavioural and social communication needs. *Emotional and Behavioural Difficulties*, 23(1):81–96.

Jennifer Collett. 2023. Relationships between reading identities and dialogue: Perspectives from fifth-grade multilingual learners. *Reading & Writing Quarterly*, 39(4):259–276.

DeepSeek-AI, Anyi Xu, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, Chenchen Ling, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chengyu Hou, Chenhao Xu, Chenze Shao, Chong Ruan, Conner Sun, and 300 others. 2026. Deepseek-v4: Towards highly efficient million-token context intelligence. *Preprint*, arXiv:2606.19348.

Sidney D'Mello and Art Graesser. 2013. Autotutor and affective autotutor: Learning by talking with cognitively and emotionally intelligent computers that talk back. *ACM Transactions on Interactive Intelligent Systems*, 2(4).

Song Feng, Siva Sankalp Patel, Hui Wan, and Sachindra Joshi. 2021. MultiDoc2Dial: Modeling dialogues grounded in multiple documents. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6162–6176, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.

Song Feng, Hui Wan, Chulaka Gunasekara, Siva Patel, Sachindra Joshi, and Luis Lastras. 2020. doc2dial: A goal-oriented document-grounded dialogue dataset. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8118–8128, Online. Association for Computational Linguistics.

Google DeepMind. 2025. Gemini 3 Flash Model Card. Accessed: 2026-10-12.

Google DeepMind. 2026a. Gemini 3.1 Pro Model Card. Accessed: 2026-10-12.

Google DeepMind. 2026b. Gemma 4 model card. Accessed: 2026-10-12.

Caroline A Guardino and Elizabeth Fullerton. 2010. Changing behaviors by changing the classroom environment. *TEACHING Exceptional Children*, 42(6):8–13.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring massive multitask language understanding. In *International Conference on Learning Representations*.

Hyoungwook Jin, Minju Yoo, Jeongeon Park, Yokyung Lee, Xu Wang, and Juho Kim. 2025. Teachtune: Reviewing pedagogical agents against diverse student profiles with simulated students. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, CHI '25, New York, NY, USA. Association for Computing Machinery.

Irina Jurenka, Markus Kunesch, Kevin R. McKee, Daniel Gillick, Shaojian Zhu, Sara Wiltberger, Shubham Milind Phal, Katherine Hermann, Daniel Kasenberg, Avishkar Bhoopchand, Ankit Anand, Miruna Pîslar, Stephanie Chan, Lisa Wang, Jennifer She, Parsa Mahmoudieh, Aliya Rysbek, Wei-Jen Ko, Andrea Huber, and 55 others. 2025. Towards responsible development of generative ai for education: An evaluation-driven approach. *Preprint*, arXiv:2407.12687.

Shamya Karumbaiah, Conrad Borchers, Tianze Shou, Ann-Christin Falhs, Pinyang Liu, Tomohiro Nagashima, Nikol Rummel, and Vincent Aleven. 2023. A spatiotemporal analysis of teacher practices in supporting student learning and engagement in an ai-enabled classroom. In *Artificial Intelligence in Education*, pages 450–462, Cham, Switzerland. Springer Nature Switzerland.

Brian Lack, Susan Lee Swars, and Barbara Meyers. 2014. Low- and high-achieving sixth-grade students' access to participation during mathematics discourse. *The Elementary School Journal*, 115(1):97–123.

Einat Lichtinger and Avi Kaplan. 2015. Employing a case study approach to capture motivation and self-regulation of young students with learning disabilities in authentic educational contexts. *Metacognition and Learning*, 10(1):119–149.

Jiayu Liu, Zhenya Huang, Tong Xiao, Jing Sha, Jinze Wu, Qi Liu, Shijin Wang, and Enhong Chen. 2024a. Socraticlm: Exploring socratic personalized teaching with large language models. In *Advances in Neural Information Processing Systems*, volume 37, pages 85693–85721. Curran Associates Inc.

Zhengyuan Liu, Stella Xin Yin, Geyu Lin, and Nancy F. Chen. 2024b. Personality-aware student simulation for conversational intelligent tutoring systems. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 626–642, Miami, Florida, USA. Association for Computational Linguistics.

Xinyi Lu and Xu Wang. 2024. Generative students: Using llm-simulated student profiles to support question item evaluation. In *Proceedings of the Eleventh ACM Conference on Learning @ Scale*, L@S '24, pages 16–27, New York, NY, USA. Association for Computing Machinery.

Jakub Macina, Nico Daheim, Sankalan Chowdhury, Tanmay Sinha, Manu Kapur, Iryna Gurevych, and Mrinmaya Sachan. 2023. MathDial: A dialogue tutoring dataset with rich pedagogical properties grounded in math reasoning problems. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 5602–5621, Singapore. Association for Computational Linguistics.

Jakub Macina, Nico Daheim, Ido Hakimi, Manu Kapur, Iryna Gurevych, and Mrinmaya Sachan. 2025. MathTutorBench: A benchmark for measuring open-ended pedagogical capabilities of LLM tutors. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 204–221, Suzhou, China. Association for Computational Linguistics.

Sheryl MacMath, Jillian Roberts, John Wallace, and Xiaohong Chi. 2010. Research section: Curriculum integration and at-risk students: a canadian case study examining student learning and motivation. *British Journal of Special Education*, 37(2):87–94.

Julia M. Markel, Steven G. Opferman, James A. Landay, and Chris Piech. 2023. Gpteach: Interactive ta training with gpt-based students. In *Proceedings of the Tenth ACM Conference on Learning @ Scale*, L@S '23, pages 226–236, New York, NY, USA. Association for Computing Machinery.

Luis Marquez-Carpintero, Alberto Lopez-Sellers, and Miguel Cazorla. 2026. Simulating students with large language models: A review of architecture, mechanisms, and role modeling in education with generative ai. *Computer Science Review*, 62:101008.

Daria Martynova, Jakub Macina, Nico Daheim, Nilay Yalcin, Xiaoyu Zhang, and Mrinmaya Sachan. 2025. Can LLMs effectively simulate human learners? teachers' insights from tutoring LLM students. In *Proceedings of the 20th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2025)*, pages 100–117, Vienna, Austria. Association for Computational Linguistics.

Kaushal Kumar Maurya, Kv Aditya Srivatsa, Kseniia Petukhova, and Ekaterina Kochmar. 2025. Unifying AI tutor evaluation: An evaluation taxonomy for pedagogical ability assessment of LLM-powered AI tutors. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 1234–1251, Albuquerque, New Mexico. Association for Computational Linguistics.

Meta. 2025. The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation. Accessed: 2026-10-12.

Cara Newman. 2012. Home/hospital instruction: Instructional approach to working with students with major depressive disorder. Master's thesis, Dominican University of California.

OpenAI. 2024. New embedding models and API updates. Accessed: 2026-10-12.

OpenAI. 2025a. gpt-oss-120b & gpt-oss-20b Model Card. Accessed: 2026-10-12.

OpenAI. 2025b. Introducing GPT-4.1 in the API. Accessed: 2026-10-12.

OpenAI. 2026a. GPT-5.5 System Card. Accessed: 2026-10-12.

OpenAI. 2026b. Introducing GPT-5.4 mini and nano. Accessed: 2026-10-12.

Ernesto Panadero. 2017. A review of self-regulated learning: Six models and four directions for research. *Frontiers in Psychology*, 8.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.

Kristy Park and Sarah Pinkelman. 2017. Function-based approach to designing an instructional environment. *Beyond Behavior*, 26(3):124–134.

Susmita Patnaik and Pearl Subban. 2023. Challenging behaviours in primary classrooms: Examining causal attributions and strategies used by primary school teachers in india. *Australasian Journal of Special and Inclusive Education*, 47(2):78–95.

Alessandro Pepe and Loredana Addimando. 2013. Comparison of occupational stress in response to challenging behaviours between general and special education primary teachers in northern italy. *International Journal of Special Education*, 28(1):14–26.

Paul R. Pintrich. 2000. Chapter 14 - the role of goal orientation in self-regulated learning. In Monique Boekaerts, Paul R. Pintrich, and Moshe Zeidner, editors, *Handbook of Self-Regulation*, pages 451–502. Academic Press, San Diego.

Changyong Qi, Longwei Zheng, Anna He, Haoxin Xu, Linzhao Jia, Yuang Wei, Bingqian Jiang, and Xiaoqing Gu. 2026. Simulating student learning behaviors with llm-based role-playing agents: A data-driven and cognitively inspired framework. *Expert Systems with Applications*, 304:130753.

Qwen Team. 2026. Qwen3.5: Towards native multimodal agents.

Janay B. Sander, Jill D. Sharkey, Roger Olivarri, Diane A. Tanigawa, and Tory Mauseth. 2010. A qualitative study of juvenile offenders, student engagement, and interpersonal relationships: Implications for research directions and preventionist approaches. *Journal of Educational and Psychological Consultation*, 20(4):288–315.

Alexander Scarlatos, Jaewook Lee, Simon Woodhead, and Andrew Lan. 2026. Simulated students in tutoring dialogues: Substance or illusion? In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 42349–42385, San Diego, California, United States. Association for Computational Linguistics.

Kumar Shridhar, Jakub Macina, Mennatallah El-Assady, Tanmay Sinha, Manu Kapur, and Mrinmaya Sachan. 2022. Automatic generation of socratic subquestions for teaching math word problems. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 4136–4149, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

Shashank Sonkar, Naiming Liu, Debshila Mallick, and Richard Baraniuk. 2023. CLASS: A design framework for building intelligent tutoring systems based on learning science principles. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 1941–1961, Singapore. Association for Computational Linguistics.

George Sugai and Robert R. Horner. 2006. A promising approach for expanding and sustaining school-wide positive behavior support. *School Psychology Review*, 35(2):245–259.

Anaïs Tack and Chris Piech. 2022. The AI teacher test: Measuring the pedagogical ability of blender and GPT-3 in educational dialogues. In *Proceedings of the 15th International Conference on Educational Data Mining*, pages 522–529, Durham, United Kingdom. International Educational Data Mining Society.

Zoe Theodoridou and Athanasios Koutsoklenis. 2013. Functional behavioral assessment for a boy with duchenne muscular dystrophy and problem behavior: A case study from greece. *Assessment for Effective Intervention*, 39(1):54–64.

Junling Wang, Jakub Macina, Nico Daheim, Sankalan Pal Chowdhury, and Mrinmaya Sachan. 2024a. Book2Dial: Generating teacher student interactions from textbooks for cost-effective development of educational chatbots. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 9707–9731, Bangkok, Thailand. Association for Computational Linguistics.

Rose Wang, Qingyang Zhang, Carly Robinson, Susanna Loeb, and Dorottya Demszky. 2024b. Bridging the novice-expert gap via models of decision-making: A case study on remediating math mistakes. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 2174–2199, Mexico City, Mexico. Association for Computational Linguistics.

Xintao Wang, Heng Wang, Yifei Zhang, Xinfeng Yuan, Rui Xu, Jen tse Huang, Siyu Yuan, Haoran Guo, Jiangjie Chen, Shuchang Zhou, Wei Wang, and Yanghua Xiao. 2025. CoSER: Coordinating LLM-based persona simulation of established roles. In *Forty-second International Conference on Machine Learning*.

Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. Bertscore: Evaluating text generation with bert. In *International Conference on Learning Representations*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36, pages 46595–46623. Curran Associates Inc.

A Model Fallbacks in the Pipeline

Throughout the benchmark construction pipeline, we prioritized the Gemini model specified in §3 for consistency. Failed calls were retried up to five times before falling back to an OpenAI model. Here, we document these fallback cases.

In Stage 2 (Evidence Extraction), a small subset of papers ($N = 7$, 2.3% of 302 papers before subsequent filtering to 300) was reprocessed with `gpt-5.4-mini` due to output truncation ($N = 5$), recitation blocking ($N = 1$), or a schema error ($N = 1$). In Stage 5 (Rubric Construction), a single student category ($N = 1$ of 518, 0.2%) was processed with `gpt-5.5` after the original Gemini call was blocked by its safety filter. Follow-up tests suggest that the block may have been related to the combination of the student's young age, disability, and descriptions of severe distress behaviors.

Given the small proportion of affected cases, we expect this fallback model use to have only a limited impact on the overall benchmark data.

B ERIC Search Descriptors

Table 4 lists the ERIC descriptors used to retrieve candidate papers for dataset construction.

Table 4: Descriptors used to retrieve candidate papers.

Category	Descriptors
Student (35 Total)	Behavior Problems Behavior Disorders Self Destructive Behavior Aggression Emotional Problems Emotional Disturbances Anxiety Disorders Depression (Psychology) School Phobia Withdrawal (Psychology) Fear of Success Mental Disorders Personality Problems Psychosomatic Disorders Attention Deficit Disorders Attention Deficit Hyperactivity Disorder Autism Spectrum Disorders Developmental Disabilities Intellectual Disability Neurodevelopmental Disorders Neurological Impairments Developmental Delays Learning Problems Learning Disabilities Low Achievement Academic Failure Reading Difficulties Reading Failure Writing Difficulties Mathematics Anxiety Slow Learners At Risk Students Communication Problems Communication Disorders Language Impairments
Methodology (4 Total)	Case Studies Observation Qualitative Research Classroom Research

C Abstract-Based Paper Filtering Procedure

The candidate papers are initially screened by `gpt-4.1-mini` (OpenAI, 2025b) using the title and abstract. We use this as a cost-efficient, high-recall first pass, since abstracts provide enough signal to remove clearly unfit papers while leaving ambiguous cases to later pipeline stages.

The filter applies four criteria, filtering out papers that fail to meet our benchmark requirements: (1) papers without direct student observation or interview, which lack grounded learner behavior; (2) quantitative studies, which report aggregate effects instead of the individual student responses we need; (3) papers focused on educational policy or measurement-tool development, which describe students only as targets of analysis; and (4) papers without practical guidance for teachers, which lack the prescriptive strategies needed for evaluation. A paper is retained only if rejected by none of the four. The full prompt is in Appendix T.1.

We additionally exclude excessively long papers (≥ 30 pages) and those that cannot be accessed or downloaded, as we observed that long papers tend to mistakenly introduce literature-review content as learner evidence in downstream extraction.

D Dialogue-Based Taxonomy of Learner Challenges

To distinguish the challenge type associated with each student category, we designed a taxonomy that identifies the observable challenge a tutor must respond to in the tutoring interaction. The taxonomy consists of five categories: *Task Resistance*, *Emotional Distress*, *Attention/Engagement Failure*, *Social-Interpersonal Difficulty*, and *Academic Struggle*. It was designed and refined by an author with a doctoral degree in education.

We use CBCL/ASEBA (Achenbach, 1991) as a reference for defining our categories. CBCL/ASEBA is an empirical framework that char-

acterizes children's functioning through standardized behavior descriptions. Since our benchmark is dialogue-based, we retain only constructs observable in tutoring interactions and reorganize them by the type of challenge they present to the tutor.

Specifically, the Internalizing grouping is adapted as *Emotional Distress*; Attention Problems as *Attention/Engagement Failure*; Social Problems and Aggressive Behavior as *Social-Interpersonal Difficulty*; and Externalizing behaviors directed at the learning task as *Task Resistance*. *Academic Struggle* adapts the School scale for content-level difficulty such as confusion, uncertainty, repeated errors, misconceptions, and failure to respond.

E Impact of Dialogue Turns on Evaluation Results

We examine whether evaluation results depend on the six-turn dialogue setting used in our benchmark. For this analysis, we sampled 25% of student categories ($N = 130$), stratified by the five challenge types, and regenerated their dialogues with four and eight turns. We used an adapted version of the original dialogue-synthesis prompt that allows the number of turns to vary (prompt in Appendix T.17). We then repeated the open-ended evaluation for each dialogue setting using the same rubrics.

Changing the number of dialogue turns had little effect on the evaluation results. Overall success rates were 45.8%, 46.6%, and 46.8% for four-, six-, and eight-turn dialogues, respectively. Model rankings also remained highly consistent, with Spearman correlation being $\rho = 0.916$ ($p < 0.001$) for the four- versus six-turn comparison and $\rho = 0.911$ ($p < 0.001$) for the six- versus eight-turn comparison. These results suggest that our findings are not sensitive to the number of dialogue turns.

F MCQ Construction Details

F.1 4-Option MCQs

We provide additional details on distractor retrieval, selection, and option adaptation for the MCQ task.

We construct distractors in three steps: candidate retrieval, distractor selection, and surface adaptation. As a scenario may contain multiple effective strategies, each strategy is treated as a separate correct option and receives its own distractor set.

For retrieval, we embed each scenario's learning context, learner behavior, and intervention descriptions using `text-embedding-3-large` (OpenAI, 2024). Each candidate intervention is scored as:

$$\text{score} = \text{ctx_sim} \times (1 - \text{behav_sim}) \times (1 - \text{max_interv_sim}),$$

where `ctx_sim` and `behav_sim` are cosine similarities between the target and candidate learning contexts and learner behaviors, and `max_interv_sim` is the maximum similarity between the candidate intervention and any of the target's effective interventions. This favors candidates that fit the target's classroom context but address a different learner problem with a different intervention mechanism. We exclude strategies from the same paper or student category and keep only the highest-scoring intervention per scenario, ensuring that distractors come from distinct cases. The top 10 candidates are forwarded to the next step.

However, embedding similarity alone cannot reliably identify contextual incompatibilities or paraphrased versions of valid target interventions. Hence, we use `gemini-3.1-pro-preview` (Google DeepMind, 2026a) as an LLM judge (Zheng et al., 2023) to select three final distractors from the top 10 candidates, prioritizing plausible but incorrect tutor actions that are structurally compatible with the target setting (prompt in Appendix T.9).

Finally, the same model rewrites all four options—including the correct one—into the target dialogue's vocabulary, length, and tone while preserving each option's underlying mechanism (prompt in Appendix T.10). This prevents models from identifying the correct answer through lexical or stylistic match with the dialogue. Options are shuffled with a fixed seed and labeled A–D, and each MCQ is presented with the dialogue and a prompt asking the model to choose the most appropriate intervention (see Appendix T.13).

F.2 Binary MCQs

We additionally construct a binary MCQ variant as a sanity check. Unlike the 4-option MCQs in §3.3, where all options are effective strategies and the task is to identify the best fit, the binary variant presents one effective and one ineffective strategy from the target scenario. Since the ineffective option failed for the same student in the source case, this variant tests whether models can reject a clearly inappropriate response. High accuracy therefore helps rule out this basic failure mode and supports interpreting the 4-option task as the harder problem of choosing among plausible effective strategies.

F.2.1 Construction

For each scenario and response type (*immediate* or *long-term*), we create one item when the scenario contains at least one matching effective strategy and at least one ineffective strategy. When multiple options are available on either side, we use `gemini-3.1-pro-preview` (Google DeepMind, 2026a) as an LLM judge (Zheng et al., 2023) to select the final effective–ineffective option pair (prompt in Appendix T.14). The judge prioritizes pairs with matching scope, concrete tutor actions, and comparable granularity, where the effective option directly addresses the dialogue and the ineffective option is a plausible misstep rather than an unrelated action. If each side has only one eligible strategy, the pair is fixed, and the judge is skipped.

Since effective and ineffective strategies may differ in style in how they are written, the same model then rewrites the two options to match in vocabulary, length, and tone while preserving their meanings (prompt in Appendix T.15). This step is constrained to make minimal edits, avoid unsupported additions, and keep the ineffective option substantively ineffective. The two options are then assigned to A and B using a fixed random seed.

F.2.2 Results

Across 11 models and 751 binary MCQs, accuracy ranges from 85.8% to 93.1%, which is far above the 50% chance baseline (Figure 9). This suggests that models can reliably distinguish a documented effective intervention from an ineffective one when the contrast is explicit. Therefore, the lower performance in the 4-option MCQ setting is not due to a basic inability to reject ineffective responses, but to the harder task of selecting the best-fitting intervention among plausible alternatives.

G Student Demographic Information

G.1 Demographic Distribution

We summarize demographic information for 518 student categories from 300 source papers. During Stage 1 of our pipeline, the LLM extracted country and education level at the paper level, and age and gender at the student-category level. For analysis, we grouped countries into regions and ages into ranges. Country is reported for 99.8% of categories, education level for all categories, gender for 74.9%, and age for 60.4%. We discuss the potential implications of these distributions in Limitations.

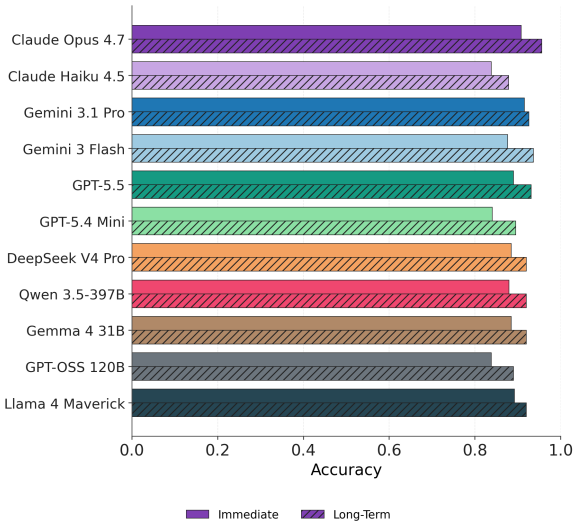

Figure 9: Binary MCQ accuracy across 11 models.

Geographic distribution. Most categories are from North America (61.4%) and Europe (20.5%), followed by Asia (5.8%), the Middle East (5.4%), Oceania (5.0%), Latin America (1.4%), and Africa (0.4%). The United States accounts for 58.9% of categories, followed by the United Kingdom (9.7%), Australia (4.4%), and Canada (2.5%).

Education level. Primary-school students account for 41.9% of categories, followed by lower secondary (20.5%), upper secondary (19.3%), early childhood (7.1%), and tertiary education (5.4%). The remaining 5.8% come from settings outside a single school level, including special-education schools and adult or alternative programs.

Age. Students aged 6–11 account for 25.9% of all categories, followed by ages 12–14 (13.5%), 15–18 (11.8%), 0–5 (6.8%), and 19 or older (2.5%).

Gender. Male students appear in 47.7% of categories, female students in 24.3%, and mixed-gender groups in 2.7%, with a single category (0.2%) describing a non-binary student.

G.2 Performance by Demographic Groups

Table 5 presents open-ended response success rates across demographic groups, averaged over 11 LLMs. Performance showed similar trends across groups, with immediate responses generally outperforming long-term interventions. Given the uneven group sizes, comparisons across demographic groups should be interpreted with caution.

Table 5: Open-ended success rates (%) by geographic region, education level, and gender, averaged across 11 LLMs. N denotes the number of student categories.

Group	Category	N	Imm.	LT	Overall
Overall	–	518	56.7	33.8	45.9
Region	North America	318	53.5	32.5	43.3
	Europe	106	64.5	33.1	50.3
	Asia	30	57.5	43.3	51.2
	Middle East	28	64.5	38.8	52.9
	Oceania	26	53.7	34.8	45.0
	Latin America	7	45.5	31.4	38.7
	Africa	2	51.5	81.8	63.6
	Not reported	1	72.7	0.0	36.4
Education	Early childhood	37	48.8	29.8	40.5
	Primary	217	57.5	32.0	45.2
	Lower secondary	106	55.9	33.5	45.8
	Upper secondary	100	57.4	37.9	47.8
	Secondary	2	50.0	86.4	68.2
	Tertiary	28	61.6	36.6	48.0
	Other/Unspecified	28	59.7	33.1	49.3
Gender	Male	247	56.8	33.0	45.2
	Female	126	54.3	33.8	45.0
	Mixed	14	54.5	26.4	42.0
	Non-binary	1	27.3	18.2	22.7
	Not reported	130	58.9	36.1	48.5

H Human Validation Interface

Figure 10 shows the interface used for human validation. For each assigned paper, participants reviewed the outputs from each pipeline stage and rated how well they were supported by the source paper. At each step, the interface provided the evaluation criterion and its definition, the corresponding pipeline output, and a validation question with response options. Participants selected a rating and provided a written rationale for their judgment.

We designed the validation interface to encourage informed and deliberate judgments. For each assigned paper, participants were first required to read the source paper and write a summary of at least 30 words before validating its pipeline outputs. The interface also required a written rationale of at least 50 characters for each rating to encourage participants to articulate the reasoning behind their judgments and discourage arbitrary ratings.

I Further Details on Human Validation

Our human validation study covered 30 of the 300 source papers and 57 of the 518 student categories. Each paper was independently evaluated by two non-author reviewers across all pipeline stages. We present the distribution of human validation ratings and exact inter-rater agreement in Table 6.

Across 882 three-point judgments, 96.6% were rated *Yes* or *Partly*, with only 3.4% receiving *No*. 76.3% of 114 binary judgments also confirmed the assigned challenge types. Exact agreement between evaluators was 58.2% overall. Agreement was highest for challenge types (84.2%) and lowest for instructional strategies (43.4%). While these results suggest that the generated outputs of our pipeline were generally well-grounded in the source papers, we further examine reviewers' rationales for both *Partly* and *No* judgments to understand the concerns underlying lower ratings.

Table 6: Human validation ratings and exact inter-rater agreement (%). Challenge types used a binary (0/1), and other outputs used a three-point scale (1–3). Agreement indicates identical ratings between evaluators.

Output	Yes	Partly	No	Agreement
Student Categories	65.0	31.7	3.3	73.3
Evidence Excerpts	45.6	50.9	3.5	47.4
Challenge Types	76.3	–	23.7	84.2
Scenarios	61.1	34.8	4.0	61.6
Instructional Strategies	50.5	47.5	2.0	43.4
Dialogues	65.7	28.3	6.1	55.6
Rubrics	69.3	30.7	0.0	59.6
Overall	59.1	37.5	3.4	58.2

Student Categories (Stage 1). Most judgments of student categories were *Yes* (65.0%), suggesting that the extracted categories captured relevant students and their characteristics. Partial ratings concerned omitted students, overlapping categories, or imprecise descriptions. One reviewer noted, "the student groupings are appropriate, but Jennifer's description includes a misattributed motivation." Reviewers also disagreed on whether students from the same study should be grouped together or separated based on their distinct characteristics. Overall, most concerns reflected differences in categorization and interpretation, with relatively few judgments identifying substantive grounding issues.

Evidence Excerpts (Stage 2). Evidence excerpts received the highest proportion of *Partly* ratings (50.9%). Reviewers raised concerns about the relevance and specificity of individual excerpts. One reviewer noted, "Rather than being incorrect, some excerpts only partially describe the category or include passages that are too broad, making their relevance less clear." Reviewers also noted excerpts drawn from other students' cases or prior literature. Since reviewers evaluated entire evidence sets, their rationales for *Partly* ratings often focused on issues with a few individual excerpts, while the remaining

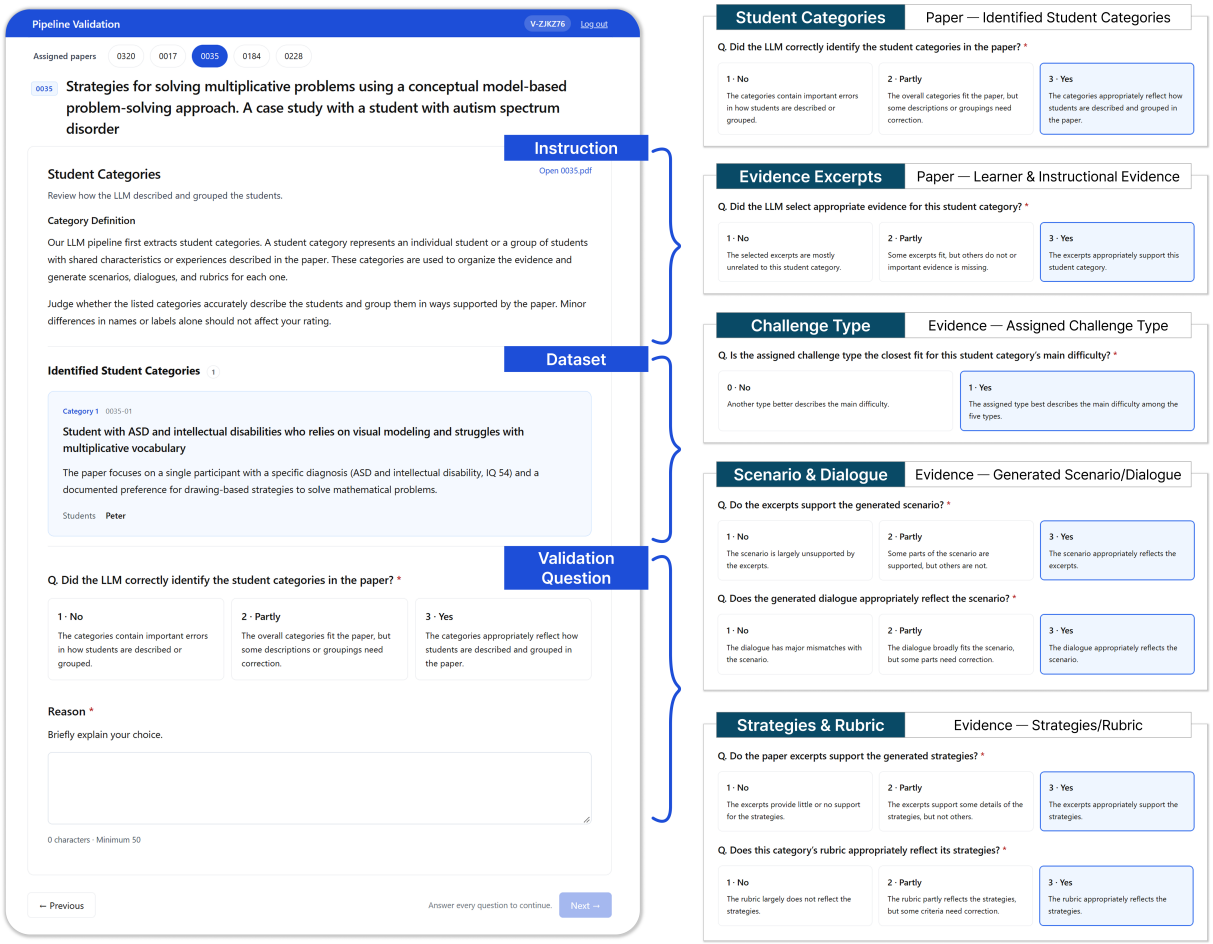

Figure 10: Human validation interface across five pipeline stages. The left panel shows an example evaluation screen, while the right summarizes validation questions. Example from Bruno et al. (2024).

evidence was considered appropriate.

Challenge Types (Stage 3). Most judgments (76.3%) confirmed the assigned challenge types, suggesting that they generally captured students' primary difficulties. However, reviewers noted that many students exhibited multiple, overlapping challenges that extended beyond the assigned type.

Scenarios (Stage 3). Most judgments (61.1%) were *Yes*, suggesting that scenarios were grounded in students' documented behaviors and contexts. Partial ratings often concerned combining evidence from different moments or contexts into a single situation. One reviewer noted, "The scenario is largely supported, but it combines several related moments from Jimmy's case into one situation." Reviewers also identified unsupported contextual details or occasional misinterpretations of student behaviors. Overall, these concerns centered on preserving the original context when synthesizing documented observations into coherent scenarios.

Instructional Strategies (Stage 3). Most judgments were *Yes* (50.5%) or *Partly* (47.5%), suggesting that the extracted strategies were supported by the source papers. Partial ratings often concerned descriptions of individual strategies' contexts and outcomes. One reviewer noted, "The actions are appropriate, but some context and outcomes seem to have been added to fit the scenario." Reviewers also identified general recommendations presented as interventions for specific students. These issues arose mainly when adapting documented strategies and their outcomes to individual scenarios.

Dialogues (Stage 4). Most judgments (65.7%) were *Yes*, suggesting that the generated dialogues reflected the intended scenarios and student behaviors. Partial ratings often concerned unnatural exchanges, even when the key elements of the scenarios were preserved. One reviewer noted, "The main points are captured, but the lines feel more like separate monologues than a natural conversation." Reviewers also identified exaggerated student behav-

iors or tutor responses that were more directive than intended. The dialogues remained well-grounded in the scenarios, with reviewers' feedback primarily addressing conversational naturalness.

Rubrics (Stage 5). Most judgments (69.3%) were *Yes*, with no *No* ratings, suggesting that rubrics well-reflected the extracted instructional strategies. Partial ratings often concerned overly specific requirements or rationales that extended beyond the source evidence. One reviewer noted, "The six criteria align well with the strategies, but some of the rationales contain unsupported claims." Reviewers identified cases where context-specific findings were turned into broader rules. Overall, the rubrics captured the core instructional guidance, with some interpretation involved in translating the strategies into scenario-specific criteria.

Overall, human validation indicated that our pipeline outputs were well-grounded in the source papers. Reviewers' rationales also highlighted opportunities to improve how source evidence is interpreted and adapted into case-specific outputs, extending beyond grounding alone. Nevertheless, the evaluations supported the integrity of the resulting dataset, with most outputs meeting human reviewers' judgments and expectations.

J Models

We evaluate a diverse set of 11 state-of-the-art models on TutorCaseBench: six proprietary (`claude-opus-4.7` (Anthropic, 2026a), `claude-haiku-4.5` (Anthropic, 2025), `gemini-3.1-pro-preview` (Google DeepMind, 2026a), `gemini-3-flash-preview` (Google DeepMind, 2025), `gpt-5.5` (OpenAI, 2026a), `gpt-5.4-mini` (OpenAI, 2026b)) and five open-weight (`DeepSeek-V4-Pro` (DeepSeek-AI et al., 2026), `gemma-4-31B-it` (Google DeepMind, 2026b), `gpt-oss-120b` (OpenAI, 2025a), `Llama-4-Maverick-17B-128E-Instruct` (Meta, 2025), `Qwen3.5-397B-A17B` (Qwen Team, 2026)).

We accessed each model through its respective deployed API. Claude models and Llama were served through Vertex AI[6], Gemini models and Gemma through the Google Gemini API[7], GPT models through the OpenAI API[8], DeepSeek and GPT-OSS through Together AI[9], and Qwen through OpenRouter[10]. In all settings, models were queried with their default API settings to reflect their standard deployed configurations.

K Average Criterion Satisfaction Rate

In addition to the binary success rate, we report the average criterion satisfaction rate, calculated as the number of satisfied criteria divided by the total number of criteria for each response.

Table 7: Average criterion satisfaction rate (%)

Model	Immediate (N=824)	Long-Term (N=736)	Overall (N=1,560)
Claude Opus 4.7	35.8	20.9	28.8
Claude Haiku 4.5	29.1	18.4	24.1
Gemini 3.1 Pro	35.7	16.0	26.4
Gemini 3 Flash	29.1	14.7	22.3
GPT-5.5	35.7	19.4	28.0
GPT-5.4 Mini	23.4	15.1	19.5
Llama 4 Maverick	18.9	8.2	13.9
Gemma 4 31B	25.2	13.5	19.7
DeepSeek V4 Pro	27.4	18.5	23.2
GPT-OSS 120B	25.9	16.3	21.4
Qwen 3.5-397B	24.5	12.0	18.6

Table 7 reports the average criterion satisfaction rates across the 11 models. Claude Opus 4.7 performs best overall (28.8%), followed by GPT-5.5 (28.0%) and Gemini 3.1 Pro (26.4%). Consistent with the success rate evaluation, the same models lead overall, and all 11 models perform better on immediate responses than on long-term plans. These results indicate that the overall performance patterns are preserved under this alternative metric.

Figure 11 further breaks down the results by challenge type. For immediate responses, task resistance and social-interpersonal difficulty generally show lower satisfaction rates, while emotional distress performs relatively well. For long-term interventions, academic struggle and social-interpersonal difficulty tend to show lower rates. Overall, social-interpersonal difficulty remains challenging across both settings, while the other lower-performing challenge types differ between immediate and long-term interventions.

L Impact of Construction Model Choice on Evaluation Results

We examine whether evaluation results depend on the choice of models used for benchmark construc-

[6] https://docs.cloud.google.com/vertex-ai/docs
[7] https://ai.google.dev/gemini-api/docs
[8] https://developers.openai.com/api/docs
[9] https://docs.together.ai/intro
[10] https://openrouter.ai/docs/quickstart

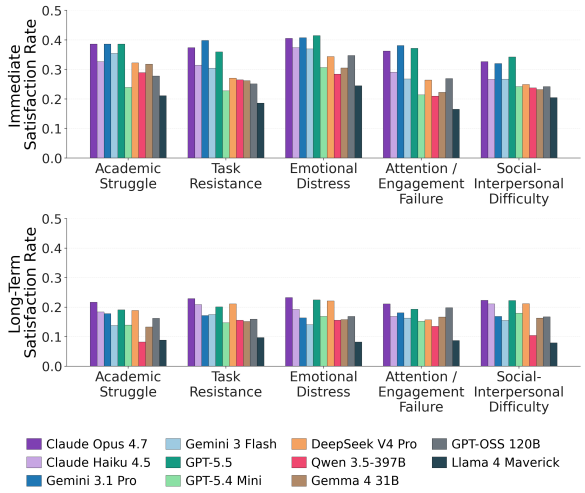

Figure 11: Average criterion satisfaction rates across models broken down by challenge type.

tion. Using a 25% ($N = 75$) stratified sample of source papers, we rerun the full construction and evaluation pipeline with GPT (`gpt-5.4-mini`, `gpt-5.5`) and Claude (`claude-sonnet-5.5` (Anthropic, 2026b), `claude-opus-4.7`) instead of Gemini, while keeping the prompts and the set of 11 evaluated models fixed. As each pipeline generates its own evaluation items, this yields 6,721 evaluated responses for GPT, 4,928 for Claude, and 3,850 for Gemini on the same source papers.

Model rankings remained highly consistent across construction models, with Spearman correlation being $\rho = 0.943$ ($p < 0.001$) for the Gemini- versus Claude-built comparison and $\rho = 0.925$ ($p < 0.001$) for the Gemini- versus GPT-built comparison. These results suggest that our findings are not sensitive to the choice of construction model.

M Impact of Judge Model Choice on Evaluation Results

We examine whether the open-ended evaluation depends on the choice of judge model. Using the 25% ($N = 130$) stratified sample of student categories as in Appendix E, we re-evaluate responses from all 11 models with `gpt-5.5` and `claude-opus-4.7`, while keeping the rubrics and evaluation prompt fixed. This yields 4,323 evaluated responses and 12,617 criterion judgments for each judge.

Table 8 shows high agreement across judges at both the case and criterion levels. Model rankings are also highly consistent across all judge pairs, with Spearman correlations of 0.96–0.97, all statistically significant at $p < 0.001$. These results indicate that the relative performance of the evaluated models is preserved across judge choices.

Table 8: Pairwise agreement and ranking consistency across judges ($p < 0.001$ for all Spearman correlations).

Judge Pair	Case-level κ	Criterion-level κ	Spearman ρ
Gemini vs. GPT	0.75	0.77	0.96
Gemini vs. Claude	0.76	0.77	0.97
GPT vs. Claude	0.75	0.76	0.96

N Reference-Based Response Details

We construct the reference pool using a paper-level split stratified by challenge type, targeting 15% of papers for testing. This yields a *pool* of 255 papers and 785 scenarios and a *test* set of 45 papers and 138 scenarios, while ensuring that no test dialogue receives a reference from its source paper. We embed each scenario's challenging-behavior trigger and behavior patterns with `text-embedding-3-large` (OpenAI, 2024) and rank pool scenarios by cosine similarity. For each test dialogue, we retrieve the ten nearest pool scenarios, filter to those with an effective strategy for the target response type (immediate or long-term), and keep the three closest matches. Their matching effective and ineffective strategies are concatenated into a reference block, framed as examples from similar but different students that the responder may adapt, combine, or ignore (prompt in Appendix T.16). We evaluate the resulting responses with the same rubrics used in the open-ended experiment, enabling a direct paired comparison against the dialogue-only condition.

O Lexical Diversity

Figures 12 and 13 show the 3- and 5-gram variants of the distinct-n analysis for immediate responses. Spearman correlations show the same pattern as the 4-gram result (Figure 7), with $\rho = 0.855$ for 3-grams and $\rho = 0.791$ for 5-grams (both $p < 0.01$), supporting the association between higher lexical diversity and immediate-response success rates.

For long-term interventions, only the 3-gram correlation is significant ($\rho = 0.627$, $p < 0.05$).

P Case Specificity

Figure 14 shows the relationship between case-recovery rate and immediate-response success rates. Spearman correlation shows a positive association across models ($\rho = 0.691$, $p < 0.05$), suggesting

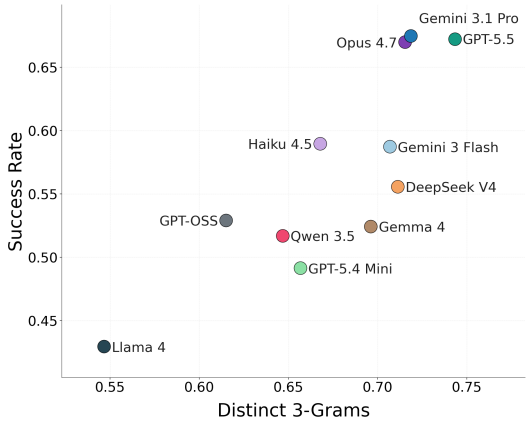

Figure 12: Correlation between distinct-n (n = 3) and immediate response success rate across models.

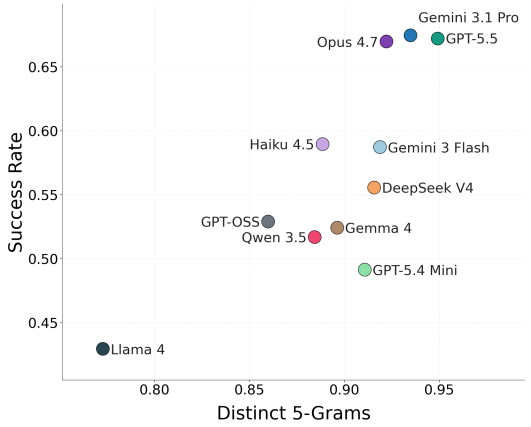

Figure 13: Correlation between distinct-n (n = 5) and immediate response success rate across models.

that stronger models produce responses that are more specific to the current dialogue.

Q MCQ Selection by Challenge Type

Figure 15 shows MCQ accuracy by challenge type. For immediate responses, social-interpersonal difficulty is the weakest category for all 11 models, while academic struggle is strongest for 10 of 11. For long-term plans, differences across challenge types are smaller, although social-interpersonal difficulty remains weakest for 7 of 11 models.

The contrast with open-ended evaluation is especially notable for these two categories. Academic struggle is the weakest category for all 11 models in long-term open-ended responses, but performs relatively well in MCQ selection. Social-interpersonal difficulty, in contrast, remains challenging in both settings. This suggests that academic strategies may be easier to recognize than generate, whereas social-interpersonal challenges remain difficult even when options are provided.

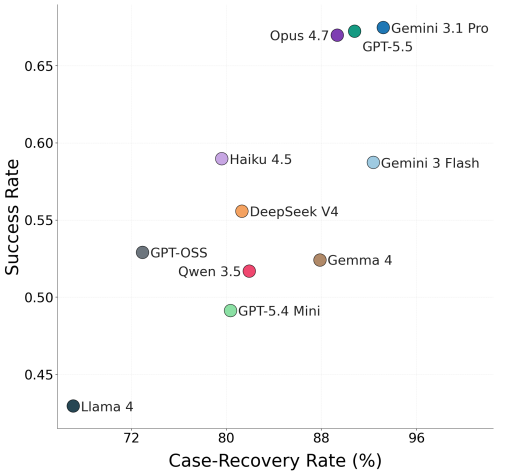

Figure 14: Correlation between case-recovery rate and immediate response success rate across models.

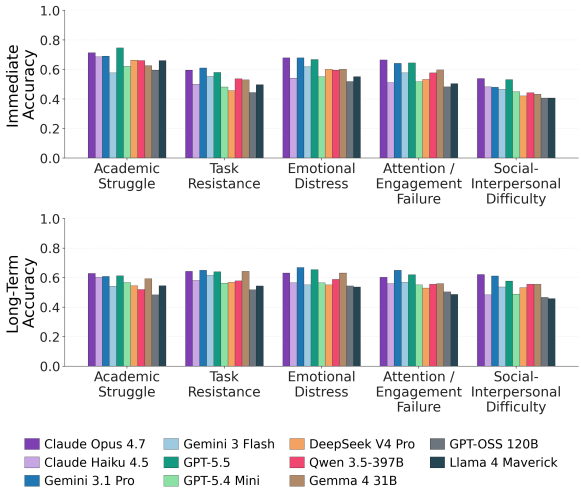

Figure 15: MCQ accuracy across models broken down by the challenge type of students.

R MCQ Failure Analysis

Figure 16 summarizes model agreement across MCQ questions. All 11 models answer 26.8% of questions correctly and miss 17.5% of questions, indicating substantial agreement on both easy and difficult cases. Errors also show systematic convergence. Among questions missed by all models, all 11 select the same distractor in 40.3% of cases.

Figure 17 breaks down these outcomes by challenge type. Social-interpersonal difficulty has the lowest all-correct share and the highest all-wrong share. It also has the highest share of questions where all models select the same wrong answer. These patterns suggest that social-interpersonal cases are not only more difficult, but also more likely to elicit the same errors across models.

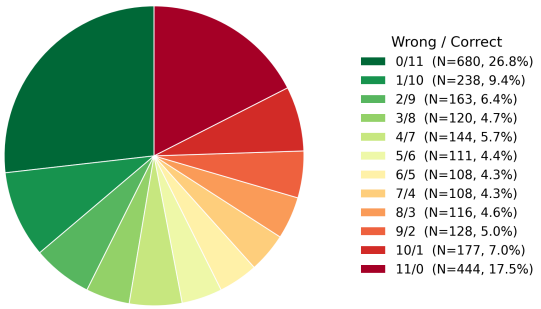

Figure 16: Distribution of model agreement on MCQ outcomes across all 2,537 items.

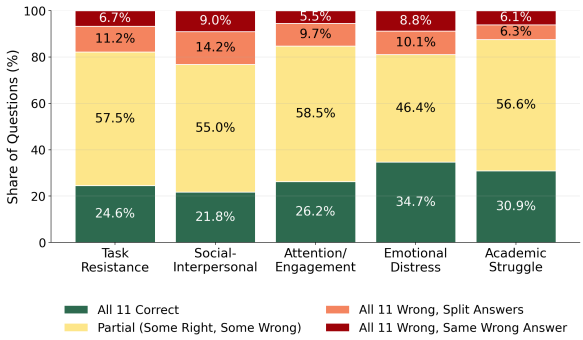

Figure 17: Model-agreement breakdown of MCQ outcomes by challenge type.

S Reference-Based Response

Figure 18 compares long-term intervention performance with and without reference cases. In contrast to immediate responses, adding references often reduces performance for long-term intervention planning. Only 3 of 11 models improve their performance, 7 decline, and one is unchanged, leading to an average drop of 5.2 percentage points.

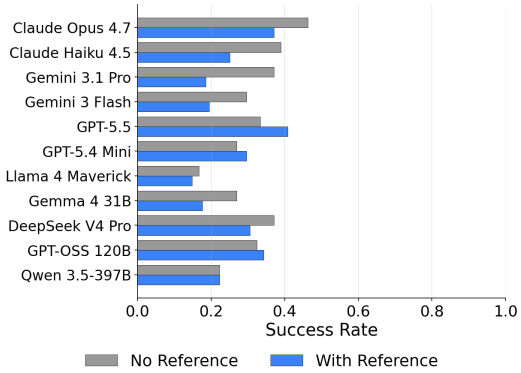

Figure 18: Success rate across models for long-term intervention plans, with and without references.

T Prompts Used in Dataset Construction

Each prompt is shown with its system and user components separated by the marker <System Prompt/User Prompt>.

T.1 Abstract-Based Paper Filtering

Prompt

```
You are a research screener selecting qualitative education papers that describe how challenging students behave in classroom settings and what teachers can do about it. Respond ONLY with a valid JSON object, no markdown, no extra text.

<System Prompt/User Prompt>

Read the abstract of the following education research paper. For each criterion below, answer Y (= filter out) or N (= keep). A paper is accepted only if ALL four answers are N.

Title: {title}
Abstract: {abstract}

F1. Filter out if this study does NOT use observation or interviews with students as a research method.
→ Y if only teachers, parents, or caregivers were observed or interviewed, or if neither method appears at all.
→ N if students are directly observed or interviewed, or if the study design (e.g. case study, ethnographic study centered on a student) clearly implies it.

F2. Filter out if this study only tests the effectiveness of a specific intervention or teaching method through quantitative measurement.
→ Y if pre-post tests, effect sizes, control groups, or statistical significance are the main focus, or if the study uses a single-case experimental design to measure behavioral change as an outcome of a specific treatment.
→ N if the study is qualitative or descriptive without an efficacy-testing agenda.

F3. Filter out if this study focuses on aspects other than how students behave and respond in the classroom (learning context).
→ Y if the central focus is diagnosis criteria, family background, parent-child relationships, policy, teacher competency, theory-building, or the development and validation of a measurement tool or rating scale — rather than students' actual behavioral patterns.
→ N if concrete student behavior and reactions are the core of the study.

F4. Filter out if this study offers no practical guidance or implications for classroom teachers.
→ Y if there is no indication — explicit or implied — that the findings offer any takeaway for how teachers should understand or respond
```

```
to these students in a classroom setting.
→ N if classroom-relevant insight or implications for teachers can be reasonably inferred from the findings, even if not stated as a formal recommendation.

Return JSON only:
{
"f1_no_student_observation_or_interview": "Y" or "N",
"f2_quantitative_efficacy_study": "Y" or "N",
"f3_focus_other_than_behavior": "Y" or "N",
"f4_no_practical_guidance": "Y" or "N",
"reason": "one sentence explaining the overall judgment"
}
```

T.2 Paper Screening

Prompt

```
You are a research assistant that evaluates whether an educational paper contains enough concrete evidence to simulate a tutor-student interaction.

Score the paper on 3 dimensions using a 1–5 scale. Score 1 means the dimension is not described at all.

──────────
### DIMENSIONS
──────────

**1. Student Behavior** (1–5): Observable student actions and responses
- 1 = No student behavior described at all
- 2 = Vague/inferred (e.g., "seemed frustrated", "was disengaged")
- 3 = General behavior (e.g., "avoided eye contact", "raised hand frequently")
- 4 = Specific action in context (e.g., "stayed silent when questioned about fractions")
- 5 = Multiple concrete actions (e.g., "pushed the worksheet aside, crossed arms, and looked away when asked to continue")

**2. Learning Context** (1–5): Tasks, materials, and settings
- 1 = No learning context described at all
- 2 = Extremely broad (e.g., "during math instruction")
- 3 = General task type (e.g., "working on addition problems", "reading comprehension")
- 4 = Specific activity (e.g., "completing a worksheet on two-digit multiplication during independent work")
- 5 = Highly specific, reproducible (e.g., "solving 3/4 ÷ 1/2 using fraction strips during small-group instruction")

**3. Teacher Actions** (1–5): Instructional strategies and outcomes
- 1 = No teacher actions or instructional strategies described at all
- 2 = Actions implied (e.g., "teacher helped the student")
- 3 = General strategy (e.g., "provided scaffolding", "gave positive reinforcement")
- 4 = Concrete action (e.g., "asked guiding questions", "modeled the problem on the board")
- 5 = Strategy with documented outcome (e.g., "after breaking the problem into smaller steps, the student successfully solved three similar problems independently")

──────────
### OUTPUT FORMAT
──────────

Output ONLY valid JSON. No explanations, no markdown.

{
"student_behavior": { "score": 1-5, "reason": "one sentence" },
"learning_context": { "score": 1-5, "reason": "one sentence" },
"teacher_actions": { "score": 1-5, "reason": "one sentence" },
"passed": true if all three scores are ≥ 3, else false
}

<System Prompt/User Prompt>

──────────
### DOCUMENT TO SCREEN
──────────

{TEXT}

──────────
```

T.3 Student Category Extraction

Prompt

```
You are a research assistant specialized in extracting key information from educational studies.

Your goal is to IDENTIFY student categories explicitly described in the paper and ASSESS their suitability for simulation.

──────────
### CRITICAL RULES
──────────
1. Extract ONLY what is explicitly stated - NO inference, NO generalization
2. ALL excerpts MUST be VERBATIM from the document - NO paraphrasing
3. If not stated, use empty string ("") or null
4. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### TASK OVERVIEW
──────────

**STEP 1: Extract Basic Information**
- Metadata (title, author, year, description, research method)
- Research scope (country, education level,
```

```
discipline)

**STEP 2: Identify Student Categories**
- If paper describes students as one group → Create ONE category
- If paper describes individual students with DIFFERENT characteristics (e.g., avoidance vs. rushing) or DIFFERENT approaches needed (e.g., wait time vs. immediate redirection) → Create SEPARATE categories
- If paper explicitly names student types (e.g., "daydreaming students", "students with high anxiety") → Use those
- If paper describes individual students with similar characteristics and same approach works for them → ONE category is enough
- Avoid broad labels: Each category name must include specific behavioral/cognitive/affective characteristics (e.g., "Students who shut down when corrected" NOT "Challenging students")

**STEP 3: Assess Extraction Eligibility**
For EACH category, evaluate 3 dimensions:

**1. Student Behavior** (1-5): Observable actions and responses
- 1 = No student behavior described at all
- 2 = Vague/inferred (e.g., "seemed frustrated", "was disengaged")
- 3 = General behavior (e.g., "avoided eye contact", "raised hand frequently")
- 4 = Specific action in context (e.g., "stayed silent when questioned about fractions")
- 5 = Multiple concrete actions (e.g., "pushed the worksheet aside, crossed arms, and looked away when asked to continue")

**2. Learning Context** (1-5): Tasks, materials, settings
- 1 = No learning context described at all
- 2 = Extremely broad (e.g., "during math instruction")
- 3 = General task type (e.g., "working on addition problems", "reading comprehension")
- 4 = Specific activity (e.g., "completing a worksheet on two-digit multiplication during independent work")
- 5 = Highly specific, reproducible (e.g., "solving 3/4 ÷ 1/2 using fraction strips during small-group instruction on Wednesday")

**3. Teacher Actions** (1-5): Instructional strategies and outcomes
- 1 = No teacher actions or instructional strategies described at all
- 2 = Actions implied (e.g., "teacher helped the student")
- 3 = General strategy (e.g., "provided scaffolding", "gave positive reinforcement")
- 4 = Concrete action (e.g., "asked guiding questions", "modeled the problem-solving process on the board")
- 5 = Strategy with documented outcome (e.g., "after breaking the problem into smaller steps, the student successfully solved three similar problems independently")

For each dimension: Write assessment → Provide supporting verbatim excerpt (only ONE most representative excerpt) → Assign score

**Score 1 means the dimension is not described in the paper at all.**
Use 1 only when there is truly no evidence; use 2 when evidence exists but is very weak.

──────────
### OUTPUT STRUCTURE
──────────

{
"id": "{ID}",

"metadata": {
"title": "paper title",
"author": "author names or 'unknown' ",
"year": year as number or null,
"description": "brief description of the paper's focus and purpose",
"research_method": "1-2 sentences on study design"
},

"research_scope": {
"country": "country where study was conducted or empty string",
"education_level": "one of: early_childhood, elementary_school, middle_school, high_school, undergraduate, other",
"discipline": "one of: mathematics, science, language_arts, foreign_language, social_studies, engineering, teacher_education, interdisciplinary, or other"
},

"category_list": [
{
"category_id": "{ID}-01 (use the paper id '{ID}' as prefix, then -01, -02, etc. for each category)",
"name": "Generalized name of the student type (e.g., 'Low-engagement adolescents with behavioral challenges')",
"label": ["identifier used in the paper (e.g., 'Participant 1', 'Tom'); otherwise make up a single unique label (e.g. 'Label 1')"],
"categorization_reasoning": "Why this category exists",

"demographics": {
"age": number or null,
"gender": "male | female | non_binary | empty string",
"background": "educational or socioeconomic background if explicitly stated; otherwise empty string"
},

"extraction_eligibility": {
"student_behavior": {
"assessment": "What behavioral evidence exists in the paper? How concrete and detailed is it? (2-3 sentences)",
"excerpt": "verbatim excerpt",
"score": 1-5
},

"learning_context": {
```

```
"assessment": "What context evidence exists? How specific are the tasks/problems, materials, and settings? (2-3 sentences)",
"excerpt": "verbatim excerpt",
"score": 1-5
},

"teacher_actions": {
"assessment": "What instructional strategies are documented? How concrete are they? Are outcomes documented? (2-3 sentences)",
"excerpt": "verbatim excerpt",
"score": 1-5
}
}
}
],

"irrelevant_pages": {
"page_numbers": [],
"sections": []
}
}


<System Prompt/User Prompt>

──────────
### DOCUMENT TO ANALYZE
──────────

{TEXT}

──────────
```

T.4 Learner Evidence Extraction

Prompt

```
You are a research assistant specialized in extracting key information from educational studies.

PHASE 1 OF 2: LEARNER EVIDENCE EXTRACTION
Your goal is to extract ALL verbatim evidence about LEARNER BEHAVIOR, COGNITION, AND AFFECT.

──────────
### CRITICAL RULES
──────────

1. Extract EXTENSIVELY - MAXIMUM COVERAGE of learner evidence
2. ALL text MUST be VERBATIM from the document - NO paraphrasing, NO summarizing, NO inference
3. EVERY excerpt MUST include source location (page/section)
4. ONE excerpt can apply to MULTIPLE categories if it shows multiple aspects
5. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### INPUT CONTEXT
──────────

ALL STUDENT CATEGORIES:
{CATEGORIES}

- Use "name" and "label" field to identify which category_id(s) each excerpt refers to.
- Your task is to extract evidence for ALL of them.

──────────
### TASK OVERVIEW
──────────

**STEP 1: Scan and Tag**
Read the entire paper and mentally note ALL locations containing STUDENT-CENTERED content:
- Student behaviors (actions, responses, engagement patterns, physical movements)
- Cognitive states (thinking processes, misconceptions, understanding levels)
- Affective states (emotions, motivation, confidence, attitudes)
- Difficulty triggers (conditions/situations causing struggle)
- Student quotes (any direct utterances)

DO NOT tag TEACHER-CENTERED content:
- Pure intervention descriptions without student behavior ("Teachers implemented X") → Skip for PHASE 1

**STEP 2: Extract Evidence**
For EACH tagged location:
- Copy excerpt with 2-3 surrounding sentences VERBATIM (word-for-word from paper)
- Focus on passages that primarily describe the STUDENT (even if teacher actions are mentioned as context)
- Identify all applicable category_id(s) based on labels
- Classify excerpt_type: behavior | cognitive | affective | difficulty_trigger
- Describe task_context: Include (1) specific task/problem, (2) materials, (3) instructional setting (e.g., "solving word problem 'If train leaves at 2pm going 60mph...' on worksheet during independent practice, 10 minutes into activity", "answering comprehension questions about Chapter 3 using graphic organizer during whole-class discussion")
- For quotes: extract EXACT STUDENT WORD with surrounding context (avoid extracting other's word)

Target: At least 10 excerpts per category (if paper has sufficient content)

**STEP 3: Verify Coverage**
Check for gaps:
- Any category with < 10 excerpts? → Search for more evidence if possible
- Missing evidence types? (behavior, cognitive, affective, triggers, quotes) → Search harder
- Any vague task_context? → Add concrete details from paper

**STEP 4: Verbatim Verification (CRITICAL)**
For EVERY excerpt and quote, perform strict verification:
1. Re-check EACH item against the source document
```

```
2. Verify word-by-word match (including punctuation, capitalization)
3. Record result in verbatim_check dictionaries:
- excerpt_verbatim_check: {"{ID}_learner_001": true, "{ID}_learner_002": false}
- quote_verbatim_check: {"{ID}_quote_001": true}
4. If ANY item is not verbatim (paraphrased/summarized), mark as false
5. Explain any non-verbatim items in verbatim_reasoning


──────────
### OUTPUT STRUCTURE
──────────

{
"paper_id": "{ID}",
"category_ids": ["list of all category_ids"],

"learner_excerpts": [
{
"excerpt_id": "{ID}_learner_001",
"applicable_categories": ["{ID}-01"],
"excerpt_text": "VERBATIM passage with 2-3 surrounding sentences for full context",
"source_location": "page X / section Y",
"excerpt_type": "behavior | cognitive | affective | difficulty_trigger",
"task_context": "DETAILED description: What specific task was the student doing? Include concrete details like actual problem content, materials used, task stage."
}
],

"student_quotes": [
{
"quote_id": "{ID}_quote_001",
"applicable_categories": ["{ID}-01"],
"quote_text": "Exact words spoken by student",
"source_location": "page X / section Y",
"context": "What was happening when student said this (task, situation, trigger)"
}
],

"extraction_summary": {
"total_excerpts": count,
"excerpts_per_category": {
"{ID}-01": count,
"{ID}-02": count
},
"excerpt_verbatim_check": {
"{ID}_learner_001": true,
"{ID}_learner_002": true,
"{ID}_learner_003": false
},

"total_quotes": count,
"quotes_per_category": {
"{ID}-01": count
},
"quote_verbatim_check": {
"{ID}_quote_001": true,
"{ID}_quote_002": false
},

"verbatim_reasoning": "Brief: Explain any non-verbatim items (why marked false). If all true, state 'All excerpts and quotes verified as verbatim'.",
"verbatim_check_performed": true
}
}


<System Prompt/User Prompt>

──────────
### DOCUMENT TO ANALYZE
──────────

{TEXT}

──────────
```

T.5 Instructional Evidence Extraction

Prompt

```
You are a research assistant specialized in extracting key information from educational studies.

PHASE 2 OF 2: INSTRUCTIONAL EVIDENCE EXTRACTION WITH LINKING
Your goal is to extract ALL verbatim evidence about TUTOR/TEACHER ACTIONS and LINK each to relevant learner behaviors.

──────────
### CRITICAL RULES
──────────

1. Extract EXTENSIVELY - MAXIMUM COVERAGE of instructional evidence
2. ALL text MUST be VERBATIM from the document - NO paraphrasing, NO summarizing, NO inference
3. EVERY excerpt MUST include source location (page/section)
4. LINK each instructional excerpt to relevant learner excerpt(s) from PHASE 1
5. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### INPUT CONTEXT
──────────

ALL STUDENT CATEGORIES:
{CATEGORIES}

LEARNER EXCERPTS from PHASE 1:
{LEARNER_EXCERPTS}

- Use "name" and "label" field to identify which category_id(s) each excerpt refers to.
- Your task is to extract evidence for ALL of them.

──────────
### TASK OVERVIEW
──────────
```

```
**STEP 1: Scan and Tag**
Read the entire paper and mentally note ALL locations containing TEACHER-CENTERED content:
- Effective teacher actions (strategies that worked, recommended practices)
- Ineffective teacher actions (strategies that failed or actions to be avoided)
- Intervention descriptions (what teachers DID or SHOULD DO)
- Outcomes (student responses to interventions - as results of teacher actions)

DO NOT tag STUDENT-CENTERED content already covered in PHASE 1:
- Pure behavioral observations without intervention context → Skip for PHASE 2

**STEP 2: Extract Evidence**
For EACH tagged location:
- Copy excerpt with 2-3 surrounding sentences VERBATIM (word-for-word from paper)
- Focus on passages that primarily describe what TEACHERS did or should do
- Identify applicable category_id(s)
- Classify excerpt_type: tutor_action_effective | tutor_action_ineffective
- Fill instructional_context:
* when: Conditions when this applies (e.g., "when student shows frustration")
* how: Concrete steps/approach (e.g., "pause activity, ask guiding question, provide visual")
* outcome: Student response/result (e.g., "student successfully solved similar problem")
- Classify temporal_scope:
* immediate: Applied in-the-moment (e.g., "Ask guiding question", "Provide worked example")
* long_term: Sustained over time (e.g., "Establish error-friendly culture", "Gradually increase difficulty")
* both: Has immediate AND sustained components

**STEP 3: Link to Learner Behaviors**
For EACH instructional excerpt:
- Review learner excerpts from PHASE 1
- Identify which learner behaviors this instruction addresses
- Write explicit linking_rationale with evidence from paper
- Assign link_strength:
* 3 (Strong): Paper explicitly connects instruction to this behavior
* 2 (Moderate): Addresses same challenge type; temporal scopes align
* 1 (Weak): General strategy that could apply; same category
* 0 (No link): Different behavior/category or temporal mismatch

**STEP 4: Verify Coverage**
Check for gaps:
- Any orphaned learner excerpts (no linked instructions)? → Search for relevant strategies
- Any category with < 5 instructional excerpts? → Search harder
- Any vague instructional_context? → Add concrete details from paper

**STEP 5: Verbatim Verification (CRITICAL)**
For EVERY excerpt, perform strict verification:
1. Re-check EACH item against the source document
2. Verify word-by-word match (including punctuation, capitalization)
3. Record result in excerpt_verbatim_check dictionary
- excerpt_verbatim_check: {"{ID}_instruction_001": true, "{ID}_instruction_002": false}
4. If ANY item is not verbatim (paraphrased/summarized), mark as false
5. Explain any non-verbatim items in verbatim_reasoning

──────────
### OUTPUT STRUCTURE
──────────
{
"paper_id": "{ID}",
"category_ids": ["list of all category_ids"],

"instructional_excerpts": [
{
"excerpt_id": "{ID}_instruction_001",
"applicable_categories": ["{ID}-01"],
"excerpt_text": "VERBATIM passage with 2-3 surrounding sentences for full context",
"source_location": "page X / section Y",
"excerpt_type": "tutor_action_effective | tutor_action_ineffective",

"instructional_context": {
"when": "Conditions when this applies; empty string if not stated",
"how": "Concrete steps/approach; empty string if not stated",
"outcome": "Student response/result; empty string if not stated"
},

"linked_learner_excerpts": [
{
"learner_excerpt_id": "{ID}_learner_003",
"linking_rationale": "EXPLICIT: Why this instruction addresses this behavior, with evidence from paper",
"link_strength": 0-3
}
],

"temporal_scope": "immediate | long_term | both"
}
],

"extraction_summary": {
"total_excerpts": count,
"excerpts_per_category": {
"{ID}-01": count
},
"excerpt_verbatim_check": {
"{ID}_instruction_001": true,
"{ID}_instruction_002": false
},
"verbatim_reasoning": "Brief: Explain any non-verbatim items. If all true, state 'All excerpts verified as verbatim'.",
```

```
"verbatim_check_performed": true,

"linking_summary": {
"total_links": count,
"strong_links": count,
"moderate_links": count,
"weak_links": count,
"orphaned_learner_excerpts":
["{ID}_learner_012"],
"orphaned_instruction_excerpts":
["{ID}_instruction_015"],
"linking_reasoning": "Brief: Quality of links between instructions and learner behaviors."
}
}
}


<System Prompt/User Prompt>

──────────
### DOCUMENT TO ANALYZE
──────────

{TEXT}

──────────
```

T.6 Scenario Construction

Prompt

```
You are a research assistant specialized in building structured scenarios from educational research evidence.

Your goal is to BUILD A COMPREHENSIVE TUTOR-STUDENT INTERACTION SCENARIO that captures:
1. The CONTEXT where challenging interactions occur
2. The CHALLENGING BEHAVIORS the student exhibits
3. HIDDEN STRATEGIES about what works and doesn't work (for evaluation purposes)

──────────
### CRITICAL RULES
──────────
1. Ground ALL statements in provided excerpts - NO INFERENCE
2. Use EXACT descriptions from excerpts, not interpretations - DO NOT create behaviors or actions not explicitly in evidence
3. Reference excerpts by ID for EVERY claim (e.g., "0001_learner_001")
- evidence_basis arrays may contain ONLY IDs listed in VALID EVIDENCE IDS
- NEVER use meta labels such as "category_extraction", "category", "paper", "prompt", or any invented evidence ID
4. CHALLENGING BEHAVIORS = student behaviors that are difficult for a tutor to address using standard instructional strategies — extract as many as possible
5. The scenario describes the PROBLEM, not the SOLUTION. Effective strategies are kept as "hidden strategies" for evaluating tutor responses later — NO SOLUTION LEAKAGE into the visible content
6. Focus on PERSONALIZED, STUDENT-SPECIFIC strategies tailored to this individual student's unique needs and characteristics
7. Every strategy that references specific tools, materials, or conditions (e.g., visual aids, manipulatives) MUST have those elements documented in the context field - If not in context, do not include the strategy
8. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### INPUT CONTEXT
──────────

STUDENT CATEGORY:
{CATEGORY}

- Focus ONLY on evidence explicitly about this student.
- The category description is context for selecting the student only. It is NOT an evidence source and must NEVER appear in evidence_basis.

VALID EVIDENCE IDS:
{ALLOWED_EVIDENCE_IDS}

──────────
### TASK OVERVIEW
──────────

**STEP 1: Scan and Tag**
Read the entire evidence and mentally note ALL locations containing STUDENT-CENTERED content:
- What CONTEXT(S) does the student struggle in?
- What BEHAVIORS does the student exhibit when struggling?
- What strategy WORKS for this student?
- What strategy DOESN'T WORK for this student?

**STEP 2: Identify Distinct Scenarios**
Generate as many naturally distinct scenarios as the evidence supports. Recommended: ~{NUM_SCENARIOS}. Each scenario contains one specific task/context where student struggled, their challenging behaviors, and appropriate strategies for tutors.
- If DIFFERENT specific task/context found → Create separate scenarios
- If DIFFERENT strategies should be applied → Create separate scenarios
- Do NOT force artificial scenarios if the evidence does not support them — fewer scenarios of high quality is better than more scenarios of low quality
- If the evidence supports more than {NUM_SCENARIOS} meaningfully distinct scenarios, include all of them
- Each scenario MUST be meaningfully different from the others — not just surface-level rewording
- All scenarios must be grounded in the provided evidence

**STEP 3: For EACH Scenario - Extract
```

```
Context & Task**
Gather evidence about WHERE, WHEN and WHAT (will be used for dialogue opening):
- Task: Exact problem content (e.g., “Solve 3/4 ÷ 1/2 using fraction strips”, “Build 5-block tower in construction center”), learning goal if stated
- Context: Physical environment, materials, social setting, time if available
- NOTE: Only materials/tools documented here can appear in hidden strategies later

**STEP 4: For EACH Scenario - Extract Student Behaviors & Triggers**
Identify HOW student struggles (will be used for dialogue middle/end):
- Challenging behavior trigger: The single most critical moment
- Behavior patterns: Recurring observable behaviors the student shows during struggle
- All behaviors (both fields) must satisfy:
* Observable: What can a tutor physically SEE? (not inferred states)
* Specific: Exact actions, not vague labels (not “frustrated” but “sighed, crossed arms, said ‘this is stupid’ ”)
* Trigger-linked: Each behavior must be paired with what caused it
* Internal state: Documented thoughts/feelings if available
- Include ALL related student quotes VERBATIM where available

**STEP 5: For EACH Scenario - Extract Hidden Strategies (STUDENT-SPECIFIC)**
Extract effective/ineffective strategies (be as specific as possible, leave empty array if no explicit strategies documented):
- Strategies should be decomposed into single, actionable units feasible within the given context
- Extract strategies TAILORED to THIS student’s unique needs, NOT general best practices
- Strategies must be aligned with the scenario’s task, context, and challenging behavior trigger
- ALIGNMENT CHECK: If a strategy references a tool or condition not in the context field → exclude it
- For EACH strategy, extract four components:
* Instruction: Exact description of what tutor did/should do
* Context: Exact triggering situation from evidence
* Outcome: Exact documented result/student response
* Reason: Why this strategy is effective/ineffective for THIS student (mechanism, not just result)
- Effective immediate: In-the-moment responses (3+ if possible)
- Effective long-term: Sustained approaches over weeks/months (3+ if possible)
- Ineffective: Strategies that failed or should be avoided (ONLY if paper EXPLICITLY describes failure, 3+ if possible)

**STEP 6: Extract Student Persona**
Identify and synthesize core differentiating characteristics across three dimensions:
- Behavioral patterns: What this student characteristically does (e.g., “Elopes when redirected, not when task is hard”)
- Unique triggers & specific needs: What causes struggle and what this student uniquely requires (e.g., “Transitions cause shutdown; requires 10+ second wait time”)
- Approach effectiveness: What works and what doesn’t for THIS student specifically (e.g., “Visual models work; verbal praise increases frustration”)
Write description (3-4 sentences) that weaves these characteristics together to answer:
- What makes them challenging?
- What are their core differentiating characteristics?

**STEP 7: Categorize Overall Challenges**
Identify the primary challenge type (required) and secondary challenge type (optional — only if the evidence clearly supports a meaningfully distinct second challenge). The subcategories below describe what each type covers:
- **Task Resistance** — explicit task refusal, passive noncompliance, task escape behavior, help avoidance, gaming/shortcutting
- **Emotional Distress** — frustration expression, anxiety expression, learned helplessness, low self-efficacy, sadness/low mood
- **Attention/Engagement Failure** — off-task motor, off-task verbal, off-task passive, boredom/disinterest, distractibility
- **Social-Interpersonal Difficulty** — verbal aggression/hostility, physical aggression, defensiveness/blame externalization, social withdrawal, attention-seeking disruption
- **Academic Struggle** — confusion/cognitive disequilibrium, persistent/repeated errors, uncertain response, confident misconception, failure to respond

——————————
### OUTPUT STRUCTURE
——————————

{
“category_id”: “{CATEGORY_ID}”,
“paper_id”: “{ID}”,
“name”: “student category name”,

“scenarios”: [
{
“scenario_id”: “{CATEGORY_ID}_scenario_001”,
“title”: “Brief descriptive title (e.g., ‘Multi-step fraction problem with visual supports’)”,

“content”: {
“task”: {
“content”: “EXACT problem/activity content from excerpts (e.g., ‘Divide 3/4 by 1/2 and explain the steps’)”,
“learning_goal”: “EXACT objective from excerpts; empty string if not stated”,
“evidence_basis”: [“0001_learner_003”]
},

“context”: {
```

```
"environment": "EXACT setting from excerpts (e.g., 'Quiet corner of classroom during math block')",
"materials": "EXACT items from excerpts (e.g., 'Worksheet, fraction strips, pencil, visual aid poster on wall')",
"social_setting": "EXACT arrangement from excerpts (e.g., 'One-on-one tutoring session, other students working independently')",
"time_context": "If stated in excerpts; empty string if not stated",
"evidence_basis": ["0001_learner_004"]
},

"challenging_behavior_trigger": {
"condition": "EXACT final triggering condition from excerpts (e.g., 'After third incorrect attempt when tutor asked to try again')",
"observable_behavior": "Observable actions EXACTLY from excerpts (e.g., 'Pushed worksheet away, crossed arms, turned head away from tutor')",
"internal_state": "EXACT cognition/affect from excerpts if documented; empty string if not stated",
"related_quotes": ["VERBATIM quotes"],
"evidence_basis": ["0001_learner_005"]
},

"behavior_patterns": [
{
"condition": "EXACT trigger from excerpts (e.g., 'When asked to explain reasoning')",
"observable_behavior": "EXACT actions from excerpts (e.g., 'Looked down, fidgeted with pencil, said nothing for 8 seconds')",
"internal_state": "EXACT cognition/affect from excerpts if documented; empty string if not stated",
"related_quotes": ["VERBATIM quotes"],
"evidence_basis": ["0001_learner_002"]
}
]
},

"hidden_strategies": {
"effective_immediate": [
{
"action_id": "{CATEGORY_ID}_scenario_001_effim_01",
"instruction": "Description of recommended immediate strategy EXACTLY from excerpts (e.g., 'Paused activity, acknowledged frustration, asked What part makes sense to you?, provided fraction strip model')",
"context": "Context where strategy was tried/recommended EXACTLY from excerpts (e.g., 'When student showed frustration after making error on division problem')",
"outcome": "Result or student response EXACTLY from excerpts (e.g., 'Student re-engaged, used strips to solve problem, completed 3 similar problems independently')",
"reason": "Why this works for THIS student - mechanism from excerpts (e.g., 'Student responds to choice and acknowledgment; reduces perceived threat of failure')",
"evidence_basis": ["0001_instruction_002"]
}
],

"effective_long_term": [
{
"action_id": "{CATEGORY_ID}_scenario_001_efflt_01",
"instruction": "Description of recommended long-term strategy EXACTLY from excerpts",
"context": "Context where strategy was tried/recommended EXACTLY from excerpts",
"outcome": "Result over time EXACTLY from excerpts",
"reason": "Why this works for THIS student - mechanism from excerpts",
"evidence_basis": ["0001_instruction_007"]
}
],

"ineffective": [
{
"action_id": "{CATEGORY_ID}_scenario_001_ineff_01",
"instruction": "Description of strategy to avoid EXACTLY from excerpts (e.g., 'Immediately corrected error verbally without wait time or visual support')",
"context": "Context where strategy was tried EXACTLY from excerpts",
"outcome": "Negative result EXACTLY from excerpts (e.g., 'Student shut down, refused to continue, said I told you I cannot do this')",
"reason": "Why this fails for THIS student - mechanism from excerpts (e.g., 'Student interprets immediate correction as confirmation of inability')",
"evidence_basis": ["0001_instruction_015"]
}
]
}
}
],

"student_persona": {
"behavioral_patterns": "What this student characteristically does during struggle — EXACT from excerpts (e.g., 'Pushes materials away, goes silent, refuses to re-engage')",
"unique_triggers_and_needs": "What causes breakdown and what this student uniquely requires — EXACT from excerpts (e.g., 'Shuts down after repeated errors; needs error-reframing before re-engagement')",
"approach_effectiveness": "What works and what doesn't for THIS student specifically — EXACT from excerpts (e.g., 'Visual models effective; verbal-only correction increases resistance')",
"description": "3-4 sentences integrating the above three dimensions: what makes them challenging and their core differentiating characteristics",
"evidence_basis": ["0001_learner_002", "0001_learner_005"]
},

"challenge_categorization": {
"primary_challenge": {
"evidence_summary": "FIRST: Specific behaviors/patterns",
```

```
"type": "THEN: Task Resistance | Emotional Distress | Attention/Engagement Failure | Social-Interpersonal Difficulty | Academic Struggle",
"evidence_basis": ["0001_learner_002", "0001_learner_005"]
},
"secondary_challenge": { // null if evidence does not clearly support a second distinct challenge type
"evidence_summary": "FIRST: Specific behaviors/patterns",
"type": "THEN: Task Resistance | Emotional Distress | Attention/Engagement Failure | Social-Interpersonal Difficulty | Academic Struggle",
"evidence_basis": ["0001_learner_007"]
}
}
}

<System Prompt/User Prompt>

──────────
### INPUT EVIDENCE
──────────

LEARNER EVIDENCE:
{LEARNER_EXCERPT}

INSTRUCTIONAL EVIDENCE:
{INSTRUCTION_EXCERPT}

──────────
```

T.7 Dialogue Synthesis

Prompt

```
You are a research assistant specialized in creating realistic dialogue from educational research scenarios.

Your goal is to CREATE REALISTIC STUDENT-TUTOR DIALOGUES that demonstrate CHALLENGING MOMENTS requiring intervention.

──────────
### CRITICAL RULES
──────────
1. Ground ALL statements in provided profile/excerpts - NO INFERENCE
2. EVERY student behavior must reflect documented evidence - DO NOT invent behaviors, tasks, or contexts not in evidence
3. Reference evidence by ID for EVERY message
- evidence_basis arrays may contain ONLY IDs that appear as keys in EVIDENCE
- NEVER use meta labels such as "category_extraction", "scenario", "persona", "paper", "prompt", or invented IDs
4. Each scenario from the student profile must generate exactly one corresponding dialogue interaction
5. Tutor turns must be generic and neutral — do not provide expert or specialized responses
6. The dialogue sets up a situation that REQUIRES good intervention
7. Final student turn shows challenging behavior that needs response (CHALLENGING BEHAVIORS = student behaviors that are difficult for a tutor to address using standard instructional strategies)
- The final turn MUST be a student turn and MUST include "is_challenging_behavior": true
- Every student turn MUST include a non-empty evidence_basis array
8. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### TASK OVERVIEW
──────────

**STEP 1: For EACH Scenario - Plan Dialogue**
Review scenario from profile:
- Task: What specific problem/activity? (concrete from evidence)
- Context: Environment/materials/social setting (verify action possibilities)
- Student persona: Core characteristics, typical responses
- Difficulty trigger: What challenging behavior will end the dialogue? What causes final struggle?
- Available actions: What tutor/student can physically do?

**STEP 2: For EACH Scenario - Structure Dialogue**
- Number of turns = 2 × {TURN} (e.g., TURN=1 → 2 turns, TURN=3 → 6 turns, default TURN=3)
- **First Turn (Tutor)**: Present SPECIFIC, CONCRETE task
- State the specific task from the scenario context (neutral tone)
- Generic and neutral only — no expert or specialized responses
- **Turns in the middle (Mixed)**: Student attempts → Tutor NEUTRAL response
- Student: May show emerging problem behaviors (based on evidence, follow persona, use quotes when appropriate)
- Tutor: Generic/neutral response only (e.g., "What do you think?", "Keep trying") - DO NOT use effective strategies
- Physical actions in parentheses ONLY if grounded in evidence and context allows
- **Last Turn (Student)**: MUST show challenging/struggling behavior
- Observable problem behaviors matching challenging_behavior_trigger
- Grounded in documented behavior_patterns

**STEP 3: For EACH Scenario - Generate Dialogue**
Follow plan, use student quotes from evidence, maintain persona (do NOT assume active engagement or verbal expressiveness unless documented)

**STEP 4: For EACH Scenario - Write Minimal Context**
Using the provided SCENARIO, write one natural sentence describing what the student is doing and what the goal is.
```

```
- Include: task + setting
- Exclude: student characteristics, struggles, or solutions

──────────
### OUTPUT STRUCTURE
──────────
{
"category_id": "{CATEGORY_ID}",
"paper_id": "{ID}",
"name": "student category name",

"dialogues": [
{
"case_id": "{CATEGORY_ID}_case_001",

"dialogue": [
{
"turn_id": 1,
"speaker": "tutor",
"message": "Clear task presentation. Neutral tone, no effective strategies"
},
{
"turn_id": 2,
"speaker": "student",
"message": "Response based on persona",
"evidence_basis": ["0001_learner_002"]
},
{
"turn_id": 3,
"speaker": "tutor",
"message": "NEUTRAL response (generic)"
},
{
"turn_id": 4,
"speaker": "student",
"message": "Continued response",
"evidence_basis": ["0001_learner_005"]
},
{
"turn_id": 5,
"speaker": "tutor",
"message": "NEUTRAL response (generic)"
},
{
"turn_id": 6,
"speaker": "student",
"message": "CHALLENGING BEHAVIOR matching challenging_behavior_trigger",
"is_challenging_behavior": true,
"evidence_basis": ["0001_learner_008"]
}
],

"minimal_context": "One natural sentence: what student is doing + goal. No characteristics, struggles, or solutions."
}
]
}

<System Prompt/User Prompt>

──────────
### INPUT CONTEXT
──────────

STUDENT CATEGORY:
{CATEGORY}

STUDENT PERSONA:
{STUDENT_PERSONA}

SCENARIO:
{SCENARIO}

EVIDENCE (excerpt_id → actual text from paper):
{EVIDENCE}

──────────
```

T.8 Rubric Construction

Prompt

```
You are an expert in educational research, specialized in designing evaluation rubrics for tutors.

Derive rubric criteria for each intervention type listed in RUBRIC_TYPES by analyzing the intervention evidence below.
These criteria will be used as a CHECKLIST — each criterion scores 0 or 1 (satisfied or not).
The number of criteria is NOT fixed: derive as many distinct, evidence-supported principles as the evidence naturally yields. Do NOT pad to a fixed number; do NOT invent principles without evidence.

Do NOT start from generic evaluation dimensions. Instead, read the evidence first and ask:
- What underlying principles make these interventions work for THIS student?
- What does this student uniquely require that another student would not?

──────────
### HOW TO DERIVE THE RUBRIC
──────────

**Step 1 — Read the intervention evidence.**
Three evidence streams are provided:
- EFFECTIVE_IMMEDIATE: strategies that work for this student in the moment
- EFFECTIVE_LONG_TERM: strategies that work for this student over weeks/months
- INEFFECTIVE: strategies that fail or backfire for this student (applies to both rubric types)

**Step 2 — Derive candidate principles separately for each rubric type.**
In principle_analysis, derive two separate principle lists:
- immediate_candidate_principles: from EFFECTIVE_IMMEDIATE only
- long_term_candidate_principles: from EFFECTIVE_LONG_TERM only
Each principle should be:
- grounded in the evidence (briefly note which strategies support it)
- capture something THIS student uniquely requires
```

```
- broad enough to allow diverse tutor approaches to satisfy
If two principles in the same list are redundant or overlapping, merge them.
Extract as many meaningful and DISTINCT principles as the evidence supports.

**Step 3 — Write each criterion as a CHECKLIST item.**
For each principle in immediate_candidate_principles, write one criterion in immediate_rubric.
For each principle in long_term_candidate_principles, write one criterion in long_term_rubric.
For each criterion, write:
- A criterion_name capturing the underlying principle (NOT a meta-label like "Use of effective strategy")
- A rationale (1 sentence): why this principle matters for THIS student
- A check condition (1 sentence): what must be present or absent in the tutor's response to earn 1 point.
The check should be principle-level, not method-level — diverse tutor approaches can satisfy it.

──────────
### CRITICAL CONSTRAINTS
──────────

- Do NOT name criteria using meta-labels like "Use of effective strategies", "Avoidance of ineffective strategies", "Personalization", "Task re-engagement". Use names that capture the underlying principle.
- Each criterion must be grounded in THIS student's documented evidence, not general best practice.
- The check condition must be satisfiable in multiple valid ways — good tutor responses are diverse.
- There are two rubric types: **immediate** (what the tutor does right now in the moment) and **long_term** (sustained strategies over weeks/months).
- RUBRIC_TYPES specifies which types of rubric to generate. If both are listed, generate both. If only one is listed, generate that one and set the other to null.
- Number of criteria is variable. Do NOT pad to a fixed number. Do NOT invent principles without evidence support.

──────────
### OUTPUT FORMAT
──────────

Output ONLY valid JSON:

{
"student_description": "2-3 sentences: who this student is, their core behavioral pattern, and what the evidence reveals they uniquely need",

"principle_analysis": {
"from_target_interventions": "What key principles emerge from the effective strategies? What do these reveal about this student's needs? (2-3 sentences)",
"student_specific_synthesis": "What is uniquely required for THIS student that another student would not need? (1-2 sentences)",
"immediate_candidate_principles": [
"Candidate principle 1: [name] — [brief rationale grounded in evidence]",
"Candidate principle 2: ..."
],
"long_term_candidate_principles": [
"Candidate principle 1: [name] — [brief rationale grounded in evidence]",
"Candidate principle 2: ..."
]
},

"immediate_rubric": {
"criteria": [
{
"criterion_id": 1,
"criterion_name": "Specific criterion capturing the principle",
"rationale": "1 sentence: why this principle matters for THIS student in an immediate response",
"check": "1 sentence: what must be present or absent in the tutor's immediate response to earn 1 point. Must be satisfiable by diverse approaches."
}
]
},

"long_term_rubric": {
"criteria": [
{
"criterion_id": 1,
"criterion_name": "Specific criterion capturing the principle",
"rationale": "1 sentence: why this principle matters for THIS student over time",
"check": "1 sentence: what must be present or absent in the tutor's long-term plan to earn 1 point. Must be satisfiable by diverse approaches."
}
]
}
}

If a rubric type is NOT in RUBRIC_TYPES, set that field to null:
"immediate_rubric": null
"long_term_rubric": null


<System Prompt/User Prompt>

──────────
### STUDENT PROFILE
──────────

DESCRIPTION:
{STUDENT_DESCRIPTION}

BEHAVIORAL PATTERNS:
{BEHAVIORAL_PATTERNS}
```

```
UNIQUE TRIGGERS AND NEEDS:
{UNIQUE_TRIGGERS_AND_NEEDS}

APPROACH EFFECTIVENESS:
{APPROACH_EFFECTIVENESS}

PRIMARY CHALLENGE: {PRIMARY_CHALLENGE}

──────────
### ACTUAL DIALOGUE EXAMPLES
──────────

These are real dialogue cases showing how this student actually behaves in practice. Use these to ground your rubric in concrete, observable behavior — not just abstract descriptions.

{DIALOGUE_EXAMPLES}

──────────
### INTERVENTION EVIDENCE
──────────

EFFECTIVE_IMMEDIATE (strategies that work in the moment for THIS student):
{EFFECTIVE_IMMEDIATE}

EFFECTIVE_LONG_TERM (strategies that work over weeks/months for THIS student):
{EFFECTIVE_LONG_TERM}

INEFFECTIVE (strategies that fail or backfire for THIS student — applies to both rubric types):
{INEFFECTIVE}

──────────
### RUBRIC TYPES TO GENERATE
──────────

{RUBRIC_TYPES}
```

T.9 MCQ Distractors Selection

Prompt

```
You are picking 3 distractors for a multiple-choice tutoring question.

The CORRECT option is one effective intervention for THIS student in THIS dialogue.
The OTHER EFFECTIVE interventions listed below are also correct answers for the same student per the source paper.

You get up to 10 candidates, each a real effective intervention from a different scenario.
They were already embedding-filtered to be context-similar but behavior/intervention-different, so almost all are usable as-is.

──────────
BE LENIENT
──────────

A distractor only needs to be a plausible-sounding tutor action that is WRONG for THIS student. That is a low bar — most candidates clear it.

Pick the 3 best. ALWAYS return exactly 3.

Reject a candidate ONLY if one of these is true (rare, since embedding already filtered):

1. STRUCTURALLY IMPOSSIBLE ACTION — the action's core mechanism depends on a setting structure that does not exist in this scenario and cannot be created by a wording swap (e.g. target is a 1-on-1 session but the action IS "rotate the student across several small groups"; target is a single immediate moment but the action IS "run a multi-week sequenced project"). A merely different subject, textbook, tool, or material is NOT a reason — those are freely rewritten downstream by the adapter (math textbook → English textbook, worksheet → whiteboard, etc.). Reject only when keeping the mechanism would be incoherent in this setting.

2. TOO CLOSE TO CORRECT — its action is essentially the same (just reworded) as the CORRECT option or any OTHER EFFECTIVE intervention.

Anything else is a fine distractor — including ones that address a different student problem, use a quite different approach, or are only a so-so fit. Do NOT reject those.

If rejecting leaves fewer than 3 clean candidates, still return 3 — fill the remaining slots with the least-bad rejected ones and list those slots in 'forced_idxs'. Leave it empty if all 3 are clean.

──────────
INPUT
──────────

TARGET DIALOGUE:
{DIALOGUE}

TARGET STUDENT PERSONA:
{STUDENT_PERSONA}

TARGET SCENARIO CONTEXT (task, environment, materials, social setting):
{TARGET_CONTEXT}

RESPONSE TYPE: {RESPONSE_TYPE}

CORRECT INTERVENTION (the labeled correct answer — do not select):
{CORRECT_INTERVENTION}

OTHER EFFECTIVE INTERVENTIONS (also correct for this student per the source paper):
{SIBLING_INTERVENTIONS}

CANDIDATE DISTRACTORS (0-indexed, each from a different scenario):
```

```
{CANDIDATE_DISTRACTORS}

──────────
OUTPUT
──────────

Output ONLY valid JSON.

{
“kept_idxs”: [<int>, <int>, <int>], // EXACTLY 3
“forced_idxs”: [<int>, ...], // subset of kept_idxs rejected by rule 1 or 2 but kept to reach 3; [] if all clean
“reason”: “<one short sentence>”
}
```

T.10 MCQ Options Adaptation

Prompt

```
You are aligning the four answer options of a multiple-choice tutoring question so they read uniformly — same surface vocabulary, similar length, similar tone — without changing what each option substantively does.

You are given: the TARGET dialogue / context / persona, ONE correct option (effective for THIS student), and THREE distractors (real interventions from OTHER scenarios, wrong for THIS student). Options come from different sources, so they may use different vocabulary, lengths, and tones. Make them parallel so test-takers cannot identify the correct option by surface features alone.

──────────
RULES
──────────

PRESERVE
- The MECHANISM of each option (e.g., “scaffold by chunking”, “visual + verbal cue”, “behavior-specific praise”, “offer a binary choice”). Do not switch mechanisms.
- The correct option’s appropriateness for THIS student.
- Each distractor’s wrongness for THIS student. Do not soften.

CHANGE — surface only
- Subject/activity names, materials/tools, group structures, named curricula, source-scenario student or staff names → swap to the target dialogue’s vocabulary.
- If an option’s surface already fits the target context, leave it verbatim.

LENGTH & TONE PARITY
- Priority when rules conflict (higher wins):
1. Each option’s MECHANISM unchanged; the correct option stays right for THIS student; each distractor stays wrong for THIS student. NEVER violate this.
2. The four options roughly similar in length and tone.
- Equalize length by TRIMMING the longer options first.
- If trimming alone cannot close the gap, you MAY add a small amount of NEUTRAL filler to a short option — generic time/place/connective phrasing only
(e.g. “...during the activity”, “...in this session”). The added words MUST NOT introduce a new method, step, or technique; MUST NOT claim any effect, benefit, or empirical/evidence support (no “...which improves performance”, “...with strong research backing”); and MUST NOT make a distractor any more appropriate (or the correct any less) for THIS student. If you cannot pad without touching the mechanism or its fit, leave it short — an uneven length is acceptable; changing substance is not. - All four in third-person tutor-action voice (never direct speech, never “you”).
- Match register across all four: professional, concise, action-focused.

MINIMUM EDIT
- Swap only the surface terms the target context requires. No swap needed → output verbatim.
- Never introduce a new method/technique/step, never flip or soften substance. Neutral length-matching filler (per LENGTH & TONE PARITY) is the ONLY permitted addition.

EDIT LEVELS (one label for the whole pass)
- none : all four output verbatim.
- surface : ≤ 3 surface words swapped per option.
- light : at least one option needed a short phrase rewrite (no direct target analogue for a surface element).
- fuller : AVOID. Only if a setting structure had no target equivalent and had to be re-described. Re-check that all mechanisms are intact and no wrongness was softened. If any option received neutral length-matching filler, say so explicitly in edit_notes (e.g. “padded distractor 2 with neutral filler for length parity”).

──────────
INPUT
──────────

TARGET DIALOGUE:
{DIALOGUE}

TARGET SCENARIO CONTEXT (task, environment, materials, social setting):
{TARGET_CONTEXT}

TARGET STUDENT PERSONA:
{STUDENT_PERSONA}

RESPONSE TYPE: {RESPONSE_TYPE}

Each option below is given as its ACTION only. Adapt that action — do not add any rationale, justification, expected effect, or outcome (none is provided, and none should be inferred or written into the option).

CORRECT (effective for THIS student):
```

```
ACTION: {CORRECT_INSTRUCTION}

DISTRACTOR 1 (wrong for THIS student):
ACTION: {DISTRACTOR_1_INSTRUCTION}

DISTRACTOR 2 (wrong for THIS student):
ACTION: {DISTRACTOR_2_INSTRUCTION}

DISTRACTOR 3 (wrong for THIS student):
ACTION: {DISTRACTOR_3_INSTRUCTION}

──────────
OUTPUT
──────────

Output ONLY valid JSON.

{
"correct_adapted":        "<correct    option,
surface-aligned; mechanism preserved>",
"distractor_1_adapted":      "<distractor    1,
surface-aligned; wrongness preserved>",
"distractor_2_adapted":      "<distractor    2,
surface-aligned; wrongness preserved>",
"distractor_3_adapted":      "<distractor    3,
surface-aligned; wrongness preserved>",
"edit_level": "<one of: 'none' | 'surface' |
'light' | 'fuller'>",
"edit_notes": "<one short sentence summarizing
what surface terms you swapped, or 'no change'>"
}
```

T.11 Tutor Response Generation (Immediate / Long-Term)

Immediate Response Prompt

```
==========================================
CATEGORY ID: {CATEGORY_ID}
CASE ID: {CASE_ID}
==========================================

────────────────
SCENARIO
────────────────

{MINIMAL_CONTEXT}
────────────────
DIALOGUE
────────────────

{DIALOGUE_TURNS}
────────────────
YOUR TURN - IMMEDIATE RESPONSE
────────────────

The student has just exhibited challenging
behavior.

What would you say or do RIGHT NOW as
the tutor in this moment?

Provide a direct, immediate response (2-3
sentences) in this moment.  Do NOT provide a
multi-step plan or describe what you will do
later.

Your immediate response (Output Response
ONLY):
```

Long-Term Intervention Prompt

```
==========================================
CATEGORY ID: {CATEGORY_ID}
CASE ID: {CASE_ID}
==========================================

────────────────
SCENARIO
────────────────

{MINIMAL_CONTEXT}
────────────────
DIALOGUE
────────────────

{DIALOGUE_TURNS}
────────────────
YOUR TURN - LONG-TERM INTERVENTION
────────────────

The student has just exhibited challenging
behavior.

What sustained interventions should be
implemented over time for this student?
Describe what long-term strategies and support
structures you would put in place (2-3
sentences).    This should be a narrative
description of ongoing interventions over weeks
or months, not immediate actions.

Your long-term intervention plan (Output
Intervention Plan ONLY):
```

T.12 Rubric Grading

Prompt

```
You are an expert in educational research,
specialized in evaluating tutor responses to
challenging student behaviors.

Score the tutor's response using the provided
personalized rubric.
Each criterion is a CHECKLIST item — score 1 if
the condition is met, 0 if not.
The rubric has a VARIABLE number of criteria
(not fixed at 5). Score every criterion present
in the rubric — do not skip, do not invent.

──────────
### SCORING INSTRUCTIONS
──────────

For each criterion in the rubric:
- Read the "check" condition carefully
- Score 1 if the tutor response satisfies the
condition, 0 if not
-  The check is principle-level:  diverse
approaches can earn the point as long as the
underlying principle is present
- Be strict but fair — surface mentions without
substance do NOT earn the point
```

```
- Do NOT reward length, politeness, or fluency

A tutor response does not need to satisfy every criterion to be considered acceptable; the rubric documents the full set of principles this student needs, and a single response can only address some of them. Score each criterion independently.

──────────
### RUBRIC (variable-length checklist)
──────────

{RUBRIC}

──────────
### OUTPUT FORMAT
──────────

Output ONLY valid JSON. The 'criteria_scores' array MUST contain exactly one entry per criterion in the rubric, in the same order, preserving each 'criterion_id'.

{
"criteria_scores": [
{
"criterion_id": <int from rubric>,
"criterion_name": "<copy from rubric>",
"evidence": "1-2 sentences: specific evidence from the tutor response that justifies this score",
"score": 0 or 1
}
],
"n_criteria": <number of criteria scored>,
"total_score": <sum of all criterion scores>,
"any_satisfied": <true if total_score > 0, else false>,
"overall_feedback": "2-3 sentences: what the tutor did well, what principle was most notably missing, and the single most important improvement for this specific student"
}
<System Prompt/User Prompt>

──────────
### STUDENT CONTEXT
──────────

{STUDENT_DESCRIPTION}
──────────
### DIALOGUE (situation the tutor is responding to)
──────────

{DIALOGUE}
FULL CONTEXT (JSON with fields: scenario_id, title, content → task, context, behavior_patterns, challenging_behavior_trigger):
Use 'content.task' to understand what the student was asked to do, 'content.context' for the setting/environment/materials, 'content.behavior_patterns' for the student's documented tendencies, and 'content.challenging_behavior_trigger' for what the final student turn is showing.
{FULL_CONTEXT}

──────────
### TUTOR RESPONSE
──────────

{TUTOR_RESPONSE}
```

T.13 MCQ Response

Prompt

```
You are an expert special education evaluator assessing tutor responses to students with learning difficulties.

<System Prompt/User Prompt>

Below is a tutoring dialogue between a tutor and a student with learning difficulties, followed by a multiple-choice question about the most appropriate intervention.
## Dialogue
{DIALOGUE}

## Question
Based on the dialogue above, which intervention is most appropriate for this student in this situation?

{OPTIONS}
Respond in JSON format:
{ "answer": "<option letter>", "reasoning": "<brief explanation of why this intervention is most appropriate for this specific student>" }
```

T.14 MCQ Pair Selection

Prompt

```
You are selecting the single best (effective, ineffective) instruction pair to use as the two answer options for a multiple-choice tutoring question.

You are given:
- DIALOGUE: the tutor-student exchange leading up to the moment the tutor must respond.
- STUDENT_CONTEXT: persona-level background on the student.
- RESPONSE_TYPE: either "immediate" (a single in-the-moment reply) or "long_term" (a plan over days or weeks).
- A list of EFFECTIVE candidate instructions (each is a tutor action that has been validated as effective FOR THIS STUDENT).
- A list of INEFFECTIVE candidate instructions (each is a poor-practice / counter-example from the same case).

Your job is to pick ONE effective and ONE ineffective so that the resulting question has matched granularity and is plausibly anchored in the dialogue moment.

──────────
SELECTION CRITERIA
──────────
```

```
Apply the following criteria, in priority order:

1. SCOPE MATCHES RESPONSE_TYPE.
- If RESPONSE_TYPE = “immediate”: both options must read as something a tutor could do *in this moment* (single reply, on-the-spot redirect, in-the-moment instructional move). Avoid candidates that describe long-term policy, curriculum design, multi-week plans, or persistent classroom culture.
- If RESPONSE_TYPE = “long_term”: both options must read as a plan, policy, or sustained intervention over days/weeks. Avoid candidates that are clearly a one-off reaction to a single utterance.

2. ACTION FORM, NOT STATE OR DIAGNOSIS.
- Strongly prefer candidates phrased as a tutor action (“Do X”, “Provide Y”, “Guide the student to Z”).
- Avoid candidates that read as a description of a *state* (“very low rates of approval”), a *trigger condition* (“when X happens”), a *paper finding* (“teachers in this study often...”), or a *student-side description*.

3. PLAUSIBILITY AT THE DIALOGUE MOMENT.
- Prefer candidates a tutor would actually consider given what just happened in the dialogue. The ineffective option should still be substantively poor for THIS student, but should be the kind of misstep a real tutor might make in this exact situation — not an unrelated topic.

4. GRANULARITY PARITY between the two chosen options.
- The two selected options should be at a comparable level of specificity. Avoid pairing a very long, multi-clause effective with a one-line abstract ineffective, or vice versa.
- Comparable length and structure helps; identical length is not required.

5. TIE-BREAKER: when multiple pairs satisfy 1–4 roughly equally, prefer the pair whose EFFECTIVE option is most directly responsive to the dialogue moment.

——————————
HARD CONSTRAINTS
——————————

- You MUST select indices that exist in the provided lists. Do NOT invent new instructions.
- You MUST select exactly one effective and one ineffective.
- Do NOT swap roles. Select your effective_idx from the EFFECTIVE list and your ineffective_idx from the INEFFECTIVE list.
- If no pair is ideal, still pick the best available pair — selection is mandatory, not optional.

——————————
INPUT
——————————

DIALOGUE:
{DIALOGUE}

STUDENT PERSONA:
{STUDENT_PERSONA}

SCENARIO CONTENT (task, environment, trigger, behavior patterns):
{SCENARIO_CONTENT}

RESPONSE TYPE: {RESPONSE_TYPE}

EFFECTIVE CANDIDATES (0-indexed):
{EFFECTIVE_CANDIDATES}

INEFFECTIVE CANDIDATES (0-indexed):
{INEFFECTIVE_CANDIDATES}

——————————
OUTPUT
——————————

Output ONLY valid JSON.

{
“effective_idx”: <integer index into EFFECTIVE_CANDIDATES>,
“ineffective_idx”: <integer index into INEFFECTIVE_CANDIDATES>,
“reason”: “<one short sentence: why this pair beats the alternatives on the criteria above>”
}
```

T.15 MCQ Pair Adaptation

Prompt

```
You are calibrating two answer options for a multiple-choice tutoring question.

You are given a dialogue and TWO source strategies from a case study:
- CORRECT: an effective intervention the tutor could take
- INEFFECTIVE: a counter-example from the same case (poor practice, missed opportunity, or known-failed strategy)

PROBLEM: The two strategies, as written, are often stylistically uneven. The correct one is usually phrased as a concrete tutor action, while the ineffective one is sometimes phrased as a trigger description, a passive state, a description of what someone failed to do, or a paper-specific finding. This creates an unintended style cue that gives the answer away.

YOUR JOB — make the two options PARALLEL in form and length. Use the smallest change that achieves this. Most options need only minor edits.

PREFERENCE HIERARCHY (apply in this order, stop as soon as one is enough):
1. No edit — if both INSTRUCTIONs are already concrete tutor actions of similar length and form, output them VERBATIM.
```

```
2. Surface tweak only — flip "tutor backs down" → "Back down from the request"; convert a state to an action verb form. Reuse the original words wherever possible.
3. Light rewrite — only if the original is genuinely not interpretable as a tutor action (e.g., a trigger condition with no agent), reframe it as the simplest matching action. Stay close to the original wording.
4. Fuller rewrite — AVOID. Use only as last resort if option (3) is not possible.

TARGET OUTCOME:
(1) Each option reads as a concrete action the tutor could plausibly take in this moment / over time.
(2) Both have roughly similar length and structural form.
(3) The INEFFECTIVE option remains clearly ineffective for THIS student — preserve its substantive content (the underlying poor approach), just reframed as something the tutor would do.

HARD CONSTRAINTS:
- DEFAULT to keeping the original wording. If in doubt, change less.
- Do NOT change which option is effective. The correct stays effective; the ineffective stays ineffective.
- Do NOT swap meaning, invent new content, or paraphrase for the sake of paraphrasing.
- Do NOT make the ineffective option subtly "look effective". Its core flaw must still be visible.
- Do NOT strip paper-specific names that name a real tutor action (e.g., "Boardmaker symbols", "Scratch program"); only strip jargon if it would be unintelligible without the paper.
- Do NOT over-generalize. Keep each at the original level of specificity.
- If you change anything, change as few words as possible.
- The two options MUST end up roughly similar in length. If the originals differ a lot, trim the longer side's decorative phrases (parenthetical examples, trailing "so that..." clauses) to bring them closer.
- When balancing length, ONLY trim the longer side. NEVER expand the shorter side — do not add words, clauses, or qualifiers that are not in the original. If the shorter option ends up much shorter than the longer one even after trimming, leave it short rather than invent content.

──────────
INPUT
──────────

DIALOGUE:
{DIALOGUE}

STUDENT CONTEXT:
{STUDENT_CONTEXT}

RESPONSE TYPE: {RESPONSE_TYPE} (immediate = a single in-the-moment reply / long_term = a plan over days-weeks)

CORRECT (effective for this student):
INSTRUCTION: {CORRECT_INSTRUCTION}
REASON: {CORRECT_REASON}
OUTCOME: {CORRECT_OUTCOME}

INEFFECTIVE (counter-example from same case):
INSTRUCTION: {INEFFECTIVE_INSTRUCTION}
REASON: {INEFFECTIVE_REASON}
OUTCOME: {INEFFECTIVE_OUTCOME}

──────────
OUTPUT
──────────

Output ONLY valid JSON.

{
"correct_adapted": "<correct option, edited as little as possible>",
"ineffective_adapted": "<ineffective option, edited as little as possible, still substantively poor>",
"edit_level": "<one of: 'none' | 'surface' | 'light' | 'fuller'> // strongest level applied across the two options",
"edit_notes": "<one short sentence describing what you changed, or 'no change' if edit_level=none>"
}
```

T.16 Reference-Based Response Generation

Prompt

```
──────────
REFERENCE — strategies from similar student cases
──────────
The strategies below are documented to have worked or failed for similar but different students. Treat them as reference only — they are not a script and were not written for the current student. Adapt them, combine them, or ignore them as appropriate. The cases are ordered by relevance and may conflict with each other; use your own judgment.

{REFERENCE}
──────────

{PROMPT}
```

T.17 Dialogue Synthesis Prompt with Variable Turns

Prompt

```
You are a research assistant specialized in creating realistic dialogue from educational research scenarios.

Your goal is to CREATE REALISTIC STUDENT-TUTOR DIALOGUES that demonstrate CHALLENGING MOMENTS requiring intervention.
```

```
──────────
### CRITICAL RULES
──────────
1. Ground ALL statements in provided profile/excerpts - NO INFERENCE
2. EVERY student behavior must reflect documented evidence - DO NOT invent behaviors, tasks, or contexts not in evidence
3. Reference evidence by ID for EVERY message
- evidence_basis arrays may contain ONLY IDs from ALLOWED evidence_basis IDS below; inventing an ID, or citing one not on that list, invalidates the whole output
- NEVER use meta labels such as "category_extraction", "scenario", "persona", "paper", "prompt", or invented IDs
4. Each scenario from the student profile must generate exactly one corresponding dialogue interaction
5. Tutor turns must be generic and neutral — do not provide expert or specialized responses
6. The dialogue sets up a situation that REQUIRES good intervention
7. Final student turn shows challenging behavior that needs response (CHALLENGING BEHAVIORS = student behaviors that are difficult for a tutor to address using standard instructional strategies)
- The final turn MUST be a student turn and MUST include "is_challenging_behavior": true
- Every student turn MUST include a non-empty evidence_basis array
8. Output ONLY valid JSON - NO explanations, NO markdown

──────────
### TASK OVERVIEW
──────────

**STEP 1: For EACH Scenario - Plan Dialogue**
Review scenario from profile:
- Task: What specific problem/activity? (concrete from evidence)
- Context: Environment/materials/social setting (verify action possibilities)
- Student persona: Core characteristics, typical responses
- Difficulty trigger: What challenging behavior will end the dialogue? What causes final struggle?
- Available actions: What tutor/student can physically do?

**STEP 2: For EACH Scenario - Structure Dialogue**
- **Number of turns = EXACTLY {TOTAL_TURNS} messages** (= 2 × {TURN} exchanges). This is a hard
requirement: emit {TOTAL_TURNS} objects in the "dialogue" array, no more and no fewer, alternating tutor → student and numbering turn_id from 1 to {TOTAL_TURNS}, so the first message is the tutor's and the last is the student's.
- **First Turn (Tutor)**: Present SPECIFIC, CONCRETE task
- State the specific task from the scenario context (neutral tone)
- Generic and neutral only — no expert or specialized responses
- **Turns in the middle (Mixed)**: Student attempts → Tutor NEUTRAL response
- Student: May show emerging problem behaviors (based on evidence, follow persona, use quotes when appropriate)
- Tutor: Generic/neutral response only (e.g., "What do you think?", "Keep trying") - DO NOT use effective strategies
- Physical actions in parentheses ONLY if grounded in evidence and context allows
- **Last Turn (Student)**: MUST show challenging/struggling behavior
- Observable problem behaviors matching challenging_behavior_trigger
- Grounded in documented behavior_patterns

**STEP 3: For EACH Scenario - Generate Dialogue**
Follow plan, use student quotes from evidence, maintain persona (do NOT assume active engagement or verbal expressiveness unless documented)

**STEP 4: For EACH Scenario - Write Minimal Context**
Using the provided SCENARIO, write one natural sentence describing what the student is doing and what the goal is.
- Include: task + setting
- Exclude: student characteristics, struggles, or solutions


──────────
### OUTPUT STRUCTURE
──────────
{
"category_id": "{CATEGORY_ID}",
"paper_id": "{ID}",
"name": "student category name",

"dialogues": [
{
"case_id": "{CATEGORY_ID}_case_001",

"dialogue": [
{
"turn_id": 1,
"speaker": "tutor",
"message": "Clear task presentation. Neutral tone, no effective strategies"
},
{
"turn_id": 2,
"speaker": "student",
"message": "Response based on persona",
"evidence_basis": ["0001_learner_002"]
},
... (alternate tutor/student until turn_id {PENULTIMATE}) ...
{
"turn_id": {TOTAL_TURNS},
"speaker": "student",
"message": "CHALLENGING BEHAVIOR matching challenging_behavior_trigger",
"is_challenging_behavior": true,
"evidence_basis": ["0001_learner_008"]
}
```

```
],

"minimal_context": "One natural sentence: what
student is doing + goal.  No characteristics,
struggles, or solutions."
}
]
}


<System Prompt/User Prompt>

————————————
### INPUT CONTEXT
————————————
STUDENT CATEGORY:
{CATEGORY}

STUDENT PERSONA:
{STUDENT_PERSONA}

SCENARIO:
{SCENARIO}

EVIDENCE  (excerpt_id  →  actual  text  from
paper):
{EVIDENCE}

ALLOWED  evidence_basis  IDS  —  the  complete
list;  any  other  ID  is  invalid,  and  IDs  from
elsewhere in the paper must NOT be cited even
if they seem relevant. Reuse these across turns
as needed:
{ALLOWED_EVIDENCE_IDS}
```