arXiv:2610.34894v1 [cs.CL] 26 Oct 2026

What Outcome RL Buys: Finishing, Reliability, Rarely Reach

Shan Yang [1]

Abstract

Outcome RL raises pass@1 but not pass@k at large k (Yue et al., 2025). But every pass@1 is scored under a token budget. Every recipe also trains under one. On physics and math problems at fixed token budgets, we compare the pass@1 of 20 released RL models (1.5B to 32B) with that of the thinking models they started from. We ask whether each pass@1 gain is changed behaviour or new knowledge. We divide each change in pass@1 into three parts. *Finishing* measures answering before the budget runs out. *Reliability* measures being right once finished. *Reach* measures problems RL solves that its base answered wrongly in all eight tries, minus problems the base solves that RL answered wrongly in all eight. We find four things. (1) Where the base often hits the token budget before answering, most of the change, gain or loss, is finishing (82% at 4k and 8k). Given more tokens, the base closes most of the finishing gap itself. (2) Where the base has enough tokens, the gain is reliability: the RL model is graded right more often, even on problems both finish. (3) By our definition, RL rarely solves what its base cannot. Reach adds at most 2.1 percentage points on any of our four physics and math sets. We checked 144 physics problems that only the RL model solved. Of these, 3.5% hold up when resampled or read. In a random sample, nine of ten new RL steps a human confirmed also appear in the base's samples. (4) In a small-scale RL run on an 8B model, the gain on new problems is finishing, with reliability flat. Supplying each problem's physical setup was not shown to improve on plain RL.

Project page: https://shanyang.me/what-outcome-rl-buys/. Code and released tables: https://github.com/shanyang-me/outcome-rl-buys.

[1]Individual Researcher. Correspondence to: Shan Yang <alexyangshan@gmail.com>.

Preprint. October 27, 2026.

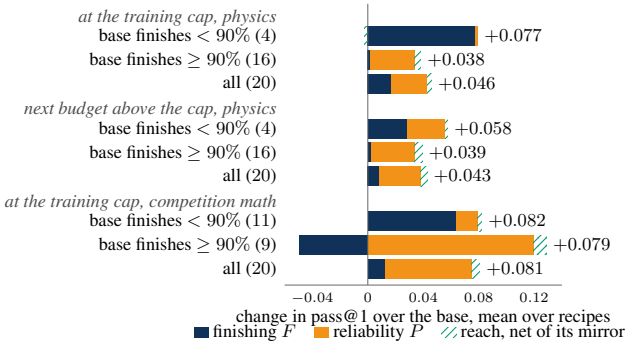

Figure 1. Outcome RL's pass@1 gain on physics and competition math, split into finishing, reliability and net reach. Each bar is the mean over the recipes in its group. Recipes are grouped by how often their base finishes at the training cap on that set, with their number in parentheses. Reliability is measured on each model's own finished derivations. A part left of zero lowers the score: the RL model hits the budget before answering more often (finishing), is right less often once finished (reliability), or solves fewer problems alone than the base does (reach).

1. Introduction

Outcome RL, which rewards only a correct final answer, is a standard way to train reasoning models. Its gains are reported as pass@1. Whether such a gain is new knowledge or changed behaviour matters. New knowledge widens what a model can solve. Changed behaviour only helps it use what it already knows. The two are hard to tell apart. Every pass@1 is scored under a token budget, so one number mixes answering in time, answering right and solving new problems. Yue et al. (2025) found that outcome RL adds no reasoning the base cannot already produce. Pass@k is the chance that at least one of k samples is right. Once k reaches the hundreds, the base scores higher than the RL model. The problems the RL model solves are nearly a subset of those its base solves. Prior work calls this narrowing *sharpening* (Huang et al., 2024; Qin et al., 2025; Wu et al., 2025c). We ask whether RL changes behaviour or knowledge at $k = 1$, the score every report rests on. Finishing in time and being right more often change behaviour. Solving problems the base cannot points to new knowledge. We count a gain as new knowledge only when the base finishes all eight samples on the problem and gets every one wrong. We count any other gain as changed behaviour, since the base either answers the problem right at least once or does not

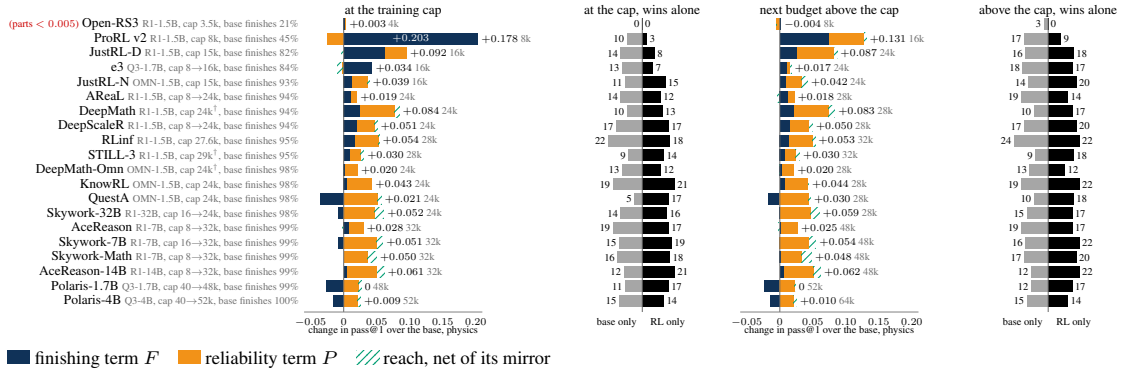

Figure 2. What outcome RL buys at each recipe's own training cap and at the next budget above it, in pass@1 on physics. Each row is one recipe, split as in Figure 1. Parts left of zero read as in Figure 1. The number at the right of each row is the change, with the budget read. Red *parts* < 0.005 marks a row whose parts are all below 0.005, too small to read. Rows are sorted by how often the base finishes at the cap on physics, least first. The lone-win panels count physics problems that only the RL model solves (black, right) or only the base (grey, left), while the other model finished all eight samples. †: the cap is not exact. DeepMath trains at a 24k cap but penalizes length from 20k. STILL-3 trains at 29k and is read at 28k.

finish every sample.

Our first question asks of released models: how much of a pass@1 gain is finishing within the token budget, and is the rest reliability or reach? Every recipe is also trained under a budget. If part of a gain is the model fitting its answer into the budget, comparisons between recipes, or between a recipe and its base, partly measure how long each model writes rather than how well it reasons. We ask the first question of thinking models on physics and math. A bounded second question follows: can training install knowledge the model does not use on its own, when it is supplied?

Our main study answers the first question by reading twenty released RL models against their bases at fixed budgets. All twenty start from thinking bases, which distillation has already widened. Yue et al. (2025) leave thinking bases open. Once the budget suffices, reliability rises on physics wherever the base already finishes. Reach stays small.

The second question tests the obvious remedy for sharpening: hand the model the knowledge its derivations lack (a hint, a teacher's derivation, or a question about the setup). We train an 8B thinking model with outcome RL on physics, supplying in five ways the physical setup of problems its base almost never solves. At this small scale, the supplied setup changes the model during training but is not shown to carry over to held-out problems beyond plain RL. Reward raises a problem's solve rate only once the model samples a success on that problem (§5). Released recipes that use the remedy, QuestA and KnowRL, reach only a little further. Their net reach is at most 0.018 (§4.3).

Contributions. (1) *A budget-aware split of the pass@1 gain* into finishing, reliability and reach (Eqs. (1)–(5)). Unlike pass@k, it separates answering in time from answering right. The split also nets reach against the reverse. (2) *A read of twenty released pairs on thinking bases*, the case Yue et al. (2025) leave open, at each recipe's own training cap and the next budget above it (§4). (3) *A comparison rule*: compare a recipe with its base at or above its training cap. (4) *A controlled probe* that supplies an 8B model, in five ways, the setup of physics problems its base almost never solves, and tests transfer to their text-only twins (§5).

2. Related Work

What RL with verifiable rewards changes. Prior work disagrees on what outcome RL changes. None splits the measured gain into finishing and the rest. On pass@k curves, RL raises pass@1, but the base wins at large k (Yue et al., 2025; Wu et al., 2025b). RL stays within what the base can sample. Yue et al. (2025) count the problems only the RL model solves at 0–1% on AIME24 and MATH500, at k up to 1,024. Follow-ups trace this lost coverage to hard prompts whose rare successes drop out of small sample groups (Zhou, 2026) and to winner-take-all reinforcement of what the base already finds likely (Nguyen et al., 2025). Others trace part of reported gains to budget mismatch and attempt inflation (Wu et al., 2025a; Yang & Chen, 2026). Studies of derivations say RL narrows the policy where fine-tuning broadens it (Matsutani et al., 2025), acts on a few tokens where derivations branch (Wang et al., 2025b), and amplifies habits the base already has (Gandhi et al., 2025). Power sampling, entropy minimisation and even spurious rewards reproduce much of RL's gain (Karan & Du, 2025; Agarwal et al., 2025; Shao et al., 2025). One training example recovers most of a full data set's gain (Wang et al., 2025d). A second RL stage shortens derivations at no cost in accuracy (Fatemi et al., 2025). Both results say RL moves behaviour more than content. Prolonged training, by

contrast, is reported to expand the solvable set (Liu et al., 2025a). Decompositions find RL improves execution within a fixed coverage (Qin et al., 2025). Fine-tuning, not RL, moves knowledge (Wu et al., 2025c). All this evidence is observational: no work supplies the missing content and tests whether RL learns it. Our reliability term is the execution gain of Qin et al. (2025) with finishing removed. We compare reliability with finishing at each budget. Our reach is a term of an exact split of the pass@1 gain, counted per problem at $k = 8$. The base-only count in reach controls for sampling noise. Problems only one model solves are resampled and read. Unlike budget corrections of reported gains (Wu et al., 2025a), we split the gain at a fixed budget. We also read each recipe at its own cap.

Objectives, samplers and replay. Other work changes how RL credits rare successes. We use one such method, replay. Pass@k objectives reshape the group advantage (Chen et al., 2025d; Walder & Karkhanis, 2025; Liu et al., 2025b). VinePPO branches continuations to estimate credit (Kazemnejad et al., 2024). RLEP replays past successes into the group (Zhang et al., 2025a). Replay shows what reward teaches once a success is sampled.

Measuring RL on thinking models. Length and answer rate are known confounds (Liu et al., 2025b; Yu et al., 2025b; Aggarwal & Welleck, 2025; Hou et al., 2025). For a fixed thinking model, accuracy at a cap rises with the cap (Muennighoff et al., 2025; Qwen Team, 2025). Any change to derivation length therefore changes the score. We separate the two with the finishing decomposition. We also fix a held-out set before training and measure its noise floor.

3. Experimental Setup

This section gives the study's scope, the three metrics that split pass@1 under a budget (§3.1) and the protocol for reading the twenty released pairs (§3.2).

Scope: thinking bases. We study thinking bases because they give RL the most to change: how long the model reasons, when it commits to an answer and how often that answer is right, not only whether it answers. A token budget makes each change visible in the score. Practice uses thinking bases too. The strongest open models on math and science, DeepSeek-R1 and its distills and Qwen3, ship in this thinking mode. Every recipe we study starts from one of them. All models studied are thinking models: the five bases (DeepSeek-R1-Distill-Qwen 1.5B and 7B, Qwen3-1.7B and 4B in thinking mode, OpenMath-Nemotron-1.5B), the 14B and 32B bases added later, and our controlled run's Qwen3-VL-8B-Thinking. Models that answer in a few hundred tokens are out of scope, except one pair read for contrast: Qwen2.5-7B and its outcome-RL version, Open-Reasoner-Zero-7B (§4.1).

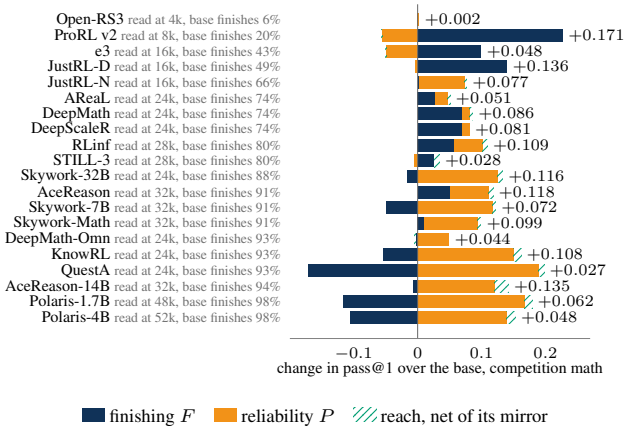

Figure 3. What outcome RL buys at each recipe's own training cap on competition math, the recipes' own domain. Rows are split as in Figure 2. Parts left of zero read as in Figure 1. Rows are sorted by how often the base finishes at the cap, least first.

Scope: physics and math. To tell changed behaviour from new knowledge, we read every pair in two domains. Math is the recipes' own domain: all twenty recipes we study train on competition math and report math results. A gain on math shows what RL buys where it was trained. A gain on physics shows what RL buys outside that domain. Physics is at most about 1% of the RL training data of every recipe that states its mix, except ProRL v2, whose mix includes STEM. Skywork-OR1-Math-7B does not state its mix. Physics problems also need modelling, which math problems rarely exercise. A physics derivation first picks the governing principle and decides which quantities and effects matter, and only then does the algebra. A wrong setup is not repaired by careful algebra later. Once they finish at 16k, the five bases are right on 85–91% of their MATH-500 derivations but only 21–56% of their physics ones. The modelling also makes physics derivations long, so the token budget binds more often than on math (§4.1).

Test sets. Each domain has a main set and a check. The main physics set is 300 text-only problems from the training corpus of our controlled run (§5), so both questions share one audited answer checker. The set keeps only problems the controlled run's base solves at least sometimes, so pass@1 can move. Filtering on the base caps how much reach can show, so 671 unfiltered physics problems check the reach result (Appendix A.6). On math, the main set is MATH-500, the recipes' standard benchmark. The math check is 276 recent contest problems, below MATH-500's ceiling, some released after every model (Appendix A.7). We call it competition math.

3.1. Metrics

We split a pass@k gain under a token budget B into three parts: *finishing* F, answering in time; *reliability* P, answer-

ing right once finished; and *reach*, reach@$k(B)$, solving problems the base leaves unsolved in k samples. A *finished* derivation stops within the budget. With k samples per problem at budget B, pass@$k(B)$ is the share of problems with at least one correct finished sample. The finishing rate fin@$k(B)$ is the share with at least one finished sample. The reliability rate pass@$k(B \mid \text{fin})$ is the share of problems with a finished sample that also have a correct one, so pass@$k(B)$ factors as

$$\text{pass@}k(B) = \text{fin@}k(B) \times \text{pass@}k(B \mid \text{fin}), \quad (1)$$

and a change splits exactly into a finishing term F and a reliability term P (reach is split off below):

$$F = \Delta\text{fin@}k(B) \times \text{pass@}k_{\text{RL}}(B \mid \text{fin}), \quad (2)$$

$$P = \text{fin@}k_{\text{base}}(B) \times \Delta\text{pass@}k(B \mid \text{fin}). \quad (3)$$

Here F takes the cross term, $\Delta\text{fin@}k(B) \times \Delta\text{pass@}k(B \mid \text{fin})$. Giving it to P, or half to each, gives the split's other two orders.

Reach, reach@$k(B)$, asks the question of Yue et al. (2025) at a given k: how much does RL score on problems its base leaves unsolved in k samples? A model leaves a problem *unsolved* when it finishes all k samples and none is correct. Requiring all k to finish rules out problems lost only to the budget. Out of n problems, let U be those the base leaves unsolved and U' those the RL model leaves unsolved. With k samples, either model can miss a problem it solves with low probability, so winning a problem alone can be chance. Reach therefore subtracts the *mirror*: what the base scores on the problems the RL model leaves unsolved.

$$\text{reach@}k(B) = \frac{1}{n}\Big(\sum_{q \in U} \text{pass@}1_{\text{RL}}(q; B) - \sum_{q \in U'} \text{pass@}1_{\text{base}}(q; B)\Big). \quad (4)$$

k only decides which problems count as unsolved. Each model is scored on them by pass@1, so reach is in pass@1 units. F and P are computed at $k = 1$ on the problems outside U, weighted by their share of the n. Reach subtracts the mirror and P adds it back, so nothing counts twice. The parts sum exactly: $\Delta\text{pass@}1(B) = F + P + \text{reach@}k(B)$. *Robust* reach ignores problems solved once by chance. Robust reach keeps only unsolved problems the other model solves in at least $k/2$ of its k samples, the sets $U_{\geq k/2} \subseteq U$ and $U'_{\geq k/2} \subseteq U'$:

$$\text{reach}_{\geq k/2}\text{@}k(B) = \frac{1}{n}\Big(\sum_{q \in U_{\geq k/2}} \text{pass@}1_{\text{RL}}(q; B) - \sum_{q \in U'_{\geq k/2}} \text{pass@}1_{\text{base}}(q; B)\Big). \quad (5)$$

base derivations that finish, physics / MATH-500	
at 4k and 8k, range over bases	12–65% / 56–90%
at 16k	82–93% / 92–97%
finishing's share of the change, 4k and 8k	
both sets	82% [77, 86]
physics / MATH-500	71% / 92%
RL's change in length vs. in finishing	$\rho = -0.93$
Open-Reasoner-Zero-7B, finishing at 2k	$0.98 \to 0.65$
base score, tight budget $\to 64k$	
problems picked at 4k (RL 0.68)	$0.44 \to 0.80$
problems picked at 8k (RL 0.65)	$0.38 \to 0.59$

Table 1. Finishing at tight budgets. Arrows read base to RL, except in the last block. Share $= \sum|F|/(\sum|F| + \sum|P|)$, with a 95% interval. ρ: Spearman's correlation over recipes.

Matched reliability is reliability measured only on problems both models finish. For the released pairs, both must finish at least seven of eight samples.

3.2. Protocol for Reading the Released Pairs

We read twenty released outcome-RL models, each against the thinking base it was trained from (Table 6). Eighteen are of up to 7B. The other two are a 14B and a 32B pair. Each RL model and its base are sampled the same way, with one prompt and one answer checker, the matcher. Each model writes $k = 8$ answers per problem at temperature $T = 0.6$. Budgets: 4k, 8k and 16k; the recipe's training cap; and the next budget above the cap on a fixed grid up to 64k. The split is read at $k = 1$, per sample. fin@$1(B)$ is the share of derivations that finish within B. pass@$1(B \mid \text{fin})$ is pass@1 on finished derivations. Reach is read at $k = 8$, as reach@8 and reach$_{\geq 4}$@8 ($k = 32$ in Appendix A.5). We *registered* predictions for ten pairs at 4k and 8k, committing them before any of their outputs existed. Reads at the cap and the next budget above it were not registered.

4. Finishing, Reliability, Rarely Reach

4.1. Finishing

When the base often hits the budget before answering, finishing carries most of the change. At 4k and 8k the bases leave many derivations unfinished. Finishing is then the larger part of the change in every recipe up to 7B except Open-RS3, whose score moves by at most 0.004 at these budgets (Figure 2). The reward explains why. An unfinished derivation scores zero whatever it contains, so at a tight budget the first thing group-relative credit can tell apart is whether a sample finished. The finishing share is larger on MATH-500 than on physics, because a finished base derivation on MATH-500 is already usually right. Figure 4 shows each cell's change tracking its finishing term.

Outcome RL moves derivation length toward what earns most reward within its cap. The effect on finishing there-

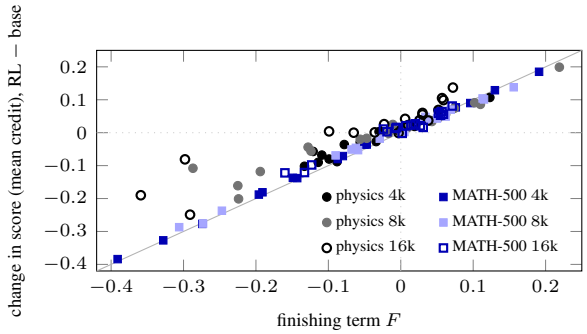

Figure 4. Change in score against the finishing term F, with Spearman $\rho = 0.97$ over all 108 points. Each point is one pair on one set at one budget. A point on the diagonal is a change carried by finishing alone. A point's vertical distance from the diagonal is the rest of the change: reliability, reach and partial credit. The correlation is 0.99 at 4k and 8k and 0.92 at 16k.

fore depends on where the base starts. At a short cap, RL teaches a thinking base to fit its derivations inside the cap, as ProRL v2 does at 8k. At a long cap, RL teaches longer derivations. The RL model then finishes less often, as Polaris-1.7B does even at its own cap. Across recipes, the change in derivation length tracks the change in finishing (Table 1). RL lengthens derivations on a short-answer base too, and the RL model then finishes less at a tight budget (Open-Reasoner-Zero-7B (Hu et al., 2025); Table 1).

Finishing is not hidden knowledge. A base that cannot finish within B tokens might not know how to derive the answer. Two results argue against this reading. First, the extra tokens go to hesitation, not derivation. In our controlled run (§5), RL halves the backtracking words ("wait", "actually") per derivation. Reliability on problems finished both before and after training does not move. Second, given more tokens, the base solves most of the problems on which RL, at a tight budget, finishes more often and scores higher. We pick those problems with four of the eight samples, then score them on the other four and on a separate 64k base run. At 64k the base passes the RL model on the problems picked at 4k and closes most of the gap on those picked at 8k (Table 1). The gap the base does not close counts as reliability, not finishing. Finishing and knowledge are not fully separable, since committing to the right approach is a skill. Our claim is narrower: RL adds few steps the base does not already write (§4.3). Finishing still pays at a tight deployment budget, as ProRL v2 at 8k shows.

4.2. Reliability, Where the Base Already Finishes

Whether the base still hits the training cap before answering decides which part grows. RL can buy finishing only where the base fails to finish (Figure 2). A recipe is *cap-bound* when its base still often hits its cap before answering. Four recipes are cap-bound: ProRL v2, Open-RS3, e3 and JustRL-DeepSeek. For ProRL v2 and e3 on physics, finishing is the whole gain (Table 7). JustRL-DeepSeek gains about a third by reliability, half under the matched split. Open-RS3's score does not measurably change at its cap. The split of Eqs. (2)–(3) can hide a reliability gain, because the problems RL newly finishes are harder and pull its reliability down. *Matched* reliability, measured only on problems both models finish, avoids that pull (§3.1). With matched reliability, the four recipes gain some reliability as well. Across all twenty recipes, the more the base leaves unfinished at the cap, the larger the finishing share.

physics at the cap, mean per recipe	F	P
four cap-bound recipes	$+0.078$	$+0.002$
matched split	$+0.058$	$+0.021$
Open-RS3, total change*	$+0.003$	
sixteen recipes that are not cap-bound	$+0.002$	$+0.032$
all twenty recipes	$+0.017$	$+0.026$
competition math, matched split		
base finishes $< 90\%$	$+0.045$	$+0.035$
base finishes $\geq 90\%$	-0.011	$+0.081$
sixteen recipes that are not cap-bound		
reliability rises	16 of 16	
reliability is larger by size	14 of 16	
matched reliability rises (clear of 0)	16 (12) of 16	
bases with $P > F$ / larger by size	7 / 6 of 7	
all twenty recipes		
mean reach	$+0.003$	
finishing share vs. unfinished at the cap	$\rho = 0.59$	
reliability rises at the next budget	18 of 20	
finishing rises at the next budget	11 of 20	
step check: RL's decisive steps the base lacks		
base right 1–4 of 8, 16k	0 of 26	
base right 5–7, 13 caps	0 of 26	
base unfinished, never right, 16 caps	1 of 57	
reverse, base's steps RL lacks, 16k	0 of 19	

Table 2. Finishing and reliability by recipe group, and the step check, on physics at the cap unless noted. ρ: Spearman's correlation over recipes. At the next budget, a rise is a gain above 0.005. *With 95% interval $[-0.007, +0.014]$.

Where the base already finishes, reliability rises in every recipe. The sixteen recipes that are not cap-bound train at caps where the base already finishes 93–100% of its derivations. In all sixteen, reliability rises on physics. By size, reliability is the larger part in fourteen of them (Figure 2). Averaged by base, reliability exceeds finishing on all seven bases, and is larger by size on six (Table 2). Finishing moves either way: seven of the sixteen finish less than their base. Math shows the same rise, on MATH-500 and on competition math, a separate set of 276 recent contest problems (Figure 3). One recipe goes the other way: e3's reliability falls on MATH-500, although its base finishes there.

Reliability still rises on problems both models finish. Matched reliability still rises in all sixteen recipes, clear

of zero in twelve (Table 7). The Polaris pairs finish less than their bases, and their matched rise is smaller than their plain rise. On competition math, the matched split narrows the gap between the two parts in both groups of Figure 1, without changing which is larger (Table 2). A registered prediction that the rise would stay under 0.03 held with both benchmarks pooled but failed on physics alone (Appendix A).

On problems where RL is right more often, the step behind its win is one the base already writes. To test whether RL's reliability gain adds a step the base lacks, we drew 40 problems at random where the base finishes all eight samples at 16k but is right in only one to four, and the RL model is right more often. A language-model reader, blind to which model wrote which, named the step that makes the RL derivation right and the base derivation wrong. A second reader looked for that step in the base's other samples. On 12 of the 40 problems, the two derivations differ only in answer form or grading. None of the 26 steps judged decisive on the other 28 is missing from the base (Table 2). The base's own wrong samples state the step too, but on 17 of the 26 they do not carry it through. The reverse search, for the base's steps in the RL model's samples, finds no step missing either. At the caps of 13 of the 16 recipes that are not cap-bound, the check finds no RL step missing where the base is right five to seven times. RL uses, more reliably, a step its base already has.

Across all twenty recipes, reliability is the part of the gain that rises consistently. At the cap, reliability is the largest part of the mean physics gain (Table 2). Across recipes, reliability and reach are clear of zero. Finishing is not, because it moves either way. Resampling the seven bases instead of problems widens the intervals of the two shares until they overlap (Appendix A). At the next budget above the cap, reliability rises in more recipes than finishing does (Table 2). Our reads cannot explain why some recipes gain more reliability than others. The recipes differ in data, steps and algorithm at once.

4.3. Rarely Reach

Of what RL solves on physics, its base already solves 85%. Problems only one model solves, *lone wins*, are about as common in each direction (Table 3). In two thirds of them, the losing model left at least one sample unfinished. The rest, where the losing model finished all eight and got all eight wrong, are the candidates for new knowledge.

A lone win can be chance. Eight samples can miss a problem a model solves with low probability. Given 24 more samples, each model solves a similar share of the problems the other got wrong in all eight (Table 3).

physics	RL	base
solves what the other solves (16k)	83%	85%
lone wins, of 5,400 pair–problems	392	457
lone wins, other wrong in all 8	128	153
solves other's all-wrong, at cap	14%	11%
with 24 more samples	23%	27%
blind: a step the other lacks (143)	80	71
upheld by an author, of decided	10/26	9/19
found in the other's samples	49/60	52/62

Table 3. Lone wins and new steps, for the RL model and its base. Rows one to three are at 16k, rows four and five at the caps of the eighteen smaller recipes. Row one is the share of the other model's solved problems this model also solves. The last three rows read 143 derivation pairs at 16k (23 on math), the RL one right and the base one wrong. They count steps one has and the other lacks, those an author upheld in a random sample of 50 pairs, and, on physics, those found in four other samples of the model lacking them.

RL lone wins at 16k on four bases	
checked	121
solved by the base, 56 more samples	87
no longer lone wins when regenerated	22
differ only in answer form	9
real physics	3
RL lone wins at the caps: real physics	2 of 23
net reach	
harder physics set (671 problems)	−0.006 to 0.018
competition math (276 problems)	−0.002 to 0.021
MATH-500, 18 recipes	−0.003 to 0.005
physics, 1.5B to 4B recipes	−0.007 to +0.007
physics, 7B and up	+0.009 to +0.013
32B recipe, matcher / answer judge	0.013 / 0.004
base never right, some samples unfinished‡	
gain share, not cap-bound / cap-bound	10% / 34%
of RL's wins there, base solves at 64k	8% / 44%
net reach if unsolved, not cap-bound	≤ 0.019
ProRL v2's net reach if unsolved¶	0.067
robust wins§: QuestA/KnowRL/ProRL v2	1 / 3 / 3
their bases	0 / 4 / 0
training steps vs. robust reach	$\rho = 0.13$
32B at its cap: gain / reliability	+0.052 / +0.047

Table 4. Lone wins checked, and reach on other sets and at other sizes, physics at each recipe's cap unless noted. §With MATH-500. The 7B-and-up range covers four of five recipes. ρ: Spearman's correlation over the 16 recipes with stated steps. ‡The 16 regenerated pairs, three of them cap-bound (not e3). ¶JustRL-DeepSeek's would be 0.013 and Open-RS3's −0.001.

Of 144 lone wins we checked, 5 are real physics. A lone win is real physics if the base's answers are wrong in substance in every sample we have, not only in form such as an option letter. On four bases at 16k, the base solves 87 of RL's 121 lone wins in 56 more samples (Table 4). On the other 34, we read new samples of both models that keep the full text. Three are real physics. Reading 23 lone wins at the caps, without resampling, finds two more.

Reach stays at or below 0.021 on all four sets. The bound holds on our physics set, on a harder, unfiltered physics set, on MATH-500 and on competition math (Table 4).

A tenth to a third of the gain falls on problems the base never gets right but leaves some samples unfinished. Eq. (4) needs all eight base samples to finish, so it counts such a problem toward finishing and reliability. At the cap on the 16 regenerated pairs, these problems carry about a tenth of the gain for the recipes that are not cap-bound and about a third for the cap-bound ones (Table 4). Given 64k tokens, the bases solve under half of the ones RL gets right. Counted as unsolved, these problems would lift net reach above 0.02 only for ProRL v2 at its 8k cap (Table 4). Our definition counts this part as behaviour. On 56 of the 57 such problems checked where a step decides, the base's own samples write RL's deciding step (Appendix A.1).

In a random sample, nine of ten new RL steps a human upheld are already in the base's samples. On problems an RL model gets right and its base gets wrong, language-model readers compared the winning RL derivation with a failed base one. Read blind, each derivation has a step the other lacks about as often. In a random sample of 50 pairs, an author upheld under half of these claims on each side (Table 3). Nine of the ten upheld RL steps appear in the base's own samples. A reader then searched four other samples of the model whose derivation lacked the step. On physics, the search finds RL-side steps in the base's samples about as often as the reverse (Table 3).

Recipes built to add knowledge solve few new problems. QuestA puts a partial reference solution in every training prompt. KnowRL adds knowledge-point hints to about four prompts in five. ProRL v2 trains 3,000 steps on data that includes physics. At its cap, each of the three robustly solves at most three problems its base gets wrong in all eight (Table 4). On the harder, unfiltered physics set, QuestA and KnowRL lead the twenty recipes in net reach (Table 10). Read one by one, 13 of their 19 robust wins there stand, and none is answer form (Appendix A.6). KnowRL trains on QuestA's problems, so the lead may come from those problems rather than the hints. Training length does not predict reach (Table 4).

Larger bases split the same way, and four of the five show more reach. The three 7B recipes split like the smaller ones. The 14B and 32B recipes train at long caps, so at 4k and 8k they finish less often than their bases and score lower. At its cap, the 32B recipe gains mostly reliability (Table 4). For four of the five recipes from 7B up, net reach at the cap is higher than for any smaller recipe. Part of this extra reach is answer form, which the matcher rejects more often for larger models. Six of the 32B recipe's nine robust wins on physics are answer form. A more lenient answer judge removes most of the 32B recipe's net reach. The judge does not lower the 7B recipes' reach, so for them the answer-form reading rests only on their robust counts (Table 9). Whether real reach grows with scale is open.

channel	in training	held-out vs. plain RL
hint in the prompt	setup claims $55 \to 92\%$	$+0.01$ [-0.01, 0.03]
teacher derivation	touched problems $12 \to 59\%$	-0.05 [-0.08, -0.03]
written setup	correct sections $0.29 \to 0.45$	-0.03 [-0.05, -0.01]
question fine-tune	99% open with questions	at the base
setup reward	the policy gamed its scorer	not measurable
plain RL, the control		*value*
plain RL, held-out score		$0.16 \to 0.26$
revisited problems: solve rate		$0.10 \to 0.28$
credited problems: finishing share		
Eq. (2) / even split		71% / 50%
twins, credited minus other		
planned: mean score		$+0.096$ [0.008, 0.182]
planned: solved once in eight		$+0.095$ [-0.054, 0.242]
adjusted for the base†		$+0.077$ [-0.004, 0.160]
on training problems		$+0.186$
minus the fine-tune's excess†		$+0.061$ [-0.029, 0.149]
both finish 7 of 8†		-0.003 [-0.048, 0.037]

Table 5. The controlled probe on Qwen3-VL-8B-Thinking. Held-out differences from plain RL are on 87 hard problems, with 95% intervals. †Chosen after seeing the data.

5. A Bounded Probe: Can Supplied Knowledge Be Installed at Small Scale?

At this scale, reward raises reliability only on the problems it lands on. Beyond those problems, the planned test finds a gain whose part beyond finishing is not robust. To see whether outcome RL installs knowledge it is given, we supply it ourselves, at small scale. We train Qwen3-VL-8B-Thinking (Bai et al., 2025) with group-relative outcome RL (GSPO; Zheng et al., 2025; Shao et al., 2024) on physics at a 12k-token cap, for up to sixty updates, one seed per treated arm and two for the control. Plain RL repeats the result of §4: its held-out gain comes from finishing, with reliability flat (Appendix B.1). The knowledge supplied is the *setup*: the physical modelling a problem needs before any algebra. The base solves at most 16 of the 160 hard training problems at least once in eight samples. A gain on problems the base finishes and still gets wrong would count as new knowledge. Five *channels* supply the setup, from a hint in the prompt to a teacher's full derivation. Each channel changes the model during training (Table 5). None of the four channels we could evaluate is shown to carry over to held-out problems beyond plain RL. The test would detect a gain a third the size of plain RL's. A smaller gain

could go undetected, such as the hint's lead over the second control seed (Table 13). The fifth channel, a setup reward, taught the policy to game its scorer, so its held-out effect could not be measured.

Reward works where it lands. Training again and again on 60 hard problems, *revisiting*, raises their solve rate from 0.10 to 0.28 (Table 5). Revisiting moves neither the other 293 hard problems nor the held-out set. On the training problems the control earned reward on, finishing and reliability both rise, so finishing's share of the gain depends on the order of the split (Table 5).

Transfer beyond those problems is not shown. We call a training problem *credited* when the control earned reward on it in updates 31 to 60. Each hard training problem has a text-only *twin*: the same physics and method on a new, much easier surface. On the planned mean score, credited problems' twins gain more than other twins, clear of zero (Table 5). The other planned measure, solving a twin at least once in eight samples, includes zero. Over the two control seeds and the three orders of the split, finishing is 26–62% of that excess. The base finishes credited twins less often than other twins. The rest is not clear of zero under any order for the first control seed, and only in Eq. (2)'s order for the second (Appendix B.13). Credited problems also differ in kind: a supervised fine-tune never trained on them shows part of the same excess. Two adjustments added after the full read, for the base and for problem kind, each leave the excess including zero. A last test, chosen after a first look, finds no difference on twins both models finish. A partial transfer cannot be excluded. Even the control's own gains are not shown to transfer beyond finishing. The probe cannot tell whether the channels failed or supplied knowledge does not transfer this way.

6. Discussion and Limitations

The two questions, answered. On the released models, RL's gain is changed behaviour by our definition, not new knowledge. Where the budget binds, finishing is most of the change (82% at 4k and 8k). Where it no longer binds, the gain is reliability, from steps the base already writes or from answer form. New knowledge, measured as reach, is rare. At our probe's scale, supplied knowledge is not shown to transfer beyond the problems reward lands on. ProRL v2, trained fifty times as long as our probe, holds two of the five real physics wins. A much larger scale may matter.

What this means in practice. Report pass@1 at a budget together with fin@1(B) at that budget. Otherwise a model that stops sooner will look like one that reasons better. Compare a recipe with its base at or above its training cap. Below the cap, a recipe trained long looks worse than it is. QuestA and the two Polaris recipes score below their bases on physics at 16k and no longer do at their own caps (Table 7; 16k cells released). On MATH-500, QuestA still loses 0.07 at its cap.

Limitations. The released-pair study and the controlled run are bounded in scale, sampling and grading, and cover thinking bases only. Each limit below names what we checked.

The probe. The controlled run has one seed per treated arm, one model family and one domain. The run's held-out set was chosen on base-model samples. The channels ran at up to the control's scale, about two visits per problem. Prolonged training is untested. Supervised distillation at scale, known to transfer (Yue et al., 2025), is untested too.

Sampling. The released pairs are read with eight samples per problem (32 or 64 on lone-win problems). The budgets are three fixed ones, the cap and the next budget above it, up to 64k. Only Polaris evaluates higher, at 90k. The pairs are read at temperature 0.6, where nine recipes report at 0.7 to 1.4. The eight pairs re-read at their authors' temperatures split the same way (Appendix A.3). A second generation of 16 pairs, three of them cap-bound (Table 8), moves each term at the cap by at most 0.007, inside every interval.

Reach. Our definition counts as behaviour the tenth to third of the gain on problems the base never gets right but leaves some samples unfinished (§4.3). Filtering the physics problems on the controlled run's base caps how much reach they can show. On MATH-500, neither model was resampled on the problems it alone fails. The MATH-500 mirror may therefore hold real losses of coverage. On both physics sets, both models were resampled (Appendices A.5 and A.6). Physics is at most about 1% of the recipes' stated RL data but may be in the bases' pretraining. We did not check our problems against any recipe's training set.

Grading. The answer matcher is strict, and stricter for larger models. On the 40 reliability problems checked by language-model readers, 12 RL wins over the base differ only in answer form or grading. Part of the measured reliability gain therefore reflects which answer forms the matcher accepts. Regraded under the fixed judge protocol of Appendix C.2, 3–6% of the matcher-rejected physics answers become right for the 1.5B and 1.7B models, and 10–17% for the 4B to 32B models. On the eight pairs with stored texts and the 16 regenerated pairs, the judge raises pass@1($B \mid$ fin) for base and RL alike. On the 16 regenerated pairs, the judge lowers the reliability term at the cap by 0.003 on average, about 8% of it. The judge moves net reach by less than 0.01 (Table 8). Closed models were trace readers (Claude Opus 5.5) and setup grader (Claude Sonnet 5). The trace readings are by one language model. One author checked the model's claimed new steps. No human checked the step checks' finds. Only the new-step comparison (143 pairs) and the step checks were read blind.

References

Agarwal, S., Zhang, Z., Yuan, L., Han, J., and Peng, H. The unreasonable effectiveness of entropy minimization in LLM reasoning. *arXiv preprint arXiv:2505.15134*, 2025.

Aggarwal, P. and Welleck, S. L1: Controlling how long a reasoning model thinks with reinforcement learning. *arXiv preprint arXiv:2503.04697*, 2025.

Agrawal, P., Samanta, A., Ghasemlou, S., Bhandari, J., Asadi, K., Jiang, D., and Modi, A. Off-context GRPO: Learning to reason on hard problems using privileged information. *arXiv preprint arXiv:2607.19313*, 2026.

Amani, M. H., Lotfi, A., Baldwin, N. M., Bengio, S., Farajtabar, M., Abbe, E., and West, R. RL for reasoning by adaptively revealing rationales. *arXiv preprint arXiv:2506.18110*, 2025.

An, C., Xie, Z., Li, X., Li, L., Zhang, J., Gong, S., Zhong, M., Xu, J., Qiu, X., Wang, M., and Kong, L. POLARIS: A post-training recipe for scaling reinforcement learning on advanced reasoning models. Blog post, https://hkunlp.github.io/blog/2025/Polaris; checkpoint POLARIS-Project/Polaris-4B-Preview, 2025. URL https://hkunlp.github.io/blog/2025/Polaris.

Bai, S., Cai, Y., Chen, R., et al. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025.

Chen, J., Fu, D., Xu, S., Chen, J., Feng, Y., Yang, Y., Yan, J., Zha, H., and Xia, R. Milestones over outcome: Unlocking geometric reasoning with sub-goal verifiable reward. *arXiv preprint arXiv:2601.05073*, 2026.

Chen, J. C.-Y., Peng, B. X., Choubey, P. K., Huang, K.-H., Zhang, J., Bansal, M., and Wu, C.-S. Nudging the boundaries of LLM reasoning. *arXiv preprint arXiv:2509.25666*, 2025a.

Chen, Y., Yang, Z., Liu, Z., Lee, C., Xu, P., Shoeybi, M., Catanzaro, B., and Ping, W. AceReason-Nemotron: Advancing math and code reasoning through reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025b.

Chen, Z., Min, Y., Zhang, B., Chen, J., Jiang, J., Cheng, D., Zhao, W. X., Liu, Z., Miao, X., Lu, Y., Fang, L., Wang, Z., and Wen, J.-R. An empirical study on eliciting and improving R1-like reasoning models. *arXiv preprint arXiv:2503.04548*, 2025c.

Chen, Z., Qin, X., Wu, Y., Ling, Y., Ye, Q., Zhao, W. X., and Shi, G. Pass@k training for adaptively balancing exploration and exploitation of large reasoning models. *arXiv preprint arXiv:2508.10751*, 2025d.

Dang, Q.-A. and Ngo, C. Reinforcement learning for reasoning in small LLMs: What works and what doesn't. *arXiv preprint arXiv:2503.16219*, 2025.

DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Fatemi, M., Rafiee, B., Tang, M., and Talamadupula, K. Concise reasoning via reinforcement learning. *arXiv preprint arXiv:2504.05185*, 2025.

Fu, W., Gao, J., Shen, X., Zhu, C., Mei, Z., He, C., Xu, S., Wei, G., Mei, J., Wang, J., Yang, T., Yuan, B., and Wu, Y. AReaL: A large-scale asynchronous reinforcement learning system for language reasoning. *arXiv preprint arXiv:2505.24298*, 2025.

Gandhi, K., Chakravarthy, A., Singh, A., Lile, N., and Goodman, N. D. Cognitive behaviors that enable self-improving reasoners, or, four habits of highly effective STaRs. *arXiv preprint arXiv:2503.01307*, 2025.

Han, T., Shi, H., Hu, J., Yang, X., Wang, Z., and Su, J. Correct is not enough: Training reasoning planners with executor-grounded rewards. *arXiv preprint arXiv:2605.03862*, 2026.

He, B., Qu, Z., Liu, Z., Chen, Y., Zuo, Y., Qian, C., Zhang, K., Chen, W., Xiao, C., Cui, G., Ding, N., and Liu, Z. JustRL: Scaling a 1.5B LLM with a simple RL recipe. *arXiv preprint arXiv:2512.16649*, 2025a.

He, C., Luo, R., Bai, Y., Hu, S., Thai, Z. L., Shen, J., Hu, J., Han, X., Huang, Y., Zhang, Y., et al. OlympiadBench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. In *Proceedings of ACL*, 2024.

He, J., Liu, J., Liu, C. Y., Yan, R., Wang, C., Cheng, P., Zhang, X., Zhang, F., Xu, J., Shen, W., et al. Skywork Open Reasoner 1 technical report. *arXiv preprint arXiv:2505.22312*, 2025b.

He, Z., Liang, T., Xu, J., Liu, Q., Chen, X., Wang, Y., Song, L., Yu, D., Liang, Z., Wang, W., Zhang, Z., Wang, R., Tu, Z., Mi, H., and Yu, D. DeepMath-103K: A large-scale, challenging, decontaminated, and verifiable mathematical dataset for advancing reasoning. *arXiv preprint arXiv:2504.11456*, 2025c.

Hou, B., Zhang, Y., Ji, J., Liu, Y., Qian, K., Andreas, J., and Chang, S. ThinkPrune: Pruning long chain-of-thought of LLMs via reinforcement learning. *arXiv preprint arXiv:2504.01296*, 2025.

Hu, J., Zhang, Y., Han, Q., Jiang, D., Zhang, X., and Shum, H.-Y. Open-Reasoner-Zero: An open source approach

to scaling up reinforcement learning on the base model. *arXiv preprint arXiv:2503.24290*, 2025.

Huang, A., Block, A., Foster, D. J., Rohatgi, D., Zhang, C., Simchowitz, M., Ash, J. T., and Krishnamurthy, A. Self-improvement in language models: The sharpening mechanism. *arXiv preprint arXiv:2412.01951*, 2024.

Karan, A. and Du, Y. Reasoning with sampling: Your base model is smarter than you think. *arXiv preprint arXiv:2510.14901*, 2025.

Kazemnejad, A., Aghajohari, M., Portelance, E., Sordoni, A., Reddy, S., Courville, A., and Le Roux, N. VinePPO: Refining credit assignment in RL training of LLMs. *arXiv preprint arXiv:2410.01679*, 2024.

Li, J. et al. QuestA: Expanding reasoning capacity in LLMs via question augmentation. *arXiv preprint arXiv:2507.13266*, 2025.

Liao, B., Dong, H., Xu, X., Monz, C., and Bian, J. Self-hinting language models enhance reinforcement learning. *arXiv preprint arXiv:2602.03143*, 2026.

Liu, M., Diao, S., Lu, X., Hu, J., Dong, X., Choi, Y., Kautz, J., and Dong, Y. ProRL: Prolonged reinforcement learning expands reasoning boundaries in large language models. *arXiv preprint arXiv:2505.24864*, 2025a.

Liu, Z., Chen, C., Li, W., Qi, P., Pang, T., Du, C., Lee, W. S., and Lin, M. Understanding R1-Zero-Like training: A critical perspective. *arXiv preprint arXiv:2503.20783*, 2025b.

Luo, M., Tan, S., Wong, J., Shi, X., Tang, W. Y., Roongta, M., Cai, C., Luo, J., Li, L. E., Popa, R. A., and Stoica, I. DeepScaleR: Surpassing O1-Preview with a 1.5B model by scaling RL. Agentica project blog; model card https://huggingface.co/agentica-org/DeepScaleR-1.5B-Preview, 2025. URL https://pretty-radio-b75.notion.site/DeepScaleR-Surpassing-O1-Preview-with-a-1-5B-Model-by-Scaling-RL-19681902c1468005bed8ca303013a4e2.

Matsutani, K., Takashiro, S., Minegishi, G., Kojima, T., Iwasawa, Y., and Matsuo, Y. RL squeezes, SFT expands: A comparative study of reasoning LLMs. *arXiv preprint arXiv:2509.21128*, 2025.

Moshkov, I., Hanley, D., Sorokin, I., Toshniwal, S., Henkel, C., Schifferer, B., Du, W., and Gitman, I. AIMO-2 winning solution: Building state-of-the-art mathematical reasoning models with OpenMathReasoning dataset. *arXiv preprint arXiv:2504.16891*, 2025.

Muennighoff, N., Yang, Z., Shi, W., Li, X. L., Fei-Fei, L., Hajishirzi, H., Zettlemoyer, L., Liang, P., Candès, E., and Hashimoto, T. s1: Simple test-time scaling. *arXiv preprint arXiv:2501.19393*, 2025.

Nguyen, P. M., La, C. D., Nguyen, D. M. H., Chawla, N. V., Nguyen, B. T., and Doan, K. D. The reasoning boundary paradox: How reinforcement learning constrains language models. *arXiv preprint arXiv:2510.02230*, 2025.

Phan, D. et al. Training chain-of-thought via latent-variable inference. *arXiv preprint arXiv:2312.02179*, 2023.

Qin, T., Park, C. F., Kwun, M., Walsman, A., Malach, E., Anand, N., Tanaka, H., and Alvarez-Melis, D. Decomposing elements of problem solving: What "Math" does RL teach? *arXiv preprint arXiv:2505.22756*, 2025.

Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Setlur, A., Yang, M. Y. R., Snell, C., Greer, J., Wu, I., Smith, V., Simchowitz, M., and Kumar, A. e3: Learning to explore enables extrapolation of test-time compute for LLMs. *arXiv preprint arXiv:2506.09026*, 2025.

Shao, R., Li, S. S., Xin, R., Geng, S., Wang, Y., Oh, S., Du, S. S., Lambert, N., Min, S., Krishna, R., Tsvetkov, Y., Hajishirzi, H., Koh, P. W., and Zettlemoyer, L. Spurious rewards: Rethinking training signals in RLVR. *arXiv preprint arXiv:2506.10947*, 2025.

Shao, Z., Wang, P., Zhu, Q., et al. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

Sheng, G. et al. HybridFlow: A flexible and efficient RLHF framework. *arXiv preprint arXiv:2409.19256*, 2024.

Walder, C. and Karkhanis, D. Pass@k policy optimization: Solving harder reinforcement learning problems. *arXiv preprint arXiv:2505.15201*, 2025.

Wang, L., Su, E., Liu, J., Li, P., Xiao, J., Zhang, W., Dai, X., Chen, X., Meng, Y., Bai, L., Ouyang, W., Tang, S., Wang, A., and Ma, X. PhysUniBench: A multimodal physics reasoning benchmark at undergraduate level. *arXiv preprint arXiv:2506.17667*, 2025a.

Wang, P. et al. Math-Shepherd: Verify and reinforce LLMs step-by-step without human annotations. *arXiv preprint arXiv:2312.08935*, 2023.

Wang, S., Yu, L., Gao, C., Zheng, C., Liu, S., Lu, R., Dang, K., Chen, X., Yang, J., Zhang, Z., et al. Beyond the 80/20 rule: High-entropy minority tokens drive effective reinforcement learning for LLM reasoning. *arXiv preprint arXiv:2506.01939*, 2025b.

Wang, X., Han, J., Jiang, Z., et al. Don't tell the answer, truly guide the reasoning during RL rollouts. *arXiv preprint arXiv:2510.09388*, 2025c. Findings of ACL 2026.

Wang, Y., Yang, Q., Zeng, Z., Ren, L., Liu, L., Peng, B., Cheng, H., He, X., Wang, K., Gao, J., Chen, W., Wang, S., Du, S. S., and Shen, Y. Reinforcement learning for reasoning in large language models with one training example. *arXiv preprint arXiv:2504.20571*, 2025d.

Wu, F., Tu, A., Xuan, W., Qi, H., Huang, X., Zeng, Q., Talaei, S., Xiao, Y., Xia, P., Tang, X., et al. Position: The hidden costs and measurement gaps of reinforcement learning with verifiable rewards. *arXiv preprint arXiv:2509.21882*, 2025a.

Wu, F., Xuan, W., Xing, D., Lu, X., Liu, M., Dong, Y., Harchaoui, Z., and Choi, Y. The invisible leash: Why RLVR may or may not escape its origin. *arXiv preprint arXiv:2507.14843*, 2025b.

Wu, J., Liu, S., Tu, H., Yu, H., Huang, X., Zou, J., Xie, C., and Zhou, Y. Knowledge or reasoning? A close look at how LLMs think across domains. *arXiv preprint arXiv:2506.02126*, 2025c.

Xi, Z. et al. Training large language models for reasoning through reverse curriculum reinforcement learning. In *ICML*, 2024. arXiv:2402.05808.

Xia, Y., Xu, C., Yao, Z., et al. Learning to hint for reinforcement learning. *arXiv preprint arXiv:2604.00698*, 2026.

Xie, W., Zhao, H., Liu, W., Zhu, Y., Chen, L., Ye, M., Chen, Z., Xu, Y., Dong, S., Wang, Z., Xu, X., Shi, K., Wu, R., Zhang, X., Shao, W., Chang, B., Duan, N., and Wang, J. Step-wise rubric rewards for LLM reasoning. *arXiv preprint arXiv:2605.17291*, 2026.

Xu, X., Xu, Q., Xiao, T., Chen, T., Yan, Y., Zhang, J., Diao, S., Yang, C., and Wang, Y. UGPhysics: A comprehensive benchmark for undergraduate physics reasoning with large language models. *arXiv preprint arXiv:2502.00334*, 2025.

Yan, J., Li, Y., Hu, Z., Wang, Z., Cui, G., Qu, X., Cheng, Y., and Zhang, Y. Learning to reason under off-policy guidance. *arXiv preprint arXiv:2504.14945*, 2025.

Yang, C. and Chen, J. RLVR is a kernel, not a function: Statistical inference for pass@k crossovers. *arXiv preprint arXiv:2609.22547*, 2026.

Yang, S. Physics-R1: An audited olympiad corpus and released verifiers for visual physics reasoning. *arXiv preprint arXiv:2605.14040*, 2026.

Yu, C., Wang, Y., Guo, Z., Lin, H., Xu, S., Zang, H., Zhang, Q., Wu, Y., Zhu, C., Hu, J., Huang, Z., Wei, M., Xie, Y., Yang, K., Dai, B., Xu, Z., Du, J., Wang, X., Fu, X., Shi, L., Liu, Z., Chen, K., Liu, W., Liu, G., Li, B., Yang, J., Yang, Z., Dai, G., and Wang, Y. RLinf: Flexible and efficient large-scale reinforcement learning via macro-to-micro flow transformation. *arXiv preprint arXiv:2509.15965*, 2025a.

Yu, L., Yang, T., Ding, S., Jin, R., Gu, N., Hao, X., Nie, S., Xiong, D., Yin, W., Sun, Y., and Wu, H. KnowRL: Boosting LLM reasoning via reinforcement learning with minimal-sufficient knowledge guidance. *arXiv preprint arXiv:2604.12627*, 2026.

Yu, Q. et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025b.

Yue, Y., Chen, Z., Lu, R., Zhao, A., Wang, Z., Yue, Y., Song, S., and Huang, G. Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model? In *Advances in Neural Information Processing Systems*, 2025.

Zelikman, E., Wu, Y., Mu, J., and Goodman, N. D. STaR: Bootstrapping reasoning with reasoning. In *NeurIPS*, 2022. arXiv:2203.14465.

Zhang, H., Fu, J., Zhang, J., Fu, K., Wang, Q., Zhang, F., and Zhou, G. RLEP: Reinforcement learning with experience replay for LLM reasoning. *arXiv preprint arXiv:2507.07451*, 2025a.

Zhang, J. et al. R1-VL: Learning to reason with multimodal large language models via step-wise group relative policy optimization. In *ICCV*, 2025b. arXiv:2503.12937.

Zhang, T. et al. The wisdom of hindsight makes language models better instruction followers. In *ICML*, 2023. arXiv:2302.05206.

Zhang, X., Wu, S., Zhu, Y., Tan, H., Yu, S., He, Z., and Jia, J. Scaf-GRPO: Scaffolded group relative policy optimization for enhancing LLM reasoning. *arXiv preprint arXiv:2510.19807*, 2025c.

Zhang, X. et al. PhysReason: A comprehensive benchmark towards physics-based reasoning. *arXiv preprint arXiv:2502.12054*, 2025d. ACL 2025.

Zheng, C., Liu, S., Li, M., Chen, X.-H., Yu, B., Gao, C., Dang, K., Liu, Y., Men, R., Yang, A., Zhou, J., and Lin, J. Group sequence policy optimization. *arXiv preprint arXiv:2507.18071*, 2025.

Zhou, R., Li, S., Zhang, A., and Leqi, L. ExPO: Unlocking hard reasoning with self-explanation-guided reinforcement learning. *arXiv preprint arXiv:2507.02834*, 2025.

Zhou, T. When RLVR shrinks the reasoning boundary: Diagnosing pass@k inversion. *arXiv preprint arXiv:2607.20543*, 2026.

A. The released pairs

How the pairs were read. The base scores zero on U, so by Eq. (1) the change is RL's score on U plus $F + P$ (§3.1). Pooled shares divide $\sum |F|$, $\sum |P|$ and $\max(0, \sum \text{reach@8}(B))$ by their total. F includes the cross term. Counting the cross term as reliability puts the finishing share below one half in only three cells, all physics at 16k. Counting only full credit moves a cell's change by up to 0.03.

Models and sampling. Table 6 lists the twenty pairs (DeepSeek-AI, 2025; Moshkov et al., 2025; Qwen Team, 2025; Luo et al., 2025; Liu et al., 2025a; Dang & Ngo, 2025; Chen et al., 2025c; He et al., 2025a;c; Yu et al., 2025a; Fu et al., 2025; He et al., 2025b; Chen et al., 2025b; Li et al., 2025; Yu et al., 2026; An et al., 2025; Setlur et al., 2025). OpenMath-Nemotron-1.5B is a supervised fine-tune of Qwen2.5-Math-1.5B, without RL. The RL data is mathematics, with code or puzzles in some recipes. Physics is at most about 1% of it in any recipe that states its mix, except ProRL v2. Each model writes eight samples on each of 300 random physics problems. Physics uses the paper's outcome reward, MATH-500 symbolic equivalence. The two generation protocols agree within 0.01 on cell means.

The base catches up on the problems behind finishing. Samples 0–3 pick problems where the RL model finishes more and scores higher. Samples 4–7 score them, so selection does not inflate the comparison. The RL model scores 0.68 at 4k, the base 0.44 at 4k and 0.80 at 64k. At 8k the three scores are 0.65, 0.38 and 0.59. The Qwen3 bases' 64k read covers physics only.

Statistics. Resampling the seven bases instead widens the at-cap intervals to [18, 59] for finishing and [38, 70] for reliability, which overlap. Reliability's lead at the next budget above the cap stays clear of zero, 34 points [19, 62]. Pooled over eighteen recipes and both benchmarks with random effects, the matched change is 0.014 [0.008, 0.021] (between-recipe heterogeneity $I^2 = 0.73$). On physics alone the matched change is 0.028 [0.013, 0.043] ($I^2 = 0.69$). The finished-length ratio is not significantly associated with the reliability term at the cap (0.35, $p = 0.15$).

The shares survive moving the cross term and removing selection. A matched split, taking the change only on problems both models finish in seven of eight samples, keeps selection out. At the cap on physics, the shares are 45% finishing, 49% reliability and 5% reach, and beyond the cap 28%, 62% and 10%. The matched split moves them toward reliability, to 37% and 57% at the cap and 19% and 71% beyond, where e3 has no matched read. A symmetric split moves them by a point. For the five recipes whose finishing term is negative at the cap, selection favours reliability: matched reliability is lower for four (Table 7). Two of the five differences, QuestA's and Skywork-7B's, are within the 0.007 by which the two generations differ. Under the matched split, the cap-bound recipes' finishing and reliability at the cap are $+0.155$ and $+0.023$ for ProRL v2 (against $+0.203$ and -0.026 plain), $+0.047$ and $+0.048$ for JustRL-DeepSeek, $+0.032$ and $+0.009$ for e3, and 0.000 and $+0.003$ for Open-RS3. The Polaris pairs fall to $+0.015$ and $+0.012$, intervals including zero.

Recipes are not independent tests. The twenty recipes start from seven bases, eight of them from R1-Distill-1.5B. At 4k and 8k, finishing is the larger part of the change on both benchmarks together in each of the 17 smaller recipes whose score changes (one-sided sign test $p = 8 \times 10^{-6}$). Counted by base, five of five gives only $p = 0.03$, so we rely on the pooled bootstrap. Where the base finishes, the base-level test is stronger. With the base as a random effect (a one-way random intercept fitted by REML, with a t interval on six degrees of freedom), reliability exceeds finishing at the cap by $+0.033$ [$+0.016$, $+0.049$] over the sixteen recipes that are not cap-bound, and by $+0.030$ [$+0.018$, $+0.043$] at the next budget. Leaving out any one base, the excess at the cap is $+0.030$ to $+0.040$, each clear of zero. The excess is positive on average for each of their seven bases (one-sided sign test $p = 0.008$). Over all twenty recipes at the cap it is $+0.014$ [-0.026, $+0.053$], because all four cap-bound recipes gain more by finishing. The length association is partly between bases: -0.78 within bases ($p = 0.006$), -0.62 ($p = 0.12$) over the eight R1-Distill-1.5B recipes. On physics, AReaL's and DeepMath's finishing-share intervals include one half. We registered no 16k prediction. At 16k the first eight pairs gave 71% [49, 84].

The registered test. Predictions for the ten added pairs were committed before their outputs existed. A same-day amendment added four recipes. H1, a finishing share above one half for each added recipe whose score changes at 4k or 8k, passes for all ten ($p = 0.001$). H1 is weak where every base often hits the budget before answering. H2, a mean matched change below 0.03 with none above 0.10, held at $+0.017$ while pooling physics with MATH-500. On MATH-500 the matched base is near ceiling. On physics alone the ten average $+0.042$ and QuestA reaches $+0.108$, outside H2's bounds. Four of five registered 8k finishing directions held. JustRL-Nemotron finishes less (-0.05), not more.

A.1. How the traces were read

Protocol. We read up to twelve problems per recipe that the RL model solves and its base, despite finishing, does not. Two

model	base	algorithm, scale	length shaping	RL data	training cap	rollout T	authors' evaluation
DeepScaleR-1.5B-Preview	R1-Distill-1.5B	GRPO with a KL loss; about 1,750 steps	none in the reward; cap raised 8k→ 16k→ 24k	40k math	8k to 24k	0.6	32k, T=0.6
ProRL v2 (Nemotron-Research-Reasoning-Qwen-1.5B)	ProRL v1, from R1-Distill-1.5B	REINFORCE++ baseline, clip-higher, dynamic sampling, KL with reference resets; 3,000 steps in all, counting ProRL v1's 2,000	cosine length penalty, switched on in scheduled cycles	136k: math, code, STEM, puzzles, instruction following	8k	n.d.	32k, T=0.6
Open-RS3	R1-Distill-1.5B	GRPO with KL; 50 steps	cosine length reward	7k math	3,584	0.7	32k, T=0.6
STILL-3-1.5B-preview	R1-Distill-1.5B	GRPO with KL and reference resets; about 210 steps (inferred)	none documented	30k math	29k	1.0	32k, T=0.6
Skywork-OR1-7B	R1-Distill-7B	GRPO without KL, adaptive entropy, rejection sampling; 1,320 steps	none; cap raised 16k→ 32k	105k math, 14k code	16k to 32k	1.0	32k, T=1.0
Skywork-OR1-Math-7B	R1-Distill-7B	as above, two updates per step; 2,160 steps	none; cap raised 8k→ 32k	Skywork set (mix not stated)	8k to 32k	1.0	32k, T=1.0
AceReason-Nemotron-7B	R1-Distill-7B	on-policy GRPO without KL; math, then code, then math	none; cap raised 8k→ 32k	49k math, 8.5k code	8k to 32k	n.d.	32k, T=0.6
Polaris-4B-Preview	Qwen3-4B	GRPO without KL or entropy loss, clip-higher; about 700 steps	none; cap raised 40k→ 52k	30k math	40k to 52k	1.4 to 1.5	90k, T=1.4

Table 6. What each released recipe did, from its model card, paper, blog or training scripts (n.d.: not documented). This part lists the eight recipes read first. The ten added are in the continuation. Rewards are outcome rewards unless a length term is listed. Outcome rewards are 0 or 1, except ±1 for STILL-3, JustRL and DeepMath and ±5 for RLinf and AReaL. A truncated response usually scores as a wrong one. A truncated response scores -2 under STILL-3, which docks a missing boxed answer, and under DeepMath's overlong penalty. Some verifiers credit a truncated response if a correct boxed answer appears before the cut. "Authors' evaluation" is the generation budget and temperature its authors report with. Figures and tables use short names. R1-1.5B to R1-32B are the DeepSeek-R1-Distill-Qwen bases, OMN-1.5B is OpenMath-Nemotron-1.5B, and Q3-1.7B and Q3-4B are the Qwen3 bases. JustRL-D and JustRL-N are JustRL-DeepSeek and JustRL-Nemotron. Skywork-7B, Skywork-Math and Skywork-32B are Skywork-OR1-7B, Skywork-OR1-Math-7B and Skywork-OR1-32B. AceReason and AceReason-14B are AceReason-Nemotron-7B and AceReason-Nemotron-14B. The other recipes go by their names without the size.

language-model readers (Claude Opus 5.5) see both derivations with model and grade labels, so the reading is not blind. The readers mark whether the RL derivation states a principle, method or reading the base's never states. A third reader settles disagreements. Of the 195 wins, 42 are grading or answer-form artifacts (full table in the released repository).

Blind reading. A blind, shuffled re-read of 143 pairs asks whether each side states a step the other never does. Weighted to the 195 wins, it finds a new step on the RL side in 46% and on the base side in 50%, assuming the 60 unmarked pairs read stand for the other 112. A failed derivation lacks steps as often as the winning one, so the labelled count measures attention as much as content.

A human check of the new steps. One author, blind to side, checked the claimed new steps in 50 of the 143 pairs, drawn at random with a fixed seed. Of 45 decided claims (12 unsure), 19 hold. There is no detectable difference between sides: 10 of 26 for RL (0.38 [0.22, 0.57]), 9 of 19 for the base (0.47 [0.27, 0.68]). The readers overcount on both sides, so the counts above are upper bounds. The symmetry the reuse finding rests on survives. The check tests only claims made, so it gives no κ.

Repertoire check. A step missing from one base sample may still be in the model, so checkers searched four other base samples on the same problem. Of the 83 steps flagged in the unblinded reading (37 from the first eight recipes, 46 from the ten added), 77 appear in one, 37 used correctly. Six appear in none. Two of the 77 were found only weakly and no finds were re-examined, so the count may lean toward finding them. Of the 10 RL-side steps the human check upheld, the base's samples hold 9. Of the 80 blind RL-side steps, the search finds 58, 32 of them used correctly in a base sample that then abandons them. The search misses 13, and 9 were not searched because their base samples were not stored. The reverse search gave each blind base-side step on a physics problem, with the same prompt, to a reader with four samples of the RL model from the 16 regenerated pairs. The reverse search finds 52 of the 62, 15 of them used correctly, and misses 10. The regenerated pairs hold no math or e3 texts, so 9 base-side steps were not searched. On physics the forward search finds 49 of 60. Beyond the verdicts, the reverse readers wrote short reports, which call five of the 52 finds borderline. Counting the five borderline finds as misses leaves 47 of 62. The forward readers returned verdicts only, so the forward search has no count of borderline finds. The human check took one author one sitting of about half an hour. The same author later checked the 15 misses the human check had not decided, blind to side and with the same two questions. Over both checks, the author

model	base	algorithm, scale	length shaping	RL data	training cap	rollout T	authors' evaluation
JustRL-DeepSeek-1.5B	R1-Distill-1.5B	GRPO without KL or entropy bonus, clip-higher; 4,380 steps	none; fixed 16k context	17k math (DAPO-Math-17k)	15k	1.0	32k, T=0.7
DeepMath-1.5B	R1-Distill-1.5B	GRPO with a KL loss, clip-higher, dynamic sampling; 1,800 steps	overlong penalty, 0 to -1 over 20k to 24k; truncated -2	103k math (DeepMath-103K)	24k	0.6	32k, T=0.6
RLinf-math-1.5B	R1-Distill-1.5B	GRPO without KL; steps n.d.	none	106k math, with 10k logic puzzles	27.6k	1.0	32k, T=0.6
AReaL-1.5B-Preview-Stage-3	R1-Distill-1.5B, through two released earlier stages	critic-free PPO with a small KL; about 435 steps over three stages	none; cap raised 8k→ 16k→ 24k	40k math (DeepScaleR set)	8k to 24k	0.6	32k, T=0.6
JustRL-Nemotron-1.5B	OpenMath-Nemotron-1.5B	as JustRL-DeepSeek; 3,440 steps	none; fixed 16k context	17k math (DAPO-Math-17k)	15k	1.0	32k, T=0.7
QuestA-Nemotron-1.5B	OpenMath-Nemotron-1.5B	GRPO without KL, dynamic sampling; a partial solution in every prompt (50%, then 25% of it); 2,000 steps	none	9k hard math (OpenR1-Math)	24k	1.0	32k, T=0.7
KnowRL-Nemotron-1.5B	OpenMath-Nemotron-1.5B	GRPO without KL, clip-higher, dynamic sampling; knowledge-point hints in 78% of prompts; 2,960 steps	none	9k math (QuestA's problems; a third multiple choice)	24k (paper), 32k (card)	1.0	32k, T=0.9
DeepMath-Omn-1.5B	OpenMath-Nemotron-1.5B	as DeepMath-1.5B; 700 steps	overlong penalty, 0 to -1 over 20k to 24k; truncated -2	103k math (DeepMath-103K)	24k	0.6	32k, T=0.6
e3-1.7B	Qwen3-1.7B	GRPO with a KL loss, clip-higher; steps n.d.	none; cap raised 8k→ 16k with harder data	15k math (DeepScaleR set)	8k to 16k	0.6	32k, T=0.6
Polaris-1.7B-Preview	Qwen3-1.7B	as Polaris-4B; about 500 steps	none; cap raised 40k→ 48k	53k math, with 4.6k logic puzzles	40k to 48k	1.5 to 1.55	90k, T=1.4
AceReason-Nemotron-14B	R1-Distill-14B	on-policy GRPO without KL or entropy loss; math, then code, then math	none; cap raised 8k→ 24k on math, 32k on code and the last math stage	49k math, 8.5k code	8k to 32k	n.d.	32k, T=0.6
Skywork-OR1-32B	R1-Distill-32B	GRPO without KL, adaptive entropy, rejection sampling; released at step 1,000 of 1,130	none; cap raised 16k→ 24k	Skywork set, as Skywork-OR1-7B	16k to 24k	1.0	32k, T=1.0

Table 6. (continued) The ten added recipes, registered before their outputs were generated. Below the rule are the 14B and 32B recipes, read in addition and not registered. QuestA and KnowRL put guidance in the training prompt. QuestA's guidance is a prefix of a reference solution. KnowRL's is knowledge points. The hints are absent from this paper's prompts. QuestA's authors also evaluate without the hints. KnowRL's authors report both. Columns as above.

upholds 2 of the 11 RL-side steps the forward search misses on physics, rejects 6 and is unsure of 3. Of the 10 base-side steps the reverse search misses, the author upholds 4, rejects 2 and is unsure of 4. The forward search's two math misses are both rejected. Counting both checks, the other model's samples hold 9 of the 11 upheld RL-side steps and 6 of the 10 upheld base-side steps on physics.

Reliability problems. The checks above cover lone wins. A second check covers problems where the gain is reliability: the base finishes all eight samples at 16k and gets one to four right, and the RL model gets more right. There are 249 such pair–problems in 15 of the 16 regenerated pairs. We drew 40 problems, two or three per pair, with a fixed seed. One blind reader per problem compared an RL derivation graded right with a base derivation graded wrong, shuffled among the reverse comparisons, and named the decisive step. A fresh reader then searched the base's right and wrong samples for that step. Twelve comparisons had no step to search: the values match, and only form, units, a missing box or grading differs. Of the 28 searched, 2 claims were judged not decisive. Of the other 26 steps, 25 are used correctly in a base sample graded right, 1 is stated only in a base sample graded wrong, and none is absent. The base's wrong samples state the step on all 26 problems: 17 without carrying it through, 9 using it correctly and failing elsewhere. In the reverse direction, 35 problems have an RL sample graded wrong that finished. Of these, 11 differ only in answer form or grading, and 24 were searched. In 5 of the 24, the search reader judged the claimed step not decisive. Of the 19 decisive base steps, 17 are used correctly in an RL sample graded right, 2 are stated only, and none is absent. A step that decides correctness should appear in most right samples, so the wrong samples are the informative part. Each comparison had one reader. Readers did not see the gold answers. The labels are the matcher's, not the judge's.

Two more strata, at each recipe's cap. The same check, forward only, ran at the caps of the 13 recipes in the second

recipe	cap	fin@1(B), base→RL	F	P	matched P	reach	change [interval]	P, MATH-500
Open-RS3	4k	0.21 → 0.21	0.002	0.001	0.003	0	0.003 [−0.007, 0.014]	−0.002
ProRL v2	8k	0.45 → 0.99	**0.203**	**−0.026**	0.023	0	0.178 [0.146, 0.211]	**−0.025**
JustRL-D	16k	0.82 → 0.98	**0.063**	**0.032**	**0.048**	−0.003	0.092 [0.068, 0.118]	0.008
e3	16k	0.84 → 0.92	**0.043**	−0.002	0.009	**−0.007**	0.034 [0.013, 0.056]	**−0.018**
JustRL-N	16k	0.93 → 0.95	**0.013**	**0.024**	**0.032**	0.003	0.039 [0.018, 0.060]	**0.014**
DeepScaleR	24k	0.94 → 0.98	**0.020**	**0.027**	**0.039***	0.004	0.051 [0.031, 0.071]	0.009
DeepMath	24k	0.94 → 0.99	**0.025**	**0.052**	**0.071***	0.007	0.084 [0.063, 0.105]	0.011
AReaL	24k	0.94 → 0.96	**0.012**	0.008	**0.017***	−0.001	0.019 [0.003, 0.036]	**0.012**
RLinf	28k	0.95 → 0.99	**0.017**	**0.036**	**0.039***	0	0.054 [0.030, 0.078]	**0.025**
STILL-3	28k	0.95 → 0.97	**0.010**	**0.016**	**0.020***	0.005	0.030 [0.014, 0.047]	0.004
QuestA	24k	0.98 → 0.81	**−0.036**	**0.051**	**0.045***	**0.006**	0.021 [0, 0.043]	**0.025**
KnowRL	24k	0.98 → 0.98	**0.006**	**0.037**	**0.033***	0	0.043 [0.022, 0.065]	**0.024**
DeepMath-Omn	24k	0.98 → 0.97	0.002	**0.018**	0.017*	−0.001	0.020 [0.002, 0.038]	**0.011**
Skywork-32B	24k	0.98 → 0.97	**−0.009**	**0.047**	**0.037**	0.013	0.052 [0.027, 0.078]	**0.010**
Skywork-7B	32k	0.99 → 0.96	**−0.008**	**0.050**	**0.052***	0.009	0.051 [0.029, 0.074]	**0.025**
Skywork-Math	32k	0.99 → 0.98	0.001	**0.036**	**0.027***	0.013	0.050 [0.023, 0.077]	**0.023**
AceReason	32k	0.99 → 1.00	**0.008**	**0.022**	0.019*	−0.002	0.028 [0.005, 0.050]	**0.016**
AceReason-14B	32k	0.99 → 1.00	**0.006**	**0.045**	**0.042***	0.010	0.061 [0.036, 0.086]	**0.018**
Polaris-1.7B	48k	0.99 → 0.90	**−0.027**	**0.023**	0.015	0.003	0 [−0.022, 0.022]	–
Polaris-4B	52k	1.00 → 0.96	**−0.016**	**0.021**	0.012	0.005	0.009 [−0.010, 0.029]	–

Table 7. What RL buys at each recipe's training cap, on physics: one row per recipe, the cells behind Figure 1 and the cap rows of Figure 2. F, P and net reach split the change in full-credit pass@1, so the change can differ from a mean-credit score by up to 0.03. **Bold**: the 95% paired bootstrap interval over problems excludes zero. F and P are computed outside the unsolved set U, so F's sign can differ from the printed change in fin@1(B). Matched P is reliability on problems both models finish seven times in eight. *: computed on the second generation of Table 8, which agrees within 0.007 on every term. Open-RS3's matched P rests on ten problems. The Polaris pairs were not read on MATH-500 at their caps. The cells at the next budget above the cap, MATH-500 in full and every interval are in the released tables.

generation that are not cap-bound. We drew 40 of the 270 pair–problems where the base finishes all eight samples and is right in five to seven, and the RL model is right more often. There, 12 comparisons differ only in answer form or grading, and 2 claims were judged not decisive. All 26 decisive RL steps are used correctly in a base sample graded right. The base's other wrong samples state 8 of them without carrying them through, use 5 and lack 2. On 11 problems the base has no other wrong sample. We also took all 32 distinct problems among the 72 pair–problems where the base is never right but leaves a sample unfinished, and the RL model is right at least once. There, 5 comparisons differ only in form or grading, and 2 claims were judged not decisive. The compared base derivation finished in all 32. Of the 25 decisive RL steps, the base's own samples state 17 without carrying them through and use 8 correctly while failing elsewhere. None is absent.

On the cap-bound recipes' Z, the base's samples write 31 of the 32 deciding steps. We drew 40 of the 68 distinct problems among the 84 Z pair–problems of ProRL v2, Open-RS3 and JustRL-DeepSeek, by the same rules. Each base sample was given whole, up to 64k tokens, with the point where the cap falls marked. The compared base derivation was cut at the cap on 15 of the 40. Of the 40 comparisons, 6 differ only in form or grading. In 1 more, the base derivation has RL's physics but is cut before it answers. One claim was judged not decisive. Within the cap, the base's own samples use 24 of the 32 decisive RL steps correctly and state 7 without carrying them through. The one step absent, even at 64k tokens, weights two hydrogen lines by $2J+1$ to get their intensity ratio. The RL derivation there takes the ratio from a rule it does not derive. Given 64k tokens, a base sample graded right uses the step on 15 of the 32, and on all 8 of Open-RS3's. Each comparison and each search in these checks had one reader, and no reader saw the gold answers.

A.2. Judge regrade

The split survives a judge regrade. Under judge protocol v2, a pinned Qwen3-32B regraded part by part every stored physics sample the frozen matcher scored below one. The judge is stricter than an earlier, lenient language-model audit of answers the matcher rejected: on 2,800 answers both read, the audit alone calls 440 right and the judge alone 55. The answer-form share of the reliability term lies between two estimates: the judge lowers the term by 8% (§6), and 12 of 40 sampled reliability wins differ only in answer form or grading (§A.1). Net reach moves by less than 0.01 in every cell (full table in the released repository). For the 32B pair it falls at the cap from 0.013 to 0.004, because 29 of the 68 problems the base left unsolved in eight samples under the matcher are solved under the judge.

The same holds on a regenerated read of sixteen of the pairs. The judge moves F by at most 0.021, P by at most 0.011 and net reach by at most 0.008 (Table 8). No reliability term changes sign. AceReason-Nemotron-14B's net reach at its cap falls from 0.016 to 0.010. For the 7B recipes the judge leaves net reach in place or raises it. For these, the answer-form

reading of §4.3 rests on the robust cells of Appendix A.4, not on the regrade.

The second generation agrees with the main read within 0.007. Split like Table 7, it moves F by at most 0.004, P by at most 0.006 and net reach by at most 0.007, each inside the bootstrap half-width. Its reliability term is lower by 0.002 on average. Only Open-RS3's reliability term changes sign, from $+0.001$ in the main read to -0.002.

recipe	cap	F	P	net reach
ProRL v2	8k	$+0.219 \to +0.240$	$-0.021 \to -0.026$	$+0.001 \to -0.001$
JustRL-D	16k	$+0.058 \to +0.063$	$+0.032 \to +0.028$	$-0.004 \to -0.007$
DeepScaleR	24k	$+0.015 \to +0.017$	$+0.040 \to +0.039$	$-0.002 \to +0.000$
DeepMath	24k	$+0.019 \to +0.021$	$+0.070 \to +0.069$	$+0.004 \to +0.006$
RLinf	28k	$+0.012 \to +0.013$	$+0.039 \to +0.038$	$+0.000 \to +0.002$
STILL-3	28k	$+0.006 \to +0.006$	$+0.023 \to +0.021$	$+0.000 \to +0.000$
AReaL	24k	$+0.009 \to +0.010$	$+0.012 \to +0.012$	$-0.005 \to -0.002$
Open-RS3	4k	$+0.002 \to +0.002$	$-0.003 \to -0.007$	$+0.000 \to +0.000$
JustRL-N	16k	$+0.006 \to +0.007$	$+0.032 \to +0.025$	$+0.002 \to +0.000$
QuestA	24k	$-0.047 \to -0.051$	$+0.062 \to +0.057$	$+0.004 \to +0.004$
KnowRL	24k	$-0.000 \to -0.000$	$+0.046 \to +0.036$	$+0.000 \to -0.001$
DeepMath-Omn	24k	$+0.000 \to +0.000$	$+0.016 \to +0.005$	$-0.008 \to -0.010$
Skywork-7B	32k	$-0.014 \to -0.016$	$+0.060 \to +0.065$	$+0.007 \to +0.011$
AceReason-14B	32k	$+0.001 \to +0.001$	$+0.066 \to +0.058$	$+0.016 \to +0.010$
Skywork-Math	32k	$-0.002 \to -0.003$	$+0.048 \to +0.049$	$+0.014 \to +0.013$
AceReason	32k	$+0.007 \to +0.007$	$+0.030 \to +0.039$	$+0.002 \to +0.010$

Table 8. A second generation of 16 pairs on physics at each recipe's cap, each cell matcher → judge (judge protocol v2). Each model wrote one 64k generation at temperature 0.6, read at the cap by truncation, eight samples per problem on the 300 physics problems. Its columns use partial credit and split the whole change without taking reach out first, unlike Table 7, which takes reach and its mirror out first and adds the mirror back to P. The 16k rows and the reliabilities behind each term are in the released tables.

A.3. Authors' temperature

For the eight pairs re-read, the split does not depend on the sampling temperature. Nine recipes report at 0.7 to 1.4 where we read at 0.6. Each of the eight pairs re-read at the authors' settings splits as at 0.6 (full table in the released repository). At its cap, Skywork-OR1-32B's net reach falls from 0.013 [-0.001, $+0.029$] to $+0.002$ [-0.010, $+0.015$], as in the regrade. The reach it showed at 0.6 was sampling noise and answer form. Polaris-1.7B at 1.4 finishes 97% of its derivations at the cap rather than 90%. Read at 1.4 against its base at 0.6, Polaris-1.7B is 0.019 [0.001, 0.037] ahead. Across the other six recipes, at-cap reliability runs from $+0.012$ to $+0.065$, against $+0.019$ to $+0.051$ in the re-read's own samples at 0.6. The lowest, Polaris-4B's $+0.012$ [-0.005, $+0.028$], has an interval that includes zero. KnowRL, which reports 0.9, is the one recipe above 0.6 not re-read.

A.4. Robust reach by recipe

Ordering. Rows are ordered by robust reach: wins in at least four of eight samples on problems the base leaves unsolved in eight.

A tenth to a third of the gain at the cap falls on problems the base never gets right but leaves some samples unfinished. Call these problems Z: outside U, with no correct base sample and at least one unfinished. Eq. (4) books RL's gain on Z to F and P. We read Z on the 16 regenerated pairs, one 64k run per model read at each budget. Their F, P and reach at the cap are each within 0.007 of Table 7. For the 13 of them that are not cap-bound, RL's pass@1 on Z is 10% of their summed gain. Counting Z as unsolved would move their net reach by at most 0.008 and their robust reach by at most two problems. For the three cap-bound recipes, Z carries 34% of the summed gain. ProRL v2 alone carries 0.066 of its 0.177 there. Counted as unsolved, Z would raise ProRL v2's net reach from 0.001 to 0.067 and JustRL-DeepSeek's from -0.004 to 0.013. Open-RS3's would stay at about zero. ProRL v2's robust reach on this read would rise from 2 to 21 problems. Split within Z by Eqs. (2)–(3), RL's gain there is 13% finishing and 87% reliability for the 13 recipes, and 60% and 40% for the three cap-bound ones. The base is never right on Z, so the reverse order of the split books all of that gain to reliability, and the even split halves the finishing share. The fourth cap-bound recipe, e3, has no second-generation texts, so its Z is not measured. Given 64k tokens, the bases solve 6 of the 72 Z problems RL gets right in the 13 recipes, and 37 of the 84 in the three. On 31 of the 84, the base finishes all eight samples and gets every one wrong. The RL model is right in at least four of eight samples on 5 of these 31. Step checks on 32 of the Z problems RL gets right in the 13 recipes and 40 in the three cap-bound ones find the base's own samples writing RL's deciding step on 56 of the 57 where a step decides (§A.1).

Most RL lone wins at 16k are sampling. On four bases (R1-Distill-Qwen 1.5B and 7B, OpenMath-Nemotron-1.5B and

Qwen3-1.7B), with 56 extra samples, the base solves 87 of the 121 physics pairs where a recipe beat a base that left the problem unsolved in eight samples. The 87 include 12 of the 18 the RL model solves in at least four of eight. Only 3 real physics wins remain among the 121. In the crossing test, which compares RL's gain on problems the base never solves in eight samples with the base's gain on problems the RL model never solves, the net difference is -0.013 (full table in the released repository).

The residue recurs across recipes on two of the three bases with three or more recipes. The residue is the RL lone wins the base finished eight times and got wrong all eight. Shared residue problems exceed independent draws on R1-Distill-7B (10 against 5.1 [2, 8]) and OpenMath-Nemotron-1.5B (8 against 2.1 [0, 5]). On R1-Distill-1.5B they fall short (11 against 16.2 [13, 20]). One problem is residue for all eight recipes. Robust cells recur the same way. A problem every recipe on a base flips is a property of the base and the matcher, not of a recipe.

Most robust cells are answer form. Of 26 robust cells read, about three fifths are answer form, such as 0.000225 g against 0.225 mg. About a quarter are near-misses. A handful are one base's systematic errors. One recurring cell is a real change in what the policy samples. No cell is a physics fact every base lacks.

		solved in 4+ of 8 samples		solved at least once	
recipe	read at	by RL, never by base	by base, never by RL	by RL	by base
Skywork-32B	24k	11 of 111	1 of 100	18	14
Skywork-Math	32k	9 of 122	2 of 114	21	16
RLinf	28k	5 of 141	1 of 180	25	28
Skywork-7B	32k	5 of 122	0 of 102	24	15
AceReason	32k	5 of 122	2 of 129	22	20
KnowRL	24k	3 of 201	4 of 199	28	23
AceReason-14B	32k	3 of 126	0 of 118	22	12
ProRL v2	8k	3 of 34	0 of 167	4	12
DeepScaleR	24k	2 of 129	0 of 176	21	24
JustRL-N	16k	2 of 158	1 of 179	18	14
Polaris-1.7B	48k	2 of 114	1 of 70	17	11
Polaris-4B	52k	2 of 86	0 of 77	14	15
QuestA	24k	1 of 201	0 of 89	22	7
DeepMath	24k	1 of 129	1 of 170	19	14
STILL-3	28k	1 of 141	0 of 144	19	12
AReaL	24k	1 of 129	0 of 150	17	15
JustRL-D	16k	1 of 80	2 of 160	10	16
DeepMath-Omn	24k	0 of 201	2 of 205	16	17
e3	16k	0 of 90	1 of 128	10	13
Open-RS3	4k	0 of 17	0 of 19	0	0

Table 9. Robust wins by recipe, with the base's. Each count is of problems one model solves while the other, finishing all eight samples, never does. "11 of 111" means the RL model solves 11 of the 111 problems its base leaves unsolved in eight. Physics and MATH-500 are pooled at each recipe's cap, eight samples a side. ProRL v2 is read at its 8k cap and Open-RS3 at 4k, just above its cap. The JustRL pairs, trained at 15k, and e3 are read at 16k. The Polaris pairs are read on physics only.

A.5. Reach at $k = 32$

Sampled 24 more times, the bases recover their lone wins at least as often as the RL models do. Each RL model wrote 24 more samples on the physics part of U (Table 9), and each base on the mirror set (full table in the released repository). The RL models solve at least one on 23% of these problems, the bases on 27%. The lone-win counts at $k = 8$ are what sampling produces. The mirror is noise of the same size. For recipes trained at 16k and below, the base reads at a larger budget, so its recovery there is an upper bound.

A.6. A harder, unfiltered physics set

Reach stays rare on problems the filter never touched. The paper's 300 physics problems were filtered to ones the controlled run's base solves sometimes, so reach there is a lower bound. On two further sets without that limit, all twenty pairs split the same way (Table 10). Net reach at the cap is -0.006 to $+0.018$, clear of zero for seven pairs, touching zero for four and below zero for one. Bases leave two to three times more problems finished in all eight samples but never solved than on the paper's sets (Table 9). Even so, the RL model solves at least four of eight on at most 10 of them.

The sets. The *harder set* has 471 closed-form problems from OlympiadBench, PHYBench, PHYSICS's olympiad tier and FrontierScience. The *unfiltered set* has 200 random problems from the paper's pools, excluding its 300. On AceReason-Nemotron-7B the judge leaves net reach unchanged.

Far from ceiling, the split is the same. The gain is small and mostly reliability where both models finish, and finishing for recipes trained at 16k or below. Reach counts only finished samples, so QuestA and Polaris-1.7B, which finish less, could show less reach. At the next budget above the cap, where both finish more, their filtered-set net reach falls (released tables).

The two knowledge-guided recipes' robust wins hold up when read. QuestA and KnowRL have 10 and 9 robust wins on this set at their caps, on 13 distinct problems. For each win, a language-model reader compared the base's eight answers and the RL model's eight with the gold. No base answer agrees with the gold, so none of the 19 is answer form. On three wins, two of QuestA's and one of KnowRL's, the RL answers the judge credited do not agree with the gold. On three more, one credited answer is wrong, which leaves three right in eight. The other 13 stand. Without the wrongly credited answers, net reach is about 0.016 for QuestA and 0.014 for KnowRL, still the largest. Other recipes' wins were not read this way.

The filter did not hide reach. Scored by the matcher alone, net reach at the cap is -0.009 to $+0.014$ on the unfiltered 200, against -0.007 to $+0.013$ on the filtered 300 (full table in the released repository). Net reach on OlympiadBench (February 2024, before every base) is no larger than on FrontierScience (December 2025, after every model) (-0.013 to $+0.027$ against -0.005 to $+0.035$), so there is no contamination signature.

With 24 samples, about half of the lone wins are reached by the other side. Over the 19 pairs resampled, six of them only in part, the base solves 244 of 448 RL wins (54%). The RL model solves 187 of 289 base wins (65%). After both are resampled, the RL model solves 162 of its wins again in its new samples while the base still fails them, and the base does the same on 70 of its own, a net of 92, about five per pair.

recipe	cap	fin@1(B), base→RL	F	P	net reach	robust net	RL solves 4+ of 8, never by base
ProRL v2	8k	0.32 → 0.99	**0.137**	**−0.025**	−0.004	−0.003	1 of 11
JustRL-D	16k	0.75 → 0.95	**0.044**	**0.013**	**−0.006**	−0.001	0 of 105
DeepScaleR	24k	0.92 → 0.97	**0.016**	**0.015**	−0.002	0	0 of 256
DeepMath	24k	0.92 → 0.98	**0.018**	**0.028**	0.001	0	0 of 256
RLinf	28k	0.94 → 0.98	**0.015**	**0.026**	0.005	0.001	2 of 284
STILL-3	28k	0.94 → 0.95	**0.006**	0.006	−0.001	0	0 of 284
AReaL	24k	0.92 → 0.96	**0.011**	0.004	0.001	0	0 of 256
Open-RS3	4k	0.12 → 0.12	0	0.001	0	0	0 of 1
JustRL-N	16k	0.89 → 0.95	**0.010**	**0.009**	**0.007**	0.004	3 of 286
QuestA	24k	0.96 → 0.79	**−0.011**	**0.023**	**0.018**	0.015	10 of 414
KnowRL	24k	0.96 → 0.96	**0.005**	**0.014**	**0.015**	0.013	9 of 414
DeepMath-Omn	24k	0.96 → 0.98	**0.005**	0.004	**0.004**	0.001	1 of 414
e3	16k	0.76 → 0.87	**0.037**	**−0.018**	−0.005	−0.001	2 of 181
Polaris-1.7B	48k	0.95 → 0.86	0.003	**0.019**	0.005	0.001	2 of 378
Polaris-4B	52k	0.99 → 0.89	**−0.028**	**0.039**	**0.008**	0.003	2 of 315
Skywork-7B	32k	0.98 → 0.94	**−0.008**	**0.026**	**0.009**	0.006	5 of 310
AceReason-14B	32k	0.98 → 0.98	**0.005**	**0.019**	0.006	0.001	4 of 287
Skywork-32B	24k	0.97 → 0.91	**−0.011**	**0.045**	**0.011**	0.006	5 of 240
Skywork-Math	32k	0.98 → 0.98	**0.004**	**0.026**	0.006	0.007	7 of 310
AceReason	32k	0.98 → 1.00	**0.009**	**0.027**	0.006	0.006	5 of 310

Table 10. The harder, unfiltered set at each recipe's cap. The set has 471 harder problems and 200 drawn from the paper's pools without the base filter. Each model writes eight samples per problem, scored by the fixed judge of Appendix C.2 with the deterministic checker and OlympiadBench's per-problem tolerances. Robust reach counts wins in at least four of eight, net of the mirror. **Bold**: the paired bootstrap interval over problems excludes zero. The last column counts the problems the base finishes eight times and never solves that the RL model solves in at least four of eight. Every interval is in the released tables.

A.7. A competition-math set

Reach stays rare in domain. Physics is out of domain for every recipe that states its mix except ProRL v2, and MATH-500 is near ceiling, so we read every pair at its cap on 276 MathArena final-answer problems (Table 11). Of these, 213 are from 2025 and 63 from AIME and HMMT February 2026, after every model. Net reach is -0.002 to $+0.021$, clear of zero for nine pairs and touching or straddling zero for eleven. The judge moves net reach by at most $+0.001$.

Registered predictions. Before any output existed we registered five. P1: every base finishes less than on physics, and finishing leads for the four cap-bound recipes. P2: reliability rises clear of zero, leading in most, in at least three quarters of recipes whose base finishes 90% or more. P3: net reach within $[-0.02, +0.03]$, clear of zero for at most half the pairs, robust reach at most 0.02. P4: no more reach on 2026 than on 2025. P5: AIME 2025 pass@1 within 0.08 of the authors' comparable numbers. P2 to P5 held. P1 failed on Open-RS3, where neither model finishes more than 6% at 4k, leaving too few finished samples to split.

Far from finishing, which term leads depends on the estimator. For the five recipes on R1-Distill-1.5B, whose base finishes 0.74 to 0.80, finishing carries more of the gain than on physics. Under the matched split, reliability is larger for four of the five. The matched finishing term is a remainder that also carries accuracy changes outside the matched set.

Reliability leads where the base finishes. In all nine recipes whose base finishes 90% or more, reliability rises clear of zero and leads under either split. On matched problems, where selection cannot act, pass@1$(B \mid \text{fin})$ still rises clear of zero for every such base. For QuestA and KnowRL it rises $+0.109$ and $+0.129$, against $+0.245$ and $+0.188$ over all. Selection therefore accounts for part of each gain. A gain clear of zero remains.

Contamination. P4, a test of base contamination, holds for every pair (full table in the released repository). More reach on 2025, the sign of recipe contamination, is clear of zero only for STILL-3 and Skywork-OR1-7B. With only 11 and 13 finished-but-unsolved problems in their 2026 strata, contamination and chance cannot be told apart.

The harness reproduces the authors' AIME 2025 numbers. Where authors report pass@1 at our temperature and within one budget step of the cap, our read is within 0.040 for all five models. The other gaps (-0.160 to -0.011) are budget or temperature effects.

With 24 samples the base solves about three fifths of the RL model's lone wins. The base solves 99 of 159 (62%). The RL model solves 31 of the 60 base wins. After both are resampled, the RL model solves 44 of its wins again in its new samples while the base still fails them, and the base does the same on 17 of its own, a net of 27, about 1.4 per pair.

recipe	cap	fin@1(B), base→RL	F	P	matched P	net reach	robust net	RL solves 4+ of 8, never by base
ProRL v2	8k	$0.20 \to 0.89$	**0.227**	**−0.055**	0.006	−0.001	0	0 of 0
JustRL-D	16k	$0.49 \to 0.87$	**0.140**	−0.004	**0.046**	0	0	0 of 8
DeepScaleR	24k	$0.74 \to 0.97$	**0.070**	0.011	**0.043**	0	0	1 of 31
DeepMath	24k	$0.74 \to 0.96$	**0.070**	0.011	**0.050**	0.005	0.004	1 of 31
RLinf	28k	$0.80 \to 0.96$	**0.057**	**0.045**	**0.069**	0.007	0.014	4 of 44
STILL-3	28k	$0.80 \to 0.89$	**0.026**	−0.006	−0.001	**0.008**	0.007	2 of 44
AReaL	24k	$0.74 \to 0.81$	**0.027**	**0.020**	**0.030**	0.004	0.004	1 of 31
Open-RS3	4k	$0.06 \to 0.06$	0	0.002	0	0	0	0 of 0
JustRL-N	16k	$0.66 \to 0.66$	0.003	**0.071**	**0.035**	0.003	0.004	1 of 15
QuestA	24k	$0.93 \to 0.63$	**−0.171**	**0.189**	**0.087**	**0.009**	0.004	1 of 56
KnowRL	24k	$0.93 \to 0.82$	**−0.054**	**0.150**	**0.102**	**0.012**	0.007	3 of 56
DeepMath-Omn	24k	$0.93 \to 0.92$	−0.003	**0.049**	**0.045**	−0.002	0	0 of 56
e3	16k	$0.43 \to 0.62$	**0.099**	**−0.049**	0.015	−0.002	−0.004	0 of 5
Polaris-1.7B	48k	$0.98 \to 0.64$	**−0.117**	**0.167**	**0.089**	**0.012**	0	1 of 99
Polaris-4B	52k	$0.98 \to 0.77$	**−0.105**	**0.139**	**0.050**	**0.014**	0.004	1 of 64
Skywork-7B	32k	$0.91 \to 0.80$	**−0.050**	**0.117**	**0.093**	**0.005**	0	0 of 54
AceReason-14B	32k	$0.94 \to 0.91$	−0.007	**0.121**	**0.101**	**0.021**	0.018	5 of 49
Skywork-32B	24k	$0.88 \to 0.83$	−0.017	**0.125**	**0.087**	**0.008**	0.004	1 of 25
Skywork-Math	32k	$0.91 \to 0.91$	0.010	**0.084**	**0.085**	**0.005**	0	0 of 54
AceReason	32k	$0.91 \to 1.00$	**0.050**	**0.061**	**0.074**	0.007	0.007	2 of 54

Table 11. The competition-math read at each recipe's cap. The set has 276 final-answer problems from eight 2025 and 2026 contests. Each model writes eight samples per problem, scored by the fixed judge of Appendix C.2 after the exact-match checker. Matched P is reliability on the problems both models finish at least seven times in eight, where selection into the finished set cannot act, weighted as P is. Matched reliability is the per-pair version of the matched split. Robust reach counts wins in at least four of eight, net of the mirror. **Bold**: the paired bootstrap interval over problems excludes zero. The last column counts the problems the base finishes eight times and never solves that the RL model solves in at least four of eight. Every interval is in the released tables.

B. The controlled run

Privileged information for hard groups. Partial solutions (Li et al., 2025; Amani et al., 2025), answer-derived hints (Chen et al., 2025a), scaffolds (Zhang et al., 2025c; Wang et al., 2025c), reverse curricula (Xi et al., 2024), hint reweighting (Agrawal et al., 2026), expert traces (Zhou et al., 2025; Yan et al., 2025) and trained hinters (Xia et al., 2026; Han et al., 2026; Liao et al., 2026) report the deployed score, not whether content was learned. These methods do not control finishing. Our channels belong to these families and claim no algorithmic novelty. Process rewards (Wang et al., 2023; Chen et al., 2026; Xie et al., 2026; Zhang et al., 2025b) need a target step, which on all-fail problems only a reference solution gives. Answer-conditioned rationales (Zelikman et al., 2022; Phan et al., 2023) and hindsight relabelling (Zhang et al., 2023) leak the reward. We supply a setup, not an answer.

B.1. The probe in full

The run. The *control* is Qwen3-VL-8B-Thinking (Bai et al., 2025) trained with outcome RL (GSPO; Zheng et al., 2025; Shao et al., 2024) on PhysReason (Zhang et al., 2025d) and UGPhysics (Xu et al., 2025; Yang, 2026) at a 12k-token cap. Each *arm* adds a *channel* supplying the *setup*, the physical modelling before any algebra: a reward, the prompt, hinted derivations, a policy-written setup, or a stronger model's questions.

Two hard sets. The *held-out hard set* is 87 PhysReason problems the base solves at most once in eight, never trained on, frozen before any arm under one registered protocol. A fresh base read scores 0.16, above what selection implies, so arms are compared with the control. The *training hard set* is 160 problems the control never solved in its first thirty steps. As in §4, the control's held-out score rises from 0.16 to 0.26 almost all by finishing (fin@1(B) 0.39 → 0.69; pass@1(B | fin) flat at 0.38).

The missing fact is often in the model, which does not ask for it. The right question, without the answer, nearly triples the training hard set's share solved at least in part in eight samples (0.062 → 0.175; gold setup 0.350). A stronger model's questions written without the reference setup give 0.200, so the gain is not only leakage. At 64 samples, 41% is solved only with help and 34% under no condition we tried. The model's own questions do no better than a generic checklist.

No channel raises the held-out score clear of zero. Only a pilot of the reward gave a valid read, since at full scale the policy gamed its scorer (§B.2; Appendices B.6–B.10). The setup hint leads the two control seeds by +0.009 and +0.019, below the detection floor (Table 13). Hinted derivations move it −0.05, beyond the seed gap. The policy-written setup moves it −0.03, near the detection floor. A +2.4-point hint-arm lead on an earlier corpus did not hold at the 16k read. The held-out problems mostly carry figures but are given as text.

Content is detected only where trained. The control lifts the training hard set's pass@8 from 0.056 to 0.344 (§B.2). On the held-out set, drawn from another source, the control's gain is almost all finishing. On the 160 easier text-only twins of its training problems it gains mostly by finishing. Content transfer is neither shown nor excluded (Appendix B.13). The channels' failure may therefore not be specific to them. Whether recurring techniques explain why guidance gains on competition mathematics is untested here.

B.2. Why the knowledge is not installed

The policy rarely samples a success. The base solves 24% of the training hard set in 64 samples (44% at a 16k cap), still rising (Figure 5, left). At a few percent per sample, a group of sixteen sees a success less than half the time. At least a quarter are within reach but need a long, improbable derivation. Gains follow a $p(1-p)$ shape: +0.16 to +0.18 next visit when solved one to six times in eight, +0.05 and +0.02 at the ends.

Finishing carries because every prompt shares it. Finishing explains 48% of within-group reward variance early and 21% once saturated. Finishing's gradient adds up across the batch. Problem-specific content is not detected to carry.

Each channel moves only what the policy already samples. At full scale the setup reward taught unit-appending for the answer matcher, a gain that was not real. Derivations hinted with the setup lean heavily on the hint (400–580 nats; Appendix B.11), so the bound of Xia et al. (2026) gives no reason to expect transfer. A teacher's derivation lifts touched problems from 12% to 59%, not their neighbours. Self-written setups succeed equally often (intraclass correlation 0.03), so rewarding them changes only format. Questioner samples that ask solve no more often.

Once a success is sampled, credit installs it, but only on that problem. On 60 of 353 hard training problems, read after

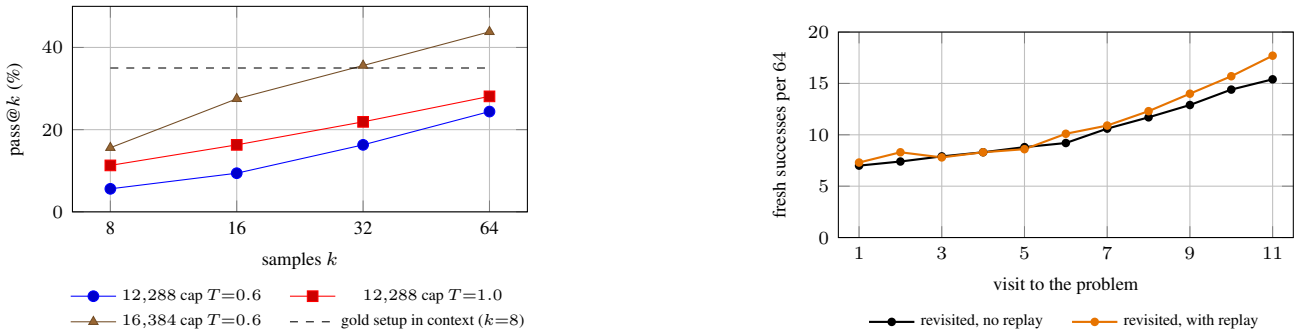

Figure 5. Left: the panel plots pass@k of the base model against sample count, on 160 training problems never solved in the control's first thirty steps. The curves are the registered decoding, a hotter sampler and a longer cap. The plotted points are read off the first k samples, not the unbiased estimates quoted in the text. The dashed line is the gold setup in context at eight samples. Right: the panel plots fresh-rollout successes per 64 on 60 hard training problems solved at least once in 32 samples at 16k. The count is taken visit by visit during training, over the first thirty updates of each run. The curves are with and without the policy's own earlier successes replayed into the group (§B.2). Replayed rows are excluded from the count.

training at 16k with eight samples, the solve rate goes from 0.10 to 0.28, or 0.31 with replay (Zhang et al., 2025a). Replay only shortens the wait (Figure 5, right). The held-out set and the 293 unrevisited problems do not move. These problems are rarely solved, so spread to sometimes-solved problems is not excluded. Twins bear on it weakly: on the 16 of the 60 with twins, the revisit arm does no better than the control (Appendix B.13). The held-out bottleneck is the first success.

B.3. Setting details

Setting. The policy is Qwen3-VL-8B-Thinking (Bai et al., 2025), given problems as text and trained with GSPO (Zheng et al., 2025) under verl (Sheng et al., 2024). Table 12 lists the arms. The reward is the fraction of gold answers matched by a boxed answer under one frozen, blind-audited matcher. The *old corpus* mixes UGPhysics (Xu et al., 2025; Yang, 2026) and PhysReason (Zhang et al., 2025d) items. Every row of *corpus v2* is solvable at least sometimes by the base.

Setup claims. The 7,398 claims are mined from answer-conditioned derivations. A blind audit found 73.8% strict precision. About 1% still leak a gold answer. A blind panel judged 149 of 150 real claims correct.

arm	channel	corpus	what differs from its control	steps × groups
outcome RL (seeds 101, 202)	—	old	mined-milestone term, inert (4.7% mass); the comparator	60 × 96
setup-reward arm (arm F)	reward	old	judged setup-waypoint reward on all-wrong groups	60 × 8
hint arm (arm H)	prompt	old	setup block, $\Pr(\text{hint}) = 1 - s/59$	60 × 96
corpus v2 control (two seeds)	—	v2	none; seed 2 changes rollout seed and data order only	60 × 96
hint on v2	prompt	v2	the hint arm's block, matched rows and draws	60 × 96
generic block	prompt	v2	content-free block, same draws	60 × 96
guided-splice arm (arm I)	loss	old	frozen base's hinted derivations spliced into dead groups	stopped at 4
teacher arm (arm J)	loss	v2	online hinted teacher spliced into dead student groups	60 × 96
proposer arm (arm K)	produce	v2	self-written setup block rewarded by a frozen executor's score	stopped at 11
setup-first arm v2 / v3 (arm M)	produce	v2 (+ claims)	self-written setup section judged against the reference (lenient / strict judge)	30 × 96 each
deep-sampling arm (arm O)	sampling	v2 hard stratum	24 × 64 from the control's step 60, 12k cap (arm O16: 16k cap, reachable stratum)	10 / 15 × 24
dense-revisit arm (arm D16)	sampling	60 learnable	24 × 64, every problem every 2.5 steps	8 × 24; matched rerun 30 updates
replay arm (arm R)	replay	60 learnable	dense revisit + the policy's own successes replayed into groups	34 × 24
wide replay arm (arm RW)	replay	v2	the control + replay (r=4 of 16)	paused at 11
questioner fine-tune (arm S)	distil	595 problems	supervised fine-tune on 210 own derivations opening with the teacher's questions (+ a no-questions control)	SFT
16k-cap arm (arm B16)	budget	v2	response cap 16,384	40 × 96
dynamic-pool arm (arm P)	sampling	v2 + reserve	sampler retiring saturated and dead problems	60 × 96

Table 12. Trained arms. Groups are prompts per step. Every arm samples 16 rollouts per prompt except the setup-reward arm (8). One seed per arm except the control and the comparator. The letters are the names the appendices use. "Learnable" problems are the 60 of the 353-problem text-only hard stratum with at least one full solve in 32 samples at 16k.

The registered endpoint. The held-out hard set is PhysReason problems never used for training, mining or selection, frozen before training. The *wide* set (136 parts) has the base right in at most one of eight samples, the *strict* subset (84) in none.

The protocol takes eight samples at temperature 0.6, capped at $12{,}288$ tokens. The *16k* variant raises the cap to $16{,}384$. The *forced* variant makes every truncated derivation answer. Greedy decoding loops on this model, and 17% of items flip between runs, so it is never used. Repeated reads vary by about ± 0.01. We do not read differences inside that band.

Benchmarks and readout. The benchmarks are PhysReason held-out, PhysUniBench-OE (Wang et al., 2025a), OlympiadBench physics (He et al., 2024) and PhysOlym-A (Yang, 2026), judged as in Appendix C.2. The setup readout labels each mined claim in a rollout established, contradicted or absent (Appendix C.1).

B.3.1. The bottleneck is the setup, and it can be supplied

Setup errors are the largest group of failures. Of 120 typed misses, 41.7% are setup or concept and 1.7% computation. Of the model's modelling decisions on 90 held-out parts, 36.8% are wrong.

The setup can be supplied. The share of 116 problems with any of four samples scoring is 0.345 unaided and 0.578 with the mined setup. Another problem's claims add only $+0.017$ $[-0.034, +0.069]$, so 90% of the lift is content. A blind grader never shown the gold separates correct from incorrect samples at AUC 0.766.

B.4. The control gains by finishing, and the training cap sets how much

The control gains by finishing more (full table in the released repository). On problems finished at both the base and step 60, pass@1$(B \mid \text{fin})$ is flat (Table 13). The finished set grows to include harder problems, so new finishes are right as often as old ones.

The training cap sets how much finishing is learned. The 16k-trained arm learns less finishing and gains a little matched pass@1$(B \mid \text{fin})$ instead (Table 13). Read back at 12k it falls below the control (0.202 against 0.242). Caveats: one run each, the 16k arm stopped at step 40, and the control's step-40 weights were not kept, so it is read at step 40 only from its 12k validation pass during training.

B.5. Statistical power

No claim rests on a difference smaller than about 0.03. Table 13's intervals resample problems only and leave out run-to-run variance. Adding the variance between the two control seeds (full table in the released repository) puts the minimum detectable effect near 0.03. Against the second seed the setup-first deficits sit at the edge of the seed gap. The 16k-cap arm's lead reads each policy at its own training cap, so it is not a same-cap contrast. Only the hint and teacher channels ran at full scale.

B.6. Channel A: reward the setup

The comparator's milestone reward is inert. The milestone reward's dense term carries 4.7% of advantage mass, flips no sign and is zero 96.7% of the time.

Design and result. All-wrong rollouts are rewarded with a frozen judge's probability that they commit to the setup claims. At one twelfth of the control's scale every trend is flat. The judge's AUC there is 0.655. Signal-to-noise is 0.65 against 2.26.

Caveat: at full scale the scorer was gamed. The rerun learned to append the units the answer matcher required, a $+99\%$ apparent gain and no real one. The rerun's absolute figures cannot be recovered. A powered rerun with a valid scorer was never done.

B.7. Channel B: supply the setup in the prompt

The hint helps in context and is not kept. The setup block appears with probability $1 - s/59$ at step s, so the final policy trains without it. Hinted rollouts lead by $+0.13$ early and do not lead late. On the 59 problems the hint revives, a hinted correct rollout establishes 62 of 68 contested claims and a late unhinted one 12 of 57.

On matched data the hint shows no gain. The hint arm is $+0.09$ $[-1.38, +1.75]$ over its control (full table in the released repository). The setup hint beats a content-free checklist of the same length only on PhysUniBench-OE ($+3.9$ $[+1.1, +6.6]$). The PhysOlym-A lead is therefore prompt scaffolding.

paired contrast (problem-level mean score, 87 problems)	diff	95% CI	$\Pr(\leq 0)$
control s60 − base (12k)	$+0.098$	$[+0.059, +0.142]$	0.000
hint on v2 s60 − control s60 (12k)	$+0.009$	$[-0.014, +0.032]$	0.232
content-free block s60 − control s60 (12k)	-0.003	$[-0.025, +0.018]$	0.600
teacher arm s60 − control s60 (12k)	-0.052	$[-0.076, -0.030]$	1.000
setup-first (lenient judge) s30 − control s30 (12k)	-0.029	$[-0.054, -0.007]$	0.995
setup-first (strict judge) s30 − control s30 (12k)	-0.036	$[-0.060, -0.015]$	1.000
16k-cap arm s40 (16k) − control s40 (12k)	$+0.038$	$[+0.014, +0.063]$	0.001
deep sampling s15 − s0 (16k, from control s60)	$+0.003$	$[-0.013, +0.019]$	0.340
replay s30 − s0 (16k, from control s60)	-0.007	$[-0.026, +0.012]$	0.766
the same arms against the second control seed			
hint on v2 s60 − control seed 2 s60 (12k)	$+0.019$	$[-0.001, +0.039]$	0.033
content-free block s60 − control seed 2 s60 (12k)	$+0.007$	$[-0.014, +0.027]$	0.243
teacher arm s60 − control seed 2 s60 (12k)	-0.041	$[-0.068, -0.018]$	1.000
setup-first v2 s30 − control seed 2 s30 (12k)	-0.021	$[-0.042, -0.000]$	0.977
setup-first v3 s30 − control seed 2 s30 (12k)	-0.027	$[-0.046, -0.010]$	0.999
matched pass@1($B \mid$ fin), finished at both reads			
control s60 − base (49 problems)	-0.000	$[-0.032, +0.034]$	
control seed 2 s60 − base (48 problems)	$+0.003$	$[-0.033, +0.041]$	
16k-cap arm s40 − base (68 problems)	$+0.035$	$[+0.007, +0.064]$	

Table 13. Clustered paired bootstrap (20,000 draws) over held-out problems. Repeated reads were taken where held-out was flat. Deep sampling at 16k has s.d. 0.003 over four reads. Replay at 16k has s.d. 0.008 over seven. The matched rows read pass@1($B \mid$ fin) on problems with a finished sample at both reads.

channel (scale)	how the setup is given	inside training	on held-out problems
reward the setup (pilot)	a judge rewards a stated setup on groups where every sample fails	no effect at the pilot's scale; at full scale the policy gamed the scorer	not tested with a valid scorer
hint in the prompt (full)	the setup written into the prompt, faded out over training	the policy uses it: setup claims stated $55 \to 92\%$	not kept once the hint is gone; a content-free block does as well
teacher derivations (full)	correct hinted derivations spliced into groups that failed	learned on the problems it touched ($12 \to 59\%$ solved)	below the control (-0.05); no gain on similar problems
policy writes the setup (half)	the policy writes a setup section, which is judged	the sections become correct (strict content $0.29 \to 0.45$)	below the control (-0.03); a section does not change which sample succeeds
fine-tune on questions (supervised)	210 derivations that open with a stronger model's questions	the habit is learned (99% open with questions)	no better than the same derivations without the questions; held-out at the base

Table 14. Five ways of supplying the setup, and what each does. Full scale is the control's sixty steps of 96 problems. Every comparison is against the control on the same problems or on the frozen held-out set.

B.8. Channel C: distil the hinted derivation

Design. When a student group is all wrong, up to $k{=}4$ correct hinted teacher rollouts replace its last rows, rescored under the unhinted prompt. There is no KL anchor or off-context correction.

Channel C teaches the problems it touches. Early-spliced problems go from 12% to 59% fully solved, against 17% to 31% for never-spliced ones. The setup readout does not change.

Channel C ends below the control, and untouched problems get worse. Forced to answer on the held-out hard set, it scores 6.2 against the control's 8.1, a gap of 1.9 [−0.2, +4.0] in the control's favour. On 371 untouched training problems the delta is −5.1 [−6.8, −3.5], flat across similarity. Outputs run about 15% longer. One seed.

B.9. Channel D: make the policy write the setup

A self-written block shows no gain. Following TraceLift (Han et al., 2026), the policy's blocks give a frozen copy −0.008 [−0.054, +0.040] over no block, against +0.050 [−0.009, +0.110] for the gold block. A block the policy can write is one the frozen copy already assumed.

Rewarding a written setup section changes the section, not accuracy. A lenient judge was gamed (75% of claims marked established against a strict grader's 34%). Under a strict judge held-out hard-set accuracy is 0.202 against the control's 0.238 at 30 updates. Within a group, a more complete section does not solve more often (+3.0 [−2.8, +5.8]). Most wrong derivations (80%) used every reference claim.

B.10. Channel E: distil the questioner

Design. We fine-tuned on 210 correct base derivations, each opened by setup questions Claude Sonnet 5 wrote. The control fine-tune drops the questions.

The habit is learned and inert. The model asks in 99.5% of samples. Asking samples solve 2.0% of the time, others 2.6%. The gate for running RL required training hard-set pass@8 $\geq$ 0.17 and held-out at or above the base. The fine-tune fails both, so no RL was run from it.

B.11. Reliance, corpus and the hard sets

Correct hinted derivations lean too heavily on the hint for the transfer bound to bite. The law $p \geq p_h \exp(-\rho_c)$ bounds unhinted success by hinted success and the reliance ρ_c of correct hinted derivations. Reliance is 400–580 nats for the setup hint and 530–600 for spliced teacher derivations. The teacher arm's deficit carries the missing off-context correction as a caveat.

The two hard sets come from different sources. The training hard set is text-only and UGPhysics-derived (156 of 160). The held-out set is PhysReason. Testing recurring content would need a held-out split where it recurs, or text-only twins. The twins were read (Appendix B.13). The split was not.

More samples, a longer cap and a hotter sampler all raise pass@k on the training hard set (Figure 5). Each doubling of k multiplies pass@k by 1.67, 1.53 and 1.39. The longer cap lets more derivations finish (0.34 → 0.53).

B.12. Replay

Replay refills failed rows with the policy's own past successes (Table 15). Replay runs on 60 hard problems, against a matched no-replay rerun.

fresh-rollout successes per 64	visit 1	2	3	6	9	11
revisited, no replay (30 updates)	7.0	7.4	7.9	9.2	12.9	15.4
revisited, with replay (30 updates)	7.3	8.3	7.8	10.1	14.0	17.7

Table 15. Revisiting 60 hard problems, with and without replay. Each cell counts successes among 64 fresh samples at one visit during training. Replayed samples are excluded.

New successes are mostly fresh derivations. A late success's nearest replayed row has median shingle Jaccard 0.039, against 0.034 for its nearest early success. The held-out hard set reads 0.329 without replay and 0.325 with it, against the control's 0.324.

B.13. Text-only twins of the training hard set

Design. Each of the 160 training hard-set problems has a twin (same physics and method, new surface) whose gold two blind solvers reproduce. The *credited* problems are the 75 where the control's first seed earned reward in steps 31–60. The committed plan compares credited and other twins' gains over the base in mean score and any-solve@8. The split below, with its matched contrast as the content-transfer test, was chosen after a first look at 113 of the 160 problems.

On the credited originals, finishing's share of the gain depends on the order of the split. Finishing and reliability both rise there (Table 16), so the term they make together is large. Eq. (2) gives that term to finishing, which then makes up 71% of the control's gain (73% for the second seed). Splitting the term evenly gives finishing 50% (51%). Giving it to reliability leaves finishing 29% (also 29%).

On the planned measure, the credited twins gain more, 26–62% of it by finishing. The mean-score excess is $+0.096$ [$+0.008$, $+0.182$], and $+0.115$ [$+0.029$, $+0.202$] for the second seed, whose credit does not define the groups. The any-solve@8 excess, $+0.095$ [-0.054, $+0.242$], includes zero. Of the pass@1 excess, $+0.093$ [$+0.005$, $+0.181$], finishing makes up 0.058 under Eq. (2), 0.048 with the cross term split evenly and 0.038 with it given to reliability. The rest is $+0.035$ [-0.011, $+0.078$], $+0.045$ [-0.021, $+0.108$] and $+0.055$ [-0.032, $+0.140$] under the three orders. For the second seed, finishing makes up 56, 41 and 26% of $+0.110$, and the rest is $+0.048$ [$+0.004$, $+0.090$], $+0.065$ [-0.001, $+0.128$] and $+0.081$ [-0.007, $+0.169$]. Mean score per finished sample falls on both groups. Regrouping by the second seed's credit is a weak check, since the credited sets share 53 of 74 problems.

The split of Eqs. (1) and (4) finds no content gain on the twins, and cannot exclude one (Table 16). Where both models finish at least seven times in eight, the credited-minus-other reliability contrast is -0.003 [-0.048, $+0.037$] (second seed: -0.008 [-0.043, $+0.021$]). That interval is as wide as the originals' contrast, $+0.039$ [$+0.006$, $+0.081$], on only 5 and 10 problems.

A never-credited model and an adjusted contrast leave the excess unresolved. The questioner fine-tune (E), never trained on these problems, shows a credited-minus-other twin difference of $+0.035$ [-0.020, $+0.091$] from the problems alone. Finishing little more than the base, E is not a clean control. The control's difference less E's is $+0.061$ [-0.029, $+0.149$]. Adjusting for the base's finishing and score gives $+0.077$ [-0.004, $+0.160$] on the twins, against $+0.186$ [$+0.132$, $+0.237$] on the originals. Neither check is clear of zero. Both checks were added after the full read.

The twins are easier than their originals (Table 16), partly perhaps from answer form, since only twins whose gold the reward can match were kept. With the base finishing under half the twins, finishing dominates (§4.1). Content needed only on the hardest problems would not show.

No channel beats the control on the twins (full table in the released repository). On credited twins the hint (B) is level, the revisit arm (§B.2) within noise, and the teacher's derivation (C), D and E lower. C's interval includes zero. Channels are compared on all twins, not only those of problems they touched.

In the twins read, the base solves 16 of the 160 originals at least once, against 9 in the earlier read. At such low rates the solved set varies by read.

control	problems	n	pass@1 base→control	fin@1(B) base→control	pass@1(B | fin) base→control	finishing F	reliability P	matched P (n)	net reach
seed 1	credited originals	75	0.02 → 0.21	0.30 → 0.75	0.08 → 0.28	0.131	0.034 [0.017, 0.055]	0.033 [0, 0.074] (5)	0.018 [0.003, 0.037]
	credited twins	75	0.28 → 0.55	0.37 → 0.76	0.77 → 0.72	0.288	−0.020 [−0.045, 0.006]	0.014 [−0.002, 0.032] (15)	−0.002 [−0.005, 0]
	other originals	85	0.01 → 0.02	0.38 → 0.75	0.02 → 0.03	0.013	0.007 [0.001, 0.014]	−0.006 [−0.016, 0.001] (10)	−0.006 [−0.015, 0.001]
	other twins	85	0.36 → 0.54	0.45 → 0.78	0.81 → 0.69	0.231	−0.053 [−0.084, −0.016]	0.017 [−0.018, 0.060] (27)	−0.004 [−0.012, 0]
seed 2	credited originals	75	0.02 → 0.23	0.30 → 0.81	0.08 → 0.29	0.154	0.036 [0.017, 0.058]	0.023 [−0.004, 0.081] (5)	0.022 [0.002, 0.055]
	credited twins	75	0.28 → 0.57	0.37 → 0.78	0.77 → 0.73	0.304	−0.019 [−0.046, 0.009]	0.009 [−0.005, 0.026] (15)	0 [0, 0]
	other originals	85	0.01 → 0.04	0.38 → 0.77	0.02 → 0.05	0.024	0.007 [−0.001, 0.016]	0 [−0.010, 0.010] (10)	0.001 [−0.003, 0.006]
	other twins	85	0.36 → 0.54	0.45 → 0.81	0.81 → 0.67	0.243	−0.059 [−0.089, −0.024]	0.017 [−0.006, 0.049] (27)	−0.009 [−0.025, 0]

Table 16. Text-only twins: each control seed against the base on the training hard set's credited and other problems and on their text-only twins. Each read takes eight samples at a 12,288-token cap, with full credit as success. The change in pass@1 splits into finishing F, reliability P and net reach as in Eq. (4). Matched P is reliability on the n problems both models finish at least seven times in eight, weighted as P is. The groups are defined by the first seed's credit in steps 31–60. Splitting the twins' change this way was chosen after a first look at 113 of the 160 problems.

B.14. Why the finishing signal transfers

Credit splits into a finishing term and a content term. Write rewards $r_j = f_j c_j$ (f_j: finished inside the budget; c_j: correct once finished). The group-relative step is proportional to $\mathrm{Cov}(r_j, \nabla \log \pi(y_j))$, which scales with $p(1-p)$. By the law of total covariance,

$$\mathrm{Cov}(r, \nabla \log \pi) = \mathrm{Cov}(f\, \mathbb{E}[c \mid f], \nabla \log \pi) + \mathbb{E}[f\, \mathrm{Cov}(c, \nabla \log \pi \mid f)].$$

Finishing is shared by every prompt, so its term adds across a batch and carries to held-out prompts. Content adds only where a problem recurs, so it installs success on its own problem, as replay shows.

C. Instruments

C.1. Setup readout rubric

Each item is a problem, its mined setup claims and one trimmed rollout. The grader (Claude Sonnet 5) marks each claim ESTABLISHED, CONTRADICTED (an incompatible setup) or ABSENT, judging only the setup, not the arithmetic or the answer.

C.2. The judge protocol

Protocol v2 fixes the matching rules in code. Graders disagree mostly on whether a bare number matches a gold with a unit. Giving units tracks training: models trained with and without the benchmark's split give units 97.2% and 64.0% of the time. A deterministic checker decides 43–55% of items. A model given the same rules decides the rest.

Differences between systems are grader-invariant; absolute scores are not. Regrading six systems on PhysReason held-out (n=526) under v2 with a pinned Qwen3-32B is 6.7–8.9 points stricter on every system. The margin between systems agrees to 0.11 points.