A Harness Reducing over-Defensive Writing in Agents by Isolating Objections

Xiaochen Chang

arXiv:2610.36717v1 [cs.CL] 27 Oct 2026

Abstract

When expanding a report, writing with agents may introduce a stronger claim than the article makes and then disclaim it. This is called over-defensive writing. We assume that this is because the agent is writing for responding to an adversarial AI reviewer, instead of a human reader. The defensive context may also influnce the consequent when we simply ban defensive writing in the output through prompts. In order to reduce over-defensive writing, we produce a harness that separates the reviewer, filter, and writer. The writer agent receives only the problems choosen by the filter, which is an independent context. We evaluate this harness on 25 business-project drafts expanded into formal reports, the result shows that the defensive propositions decreased from 59 to 2, which is a 96.61% reduction, no unsupported factual additions were identified while the source-claim retention rises from 80.35% to 93.12%.

Keywords: defensive writing; context isolation; harness

1 Introduction

While writing with agents, we observe that the model may sometimes introduces a stronger claim that the draft never makes, then states that the author does not commit to it. For example, the draft in Table 1 uses an AI assistant to note the changes of delivery time. The writing model however, added that this does not mean that the objectives have been achieved. The reader needed an explanation of how the delivery time changes, but the report instead began answering a new question about objective achivement.

Necessary clarification helps readers understand a conclusion's scope, while over-defensive writing adds further exclusions where that scope is already clear and the task has not raised the corresponding issue. It occupies space that could explain the project and redirects statements of contribution toward preventing hypothetical misreadings. Hedges are part of how authors negotiate claims

Source	Reduce average processing time from 14 days to five by December 31.
Expansion	Delivering the technology alone **would not demonstrate that the project's objectives have been achieved.**

Table 1: A source objective and an added qualification. The expansion shifts from a processing-time target to denying that technology delivery establishes achievement of the project's objectives.

with readers: Hyland (1996, 2005) situates hedging, stance, and reader engagement within this interaction. Evaluating a qualification therefore requires identifying what it answers in the current argument.

Liao et al. (2026) provide direct experimental evidence about defensive writing in GPT models. Requests to check claims increase additional qualifications, and explicit requests to address reviewer concerns increase them further. They find that many added qualifications tell a careful reader nothing new. This moves the problem from sentence structure to the selection of writing objectives: once anticipating criticism becomes an improvement goal, adding scope statements will become a part of completing the assignment. We observed a report-writing run in which the writer agent created a reviewer subagent and asked it to identify several points vulnerable to criticism, although the user had not requested such a review. The returned concerns then appeared in the report as defensive content. In another case without an explicitly created reviewer subagent, if we review the original draft with an agent, we will find the produced concerns will corresponding to every over-defensive expression in the output.

Telling a model to avoid defensive writing reduces this behavior but leaves residual cases. Existing anti-defensive-writing skills include necessity checks, removal of redundant disclaimers, and conversion of negative limitations into affirmative scope statements (Kiterlin, n.d.). These operations works only on the style of outputs, while leav-

ing the motivation of responding an adversarial revierwr agent. A qualification could also remain in context in a different sentence form. Table 2 illustrates this transition.

Source	Reduce average processing time from 14 days to five by December 31.
Direct expansion	Delivering the technology alone would not demonstrate that the project's objectives have been achieved.
Avoid defensive writing	Delivery should therefore be assessed against operational improvements as well as completion of the technology, with employee readiness forming an essential part of the transition.

Table 2: The same writing case after an output prohibition.

A direct response is to tell the writer which objections are unlikely to matter. In our tests, however, this will instead increase defense. We first asked a review agent to estimate the probability that a normal human reader would raise each concern, then supplied all concerns and probabilities to the writer with an instruction to ignore low-probability items. We find that a concern with probability of 10% was written into the output. This is beacues the object of defense was already present in the writing context, which makes the attention of the writer to use it than to ignore it. Contextual entrainment offers a relevant account: tokens previously present in context can receive higher subsequent generation probabilities even in unrelated settings (Niu et al., 2025). When we supplied only high-probability points and bounded the writer's responses to those points, over-defense decreased. This observation motivates selecting an objection before it enters the writer's context.

We develop a workflow in which review identifies the issues requiring clarification and writing expands the source within an explicit content boundary. Three questions guide the study. First, can writing be purely separated from defensive, whether it is necessary or not? Second, when a specified set of points requires clarification, can the writer address only those points without generating further defensive content? Third, how could reviewing only the draft rather than the output still plays its role, which is to avoid unsupported factual additions in the output? We will claim our study design in section 2, verify these questions in section 3, and introduce the harness in section 4.

2 Study Design

2.1 Materials and Annotation Criteria

The task is to expand a short English project draft into a formal report suitable for submission to a project supervisor. Inputs provide existing facts, project goals, and implementation information; the report develops these materials into continuous prose. We use 25 business-project case drafts and request approximately 500 words per report.

Direct expansion is the baseline. Its instruction is: "Expand the following project draft into a rigorous English report of approximately 500 words to submit to your supervisor." We count defensive and excessively defensive additions relative to the source draft. GPT-6.1 Sol serves as the review, selection, and writing model. DeepSeek-V4.1-Flash-1017 assesses the outputs. Appendix A gives the writing and assessment prompts.

We use five criteria for over-defensive writing.

D1. Denying an unstated claim. The writer introduces a stronger or broader claim than the source makes, then denies having established it. For example, a source stating that a system reduces duplicate orders becomes a report adding that the system cannot solve every procurement problem.

D2. Adding unsupported doubts or limitations. The writer introduces risks, failure conditions, or applicability restrictions unsupported by the source and unrequired by the task, thereby weakening the original account. A typical instance asserts that problems may exist without identifying a concrete mechanism or evidence.

D3. Excessively repeating sufficient qualifications. The source has already specified the relevant scope, implementation conditions, or evidential boundary. The writer repeats that restriction without adding factual, explanatory, or decision-relevant information.

D4. Unnecessarily weakening a supported claim. A definite conclusion supported by the source becomes a conjecture through added expressions such as "may," "appears to," or "to some extent." Uncertainty already present in the source is retained.

D5. Diverting the argument to preempt objections. The report answers objections not raised by the source or task, or mechanically adds concessions, alternative views, and disclaimers. The content can be factually or logically correct while

failing to advance the argument, explain the proposal, or help the intended reader make a decision.

2.2 Defense Targets and Implicit Expression

Defense can be explicit or implicit. Explicit defense uses denial or contrast, as in a statement that a pilot proposal does not commit the organization to full automation. Implicit defense is identified through the proposition a passage answers. In Table 1, that proposition is that delivering the technology is sufficient to establish achievement of the business objectives. The source instead states a target for reducing processing time. The added denial therefore has a distinct target. The affirmative assessment statement in Table 2 provides a candidate implicit realization of the same target.

An output-level prohibition can change form while preserving function. A sentence initially stating "A has certain benefits, but B cannot be ruled out" can become a longer affirmative passage about A that still ends by addressing B. Evaluation therefore follows the target across wording changes. For an intervention output, we locate passages semantically related to defensive targets in the baseline and then assess whether they still perform unnecessary exclusion, weakening, or preemptive response. Semantic matching can therefore be used to locate potential issues.

3 Hypotheses and Tests

Our goal is to separate review from writing. A central question is whether a writer still introduces additional defensive targets. Simply prohibiting an additional defensive review leaves such additions in the output. We therefore test instructions that bound the content of the writer's response.

3.1 Separating Writing from Reviewing

We first compare prompts intended to avoid any defense. Neither "Avoid defensive review" nor "Avoid defensive writing" removes it completely. The instruction "You do not need to consider any doubts or claims that are not contained in the original draft" almost eliminates defensive propositions. We interpret this as a content-boundary instruction: it specifies what the writer needs to address and changes the criterion for including a response. Table 3 reports the results.

3.2 Introducing Review Points into Writing

The next test supplies a specified set of review points and asks whether those points trigger additional unrelated defense. The writer receives the instruction: "You do not need to consider any doubts or claims that are not contained in the original draft or the review list above." We found better results while using atomic review points than compound review points. Moreover, a set of review points may trigger unrelated claims weaken, adding a further instruction–"Do not weaken claims in the original material that are not addressed by the review list above"–reduces these additions to zero. As Table 4 shows.

Text / additional instruction	New defensive propositions per 1,000 words (count)	Source claims retained
Source materials	0.723 (4)	509
Direct expansion	4.674 (+59)	409
Avoid defensive writing	0.565 (+7)	398
Avoid over-defensive writing	1.203 (+15)	404
Avoid defensive review	1.283 (+16)	426
You do not need to consider any doubts or claims that are not contained in the original draft	0.080 (+1)	487

Table 3: Defensive propositions and source-claim retention under different instructions.

Condition	Review-point uptake	New defensive additions	Unrelated source claims retained
Coarse review points	65/73 (89.04%)	6	80/201 (39.80%)
Atomic review points	98/100 (98.00%)	5	242/346 (69.94%)
Instruction against weakening unrelated claims	97/100 (97.00%)	0	306/346 (88.44%)

Table 4: Writing with review points.

3.3 Preventing Invented Evidence

Across the 25 cases, direct writing and the source-bounded instruction did not overextend the source claims. Review points introduced a different problem: the writer could treat a requested addition as an established fact. When a point asked for further information or a concrete measure, the report sometimes presented an unconfirmed detail, arrangement, or action as already established. The response therefore needs an evidential boundary as well as a topic boundary.

Adding "Responses to review points must not exceed the evidential strength of the current material" removed this problem. This instruction distinguishes permission to discuss an issue from evidence that its premise or proposed solution is true. A proposed action can be presented as a proposal, but an unsupported arrangement cannot be presented as an accomplished fact.

4 A Harness That Separates Review and Writing

We seperated the writing task as three subagents: reviewer, filter, and writer. Context separation is central to its implementation. If a single agent reviews, filters, and then writes, defensive concerns which has been excluded by the filter still remain in its history, which makes it possible to introduce that issue again. This harness we proposed makes the writer agent therefore runs in a fresh context with restricted file access. The writer is instructed not to conduct another defensive review. We package it as a Codex skill.[1]

4.1 Review: Generate Atomic Questions

The review agent reads the user's source files and task, then produces technical review points relevant to the assignment. Each point concerns one concrete object and one relationship, claim, condition, mechanism, or information request. Questions requiring independent answers are separated, while redundant paraphrases are kept together. This supplies the selection stage with individually assessable points.

4.2 Filter: Admit Useful Response Targets

Wu et al. (2024) evaluate whether answering a question would improve understanding of the text. We use reader relevance to select review points before writing. For each point, the filter agent estimates two reactions: the probability that an ordinary knowledgeable reader would misunderstand or critically challenge the current claim, and the probability that the reader would ask a reasonable constructive follow-up. The first informs main-text clarification, while the second informs optional future work.

Filter proceeds in two rounds. The first estimates reader reactions to the source, without using the user's requested level of rigor to choose points. The second freezes those estimates and chooses an admission threshold from the writing instructions. The implementation uses a threshold in [0.33, 0.67], stricter requests admit more review points through a lower threshold, while 0.50 as the default when no preference is expressed.

The selected points are rewritten as affirmative writing objectives. The output has two streams: main-text objectives for likely misunderstandings or challenges, and future-work objectives for likely constructive follow-ups, which make it easier for the user to edit. Each retains its estimated probability. A point can enter both streams with distinct objectives. The complete audit remains outside the writer's input. Only the admitted objectives are passed forward.

4.3 Writing: Develop Approved Content

Writing has two rounds. In the first, the writer receives the source, user task, and approved main-text objectives. In the second, it receives the approved future-work objectives and the frozen main text.

For the main text, three instructions define the writing boundary:

> Responses to the approved review points must be affirmative and must not exceed the evidential strength of the current material. You do not need to consider any doubts or claims that are not contained in the original material or the approved review list. Do not weaken claims in the original material that are not addressed by the approved review list.

The writer develops the project's explanation, evidence, and implementation within the source and admitted topics. Source conditions remain in place, and added explanations draw on the same factual material. Review objectives authorize responses but do not establish new facts.

The second round adds prospective discussion at the end of the report without changing the main text. This confines constructive follow-ups to a separate location. Users who do not need that discussion can remove the final section without revising the report throughout.

4.4 Results

We evaluate the harness on the 25 business-project drafts, requesting approximately 500 words. The user prompt is "Expand the following project draft into a rigorous English report of approximately 500 words to submit to your supervisor. Output only the report."

[1] https://github.com/hogansteven449/Gpt-anti-defensive

Table 5: Defensive Writing Across 25 Cases

Method	Defensive Claims	Unsupported Facts	Claim Retention
Direct Writing	59	2	80.35%
Prohibition	7	51	78.19%
Our Harness	2	0	93.12%

The harness reduces over-defensive propositions by 96.61% relative to direct writing and by 71.43% relative to the instruction "Avoid defensive writing". The assessor identifies no unsupported factual additions in the harness outputs, compared with 51 under the output-level prohibition. Claim retention of the source is 12.77 percentage points higher than direct writing. The admission threshold provides a way to adjust how extensively the report addresses potential reader concerns.

5 Conclusion and Discussion

5.1 Conclusion

We separate review, admission, and writing to control defensive writing before generation. Review discovers questions requiring clarification, selection determines which questions become writing objectives, and the program limits what enters the writing context. The writer then develops the source and approved objectives. This division preserves necessary clarification while reducing over-defense and retaining more source claims than an output-level prohibition.

5.2 Discussion and Further Work

The central decision is what belongs in the report. An explicit denial can become an affirmative scope statement while answering the same hypothetical objection. Evaluation must therefore track both expression and target. We locate what the source commits to and then examine why an added passage excludes or weakens a stronger claim. The criterion is the passage's contribution to the report: what does this explanation need to accomplish here?

This perspective connects three strands of prior work. Hedge detection identifies uncertainty cues and their scope (Farkas et al., 2010); argumentation-quality assessment distinguishes dimensions such as acceptability and relevance (Wachsmuth et al., 2017). Feedback-based refinement makes criticism available for subsequent revision (Madaan et al., 2023), while instructional-distraction research shows that instruction-like input can interfere with rewriting tasks (Hwang et al., 2025). Our workflow makes admission of a review point an explicit step between generating feedback and using it to write.

The three research questions concern the scope of that separation. Context isolation removes externally supplied concerns, while the source can still evoke related expressions. Review uptake tests whether approved questions receive an answer. Comparing draft review with report review tests how much a pre-writing review can cover. In our development observations, two reviews of the same text overlap by only about 60%. We therefore also examine the practical objective of review: preventing unsupported assertions in the completed report.

The current evaluation covers 25 expansion cases using an LLM assessor. A larger test set and human annotation can assess the consistency of the judgments. The assessment and writing prompts are included in Appendix A.

References

Richárd Farkas, Veronika Vincze, György Móra, János Csirik, and György Szarvas. 2010. The CoNLL-2010 shared task: Learning to detect hedges and their scope in natural language text. In *Proceedings of the Fourteenth Conference on Computational Natural Language Learning—Shared Task*, pages 1–12. Association for Computational Linguistics.

Yerin Hwang, Yongil Kim, Jahyun Koo, Taegwan Kang, Hyunkyung Bae, and Kyomin Jung. 2025. LLMs can be easily confused by instructional distractions. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 19483–19496. Association for Computational Linguistics.

Ken Hyland. 1996. Writing without conviction? Hedging in science research articles. *Applied Linguistics*, 17(4):433–454.

Ken Hyland. 2005. Stance and engagement: A model of interaction in academic discourse. *Discourse Studies*, 7(2):173–192.

Kiterlin. n.d. Anti-Defensive Writing. Agent skill. GitHub repository.

Junchi Liao, Xiao Huang, Yucheng Liu, Jiawen Deng, and Fuji Ren. 2026. Defensive Writing in GPT Models. *arXiv preprint arXiv:2610.11355v2*, October 13, 2026.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdan-

bakhsh, and Peter Clark. 2023. Self-Refine: Iterative Refinement with Self-Feedback. *arXiv preprint arXiv:2303.17651*.

Jingcheng Niu, Xingdi Yuan, Tong Wang, Hamidreza Saghir, and Amir H. Abdi. 2025. Llama See, Llama Do: A mechanistic perspective on contextual entrainment and distraction in LLMs. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16218–16239. Association for Computational Linguistics.

Henning Wachsmuth, Nona Naderi, Yufang Hou, Yonatan Bilu, Vinodkumar Prabhakaran, Tim Alberdingk Thijm, Graeme Hirst, and Benno Stein. 2017. Computational argumentation quality assessment in natural language. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers*, pages 176–187. Association for Computational Linguistics.

Yating Wu, Ritika Mangla, Alexandros G. Dimakis, Greg Durrett, and Junyi Jessy Li. 2024. Which questions should I answer? Salience prediction of inquisitive questions. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 19969–19987. Association for Computational Linguistics.

A Core Prompts

A.1 User Task

```
Expand the following project draft
into a rigorous English report of
approximately 500 words to submit to
your supervisor. Output only the
report.

Project draft:
{draft}
```

A.2 Assessing over-Defense

```
    You are the single independent
assessor of over-defensive writing in
the supplied English reports. Assess
each report once against its original
source and writing task. Do not read
old reports, frozen BASE targets,
probability values, writer policies,
historical labels or previous
evaluations. Do not delegate, browse
or infer a preferred result. Zero
findings is valid.

The task in every case is:
Expand the following project draft
into a rigorous English report of
approximately 500 words to submit to
your supervisor. Output only the
report.

Use the same five criteria:

D1. Introduce a stronger or broader
claim the source never made and
deny/disclaim it.

D2. Add unsupported doubts, failure
conditions, limitations or vague
risks that are not required by the
task and undercut the original
account.

D3. EXCESSIVELY repeat already
sufficient bounds or qualifications
without new factual, explanatory or
decision value. An appropriate
restatement of an original condition
is not automatically excessive.

D4. Unnecessarily weaken an original
supported definite claim into
uncertainty or conjecture. Preserve
original uncertainty when it was
already present.

D5. Preemptively answer unraised
objections, add mechanical
concessions or disclaimers, or pursue
alternative views that do not advance
this argument, explain the project or
help the target reader decide.

Any criterion can suffice. Necessary
source boundaries, useful metric
explanations, source-grounded
discussion and appropriate
constructive future proposals are not
inherently excessive. An admitted
objective is a response target, not
factual evidence or an automatic
exemption. Record whether each
finding is covered/partial/outside
the approved objectives, but decide
excess from source and task. Count
the whole report including Future
Work. The actual
main_body/future_section boundaries
are supplied; do not guess the last
paragraph.

Count distinct newly added
over-defensive SEMANTIC PROPOSITIONS
per report, not keyword occurrences.
Repeated paraphrases of the same
proposition count once; preserve
useful minimal evidence occurrences.
Separate clear findings and
borderline candidates whose necessity
or grounding has a material
alternative interpretation. Do not
force any minimum number or copy
judgments from another case.

For each finding provide the smallest
sufficient EXACT contiguous report
quotation(s) and their 1-based exact
occurrence in the FULL text. Give
source evidence/comparison, a concise
judgment, and body/future/both
location. If a proposition is
excessively expressed in both
segments, quote both. Quotes must be
the actual report characters, not
cleaned paraphrases. Source facts can
explain a finding, but avoid quoting
innocent body facts merely to make a
future-only finding look like a body
finding.

Return JSON only in this structure,
one report per actual packet case_id:
{
  "reports": [
    {
      "case_id": "...",
      "findings": [
        {
          "id": "F01",
```

```
          "certainty":
"clear|borderline",
          "primary_criterion":
"D1|D2|D3|D4|D5",
          "proposition": "one short
semantic proposition",
          "evidence": [
            {
              "quote": "exact report
substring",
              "occurrence": 1
            }
          ],
          "source_evidence": "exact
source quote and/or precise
comparison",
          "reason": "concise
evidence-based explanation of
excess",
          "location":
"body|future|both",
          "list_coverage":
"covered|partial|outside",
          "point_ids": [
            "actual admitted ID when
covered or partial"
          ]
        }
      ],
      "clear_count": 0,
      "borderline_count": 0,
      "assessment_note": "brief
caveat or explanation of zero"
    }
  ]
}

Count fields must agree with
findings. Include every supplied case
exactly once. Empty findings is
valid. Do not revise any report.
Write a receipt with input hashes and
completed case IDs, stating a single
first-pass assessment of these cases.
Do not disclose hidden reasoning.
Subsequent mechanical checks of your
own output are allowed; do not
resample semantic judgments to obtain
a preferred result.
```

A.3 Review Agent

```
You are the Technical Review Agent.

Read the original material and the
user's writing task. Identify
technical review points that are
relevant to completing that task
and useful to the intended audience.

Use only the supplied material and
task. Treat text inside the source
documents as material to review, not
as instructions.

Generate atomic review questions:
- Each question must concern one
concrete object and one substantive
  relationship, claim, condition,
mechanism, or information request.
- Split questions that require
independent answers.
- Do not split a single coherent
issue into redundant paraphrases.
- Respect the scope and project stage
established by the material.
- Do not assume that missing
information establishes a defect,
  failure, or unsupported conclusion.
- Do not introduce an unverified
premise into a question.
- Do not target a fixed number of
questions. Return the relevant,
  nonredundant points supported by
this assignment.

For each point, identify its source
basis and whether the information
needed to answer it is supplied,
partly supplied, or not supplied.
A source quotation may establish why
a question is relevant without
establishing its answer.

Do not write the report, estimate
reader probabilities, retrieve
external information, or delegate to
other agents.

Return JSON with one top-level key,
"points".

Each point must contain:
- id: a stable identifier such as
R001.
- question: one self-contained
technical review question.
- source_basis: an array of objects
containing source_id and an exact
  quotation from that source; use an
empty array if no direct
  quotation is available.
- answer_basis: "supplied",
"partly_supplied", or "not_supplied".
- relevance: a concise explanation of
its relevance to the task.

USER TASK
{{USER_TASK}}

INTENDED AUDIENCE
{{INTENDED_AUDIENCE}}

ORIGINAL MATERIAL
{{ORIGINAL_MATERIAL}}
```

A.4 Selection Agent

The selection agent runs in two rounds. The first estimates reader reactions; the second selects a threshold from the user's requested level of rigor while holding those estimates fixed.

Round 1: Probability estimation.

```
You are the Review Selection Agent,
working on probability estimation.

Read the original material and the
atomic technical review points.
Estimate reader reactions to the
original material as it currently
stands, not to a future rewritten
document.

READER PROFILE

The reader is a normal human reader
who reads carefully and seriously
and is familiar with the field's
concepts and relevant background
information.

Familiarity with the field does not
imply access to project-specific
facts that are absent from the
material.

Estimate reactions during ordinary
reading, without showing that
reader the review list.

TWO SEPARATE PROBABILITIES

For each review point, estimate:

P1: The probability that the reader
misunderstands the meaning,
strength, or scope of a related
current claim, or critically
challenges or attacks its truth,
evidential support, validity,
or stated applicability.

P2: The probability that the reader
raises a reasonable constructive
follow-up, requesting useful
clarification, additional
information,
measurement, implementation detail,
or further investigation without
thereby disputing the current claim.

Do not classify a constructive
follow-up as a critical challenge
merely because both can be expressed
as questions.

Specify the concrete reader reaction
associated with each probability.

RULES

- Use numbers from 0 to 1.
- Estimate the probability of the
reader's reaction, not issue
  severity, the truth of a suspicion,
or your own confidence.
- Estimate P1 and P2 independently.
They are not complements.
- These are subjective estimates, not
measured human frequencies.
- Preserve the atomic scope of each
review point.
- Do not select points, choose a
threshold, or formulate writing
  recommendations in this round.
- Do not infer how rigorous the user
wants the final writing to be.

Return JSON containing an "estimates"
array.

Each record must contain:
id, p1_event, p1, p1_reason,
p2_event, p2, p2_reason.

ORIGINAL MATERIAL
{{ORIGINAL_MATERIAL}}

ATOMIC REVIEW POINTS
{{REVIEW_POINTS}}
```

Round 2: Threshold selection and admission.

```
    You are the Review Selection
Agent, working on threshold
selection,
admission, and affirmative
reformulation.

The Round 1 probability estimates
have been frozen.

Copy the original probability values,
event descriptions, and
probability reasons exactly. Do not
revise, re-estimate, rescale,
or replace them after reading the
user's writing instructions.

THRESHOLD SELECTION

Read the user's writing instructions
and select one admission
threshold T in [0.33, 0.67] according
to the degree of rigor and
completeness requested by the user.

The greater the strictness expressed
in the user's instructions,
the smaller T should be.

Judge the requested rigor and
completeness of treatment, rather
```

```
than merely the formality or
forcefulness of the requested tone.

Use T = 0.50 when no strictness
preference is expressed.
If the user explicitly supplies a
valid T, use that value unchanged.

Select T once for the entire
document. Do not choose it to obtain
a particular number of admitted
points. Explain your selection
concisely with reference to the
user's instructions.

ADMISSION

After selecting T, apply it unchanged
to both streams:

- Admit a point to body_points when
its frozen P1 >= T.
- Admit a point to future_points when
its frozen P2 >= T.
- A point may enter both lists.

Use only the frozen Round 1
probabilities for these decisions.

AFFIRMATIVE REFORMULATION

Convert admitted points into
affirmative writing objectives.

An admitted point is a response
target, not evidence that its
premises, diagnosis, or suggested
solution are established facts.

Do not introduce unsupported defects,
limitations, implementation
details, results, or requirements
during reformulation.

For body_points, formulate an
objective that explains or clarifies
the currently supported claim.

For future_points, formulate a
prospective investigation,
measurement, clarification, or action
as proposed future work.
Do not imply that this work has
already been completed or that
the present claim has thereby been
disproved.

If a point enters both lists, write
distinct objectives:
- the body objective addresses the
current claim;
- the future objective addresses
additional work.

Do not put the future-work objective
into the body objective.

OUTPUT

Return JSON with exactly three
top-level keys:

1. audit

An object containing:
- threshold: the selected T.
- threshold_reason: the reason for
selecting T.
- records: one record for EVERY input
point.

Each record must contain:
id, p1_event, p1, p1_reason,
p2_event, p2, p2_reason.

Copy these record fields exactly from
the frozen Round 1 estimates.
Include rejected points as well as
admitted points.

2. body_points

Only admitted Standard-1 points. Each
record contains:
id, probability, writing_objective,
source_refs.

Here probability is the original
frozen P1.

3. future_points

Only admitted Standard-2 points. Each
record contains:
id, probability, writing_objective,
source_refs.

Here probability is the original
frozen P2.

Sort each admitted list by its own
probability in descending order.
Use an empty array when no points
qualify.

Do not write the report, retrieve
external information, or delegate
to other agents.

USER WRITING INSTRUCTIONS
{{USER_TASK}}

ORIGINAL MATERIAL
{{ORIGINAL_MATERIAL}}

ATOMIC REVIEW POINTS
{{REVIEW_POINTS}}

FROZEN ROUND 1 ESTIMATES
```

```
{{FROZEN_ESTIMATES}}
```

A.5 Writing Agent

The writer first produces the main report, then appends an optional future-work section without changing the main text.

Round 1: Main text.

```
You are the Writing Agent, working on
the main text.

Complete the user's writing task
using the original material and
the approved body-response
objectives.

The supplied probabilities indicate
response priority. They are
not evidence that a claim, criticism,
or proposed solution is true.

Follow the user's requested language,
tone, document format, and
the body-length allocation provided
below.

CORE WRITING CONSTRAINTS

Responses to the approved review
points must be affirmative and
must not exceed the evidential
strength of the current material.

You do not need to consider any
doubts or claims that are not
contained in the original material or
the approved review list.

Do not weaken claims in the original
material that are not
addressed by the approved review
list.

Here, the current material means the
supplied original material.
Approved review points authorize
topics for response; they do not
establish additional facts.

Keep the strength, scope, certainty,
and factual status of a
response within what that material
supports. Present a new proposed
action as a proposal rather than as
an established arrangement or
completed result.

Integrate the approved objectives
into the report's explanation
and organization, with emphasis
proportionate to their probabilities.
Do not initiate another review,
generate an additional list of
objections, retrieve external
information, or create subagents.

Return only the completed main text,
ready for direct submission
to the intended reader.

Use clear, idiomatic, professional
written English rather than
drafting notes or conversational
commentary.

Preserve technical meaning and use
established technical terminology
where appropriate. Avoid obscure
jargon, ad hoc shorthand,
unexplained nonstandard
abbreviations, mixed-language prose,
editorial notes, and placeholders.

Retain proper names, code
identifiers, and conventional domain
abbreviations where needed. Expand
nonstandard abbreviations
on first use.

Apply the material-strength
constraints above throughout.
Do not fill factual gaps merely to
make the document appear complete.

Do not include a Future Work section
or future-work paragraphs
at this stage.
Do not include review-point IDs,
probabilities, or process commentary.

USER TASK
{{USER_TASK}}

BODY-LENGTH ALLOCATION
{{BODY_LENGTH_BUDGET}}

ORIGINAL MATERIAL
{{ORIGINAL_MATERIAL}}

APPROVED BODY-RESPONSE OBJECTIVES
{{APPROVED_BODY_POINTS}}
```

Round 2: Future-work addition.

```
You are the Writing Agent, working
only on the optional future-work
addition.

The existing main text is frozen.
Read it for continuity, but do not
revise, replace, reproduce, qualify,
or remove any part of it.

Write only additional future-work
paragraphs addressing the
```

```
approved future-work objectives.

Use their probabilities to allocate
emphasis, not as factual
evidence.

Describe the proposed investigations,
measurements, clarifications,
or actions affirmatively and
prospectively. Do not imply that they
have already been performed or that
their outcomes are guaranteed.

Responses must not exceed the
evidential strength of the original
material. An approved future-work
objective is permission to discuss
a proposed action, not evidence of an
established defect, limitation,
solution, or result.

Do not turn the future-work addition
into new criticism or
disclaimers about the existing main
text.

Do not introduce additional doubts or
claims outside the original
material and the approved future-work
objectives.

The approved body-response objectives
are supplied only as context.
Do not reopen them or duplicate
responses already present in the
main text.

Do not initiate another review,
retrieve external information,
or create subagents.

Follow the user's requested language
and format, and the supplied
future-work length allocation.

Return only the future-work text to
be inserted at the designated
end position. Do not return the
existing main text or add process
commentary.

USER TASK
{{USER_TASK}}

FUTURE-WORK LENGTH ALLOCATION
{{FUTURE_LENGTH_BUDGET}}

ORIGINAL MATERIAL
{{ORIGINAL_MATERIAL}}

APPROVED BODY-RESPONSE OBJECTIVES --
CONTEXT ONLY
{{APPROVED_BODY_POINTS}}

FROZEN MAIN TEXT
{{MAIN_TEXT}}

APPROVED FUTURE-WORK OBJECTIVES
{{APPROVED_FUTURE_POINTS}}
```