Where Should the Evaluation Budget Go in LLM-Driven Agent Scaffold Evolution?

Sichun Luo[1], Yi Huang[2], Mukai Li[1], Haochen Luo[1], Guanzhi Deng[3], Haibo Wang[4]
Jian Xu[5], Zefa Hu[2], Junlan Feng[2], Qi Liu[1]
[1]The University of Hong Kong [2]JIUTIAN Research, China Mobile
[3]City University of Hong Kong [4]Carnegie Mellon University [5]Tsinghua University

arXiv:2610.37276v1 [cs.LG] 27 Oct 2026

Abstract

Large language model (LLM)-driven agent scaffold evolution improves agent workflows through iterative candidate generation and evaluation. Evaluation both guides subsequent search and determines the final scaffold, but evaluating every candidate on the full task set incurs substantial execution costs. We investigate how evaluation budgets can be allocated across these two roles. Our experiments show that sparse feedback can sustain productive evolution, while additional targeted evaluation improves final selection. Replay analyses further show that shared informative items and a common scoring reference support comparisons among closely matched candidates. Motivated by these findings, we propose **EvalAlloc**, an evaluation allocation framework that combines sparse search-time feedback with targeted final confirmation. A short fully evaluated warm-up establishes a shared informative core and calibrates a response model whose item parameters remain fixed during search. After search, EvalAlloc refits the model and applies paired successive halving, reusing accumulated observations and executing only missing evaluations. Across four benchmarks and two evolution frameworks, EvalAlloc achieves higher mean performance than full evaluation in most settings while using approximately half as many evaluation cells.

1 Introduction

A scaffold for a large language model (LLM) agent is an executable program that defines how an agent coordinates LLM calls, maintains intermediate state, and interacts with tools and external information sources [13, 26, 31]. For Web-facing agents, scaffolds govern how agents acquire and use information through Web searches, document retrieval, and interactions with online services [22, 36]. While these scaffolds have traditionally been designed and refined manually, recent advances have enabled their automated optimization [2, 13, 34]. For example, ShinkaEvolve [17] iteratively proposes scaffold modifications and uses evaluation feedback to guide subsequent evolution.

However, evaluating candidate scaffolds can incur substantial execution costs. A single task instance may involve multiple LLM calls, retrieval steps, and tool interactions [21, 32], making full evaluation across many candidates expensive. Although parallel evaluation and asynchronous program generation can overlap computation [28], feedback-dependent search decisions still require evaluation results. Existing approaches reduce evaluation costs through sparse score estimation and evaluation cascades [8, 9], treating evaluation primarily as a fitness-estimation problem. In scaffold evolution, however, evaluation supports two distinct decisions: it provides fitness feedback that shapes archive updates and subsequent proposals [33], and it determines which scaffold is ultimately returned. These decisions impose different requirements: search benefits from inexpensive yet comparable feedback, whereas final selection requires reliable discrimination among competitive candidates. This raises a central question: *How should a limited evaluation budget be allocated across these two decisions to improve the quality of the ultimately selected scaffold?*

To answer this question, we investigate how evaluation allocation affects evolutionary search and final scaffold selection. Our investigation yields three findings. First, sparse evaluation feedback can sustain productive evolution, whereas removing feedback substantially degrades search quality. Second, shared informative items improve comparisons between closely matched candidates, while maintaining a common scoring reference helps preserve comparability across evolutionary iterations. Third, search improvements are concentrated in early iterations, whereas final selection remains sensitive to uncertainty among top candidates. These findings suggest that evaluation should provide inexpensive yet comparable feedback during search, while reserving additional observations for final selection.

Motivated by these findings, we propose **EvalAlloc**, an evaluation allocation framework that separates sparse search-time feedback from targeted final selection. During search, a short fully evaluated warm-up identifies a shared informative core and calibrates a response model. Subsequent candidates are evaluated on sparse item sets centered on this core, while the model's item parameters remain fixed to maintain a common scoring reference for archive updates and parent selection. After search, EvalAlloc applies warm-started paired successive halving (SH), reusing accumulated observations and performing only missing evaluations. A refitted response model further supports comparisons during later selection rounds. We evaluate EvalAlloc on four benchmarks and two evolution frameworks, with three independent runs per setting. Compared with full evaluation, EvalAlloc uses approximately half as many evaluation cells while achieving competitive or higher mean performance across the evaluated settings. Ablation studies further show that informative item allocation and final selection are the main contributors to its effectiveness. These results demonstrate the benefits of separating sparse search-time feedback from targeted final evaluation.

Our main contributions are summarized as follows:

- We investigate evaluation allocation as a decision-dependent problem in LLM-driven scaffold evolution, demonstrating that search-time feedback and final scaffold selection impose distinct requirements on evaluation resources.
- We propose EvalAlloc, an evaluation allocation framework that combines a shared informative core, a fixed search-time

response model, and warm-started paired successive halving to allocate evaluation effort across evolutionary search and final selection.
- We evaluate EvalAlloc across four benchmarks and two evolution frameworks, complemented by controlled replay analyses, demonstrating competitive final performance with substantially fewer evaluation cells.

2 Related Work

LLM-Driven Agent and Program Evolution. Recent work has explored the automated design and optimization of agent scaffolds and executable workflows using LLMs [23]. ADAS [13] uses a meta-agent to generate agent programs based on previous discoveries, while AFlow [34] employs Monte Carlo Tree Search to optimize code-represented workflows. Other approaches, including ShinkaEvolve [17] and AdaEvolve [4], iteratively generate, evaluate, and refine programs using evaluation feedback to guide subsequent search. These approaches primarily focus on improving candidate generation and evolutionary optimization. Our work instead examines the evaluation process underlying this feedback-driven search, studying how limited evaluation resources can support both evolutionary decisions and final scaffold selection.

Efficient LLM Evaluation. Reducing the cost of evaluating LLMs has received growing attention. Methods such as tinyBenchmarks [24] use carefully selected examples and statistical estimation to approximate full-benchmark performance, while SparseEval [35] optimizes representative items for efficient score estimation and ranking. CollabEval [8] exploits shared information across evaluations through matrix completion to improve statistical efficiency. Growing Pains [12] uses fixed item-response calibration and anchor items to maintain score comparability across evaluation periods. These studies demonstrate the effectiveness of sparse observations and response modeling. However, in scaffold evolution, evaluation estimates also influence archive updates and subsequent candidate generation. Therefore, we examine not only estimation accuracy, but also how shared informative observations and stable scoring references support reliable evolutionary decisions.

Evaluation Allocation and Candidate Selection. Evaluation allocation is closely related to sequential resource allocation and best-arm identification. Successive halving [14] and Hyperband [18] progressively concentrate evaluation resources on promising candidates, while racing-based methods such as irace [20] eliminate weaker configurations through sequential comparisons. Common random numbers [10] provide a classical motivation for evaluating alternatives under shared conditions to reduce the variance of pairwise comparisons. Recent analyses of code evolution [9] also highlight the importance of reliable candidate evaluation and selection. Building on these ideas, our work studies evaluation allocation across two distinct stages: providing inexpensive yet comparable feedback during evolution and reusing accumulated observations for reliable final scaffold selection.

3 Problem Formulation

Agent Scaffold Evaluation. Let c denote an executable agent scaffold and $\mathcal{D}_s = \{x_i\}_{i=1}^N$ a set of evaluation tasks. Executing c on x_i produces a score $Y_{ci}^{(r)} \in [0,1]$, where r indexes repeated executions. Each scaffold–item execution constitutes an evaluation *cell*. The expected fitness of a scaffold is $F(c) = \frac{1}{N}\sum_{i=1}^N \mathbb{E}[Y_{ci}^{(r)}]$. Evaluating each item once yields the full-set reference score $F_{\text{ref}}(c) = \frac{1}{N}\sum_{i=1}^N Y_{ci}^{(1)}$.

Table 1: Benchmark splits. SEARCH is used by the evaluators during evolution and selection, DEV only for the audit, and TEST only to score the selected programs.

Benchmark	SEARCH	DEV	TEST
HotpotQA	300	150	300
HoVer	300	150	300
DROP	300	150	300
IFBench	300	150	294

LLM-Driven Agent Scaffold Evolution. At iteration t, an evolution framework proposes a candidate scaffold c_t based on its current archive $\mathcal{A}_t$. An evaluator selects an item subset $S_t \subseteq \{1,\dots,N\}$ and returns an estimated fitness $\widehat{F}_t(c_t)$ using the available observations. This feedback informs archive updates and subsequent candidate generation. Evaluation allocation therefore influences the search trajectory and the resulting candidate pool. After search, a final selector returns a scaffold c_{pick} from the generated candidates, potentially collecting additional observations. We distinguish two evaluation stages: *search-time evaluation*, which provides fitness feedback for continued evolution, and *final selection*, which identifies the scaffold to return. Observations collected during search can be reused during final selection.

Cost–Performance Objective. Let π denote an evaluation allocation policy governing observations collected during both stages and the final scaffold selection. The total cost C_{eval}^{π} counts executed evaluation cells during search and final selection. Given an evaluation budget B, our objective is

$$\max_{\pi} \ \mathbb{E}\left[F(c_{\text{pick}}^{\pi})\right] \quad \text{s.t.} \quad \mathbb{E}[C_{\text{eval}}^{\pi}] \le B, \tag{1}$$

where c_{pick}^{π} is the scaffold returned under policy π, and the expectation accounts for randomness in candidate generation, evaluation, and allocation.

4 Investigation

4.1 Setup

We study four benchmarks: HotpotQA [30], HoVer [15], DROP [7], and IFBench [25]. Each benchmark is divided into SEARCH (300 items), DEV (150 items), and TEST (300 items, except 294 for IFBench), as shown in Table 1. SEARCH is used for evolution and final selection, DEV for auxiliary audits, and TEST exclusively for evaluating the selected scaffolds. We use two evolution frameworks: ShinkaEvolve [17] and AdaEvolve [4], the latter implemented using SkyDiscover [19]. Each run makes 80 proposals from the same initial program, with three seeds per setting. We use Qwen3.5-9B [27] as the executor, with temperature zero, thinking disabled, and at most 16,384 output tokens, and DeepSeek-V4-Flash-0731 [29] as the proposer. Both models are served locally with vLLM [16] on a single node with eight NVIDIA H20 GPUs.

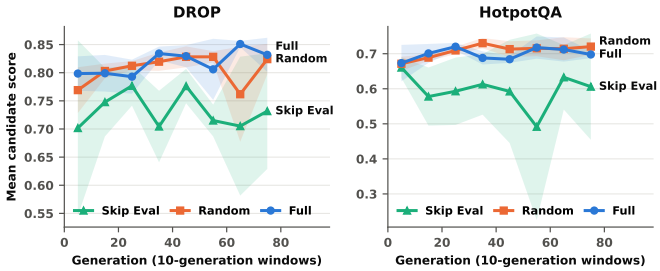

Figure 1: Evolution trajectories under different evaluation strategies on DROP and HotpotQA. Curves show mean candidate scores on a shared audit set over successive 10-generation windows.

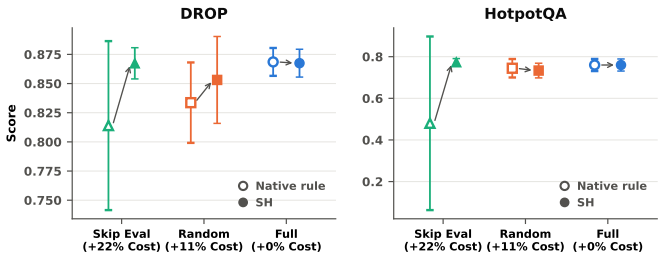

Figure 2: Effect of search-time feedback and final selection on DROP and HotpotQA. We compare the native selection rule with successive halving (SH) across three search strategies. Percentages indicate the additional evaluation cost of SH relative to the full-evaluation budget.

Most analyses in this section replay evaluators on the 24 *Full* runs (four benchmarks, two frameworks, three seeds), in which every candidate is evaluated on all SEARCH items, so the full-SEARCH mean of each candidate is known. A replay keeps the candidate sequence fixed and varies only which cells the evaluator observes. We measure two search-time quantities. *Close-pair agreement* is the share of candidate pairs whose full-SEARCH means differ by at most .03 that the evaluator orders correctly; these are the comparisons that archive updates face. *Archive regret* is, averaged over the run, the gap between the mean full-SEARCH score of the ten best candidates generated so far and that of the ten candidates the evaluator ranks highest.

4.2 Search Needs Feedback, but Not Full Evaluation

Evaluation feedback plays an important role in guiding evolutionary search. Figure 1 compares the evolution trajectories under three evaluation strategies on DROP and HotpotQA, using a shared audit set. Without evaluation feedback, candidate quality remains relatively low and shows limited improvement. In contrast, random partial evaluation maintains candidate quality close to Full, suggesting that effective evolution does not require evaluating every candidate on the entire SEARCH set.

Search improvements are also concentrated in early iterations. Although this does not establish the optimal temporal allocation, it motivates investing in early calibration and using sparse feedback for subsequent candidates.

> ***Takeaway 1.*** Evolutionary search requires evaluation feedback, but sparse evaluation can sustain candidate improvement without the cost of full evaluation.

4.3 Comparability Depends on Shared Informative Observations

We conduct offline replays on 24 Full evolution trajectories (four benchmarks, two frameworks, and three seeds), with three replays per trajectory. Each replay fixes the generated candidate sequence and varies only the evaluation cells observed by the evaluator. The first $W = 10$ candidates are fully evaluated, while subsequent candidates are evaluated on 60 of the 300 SEARCH items. All policies use the same evaluation budget, with full-SEARCH scores serving as the reference.

We assess three metrics: Spearman correlation with full-SEARCH scores, close-pair agreement for candidates whose reference scores differ by at most .03, and archive regret, defined as the average fitness gap between the true and estimated top-10 archives. These metrics capture both score estimation quality and the evolutionary decisions supported by limited feedback.

Informativeness and sharing. Table 2 shows that shared informative items yield more reliable close-candidate comparisons than random or unshared items. For the unshared informative policy, each candidate samples 60 items from the 120 most informative ones, rather than using a common item set. Despite drawing from informative items, this policy produces lower close-pair agreement and higher archive regret.

Pairing provides an explanation for the benefit of sharing. For k items drawn independently from the task distribution, the variance of the score difference between two candidates evaluated on shared items is $\mathrm{Var}(\widehat{\Delta}_{cc'}) = \frac{\sigma_c^2+\sigma_{c'}^2-2\sigma_{cc'}}{k}$, where $\sigma_{cc'}$ denotes their response covariance. With independently sampled item sets, the covariance term vanishes. Because evolved candidates often exhibit correlated responses, shared items can reduce the variance of pairwise comparisons.

Score comparability. Accurate individual rankings do not necessarily translate into reliable archive decisions. Although raw means achieve similar Spearman correlation to model-completed scores, they produce substantially higher archive regret. This is because sparse candidates are evaluated on informative items, whereas warm-up candidates are evaluated on the full SEARCH set. The fixed response model places both on a common full-SEARCH scale, avoiding inconsistent archive comparisons and repeated rescoring during evolution.

> ***Takeaway 2.*** Reliable evolutionary decisions require informative, shared observations and comparable fitness scores, rather than accurate individual estimates alone.

Table 2: Offline comparison of item policies and fitness estimators under matched evaluation budgets. Results are averaged over 24 Full trajectories, with three replays per trajectory.

Item policy	Estimator	Spearman	Close-pair agreement	Archive regret
Shared informative	Model	.862	.732	.0043
Shared informative	Raw	.865	.698	.0102
Unshared informative	Model	.795	.674	.0056
Shared random	Model	.750	.649	.0071
Independent random	Model	.717	.626	.0078

4.4 Final Selection Needs Confirmation

Final scaffold selection requires more reliable comparisons than those needed to guide evolutionary search. Figure 2 compares the native selection rule with successive halving (SH) on the same candidate pools under different search-time evaluation strategies. SH substantially improves final performance when search feedback is absent, recovering strong candidates that the native selection rule may overlook. With random partial evaluation, the gains are smaller and less consistent, while full evaluation leaves little room for further improvement. These results highlight the importance of final confirmation when search-time evaluation is limited.

> ***Takeaway 3.*** Final selection requires reliable comparisons beyond search-time fitness estimates, particularly when evaluation feedback is sparse or absent.

5 Method

Figure 3 illustrates EvalAlloc, an evaluation allocation framework that distributes evaluation effort across three stages: *Core Calibration*, *Shared-Core Evolution*, and *Warm-Started Selection*. EvalAlloc decouples the two roles of evaluation: search-time feedback only needs to be cheap and comparable enough to guide evolution, whereas the final choice needs precise comparisons among the strongest candidates. It therefore splits one evaluation budget into a sparse search stage and a separate selection stage that reuses every observation collected during search. The key idea is to invest a small evaluation budget upfront to establish a reusable reference, use shared informative items to provide comparable feedback throughout evolution, and concentrate additional evaluations on final scaffold selection. EvalAlloc operates as a drop-in evaluator, leaving the native framework's proposal, archive, and parent-selection policies unchanged.

5.1 Core Calibration

EvalAlloc begins by fully evaluating the first W valid candidates on all N SEARCH items. This upfront investment serves two purposes: identifying informative items for subsequent evaluation and calibrating a response model for estimating full-set fitness from sparse observations. The warm-up candidates receive their directly observed full-set scores, while their item-level outcomes form a dense response matrix for calibration.

Identifying informative items. We estimate item informativeness from accumulated evaluation outcomes. Let a_i denote the number of full-credit outcomes and n_i the number of observed outcomes for item i. Using a Beta posterior for its full-credit probability,

$$p_i \sim \mathrm{Beta}(a_i + 1, n_i - a_i + 1), \tag{2}$$

we define a conservative informativeness estimate:

$$I_i = 4\left(\mathbb{E}[X_i] - z\sqrt{\mathrm{Var}(X_i)}\right), \qquad X_i = p_i(1 - p_i). \tag{3}$$

The variance term favors items whose outcomes vary across candidates, while the uncertainty penalty discourages selecting items with insufficient evidence. Fractional outcomes below full credit count as failures for these acquisition statistics. The posterior moments are available in closed form.

Response model. We fit a low-rank logistic response model following a multidimensional item response theory (MIRT) formulation:

$$q_{ci} = \sigma\left(\mu + \eta_i + b_c + u_i^\top v_c\right), \qquad u_i, v_c \in \mathbb{R}^d, \tag{4}$$

where μ is a global intercept, (η_i, u_i) represent item-specific parameters, and (b_c, v_c) capture candidate-specific response characteristics. The model is fitted on the fully evaluated warm-up candidates by minimizing regularized fractional cross-entropy. Their fitted candidate parameters also define a Gaussian prior for subsequent candidates.

During evolution, the global intercept, item parameters, and candidate prior remain fixed. For each new candidate, only (b_c, v_c) are inferred from its sparse observations by MAP estimation under this prior. This maintains a common scoring reference across candidates without repeatedly refitting the model or rescoring archived candidates.

5.2 Shared-Core Evolution

After calibration, each new candidate is evaluated on a small subset of k SEARCH items ($k = 60$, i.e., 20% of $N = 300$). Rather than independently selecting items for each candidate, EvalAlloc combines a fixed informative core with a small adaptive component.

Shared informative core. Using the warm-up informativeness estimates, we select the top k_0 items to form a fixed core S_0. Every subsequent candidate is evaluated on this same core, creating a common basis for cross-candidate comparisons. This design follows the intuition of common random numbers (CRN): exposing candidates to shared evaluation conditions helps reduce variation arising from differences in the evaluated item sets.

To retain limited adaptability, the remaining $k - k_0$ items are candidate-specific. For each candidate c, we sample a set A_c from outside the core, using Gumbel-top-k sampling with logits I_i/τ computed from the current informativeness scores. With a small temperature, this amounts to taking the currently most informative items with random tie-breaking. The resulting evaluation set is

$$S_c = S_0 \cup A_c, \qquad |S_0| = k_0, \quad |S_c| = k. \tag{5}$$

The shared core remains unchanged, whereas the acquisition statistics are updated after each candidate using newly observed outcomes. This allows the adaptive component to respond to changes in the evolving candidate population without sacrificing the common comparison basis.

Concretely, for two candidates scored on the same k_0 items, the variance of their score difference is $\mathrm{Var}(\bar{Y}_a - \bar{Y}_b) = (\sigma_a^2 +$

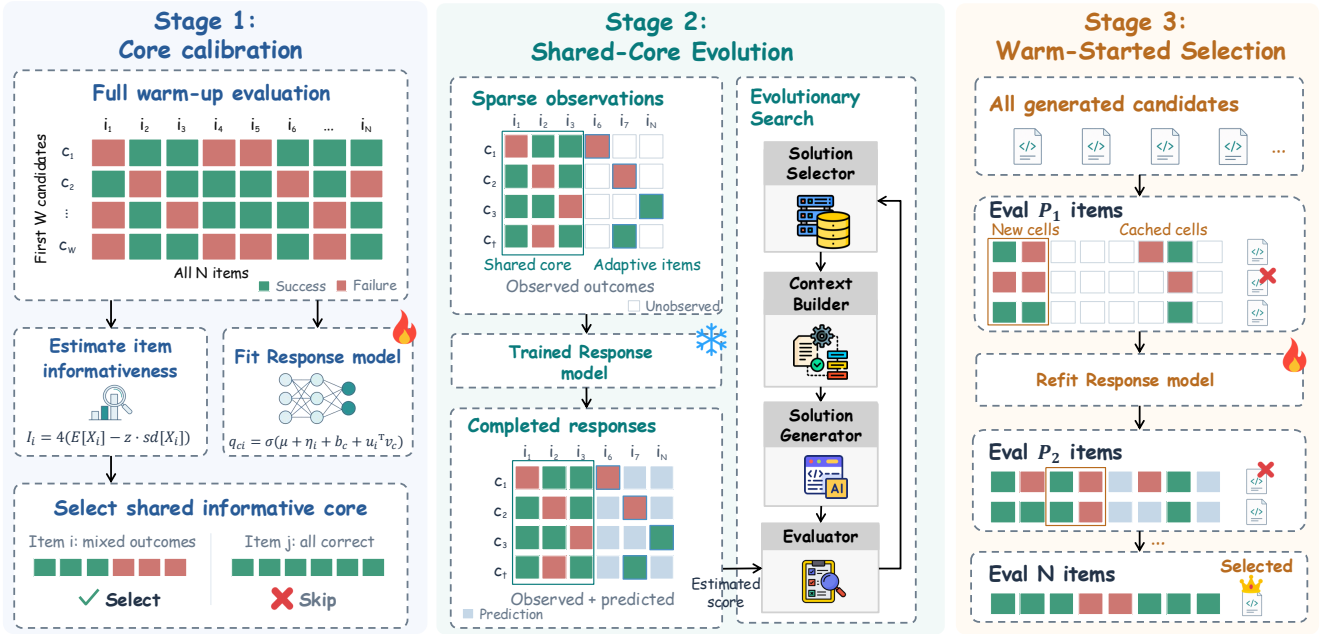

Figure 3: Overview of EvalAlloc. The framework consists of three stages: (1) core calibration using fully evaluated warm-up candidates, (2) sparse search-time evaluation with a shared informative core and fixed response model, and (3) warm-started successive halving that reuses accumulated observations for final scaffold selection.

$\sigma_b^2 - 2\,\mathrm{Cov}(Y_a, Y_b))/k_0$, whereas independently drawn item sets remove the covariance term. Because evolved scaffolds are strongly correlated—competitive candidates agree on most items—a shared core makes pairwise comparisons less noisy at the same budget.

Comparable fitness feedback. After evaluating candidate c on S_c, we infer its candidate-specific response parameters under the fixed warm-up model. The fitted model then predicts responses on unobserved items, producing a completed estimate of full-set fitness:

$$\widehat{F}(c) = \frac{1}{N}\left(\sum_{i\in S_c} Y_{ci} + \sum_{i\notin S_c} q_{ci}\right). \tag{6}$$

Here, Y_{ci} denotes the observed outcome and q_{ci} the model prediction using the inferred parameters of candidate c.

Combining observed and predicted outcomes places sparse evaluations on the full-SEARCH fitness scale expected by the native evolution framework, which also contains the fully evaluated warm-up candidates. Keeping the response model fixed also avoids retrospectively changing the scoring reference for archived candidates. The resulting fitness estimates are returned directly to the framework for its existing archive updates and subsequent proposals. All observed outcomes are cached for reuse during final selection.

5.3 Warm-Started Selection

Search-time evaluation aims to provide inexpensive feedback for continued evolution, whereas final selection requires distinguishing the strongest candidates. After search, EvalAlloc therefore runs successive halving (SH) over all valid candidates, using additional observed outcomes rather than relying on search-time fitness estimates. Candidates with identical code (the same abstract syntax tree after removing docstrings) are merged first.

We construct a shared random permutation of the SEARCH items and define increasingly large cumulative prefixes $P_1 \subset P_2 \subset \cdots \subset P_R$ with $P_R = \mathcal{D}_s$.

In each round, surviving candidates are evaluated on the current prefix, and the top candidates are retained according to their round scores. The evaluation budget thus becomes increasingly concentrated on competitive candidates.

Warm-starting from search observations. Rather than discarding evaluations collected during evolution, we reuse all available outcomes. Let O_c denote the items already observed for candidate c during search, and let $U_c^r = O_c \cup P_r$. Only the missing observations in U_c^r are executed; cached outcomes are reused without additional cost. In the first round, the score is the observed mean over this union:

$$\widetilde{F}_1(c) = \frac{1}{|U_c^1|}\sum_{i\in U_c^1} Y_{ci}. \tag{7}$$

Regression-adjusted scores. Because O_c contains informative and partly candidate-specific items, union means of different candidates are computed on item sets of different difficulty. From the second round on, EvalAlloc corrects this with a regression (difference) estimator [3, 6]:

$$\widetilde{F}_r(c) = \frac{1}{|U_c^r|}\sum_{i\in U_c^r} Y_{ci} + \lambda_r\left(\frac{1}{N}\sum_{i=1}^{N}\tilde{q}_{ci} - \frac{1}{|U_c^r|}\sum_{i\in U_c^r}\tilde{q}_{ci}\right), \tag{8}$$

Table 3: Test performance of the program selected under each evaluation strategy. For each benchmark, Score reports the TEST mean over three seeds, with standard deviation shown as a subscript. Cost (%) is the evaluation spent during search and final selection, as a share of evaluating every proposal on all 300 search items (80×300 cells for ShinkaEvolve, 81×300 for AdaEvolve); programs that fail to load or error on every evaluated item cost nothing. Average reports the unweighted mean over the four benchmarks. The highest Score per benchmark and on average within each framework is bold.

Framework	Method	DROP		HotpotQA		HoVer		IFBench		Average	
		Score ↑	Cost ↓	Score ↑	Cost ↓	Score ↑	Cost ↓	Score ↑	Cost ↓	Score ↑	Cost ↓
ShinkaEvolve	Skip Eval	$0.7674_{\pm 0.1046}$	0.0	$0.4633_{\pm 0.4014}$	0.0	$0.3933_{\pm 0.0961}$	0.0	$0.3917_{\pm 0.0741}$	0.0	0.5039	0.0
	Successive Halving	$0.8738_{\pm 0.0248}$	22.1	$0.7089_{\pm 0.0139}$	22.0	$0.6844_{\pm 0.0948}$	22.5	$0.4433_{\pm 0.0098}$	21.2	0.6776	22.0
	Random	$0.8629_{\pm 0.0083}$	42.5	$0.7244_{\pm 0.0158}$	45.6	$0.7156_{\pm 0.0443}$	43.8	$0.4603_{\pm 0.0332}$	37.7	0.6908	42.4
	Cascade	$0.8643_{\pm 0.0139}$	39.6	$0.7178_{\pm 0.0135}$	53.3	$\mathbf{0.7289}_{\pm 0.0375}$	54.8	$0.4683_{\pm 0.0214}$	38.5	0.6948	46.6
	Full	$0.8832_{\pm 0.0155}$	77.1	$0.7133_{\pm 0.0067}$	92.5	$0.7267_{\pm 0.0176}$	92.5	$0.4660_{\pm 0.0290}$	75.4	0.6973	84.4
	EvalAlloc	$\mathbf{0.8907}_{\pm 0.0164}$	43.3	$\mathbf{0.7411}_{\pm 0.0234}$	46.0	$0.7278_{\pm 0.0847}$	46.9	$\mathbf{0.4739}_{\pm 0.0246}$	39.2	**0.7084**	43.8
AdaEvolve	Skip Eval	$0.7720_{\pm 0.0833}$	0.0	$0.6678_{\pm 0.0287}$	0.0	$0.4933_{\pm 0.1389}$	0.0	$0.4303_{\pm 0.0398}$	0.0	0.5908	0.0
	Successive Halving	$0.8559_{\pm 0.0480}$	22.8	$0.6922_{\pm 0.0402}$	23.2	$0.5622_{\pm 0.1155}$	23.1	$0.4478_{\pm 0.0354}$	23.1	0.6395	23.1
	Random	$0.8759_{\pm 0.0098}$	49.2	$0.7167_{\pm 0.0384}$	49.8	$0.7178_{\pm 0.0877}$	49.6	$0.4473_{\pm 0.0173}$	48.6	0.6894	49.3
	Cascade	$0.8204_{\pm 0.0398}$	58.6	$0.7378_{\pm 0.0356}$	53.8	$\mathbf{0.7922}_{\pm 0.1091}$	60.3	$0.4518_{\pm 0.0221}$	51.4	0.7006	56.0
	Full	$0.8840_{\pm 0.0155}$	100.0	$0.7267_{\pm 0.0296}$	98.8	$0.7733_{\pm 0.0899}$	100.0	$0.4524_{\pm 0.0123}$	97.9	0.7091	99.2
	EvalAlloc	$\mathbf{0.8881}_{\pm 0.0087}$	46.3	$\mathbf{0.7456}_{\pm 0.0310}$	46.9	$0.7844_{\pm 0.0619}$	47.3	$\mathbf{0.4722}_{\pm 0.0513}$	45.3	**0.7226**	46.4

where $\tilde{q}_{ci}$ is the prediction of a refitted response model. The second term adds the predicted difference between the full SEARCH set and the observed items, so $\lambda_r = 0$ recovers the union mean and $\lambda_r = 1$ the classical difference estimator. We estimate λ_r from the data as the pooled within-candidate regression slope of observed outcomes on predictions over the surviving candidates C_r,

$$\lambda_r = \frac{\sum_{c \in C_r} \sum_{i \in U_c^r} (Y_{ci} - \bar{Y}_c)(\tilde{q}_{ci} - \bar{\tilde{q}}_c)}{\sum_{c \in C_r} \sum_{i \in U_c^r} (\tilde{q}_{ci} - \bar{\tilde{q}}_c)^2}, \tag{9}$$

with means taken over U_c^r and the result clipped to $[0, 1.5]$. This slope approximates the variance-minimizing coefficient of the difference estimator and automatically shrinks the correction when predictions track the observed outcomes poorly.

Cross-fitted refit. After the first round, we refit the response model on all search-time and round-1 observations. To avoid fitting the item parameters that correct a candidate on that candidate's own observations, we use two-fold cross-fitting [5]: the sparse (non-warm-up) candidates are randomly split into two folds, and each fold is scored by a model fitted on the warm-up candidates and the other fold. Warm-up candidates are fully observed, so their correction term is zero. Likewise, once the last round observes all N items, $U_c^R = \mathcal{D}_s$ and Eq. (8) reduces to the observed full-SEARCH mean, so the final pick is always decided by observed outcomes.

Evaluation cost. For M valid candidates, the total evaluator cost in executed cells is

$$C_{\text{eval}} = WN + (M - W)k + C_{\text{SH,new}}, \tag{10}$$

where the first two terms account for calibration and sparse search, and $C_{\text{SH,new}}$ counts only additional executions required by successive halving. With 80 valid proposals, the search stage costs $(WN + 70k)/(80N) = 30\%$ of full evaluation. Thus, EvalAlloc reuses evaluation effort across stages: dense observations establish the scoring reference, shared sparse observations guide evolution, and cached evaluations reduce the incremental cost of final confirmation.

6 Experiments

6.1 Setup

Baselines. We compare EvalAlloc against five evaluation strategies with different budget allocations. **Skip Eval** performs no evaluation during either search or final selection, and select the last valid evolved program. **Successive Halving (SH)** [14] provides no evaluation feedback during search and applies successive halving to the generated candidate pool for final selection. **Random** evaluates each candidate on a randomly sampled 50% subset of SEARCH items. **Cascade** [2, 23] first evaluates candidates on 20% of SEARCH items and promotes candidates meeting a top-25% threshold to full evaluation, while allocating no additional evaluations to the remaining candidates. **Full** evaluates every candidate on all SEARCH items.

Implementation Details. We use the same benchmarks, evolution frameworks, and model configurations as in Section 4.1. Unless otherwise specified, EvalAlloc allocates 30% of the full-evaluation budget to search, with the first $W = 10$ valid candidates fully evaluated and each subsequent candidate evaluated on $k = 60$ items (20% of SEARCH). The shared informative core contains $k_0 = 54$ items (90% of k), and the Gumbel sampling temperature is $\tau = 0.001$. For final selection, we use $R = 4$ successive-halving rounds with cumulative item prefixes of 30, 60, 120, and 300, retaining 30, 15, 8, and 1 candidates, respectively. For the response model, we set the latent dimension to $d = 2$, the regularization coefficient to $\rho = 1$, and the informativeness uncertainty penalty to $z = 2$.

6.2 Main Results

Table 3 compares the TEST performance and evaluation cost of different strategies. Across both evolution frameworks and all four benchmarks, EvalAlloc matches or slightly exceeds the mean

Table 4: Ablation of EvalAlloc on ShinkaEvolve. For each benchmark, we report the score over three runs, with standard deviation shown as a subscript. Each variant removes one component and keeps the rest unchanged. Δ is the change in the four-benchmark average relative to EvalAlloc. Cost (%) is the evaluation spent during search and final selection. Best value per column is bold.

Variant	DROP	HotpotQA	HoVer	IFBench	Avg.	Δ	Cost
EvalAlloc	$0.8907_{\pm 0.0164}$	$\mathbf{0.7411}_{\pm 0.0234}$	$\mathbf{0.7278}_{\pm 0.0847}$	$0.4739_{\pm 0.0246}$	**0.7084**	–	43.8%
w/o selection	$0.8755_{\pm 0.0157}$	$0.7278_{\pm 0.0139}$	$0.6922_{\pm 0.0683}$	$0.4688_{\pm 0.0363}$	0.6911	−0.0173	27.4%
w/o response model	$\mathbf{0.8946}_{\pm 0.0113}$	$0.7222_{\pm 0.0139}$	$0.6800_{\pm 0.0361}$	$\mathbf{0.4892}_{\pm 0.0513}$	0.6965	−0.0119	44.2%
w/o informative items	$0.8866_{\pm 0.0028}$	$0.6967_{\pm 0.0260}$	$0.6689_{\pm 0.0369}$	$0.4649_{\pm 0.0326}$	0.6793	−0.0291	43.4%

TEST score of Full while using approximately half as many evaluation cells. It also outperforms Random and Cascade in cross-benchmark average performance at comparable or lower evaluation cost. These results demonstrate that allocating evaluations selectively can provide a better cost–performance tradeoff than uniformly evaluating candidates. The evaluation costs differ between the two frameworks partly due to their handling of invalid proposals. AdaEvolve retries program generation when a proposal is invalid, whereas ShinkaEvolve discards invalid proposals without replacement. Since invalid programs incur no evaluation cost, ShinkaEvolve evaluates fewer valid candidates, resulting in lower realized evaluation costs, particularly under Full evaluation.

The no-evaluation baselines further highlight the complementary roles of search feedback and final selection. Applying SH to candidates generated without evaluation feedback substantially improves final performance, but still falls short of EvalAlloc. This suggests that final selection can recover promising candidates even from unguided search, while informative feedback during evolution provides additional benefits. These findings support combining sparse search-time evaluation with targeted final confirmation.

6.3 Ablation Study

We ablate three components of EvalAlloc on ShinkaEvolve across all four benchmarks. *W/o selection* skips final confirmation and selects the candidate with the highest refitted response-model estimate from the same search runs. *W/o response model* uses raw observed means for search feedback and raw union means for final selection. *W/o informative items* replaces informative item selection with random sampling while preserving the shared core and evaluation budget. The latter two variants require new evolution runs because they change the search-time feedback.

Table 4 shows that removing final confirmation consistently reduces TEST performance, highlighting the importance of additional evaluation for selecting among competitive candidates. Replacing informative items with random ones produces the largest average degradation, suggesting that sharing evaluation items alone is insufficient without considering their informativeness. Removing the response model has a smaller and less consistent effect, as the shared core already provides substantial comparability among candidates. The response model therefore primarily serves to calibrate sparse feedback onto a common scoring scale. Overall, these results identify informative item selection and final confirmation as the main contributors to EvalAlloc.

6.4 Analysis

Budget Allocation in Warm-up Stage. Figure 4 examines the effect of warm-up length W on DROP with ShinkaEvolve under a fixed search-stage budget. A longer warm-up provides more observations for calibration but leaves fewer evaluation items for subsequent candidates. The default $W = 10$ achieves the highest mean performance, although the differences are within seed-to-seed variation. All three configurations outperform Random at comparable evaluation cost. Offline replays further show that an excessively long warm-up reduces search-time ranking accuracy as later evaluations become too sparse. These results highlight the trade-off between upfront calibration and informative feedback during subsequent evolution.

Budget Allocation in Search Stage. Figure 5 examines the effect of search-time evaluation budget on DROP with ShinkaEvolve. We fix the warm-up at ten fully evaluated candidates and vary the target search budget from 20% to 40%. The default 30% allocation achieves the highest observed mean TEST performance, while the 20% configuration offers comparable performance at lower cost. Increasing the search-time budget does not consistently improve final performance in this setting, suggesting that sparse feedback can be sufficient to guide productive evolution.

Budget Allocation in Selection Stage. Figure 6 examines the effect of final confirmation budget while reusing the same search runs. Compared with direct selection based on search-stage estimates, successive halving (SH) improves TEST performance across all four benchmarks. Applying SH to the top 30 candidates captures most of the improvement at lower cost, while considering all candidates provides a modest additional gain. These results highlight the value of targeted final confirmation and the trade-off between selection cost and final scaffold quality.

Evaluation cost. Table 5 compares evaluation-cell and executor-token costs relative to Full. EvalAlloc reduces token consumption across all benchmark–framework settings, although token savings are not always proportional to cell savings. On ShinkaEvolve, it uses 52.13% of Full's cells but 64.76% of its tokens, whereas its token share on AdaEvolve is 45.47%. The latter average is influenced by an unusually token-intensive HoVer Full run.

To examine this discrepancy without confounding differences in search trajectories, we analyze the fully evaluated candidate pools from the Full runs. Higher-scoring candidates tend to consume more tokens per item, with average within-run Spearman correlations of 0.25 and 0.36 for ShinkaEvolve and AdaEvolve, respectively. Likewise, the 60 items with the highest response variability consume

, ,

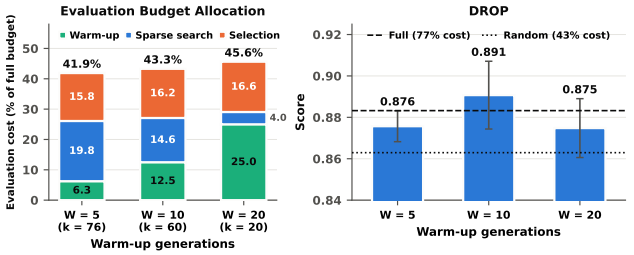

Figure 4: Effect of warm-up length on DROP with ShinkaEvolve under a fixed search-stage budget. Left: evaluation cost breakdown for different numbers of fully evaluated warm-up candidates. Right: corresponding performance, with Full and Random shown as references.

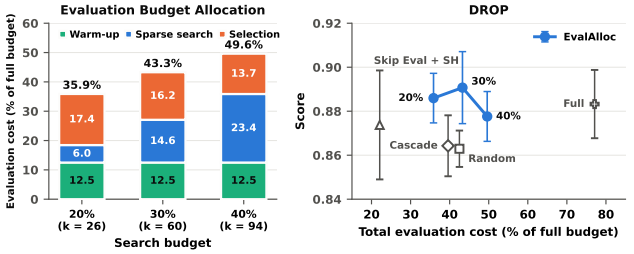

Figure 5: Effect of search-time evaluation budget on DROP with ShinkaEvolve. Left: evaluation cost breakdown across warm-up, sparse search, and final selection. Right: Performance versus total evaluation cost under different search budgets.

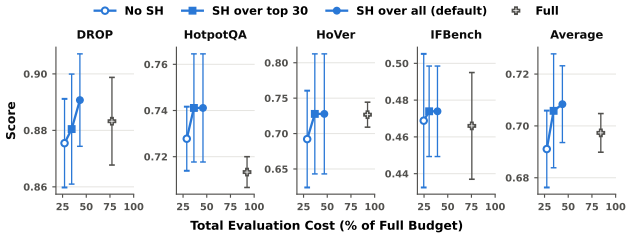

Figure 6: Effect of final selection strategies on TEST performance and evaluation cost across four benchmarks and their average. We compare selection without SH, SH over the top 30 candidates, and SH over all candidates, using the same search runs. Full evaluation is shown as a reference.

1.08 and 1.13 times the average tokens per item. These patterns suggest that concentrating evaluations on informative items and competitive candidates can favor more expensive execution cells. Thus, evaluation-cell counts measure the number of executions, while token counts additionally reflect heterogeneity in execution cost.

Other Proposers and Executor. We further evaluate alternative model configurations on ShinkaEvolve in Figure 7. We replace the proposer with gpt-oss-120b [1] on IFBench, or the executor with

Table 5: Evaluation-cell and executor-token (prompt and completion) costs of search and final selection, relative to Full (100%). Each entry averages the per-benchmark shares over the four benchmarks; programs that fail to load or error on every evaluated item cost nothing.

Method	ShinkaEvolve		AdaEvolve	
	Cells	Tokens	Cells	Tokens
Full	100.00%	100.00%	100.00%	100.00%
Successive Halving	26.26%	27.50%	23.26%	15.10%
Random	50.44%	55.18%	49.69%	43.66%
Cascade	54.83%	66.32%	56.47%	54.50%
EvalAlloc	52.13%	64.76%	46.83%	45.47%

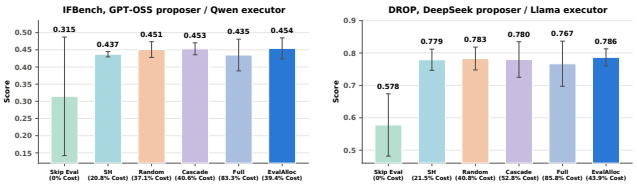

Figure 7: Performance with alternative proposer and executor models on ShinkaEvolve. Left: IFBench with a gpt-oss-120b proposer. Right: DROP with a Llama-3.1-8B-Instruct executor. Bars show the mean over three seeds, with standard deviation as error bars. Percentages indicate evaluation cost relative to the nominal full budget.

Llama-3.1-8B-Instruct [11] on DROP. In both settings, EvalAlloc achieves competitive TEST performance with substantially fewer evaluation cells than Full. The relative performance of evaluation strategies varies across model configurations, suggesting that the benefits of evaluation allocation can depend on the underlying models.

7 Conclusion

Evaluation influences both evolutionary search and final scaffold selection. Our investigation shows that sparse informative feedback can guide evolution, while shared evaluation items and a stable scoring reference support comparable decisions across candidates. Motivated by these findings, we propose EvalAlloc, which combines sparse search-time feedback with warm-started final confirmation. Across four benchmarks and two evolution frameworks, EvalAlloc achieves competitive performance while using approximately half as many evaluation cells as Full. These results highlight the importance of allocating evaluation effort according to its distinct roles in search and final selection.

References

[1] Sandhini Agarwal, Lama Ahmad, Jason Ai, Sam Altman, Andy Applebaum, Edwin Arbus, Rahul K Arora, Yu Bai, Bowen Baker, Haiming Bao, et al. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.

[2] Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, et al. Gepa: Reflective prompt evolution can outperform reinforcement learning. In *International Conference on Learning Representations*, volume 2026, pages 8479–8565, 2026.

[3] Anastasios N Angelopoulos, John C Duchi, and Tijana Zrnic. Ppi++: Efficient prediction-powered inference. *The Annals of Applied Statistics*, 20(3):2235–2252, 2026.

[4] Mert Cemri, Shubham Agrawal, Akshat Gupta, Shu Liu, Audrey Cheng, Qiuyang Mang, Ashwin Naren, Lutfi Eren Erdogan, Koushik Sen, Matei Zaharia, et al. Adaevolve: Adaptive llm driven zeroth-order optimization. *arXiv preprint arXiv:2602.20133*, 2026.

[5] Victor Chernozhukov, Denis Chetverikov, Mert Demirer, Esther Duflo, Christian Hansen, Whitney Newey, and James Robins. Double/debiased machine learning for treatment and structural parameters, 2018.

[6] William Gemmell Cochran and William G Cochran. *Sampling techniques*, volume 3. wiley, New York, 1977.

[7] Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. Drop: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2368–2378, 2019.

[8] Adam Fisch, Daniel Deutsch, Joshua Maynez, Alekh Agarwal, Jonathan Berant, William Cohen, Amir Globerson, and Jacob Eisenstein. Collabeval: Statistically efficient collaborative model evaluation via matrix completion. *arXiv preprint arXiv:2607.05046*, 2026.

[9] Yonatan Gideoni, Sebastian Risi, and Yarin Gal. Simple baselines are competitive with code evolution. *arXiv preprint arXiv:2602.16805*, 2026.

[10] Paul Glasserman and David D Yao. Some guidelines and guarantees for common random numbers. *Management Science*, 38(6):884–908, 1992.

[11] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

[12] Eliya Habba, Itay Itzhak, Asaf Yehudai, Yotam Perlitz, Elron Bandel, Michal Shmueli-Scheuer, Leshem Choshen, and Gabriel Stanovsky. Growing pains: Extensible and efficient llm benchmarking via fixed parameter calibration. *arXiv preprint arXiv:2604.12843*, 2026.

[13] Shengran Hu, Cong Lu, and Jeff Clune. Automated design of agentic systems. In *International Conference on Learning Representations*, volume 2025, pages 21344–21377, 2025.

[14] Kevin Jamieson and Ameet Talwalkar. Non-stochastic best arm identification and hyperparameter optimization. In *Artificial intelligence and statistics*, pages 240–248. PMLR, 2016.

[15] Yichen Jiang, Shikha Bordia, Zheng Zhong, Charles Dognin, Maneesh Singh, and Mohit Bansal. Hover: A dataset for many-hop fact extraction and claim verification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3441–3460, 2020.

[16] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th symposium on operating systems principles*, pages 611–626, 2023.

[17] Robert Lange, Yuki Imajuku, and Edoardo Cetin. Shinkaevolve: Towards open-ended and sample-efficient program evolution. In *International Conference on Learning Representations*, volume 2026, pages 74026–74078, 2026.

[18] Lisha Li, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. Hyperband: A novel bandit-based approach to hyperparameter optimization. *Journal of machine learning research*, 18(185):1–52, 2018.

[19] Shu Liu, Mert Cemri, Shubham Agarwal, Alexander Krentsel, Ashwin Naren, Qiuyang Mang, Zhifei Li, Akshat Gupta, Monishwaran Maheswaran, Audrey Cheng, et al. Skydiscover: A flexible, adaptive framework for ai-driven scientific and algorithmic discovery. In *Proceedings of the ACM Conference on AI and Agentic Systems*, pages 1223–1227, 2026.

[20] Manuel López-Ibáñez, Jérémie Dubois-Lacoste, Leslie Pérez Cáceres, Mauro Birattari, and Thomas Stützle. The irace package: Iterated racing for automatic algorithm configuration. *Operations Research Perspectives*, 3:43–58, 2016.

[21] Grégoire Mialon, Clémentine Fourrier, Thomas Wolf, Yann LeCun, and Thomas Scialom. Gaia: a benchmark for general ai assistants. In *International Conference on Learning Representations*, volume 2024, pages 9025–9049, 2024.

[22] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2021.

[23] Alexander Novikov, Ngân Vũ, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, Sergey Shirobokov, Borislav Kozlovskii, Francisco JR Ruiz, Abbas Mehrabian, et al. Alphaevolve: A coding agent for scientific and algorithmic discovery. *arXiv preprint arXiv:2506.13131*, 2025.

[24] Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. tinybenchmarks: evaluating LLMs with fewer examples. In *Forty-first International Conference on Machine Learning*, 2024. URL https://openreview.net/forum?id=qAml3FpfhG.

[25] Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hanna Hajishirzi. Generalizing verifiable instruction following. *Advances in Neural Information Processing Systems*, 38, 2026.

[26] Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, et al. Toolllm: Facilitating large language models to master 16000+ real-world apis. In *International Conference on Learning Representations*, volume 2024, pages 9695–9717, 2024.

[27] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL https://qwen.ai/blog?id=qwen3.5.

[28] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M Pawan Kumar, Emilien Dupont, Francisco JR Ruiz, Jordan S Ellenberg, Pengming Wang, Omar Fawzi, et al. Mathematical discoveries from program search with large language models. *Nature*, 625(7995):468–475, 2024.

[29] Anyi Xu, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, Chenchen Ling, et al. Deepseek-v4: Towards highly efficient million-token context intelligence. *arXiv preprint arXiv:2606.19348*, 2026.

[30] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D Manning. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 conference on empirical methods in natural language processing*, pages 2369–2380, 2018.

[31] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik R Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations*, 2023. URL https://openreview.net/forum?id=WE_vluYUL-X.

[32] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. In *International Conference on Learning Representations*, volume 2025, pages 9965–10017, 2025.

[33] Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. Darwin gödel machine: Open-ended evolution of self-improving agents. In *International Conference on Learning Representations*, volume 2026, pages 104223–104294, 2026.

[34] Jiayi Zhang, Jinyu Xiang, Zhaoyang Yu, Fengwei Teng, Xionghui Chen, Jiaqi Chen, Mingchen Zhuge, Xin Cheng, Sirui Hong, Jinlin Wang, et al. Aflow: Automating agentic workflow generation. In *International Conference on Learning Representations*, volume 2025, pages 34040–34077, 2025.

[35] Taolin Zhang, Hang Guo, Wang Lu, Tao Dai, Shu-Tao Xia, and Jindong Wang. Sparseeval: Efficient evaluation of large language models by sparse optimization. In *The Fourteenth International Conference on Learning Representations*, 2026. URL https://openreview.net/forum?id=CZAzAedGSV.

[36] Shuyan Zhou, Frank F Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, et al. Webarena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations*, volume 2024, pages 15585–15606, 2024.