arXiv:2611.04169v1 [cs.CL] 3 Nov 2026

Length Is Cheap, Competition Is Not: Attention Predicts Context Rot in Retrieval

Yinuo Liu

Department of Computer Science and Engineering
University of California San Diego, La Jolla, CA, USA

ORCID: 0009-0006-1114-4781

Abstract

Why do language models fail on long inputs? The common answer is length. We show that, for retrieval, length is cheap and competition is not. On five open models (0.5B to 3B parameters, four families), neutral filler costs nothing: accuracy stays at 0.97 to 1.00, up to 32K tokens on Qwen2.5-1.5B. Statements in the format of the evidence, placed at the same positions, drop it to 0.11 to 0.21. A single statement that gives the queried key a rival value multiplies the loss by 9 to 110, and every additional rival adds to it. Inside the model, one quantity tracks the failure: the attention mass left on the evidence. Across 126,473 trials it explains up to 58% of the variance in loss; log length explains at most 8%. The masked-filler protocol behind the claim that length alone hurts reproduces on MMLU (11 to 18 points lost at 30K tokens), but fails a visible control in retrieval: masking costs 46 accuracy points at 32K tokens, while the same filler left visible costs none. In retrieval, context rot is a competition for attention. We release the probe and all 126K trial-level measurements.

Keywords: context rot, long-context language models, attention dilution, distractors, knowledge conflicts, needle in a haystack, retrieval-augmented generation.

Code github.com/yliusci/length-is-cheap Data huggingface.co/datasets/yliusci/length-is-cheap

1 Introduction

Thirty-two thousand tokens of neutral text cost a 1.5B-parameter model nothing. A few sentences that contradict the evidence cost it most of its accuracy. This paper is about the gap between those two facts.

Language models now accept hundreds of thousands of tokens [1–7], yet they use far fewer [8–10]. Retrieval accuracy declines as inputs grow [8, 11–15], a degradation now called *context rot* [16, 17]. The prevailing explanation is length itself: Du et al. [18] hid the added tokens from attention, kept their positions, and still measured a loss. If length is the cause, the remedy is architectural, and every added token is a tax.

Length, however, never grows alone. Every added token is also added content, and in realistic inputs much of that content resembles what the model is looking for. Softmax attention [19] makes the distinction consequential. The weight on the evidence is

$$a^* = \frac{e^{s^*}}{e^{s^*} + \sum_{j \neq *} e^{s_j}},$$

where s^* is the attention score of the evidence and s_j the scores of all other keys. A thousand dissimilar tokens with low s_j add little to the denominator; a handful of near-copies of the evidence can halve a^*. If this is what happens inside the model, context rot is a problem of competition, not of size, the remedy is to curate what enters the context, and a^* should predict the loss better than length does.

We test this prediction directly. In a key-value retrieval task, we vary context length from 512 to 32K tokens, fill the same positions either with neutral material or with statements that share the format of the evidence, and measure on every trial both the log-probability of the correct answer and the attention mass a^* that the model places on the evidence. We study five open models from 0.5B to 3B parameters from four families. Figure 1 summarizes the outcome: at equal length, the loss follows what the context contains and how much attention the evidence keeps, not how many tokens surround it.

Our contributions are:

1. **Length is cheap; competition is not.** With neutral filler, accuracy stays between 0.97 and 1.00 on four of five models, up to 4K tokens on all four and up to 32K on Qwen2.5-1.5B. Filling the same positions with statements in the evidence template drives accuracy to 0.11 to 0.21 at 8K to 32K tokens. The loss concentrates on statements that give the queried key a rival value: one such statement multiplies the loss by 9 to 110, and the loss grows with their number at fixed length.
2. **Attention mass on the evidence predicts the loss.** Across 126,473 trials on five models, $\log(1/a^*)$ explains 35 to 58% of the variance in loss, against 2 to 8% for log length, and the gap persists from 8K to 32K tokens.
3. **Masking is not a length baseline.** Replicating the masked-filler protocol of Du et al. [18], MMLU accuracy falls by 11 to 18 points at 30K masked tokens on Llama-3.2-3B and Qwen2.5-7B. In retrieval, the same masking costs 46 points of accuracy at 32K tokens while the identical filler left visible costs none.
4. **An open probe.** We release the harness, the attention probe, and all trial-level data: code at `https://github.com/yliusci/length-is-cheap`, data at `https://huggingface.co/datasets/yliusci/length-is-cheap`.

2 Related Work

Long-context degradation. Accuracy falls as inputs grow, well before the nominal context limit [20–27]. Liu et al. [11] showed that retrieval depends on where evidence sits [see also 28–33]; RULER [8] found effective lengths far below advertised ones; NoLiMa [34] showed that removing lexical overlap between question and evidence accelerates the decline. Hong et al. [16] named the phenomenon *context rot* across 18 frontier models. These studies vary length and content together, so they cannot say which of the two is responsible.

Knowledge conflicts. When a context contains statements that disagree, models often follow the wrong one [35; 36; 37; 38; see 39, for a survey]. Proactive interference, in which earlier values of a key disrupt retrieval of the latest one, drives accuracy log-linearly toward zero as updates accumulate [40]. This literature treats conflicts as a reliability problem, separate from length. We show that, in long-context retrieval, rivals carry most of the loss that grows with length.

Length or content? Two lines of work give opposite answers. Du et al. [18] masked every distracting token from attention while preserving position indices and still observed drops of at least 7.9% at 30K tokens, concluding that length alone hurts. Conversely, studies that hold length fixed and vary what fills it find that semantically competing material dominates [41, 42]: hard distractors

Table 1: Controls in the closest prior work.

Study	Length fixed, content varied	Visible neutral filler	Masked filler	Attention vs. loss, per trial
Du et al. [18]	no	whitespace	yes	no
Gao et al. [43]	yes	repeated text	no	head-level margins
Repantis et al. [44]	yes	no	no	no
Hong et al. [16]	partial	no	no	no
Ours	**yes**	**three text types**	**yes, replicated**	**yes, 126K trials**

collapse accuracy within the first 10% of the context [43], competing passages shorten retention at constant passage count [44; see also 45; 46], and near-miss negatives, not corpus size, bound retrieval at scale [47, 48]. Earlier, Levy et al. [15] and Shi et al. [49] showed that similar padding or a single irrelevant sentence suffices to degrade reasoning in short contexts [50–52]. Only Gao et al. [43] inspect attention, through logit margins on retrieval heads; none relates the attention on the evidence to the loss trial by trial, so none can tell how much of the failure competitors explain by drawing attention away from the evidence.

Attention dilution theory. Softmax attention assigns each key a share that shrinks as keys accumulate. Chiang and Cholak [53] and Veličković et al. [54] proved that this dispersion limits sharp retrieval as length grows [55–57]; Nakanishi [58] and Chen et al. [59] derived length-dependent temperature corrections [60, 61]; Ghaffari et al. [62] modeled distractors as extreme-value competitors whose maximum score grows with their number, requiring the evidence margin to grow as $\sqrt{\log N}$. Tworkowski et al. [63] first described this *distraction issue*: relevant keys become a vanishing fraction of all keys, and similar keys compete. These accounts predict that degradation should track the attention mass left on the evidence. That prediction has not been tested trial by trial on measured attention in pretrained models.

Our position. We do not claim the phenomenon, the distractor effect, or the dilution theory. We close the loop between them, measuring behavior, content and attention on the evidence in the same trials (Table 1). Du et al. [18] also used a visible baseline, but filled it with whitespace, which we find harmful on MMLU as they did; with neutral text instead, length is free in retrieval. Their masked condition, which we replicate, is the only evidence for a pure length effect that remains, and we show that it fails a visible control.

3 Setup

One principle governs the design: every comparison holds length and evidence position fixed and changes only what surrounds the evidence. We call this the *visible-control protocol*. It needs forward passes only and applies to any open-weight model.

Task. Each context hides one *evidence* statement, "The secret code of the river is 4821.", where the key (here, river) is a common noun and the value a four-digit number. The context ends with the query "Question: what is the secret code of the river? Answer: The secret code of the river is", and we score the log-probability ℓ of the correct value under teacher forcing, summed over its tokens. Tokenizers split the value differently (five tokens for Qwen and SmolLM2, three for OLMo and Llama), so we compare losses within models, not across them. The evidence sits at a depth drawn uniformly over the context.

Conditions. For each trial we build four contexts that share the evidence, the query, and the total length L:

Condition	Content around the evidence
short	nothing
neutral	random numbers, random words, or one repeated sentence
competing	statements in the evidence template, keys drawn uniformly from 51 nouns
masked	filler hidden from attention, positions preserved

We define the loss of a condition c as $\Delta_c = \log(p_{\text{short}}/p_c)$, in nats, where $p_c = e^{\ell_c}$ is the probability of the correct value in condition c. The *neutral* condition isolates length; the difference between *competing* and *neutral* isolates content at fixed length.

Because keys are drawn uniformly, some competing statements share the queried key and assign it a different value. These *conflicts* are the strongest competitors the filler contains. Their number n_{conf} grows with length (mean 0.8 at 512 tokens, 3.2 at 2,048, 50.5 at 32,768; none in 42 to 48% of 512-token contexts and in no context of 4,096 tokens or more). We recover n_{conf} for every trial by regenerating its context from the recorded seed, which lets us separate competitors that share the key from those that only share the format. Separately, we insert 0 to 5 explicit distractors that share the format but never the key.

Masking. In *masked*, a 2D attention mask removes every filler token from attention for all subsequent queries while position indices keep counting through them, so the evidence stays at its original distance from the query and every competitor becomes invisible. The first tokens of the sequence remain visible, preserving the attention sink [64–70].

Models and lengths. We test Qwen2.5-0.5B-Instruct and Qwen2.5-1.5B-Instruct [7], SmolLM2-1.7B-Instruct [71], OLMo-2-0425-1B-Instruct [72] and the base Llama-3.2-3B [6], at $L \in \{512, 1024, 2048\}$ tokens for all models and up to 32,768 tokens within each model's context window. All runs use float16 weights and PyTorch SDPA attention, on one Apple workstation (64 GB unified memory) and one RTX 4090 laptop GPU (16 GB).

Attention mass on the evidence. For each configuration we select, in the *short* context, the eight (layer, head) pairs that place the most attention on the evidence tokens [73–75]. In every condition, a^* is the attention that these heads send from the final query token to the evidence tokens, summed over evidence tokens and averaged over heads. We do not renormalize a^* to exclude the attention sink. We compare predictors of Δ by the coefficient of determination R^2 of univariate least-squares fits on a^*, $\log(1/a^*)$ and $\log L$.

Replication of Du et al. [18]. We reproduce their masked-filler protocol on 1,000 MMLU questions with the base Llama-3.2-3B and Qwen2.5-7B-Instruct, inserting 3,750, 7,500, 15,000 or 30,000 masked tokens between the question and its answer options, following their layout, in which the masked positions separate the evidence from the question. We implement masking by deleting the filler tokens and advancing the position indices by the same amount, which we verified to match explicit attention masking on Llama. We report accuracy with Wilson 95% intervals and paired McNemar tests against the unpadded baseline.

4 Results

One comparison condenses the paper. On Qwen2.5-1.5B at 32K tokens, with length and evidence position fixed, the median cost of neutral numeric text is 0.004 nats, the cost of masking it is 1.20 nats, and the cost of competing statements is 4.76 nats.

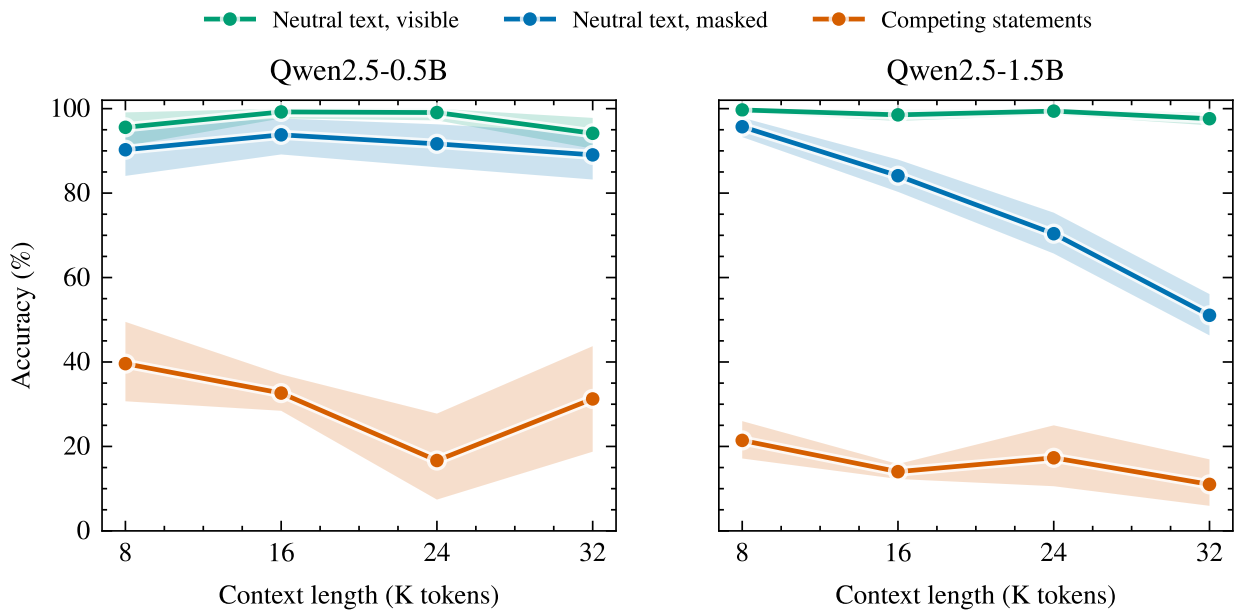

Figure 1: Retrieval accuracy against context length, with length and evidence position fixed. Neutral text left visible (green) costs almost nothing; the same text masked at preserved positions (blue) costs up to 47 points on Qwen2.5-1.5B; competing statements (orange) cost 56 to 87 points at every length. Mean over trials, pooled over the three neutral fillers; shaded bands are 95% bootstrap intervals.

Table 2: Qwen2.5-1.5B, neutral filler (numbers). Accuracy with the filler visible and with the same filler masked at preserved positions; median masked loss Δ_{masked} in nats.

L	8K	16K	24K	32K
Accuracy, visible	1.00	0.99	1.00	0.98
Accuracy, masked	0.93	0.87	0.73	0.52
Δ_{masked}	0.03	0.11	0.44	1.20

4.1 Length is cheap

Visible neutral filler barely moves the target. On Qwen2.5-1.5B, the median loss stays at or below 0.006 nats at every length from 8K to 32K tokens and for all three neutral fillers, and accuracy stays between 0.96 and 1.00 (Table 2). Qwen2.5-0.5B, SmolLM2 and Llama-3.2-3B show the same behavior up to 4K tokens. OLMo-2-1B is the exception: it loses 0.57 nats with neutral filler at 2,048 tokens (accuracy 0.91).

4.2 Competition is not

Filling the same positions with competing statements instead of neutral text turns a free context into a failing one. On Qwen2.5-1.5B, the median loss reaches 2.82, 3.94, 4.03 and 4.76 nats at 8K, 16K, 24K and 32K tokens, and accuracy falls to 0.11 to 0.21; on Qwen2.5-0.5B, the loss reaches 1.79 to 4.07 nats and accuracy 0.17 to 0.40. Length is the same in both conditions; only content differs (Figure 1).

The loss concentrates on conflicts (Table 3). Across five models, contexts without any statement

Table 3: Mean loss (nats) by presence of a same-key conflict, $L \leq 2{,}048$.

Model	No conflict	At least one
Qwen2.5-0.5B	0.175	1.812
Qwen2.5-1.5B	0.020	2.238
SmolLM2-1.7B	0.022	1.381
OLMo-2-1B	0.058	2.796
Llama-3.2-3B	0.344	3.239

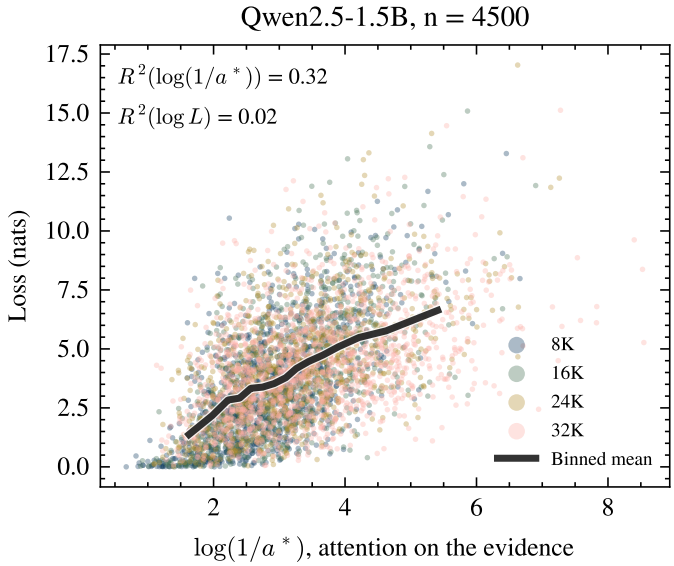

Figure 2: Loss Δ (nats) against $\log(1/a^*)$ for 4,500 trials of Qwen2.5-1.5B with competing filler at 8K to 32K tokens, colored by length. Black: mean loss in 15 quantile bins of $\log(1/a^*)$. Lengths overlap completely, while attention orders the loss.

that shares the queried key lose 0.02 to 0.34 nats; contexts with at least one lose 1.38 to 3.24 nats. Accuracy is 0.97 to 1.00 without conflict and 0.25 to 0.52 with one. At fixed length the loss grows with the number of conflicts: on Qwen2.5-1.5B at 2,048 tokens it is 0.03, 1.37, 2.24, 2.48, 2.71, 3.26 and 3.45 nats for 0 to 6 or more. Competitors that share only the format matter far less: the number of explicit distractors explains at most 0.3% of the variance in loss ($R^2 \leq 0.003$).

4.3 One quantity predicts the failure: attention on the evidence

Across 126,473 trials, $\log(1/a^*)$ predicts the loss on every model, and log length barely does (Table 4). $R^2(\log(1/a^*))$ ranges from 0.35 to 0.58, against 0.02 to 0.08 for $\log L$; a bootstrap confidence interval on the difference excludes zero for each model, and the ordering holds within each (length, depth) cell. Attention on the evidence collapses from 0.51 to 0.77 in the short context to 0.17 to 0.31 with competing filler. The relation extends to long contexts: from 8K to 32K tokens, $R^2(\log(1/a^*))$ is 0.32 on Qwen2.5-1.5B and 0.33 on Llama-3.2-3B, against 0.02 and 0.03 for $\log L$; on SmolLM2-1.7B, up to its 8K window, it is 0.45 against 0.18.

a^* summarizes the effect of conflicts rather than adding to it: once $\log(1/a^*)$ is known, n_{conf} raises R^2 by at most 0.04. Within contexts free of conflicts, where losses are small, $\log(1/a^*)$ explains only 2 to 10% of the variance. The relation has a threshold: on Qwen2.5-1.5B the loss stays near

Table 4: Predictors of the loss, $L \in \{512, 1024, 2048\}$. Slope of Δ on $\log(1/a^*)$ with 95% CI.

Model	n	$R^2(a^*)$	$R^2(\log(1/a^*))$	$R^2(\log L)$	Slope
Qwen2.5-0.5B	32,318	0.23	0.35	0.03	1.40 [1.38, 1.42]
Qwen2.5-1.5B	40,632	0.43	0.58	0.05	2.22 [2.21, 2.24]
SmolLM2-1.7B	9,325	0.30	0.41	0.07	1.68 [1.64, 1.72]
OLMo-2-1B	42,005	0.26	0.35	0.02	3.97 [3.92, 4.02]
Llama-3.2-3B	2,193	0.36	0.45	0.08	2.75 [2.63, 2.88]

Table 5: MMLU accuracy (%) with masked filler, $n = 1{,}000$ per cell, Wilson 95% intervals.

Masked tokens	0	3,750	7,500	15,000	30,000
Llama-3.2-3B	51.7 [48.6, 54.8]	44.7 [41.6, 47.8]	42.4 [39.4, 45.5]	39.3 [36.3, 42.4]	40.5 [37.5, 43.6]
Qwen2.5-7B	65.8 [62.8, 68.7]	62.6 [59.6, 65.5]	54.5 [51.4, 57.6]	49.4 [46.3, 52.5]	48.1 [45.0, 51.2]

zero until a^* falls to about 0.35, then rises steeply. The model tolerates losing about half of the attention it gives the evidence in a short context, then fails abruptly. Slopes exceed 1 on every model and differ across models, so we report no universal law; the loss grows faster than the removal of evidence mass alone would imply under a single-head softmax reading.

4.4 Masking is not a length control

Masking reproduces the length-alone effect of Du et al. [18]. On MMLU, accuracy falls with the number of masked tokens on both models, plateauing beyond 15,000 on Llama-3.2-3B (Table 5; all paired McNemar tests against the baseline $p \leq 0.007$). From 7,500 tokens onward, Qwen2.5-7B loses 11.3 to 17.7 points, against 15.5 to 21.1 for Llama-3.1-8B in Du et al. [18, Table 3], so the effect is not specific to one model family.

The same intervention degrades retrieval, whereas the filler it hides does not. In Table 2, the masked context loses 1.20 nats and 46 points of accuracy at 32K tokens, while the identical filler left visible loses nothing. Across the three neutral fillers at 32K, accuracy is 0.96 to 0.99 when visible and 0.48 to 0.54 when masked. A context of masked tokens at preserved positions is therefore not a context of irrelevant tokens: it presents the model with long positional gaps that it never encounters in training [9, 76–80], and its cost should not be read as the cost of length. The effect is model-dependent: on Qwen2.5-0.5B, masking costs 0.1 to 0.4 nats.

MMLU shows the same ordering at 4,096 tokens on Qwen2.5-1.5B: 59.1% without filler, 57.5% with visible neutral filler, 53.9% with the filler masked. Length is not free in every form, however: visible whitespace padding lowers the same model to 36.0% at 4,096 tokens, an effect absent at 2,048. We have no visible-filler arm on MMLU beyond 4K tokens, so the drops in Table 5 cannot be attributed to the positional gap alone.

5 Discussion

Curate, do not truncate. In retrieval, thousands of neutral tokens cost nothing, while a few statements that compete with the evidence cost most of the accuracy. The effective budget of a context is therefore not its token count but the number of statements that compete for the query. For retrieval-augmented generation [81, 82], agent memory [83] and long-document pipelines, this reverses a common priority: deduplicating and resolving conflicting statements before they reach

the model [84–88] should matter more than shortening the context.

Measure attention, not only accuracy. Conflicts, length and depth reach the loss through one measurable quantity: once $\log(1/a^*)$ is known, the number of conflicts adds at most 0.04 to R^2. Softmax dilution accounts predict exactly this [62, 63, 89]; we observe it on measured attention, trial by trial. Because a^* requires only a forward pass, it can serve as a diagnostic for any open model and any retrieval benchmark. The threshold (near $a^* \approx 0.35$ on Qwen2.5-1.5B, at a model-dependent point elsewhere) and slopes above 1 indicate a nonlinear readout: the model absorbs a moderate loss of evidence mass, then fails abruptly. We predict that methods which sharpen attention at long range [58, 59, 90–100] help most when they keep a^* above it.

Report a visible control. Masking tokens at preserved positions creates long positional gaps absent from pretraining [101–103]. In retrieval, this intervention costs 46 accuracy points where the same filler left visible costs none, so its cost is not the cost of length. We recommend that any claim about the effect of length include a visible neutral-text control at the same positions. Whitespace padding on MMLU shows that some forms of length do hurt reasoning, so the two accounts are not exclusive: competition dominates retrieval, and length effects remain possible in reasoning.

6 Limitations

a^* is measured in the same forward pass as the loss and we ran no intervention on it, so our evidence is predictive, not causal. The predictor analysis covers 512 to 2,048 tokens on five models and 8K to 32K on two. In our competing filler, the number of conflicts grows with length, and every context of 4,096 tokens or more contains at least one, so the effect of format-only competitors above 2K tokens is not measured; the conflict analysis was not preregistered. a^* is read at a single query token, from eight heads, without removing the attention sink. Value tokenization differs across models, so losses are comparable within but not across models. All results come from one synthetic retrieval task and models of at most 3B parameters, except the 7B model in the MMLU replication, which has no visible-filler arm beyond 4K tokens.

7 Conclusion

In retrieval, long contexts do not fail because they are long. They fail when what they contain competes with the evidence, and the attention that the evidence keeps tells us when. Context engineering, not context length, is the lever.

References

[1] Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending Context Window of Large Language Models via Positional Interpolation. *arXiv preprint arXiv:2306.15595*, 2023.

[2] Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient Context Window Extension of Large Language Models. *arXiv preprint arXiv:2309.00071*, 2023.

[3] Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, et al. LongRoPE: Extending LLM Context Window Beyond 2 Million Tokens. *arXiv preprint arXiv:2402.13753*, 2024.

[4] Amirkeivan Mohtashami and Martin Jaggi. Landmark Attention: Random-Access Infinite Context Length for Transformers. *arXiv preprint arXiv:2305.16300*, 2023.

[5] Yuhuai Wu, Markus N. Rabe, DeLesley Hutchins, and Christian Szegedy. Memorizing Transformers. *arXiv preprint arXiv:2203.08913*, 2022.

[6] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, et al. The Llama 3 Herd of Models. *arXiv preprint arXiv:2407.21783*, 2024.

[7] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, et al. Qwen2.5 Technical Report. *arXiv preprint arXiv:2412.15115*, 2024.

[8] Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekesh, Fei Jia, et al. RULER: What's the Real Context Size of Your Long-Context Language Models? *arXiv preprint arXiv:2404.06654*, 2024.

[9] Chenxin An, Jun Zhang, Ming Zhong, Lei Li, Shansan Gong, Yao Luo, et al. Why Does the Effective Context Length of LLMs Fall Short? *arXiv preprint arXiv:2410.18745*, 2024.

[10] Taiming Lu, Muhan Gao, Kuai Yu, Adam Byerly, and Daniel Khashabi. Insights into LLM Long-Context Failures: When Transformers Know but Don't Tell. *arXiv preprint arXiv:2406.14673*, 2024.

[11] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, et al. Lost in the Middle: How Language Models Use Long Contexts. *arXiv preprint arXiv:2307.03172*, 2023.

[12] Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, et al. LongBench: A Bilingual, Multitask Benchmark for Long Context Understanding. *arXiv preprint arXiv:2308.14508*, 2023.

[13] Xinrong Zhang, Yingfa Chen, Shengding Hu, Zihang Xu, Junhao Chen, Moo Khai Hao, et al. ∞Bench: Extending Long Context Evaluation Beyond 100K Tokens. *arXiv preprint arXiv:2402.13718*, 2024.

[14] Yuri Kuratov, Aydar Bulatov, Petr Anokhin, Ivan Rodkin, Dmitry Sorokin, Artyom Sorokin, et al. BABILong: Testing the Limits of LLMs with Long Context Reasoning-in-a-Haystack. *arXiv preprint arXiv:2406.10149*, 2024.

[15] Mosh Levy, Alon Jacoby, and Yoav Goldberg. Same Task, More Tokens: the Impact of Input Length on the Reasoning Performance of Large Language Models. *arXiv preprint arXiv:2402.14848*, 2024.

[16] Kelly Hong, Anton Troynikov, and Jeff Huber. Context Rot: How Increasing Input Tokens Impacts LLM Performance. Chroma Technical Report, `https://research.trychroma.com/context-rot`, 2025.

[17] Shijie Xia, Yikun Wang, Zhen Huang, and Pengfei Liu. Diagnosing and Mitigating Context Rot in Long-horizon Search. *arXiv preprint arXiv:2606.29718*, 2026.

[18] Yufeng Du, Minyang Tian, Srikanth Ronanki, Subendhu Rongali, Sravan Bodapati, Aram Galstyan, et al. Context Length Alone Hurts LLM Performance Despite Perfect Retrieval. *arXiv preprint arXiv:2510.05381*, 2025.

[19] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, et al. Attention Is All You Need. *arXiv preprint arXiv:1706.03762*, 2017.

[20] Chenxin An, Shansan Gong, Ming Zhong, Xingjian Zhao, Mukai Li, Jun Zhang, et al. L-Eval: Instituting Standardized Evaluation for Long Context Language Models. *arXiv preprint arXiv:2307.11088*, 2023.

[21] Uri Shaham, Maor Ivgi, Avia Efrat, Jonathan Berant, and Omer Levy. ZeroSCROLLS: A Zero-Shot Benchmark for Long Text Understanding. *arXiv preprint arXiv:2305.14196*, 2023.

[22] Yushi Bai, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, et al. LongBench v2: Towards Deeper Understanding and Reasoning on Realistic Long-context Multitasks. *arXiv preprint arXiv:2412.15204*, 2024.

[23] Howard Yen, Tianyu Gao, Minmin Hou, Ke Ding, Daniel Fleischer, Peter Izsak, et al. HELMET: How to Evaluate Long-Context Language Models Effectively and Thoroughly. *arXiv preprint arXiv:2410.02694*, 2024.

[24] Kiran Vodrahalli, Santiago Ontanon, Nilesh Tripuraneni, Kelvin Xu, Sanil Jain, Rakesh Shivanna, et al. Michelangelo: Long Context Evaluations Beyond Haystacks via Latent Structure Queries. *arXiv preprint arXiv:2409.12640*, 2024.

[25] Rishabh Agarwal, Avi Singh, Lei M. Zhang, Bernd Bohnet, Luis Rosias, Stephanie Chan, et al. Many-Shot In-Context Learning. *arXiv preprint arXiv:2404.11018*, 2024.

[26] Kaijian Zou, Muhammad Khalifa, and Lu Wang. On Many-Shot In-Context Learning for Long-Context Evaluation. *arXiv preprint arXiv:2411.07130*, 2024.

[27] Weiwei Wang, Jiyong Min, and Weijie Zou. Intelligence Degradation in Long-Context LLMs: Critical Threshold Determination via Natural Length Distribution Analysis. *arXiv preprint arXiv:2601.15300*, 2026.

[28] Cheng-Yu Hsieh, Yung-Sung Chuang, Chun-Liang Li, Zifeng Wang, Long T. Le, Abhishek Kumar, et al. Found in the Middle: Calibrating Positional Attention Bias Improves Long Context Utilization. *arXiv preprint arXiv:2406.16008*, 2024.

[29] Zhenyu Zhang, Runjin Chen, Shiwei Liu, Zhewei Yao, Olatunji Ruwase, Beidi Chen, et al. Found in the Middle: How Language Models Use Long Contexts Better via Plug-and-Play Positional Encoding. *arXiv preprint arXiv:2403.04797*, 2024.

[30] Ziqi Wang, Hanlin Zhang, Xiner Li, Kuan-Hao Huang, Chi Han, Shuiwang Ji, et al. Eliminating Position Bias of Language Models: A Mechanistic Approach. *arXiv preprint arXiv:2407.01100*, 2024.

[31] Alexander Peysakhovich and Adam Lerer. Attention Sorting Combats Recency Bias In Long Context Language Models. *arXiv preprint arXiv:2310.01427*, 2023.

[32] Xinyi Wu, Yifei Wang, Stefanie Jegelka, and Ali Jadbabaie. On the Emergence of Position Bias in Transformers. *arXiv preprint arXiv:2502.01951*, 2025.

[33] Borun D Chowdhury. Lost in the Middle at Birth: An Exact Theory of Transformer Position Bias. *arXiv preprint arXiv:2603.10123*, 2026.

[34] Ali Modarressi, Hanieh Deilamsalehy, Franck Dernoncourt, Trung Bui, Ryan A. Rossi, Seunghyun Yoon, et al. NoLiMa: Long-Context Evaluation Beyond Literal Matching. *arXiv preprint arXiv:2502.05167*, 2025.

[35] Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. Entity-Based Knowledge Conflicts in Question Answering. *arXiv preprint arXiv:2109.05052*, 2021.

[36] Hung-Ting Chen, Michael J. Q. Zhang, and Eunsol Choi. Rich Knowledge Sources Bring Complex Knowledge Conflicts: Recalibrating Models to Reflect Conflicting Evidence. *arXiv preprint arXiv:2210.13701*, 2022.

[37] Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. Adaptive Chameleon or Stubborn Sloth: Revealing the Behavior of Large Language Models in Knowledge Conflicts. *arXiv preprint arXiv:2305.13300*, 2023.

[38] Kevin Wu, Eric Wu, and James Zou. ClashEval: Quantifying the tug-of-war between an LLM's internal prior and external evidence. *arXiv preprint arXiv:2404.10198*, 2024.

[39] Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, et al. Knowledge Conflicts for LLMs: A Survey. *arXiv preprint arXiv:2403.08319*, 2024.

[40] Chupei Wang and Jiaqiu Vince Sun. Unable to Forget: Proactive Interference Reveals Working Memory Limits in LLMs Beyond Context Length. *arXiv preprint arXiv:2506.08184*, 2025.

[41] Owen Bianchi, Mathew J. Koretsky, Maya Willey, Chelsea X. Alvarado, Tanay Nayak, Adi Asija, et al. Hidden in the Haystack: Smaller Needles are More Difficult for LLMs to Find. *arXiv preprint arXiv:2505.18148*, 2025.

[42] Hui Dai, Dan Pechi, Xinyi Yang, Garvit Banga, and Raghav Mantri. DENIAHL: In-Context Features Influence LLM Needle-In-A-Haystack Abilities. *arXiv preprint arXiv:2411.19360*, 2024.

[43] Muhan Gao, Zih-Ching Chen, and Kuan-Hao Huang. The First Drop of Ink: Nonlinear Impact of Distracting Information in Long-Context Reasoning. *arXiv preprint arXiv:2605.10828*, 2026.

[44] Vyzantinos Repantis, Ameya Gawde, Harshvardhan Singh, Rohit Alekar, Cien Zhang, Svetlana Karslioglu, et al. Separating Semantic Competition from Context Length in RAG Reading. *arXiv preprint arXiv:2605.27294*, 2026.

[45] Florin Cuconasu, Giovanni Trappolini, Federico Siciliano, Simone Filice, Cesare Campagnano, Yoelle Maarek, et al. The Power of Noise: Redefining Retrieval for RAG Systems. *arXiv preprint arXiv:2401.14887*, 2024.

[46] Ori Yoran, Tomer Wolfson, Ori Ram, and Jonathan Berant. Making Retrieval-Augmented Language Models Robust to Irrelevant Context. *arXiv preprint arXiv:2310.01558*, 2023.

[47] Tianwei Lin, Zuyi Zhou, Xinda Zhao, Chenke Wang, Xiaohong Li, Yu Chen, et al. Beyond the Needle's Illusion: Decoupled Evaluation of Evidence Access and Use under Semantic Interference at 326M-Token Scale. *arXiv preprint arXiv:2601.20276*, 2026.

[48] Mufei Li, Dongqi Fu, Limei Wang, Si Zhang, Hanqing Zeng, Kaan Sancak, et al. Haystack Engineering: Context Engineering for Heterogeneous and Agentic Long-Context Evaluation. *arXiv preprint arXiv:2510.07414*, 2025.

[49] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed Chi, et al. Large Language Models Can Be Easily Distracted by Irrelevant Context. *arXiv preprint arXiv:2302.00093*, 2023.

[50] Robin Jia and Percy Liang. Adversarial Examples for Evaluating Reading Comprehension Systems. *arXiv preprint arXiv:1707.07328*, 2017.

[51] Minglai Yang, Ethan Huang, Liang Zhang, Mihai Surdeanu, William Wang, and Liangming Pan. How Is LLM Reasoning Distracted by Irrelevant Context? An Analysis Using a Controlled Benchmark. *arXiv preprint arXiv:2505.18761*, 2025.

[52] Yanbo Wang, Zixiang Xu, Yue Huang, Chujie Gao, Siyuan Wu, Jiayi Ye, et al. Adaptive Distraction: Probing LLM Contextual Robustness with Automated Tree Search. *arXiv preprint arXiv:2502.01609*, 2025.

[53] David Chiang and Peter Cholak. Overcoming a Theoretical Limitation of Self-Attention. *arXiv preprint arXiv:2202.12172*, 2022.

[54] Petar Veličković, Christos Perivolaropoulos, Federico Barbero, and Razvan Pascanu. Softmax is not Enough (for Sharp Size Generalisation). *arXiv preprint arXiv:2410.01104*, 2024.

[55] Michael Hahn. Theoretical Limitations of Self-Attention in Neural Sequence Models. *arXiv preprint arXiv:1906.06755*, 2019.

[56] Timur Mudarisov, Mikhail Burtsev, Tatiana Petrova, and Radu State. Limitations of Normalization in Attention Mechanism. *arXiv preprint arXiv:2508.17821*, 2025.

[57] Pavlo Vasylenko, Hugo Pitorro, André F. T. Martins, and Marcos Treviso. Long-Context Generalization with Sparse Attention. *arXiv preprint arXiv:2506.16640*, 2025.

[58] Ken M. Nakanishi. Scalable-Softmax Is Superior for Attention. *arXiv preprint arXiv:2501.19399*, 2025.

[59] Shi Chen, Zhengjiang Lin, Yury Polyanskiy, and Philippe Rigollet. Critical attention scaling in long-context transformers. *arXiv preprint arXiv:2510.05554*, 2025. ICLR 2026.

[60] Kewei Li, Yanwen Kong, Yiping Xu, Jianlin Su, Lan Huang, Ruochi Zhang, et al. Information Entropy Invariance: Enhancing Length Extrapolation in Attention Mechanisms. *arXiv preprint arXiv:2501.08570*, 2025.

[61] Ben Anson, Xi Wang, and Laurence Aitchison. Scale-invariant Attention. *arXiv preprint arXiv:2505.17083*, 2025.

[62] Meysam Ghaffari, Nina Fatehi, Bhaskar Sen, Nasim Sabetpour, and Carlos Morato. Context Poisoning as Extreme-Value Attention Interference in Long-Context Language Models. *arXiv preprint arXiv:2609.22101*, 2026.

[63] Szymon Tworkowski, Konrad Staniszewski, Mikołaj Pacek, Yuhuai Wu, Henryk Michalewski, and Piotr Miłoś. Focused Transformer: Contrastive Training for Context Scaling. *arXiv preprint arXiv:2307.03170*, 2023.

[64] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient Streaming Language Models with Attention Sinks. *arXiv preprint arXiv:2309.17453*, 2023.

[65] Xiangming Gu, Tianyu Pang, Chao Du, Qian Liu, Fengzhuo Zhang, Cunxiao Du, et al. When Attention Sink Emerges in Language Models: An Empirical View. *arXiv preprint arXiv:2410.10781*, 2024.

[66] Mingjie Sun, Xinlei Chen, J. Zico Kolter, and Zhuang Liu. Massive Activations in Large Language Models. *arXiv preprint arXiv:2402.17762*, 2024.

[67] Federico Barbero, Álvaro Arroyo, Xiangming Gu, Christos Perivolaropoulos, Michael Bronstein, Petar Veličković, et al. Why do LLMs attend to the first token? *arXiv preprint arXiv:2504.02732*, 2025.

[68] Yuval Ran-Milo. Attention Sinks Are Provably Necessary in Softmax Transformers: Evidence from Trigger-Conditional Tasks. *arXiv preprint arXiv:2603.11487*, 2026.

[69] Sara Rizwan, Samaanah Abdus Salam, and Mohammed Mudassir Uddin. Do New Attention Mechanisms Actually Fix Attention Sinks at Million-Token Context? *arXiv preprint arXiv:2609.08574*, 2026.

[70] Zunhai Su, Hengyuan Zhang, Wei Wu, Yifan Zhang, Yaxiu Liu, He Xiao, et al. Attention Sink in Transformers: A Survey on Utilization, Interpretation, and Mitigation. *arXiv preprint arXiv:2604.10098*, 2026.

[71] Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, Gabriel Martín Blázquez, Guilherme Penedo, Lewis Tunstall, et al. SmolLM2: When Smol Goes Big: Data-Centric Training of a Small Language Model. *arXiv preprint arXiv:2502.02737*, 2025.

[72] Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, et al. 2 OLMo 2 Furious. *arXiv preprint arXiv:2501.00656*, 2024.

[73] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, et al. In-context Learning and Induction Heads. *arXiv preprint arXiv:2209.11895*, 2022.

[74] Wenhao Wu, Yizhong Wang, Guangxuan Xiao, Hao Peng, and Yao Fu. Retrieval Head Mechanistically Explains Long-Context Factuality. *arXiv preprint arXiv:2404.15574*, 2024.

[75] Wuwei Zhang, Fangcong Yin, Howard Yen, Danqi Chen, and Xi Ye. Query-Focused Retrieval Heads Improve Long-Context Reasoning and Re-ranking. *arXiv preprint arXiv:2506.09944*, 2025.

[76] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced Transformer with Rotary Position Embedding. *arXiv preprint arXiv:2104.09864*, 2021.

[77] Ofir Press, Noah A. Smith, and Mike Lewis. Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation. *arXiv preprint arXiv:2108.12409*, 2021.

[78] Amirhossein Kazemnejad, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Payel Das, and Siva Reddy. The Impact of Positional Encoding on Length Generalization in Transformers. *arXiv preprint arXiv:2305.19466*, 2023.

[79] Jie Wang, Tao Ji, Yuanbin Wu, Hang Yan, Tao Gui, Qi Zhang, et al. Length Generalization of Causal Transformers without Position Encoding. *arXiv preprint arXiv:2404.12224*, 2024.

[80] Yufeng Du, Phillip Harris, Minyang Tian, Eliu A Huerta, Srikanth Ronanki, Subendhu Rongali, et al. RoPE Distinguishes Neither Positions Nor Tokens in Long Contexts, Provably. *arXiv preprint arXiv:2605.15514*, 2026.

[81] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, et al. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *arXiv preprint arXiv:2005.11401*, 2020.

[82] Seongwoong Shim, Myunsoo Kim, Jae Hyeon Cho, and Byung-Jun Lee. Beyond RAG vs. Long-Context: Learning Distraction-Aware Retrieval for Efficient Knowledge Grounding. *arXiv preprint arXiv:2509.21865*, 2025.

[83] Vardhan Dongre, Joseph Hsieh, Viet Dac Lai, Seunghyun Yoon, Trung Bui, and Dilek Hakkani-Tür. When Attention Closes: How LLMs Lose the Thread in Multi-Turn Interaction. *arXiv preprint arXiv:2605.12922*, 2026.

[84] Nadezhda Chirkova, Thibault Formal, Vassilina Nikoulina, and Stéphane Clinchant. Provence: efficient and robust context pruning for retrieval-augmented generation. *arXiv preprint arXiv:2501.16214*, 2025.

[85] Zijun Wu, Bingyuan Liu, Ran Yan, Lei Chen, and Thomas Delteil. Reducing Distraction in Long-Context Language Models by Focused Learning. *arXiv preprint arXiv:2411.05928*, 2024.

[86] Daewon Choi, Jimin Lee, Jihoon Tack, Woomin Song, Saket Dingliwal, Sai Muralidhar Jayanthi, et al. Think Clearly: Improving Reasoning via Redundant Token Pruning. *arXiv preprint arXiv:2507.08806*, 2025.

[87] Ngoc Bui, Hieu Trung Nguyen, Arman Cohan, and Rex Ying. Make Each Token Count: Towards Improving Long-Context Performance with KV Cache Eviction. *arXiv preprint arXiv:2605.09649*, 2026.

[88] Mohsen Arjmandi. Distractor-Aware Truncation: Disentangling Context-Length Effects from Signal Loss in Long-Context LLM Benchmarks. *arXiv preprint arXiv:2608.03297*, 2026.

[89] Shangding Gu. Long Context, Less Focus: A Scaling Gap in LLMs Revealed through Privacy and Personalization. *arXiv preprint arXiv:2602.15028*, 2026.

[90] Changhun Lee, Minsang Seok, Jun-gyu Jin, Younghyun Cho, and Eunhyeok Park. SEAL: Scaling to Emphasize Attention for Long-Context Retrieval. *arXiv preprint arXiv:2501.15225*, 2025.

[91] Xi Ye, Wuwei Zhang, Fangcong Yin, Howard Yen, and Danqi Chen. DySCO: Dynamic Attention-Scaling Decoding for Long-Context Language Models. *arXiv preprint arXiv:2602.22175*, 2026.

[92] Rachit Bansal, Aston Zhang, Rishabh Tiwari, Lovish Madaan, Sai Surya Duvvuri, Devvrit Khatri, et al. Let's (not) just put things in Context: Test-Time Training for Long-Context LLMs. *arXiv preprint arXiv:2512.13898*, 2025.

[93] Zehua Pei, Hui-Ling Zhen, Xianzhi Yu, Sinno Jialin Pan, Mingxuan Yuan, and Bei Yu. FocuSFT: Bilevel Optimization for Dilution-Aware Long-Context Fine-Tuning. *arXiv preprint arXiv:2605.09932*, 2026.

[94] Xuechen Zhang, Xiangyu Chang, Mingchen Li, Amit Roy-Chowdhury, Jiasi Chen, and Samet Oymak. Selective Attention: Enhancing Transformer through Principled Context Control. *arXiv preprint arXiv:2411.12892*, 2024.

[95] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Selective Attention Improves Transformer. *arXiv preprint arXiv:2410.02703*, 2024.

[96] Tianzhu Ye, Li Dong, Yuqing Xia, Yutao Sun, Yi Zhu, Gao Huang, et al. Differential Transformer. *arXiv preprint arXiv:2410.05258*, 2024.

[97] Zihan Qiu, Zekun Wang, Bo Zheng, Zeyu Huang, Kaiyue Wen, Songlin Yang, et al. Gated Attention for Large Language Models: Non-linearity, Sparsity, and Attention-Sink-Free. *arXiv preprint arXiv:2505.06708*, 2025.

[98] Xingyue Huang, Xueying Ding, Mingxuan Ju, Yozen Liu, Neil Shah, and Tong Zhao. Threshold Differential Attention for Sink-Free, Ultra-Sparse, and Non-Dispersive Language Modeling. *arXiv preprint arXiv:2601.12145*, 2026. ACL 2026.

[99] Zhixuan Lin, Evgenii Nikishin, Xu Owen He, and Aaron Courville. Forgetting Transformer: Softmax Attention with a Forget Gate. *arXiv preprint arXiv:2503.02130*, 2025.

[100] Shawn Tan, Songlin Yang, Aaron Courville, Rameswar Panda, and Yikang Shen. Scaling Stick-Breaking Attention: An Efficient Implementation and In-depth Study. *arXiv preprint arXiv:2410.17980*, 2024.

[101] Zikai Xiao, Ziyang Wang, Wen Ma, Yan Zhang, Wei Shen, Yan Wang, et al. Mitigating Posterior Salience Attenuation in Long-Context LLMs with Positional Contrastive Decoding. *arXiv preprint arXiv:2506.08371*, 2025.

[102] Bowen Yang, Bharat Venkitesh, Dwarak Talupuru, Hangyu Lin, David Cairuz, Phil Blunsom, et al. Rope to Nope and Back Again: A New Hybrid Attention Strategy. *arXiv preprint arXiv:2501.18795*, 2025.

[103] Yoav Gelberg, Koshi Eguchi, Takuya Akiba, and Edoardo Cetin. Extending the Context of Pretrained LLMs by Dropping Their Positional Embeddings. *arXiv preprint arXiv:2512.12167*, 2025.