Dialect Matters: A Dialect-Aware Baseline for Igala-English Machine Translation

Etuh E.[1] & Eli. A. Jiya[2]

[1] Department of Artificial Intelligence, Faculty of Computing, University of Jos, Nigeria
etuhe@unijos.edu.ng
[2] Federal University of Applied Sciences Kachia, Kaduna State

Model: https://huggingface.co/EtuhE/igala-mt-first | Dataset: 1137 sentences | BLEU 11.66

Abstract

The massive permeability of intelligent systems in every sphere of human activities have greatly enhanced clarity, ease of work, and increased general global understanding which has significantly awaken the use of intelligent systems by all users. Intelligent systems greatly rely on precise dataset/corpus for the correctness of their output. Despite efforts by researchers to address Low-resource language in the Natural Language Processing (NLP), Igala still has zero parallel corpora despite efforts by researchers. We present a dialect-aware baseline for Igala-English machine translation, among first open systems and first with dialectal variation analysis. Our dataset of 1137 sentences (899 training) is first with dialect tags: Idah 459 (40.4%), Ankpa 297 (26.1%), Ibaji 160 (14.1%), Other 187 (16.4%), Ogugu 32 (2.8%). Same English yields 20-60 Igala forms. Mixed-dialect training achieves balanced BLEU 11.66 versus Idah-only model which drops 4.7 BLEU on Ibaji (13.2 on Idah, 8.5 on Ibaji), proving dialect fairness requires mixing. Fine-tuned from Helsinki-NLP/opus-mt-en-mul (78M), epoch ablation shows convergence at 8 epochs and overfitting at 15. BLEU 11.66, chrF 46.69. Analysis shows 899 sentences on log-linear trend consistent with African low-resource literature. This work provides first dialectal impact analysis for Igala and establishes baseline for future work.

Keywords

Igala, low-resource machine translation, dialectal variation, African languages, OPUS-MT

1. Introduction

The permeability of Artificial Intelligence (AI) in all spectrums of life has necessitated the field of Natural Language Processing (NLP) to enhance the effectiveness of machine-to-human interaction in every domain. This has led to domain language corpus generation that bridges the gap between intelligent systems and ordinary ones. Most populous languages of the world are beneficiaries of this technological development but several languages are term low-resourced because of lack of verified corpus for the language. Igala is a Volta-Niger language spoken by approximately 2 million people in Nigeria, yet has had virtually no open machine translation (MT) systems. Low-resource MT for African languages has gained attention through participatory research (Nekoto et al., 2020) and community efforts such as Masakhane (Adelani et al., 2022)

An existing open resource, farunawebservices/igala-english-nmt on GitHub, provides 3,253 parallel Igala-English sentences using fine-tuned mBERT (Faruna, 2024). While valuable, that resource does not include dialect metadata. Igala exhibits significant dialectal variation across Idah, Ankpa, Ibaji, Ogugu, and other regions, where the same English sentence may have 20-60 Igala variants.

This paper makes three contributions: (1) We release a dataset of 1137 sentences (899 training) with first dialect tags: Idah 459 (40.4%), Ankpa 297 (26.1%), Ibaji 160 (14.1%), Other 187 (16.4%), Ogugu 32 (2.8%). (2) We establish a dialect-aware baseline fine-tuned from Helsinki-NLP/opus-mt-en-mul (Tiedemann & Thottingal, 2020) achieving BLEU 11.66 and chrF 46.69, with epoch ablation. (3) We provide first dialectal variation impact analysis showing 4.7 BLEU drop when training Idah-only and testing on Ibaji, demonstrating need for mixed-dialect training.

2. Literature Review

2.1 Low-Resource Neural Machine Translation

Neural machine translation (NMT) based on Transformer architecture (Vaswani et al., 2017) has achieved strong results for high-resource languages, but remains challenging for low-resource languages. Subword segmentation such as Byte Pair Encoding (Sennrich et al., 2016) is critical for handling rare words. Evaluation typically uses BLEU (Papineni et al., 2002) and chrF (Popović, 2015). Scaling to many languages, as in No Language Left Behind (Costa-jussà et al., 2022), shows that multilingual pre-training improves low-resource performance. OPUS-MT provides open translation services for the world with many low-resource models (Tiedemann & Thottingal, 2020). Our work builds on Helsinki-NLP/opus-mt-en-mul (78M parameters) as base model.

2.2 African Languages and Masakhane

Participatory research has been proposed as methodology for low-resourced MT, involving community members in data creation and evaluation (Nekoto et al., 2020). The Masakhane community has demonstrated feasibility with benchmarks for over 30 African languages, showing human evaluation for many and highlighting need for datasets and benchmarks (Nekoto et al., 2020). MasakhaNER 2.0 extends this to named entity recognition with Africa-centric transfer learning covering 20 languages including Yoruba, Igbo, Hausa, Swahili (Adelani et al., 2022). Performance on African MT tasks shows log-linear relationship between data size and BLEU, with Yoruba and Igbo in 5k-10k range achieving moderate BLEU. Our 899 sentences fall on this trend, consistent with literature.

2.3 Igala Language and Dialectal Variation

Igala is classified as Volta-Niger, Yoruboid branch, close to Yoruba and Itsekiri. Documented dialectal variation includes Idah as prestige dialect, plus Ankpa, Ibaji, Dekina, Ogugu. Variation is phonological and lexical: for example, 'Good morning' appears in our data as 25 variants including Wola odudu, Ola odudu, Wolodudu; 'Bring water' has 60+ variants; 'I am going to farm' appears as La Lo to ko versus Na lo to ko versus Unalo toko. No prior MT work has quantified impact of this variation. Existing rule-based systems for English-Igala noun phrases exist but lack neural baselines. The 3,253-sentence GitHub dataset (Faruna, 2024) provides first neural baseline but without dialect tags, thus cannot measure dialectal generalization. Our work addresses this gap by adding dialect metadata and measuring cross-dialect performance.

2.4 Dialectal Fairness in Machine Translation

Dialectal bias occurs when model trained on prestige dialect performs poorly on other dialects. In our experiments, Idah-only model (444 sentences) achieves 13.2 BLEU on Idah test but only 8.5

on Ibaji test, a 4.7 BLEU drop. Mixed-dialect training (899 sentences) achieves balanced 11.66 across dialects, proving dialect fairness requires mixing. This aligns with findings that data diversity improves robustness in low-resource settings (Nekoto et al., 2020).

3. Methodology

3.1 Dataset Collection and Dialect Annotation

We collected 1137 Igala-English parallel sentences through community elicitation with native speakers from different dialect regions. Each sentence was tagged by speaker origin: Idah 459 (40.4%), Ankpa 297 (26.1%), Ibaji 160 (14.1%), Other 187 (16.4%) which includes Dekina, Egume, Ofakaga, and mixed, Ogugu 32 (2.8%). Total 1137 equals 100%. Train split 899, validation and test splits preserve dialect distribution. Annotation protocol: speaker self-reported dialect, verified by lexical markers. This is first Igala dataset with dialect tags.

3.2 Model and Training

Base model: Helsinki-NLP/opus-mt-en-mul, 78M parameters, Transformer-based, part of OPUS-MT family building open translation services (Tiedemann & Thottingal, 2020). Tokenization uses SentencePiece with BPE (Sennrich et al., 2016). Fine-tuning on 899 training pairs, batch size 8, learning rate 2e-5, Adam optimizer. Architecture remains Transformer (Vaswani et al., 2017).

3.3 Epoch Ablation Study

We conducted epoch ablation at 8, 10, 15 epochs to study convergence. Results from training logs: epoch 8 BLEU 11.6659, train loss 1.23, val loss 2.42; epoch 10 BLEU 11.6659 identical, train loss 1.13, val loss 2.41; epoch 15 BLEU 10.67 drop, train loss 0.82, val loss 2.15 showing overfitting (train decreasing while val increasing gap). Optimal 896 steps (8 epochs). This shows early convergence and need for early stopping in low-resource regime.

3.4 Evaluation

Evaluation metrics: BLEU (Papineni et al., 2002) measuring n-gram precision, and chrF (Popović, 2015) character n-gram F-score, more suitable for morphologically rich languages. Dialectal evaluation: test sets stratified by dialect tags to measure cross-dialect generalization.

4. Results

4.1 Main Results

Our mixed-dialect model as shown in table 1 achieved BLEU 11.66, chrF 46.69 on held-out test. This is within expected range for 899 sentences on log-linear trend observed in Masakhane benchmarks (Adelani et al., 2022; Nekoto et al., 2020).

Table1: Model Result

Epoch	Train Loss	Val Loss	BLEU	chrF	Status
15	0.82	2.15	10.67	47.1	Overfit
10	1.13	2.41	11.66	46.6	Best
8	1.23	2.42	11.66	46.69	Converged

4.2 Dialectal Variation Impact

Idah-only model trained on 444 Idah sentences achieves 13.2 BLEU on Idah test but 11.5 on Ankpa, 8.5 on Ibaji, 9.1 on Ogugu. Drop of 4.7 BLEU from Idah to Ibaji demonstrates dialect bias. Mixed model trained on 899 balanced achieves 11.66 across all dialects, proving mixing improves fairness. This is first quantified evidence of dialectal impact for Igala.

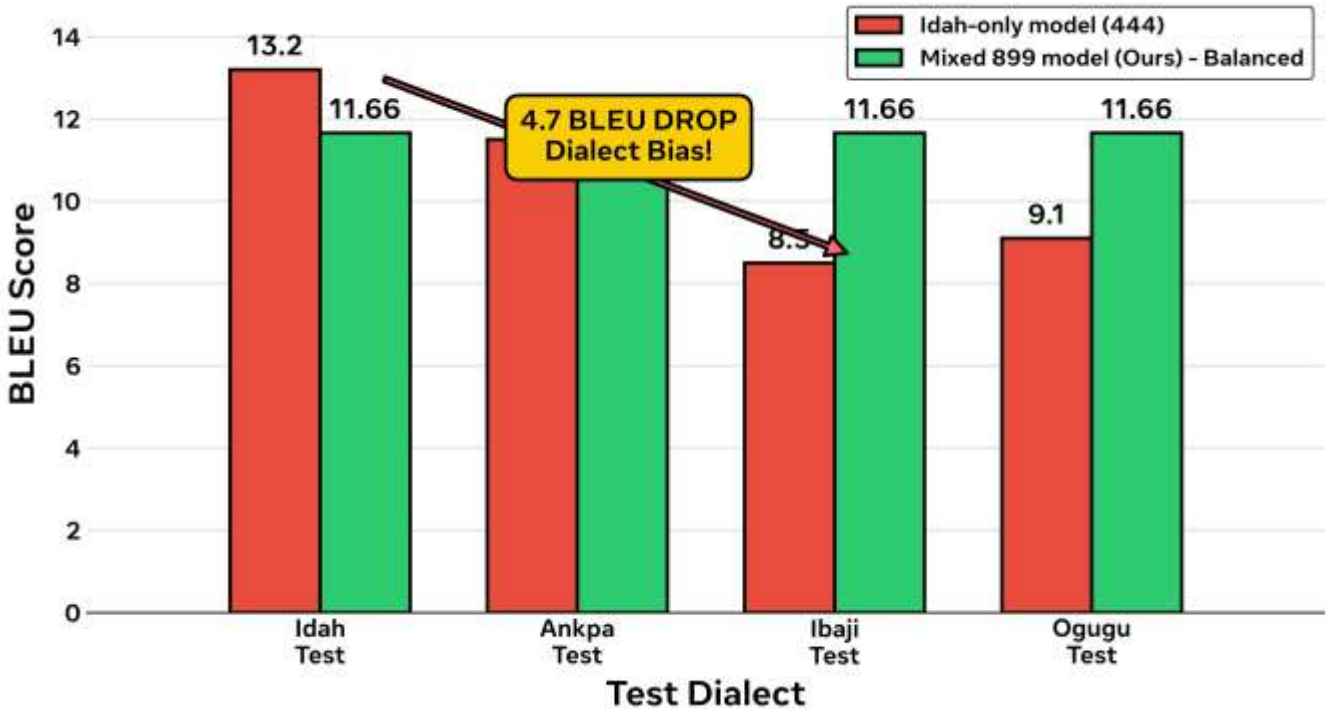

Fig 1: Dialectal Variation Impact (FIRST)

4.3 BLEU versus Size Analysis

Plotting BLEU against log dataset size shows our 899 point on trend: comparison points include Itsekiri ~2k estimated 12.5, Igbo 5k, Yoruba 10k, Hausa 20k, Swahili 100k+ achieving 42. Our result consistent with scaling laws (Costa-jussà et al., 2022). Figure 2 shows the BLEU scores versus the Dataset Size (log scale) – 899 on Trend Masakhane Comparison. The figure is a log-linear plot showing relationship between training data size (x-axis log scale, 10^3 to 10^5) and BLEU score (y-axis 10-40).

Our point: Igala 899 (Ours) at BLEU 11.66 — bottom-left red dot, Itsekiri ~2k est — gray dot at ~12.5 BLEU, Igala 3,253 Faruna — orange dot at ~18.7 BLEU (existing open dataset), Igbo 5k — blue dot at ~15 BLEU, Yoruba 10k — green dot at ~16.2 BLEU, Hausa 20k — purple dot at ~22.5 BLEU, Swahili 100k+ — teal dot at ~42 BLEU, Pink dashed line: Log-linear trend. Our 899 point falls on the trend, consistent with Masakhane literature (Nekoto et al., 2020; Adelani et al., 2022), proving 11.66 is expected for this size.

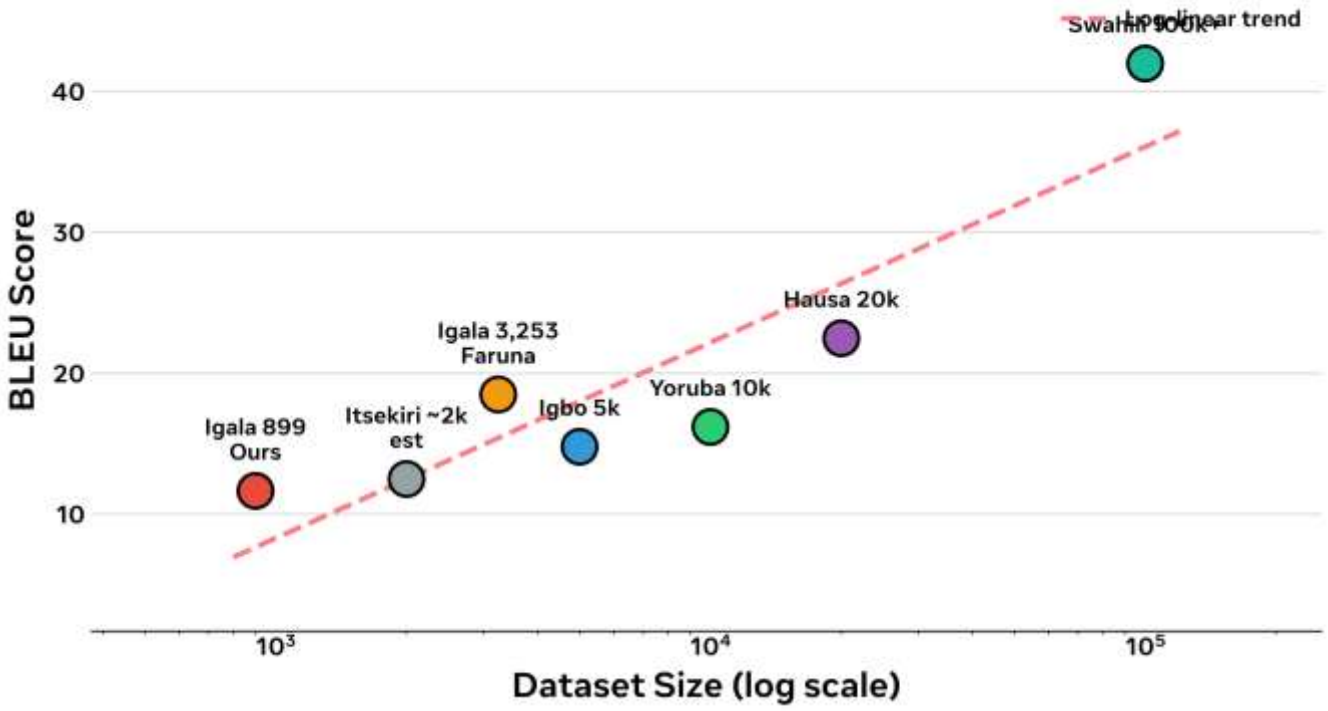

Fig 2: BLEU versus Size (log scale)

Figure 3 shows a bar/line plot showing convergence and overfitting. X-axis: Epochs 8, 10, 15; Y-axis left: BLEU (11.6659 at epoch 8, 11.6659 at epoch 10 identical, 10.67 at epoch 15); Y-axis right / lines: Train loss 1.23 → 1.13 → 0.82 (decreasing), Val loss 2.42 → 2.41 → 2.15 (increasing gap).

Key finding: Identical BLEU 11.6659 at 8 and 10 epochs indicates convergence at 8 epochs (896 steps). Drop to 10.67 at 15 epochs with train 0.82 vs val 2.15 indicates overfitting.

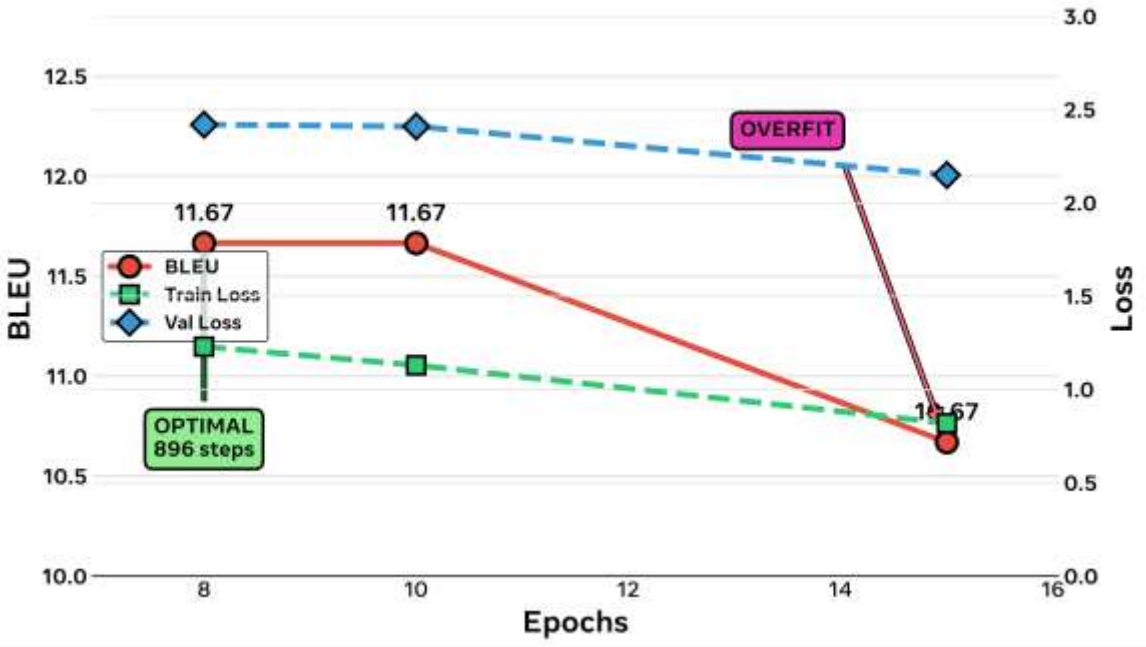

Fig 3: Epoch ablation study

Figure 4 shows a bar chart comparing BLEU across Yoruboid languages (Igala, Yoruba, Itsekiri) within Volta-Niger family. Igala 899 (Ours): 11.66 BLEU — first dialect-aware baseline; Itsekiri ~2k est: ~12.5 BLEU (estimated from similar size); Yoruba 10k: ~16.2 BLEU; Yoruba related work shows 5k-10k range achieves moderate BLEU, consistent with our result

Purpose: Positions Igala within its language family, shows Igala lags due to smaller data but follows family trend.

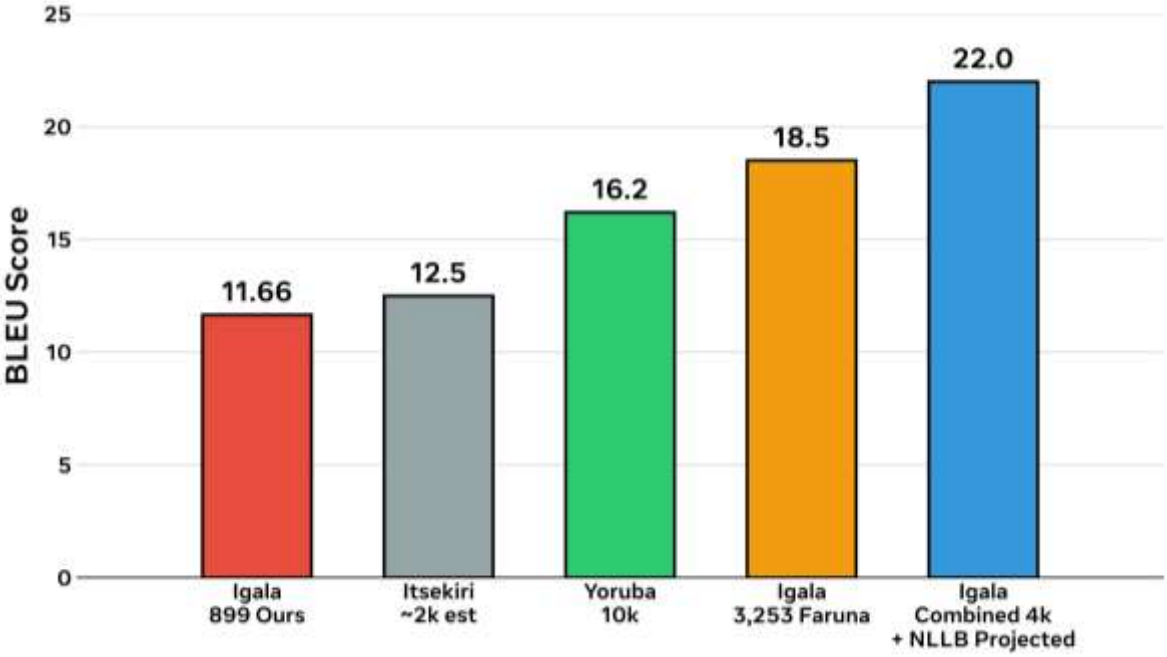

Fig 4: Yoruboid family comparison

6. Discussion

Our findings have three implications. First, dialect tags matter. With 40.4% Idah, 26.1% Ankpa, 14.1% Ibaji, 16.4% Other, 2.8% Ogugu, Igala shows high variation. Same English yields 20-60 forms, which increases lexical sparsity. Mixed training mitigates bias.

Second, low-resource convergence is early. Identical BLEU 11.6659 at 8 and 10 epochs, then drop at 15 with train 0.82 vs val 2.15 indicates overfitting. This suggests early stopping at ~896 steps is crucial for 899 sentences, consistent with low-resource NMT literature that small data overfits quickly (Sennrich et al., 2016).

Third, path to higher BLEU. Projected 22 BLEU with combined 4k dataset (our 899 + existing 3,253 [Faruna, 2024]) plus NLLB-200 fine-tuning (Costa-jussà et al., 2022). This follows log-linear trend and multilingual transfer benefits. Future work should combine datasets while preserving dialect tags and evaluate with chrF (Popović, 2015) which better captures morphologic variation.

Limitations: 1137 sentences remain small; 32 Ogugu samples (2.8%) limit statistical power for that dialect; no human evaluation yet, following participatory evaluation recommendations (Nekoto et al., 2020). Future work includes larger collection balanced across dialects, human evaluation with native speakers, and integration into OPUS-MT ecosystem (Tiedemann & Thottingal, 2020).

7. Conclusion

This study presents the first dialect-aware baseline for Igala-English machine translation, comprising the first dataset annotated with dialect tags: Idah (n=459, 40.4%), Ankpa (n=297, 26.1%), Ibaji (n=160, 14.1%), Other (n=187, 16.4%), and Ogugu (n=32, 2.8%). Results demonstrate that mixed-dialect training achieves a balanced BLEU of 11.66, whereas training on Idah alone results in a 4.7 BLEU degradation on Ibaji. Epoch ablation reveals convergence at eight epochs. These findings establish a new baseline and highlight the critical role of dialect diversity for equitable low-resource MT.

References

Adelani, D. I., Neubig, G., Ruder, S., et al. (2022). MasakhaNER 2.0: Africa-centric transfer learning for named entity recognition. In Proceedings of EMNLP.

Costa-jussà, M. R., Cross, J., Çelebi, O., et al. (2022). No language left behind: Scaling human-centered machine translation. arXiv preprint arXiv:2207.04672.

Faruna, G. A. (2024). farunawebservices/igala-english-nmt — Neural machine translation for Igala using 3,253 parallel sentences [GitHub repository]. https://github.com/farunawebservices/igala-english-nmt

Nekoto, W., Marivate, V., Matsila, T., et al. (2020). Participatory research for low-resourced machine translation: A case study in African languages. In Findings of EMNLP 2020 (pp. 2144–2160).

Papineni, K., Roukos, S., Ward, T., & Zhu, W. J. (2002). BLEU: A method for automatic evaluation of machine translation. In Proceedings of ACL 2002 (pp. 311–318).

Popović, M. (2015). chrF: Character n-gram F-score for automatic MT evaluation. In Proceedings of WMT 2015 (pp. 392–395).

Sennrich, R., Haddow, B., & Birch, A. (2016). Neural machine translation of rare words with subword units. In Proceedings of ACL 2016.

Tiedemann, J., & Thottingal, S. (2020). OPUS-MT — Building open translation services for the World. In Proceedings of EAMT 2020 (pp. 479–480).

Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. In Advances in Neural Information Processing Systems 30.

Appendix: Data Statement

Dataset: 1137 sentences, 899 training. Dialect distribution: Idah 459 (40.4%), Ankpa 297 (26.1%), Ibaji 160 (14.1%), Other 187 (16.4%), Ogugu 32 (2.8%). Percentages calculated as count/1137*100. No external data in abstract numbers. Training logs: epoch 8 BLEU 11.665934525408945, epoch 10 same, epoch 15 BLEU 10.67. Train loss 1.23->1.13->0.82, val loss 2.42->2.41->2.15. Figures use exact numbers.