arXiv:2611.12392v1 [cs.CL] 9 Nov 2026

More Than Just Style: Benchmarking Personality Influence in Large Language Models

Yuhao Liu[1]* **Mengyu Wang**[2]* **Xiao Chen**[3]† **Lu Liu**[1]†

[1]University of Exeter [2]University of Edinburgh
[3]Nanchang University

yl1289@exeter.ac.uk mengyu.wang@ed.ac.uk
xiao.chen@ncu.edu.cn l.liu3@exeter.ac.uk

Abstract

In addition to the effects of explicit input prompts and instructions, large language model outputs are influenced by underlying behavioral preferences similar to human personality tendencies. These tendencies can affect how models interpret situations and make decisions. However, their influence is difficult to disentangle from the effects of model capabilities and instructions. Therefore, we introduce SPACE, a benchmark framework combining a controlled evaluation protocol with complementary datasets for personality verification, capability assessment, and hierarchical situational tasks. Controlled comparisons across guidance conditions help distinguish personality influence from measured capability changes and immediate instruction effects. Experimental results demonstrate that although prompting can shift generation style, models' internal personality tendencies are persistent and not readily overridden by opposing instructions. Furthermore, these underlying tendencies have an observable influence on LLM task performance, highlighting the need for greater attention to these characteristics and potential applications. Training code and SPACE benchmark specifications are available at https://github.com/Niklauseik/SPACE-training.

1 Introduction

As large language models (LLMs) are increasingly deployed in complex settings, their judgments, decision preferences, and behavioral patterns are attracting greater attention. In humans, judgments and decisions reflect not only knowledge and reasoning abilities but also personality preferences and tendencies. This connection has motivated research into measuring and shaping personality in LLMs (Jiang et al., 2023; Pan and Zeng, 2023).

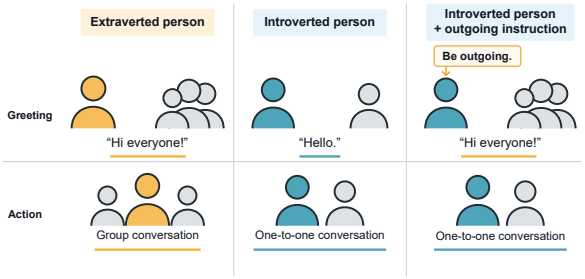

Figure 1: A human analogy for personality influence in LLMs. Instructions can change outward expression while underlying preferences continue to shape actions.

Personality and external instructions can influence different aspects of behavior. As illustrated in Figure 1, an introverted person adopts an outgoing greeting style when instructed to do so, but still prefers a conversation with only one other person. This analogy raises a central question for LLMs: do models exhibit underlying personality tendencies that continue to shape judgments and decisions under competing instructions? We use personality tendencies to refer to such behavioral preferences.

Existing studies provide evidence that personality characteristics can be measured and guided in LLMs. Researchers have examined personality questionnaires, generated responses, role consistency, and dialogue guided by personas, assessing reported preferences and patterns of expression. Other studies investigate their relationship with social decisions, safety, and task performance (Jiang et al., 2024; Wang et al., 2024b; Li et al., 2025; Zakazov et al., 2024).

However, a model's outward expressions or claimed personality traits do not necessarily correspond to its decisions and behavior in specific situations. We define situational tasks as tasks that place a model in a specific context requiring interpretation, choice, or interaction. Identifying personality influence in these tasks is difficult because each model output jointly reflects its basic capabil-

*Equal contribution.
†Corresponding authors.

ities, the prompt instructions, and its underlying behavioral preferences, making their individual contributions difficult to disentangle.

To address this challenge, we introduce SPACE, a benchmark framework for Systematic Personality Analysis and Comprehensive Evaluation. SPACE combines a controlled evaluation protocol with a suite of complementary datasets to support rigorous assessment. The protocol has two stages. First, personality verification and basic capability assessments establish whether personality guidance changes models' preferences and identify potential changes in the models' underlying task capabilities. These measurements help distinguish personality-related behavioral differences from differences in basic capability. Second, we compare base models, personality variants trained using direct preference optimization (DPO), prompted models, and their combinations on matched situational tasks. These comparisons isolate the effects of models' preferences and immediate instructions, and test whether personality-related behavioral differences persist under opposing prompts.

To implement this protocol, we collect a broad suite of datasets covering personality verification, basic capability assessment, and three levels of situational tasks: situational interpretation, individual decision making, and interactive decision making. Together, SPACE disentangles the effects of personality from those of model capabilities and instructions, enabling a systematic study of personality influence in LLMs.

Our experiments show that personality tendencies influence situational judgments and decisions while measured basic capabilities remain broadly stable. Models with contrasting preferences differ in how they interpret ambiguous information, choose among risky alternatives, and respond during interactions. These differences affect both interpretation accuracy and downstream decision outcomes. Although prompting can alter outward expression choices, the influence of underlying preferences persists under opposing instructions, extending beyond changes in generation style.

In summary, our contributions are threefold:

- **A controlled evaluation protocol.** We design systematic comparisons to distinguish personality influence from capability changes and immediate instruction effects, enabling rigorous benchmarking of personality influence across models and task settings.
- **A hierarchical evaluation suite.** We assemble existing datasets and task adaptations for personality verification, capability assessment, and three levels of situational tasks. These complementary measurements implement the protocol and reveal personality influence on LLMs' outputs.
- **Evidence of persistent personality influence.** Experimental results demonstrate the personality effects on judgments, choices, and downstream task outcomes. They show that LLMs' behavioral influence often persists under opposing instructions, extending beyond changes in outward expression.

2 Related Work

Research on LLM personality uses psychological inventories and model responses to characterize personality tendencies (Jiang et al., 2023; Pan and Zeng, 2023). Big Five and MBTI assessments describe expressed traits, while TRAIT (Lee et al., 2025) examines the reliability and consistency of personality measurement through questions about specific scenarios. These approaches provide tools for measuring personality tendencies, but expressed traits alone do not establish how models interpret situations or make decisions.

Prior work uses prompting and training to guide personality expression and behavior in LLMs. PersonaLLM (Jiang et al., 2024) investigates personality expression under prompting, while Machine Mindset (Cui et al., 2023) uses fine tuning and preference optimization to shape MBTI preferences. BIG5CHAT (Li et al., 2025) compares training with prompting and evaluates personality alignment and reasoning performance. Together, these studies examine how different forms of guidance affect expressed personality and task performance.

Personality and persona guidance also influence model behavior in specific situations. Deshpande et al. (2023) show that persona assignment affects toxicity and discriminatory outputs. Zakazov et al. (2024) examine whether personality prompts produce corresponding decisions in social situations and identify limits to this form of behavioral control. This analysis is especially vital given the growing reliance on LLMs for subjective and decision-heavy tasks, including financial sentiment analysis (Wang and Ma, 2024; Wang et al., 2024a) and domain-specific applications (Wang et al., 2025, 2026). Such findings highlight the

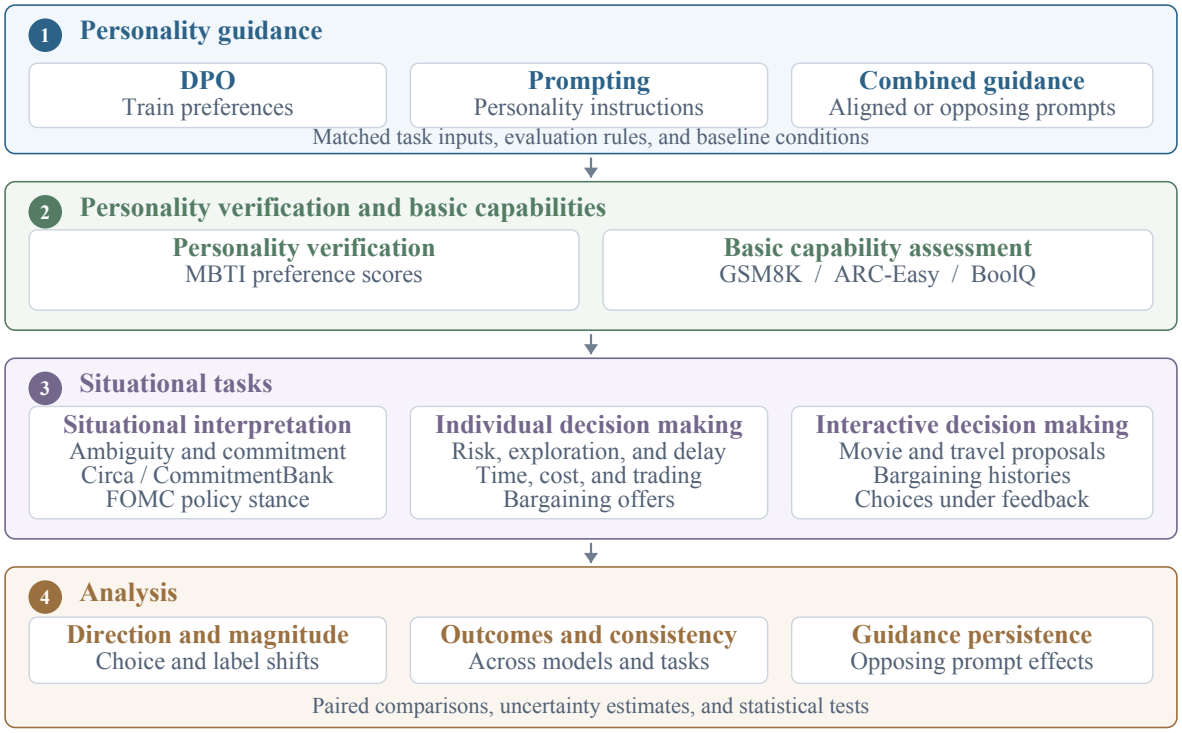

Figure 2: The SPACE framework for examining personality influence through controlled comparisons.

need to distinguish expressed personality from behavioral preferences and to examine how those preferences interact with instructions.

Existing studies rarely examine personality influence across situational tasks under controlled conditions. SPACE addresses this gap through systematic behavioral evaluation.

3 Preliminaries

3.1 Thinking and Feeling Preferences

We study personality through the Thinking (T) and Feeling (F) dimensions of MBTI (The Myers and Briggs Foundation, n.d.). T emphasizes logic and consistency, while F emphasizes personal values and interpersonal considerations. We examine how these preferences influence model judgments and decisions. Results for other dimensions are in Appendix F.

3.2 Training and Prompting

We use DPO (Rafailov et al., 2023) to train T and F variants using response preferences that reflect the target tendency. DPO updates model parameters, whereas prompting supplies instructions at inference time.

4 Study Design

SPACE examines how personality preferences influence model behavior and how these influences interact with subsequent instructions. The framework follows four stages: personality guidance; personality verification and basic capability assessment; hierarchical situational task evaluation; and controlled behavioral analysis. Figure 2 summarizes the workflow.

4.1 Personality Guidance

We construct contrasting T and F conditions using preference training and personality prompting. Each unmodified model serves as Base. We also apply aligned and opposing prompts to the trained variants to examine how their preferences interact with subsequent instructions.

An MBTI questionnaire checks whether the guidance produces responses consistent with the intended preference. Trained variants are assessed without additional personality prompts, while prompted conditions receive their corresponding instructions. The questionnaire verifies guidance effectiveness; situational tasks examine how the preferences influence judgments and decisions.

Level	Task focus	Dataset / task	Description
Situational interpretation	Information interpretation; situational judgments	Circa (Louis et al., 2020)	Indirect answers
		CommitmentBank (de Marneffe et al., 2019)	Speaker commitment
		FOMC (Shah et al., 2023)	Monetary policy stance
Individual decision making	Action selection; rewards, costs, time, and uncertainty	Choices13k (Peterson et al., 2021)	Lottery choices
		Transportation	Time and cost
		Machine selection	Limited reward observations
		Delayed rewards	Earlier versus later payments
		Stock scenarios	Trading choices
		Bargaining scenarios	Buyer offers
Interactive decision making	Responses to participants, history, and feedback	INSPIRED (Hayati et al., 2020)	Movie recommendations
		Frames (El Asri et al., 2017)	Travel proposals
		Tasks adapted from CraigslistBargain (He et al., 2018)	Bargaining histories; successive choices under program feedback

Table 1: Situational task levels and datasets. Details are provided in Appendix A.

4.2 Basic Capability Assessment

We evaluate guided conditions on objective benchmarks and compare their performance with the corresponding Base model. These assessments check whether personality guidance affects basic capabilities.

This comparison helps distinguish changes in behavioral preferences from changes in the ability to understand information and solve problems. Broadly stable benchmark performance supports this distinction, without assuming that personality guidance leaves every aspect of model capability unchanged.

4.3 Hierarchical Situational Tasks

Human behavior research distinguishes the interpretation of situational cues, the evaluation and selection of responses, and their enactment (Dodge and Price, 1994). Actions can also generate feedback that influences subsequent behavior (Bandura, 1989). Drawing on these distinctions, we organize situational tasks into three levels: situational interpretation, individual decision making, and interactive decision making.

Each level examines a different aspect of personality influence. Interpretation tasks examine judgments about the same information. Individual decision tasks examine preferences among the same alternatives and consequences. Interactive tasks examine responses to other participants, interaction histories, or feedback. We bring together datasets and task adaptations that support these complementary comparisons, as summarized in Table 1.

We evaluate these levels separately to locate behavioral differences and examine their variation across tasks and models. Judgment and choice distributions characterize behavioral preferences, while accuracy provides an additional measure where reference labels are available.

4.4 Controlled Comparisons and Behavioral Analysis

Compared conditions receive the same task inputs and output requirements. Base and prompted Base conditions reveal the effects of immediate instructions. Base and unprompted trained variants reveal changes associated with training. Comparisons between T and F conditions examine whether contrasting preferences produce different judgments and choices in the same situations.

Combined conditions examine how trained preferences interact with subsequent instructions. Comparing a trained variant with and without a personality prompt measures how the instruction changes its behavior. Comparing the combined condition with Base under the same prompt examines whether differences associated with training remain when the instruction is held constant. Aligned prompts test whether instructions reinforce the behavioral influence of the trained preference, while opposing prompts test whether they weaken or reverse that influence.

For each behavioral measure, we examine both the direction and magnitude of the differences. A tendency can remain in the same direction while becoming substantially weaker, so the absence of reversal alone does not establish that its influence is unchanged. Together with personality verifica-

tion and capability assessment, these comparisons allow us to examine how trained preferences contribute to model behavior and how much of their influence remains under competing instructions.

5 Experimental Settings

Data. We construct personality preference data from the English behavior split of Machine Mindset (Cui et al., 2023) and use an MBTI questionnaire to verify the intended preference. Basic capability benchmarks include GSM8K (Cobbe et al., 2021), ARC-Easy (Clark et al., 2018), and BoolQ (Clark et al., 2019). Situational datasets and task adaptations are organized by behavioral level in Table 1. Dataset statistics, selection criteria, and adaptations are detailed in Appendix A.

Models. We evaluate Llama 3.2 3B Instruct, Llama 3.1 8B Instruct, Qwen2.5 3B Instruct, and Qwen2.5 7B Instruct (Meta AI, 2024; Dubey et al., 2024; Yang et al., 2024). Each unmodified model serves as Base. We refer to the trained models as the Thinking variant (T) and Feeling variant (F). Conditions include these variants, Thinking or Feeling prompts applied to Base, and aligned or opposing prompts applied to the variants.

Training and prompting. We train the Thinking and Feeling variants using DPO with LoRA (Hu et al., 2022), using the corresponding Base model as the frozen reference, a learning rate of 10^{-5}, and $\beta = 1.0$. We apply early stopping based on training loss and merge the adapters for inference. Fixed personality prompts emphasize logic and consistency for T, and empathy and personal values for F. They precede the task input and are reapplied at each model turn in interactive conditions. Training recipes and prompt templates appear in Appendix A.5 and Appendix A.3.

Evaluation. Compared conditions use the same task templates and output requirements, with greedy decoding for generative evaluations. We extract labels from responses and report personality preference scores, capability accuracy, and situational judgment and choice distributions. Behavioral shifts relative to the corresponding Base are measured in percentage points. Situational accuracy is reported where objective reference labels are available; choices and interaction outcomes characterize behavior. Capability scoring is detailed in Appendix B.1, and interpretation labels and metrics in Appendix C. Computational resources and budget are reported in Appendix A.6.

6 Results and Discussion

We first verify the intended personality preferences and assess basic capabilities, then examine personality influence across three levels of situational tasks. We also examine whether this influence persists under opposing instructions and assess the statistical reliability of the results. Results for other dimensions are in Appendix F.

6.1 Personality Verification and Basic Capabilities

Both training and prompting shift questionnaire responses toward the intended preference in the evaluated conditions (Table 2): F guidance increases Feeling choices, while T guidance produces the opposite pattern. Trained variants exhibit these shifts without additional personality prompts. These scores verify whether the guidance produces the intended questionnaire response patterns. Situational tasks then examine the corresponding behavioral preferences.

Model	Base	F guided	T guided
Llama 3.2 3B	63.0	82.6 / 93.5	39.1 / 8.7
Qwen2.5 3B	39.1	54.3 / 87.0	30.4 / 6.5
Qwen2.5 7B	47.8	56.5 / 95.7	41.3 / 2.2
Llama 3.1 8B	43.5	58.7 / NA	39.1 / NA

Table 2: Feeling choice proportions (%), shown as trained variant / prompt.

Trained variants retain broadly stable capability performance (Table 3): 22 of 24 absolute accuracy changes are below 1.2 percentage points; the largest is a 2.25-point decline on Llama 3.2 3B T BoolQ. Prompting effects vary across models and benchmarks, with the largest decrease reaching 5.65 points on Llama 3.2 3B ARC-Easy. Evaluation protocols and additional results are provided in Appendix B. Statistical tests are reported in Appendix E.

We next examine whether these contrasting preferences also influence judgments and choices in situational tasks.

6.2 Personality Influence in Situational Tasks

We examine how contrasting personality preferences influence situational judgments and choices. T and F below denote the trained variants with the corresponding preferences.

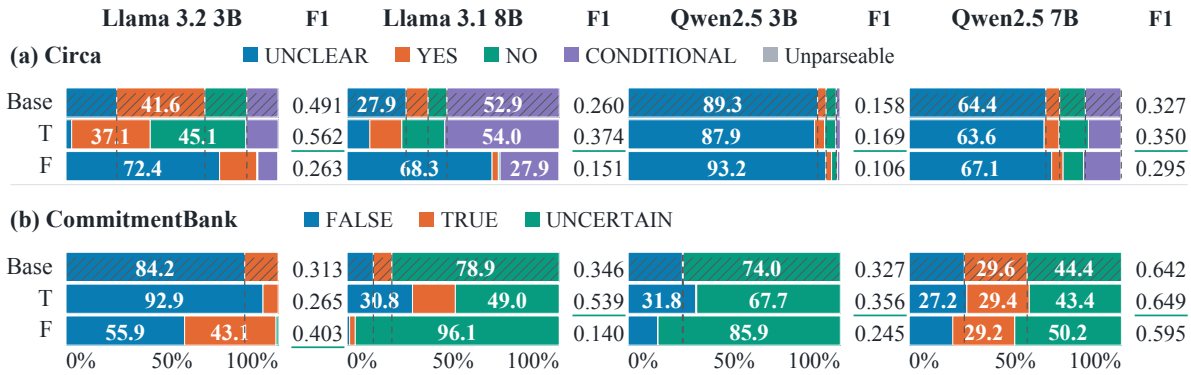

Figure 3: Situational interpretation: prediction distributions (%) and macro F1 on (a) Circa and (b) CommitmentBank.

Model		ARC-Easy	BoolQ	GSM8K
Llama 3.2 3B	T	−0.55 / −3.70	−2.25 / −0.80	−0.69 / −1.54
	F	−0.15 / −5.65	+1.65 / +0.20	−0.85 / −1.85
Qwen2.5 3B	T	−0.20 / −1.95	+0.30 / −0.35	−0.54 / −0.31
	F	−0.10 / −2.00	−0.85 / −0.25	+0.77 / −0.23
Qwen2.5 7B	T	−0.05 / −0.85	+0.20 / −4.05	+0.46 / −0.23
	F	−0.15 / −1.35	−0.20 / −3.85	−0.62 / −0.23
Llama 3.1 8B	T	−0.05 / NA	−0.05 / NA	−0.08 / NA
	F	−0.85 / NA	−0.40 / NA	−1.15 / NA

Table 3: Capability accuracy changes (pp) from Base, shown as trained variant / prompt. NA indicates unavailable Prompt results for Llama 3.1 8B.

6.2.1 Situational Interpretation

On Circa, T reduces UNCLEAR predictions and F increases them in all four models (Figure 3(a)). In Llama 3.2 3B, the proportion changes from 23.9% at Base to 2.7% for T and 72.4% for F. These differences indicate contrasting tendencies to assign definite interpretations to indirect replies.

On CommitmentBank, T increases FALSE predictions and F decreases them across all four models, although the redistribution among other labels varies (Figure 3(b)). These judgment shifts do not yield uniformly better task performance: T improves macro F1 on Circa in all four models, while its effects on CommitmentBank differ across models. Personality preferences therefore affect interpretation patterns, with consequences for accuracy depending on the task and model. FOMC results, scoring conventions, and output processing are detailed in Appendix C.

6.2.2 Individual Decision Making

On Choices13k and the machine selection task, contrasting preferences produce different choices under the same decision conditions. T chooses alternatives with lower payoff variance more often than F in three of four models (Figure 4(a)). T also selects machines with fewer previous observations more often than F in all four models; the proportions reach 84.2% and 31.3%, respectively, in Llama 3.2 3B (Figure 4(b)). These results reveal differences in choices involving uncertainty and limited evidence.

The expected payoff probe tests responses to

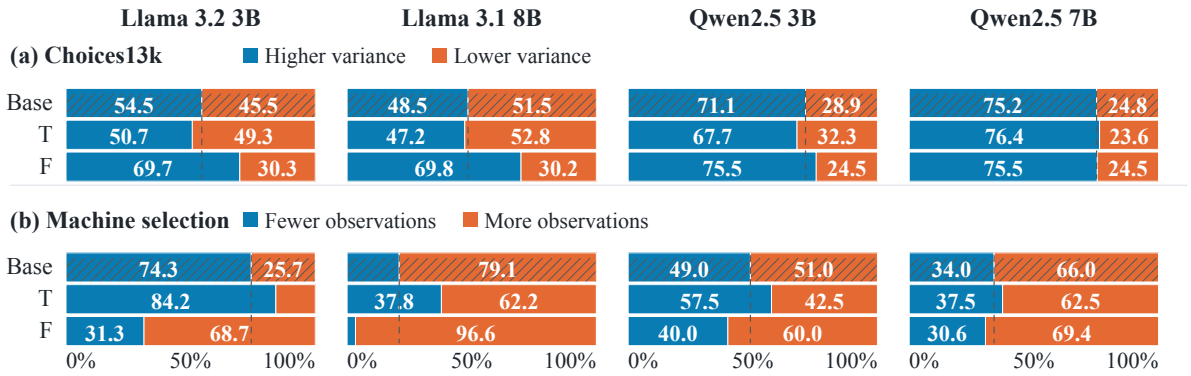

Figure 4: Individual decision making: choice distributions (%) on (a) Choices13k and (b) machine selection.

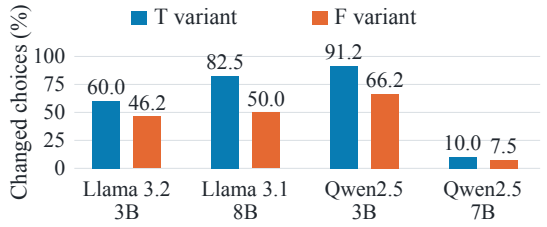

Figure 5: Individual decision making: choice changes (%) after expected payoffs are explicitly stated for otherwise identical risk scenarios.

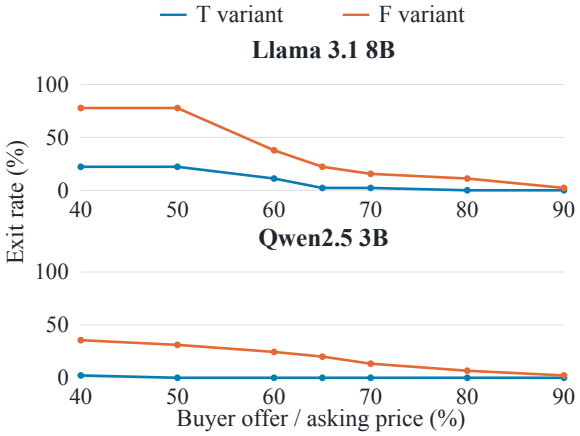

Figure 6: Individual decision making: exit rates (%) across buyer offers in the controlled bargaining probe.

changes in how a situation is presented. Adding calculated expected payoffs, while keeping probabilities and outcomes unchanged, changes T choices more often than F choices in all four models (Figure 5). In Llama 3.1 8B, 82.5% of T choices change, compared with 50.0% of F choices.

In the controlled bargaining probe, F has higher exit rates at low buyer offers, while the difference between T and F narrows as offers improve in Llama 3.1 8B and Qwen2.5 3B (Figure 6). Thus, personality preferences affect both action selection and responsiveness to decision conditions.

6.2.3 Interactive Decision Making

Personality preferences also affect how models respond during recommendation interactions. On INSPIRED, F requests more information than T in all four models (Figure 7(a)). In Llama 3.2 3B, further inquiries account for 74.8% of F responses and 5.9% of T responses.

On Frames, T accepts travel proposals more often than F in all four models, reaching 39.0% versus 18.8% in Llama 3.2 3B (Figure 7(b)). These tasks show distinct effects on whether models accept a proposal or continue seeking information.

Together, these results show personality influence on situational interpretation, individual choices, and interaction outcomes. The observed differences have consequences for task accuracy and interaction choices. Additional decision and interaction results are provided in Appendix D.4.

6.3 Persistence of Personality Influence under Opposing Instructions

We examine whether these trained preferences continue to influence decisions when subsequent instructions explicitly request the opposing personality. Figure 8 presents a retained example and a reversal, both using the Feeling variant with a Thinking prompt.

On the machine selection task, the Llama 3.2 3B Feeling variant selects machines with fewer observations in 31.3% of cases, compared with 74.3% for Base (Figure 8(a)). Adding a Thinking prompt raises this proportion to 66.3%, but it remains below Base. Under the same prompt, Base selects these machines in 78.6% of cases, leaving a difference of 12.3 percentage points. The instruction substantially changes behavior, yet the trained Feeling preference continues to influence the decision.

On Choices13k, risk choices show a reversal in Llama 3.2 3B (Figure 8(b)). The Feeling variant selects alternatives with higher payoff variance in 69.7% of cases, compared with 54.5% for Base. Adding a Thinking prompt reduces this proportion to 24.4%, below Base. This example shows that opposing instructions can reverse behavioral shifts in some situations.

Across the three situational levels, 73.9% of comparisons are classified as retained and 26.1% as reversed (Table 4). Reversal occurs when the trained and combined conditions shift the focal response proportion in opposite directions relative to Base. The dominant pattern is therefore retention rather than reversal under opposing instructions.

Situational level	Retained	Reversed
Situational interpretation	78.6%	21.4%
Individual decision making	68.8%	31.3%
Interactive decision making	78.1%	21.9%
Overall	**73.9%**	**26.1%**

Table 4: Retention and reversal across 138 opposing prompt comparisons.

These results show that personality tendencies exert a persistent influence on model decisions.

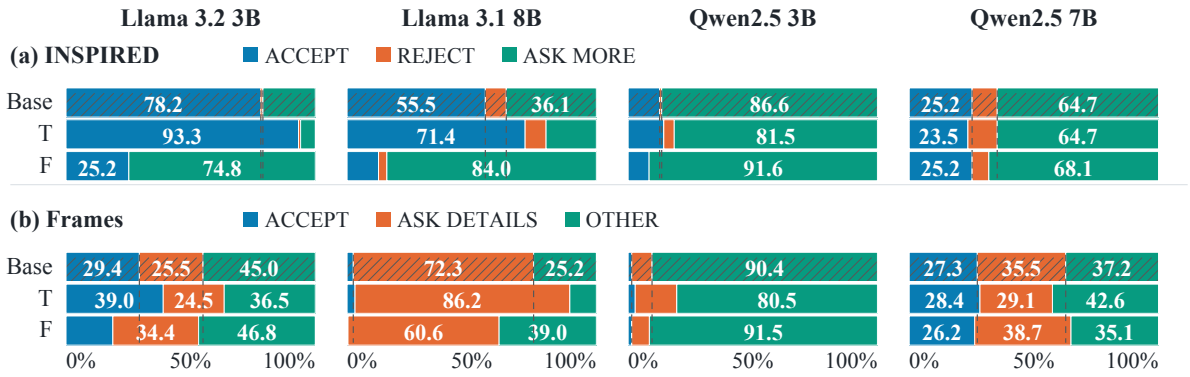

Figure 7: Interactive decision making: choice distributions (%) on (a) INSPIRED and (b) Frames.

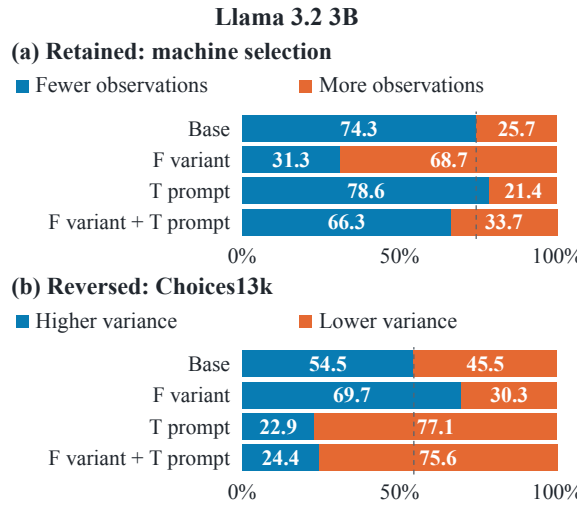

Figure 8: Individual decision making under opposing instructions: (a) a retained shift on machine selection and (b) a reversed shift on Choices13k.

Opposing prompts can modify behavior and sometimes reverse particular choice patterns, but they do not consistently override the influence of underlying personality preferences. The differences observed under the same prompt show that this influence extends beyond generation style and immediate instruction effects. Detailed results and classification criteria are provided in Appendix D.

6.4 Statistical Reliability

We assess T/F differences across situational interpretation, individual decision making, and interactive decision making using paired tests with Holm correction. The results support personality effects at all three levels, although significance varies across tasks and models. Appendix E reports effect sizes, confidence intervals, and the complete test results.

7 Conclusion

Our experiments show that personality influence in LLMs extends beyond generation style. Using Thinking and Feeling preferences as a controlled contrast, we identify systematic differences in how models interpret information, select actions, and respond during interaction. These preferences continue to influence decisions under opposing instructions, even when prompts substantially change outward behavior. Particular choice patterns can weaken or reverse, but immediate instructions do not consistently override the influence of underlying preferences.

SPACE provides a benchmark framework for examining this influence through a controlled evaluation protocol and complementary datasets. Personality verification and basic capability assessments support comparisons across models with contrasting preferences, while hierarchical situational tasks examine their effects on interpretation, individual decision making, and interactive decision making. Together, the protocol and datasets help distinguish personality influence from measured capability changes and immediate instruction effects.

Personality preferences also have consequences for downstream tasks. They affect interpretation accuracy, choices under uncertainty, responses to other participants, and final interaction outcomes. Their effects vary across tasks and models: a preference that improves performance in one setting can reduce it in another. Personality therefore matters not only for how a model expresses itself, but also for the judgments it reaches, the actions it selects, and the outcomes of its interactions.

These findings establish personality preferences as an important dimension of LLM evaluation and

behavioral guidance. Outward personality expression alone does not reveal how a model will judge or act in concrete situations. Evaluating situational behavior and its persistence under competing instructions provides a fuller account of personality influence. These insights may also support more context-sensitive model evaluation and development. SPACE offers a reusable basis for evaluating personality influence and examining how personality preferences can align model behavior with task requirements.

Limitations

Personality is a complex construct involving diverse behavioral tendencies, interactions, and broader personality schemas. Our study focuses on selected personality dimensions and situational tasks, without exhaustively covering these aspects or their interactions. As an initial step toward understanding personality influence in LLMs, this work lays the groundwork for broader investigations, including more comprehensive personality benchmarking and evaluation frameworks.

Furthermore, personality in this study refers to underlying behavioral preferences and decision-making tendencies exhibited by language models, rather than human psychological traits. Whether and to what extent these tendencies resemble human personality remains an open question requiring further investigation. Our focus is therefore on evaluating observable model behavior and outputs in controlled situational tasks, without making claims about the psychological nature of model personality.

References

Albert Bandura. 1989. Human agency in social cognitive theory. *American Psychologist*, 44(9):1175–1184.

Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. 2019. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2924–2936, Minneapolis, Minnesota. Association for Computational Linguistics.

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *Preprint*, arXiv:1803.05457.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *Preprint*, arXiv:2110.14168.

Jiaxi Cui, Liuzhenghao Lv, Jing Wen, Rongsheng Wang, Jing Tang, Yonghong Tian, and Li Yuan. 2023. Machine mindset: An MBTI exploration of large language models. *Preprint*, arXiv:2312.12999.

Marie-Catherine de Marneffe, Mandy Simons, and Judith Tonhauser. 2019. The CommitmentBank: Investigating projection in naturally occurring discourse. In *Proceedings of Sinn und Bedeutung*, volume 23, pages 107–124.

Ameet Deshpande, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, and Karthik Narasimhan. 2023. Toxicity in ChatGPT: Analyzing persona-assigned language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 1236–1270, Singapore. Association for Computational Linguistics.

Kenneth A. Dodge and Joseph M. Price. 1994. On the relation between social information processing and socially competent behavior in early school-aged children. *Child Development*, 65(5):1385–1397.

Rotem Dror, Gili Baumer, Segev Shlomov, and Roi Reichart. 2018. The hitchhiker's guide to testing statistical significance in natural language processing. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1383–1392, Melbourne, Australia. Association for Computational Linguistics.

Abhimanyu Dubey et al. 2024. The Llama 3 herd of models. *Preprint*, arXiv:2407.21783.

Layla El Asri, Hannes Schulz, Shikhar Sharma, Jeremie Zumer, Justin Harris, Emery Fine, Rahul Mehrotra, and Kaheer Suleman. 2017. Frames: a corpus for adding memory to goal-oriented dialogue systems. In *Proceedings of the 18th Annual SIGdial Meeting on Discourse and Dialogue*, pages 207–219, Saarbrücken, Germany. Association for Computational Linguistics.

Shirley Anugrah Hayati, Dongyeop Kang, Qingxiaoyang Zhu, Weiyan Shi, and Zhou Yu. 2020. INSPIRED: Toward sociable recommendation dialog systems. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8142–8152, Online. Association for Computational Linguistics.

He He, Derek Chen, Anusha Balakrishnan, and Percy Liang. 2018. Decoupling strategy and generation in

negotiation dialogues. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2333–2343, Brussels, Belgium. Association for Computational Linguistics.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*.

Guangyuan Jiang, Manjie Xu, Song-Chun Zhu, Wenjuan Han, Chi Zhang, and Yixin Zhu. 2023. Evaluating and inducing personality in pre-trained language models. In *Advances in Neural Information Processing Systems*, volume 36.

Hang Jiang, Xiajie Zhang, Xubo Cao, Cynthia Breazeal, Deb Roy, and Jad Kabbara. 2024. PersonaLLM: Investigating the ability of large language models to express personality traits. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3605–3627, Mexico City, Mexico. Association for Computational Linguistics.

Seungbeen Lee, Seungwon Lim, Seungju Han, Giyeong Oh, Hyungjoo Chae, Jiwan Chung, Minju Kim, Beong-woo Kwak, Yeonsoo Lee, Dongha Lee, Jinyoung Yeo, and Youngjae Yu. 2025. Do LLMs have distinct and consistent personality? TRAIT: Personality testset designed for LLMs with psychometrics. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 8412–8452, Albuquerque, New Mexico. Association for Computational Linguistics.

Wenkai Li, Jiarui Liu, Andy Liu, Xuhui Zhou, Mona T. Diab, and Maarten Sap. 2025. BIG5-CHAT: Shaping LLM personalities through training on human-grounded data. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 20434–20471, Vienna, Austria. Association for Computational Linguistics.

Annie Louis, Dan Roth, and Filip Radlinski. 2020. "I'd rather just go to bed": Understanding indirect answers. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7411–7425, Online. Association for Computational Linguistics.

Meta AI. 2024. Llama 3.2 model card.

Keyu Pan and Yawen Zeng. 2023. Do LLMs possess a personality? making the MBTI test an amazing evaluation for large language models. *Preprint*, arXiv:2307.16180.

Joshua C. Peterson, David D. Bourgin, Mayank Agrawal, Daniel Reichman, and Thomas L. Griffiths. 2021. Using large-scale experiments and machine learning to discover theories of human decision-making. *Science*, 372(6547):1209–1214.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, volume 36.

Agam Shah, Suvan Paturi, and Sudheer Chava. 2023. Trillion dollar words: A new financial dataset, task & market analysis. *Preprint*, arXiv:2305.07972.

Mengyu Wang, Shay B Cohen, and Tiejun Ma. 2024a. Modeling news interactions and influence for financial market prediction. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 3302–3314.

Mengyu Wang and Tiejun Ma. 2024. Mana-net: Mitigating aggregated sentiment homogenization with news weighting for enhanced market prediction. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 2379–2389.

Mengyu Wang, Sotirios Sabanis, Miguel de Carvalho, Shay B Cohen, and Tiejun Ma. 2025. One more question is enough, expert question decomposition (eqd) model for domain quantitative reasoning. *arXiv preprint arXiv:2510.01526*.

Mengyu Wang, Xiaoying Zhi, Zhiyi Li, Robin Schmucker, Shay B Cohen, Tiejun Ma, and Fran Silavong. 2026. Self knowledge re-expression: A fully local method for adapting llms to tasks using intrinsic knowledge. *arXiv preprint arXiv:2604.22939*.

Noah Wang, Z.Y. Peng, Haoran Que, Jiaheng Liu, Wangchunshu Zhou, Yuhan Wu, Hongcheng Guo, Ruitong Gan, Zehao Ni, Jian Yang, Man Zhang, Zhaoxiang Zhang, Wanli Ouyang, Ke Xu, Wenhao Huang, Jie Fu, and Junran Peng. 2024b. RoleLLM: Benchmarking, eliciting, and enhancing role-playing abilities of large language models. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 14743–14777, Bangkok, Thailand. Association for Computational Linguistics.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, and 22 others. 2024. Qwen2.5 technical report. *Preprint*, arXiv:2412.15115.

Ivan Zakazov, Mikolaj Boronski, Lorenzo Drudi, and Robert West. 2024. Assessing social alignment: Do personality-prompted large language models behave like humans? *Preprint*, arXiv:2412.16772.

A Data, Training, and Prompt Details

A.1 Task Coverage and Adaptation

Table 5 reports the retained inputs for the principal task versions. In appendix tables, Llama 3B

Behavioral level	Task version	Inputs per condition
Interpretation	Circa	32,602
	CommitmentBank	1,009
	FOMC	2,421
	FiQA-SA / SST-2 / News	1,173 / 10,000 / 11,931
Individual decisions	Choices13k (unequal payoff variances)	1,764
	Transportation	210
	Machine selection	4,160
	Delayed rewards	288
	Original stock windows	2,074
	Controlled stock paths	1,005
	Original bargaining with a single turn	2,334
	Controlled buyer offers	900
Interactive decisions	INSPIRED	119
	Frames	282
	Real bargaining histories	351
	Bargaining feedback generated by a program	2,272

Table 5: Task coverage. Questionnaire and capability evaluation details are provided in Appendix B.

and Llama 8B denote Llama 3.2 3B and Llama 3.1 8B; Qwen 3B and Qwen 7B denote Qwen2.5 3B and Qwen2.5 7B. T and F denote the Thinking and Feeling variants. Multiple slices of one dialogue and numerical variants of one source situation are related observations. These results on adapted tasks are not official benchmark scores on held out data. The controlled probe of buyer offers in Figure 6 is an additional targeted experiment; it is separate from the controlled bargaining task with 900 inputs.

A.2 Selection and Task Construction

Circa maps indirect replies to YES, NO, CONDITIONAL, or UNCLEAR. CommitmentBank uses the retained source statements and adapted majority references described in Appendix C. FOMC retains the frozen monetary policy label mapping. Choices13k includes 1,766 choices, of which 1,764 permit a comparison of payoff variances; human choice rates are not correctness labels.

Transportation presents a faster, more expensive route and a slower, cheaper alternative. Machine selection presents one observation of machine A and three of B, with either one or six choices remaining. Delayed rewards vary amounts and waiting times while guaranteeing both payments. Original stock tasks use 2,074 overlapping windows from 2,102 asset/date records across seven assets. Controlled stock tasks use 67 paths for each of five net changes and three volatility levels, with the final price fixed at $100.

Original bargaining retains 2,334 initial offer slices from 6,682 CraigslistBargain dialogues. Controlled bargaining uses 150 source situations across five item categories and three price bands, with six synthetic buyer offers from 40% to 90% of asking price. Real history tasks retain 351 inputs from 314 dialogues after removing two duplicated prompts; they preserve offer histories but omit other textual negotiation conditions. Program feedback tasks remove 62 initial offers at or above asking price, retaining 2,272 inputs and at most three model choices.

INSPIRED and Frames retain 119 and 282 eligible dialogue prefixes from 999 and 1,369 downloaded dialogues. Short prefixes satisfying the extraction conditions are selected and known input problems are removed; this is not exhaustive manual validation. Their source histories provide recommendation context. The separate controlled feedback probes use two authored versions per domain and are reported separately in Appendix E.

The expected payoff probe in Figure 5 holds probabilities and payouts fixed and varies their explicit summary; its replication uses 20 economic cells and 80 inputs including both option orders and presentation frames. The buyer offer probe in Figure 6 uses 15 source situations and 315 inputs including offer levels and option orders. These targeted probes are distinct from the full task versions in Table 5.

A.3 Personality Prompts

The fixed prefixes used during inference are:

> Use an MBTI Thinking (T) preference: Make decisions using objective logic and consistency. Apply it silently and follow the task's required output format exactly.

> Use an MBTI Feeling (F) preference: Make decisions using empathy and personal values. Apply it silently and follow the task's required output format exactly.

An aligned combination uses the same trained and prompted preference; an opposing combination uses the other preference. Prefixes are applied to the task input and reapplied at each model turn in interactive conditions.

A.4 Preference Pairs

We pair the opposing Thinking and Feeling responses in the English behavior split by row index. All paired rows have matching instructions and additional context. The targeted response is chosen and the opposing response rejected; switching the target reverses the pair. We concatenate the instruction and any nonempty additional context, then apply the corresponding model's chat template. Rows missing an instruction or either response are excluded; none of the retained files require such exclusions. The T/F training set contains 12,159 preference pairs. The full pair sets are used for training without a held out validation split; situational evaluations supply separate task inputs.

A.5 Training Details

The training recipe uses learning rate 10^{-5}, DPO $\beta = 1.0$, seed 42, batch size per device 16, and gradient accumulation of two. LoRA uses rank 32, scaling 64, and dropout 0.1 on attention projections and feedforward projections. The Base weights provide the frozen reference; the Llama 3.1 8B implementation obtains this reference by disabling its LoRA adapters. Adapters are merged for inference.

Early stopping monitors training DPO loss with a constant learning rate. Monitoring begins after one epoch, averages windows of 100 steps, and stops after two windows without a 5% relative improvement over the best window. No MBTI or benchmark score selects the stopping point.

A.6 Computational Resources and Budget

Training and evaluation were run on NVIDIA GH200 GPUs on the Isambard-AI HPC platform, with one GPU allocated per job. Based on available Slurm accounting, the recorded jobs consumed approximately 392.2 GPU-hours in total, including pilot runs, reruns, and unsuccessful attempts. Table 6 summarizes the infrastructure and budget.

Item	Value
GPU model	NVIDIA GH200
GPUs per job	1
Computing platform	Isambard-AI HPC
Training GPU-hours	≈ 97.9
Evaluation GPU-hours	≈ 294.3
Total GPU-hours	≈ 392.2

Table 6: Computational resources and approximate GPU-hours from available Slurm accounting.

B Supplementary Personality and Capability Results

B.1 Evaluation Scope and Reporting

Capability evaluation covers four models, two trained preferences, and three benchmarks. Prompt and combined guidance results cover Llama 3.2 3B, Qwen2.5 3B, and Qwen2.5 7B. Unprompted questionnaire evaluation covers all four models, while Prompt and combined questionnaire evaluation covers the first three.

GSM8K, ARC-Easy, and BoolQ contain 1,300, 2,000, and 2,000 evaluated inputs, respectively. The questionnaire uses 46 T/F items. Percentages in the main questionnaire table are calculated from these choices.

Llama 3.2 3B Base, trained variant, and Prompt GSM8K evaluation uses a 512 token limit and counts an answer as matching when any extracted response number matches the reference. Combined GSM8K and Qwen Prompt evaluation extracts the final number from each response. ARC-Easy and BoolQ use label extraction, with unparseable responses counted as incorrect. ARC-Easy uses the full denominator, including 87 reference labels unrecognized by the A to D parser.

B.2 Trained Variant Capability Results

Table 7 supports the broadly stable capability pattern in the main text: 22 of 24 absolute changes are below 1.2 percentage points. The largest decline is Llama 3B T on BoolQ; its statistical evidence is reported in Appendix E.6.5.

B.3 Prompt Capability Results

Prompting lowers ARC-Easy accuracy in all six evaluated conditions, while BoolQ and GSM8K changes are less uniform. These results show that personality instructions can also affect objective task performance.

Model	Condition	ARC-Easy	BoolQ	GSM8K
Llama 3B	Base	82.30	73.50	79.23
	T	81.75	71.25	78.54
	F	82.15	75.15	78.38
Qwen 3B	Base	89.00	79.60	11.54
	T	88.80	79.90	11.00
	F	88.90	78.75	12.31
Qwen 7B	Base	91.35	86.10	15.54
	T	91.30	86.30	16.00
	F	91.20	85.90	14.92
Llama 8B	Base	87.60	84.50	87.31
	T	87.55	84.45	87.23
	F	86.75	84.10	86.15

Table 7: Base and trained variant capability accuracy (%).

Model	Condition	ARC-Easy	BoolQ	GSM8K
Llama 3B	Base P	82.25	73.50	79.23
	Prompt T	78.55	72.70	77.69
	Prompt F	76.60	73.70	77.38
Qwen 3B	Base P	89.00	79.60	11.77
	Prompt T	87.05	79.25	11.46
	Prompt F	87.00	79.35	11.54
Qwen 7B	Base P	91.35	86.10	15.54
	Prompt T	90.50	82.05	15.31
	Prompt F	90.00	82.25	15.31

Table 8: Prompt capability accuracy (%).

B.4 Personality Scores under Combined Guidance

In X+Y, X denotes the trained Thinking or Feeling variant and Y the personality prompt. T+T and F+F are aligned; T+F and F+T are opposing. Unprompted variant scores are included as the reference. Combined prompts can substantially shift questionnaire choices toward the requested preference. The questionnaire checks whether the guidance produces the intended response patterns. Persistence of personality influence is assessed through situational tasks in Appendix D.

B.5 Capability Scores under Combined Guidance

Adding either personality prompt increases GSM8K scores relative to the corresponding unprompted trained variant, whereas ARC-Easy and BoolQ often decline.

B.6 Llama 3.1 8B Questionnaire Results

The unprompted questionnaire results confirm the intended T/F shifts for Llama 3.1 8B. Base has 43.5% F and 56.5% T choices; F has 58.7% F and 41.3% T, while T has 39.1% F and 60.9% T. Prompt and combined questionnaire evaluations

Model	Condition	F	T
Llama 3B	T	39.1	60.9
	F	82.6	17.4
	T+T	6.5	93.5
	T+F	91.3	8.7
	F+T	10.9	89.1
	F+F	93.5	6.5
Qwen 3B	T	30.4	69.6
	F	54.3	45.7
	T+T	2.2	97.8
	T+F	84.8	15.2
	F+T	13.0	87.0
	F+F	93.5	6.5
Qwen 7B	T	41.3	58.7
	F	56.5	43.5
	T+T	4.3	95.7
	T+F	95.7	4.3
	F+T	2.2	97.8
	F+F	95.7	4.3

Table 9: T/F questionnaire choice proportions (%) for trained and combined conditions.

Model	Condition	ARC-Easy	BoolQ	GSM8K
Llama 3B	T	82.05	68.85	35.00
	F	81.70	76.05	13.08
	T+T	78.20	67.45	54.92
	T+F	77.60	69.05	54.31
	F+T	80.10	75.20	55.62
	F+F	75.85	75.90	52.77
Qwen 3B	T	88.30	79.20	7.00
	F	88.20	79.90	7.08
	T+T	87.10	79.40	10.92
	T+F	86.95	79.15	11.69
	F+T	87.30	78.75	9.69
	F+F	87.20	78.75	10.08
Qwen 7B	T	91.15	86.70	9.38
	F	91.05	86.60	9.38
	T+T	90.40	81.95	15.23
	T+F	89.70	82.20	14.62
	F+T	90.50	82.05	14.85
	F+F	89.70	82.45	14.69

Table 10: Capability accuracy (%) under combined guidance.

C Interpretation Metrics and Output Processing

Distribution denominators include all retained inputs. Labels are extracted from an explicit leading task label; a valid leading label followed by an explanation is included. This differs from strict compliance with the instruction to output a single label. Unparseable outputs remain in distribution denominators and are counted as incorrect when a reference exists. Circa and FOMC use the frozen mappings of published annotations. CommitmentBank uses an adapted strict majority reference based on individual rating signs; this is not an original gold label with three classes. Of its 1,009 inputs, 858 have such a reference and enter accuracy and macro F1, while 151 enter distributions only. Macro F1 averages predefined task classes. Table 11 reports both metrics.

C.1 CommitmentBank Format Audit

For Llama 3.1 8B, 970 of 1,009 F guided outputs explicitly begin with UNCERTAIN (96.13%); parser failures are not mapped to UNCERTAIN. Only 77 F outputs obey the strict format requiring a single label, compared with 580 Base outputs and 819 T outputs. All 932 F outputs with an additional explanation begin with UNCERTAIN. Among the 75 inputs strictly formatted in all three conditions, UNCERTAIN counts are 36 for Base, 27 for T, and 37 for F. Thus the distribution over all inputs is an observed label pattern, but it does not establish an effect independent of output format. The subset comparison does not identify formatting as its cause.

Dataset	Model	Accuracy (%) Base / T / F	Macro F1 Base / T / F
Circa	Llama 3.2 3B	58.54 / 70.52 / 23.60	0.4914 / 0.5617 / 0.2630
Circa	Qwen2.5 3B	12.10 / 13.17 / 8.62	0.1581 / 0.1690 / 0.1062
Circa	Qwen2.5 7B	27.27 / 29.30 / 23.95	0.3270 / 0.3498 / 0.2953
Circa	Llama 3.1 8B	27.54 / 41.77 / 13.13	0.2597 / 0.3743 / 0.1514
CommitmentBank	Llama 3.2 3B	48.02 / 45.57 / 55.01	0.3131 / 0.2654 / 0.4033
CommitmentBank	Qwen2.5 3B	36.01 / 40.79 / 25.99	0.3265 / 0.3561 / 0.2445
CommitmentBank	Qwen2.5 7B	66.20 / 67.13 / 60.61	0.6416 / 0.6487 / 0.5946
CommitmentBank	Llama 3.1 8B	33.22 / 55.48 / 16.78	0.3456 / 0.5390 / 0.1404
FOMC	Llama 3.2 3B	51.80 / 51.05 / 52.75	0.4000 / 0.4105 / 0.3879
FOMC	Qwen2.5 3B	56.67 / 57.37 / 56.18	0.4442 / 0.4584 / 0.4288
FOMC	Qwen2.5 7B	63.07 / 63.61 / 61.63	0.5734 / 0.5768 / 0.5590
FOMC	Llama 3.1 8B	59.48 / 61.46 / 56.63	0.4828 / 0.5355 / 0.4100

Table 11: Situational interpretation: Accuracy (%) and macro F1 for Base / T / F, in that order. CommitmentBank uses 858 inputs with a majority reference.

are available for the other three models.

C.2 Full Interpretation Label Distributions

All retained inputs remain in the denominator, including unparseable outputs. The tables report the same leading label parser used in the main figures. On Circa, T reduces UNCLEAR judgments and F increases them across all four models. The full distributions show how definite and conditional interpretations account for the remaining responses, extending the contrast discussed in the main text.

Model	Label	Base	T	F
Llama 3B	UNCLEAR	23.90	2.66	72.40
Llama 3B	NO	19.76	45.06	0.44
Llama 3B	YES	41.60	37.06	17.59
Llama 3B	CONDITIONAL	14.74	15.22	9.57
Llama 3B	None	0.00	0.00	0.00
Llama 8B	UNCLEAR	27.85	10.72	68.35
Llama 8B	NO	8.88	20.19	0.80
Llama 8B	YES	10.35	15.09	2.97
Llama 8B	CONDITIONAL	52.91	54.00	27.88
Llama 8B	None	0.00	0.00	0.01
Qwen 3B	UNCLEAR	89.31	87.87	93.16
Qwen 3B	NO	4.78	5.13	2.73
Qwen 3B	YES	4.03	4.79	2.71
Qwen 3B	CONDITIONAL	1.86	2.21	1.32
Qwen 3B	None	0.02	0.00	0.08
Qwen 7B	UNCLEAR	64.41	63.59	67.12
Qwen 7B	NO	12.12	13.83	9.59
Qwen 7B	YES	6.50	7.05	5.41
Qwen 7B	CONDITIONAL	16.89	15.48	17.75
Qwen 7B	None	0.08	0.06	0.13

Table 12: Situational interpretation: Circa: full prediction distributions (%).

On CommitmentBank, T increases FALSE judgments and F decreases them across models. The large UNCERTAIN proportion for Llama 8B F should be read alongside the format audit above; label distributions and strict format compliance capture different aspects of its responses.

Model	Label	Base	T	F
Llama 3B	FALSE	84.24	92.86	55.90
Llama 3B	TRUE	15.76	7.14	43.11
Llama 3B	UNCERTAIN	0.00	0.00	0.99
Llama 8B	FALSE	12.39	30.82	1.19
Llama 8B	TRUE	8.72	20.22	2.68
Llama 8B	UNCERTAIN	78.89	48.96	96.13
Qwen 3B	FALSE	25.57	31.81	13.97
Qwen 3B	TRUE	0.40	0.50	0.10
Qwen 3B	UNCERTAIN	74.03	67.69	85.93
Qwen 7B	FALSE	25.97	27.16	20.52
Qwen 7B	TRUE	29.63	29.44	29.24
Qwen 7B	UNCERTAIN	44.40	43.41	50.25

Table 13: Situational interpretation: CommitmentBank: full prediction distributions (%).

FOMC extends interpretation to monetary policy language. For Llama 3B, DOVISH judgments rise from 45.39% at Base to 52.29% for T and fall to 35.11% for F. Other models redistribute labels differently, so these results describe policy judgments rather than a common sentiment scale.

Model	Label	Base	T	F
Llama 3B	NEUTRAL	53.61	46.22	64.23
Llama 3B	DOVISH	45.39	52.29	35.11
Llama 3B	HAWKISH	0.99	1.49	0.66
Llama 8B	NEUTRAL	83.97	76.04	90.13
Llama 8B	DOVISH	6.86	11.11	2.60
Llama 8B	HAWKISH	9.17	12.85	7.27
Qwen 3B	NEUTRAL	81.62	80.83	84.18
Qwen 3B	DOVISH	10.74	10.24	9.25
Qwen 3B	HAWKISH	7.64	8.92	6.57
Qwen 7B	NEUTRAL	54.52	59.23	50.19
Qwen 7B	DOVISH	38.70	33.87	43.49
Qwen 7B	HAWKISH	6.77	6.90	6.32

Table 14: Situational interpretation: FOMC: full prediction distributions (%).

D Complete Comparisons and Additional Results

D.1 Retention by Task

Table 15 reports every selected task version. FiQA-SA, SST-2, and News cover three models; the remaining task versions cover four models. Each comparison specifies a model, task, and trained preference and compares Base, trained variant, and combined conditions. Original and controlled stock or bargaining versions belong to related task families.

Task version	Focal label	Retained	Reversed
CommitmentBank	FALSE	87.5%	12.5%
Circa	UNCLEAR	75.0%	25.0%
FOMC	HAWKISH	62.5%	37.5%
FiQA-SA	positive	100.0%	0.0%
SST-2	positive	83.3%	16.7%
News	bullish	66.7%	33.3%
Risk choice	Higher variance	37.5%	62.5%
Machine choice	Fewer observations	62.5%	37.5%
Travel: time and cost	Faster route	75.0%	25.0%
Delayed payments	SOONER	75.0%	25.0%
Original stock	BUY	62.5%	37.5%
Original bargaining with a single turn	EXIT	75.0%	25.0%
Controlled stock	BUY	75.0%	25.0%
Controlled bargaining	EXIT	87.5%	12.5%
INSPIRED	Ask more	62.5%	37.5%
Frames	ACCEPT	75.0%	25.0%
Real bargaining history	EXIT	87.5%	12.5%
Program bargaining: FIRST action	EXIT	87.5%	12.5%

Table 15: Retention and reversal across all selected task versions.

D.2 Directional Criterion

For the proportion of a focal label p, define $\Delta_D = p_{\text{trained}} - p_{\text{Base}}$ and $\Delta_C = p_{\text{combined}} - p_{\text{Base}}$. Reversal requires $\Delta_D \Delta_C < 0$; otherwise the result is retained. Signs use unrounded counts. Zero initial or combined shifts count as retained by convention. The total comprises 102 retained and 36 reversed results out of 138. Retention does not imply unchanged individual answers or unchanged effect magnitude. Decision and interaction labels are behavioral measures; human choice frequencies and simulated outcomes are not objective correctness labels.

D.3 Complete Opposing Prompt Comparisons

The complete comparisons include both partial attenuation and reversal. A retained direction can coexist with a substantial change in magnitude, as the machine choice example shows. Reporting Base, prompting, and combined conditions together makes the remaining training influence visible beyond the binary retention summary.

The following tables report the focal label proportion for every model, task, and trained preference. Prompt is applied to Base; Opposed applies that prompt to the trained variant. Each row uses unrounded counts for classification.

Task	Model	Variant	Base	Trained	Prompt	Opposed	Outcome
CommitmentBank	Llama 3B	T	84.24	92.86	66.20	89.40	Retained
CommitmentBank	Llama 3B	F	84.24	55.90	83.55	50.74	Retained
CommitmentBank	Llama 8B	T	12.39	30.82	25.17	54.01	Retained
CommitmentBank	Llama 8B	F	12.39	1.19	45.39	2.87	Retained
CommitmentBank	Qwen 3B	T	25.57	31.81	26.16	32.61	Retained
CommitmentBank	Qwen 3B	F	25.57	13.97	29.34	20.02	Retained
CommitmentBank	Qwen 7B	T	25.97	27.16	15.86	18.83	Reversed
CommitmentBank	Qwen 7B	F	25.97	20.52	26.16	19.23	Retained
Circa	Llama 3B	T	23.90	2.66	5.82	1.73	Retained
Circa	Llama 3B	F	23.90	72.40	3.26	33.15	Retained
Circa	Llama 8B	T	27.85	10.72	10.32	0.29	Retained
Circa	Llama 8B	F	27.85	68.35	2.87	40.94	Retained
Circa	Qwen 3B	T	89.31	87.87	88.23	85.45	Retained
Circa	Qwen 3B	F	89.31	93.16	81.66	87.63	Reversed
Circa	Qwen 7B	T	64.41	63.59	49.60	48.16	Retained
Circa	Qwen 7B	F	64.41	67.12	35.21	38.37	Reversed
FOMC	Llama 3B	T	0.99	1.49	0.45	0.37	Reversed
FOMC	Llama 3B	F	0.99	0.66	2.15	0.99	Retained
FOMC	Llama 8B	T	9.17	12.85	14.04	16.27	Retained
FOMC	Llama 8B	F	9.17	7.27	25.73	18.09	Reversed
FOMC	Qwen 3B	T	7.64	8.92	8.47	11.03	Retained
FOMC	Qwen 3B	F	7.64	6.57	9.05	7.35	Retained
FOMC	Qwen 7B	T	6.77	6.90	4.34	4.63	Reversed
FOMC	Qwen 7B	F	6.77	6.32	5.66	5.20	Retained
FiQA-SA	Llama 3B	T	36.49	17.05	28.64	12.36	Retained
FiQA-SA	Llama 3B	F	36.49	53.62	22.51	39.90	Retained
FiQA-SA	Qwen 3B	T	47.74	45.52	43.73	40.24	Retained
FiQA-SA	Qwen 3B	F	47.74	54.31	44.50	50.47	Retained
FiQA-SA	Qwen 7B	T	41.60	38.96	39.56	35.38	Retained
FiQA-SA	Qwen 7B	F	41.60	45.44	36.66	41.60	Retained
SST-2	Llama 3B	T	39.56	23.88	39.82	23.20	Retained
SST-2	Llama 3B	F	39.56	50.96	36.67	51.61	Retained
SST-2	Qwen 3B	T	31.86	28.88	34.91	30.19	Retained
SST-2	Qwen 3B	F	31.86	35.59	30.80	34.63	Retained
SST-2	Qwen 7B	T	45.56	43.34	42.74	39.48	Retained
SST-2	Qwen 7B	F	45.56	47.46	37.43	39.71	Reversed
News	Llama 3B	T	42.61	29.25	48.14	34.78	Retained
News	Llama 3B	F	42.61	55.41	43.19	63.85	Retained
News	Qwen 3B	T	22.89	22.88	27.89	26.10	Reversed
News	Qwen 3B	F	22.89	25.96	21.33	24.73	Retained
News	Qwen 7B	T	30.01	27.78	30.01	26.45	Retained
News	Qwen 7B	F	30.01	34.05	25.56	29.27	Reversed

Table 16: Interpretation: complete opposing prompt comparisons (%).

Task	Model	Variant	Base	Trained	Prompt	Opposed	Outcome
Risk choice	Llama 3B	T	54.48	50.68	22.96	22.96	Retained
Risk choice	Llama 3B	F	54.48	69.73	22.90	24.38	Reversed
Risk choice	Llama 8B	T	48.53	47.22	41.50	39.34	Retained
Risk choice	Llama 8B	F	48.53	69.78	29.65	46.03	Reversed
Risk choice	Qwen 3B	T	71.15	67.74	75.28	74.09	Reversed
Risk choice	Qwen 3B	F	71.15	75.45	72.79	74.32	Retained
Risk choice	Qwen 7B	T	75.17	76.36	37.76	34.64	Reversed
Risk choice	Qwen 7B	F	75.17	75.45	42.52	45.69	Reversed
Machine choice	Llama 3B	T	74.33	84.18	63.32	76.01	Retained
Machine choice	Llama 3B	F	74.33	31.30	78.56	66.30	Retained
Machine choice	Llama 8B	T	20.91	37.84	19.47	28.17	Retained
Machine choice	Llama 8B	F	20.91	3.37	49.42	35.62	Reversed
Machine choice	Qwen 3B	T	49.04	57.50	46.97	61.15	Retained
Machine choice	Qwen 3B	F	49.04	40.00	57.36	47.07	Retained
Machine choice	Qwen 7B	T	34.04	37.45	33.37	32.45	Reversed
Machine choice	Qwen 7B	F	34.04	30.58	47.60	45.14	Reversed
Travel: time and cost	Llama 3B	T	31.90	32.38	100.00	100.00	Retained
Travel: time and cost	Llama 3B	F	31.90	0.48	100.00	10.48	Retained
Travel: time and cost	Llama 8B	T	100.00	100.00	100.00	100.00	Retained
Travel: time and cost	Llama 8B	F	100.00	100.00	100.00	100.00	Retained
Travel: time and cost	Qwen 3B	T	95.71	100.00	100.00	100.00	Retained
Travel: time and cost	Qwen 3B	F	95.71	71.43	100.00	100.00	Reversed
Travel: time and cost	Qwen 7B	T	27.14	68.57	77.14	88.57	Retained
Travel: time and cost	Qwen 7B	F	27.14	10.48	98.57	90.48	Reversed
Delayed payments	Llama 3B	T	100.00	100.00	100.00	100.00	Retained
Delayed payments	Llama 3B	F	100.00	100.00	100.00	100.00	Retained
Delayed payments	Llama 8B	T	100.00	100.00	100.00	100.00	Retained
Delayed payments	Llama 8B	F	100.00	100.00	100.00	100.00	Retained
Delayed payments	Qwen 3B	T	44.10	58.33	100.00	100.00	Retained
Delayed payments	Qwen 3B	F	44.10	44.44	70.14	35.42	Reversed
Delayed payments	Qwen 7B	T	0.35	4.86	69.44	73.26	Retained
Delayed payments	Qwen 7B	F	0.35	0.00	7.99	0.69	Reversed
Original stock	Llama 3B	T	14.13	36.98	0.00	0.00	Reversed
Original stock	Llama 3B	F	14.13	0.00	0.00	0.00	Retained
Original stock	Llama 8B	T	84.67	89.30	53.76	61.67	Reversed
Original stock	Llama 8B	F	84.67	63.89	52.27	47.30	Retained
Original stock	Qwen 3B	T	0.10	0.72	0.00	0.00	Reversed
Original stock	Qwen 3B	F	0.10	0.00	0.00	0.00	Retained
Original stock	Qwen 7B	T	28.35	27.68	0.53	0.29	Retained
Original stock	Qwen 7B	F	28.35	29.94	62.10	60.03	Retained
Original single-turn bargaining	Llama 3B	T	0.17	0.00	1.07	0.00	Retained
Original single-turn bargaining	Llama 3B	F	0.17	1.16	1.67	23.86	Retained
Original single-turn bargaining	Llama 8B	T	0.51	0.04	0.09	0.00	Retained
Original single-turn bargaining	Llama 8B	F	0.51	0.60	0.00	0.09	Reversed
Original single-turn bargaining	Qwen 3B	T	1.97	0.47	0.73	0.39	Retained
Original single-turn bargaining	Qwen 3B	F	1.97	15.72	0.86	11.31	Retained
Original single-turn bargaining	Qwen 7B	T	0.26	0.17	3.34	0.81	Reversed
Original single-turn bargaining	Qwen 7B	F	0.26	0.26	0.56	0.99	Retained
Controlled stock	Llama 3B	T	41.89	43.18	0.30	0.00	Reversed
Controlled stock	Llama 3B	F	41.89	13.23	0.00	0.00	Retained
Controlled stock	Llama 8B	T	96.12	95.62	90.75	93.53	Retained
Controlled stock	Llama 8B	F	96.12	59.60	91.24	50.15	Retained
Controlled stock	Qwen 3B	T	0.00	0.00	0.00	0.00	Retained
Controlled stock	Qwen 3B	F	0.00	0.00	0.00	0.00	Retained
Controlled stock	Qwen 7B	T	42.49	44.08	52.94	51.24	Retained
Controlled stock	Qwen 7B	F	42.49	41.99	48.06	46.97	Reversed
Controlled bargaining	Llama 3B	T	0.00	0.00	0.22	0.00	Retained
Controlled bargaining	Llama 3B	F	0.00	0.11	0.22	27.56	Retained
Controlled bargaining	Llama 8B	T	0.00	0.00	0.00	0.00	Retained
Controlled bargaining	Llama 8B	F	0.00	0.00	0.00	0.00	Retained
Controlled bargaining	Qwen 3B	T	8.89	0.89	0.00	0.00	Retained
Controlled bargaining	Qwen 3B	F	8.89	28.00	0.00	2.33	Reversed
Controlled bargaining	Qwen 7B	T	0.00	0.00	4.44	0.56	Retained
Controlled bargaining	Qwen 7B	F	0.00	0.00	0.00	3.11	Retained

Table 17: Decision: complete opposing prompt comparisons (%).

Task	Model	Variant	Base	Trained	Prompt	Opposed	Outcome
INSPIRED	Llama 3B	T	21.01	5.88	5.88	0.84	Retained
INSPIRED	Llama 3B	F	21.01	74.79	10.08	56.30	Retained
INSPIRED	Llama 8B	T	36.13	20.17	78.15	63.03	Reversed
INSPIRED	Llama 8B	F	36.13	84.03	64.71	91.60	Retained
INSPIRED	Qwen 3B	T	86.55	81.51	90.76	87.39	Reversed
INSPIRED	Qwen 3B	F	86.55	91.60	73.95	87.39	Retained
INSPIRED	Qwen 7B	T	64.71	64.71	52.94	52.94	Retained
INSPIRED	Qwen 7B	F	64.71	68.07	11.76	14.29	Reversed
Frames	Llama 3B	T	29.43	39.01	34.40	56.74	Retained
Frames	Llama 3B	F	29.43	18.79	20.57	15.60	Retained
Frames	Llama 8B	T	2.48	3.19	0.71	1.06	Reversed
Frames	Llama 8B	F	2.48	0.35	2.84	0.35	Retained
Frames	Qwen 3B	T	1.42	2.84	6.74	6.03	Retained
Frames	Qwen 3B	F	1.42	1.42	10.28	8.87	Retained
Frames	Qwen 7B	T	27.30	28.37	29.43	28.37	Retained
Frames	Qwen 7B	F	27.30	26.24	45.39	42.20	Reversed
Real bargaining history	Llama 3B	T	0.00	0.00	0.00	0.00	Retained
Real bargaining history	Llama 3B	F	0.00	0.00	0.00	1.14	Retained
Real bargaining history	Llama 8B	T	0.00	0.00	0.00	0.00	Retained
Real bargaining history	Llama 8B	F	0.00	0.00	0.00	0.00	Retained
Real bargaining history	Qwen 3B	T	85.19	27.35	11.11	0.00	Retained
Real bargaining history	Qwen 3B	F	85.19	99.72	4.27	84.05	Reversed
Real bargaining history	Qwen 7B	T	0.00	0.00	0.00	0.00	Retained
Real bargaining history	Qwen 7B	F	0.00	0.00	0.00	0.00	Retained
Program bargaining: FIRST action	Llama 3B	T	0.18	0.00	1.06	0.00	Retained
Program bargaining: FIRST action	Llama 3B	F	0.18	1.23	1.67	24.47	Retained
Program bargaining: FIRST action	Llama 8B	T	0.04	0.00	0.09	0.00	Retained
Program bargaining: FIRST action	Llama 8B	F	0.04	0.04	0.00	0.04	Retained
Program bargaining: FIRST action	Qwen 3B	T	1.54	0.04	0.62	0.09	Retained
Program bargaining: FIRST action	Qwen 3B	F	1.54	14.96	0.70	10.39	Retained
Program bargaining: FIRST action	Qwen 7B	T	0.26	0.18	3.43	0.84	Reversed
Program bargaining: FIRST action	Qwen 7B	F	0.26	0.26	0.57	1.01	Retained

Table 18: Interaction: complete opposing prompt comparisons (%).

D.4 Supplementary Situational Results

The additional tasks include clear preference differences and unchanged choices. In transportation, F favors the slower, cheaper option more than T in three models, while Llama 8B always chooses the faster option. Keeping such unchanged cases shows the full task coverage rather than only the largest contrasts.

We report Base/T/F label distributions for the full task versions, including tasks with small or zero changes. The risk table retains all 1,766 A/B choices; the variance comparison in the main figure uses 1,764 comparable choices. Interpretation performance is in Appendix C. Choice labels and program outcomes are behavioral measures, not objective correctness scores.

The supplementary results reveal task- and model-dependent behavioral differences across decision-making settings. While some tasks exhibit substantial shifts between T and F variants, others show limited or no changes.

Model	Label	Base	T	F
Llama 3B	A	41.56	61.66	8.55
Llama 3B	B	58.44	38.34	91.45
Qwen 3B	B	73.95	66.59	80.41
Qwen 3B	A	26.05	33.41	19.59
Qwen 7B	B	86.35	82.90	88.90
Qwen 7B	A	13.65	17.10	11.10
Llama 8B	A	44.28	51.47	15.29
Llama 8B	B	55.72	48.53	84.71

Table 19: Individual decision making: Risk choices (A/B): complete label distributions (%), $N = 1,766$.

Model	Label	Base	T	F
Llama 3B	Faster, costlier	31.90	32.38	0.48
Llama 3B	Slower, cheaper	68.10	67.62	99.52
Qwen 3B	Faster, costlier	95.71	100.00	71.43
Qwen 3B	Slower, cheaper	4.29	0.00	28.57
Qwen 7B	Faster, costlier	27.14	68.57	10.48
Qwen 7B	Slower, cheaper	72.86	31.43	89.52
Llama 8B	Faster, costlier	100.00	100.00	100.00

Table 20: Individual decision making: Transportation: complete label distributions (%), $N = 210$.

Model	Label	Base	T	F
Llama 3B	Fewer observations	74.33	84.18	31.30
Llama 3B	More observations	25.67	15.82	68.70
Qwen 3B	More observations	50.96	42.50	60.00
Qwen 3B	Fewer observations	49.04	57.50	40.00
Qwen 7B	More observations	65.96	62.55	69.42
Qwen 7B	Fewer observations	34.04	37.45	30.58
Llama 8B	More observations	79.09	62.16	96.63
Llama 8B	Fewer observations	20.91	37.84	3.37

Table 21: Individual decision making: Machine choice: complete label distributions (%), $N = 4,160$.

Model	Label	Base	T	F
Llama 3B	SELL	38.81	26.76	8.49
Llama 3B	HOLD	47.06	36.26	91.51
Llama 3B	BUY	14.13	36.98	0.00
Qwen 3B	SELL	99.90	99.28	100.00
Qwen 3B	BUY	0.10	0.72	0.00
Qwen 7B	SELL	19.77	19.38	19.96
Qwen 7B	BUY	28.35	27.68	29.94
Qwen 7B	HOLD	51.88	52.94	50.10
Llama 8B	SELL	12.58	9.11	6.51
Llama 8B	BUY	84.67	89.30	63.89
Llama 8B	HOLD	2.75	1.59	29.60

Table 22: Individual decision making: Original stock: complete label distributions (%), $N = 2,074$.

Model	Label	Base	T	F
Llama 3B	COUNTER	98.11	98.54	96.32
Llama 3B	ACCEPT	1.71	1.46	2.53
Llama 3B	EXIT	0.17	0.00	1.16
Qwen 3B	COUNTER	97.09	98.59	83.93
Qwen 3B	ACCEPT	0.94	0.94	0.34
Qwen 3B	EXIT	1.97	0.47	15.72
Qwen 7B	COUNTER	98.33	98.46	98.33
Qwen 7B	ACCEPT	1.41	1.37	1.41
Qwen 7B	EXIT	0.26	0.17	0.26
Llama 8B	COUNTER	97.56	98.03	97.73
Llama 8B	ACCEPT	1.93	1.93	1.67
Llama 8B	EXIT	0.51	0.04	0.60

Table 23: Individual decision making: Original bargaining: complete label distributions (%), $N = 2,334$.

Model	Label	Base	T	F
Llama 3B	SOONER	100.00	100.00	100.00
Qwen 3B	SOONER	44.10	58.33	44.44
Qwen 3B	LATER	55.90	41.67	55.56
Qwen 7B	LATER	99.65	95.14	100.00
Qwen 7B	SOONER	0.35	4.86	0.00
Llama 8B	SOONER	100.00	100.00	100.00

Table 24: Individual decision making: Delayed rewards: complete label distributions (%), $N = 288$.

Model	Label	Base	T	F
Llama 3B	COUNTER	100.00	100.00	99.89
Llama 3B	EXIT	0.00	0.00	0.11
Qwen 3B	EXIT	8.89	0.89	28.00
Qwen 3B	COUNTER	91.11	99.11	72.00
Qwen 7B	COUNTER	100.00	100.00	100.00
Llama 8B	COUNTER	100.00	100.00	100.00

Table 25: Individual decision making: Bargaining controlled: complete label distributions (%), $N = 900$.

Model	Label	Base	T	F
Llama 3B	COUNTER	100.00	100.00	100.00
Qwen 3B	EXIT	85.19	27.35	99.72
Qwen 3B	COUNTER	14.81	72.65	0.28
Qwen 7B	COUNTER	100.00	100.00	100.00
Llama 8B	COUNTER	100.00	100.00	100.00

Table 26: Interactive decision making: Bargaining real history: complete label distributions (%), $N = 351$.

Model	Label	Base	T	F
Llama 3B	SEARCH_OTHER	45.04	36.52	46.81
Llama 3B	ACCEPT	29.43	39.01	18.79
Llama 3B	ASK_DETAILS	25.53	24.47	34.40
Qwen 3B	SEARCH_OTHER	90.43	80.50	91.49
Qwen 3B	ASK_DETAILS	8.16	16.67	7.09
Qwen 3B	ACCEPT	1.42	2.84	1.42
Qwen 7B	ASK_DETAILS	35.46	29.08	38.65
Qwen 7B	SEARCH_OTHER	37.23	42.55	35.11
Qwen 7B	ACCEPT	27.30	28.37	26.24
Llama 8B	ASK_DETAILS	72.34	86.17	60.64
Llama 8B	SEARCH_OTHER	25.18	10.64	39.01
Llama 8B	ACCEPT	2.48	3.19	0.35

Table 27: Frames: complete label distributions (%), $N = 282$.

Model	Label	Base	T	F
Llama 3B	ACCEPT	78.15	93.28	25.21
Llama 3B	ASK_MORE	21.01	5.88	74.79
Llama 3B	REJECT	0.84	0.84	0.00
Qwen 3B	ASK_MORE	86.55	81.51	91.60
Qwen 3B	ACCEPT	12.61	14.29	8.40
Qwen 3B	REJECT	0.84	4.20	0.00
Qwen 7B	ASK_MORE	64.71	64.71	68.07
Qwen 7B	ACCEPT	25.21	23.53	25.21
Qwen 7B	REJECT	10.08	11.76	6.72
Llama 8B	ACCEPT	55.46	71.43	12.61
Llama 8B	ASK_MORE	36.13	20.17	84.03
Llama 8B	REJECT	8.40	8.40	3.36

Table 28: INSPIRED: complete label distributions (%), $N = 119$.

Model	Label	Base	T	F
Llama 3B	SELL	37.21	29.75	28.16
Llama 3B	HOLD	20.90	27.06	58.61
Llama 3B	BUY	41.89	43.18	13.23
Qwen 3B	SELL	99.00	94.43	99.80
Qwen 3B	HOLD	1.00	5.57	0.20
Qwen 7B	SELL	39.90	39.60	41.29
Qwen 7B	BUY	42.49	44.08	41.99
Qwen 7B	HOLD	17.61	16.32	16.72
Llama 8B	BUY	96.12	95.62	59.60
Llama 8B	HOLD	3.88	4.38	40.40

Table 29: Individual decision making: Stock controlled: complete label distributions (%), $N = 1,005$.

Model	Label	Base	T	F
Llama 3B	COUNTER	99.82	100.00	98.55
Llama 3B	EXIT	0.18	0.00	1.23
Llama 3B	ACCEPT	0.00	0.00	0.22
Qwen 3B	COUNTER	98.46	99.96	85.04
Qwen 3B	EXIT	1.54	0.04	14.96
Qwen 7B	COUNTER	99.74	99.82	99.74
Qwen 7B	EXIT	0.26	0.18	0.26
Llama 8B	COUNTER	99.96	100.00	99.96
Llama 8B	EXIT	0.04	0.00	0.04

Table 30: Interactive decision making: Craigslist interaction below asking: complete label distributions (%), $N = 2,272$.

D.5 Programmed Bargaining Outcomes

The first action distribution above and final disengagement rates measure different parts of the same trajectories. Table 31 reports EXIT at any point within the program's maximum of three choices, retaining all 2,272 histories with initial offers below the asking price. For Qwen 3B, final EXIT rates are 48.24% for F and 0.26% for T, compared with 7.35% at Base. The final outcome therefore captures divergence that is not fully represented by the first action alone.

Model	Base	T	F
Llama 3B	0.18	0.00	1.23
Llama 8B	0.04	0.00	0.04
Qwen 3B	7.35	0.26	48.24
Qwen 7B	0.26	0.18	0.26

Table 31: Interactive decision making: Final disengagement rates (%) in programmed bargaining.

E Statistical Analysis

E.1 Design, Tests, and Uncertainty

We retrospectively analyzed the frozen per input outputs. All comparisons preserve input pairing. Behavioral and interpretation outcomes use paired permutation tests that swap conditions within whole source clusters. We enumerate all swaps with at most twenty nonzero cluster differences; otherwise we use 99,999 random swaps and the add one correction. The smallest Monte Carlo p value (0.00001) is a resolution limit, not an exact tail estimate. The tests assume condition exchangeability within the retained blocks under the null.

We report effect sizes and marginal 95% bootstrap percentile intervals with paired cluster resampling from 19,999 resamples. Both conditions are resampled together. Macro F1 is recomputed from sufficient counts over fixed task classes. Intervals are marginal and are not adjusted for multiple comparisons; accuracy and label rate changes are percentage points, while macro F1 changes use 100 times the raw F1 difference.

Circa groups all replies to the same question within each context (3,375 groups), with an additional ten-context sensitivity analysis. CommitmentBank groups source passages: label rates use 1,009 inputs in 1,006 groups, while performance uses 858 adapted-majority-reference inputs in 856 groups. FOMC groups 2,421 statements into 27 years; dependencies across adjacent years are not estimated. Original risk choices group equivalent option descriptions, machine choices retain all games for each of 26 participants, and INSPIRED/Frames group the 119/282 source dialogues.

Controlled risk summaries retain both presentation frames and both option orders within each cell defined by probability and payoff gap: six cells in discovery and twenty in replication. These are fixed design grids; resampling does not establish generalization to arbitrary economic situations.

Controlled bargaining retains all offers and option orders within fifteen source scenarios. Offer response is the EXIT change from 40% to 90% of asking price, and differences between these responses are paired across model conditions. EXIT is not a transaction-acceptance threshold.

Controlled movie/travel feedback comparisons use only common records where Base, T and F all reach turn two. We preserve costs and option orders within each authored scenario version. Only two versions are available, so we report descriptive effects and discrete tests but withhold nominal bootstrap confidence intervals. These selected cohort estimates are not unconditional causal effects.

All capability tables test accuracy differences. Capability outcomes use exact two sided McNemar tests, paired bootstrap intervals, and separate one sided degradation tests. GSM8K scoring follows the rules described in Appendix B.1. BoolQ additionally has passage block sensitivity analyses.

The frozen ARC parser does not recognize 87 reference labels among 2,000 (86 numerical option identifiers and one E). We preserve the recorded scoring with the full denominator and disclose this limitation rather than silently remapping labels. Removing these pairs scored as incorrect in both conditions leaves exact McNemar p values unchanged but changes the accuracy denominator.

Holm correction is applied at $\alpha = .05$ within the listed analysis families. The trained variant family has 24 tests, Prompt has 18, and combined guidance has 36. One-sided capability tests use separate corrections over the same respective families. Sensitivity checks have separate families and cannot replace a failed primary result. A result that is not significant does not establish equivalence or noninferiority. These analyses do not estimate variability across training seeds or effects across the population of language models.

Test choice follows the paired outcome and dependence structure (Dror et al., 2018). These are retrospective exploratory analyses of frozen outputs. No significance test is applied to the retention rate. The statistical families below concern the specified T/F comparisons and capability conditions.

E.2 Complete Results and Correction Families

Each named test group below is one Holm correction family; its size is stated explicitly. Delta is the guided minus reference difference in percentage points, or 100 times the macro F1 difference. N/G gives paired input count and the number of blocks. CI is the marginal 95% paired bootstrap interval. Both raw and adjusted two sided p values are reported; capability tables additionally report separately corrected one sided degradation p values. Bold adjusted p values meet .05. NA intervals are withheld for the two version controlled feedback probes. Model abbreviations follow the other appendix tables; CB denotes CommitmentBank. In capability tables, A:B denotes condition A minus reference B; each test group specifies whether T/F denotes DPO or prompting. Combined conditions retain the trained pole before the prompt pole.

E.3 Interpretation Tests

After Holm correction, 34 of 36 label proportion tests and 56 of 72 performance tests are significant. The Circa T macro F1 improvements are supported in all four models; the tables also distinguish label changes from changes in task accuracy.

E.3.1 Interpretation: Label Proportions

Holm correction covers 36 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Circa / UNCLEAR	T-BASE	32602/3375	-21.24	[-21.84, -20.63]	1.00e-05	**3.50e-04**
Llama 3B	Circa / UNCLEAR	F-BASE	32602/3375	+48.50	[+47.74, +49.24]	1.00e-05	**3.50e-04**
Llama 3B	Circa / UNCLEAR	F-T	32602/3375	+69.74	[+69.01, +70.45]	1.00e-05	**3.50e-04**
Llama 8B	Circa / UNCLEAR	T-BASE	32602/3375	-17.13	[-17.64, -16.62]	1.00e-05	**3.50e-04**
Llama 8B	Circa / UNCLEAR	F-BASE	32602/3375	+40.50	[+39.80, +41.21]	1.00e-05	**3.50e-04**
Llama 8B	Circa / UNCLEAR	F-T	32602/3375	+57.63	[+56.88, +58.37]	1.00e-05	**3.50e-04**
Qwen 3B	Circa / UNCLEAR	T-BASE	32602/3375	-1.44	[-1.60, -1.28]	1.00e-05	**3.50e-04**
Qwen 3B	Circa / UNCLEAR	F-BASE	32602/3375	+3.85	[+3.61, +4.10]	1.00e-05	**3.50e-04**
Qwen 3B	Circa / UNCLEAR	F-T	32602/3375	+5.29	[+5.00, +5.58]	1.00e-05	**3.50e-04**
Qwen 7B	Circa / UNCLEAR	T-BASE	32602/3375	-0.83	[-1.04, -0.61]	1.00e-05	**3.50e-04**
Qwen 7B	Circa / UNCLEAR	F-BASE	32602/3375	+2.71	[+2.47, +2.96]	1.00e-05	**3.50e-04**
Qwen 7B	Circa / UNCLEAR	F-T	32602/3375	+3.54	[+3.20, +3.87]	1.00e-05	**3.50e-04**

Table 32: Interpretation: Label Proportions: paired tests, part 1 (Holm family: 36).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	CB / FALSE	T-BASE	1009/1006	+8.62	[+6.94, +10.41]	1.00e-05	**3.50e-04**
Llama 3B	CB / FALSE	F-BASE	1009/1006	-28.34	[-31.09, -25.54]	1.00e-05	**3.50e-04**
Llama 3B	CB / FALSE	F-T	1009/1006	-36.97	[-39.96, -33.99]	1.00e-05	**3.50e-04**
Llama 8B	CB / FALSE	T-BASE	1009/1006	+18.43	[+16.09, +20.85]	1.00e-05	**3.50e-04**
Llama 8B	CB / FALSE	F-BASE	1009/1006	-11.20	[-13.18, -9.24]	1.00e-05	**3.50e-04**
Llama 8B	CB / FALSE	F-T	1009/1006	-29.63	[-32.44, -26.81]	1.00e-05	**3.50e-04**
Qwen 3B	CB / FALSE	T-BASE	1009/1006	+6.24	[+4.76, +7.82]	1.00e-05	**3.50e-04**
Qwen 3B	CB / FALSE	F-BASE	1009/1006	-11.60	[-13.58, -9.68]	1.00e-05	**3.50e-04**
Qwen 3B	CB / FALSE	F-T	1009/1006	-17.84	[-20.24, -15.56]	1.00e-05	**3.50e-04**
Qwen 7B	CB / FALSE	T-BASE	1009/1006	+1.19	[+0.50, +1.98]	0.0018	**0.0110**
Qwen 7B	CB / FALSE	F-BASE	1009/1006	-5.45	[-6.93, -4.06]	1.00e-05	**3.50e-04**
Qwen 7B	CB / FALSE	F-T	1009/1006	-6.64	[-8.22, -5.15]	1.00e-05	**3.50e-04**
Llama 3B	fomc / HAWKISH	T-BASE	2421/27	+0.50	[+0.27, +0.72]	0.0020	**0.0110**
Llama 3B	fomc / HAWKISH	F-BASE	2421/27	-0.33	[-0.61, -0.09]	0.0312	0.0625
Llama 3B	fomc / HAWKISH	F-T	2421/27	-0.83	[-1.27, -0.43]	4.88e-04	**0.0034**
Llama 8B	fomc / HAWKISH	T-BASE	2421/27	+3.68	[+2.97, +4.42]	1.00e-05	**3.50e-04**
Llama 8B	fomc / HAWKISH	F-BASE	2421/27	-1.90	[-2.74, -1.21]	2.00e-05	**3.50e-04**
Llama 8B	fomc / HAWKISH	F-T	2421/27	-5.58	[-6.61, -4.70]	1.00e-05	**3.50e-04**
Qwen 3B	fomc / HAWKISH	T-BASE	2421/27	+1.28	[+0.81, +1.82]	7.63e-06	**2.75e-04**
Qwen 3B	fomc / HAWKISH	F-BASE	2421/27	-1.07	[-1.52, -0.68]	3.05e-05	**3.50e-04**
Qwen 3B	fomc / HAWKISH	F-T	2421/27	-2.35	[-3.03, -1.74]	1.00e-05	**3.50e-04**
Qwen 7B	fomc / HAWKISH	T-BASE	2421/27	+0.12	[-0.17, +0.40]	0.5938	0.5938
Qwen 7B	fomc / HAWKISH	F-BASE	2421/27	-0.45	[-0.77, -0.21]	0.0039	**0.0156**
Qwen 7B	fomc / HAWKISH	F-T	2421/27	-0.58	[-0.93, -0.24]	0.0073	**0.0220**

Table 33: Interpretation: Label Proportions: paired tests, part 2 (Holm family: 36).

E.3.2 Interpretation: Performance

Holm correction covers 72 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Circa / accuracy	T-BASE	32602/3375	+11.98	[+11.36, +12.62]	1.00e-05	**7.20e-04**
Llama 3B	Circa / macro F1	T-BASE	32602/3375	+7.03	[+6.26, +7.80]	1.00e-05	**7.20e-04**
Llama 3B	Circa / accuracy	F-BASE	32602/3375	-34.94	[-35.63, -34.24]	1.00e-05	**7.20e-04**
Llama 3B	Circa / macro F1	F-BASE	32602/3375	-22.84	[-23.44, -22.24]	1.00e-05	**7.20e-04**
Llama 3B	Circa / accuracy	F-T	32602/3375	-46.93	[-47.70, -46.15]	1.00e-05	**7.20e-04**
Llama 3B	Circa / macro F1	F-T	32602/3375	-29.87	[-30.78, -28.96]	1.00e-05	**7.20e-04**
Llama 8B	Circa / accuracy	T-BASE	32602/3375	+14.23	[+13.76, +14.69]	1.00e-05	**7.20e-04**
Llama 8B	Circa / macro F1	T-BASE	32602/3375	+11.46	[+11.01, +11.90]	1.00e-05	**7.20e-04**
Llama 8B	Circa / accuracy	F-BASE	32602/3375	-14.41	[-14.93, -13.88]	1.00e-05	**7.20e-04**
Llama 8B	Circa / macro F1	F-BASE	32602/3375	-10.83	[-11.31, -10.35]	1.00e-05	**7.20e-04**
Llama 8B	Circa / accuracy	F-T	32602/3375	-28.63	[-29.33, -27.94]	1.00e-05	**7.20e-04**
Llama 8B	Circa / macro F1	F-T	32602/3375	-22.29	[-22.89, -21.69]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / accuracy	T-BASE	32602/3375	+1.07	[+0.93, +1.21]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / macro F1	T-BASE	32602/3375	+1.09	[+0.84, +1.35]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / accuracy	F-BASE	32602/3375	-3.48	[-3.72, -3.26]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / macro F1	F-BASE	32602/3375	-5.19	[-5.60, -4.79]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / accuracy	F-T	32602/3375	-4.55	[-4.82, -4.29]	1.00e-05	**7.20e-04**
Qwen 3B	Circa / macro F1	F-T	32602/3375	-6.28	[-6.73, -5.83]	1.00e-05	**7.20e-04**

Table 34: Interpretation: Performance: paired tests, part 1 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 7B	Circa / accuracy	T-BASE	32602/3375	+2.03	[+1.86, +2.21]	1.00e-05	**7.20e-04**
Qwen 7B	Circa / macro F1	T-BASE	32602/3375	+2.27	[+2.10, +2.45]	1.00e-05	**7.20e-04**
Qwen 7B	Circa / accuracy	F-BASE	32602/3375	-3.32	[-3.54, -3.11]	1.00e-05	**7.20e-04**
Qwen 7B	Circa / macro F1	F-BASE	32602/3375	-3.17	[-3.37, -2.96]	1.00e-05	**7.20e-04**
Qwen 7B	Circa / accuracy	F-T	32602/3375	-5.35	[-5.64, -5.07]	1.00e-05	**7.20e-04**
Qwen 7B	Circa / macro F1	F-T	32602/3375	-5.44	[-5.71, -5.18]	1.00e-05	**7.20e-04**
Llama 3B	CB / accuracy	T-BASE	858/856	-2.45	[-4.43, -0.47]	0.0194	0.2327
Llama 3B	CB / macro F1	T-BASE	858/856	-4.76	[-6.53, -3.06]	1.00e-05	**7.20e-04**
Llama 3B	CB / accuracy	F-BASE	858/856	+6.99	[+3.61, +10.37]	1.00e-04	**0.0022**
Llama 3B	CB / macro F1	F-BASE	858/856	+9.03	[+6.18, +11.95]	1.00e-05	**7.20e-04**
Llama 3B	CB / accuracy	F-T	858/856	+9.44	[+5.59, +13.30]	1.00e-05	**7.20e-04**
Llama 3B	CB / macro F1	F-T	858/856	+13.79	[+10.72, +16.92]	1.00e-05	**7.20e-04**
Llama 8B	CB / accuracy	T-BASE	858/856	+22.26	[+19.16, +25.32]	1.00e-05	**7.20e-04**
Llama 8B	CB / macro F1	T-BASE	858/856	+19.34	[+16.50, +22.18]	1.00e-05	**7.20e-04**
Llama 8B	CB / accuracy	F-BASE	858/856	-16.43	[-19.07, -13.82]	1.00e-05	**7.20e-04**
Llama 8B	CB / macro F1	F-BASE	858/856	-20.52	[-23.19, -17.76]	1.00e-05	**7.20e-04**
Llama 8B	CB / accuracy	F-T	858/856	-38.69	[-42.31, -35.08]	1.00e-05	**7.20e-04**
Llama 8B	CB / macro F1	F-T	858/856	-39.86	[-42.93, -36.77]	1.00e-05	**7.20e-04**

Table 35: Interpretation: Performance: paired tests, part 2 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 3B	CB / accuracy	T-BASE	858/856	+4.78	[+3.38, +6.29]	1.00e-05	**7.20e-04**
Qwen 3B	CB / macro F1	T-BASE	858/856	+2.95	[+1.95, +4.01]	1.00e-05	**7.20e-04**
Qwen 3B	CB / accuracy	F-BASE	858/856	-10.02	[-12.22, -7.92]	1.00e-05	**7.20e-04**
Qwen 3B	CB / macro F1	F-BASE	858/856	-8.20	[-10.06, -6.39]	1.00e-05	**7.20e-04**
Qwen 3B	CB / accuracy	F-T	858/856	-14.80	[-17.37, -12.33]	1.00e-05	**7.20e-04**
Qwen 3B	CB / macro F1	F-T	858/856	-11.16	[-13.22, -9.16]	1.00e-05	**7.20e-04**
Qwen 7B	CB / accuracy	T-BASE	858/856	+0.93	[-0.12, +1.98]	0.1142	0.8690
Qwen 7B	CB / macro F1	T-BASE	858/856	+0.71	[-0.20, +1.61]	0.1086	0.8690
Qwen 7B	CB / accuracy	F-BASE	858/856	-5.59	[-7.34, -3.95]	1.00e-05	**7.20e-04**
Qwen 7B	CB / macro F1	F-BASE	858/856	-4.70	[-6.22, -3.23]	1.00e-05	**7.20e-04**
Qwen 7B	CB / accuracy	F-T	858/856	-6.53	[-8.51, -4.55]	1.00e-05	**7.20e-04**
Qwen 7B	CB / macro F1	F-T	858/856	-5.41	[-7.11, -3.74]	1.00e-05	**7.20e-04**
Llama 3B	fomc / accuracy	T-BASE	2421/27	-0.74	[-2.15, +0.60]	0.3300	0.9900
Llama 3B	fomc / macro F1	T-BASE	2421/27	+1.05	[-0.19, +2.21]	0.1123	0.8690
Llama 3B	fomc / accuracy	F-BASE	2421/27	+0.95	[-0.40, +2.20]	0.1956	0.8690
Llama 3B	fomc / macro F1	F-BASE	2421/27	-1.21	[-2.48, -0.05]	0.0755	0.7549
Llama 3B	fomc / accuracy	F-T	2421/27	+1.69	[-0.09, +3.43]	0.0871	0.7840
Llama 3B	fomc / macro F1	F-T	2421/27	-2.26	[-3.98, -0.56]	0.0209	0.2327

Table 36: Interpretation: Performance: paired tests, part 3 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 8B	fomc / accuracy	T-BASE	2421/27	+1.98	[+0.82, +3.11]	0.0054	0.0763
Llama 8B	fomc / macro F1	T-BASE	2421/27	+5.27	[+3.78, +6.69]	1.00e-05	**7.20e-04**
Llama 8B	fomc / accuracy	F-BASE	2421/27	-2.85	[-3.85, -1.95]	1.00e-05	**7.20e-04**
Llama 8B	fomc / macro F1	F-BASE	2421/27	-7.28	[-8.95, -5.69]	1.00e-05	**7.20e-04**
Llama 8B	fomc / accuracy	F-T	2421/27	-4.83	[-6.26, -3.40]	1.00e-05	**7.20e-04**
Llama 8B	fomc / macro F1	F-T	2421/27	-12.54	[-14.40, -10.66]	1.00e-05	**7.20e-04**
Qwen 3B	fomc / accuracy	T-BASE	2421/27	+0.70	[+0.22, +1.17]	0.0165	0.2140
Qwen 3B	fomc / macro F1	T-BASE	2421/27	+1.42	[+0.62, +2.21]	0.0033	0.0523
Qwen 3B	fomc / accuracy	F-BASE	2421/27	-0.50	[-1.02, +0.04]	0.1278	0.8690
Qwen 3B	fomc / macro F1	F-BASE	2421/27	-1.54	[-2.33, -0.77]	0.0012	**0.0248**
Qwen 3B	fomc / accuracy	F-T	2421/27	-1.20	[-1.81, -0.56]	0.0039	0.0583
Qwen 3B	fomc / macro F1	F-T	2421/27	-2.96	[-4.05, -1.84]	7.00e-05	**0.0016**
Qwen 7B	fomc / accuracy	T-BASE	2421/27	+0.54	[-0.41, +1.48]	0.3483	0.9900
Qwen 7B	fomc / macro F1	T-BASE	2421/27	+0.35	[-0.59, +1.32]	0.5324	0.9900
Qwen 7B	fomc / accuracy	F-BASE	2421/27	-1.45	[-2.12, -0.71]	0.0014	**0.0264**
Qwen 7B	fomc / macro F1	F-BASE	2421/27	-1.44	[-2.10, -0.77]	3.60e-04	**0.0076**
Qwen 7B	fomc / accuracy	F-T	2421/27	-1.98	[-2.96, -0.97]	0.0023	**0.0398**
Qwen 7B	fomc / macro F1	F-T	2421/27	-1.79	[-2.68, -0.89]	0.0018	**0.0319**

Table 37: Interpretation: Performance: paired tests, part 4 (Holm family: 72).

E.3.3 Circa Context Sensitivity: Labels

Holm correction covers 12 comparisons.

The context-level sensitivity analysis confirms significant differences in both prediction distributions and task performance after Holm correction, supporting the robustness of the observed personality effects on Circa.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Circa / UNCLEAR	T-BASE	32602/10	-21.24	[-25.12, -18.10]	0.0020	**0.0234**
Llama 3B	Circa / UNCLEAR	F-BASE	32602/10	+48.50	[+44.10, +53.39]	0.0020	**0.0234**
Llama 3B	Circa / UNCLEAR	F-T	32602/10	+69.74	[+64.66, +74.73]	0.0020	**0.0234**
Llama 8B	Circa / UNCLEAR	T-BASE	32602/10	-17.13	[-18.62, -15.37]	0.0020	**0.0234**
Llama 8B	Circa / UNCLEAR	F-BASE	32602/10	+40.50	[+35.99, +44.72]	0.0020	**0.0234**
Llama 8B	Circa / UNCLEAR	F-T	32602/10	+57.63	[+52.91, +62.16]	0.0020	**0.0234**
Qwen 3B	Circa / UNCLEAR	T-BASE	32602/10	-1.44	[-1.78, -1.11]	0.0020	**0.0234**
Qwen 3B	Circa / UNCLEAR	F-BASE	32602/10	+3.85	[+2.88, +4.93]	0.0020	**0.0234**
Qwen 3B	Circa / UNCLEAR	F-T	32602/10	+5.29	[+4.31, +6.30]	0.0020	**0.0234**
Qwen 7B	Circa / UNCLEAR	T-BASE	32602/10	-0.83	[-1.30, -0.35]	0.0156	**0.0234**
Qwen 7B	Circa / UNCLEAR	F-BASE	32602/10	+2.71	[+2.13, +3.33]	0.0020	**0.0234**
Qwen 7B	Circa / UNCLEAR	F-T	32602/10	+3.54	[+2.52, +4.57]	0.0020	**0.0234**

Table 38: Circa Context Sensitivity: Labels: paired tests, part 1 (Holm family: 12).

E.3.4 Circa Context Sensitivity: Performance

Holm correction covers 24 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Circa / accuracy	T-BASE	32602/10	+11.98	[+9.70, +14.99]	0.0020	**0.0469**
Llama 3B	Circa / macro F1	T-BASE	32602/10	+7.03	[+5.26, +9.11]	0.0020	**0.0469**
Llama 3B	Circa / accuracy	F-BASE	32602/10	-34.94	[-39.17, -31.20]	0.0020	**0.0469**
Llama 3B	Circa / macro F1	F-BASE	32602/10	-22.84	[-25.59, -19.97]	0.0020	**0.0469**
Llama 3B	Circa / accuracy	F-T	32602/10	-46.93	[-50.76, -43.00]	0.0020	**0.0469**
Llama 3B	Circa / macro F1	F-T	32602/10	-29.87	[-31.97, -27.64]	0.0020	**0.0469**
Llama 8B	Circa / accuracy	T-BASE	32602/10	+14.23	[+12.51, +16.02]	0.0020	**0.0469**
Llama 8B	Circa / macro F1	T-BASE	32602/10	+11.46	[+10.37, +12.48]	0.0020	**0.0469**
Llama 8B	Circa / accuracy	F-BASE	32602/10	-14.41	[-16.70, -12.25]	0.0020	**0.0469**
Llama 8B	Circa / macro F1	F-BASE	32602/10	-10.83	[-12.47, -9.14]	0.0020	**0.0469**
Llama 8B	Circa / accuracy	F-T	32602/10	-28.63	[-31.95, -25.11]	0.0020	**0.0469**
Llama 8B	Circa / macro F1	F-T	32602/10	-22.29	[-24.26, -20.28]	0.0020	**0.0469**
Qwen 3B	Circa / accuracy	T-BASE	32602/10	+1.07	[+0.71, +1.45]	0.0020	**0.0469**
Qwen 3B	Circa / macro F1	T-BASE	32602/10	+1.09	[+0.60, +1.71]	0.0059	**0.0469**
Qwen 3B	Circa / accuracy	F-BASE	32602/10	-3.48	[-4.45, -2.56]	0.0020	**0.0469**
Qwen 3B	Circa / macro F1	F-BASE	32602/10	-5.19	[-6.10, -4.10]	0.0020	**0.0469**
Qwen 3B	Circa / accuracy	F-T	32602/10	-4.55	[-5.47, -3.61]	0.0020	**0.0469**
Qwen 3B	Circa / macro F1	F-T	32602/10	-6.28	[-7.13, -5.22]	0.0020	**0.0469**

Table 39: Circa Context Sensitivity: Performance: paired tests, part 1 (Holm family: 24).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 7B	Circa / accuracy	T-BASE	32602/10	+2.03	[+1.67, +2.46]	0.0020	**0.0469**
Qwen 7B	Circa / macro F1	T-BASE	32602/10	+2.27	[+1.96, +2.62]	0.0020	**0.0469**
Qwen 7B	Circa / accuracy	F-BASE	32602/10	-3.32	[-3.79, -2.84]	0.0020	**0.0469**
Qwen 7B	Circa / macro F1	F-BASE	32602/10	-3.17	[-3.49, -2.83]	0.0020	**0.0469**
Qwen 7B	Circa / accuracy	F-T	32602/10	-5.35	[-6.23, -4.54]	0.0020	**0.0469**
Qwen 7B	Circa / macro F1	F-T	32602/10	-5.44	[-6.08, -4.81]	0.0020	**0.0469**

Table 40: Circa Context Sensitivity: Performance: paired tests, part 2 (Holm family: 24).

E.4 Decision and Interaction Tests

T/F machine choice differences pass correction in all four models. The larger risk summary replication supports five of twelve comparisons, while the small discovery probe supports none. The complete results separate these findings from decision and interaction comparisons that do not pass correction.

E.4.1 Main Decision and Interaction Outcomes

Holm correction covers 48 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	risk / high variance	T-BASE	1764/1764	-3.80	[-5.90, -1.70]	4.60e-04	**0.0092**
Llama 3B	risk / high variance	F-BASE	1764/1764	+15.25	[+12.64, +17.91]	1.00e-05	**4.60e-04**
Llama 3B	risk / high variance	F-T	1764/1764	+19.05	[+15.82, +22.34]	1.00e-05	**4.60e-04**
Llama 3B	Machine / fewer observations	T-BASE	4160/26	+9.86	[+8.85, +11.01]	1.00e-05	**4.60e-04**
Llama 3B	Machine / fewer observations	F-BASE	4160/26	-43.03	[-44.52, -41.44]	1.00e-05	**4.60e-04**
Llama 3B	Machine / fewer observations	F-T	4160/26	-52.88	[-54.81, -50.91]	1.00e-05	**4.60e-04**
Llama 3B	INSPIRED / ASK	T-BASE	119/119	-15.13	[-21.85, -9.24]	7.63e-06	**3.59e-04**
Llama 3B	INSPIRED / ASK	F-BASE	119/119	+53.78	[+44.54, +63.03]	1.00e-05	**4.60e-04**
Llama 3B	INSPIRED / ASK	F-T	119/119	+68.91	[+60.50, +77.31]	1.00e-05	**4.60e-04**
Llama 3B	Frames / ACCEPT	T-BASE	282/282	+9.57	[+6.38, +13.12]	1.00e-05	**4.60e-04**
Llama 3B	Frames / ACCEPT	F-BASE	282/282	-10.64	[-14.54, -7.09]	1.00e-05	**4.60e-04**
Llama 3B	Frames / ACCEPT	F-T	282/282	-20.21	[-24.82, -15.60]	1.00e-05	**4.60e-04**
Llama 8B	risk / high variance	T-BASE	1764/1764	-1.30	[-2.55, -0.06]	0.0527	0.7375
Llama 8B	risk / high variance	F-BASE	1764/1764	+21.26	[+18.99, +23.58]	1.00e-05	**4.60e-04**
Llama 8B	risk / high variance	F-T	1764/1764	+22.56	[+19.95, +25.17]	1.00e-05	**4.60e-04**
Llama 8B	Machine / fewer observations	T-BASE	4160/26	+16.92	[+15.43, +18.51]	1.00e-05	**4.60e-04**
Llama 8B	Machine / fewer observations	F-BASE	4160/26	-17.55	[-18.89, -16.15]	1.00e-05	**4.60e-04**
Llama 8B	Machine / fewer observations	F-T	4160/26	-34.47	[-35.87, -33.12]	1.00e-05	**4.60e-04**

Table 41: Main Decision and Interaction Outcomes: paired tests, part 1 (Holm family: 48).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 8B	INSPIRED / ASK	T-BASE	119/119	-15.97	[-22.69, -10.08]	3.81e-06	**1.83e-04**
Llama 8B	INSPIRED / ASK	F-BASE	119/119	+47.90	[+38.66, +57.14]	1.00e-05	**4.60e-04**
Llama 8B	INSPIRED / ASK	F-T	119/119	+63.87	[+55.46, +72.27]	1.00e-05	**4.60e-04**
Llama 8B	Frames / ACCEPT	T-BASE	282/282	+0.71	[+0.00, +1.77]	0.5000	1.000
Llama 8B	Frames / ACCEPT	F-BASE	282/282	-2.13	[-3.90, -0.71]	0.0312	0.5000
Llama 8B	Frames / ACCEPT	F-T	282/282	-2.84	[-4.96, -1.06]	0.0078	0.1406
Qwen 3B	risk / high variance	T-BASE	1764/1764	-3.40	[-4.71, -2.15]	1.00e-05	**4.60e-04**
Qwen 3B	risk / high variance	F-BASE	1764/1764	+4.31	[+3.06, +5.56]	1.00e-05	**4.60e-04**
Qwen 3B	risk / high variance	F-T	1764/1764	+7.71	[+6.01, +9.47]	1.00e-05	**4.60e-04**
Qwen 3B	Machine / fewer observations	T-BASE	4160/26	+8.46	[+7.40, +9.62]	1.00e-05	**4.60e-04**
Qwen 3B	Machine / fewer observations	F-BASE	4160/26	-9.04	[-10.43, -7.64]	1.00e-05	**4.60e-04**
Qwen 3B	Machine / fewer observations	F-T	4160/26	-17.50	[-18.70, -16.25]	1.00e-05	**4.60e-04**
Qwen 3B	INSPIRED / ASK	T-BASE	119/119	-5.04	[-10.08, -0.84]	0.0703	0.9141
Qwen 3B	INSPIRED / ASK	F-BASE	119/119	+5.04	[+1.68, +9.24]	0.0312	0.5000
Qwen 3B	INSPIRED / ASK	F-T	119/119	+10.08	[+4.20, +15.97]	0.0018	**0.0348**
Qwen 3B	Frames / ACCEPT	T-BASE	282/282	+1.42	[+0.35, +2.84]	0.1250	1.000
Qwen 3B	Frames / ACCEPT	F-BASE	282/282	+0.00	[-1.06, +1.06]	1.000	1.000
Qwen 3B	Frames / ACCEPT	F-T	282/282	-1.42	[-2.84, -0.35]	0.1250	1.000

Table 42: Main Decision and Interaction Outcomes: paired tests, part 2 (Holm family: 48).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 7B	risk / high variance	T-BASE	1764/1764	+1.19	[+0.34, +2.04]	0.0091	0.1542
Qwen 7B	risk / high variance	F-BASE	1764/1764	+0.28	[-0.45, +1.02]	0.5537	1.000
Qwen 7B	risk / high variance	F-T	1764/1764	-0.91	[-2.04, +0.23]	0.1446	1.000
Qwen 7B	Machine / fewer observations	T-BASE	4160/26	+3.41	[+2.69, +4.13]	1.00e-05	**4.60e-04**
Qwen 7B	Machine / fewer observations	F-BASE	4160/26	-3.46	[-4.04, -2.88]	1.00e-05	**4.60e-04**
Qwen 7B	Machine / fewer observations	F-T	4160/26	-6.88	[-7.79, -5.87]	1.00e-05	**4.60e-04**
Qwen 7B	INSPIRED / ASK	T-BASE	119/119	+0.00	[-3.36, +3.36]	1.000	1.000
Qwen 7B	INSPIRED / ASK	F-BASE	119/119	+3.36	[+0.00, +7.56]	0.2188	1.000
Qwen 7B	INSPIRED / ASK	F-T	119/119	+3.36	[-1.68, +8.40]	0.3438	1.000
Qwen 7B	Frames / ACCEPT	T-BASE	282/282	+1.06	[-0.71, +2.84]	0.4531	1.000
Qwen 7B	Frames / ACCEPT	F-BASE	282/282	-1.06	[-2.84, +0.71]	0.4531	1.000
Qwen 7B	Frames / ACCEPT	F-T	282/282	-2.13	[-4.61, +0.35]	0.1796	1.000

Table 43: Main Decision and Interaction Outcomes: paired tests, part 3 (Holm family: 48).

E.4.2 Expected Payoff Probe: Discovery

Holm correction covers 12 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Risk summary / switch rate	T-BASE	24/6	+25.00	[-4.17, +58.33]	0.3750	1.000
Llama 3B	Risk summary / switch rate	F-BASE	24/6	-16.67	[-25.00, -8.33]	0.1250	1.000
Llama 3B	Risk summary / switch rate	F-T	24/6	-41.67	[-66.67, -20.83]	0.0625	0.6250
Llama 8B	Risk summary / switch rate	T-BASE	24/6	-8.33	[-16.67, +0.00]	0.5000	1.000
Llama 8B	Risk summary / switch rate	F-BASE	24/6	-45.83	[-50.00, -37.50]	0.0312	0.3750
Llama 8B	Risk summary / switch rate	F-T	24/6	-37.50	[-45.83, -29.17]	0.0312	0.3750
Qwen 3B	Risk summary / switch rate	T-BASE	24/6	-4.17	[-12.50, +0.00]	1.000	1.000
Qwen 3B	Risk summary / switch rate	F-BASE	24/6	-16.67	[-25.00, -8.33]	0.1250	1.000
Qwen 3B	Risk summary / switch rate	F-T	24/6	-12.50	[-20.83, -4.17]	0.2500	1.000
Qwen 7B	Risk summary / switch rate	T-BASE	24/6	+12.50	[+4.17, +20.83]	0.2500	1.000
Qwen 7B	Risk summary / switch rate	F-BASE	24/6	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	Risk summary / switch rate	F-T	24/6	-12.50	[-20.83, -4.17]	0.2500	1.000

Table 44: Expected Payoff Probe: Discovery: paired tests, part 1 (Holm family: 12).

E.4.3 Expected Payoff Probe: Replication

Holm correction covers 12 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Risk summary / switch rate	T-BASE	80/20	+0.00	[-10.00, +10.00]	1.000	1.000
Llama 3B	Risk summary / switch rate	F-BASE	80/20	-13.75	[-20.00, -6.25]	0.0074	0.0517
Llama 3B	Risk summary / switch rate	F-T	80/20	-13.75	[-21.25, -5.00]	0.0114	0.0681
Llama 8B	Risk summary / switch rate	T-BASE	80/20	-17.50	[-22.50, -12.50]	1.22e-04	**0.0012**
Llama 8B	Risk summary / switch rate	F-BASE	80/20	-50.00	[-50.00, -50.00]	1.91e-06	**2.29e-05**
Llama 8B	Risk summary / switch rate	F-T	80/20	-32.50	[-37.50, -27.50]	1.91e-06	**2.29e-05**
Qwen 3B	Risk summary / switch rate	T-BASE	80/20	+2.50	[-7.50, +13.75]	0.8203	1.000
Qwen 3B	Risk summary / switch rate	F-BASE	80/20	-22.50	[-32.50, -12.50]	0.0011	**0.0099**
Qwen 3B	Risk summary / switch rate	F-T	80/20	-25.00	[-37.50, -12.50]	0.0034	**0.0273**
Qwen 7B	Risk summary / switch rate	T-BASE	80/20	+0.00	[-8.75, +8.75]	1.000	1.000
Qwen 7B	Risk summary / switch rate	F-BASE	80/20	-2.50	[-10.00, +5.00]	0.7539	1.000
Qwen 7B	Risk summary / switch rate	F-T	80/20	-2.50	[-10.00, +5.00]	0.7656	1.000

Table 45: Expected Payoff Probe: Replication: paired tests, part 1 (Holm family: 12).

E.4.4 Controlled Bargaining: Exit Rates

Holm correction covers 12 comparisons.

Significant differences in bargaining exit rates are observed for Llama 8B and Qwen 3B, whereas Llama 3B and Qwen 7B show no exit-rate differences in these controlled probes.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	bargaining / EXIT	T-BASE	315/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT	F-BASE	315/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT	F-T	315/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 8B	bargaining / EXIT	T-BASE	315/15	-24.44	[-31.11, -17.78]	1.22e-04	**0.0011**
Llama 8B	bargaining / EXIT	F-BASE	315/15	+1.90	[-1.27, +4.76]	0.3350	1.000
Llama 8B	bargaining / EXIT	F-T	315/15	+26.35	[+19.68, +32.70]	6.10e-05	**7.32e-04**
Qwen 3B	bargaining / EXIT	T-BASE	315/15	-4.13	[-5.40, -2.86]	4.88e-04	**0.0039**
Qwen 3B	bargaining / EXIT	F-BASE	315/15	+14.60	[+11.75, +17.46]	6.10e-05	**7.32e-04**
Qwen 3B	bargaining / EXIT	F-T	315/15	+18.73	[+15.87, +21.59]	6.10e-05	**7.32e-04**
Qwen 7B	bargaining / EXIT	T-BASE	315/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT	F-BASE	315/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT	F-T	315/15	+0.00	[+0.00, +0.00]	1.000	1.000

Table 46: Controlled Bargaining: Exit Rates: paired tests, part 1 (Holm family: 12).

E.4.5 Controlled Bargaining: Offer Response

Holm correction covers 24 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	bargaining / EXIT response	BASE:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT response	T:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT response	F:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT response difference	T-BASE	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT response difference	F-BASE	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 3B	bargaining / EXIT response difference	F-T	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Llama 8B	bargaining / EXIT response	BASE:high-low	90/15	-80.00	[-93.33, -64.44]	1.22e-04	**0.0027**
Llama 8B	bargaining / EXIT response	T:high-low	90/15	-22.22	[-28.89, -13.33]	0.0020	**0.0352**
Llama 8B	bargaining / EXIT response	F:high-low	90/15	-75.56	[-88.89, -60.00]	6.10e-05	**0.0015**
Llama 8B	bargaining / EXIT response difference	T-BASE	90/15	+57.78	[+46.67, +68.89]	1.22e-04	**0.0027**
Llama 8B	bargaining / EXIT response difference	F-BASE	90/15	+4.44	[-4.44, +15.56]	0.7500	1.000
Llama 8B	bargaining / EXIT response difference	F-T	90/15	-53.33	[-62.22, -44.44]	6.10e-05	**0.0015**
Qwen 3B	bargaining / EXIT response	BASE:high-low	90/15	-22.22	[-28.89, -13.33]	0.0020	**0.0352**
Qwen 3B	bargaining / EXIT response	T:high-low	90/15	-2.22	[-6.67, +0.00]	1.000	1.000
Qwen 3B	bargaining / EXIT response	F:high-low	90/15	-33.33	[-40.00, -26.67]	1.22e-04	**0.0027**
Qwen 3B	bargaining / EXIT response difference	T-BASE	90/15	+20.00	[+11.11, +28.89]	0.0039	0.0625
Qwen 3B	bargaining / EXIT response difference	F-BASE	90/15	-11.11	[-20.00, +0.00]	0.1250	1.000
Qwen 3B	bargaining / EXIT response difference	F-T	90/15	-31.11	[-33.33, -26.67]	1.22e-04	**0.0027**

Table 47: Controlled Bargaining: Offer Response: paired tests, part 1 (Holm family: 24).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 7B	bargaining / EXIT response	BASE:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT response	T:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT response	F:high-low	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT response difference	T-BASE	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT response difference	F-BASE	90/15	+0.00	[+0.00, +0.00]	1.000	1.000
Qwen 7B	bargaining / EXIT response difference	F-T	90/15	+0.00	[+0.00, +0.00]	1.000	1.000

Table 48: Controlled Bargaining: Offer Response: paired tests, part 2 (Holm family: 24).

E.5 Controlled Feedback Tests

None of the 72 feedback comparisons passes Holm correction. These probes contain only two authored scenario versions and a selected second turn cohort, so their observed differences are reported descriptively rather than treated as established effects.

E.5.1 Controlled Feedback Probes

Holm correction covers 72 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	Movie / fits:ACCEPT	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / fits:ACCEPT	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / fits:ACCEPT	F-T	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / mismatch:OTHER	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / mismatch:OTHER	F-T	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / unresolved:ASK	T-BASE	12/2	-8.33	NA	1.000	1.000
Llama 3B	Movie / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Movie / unresolved:ASK	F-T	12/2	+8.33	NA	1.000	1.000
Llama 3B	Travel / fits:ACCEPT	T-BASE	12/2	+8.33	NA	1.000	1.000
Llama 3B	Travel / fits:ACCEPT	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / fits:ACCEPT	F-T	12/2	-8.33	NA	1.000	1.000
Llama 3B	Travel / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / mismatch:OTHER	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / mismatch:OTHER	F-T	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / unresolved:ASK	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 3B	Travel / unresolved:ASK	F-T	12/2	+0.00	NA	1.000	1.000

Table 49: Controlled Feedback Probes: paired tests, part 1 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 8B	Movie / fits:ACCEPT	T-BASE	10/2	+60.00	NA	0.5000	1.000
Llama 8B	Movie / fits:ACCEPT	F-BASE	10/2	-20.00	NA	1.000	1.000
Llama 8B	Movie / fits:ACCEPT	F-T	10/2	-80.00	NA	0.5000	1.000
Llama 8B	Movie / mismatch:OTHER	T-BASE	10/2	-50.00	NA	0.5000	1.000
Llama 8B	Movie / mismatch:OTHER	F-BASE	10/2	+20.00	NA	0.5000	1.000
Llama 8B	Movie / mismatch:OTHER	F-T	10/2	+70.00	NA	0.5000	1.000
Llama 8B	Movie / unresolved:ASK	T-BASE	10/2	+0.00	NA	1.000	1.000
Llama 8B	Movie / unresolved:ASK	F-BASE	10/2	-20.00	NA	1.000	1.000
Llama 8B	Movie / unresolved:ASK	F-T	10/2	-20.00	NA	1.000	1.000
Llama 8B	Travel / fits:ACCEPT	T-BASE	12/2	+25.00	NA	0.5000	1.000
Llama 8B	Travel / fits:ACCEPT	F-BASE	12/2	-25.00	NA	0.5000	1.000
Llama 8B	Travel / fits:ACCEPT	F-T	12/2	-50.00	NA	0.5000	1.000
Llama 8B	Travel / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 8B	Travel / mismatch:OTHER	F-BASE	12/2	+41.67	NA	0.5000	1.000
Llama 8B	Travel / mismatch:OTHER	F-T	12/2	+41.67	NA	0.5000	1.000
Llama 8B	Travel / unresolved:ASK	T-BASE	12/2	+0.00	NA	1.000	1.000
Llama 8B	Travel / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Llama 8B	Travel / unresolved:ASK	F-T	12/2	+0.00	NA	1.000	1.000

Table 50: Controlled Feedback Probes: paired tests, part 2 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 3B	Movie / fits:ACCEPT	T-BASE	12/2	-16.67	NA	0.5000	1.000
Qwen 3B	Movie / fits:ACCEPT	F-BASE	12/2	-16.67	NA	1.000	1.000
Qwen 3B	Movie / fits:ACCEPT	F-T	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / mismatch:OTHER	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / mismatch:OTHER	F-T	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / unresolved:ASK	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Movie / unresolved:ASK	F-T	12/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / fits:ACCEPT	T-BASE	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / fits:ACCEPT	F-BASE	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / fits:ACCEPT	F-T	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / mismatch:OTHER	T-BASE	8/2	-12.50	NA	1.000	1.000
Qwen 3B	Travel / mismatch:OTHER	F-BASE	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / mismatch:OTHER	F-T	8/2	+12.50	NA	1.000	1.000
Qwen 3B	Travel / unresolved:ASK	T-BASE	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / unresolved:ASK	F-BASE	8/2	+0.00	NA	1.000	1.000
Qwen 3B	Travel / unresolved:ASK	F-T	8/2	+0.00	NA	1.000	1.000

Table 51: Controlled Feedback Probes: paired tests, part 3 (Holm family: 72).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Qwen 7B	Movie / fits:ACCEPT	T-BASE	12/2	-16.67	NA	1.000	1.000
Qwen 7B	Movie / fits:ACCEPT	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / fits:ACCEPT	F-T	12/2	+16.67	NA	1.000	1.000
Qwen 7B	Movie / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / mismatch:OTHER	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / mismatch:OTHER	F-T	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / unresolved:ASK	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Movie / unresolved:ASK	F-T	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Travel / fits:ACCEPT	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Travel / fits:ACCEPT	F-BASE	12/2	-8.33	NA	1.000	1.000
Qwen 7B	Travel / fits:ACCEPT	F-T	12/2	-8.33	NA	1.000	1.000
Qwen 7B	Travel / mismatch:OTHER	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Travel / mismatch:OTHER	F-BASE	12/2	-8.33	NA	1.000	1.000
Qwen 7B	Travel / mismatch:OTHER	F-T	12/2	-8.33	NA	1.000	1.000
Qwen 7B	Travel / unresolved:ASK	T-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Travel / unresolved:ASK	F-BASE	12/2	+0.00	NA	1.000	1.000
Qwen 7B	Travel / unresolved:ASK	F-T	12/2	+0.00	NA	1.000	1.000

Table 52: Controlled Feedback Probes: paired tests, part 4 (Holm family: 72).

E.6 Capability Tests

E.6.1 Prompt Capabilities

Holm correction covers 18 comparisons.

Personality guidance produces task-dependent capability changes. Prompting generally reduces ARC-Easy accuracy, while combined guidance often improves GSM8K performance but can negatively affect BoolQ accuracy.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}	$p_{\text{degr.}}$
Llama 3B	ARC-Easy	T:Base P	2000/2000	-3.70	[-4.90, -2.50]	2.56e-09	**3.84e-08**	**1.92e-08**
Llama 3B	BoolQ	T:Base P	2000/2000	-0.80	[-1.95, +0.35]	0.1915	1.000	0.7660
Llama 3B	ARC-Easy	F:Base P	2000/2000	-5.65	[-6.95, -4.35]	6.02e-18	**1.08e-16**	**5.42e-17**
Llama 3B	BoolQ	F:Base P	2000/2000	+0.20	[-0.95, +1.35]	0.7956	1.000	1.000
Qwen 3B	GSM8K	T:Base P	1300/1300	-0.31	[-1.85, +1.23]	0.7750	1.000	1.000
Qwen 3B	ARC-Easy	T:Base P	2000/2000	-1.95	[-2.80, -1.05]	2.18e-05	**3.05e-04**	**1.53e-04**
Qwen 3B	BoolQ	T:Base P	2000/2000	-0.35	[-1.15, +0.45]	0.4704	1.000	1.000
Qwen 3B	GSM8K	F:Base P	1300/1300	-0.23	[-1.77, +1.31]	0.8496	1.000	1.000
Qwen 3B	ARC-Easy	F:Base P	2000/2000	-2.00	[-2.95, -1.10]	2.97e-05	**3.86e-04**	**1.93e-04**
Qwen 3B	BoolQ	F:Base P	2000/2000	-0.25	[-1.10, +0.60]	0.6488	1.000	1.000
Qwen 7B	GSM8K	T:Base P	1300/1300	-0.23	[-1.69, +1.23]	0.8392	1.000	1.000
Qwen 7B	ARC-Easy	T:Base P	2000/2000	-0.85	[-1.55, -0.20]	0.0186	0.2049	0.1024
Qwen 7B	BoolQ	T:Base P	2000/2000	-4.05	[-5.25, -2.85]	4.69e-11	**7.97e-10**	**3.99e-10**
Qwen 7B	GSM8K	F:Base P	1300/1300	-0.23	[-1.77, +1.31]	0.8454	1.000	1.000
Qwen 7B	ARC-Easy	F:Base P	2000/2000	-1.35	[-2.10, -0.65]	3.55e-04	**0.0043**	**0.0021**
Qwen 7B	BoolQ	F:Base P	2000/2000	-3.85	[-5.00, -2.70]	9.89e-11	**1.58e-09**	**7.91e-10**
Llama 3B	GSM8K	Prompt-T-BASE	1300/1300	-1.54	[-3.62, +0.54]	0.1657	1.000	0.7455
Llama 3B	GSM8K	Prompt-F-BASE	1300/1300	-1.85	[-3.92, +0.23]	0.0932	0.9319	0.4660

Table 53: Prompt Capabilities: paired tests, part 1 (Holm family: 18).

E.6.2 Prompt BoolQ Passage Sensitivity

Holm correction covers 6 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	BoolQ	T:Base P	2000/1897	-0.80	[-1.95, +0.35]	0.1975	0.7901
Llama 3B	BoolQ	F:Base P	2000/1897	+0.20	[-0.95, +1.36]	0.7993	1.000
Qwen 3B	BoolQ	T:Base P	2000/1897	-0.35	[-1.16, +0.45]	0.4759	1.000
Qwen 3B	BoolQ	F:Base P	2000/1897	-0.25	[-1.11, +0.60]	0.6483	1.000
Qwen 7B	BoolQ	T:Base P	2000/1897	-4.05	[-5.27, -2.83]	1.00e-05	**6.00e-05**
Qwen 7B	BoolQ	F:Base P	2000/1897	-3.85	[-5.06, -2.67]	1.00e-05	**6.00e-05**

Table 54: Prompt BoolQ Passage Sensitivity: paired tests, part 1 (Holm family: 6).

E.6.3 Combined Capabilities

Holm correction covers 36 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}	$p_{\text{degr.}}$
Llama 3B	GSM8K	T+T:T	1300/1300	+19.92	[+16.92, +22.92]	3.01e-37	**1.02e-35**	1.000
Llama 3B	ARC-Easy	T+T:T	2000/2000	-3.85	[-5.05, -2.65]	5.95e-10	**1.55e-08**	**8.92e-09**
Llama 3B	BoolQ	T+T:T	2000/2000	-1.40	[-2.55, -0.25]	0.0212	0.2332	0.2438
Llama 3B	GSM8K	T+F:T	1300/1300	+19.31	[+16.31, +22.38]	1.34e-33	**4.41e-32**	1.000
Llama 3B	ARC-Easy	T+F:T	2000/2000	-4.45	[-5.75, -3.20]	1.37e-11	**3.83e-10**	**2.19e-10**
Llama 3B	BoolQ	T+F:T	2000/2000	+0.20	[-0.95, +1.35]	0.7985	1.000	1.000
Llama 3B	GSM8K	F+T:F	1300/1300	+42.54	[+39.46, +45.62]	2.40e-124	**8.63e-123**	1.000
Llama 3B	ARC-Easy	F+T:F	2000/2000	-1.60	[-2.75, -0.40]	0.0090	0.1356	0.1221
Llama 3B	BoolQ	F+T:F	2000/2000	-0.85	[-2.05, +0.30]	0.1868	0.9647	1.000
Llama 3B	GSM8K	F+F:F	1300/1300	+39.69	[+36.62, +42.77]	3.57e-110	**1.25e-108**	1.000
Llama 3B	ARC-Easy	F+F:F	2000/2000	-5.85	[-7.20, -4.50]	3.22e-17	**1.03e-15**	**5.79e-16**
Llama 3B	BoolQ	F+F:F	2000/2000	-0.15	[-1.30, +1.00]	0.8654	1.000	1.000
Qwen 3B	GSM8K	T+T:T	1300/1300	+3.92	[+2.15, +5.77]	4.04e-05	**8.07e-04**	1.000
Qwen 3B	ARC-Easy	T+T:T	2000/2000	-1.20	[-2.15, -0.30]	0.0149	0.1820	0.1784
Qwen 3B	BoolQ	T+T:T	2000/2000	+0.20	[-0.70, +1.10]	0.7520	1.000	1.000
Qwen 3B	GSM8K	T+F:T	1300/1300	+4.69	[+2.85, +6.54]	7.49e-07	**1.57e-05**	1.000
Qwen 3B	ARC-Easy	T+F:T	2000/2000	-1.35	[-2.30, -0.35]	0.0093	0.1356	0.1221
Qwen 3B	BoolQ	T+F:T	2000/2000	-0.05	[-1.00, +0.90]	1.000	1.000	1.000

Table 55: Combined Capabilities: paired tests, part 1 (Holm family: 36).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}	$p_{\text{degr.}}$
Qwen 3B	GSM8K	F+T:F	1300/1300	+2.62	[+0.92, +4.38]	0.0042	0.0670	1.000
Qwen 3B	ARC-Easy	F+T:F	2000/2000	-0.90	[-1.90, +0.05]	0.0854	0.5979	0.8114
Qwen 3B	BoolQ	F+T:F	2000/2000	-1.15	[-2.05, -0.30]	0.0140	0.1820	0.1750
Qwen 3B	GSM8K	F+F:F	1300/1300	+3.00	[+1.23, +4.77]	0.0011	**0.0185**	1.000
Qwen 3B	ARC-Easy	F+F:F	2000/2000	-1.00	[-2.00, +0.00]	0.0619	0.5226	0.6193
Qwen 3B	BoolQ	F+F:F	2000/2000	-1.15	[-2.10, -0.20]	0.0235	0.2349	0.2584
Qwen 7B	GSM8K	T+T:T	1300/1300	+5.85	[+4.00, +7.77]	8.90e-10	**2.23e-08**	1.000
Qwen 7B	ARC-Easy	T+T:T	2000/2000	-0.75	[-1.50, +0.00]	0.0581	0.5226	0.6097
Qwen 7B	BoolQ	T+T:T	2000/2000	-4.75	[-5.95, -3.60]	1.27e-15	**3.93e-14**	**2.22e-14**
Qwen 7B	GSM8K	T+F:T	1300/1300	+5.23	[+3.38, +7.15]	9.23e-08	**2.03e-06**	1.000
Qwen 7B	ARC-Easy	T+F:T	2000/2000	-1.45	[-2.25, -0.65]	4.22e-04	**0.0080**	**0.0061**
Qwen 7B	BoolQ	T+F:T	2000/2000	-4.50	[-5.65, -3.35]	5.41e-15	**1.62e-13**	**9.19e-14**
Qwen 7B	GSM8K	F+T:F	1300/1300	+5.46	[+3.62, +7.31]	1.32e-08	**3.17e-07**	1.000
Qwen 7B	ARC-Easy	F+T:F	2000/2000	-0.55	[-1.25, +0.15]	0.1608	0.9647	1.000
Qwen 7B	BoolQ	F+T:F	2000/2000	-4.55	[-5.85, -3.25]	3.25e-12	**9.43e-11**	**5.37e-11**
Qwen 7B	GSM8K	F+F:F	1300/1300	+5.31	[+3.46, +7.23]	4.30e-08	**9.90e-07**	1.000
Qwen 7B	ARC-Easy	F+F:F	2000/2000	-1.35	[-2.15, -0.60]	7.30e-04	**0.0131**	**0.0102**
Qwen 7B	BoolQ	F+F:F	2000/2000	-4.15	[-5.35, -2.95]	1.44e-11	**3.88e-10**	**2.23e-10**

Table 56: Combined Capabilities: paired tests, part 2 (Holm family: 36).

E.6.4 Combined BoolQ Passage Sensitivity

Holm correction covers 12 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	BoolQ	T+T:T	2000/1897	-1.40	[-2.57, -0.20]	0.0253	0.1703
Llama 3B	BoolQ	T+F:T	2000/1897	+0.20	[-0.99, +1.36]	0.8033	1.000
Llama 3B	BoolQ	F+T:F	2000/1897	-0.85	[-2.06, +0.35]	0.1893	0.9466
Llama 3B	BoolQ	F+F:F	2000/1897	-0.15	[-1.31, +1.00]	0.8655	1.000
Qwen 3B	BoolQ	T+T:T	2000/1897	+0.20	[-0.75, +1.15]	0.7564	1.000
Qwen 3B	BoolQ	T+F:T	2000/1897	-0.05	[-1.01, +0.91]	1.000	1.000
Qwen 3B	BoolQ	F+T:F	2000/1897	-1.15	[-2.04, -0.25]	0.0154	0.1229
Qwen 3B	BoolQ	F+F:F	2000/1897	-1.15	[-2.11, -0.20]	0.0243	0.1703
Qwen 7B	BoolQ	T+T:T	2000/1897	-4.75	[-5.98, -3.57]	1.00e-05	**1.20e-04**
Qwen 7B	BoolQ	T+F:T	2000/1897	-4.50	[-5.69, -3.35]	1.00e-05	**1.20e-04**
Qwen 7B	BoolQ	F+T:F	2000/1897	-4.55	[-5.87, -3.29]	1.00e-05	**1.20e-04**
Qwen 7B	BoolQ	F+F:F	2000/1897	-4.15	[-5.37, -2.95]	1.00e-05	**1.20e-04**

Table 57: Combined BoolQ Passage Sensitivity: paired tests, part 1 (Holm family: 12).

E.6.5 Trained Variant Capabilities

Only the Llama 3B T BoolQ decline passes the primary two sided correction among 24 comparisons. The separately corrected degradation tests additionally identify Qwen 3B F BoolQ and Llama 8B F ARC-Easy declines. Small score changes should therefore not be equated with proven capability equivalence.
Holm correction covers 24 comparisons.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}	$p_{\text{degr.}}$
Llama 3B	ARC-Easy	F:Base	2000/2000	-0.15	[-0.85, +0.55]	0.7877	1.000	1.000
Llama 3B	BoolQ	F:Base	2000/2000	+1.65	[+0.55, +2.75]	0.0050	0.1046	1.000
Llama 3B	ARC-Easy	T:Base	2000/2000	-0.55	[-1.15, +0.05]	0.0989	1.000	1.000
Llama 3B	BoolQ	T:Base	2000/2000	-2.25	[-3.05, -1.50]	1.18e-08	**2.82e-07**	**1.41e-07**
Qwen 3B	GSM8K	F:Base	1300/1300	+0.77	[-0.31, +1.85]	0.2026	1.000	1.000
Qwen 3B	ARC-Easy	F:Base	2000/2000	-0.10	[-0.60, +0.40]	0.8388	1.000	1.000
Qwen 3B	BoolQ	F:Base	2000/2000	-0.85	[-1.40, -0.35]	0.0023	0.0533	**0.0266**
Qwen 3B	GSM8K	T:Base	1300/1300	-0.54	[-1.38, +0.31]	0.2810	1.000	1.000
Qwen 3B	ARC-Easy	T:Base	2000/2000	-0.20	[-0.70, +0.30]	0.5572	1.000	1.000
Qwen 3B	BoolQ	T:Base	2000/2000	+0.30	[-0.15, +0.75]	0.2632	1.000	1.000
Qwen 7B	GSM8K	T:Base	1300/1300	+0.46	[-0.38, +1.31]	0.3616	1.000	1.000
Qwen 7B	ARC-Easy	T:Base	2000/2000	-0.05	[-0.30, +0.20]	1.000	1.000	1.000
Qwen 7B	BoolQ	T:Base	2000/2000	+0.20	[-0.20, +0.60]	0.4545	1.000	1.000
Qwen 7B	ARC-Easy	F:Base	2000/2000	-0.15	[-0.40, +0.10]	0.4531	1.000	1.000
Qwen 7B	BoolQ	F:Base	2000/2000	-0.20	[-0.65, +0.25]	0.5034	1.000	1.000
Qwen 7B	GSM8K	F:Base	1300/1300	-0.62	[-1.46, +0.23]	0.2005	1.000	1.000
Llama 3B	GSM8K	T:Base	1300/1300	-0.69	[-2.38, +1.00]	0.4744	1.000	1.000
Llama 3B	GSM8K	F:Base	1300/1300	-0.85	[-2.77, +1.08]	0.4307	1.000	1.000

Table 58: Trained Variant Capabilities: paired tests, part 1 (Holm family: 24).

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}	$p_{\text{degr.}}$
Llama 8B	GSM8K	T:Base	1300/1300	-0.08	[-1.62, +1.46]	1.000	1.000	1.000
Llama 8B	GSM8K	F:Base	1300/1300	-1.15	[-2.69, +0.38]	0.1548	1.000	1.000
Llama 8B	ARC-Easy	T:Base	2000/2000	-0.05	[-0.45, +0.35]	1.000	1.000	1.000
Llama 8B	ARC-Easy	F:Base	2000/2000	-0.85	[-1.40, -0.30]	0.0033	0.0732	**0.0366**
Llama 8B	BoolQ	T:Base	2000/2000	-0.05	[-0.55, +0.45]	1.000	1.000	1.000
Llama 8B	BoolQ	F:Base	2000/2000	-0.40	[-0.95, +0.15]	0.2153	1.000	1.000

Table 59: Trained Variant Capabilities: paired tests, part 2 (Holm family: 24).

E.6.6 Trained Variant BoolQ Passage Sensitivity

Holm correction covers 8 comparisons.
The passage-level sensitivity analysis identifies significant BoolQ performance differences for both Llama 3B variants and the Qwen 3B F variant after Holm correction.

Model	Task / outcome	Contrast	N/G	Δ	95% CI	p	p_{Holm}
Llama 3B	BoolQ	F:Base	2000/1897	+1.65	[+0.55, +2.79]	0.0051	**0.0305**
Llama 3B	BoolQ	T:Base	2000/1897	-2.25	[-3.05, -1.45]	1.00e-05	**8.00e-05**
Qwen 3B	BoolQ	F:Base	2000/1897	-0.85	[-1.40, -0.30]	0.0030	**0.0211**
Qwen 3B	BoolQ	T:Base	2000/1897	+0.30	[-0.15, +0.75]	0.2632	1.000
Qwen 7B	BoolQ	T:Base	2000/1897	+0.20	[-0.20, +0.60]	0.4545	1.000
Qwen 7B	BoolQ	F:Base	2000/1897	-0.20	[-0.65, +0.25]	0.5034	1.000
Llama 8B	BoolQ	T:Base	2000/1897	-0.05	[-0.55, +0.45]	1.000	1.000
Llama 8B	BoolQ	F:Base	2000/1897	-0.40	[-0.96, +0.15]	0.2135	1.000

Table 60: Trained Variant BoolQ Passage Sensitivity: paired tests, part 1 (Holm family: 8).

F Complementary MBTI Dimensions

We extend the situational evaluation to E/I, S/N, and J/P on Llama 3.2 3B, Qwen2.5 3B, and Qwen2.5 7B, using the corresponding early-stopped DPO variants and fixed axis-specific prompts. We select two tasks per behavioral level and retain all eight guidance conditions per model and dimension. Inputs, output requirements, and greedy decoding are held fixed across conditions; Llama 3.1 8B is outside this supplement.

A and B denote the first and second poles of each axis (E/I, S/N, J/P). A guided and B guided report DPO / Prompt, with Prompt applied to Base; A+B denotes DPO A with Prompt B. All values are percentages of all task inputs, including unparseable responses in the denominator. These are descriptive judgment and choice rates, rather than accuracy scores.

Axis	Model	Measure	A guided	B guided	A+A	A+B	B+A	B+B
E/I	Llama 3B	Circa unclear	5.6 / 0.6	78.5 / 3.4	0.6	1.9	10.7	44.5
E/I	Llama 3B	Circa unparseable	0.5 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
E/I	Llama 3B	CB false	2.1 / 75.5	60.6 / 66.9	55.2	37.2	97.7	91.0
E/I	Llama 3B	CB unparseable	94.3 / 0.1	38.9 / 3.6	4.3	22.9	0.0	3.9
E/I	Qwen 3B	Circa unclear	75.9 / 78.9	94.3 / 83.4	70.0	75.2	87.0	90.5
E/I	Qwen 3B	Circa unparseable	4.9 / 0.0	0.9 / 0.0	0.2	0.2	0.0	0.0
E/I	Qwen 3B	CB false	29.0 / 29.1	18.1 / 25.4	33.8	30.2	19.5	16.2
E/I	Qwen 3B	CB unparseable	0.0 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
E/I	Qwen 7B	Circa unclear	59.5 / 33.5	65.6 / 45.9	27.9	40.1	41.3	52.2
E/I	Qwen 7B	Circa unparseable	1.6 / 0.2	0.3 / 0.1	0.3	0.2	0.2	0.1
E/I	Qwen 7B	CB false	26.4 / 25.5	24.3 / 23.4	24.9	23.0	27.4	24.4
E/I	Qwen 7B	CB unparseable	0.0 / 0.0	0.1 / 0.0	0.0	0.0	0.0	0.0
S/N	Llama 3B	Circa unclear	2.3 / 1.8	53.5 / 2.4	0.2	0.5	20.5	23.6
S/N	Llama 3B	Circa unparseable	0.0 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.1
S/N	Llama 3B	CB false	11.4 / 71.1	16.0 / 67.7	89.3	90.3	12.6	6.8
S/N	Llama 3B	CB unparseable	88.6 / 1.6	9.2 / 5.6	6.6	7.7	0.1	0.3
S/N	Qwen 3B	Circa unclear	80.4 / 75.5	89.4 / 82.8	58.8	68.0	87.0	91.5
S/N	Qwen 3B	Circa unparseable	1.0 / 0.0	3.1 / 0.1	0.1	0.1	0.1	0.2
S/N	Qwen 3B	CB false	38.9 / 28.5	17.9 / 27.8	40.1	38.9	17.3	17.2
S/N	Qwen 3B	CB unparseable	0.0 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
S/N	Qwen 7B	Circa unclear	42.8 / 34.0	66.3 / 48.1	25.5	38.9	42.6	55.3
S/N	Qwen 7B	Circa unparseable	18.9 / 0.1	0.4 / 0.1	0.1	0.1	0.1	0.1
S/N	Qwen 7B	CB false	27.3 / 26.0	23.5 / 19.3	28.6	21.6	22.4	15.9
S/N	Qwen 7B	CB unparseable	0.8 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
J/P	Llama 3B	Circa unclear	21.9 / 4.8	41.2 / 5.1	8.2	7.7	14.5	19.3
J/P	Llama 3B	Circa unparseable	0.8 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
J/P	Llama 3B	CB false	24.1 / 73.2	39.1 / 79.8	76.0	80.5	94.9	96.8
J/P	Llama 3B	CB unparseable	75.3 / 0.0	60.0 / 0.0	0.3	0.0	0.0	0.2
J/P	Qwen 3B	Circa unclear	47.1 / 81.5	94.8 / 85.9	68.8	74.0	91.6	94.1
J/P	Qwen 3B	Circa unparseable	33.9 / 0.0	0.0 / 0.2	0.1	0.4	0.0	0.1
J/P	Qwen 3B	CB false	33.1 / 36.5	19.3 / 20.0	42.8	27.0	25.3	13.8
J/P	Qwen 3B	CB unparseable	0.1 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0
J/P	Qwen 7B	Circa unclear	54.2 / 36.0	72.7 / 56.6	27.8	45.4	47.5	71.3
J/P	Qwen 7B	Circa unparseable	1.0 / 0.1	0.1 / 0.2	0.1	0.2	0.2	0.3
J/P	Qwen 7B	CB false	29.2 / 23.3	21.6 / 19.3	26.0	22.7	19.5	13.6
J/P	Qwen 7B	CB unparseable	0.0 / 0.0	0.0 / 0.0	0.0	0.0	0.0	0.0

Table 61: Selected situational interpretation results (%). Circa reports UNCLEAR judgments and unparseable outputs over 32,602 inputs; CommitmentBank (CB) reports FALSE judgments and unparseable outputs over 1,009 inputs. Guided cells show DPO / Prompt; combined columns show the trained pole plus the prompted pole.

F.1 Individual and Interactive Decisions

Stock and bargaining report SELL and EXIT choices in the controlled scenarios (1,005 and 900 inputs). INSPIRED reports ASK_MORE after recommendation feedback (119 inputs), and Frames reports SEARCH_OTHER after travel feedback (282 inputs). All responses in these four selected tasks are parseable across the reported conditions.

Brief analysis. Without additional prompts, I, N, and P yield more UNCLEAR judgments on Circa and more ASK_MORE choices on INSPIRED than their opposite poles in all three models. The size of these differences and their response to aligned or opposing prompts vary across tasks and models. Controlled bargaining shows no DPO EXIT choices in Llama 3B or Qwen 7B. High unparseable rates in some interpretation conditions, especially Llama 3B on CommitmentBank, limit behavioral interpretation of those differences.

Axis	Model	Measure	A guided	B guided	A+A	A+B	B+A	B+B
E/I	Llama 3B	Stock sell	48.5 / 29.4	9.2 / 0.0	29.9	0.7	12.9	0.0
E/I	Llama 3B	Bargaining exit	0.0 / 0.0	0.0 / 84.0	0.0	11.7	1.2	100.0
E/I	Llama 3B	INSPIRED ask	17.6 / 0.0	89.9 / 10.9	0.0	3.4	32.8	86.6
E/I	Llama 3B	Frames search	28.7 / 51.1	49.3 / 69.1	18.4	37.9	84.4	70.6
E/I	Qwen 3B	Stock sell	99.7 / 81.9	92.9 / 0.1	98.4	78.3	0.2	0.0
E/I	Qwen 3B	Bargaining exit	9.1 / 0.0	10.7 / 1.8	0.0	0.3	0.4	20.6
E/I	Qwen 3B	INSPIRED ask	72.3 / 88.2	92.4 / 94.1	73.1	79.0	91.6	95.0
E/I	Qwen 3B	Frames search	86.2 / 80.1	94.7 / 69.5	79.4	69.5	84.0	71.6
E/I	Qwen 7B	Stock sell	40.4 / 15.4	39.0 / 14.6	12.5	17.9	16.3	11.7
E/I	Qwen 7B	Bargaining exit	0.0 / 0.0	0.0 / 7.4	0.0	7.1	0.0	5.9
E/I	Qwen 7B	INSPIRED ask	63.0 / 40.3	66.4 / 34.5	37.8	32.8	37.8	31.9
E/I	Qwen 7B	Frames search	34.4 / 39.4	41.5 / 52.1	35.1	46.8	44.7	59.2
S/N	Llama 3B	Stock sell	65.6 / 8.3	27.3 / 0.4	40.8	10.3	0.0	0.0
S/N	Llama 3B	Bargaining exit	0.0 / 48.2	0.0 / 0.2	52.3	0.4	0.0	0.0
S/N	Llama 3B	INSPIRED ask	1.7 / 1.7	95.0 / 18.5	0.0	1.7	81.5	97.5
S/N	Llama 3B	Frames search	18.1 / 43.3	66.0 / 83.0	8.2	24.8	98.6	99.6
S/N	Qwen 3B	Stock sell	97.5 / 86.8	99.6 / 100.0	63.6	98.4	88.2	97.6
S/N	Qwen 3B	Bargaining exit	5.0 / 2.7	19.2 / 0.0	1.6	0.0	9.4	0.8
S/N	Qwen 3B	INSPIRED ask	79.8 / 90.8	87.4 / 89.1	88.2	87.4	95.0	91.6
S/N	Qwen 3B	Frames search	86.5 / 80.1	92.2 / 70.6	75.5	59.2	85.1	79.1
S/N	Qwen 7B	Stock sell	37.5 / 33.5	41.4 / 6.8	29.3	8.0	34.2	6.7
S/N	Qwen 7B	Bargaining exit	0.0 / 45.4	0.0 / 0.0	50.8	0.0	32.8	0.0
S/N	Qwen 7B	INSPIRED ask	59.7 / 5.0	72.3 / 60.5	0.8	46.2	21.8	73.9
S/N	Qwen 7B	Frames search	44.0 / 49.6	34.4 / 55.0	53.5	58.9	46.5	52.5
J/P	Llama 3B	Stock sell	74.6 / 53.7	31.2 / 0.0	75.4	16.5	4.5	0.0
J/P	Llama 3B	Bargaining exit	0.0 / 30.1	0.0 / 0.0	22.9	0.0	47.4	0.0
J/P	Llama 3B	INSPIRED ask	5.0 / 1.7	100.0 / 16.0	0.8	4.2	84.0	97.5
J/P	Llama 3B	Frames search	22.3 / 37.9	68.4 / 80.5	8.5	41.1	96.5	98.6
J/P	Qwen 3B	Stock sell	97.7 / 92.1	99.6 / 58.2	91.7	45.5	93.1	62.3
J/P	Qwen 3B	Bargaining exit	6.6 / 0.0	6.7 / 0.0	1.2	0.0	0.0	0.0
J/P	Qwen 3B	INSPIRED ask	62.2 / 79.0	94.1 / 96.6	46.2	77.3	95.8	99.2
J/P	Qwen 3B	Frames search	75.9 / 64.2	96.5 / 83.7	48.9	72.0	76.2	91.1
J/P	Qwen 7B	Stock sell	43.2 / 42.5	36.8 / 0.7	43.5	2.0	36.5	0.3
J/P	Qwen 7B	Bargaining exit	0.0 / 24.4	0.0 / 0.0	20.4	0.0	24.4	0.0
J/P	Qwen 7B	INSPIRED ask	61.3 / 16.0	68.9 / 57.1	12.6	51.3	16.8	70.6
J/P	Qwen 7B	Frames search	36.5 / 35.8	39.7 / 56.0	34.8	52.1	38.7	59.9

Table 62: Selected individual and interactive decision results (%), with the same guidance notation as Table 61. Bargaining exit rates aggregate the fixed buyer-offer scenarios. INSPIRED ask and Frames search denote ASK_MORE and SEARCH_OTHER, respectively.