arXiv:2611.14364v1 [cs.CL] 10 Nov 2026

ARE-bench: Benchmarking Agents' Research Ability in Real-World Researcher Sessions and Workspaces

Shanghai Artificial Intelligence Laboratory, Fudan University, East China Normal University, Shanghai Jiao Tong University, Renmin University of China, Peking University

Core Contributor: Jiahang Lin†, Wanying Ren†, Wanghan Xu†, Huayu Sha, Yang Shu, Zhonghang Lu, Chengjun Pan, Yi Dong

Contributor: Shichun Liu, Yuyang Hu, Shihan Dou, Honglin Guo, Huijie Lv , Pujun Zheng, Zhe Sun, Yuankai Ying, Yiding Wang, Jiahao Chen, Zhixiong Yang

Advisor: Dingbo Yuan, Yicheng Zou, Zhecan Wang, Guoxiu He, Xiaoyu Xing ,Qi Zhang✉, Tao Gui✉, Xuanjing Huang

† Project leaders. ✉ Corresponding authors.

Research agents achieve increasingly strong benchmark scores, yet these scores do not necessarily translate into reliable delivery in ongoing research projects. Many research-agent benchmarks emphasize standalone tasks and relatively self-contained workflows in clean or curated workspaces. However, researchers ask agents to complete one step within an ongoing project, where requests depend on earlier session decisions and workspaces contain unfinished files, outdated versions, and failed attempts. We introduce ARE-bench to evaluate whether research agents advance ongoing projects while respecting prior decisions and existing workspace state, grounded in **real session context** and **in-progress workspaces**. ARE-bench includes 75 tasks built on agent sessions and full workspaces from recruited researchers' own projects, spanning seven research stages, 33 requirements, and 150 sub-requirements, with 1,173 binary checklist items in total. Each task is anchored at a real user request, preserves its requirements, and summarizes the earlier requests still in effect as background. The complete workspace is replayed to that boundary without cleanup, preserving artifacts from earlier turns. Each checklist is compiled from an intent-aligned rubric and credits any valid solution. Across seven frontier models, the strongest passes every checklist item on only 66.2% of tasks, and Design and Execution rank among the hardest stages for every model. Three models pass about 69% of tasks when only delivery contracts and general norms are checked, but just 28.4% to 40.0% once historical and material or workspace requirements are also checked. Evaluating delivery without these contextual requirements can therefore overestimate research agents and mask differences between them.

Date: November 11, 2026
Project Page: https://arebench.com

1 Introduction

High research-agent benchmark scores can give an incomplete picture of reliable delivery in real research projects. For example, Qwen reports a score of 93.0% for Qwen3.8-Max on PaperBench Code-Dev (Qwen Team, 2026). This variant assesses code development without checking execution or whether the paper's empirical results are reproduced (Starace et al., 2025). Even a near-ceiling score on this assessment therefore leaves important aspects of research delivery untested.

Many established research-agent benchmarks emphasize standalone objectives, such as Kaggle competitions (Chan et al., 2024), research engineering (Wijk et al., 2024), or full paper replication (Starace et al., 2025). In

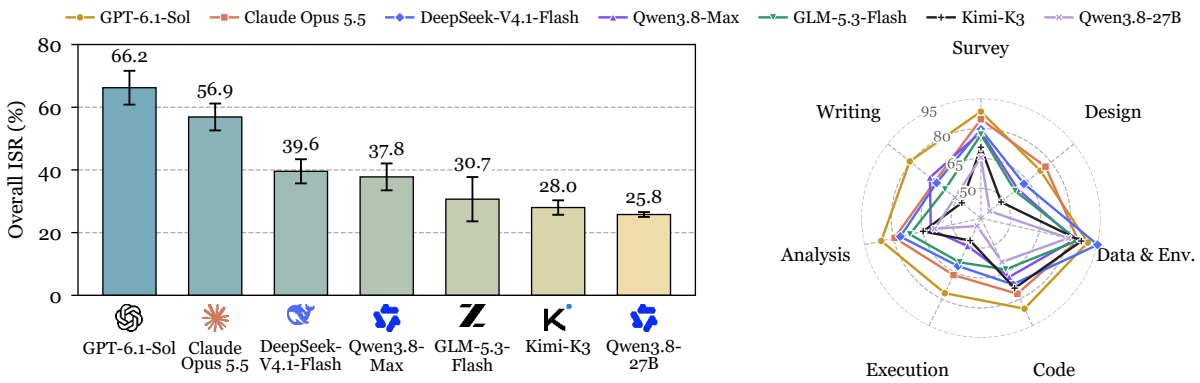

Figure 1 ISR of seven models on 75 tasks, averaged over three rounds. Left: overall ISR, with error bars for the sample standard deviation across rounds. Right: ISR by research stage, with the radial axis from 35% to 95%.

ongoing projects, researchers also ask agents to handle an individual request, such as repairing a script after a failed run or adding a table to a paper. Success then depends on both the requested change and the prior session decisions and workspace artifacts that the change must respect. Standalone objectives and curated starting states can leave this ability to advance an existing project underexamined, even when they rigorously assess their intended outcomes. We therefore evaluate whether research agents advance ongoing projects while respecting prior decisions and existing workspace state, along two dimensions, **real session context** and **in-progress workspaces**.

First, **real session context** supports evaluating whether agents carry earlier decisions into the current delivery. Real research sessions switch between task types (Slinko et al., 2026), returning to literature review after experiments or adding experiments during paper writing (Figure 2, top). A request to fix a failed experiment may still require preserving methodological choices established during an earlier design discussion. Summarizing still-effective earlier requests as Background makes these constraints available for evaluating whether the agent integrates them with the current request.

Second, **in-progress workspaces** support evaluating whether agents change existing artifacts while preserving their scientific meaning. Clean or curated starting states simplify this setting, whereas ongoing projects can contain temporary files, outdated versions, and failed attempts that require checking which artifacts and conventions remain applicable (Figure 2, bottom). For example, in our CylindFormer repair task, two executions both resume training with nonzero loss and parameter updates, yet pass 20/20 and 16/20 checklist items, respectively. The latter changes coordinate handling, overlap labels, and matching losses required by the reference paper in the workspace (Figure 7). Restoring execution alone therefore does not establish that a repair preserves the research method.

Recent benchmarks also incorporate real developer sessions (Lv et al., 2026; Jha et al., 2026) or fine-grained research tasks (Wang et al., 2026a). Our focus is on jointly evaluating the still-effective requirements of a real research request and the in-progress workspace in which that request arose.

We introduce ARE-bench to evaluate this ability. ARE-bench includes 75 tasks spanning seven research stages (Survey, Design, Data & Env., Code, Execution, Analysis, and Writing), 33 requirements, and 150 sub-requirements, with 1,173 binary checklist items in total. We collect agent sessions and full workspaces from recruited researchers' own projects. For session context, we select a real user request as the task boundary, retain its requirements, and use the Background section to summarize the earlier requests still in effect, covering 10.3 earlier requests on average. For in-progress workspaces, we replay the complete workspace to that boundary without cleaning, preserving 41.1 files per task on average, with agent-written files from earlier turns in 65.3% of tasks. The actor receives only the instruction and workspace, and an intent-aligned rubric compiles into the checklist that credits any valid solution.

We evaluate seven frontier models with three rollouts per task by instance success rate (ISR), the fraction of tasks whose checklist items all pass. GPT-6.1-Sol and Claude Opus 5.5 reach 66.2% and 56.9% ISR, while the other five models stay below 40% and Qwen3.8-27B reaches only 25.8% (Figure 1, left). Checklist item

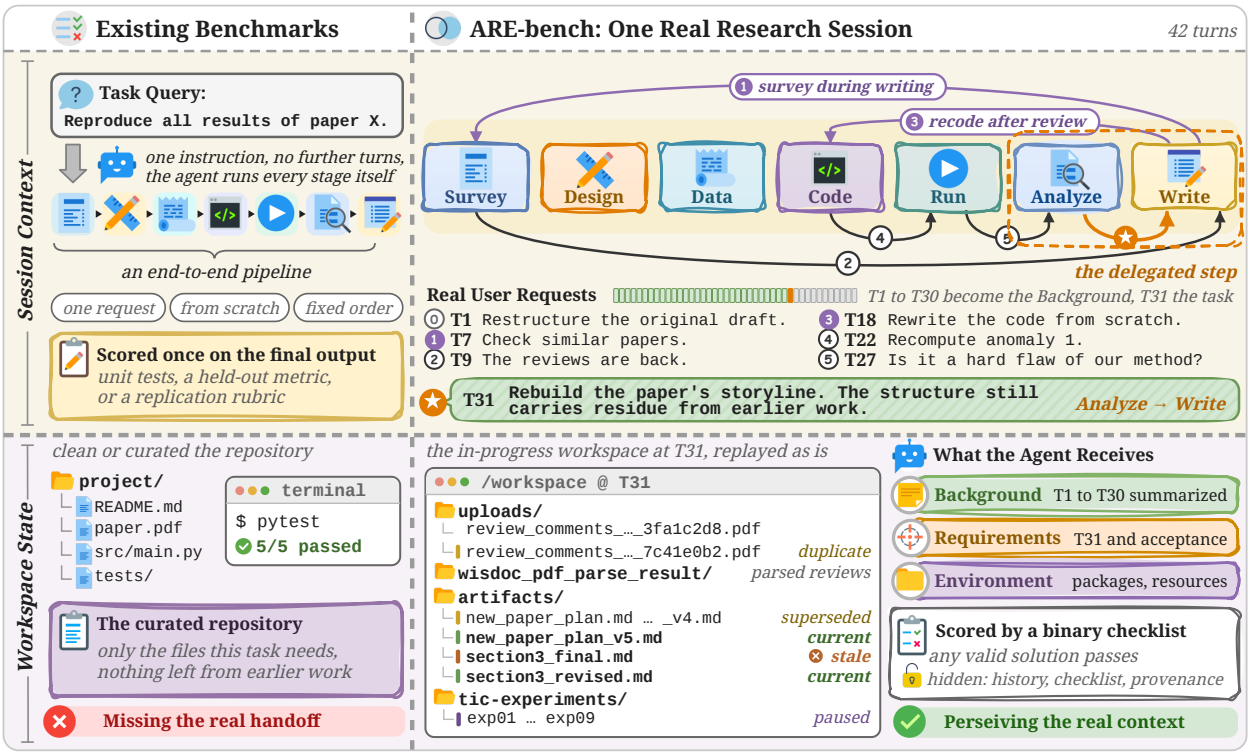

Figure 2 An illustrative contrast between a standalone research task and one real research session in ARE-bench. Top: the standalone task has an end-to-end objective, while the real session moves back and forth between research stages. The task is cut at request T31, with still-effective requests from T1 to T30 summarized in Background. Bottom: a clean or curated starting repository is contrasted with the replayed in-progress workspace, including superseded drafts, misleadingly named files, and paused runs.

success rate (CSR), which gives partial credit, spans only 13.1 points across models, whereas ISR spans 40.4 points, so models satisfy most items but differ in meeting all of them. For the same actor deliveries, DeepSeek-V4.1-Flash, GLM-5.3-Flash, and Kimi-K3 each reach about 69% ISR when only delivery contracts and general norms are checked. Including historical and material or workspace requirements lowers their ISR to 40.0%, 31.6%, and 28.4%. Under the full checklist, the gap between DeepSeek-V4.1-Flash and Kimi-K3 grows from 0.44 to 11.56 points, showing that contextual requirements reveal delivery gaps masked by narrower checks. For session context, real sessions move between research stages, but stage ability is uneven across models. Kimi-K3 ranks third on Code but last on Writing, Qwen3.8-Max ranks fourth overall but second on Writing, and Claude Opus 5.5 and DeepSeek-V4.1-Flash lead Design and Data & Env., so an evaluation covering a single stage would rank models differently. Design and Execution rank among the three hardest stages for every model, where the best and worst models differ by 35.8 and 37.4 points, compared with 15.0 points on the easiest stage, Data & Env. (Figure 1, right). Claude Opus 5.5 also costs 2.6 times as much per run as GPT-6.1-Sol yet trails it by 9.3 points in ISR.

Our contributions are:

- We introduce ARE-bench to evaluate whether research agents advance ongoing projects while respecting prior decisions and existing workspace state, with 75 tasks drawn from real researcher sessions and paired with replayed in-progress workspaces.
- We map each checklist item to one of 207 sub-requirements across seven stages, covering 150 of them and all 33 requirements, and aggregate these labels into a requirement profile for each task.
- We evaluate seven frontier models and find that the strongest reaches only 66.2% ISR, that requirements carried over from the session and workspace expose failures and model differences hidden by checks limited to delivery contracts and general norms, and that stage rankings diverge from overall rankings.

2 Related Work

Research-agent benchmarks. Research-agent benchmarks cover scientific work at different scopes, with success defined by the outcomes each task requires. For specialized research activities, RE-Bench and MLE-bench assess research-engineering and competition performance (Wijk et al., 2024; Chan et al., 2024); ScienceAgentBench and SWE-bench Science evaluate scientific programming and software tasks (Chen et al., 2024; Xu et al., 2026c), while SUPER and ResearchEnvBench target repository execution and runtime preparation (Bogin et al., 2024; Wang et al., 2026b). AARRI-Bench evaluates granular research scenarios involving drafts, logs, and other research artifacts (Wang et al., 2026a). These settings tie acceptance to a specified engineering or research deliverable. Broader coverage includes AstaBench's heterogeneous scientific suite, including literature-query tasks drawn from user logs, with member-specific metrics (Bragg et al., 2025), alongside benchmarks organized around reproduction or experimental outcomes. PaperBench and CORE-Bench assess paper replication and computational reproduction (Starace et al., 2025; Siegel et al., 2024); ResearchGym evaluates research outcomes against baselines, while EXP-Bench assesses experiment execution and conclusions (Garikaparthi et al., 2026; Kon et al., 2025). For scientific-output evaluation, ResearchClawBench anchors rubric-based judgments to a hidden reference study (Xu et al., 2026b). Together, these benchmarks move from specialized research activities to broader scientific workflows while retaining task-specific acceptance targets. ARE-bench takes the current research delegation as its acceptance boundary, compiling requested deliverables and applicable inherited scientific constraints into a checklist. Compared with full-experiment or reference-study targets, this scope supports diagnosing unmet requirements and accepting alternative valid implementations without requiring completion of the surrounding study.

Stateful and interaction-derived benchmarks. An evaluation's starting point determines which prior work an agent must continue. Existing benchmarks provide several forms of task state: some Terminal-Bench tasks start with task-specific repository changes (Merrill et al., 2026); OSWorld configures application and file states, while τ-bench evaluates tool-mediated conversations against database goals (Xie et al., 2024; Yao et al., 2024). LongDS-Bench accumulates analytical state across notebook-derived dialogues (Xu et al., 2026a). For usage-derived tasks, the additional question is how a recorded request and its surrounding work become the evaluation input. SWE-bench pairs issue reports with repository snapshots (Jimenez et al., 2023). At session entry, K-Bench samples real first-turn scientific requests, while SWE-Together starts from a coding session's initial request (Brueckner et al., 2026; Wu et al., 2026). Other benchmarks reconstruct the starting point: RealClawBench rewrites requests into standalone instructions with sanitized standalone workspaces (Lv et al., 2026), and REAP backs out accepted solution changes from a current monorepo (Jha et al., 2026). DuMateBench most directly addresses historical continuation by selecting a local objective, retaining visible history, and restoring recoverable pre-cutoff files; it combines deterministic checks with artifact rubrics (Niu et al., 2026). ARE-bench focuses on subsequent requests in real research projects, pairing recovered in-progress workspaces with Background summaries of still-valid prior context, current Requirements, and Environment. Compared with DuMateBench's general-workflow setting, its research-specific requirements and mapping of checklist items to experiment execution, result analysis, and writing support diagnosis of unmet work by research activity.

Positioning. Table 1 summarizes these comparisons. The table separates a genuine work request from its provenance in a recorded session, then distinguishes inherited workspace progress from the broader research-workflow scope. ARE-bench combines all four binary properties: each task is a real research request selected from a recorded session, starts from an in-progress workspace, and belongs to a benchmark covering the full research lifecycle. Its rubric scoring translates the active delegation into deliverable checks, linking completion to requested outcomes and locating unmet requirements by research activity.

Benchmark	Real request	Session-derived	In-progress workspace	Research workflow	Scoring
SWE-bench (Jimenez et al., 2023)	✔	✘	✘	✘	Tests
Terminal-Bench (2.0) (Merrill et al., 2026)	✔	✘	✔	✘	Varies
OSWorld (Xie et al., 2024)	✔	✘	✔	✘	Varies
τ-bench (Yao et al., 2024)	✔	✘	✔	✘	Varies
MLE-bench (Chan et al., 2024)	✘	✘	✘	✘	Task metric
ScienceAgentBench (Chen et al., 2024)	✘	✘	✔	✘	Varies
PaperBench (Starace et al., 2025)	✘	✘	✘	✘	Rubric
RE-Bench (Wijk et al., 2024)	✘	✘	✘	✘	Task metric
CORE-Bench (Siegel et al., 2024)	✘	✘	✔	✘	Answer check
AARRI-Bench (Wang et al., 2026a)	✔	✘	✔	✔	Varies
ResearchGym (Garikaparthi et al., 2026)	✘	✘	✘	✘	Task metric
EXP-Bench (Kon et al., 2025)	✘	✘	✘	✘	Rubric + Answer check + Tests
AstaBench (Bragg et al., 2025)	✔	✔	✘	✔	Varies
SWE-bench Science (Xu et al., 2026c)	✔	✘	✘	✘	Tests
SUPER (Bogin et al., 2024)	✘	✘	✔	✘	Varies
ResearchEnvBench (Wang et al., 2026b)	✘	✘	✘	✘	Tests
ResearchClawBench (Xu et al., 2026b)	✘	✘	✔	✔	Rubric
K-Bench (Brueckner et al., 2026)	✔	✔	✔	✘	Rubric
DuMateBench (Niu et al., 2026)	✔	✔	✔	✘	Varies
SWE-Together (Wu et al., 2026)	✔	✔	✘	✘	Rubric
RealClawBench (Lv et al., 2026)	✔	✔	✘	✘	Tests
REAP/Harvest (Jha et al., 2026)	✔	✔	✘	✘	Tests
LongDS-Bench (Xu et al., 2026a)	✘	✘	✘	✘	Answer check
ARE-bench	✔	✔	✔	✔	**Rubric**

Table 1 Comparison of benchmark designs for research-project continuation. **Real request** denotes a documented genuine work need, regardless of source. **Session-derived** denotes a task or request whose provenance is an actual recorded work session; earlier conversation transfer is optional. **In-progress workspace** denotes task-relevant files or application state containing prior activity or intermediate progress available when evaluation starts; a base repository, raw data, dependency environment, or empty scaffold alone does not qualify. **Research workflow** denotes suite-level coverage of Survey, Design, Data & Environment, Code, Execution, Analysis, and Writing. **Scoring** names the primary acceptance basis: Tests, State check, Answer check, Task metric, Rubric, a stable combination written with '+', or Varies when mechanisms differ across tasks or sub-benchmarks. A tick marks a documented feature and a cross means it is not established under these definitions.

3 ARE-bench

3.1 Task Formulation

Each task is anchored at a user request in an ongoing research session. At cut point c, the raw task is

$$\tau_c = (q_c, H_{<c}, W_c), \tag{1}$$

where q_c is the current request, $H_{<c}$ is the history of earlier user requests, and W_c is the workspace immediately before the agent responds. Construction maps τ_c to the released task

$$b_c = (I_c, W_c, R_c), \tag{2}$$

where I_c is the instruction and R_c is the rubric and provenance added during construction. The benchmark is $B = \{b_c\}$.

We define the instruction as three sections,

$$I_c = \left(I_c^{\text{bg}}, I_c^{\text{req}}, I_c^{\text{env}}\right), \tag{3}$$

with Background derived from the request history as

$$I_c^{\text{bg}} = \text{Compress}(H_{<c}), \tag{4}$$

where Compress retains context relevant to q_c that remains valid at cut point c and rewrites it as a state summary. Requirements I_c^{req} preserves q_c and specifies delivery paths and acceptance conditions. Environment I_c^{env} describes available resources and execution limits.

The actor receives only I_c and W_c and produces a delivery y_c. The rubric compiles into binary checklist items aligned with user intent. Let K_c be the set of verifiable checklist items, excluding non-triggered conditional ones, and $r_{c,k} \in \{0, 1\}$ indicate whether checklist item k is satisfied.

3.2 Benchmark Construction

Figure 3 shows the three construction roles. The Curator selects researcher sessions, the Builder turns a request in a selected session into a task, and the Validator tests each task before release.

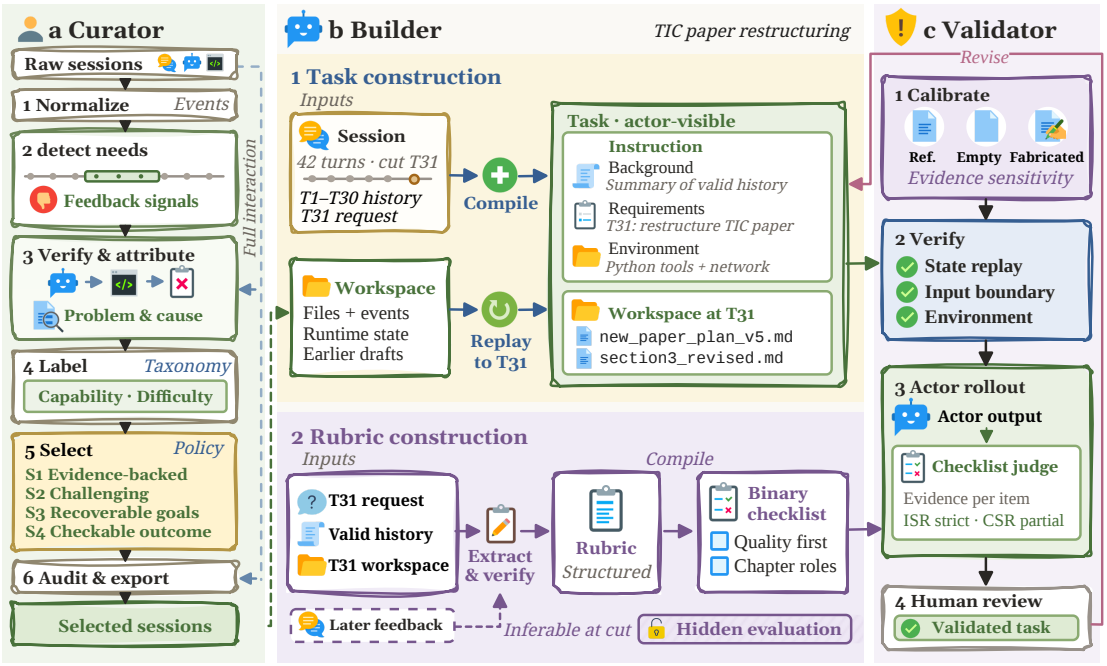

Figure 3 Overview of ARE-bench construction, illustrated by a paper restructuring task cut at request T31 of a 42-turn session. (a) The Curator normalizes raw researcher sessions, verifies requirements against the full interaction, labels them with the requirement taxonomy, and selects sessions by rules S1 to S4. (b) The Builder compiles the target request and still-valid history into the actor-visible instruction, replays the workspace to that request, and compiles a hidden rubric into a binary checklist. (c) The Validator calibrates the judge, verifies the task package, runs actors, and sends each task to human review, which returns failing tasks to the Builder.

3.2.1 Curator

The Curator filters raw researcher sessions in six steps. It normalizes each session into ordered turns and events that keep their source positions. Reading user turns alone, it splits the session into episodes, recovers goals and requirements with their revisions, and flags feedback signals such as rejections or requests to redo work. It then verifies each requirement against the full interaction, including agent replies, tool results, and produced files, and attributes each confirmed problem to its cause. Verified requirements are labeled with the research requirement taxonomy of Appendix B and rated for difficulty, and each label cites quoted evidence. After schema and evidence prechecks, an episode is selected only if it satisfies four rules:

1. **S1 Evidence-backed**: at least one primary requirement label is supported by evidence.
2. **S2 Challenging**: the task has medium or high difficulty.
3. **S3 Recoverable goals**: the goals and requirements can be recovered coherently from the session.
4. **S4 Checkable outcome**: completion criteria can be checked from the record, whether or not the original agent succeeded.

Episodes with an unresolved rule are held for review, and other failing episodes are rejected. A final audit checks every quote against the source session and exports each session that contains a selected episode.

3.2.2 Builder

Task construction. The Builder picks a target request q_c at a turn where a core requirement of the selected episode is carried out. Background summarizes earlier requests still in effect as the state at cut point c. It takes facts from earlier agent turns but omits their unexecuted plans, which would reveal solutions. Requirements retains q_c with its delivery paths and explicit acceptance conditions, and Environment states the available tools and network access. The workspace is reconstructed by replaying the platform event stream $\mathcal{A}$ up to the cut point,

$$W_c = \text{Replay}(\mathcal{A}, c), \tag{5}$$

and cross-checked against the transcript and the final disk state. A cut point is discarded if it falls inside a file write, if replay disagrees with the transcript, or if a required dependency cannot be recovered. The environment contains the platform baseline and the packages available at the cut point, and it keeps network access when the source environment had it.

Rubric construction. The rubric follows the user's intended outcome rather than the original agent's output. The Builder extracts intent from q_c, the still-valid history, and the workspace at the cut point. Later user feedback often reveals requirements the user held without stating, so the Builder uses it to sharpen such requirements and keeps an item only if a competent researcher could infer it at the cut point. Requirements raised only in later turns and personal preferences are excluded. Each item checks one deliverable or acceptance condition with a binary decision rule and an evidence requirement. Items do not prescribe implementation paths, so any valid alternative passes. Items stated by the user appear in Requirements, while items a competent researcher would infer stay out of the instruction so that it does not become an answer checklist. Each item also receives one sub-requirement label from the taxonomy, and a task's requirement profile is the union of the parent requirements of its items. The rubric and a provenance package holding raw wording, source turns, path mappings, replay logs, and anonymization decisions form the hidden evaluation set. Appendix E walks through a complete task.

3.2.3 Validator

The Validator tests each task in four steps and returns failing tasks to the Builder for revision. Calibration scores three deliveries with the checklist judge. The original agent's output, when available, tests whether each verdict cites evidence present in the delivery. It calibrates the judge rather than serving as an answer key, because later turns often show that it fell short. An empty delivery should score near zero, and a fabricated delivery that looks complete but contains false content should score clearly below actor runs. Verification then confirms three properties. The replayed state matches the audited state for every file the task uses, the instruction and workspace contain no information from after the cut point, and the released environment installs and imports the required dependencies. In the actor rollout, an agent runs the task and the judge scores each item with cited evidence, yielding ISR and CSR. When every item passes, the Builder rechecks rubric coverage and strictness but does not add uninferable items to lower the score. In human review, the task author checks every failed item. A task proceeds to peer review only when each failure is a valid item that the actor genuinely missed. Otherwise the instruction or rubric is revised and the task is rerun.

3.3 Benchmark Statistics

Table 18 summarizes ARE-bench. Each of the 75 tasks compiles into 15.6 checklist items on average. Three properties distinguish the task distribution. First, every task starts from an in-progress workspace: the replayed workspace holds 41.1 files on average, 65.3% of tasks contain agent-written files from earlier turns, and the Background section compresses 36.0 earlier user turns into a state summary that cites on average 10.3 earlier requests still in effect. Second, each task is a single local step with 2.5 deliverables on average. Third, the checklist items span all seven research stages.

These tasks contain 1,173 checklist items covering 150 sub-requirements and all 33 requirements. Table 2 summarizes this coverage across the seven research stages, and Table 14 in Appendix B presents the corresponding requirement hierarchy. A task can contribute to multiple stages, so the task counts across rows overlap.

Research stage	Requirements	Sub-requirements	Checklist items	Tasks
Survey	2	20	111	26
Design	3	25	268	41
Data & Env.	5	16	45	20
Code	5	23	217	32
Execution	7	26	173	33
Analysis	4	13	92	34
Writing	7	27	267	31
Total	33	150	1,173	75

Table 2 Observed requirement coverage across the 75 tasks. Tasks count distinct tasks within each stage, and the total counts their union.

3.4 Evaluation Protocol

For each task, an actor receives the instruction and the restored workspace, works under the declared tool and time limits, and submits a delivery y_c containing the artifacts listed in the instruction. The checklist evaluator reads y_c and the minimum input source needed to verify a claimed change. It then assigns each $r_{c,k}$ using the checklist item's evidence rule. Missing artifacts, unresolvable claims, and unsupported conclusions receive zero for the affected checklist item.

The evaluation preserves both the request context and the project's in-progress state. Background carries forward earlier user requests that remain in effect, Requirements specifies the current deliverables and acceptance conditions, and Environment states the available resources and limits. The restored workspace retains unfinished artifacts, outdated versions, and prior failed attempts, requiring the actor to interpret the current request against existing project materials. Checklist items and provenance remain hidden from the actor and support scoring and audit.

Each model runs three rollouts per task, and we report the mean and sample standard deviation of each metric across the three rounds. Each run records its model identifier, actor harness, task resources, and time limit. A run that reaches its time limit is scored on the artifacts delivered by then.

3.5 Metrics

Following the checklist-based definitions of Ding et al. (2026), let $N = |B|$ and $n_c = |K_c|$, with K_c and $r_{c,k}$ as defined above. The Instance Success Rate (ISR) is strict: a task is successful only when every verifiable checklist item passes.

$$\mathrm{ISR} = \frac{1}{N} \sum_{c \in B} \mathbf{1} \left[\bigwedge_{k \in K_c} (r_{c,k} = 1) \right]. \tag{6}$$

The Checklist Item Success Rate (CSR) gives partial credit at checklist-item level:

$$\mathrm{CSR} = \frac{1}{N} \sum_{c \in B} \frac{1}{n_c} \sum_{k \in K_c} r_{c,k}. \tag{7}$$

ISR captures holistic satisfaction under conjunctions of heterogeneous requirements, whereas CSR diagnoses which requirements are most often violated. We report peak context and estimated API cost per run separately. Tasks with no verifiable checklist items are invalid construction outputs and are reported as construction errors rather than assigned a score.

4 Experiments

4.1 Experimental Setup

We evaluate agents on continuing user requests in real research workspaces under the protocol in Section 3.4. Actors receive the instruction and restored workspace, and deliveries are scored against task-specific binary

checklists.

Tasks and evaluated models. The main evaluation comprises 75 tasks: 67 CPU tasks, 8 GPU tasks. Its 1,173 checklist items cover the seven research stages defined above: Survey, Design, Data & Env., Code, Execution, Analysis, and Writing. We evaluate the seven models in Table 3 using the actor harness recorded for each model in the archive.

Evaluated model	Archived model identifier	Actor harness
GPT-6.1-Sol	`gpt-6.1-sol`	Codex
Claude Opus 5.5	`claude-opus-5-5`	Claude Code
DeepSeek-V4.1-Flash	`deepseek-v4.1-flash`	Claude Code
Qwen3.8-Max	`qwen3.8-max`	Claude Code
GLM-5.3-Flash	`glm-5.3-flash`	Claude Code
Kimi-K3	`kimi-k3`	Claude Code
Qwen3.8-27B	`qwen3.8-27b`	Claude Code

Table 3 Evaluated models and actor harnesses. Model identifiers follow the main evaluation archive. The specific historical GLM deployment identifiers are retained in individual run records.

Execution environments and compute resources. CPU tasks run in E2B sandboxes or Docker containers, while GPU and research tasks use Docker. Table 4 reports the declared task resources and actor time limits recorded for each batch. Resources are configured per task to support the existing workspace and its dependencies. GPU and research batches use Codex 0.118.0 or Claude Code 2.1.283.

Task batch	Tasks	CPUs	Memory (MB)	GPUs	Configured actor time limit
CPU tasks	67	2–16	4,096–49152	0	30 minutes–6 hours
GPU tasks	8	8–32	32,768–262,144	1–4	2, 3, 4, 6, or 48 hours, depending on the task

Table 4 Declared task resources and configured time limits. CPU, memory, and GPU counts come from the task configuration bound to each selected source. Time limits are calculated from overrides and multipliers in the run configuration and specify the actor execution budget. The physical hosts and GPU models remain to be verified.

Rollouts and metrics. Each model performs three rollouts per task, with each rollout scored independently against the checklist. A GLM-5.3 judge issues the verdicts. Following the preceding definitions, **Instance Success Rate (ISR)** requires all applicable checklist items for a task to pass. **Checklist Item Success Rate (CSR)** first computes the fraction of passing items within each task and then averages equally across tasks. The main table computes ISR and CSR separately for each round and reports the three-round mean and sample standard deviation, corresponding to single-attempt pass@1 evaluation.

The experiments address the following questions.

RQ1: Can current models fully complete user tasks in real research workspaces?
Integrating the current request with historical context and workspace requirements

RQ2: Why do agents fail to complete research deliveries in real research workspaces?
Correctness bottlenecks and omissions of historical, material, and workspace requirements

RQ3: Which research stages are most difficult?
Stage difficulty and the distribution of model capabilities

RQ4: Do larger contexts and higher costs correspond to greater strict success?
The relationship between resource use and success

RQ5: How does changing the LLM judge affect binary verdicts and model comparisons?
Item agreement and ranking sensitivity

4.2 Main Results

4.2.1 RQ1: Can current models fully complete user tasks in real research workspaces?

Model	Overall ISR ↑	Overall CSR ↑
GPT-6.1-Sol	66.22 ± 5.39	92.72 ± 2.08
Claude Opus 5.5	56.89 ± 4.29	90.20 ± 1.83
DeepSeek-V4.1-Flash	39.56 ± 3.85	89.26 ± 1.72
Qwen3.8-Max	37.78 ± 4.29	85.45 ± 0.90
GLM-5.3-Flash	30.67 ± 7.06	86.66 ± 0.68
Kimi-K3	28.00 ± 2.31	85.33 ± 0.50
Qwen3.8-27B	25.78 ± 0.77	79.66 ± 0.24

Table 5 Overall performance of seven models on 75 tasks. ISR and CSR are percentages, reported as the three-round mean ± sample standard deviation, with higher values indicating better performance. Darker blue indicates higher percentages on a shared quadratic scale.

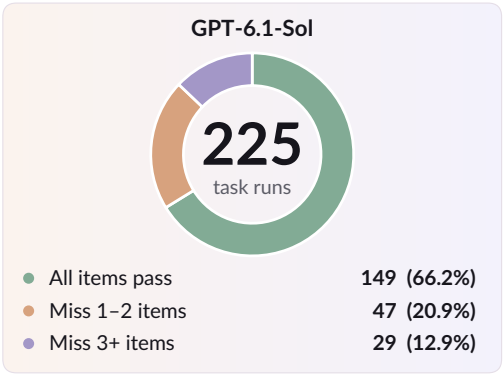

Figure 4 Breakdown of GPT-6.1-Sol outcomes. The unit is a task run, and a run may fail multiple checklist items.

Current models remain unreliable at completing research requests in real workspaces. As Table 5 shows, even the strongest model, GPT-6.1-Sol, achieves only about 66.2% strict task success, leaving roughly one-third of runs unable to satisfy all task requirements. Claude Opus 5.5 reaches 56.9%, all remaining models stay below 40%, and Qwen3.8-27B achieves only about 25.8%. Real research requests build on prior work. Background summarizes earlier requests still in effect, and Requirements specifies the current delivery and acceptance conditions, whose project-specific meaning must be interpreted through the restored code, papers, configurations, and logs. Full completion therefore requires models to situate the current request within its project context and identify the requirements that the current changes must continue to satisfy.

CSR exceeds overall ISR for every model, ranging from 79.66% to 92.72%. Qwen3.8-Max's higher ISR but lower CSR than GLM-5.3-Flash is consistent with these definitions: ISR measures fully successful deliveries, while CSR measures average checklist completion across tasks, so the metrics need not rank models identically. The per-run failure distribution clarifies this gap: of GPT-6.1-Sol's 76 failed runs, 47 miss only one or two checklist items, accounting for 61.84%, and this share ranges from 53.09% to 72.06% across the seven models. Many failures therefore occur after the model has made progress and satisfied most checklist items. **Full delivery requires models to jointly address the current request and project requirements still in effect.** Even after resolving the most prominent issue in the current turn, overlooking prior agreements in Background or methodological requirements carried by the workspace can leave critical gaps. RQ2 examines how these gaps manifest in archived deliveries. Figure 4 breaks down the strongest model's failures.

4.2.2 RQ2: Why do agents fail to complete research deliveries in real research workspaces?

Model	Completeness ISR	Correctness ISR	Constraint compliance ISR	Failed runs with correctness errors
GPT-6.1-Sol	83.11 ± 5.39	71.11 ± 5.05	93.78 ± 0.77	65/76 (85.5%)
Claude Opus 5.5	87.56 ± 3.08	62.67 ± 3.53	89.33 ± 4.62	84/97 (86.6%)
DeepSeek-V4.1-Flash	82.67 ± 2.67	47.11 ± 1.54	88.89 ± 5.05	119/136 (87.5%)
Qwen3.8-Max	79.11 ± 2.04	43.56 ± 8.57	84.89 ± 0.77	127/140 (90.7%)
GLM-5.3-Flash	80.44 ± 0.77	39.11 ± 3.08	81.78 ± 1.54	137/156 (87.8%)
Kimi-K3	76.44 ± 1.54	34.22 ± 2.04	83.56 ± 1.54	148/162 (91.4%)
Qwen3.8-27B	72.00 ± 1.33	31.11 ± 0.77	79.11 ± 2.04	155/167 (92.8%)

Table 6 ISR by completeness, correctness, and constraint compliance. Each dimension requires all its applicable checklist items to pass. The final column reports the count and share of failed runs containing correctness errors. Error categories can overlap. Darker blue indicates higher ISR on a shared quadratic percentage scale; darker orange indicates a greater share of failures with correctness errors.

RQ2.1: Which checklist items most often block full delivery? Table 6 identifies correctness as the lowest-scoring dimension for all seven models. Correctness ISR ranges from 31.11% to 71.11%, compared with 72.00%–87.56% for completeness and 79.11%–93.78% for constraint compliance. The same pattern holds for CSR, with correctness lowest for every model. At the run level, 85.5%–92.8% of failures include an unmet correctness item. DeepSeek-V4.1-Flash reaches 93.16% completeness CSR but only 47.11% correctness ISR and 39.56% overall ISR, illustrating the gap between delivery coverage and content reliability.

Research requirements must hold throughout implementation and validation, as local functional recovery can still lose methodological meaning. RQ1 concerns whether models integrate the context needed for a task, while Table 6 shows the difficulty of accurately realizing its requirements in research artifacts. The current request may ask only for training repair or framework adaptation, but the model must also preserve the original method and engineering conventions using Background and the papers, code, and design documents in the workspace. The archived examples in Table 7 show the consequences of failing to implement these requirements. The CylindFormer repair retains the KPConv name but omits neighborhood and kernel-point aggregation, while the Graph RAG adaptation implements a local workflow without isolating indexes across datasets. Delivering files or code can satisfy some completeness requirements while deviations between the method and its implementation fail correctness items.

These examples illustrate the difficulty of full delivery through method preservation, requirement implementation, and evidence checking. Models must interpret Background and Requirements together and verify that the requirements still hold against existing workspace artifacts after making changes. For example, resuming training also requires checking the original method's key computations, and adapting a framework requires checking its behavior across datasets. Reliable continuation means advancing the current work while maintaining the project's existing requirements.

Failure manifestation	Effect on research delivery	Archived example
Task degradation	A simpler implementation replaces the original method, retaining its name while changing the required capability	In the round 1 DeepSeek repair of CylindFormer, the backbone named KPConv applies only Linear, LayerNorm, and ReLU to individual points, without neighborhood or kernel-point aggregation, failing O4
Insufficient investigation and evidence checking	Cited work does not support the mechanism or subject discussed, leaving conclusions without reliable evidence	In the round 1 Qwen3.8-27B revision of a cuproptosis review, ferroptosis-resistance research is described as a cuproptosis mechanism and cancer subtypes are conflated, failing O6
Requirement omission	Existing constraints are not implemented, affecting subsequent use despite local functional progress	In the round 2 DeepSeek adaptation of Graph RAG, datasets share a global index directory instead of isolated indexes, failing O11
Inadequate validation design	Missing comparison rules prevent attribution of improvements to the factor under study	In the round 1 GLM-5.3-Flash plan for 4D spacetime tubes, costs are recorded but executable rules for matching inference budgets across comparison arms are absent, failing O6
Conclusions exceeding the evidence	A limited diagnostic result is extended into an untested causal explanation	In the round 2 GLM-5.3-Flash diagnosis of network-security generalization, adversarial validation is treated as a test of conditional distributions, and a failure to detect differences is attributed to a single cause, failing O13

Table 7 Failure manifestations and archived examples. These selected qualitative examples explain why checklist items fail. They do not constitute a uniform classification or frequency analysis of all failures.

RQ2.2: What delivery gaps are revealed by historical and material or workspace requirements? Real research requests often briefly describe only the current work, leaving some requirements and constraints implicit in prior agreements and project materials rather than restating them in the current turn. In ARE-bench, Background summarizes historical constraints, while papers, code, configurations, and existing results in the workspace clarify the specific meaning of methods and acceptance conditions. The agent must therefore interpret the current request within the project context and identify requirements that its changes must continue to respect. To examine these gaps, we hold actor deliveries and item-level judge verdicts fixed and change only the checklist items included in ISR. We first check delivery contracts and general professional norms, then add historical and material or workspace requirements, and finally check the complete checklist. Table 8 contrasts three models with similar scores under the first scope.

Model	Delivery contracts + general norms	Add historical and material or workspace requirements	Full checklist
DeepSeek-V4.1-Flash	69.33	40.00	39.56
GLM-5.3-Flash	69.33	31.56	30.67
Kimi-K3	69.78	28.44	28.00

Table 8 ISR (%) as the scope of checklist evaluation expands. Every model covers the same 75 tasks with three rollouts per task, and values are three-round means. Each column requires all applicable items within its scope to pass, reusing the same deliveries and item verdicts. Delivery contracts and general norms correspond to `contract + norm`. The second scope adds `inherited + anchored`, while the full checklist also includes `latent` items discovered through subsequent rework but inferable at the task cut point. Darker blue indicates higher ISR on a shared quadratic percentage scale.

As Table 8 shows, all three models achieve approximately 69% ISR when only delivery contracts and general norms are checked. Adding historical and material or workspace requirements reduces DeepSeek, GLM, and Kimi to 40.00%, 31.56%, and 28.44%, respectively. Among runs passing the first group of checks, 66, 85, and

93 still fail at least one historical or material or workspace requirement, accounting for 48.5%, 54.5%, and 57.4% of each model's total failed runs. DeepSeek and Kimi differ by only 0.44 percentage points under the first scope, compared with 11.56 points under the full checklist.

An important bottleneck in real research delivery is satisfying implicit requirements that remain binding through project context but are not restated in the current turn. Roughly half of the failed runs for these three models satisfy delivery contracts and general norms yet are blocked by historical or material or workspace requirements. Thus, even a model that performs the action described in the current request may fail to deliver a usable result because it departs from prior research agreements, material-grounded methods, or the computations the project requires. Reliable research continuation requires the agent to integrate these distributed constraints and preserve them throughout modification and validation.

This capability requirement connects back to the distinction in task granularity and starting state established earlier. Compared with a setting that pursues a standalone end-to-end objective from an initial state, ARE-bench starts at a user request within an ongoing real project, retaining historical agreements still in effect in Background and existing artifacts in the cut-point workspace. **Evaluation concerns both achieving the current objective and maintaining the project's existing research requirements while doing so.** Table 8 shows that this dimension reveals capability differences that are difficult to observe when only delivery contracts and general norms are checked.

4.3 Stage Analysis

4.3.1 RQ3: Which research stages are most difficult?

Model	Survey	Design	Data & Env.	Code	Execution	Analysis	Writing
GPT-6.1-Sol	88.46 ± 6.66	73.17 ± 4.88	90.00 ± 0.00	85.42 ± 9.55	76.77 ± 7.00	86.27 ± 6.12	80.65 ± 5.59
Claude Opus 5.5	84.62 ± 3.85	76.42 ± 2.82	85.00 ± 5.00	77.08 ± 7.86	66.67 ± 3.03	79.41 ± 2.94	64.52 ± 11.63
DeepSeek-V4.1-Flash	79.49 ± 4.44	62.60 ± 2.82	95.00 ± 0.00	71.88 ± 5.41	61.62 ± 4.63	76.47 ± 0.00	63.44 ± 6.72
Qwen3.8-Max	78.21 ± 4.44	58.54 ± 7.32	83.33 ± 12.58	67.71 ± 6.51	50.51 ± 4.63	60.78 ± 1.70	67.74 ± 5.59
GLM-5.3-Flash	76.92 ± 3.85	56.91 ± 6.14	85.00 ± 5.00	63.54 ± 1.80	59.60 ± 1.75	71.57 ± 6.12	58.06 ± 5.59
Kimi-K3	70.51 ± 4.44	47.97 ± 2.82	86.67 ± 7.64	73.96 ± 4.77	47.47 ± 3.50	64.71 ± 0.00	47.31 ± 3.72
Qwen3.8-27B	65.38 ± 6.66	40.65 ± 2.82	80.00 ± 5.00	59.38 ± 5.41	39.39 ± 0.00	58.82 ± 2.94	51.61 ± 0.00

Table 9 ISR across seven research stages. A task passes a stage only if all its applicable checklist items in that stage pass. Tasks without items in the stage are excluded from its denominator. Values are three-round means ± sample standard deviations, and one task may involve multiple stages. Darker blue indicates higher ISR on a shared quadratic percentage scale.

Research stage	Tasks covered	Checklist items	Seven-model mean stage ISR	Best-to-worst model gap
Survey	26	111	77.66	23.08
Design	41	268	59.47	35.77
Data & Env.	20	45	86.43	15.00
Code	32	217	71.28	26.04
Execution	33	173	57.43	37.37
Analysis	34	92	71.15	27.45
Writing	31	267	61.90	33.33

Table 10 Stage coverage and descriptive difficulty. The seven-model mean weights models equally, and the best-to-worst gap is in percentage points. Task counts indicate stage coverage in the benchmark, while effective denominators follow the scored records in each round. Darker blue indicates higher mean ISR, and darker orange indicates larger model gaps, on a shared quadratic percentage scale.

Tables 9 and 10 show that Design and Execution are among the three most difficult stages for every model. The seven-model mean ISR is 86.43% for Data & Env. and 71.28% for Code, falling to 59.47% for Design

and 57.43% for Execution. Design and Execution also show the largest gaps between models, at 35.77 and 37.37 percentage points, respectively. Stage strengths differ: GPT-6.1-Sol leads five stages, Claude Opus 5.5 leads Design, and DeepSeek-V4.1-Flash leads Data & Env. Kimi-K3 ranks third on Code at 73.96%, but its Execution and Writing scores are only 47.47% and 47.31%, ranking second-last and last, respectively.

Strength in code implementation does not yet ensure effective research work. Dependency installation and code executability provide relatively direct verification feedback. Design and Execution additionally require judging whether an experiment can answer the research question and taking appropriate action based on existing scripts, logs, and checkpoints. These stage differences indicate that implementation strength must be assessed separately in subsequent execution, validation, and reporting. A single score cannot capture this capability distribution. Stages come from naturally occurring research sessions, so this interpretation concerns the evaluated task distribution. Future ablations can further distinguish the contributions of context and workspace state.

4.4 Session Position and Cut-Point Selection

4.4.1 Supplementary question: Are tasks near the session midpoint more challenging than those at adjacent positions?

Position relative to the complete session midpoint	CSR (%)
Before the midpoint, original segment 2	95.92
Near the midpoint, original segment 3	91.63
After the midpoint, original segment 4	97.66

Table 11 CSR at and around the midpoint of the complete session. Each of 10 sessions is divided into five consecutive segments by user-turn count. We show the original segments 2–4, each containing 10 tasks, for 30 tasks in total. GLM-5.3 runs on the instruction and restored workspace at each cut point, and results weight tasks equally. Blue intensity is scaled over the displayed CSR range, with darker cells indicating higher values.

As Table 11 shows, CSR near the midpoint is 91.63%, lower than the preceding and following segments by 4.29 and 6.03 percentage points, respectively. Within these three adjacent segments, tasks in the middle of the session are more challenging. Models can already use the state summary in Background and artifacts in the restored workspace, but must still determine how earlier decisions constrain the current request and advance the research accordingly.

The midpoint of a complete session provides candidate cut points that remain challenging despite the available research state. This observation gives empirical support to the construction practice of typically seeking candidate requests near the complete user session's midpoint. Candidates come from sessions that pass Curator rules S1 to S4 and must also pass the Builder's replay checks. After selection, earlier requests still in effect are summarized in Background, the current request and delivery and acceptance conditions are retained in Requirements, and the workspace is replayed to the cut point. The experiment therefore examines an in-progress research state consistent with the construction procedure.

4.5 Cost Analysis

4.5.1 RQ4: Do larger contexts and higher costs correspond to greater strict success?

Figure 5 shows that neither peak context nor cost aligns with strict success across models. DeepSeek-V4.1-Flash reaches a peak context of 132.55K, close to the 142.87K of GPT-6.1-Sol, but achieves 39.56% ISR compared with 66.22%. Claude Opus 5.5 and Qwen3.8-27B reach similar peak contexts of 104.71K and 108.02K, yet their ISR differs by 31.11 percentage points. Under the frozen prices, GPT-6.1-Sol costs about USD 2.17 per run and Claude Opus 5.5 about USD 5.69, roughly 2.62 times as much for an ISR lower by 9.33 percentage points. GLM-5.3-Flash and DeepSeek-V4.1-Flash offer 30.67% and 39.56% ISR at approximately USD 0.20 and USD 0.53, respectively.

Larger contexts and higher costs do not automatically produce more reliable research deliveries. Under this

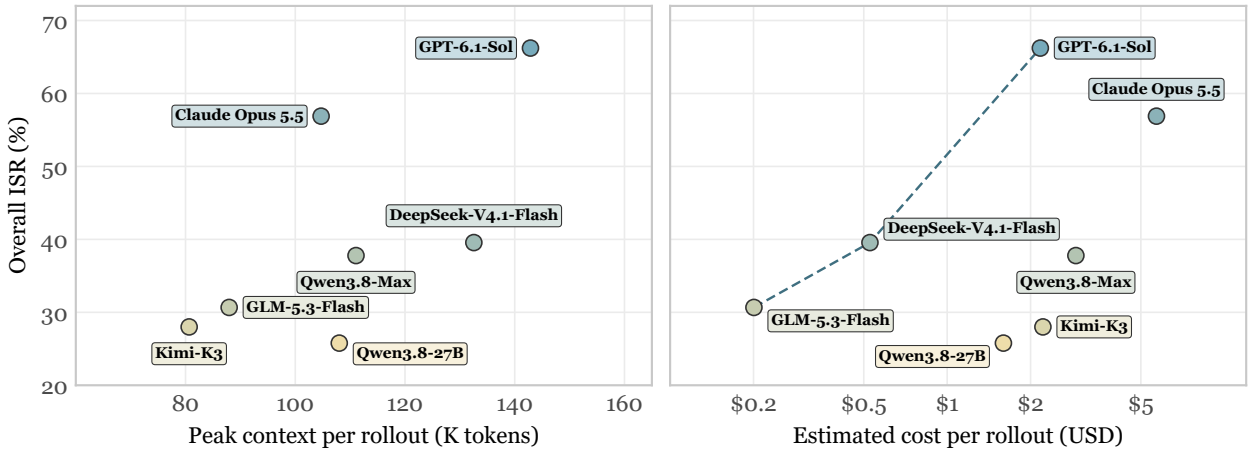

Figure 5 Mean peak context, estimated cost, and ISR per run. Peak context is the largest single API call input within a run, and K denotes thousands. Estimates apply the manuscript's frozen prices dated 2026-11-02 to the measured uncached input, cache reads, cache writes, and output tokens of every call, covering model API charges. Cache writes use the ordinary input price, except that Claude Opus 5.5 uses its cache-write price of USD 5 per million tokens. DeepSeek's context and cost use 222 runs with token records, while its ISR uses all 225 scored runs. Other models use their available records.

task distribution and these frozen prices, GLM-5.3-Flash, DeepSeek-V4.1-Flash, and GPT-6.1-Sol form the cost-success Pareto frontier, offering choices at different cost levels. The additional resources spent on longer trajectories constitute progress only when they satisfy research requirements. These results suggest that the value of resource use lies in helping models check constraints, complete validation, and resolve remaining errors. Cross-model comparisons reveal this efficiency difference, while the benefit of increasing a single model's budget requires a separate experiment.

4.6 Judge Agreement and Result Sensitivity

4.6.1 RQ5: How does changing the LLM judge affect binary verdicts and model comparisons?

DeepSeek verdict / Qwen verdict	Pass	Fail	Total
Pass	8,043	67	8,110
Fail	648	1,116	1,764
Total	8,691	1,183	9,874

Table 12 DeepSeek and Qwen binary verdicts on matched inputs and checklist items. Starting from 657 matched scoring pairs, matching checklist-item IDs and content hashes yields 9,874 paired binary verdicts. The judges are `deepseek-v4.1-flash-1017` and `qwen3.8-flash`. These results cover a subset of tasks. Darker blue indicates larger joint counts; marginal totals use neutral shading.

On the same 20-task common-coverage subset, we compare cross-judge ranking agreement for ISR and CSR. The mean Spearman correlation across the three judge pairs is 0.9524 for CSR, exceeding 0.9028 for ISR, so Figure 6 uses CSR ranks. Correlations preserve ISR ties rather than arbitrarily ordering models with equal scores.

Table 12 contains 9,159 agreeing binary verdicts, giving 92.76% agreement and Cohen's κ of 0.717 after adjustment for chance agreement using marginal pass rates. Of 715 disagreements, 648, or 90.63%, are DeepSeek fail and Qwen pass, indicating a directional difference in acceptance. In Figure 6, GPT-6.1-Sol ranks first under all three scoring sources. Spearman rank correlation is 0.9643 between DeepSeek and Qwen, and 0.9286 and 0.9643 between the historical reference and the two judges, respectively. Rank changes mainly involve adjacent models.

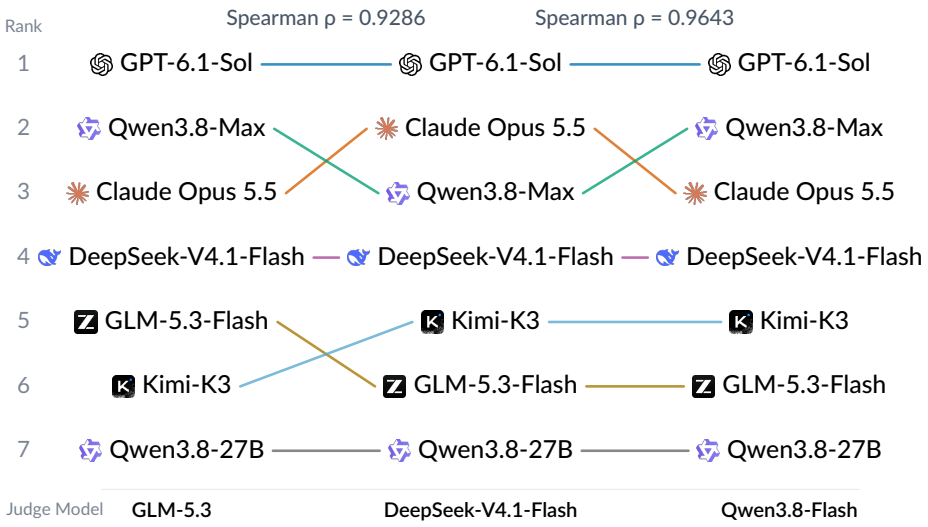

Figure 6 Cross-judge CSR ranks on the 20-task common-coverage subset. The columns show the historical default GLM reference, DeepSeek judge, and Qwen judge. Nodes use model logos and short names, and the same actor model is linked by the same color. More stable ranks produce more horizontal lines. The subset contains 147 actor trajectories, 21 per model. Scores first average included rounds within each task and then weight tasks equally. Current DeepSeek and Qwen inputs match, while historical default scores are a reference only. Tasks have different numbers of included rounds, and the subset retains the stage package's source selection and policy of deferring low-scoring sources.

Changing the judge preserves most binary verdicts and the model ranking. Disagreements concentrate in one direction and mainly reorder adjacent models. Independent human annotations are the next step for attributing disagreements and completing the evidence on judge reliability.

4.7 Case Study

4.7.1 Case question: Does resuming training mean that a research repair is complete?

Evaluation aspect	Successful execution, 20/20	Execution with unmet items, 16/20
Training progress	Sustained nonzero loss and parameter updates	Also sustained training, nonzero loss, and parameter updates
O7, fusion coordinates	Projection indexes agree with the actual feature grid	Inconsistent coordinates clip some samples to boundary columns
O9, overlap labels	Separate 2D and 3D overlap checks retain positive, negative, and ignored regions	A single projection ratio replaces independent 2D and 3D overlap tests
O11, coarse matching loss	L2 distance, overlap weights, and specified settings are retained	Squared distance replaces L2 without corresponding boundary-parameter changes, and weights are altered
O12, fine matching loss	Bidirectional loss and the coarse/fine loss combination are retained	The changed distance and weight expressions are reused
Training and delivery items	Archived review accepts checkpoints, limited validation, and the report	The other 16 items pass, including training recovery, checkpoints, validation, and the report

Table 13 Two executions of the CylindFormer training-repair task. The task ID is `cylindformer-training-repair`. Both executions start at the same repair boundary and must resume training according to the reference paper. Scores follow the archived review. Appendix E contains the complete inputs, restored workspace, and 20-item checklist.

In Table 13, both executions resume image-to-point-cloud registration training with nonzero loss and parameter updates, but pass 20/20 and 16/20 checklist items, respectively. The execution with unmet items has four methodological deviations in projection and feature-grid coordinates, 2D/3D overlap labels, coarse matching loss, and fine matching loss. The successful execution resumes training while preserving the paper's method. Figure 7 connects their shared execution progress to their different methodological validity.

Program recovery and faithful repair of the research method are distinct outcomes that require separate verification. Nonzero loss and training progress provide evidence of execution recovery, while checking coordinates, labels, and objectives establishes whether the repair preserves methodological meaning. This comparison grounds RQ2's correctness bottleneck in specific failures and shows how restored workspaces and retained research constraints jointly support evaluation of requests within ongoing research.

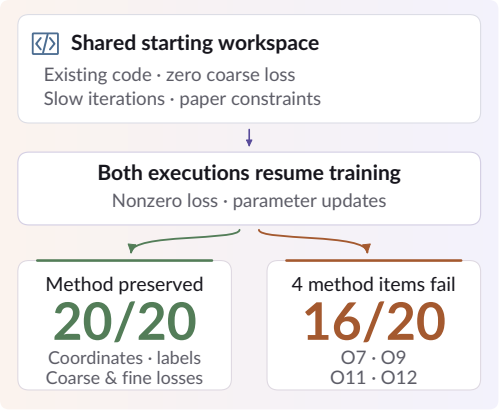

Figure 7 Divergent executions of the same training-repair request. The comparison concerns method consistency during recovery. Differences in final registration accuracy still require separate measurement.

5 Conclusion

ARE-bench evaluates the next research action inside an evolving project. Every task starts from the replayed mid-session workspace with unfinished files and failed attempts preserved, while its task brief gives the actor a self-contained and auditable specification. The ISR and CSR checklist metrics expose strict completion and partial compliance, while separate cost metrics expose efficiency trade-offs. Replay checks and anonymization make the benchmark suitable for comparing research agents without treating an original trajectory as a gold answer. The release protocol supports filling the result schema as model runs become available while keeping the benchmark definition fixed.

Limitations

The source sessions may overrepresent particular projects, tools, or researcher populations. Cut-point sampling from existing sessions does not guarantee coverage of all research workflows. Replay and dependency recovery can lose context that was available to the original researcher, and anonymization can remove clues needed for faithful interpretation. Rubric writing remains a human judgment process, and a checklist evaluator can still miss evidence or over-credit a plausible claim. GPT-6.1-Sol runs with the Codex harness and the other models with Claude Code, so cross-model differences reflect the model and harness together. The benchmark therefore publishes provenance records for audit, but its scores should be interpreted as evidence about the released task distribution rather than as a universal measure of scientific ability.

Ethics Statement

Session data must be collected under an appropriate authorization or consent arrangement. Before release, identifiers, private paths, credentials, unpublished data, and organization-specific metadata are removed or replaced with stable aliases. The provenance package is access-controlled and is not included in actor inputs. We release only the minimum workspace content needed to reproduce each task and document residual disclosure risks, withdrawal procedures, and the scope of permitted reuse.

Appendix

A Task and Rubric Example

A task contains an instruction with Background, Requirements, and Environment sections. A representative requirements block asks the agent to update an existing analysis, save the result to an absolute path, preserve a named preprocessing choice, and report the command used to reproduce the artifact. Its rubric compiles into one binary checklist item for each deliverable, required transformation, preservation constraint, and reproducibility condition. The actor never sees the rubric text.

B Research Requirement Taxonomy

We use a research requirement taxonomy to characterize the kinds of work researchers request from agents. The taxonomy serves as a reference for benchmark construction and organizes research work into 33 requirements and 207 sub-requirements, grouped into seven research stages. Requirements describe broader areas of work, such as literature review, experiment design, and code testing. Sub-requirements specify individual user needs using action–object descriptions, such as searching for relevant literature or writing unit tests. Table 14 presents the portion of this taxonomy represented in our benchmark: 150 sub-requirements covering all 33 requirements.

The table organizes research requirements into three levels: research stage, requirement, and sub-requirement. Two additional attributes characterize each sub-requirement. *Capability primitives* are drawn from the cognitive process dimension of the revised Bloom's taxonomy (Anderson and Krathwohl, 2001; Krathwohl, 2002) and describe the cognitive operations involved, such as differentiating under Analyze, checking under Evaluate, and planning under Create. *Work object* identifies the material these operations act on, such as code, data, or research documents.

Each checklist item is associated with one sub-requirement, which identifies its parent requirement and research stage. A task can therefore involve several requirements and stages through its checklist items.

Table 14 Research requirement taxonomy represented in the benchmark. Each row describes a sub-requirement, its parent requirement and research stage, and the associated capability primitives and work object.

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Survey	Literature review	Search for relevant literature	Apply: implement; Analyze: differentiate	Evidence
		Verify bibliographic information	Evaluate: critique	Evidence
		Extract key points from literature	Analyze: differentiate; Understand: summarize	Evidence
		Verify literature content	Evaluate: check	Evidence
		Assess literature quality	Evaluate: critique	Evidence
		Establish evidence traceability	Analyze: organize; Evaluate: check	Evidence
		Synthesize external evidence	Understand: summarize; Understand: compare; Analyze: organize	Evidence
		Revise research questions	Apply: implement; Understand: explain; Analyze: organize	Questions
		Assess method applicability	Evaluate: critique	Methods
		Write a literature review report	Understand: summarize; Understand: explain; Analyze: organize	Documents

Continued on next page.

Table 14 Research requirement taxonomy (continued).

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Survey (cont.)	Problem definition	Generate candidate research questions	Create: generate; Understand: explain	Questions
		Define core concepts	Understand: explain; Understand: exemplify; Analyze: organize	Questions
		Define research variables	Understand: explain; Analyze: organize; Create: plan	Questions
		Identify problem assumptions	Understand: infer; Analyze: attribute; Analyze: differentiate	Questions
		Formulate research hypotheses	Understand: infer; Create: generate	Questions
		Assess research novelty	Understand: compare; Analyze: differentiate; Evaluate: critique	Questions
		Assess research feasibility	Analyze: differentiate; Evaluate: critique	Questions
		Identify research ethics risks	Understand: classify; Evaluate: critique	Questions
		Identify research constraints	Remember: recall; Understand: classify; Analyze: differentiate	Projects
		Write a problem definition document	Understand: summarize; Understand: explain; Analyze: organize	Documents
Design	Resource estimation	Identify resource requirements	Understand: classify; Analyze: differentiate; Analyze: organize	Plans
		Estimate data needs and costs	Apply: implement; Apply: execute; Understand: infer	Plans
		Estimate compute needs and costs	Apply: implement; Apply: execute; Understand: infer	Plans
		Estimate the research timeline	Apply: implement; Understand: infer	Plans
		Plan research execution	Create: plan	Plans
	Experiment design	Define experimental subjects and units of analysis	Understand: explain; Analyze: organize	Plans
		Operationalize experimental variables	Understand: explain; Analyze: organize; Create: plan	Plans
		Design experimental controls	Analyze: differentiate; Create: plan	Plans
		Design ablation experiments	Analyze: differentiate; Create: plan	Plans
		Design repeated experiments	Analyze: organize; Create: plan	Plans
		Define evaluation metrics	Analyze: organize; Create: plan	Plans
		Define research acceptance criteria	Analyze: organize; Create: plan	Plans
		Design sensitivity analyses	Understand: compare; Analyze: differentiate; Create: plan	Plans
		Design experimental procedures	Analyze: organize; Create: plan	Plans
		Design experiment recording standards	Analyze: organize; Create: plan	Plans
		Design experiment stopping rules	Analyze: organize; Evaluate: check; Create: plan	Plans
		Identify threats to experimental validity	Understand: classify; Understand: infer; Analyze: differentiate	Plans

Continued on next page.

Table 14 Research requirement taxonomy (continued).

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Design (cont.)	Experiment design (cont.)	Assess experiment design feasibility	Analyze: differentiate; Evaluate: critique	Plans
		Verify experiment designs	Evaluate: check; Analyze: organize	Plans
		Write an experiment design document	Understand: summarize; Understand: explain; Analyze: organize	Documents
	Method evaluation	Select experiment evaluation metrics	Analyze: differentiate; Evaluate: critique	Plans
		Design method evaluation protocols	Create: plan	Plans
		Implement experiment evaluation metrics	Apply: implement	Code
		Configure controlled experiments	Analyze: organize; Apply: implement	Experiments
		Configure cross-model experiments	Analyze: organize; Apply: implement	Experiments
Data & Env.	Environment setup	Install project dependencies	Apply: implement	Environment
		Resolve dependency conflicts	Apply: implement; Evaluate: check	Environment
		Verify the computing environment	Evaluate: check	Environment
	Resource acquisition	Acquire model assets	Apply: execute	Assets
		Acquire research data	Apply: execute	Data
		Verify research data identity	Evaluate: check	Data
		Acquire compute resources	Evaluate: critique; Apply: implement	Environment
	Data analysis	Analyze data structures	Analyze: organize	Data
		Assess data quality	Evaluate: critique	Data
		Assess statistical power	Evaluate: critique	Plans
	Data preprocessing	Convert data formats	Understand: explain; Apply: execute	Data
		Handle missing values	Apply: implement; Evaluate: check	Data
		Handle outliers	Apply: implement; Evaluate: check	Data
		Create training, validation, and test splits	Analyze: organize	Data
	Feature engineering	Extract data features	Analyze: differentiate	Data
		Compare features across groups	Understand: compare	Data
Code	Repository comprehension	Analyze code implementations	Analyze: organize	Code
		Verify code modules	Evaluate: check	Code
		Verify code completeness	Evaluate: check	Code

Continued on next page.

Table 14 Research requirement taxonomy (continued).

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Code (cont.)	Code framework design	Derive research methods	Create: plan	Methods
		Adapt research methods	Apply: implement	Methods
		Implement experiment scaffolding	Apply: implement	Code
		Design model architectures	Create: plan	Methods
		Design algorithm mechanisms	Create: plan	Methods
		Design loss functions	Create: plan	Methods
		Design code implementations and boundaries	Create: plan	Code
		Add experiment instrumentation logs	Apply: implement	Code
	Algorithm implementation	Implement algorithm modules	Apply: implement	Code
		Implement algorithm variants	Apply: implement	Code
		Integrate open-source components	Analyze: organize; Apply: implement	Code
		Verify algorithm implementation equivalence	Evaluate: check	Code
	Code change management	Manage code version changes	Apply: execute; Analyze: organize	Code
		Modify code across modules	Analyze: organize; Apply: implement	Code
		Refactor code	Analyze: organize; Apply: implement	Code
		Verify code changes	Evaluate: check	Code
	Code testing	Write unit tests	Analyze: differentiate; Apply: implement	Code
		Write integration tests	Analyze: organize; Apply: implement	Code
		Write end-to-end tests	Analyze: organize; Apply: implement	Code
		Run regression tests	Apply: execute; Evaluate: check	Code
Execution	Experiment configuration	Summarize experiment configurations	Understand: summarize; Analyze: organize	Experiments
		Summarize training code	Understand: summarize	Code
		Configure experiment runs	Apply: implement	Experiments
	Execution scheduling	Run experiments	Apply: execute	Experiments
	Run management and monitoring	Resume experiment runs	Remember: recall; Apply: implement; Evaluate: check	Experiments
		Organize model checkpoints	Analyze: organize; Apply: execute	Experiments
		Record experiment metrics	Apply: execute	Experiments
		Record experiment logs	Apply: execute	Experiments
		Organize experiment artifacts	Analyze: organize	Experiments

Continued on next page.

Table 14 Research requirement taxonomy (continued).

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Execution (cont.)	Experiment diagnosis	Describe training trends	Understand: summarize; Understand: infer	Experiments
		Verify experiment run status	Evaluate: check	Experiments
		Compare performance across data slices	Understand: compare	Results
		Compare effects of experiment settings	Understand: compare	Results
		Diagnose method failures	Analyze: differentiate; Understand: infer	Experiments
		Diagnose experiment anomalies	Analyze: differentiate; Understand: infer	Experiments
		Diagnose experiment performance bottlenecks	Analyze: differentiate; Understand: infer; Understand: compare	Experiments
		Diagnose code defects	Analyze: differentiate; Understand: infer	Code
		Recommend when to stop experiments	Evaluate: critique	Experiments
	Experiment repair	Fix code defects	Apply: implement; Evaluate: check	Code
		Verify experiment retry results	Evaluate: check	Experiments
		Assess experiment repair plans	Evaluate: critique	Experiments
	Experiment tuning	Identify data contamination	Understand: classify; Analyze: differentiate	Data
		Search hyperparameter combinations	Apply: implement; Evaluate: critique	Experiments
		Run ablation experiments	Apply: implement; Evaluate: check	Experiments
	Experiment verification	Run sensitivity analyses	Apply: implement; Evaluate: check	Experiments
		Verify experiment reproducibility	Evaluate: check	Experiments
Analysis	Result statistics	Compute descriptive statistics	Apply: execute	Results
		Compute statistical test results	Apply: execute	Results
		Compute effect sizes	Apply: execute	Results
		Estimate result uncertainty	Apply: execute; Understand: infer	Results
	Result comparison	Evaluate algorithm variants	Evaluate: critique	Results
		Compare research methods	Understand: compare	Results
	Result interpretation	Explain model decision mechanisms	Understand: explain; Understand: infer	Results
		Verify research hypotheses	Evaluate: check	Results
		Verify the limits of research conclusions	Evaluate: check; Evaluate: critique	Results
		Report negative results	Understand: summarize; Evaluate: critique	Results
		Define the scope of research applicability	Analyze: organize; Create: plan	Results
	Result visualization	Generate result tables	Analyze: organize; Apply: execute	Results
		Visualize research results	Understand: explain; Apply: implement	Results

Continued on next page.

Table 14 Research requirement taxonomy (continued).

Research stage	Requirement	Sub-requirement	Capability primitives	Work object
Writing	Paper structure planning	Verify paper argumentation	Evaluate: check	Documents
		Plan paper outlines	Create: plan	Documents
	Related work writing	Write reference lists	Apply: implement	Documents
		Format citations	Apply: execute	Documents
	Methods and experiments writing	Write methods sections	Understand: explain; Analyze: organize	Documents
		Write experiment setup sections	Understand: explain; Analyze: organize	Documents
		Verify consistency of method descriptions	Evaluate: check	Documents
		Verify consistency between notation and implementation	Evaluate: check	Documents
		Write result analyses	Understand: summarize; Analyze: organize	Documents
	Figure and table preparation	Generate paper tables	Analyze: organize; Apply: execute	Documents
		Annotate paper figures and tables	Understand: explain	Documents
	Paper revision	Revise academic paragraphs	Apply: implement	Documents
		Integrate external evidence	Analyze: organize; Apply: implement	Documents
		Adjust paper length	Apply: implement	Documents
		Add factual content	Analyze: organize; Apply: implement	Documents
		Merge paper revisions	Analyze: organize; Apply: implement	Documents
		Align paper sections	Analyze: organize	Documents
		Verify paper cross-references	Evaluate: check	Documents
		Verify consistency of result statements	Evaluate: check	Documents
		Verify research evidence chains	Evaluate: check	Documents
	Reviewer response	Revise content addressing review comments	Apply: implement	Documents
		Write reviewer responses	Understand: explain; Apply: implement	Documents
	Research release	Verify submission formatting	Evaluate: critique	Documents
		Organize submission materials	Analyze: organize	Assets
		Write research reproduction documentation	Understand: explain; Analyze: organize	Documents
		Package research assets	Analyze: organize; Apply: execute	Assets
		Verify research assets	Evaluate: check; Evaluate: critique	Assets

C Starting Workspace State by Research Stage

Table 15 lists the workspace content that a task cut at each research stage typically carries over from earlier turns. Every listed state is replayed from the source session rather than authored for the task, so a task cut inside Execution inherits run scripts and logs rather than a clean environment, and a task cut inside Writing inherits an existing draft rather than a blank document.

Research stage	Representative starting workspace state W_c
Survey	Partial reading notes and a candidate research question, with no experiment code yet written
Design	An existing problem statement and prior baseline results that the new design must extend or contrast
Data & Env.	A partially installed environment and raw or partially downloaded data left from an earlier setup attempt
Code	A repository with a partial implementation, stub functions, and tests that do not yet pass
Execution	Configured run scripts together with logs or checkpoints from an interrupted or failed run
Analysis	Raw experiment outputs and an unfinished or outdated analysis script
Writing	A paper draft with existing sections, figures, or reviewer comments that the current request must incorporate

Table 15 Representative starting workspace state W_c for a task cut at each research stage. Every state is replayed from the source session rather than authored for the task, so earlier work is present regardless of stage.

D Cut-Point Selection, Replay, and Anonymization

The builder records the source episode, user-turn index, file-event range, process-event range, dependency snapshot, and validation hash. Replay is accepted only when the reconstructed tree matches the audited state for all files declared relevant to the task. Paths, usernames, hostnames, keys, and dataset identifiers are replaced by deterministic aliases. A second reviewer checks that the replacement preserves references used in the instruction and does not leak the source project.

E Case Study: Repairing an Interrupted Training Experiment

This case follows a request to repair an unfinished image–point-cloud registration experiment, which matches photographs to a three-dimensional scene to estimate camera pose. Training has started, but the coarse-matching loss stays at zero and each iteration takes about 28 seconds. The repair must preserve the supplied paper's method.

From a conversation to a task boundary. Figure 8 shows the real session. We cut after the fourth user turn, the repair request. Earlier decisions supply Background; the implementation at this boundary supplies the repair starting point. Figure 9 separates the complete task instruction from the restored workspace and prepared execution resources. Assessment material is shown separately from actor input.

A research session, cut at a repair request

User Please allocate an RTX PRO 6000 GPU, configure the environment, and download the full dataset. Then reproduce the main experiments in the paper I uploaded, including the results on its main datasets. The method must match the paper's description.

Assistant The repository contains only a README; the source code has not been released. We can either (1) implement CylindFormer from scratch, following the paper's detailed method and training settings, or (2) wait for the official code release.

User Go ahead with the first option.

Assistant I will implement CylindFormer from the paper. First, I will install the dependencies, then build the complete code.

User What is the current status?

Assistant Training on 7-Scenes is running, but Lc remains zero because the overlap matrix has not been computed. Each iteration takes about 28 seconds. Should we fix the coarse-matching loss and speed problems, then continue training?

User Fix the issues, then continue training.

★ **Task cut: after the repair request**

Later original-session turns Excluded from task construction.

Assistant I will diagnose the speed bottleneck first, then fix the issues one by one. ... DINOv2 takes only 9 ms; the bottleneck is elsewhere. I will continue investigating.

User What is the current status?

Assistant 7-Scenes has been downloaded, but most per-scene sequence ZIPs still need to be extracted. The revised code is ready but has not yet been tested. Next, we need to finish extraction and run a speed test before restarting training. Shall I continue extracting the data?

User Yes, go ahead.

Figure 8 Session with all six user turns and brief assistant excerpts. The star marks the cut after the repair request. Later turns appear below the boundary and are excluded from task construction.

Actor input at the repair boundary

1. Task instruction `instruction.md`

Background

You are helping me reproduce CylindFormer's image–point-cloud registration experiments. The method must follow the uploaded paper and use the full dataset. When checked, the official repository contained only a README, so we chose to implement the method from scratch based on the paper. Training on 7-Scenes has started, but the logs show that the coarse-matching loss Lc stays at zero and a full iteration takes about 28 seconds, making progress impractically slow. Earlier explanations for the slowdown have not yet been verified. A GPU machine is configured, but the original environment has no usable checkpoint from completed training. The immediate focus is repairing and resuming the ongoing 7-Scenes training.

Requirements

Repair the current training problems, then continue training while keeping the method consistent with `/workspace/materials/paper.pdf`. The existing implementation is the starting point for repair and needs to be checked. The current task is to resume this training; completing the paper's main experiments on every dataset is not required in this round.

Deliver a repaired project that can be used for further work to `/workspace/output/`. Provide an index locating the source code, configurations, reproduction commands, runtime logs, model checkpoints, validation-pair indices, and per-pair predictions and evaluation records. Explain the repairs and the current training and validation progress. You may choose the internal filenames and organization. Preserve the original paper, raw data, and historical logs; code may be modified and new artifacts added. Do not read prohibited hidden assessment materials or reference answers, or modify the scoring program.

Environment

Execution. The working directory is `/workspace`. Internet access is available, and missing tools may be installed with `pip install`. The evaluation infrastructure supplies one CUDA GPU with at least 80 GiB of memory and CUDA 12.8 compatibility.

Code and data. The code is in the workspace project. The original project and data paths are mapped to this project and `/opt/data/7scenes`, respectively. The full raw 7-Scenes dataset has been restored from the public sources used in the session, including all seven scenes and official split files. Its file inventory is not guaranteed to match the incompletely exported historical runtime file by file. The paper's exact internal pair IDs are not provided.

Dependencies. Python 3.12 provides PyTorch 2.8.0/CUDA 12.8, torchvision 0.23.0, torchaudio 2.8.0, numpy, scipy, pandas, opencv-python-headless, open3d, h5py, einops, pytorch-lightning, tensorboard, pyyaml, tqdm, matplotlib, imageio, scikit-learn, timm, transformers, lxml, openpyxl, requests, pdfminer.six, pdfplumber, and pikepdf. Fixed versions are recorded in `/workspace/materials/environment.json`. Historical package versions were not fully recorded; these are compatible fixed dependencies rather than an exact restoration of every original version.

Pretrained resources. DINOv2-S/14 and ResNet-34 pretrained weights are cached in `/opt/torch/hub`. The unknown historical DINOv2 source-cache version is replaced with a fixed revision documented in the environment description. Additional dependencies may be installed as needed.

2. Restored workspace

`/workspace/`: 18 project files and four materials files—unfinished code, supplied paper and historical training log. Full directory view follows.

3. Prepared execution resources

Full 7-Scenes data, DINOv2-S/14 and ResNet-34 weights, and installed dependencies.

1 GPU; 8 CPUs; 64 GiB RAM; 160 GiB storage; 4 h agent limit.

Assessment Separate from actor input. `tests/rubric.json`: 20 outcome checks, listed in full in the outcome rubric table. Detailed method checks also use the supplied paper.

Figure 9 Task inputs and assessment. Background, Requirements, and Environment form `instruction.md`; the other two input regions describe the provided files and execution resources. Paths are normalized. The complete outcome rubric appears in Table 16.

Constructing the task input. Recorded file operations restore the unfinished data-loading, model, loss, and training code. The paper defines the method to preserve; the diagnostic log records the observed failure. Public data and pretrained weights are restored separately, with compatible dependencies. Figure 10 shows every recovered project and materials file and the prepared resource directories.

What the agent receives at the boundary

Restored project state and separately prepared resources support the repair.

Restored workspace

18 code files + 4 method/diagnostic files

```
/workspace/
├── CylindFormer/
│   ├── config.py
│   ├── datasets/
│   │   ├── __init__.py
│   │   ├── prepare_data.py
│   │   ├── rgbd_scenes.py
│   │   └── seven_scenes_loader.py
│   ├── evaluate.py
│   ├── extract_datasets.sh
│   ├── losses/
│   │   ├── __init__.py
│   │   └── circle_loss.py
│   ├── models/
│   │   ├── __init__.py
│   │   ├── backbones.py
│   │   ├── cylindformer.py
│   │   └── cylindrical_transformer.py
│   ├── run_experiment.sh
│   ├── test_pipeline.py
│   ├── train.py
│   ├── train_7scenes.py
│   └── utils/
│       └── __init__.py
└── materials/
    ├── environment.json
    ├── paper.md
    ├── paper.pdf
    └── training_before_repair.log
```

Prepared resources

Data and pretrained models for execution

Dataset

```
/opt/data/7scenes/
├── chess/
├── fire/
├── heads/
├── office/
├── pumpkin/
├── redkitchen/
├── stairs/
└── restored_input_inventory.json
```

Each scene provides `TrainSplit.txt` and `TestSplit.txt`, with color, depth and pose files under `seq-XX/` (`frame-XXXXXX.*`).

Pretrained models

```
/opt/torch/hub/
├── facebookresearch_dinov2_main/
├── checkpoints/
│   ├── DINOv2-S/14 weights
│   └── ResNet-34 weights
└── trusted_list
```

Cached model code and weights; checkpoint entries identify the models.

Expected outputs `/workspace/output/`

Agent-generated source, configuration, reproduction commands, logs, checkpoints, validation pairs, predictions and an index.

Figure 10 Directory view of the starting inputs: all 18 project files and four materials files, plus separately prepared data and model resources. Repeated raw-data samples use a filename pattern; weight entries name their models. The output directory is the requested handoff destination.

Comparing two executions. Both executions continue training, but progress alone does not establish a faithful repair. Figure 11 connects the repairs and recorded behavior to the archived outcome assessments. The successful execution passes all 20 checks. The failed execution passes 16, including training recovery, but fails four method checks involving feature coordinates, overlap labels, and loss computation.

The same repair request, two different outcomes

Both executions resume training. The method checks explain the different assessments.

Shared starting point

Existing implementation + supplied paper + zero coarse loss + slow iterations

↙ ↘

Successful execution: 20/20	**Failed execution: 16/20**
Repair Uses a stable reformulation of the specified loss, removes expensive matching loops and retains pretrained image features.	**Repair** Changes parts of the implementation and resumes training, but introduces or retains four method deviations.
Observed training Sustained nonzero coarse loss and parameter updates.	**Observed training** Successive epochs, nonzero losses and parameter updates.
Method checks **O7:** fusion uses consistent feature coordinates. **O9:** separate 2D and 3D overlap tests retain positive, negative and ignored pairs. **O11:** L2 distance, overlap weighting and loss settings are preserved. **O12:** bidirectional fine circle loss and the coarse/fine combination are retained.	**Failed method checks** **O7:** inconsistent projection/feature-grid coordinates clamp samples to boundary columns. **O9:** one depth-consistent projection ratio replaces the separate 2D/3D overlap tests. **O11:** squared distance replaces L2 distance with unchanged margins; weighting also changes. **O12:** the fine loss shares the changed distance and weighting expressions.
Handoff The archived assessment credits checkpoint usability, limited validation and a report of remaining work. A separate debug test saves, reloads and continues from a checkpoint.	**Handoff** The archived assessment credits the other 16 checks, including training recovery, checkpoint, validation and reporting checks.
Progress with method consistency	**Progress with method deviations**

Reading the comparison Outcome scores are archived assessments, not new experiments. The failed run demonstrates method deviations despite training progress; these records do not measure a resulting drop in final registration accuracy.

Figure 11 Two independent executions of the same task. Green and red backgrounds mark the decisive findings. The scores summarize archived outcome assessments; the comparison establishes method consistency or deviation, rather than a measured change in final registration accuracy.

What the comparison shows. Restoring the existing code makes the short repair request a continuation of an unfinished experiment. The successful case shows that a usable repair can preserve its method while moving training forward. The failed case shows why observing training is insufficient: the experiment runs, but parts of the method no longer match the paper. Retaining both the project state and the research constraints makes this distinction assessable.

Complete outcome rubric. The outcome rubric checks a usable handoff, data and method implementation, actual training, and evidence-linked validation. Its source labels distinguish explicit requirements from checks inferred using the paper, prior context, or ordinary research and usability expectations. Detailed method checks draw on the supplied paper and are not all individually enumerated in Requirements.

Table 16 Complete outcome rubric. Each item is a binary check; source labels identify explicit or inferred grounding.

ID	Criterion	Source
O1	Deliver a reusable repaired project in /workspace/output/, with an index locating source code, configuration, and reproduction entry points. Internal filenames and project organization are unrestricted.	Explicit; Requirements
O2	Construct image–point-cloud training pairs with real relative poses and filter them using the paper's at-least-50% overlap rule. Pairing only an image with its own back-projected depth, or declaring min_overlap without computing and applying it, does not meet this check. Verifiable equivalent pairing is accepted; unavailable paper-internal pair IDs need not be restored.	Inferred; paper
O3	Separate training and evaluation according to the official 7-Scenes sequence splits. Validation must not train the model or select it using test performance. When the paper's internal validation pairs are unavailable, an independent validation set constructed from training sequences is accepted if its scope is explained.	Inferred; paper
O4	The point-cloud backbone must aggregate neighborhood features using kernel points and produce multiscale coarse and fine features, consistent with the paper's three-stage indoor KPConv-FPN. A per-point MLP merely named KPConv does not meet this check. Correct open-source components or equivalent reimplementations are accepted.	Inferred; paper
O5	Retain the paper's ResNet-FPN and DINOv2 image-feature paths. DINOv2 must load pretrained weights and contribute to model output; disabling it or silently replacing it with image average pooling does not meet this check. Freezing or caching unchanged pretrained features is accepted.	Inferred; prior context and paper
O6	Cylindrical self-attention must enhance the key with local fusion information, implementing the paper's Q(K+L)-transpose attention semantics. Disconnected local information or plain self-attention does not meet this check. Tensor implementations with the same mathematical meaning are accepted.	Inferred; paper
O7	Cross-modal fusion must pass information through cylindrical-projection local aggregation, adaptive fusion with global features, and cross-attention. Coordinates and indices must match the actual feature layout. An unused fusion module does not meet this check; perspective projection for supervision does not replace cylindrical projection inside the model.	Inferred; paper
O8	Fine matching must use local image and point-cloud patches selected by coarse correspondences and generate pixel–point matches using mutual top-k. Ignoring coarse matches and directly matching the full image against all points does not preserve the coarse-to-fine method.	Inferred; paper
O9	Coarse positive and negative labels must use ground-truth geometry to compute both 2D and 3D patch overlap. A pair is positive when both exceed 30%, negative when both are below 20%; the intermediate region must not be treated as negative. Learned feature similarity or grids near a single projected point do not replace patch overlap.	Inferred; paper
O10	Fine labels must follow real pixel–point geometry: positive when the 3D distance is below 3.75 cm and 2D distance below 8 px; negative when the 3D distance exceeds 10 cm or 2D distance exceeds 12 px. Coordinate transforms, camera intrinsics, and pixel scales must agree with the data.	Inferred; paper
O11	Coarse supervision must implement bidirectional overlap-aware circle loss with feature L2 distance, overlap weighting, gamma=40, and positive/negative margins 0.1/1.4. Numerically stable equivalent expressions are accepted. Similarity mean-squared error, constants, or a custom anti-collapse objective cannot substitute merely to produce nonzero loss. Where the paper's formula and prose are ambiguous, a self-consistent interpretation supported by the materials is accepted.	Inferred; paper

Continued on next page

Table 16 continued

ID	Criterion	Source
O12	Fine supervision must implement bidirectional standard circle loss with gamma=24 and positive/negative margins 0.1/1.4, combined with coarse supervision using beta=1. Numerically stable rewriting is accepted; replacing the training objective is not.	Inferred; paper
O13	A training batch with valid geometric supervision must propagate loss gradients to the trainable descriptors and matching heads actually used at inference, and update those parameters. Adding requires_grad to an unrelated constant, bypassing an untrained matching head, or detaching the entire path does not meet this check. Individual batches without valid positives or negatives may be explicitly skipped; loss need not be positive at every step.	Inferred; research expectations
O14	Remove the performance bottleneck blocking continued training. Measured full-step time must improve substantially under comparable hardware, input size, and training responsibilities, with evidence identifying the removed overhead. Omitting supervision, reducing task scope, or timing only a forward pass is insufficient. No fixed speed or improvement factor from later historical work is required.	Inferred; prior context
O15	Resume repaired 7-Scenes training on the GPU, process successive traceable real pairs, update parameters, and retain progress from this run. Code alone, a launch command, or one random-tensor forward pass is insufficient. This checks training recovery; completing all 30 epochs or the RGB-D Scenes V2 main experiment is not required.	Inferred; prior context
O16	The delivered checkpoint from this training must be compatible with the repaired model, restore learned parameters, and support continued training. Silently falling back to random initialization after a load failure is insufficient. Any appropriate state-saving format is accepted; no framework or internal filename is mandated.	Inferred; usability expectations
O17	The repaired inference path must generate pixel–point matches from learned descriptors and use PnP-RANSAC to estimate the transform from point-cloud coordinates to the image camera frame. Ground-truth poses may supervise or evaluate the model, but must not replace inference estimates.	Inferred; paper
O18	Compute registration metrics with the paper's meanings: IR is the fraction of correspondences with 3D distance below 0.05 m after the ground-truth transform; FMR is the fraction of pairs with IR above 0.1; RR is the fraction with predicted/ground-truth transform RMSE below 0.1 m. Keep failed pairs in denominators. Placeholders or validation loss alone are insufficient. Small-scale recovery validation is accepted but must not be called full paper-test-set reproduction.	Inferred; paper
O19	Deliver this run's logs, checkpoint, validation-pair index, per-pair predictions and evaluation records, and reproduction commands. Records must link validation results to the current model and actual data. Any readable organization is accepted; checkpoints may be binary.	Explicit; Requirements
O20	The repair report must describe completed training and validation stages consistently with delivered evidence, distinguishing program recovery, method consistency, and full main-experiment reproduction. Nonzero loss or completion of a few epochs alone does not establish the paper's registration results.	Inferred; research expectations

F Judge Prompt

The evaluator prompt states the task requirements, the checklist item, the evidence available to inspect, and the binary decision rule. It requires a short evidence citation for every decision and forbids inferring an unobserved result from a claimed completion.

G Full Results and Hyperparameters

The machine-readable release stores one record per task and rollout with the model identifier, score vector, delivery manifest status, wall-clock time, token counts, tool calls, and steps. Aggregate tables are generated by grouping these records by model, checklist dimension, and research stage. The run manifest is versioned with the task manifest so that a later model evaluation cannot silently alter the benchmark denominator.

Setting	Value
Actor harness	Codex for GPT-6.1-Sol, Claude Code for other models
Harness version	Codex 0.118.0 or Claude Code 2.1.283 on RJob, varying by run on E2B
Judge	GLM-5.3 (GLM-5.3-Flash for a small subset of runs), one call per checklist item
Execution backend	E2B for 920 CPU runs, RJob/Docker for the other 655 runs
Harbor version	0.23.0 in the 650 readable run protocols
Rollouts per task	3
Time limit	1.5 × task timeout for 920 CPU runs, task-specific overrides otherwise
Infrastructure retry	up to 2, actor timeout not retried

Table 17 Run configuration.

Metric	Value
Tasks	75
Checklist items	1,173
Avg. checklist items per task	15.6
Research stages covered	7
Avg. earlier requests in Background	10.3
Avg. workspace files at cut point	41.1
Avg. deliverables per task	2.5
Avg. instruction length (characters)	1,677

Table 18 Overall statistics of ARE-bench.

References

Lorin W. Anderson and David R. Krathwohl, editors. 2001. *A Taxonomy for Learning, Teaching, and Assessing: A Revision of Bloom's Taxonomy of Educational Objectives*, complete edition. Longman, New York.

Ben Bogin, Kejuan Yang, Shashank Gupta, Kyle Richardson, Erin Bransom, Peter Clark, Ashish Sabharwal, and Tushar Khot. 2024. SUPER: Evaluating agents on setting up and executing tasks from research repositories. *arXiv preprint arXiv:2409.07440*.

Jonathan Bragg, Mike D'Arcy, Nishant Balepur, Dan Bareket, Bhavana Dalvi, Sergey Feldman, Dany Haddad, Jena D. Hwang, Peter Jansen, Varsha Kishore, Bodhisattwa Prasad Majumder, Aakanksha Naik, Sigal Rahamimov, Kyle Richardson, Amanpreet Singh, Harshit Surana, Aryeh Tiktinsky, Rosni Vasu, Guy Wiener, and 20 others. 2025. AstaBench: Rigorous benchmarking of AI agents with a scientific research suite. *arXiv preprint arXiv:2510.21652*.

Aubrey M. Brueckner, Darshil Patel, Yuhuan He, and Timothy Kassis. 2026. K-Bench: Measuring model performance on real scientific agent requests. *arXiv preprint arXiv:2608.21601v2*.

Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, Evan Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, Lilian Weng, and Aleksander Mądry. 2024. MLE-bench: Evaluating machine learning agents on machine learning engineering. *arXiv preprint arXiv:2410.07095*.

Ziru Chen, Shijie Chen, Yuting Ning, Qianheng Zhang, Boshi Wang, Botao Yu, Yifei Li, Zeyi Liao, Chen Wei, Zitong Lu, Vishal Dey, Mingyi Xue, Frazier N. Baker, Benjamin Burns, Daniel Adu-Ampratwum, Xuhui Huang, Xia Ning, Song Gao, Yu Su, and Huan Sun. 2024. ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery. *arXiv preprint arXiv:2410.05080*.

Deming Ding, Shichun Liu, Enhui Yang, Jiahang Lin, Ziying Chen, Shihan Dou, Honglin Guo, Weiyu Cheng, Pengyu Zhao, Chengjun Xiao, Qunhong Zeng, Qi Zhang, Xuanjing Huang, Qidi Xu, and Tao Gui. 2026. OctoBench: Benchmarking scaffold-aware instruction following in repository-grounded agentic coding. *arXiv preprint arXiv:2601.10343*.

Aniketh Garikaparthi, Manasi Patwardhan, and Arman Cohan. 2026. ResearchGym: Evaluating language model agents on real-world AI research. *arXiv preprint arXiv:2602.15112*.

Smriti Jha, Matteo Paltenghi, Chandra Maddila, Vijayaraghavan Murali, Shubham Ugare, and Satish Chandra. 2026. REAP: Automatic curation of coding agent benchmarks from interactive production usage. *arXiv preprint arXiv:2604.01527*.

Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2023. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*.

Patrick Tser Jern Kon, Jiachen Liu, Xinyi Zhu, Qiuyi Ding, Jingjia Peng, Jiarong Xing, Yibo Huang, Yiming Qiu, Jayanth Srinivasa, Myungjin Lee, Mosharaf Chowdhury, Matei Zaharia, and Ang Chen. 2025. EXP-Bench: Can AI conduct AI research experiments? *arXiv preprint arXiv:2505.24785*.

David R. Krathwohl. 2002. A revision of Bloom's taxonomy: An overview. *Theory Into Practice*, 41(4):212–218.

Zongwei Lv, Yaoming Li, Zhewen Tan, Yilun Yao, Yuxuan Tian, Lin Sun, Xiangzheng Zhang, Weihong Lin, Tong Yang, and Guangxiang Zhao. 2026. RealClawBench: Live OpenClaw benchmarks from real Developer-Agent sessions. *arXiv preprint arXiv:2606.03889*.

Mike A. Merrill, Alexander G. Shaw, Nicholas Carlini, Boxuan Li, Harsh Raj, Ivan Bercovich, Lin Shi, Jeong Yeon Shin, Thomas Walshe, E. Kelly Buchanan, Junhong Shen, Guanghao Ye, Haowei Lin, Jason Poulos, Maoyu Wang, Marianna Nezhurina, Jenia Jitsev, Di Lu, Orfeas Menis Mastromichalakis, and 66 others. 2026. Terminal-Bench: Benchmarking agents on hard, realistic tasks in command line interfaces. *arXiv preprint arXiv:2601.11868*.

Zechun Niu, Yukun Zhao, Jiaxin Zhang, Xu Shen, Jinhua Si, Han Tian, Can Xu, Yunfan Song, Jiaxin Mao, Yansong Gao, Yuchen Li, Jianmin Wu, Lingyong Yan, Shuaiqiang Wang, and Dawei Yin. 2026. DuMateBench: Evaluating autonomous agents in complex real-world workflows. *arXiv preprint arXiv:2608.26546v1*.

Qwen Team. 2026. Qwen3.8-Max model card. `https://huggingface.co/Qwen/Qwen3.8-2.4T-A95B-FP8`. Accessed November 6, 2026. PaperBench evaluated in BasicAgent Code-Dev mode.

Zachary S. Siegel, Sayash Kapoor, Nitya Nadgir, Benedikt Stroebl, and Arvind Narayanan. 2024. CORE-Bench: Fostering the credibility of published research through a computational reproducibility agent benchmark. *arXiv preprint arXiv:2409.11363*.

Igor Slinko, Yaroslav Golubev, and Sergey Titov. 2026. Coding-Agent benchmarks should match their users' task flows. *arXiv preprint arXiv:2610.09633*.

Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. 2025. PaperBench: Evaluating AI's ability to replicate AI research. *arXiv preprint arXiv:2504.01848*.

Jiayu Wang, Weijiang Lv, Bowen Fu, Jing Fu, Jiayi Song, Lingyu Zhang, Lanxuan Xue, Luodi Chen, Zepeng Xin, Kaiyu Li, and Xiangyong Cao. 2026a. Act as a real researcher: A suite of benchmarks evaluating frontier LLMs and agentic harnesses in research lifecycle. *arXiv preprint arXiv:2606.07462*.

Yubang Wang, Chenxi Zhang, Bowen Chen, Zezheng Huai, Zihao Dai, Xinchi Chen, Yuxin Wang, Yining Zheng, Jingjing Gong, and Xipeng Qiu. 2026b. ResearchEnvBench: Benchmarking agents on environment synthesis for research code execution. *arXiv preprint arXiv:2603.06739*.

Hjalmar Wijk, Tao Lin, Joel Becker, Sami Jawhar, Neev Parikh, Thomas Broadley, Lawrence Chan, Michael Chen, Josh Clymer, Jai Dhyani, Elena Ericheva, Katharyn Garcia, Brian Goodrich, Nikola Jurkovic, Holden Karnofsky,

Megan Kinniment, Aron Lajko, Seraphina Nix, Lucas Sato, and 4 others. 2024. RE-Bench: Evaluating frontier AI R&D capabilities of language model agents against human experts. *arXiv preprint arXiv:2411.15114*.

Yifan Wu, Zhuokai Zhao, Songlin Li, Ho Hin Lee, Jiacheng Zhu, Shirley Wu, Tianhe Yu, Serena Li, Lizhu Zhang, Xiangjun Fan, and Shengzhi Li. 2026. SWE-Together: Evaluating coding agents in interactive user sessions. *arXiv preprint arXiv:2606.29957*.

Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. 2024. OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments. *arXiv preprint arXiv:2404.07972*.

Kewei Xu, Xiaoben Lu, Shuofei Qiao, Zihan Ding, Haoming Xu, Lei Liang, and Ningyu Zhang. 2026a. LongDS-Bench: On the failure of long-horizon agentic data analysis. *arXiv preprint arXiv:2605.30434v3*.

Wanghan Xu, Shuo Li, Tianlin Ye, Qinglong Cao, Yixin Chen, Hengjian Gao, Yiheng Wang, Qi Li, Kun Li, Sheng Xu, Shengdu Chai, Fangchen Yu, Xiangyu Zhao, Zhangrui Zhao, Weijie Ma, Zijie Guo, Koutian Wu, Haoyu Zhou, Haoxiang Yin, and 32 others. 2026b. ResearchClawBench: A benchmark for end-to-end autonomous scientific research. *arXiv preprint arXiv:2606.07591*.

Zhipeng Xu, Jiahao Lu, Yining Zheng, Yuxin Wang, and Xipeng Qiu. 2026c. SWE-bench Science: Can coding agents resolve engineering tasks in science? *arXiv preprint arXiv:2608.19799*.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*.