arXiv:2611.14477v1 [cs.IR] 10 Nov 2026

FERN: Cross-Benchmark Table Retrieval for Text-to-SQL with Frozen Retrievers and Learned Ranking

Arjun Rao
arjun.rao01@sap.com

Nishant Gopinath
nishant.gopinath@sap.com

November 2026

Abstract

Text-to-SQL systems over large pooled catalogs must first retrieve the tables a question needs. We compare against the published table recall of RASL, which uses Claude 3.5 models to extract keywords and select tables, on pooled BIRD, Spider, and Fiben catalogs. FERN (Feature-Engineered Ranking with No generative LLM) instead fuses three frozen retrievers into 100 candidates, describes each with up to 120 features, and reorders them with LightGBM rankers fitted once on BIRD's training questions. On Spider, where it was not fitted, its candidate pool alone reaches RASL's published recall at fifteen tables (98.0), and the ranker raises the share of questions with every gold table in the top five (complete@5) from 83.9% to 95.6%. At five tables, FERN's recall is within one point of RASL's on Spider (96.7 vs 97.0) and on BIRD with its evidence hints (96.6 vs 97.5), and 8.2 points below on BIRD without them (89.3). On 121 real BEAVER questions, without BEAVER fitting, its recall at fifteen (82.1) falls between two BEAVER-trained retrievers (80.8, 84.5), but its complete@15 is below both (49.6 vs 54.5 and 62.0). On Fiben, it falls far short at fifteen tables (54.2 vs 69.2, different question sets), although its recall at fifty is numerically close to RASL's question-only retrieval result (80.7 vs 79.4). On BIRD dev with evidence, non-retrieval features raise complete@5 by 7.7 points over retrieval signals alone. On BIRD dev, median latency is 208 ms per question on one A100, without hosted model calls.

Keywords: text-to-SQL; schema linking; table retrieval; learning to rank; hybrid retrieval; feature engineering

1 Introduction

A text-to-SQL system can write a faithful query only if it has the tables the query needs. Over a pooled catalog with hundreds of tables from many databases, a retrieval stage has to find those few tables first: a missing table makes the query wrong, and extra tables consume the generator's context. This step, schema linking or table retrieval, increasingly uses large language models to extract keywords, rewrite questions, or read candidate schemas and choose (Pourreza & Rafiei, 2023; Talaei et al., 2024; Eben et al., 2025). We use RASL's primary table-retrieval results (Eben et al., 2025) as a published reference on pooled BIRD, Spider, and Fiben catalogs. RASL retrieves schema entities with keywords extracted by Claude 3.5 Haiku and selects tables with Claude 3.5 Sonnet v2.

Generative linking has costs beyond accuracy. RASL uses hosted models for keyword extraction and table selection; its end-to-end input-token estimate is 1.85 US dollars per 100 BIRD queries, including 3 cents in the category it labels retrieval (Eben et al., 2025). In enterprise settings, sending

schema and question text to an external service may not be acceptable. Retrieval research has long combined many weak relevance signals with a learned ranker (Qin & Liu, 2013; Burges, 2010), and such a ranker needs ingredients that can run locally: lexical search, off-the-shelf embedding models, and schema statistics.

This paper asks: how far can a ranker trained once on a public benchmark take table retrieval across pooled catalogs, using frozen retrievers and no generative model in the retrieval path? We answer with FERN (Feature-Engineered Ranking with No generative LLM). Three frozen retrievers propose 100 candidates; each candidate is described by features computed from the question, the catalog, and the retrievers' scores; LightGBM rankers (Ke et al., 2017) fitted once on BIRD's training questions order them. We compare with RASL's published numbers on the same pooled catalogs, with no database identifier given, and apply the BIRD-fitted rankers unchanged to Spider, Fiben, and three pooled BEAVER warehouses (Chen et al., 2024b).

The results, in brief. At a nominal budget of five tables, FERN's recall is within 0.3 points of RASL's on Spider and 0.9 points on BIRD with evidence (BIRD's per-question hints that name relevant columns and values); when FERN returns fifteen tables, it is about one point higher on both (on question-only BIRD it is 0.7 points below RASL's BIRD figure). Spider is the cleanest of these comparisons: FERN was never fitted on it, no design decision used it, and it involves no evidence ambiguity. Without BIRD's evidence hints, FERN's recall at five falls by 7.3 points. FERN's own candidate pool already reaches RASL's Spider recall at fifteen (98.0); the ranker's value lies mainly at small budgets and in complete table sets. On BEAVER, FERN's recall at fifteen falls between two retrievers trained on BEAVER, but fewer questions get a complete table set (every gold table in the list). On Fiben, its recall is close to RASL's question-only retrieval result at fifty tables, but substantially lower than RASL's selected-table recall at fifteen (Section 7).

Contributions.

1. FERN is a table retriever that combines three frozen retrievers with gradient-boosted rankers over schema-specific features, in three variants for catalogs with or without evidence, cell values, and column types. It runs in about 0.2 seconds per question on one GPU with no hosted model calls.
2. A comparison with published LLM-based schema linking (RASL) on its pooled BIRD, Spider, and Fiben catalogs. We made design choices on BIRD training data using rules stated in advance, with two documented exceptions; we fitted the rankers once; Spider and Fiben informed no design decisions.
3. An analysis, in the frozen configuration, of which feature groups matter at which table budget. On BIRD dev with evidence, rankers over retrieval signals alone or over the other features alone each reach about 98 recall at fifteen tables; adding the non-retrieval features to the retrieval signals raises complete@5 by 7.7 points. In the exploratory complete@5 ablations, removing database-support or join-graph features lowered the point estimate by 0.8 points each, with unadjusted intervals above zero; the other intervals touched or included zero.
4. Transfer results on a second benchmark, an opaque financial schema, and an enterprise benchmark with three warehouses, separating recall from complete-set recovery and showing, on these benchmarks, where a ranker trained on public data is enough and where it is not.

2 Related work

Schema linking with language models. Earlier text-to-SQL parsers linked question spans to schema elements by name and value matching inside the model (Guo et al., 2019; Wang, B. et al., 2020). RESDSQL ranks schema items with a learned cross-encoder before decoding (Li, H. et al., 2023), and CodeS combines schema-item classification with value retrieval (Li, H. et al., 2024). With large language models, linking became a prompting step: DIN-SQL links the schema with an LLM before generation (Pourreza & Rafiei, 2023) and CHESS retrieves entities and context in an LLM-driven pipeline (Talaei et al., 2024). Maamari et al. (2024) argue that strong reasoning models reduce the need for explicit schema linking when the database is known and small. These methods mostly assume a given database. Our setting pools many databases without identifying the target, as in RASL (Eben et al., 2025), which indexes schema entities with descriptions, retrieves with LLM-extracted keywords, calibrates relevance on 200 training questions per dataset where a training split exists (not Fiben), and lets an LLM choose the tables. LinkAlign (Wang, Y. et al., 2025) also addresses linking across many databases.

Retrieval for large catalogs. CRUSH4SQL asks an LLM to hallucinate a minimal schema and retrieves real tables that match it (Kothyari et al., 2023). DBCopilot routes questions to databases and tables with a generative model trained on synthetic questions generated over the target databases (Wang, T. et al., 2025); in RASL's Table 2, it reports 94.6 recall at fifteen on BIRD and 97.6 on Spider. Dense table retrieval embeds questions and tables in a shared space (Herzig et al., 2021). Chen, Zhang & Roth (2024a) show that join-aware retrieval matters when questions need several tables. CORE-T (Soliman et al., 2026) targets coherent multi-table retrieval over pooled catalogs without database identifiers and also evaluates on BEAVER; it reports table-selection F1 rather than recall@k on RASL's catalogs, so we do not compare with it numerically. CSR-RAG (Singh et al., 2026) combines similarity to past questions on the target catalog, a schema knowledge graph, and hypergraph ranking, without an LLM in retrieval; FERN uses no target-catalog history and is fitted once on a public benchmark. NIFT-SQL (Rao & Gopinath, 2026) fine-tunes a necessity scorer and a retriever on a customer's own questions and fuses them; it asks what target-catalog supervision buys. FERN asks the complementary question: what a ranker fitted once on public data delivers with no target-domain fitting. Both use database-level support features in a learned fusion; FERN differs in using only frozen retrievers, a larger schema-specific feature set, cross-benchmark transfer, and a comparison with LLM-based linking.

Learning to rank and hybrid retrieval. Learning to rank over hand-built relevance features is classic information retrieval (Qin & Liu, 2013), with gradient-boosted trees such as LambdaMART among the strongest rankers on such features (Burges, 2010; Qin, Z. et al., 2021). For tables, Zhang & Balog (2018) already rank web tables with supervised learning over lexical and semantic features; FERN differs in its target (the set of tables a SQL query needs), its pooled multi-database catalogs, and its evaluation of transfer across benchmarks. Hybrid lexical and dense retrieval is a standard first stage, commonly fused by reciprocal rank fusion (Cormack et al., 2009; Bruch et al., 2023). We claim no new ranking algorithm. The contribution is empirical: how far this recipe reaches against LLM-based linking, at which budgets, and where it stops.

3 Problem setting

Task. Given a natural-language question and a pooled catalog of tables from many databases, return a ranked list of tables such that the tables the gold SQL query uses appear near the top. The target database is not given.

Catalogs. Table 1 lists the catalogs. For BIRD (Li, J. et al., 2023) and Spider (Yu et al., 2018), we pool every database in the training and development splits into one catalog, as in RASL; Fiben (Sen et al., 2020) is a single financial database. Our table and column counts for BIRD, Spider, and Fiben match RASL's Table 1. BEAVER's current public release pools three warehouses (dw, nova, neutron); its 121 real questions come from query logs of the dw warehouse, and we score a random sample (seed 7) of 1,000 of its 7,857 synthetic questions.

Table 1. Catalogs and test questions.

Dataset	Databases	Tables	Columns	Test questions	Gold tables per question
BIRD dev	80	597	4,337	1,534	1.94
Spider dev	166	876	4,503	1,034	1.51
Fiben	1	152	374	298 (RASL: 279)	3.82
BEAVER, three warehouses	3	381		121 real; 1,000 synthetic	3.93; 3.52

Gold tables. The gold tables of a question are the catalog tables, within the question's database, referenced in its gold SQL, extracted with the sqlglot parser (sqlglot, 2024); names introduced by common table expressions are excluded. Fiben's gold SQL is parsed as PostgreSQL, the others as SQLite. On Fiben, 298 of 300 questions parse. RASL reports 279 Fiben questions without listing which, so the Fiben comparison uses different question sets (Section 6.1 bounds the effect). BEAVER provides gold table lists. Extraction matters at this margin: a table-name regular expression instead of the parser changes the gold set for 25 of 1,534 BIRD dev questions (1.6%), all because a table named `order` matches the keyword `ORDER BY`.

Metrics. For a question with gold set G and top-k list T_k, recall@k is $|G \cap T_k|/|G|$, macro-averaged over questions as in RASL. We also report complete@k, the share of questions with $G \subseteq T_k$, because a query that misses any required table is wrong. A complete table set is necessary for a correct query but does not guarantee one. RASL's N is a maximum budget: its LLM selects only the tables it judges relevant, and on BIRD and Spider it never predicts more than 15 (Eben et al., 2025). FERN returns exactly fifteen tables at k = 15, whereas RASL may return fewer. Comparisons at five and fifteen therefore use the same nominal caps but may yield different realized output sizes. Intervals are 95% percentile bootstrap intervals over questions (2,000 resamples, seed 0); Section 6.1 adds intervals that resample whole databases. They describe uncertainty only in our estimates. RASL's per-question results are not available, so differences from RASL are point differences, not tests.

Table 2. Inputs available to each system. FERN-ev, FERN-q and FERN-schema are defined in Section 4.

	RASL full	FERN-ev	FERN-q	FERN-schema
Supervision	200 calibration questions per dataset (none on Fiben)	9,394 labeled BIRD train questions, fitted once	same	same
BIRD evidence	not stated; its examples include evidence-like clauses	yes	no	no
Cell values	not indexed	yes	yes	no
Column and value descriptions (BIRD)	yes	no	no	no
Generative model calls per question	2 (keywords, table selection)	0	0	0
Embedding models	Cohere Embed v3	Qwen3-Embedding-0.6B, BGE-large	same	same
Output size	at most N, chosen by the LLM	exactly k	exactly k	exactly k

4 Method

FERN has four stages.

1. Frozen retrieval. Each table is represented by a full card, `table (database db): col1, col2 (natural name), ...`, where underscores in the database name become spaces and a column's natural-language name (BIRD's and Spider's `column_names`) is added in parentheses when it differs from the stored name after removing spaces, underscores, and case; and by a name card, `table words (database db)`, with camelCase and underscores split. Cards contain no types, keys, or values and are not truncated; BGE-large reads at most 512 tokens, so for very wide tables its signal sees only the first part of the card. Five signals score every table: BM25 (Robertson & Zaragoza, 2009) on full cards and on name cards (rank_bm25 BM25Okapi, k1 = 1.5, b = 0.75; tokens are lower-cased alphanumeric runs after splitting camelCase and underscores, with no stemming or stop words); Qwen3-Embedding-0.6B (Zhang, Y. et al., 2025) on full and name cards, with its built-in query instruction for questions; and BGE-large-en-v1.5 (Xiao et al., 2024) on full cards, with its retrieval prefix for questions. Embeddings are L2-normalized and compared by dot product, in bfloat16. The question text is the question followed by BIRD's evidence in FERN-ev and the question alone otherwise.
2. Candidate pool. The pool score of a table is $\sum_s 1/(60+\text{rank}_s)$, summed over BM25, Qwen3, and BGE on full cards, with zero-based ranks over the whole catalog (Cormack et al., 2009). The 100 tables with the highest pool score are the candidates (average recall at 100 is 99.4 on BIRD dev). Ties are broken by array order.
3. Features. Table 3 describes each candidate's features. We compute ranks, z-scores, percentiles, and gaps over the 100 candidates for each question.
4. Ranking. A LightGBM ranker with the lambdarank objective, trained with each question's candidates as one group (pairwise gradients) and applied pointwise, scores the candidates, which are then sorted by that score.

Table 3. Features. The six signals in the score and agree families are the pool score, BM25 (full and name cards), Qwen3 (full and name cards), and BGE. FERN-ev uses all 120.

Family	n	Content
score	30	each signal's raw score, rank, z-score, gap to the best candidate and percentile
agree	9	minimum, maximum, mean and standard deviation of the six ranks; how many signals place the table in their top 5 and top 15; rank differences Qwen3 minus BM25, Qwen3 minus BGE, name card minus full card
db	14	for the pool score and for Qwen3: the database's best score, mean score and probability share (softmax of the candidates' z-scores, summed by database), the database's rank by best score, the table's rank within its database, its gap to the database's best; plus the number of candidates from the database and its table count
lex	8	overlap of question tokens with table-name tokens and with all name tokens of the table, IDF-weighted overlap, overlap count, best fuzzy match, fuzzy matches above 0.8, natural column names found in the question, table name found in the question
cov	3	share of question tokens covered by the table's tokens; tokens it adds beyond the top candidate and beyond the top three
ev	3	evidence contains the table name, how many column names, how many natural column names (zero without evidence)
val	6	question literals found among the table's cell values; share of literals matched; columns with a match; literals matched only by this table within its database; question numbers inside a numeric column's minimum-maximum range; share in range
type	7	numeric columns, share numeric, date-like columns; comparison, aggregation, time and percentage questions crossed with numeric or date-like columns
meta	6	columns, primary-key columns, foreign-key columns, log(1 + rows), junction table (at least two foreign-key columns and at most four columns), name length
graph	7	foreign-key degree, PageRank, join distance to the top candidate, neighbors among the top five, bridge (neighbor of two or more of the top five), shared non-generic column names with the top candidate and with the top five
q	11	question words, content tokens, literals, numbers, presence of evidence, entropy of the softmax of pool-score z-scores, gap between the top two pool scores, aggregation, comparison, percentage and time flags
emb	16	element-wise product of the Qwen3 question and full-card embeddings, projected to 16 dimensions

Feature details. Tokens for lex, cov, and IDF are lower-cased alphanumeric runs after camelCase and underscore splitting, with a fixed list of about 60 function words removed and a final s stripped from words longer than three characters; IDF is log(N / (1 + df)) over the catalog's tables. Fuzzy matching compares each question token of at least four characters with every catalog token by normalized Indel similarity (rapidfuzz). Literals are quoted strings of 2 to 60 characters, capitalized spans of the question with leading function words removed, and numbers. Cell values are indexed once per catalog from the databases' SQLite files: up to 20,000 distinct non-null values per text column, lower-cased, 2 to 60 characters, matched exactly; and the minimum and maximum of numeric columns. A column is date-like if its type contains date or time, or its name contains date, year, or time; fixed regular expressions detect question types. PageRank (Page et al., 1999; damping 0.85, 30 iterations) and join distances (up to five hops, nine otherwise) use declared foreign keys, undirected, within each database. The top candidate has the highest pool score. The projection for emb is a PCA with 16 components fitted on every ninth candidate row of BIRD train (with evidence) and reused for all variants. We will release the complete feature code with the materials listed under Data and code availability. The pipeline is deterministic: rerunning it from scratch reproduced every reported retrieval number exactly.

Variants and fitting. Three rankers are fitted once on 9,394 BIRD training questions whose gold SQL parses and names a catalog table (RASL reports 9,427 training questions). FERN-ev uses the question with evidence and all 120 features. FERN-q uses the question alone and drops the evidence features and the evidence flag (116). FERN-schema also drops cell values, column types, and row counts (102) for catalogs whose values or types are not available to us. Labels are 1 for gold tables among the candidates and 0 otherwise. The authors set hyperparameters once before the development runs, using common LightGBM settings, and never tuned them: 600 trees, learning rate 0.05, 63 leaves, at least 50 rows per leaf, row subsampling 0.8 every iteration, column subsampling 0.7, L2 regularization 1.0, seed 0, other parameters at LightGBM defaults.

Which variant runs where. BIRD with evidence uses FERN-ev; BIRD without evidence and Spider use FERN-q (Spider has values but no evidence). Fiben and BEAVER use FERN-schema: our copies have no cell values, BEAVER provides no column types, and Fiben's declared foreign keys (159) and BEAVER's declared keys, where they exist, feed the graph features.

5 Protocol

Rules stated in advance. The candidate pool and the set of signals were chosen by rules fixed before running them and applied with five-fold cross-validation (folds by database) on a 3,000-question subset of BIRD train: the pool with the highest complete@100, with ties within 0.2 points going to fewer retrievers; and a signal kept only if removing it costs at least 0.3 points of complete@5 or complete@15, with a joint-removal check. The rule removed a best-column dense signal (Appendix A).

Two documented exceptions. (1) The pool rule selects Qwen3 + BGE (complete@100 98.4) over BM25 + Qwen3 + BGE (98.3). We overrode it to include lexical retrieval, arguing that exact matching complements dense retrieval for codes and abbreviations. The override was made after an earlier BEAVER transfer check, in a different configuration, had been seen, but before Spider and Fiben were scored; it should be read as informed by that check. (2) A fine-tuned decision model was evaluated on BIRD dev during development and left out of FERN because it added little (Appendix B). That decision used BIRD dev.

What each test set has informed. Spider and Fiben informed no decision. Fiben was evaluated with FERN-schema. Spider was evaluated with FERN-q and, after the method was frozen, with FERN-schema for the analysis without cell values in Section 6.1; neither Spider evaluation informed a design decision. That analysis also uses a fourth ranker, FERN-ev refitted on BIRD train without cell values, column types, or row counts. BIRD dev informed the second exception and the development studies of Appendix B. BEAVER informed the first exception. BIRD dev and BEAVER results are therefore development-exposed; Spider is the cleanest comparison with RASL. We exclude evaluation questions and labels from fitting; each question is available at inference, and catalogs are visible as the task requires: BM25 statistics, IDF, and the value index are built over the catalog being searched.

Comparison with RASL. We take RASL's retrieval results from its published Tables 2 and 5 and its cost figures from its cost analysis (Eben et al., 2025); we did not rerun it. We use the same named development splits and reported catalog counts for BIRD and Spider; exact question-ID and gold-label parity with RASL is not verified. Fiben uses a different question set. Table 2 lists the input differences.

Leakage, by provenance. Table 4 lists every input and the stages that see it. The feature code

receives the question, the evidence (FERN-ev only), the catalog, the cell values, and the frozen scores; candidate identities come from the whole catalog, and the gold SQL, gold tables, and gold database are never passed to it. Labels are attached after features are computed. PCA and the rankers are fitted only on BIRD train. As a coarse sanity check, not a proof, no single feature separates gold from non-gold candidates on BIRD dev with an AUC above 0.98; the strongest, the pool score's gap to the best candidate, reaches 0.978, as expected of the main retrieval signal.

Table 4. Inputs and the stages that use them.

Input	Source	Used by
Question text	the benchmark question	retrievers, features
Evidence text	BIRD's evidence field	FERN-ev only: appended to the question for all retrievers and for every feature computed from question text (lex, cov, val literals, type and q flags, emb), plus the ev features and the evidence flag
Schema: names, types, keys	the catalog	retrievers (names only), features (names, types and keys)
Cell values, row counts	the catalog's SQLite files, never the gold SQL	val and meta features, FERN-ev and FERN-q
Retriever scores	frozen models	score, agree and db features
Gold tables	parsed from gold SQL for BIRD, Spider and Fiben; supplied by BEAVER for BEAVER	fitting the rankers (BIRD train) and evaluation only

6 Results

6.1 Comparison with published LLM-based linking

Table 5. Table recall on the pooled catalogs. RASL full uses Claude 3.5 Haiku keywords and Claude 3.5 Sonnet v2 table selection and returns at most N tables; RASL retriever ranks by retrieval with the Haiku keywords, without LLM table selection. The BM25, DBCopilot, RASL retriever, and RASL full rows are from RASL's Table 2. Differences are FERN minus RASL full, in points. FERN intervals: 95% bootstrap over questions.

	BIRD R@5	BIRD R@15	Spider R@5	Spider R@15	Fiben R@5	Fiben R@15
BM25 (RASL's Table 2)	68.3	82.8	86.5	93.9	33.3	38.6
DBCopilot	85.8	94.6	91.6	97.6	41.1	56.9
RASL retriever	70.5	92.6	72.1	94.1	34.6	64.5
RASL full	97.5	97.8	97.0	98.0	69.1	69.2
FERN pool alone	86.8	96.2	89.6	98.0	24.8	39.3
FERN	96.6 [95.9, 97.4]	98.9 [98.4, 99.2]	96.7 [95.6, 97.6]	99.1 [98.6, 99.6]	36.7 [33.7, 39.5]	54.2 [50.8, 57.4]
FERN minus RASL full	-0.9	+1.1	-0.3	+1.1	-32.4	-15.0
FERN, BIRD question only	89.3 [88.0, 90.6]	97.1 [96.4, 97.7]				

FERN variants: FERN-ev on BIRD, FERN-q on question-only BIRD and on Spider, FERN-schema on Fiben. Fiben: FERN 298 questions, RASL 279.

At a nominal budget of five tables, FERN is 0.3 points below RASL's published recall on Spider and 0.9 points below on BIRD with evidence. When FERN returns fifteen tables, it is 1.1 points higher on both; RASL returns at most fifteen, so this compares a full list of fifteen with a possibly shorter one. Whether RASL used BIRD's evidence is not stated: its error-analysis examples (Figures 3 and 5 and Appendix E of Eben et al., 2025) show questions with evidence-like clauses appended, which supports FERN-ev as the closer comparison. Without evidence, FERN's recall at five falls to 89.3, 8.2 points below RASL, and its recall at fifteen to 97.1, 0.7 points below. For deployments without per-question evidence hints, the question-only row is more relevant.

FERN's candidate pool alone already reaches RASL's Spider recall at fifteen (98.0) and is above every baseline without LLM table selection in RASL's comparison at fifteen tables, on Spider and on BIRD with evidence (the question-only BIRD pool, 94.2, is below DBCopilot's 94.6). Much of the gap between earlier non-LLM baselines and RASL is thus closed by the frozen candidate pool; the ranker earns its place at small budgets and on complete sets (Table 6). The BM25 rows of Tables 6 and 8 rank FERN's 100 candidates by their BM25 score; they show what BM25 contributes as a signal, not BM25 as a standalone retriever, and are not comparable with RASL's BM25 row. Over the whole catalog, our word-level BM25 is still weaker than the BM25 row of RASL's Table 2, which RASL adopts from DBCopilot's reported results with BM25 parameters optimized on training data (Eben et al., 2025, Appendix F); we did not tune ours.

On Fiben, we evaluate the two systems on different question sets. If RASL's 279 questions are a subset of our 298, FERN's recall on them is bounded by the two extremes for the 19 others: 51.1 to 57.9 at fifteen and 32.4 to 39.2 at five, far below RASL's 69.2 and 69.1 in either case.

Questions are clustered in databases, so we also resample whole databases (11 on BIRD dev, 20 on Spider dev). The intervals widen to [94.3, 98.5] for BIRD recall at five and [97.7, 99.7] at fifteen, and to [94.2, 98.3] and [98.3, 99.7] on Spider. These are approximate given so few clusters, and

they describe the sensitivity of FERN's own estimates; they do not test the difference from RASL, which we report only as signed point differences.

Table 6. Recall and complete@k against the candidate pool on the same catalogs. "BM25 rerank of pool" ranks the 100 candidates by BM25 score. C@k is complete@k.

Dataset	Method	R@5	R@10	R@15	C@5	C@10	C@15
BIRD, evidence	BM25 rerank of pool	50.1	66.0	73.6	31.4	47.0	56.1
	pool	86.8	93.6	96.2	75.0	87.4	91.9
	FERN-ev	96.6	98.5	98.9	93.9	96.8	97.7
BIRD, question only	BM25 rerank of pool	30.0	48.1	58.5	16.3	30.4	40.1
	pool	81.5	90.4	94.2	67.4	81.6	88.3
	FERN-q	89.3	95.1	97.1	83.3	91.9	95.0
Spider	BM25 rerank of pool	65.6	74.2	78.3	56.3	66.4	71.7
	pool	89.6	96.6	98.0	83.9	94.2	96.7
	FERN-q	96.7	98.4	99.1	95.6	98.2	98.9
Fiben	BM25 rerank of pool	25.4	34.6	39.4	1.7	3.4	4.0
	pool	24.8	34.2	39.3	0.7	2.7	4.0
	FERN-schema	36.7	44.4	54.2	3.7	7.4	15.1

Reranking the same pool raises complete@5 by 18.9 points on BIRD with evidence (75.0 to 93.9) and by 11.7 on Spider (83.9 to 95.6), and complete@15 by 5.8 and 2.2.

Table 7. Recall@k and complete@k for k = 1 to 8, 10, 15 and 20 (R / C).

k	BIRD, FERN-ev	BIRD, FERN-q	Spider, FERN-q	BIRD pool alone (R)
1	54.7 / 21.6	45.8 / 18.1	64.6 / 44.8	47.6
2	84.4 / 69.7	70.7 / 51.6	86.2 / 77.4	68.9
3	92.5 / 85.5	81.4 / 69.8	93.2 / 89.7	78.1
4	95.4 / 91.3	86.3 / 78.3	95.6 / 93.9	84.0
5	96.6 / 93.9	89.3 / 83.3	96.7 / 95.6	86.8
6	97.4 / 95.4	91.0 / 85.7	96.9 / 95.9	89.1
7	97.8 / 95.9	92.6 / 88.2	97.1 / 96.2	90.8
8	98.2 / 96.4	93.6 / 89.6	97.5 / 97.0	92.0
10	98.5 / 96.8	95.1 / 91.9	98.4 / 98.2	93.6
15	98.9 / 97.7	97.1 / 95.0	99.1 / 98.9	96.2
20	99.1 / 98.4	97.8 / 96.1	99.3 / 99.1	97.0

FERN first reaches or exceeds RASL's published recall at five at k = 7, on Spider and on BIRD with evidence. On BIRD without evidence, it is below that reference at k = 15 and above it at k = 20. Because realized list sizes and inputs differ, this is an operational comparison, not a measure of what LLM selection adds.

Without cell values, column types, or row counts. RASL does not index database values, while FERN-ev and FERN-q do. Removing values, column types, and row counts together, with separately fitted rankers, changes recall at five and fifteen by 0.3 and 0.2 points on BIRD with evidence (FERN-ev refitted, 106 features: 96.3 and 98.7) and by 0.8 and 0.1 on Spider (the already-fitted FERN-schema: 95.9 and 99.0). FERN's results are therefore robust to removing these inputs together; this does not establish parity with RASL's schema representation or evidence use. Without them, FERN is 1.2 points below RASL's published recall at five tables on BIRD and 1.1 on Spider.

6.2 Transfer to an enterprise benchmark

BEAVER (Chen et al., 2024b) is built from real enterprise data warehouses. We apply FERN-schema as fitted on BIRD; no BEAVER questions are used for fitting. The comparison rows come from NIFT-SQL's confirmation run on the same release and the same 121 real questions from the dw warehouse (Rao & Gopinath, 2026), and both are trained on BEAVER's synthetic questions. NIFT-SQL fuses a necessity scorer with hybrid search over a retriever fine-tuned on BEAVER; the second row is the retriever fine-tuned on the same training and validation questions, the label-matched baseline that NIFT-SQL reports. In that run, the full system recovered fewer complete sets on the real questions than this retriever alone (54.5 vs 62.0). From the NIFT-SQL study, we reuse only its prepared BEAVER data files, its declared-key list, and these two comparison rows; the candidate pool, features, rankers, and all other experiments are new. Because we do not have NIFT-SQL's per-question outputs, the comparisons with these rows are descriptive, without paired intervals. NIFT-SQL reports the real-question complete@15 values of these two rows (Rao & Gopinath, 2026, Table 9); recall@5, recall@15, and complete@5 come from the same run and are not reported there.

Table 8. BEAVER, real questions (121) and a sample of synthetic questions (1,000). "BM25 rerank of pool" ranks the 100 candidates by BM25 score. Rows trained on BEAVER come from the NIFT-SQL confirmation run; their complete@15 values are published in Rao & Gopinath (2026, Table 9).

Questions	Method	Trained on BEAVER	R@5	R@15 [95% CI]	C@5	C@15
121 real	BM25 rerank of pool	no	48.9	70.4	16.5	29.8
	pool	no	55.3	78.8	20.7	39.7
	FERN-schema	no	57.0	82.1 [78.3, 85.5]	22.3	49.6
	NIFT-SQL	yes	56.6	80.8	24.0	54.5
	retriever fine-tuned on BEAVER (label-matched)	yes	53.0	84.5	22.3	62.0
1,000 synthetic	BM25 rerank of pool	no	53.9	83.9	13.7	53.3
	pool	no	54.6	88.0	16.5	61.0
	FERN-schema	no	61.7	88.9 [88.0, 89.9]	22.5	63.3

On the real questions, FERN improves on its pool by 3.3 points of recall@15 and 9.9 points of complete@15. Its recall@15 lies between the two BEAVER-trained retrievers (80.8 and 84.5); its complete@15 is 5.0 and 12.4 points lower (6 and 15 questions). With almost four gold tables per question, average recall can look strong while a query still misses a required table, so complete@15 is the more relevant number. With 121 questions, these differences are descriptive: a complete@15 near 50% has a 95% interval of about plus or minus nine points, and a 5.0-point difference is six questions. Both target-trained baselines recovered more complete table sets on these questions.

6.3 Where the signal comes from

We split the 120 features into retrieval signals (the score, agree, and db families, 53 features) and the other families (67), with the same hyperparameters. We fitted each ranker on all of BIRD train and evaluated it on BIRD dev; separately, we ran five-fold cross-validation grouped by database within BIRD train (all 9,394 questions). The PCA projection of the emb family, fitted once on all of BIRD train, is shared across folds; it is unsupervised but sees the held-out folds' unlabeled rows, and the retrieval-only ranker does not use it, so the cross-validation is a development diagnostic rather than a fully held-out estimate. We also remove one family at a time (BIRD dev). The split is not perfectly clean: some of the other features use the pool's top candidates as reference points (coverage beyond the top table, join distance to it, the entropy of pool scores).

Table 9. Feature groups in the frozen configuration, FERN-ev inputs. CV: database-grouped cross-validation within BIRD train with a shared unsupervised PCA (development diagnostic). The last two rows give all features minus the subset, paired over questions, with 95% bootstrap intervals; differences are computed from unrounded per-question values.

Features	CV R@5	CV C@5	CV R@15	CV C@15	dev R@5	dev C@5	dev R@15	dev C@15
retrieval signals only (53)	78.6	70.0	90.7	86.8	91.9	86.2	98.2	96.6
other families only (67)	86.4	77.1	94.6	89.9	95.4	91.9	98.5	97.1
all (120)	89.2	82.1	95.8	92.8	96.6	93.9	98.9	97.7
all minus retrieval-only		+12.1 [11.4, 12.9]		+6.0 [5.5, 6.5]		+7.7 [6.3, 9.1]	+0.6 [0.3, 1.0]	
all minus other-only		+5.0 [4.5, 5.6]		+2.9 [2.5, 3.3]		+2.0 [1.0, 2.9]	+0.4 [0.1, 0.7]	

Both groups reach high recall at fifteen; combining them matters at five. On BIRD dev, a ranker over retrieval signals alone reaches 98.2 recall at fifteen and one over the other families 98.5; all features reach 98.9 (differences +0.6 and +0.4 computed before rounding, both intervals above zero). The large effect is at five tables: combining the groups raises complete@5 by 7.7 points over retrieval signals alone and by 2.0 over the other families. The database-grouped cross-validation within BIRD train, a development diagnostic, shows the same ordering (12.1 and 5.0 points). On BIRD with evidence, the non-retrieval features pay off most when only a few tables can be passed to a generator; we did not run these feature-group comparisons on the transfer datasets.

Table 10. Removing one family (BIRD dev; all features minus the reduced set, 95% paired bootstrap, computed from unrounded values; twelve exploratory comparisons, not corrected for multiplicity).

Family removed	dev C@5	dev R@15
db (14)	+0.8 [0.1, 1.6]	+0.2 [0.1, 0.4]
graph (7)	+0.8 [0.1, 1.6]	-0.1 [-0.2, 0.1]
val (6)	+0.7 [0.0, 1.4]	+0.1 [-0.1, 0.4]
lex (8)	+0.3 [-0.4, 1.0]	+0.1 [-0.1, 0.3]
score, agree, cov, ev, type, meta, q, emb	between -0.3 and +0.1, all intervals include zero	between -0.1 and 0.0

Few single families have a measurable effect. Removing database or graph features lowered BIRD dev complete@5 by 0.8 points; removing value features lowered it by 0.7, with the interval touching zero. For the other nine families, the exploratory complete@5 intervals included zero. These ablations do not establish whether that reflects redundancy between families (the score family, for example, overlaps with the agreement and database families, which are built from the same scores) or whether the family contains little information. In the fitted ranker, the most important features by gain are the pool score's gap to the best candidate (49.5% of the total gain), the pool rank (13.3%) and the best pool score in the table's database (3.5%).

Appendix B reports the development studies that led to this configuration, including a fine-tuned decision model that added little.

6.4 Cost

Table 11. Measured retrieval latency per question, FERN-ev, batch size 1, one NVIDIA A100 40GB, 100 BIRD dev questions after 5 warm-up questions; indexes and card embeddings precomputed.

Stage	Median (ms)	90th percentile (ms)
embed the question with Qwen3 and BGE (GPU)	71.9	87.6
BM25, two views (CPU)	6.6	10.5
dense scores and pool	0.4	12.8
features (CPU, Python)	122.7	146.0
ranker (CPU)	6.4	10.2
total	208.3	246.9

Stage medians need not add up to the median total. FERN makes no hosted model calls and sends no schema outside the system. Together, the two embedders have under one billion parameters. Fitting the three rankers takes about three and a half minutes on a CPU; computing features for the 9,394 training questions takes 11 minutes with evidence and 4 without (in bulk, without the per-question overheads included in Table 11). The features are plain Python and have not been optimized. RASL estimates 1.85 US dollars per 100 BIRD queries and 0.80 per 100 Spider queries in its input-token cost analysis, which includes SQL generation, with 3 cents per 100 in the category it labels retrieval (Eben et al., 2025); it does not report latency, so we make no latency comparison.

7 Discussion

BIRD and Spider. On these benchmarks, whose table and column names are mostly readable English, frozen retrievers already place almost every gold table in the candidate pool. The study does not isolate readability as the variable that explains the difference from Fiben. What remains is ordering the candidates. At a nominal budget of five, FERN's recall is 0.3 points below RASL's published value on Spider and 0.9 points below on BIRD with evidence; the comparison does not isolate the contribution of LLM selection. We do not know why RASL is ahead on BIRD at five tables; one hypothesis is the LLM's ability to judge which of several plausible tables a question needs, another is RASL's use of BIRD's column and value descriptions.

Evidence. BIRD's evidence hints often name the exact columns a question needs. With evidence, FERN's recall at five is 7.3 points higher (96.6 vs 89.3). The 7.3-point difference compares two separately fitted variants that differ in both the query text and the evidence-specific features; it is not decomposed between them. For deployments without per-question evidence hints, the question-only results are more relevant, including the 8.2-point gap at five tables to RASL's published number.

Opaque schemas. Fiben's tables have run-together uppercase names (for example, `ACCOUNTSPAYABLE ANDACCRUEDLIABILITIES`), 2.5 columns per table, one database, and almost four gold tables per question. Reranking the pool by BM25 and the pool itself both reach only 39 recall at fifteen; FERN improves on its pool by 15 points but remains far below RASL. The results expose limitations in candidate retrieval and in ranking under a small output budget. The fused pool is weak: 67.0 recall at fifty tables, against 79.4 for RASL's question-only retriever and 90.6 for its pool using the question plus LLM-extracted keywords (Eben et al., 2025, Table 5), and its average recall at 100 is capped at 82.8; standalone BM25 over the whole catalog reaches only 7.0 at fifteen.

Fiben has only 152 tables, so fifty is a third of the catalog. After reranking its 100 candidates, FERN reaches 80.7 recall when returning fifty tables, but only 54.2 when returning fifteen, a drop of 26.5 points. RASL's question-only configuration reports 79.4 recall for its initial fifty-table pool and 67.5 after LLM selection with a maximum output of fifteen, a drop of 11.9 points. These figures suggest that retaining relevant tables under a small output budget is a substantial difficulty for FERN. They do not isolate the effect of LLM selection, because the systems use different question sets, embeddings, and retrieval representations; RASL retrieves individual schema entities with Cohere embeddings. Fiben also differs from BIRD and Spider in several ways at once: its names are unsegmented, our copy has no cell values, it has a single database (which removes discrimination between databases but preserves within-database ranking signals), and its questions need many tables. We cannot tell which of these matters.

Enterprise use. On the 121 real dw questions, with no BEAVER question used for fitting (although BEAVER results informed one design decision; Section 5), FERN's recall at fifteen fell between the two target-trained point estimates, while its complete@15 was lower than both. Both baselines were trained on BEAVER's synthetic questions (Table 8); a retriever fine-tuned this way could enter FERN as another signal, which we have not tested. We see public-data fitting and target-catalog supervision as stages of a deployment: a ranker fitted on public data first, supervision on the target catalog once data exists.

8 Limitations

- Published reference numbers. We did not rerun RASL; its per-question results are unavailable, and its inputs differ (Table 2): it uses BIRD's descriptions, not cell values, and returns at most N tables. Differences from RASL carry no interval. The Fiben question sets differ.
- Development exposure. BIRD dev and BEAVER informed two design decisions (Section 5); Spider and Fiben did not.
- Supervision. The rankers are fitted on 9,394 labeled BIRD questions, far more than RASL's 200 calibration questions. FERN avoids target-domain fitting, not supervision. Pretraining of the frozen embedders on public benchmarks cannot be ruled out, as for any pretrained model.
- Evidence and values. FERN's best BIRD numbers use evidence and cell values. Without evidence, recall at five is 8.2 points below RASL's published number.
- Gold labels. Gold-table extraction changes 1.6% of BIRD dev gold sets, the same order as the differences from RASL; RASL does not describe its extraction. BIRD's gold SQL itself contains annotation errors (Wretblad et al., 2024; Jin et al., 2026), which can affect both FERN's and RASL's reported recall.
- Candidate ceiling. Gold tables outside the 100 candidates cannot be recovered (average recall at 100 is 99.4 on BIRD dev and 82.8 on Fiben).
- End-to-end scope. We evaluate table retrieval only, not database routing or the SQL generated from the retrieved tables.
- Benchmarks. We evaluate RASL's three benchmarks and one enterprise benchmark. Larger enterprise suites with thousands of tables need candidate-pool decisions of their own and are left to future work.

9 Conclusion

Against RASL's published results on pooled catalogs, rankers fitted once on BIRD over frozen retrievers come within one point of its published recall at five tables on Spider and on BIRD with evidence, at about 0.2 seconds per question with no hosted model calls; without evidence, the gap on BIRD at five tables is eight points. On Spider, the candidate pool of frozen retrievers (BM25, Qwen3 and BGE) alone already reaches RASL's recall at fifteen, and the ranker's main contribution is complete table sets at small budgets. On the real questions of an enterprise benchmark, the rankers' recall at fifteen falls between that of two retrievers trained on it, while their complete@15 is lower than both. On Fiben, their recall is numerically close to RASL's question-only retrieval result at fifty tables but substantially lower than RASL's selected tables at fifteen; both candidate retrieval and ranking under a small budget remain limitations. FERN was competitive in reported recall on Spider and on BIRD with evidence, which makes it a strong baseline on these benchmarks; Fiben shows where it is not enough.

Data and code availability

The code, frozen configuration, catalog and question manifests, fitted rankers and per-question results will be released upon journal publication. BEAVER requires accepting its terms on Hugging Face.

References

Bruch, S., Gai, S., & Ingber, A. (2023). An analysis of fusion functions for hybrid retrieval. ACM Transactions on Information Systems, 42(1), Article 20.

Burges, C. J. C. (2010). From RankNet to LambdaRank to LambdaMART: An overview. Microsoft Research Technical Report MSR-TR-2010-82.

Chen, P. B., Zhang, Y., & Roth, D. (2024a). Is table retrieval a solved problem? Exploring join-aware multi-table retrieval. In Proceedings of ACL 2024.

Chen, P. B., Yang, D., Li, W., Wenz, F., Zhang, Y., Tatbul, N., Cafarella, M., Demiralp, Ç., & Stonebraker, M. (2024b). BEAVER: An enterprise benchmark for text-to-SQL. arXiv:2409.02038v3.

Cormack, G. V., Clarke, C. L. A., & Büttcher, S. (2009). Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In Proceedings of SIGIR 2009.

Eben, J., Ahmad, A., & Lau, S. (2025). RASL: Retrieval augmented schema linking for massive database text-to-SQL. KDD 2025 Workshop on Structured Knowledge for Large Language Models (SKnowLLM). arXiv:2507.23104.

Guo, J., Zhan, Z., Gao, Y., Xiao, Y., Lou, J.-G., Liu, T., & Zhang, D. (2019). Towards complex text-to-SQL in cross-domain database with intermediate representation. In Proceedings of ACL 2019.

Herzig, J., Müller, T., Krichene, S., & Eisenschlos, J. M. (2021). Open domain question answering over tables via dense retrieval. In Proceedings of NAACL 2021.

Jin, T., Choi, Y., Zhu, Y., & Kang, D. (2026). Pervasive annotation errors break text-to-SQL benchmarks and leaderboards. arXiv:2601.08778v3.

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. In Advances in Neural Information Processing Systems 30, 3146 to 3154.

Kothyari, M., Dhingra, D., Sarawagi, S., & Chakrabarti, S. (2023). CRUSH4SQL: Collective retrieval using schema hallucination for Text2SQL. In Proceedings of EMNLP 2023.

Li, H., Zhang, J., Li, C., & Chen, H. (2023). RESDSQL: Decoupling schema linking and skeleton parsing for text-to-SQL. In Proceedings of AAAI 2023.

Li, H., Zhang, J., Liu, H., Fan, J., Zhang, X., Zhu, J., Wei, R., Pan, H., Li, C., & Chen, H. (2024). CodeS: Towards building open-source language models for text-to-SQL. Proceedings of the ACM on Management of Data (SIGMOD 2024).

Li, J., Hui, B., Qu, G., Yang, J., Li, B., Li, B., Wang, B., Qin, B., Cao, R., Geng, R., Huo, N., Zhou, X., Ma, C., Li, G., Chang, K. C. C., Huang, F., Cheng, R., & Li, Y. (2023). Can LLM already serve as a database interface? A BIg bench for large-scale database grounded text-to-SQLs. In Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track.

Maamari, K., Abubaker, F., Jaroslawicz, D., & Mhedhbi, A. (2024). The death of schema linking? Text-to-SQL in the age of well-reasoned language models. arXiv:2408.07702.

Page, L., Brin, S., Motwani, R., & Winograd, T. (1999). The PageRank citation ranking: Bringing order to the web. Stanford InfoLab Technical Report.

Pourreza, M., & Rafiei, D. (2023). DIN-SQL: Decomposed in-context learning of text-to-SQL with self-correction. In Advances in Neural Information Processing Systems 36.

Qin, T., & Liu, T.-Y. (2013). Introducing LETOR 4.0 datasets. arXiv:1306.2597.

Qin, Z., Yan, L., Zhuang, H., Tay, Y., Pasumarthi, R. K., Wang, X., Bendersky, M., & Najork, M. (2021). Are neural rankers still outperformed by gradient boosted decision trees? In Proceedings of ICLR 2021.

Rao, A., & Gopinath, N. (2026). NIFT-SQL: Customer-supervised table scoring for enterprise text-to-SQL retrieval. arXiv preprint arXiv:2611.12731.

Robertson, S., & Zaragoza, H. (2009). The probabilistic relevance framework: BM25 and beyond. Foundations and Trends in Information Retrieval, 3(4), 333 to 389.

Sen, J., Lei, C., Quamar, A., Özcan, F., Efthymiou, V., Dalmia, A., Stager, G., Mittal, A., Saha, D., & Sankaranarayanan, K. (2020). ATHENA++: Natural language querying for complex nested SQL queries. Proceedings of the VLDB Endowment, 13(11), 2747 to 2759.

Singh, R., Boškov, N., Drabeck, L., Gudal, A., & Khan, M. A. (2026). CSR-RAG: An efficient retrieval system for text-to-SQL on the enterprise scale. arXiv:2601.06564.

Soliman, H., Gupta, V., Roth, D., & Gurevych, I. (2026). CORE-T: COherent REtrieval of Tables for Text-to-SQL. arXiv:2601.13111v3.

sqlglot (2024). SQL parser and transpiler. https://github.com/tobymao/sqlglot.

Talaei, S., Pourreza, M., Chang, Y.-C., Mirhoseini, A., & Saberi, A. (2024). CHESS: Contextual harnessing for efficient SQL synthesis. arXiv:2405.16755.

Wang, B., Shin, R., Liu, X., Polozov, O., & Richardson, M. (2020). RAT-SQL: Relation-aware schema encoding and linking for text-to-SQL parsers. In Proceedings of ACL 2020.

Wang, T., Chen, X., Lin, H., Han, X., Sun, L., Wang, H., & Zeng, Z. (2025). DBCopilot: Natural language querying over massive databases via schema routing. arXiv:2312.03463v3.

Wang, Y., Liu, P., & Yang, X. (2025). LinkAlign: Scalable schema linking for real-world large-scale multi-database text-to-SQL. In Proceedings of EMNLP 2025, 977 to 991.

Warner, B., Chaffin, A., Clavié, B., Weller, O., Hallström, O., Taghadouini, S., Gallagher, A., Biswas, R., Ladhak, F., Aarsen, T., Cooper, N., Adams, G., Howard, J., & Poli, I. (2024). Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference. arXiv:2412.13663.

Wretblad, N., Riseby, F. G., Biswas, R., Ahmadi, A., & Holmström, O. (2024). Understanding the effects of noise in text-to-SQL: An examination of the BIRD-Bench benchmark. In Proceedings of ACL 2024 (Volume 2: Short Papers), 356 to 369.

Xiao, S., Liu, Z., Zhang, P., Muennighoff, N., Lian, D., & Nie, J.-Y. (2024). C-Pack: Packed resources for general Chinese embeddings. In Proceedings of SIGIR 2024.

Yu, T., Zhang, R., Yang, K., Yasunaga, M., Wang, D., Li, Z., Ma, J., Li, I., Yao, Q., Roman, S., Zhang, Z., & Radev, D. (2018). Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task. In Proceedings of EMNLP 2018.

Zhang, S., & Balog, K. (2018). Ad hoc table retrieval using semantic similarity. In Proceedings of The Web Conference (WWW) 2018.

Zhang, Y., Li, M., Long, D., Zhang, X., Lin, H., Yang, B., Xie, P., Yang, A., Liu, D., Lin, J., Huang, F., & Zhou, J. (2025). Qwen3 Embedding: Advancing text embedding and reranking through foundation models. arXiv:2506.05176.

Appendix A: Design selection on BIRD train

Table A1. Candidate pool options, five-fold cross-validation (folds by database) on a 3,000-question subset of BIRD train. Top 100 tables by reciprocal rank fusion. "All six views" includes the best-column signal later removed.

Signals fused	R@100	C@100
BM25	93.7	87.2
Qwen3	98.8	97.5
BGE	98.4	96.8
BM25 + Qwen3	98.9	97.7
BM25 + BGE	98.4	96.5
Qwen3 + BGE (selected by the rule)	99.3	98.4
BM25 + Qwen3 + BGE (used; Section 5, exception 1)	99.2	98.3
all six views	99.3	98.4

Table A2. Removing one signal from the ranker, same cross-validation (complete@5 / complete@15). We also checked removing every "drop" signal together; we dropped only one signal.

Ranker	C@5	C@15	Decision
all signals	82.6	92.7	
without BM25, full card	82.1	92.9	keep
without BM25, name card	81.8	93.1	keep
without Qwen3, full card	81.9	92.5	keep
without Qwen3, name card	81.6	92.8	keep
without Qwen3, best column	82.4	92.8	drop
without BGE	82.2	92.7	keep

The selected and used pools in Table A1 differ by 0.1 points in both metrics, and in Table A2 each single-signal removal changes complete@5 and complete@15 by at most 1.0 point relative to the all-signals ranker; both are close to cross-validation noise with 3,000 questions, and their value is that a stated rule chose each component.

Appendix B: Development studies (history)

These studies were run during development on BIRD dev, in a configuration that differs from the final one: 139 features (the final families plus a fine-tuned decision model and the best-column signal), a BM25 + Qwen3 pool, rankers fitted on 3,000 BIRD training questions with out-of-fold decision-model scores, and, where stated, name-regex labels. They motivated the design; they do not test it.

Table B1. Development results on BIRD dev with evidence. "Stacked" is a stacking ensemble of a random forest (200 trees), a LightGBM classifier (500 trees), and a logistic regression, combined by logistic regression with three-fold cross-fitting.

Ranker	Labels	Features	R@15	C@5	C@15
hybrid search (BM25 + Qwen3)	sqlglot		94.2	69.4	88.5
12 score features, stacked	regex	12	97.9	87.0	96.1
all features with the decision model, stacked	regex	139	98.7	93.0	97.5
retrieval signals only, lambdarank	sqlglot	58	98.1	86.4	96.5
other features only, lambdarank	sqlglot	66	98.1	91.5	96.3
all features without the decision model, stacked	sqlglot	124	98.8	93.0	97.6
all features with the decision model, stacked	sqlglot	139	98.8	93.7	97.8

The 12 score features were the hybrid score and rank, the dense and BM25 scores, the decision model's score and rank, and six database features. The decision model was OpenDecider-nano, a 0.4B ModernBERT-based encoder (Warner et al., 2024), fine-tuned as a yes/no table scorer on 2,500 questions per fold. Including it changed recall@15 by +0.1 (95% paired interval -0.1 to +0.2), complete@15 by +0.3 (0.0 to +0.6), and complete@5 by +0.8 (+0.2 to +1.4), all computed from unrounded per-question values, at 336 of 570 ms per question; it was left out of FERN (Section 5, exception 2). BIRD is close to the ceiling at fifteen tables, and this scorer was trained on few questions, so this does not show that learned scorers are unhelpful in general; on BEAVER, retrievers trained on the target catalog recovered more complete table sets than FERN (Table 8).