IsharaBART: An efficient BART-Based Framework for Low-Resource Bangla Sign Language Translation

Aonmoy Das, Ananna Dev Aishi, Md. Fazlul Kader*

Department of Electrical and Electronic Engineering, University of Chittagong, Chittagong-4331, Bangladesh.

arXiv:2611.16359v1 [cs.CL] 12 Nov 2026

Abstract

Millions of Bangla Sign Language (BdSL) users lack reliable automatic translation tools. Existing systems are limited by data scarcity, unidirectional design, and computationally expensive architectures that are not suitable for low-resource deployment. We introduce IsharaBART, a bidirectional generative framework for gloss-sentence translation in BdSL. IsharaBART uses Bidirectional and Auto-Regressive Transformers (BART)-based conditional language modeling with denoising learning objectives. We developed a parallel dataset of 1,000 professionally annotated gloss-sentence pairs, balanced across tenses and structures. We expanded the dataset to 4,831 pairs through four augmentation techniques: rule-based tense transformations, masked token substitutions, back-translation, and one-shot prompting. A compact denoising sequence-to-sequence model was pre-trained from scratch on monolingual Bangla text from six different corpora, then fine-tuned on the annotated and augmented datasets. We compare IsharaBART against four larger baseline models under identical conditions, where it achieves comparable or superior bidirectional performance. On the gloss-to-sentence (G2S) task, it attains Bilingual Evaluation Understudy (BLEU-4) score up to 51.7, and on the sentence-to-gloss (S2G) task, it achieves 61.40, outperforming multilingual models despite having fewer than 10% of the parameters. Ablation studies demonstrate that pre-training provides 62.5–90.4% faster convergence and 11.5-118% higher BLEU-4 than scratch. Low-rank adaptation (LoRA) reduces trainable parameters by 99.3%, but with 28.0-52.2% loss of BLEU-4 compared to full-parameter fine-tuning. IsharaBART establishes a scalable and deployable pipeline for Bangla sign language bidirectional translation, enhancing accessibility in education, healthcare, and social contexts for the deaf and speech-impaired community, while providing a blueprint for similar low-resource languages like Bangla.

Keywords: Sign language translation, Bangla sign language, Gloss-sentence translation, Data augmentation, Transformer

*Corresponding author

Email addresses: aonmoy@std.cu.ac.bd (Aonmoy Das), aishi.eee@std.cu.ac.bd (Ananna Dev Aishi), f.kader@cu.ac.bd (Md. Fazlul Kader)

1. Introduction

1.1. Sign Language and Sign Language Translation

Sign language stands as one of the most ancient and instinctive method of communication, which serves as the primary mode of communication for people with hearing or speech impairments. According to the World Health Organization, approximately 466 million people worldwide (6.1% of the global population) experience hearing loss, with projections indicating this number will surpass 900 million by 2050 [1]. These sign language users face persistent difficulties in accessing digital services, educational resources, and social participation due to the lack of effective and automatic translation systems. Sign languages are independent linguistic systems with their own lexicons, grammar, and syntax, conveying meaning through manual features such as hand shape, movement, and hand position, together with non-manual features such as facial expression, mouthing, and eyebrow movement [2, 3]. Because sign languages rarely have a standardized written form, most technology built for the deaf community must work directly with video or motion data rather than with text, and this is the central technical difficulty that the present work addresses.

SLR methods have progressed from small-vocabulary classifiers trained on hand-crafted features and depth sensors, to deep networks that process body keypoints or raw video using convolutional and recurrent layers, and more recently to transformer-based encoders that better capture long-range temporal structure [4]. SLT methods build on this progress in two broad ways. Cascading approaches first recognize glosses from a video and then translate the gloss sequence into text, while end-to-end approaches translate directly from video features to text, sometimes still using glosses as an auxiliary training signal rather than an explicit intermediate step [4, 3].

1.2. Bangla Sign Language

For Bangla Sign Language (BdSL), which is spoken by the deaf and speech-impaired community in Bangladesh and parts of India, representing one of the world's biggest linguistic regions, the challenges are significantly serious, as technological solutions remain severely underdeveloped compared to other high-resource sign languages such as American Sign Language (ASL) [5], Chinese Sign Language (CSL) [6] or German Sign Language (DGS) [7]. Hearing loss affects a substantial share of the population in the region: recent estimates place significant hearing impairment in Bangladesh at close to 9.6% of the population, or approximately 13.7 million people [8].

BdSL is linguistically distinct from spoken Bangla, with its own sign order and grammatical structure. Therefore, gloss–sentence translation requires cross-representation generation rather than simple word reordering. BdSL is also considerably more low-resource than languages such as ASL, DGS, or CSL. Large parallel corpora with video, gloss, and text annotations, do not yet exist for BdSL, and most prior BdSL work has focused on isolated alphabets or a small set of static signs rather than sentence-level translation [8]. Most existing BdSL resources are small and recent. Regional dialects and individual differences in signing style add a further layer of variability, with the same word sometimes signed differently across signers and regions. On top of this, there is no widely agreed standard for how BdSL should be glossed, a problem BdSL shares with sign languages in general [3], which makes it difficult to combine data from different sources or compare results across studies.

1.3. Motivation and Research Gap

BdSL recognition and translation research has mostly been in very resource limited settings. Most of the past recognition works have focused on individual sign gestures of words, vowels and consonants, and numbers [9, 10, 11, 12]. These works have been highly effective on isolated sign recognition, but their approach do not address the recognition and translation of continuous sign language gestures. Recent works have started to address continuous translation [13, 8]. These studies are limited to small and shallow architectures with limited capacity to generalize. The datasets used in these studies are limited by very small vocabularies. Moreover, a fundamental requirement for building effective sign language communication tools is bidirectional translation. Past works do not address bidirectional translation and evaluation. To the best of our knowledge, these systems do not utilize modern pre-trained language models or use architectures that could enable both G2S and sentence-to-gloss (S2G) generation.

These research gaps come from three fundamental challenges in low resource languages like BdSL. First, BdSL does not yet have large scale gloss-sentence pair parallel dataset. Such high quality dataset are essential to training high performing models. Second, existing studies use large, resource heavy multilingual models. These models may not have optimal tokenization or scoped abilities specific to BdSL gloss-sentence translation. Third, real-world application of translation tools require small, compact models that are capable of running on edge devices and mobile platforms with limited computational resources.

1.4. Contributions of This Study

In this work, we introduce IsharaBART, a bidirectional generative framework specifically designed for BdSL gloss-sentence translation that addresses the aforementioned fundamental gaps through these key contributions:

1. We developed IsharaBART, a small (56M) denoising seq2seq model pre-trained on monolingual Bangla text from six different corpora. IsharaBART achieved competitive bidirectional gloss-sentence translation performance while using substantially fewer parameters (78-90% fewer) than the baseline models.
2. We developed a parallel gloss-sentence dataset for BdSL, which contains 1,000 professionally annotated gloss-sentence pairs. The data is balanced across tenses and syntactic structures. The dataset also contains 3,831 synthetically augmented pairs using four complementary augmentation techniques.
3. We proposed a comprehensive framework for sign language gloss-sentence data augmentation and data validation that expands our 1,000-pair dataset to 4,831 pairs while preserving structural, semantic, and lexical consistency.
4. We conducted a detailed bidirectional evaluation of IsharaBART by comparing against four baseline models and different architectures under controlled experimental conditions (identical datasets, training protocols, and evaluation metrics). It is demonstrated that monolingual pre-training enables dramatic convergence speedups (62.5-90.4% faster) in gloss-sentence translation tasks.
5. We performed ablation studies comparing full fine-tuning against Low-rank adaptation (LoRA) to provide relevant information for deployment scenarios with different computational resources.

1.5. Organization of This Paper

The remainder of this paper is structured in the following way. Section 2 provides a detailed literature review in general sign language and BdSL translation. Section 3 describes the methodology used in the study. Section 4 describes the experimental design and evaluation metrics. Section 5 reports and provides analysis of the results from the experiments and ablation studies. Section 5.7 further discusses the findings of the study and provides a more abstract implications of those results. Finally, Section 6 outlines the limitations of the study and provides potential future research directions to address those limitations.

2. Literature review

Research on automatic sign language processing is broadly categorized into two different tasks [3]. Sign Language Recognition (SLR) identifies signs in a video. It usually outputs these gestures as glosses or words in the order they appear in the video. Sign Language Translation (SLT) takes these sequence of glosses and outputs meaningful sentences to complete communication [3]. The reverse direction of generating gesture sequences from spoken texts also falls under this category.

2.1. SLR Evolution

SLR has progressed from classification of gestures by simpler hand centric feature extraction to more complex and deeper multimodal deep learning [14]. Early works used depth-based CNNs for fingerspelling [15], CNNs for 3D pose estimation of hands [16], RGB-depth-skeleton fusion [17], and restricted Boltzmann machines for fusing multidimentional features [18]. Later studies used 3D residual CNNs and BiLSTMs for continuous sign recognition [19]. Hasanov et al. studied isolated Azerbaijani fingerspelling [20]. Aiswarya et al. used keyframe selection to process continuous sign videos [21]. Hao et al. combined sensing-based CNNs with BiLSTM-based attention [22]. Renjith et al. used Transformers to combine motion and visual features [23]. Barrister et al. used multiscale Transformer and convolutional features to detect Mizo signs [24]. Wu et al. added sign cues to trimodal speech processing [25]. Talaat et al. developed a sign-language translation system that generates Arabic avatars in real time [26]. However, most studies focus on isolated-sign recognition or specific applications. Few address translation in both directions between sign-language glosses and sentences.

2.2. Transformer-Based SLT

Researchers moved from CNN-LSTM models to transformer-based models for continuous SLR and translation. Camgoz et al. [27] used transformers for end-to-end sign language translation. They tested their model on the RWTH-PHOENIX-Weather-2014T dataset. The dataset included 1,066 unique signs from 9 signers in unconstrained settings. Their model performed better than earlier gloss-to-text methods. It improved the BLEU-4 score from 9.58 to 21.80. Chaudhary et al. [28] proposed SignNet II. It uses a transformer-based architecture for two-way sign language translation. The model used a dual learning mechanism to train sign-to-text (SLT) and text-to-sign (TSL) networks together. This approach performed better than singly-trained networks and achieved noteworthy improvement in

BLEU-1 to BLEU-4 scores in both translation directions. Hao et al. [29] proposed a Convolutional Vision Transformer (CvT) model combining hierarchical convolutional tokenization with transformer self-attention mechanisms. They evaluated the model on two publicly available sign language datasets; a digit dataset (1,712 images, 10 classes) and an alphabet-symbol dataset (87,000 images, 29 classes). The CvT model achieved 99% accuracy, outperforming traditional CNNs and competing vision transformer (ViT) models. In [30], the authors proposed a transformer-based SLT system that relies exclusively on 2D body keypoints extracted by MediaPipe Holistic instead of raw video frames or pixel embeddings. They introduced spatial–temporal augmentation on keypoints, convolutional embeddings and CTC loss with positional encoding for Sign2Gloss. They also fine-tuned BART for Gloss2Text. On Phoenix14T, they achieved competitive gloss-based BLEU-4 of 21.30 and gloss-free BLEU-4 of 15.70 (test set), while on How2Sign the gloss-free model reached only 3.87 BLEU-4. Although the keypoint-only pipeline significantly reduces computational cost and improves deployability, performance remains limited by the narrow weather-news domain of the datasets. Jin et al. [31] proposed a sign language production (SLP) method that employs Locality Sensitive Hashing-Clustering to automatically group similar embedded word vectors. This helps the model capture long-distance dependencies in text-to-pose sequence generation. They also used multi-scale convolutional feature extraction module to capture local gloss features. Their transformer-based architecture was evaluated on the PHOENIX14T dataset and demonstrated improvements over baseline SLP methods. However, the model is restricted to a single high-resource dataset and focuses solely on the production direction, leaving the inverse G2S translation task and low-resource sign languages entirely unaddressed.

2.3. BdSL-Specific Research and Challenges

Research on BdSL recognition has progressed from CNN-based isolated alphabets to transformer-based real-time translation systems, establishing critical datasets and benchmarks for the Bangladeshi Deaf community. Authors in [10] trained a CNN on the Ishara-Bachon digit dataset, achieving 95% accuracy for classifying BdSL digits. This model was among the first deep-learning models for BdSL and used grayscale digit images. The primary limitation was that it was tested only on static, isolated digit images from a small dataset. Hossen et al. [9] developed a transfer-learning deep convolutional neural network (DCNN) to recognize 37 static BdSL alphabet signs. Using 1,147 images of 31 signers, fine-tuning a pre-trained network yielded 96.3% training accuracy but only 84.7% on validation. The model showed signs of overfitting, as indicated by the substantial difference between its training and validation accuracies. In [32], the researchers applied transfer learning with YOLOv5 for BdSL recognition. They extracted frames from a large BdSL video collection and annotated 24.2 GB of data. Using pre-trained YOLOv5, the system achieved 91.6% average classification accuracy, but the performance depends on video frame quality and annotation density. Deep detectors were deployed on an edge device in [11]. Using the Okkhornama BdSL image set (3,760 images, 49 classes), their YOLOv7-Tiny model achieved real-time inference on Nano. The limitation was that the dataset was small, and modest mAP meant many classes were still missed. Raihan et al. [12] proposed a CNN with squeeze-and-excitation (CNN–SE) for BdSL recognition. They trained on the KU-BdSL alphabet dataset (1,500 images of 30 signs) and incorporated a smartphone app. This model achieved 99.86% test

Table 1: Evolution of SLR and SLT: from hand-centric to transformer-based approaches.

Study	Year	Approach	Dataset/Language	Key Achievement	Limitations
Phase I: Hand-Centric & Multi-Modal Foundations (2015-2019)					
Kang et al. [15]	2015	CNN + Depth Maps	31 ASL alphabets/digits	99.99% (same signer), 83.58-85.49% (cross-subject)	Isolated gestures only; no temporal continuity
Oberweger et al. [16]	2015	CNN for 3D Joints	Depth-based hand pose	10-20mm mean joint error	Requires depth calibration; fails with occlusions
Huang et al. [17]	2015	3D CNN Multi-Modal	RGB-Depth-Skeleton	94.2% accuracy; first multi-modal fusion	Static feature fusion; limited to short sequences
Rastgoo et al. [18]	2018	RBM Fusion	Massey/NYU/ASL datasets	11-28% improvement over prior work	Superseded by LSTM-based methods
Wei et al. [19]	2019	3D ResNet + BiLSTM	CSL	95.2% precision; sentence-level translation	Rule-based grammar; limited scalability
Phase II: Transformer Revolution (2020-2023)					
Camgöz et al. [27]	2020	Sign Transformers	PHOENIX-2014T (1,066 signs)	21.80 BLEU-4 vs. 9.58 prior SOTA	Gloss-dependent; computationally intensive
Chaudhary et al. [28]	2023	SignNet II (Dual)	Bidirectional SLT-TSL	Outperforms unidirectional on all BLEU metrics	Requires parallel sign-text training data
Hao et al. [29]	2025	CVT	SL digits (10 classes), alphabets (29 classes)	99% accuracy; beats CNNs and ViTs	Tested on isolated static images only
Phase III: BdSL-Specific Research (2018-2025)					
Hossen et al. [9]	2018	Transfer Learning DCNN	1,147 images, 31 signers (BdSL)	96.3% train, 84.7% validation	Severe overfitting; small dataset
Islam et al. [10]	2019	CNN	Ishara-Bachon digits	95% digit classification	Grayscale only; no dynamic gestures
Karim et al. [32]	2023	YOLOv5 Transfer	24.2GB BdSL video frames	91.6% average accuracy	Quality-dependent; annotation-intensive
Siddique et al. [11]	2023	YOLOv7-Tiny Edge	Okkhornama (3,760 images, 49 classes)	Real-time on Jetson Nano	Low mAP; many missed classes
Raihan et al. [12]	2024	CNN-SE	KU-BdSL (1,500 images, 30 signs)	99.86% test accuracy	Dataset overfitting; limited variability
Das et al. [13]	2025	Stacked LSTM	BdSL words (100 words, 500 videos)	89% test accuracy	LSTM scalability; no large language model (LLM) integration
Dhrubo et al. [8]	2025	Transformer + Landmarks	102 words, 1,200 sentences (BdSL)	Highest reported BdSL accuracy	Fixed vocabulary; landmark-dependent

accuracy on KU-BdSL. The near-perfect accuracy likely reflects limited variability in KU-BdSL, real-world performance might drop under varied lighting or new signers. The authors in [13] constructed a custom BdSL dataset and proposed a deep LSTM-based BdSL translator that modeled temporal sign sequences using stacked LSTM layers without relying on gloss-level supervision. Their system achieved strong translation accuracy on constrained sentence vocabularies. However, the approach is limited by scalability issues inherent to LSTMs and does not leverage large pre-trained language models or Transformer-based attention mechanisms. In [8], researchers developed a real-time sign-to-text BdSL translator. They curated a 102-word dataset with 1,200 sentences and extracted hand/body keypoints. A custom Transformer model (with efficient landmark encoding) mapped sign sequences to Bangla text, reportedly achieving the highest accuracy to date on this benchmark. Key limitations are that the system uses a small, fixed vocabulary of everyday words and relies on robust landmark tracking.

The literature on SLT shows a clear trajectory as depicted in Table 1 from isolated sensor-based recognition and small static datasets toward transformer-based end-to-end translation, but critical gaps remain for low-resource languages. Especially, the work on BdSL has largely been confined to small, isolated image/video datasets. Most studies focus on sign-to-text translation in one direction. They do not use large pre-trained language models to support both gloss-to-sentence and sentence-to-gloss translation. Some studies use simple models to process sign sequences. Others use large multilingual models. These models may not work well on devices with limited resources.

Our work addresses these problems. We develop a bidirectional generative framework for Bangla Sign Language (BdSL). We use a Bidirectional and Autoregressive Transformer (BART) model and we pre-train the model on multiple large Bangla text corpus. The model is then fine-tuned for both G2S and S2G translation direction. For pre-training we use conditional language modeling with specific denoising learning objectives. Our approach goes further than simple sequence mapping and uses training method that is very similar to the structural properties of BdSL. Our approach also attemps to reduce vocabulary limits and overfitting in earlier LSTM and transformer models. We describe each step, from data preparation to model training and evaluation. Our primary goal is to build a lightweight system that can produce gloss-sentence translation with high quality in both directions.

3. Methodology

In this section, we discuss the proposed methodology for the successful development of a language model for a bidirectional gloss-sentence translation system targeted towards BdSL. The proposed methodology to research and develop the primary language model responsible for sign language gloss to sentence synthesis is depicted in Figure 1, and a step-by-step description of the methodology is provided in the following subsections.

3.1. Data Acquisition

3.1.1. Pre-training Dataset

The pre-training corpus was assembled from six complementary Bangla text collections: Sangraha [33], OSCAR (Bangla) [34], C4 (Bangla) [35], Bangla Wikipedia [36], Bangla books review collection [37], and a Bangla dialogues [38] dataset. These sources were selected to provide broad coverage of written and conversational registers, genres, and domain styles necessary for robust language modeling in downstream sign-language tasks. To maintain efficient I/O for large-scale training, each cleaned dataset was tokenized and serialized into contiguous binary files. Separate files were produced for training and validation splits to enable streaming consumption by standard pre-training toolchains while preserving document boundaries. All corpora were tokenized with the trained tokenizer prior to serialization. The tokenized corpora are stored as contiguous binary streams (one stream per dataset per split) to support deterministic, high-throughput sampling during pre-training. A SentencePiece byte pair encoding tokenizer was trained from scratch on a 60 GB Bangla corpus. The resulting tokenizer has a vocabulary of 52,001 subwords, averaging 1.33 tokens per word (127.5% of ideal tokenization efficiency) and a character-to-token compression ratio of 2.96.

3.1.2. Fine-tuning Dataset

We created the fine-tuning dataset for G2S and S2G tasks in BdSL. First, we used Sketch Engine [39] and the bntenten21 corpus to find common Bangla words. We analyzed word frequency and part-of-speech (POS) labels. We ranked the words by frequency and checked their POS labels. We then selected 100 words. The list includes verbs (30%), adjectives (20%), nouns (20%), adverbs (10%), pronouns (10%), and conjunctions (10%).

We wrote ten sentences for each word. These sentences covered present, past, and future tenses in a 4:3:3 ratio. We also included different sentence structures. A professional BdSL interpreter reviewed all sentence pairs. We mainly used the BanglaTense [40] and

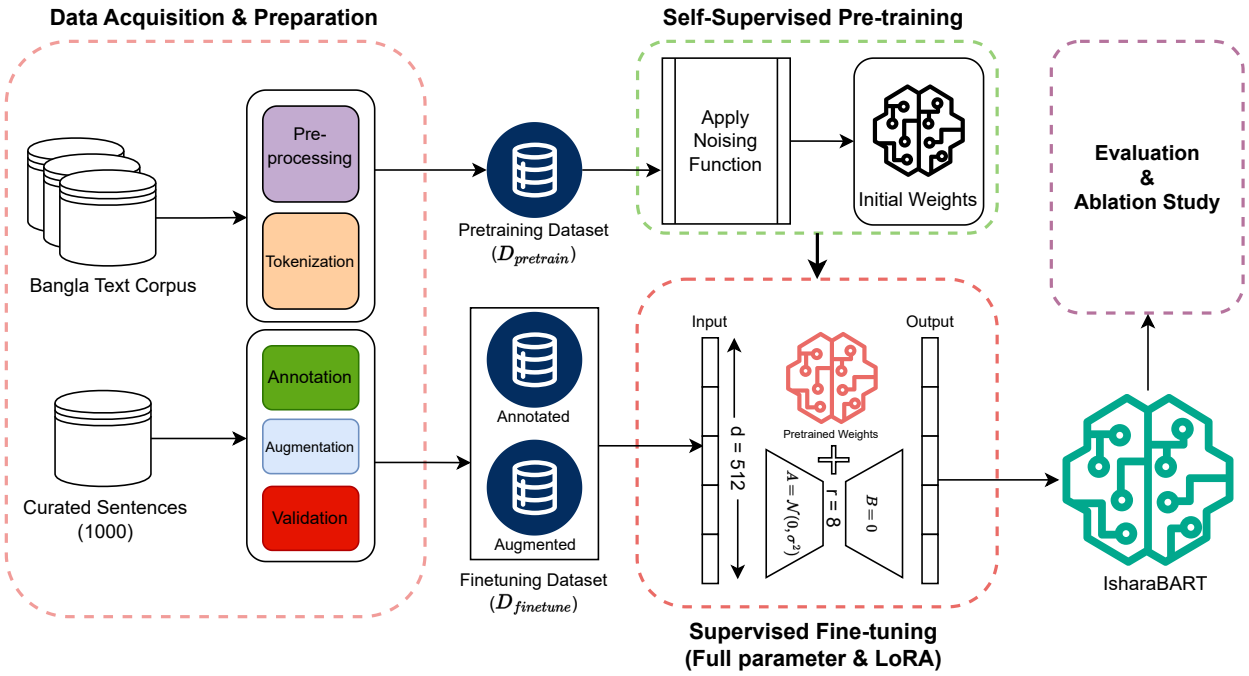

Figure 1: Proposed methodology for developing IsharaBART.

BanglaSarc [41] datasets to collect sentences. We modified some sentences to keep their grammar consistent and make them suitable for gloss annotation.

3.2. Data Augmentation

To address the lack of data in BdSL research, we used four complementary augmentation techniques to expand our annotated dataset of 1,000 gloss-sentence pairs. These methods were designed to introduce lexical, syntactic, and semantic variations while preserving the core G2S mappings essential for sequence-to-sequence learning. The augmented data was manually filtered for grammatical correctness and semantic coherence, which resulted in an augmented dataset size of 3,831 pairs and a total of 4831 pairs (annotated + augmented). The final dataset contains 1,000 annotated pairs (20.7%), 951 pairs from tense transformation (19.7%), 881 pairs from masked token substitution (18.2%), 999 pairs from back-translation (20.7%), and 1,000 pairs from one-shot prompting (20.7%). Each method contributes a similar number of pairs.

We use four methods to expand the annotated dataset. These methods are based on existing data augmentation techniques. We describe our changes below.

Rule-based tense transformation changes the tense of each Bangla sentence using predefined grammar rules. For example, it can change a present-tense sentence to future tense. We also update the matching gloss tokens to keep both sides aligned. Previous studies used rules to generate gloss data for sign language translation [5, 42]. However, they generated glosses from monolingual spoken text. Our method starts with an existing aligned sentence and gloss pair. It changes both sides of the pair. We designed the tense rules specifically for this study. **Masked token substitution** masks one content word (noun, verb, or adjective) in each sentence and replaces it with the top grammatically correct predic-

Table 2: Model and prompt template used for one-shot prompting augmentation.

Model	Provider	Details
llama-3.3-70b-versatile	Groq API	70B parameters; temperature=0.85 for diversity; max_tokens=150; top_p=1

Prompt Template

I will provide you an example of a BdSL Gloss and its corresponding Bangla Sentence.

Example Input:
Sentence: {Example Sentence}
Gloss: {Example Gloss}

Task: Generate a completely NEW Gloss and Sentence pair.

1. Structure: Follow the same grammatical style as the example.
2. Content: The meaning MUST be different (change the subject, object, or verb).
3. Format: Output ONLY the result in this exact format:
Sentence: [New Sentence]
Gloss: [New Gloss]

tion of BanglaBERT [43], updating the corresponding gloss token. This follows contextual augmentation, where words are replaced by those a language model predicts from context [44], and its use on gloss data with a masked language model [45]. **Back-translation paraphrasing** translates each sentence to English with NLLB-200 [46] (1.3B parameters) and back to Bangla with IndicTrans2 [47] (1B parameters), and keeps the original gloss. Back-translation is a standard augmentation method for low-resource translation [48], and round-trip paraphrases paired with the original gloss have been used for Korean Sign Language [45]. Back-translation that requires a trained gloss-to-text model is unreliable when gloss data is scarce [42]. Our round trip only uses pre-trained text translation models, so it needs no gloss data. Failed translations and paraphrases that changed the meaning were removed during manual filtering. **One-shot prompt-based synthesis** uses Llama 3.3-70B [49] to generate new gloss-sentence pairs. Each prompt contains a single annotated pair and asks for a new pair with altered content and a similar grammatical style (Table 2). Prompting an LLM with seed examples to create new training data follows [50] and has been used to augment data for low-resource languages [51]. Ours is a minimal one-example variant. We use a 70B model because this kind of generation is expected to work better with larger models [50]. Since such outputs can contain invalid examples [50, 51], all generated pairs were parsed and validated for BdSL compatibility. Table 3 provides illustrative examples from each technique.

Table 3: Examples of the augmentation strategies used to generate additional gloss–sentence pairs. Colored text marks the token substituted by each technique (red: original, blue: replacement).

Technique	Step-by-Step Process and Example
Rule-Based Transformation	**Original:** Sentence: তুমি তাড়াতাড়ি কাজটি করো। Gloss: তুমি তাড়াতাড়ি কাজ এটি কর **Step 1:** Identify verb tense/mood (করো - imperative present) **Step 2:** Apply tense transformation rule (imperative → future) **Step 3:** Transform verb: করো → করবে **Step 4:** Update gloss accordingly: কর → করা হবে **Augmented:** Sentence: তুমি তাড়াতাড়ি কাজটি করবে। Gloss: তুমি তাড়াতাড়ি কাজ এটি করা হবে
Masked Token Substitution	**Original:** Sentence: সে পরিশ্রমে ভাল রেজাল্ট করেছে। Gloss: সে কাজ পরিশ্রম অনেক ভাল রেজাল্ট পাওয়া **Step 1:** Randomly select a content word (রেজাল্ট) **Step 2:** Mask the token: সে পরিশ্রমে ভাল [MASK] করেছে। **Step 3:** Use BanglaBERT to predict: [MASK] → ফল **Step 4:** Substitute in sentence and update gloss **Augmented:** Sentence: সে পরিশ্রমে ভাল ফল করেছে। Gloss: সে কাজ পরিশ্রম ভাল ফল পাওয়া
Back-Translation	**Original:** Sentence: তার সাথে আমার প্রায়ই দেখা হয়। **Step 1:** Translate to English using NLLB200: "I see him often." **Step 2:** Translate back to Bangla using IndicTrans2: আমি প্রায়ই তাকে দেখি। **Step 3:** Generate gloss for paraphrased sentence **Augmented:** Sentence: আমি প্রায়ই তাকে দেখি। (paraphrased) Gloss: তার সাথে আমার হঠাৎ হঠাৎ দেখা হয়
One-Shot Prompt	**Seed Example:** Sentence: তুমি তাড়াতাড়ি কাজটি করো। Gloss: তুমি তাড়াতাড়ি কাজ এটি কর **Step 1:** Provide seed pair to LLM (Llama-3.3-70B) **Step 2:** Prompt LLM **Step 3:** LLM generates by changing subject/object/verb **Generated:** Sentence: সে ধীরে ধীরে খাবারটি খায়। Gloss: সে ধীরে ধীরে খাবার এটি খা

3.3. Dataset Validation

We first evaluate the quality of the fine-tuning dataset through statistical analysis. The parameters we mainly analyse are sentence length, gloss-to-text length ratios, semantic embeddings, and lexical diversity.

The annotated and augmented datasets have very similar lengths. The mean number of words in the annotated and augmented gloss-sentence samples are 5.45 ($SD = 1.37$) and 5.51 ($SD = 1.55$) respectively (Figure 2a). Kernel density estimation (KDE) [52] shows that both datasets have similar gloss-to-text length ratio (GTR). The GTR distribution peaks at about 1.0 (Figure 2b). We also project the BanglaBERT embeddings [43] of the samples in both dataset into two dimensions using UMAP [53]. The annotated and augmented samples show a centroid separation of 1.04 units. The UMAP also shows an overlap of 25.0% (Figure 2c). These indicates small overlapping but distinct semantic characteristics. Heaps' Law [54], $V = K \cdot N^{\beta}$, shows vocabulary growth. The estimated β values are 0.622 for the annotated dataset (vocabulary: 1,027), 0.575 for the augmented dataset (1,859), and 0.553 for the combined dataset (Figure 2d).

3.3.1. Human Evaluation of Dataset

We performed a human quality check on 300 randomly sampled gloss–sentence pairs drawn from the annotated and augmented dataset. Two BdSL signers independently rated each pair for adequacy (0–100) and fluency (1=Bad, 2=Moderate, 3=Good). The resultant statistics are in Table 4. They both received professional BdSL training from Krishti, a non-profit organization. Regarding their educational background, one evaluator was pursuing a secondary school qualification. The other evaluator had completed their higher secondary education. The evaluation was conducted in person. A trained sign language interpreter was present during the evaluation. The interpreter helped the evaluators communicate and understand the instructions. Before starting the evaluation, we explained two terms, i.e. adequacy and fluency, to the evaluators. The evaluators assigned separate scores for adequacy and fluency on their respective scales. We recorded the scores given by the evaluators. For each sample, we calculated the average separately for adequacy and fluency. No third-party adjudication or any tie-breaking method was used to resolve the disagreements between two evaluators. The mean adequacy is 99.21 (SD = 7.11) with a median of 100. 99.3% of samples reach the high-quality adequacy threshold ($\geq 75\%$). Fluency checks are equally positive. The overall mean fluency is 2.99 (SD = 0.11) and 98.0% of items were rated "Good" by both evaluators. When considering both criteria, 97.7% of pairs satisfy adequacy $\geq 75\%$ and fluency = Good. Inter-annotator agreement analysis supports the reliability of these ratings. Pearson correlation for adequacy is $r = 0.728$ ($p<0.001$) and the mean absolute error is only 0.75 points on the 0–100 DA scale. For fluency, weighted Cohen's $\kappa = 0.495$ with exact agreement 98.7% and adjacent agreement 100%, indicating moderate κ. The intraclass correlation (ICC) for adequacy is 0.465, consistent with significant consistency given the large sample.

3.4. Initial Model Selection

For the G2S and S2G synthesis tasks, we selected a compact encoder–decoder BART [55] model. The architecture was configured from scratch to be small (56M parameters) and monolingual for Bangla, rather than adopting large pre-trained checkpoints. BART

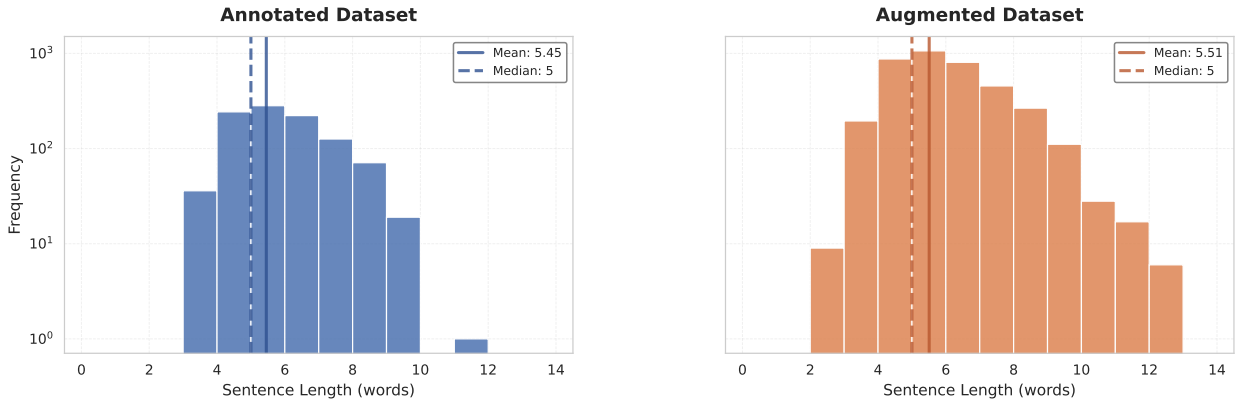

(a) Sentence length distributions.

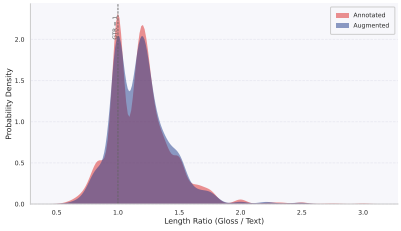

(b) Gloss-to-text length ratios.

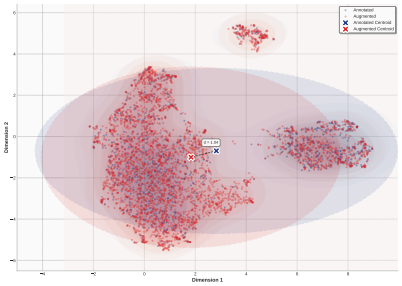

(c) Semantic space projection.

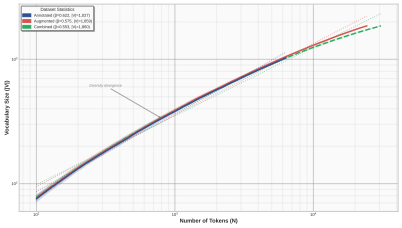

(d) Lexical diversity under Heaps' Law.

Figure 2: Dataset validation using (a) sentence length distributions, (b) gloss-to-text length ratios, (c) UMAP projections of BanglaBERT embeddings, and (d) Heaps' Law vocabulary growth.

(encoder–decoder, denoising) is explicitly designed for sequence-to-sequence tasks and the denoising training is particularly helpful for learning robust mappings between two modalities of text (glosses and sentences). Although pre-trained multilingual and English-centric checkpoints (e.g., large GPT or BART variants) offer strong priors, we deliberately refrain from using them for this study because pre-trained multilingual models contain many languages and suboptimal tokenization/representations for Bangla. Training a compact, monolingual model from scratch allows us to tailor the vocabulary, subword segmentation, and inductive bias specifically to Bangla morphology and the glossing conventions used in our annotation pipeline. Large pre-trained models are memory and compute-intensive at inference, complicating deployment on resource-constrained hardware (edge devices, low-cost servers). Our target use-cases (interactive annotation tools and real-time synthesis) require small, low-latency models. Table 5 provides a comprehensive summary of the architecture of the model selected.

3.5. Pre-training

Pre-training was carried out from scratch on the six tokenized corpora. To prevent large sources from dominating the pre-training we employed temperature-controlled sampling [56] across the six corpora. Let s_i denote the (byte or token) size of corpus i and $S = \sum_j s_j$. The base sampling probability is

$$p_i = \frac{s_i}{S}, \tag{1}$$

Table 4: Human evaluation results for the annotated + augmented dataset (300 samples). [a]Adequacy $\geq 75\%$ AND Fluency = Good (both evaluators). r = Pearson correlation, κ_w = Weighted Cohen's Kappa, ρ = Spearman's rho. MAE = Mean Absolute Error, Exact = Exact agreement percentage. $^*p < 0.05$, $^{**}p < 0.01$, $^{***}p < 0.001$.

Metric	Evaluator 1	Evaluator 2	Overall	Agreement
Adequacy				
Mean ± SD	98.83 ± 8.09	99.58 ± 5.95	99.21 ± 7.11	$r = 0.728^{***}$
Median	100	100	100	MAE = 0.75
High Quality ($\geq 75\%$)	298 (99.3%)	298 (99.3%)	298 (99.3%)	—
Fluency				
Mean ± SD	2.98 ± 0.14	2.99 ± 0.08	2.99 ± 0.11	$\kappa_w = 0.495$
Median	3.0	3.0	3.0	$\rho = 0.573^{***}$
Fluency Distribution				
Bad (1)	0 (0.0%)	0 (0.0%)	0 (0.0%)	—
Moderate (2)	6 (2.0%)	2 (0.7%)	8 (1.3%)	—
Good (3)	294 (98.0%)	298 (99.3%)	592 (98.7% of 600 ratings)	Exact: 98.7%
Overall Quality Assessment				
Both criteria met[a]	293 (97.7%)		Adeq–Flu corr.	$\rho = 0.391^{***}$

and the temperature-adjusted probability used to draw training blocks

$$\tilde{p}_i \;=\; \frac{p_i^{\,1/T}}{\sum_j p_j^{\,1/T}}, \tag{2}$$

where $T > 0$ is the sampling temperature (higher T flattens the distribution and increases sampling from smaller corpora).

3.5.1. General Pre-training Framework

Let $X = (x_1, x_2, \ldots, x_T)$ denote the original Bangla token sequence sampled from the tokenized pre-training data $\mathcal{D}$, where $x_t \in \mathcal{V}$ is a token from the vocabulary. We define a stochastic corruption function $C : \mathcal{V}^T \to \mathcal{V}^{T'}$, which is a composite function of document-level and token-level noising that transforms X into a noisy version $\tilde{X}$. The IsharaBART model parameterized by θ, is trained as a denoising autoencoder to reconstruct the original X given $\tilde{X}$. The training objective minimizes the negative log-likelihood

$$\mathcal{L}(\theta) = -\mathbb{E}_{x\sim\mathcal{D},\tilde{x}\sim C(x)} \left[\sum_{t=1}^{|x|} \log P(x_t \mid \tilde{x}, x_{<t}; \theta) \right] \tag{3}$$

Figure 3 illustrates the general pre-training framework of the initial model with a step-by-step depiction of the full workflow.

3.5.2. Document-Level Noising

Sentence permutation. To train the model to learn logical discourse and sentence relationships, we first apply sentence permutation. This is applied with probability 1.0 [55], shuffling sentences to disrupt global structure while preserving local coherence. Let the input

Table 5: Model configuration used for the gloss sequence to sentence modeling.

Property	**BART**
Model family	Seq2Seq denoising LM
Architecture	Encoder + decoder
Encoder layers	4
Decoder layers	4
Hidden dimension	512
Attention heads	8 / 8 (enc / dec)
FFN dimension	2048
Context size	512
Vocabulary	52001
Approx. parameters	∼56M
Training style	Denoising / conditional generation

sequence X be segmented into M sentences based on Bangla-specific delimiters (period and question mark). We define the segmented set as $\mathcal{S} = \{S_1, S_2, \ldots, S_M\}$. We sample a random permutation $\sigma \in \Sigma_M$ (the set of all permutations of length M) with probability p_{sent}. The document-level corrupted sequence $\tilde{X}_{doc}$ is obtained by concatenating the permuted sentences

$$\tilde{X}_{doc} = S_{\sigma(1)} \oplus S_{\sigma(2)} \oplus \cdots \oplus S_{\sigma(M)} \tag{4}$$

where $\oplus$ denotes sequence concatenation.

3.5.3. Token-Level Noising

After sentence permutation, we apply one of four distinct token-level noise transformations to $\tilde{X}_{doc}$. Let $Z = (z_1, \ldots, z_L)$ be the tokens of $\tilde{X}_{doc}$. We select a transformation $f \in \{f_{infill}, f_{del}, f_{mask}, f_{perm}\}$ based on a categorical distribution $\lambda_{choice} = [0.3, 0.2, 0.2, 0.3]$.

Text infilling (f_{infill}). This objective teaches the model to predict spans of missing text. We sample a set of span lengths from a Poisson distribution, where the length $l \sim \text{Poisson}(\lambda = 3)$. For a chosen span of tokens starting at index i with length l, denoted as $z_{i:i+l-1}$, we replace the entire span with a single unique mask token $[MASK]$. The corrupted sequence is

$$\tilde{X} = (\ldots, z_{i-1}, [MASK], z_{i+l}, \ldots) \tag{5}$$

Unlike standard single token masking, a span of length l is replaced by exactly one mask token, forcing the model to learn how many tokens are missing.

Token deletion (f_{del}). To improve the model's robustness against dropped inputs (common in SLT), we delete tokens without leaving a trace. Given a sequence of length L, we calculate the number of tokens to remove as $N = \lfloor p_{del} \cdot L \rfloor$. We uniformly sample a subset of indices $\mathcal{I}_{del} \subset \{1, \ldots, L\}$ of size N without replacement. The corrupted sequence is formed by retaining only the tokens whose indices are not in the deletion set

$$\tilde{X} = (x_t \mid t \notin \mathcal{I}_{del}) \tag{6}$$

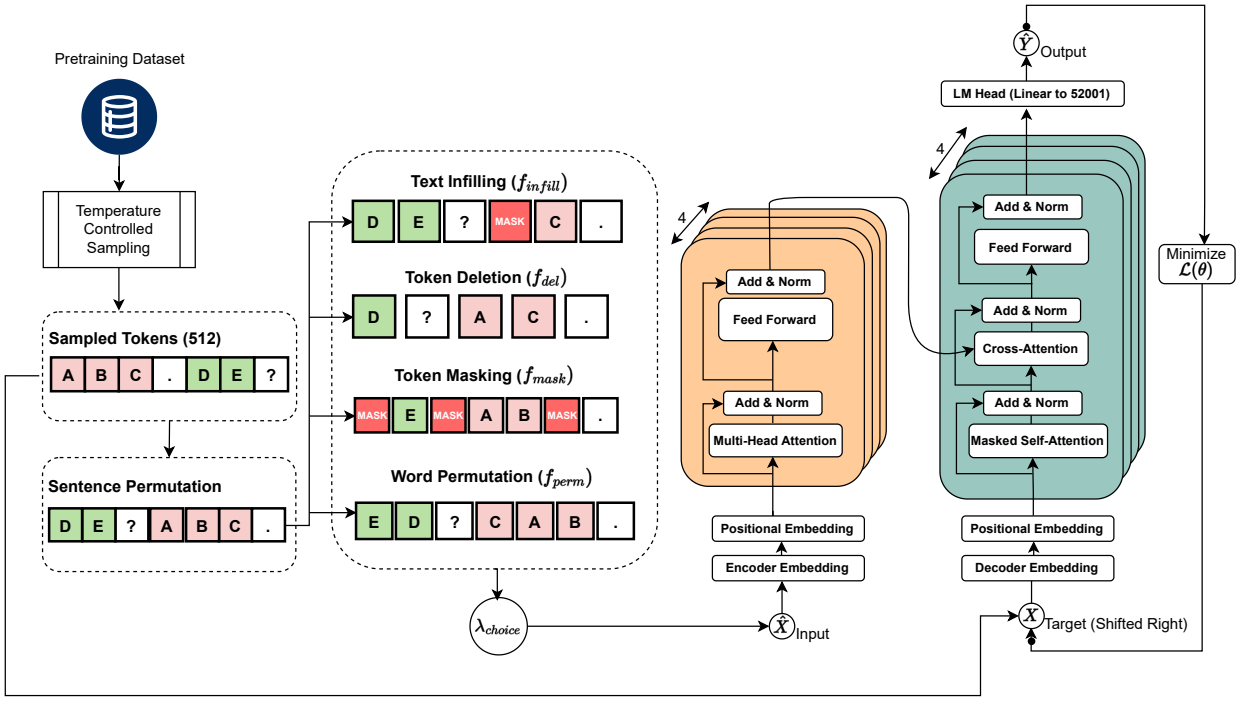

Figure 3: Proposed framework for the pre-training with denoising objectives.

The model learns to identify positions where inputs are missing and insert the correct tokens.

Token masking (f_{mask}). Following the standard single token masking objective, we randomly sample a set of indices $\mathcal{M}$ with probability p_{mask}. Tokens at these indices are replaced with the $[MASK]$ token

$$\tilde{x}_t = \begin{cases} [MASK] & \text{if } t \in \mathcal{M} \\ z_t & \text{otherwise} \end{cases} \tag{7}$$

Word permutation (f_{perm}). To model the flexible word order often found in sign language gloss sequence, we introduce an additional noising technique to shuffle words within sentence boundaries. For a given sentence $S_j \subset Z$ consisting of tokens $(w_1, \ldots, w_k)$, we sample a random permutation π of $\{1, \ldots, k\}$. The sentence is transformed to

$$\tilde{S}_j = (w_{\pi(1)}, w_{\pi(2)}, \ldots, w_{\pi(k)}) \tag{8}$$

Crucially, punctuation at the end of sentence is excluded from the shuffle to preserve sentence boundaries.

3.6. Fine-tuning

After pre-training, we fine-tune the pre-trained weights of IsharaBART on two separate downstream tasks: G2S translation and S2G translation, enabling bidirectional translation capability. For ablation, we employ two methods of fine-tuning: full parameter fine-tuning and LoRA [57]. For G2S task, we fine-tune by minimizing the negative log-likelihood of the

target Bangla sentence conditioned on the corresponding input gloss sequence and previously generated tokens, following standard sequence-to-sequence training for encoder-decoder language models. For the S2G task, the input is the Bangla sentence while the target becomes the corresponding gloss sequence. Given a parallel dataset $\mathcal{D}_{\text{finetune}}$ of gloss-sentence pairs (x, y), the model parameters θ are optimized by minimizing the expected token-level loss

$$\mathcal{L}_{\text{finetune}}(\theta) = -\mathbb{E}_{(x,y)\sim\mathcal{D}_{\text{finetune}}} \left[\sum_{i=1}^{|y|} \log P(y_i \mid x, y_{<i}; \theta) \right] \tag{9}$$

where x denotes the input token sequence, $y = (y_1, \ldots, y_{|y|})$ the corresponding target token sequence, and $P(y_i \mid x, y_{<i}; \theta)$ the model's conditional probability of the i-th target token given the source and all preceding target tokens.

4. Experimental Setup

We evaluate IsharaBART on bidirectional G2S and S2G translation tasks for BdSL, assessing the impact of data augmentation and comparing against strong multilingual and monolingual baselines. Our experiments are designed to answer two key questions: (1) Does pre-training a significantly smaller BART model with denoising objectives on monolingual Bangla text improve performance on low-resource downstream task of gloss translation? (2) How does synthetic data augmentation affect model generalization in this specialized domain?

4.1. Downstream Tasks

We fine-tune all models on two complementary translation directions. **G2S** takes a gloss sequence representing sign language structure as input and generates the corresponding natural Bangla sentence. **S2G** is the inverse task, mapping natural Bangla text to gloss sequences.

4.2. Baseline Models

We compare IsharaBART against four strong baseline architectures representing diverse pre-training objectives, architectural choices. **mBART-50** [58] is a multilingual denoising autoencoder with an encoder–decoder Transformer architecture and approximately 611M parameters in the mBART-large-50 configuration, pre-trained on 50 languages including Bangla using a BART-style denoising objective, making it a natural baseline for our monolingual BART-style pre-training while controlling for architecture and objective but varying pre-training language coverage. **mT5-Small** [59] is a multilingual text-to-text Transformer with an encoder–decoder architecture and roughly 300M parameters, pre-trained on a massively multilingual corpus using a span-corruption objective in the T5 framework. **BanglaT5** [38] is a monolingual adaptation of the T5 architecture pre-trained exclusively on Bangla corpora, with approximately 247M parameters, serving as a direct monolingual baseline that contrasts our BART-style denoising pre-training with a T5-style span-corruption objective under a Bangla-only setting and at a substantially larger parameter scale than IsharaBART. **BanglaGPT** [60] is a decoder-only autoregressive Transformer model for Bangla with 124M

trainable parameters, corresponding to a GPT-2-style medium configuration adapted to Bangla, enabling us to assess how a pure generative pre-training paradigm compares to encoder-decoder architectures on gloss translation at a smaller parameter scale.

4.3. Experimental Conditions

IsharaBART is pre-trained on the six monolingual Bangla corpora using the denoising autoencoder objective described in Section 3. We conduct four fine-tuning experiments for each of the five models: G2S and S2G, each on annotated data only and on annotated + augmented data, yielding a $5 \times 2 \times 2$ design (20 fine-tuned models) that enables comprehensive analysis of data augmentation effects and task asymmetry across all five architectures. In addition, for IsharaBART alone, we separately ablate monolingual pre-training by comparing the pre-trained model against an identical architecture trained from scratch, isolating the contribution of pre-training independent of the cross-model comparison above. Detailed hyperparameters are provided in Appendix A.

4.4. Evaluation Metrics

We report recall-oriented understudy for gisting evaluation (ROUGE-1/2/L) [61], BLEU-1/2/3/4 [62], and crosslingual optimized metric for evaluation of translation (COMET) [63] scores as our primary evaluation metrics. ROUGE checks how much of the critical information from human-written text is successfully captured in the generated text. BLEU evaluates how many words or phrases in the generated text match reference text. We complement them with COMET, which evaluates the semantic meaning and contextual quality of the generated text. Model selection is performed based on BLEU-4 on the validation set, and all reported results are on the held-out test set.

5. Results and Analysis

The experiments are designed to isolate three questions: whether monolingual Bangla pre-training benefits low-resource gloss–sentence translation, whether synthetic augmentation improves generalization, and whether the resulting model offers a favorable accuracy–efficiency trade-off. Across the controlled comparisons, three consistent findings emerge. First, monolingual pre-training substantially improves performance when only 1,000 annotated pairs are available and reduces the number of epochs required to reach peak validation BLEU-4. Second, augmentation improves translation quality in both directions, with a particularly large effect on S2G. Third, the 56M-parameter IsharaBART remains competitive with substantially larger baselines, achieving its strongest relative advantage on the bidirectional S2G task.

5.1. Pre-training Ablation and Convergence

We first isolate the contribution of monolingual Bangla pre-training by comparing IsharaBART initialized from the proposed pre-trained checkpoint with the same architecture trained from scratch under identical downstream conditions.

The pre-training of IsharaBART, which was conducted on monolingual Bangla text, showed fast convergence. This is shown in Figure 4. Convergence was achieved around 5,000 steps, as evidenced by the plateau in both losses. Further iterations yielded diminishing

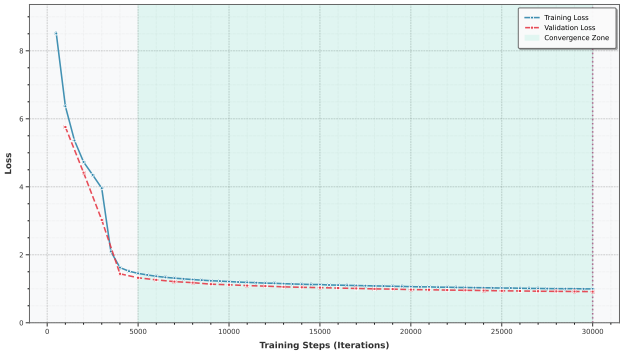

Figure 4: IsharaBART pre-training loss convergence over 30k steps. Training loss (blue) and validation loss (red) exhibit a rapid descent through 5k iterations, stabilizing with negligible improvement.

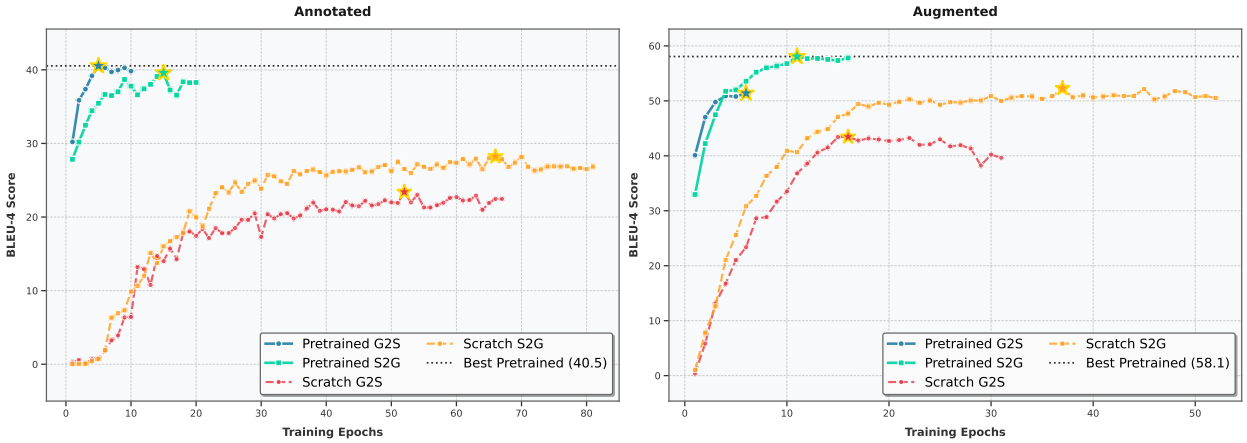

Figure 5: BLEU-4 convergence curves comparing pre-trained IsharaBART against from-scratch baselines across annotated and augmented datasets for bidirectional translation tasks. Stars indicate peak performance. The dotted horizontal line marks the best pre-trained performance for reference.

returns (<0.1 reduction per 1,000 steps). The pre-training was conducted for a total of 22.3 GPU hours on an NVIDIA A100 (80GB) with an effective batch size of 256. Given the limitation of computational resource, going beyond 30,000 steps was deemed impractical. It would have exceeded available resources without significant gains in performance.

On the annotated dataset, pre-training increases G2S BLEU-4 from 19.83 to 43.30 and S2G BLEU-4 from 23.80 to 39.39, corresponding to relative improvements of 118.4% and 65.5%, respectively. The gains remain after augmentation: BLEU-4 increases from 45.30 to 51.74 for G2S and from 55.10 to 61.40 for S2G. Thus, pre-training is most valuable when parallel gloss-sentence data is scarce, but its benefit continues even after the training set is expanded.

Pre-training also reduces the number of epochs required to reach peak validation BLEU-4 (Figure 5, Figure 6). On the annotated data, G2S reaches peak performance in 5 epochs compared with 52 for the scratch model, while S2G requires 15 rather than 66 epochs, corresponding to 90.4% and 77.3% fewer epochs. With augmented data, the corresponding reductions are 62.5% for G2S (6 vs. 16 epochs) and 70.3% for S2G (11 vs. 37 epochs). These results indicate that monolingual Bangla pre-training improves both the quality and the

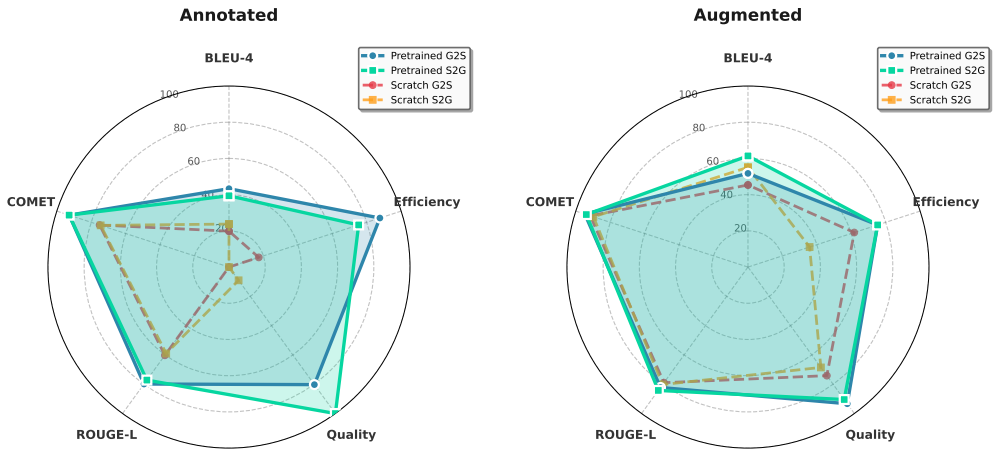

Figure 6: Multi-dimensional performance comparison between pre-trained IsharaBART and from-scratch baselines. The charts visualize five key metrics normalized to a 0–100 scale: BLEU-4, COMET, ROUGE-L, Quality (min-max normalized inverse loss), and Efficiency (percentage reduction in training epochs).

sample efficiency of downstream adaptation.

5.2. G2S Translation

On the annotated dataset, BanglaT5 achieves the highest scores (Table 6). It gets 48.91 BLEU-4, 0.95 COMET, and 81.71 ROUGE-L. IsharaBART gets 43.30 BLEU-4, 0.93 COMET, and 78.32 ROUGE-L. However, IsharaBART has only 56M parameters, while BanglaT5 has 247M.

Data augmentation reduces this gap. IsharaBART reaches 51.74 BLEU-4 and 0.95 COMET after augmentation. BanglaT5 reaches 52.24 BLEU-4. The difference is only 0.50 BLEU-4 points. IsharaBART also uses 77.3% fewer parameters. These results show that IsharaBART can achieve scores close to those of a much larger model with less computational capacity.

Augmentation raises IsharaBART's G2S BLEU-4 score from 43.30 to 51.74. This is an increase of 8.44 points. The baseline models also improve after augmentation. This suggests that the added data helps different models, not just IsharaBART. The results also show that synthetic data can help reduce the performance gap between smaller and larger models.

We also present qualitative examples from the test set to show the performance of IsharaBART on the G2S translation task. These examples are tabularized in Table 7.

5.3. S2G Translation

The S2G results show a different pattern (Table 8). On the annotated dataset, IsharaBART achieves the highest BLEU-4 (39.39) and COMET (0.94). It closely surpasses BanglaT5 at 39.29 BLEU-4 and 0.93 COMET. After augmentation, IsharaBART reaches 61.40 BLEU-4, the highest score among all five models, while mBART-50 reaches 60.82 and BanglaT5 51.57. IsharaBART therefore outperforms the 611M-parameter mBART-50 by 0.58 BLEU-4 while using approximately 91% fewer parameters.

The effect of augmentation is particularly strong for S2G. IsharaBART improves from 39.39 to 61.40 BLEU-4, a gain of 22.01 points. This is considerably larger than the corresponding G2S gain of 8.44 points. This suggests that additional lexical and structural

Table 6: Translation results for G2S on the annotated and augmented datasets.

Model	Params	R-1	R-2	R-L	B-1	B-2	B-3	B-4	COMET
G2S – Annotated Dataset (1,000 pairs)									
mBART-50	611M	77.86	63.36	76.63	65.15	54.64	46.39	39.93	0.92
mT5-Small	300M	77.76	62.37	76.20	63.78	52.78	44.76	38.76	0.90
BanglaT5	247M	**82.94**	**70.65**	**81.71**	72.31	**62.84**	**55.16**	**48.91**	**0.95**
BanglaGPT	124M	61.27	45.16	60.76	41.88	31.69	23.83	17.82	0.78
IsharaBART (ours)	56M	79.86	65.22	78.32	**72.73**	59.29	50.05	43.30	0.93
G2S – Augmented Dataset (4,831 pairs)									
mBART-50	611M	81.33	69.33	80.40	71.37	62.36	55.25	49.12	0.94
mT5-Small	300M	83.28	72.06	82.25	73.62	64.77	57.56	51.78	0.94
BanglaT5	247M	**83.74**	**72.78**	**82.60**	74.38	65.40	**58.20**	**52.24**	**0.95**
BanglaGPT	124M	72.36	57.72	71.29	56.87	46.32	37.97	31.26	0.88
IsharaBART (ours)	56M	82.22	70.51	80.95	**77.97**	**66.11**	57.98	51.74	**0.95**

variation is especially valuable when learning the more specialized mapping from natural Bangla sentences to gloss sequences. The difference between the two directions also indicates that S2G remains the more difficult task for parameter-efficient fine-tuning. We also present qualitative examples from the test set that are tabularized in Table 9.

IsharaBART demonstrates that competitive or superior performance is achievable with far fewer parameters (56M) than standard multilingual models (300-611M) when pre-training is carefully targeted to the task language. However, the task-dependent behavior suggests that no single architecture or pre-training strategy uniformly excels across all aspects of bidirectional gloss translation, and that the choice of model should be guided by the specific deployment scenario.

5.4. Human Evaluation

We conducted a human evaluation of IsharaBART in both G2S and S2G translation settings by the same two BdSL signers described in Section. They independently assessed 50 randomly sampled outputs for adequacy and fluency using a previously defined five point rating scale.A professionally trained sign language interpreter assisted us to facilitate communication. We conducted the evaluation in person and blind to the identity of the model. We calculated the mean of the two ratings for each criterion and evaluated sample. The results are summarized in Table 10. The results therefore reflect the average of the two evaluators' individual assessments.

Overall, IsharaBART demonstrates strong semantic quality across both directions. The system achieves a high mean adequacy score of 88.8, with a median of 100 and 84% of samples rated as high quality ($\geq$ 75%) (Table 10). Similarly, S2G translation achieves a mean adequacy of 87.8, with 76% of outputs meeting the high-quality threshold. These results indicate that the model consistently preserves the core meaning of BdSL inputs.

In both directions, the median fluency is 3 (Good), with 77% of G2S and 73% of S2G outputs rated as fluent. The slightly lower fluency in S2G translation is expected. Gloss generation is more constrained and sensitive to ordering and deletion. Importantly, adequacy–fluency correlations are strong and statistically significant in both settings ($\rho = 0.694$ for

Table 7: Qualitative G2S translation examples across three performance levels. English translations of input glosses are approximate and shown in uppercase.

Input (Gloss)	Ground Truth (Bangla)	Predicted (Bangla)
Excellent Performance		
তুমি তুমি আস্তে আস্তে বড় হবে, ভালো মানুষ কাজ কাজ চেষ্টা করবে	তুমি বড় হয়ে ভালো মানুষ হওয়ার চেষ্টা করবে।	তুমি বড় হয়ে ভালো মানুষ হওয়ার চেষ্টা করবে।
(YOU YOU SLOWLY SLOWLY GROW WILL, GOOD PERSON WORK WORK TRY WILL)	*(You will try to become a good person when you grow up.)*	*(You will try to become a good person when you grow up.)*
যখন তুমি আস্তে আস্তে বড় হবে, তখন তুমি কাজ কি হবে	যখন তুমি বড় হবে, তখন তুমি কি করবে?	যখন তুমি বড় হবে, তখন তুমি কি করবে?
(WHEN YOU SLOWLY SLOWLY GROW WILL, THEN YOU WORK WHAT WILL)	*(When you grow up, what will you do?)*	*(When you grow up, what will you do?)*
শিক্ষক ছাত্র নতুন বই পড়া বলা না	শিক্ষক ছাত্রদের নতুন বই পড়তে বলবেন না।	শিক্ষক ছাত্রদের নতুন বই পড়তে বলবেন না।
(TEACHER STUDENT NEW BOOK READ TELL NOT)	*(The teacher did not tell the students to read the new book.)*	*(The teacher did not tell the students to read the new book.)*
আমার হাত একটি চিঠি আছে	আমার কাছে একটি চিঠি আছে।	আমার কাছে একটি চিঠি আছে।
(MY HAND ONE LETTER HAVE)	*(I have a letter.)*	*(I have a letter.)*
ছাত্র খেলা জন্য মাঠ যাচ্ছিল	ছাত্ররা খেলার জন্য মাঠে যাচ্ছিল।	ছাত্ররা খেলার জন্য মাঠে যাচ্ছিল।
(STUDENT PLAY FOR FIELD GOING-PAST)	*(The students were going to the field to play.)*	*(The students were going to the field to play.)*
Moderate Performance		
তখন আকাশ মেঘ মেঘ ছিল না	তখন আকাশ মেঘলা ছিল না।	তখন আকাশ অন্ধকার ছিল না।
(THEN SKY CLOUD CLOUD WAS NOT)	*(The sky was not cloudy at that time.)*	*(The sky was not dark at that time.)*
তুমি তার বিষয়ে বোঝা সহজ	আপনি ওকে ভালো করে বোঝেন।	তুমি তাকে ভালোভাবে বুঝতে।
(YOU HIS/HER ABOUT UNDERSTAND EASY)	*(You understand him/her well.)*	*(You used to understand him/her well.)*
যখন সে কান্না ছিল তখন আমি তার হাত ধরা ছিল	যখন সে কাঁদছিল, আমি তার হাত ধরেছিলাম।	সে যখন কাঁদছিল, তখন আমি তার হাত ধরেছিলাম।
(WHEN HE/SHE CRY WAS, THEN I HIS/HER HAND HOLD WAS)	*(When he/she was crying, I was holding his/her hand.)*	*(When he/she was crying, then I was holding his/her hand.)*
আমি ক্লাস শেষ পরে খাওয়া হবে	ক্লাসের পর খাবো।	আমি ক্লাস শেষ করে খাব।
(I CLASS FINISH AFTER EAT WILL)	*(Will eat after class.)*	*(I will eat after finishing class.)*
নাফিস তার আগে অভ্যাস বদলানো চাওয়া ছিল	নাফিস তার অভ্যাস বদলাতে চেয়েছিল।	নাফিস তার অভ্যাস পরিবর্তন করতে চেয়েছিল।
(NAFIS HIS/HER BEFORE HABIT CHANGE WANT WAS)	*(Nafis wanted to change his habit.)*	*(Nafis wanted to change his habit.)*
Poor Performance (Limitations)		
প্রথমে আমি যাওয়া পরে পড়া বসা	প্রথমে আমি গেলাম, তারপর পড়তে বসলাম।	প্রথমে আমি পড়ি, তারপর পড়তে বসলাম।
(FIRST I GO AFTER STUDY SIT)	*(First, I went, then sat down to study.)*	*(First, I studied, then sat down to study.)*
আমি সংবাদ পত্র পড়ি	আমি খবরের কাগজ পড়ি।	আমি পত্র পড়ি।
(I NEWS PAPER READ)	*(I read the newspaper.)*	*(I read a letter.)*
আমি যাকে বন্ধু চিন্তা ছিল সে একজন প্রতারক	বন্ধু যাকে ভাবছি সে আসলে একজন প্রতারক।	যে বন্ধুকে আমি প্রতারক মনে করি।
(I WHO FRIEND THINK WAS, HE ONE FRAUD)	*(The friend I thought I knew was actually a fraud.)*	*(The friend whom I consider a fraud.)*
তুমি কম্পাস ওখানে একা চলা হবে	আপনি একা হেঁটে ক্যাম্পাসে যাবেন।	তুমি কম্পাস ওখানে চলবে একলা।
(YOU CAMPUS THERE ALONE WALK WILL)	*(You will go to the campus alone on foot.)*	*(You will go to campus alone.)*
তুমি বৃষ্টি থামা পরে কাজ শুরু করো	তুমি বৃষ্টি থামলে কাজ শুরু করে দাও।	আপনি বৃষ্টি থামিয়ে কাজ শুরু করুন।
(YOU RAIN STOP AFTER WORK START DO)	*(You will start working after the rain stops.)*	*(You will start working after stopping the rain.)*

G2S and $\rho = 0.636$ for S2G). This suggests that semantically accurate translations are also perceived as more fluent by human evaluators. Inter-annotator agreement ranges from moderate to substantial, with Pearson correlations of 0.574 (G2S) and 0.450 (S2G) for adequacy, and weighted kappa values of 0.364 and 0.443 for fluency. IsharaBART satisfies both criteria for 66% of G2S outputs and 56% of S2G outputs when jointly considering adequacy and fluency.

5.5. Fine-tuning Ablation Study

To evaluate parameter-efficient fine-tuning for SLT, we compared full fine-tuning against LoRA across all experimental configurations. Table 11 presents comprehensive results.

LoRA provides a significant parameter-efficiency advantage for IsharaBART. LoRA reduces the number of trainable parameters to approximately 0.7% of the full model. However, this reduction in trainable parameters comes with a consistent decrease in translation performance across all evaluated settings. As shown in Table 11, BLEU-4 decreases by 28.0–52.2% relative to full fine-tuning. The degradation is particularly significant for S2G, where on annotated data performance decreases from 39.39 to 22.15 BLEU-4 (-43.8%). On the augmented setting BLEU-4 decreases from 61.40 to 29.37 (-52.2%). In contrast, the corresponding G2S reductions are smaller, at 28.0% and 29.8%, respectively. Thus, the results indicate a clear accuracy–parameter-efficiency trade-off, with the cost of parameter-efficient adaptation depending strongly on the translation direction.

This behavior is consistent with recent empirical evidence that LoRA does not necessarily recover the performance of full fine-tuning when the downstream task requires substantial

Table 8: Translation results for S2G on the annotated and augmented datasets.

Model	Params	R-1	R-2	R-L	B-1	B-2	B-3	B-4	COMET
S2G – Annotated Dataset (1,000 pairs)									
mBART-50	611M	78.84	63.35	76.74	70.55	56.89	45.96	37.03	0.92
mT5-Small	300M	74.88	60.78	73.60	65.35	52.81	43.23	35.96	0.87
BanglaT5	247M	**80.06**	**65.41**	**78.54**	**73.60**	**60.20**	**49.09**	39.29	0.93
BanglaGPT	124M	58.60	43.38	58.01	45.97	33.23	23.54	16.19	0.77
IsharaBART (ours)	56M	77.16	63.01	75.78	69.08	56.71	47.18	**39.39**	**0.94**
S2G – Augmented Dataset (4,831 pairs)									
mBART-50	611M	**84.65**	**76.10**	**83.78**	80.91	72.60	66.16	60.82	0.94
mT5-Small	300M	83.70	74.85	82.78	79.67	71.19	64.42	58.92	0.93
BanglaT5	247M	82.46	71.65	81.01	77.35	66.82	58.42	51.57	0.94
BanglaGPT	124M	71.87	58.75	70.54	63.20	51.25	42.75	35.96	0.88
IsharaBART (ours)	56M	84.24	75.67	83.07	**81.06**	**72.93**	**66.55**	**61.40**	**0.94**

parameter adaptation. Under standard low-rank configurations, LoRA substantially underperforms full fine-tuning in several adaptation settings, and the parameter updates learned by full fine-tuning have ranks 10–100 times larger than typical LoRA configurations [64]. LoRA and full fine-tuning produce structurally different modifications to pretrained weights [65]. Vanilla LoRA can retain substantially less task-specific translation performance than full fine-tuning. This indicates that the LoRA constraint may prevent the model from making sufficiently expressive parameter updates [66, 67].

The difference in performance for different translation direction provides an additional insight. LoRA performs considerably better for G2S than for S2G. This suggests that the required task-specific adaptation may differ between the two directions. A plausible explanation is that G2S maps gloss sequences to natural Bangla sentences, an output space that is closer to the natural-language used during pre-training. In contrast, S2G maps natural Bangla sentences into gloss sequences, which has a more specialized and structurally difficult representation.

5.6. Dataset Comparison

To compare our dataset within the broader landscape of SLT resources, we compare it with prominent benchmarks and recent BdSL datasets. Table 12 summarizes key attributes such as size, domain, and SOTA performance metrics for G2S and S2G tasks where available. Our dataset has a modest scale of approximately 4,800 pairs (including augmentation). It achieves superior BLEU-4 scores (51.7 for G2S and 61.4 for S2G in augmented data) compared to larger benchmarks. This highlights the efficiency of our approach in low-resource settings, particularly for under-resourced languages like BdSL. In terms of BdSL, datasets like IsharaKhobor and Bangla-SGP provide valuable but domain-limited parallels.

5.7. Discussion

The results show that monolingual Bangla pre-training and data augmentation can improve bidirectional gloss-sentence translation. However, IsharaBART does not outperform all baselines on every metric. On G2S, BanglaT5 achieves higher BLEU-2, BLEU-3, and

Table 9: Qualitative S2G translation examples across three performance levels. English translations of input sentences are shown below the Bangla, and gloss translations are approximate uppercase renderings.

Input (Bangla)	Ground Truth (Gloss)	Predicted (Gloss)
Excellent Performance		
সে কাজে হাত দেয়ার আগেই কাজ শেষহয়ে গিয়েছিল।	সে কাজ হাত শুরু আগেই কাজ শেষ ছিল	সে কাজ হাত শুরু আগেই কাজ শেষ ছিল
(He/she had finished the work before even starting it.)	*(HE/SHE WORK HAND START BEFORE WORK FINISHED WAS)*	*(HE/SHE WORK HAND START BEFORE WORK FINISHED WAS)*
শিক্ষক ছাত্রদের নতুন বই পড়তে বলবেন না।	শিক্ষক ছাত্র নতুন বই পড়া বলা না	শিক্ষক ছাত্র নতুন বই পড়া হবে না
(The teacher will not ask the students to read the new book.)	*(TEACHER STUDENT NEW BOOK READ TELL NOT)*	*(TEACHER STUDENT NEW BOOK READ WILL NOT)*
ডাক্তার কোনো নতুন ওষুধ দেবেন না।	ডাক্তার কোন নতুন ওষুধ লিখে দেওয়া না	ডাক্তার কোন নতুন ওষুধ লিখে দেওয়া না
(The doctor will not prescribe any new medicine.)	*(DOCTOR ANY NEW MEDICINE PRESCRIBE GIVE NOT)*	*(DOCTOR ANY NEW MEDICINE PRESCRIBE GIVE NOT)*
নাফিস তার অভ্যাস বদলাতে চেয়েছিল।	নাফিস তার আগে অভ্যাস বদলানো চাওয়া ছিল	নাফিস তার আগে অভ্যাস বদলানো চাওয়া ছিল
(Nafis wanted to change his habit.)	*(NAFIS HIS/HER BEFORE HABIT CHANGE WANT WAS)*	*(NAFIS HIS/HER BEFORE HABIT CHANGE WANT WAS)*
যখন সে কাঁদছিল, আমি তার হাত ধরেছিলাম।	যখন সে কান্না ছিল তখন আমি তার হাত ধরা ছিল	যখন সে কান্না ছিল তখন আমি তার হাত ধরা ছিল
(When he/she was crying, I was holding his/her hand.)	*(WHEN HE/SHE CRY WAS THEN I HIS/HER HAND HOLD WAS)*	*(WHEN HE/SHE CRY WAS THEN I HIS/HER HAND HOLD WAS)*
Moderate Performance		
যখন তুমি বড় হবে, তখন তুমি কি করবে?	যখন তুমি আস্তে আস্তে বড় হবে, তখন তুমি কাজ কি হবে	যখন তুমি আস্তে আস্তে বড় হবে, তখন তুমি কি কাজ হবে
(When you grow up, what will you do?)	*(WHEN YOU SLOWLY SLOWLY GROW WILL THEN YOU WORK WHAT WILL)*	*(WHEN YOU SLOWLY SLOWLY GROW WILL THEN YOU WHAT WORK WILL)*
আমাকে ঐ রেস্তোরাঁয় খেতে মানা করেছে।	আমাকে ওই রেস্তোরাঁয় খাওয়া করা নিষেধ বলা হয়েছে	আমাকে ওই রেস্তোরাঁয় খাওয়া নিষেধ বলা হয়েছে
(I was told not to eat at that restaurant.)	*(TO ME THAT RESTAURANT EAT DO FORBIDDEN SAY HAS BEEN)*	*(TO ME THAT RESTAURANT EAT FORBIDDEN SAY HAS BEEN)*
তুমি আমাকে ছেড়ে চলে যেতে পারবে না।	তুমি আমাকে বাদ যাওয়া পারবে না	তুমি আমাকে বাদ চলে যাওয়া পারবে না
(You cannot leave me.)	*(YOU ME LEAVE GO CAN NOT)*	*(YOU ME LEAVE GO GO CAN NOT)*
আমি বুঝতে পারলাম না আপনি কি বলতে চাইছেন।	আমি বুঝা নাই আপনি কি বলা চাওয়া ছিল	আমি বুঝা পারা না আপনি কি বলা চাওয়া আছে
(I could not understand what you wanted to say.)	*(I UNDERSTAND NOT YOU WHAT SAY WANT WAS)*	*(I UNDERSTAND CAN NOT YOU WHAT SAY WANT IS)*
আপনি আসার সময়ে আমি ডিনার প্রস্তুত করে রেখেছিলাম।	আপনি আসার আগে আমি খাবার তৈরি করে রাখা ছিল	আপনি আসার আগে আমি রাখ খাবার তৈরি করে রেখেছিলাম
(I had prepared dinner by the time you arrived.)	*(YOU ARRIVE BEFORE I FOOD PREPARE KEEP WAS)*	*(YOU ARRIVE BEFORE I KEEP FOOD PREPARE KEEP WAS)*
Poor Performance (Limitations)		
আমরা ছোটবেলায় পুলিশ সাজতাম।	আমরা ছোট সময় পুলিশ সাজা ছিল	আমরা পুলিশ অফিসার ছিল
(We used to dress up as police officers when we were kids.)	*(WE CHILDHOOD POLICE DRESS UP USED TO)*	*(WE POLICE OFFICER WAS)*
ক্লাসের পর খাবো।	আমি ক্লাস শেষ পরে খাওয়া হবে	ক্লাস ভিতর কি হবে
(I will eat after class.)	*(I CLASS FINISH AFTER EAT WILL)*	*(CLASS INSIDE WHAT WILL)*
তুমি বৃষ্টি থামলে কাজ শুরু করে দাও।	তুমি বৃষ্টি থামা পরে কাজ শুরু করো	তুমি বৃষ্টি থামা পরে তাকে শুরু করো
(You will start working after the rain stops.)	*(YOU RAIN STOP AFTER WORK START DO)*	*(YOU RAIN STOP AFTER THAT START DO)*
শিক্ষার্থীদের বই পড়ুন নিয়মিত।	শিক্ষার্থী বই পড় নিয়মিত নিয়মিত	শিক্ষার্থী ওরা বই পড়া প্রতিদিন
(Read the students' books regularly.)	*(STUDENT BOOK READ REGULARLY REGULARLY)*	*(STUDENT THEY BOOK READ EVERY DAY)*
আমাকে মা জামাটা কিনে দিতো, যদি আমি চাইতাম।	যদি আমি জামা কেনা চাই, মা আমাকে জামা কিনে দেওয়া হবে	যদি আমি জামা চাই মা আমাকে জামা কিনে দেওয়া হবে
(My mother would buy me the dress if I wanted it.)	*(IF I DRESS BUY WANT MOTHER ME DRESS BUY GIVE WILL)*	*(IF I DRESS WANT MOTHER ME DRESS BUY GIVE WILL)*

BLEU-4 scores. After augmentation, IsharaBART reaches a similar BLEU-4 score (51.74 versus 52.24). However, IsharaBART uses 56M parameters, while BanglaT5 uses 247M. On S2G, IsharaBART achieves the highest BLEU-4 score among the evaluated models in both data settings. With augmented data, it reaches 61.40 BLEU-4. The score for mBART-50 is 60.82, despite its 611M parameters. These results show that IsharaBART offers a good balance between model size and translation performance.

The pre-training ablation provides further insight. Pre-training gives the largest relative gains when the model has only 1,000 annotated pairs. These gains remain after data augmentation. Pre-training also reduces the number of epochs needed to reach the best validation BLEU-4 score. These results suggest that monolingual Bangla pre-training helps translation when labeled data are limited.

Data augmentation also improves performance, but the gains differ by translation direction. For IsharaBART, BLEU-4 increases by 8.44 points on G2S and 22.01 points on S2G. One possible reason is that S2G requires the model to generate gloss sequences from natural Bangla sentences. The model must learn specific output patterns for this task. The annotated dataset may not provide enough examples of these patterns. Augmentation may help the model learn them. However, this explanation remains a hypothesis. The experiments do not identify the exact linguistic features behind the larger S2G gain.

This study extends previous BdSL research by evaluating both G2S and S2G translation in the same experimental framework. It also examines monolingual pre-training and targeted data augmentation. The results show that the two directions can produce different outcomes.

Table 10: Human evaluation of IsharaBART for G2S and S2G translation.Two BdSL signers evaluated 50 randomly sampled translations on adequacy (0–100 scale) and fluency (1=Bad, 2=Moderate, 3=Good). [a]Adequacy $\geq$ 75% AND Fluency = Good (both evaluators). r = Pearson correlation, κ_w = Weighted Cohen's Kappa, ρ = Spearman's rho. MAE = Mean Absolute Error, Exact = Exact agreement percentage. $^*p < 0.05$, $^{**}p < 0.01$, $^{***}p < 0.001$.

	G2S				S2G			
Metric	**Eval. 1**	**Eval. 2**	**Overall**	**Agreement**	**Eval. 1**	**Eval. 2**	**Overall**	**Agreement**
Adequacy								
Mean ± SD	85.0 ± 30.3	92.5 ± 22.7	88.8 ± 26.9	$r = 0.574^{***}$	83.5 ± 30.1	92.0 ± 21.1	87.8 ± 26.2	$r = 0.450^{***}$
Median	100	100	100	MAE = 11.5	100	100	100	MAE = 12.5
High quality ($\geq$ 75%)	42 (84%)	46 (92%)	42 (84%)	—	38 (76%)	38 (76%)	38 (76%)	—
Fluency								
Mean ± SD	2.64 ± 0.66	2.78 ± 0.46	2.71 ± 0.57	$\kappa_w = 0.364^*$	2.66 ± 0.69	2.56 ± 0.70	2.61 ± 0.69	$\kappa_w = 0.443$
Median	3.0	3.0	3.0	$\rho = 0.358^*$	3.0	3.0	3.0	$\rho = 0.460^{***}$
Fluency distribution								
Bad (1)	5 (10%)	1 (2%)	6 (6%)	—	6 (12%)	6 (12%)	12 (12%)	—
Moderate (2)	8 (16%)	9 (18%)	17 (17%)	—	5 (10%)	10 (20%)	15 (15%)	—
Good (3)	37 (74%)	40 (80%)	77 (77%)	Exact: 78%	39 (78%)	34 (68%)	73 (73%)	Exact: 70%
Overall quality								
Both criteria met[a]	—	—	33 (66%)	$\rho = 0.694^{***}$	—	—	28 (56%)	$\rho = 0.636^{***}$

After augmentation, IsharaBART comes close to BanglaT5 on G2S. It achieves the highest S2G BLEU-4 score among the evaluated models. These differences show why studies should evaluate both directions. Results from one direction may not reflect performance on the other. Model selection should also consider the target task and available computing resources.

Human evaluation provides further insight into translation quality. The mean adequacy scores are 88.8 for G2S and 87.8 for S2G. However, only 66% of G2S outputs and 56% of S2G outputs meet both adequacy and fluency criteria. This suggests that the model often preserves the intended meaning. However, producing fluent and consistent translations remains difficult. Inter-rater agreement was modest to moderate. The evaluation also included only 50 outputs per direction. These results show promise, but they do not establish reliable translation in every case.

The LoRA experiments show another trade-off. LoRA reduces the number of trainable parameters by 99.3%. However, BLEU-4 decreases by 28.0–52.2% compared with full fine-tuning. The decline is greater for S2G than for G2S under the tested configuration. This suggests that LoRA does not fully match full fine-tuning in these experiments.

6. Conclusion and Future Work

In this study, we presented IsharaBART, a bidirectional translation model for BdSL. It translates between gloss sequences and Bangla sentences in both directions. Our results show that monolingual pre-training and data augmentation improve translation quality. The model also reduces training time by 62.5% to 90.4%. It maintains good structural, semantic, and lexical consistency. We compared IsharaBART with four baseline models under the same conditions. The results show that our approach achieves comparable or better translation quality. Its small size and faster training may support accessibility tools for education, healthcare, and daily communication.

However, this study has some limitations. The annotated and augmented dataset is small compared to high-resource benchmarks. This may limit the model's ability to handle rare

Table 11: Comparison of full fine-tuning and low-rank adaptation (LoRA) for IsharaBART across annotated and augmented data. BLEU-4 and ROUGE-L are reported on a 0–100 scale, while COMET is reported on its native 0–1 scale. Changes are reported relative to full fine-tuning as absolute and relative differences; negative values indicate degradation under LoRA.

Dataset	Translation Direction	Fine-Tuning Strategy	BLEU-4	COMET	ROUGE-L
Annotated	G2S	Full fine-tuning	43.30	0.9279	78.32
		LoRA	31.17	0.8973	72.27
		Change	−12.13 *(−28.01%)*	−0.0306 *(−3.30%)*	−6.05 *(−7.72%)*
Annotated	S2G	Full fine-tuning	39.39	0.9283	75.78
		LoRA	22.15	0.9221	68.97
		Change	−17.24 *(−43.77%)*	−0.0062 *(−0.67%)*	−6.81 *(−8.99%)*
Augmented	G2S	Full fine-tuning	51.74	0.9471	80.95
		LoRA	36.31	0.9199	73.06
		Change	−15.43 *(−29.82%)*	−0.0272 *(−2.87%)*	−7.89 *(−9.75%)*
Augmented	S2G	Full fine-tuning	61.40	0.9388	83.07
		LoRA	29.37	0.9185	69.55
		Change	−32.03 *(−52.17%)*	−0.0203 *(−2.16%)*	−13.52 *(−16.27%)*

Table 12: Overview of major gloss-sentence datasets and associated SOTA BLEU-4 performance for G2S and S2G tasks. IsharaBART achieves substantially higher scores than prior Bangla BdSL resources despite comparable scale, highlighting the effectiveness of Bangla-specific pre-training and augmentation.

Dataset	Language	Size	Domain	G2S	S2G
PHOENIX-14T [7, 68]	German (DGS)	8,257	Weather	18.75	20.68
How2Sign [69]	English (ASL)	35,000	Instructional	8.0	—
iSign [70]	English (ISL)	118,000	Open	17.6	—
IsharaKhobor [71]	Bangla (BdSL)	5,642	News	—	—
Bangla-SGP [72]	Bangla (BdSL)	4,000	Open	27.3	—
Ours	**Bangla (BdSL)**	**4,831**	**Open**	**51.7**	**61.4**

sentence structures and glosses. The model also depends on accurate gloss sequences from a vision-based sign recognition system. We did not study this component. Recognition errors may occur due to poor lighting or differences between signers. Our monolingual approach may also miss benefits from multilingual training.

IsharaBART has a compact architecture and needs fewer training epochs. These features may help on devices with limited resources. However, we did not measure inference time, memory use, or energy consumption. We also did not test the model on mobile or edge devices. Its efficiency in real-world settings remains unclear.

The model translates gloss sequences and Bangla text. It does not process sign-language videos directly or generate visual signs. A complete translation system would need a sign recognition model and a method to generate sign-language output. IsharaBART could serve as the text translation component of such a system.

The approach may also work for other low-resource sign languages, but we have not tested this. Future work should expand the dataset and include more signers and regional gloss variations. Larger and more diverse groups of human evaluators should assess the

translations. Human evaluations of the baseline models would also show whether the automatic scores reflect real improvements. Tests of inference speed, memory use, and energy consumption are also needed. Integrating video-based sign recognition would help assess the model in real-world communication systems.

Future work should expand the dataset through annotations from the deaf community. We should also explore end-to-end video-to-text models that do not depend on gloss inputs. Hybrid multilingual and monolingual pre-training may help the model handle linguistic variation. Another direction is to test span-corruption and hybrid denoising objectives. These methods may help close the remaining BLEU-3 and BLEU-4 gaps in G2S translation, especially compared with BanglaT5. The goal is to improve translation quality without increasing model size or training time. Finally, real-world deployment studies could help measure the model's practical value and guide further improvements for other low-resource sign languages.

Disclosure statement

The author(s) declare(s) no potential conflict of interest regarding the publication of this paper.

Acknowledgment

This work was supported by the ICT fellowship (FY 2024–2025) under grant code 56.00.0000.052.33.0003.24–115 provided by the Information and Communication Technology Division, Ministry of Posts, Telecommunications and Information Technology, Government of the People's Republic of Bangladesh.

References

[1] B. O. Olusanya, A. C. Davis, H. J. Hoffman, Hearing loss grades and the international classification of functioning, disability and health, Bulletin of the World Health Organization 97 (10) (2019) 725—728. doi:10.2471/blt.19.230367.

[2] J. Guo, P. Li, T. Cohn, Bridging sign and spoken languages: Pseudo gloss generation for sign language translation, in: Advances in Neural Information Processing Systems 38, Neural Information Processing Systems Foundation, Inc. (NeurIPS), San Diego, California, USA and Mexico City, Mexico, 2025, pp. 86134—86162. doi:10.52202/085713-2598.

[3] M. De Coster, D. Shterionov, M. Van Herreweghe, J. Dambre, Machine translation from signed to spoken languages: state of the art and challenges, Universal Access in the Information Society 23 (3) (2024) 1305—1331. doi:10.1007/s10209-023-00992-1.

[4] S. Tan, N. Khan, Z. An, Y. Ando, R. Kawakami, K. Nakadai, A review of deep learning-based approaches to sign language processing, Advanced Robotics 38 (23) (2024) 1649––1667. doi:10.1080/01691864.2024.2442721.

[5] A. Othman, M. Jemni, English-ASL gloss parallel corpus 2012: ASLG-PC12, in: Proceedings of the LREC2012 5th Workshop on the Representation and Processing of Sign Languages: Interactions between Corpus and Lexicon, European Language Resources Association (ELRA), Istanbul, Turkey, 2012, pp. 151–154.

[6] H. Zhou, W. Zhou, W. Qi, J. Pu, H. Li, Improving sign language translation with monolingual data by sign back-translation, in: 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), IEEE, Nashville, TN, USA, 2021, pp. 1316—1325. doi:10.1109/CVPR46437.2021.00137.

[7] N. C. Camgoz, S. Hadfield, O. Koller, H. Ney, R. Bowden, Neural sign language translation, in: 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition, IEEE, Salt Lake City, UT, 2018, pp. 7784—7793. doi:10.1109/CVPR.2018.00812.

[8] T. H. Dhrubo, M. T. Reza, S. Sahariar, W. A. Mukto, Transformer based sign-to-text translation for bangladeshi sign language, Scientific Reports 16 (1) (2026) 1154. doi:10.1038/s41598-025-30856-y.

[9] M. Hossen, A. Govindaiah, S. Sultana, A. Bhuiyan, Bengali sign language recognition using deep convolutional neural network, in: 2018 Joint 7th International Conference on Informatics, Electronics & Vision (ICIEV) and 2018 2nd International Conference on Imaging, Vision & Pattern Recognition (icIVPR), IEEE, Kitakyushu, Japan, 2018, pp. 369—373. doi:10.1109/ICIEV.2018.8640962.

[10] M. S. Islam, S. S. S. Mousumi, N. A. Jessan, A. S. A. Rabby, S. Abujar, S. A. Hossain, Ishara-bochon: The first multipurpose open access dataset for bangla sign language isolated digits, in: Recent Trends in Image Processing and Pattern Recognition, Vol. 1035, Springer Singapore, Singapore, 2019, pp. 420—428. doi:10.1007/978-981-13-9181-1_37.

[11] S. Siddique, S. Islam, E. E. Neon, T. Sabbir, I. T. Naheen, R. Khan, Deep learning-based bangla sign language detection with an edge device, Intelligent Systems with Applications 18 (2023) 200224. doi:10.1016/j.iswa.2023.200224.

[12] M. J. Raihan, M. I. Labib, A. A. J. Jim, J. J. Tiang, U. Biswas, A.-A. Nahid, Bengali-sign: A machine learning-based bengali sign language interpretation for deaf and non-verbal people, Sensors 24 (16) (2024) 5351. doi:10.3390/s24165351.

[13] A. Das, A. D. Aishi, M. U. Toha, M. F. Kader, Bangla sign language translator for deaf and speech impaired people using deep lstm, International Journal of Speech Technology 28 (2) (2025) 351—368. doi:10.1007/s10772-025-10184-2.

[14] Y. Zhang, Y. Han, Z. Zhu, X. Jiang, Y. Zhang, Artificial intelligence in sign language recognition: A comprehensive bibliometric and visual analysis, Computers and Electrical Engineering 120 (2024) 109854. doi:10.1016/j.compeleceng.2024.109854.

[15] B. Kang, S. Tripathi, T. Q. Nguyen, Real-time sign language fingerspelling recognition using convolutional neural networks from depth map, in: 2015 3rd IAPR Asian Conference on Pattern Recognition (ACPR), IEEE, Kuala Lumpur, Malaysia, 2015, pp. 136—140. doi:10.1109/ACPR.2015.7486481.

[16] M. Oberweger, P. Wohlhart, V. Lepetit, Training a feedback loop for hand pose estimation, in: 2015 IEEE International Conference on Computer Vision (ICCV), IEEE, Santiago, Chile, 2015, pp. 3316—3324. doi:10.1109/ICCV.2015.379.

[17] J. Huang, W. Zhou, H. Li, W. Li, Sign language recognition using 3d convolutional neural networks, in: 2015 IEEE International Conference on Multimedia and Expo (ICME), IEEE, Turin, Italy, 2015, pp. 1—6. doi:10.1109/ICME.2015.7177428.

[18] R. Rastgoo, K. Kiani, S. Escalera, Multi-modal deep hand sign language recognition in still images using restricted boltzmann machine, Entropy 20 (11) (2018) 809. doi:10.3390/e20110809.

[19] C. Wei, W. Zhou, J. Pu, H. Li, Deep grammatical multi-classifier for continuous sign language recognition, in: 2019 IEEE Fifth International Conference on Multimedia Big Data (BigMM), IEEE, Singapore, Singapore, 2019, pp. 435—442. doi:10.1109/BigMM.2019.00027.

[20] J. Hasanov, N. Alishzade, A. Nazimzade, S. Dadashzade, T. Tahirov, Development of a hybrid word recognition system and dataset for the azerbaijani sign language dactyl alphabet, Speech Communication 153 (2023) 102960. doi:10.1016/j.specom.2023.102960.

[21] M. S. Aiswarya, R. Arockia Xavier Annie, Keyframe extraction algorithm for continuous sign-language videos using angular displacement and sequence check metrics, International Journal of Intelligent Systems 2024 (2024) 4725216. doi:10.1155/2024/4725216.

[22] W. Hao, C. Hou, Z. Zhang, X. Zhai, L. Wang, G. Lv, A sensing data and deep learning-based sign language recognition approach, Computers and Electrical Engineering 118 (2024) 109339. doi:10.1016/j.compeleceng.2024.109339.

[23] S. Renjith, A. Varghese, M. Rashmi, P. S.S., Transformer-based motion-visual integrated fusion for isolated sign language recognition, Computers and Electrical Engineering 130 (2026) 110902. doi:10.1016/j.compeleceng.2025.110902.

[24] B. R, A. K. V. D, A. K. V. D, Mizo hand sign language detection using a multi-scale transformer-based hybrid feature extractor and fusion network, Computers and Electrical Engineering 131 (2026) 110916. doi:10.1016/j.compeleceng.2025.110916.

[25] J. Wu, X. Li, T. Li, F. Meng, Y. Kong, G. Yang, L. Senhadji, H. Shu, Cslnspeech: Solving the extended speech separation problem with the help of chinese sign language, Speech Communication 165 (2024) 103131. doi:10.1016/j.specom.2024.103131.

[26] F. M. Talaat, W. El-Shafai, N. F. Soliman, A. D. Algarni, F. E. Abd El-Samie, A. I. Siam, Real-time arabic avatar for deaf-mute communication enabled by deep learning sign language translation, Computers and Electrical Engineering 119 (2024) 109475. doi:10.1016/j.compeleceng.2024.109475.

[27] N. C. Camgoz, O. Koller, S. Hadfield, R. Bowden, Sign language transformers: Joint end-to-end sign language recognition and translation, in: 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), IEEE, Seattle, WA, USA, 2020, pp. 10020—10030. doi:10.1109/CVPR42600.2020.01004.

[28] L. Chaudhary, T. Ananthanarayana, E. Hoq, I. Nwogu, Signnet ii: A transformer-based two-way sign language translation model, IEEE Transactions on Pattern Analysis and Machine Intelligence 45 (11) (2023) 12896–12907. doi:10.1109/TPAMI.2022.3232389.

[29] J. Hao, H. Pan, A novel deep transformer based cvt model for sign language recognition in visual communication, Scientific Reports 16 (1) (2026) 1823. doi:10.1038/s41598-025-31558-1.

[30] W. F. Maia, A. M. Lopes, S. A. David, Automatic sign language to text translation using mediapipe and transformer architectures, Neurocomputing 642 (2025) 130421. doi:10.1016/j.neucom.2025.130421.

[31] H. Jin, S. Zhang, Z. Yang, Q. Han, J. Cao, Exploiting locality sensitive hashing - clustering and gloss feature for sign language production, Speech Communication 171 (2025) 103227. doi:10.1016/j.specom.2025.103227.

[32] M. Karim, M. Ali, M. F. Hasan, M. Fatema-Tuj-Jahra, S. S. Sultana, D. M. Farid, Bangla sign language recognition using yolov5, in: 2023 IEEE 8th International Conference for Convergence in Technology (I2CT), IEEE, Lonavla, India, 2023, pp. 1—7. doi:10.1109/I2CT57861.2023.10126419.

[33] M. S. U. R. Khan, P. Mehta, A. Sankar, U. Kumaravelan, S. Doddapaneni, S. B, V. G, S. Jain, A. Kunchukuttan, P. Kumar, R. Dabre, M. Khapra, Indicllmsuite: A blueprint for creating pre-training and fine-tuning datasets for indian languages, in: Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Association for Computational Linguistics, Bangkok, Thailand, 2024, pp. 15831—15879. doi:10.18653/v1/2024.acl-long.843.

[34] J. Abadji, P. Ortiz Suarez, L. Romary, B. Sagot, Towards a cleaner document-oriented multilingual crawled corpus, in: Proceedings of the Thirteenth Language Resources and Evaluation Conference, European Language Resources Association, Marseille, France, 2022, pp. 4344–4355. doi:10.48550/ARXIV.2201.06642.

[35] J. Dodge, M. Sap, A. Marasović, W. Agnew, G. Ilharco, D. Groeneveld, M. Mitchell, M. Gardner, Documenting large webtext corpora: A case study on the colossal clean crawled corpus, in: Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, Association for Computational Linguistics, Online and

Punta Cana, Dominican Republic, 2021, pp. 1286—1305. doi:10.18653/v1/2021.emnlp-main.98.

[36] I. Johnson, L.-A. Kaffee, M. Redi, Wikimedia data for ai: a review of wikimedia datasets for nlp tasks and ai-assisted editing, in: Proceedings of the First Workshop on Advancing Natural Language Processing for Wikipedia, Association for Computational Linguistics, Miami, Florida, USA, 2024, pp. 91—101. doi:10.18653/v1/2024.wikinlp-1.14.

[37] M. Kabir, O. Bin Mahfuz, S. R. Raiyan, H. Mahmud, M. K. Hasan, Banglabook: A large-scale bangla dataset for sentiment analysis from book reviews, in: Findings of the Association for Computational Linguistics: ACL 2023, Association for Computational Linguistics, Toronto, Canada, 2023, pp. 1237—1247. doi:10.18653/v1/2023.findings-acl.80.

[38] A. Bhattacharjee, T. Hasan, W. U. Ahmad, R. Shahriyar, Banglanlg and banglat5: Benchmarks and resources for evaluating low-resource natural language generation in bangla, in: Findings of the Association for Computational Linguistics: EACL 2023, Association for Computational Linguistics, Dubrovnik, Croatia, 2023, pp. 726—735. doi:10.18653/v1/2023.findings-eacl.54.

[39] A. Kilgarriff, V. Baisa, J. Bušta, M. Jakubíček, V. Kovář, J. Michelfeit, P. Rychlý, V. Suchomel, The sketch engine: ten years on, Lexicography 1 (1) (2014) 7—36. doi:10.1007/s40607-014-0009-9.

[40] M. H. I. Bijoy, U. Ayman, M. M. Islam, Banglatense: A large-scale dataset of bangla sentences categorized by tense: Past, present, and future, Data in Brief 59 (2025) 111400. doi:10.1016/j.dib.2025.111400.

[41] T. S. Apon, R. Anan, E. A. Modhu, A. Suter, I. J. Sneha, M. G. R. Alam, Banglasarc: A dataset for sarcasm detection, in: 2022 IEEE Asia-Pacific Conference on Computer Science and Data Engineering (CSDE), IEEE, Gold Coast, Australia, 2022, pp. 1—5. doi:10.1109/CSDE56538.2022.10089322.

[42] A. Moryossef, K. Yin, G. Neubig, Y. Goldberg, Data augmentation for sign language gloss translation, in: Proceedings of the 1st International Workshop on Automatic Translation for Signed and Spoken Languages (AT4SSL), Association for Machine Translation in the Americas, 2021, pp. 1–11.

[43] A. Bhattacharjee, T. Hasan, W. Ahmad, K. S. Mubasshir, M. S. Islam, A. Iqbal, M. S. Rahman, R. Shahriyar, Banglabert: Language model pretraining and benchmarks for low-resource language understanding evaluation in bangla, in: Findings of the Association for Computational Linguistics: NAACL 2022, Association for Computational Linguistics, Seattle, United States, 2022, pp. 1318—1327. doi:10.18653/v1/2022.findings-naacl.98.

[44] S. Kobayashi, Contextual augmentation: Data augmentation by words with paradigmatic relations, in: Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technolo-

gies, Volume 2 (Short Papers), New Orleans, Louisiana, 2018, pp. 452–457. doi:10.18653/v1/N18-2072.

[45] J. Y. Jang, H. Park, S. Shin, S. Shin, B. Yoon, G. Gweon, Automatic gloss-level data augmentation for sign language translation, in: Proceedings of the Thirteenth Language Resources and Evaluation Conference, European Language Resources Association, Marseille, France, 2022, pp. 6808–6813.

[46] N. Team, M. R. Costa-jussa, J. Cross, et al., No language left behind: Scaling human-centered machine translation, arXiv preprint arXiv:2207.04672 (2022). doi:10.48550/arXiv.2207.04672.

[47] J. Gala, P. A. Chitale, A. K. Raghavan, V. Gumma, S. Doddapaneni, A. K. M, J. A. Nawale, A. Sujatha, R. Puduppully, V. Raghavan, P. Kumar, M. M. Khapra, R. Dabre, A. Kunchukuttan, Indictrans2: Towards high-quality and accessible machine translation models for all 22 scheduled indian languages, Transactions on Machine Learning Research (2023). doi:10.48550/arXiv.2305.16307.

[48] R. Sennrich, B. Haddow, A. Birch, Improving neural machine translation models with monolingual data, in: Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Berlin, Germany, 2016, pp. 86–96. doi:10.18653/v1/P16-1009.

[49] L. Team, The llama 3 herd of models, arXiv preprint arXiv:2407.21783 (2024). doi:10.48550/arXiv.2407.21783.

[50] Y. Wang, Y. Kordi, S. Mishra, A. Liu, N. A. Smith, D. Khashabi, H. Hajishirzi, Self-instruct: Aligning language models with self-generated instructions, in: Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Toronto, Canada, 2023, pp. 13484–13508. doi:10.18653/v1/2023.acl-long.754.

[51] C. Whitehouse, M. Choudhury, A. F. Aji, LLM-powered data augmentation for enhanced cross-lingual performance, in: Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, Singapore, 2023, pp. 671–686. doi:10.18653/v1/2023.emnlp-main.44.

[52] S. Węglarczyk, Kernel density estimation and its application, ITM Web of Conferences 23 (2018) 00037. doi:10.1051/itmconf/20182300037.

[53] L. McInnes, J. Healy, N. Saul, L. Großberger, Umap: Uniform manifold approximation and projection, Journal of Open Source Software 3 (29) (2018) 861. doi:10.21105/joss.00861.

[54] D. C. van Leijenhorst, T. P. van der Weide, A formal derivation of heaps' law, Information Sciences 170 (2–4) (2005) 263—272. doi:10.1016/j.ins.2004.03.006.

[55] M. Lewis, Y. Liu, N. Goyal, M. Ghazvininejad, A. Mohamed, O. Levy, V. Stoyanov, L. Zettlemoyer, Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension, in: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, Association for Computational Linguistics, Online, 2020, pp. 7871—7880. doi:10.18653/v1/2020.acl-main.703.

[56] A. CONNEAU, G. Lample, Cross-lingual language model pretraining, in: H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, R. Garnett (Eds.), Advances in Neural Information Processing Systems, Vol. 32, Curran Associates, Inc., 2019. doi:10.48550/arXiv.1901.07291.

[57] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, W. Chen, LoRA: Low-rank adaptation of large language models, in: The Tenth International Conference on Learning Representations, Virtual, 2022. doi:10.48550/arXiv.2106.09685.

[58] Y. Liu, J. Gu, N. Goyal, X. Li, S. Edunov, M. Ghazvininejad, M. Lewis, L. Zettlemoyer, Multilingual denoising pre-training for neural machine translation, Transactions of the Association for Computational Linguistics 8 (2020) 726—742. doi:10.1162/tacl_a_00343.

[59] L. Xue, N. Constant, A. Roberts, M. Kale, R. Al-Rfou, A. Siddhant, A. Barua, C. Raffel, mt5: A massively multilingual pre-trained text-to-text transformer, in: Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Association for Computational Linguistics, Online, 2021, pp. 483—498. doi:10.18653/v1/2021.naacl-main.41.

[60] M. S. Salim, H. Murad, D. Das, F. Ahmed, Banglagpt: A generative pretrained transformer-based model for bangla language, in: 2023 International Conference on Information and Communication Technology for Sustainable Development (ICICT4SD), IEEE, Dhaka, Bangladesh, 2023, pp. 56—59. doi:10.1109/ICICT4SD59951.2023.10303383.

[61] C.-Y. Lin, ROUGE: A package for automatic evaluation of summaries, in: Text Summarization Branches Out, Association for Computational Linguistics, Barcelona, Spain, 2004, pp. 74–81.

[62] K. Papineni, S. Roukos, T. Ward, W.-J. Zhu, Bleu: a method for automatic evaluation of machine translation, in: Proceedings of the 40th Annual Meeting on Association for Computational Linguistics - ACL '02, Association for Computational Linguistics, Philadelphia, Pennsylvania, 2002, pp. 311–318. doi:10.3115/1073083.1073135.

[63] R. Rei, C. Stewart, A. C. Farinha, A. Lavie, Comet: A neural framework for mt evaluation, in: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), Association for Computational Linguistics, Online, 2020, pp. 2685—2702. doi:10.18653/v1/2020.emnlp-main.213.

[64] D. Biderman, J. Portes, J. J. G. Ortiz, M. Paul, P. Greengard, C. Jennings, D. King, S. Havens, V. Chiley, J. Frankle, C. Blakeney, J. P. Cunningham, LoRA learns less and

forgets less, Transactions on Machine Learning Research (2024). doi:10.48550/arXiv.2405.09673.

[65] R. Shuttleworth, J. Andreas, A. Torralba, P. Sharma, Lora vs full fine-tuning: An illusion of equivalence, in: Advances in Neural Information Processing Systems 38, Neural Information Processing Systems Foundation, Inc. (NeurIPS), San Diego, California, USA and Mexico City, Mexico, 2025, pp. 193421—193456. doi:10.52202/085713-5808.

[66] T. Su, X. Peng, S. Thillainathan, D. Guzmán, S. Ranathunga, E.-S. Lee, Unlocking parameter-efficient fine-tuning for low-resource language translation, in: Findings of the Association for Computational Linguistics: NAACL 2024, Association for Computational Linguistics, Mexico City, Mexico, 2024, pp. 4217—4225. doi:10.18653/v1/2024.findings-naacl.263.

[67] Y. Zeng, K. Lee, The expressive power of low-rank adaptation, in: The Twelfth International Conference on Learning Representations, OpenReview.net, Vienna, Austria, 2024. doi:10.48550/arXiv.2310.17513.

[68] S. M. Abdullah, A. Paul, S. Rayana, A. Kabir, Z. Masud, State-of-the-art translation of text-to-gloss using mbart : A case study of bangla (2025). arXiv:2504.02293, doi:10.48550/arXiv.2504.02293.

[69] A. Duarte, S. Palaskar, L. Ventura, D. Ghadiyaram, K. DeHaan, F. Metze, J. Torres, X. Giro-i Nieto, How2sign: A large-scale multimodal dataset for continuous american sign language, in: 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), IEEE, Nashville, TN, USA, 2021, pp. 2735—2744. doi:10.1109/CVPR46437.2021.00276.

[70] A. Joshi, R. Mohanty, M. Kanakanti, A. Mangla, S. Choudhary, M. Barbate, A. Modi, isign: A benchmark for indian sign language processing, in: Findings of the Association for Computational Linguistics ACL 2024, Association for Computational Linguistics, Bangkok, Thailand and virtual meeting, 2024, pp. 10827—10844. doi:10.18653/v1/2024.findings-acl.643.

[71] H. A. Rubaiyeat, H. Mahmud, M. K. Hasan, Bangla sign language translation: Dataset creation challenges, benchmarking and prospects (2025). arXiv:2511.21533.

[72] N. Saha, R. Shahriyar, N. A. Roudra, S. Sakib, A. A. Rasel, Introducing a Bangla sentence – gloss pair dataset for Bangla sign language translation and research, in: Proceedings of the Fifteenth Language Resources and Evaluation Conference, ELRA Language Resource Association, Palma de Mallorca, Spain, 2026, pp. 10457–10466. doi:10.63317/38qenrwzegr9.

Appendix A. Training Configuration and Hyperparameters

We provide complete training configurations for IsharaBART pre-training and fine-tuning experiments to ensure reproducibility. All models were trained using the Hugging Face Transformers library with PyTorch backend.

Appendix A.1. IsharaBART Pre-training

IsharaBART was pre-trained on a monolingual Bangla corpus using a denoising autoencoder objective with custom corruption strategies. Table A.13 summarizes the key hyperparameters. We use polynomial learning rate decay as recommended in the original BART paper, with a relatively high initial learning rate (5×10^{-4}) to accommodate the smaller model size and focused pre-training corpus. Training was conducted in bfloat16 precision for numerical stability and computational efficiency. The effective batch size of 256 (64 per device $\times$ 4 gradient accumulation steps) balances training stability with memory constraints of encoder–decoder architectures.

Appendix A.2. Fine-tuning Configuration

All models (IsharaBART and baselines) were fine-tuned on the gloss translation tasks using identical or carefully matched configurations to ensure fair comparison. Table A.14 presents the hyperparameters used for each model and Tale A.15 present the LoRA configuration for parameter-efficient fine-tuning of IsharaBART. Configuration differences across models reflect preliminary validation experiments and model-specific characteristics: mT5-Small benefits from a higher learning rate (1×10^{-3}) due to its architecture and pre-training objective, while BanglaGPT requires a lower learning rate (1×10^{-5}) and larger gradient accumulation to stabilize training in the decoder-only setting.

For all encoder–decoder models (IsharaBART, mBART-50, mT5-Small, BanglaT5), we use beam search with 4 beams during generation for evaluation. Early stopping with patience of 3–5 epochs prevents overfitting on the low-resource fine-tuning datasets. Model selection is based on validation BLEU-4 for encoder–decoder models and validation loss for BanglaGPT, reflecting the different training objectives and generation strategies.

All experiments were conducted using Hugging Face Transformers (v4.57) and PyTorch 2.9 on NVIDIA A100 GPU for pre-training and Tesla T4 GPU for fine-tuning. For pre-training, we implemented a custom data collator that applies BART-style corruption including text infilling, sentence permutation, and token masking, with special handling for Bangla punctuation tokens (danda and question marks). Fine-tuning used the standard `DataCollatorForSeq2Seq` for dynamic padding.

Table A.13: IsharaBART pre-training configuration. All training used AdamW optimizer with $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$.

Hyperparameter	**Value**
Training Steps	30,000
Batch Size (per device)	64
Gradient Accumulation	4 steps
Effective Batch Size	256
Learning Rate	5×10^{-4}
LR Schedule	Polynomial decay
Warmup Steps	3,000
Weight Decay	0.1
Max Gradient Norm	1.0
Precision	bfloat16
Gradient Checkpointing	False
Random Seed	42
Evaluation Strategy	Every 1,000 steps
Save Strategy	Every 5,000 steps
Model Selection	Best validation loss

Table A.14: Fine-tuning configuration for all models. ES = Early Stopping patience (epochs). All models use AdamW optimizer except where noted.

Hyperparameter	**IsharaBART**	**mBART-50**	**mT5-Small**	**BanglaT5**	**BanglaGPT**
Epochs	20	20	20	20	20
Batch Size (per device)	16	16	16	16	16
Gradient Accumulation	2	2	2	2	4
Effective Batch Size	32	32	32	32	64
Learning Rate	1×10^{-4}	1×10^{-4}	1×10^{-3}	1×10^{-4}	1×10^{-5}
LR Schedule	Linear decay	Linear decay	Linear decay	Linear decay	Linear decay
Weight Decay	0.01	0.01	0.01	0.01	0.01
Max Gradient Norm	1.0	1.0	1.0	1.0	1.0
Precision	fp32	fp32	fp32	fp32	fp32
Beam Size	4	4	4	4	–
Generation Max Length	128	128	128	128	128
ES Patience	5	5	5	5	5
Model Selection Metric	BLEU-4	BLEU-4	BLEU-4	BLEU-4	Validation loss

Hyperparameter	Value
Task type	Seq2Seq language modeling
LoRA rank (r)	8
LoRA scaling factor (α)	16
LoRA dropout	0.1
Target modules	q_proj, k_proj, v_proj, out_proj

Table A.15: LoRA configuration used for parameter-efficient fine-tuning.