Low Language Transfer Speech-to-Text

Husein Zolkepli

husein.zolkepli@scicom.com.my

Scicom (MSC) Berhad, Malaysia

arXiv:2611.16260v1 [cs.CL] 12 Nov 2026

ABSTRACT

A multilingual speech recogniser has heard much speech in its own languages. We ask how much of that it brings to a language it was never trained on, when it is fine-tuned on the little that language has. We score 28 open multilingual models untouched on Low-Language-TTS, 50 low-resource languages in 11 writing systems, and fine-tune 24 of them on at most 1,000 clips a language, 82 hours in all, with a dev set choosing every epoch. Fine-tuning lowers every model's macro character error rate (CER), by 16.6 to 53.1 points, to 9.5 at best, and a word language model takes three CTC models 1.2 to 1.4 points lower; WER is still 34.4 at best. On the 12 languages that no model lists, every model starts at 46 or worse and the best ends at 11.7. The same architectures from random weights stay 60 points or more behind, so the gain is transfer, and an encoder left untouched already gives a Whisper or Qwen decoder most of what it needs. It is the language, not the corpus: on FLEURS recordings that ours do not share, the median model keeps 79% of its gain. It costs the languages the models knew, 5.7 points of CER on English, Spanish, French and German in the median model. Low-rank adapters on 4% of the weights come within 1.1 points of tuning everything, and on 2 of 3 models forget less. Seq2seq and language-model decoders are done within three to five epochs, while a new CTC head is still improving at twenty or thirty. Data beats epochs: 25 clips a language take Qwen3-ASR 0.6B from 60.9 to 39.3, and at equal clips presented, ten times the data is worth 9.7 points. A model that first learned forty of the languages learns the other ten no faster. On a public Malay benchmark, untouched Whisper large-v3 scores 8.9 when it detects the language and 14.9 when told the clip is Malay, because it then translates the English that speakers mix in; trained on the Scicom telephony dataset it scores 6.0.

1 Introduction

Speech recognisers are trained on the languages that have data, and most languages do not. A multilingual pretrained model has heard hundreds of thousands of hours of speech in its own languages. This paper asks how much of that it brings to a language it was never told about, when someone fine-tunes it on the little that language has.

We take 29 open multilingual models and score the 28 that transcribe as downloaded, untouched, on Low-Language-TTS [Malaysia-AI, 2026a], 50 of the lowest-resource languages of a large public corpus. We then fine-tune 24 of them on the same training set, at most 1,000 clips a language, and score them again. Four controls say what the gain is made of: the same architectures from random weights, each half of a model frozen in turn, each language trained alone, and the 12 languages that no model lists, where none starts ahead (Section 7.6). Three questions from practice get their own sections: whether more epochs

help (Section 8), how few clips are enough (Section 9), and whether Common Voice would have done (Section 10). Then we test the result on another corpus and on the languages the models already knew (Section 11), ask whether a model that has learned forty of the languages learns the next one faster (Section 7.8), and compare small adapters and a word language model with tuning everything (Section 12). A last section leaves the long tail for Malay, where the same models train on telephone calls alone and are scored on speech of every other kind (Section 13).

2 Background

Transfer in speech recognition. Adapting a network trained on one language to another is old and cheap: a small change to an English model's weights was enough for German [Kunze et al., 2017]. Training on many languages first helps more [Dalmia et al., 2018], and how similar the target is to the pretraining languages matters along several dimensions at once [Adams et al., 2019]. Meta-learning the starting point has also been tried [Hsu et al., 2019].

Multilingual models. XLS-R learned speech representations from 128 languages without labels [Babu et al., 2021]. MMS scaled this to 1,107 languages, on readings of religious texts [Pratap et al., 2023], and Omnilingual ASR to over 1,600 [Omnilingual ASR team et al., 2025]. Whisper [Radford et al., 2022] lists 99 languages. Recent models list fewer and transcribe them better: Parakeet TDT [Xu et al., 2023, Rekesh et al., 2023, NVIDIA, 2025b] and Canary [NVIDIA, 2025a] 25, Nemotron 3.5 40 [NVIDIA, 2026b,c], Qwen3-ASR 30 [Shi et al., 2026], Voxtral 8 [Liu et al., 2025], Granite Speech 6 [Saon et al., 2025, IBM, 2026], Cohere Transcribe 14 [Mack et al., 2026]. Several put a language model behind the audio encoder [Chen et al., 2023]. General assistants now transcribe when asked: Gemma 4 takes audio in three of its sizes [Gemma Team, 2026], and LFM3-Audio lists 18 languages for speech [Liquid AI, 2026].

Fine-tuning on one language at a time. Most low-resource work tunes one model on one language or a few: Whisper on Amharic [Gete et al., 2025], Northern Kurdish [Abdullah et al., 2024] and Malasar [Pillai et al., 2024], with data augmentation for Gronings and West Frisian [Bartelds et al., 2023]. Studies of data size find most of the gain early, on African languages [Akera et al., 2025, Nahabwe et al., 2025], and pretrained encoders differ in how well they hold up with little data [Yang et al., 2023]. The nearest comparison is the benchmark of Whisper, XLS-R, MMS and w2v-BERT on 13 African languages [Nahabwe et al., 2025].

Adapters, forgetting and decoding. A network trained on something new loses some of what it knew [Kirkpatrick et al., 2016]. Low-rank adapters train a few added weights and leave the rest as it was [Hu et al., 2021], MMS adds each language through a small adapter of its own [Pratap et al., 2023], and adapters have been used to keep a speech recogniser from forgetting [Vander Eeckt and Van hamme, 2022]. A CTC model's output can also be searched with a word language model [Heafield, 2011, Kensho Technologies, 2021].

What this paper adds. One held-out test set of 50 languages in 11 writing systems. Every family of current decoders on the same training data, including models whose vocabulary cannot spell the language at all. Random-weight controls that measure the transfer itself rather than the final error. A check of the gain on another corpus, of what it costs the languages the models knew, and of what adapters and a word language model change.

3 Data

3.1 THE TEST SET

We score on Low-Language-TTS [Malaysia-AI, 2026a]: 1,250 clips and 2.56 hours from 514 speakers, 25 clips for each of 50 of the lowest-resource languages of a 121.8M-row corpus [Malaysia-AI, 2026b]. A language counts only if some subset of the corpus was collected for it, because the rare end of a language detector's histogram is noise [Kargaran et al., 2023]. Figures 1 and 2 show what that leaves: 13 language families, and 14 languages outside the Latin script. The study never chooses anything on this set. Some of the corpus's subsets repackage FLEURS [Conneau et al., 2022]: 151 test clips in 7 languages are FLEURS recordings, among them every test clip of Estonian, Latvian, Occitan and Slovenian, and 18 of them come from FLEURS's training split, so a model trained on it has heard them (Section 11).

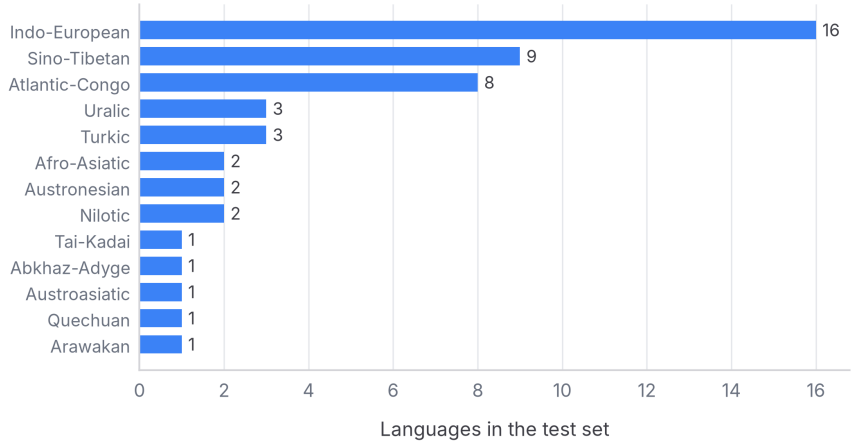

FIGURE 1 The test set's languages by language family.

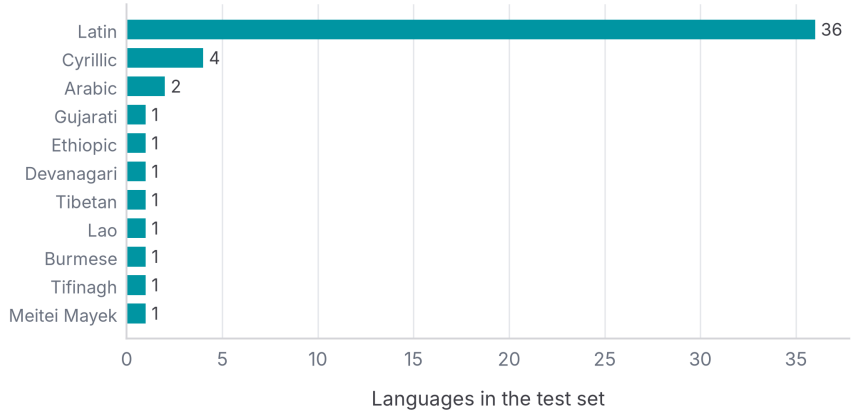

FIGURE 2 The test set's languages by writing system.

3.2 THE TRAINING SET

Common Voice 22 [Ardila et al., 2019, Mozilla, 2025] was the first candidate, and it does not cover the study: of its 137 locales only 13 are among our 50 languages, so 37 have nothing there. Training data therefore comes from the corpus the test set was cut from, under the rules of Figure 3; Section 10 adds Common Voice as a comparison.

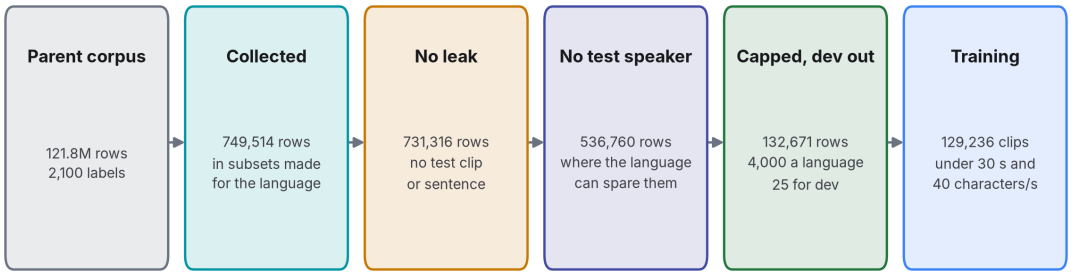

FIGURE 3 From the parent corpus to the training set. A test or dev clip's text leaves training in every subset, since Bible corpora and augmented copies repeat their sentences. No test clip, test sentence or identical decoded audio remains in training or dev; a script checks all three.

That leaves 129,236 clips and 279 hours from 27,206 speakers (Figure 4). In 45 languages no test speaker is in training. The other 5 (Crimean Tatar, Goan Konkani, Gurani, Meitei and Wayuu) have no other speakers to train on. Dev has 1,250 clips, 25 a language, gated on audio length as the test set was: some corpora pair a long sentence with a second of audio. It takes each language's smallest speakers first, because they cost training least, which makes it harder than test (Section 14). Within a language, speakers take turns in the order of the clips, so a smaller training set is always a prefix of a larger one. The main runs use at most 1,000 clips a language: 42,816 clips and 82 hours in all. Fetching the test clips again through the code that fetched the training audio returns the same samples for all 1,250, and a training clip's audio and its speech tokens, which the parent corpus keeps in separate archives, agree on length to within 20% on all but 1 of the 125,984 clips that have tokens.

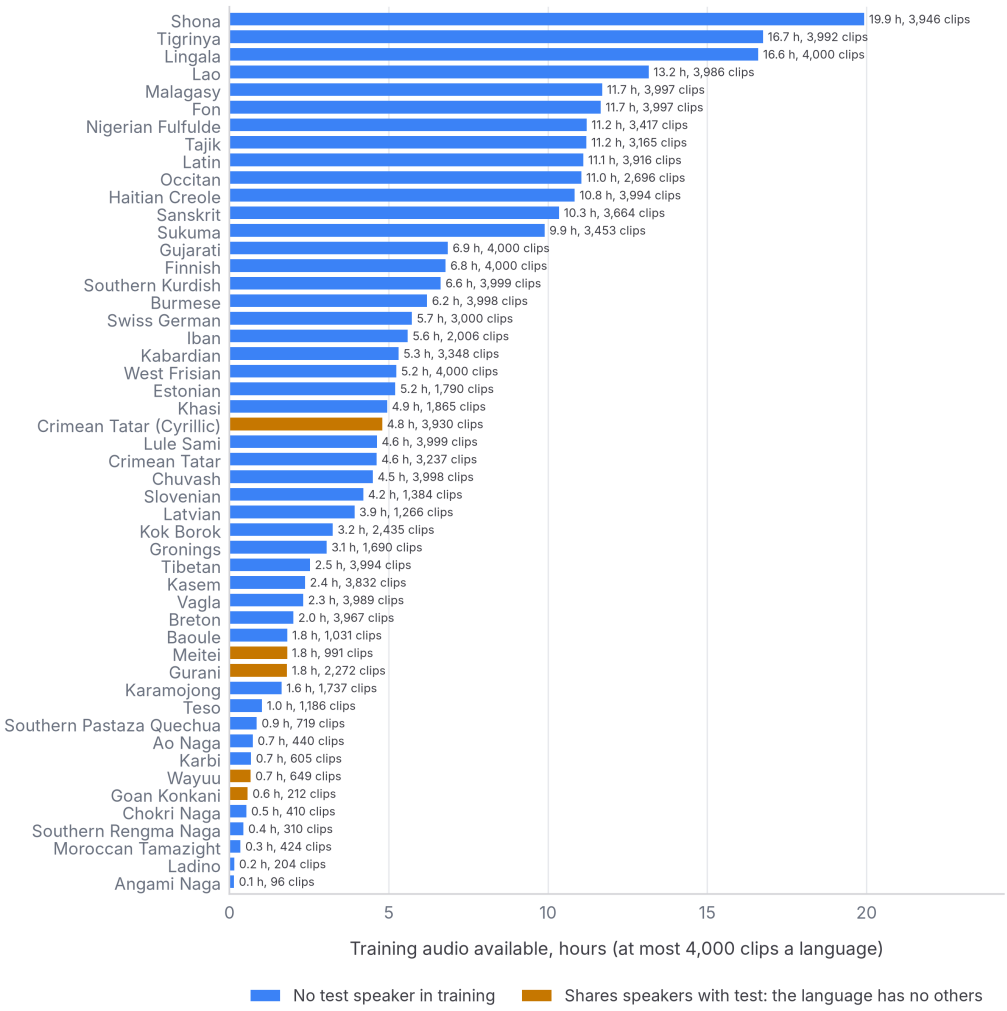

FIGURE 4 Training audio per language, after every rule of Figure 3.

3.3 THE METRIC

The headline is the macro character error rate: the CER of each language over its 25 clips, then the mean over languages, so that each language counts once. Text is normalised the same way for every model (NFKC, lower case, punctuation and symbols removed, combining marks kept). Characters are the fair unit because 11 writing systems do not agree on what a word is; WER is reported beside it.

We cap each clip's errors at its own length. A decoder stuck in a loop then counts as a complete failure of that clip, not as twenty failures. Without the cap, a few loops moved a language by tens of points from one epoch to the next and hid everything else (Appendix A). Intervals resample speakers within each language, because one speaker's clips are not independent; they take the 50 languages as given, and Section 6.4 draws the languages again.

4 Models

The study's one criterion for a model is that it is multilingual by its own card. We read the 400 most-downloaded speech recognition repositories on Hugging Face on the day of the study, removed conversions, quantisations, third-party fine-tunes and single-language models, and

kept what runs in current Transformers or its own open package (Table 1). Between them they cover the families of decoder in use today. Two general assistants are not on that list, because the Hub files them under other tasks, and were added for comparison: Gemma 4, in the three sizes that take audio, and LFM3-Audio.

TABLE 1 The models. Languages listed is what each card says it transcribes; w2v-BERT 2.0 is a self-supervised encoder pretrained on that many languages, without a transcription head. MMS as downloaded swaps in an adapter and a head for the language it is told; fine-tuned here, it has one head for all 50. Gemma 4's parameters are the stored ones; its card counts E2B and E4B as 2.3B and 4.5B effective, without the token tables each layer looks up.

MODEL	DECODER	PARAMETERS	LANGUAGES LISTED	FINE-TUNED HERE	LICENCE
Whisper large-v3	seq2seq	1.54B	99	yes	Apache-2.0
Whisper large-v3-turbo	seq2seq	0.81B	99	yes	MIT
Whisper large-v2	seq2seq	1.54B	99	yes	Apache-2.0
Whisper medium	seq2seq	0.76B	99	yes	Apache-2.0
Whisper small	seq2seq	0.24B	99	yes	Apache-2.0
Whisper base	seq2seq	0.07B	99	yes	Apache-2.0
Whisper tiny	seq2seq	0.04B	99	yes	Apache-2.0
Qwen3-ASR 1.7B	language model	2.04B	30	yes	Apache-2.0
Qwen3-ASR 0.6B	language model	0.78B	30	yes	Apache-2.0
Nemotron 3.5 Transcribe 30B	language model	32.36B	40	yes	OpenMDW-1.1
Parakeet TDT 0.6B v3	transducer	0.63B	25	yes	CC-BY-4.0
Nemotron 3.5 streaming 0.6B	transducer	0.64B	40	yes	OpenMDW-1.1
Canary 1B v2	seq2seq	0.96B	25	no	CC-BY-4.0
w2v-BERT 2.0	CTC	0.58B	143	yes	MIT
MMS 1B, a head per language	CTC	0.96B	1,198	yes	CC-BY-NC-4.0
omniASR CTC 300M	CTC	0.33B	1,600+	yes	Apache-2.0
omniASR CTC 1B	CTC	0.98B	1,600+	yes	Apache-2.0
omniASR CTC 7B	CTC	6.50B	1,600+	yes	Apache-2.0
omniASR LLM 1B	language model	2.28B	1,600+	no	Apache-2.0
Voxtral Mini 3B	language model	4.68B	8	yes	Apache-2.0
Granite Speech 5.0 2B	language model	2.31B	6	yes	Apache-2.0
Granite Speech 5.0 470M	CTC	0.47B	7	no	Apache-2.0
Fun-ASR Nano	language model	0.83B	4	no	Apache-2.0
VibeVoice-ASR	language model	8.33B	50+	no	MIT
Cohere Transcribe 2B	seq2seq	2.07B	14	yes	Apache-2.0
LFM3-Audio 2B	language model	1.98B	18	yes	LFM Open v1.0
Gemma 4 E2B	language model	5.10B	35+	yes	Apache-2.0
Gemma 4 E4B	language model	7.94B	35+	yes	Apache-2.0
Gemma 4 12B	language model	11.96B	35+	yes	Apache-2.0

A model cannot fine-tune on a language until it can accept the language and spell it, and each family needs something different (Figure 5). Models with a language model or a byte-level vocabulary can write any script already. The transducers cannot: 10 of our languages are mostly outside Parakeet's vocabulary and 7 outside Nemotron's. We add a row per missing character and keep every pretrained weight [NVIDIA, 2026a]. MMS as downloaded is different again: a small adapter in every layer and a head for each of its languages, swapped in for the language the user names. We load its English adapter once, replace the heads with one head for all 50 languages and train every weight, so the fine-tuned MMS, like the other CTC models, has one head and is never told the language.

Family	Models	A new language enters as	A new script enters as
Seq2seq	Whisper, Cohere Transcribe	a language token (Cohere: an unused one), from the mean of the others	nothing: a byte-level vocabulary spells any script
Language model	Qwen3-ASR, Nemotron 30B, Voxtral, Granite, Gemma 4, LFM3-Audio	its name in the prompt (Voxtral: a short code)	nothing: the tokenizer spells any script
Transducer	Parakeet TDT, Nemotron streaming	Parakeet: nothing. Nemotron: a prompt slot, a copy of auto	new rows in the decoder and the joint head, each a copy of the blank's, below it
CTC, fresh head	w2v-BERT, MMS	nothing: no language input	a new character head for all 50 languages
CTC, own head	omniASR	nothing: no language input	9,812 characters already; new rows for the few missing

FIGURE 5 What each family needs before it can fine-tune on a language it does not know. Every pretrained weight is kept in every family, except MMS's CTC heads, which are replaced; w2v-BERT comes without one.

A new row has to start silent, so that the untrained extension is still the pretrained model, and that took two attempts. Zero weights and a bias below every pretrained row is the obvious start, and it is not silent: this head's logits sink far below its biases, so the constant wins most frames. Untrained, Parakeet repeated its first new character on 221 of the 1,250 test clips and Nemotron on 1,149. A row that starts as a copy of the blank's, ten below it, cannot win a frame, because it is below the blank whatever the frame holds. With that start the untrained model's output is the pretrained model's on 1,208 of those clips for Parakeet and 1,250 for Nemotron. Section 7.4 follows both starts through training.

5 Untouched

As downloaded, a language on a model's own list is forced; any other is left to the model, as a user would have to, since it cannot be told about a language it was never trained on (Canary and Cohere Transcribe have no detection, so they get English). Gemma 4 is the exception: its card lists no languages for speech and the prompt it gives for transcription takes any language's name, so it is told every language.

Models that list up to a hundred languages sit at a macro CER of 50.3 to 84.9 (Figure 6): they write something, often in another script (Section 7). Gemma 4's two sizes with an audio encoder are lower, at 39.8 to 43.5, and the name is part of that: with no language in the prompt they score 47.7 to 52.2, 7.9 to 8.7 points worse. Its 12B, which reads the waveform with no encoder, is at 72.9, and at 75.8 with no language in the prompt: on English, Spanish, French and German it scores 2.2, but a language it has not heard it writes as one it has, Finnish as Arabic or Spanish, Gujarati in Hindi's script. The two systems built for the long tail are far ahead: omniASR at 27.8 to 31.5 on all 50, and MMS at 14.7 on the 26 languages it has a head for. On the rest, MMS cannot detect a language any more than Canary can, so it gets English's adapter and head, and over all 50 it scores 30.6. MMS was trained on readings of religious texts [Pratap et al., 2023], and some of our languages' clips are the same kind of reading. Appendix A shows why the cap is there and has every score with and without it.

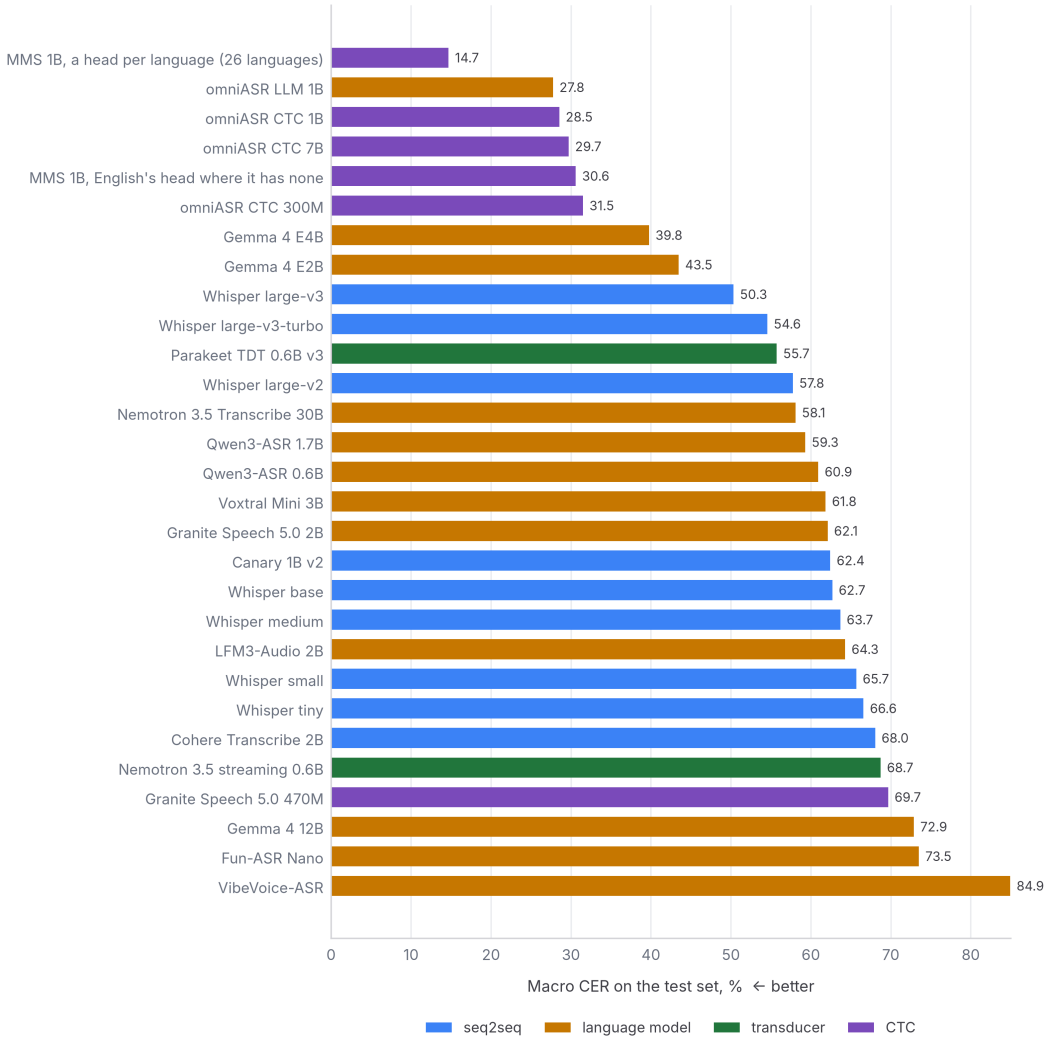

FIGURE 6 Every model as it was downloaded, before any fine-tuning, on the test set. Mean over 50 languages unless stated; the fine-tuned scores are in Section 6.2. Gemma 4 is told every language's name, the others only the languages they list.

6 Fine-tuning

6.1 THE RECIPE

Every model trains the same way: AdamW with weight decay 0.01, its pretrained weights in 32-bit floats under bf16 autocast, a learning rate chosen for it on dev (Appendix C), a warm-up of 200 steps and then a constant rate, gradient clipping at 1, and greedy decoding. All 50 languages train together on at most 1,000 clips each, 82 hours in all, for ten epochs, except Whisper large-v2 (5), Voxtral (4), Nemotron 30B (4), Gemma 4 E2B (5), Gemma 4 E4B (5) and Gemma 4 12B (4). After every epoch we score dev and test; dev alone chooses the epoch. The targets are the corpus's transcripts with their case and punctuation, except for the CTC models below. Every pretrained weight a transcript runs is kept and trains, with three exceptions: Nemotron 30B's first run freezes most of the model, Gemma 4's token tables for each layer stay frozen, and MMS's heads are replaced. Figure 7 shows what happens to each part of each family, and the paragraphs below go through them.

Whisper. Encoder and decoder train together, cross-attention included, and nothing is frozen.

Its byte-level vocabulary already spells every script, so the one addition is a language token for each of the 33 languages Whisper has none for, started at the mean of its language tokens; the other 17 keep their own. The token goes where Whisper expects one, after the start-of-transcript token, and the loss counts only the transcript and the end token. Whisper's output layer is its token table, so a new token is also a new output.

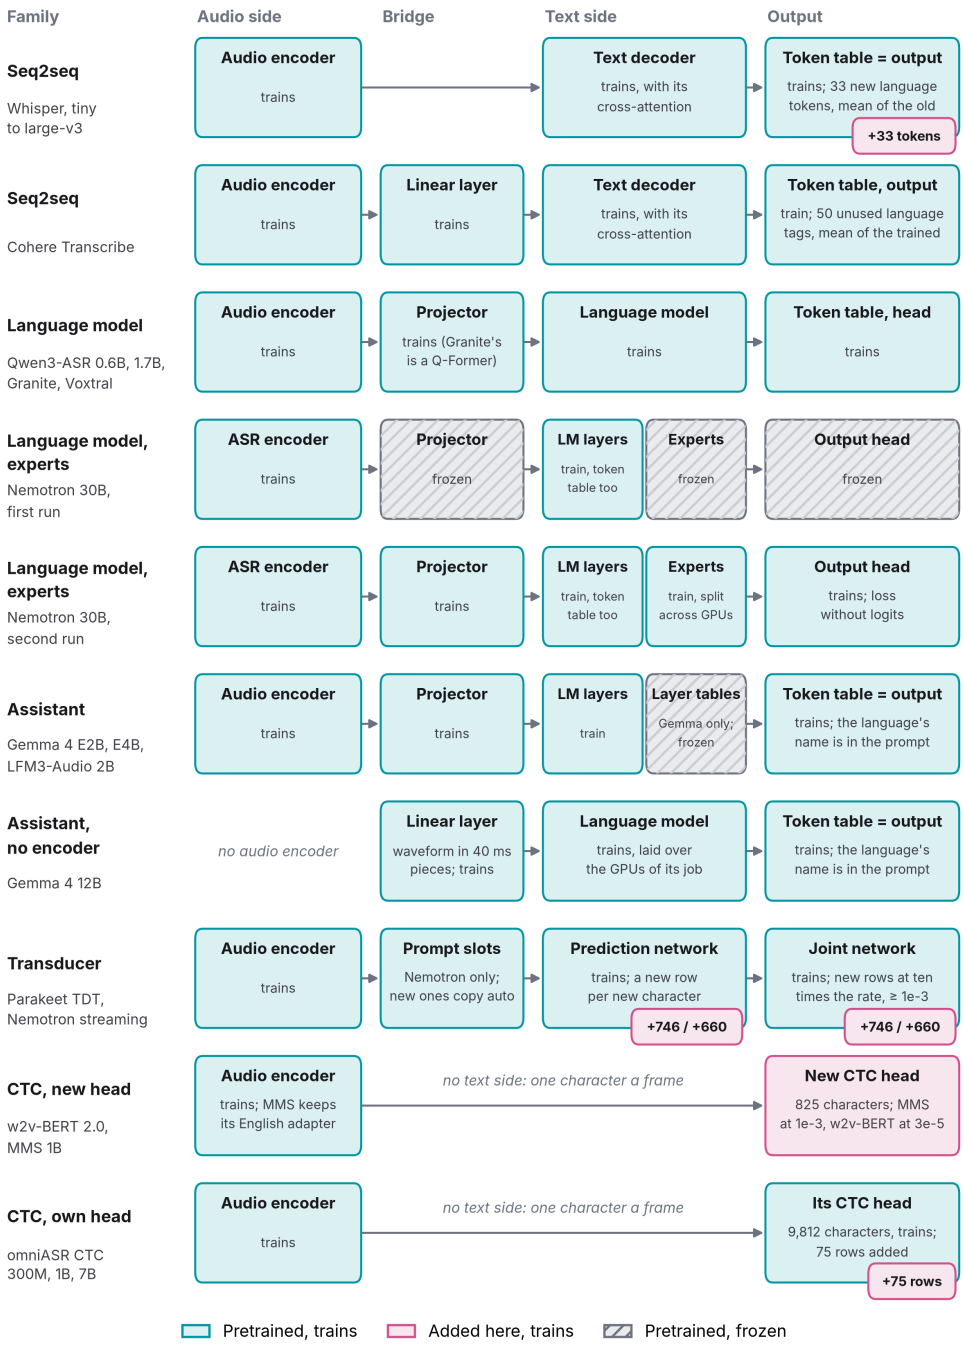

FIGURE 7 What the recipe does to each part of each family. Nemotron 30B ran twice: first with its experts, projector and output head frozen, then with every part training. A pink badge is what fine-tuning adds to a pretrained part; on the transducers it is Parakeet's count of new rows, then Nemotron streaming's.

Cohere Transcribe. The other seq2seq model is a Conformer encoder of 1.90B parameters

under a small Transformer decoder, trained from scratch on 14 languages, none of them ours [Mack et al., 2026]. It trains whole, and its SentencePiece vocabulary falls back to bytes, so it spells every script and nothing is added. Its prompt names the language with a tag, and its table holds tags that were never trained: 21 of our languages take the tag of their two-letter code and the other 29 a spare token, each started at the mean of the 14 trained tags, as a new Whisper token is. Its output layer is not tied to the token table, so the two train side by side.

Language-model decoders. Qwen3-ASR, Granite and Voxtral train whole: the audio encoder, the projector that carries its output into the language model (a Q-Former in Granite), the language model and its output head. Their tokenizers spell every script, so nothing is added. The language enters as text in the prompt, its name for Qwen3-ASR and Granite and a short code for Voxtral, and the loss counts only the transcript. Nemotron 30B is a mixture of experts, and its first run trains only part of it: its 29.8B expert weights stay frozen in bf16, with the output head, the parts that transcription never runs (speaker diarisation, word timings) and the 3.4M-parameter projector. Its ASR encoder and every language-model weight outside the experts train in fp32: 2.02B parameters, 6% of the model (Figure 8). Our freezing rule matched names, and it was meant to train the projector and freeze the token table; it did the reverse for both. A second run trains every weight that transcription uses, 32.2B: the experts split across GPUs by expert parallelism, everything else sharded with FSDP2, and the loss computed without the logits by a fused cross-entropy in fp32 [Scicom AI Enterprise Organization, 2026] after Cut Cross-Entropy [Wijmans et al., 2024] (Figure 9). Dev prefers it, so it is Nemotron 30B's main run from here on: at epoch 2 of four it scores 15.8 on the test set, 0.2 points below the first run at its epoch 3 (-0.2 to 0.6). Section 6.8 trains its parts one at a time. Voxtral and Nemotron 30B's first run train with gradient checkpointing.

General assistants. Gemma 4 and LFM3-Audio are chat models that transcribe when asked [Gemma Team, 2026, Liquid AI, 2026]. Each gets the prompt its card gives for speech recognition, with the language's name in it, and nothing is added: their tokenizers spell every script. The loss counts the transcript and the token that ends the turn. The parts a transcript never runs stay frozen, Gemma's vision parts and LFM3-Audio's speech output. Gemma E2B and E4B also keep a token table for every layer, 2.35B and 2.82B parameters, 46% and 35% of what is stored. They are lookups, Google counts the models without them, and we leave them frozen, so 2.59B and 4.95B train (Figure 8). Gemma 12B has no audio encoder: a linear layer hands the waveform to the language model in 40 ms pieces, and all 11.91B of it train, laid layer by layer over the GPUs of its job. It and E4B train with gradient checkpointing.

Transducers. Parakeet and Nemotron streaming train whole. Their SentencePiece vocabularies come from their own languages, so every character the training transcripts need and the vocabulary lacks gets a new row in the prediction network's embedding and in the joint network's output: 746 for Parakeet, 660 for Nemotron streaming. The rows start silent (Section 4) and train without weight decay at ten times the model's rate and at least 1e-3, 3e-3 for Parakeet and 1e-3 for Nemotron streaming. Parakeet has no language input. Nemotron streaming takes the language as a one-hot prompt slot, and a language it does not list gets an unused slot that starts as a copy of its "auto" slot. They train on the TDT and the RNN-T loss.

CTC models. A CTC model has no decoder: its head writes a character or a blank for every frame of the encoder's output. w2v-BERT 2.0 comes without a head and MMS's heads belong to its own languages (Section 4), so each gets a new linear head over the 825 characters

of the training transcripts. For these two the transcripts are lower-cased and stripped of punctuation, as is usual for a CTC head with no language model to put them back. The encoder trains with the head, every weight of it, MMS's English adapter included. MMS's head trains at 1e-3. w2v-BERT's sits behind a LayerNorm, because the encoder's output is so small that a plain new head stayed on the blank, and trains at the encoder's rate. omniASR keeps its own head of 9,812 characters and adds a row for each of the 75 it lacks, all at the model's rate. Every model is scored on the same normalised text (Section 3.3), so case and punctuation never count.

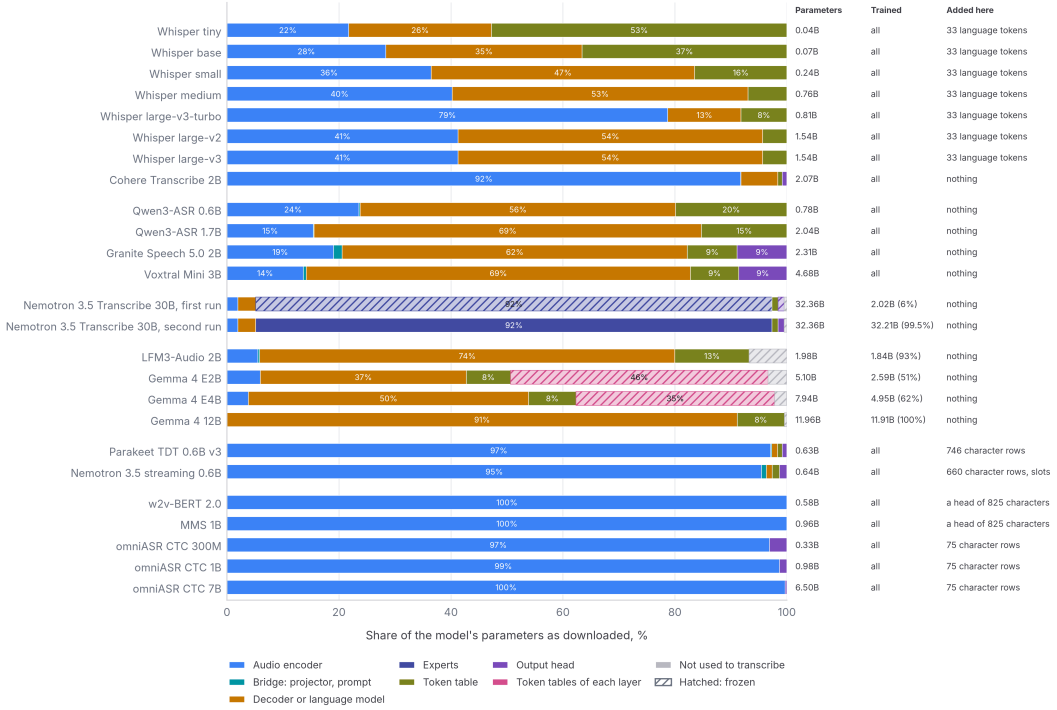

FIGURE 8 Every fine-tuned model's parameters by part, as a share of the model as downloaded; hatched parts stay frozen. Nemotron 30B's second run trains every part but those transcription never runs. Trained is what fine-tuning updates; added here is what it adds.

We add no augmentation of our own. Each model keeps the dropout and layer drop its own configuration switches on in training, and MMS's configuration also masks 5% of the encoder's frames in spans of 10, a form of SpecAugment (Table 2). A batch holds a fixed amount of padded audio, or a fixed number of clips for Whisper and Voxtral, whose encoders take 30-second windows.

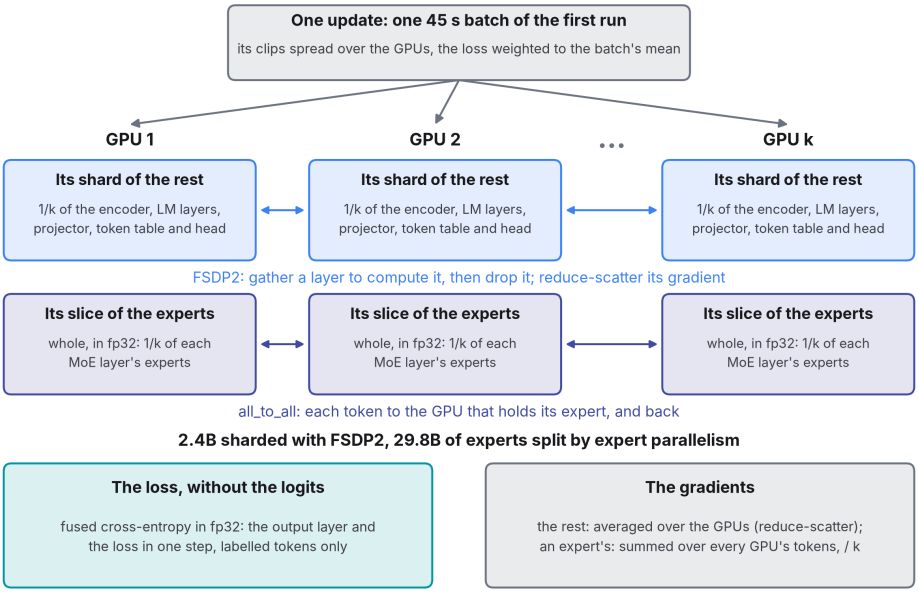

FIGURE 9 Nemotron 30B's second run on the GPUs of one job: the experts split by expert parallelism, everything else sharded by FSDP2, and the loss computed without the logits [Scicom AI Enterprise Organization, 2026]. Each update is the first run's: the same batch, its clips spread over the GPUs.

TABLE 2 What each fine-tuned model trains with. Rate is the model's own, chosen on dev (Appendix C). Added is what the training set needed and the model lacked; it trains at the model's rate unless another is given. Batch is clips, or seconds of padded audio. The last column is the dropout, layer drop and masking the model's own configuration switches on in training; we add nothing.

MODEL	RATE	ADDED, ITS RATE	LANGUAGE INPUT	BATCH	ITS OWN CONFIGURATION
Whisper tiny	1e-4	33 language tokens	language token	64 clips	none
Whisper base	1e-4	33 language tokens	language token	64 clips	none
Whisper small	3e-5	33 language tokens	language token	48 clips	none
Whisper medium	1e-5	33 language tokens	language token	32 clips	none
Whisper large-v3-turbo	1e-5	33 language tokens	language token	24 clips	none
Whisper large-v2	1e-5	33 language tokens	language token	24 clips	none
Whisper large-v3	1e-5	33 language tokens	language token	24 clips	none
Cohere Transcribe 2B	3e-5	none (50 unused tags)	language token	180 s	none
Qwen3-ASR 0.6B	1e-5	–	name in prompt	240 s	none
Qwen3-ASR 1.7B	1e-5	–	name in prompt	180 s	none
Granite Speech 5.0 2B	1e-5	–	name in prompt	90 s	dropout 0.1
Voxtral Mini 3B	3e-6	–	code in prompt	8 clips	none
Nemotron 3.5 Transcribe 30B	1e-5	–	name in instruction	45 s	dropout 0.1
LFM3-Audio 2B	1e-5	–	name in prompt	180 s	none
Gemma 4 E2B	1e-5	–	name in prompt	90 s	none
Gemma 4 E4B	3e-6	–	name in prompt	60 s	none
Gemma 4 12B	3e-6	–	name in prompt	60 s	none
Parakeet TDT 0.6B v3	3e-4	746 character rows, 3e-3	–	240 s	dropout 0.1, layer drop 0.1
Nemotron 3.5 streaming 0.6B	1e-4	660 character rows, 1e-3; slots	prompt slot	240 s	dropout 0.1, layer drop 0.1
w2v-BERT 2.0	3e-5	head of 825 characters + LayerNorm, 3e-5	–	240 s	dropout 0.1, layer drop 0.1
MMS 1B	1e-5	head of 825 characters, 1e-3	–	240 s	dropout 0.05, layer drop 0.05, masking 5% in spans of 10
omniASR CTC 300M	3e-5	75 character rows	–	240 s	dropout 0.1, layer drop 0.1
omniASR CTC 1B	3e-5	75 character rows	–	240 s	dropout 0.1, layer drop 0.1
omniASR CTC 7B	1e-5	75 character rows	–	20 s × 6	dropout 0.1, layer drop 0.1

6.2 BEFORE AND AFTER

Table 3 and Figure 10 are the paper's main result, the 21 models fine-tuned with the main recipe.

TABLE 3 The main runs, best first. CER and WER on the test set at the epoch dev chose, capped per clip, macro over 50 languages. Hours are for one training epoch. Dev also chose how the transducers' new rows start (Section 7.4): for Parakeet, a copy of the blank's row, ten below it, at 3e-4 rather than 1e-4; for Nemotron streaming, the first start.

MODEL	PARAMETERS	UNTOUCHED	FINE-TUNED	95% CI	WER	BEST EPOCH	RATE	HOURS / EPOCH
MMS 1B	0.97B	30.6	**11.0**	10.6–11.6	39.3	10 of 10	1e-5	0.52
omniASR CTC 1B	0.98B	28.5	**12.0**	11.5–12.6	43.1	7 of 10	3e-5	0.51
w2v-BERT 2.0	0.58B	–	**12.7**	12.3–13.3	44.9	9 of 10	3e-5	0.30
Whisper large-v3	1.54B	50.3	**13.5**	13.0–14.2	41.5	8 of 10	1e-5	1.11
Whisper large-v2	1.54B	57.8	**14.2**	13.7–14.9	43.2	4 of 5	1e-5	1.09
Whisper large-v3-turbo	0.81B	54.6	**14.3**	13.6–15.1	43.0	9 of 10	1e-5	0.82
Whisper medium	0.76B	63.7	**14.8**	14.2–15.5	45.0	10 of 10	1e-5	0.59
Cohere Transcribe 2B	2.07B	68.0	**15.6**	15.0–16.3	45.8	10 of 10	3e-5	0.34
Nemotron 3.5 Transcribe 30B	32.36B	58.1	**15.8**	15.2–16.5	45.3	2 of 4	1e-5	–
Voxtral Mini 3B	4.68B	61.8	**16.3**	15.7–17.1	45.7	3 of 4	3e-6	2.93
Gemma 4 E4B	7.94B	39.8	**16.6**	16.0–17.3	47.1	2 of 5	3e-6	1.90
Qwen3-ASR 1.7B	2.04B	59.3	**16.6**	16.0–17.3	47.9	4 of 10	1e-5	0.47
Whisper small	0.24B	65.7	**16.9**	16.4–17.6	49.2	9 of 10	3e-5	0.20
LFM3-Audio 2B	1.98B	64.3	**17.8**	17.2–18.6	48.7	3 of 10	1e-5	0.44
Qwen3-ASR 0.6B	0.78B	60.9	**18.2**	17.5–18.9	51.9	4 of 10	1e-5	0.26
Gemma 4 E2B	5.10B	43.5	**18.5**	17.8–19.4	49.2	5 of 5	1e-5	0.81
Gemma 4 12B	11.96B	72.9	**19.8**	19.1–20.6	51.8	2 of 4	3e-6	4.03
Parakeet TDT 0.6B v3	0.63B	55.7	**21.0**	20.2–21.9	52.6	9 of 10	3e-4	0.80
Granite Speech 5.0 2B	2.31B	62.1	**22.5**	21.8–23.4	58.1	7 of 10	1e-5	0.68
Whisper tiny	0.04B	66.6	**25.7**	24.9–26.5	63.9	5 of 10	1e-4	0.04
Nemotron 3.5 streaming 0.6B	0.64B	68.7	**35.9**	35.3–36.5	78.2	10 of 10	1e-4	0.35

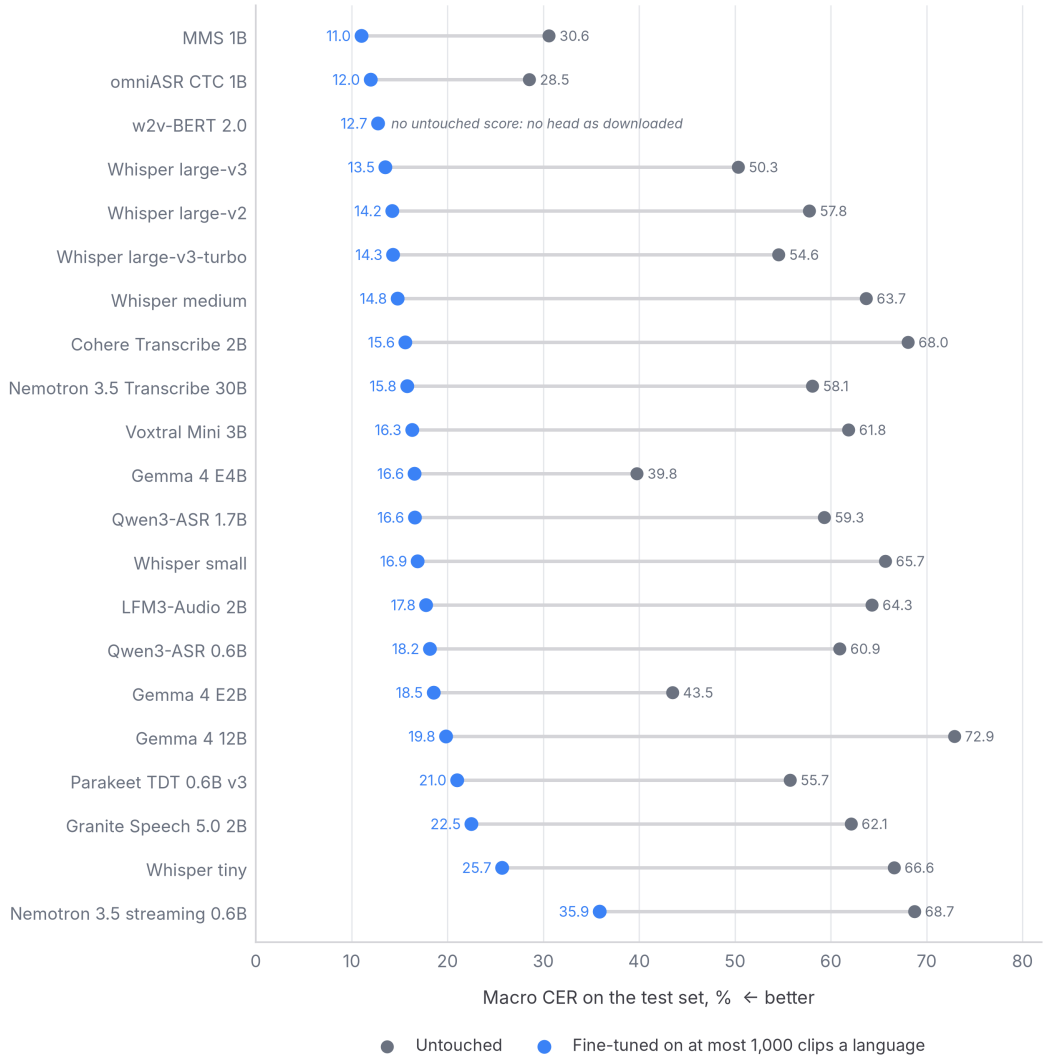

FIGURE 10 Every model before and after fine-tuning on at most 1,000 clips a language, 50 languages together. The epoch is the one dev chose. MMS untouched uses its own adapter and head for each language that has one and English's for the rest (Section 5); w2v-BERT has no head as downloaded, so no untouched score.

The rest of Figure 6 is 5 models we only scored as downloaded (Canary 1B v2, omniASR LLM 1B, Granite Speech 5.0 470M, Fun-ASR Nano and VibeVoice-ASR) and the three extra sizes trained for Section 6.7, Whisper base and omniASR CTC 300M and 7B. Every model gains, by 16.6 to 53.1 points of macro CER, and the best reaches 11.0 (MMS 1B); omniASR CTC 7B, trained with the same recipe for Section 6.7, goes lower still, to 9.5. Another seed moves little: the 5 models we trained again, 10 times in all, land within 0.4 points of their scores here (some draws are the first ten epochs of the long runs in Section 8). The first three keep their order on every seed: over 3 seeds each, MMS scores 11.0 to 11.3, omniASR CTC 1B 11.6 to 12.0 and w2v-BERT 12.4 to 12.8. A language has about 98 minutes of speech on average, trained alongside 49 others, and that takes a recogniser that could not write it to one that gets 89% of its characters right. Words are harder: the lowest WER among them is 39.3 (MMS 1B) and omniASR CTC 7B's 35.2, so this is a draft to correct, not a transcript. Figure 11 shows what such numbers look like on the page: for six languages, from easy to hard, the clip at the median of the language's errors.

The yardstick is MMS as downloaded, with an adapter and a head for each language it knows, built for the long tail and trained on far more of it. On the 26 of our languages it has a head for it scores 14.7 as downloaded. On those same languages, fine-tuned on this small set, 5 models do better: MMS 1B at 10.0, omniASR CTC 1B at 11.1, w2v-BERT 2.0 at 12.5, Whisper large-v3 at 13.8 and Whisper large-v2 at 14.5. MMS 1B in that list is the same model with its heads replaced by one trained on this set and never told the language, so on this test set a little speech from the test's own corpora does more for MMS than its own training data, mostly readings of religious texts. Fine-tuning a general model on a little of a language does as well as a model built for that language. Section 7.6 takes the other side of the same question, the 12 languages that neither MMS nor any other model here lists.

Haitian Creole — 1.2% of this clip's characters wrong after fine-tuning

Reference	Si yon timoun gen pyas, li pa dwe jwe ak lòt timoun pou evite li pwopaje enfeksyon an.
Untouched	Si yon timoun gen pias, li pa dwe jwe ak lot timoun pou evite li propaje infeksyon.
Fine-tuned	Si yon timoun gen pyas, li pa dwe jwè ak lòt timoun pou evite li pwopaje enfeksyon an.

Khasi — 1.7% of this clip's characters wrong after fine-tuning

Reference	Haba u Jakob u la wan na lyngkha ha ka por janmiet, ka Leah ka leit shaw ïa u bad ka ong, Phi dei ban thiah bad nga mynta ka miet, namar nga la siewdor ïa phi da ki mandriek jong u khun jong nga. Kumta u ïathiah shongkha bad ka ha kata ka miet.
Untouched	Kami berterima kasih kerana menonton.
Fine-tuned	Haba u Jakob u la wan na lyngkha ha ka por janmiet, ka Leia ka leit shaw ïa u bad ka ong, 'Phi dei ban thiah bad nga mynta ka miet, namar nga la siewdor ïa phi da ki mandrik jong u khun jong nga. Kumta u ïa thiah shongkha bad ka ha kata ka miet.'

Latvian — 3.6% of this clip's characters wrong after fine-tuning

Reference	Lielāko daļu rajonu apkalpo mazi japāņu piekrastes autobusi kas ir ērti un izturīgi.
Untouched	Lielāko daļu rajonu apkalpo mazi Japāņu piekrests autobus, kas ir ērti un isturīgi.
Fine-tuned	Lielāko daļu rajonu apkalpo maza japāņu piekrasts autobuss kas ir ērti un izturīgi.

Fon — 17.8% of this clip's characters wrong after fine-tuning

Reference	E nɔ gu bo tɛnmɛ tɛnmɛ lɛ bo nɔ sɔ do akpo mɛ.
Untouched	E no gubu teme temele, bono sodo ak bume.
Fine-tuned	E nɔ go bo temɛn temɛn lɛ bo nɔ sɔ ɖo akpo mɛ.

Breton — 28.2% of this clip's characters wrong after fine-tuning

Reference	Plijout a ra din sellet ouzh ar skinwel.
Untouched	Plijout a radin sellet ou zerskrin wel.
Fine-tuned	Pleijout a xradin selet ou zarskrin well.

Nigerian Fulfulde — 32.2% of this clip's characters wrong after fine-tuning

Reference	A wawa wiide o ɗumajum ɗum ɗo henna 120 jours lebbi nayyi goppa ha fin juillet walla ha nder hakkude juillet jaaha gawoya.
Untouched	Awa video mnayaj mnayaj mnayaj

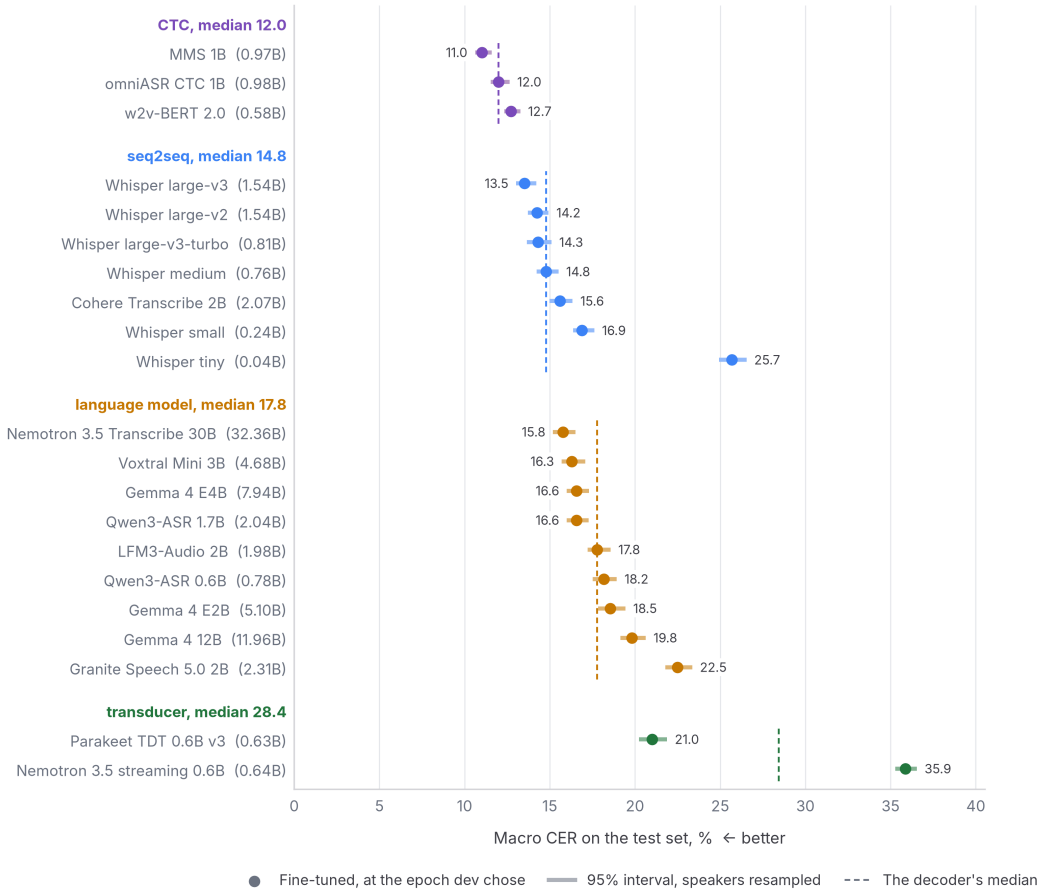

FIGURE 12 The main runs of Table 3 grouped by decoder, best first within each: macro CER on the test set at the epoch dev chose, with its 95% interval. Parameters are beside each name.

The decoder is not all that differs, so this ranks the models as they come rather than testing the decoders. The CTC models are also the three whose cards list the most languages, 143 to 1,600+ against at most 99 for the rest, though they still lead, by less, on the 12 languages that none of them lists (Section 7.6). The transducers trail on the Latin script too and fall furthest behind outside it (Section 7.3), where the rows we add have no pretrained letter behind them. The one pair built on the same pretrained encoder with different decoders, omniASR's CTC 1B and LLM 1B, scores 28.5 and 27.8 untouched, which suggests the encoder sets the order more than the decoder does; only the CTC variant was fine-tuned here, so the same test after training is still open.

6.4 HOW FIRM THE ORDER IS

Every interval so far draws the speakers again and keeps the 50 languages, so it says how firm a difference is on these languages. Whether it would hold on another fifty is a second question, and Figure 13 asks both. It draws the test set again 5,000 times, the speakers alone and then the languages with their speakers, and counts how often the better model of each pair stays ahead.

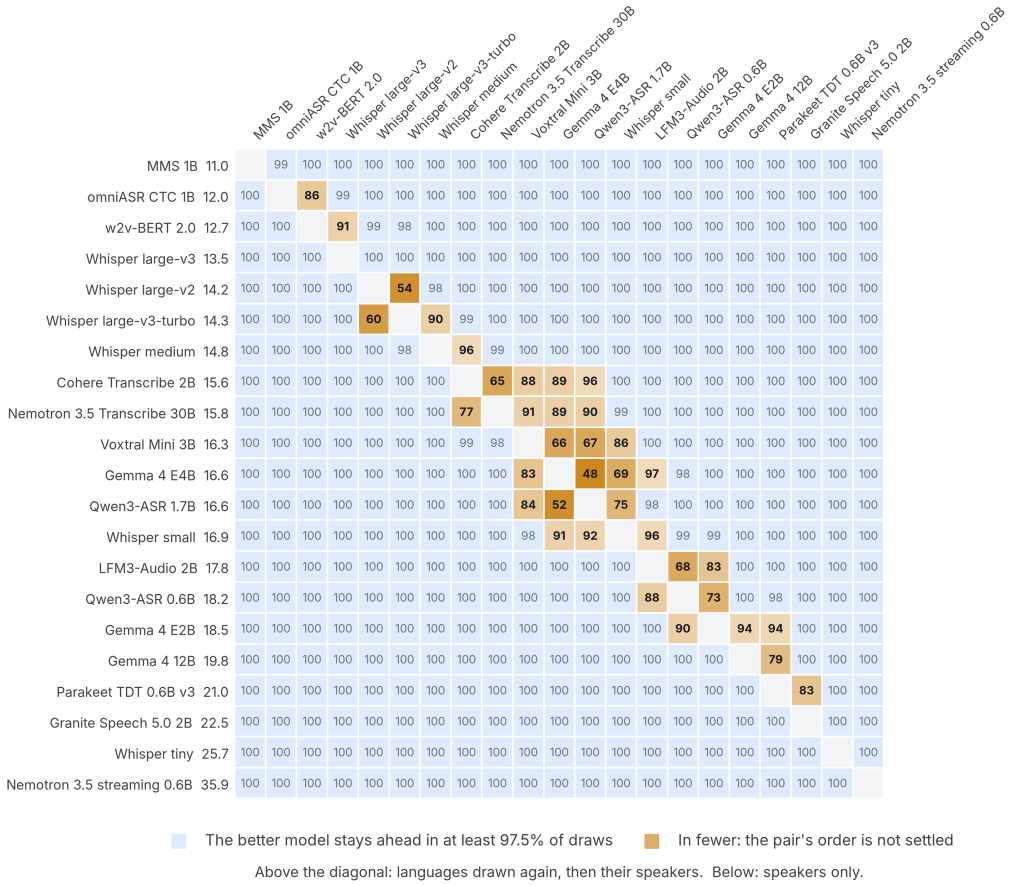

FIGURE 13 The main runs of Table 3, best first, with their macro CER. Each cell is a pair of models: the percentage of 5,000 draws of the test set in which the one ahead on the test set stays ahead. Below the diagonal a draw resamples the speakers inside each language, as the paper's intervals do; above it, the 50 languages are resampled first.

With the speakers alone, 201 of the 210 pairs keep their order in at least 97.5% of the draws; with the languages drawn too, 183. The pairs that come loose are never more than 4 places apart, and MMS 1B is first in 99% of the draws. One of them is the pair Section 6.3 rests on: w2v-BERT 2.0 stays ahead of Whisper large-v3 in 91% of the draws (a difference of -0.5 to 1.9), so on another fifty languages a Whisper may well pass the last CTC model.

6.5 TOLD THE LANGUAGE

Before any training a model can already be told what it is listening to, and Table 4 shows how little that does: the middle column has no training at all. The name of a language in Qwen3-ASR 0.6B's prompt is worth 9.5 points, mostly the script: told the language, it writes the Latin-script languages in Latin on 93% of clips instead of 82%, while the other scripts barely move. The name is worth 6.8 points to LFM3-Audio and 7.9 to 8.7 to Gemma 4's two sizes with an audio encoder. A new Whisper language token, the mean of the others, has little to say yet, and Cohere Transcribe's untrained tag, the mean of its fourteen, does harm: 79.1 against 68.0 for the model told English. Telling is not teaching: the right-hand column is where the gain is.

TABLE 4 The test set's macro CER with nothing trained, and after. The middle column is the first evaluation of each main run: the language's name in the prompt, or a new language token for Whisper and Cohere Transcribe. †Gemma 4 is told every language as downloaded (Section 5), so its middle column is its untouched score and its first the model with no language in the prompt.

MODEL	UNTOUCHED	TOLD THE LANGUAGE, UNTRAINED	FINE-TUNED
Whisper large-v3	50.3	45.1	13.5
Whisper large-v3-turbo	54.6	49.1	14.3
Whisper large-v2	57.8	56.4	14.2
Whisper medium	63.7	76.4	14.8
Whisper small	65.7	61.9	16.9
Whisper tiny	66.6	60.9	25.7
Qwen3-ASR 1.7B	59.3	50.9	16.6
Qwen3-ASR 0.6B	60.9	51.5	18.2
Nemotron 3.5 Transcribe 30B	58.1	54.4	15.8
Voxtral Mini 3B	61.8	57.3	16.3
Granite Speech 5.0 2B	62.1	62.6	22.5
Cohere Transcribe 2B	68.0	79.1	15.6
LFM3-Audio 2B	64.3	57.5	17.8
Gemma 4 E2B	52.2†	43.5	18.5
Gemma 4 E4B	47.7†	39.8	16.6
Gemma 4 12B	75.8†	72.9	19.8

6.6 LANGUAGE BY LANGUAGE

Some languages are almost solved (Figures 14 and 15), the five lowest across any model being Haitian Creole (1), Khasi (2), Finnish (2), Southern Pastaza Quechua (2), Estonian (3). Others resist every model, the five highest being Tibetan (37), Nigerian Fulfulde (30), Breton (25), Kok Borok (22), Lule Sami (20); the numbers are each language's best CER across all models. When every model fails on a language in the same way, the cause is more likely to be that language's data than the models.

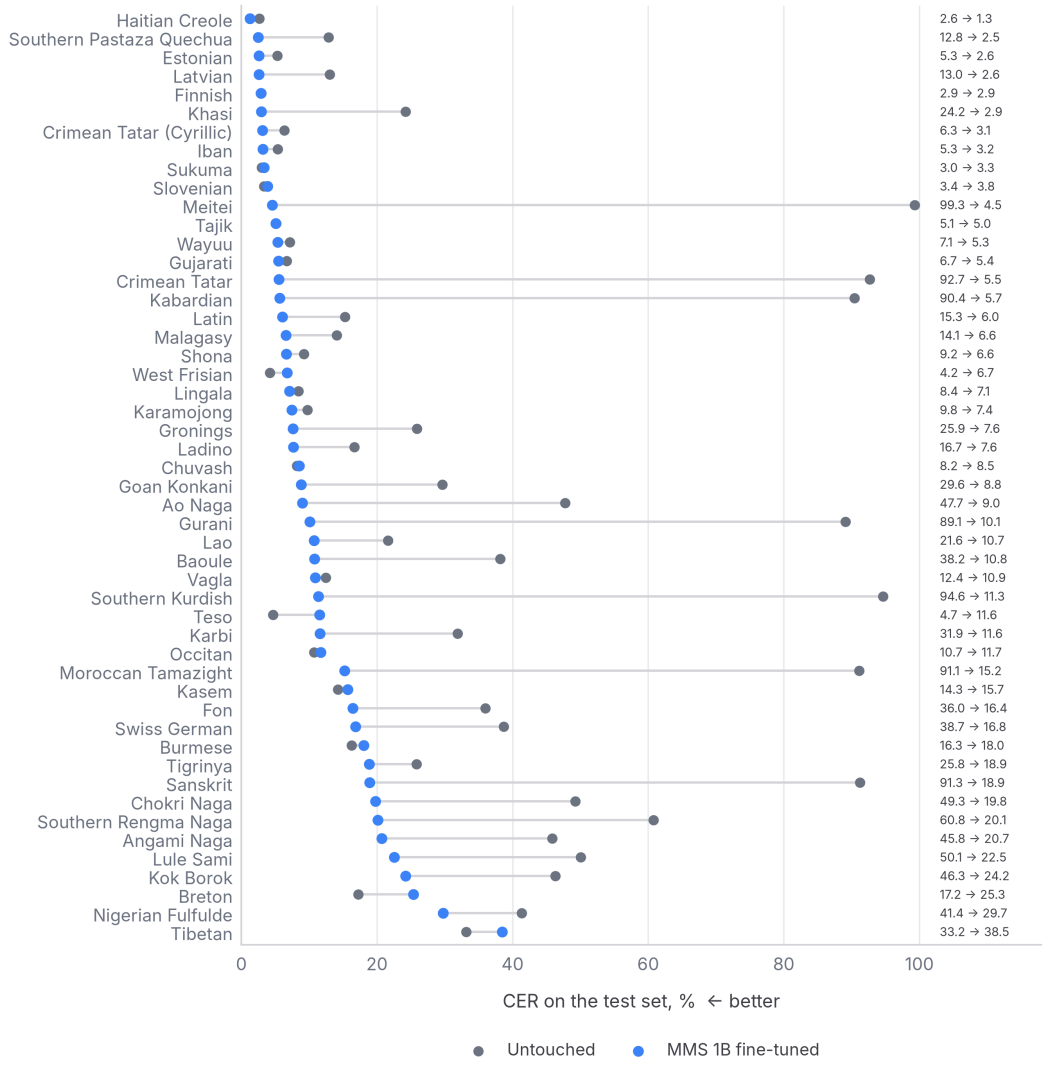

FIGURE 14 The strongest model of the main result, MMS 1B, language by language, before and after.

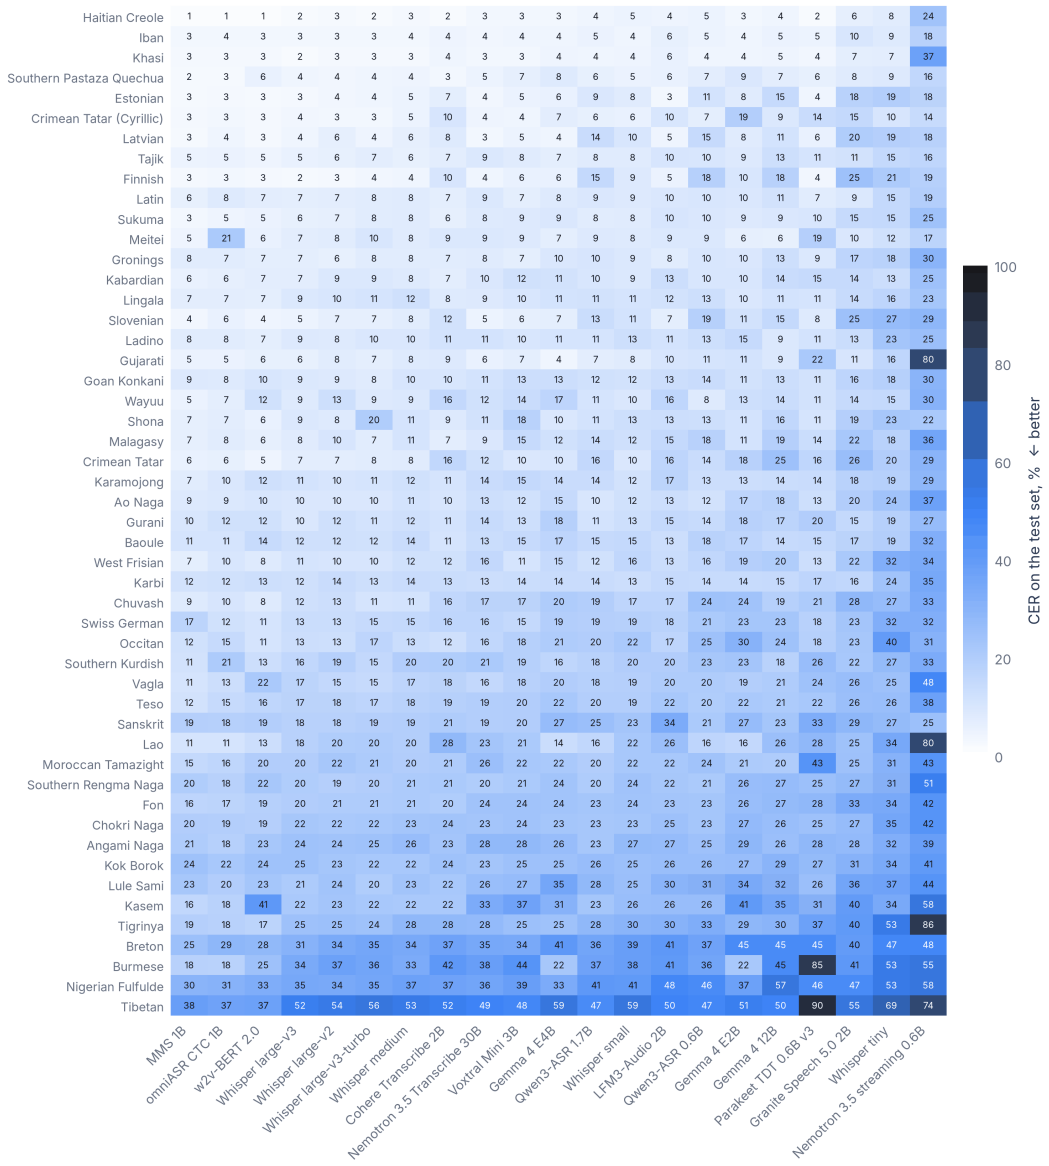

FIGURE 15 Every fine-tuned model on every language. A row that stays dark is a language no model learns; a column that stays dark is a model that cannot write that language's script.

What goes with a hard language is how it was recorded more than how much of it there is (Figure 16). By each language's CER in the median model, the half of the languages with more speakers in their training clips is 6.7 points harder than the half with fewer (1.7 to 12.1, languages resampled), the half with shorter clips 7.4 harder than the half with longer ones (2.4 to 12.7), and the 45 languages whose test speakers are new 6.5 harder than the 5 whose test speakers also train (2.9 to 10.4). A language read by a few voices in long clips is easy; one gathered from many speakers a few seconds at a time is hard. The number of clips says nothing (Spearman's ρ -0.09), and neither the script nor being on no model's list separates the languages by more than its interval. Together the six properties leave 64% of the variance between languages unexplained.

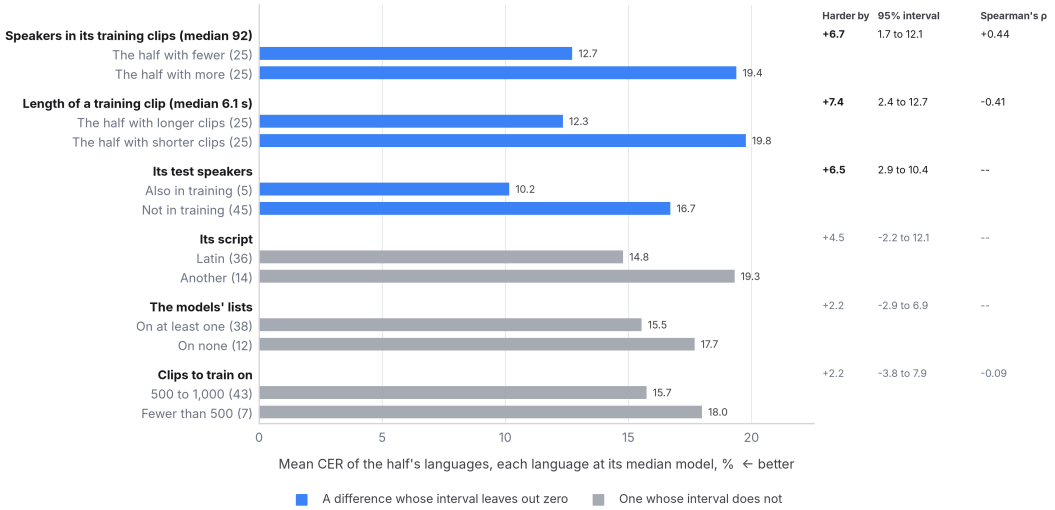

FIGURE 16 The 50 languages split in two by one property at a time, the harder half second: the mean test CER of each half's languages, each language at its median over the main models. Harder by is the difference, with an interval from resampling the languages; Spearman's ρ is between the property, where it is a number, and the language's CER.

6.7 MODEL SIZE

A larger model transfers more within a family of one design (Figure 17). Whisper large-v3 ends 12.2 points below Whisper tiny, 8.8 of them already at small, and Qwen3-ASR 1.7B is 1.6 below the 0.6B. omniASR CTC 1B is 1.3 below the 300M, and the 7B 2.4 below the 1B (2.2 to 2.7): at 9.5 it is the best model we trained, 1.5 points below MMS 1B (1.2 to 1.7). Size counts for more when data is short: on dev, each omniASR size leads the one below it by 3.7 and 3.7 points on the sweep's budget of three epochs on 200 clips a language (Appendix C), and by 1.8 and 3.0 after the full recipe. A change of design can undo it: Gemma 4's 12B, which reads the waveform with no audio encoder, ends at 19.8, behind both smaller sizes, which have one (18.5 and 16.6).

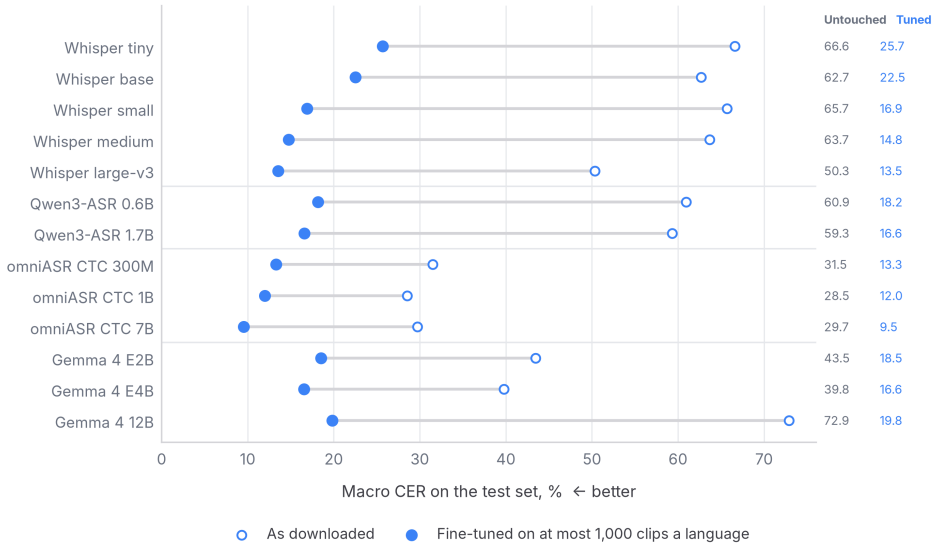

FIGURE 17 Four families by size, as downloaded and fine-tuned with the same recipe. Macro CER on the test set at the epoch dev chose. Gemma 4's 12B is its one size with no audio encoder.

6.8 WHICH WEIGHTS TRAIN

Nemotron 3.5 Transcribe 30B is the one mixture of experts here, and its routed experts are 29.4B of the 32.21B weights its second run trains. To see which parts are worth training, we trained it 18 more ways beside its two runs for 2 epochs, the second run's recipe with some parts frozen, on another seed, at another rate for all of it or for the encoder alone, or with the rate falling to zero over the two epochs, each scored at the epoch dev chooses (Figure 18).

The audio encoder has to train: freezing it costs 0.8 points with the routed experts training (0.3 to 1.2) and 1.8 with them frozen (1.3 to 2.2). So does the language model, since the audio side alone ends 3.1 points behind (2.6 to 3.5). The projector does not matter: freezing it moves a run by at most 0.2 points, while two seeds end 0.3 apart with everything training and 0.2 with the routed experts frozen. Nor do the routed experts much: freezing them costs 0.1 points (-0.3 to 0.5) and leaves 2.84B weights to train, which need no expert parallelism. The first run, which also froze the shared experts and the output head, ends 0.5 behind everything training (0.1 to 0.9), and the shared experts account for it: freezing them with the routed ones costs 0.4 (0.0 to 0.8) and ends level with the first run (-0.4 to 0.4). A lower rate for everything, 3e-6, ends 0.3 points below 1e-5 (-0.1 to 0.7), about as far as two seeds end apart; with the routed experts frozen, 3e-6 ends 0.4 behind 1e-5 (0.0 to 0.7) and 3e-5 3.3 behind (2.8 to 3.7). The encoder wants no faster a rate than the rest: at 3e-5 it ends 1.2 behind (0.9 to 1.5), at 1e-4 7.6.

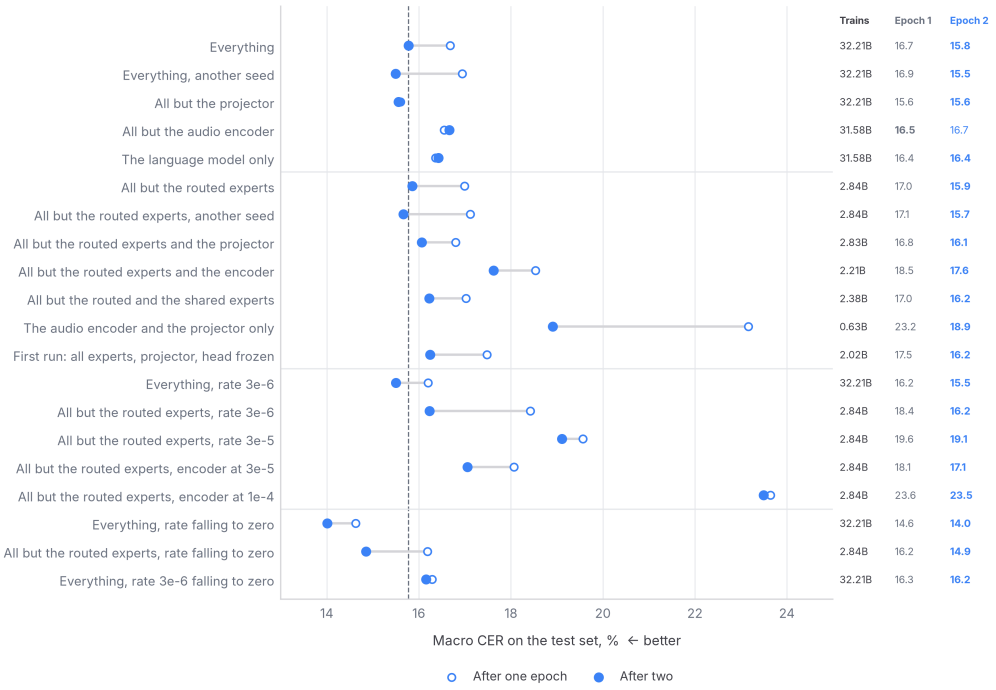

FIGURE 18 Nemotron 3.5 Transcribe 30B trained 2 epochs, part by part: test macro CER after one epoch (hollow) and two (filled), the weights each variant trains, and in bold the epoch dev chose. The rules separate the variants that train the routed experts, those that keep them as downloaded, those that change a rate, and those whose rate falls to zero over the two epochs. The dashed line is everything training, after two epochs.

The schedule gives what the parts do not: with the rate falling to zero over the two epochs, as in Section 8, everything training ends 1.8 points lower than at a constant rate (1.4 to 2.1), at 14.0, and the run with the routed experts frozen 1.0 lower (0.7 to 1.3), at 14.9: both below

every variant at a constant rate (the best of them 15.5). The routed experts, worth 0.1 at a constant rate, are worth 0.8 once the rate falls (0.5 to 1.1). The fall needs a high start: from 3e-6, everything training ends 2.1 behind the fall from 1e-5 (1.8 to 2.5) and 0.7 behind 3e-6 held constant (0.3 to 1.0).

A longer fall does not help. Planned for four epochs instead of two, everything training ends at 14.5 at the epoch dev chooses, its last, 0.4 behind the two-epoch run (0.2 to 0.8) and 1.3 below the main run, which held the rate constant over its four (0.9 to 1.6). Two epochs with the rate falling to zero are the best recipe we found for this model, at 14.0; the main result keeps the constant rate that every other model trained at.

7 What transfers

Studies that need many runs use a small model in place of a large one: Whisper small, which trains an epoch in 0.20 hours against 0.82 for large-v3-turbo, and Qwen3-ASR 0.6B rather than 1.7B. Whisper large-v3-turbo repeats the studies of the next language (Section 7.8), of training longer (Section 8) and of data size (Section 9) to check that a larger model answers the same way.

7.1 PRETRAINED AGAINST RANDOM WEIGHTS

The direct measure of transfer is the same architecture, on the same data, without its pretraining (Figure 19). From random weights every model ends 60 to 77 points behind its pretrained twin, and Parakeet never stops looping: after 8 epochs, 1,145 of the 1,250 test clips still loop. These runs had an untuned rate of 1e-4 and, but for Parakeet's, twice the main runs' epochs, so a better recipe would close some of the gap, though hardly much of it: 82 hours is far too little to learn speech from nothing. Nearly everything a fine-tuned model gets right here, it owes to pretraining on other languages.

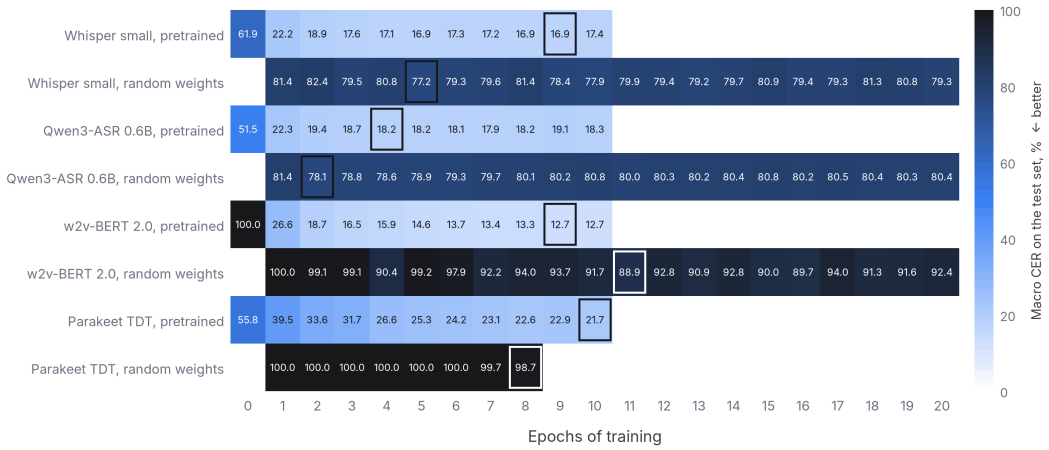

FIGURE 19 The same architecture and the same data, from pretrained and from random weights, on the test set after every epoch. Epoch 0 is the pretrained model before training, told the language where it can be (Table 4); w2v-BERT's head is new, so it starts at 100. A box marks the epoch dev chose; blank cells were not run.

7.2 ENCODER OR DECODER

To see which half carries the transfer, we freeze each half in turn for six epochs, leaving only the rows a new character needs (Figure 20). On a model whose decoder attends to all of the

audio, either half alone learns most of a new language: Qwen3-ASR 0.6B reaches 18.2 with everything training, 21.0 with its encoder frozen and 19.6 with its decoder frozen, and Whisper small behaves the same. So an encoder that has never heard these languages already gives such a decoder most of what it needs. Without such a decoder it is not enough. w2v-BERT's new head on its frozen encoder stays at 77.6, and Parakeet with its encoder frozen at 53.5, hardly below the untouched model: a linear head cannot read a new language's letters straight out of a self-supervised representation, and a transducer's decoder sees the audio only one frame at a time. In these the encoder has to learn the new language.

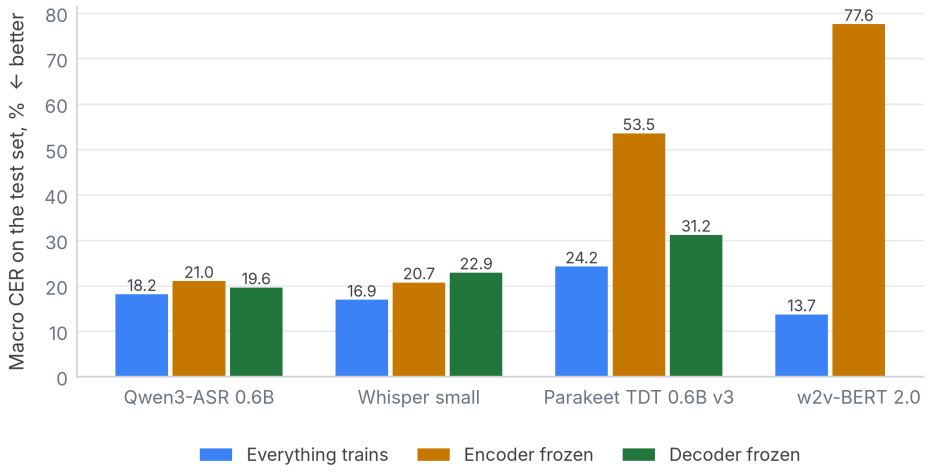

FIGURE 20 Which half has to learn, at the epoch dev chose among the first six. The frozen half keeps its pretrained weights, apart from new rows for new characters. Parakeet's decoder is its prediction network and joint; w2v-BERT has no decoder to freeze.

7.3 THE SCRIPT

The models that can already spell any script end close together on the Latin script and on the others (Figure 21), while the transducers end at 38.2 outside Latin against 24.6 on it: their vocabularies were built from their own languages, and the rows we add for a new script have no pretrained letter behind them (Section 7.4). The script is also the first thing that moves (Figure 22). Untouched, a model writes the reference's script on 0 to 78% of the clips outside Latin; after a single epoch 13 of the models write it on every clip.

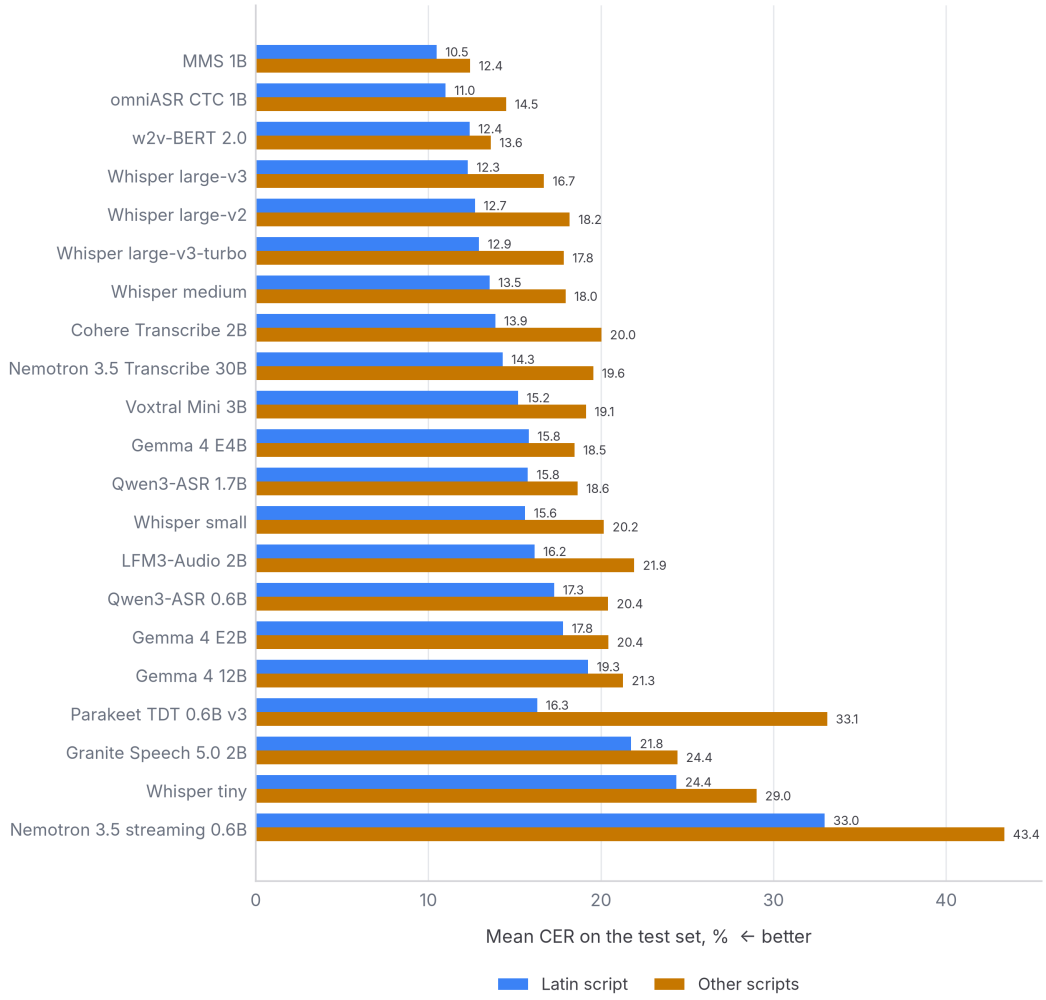

FIGURE 21 Fine-tuned models on the Latin-script languages and on the rest.

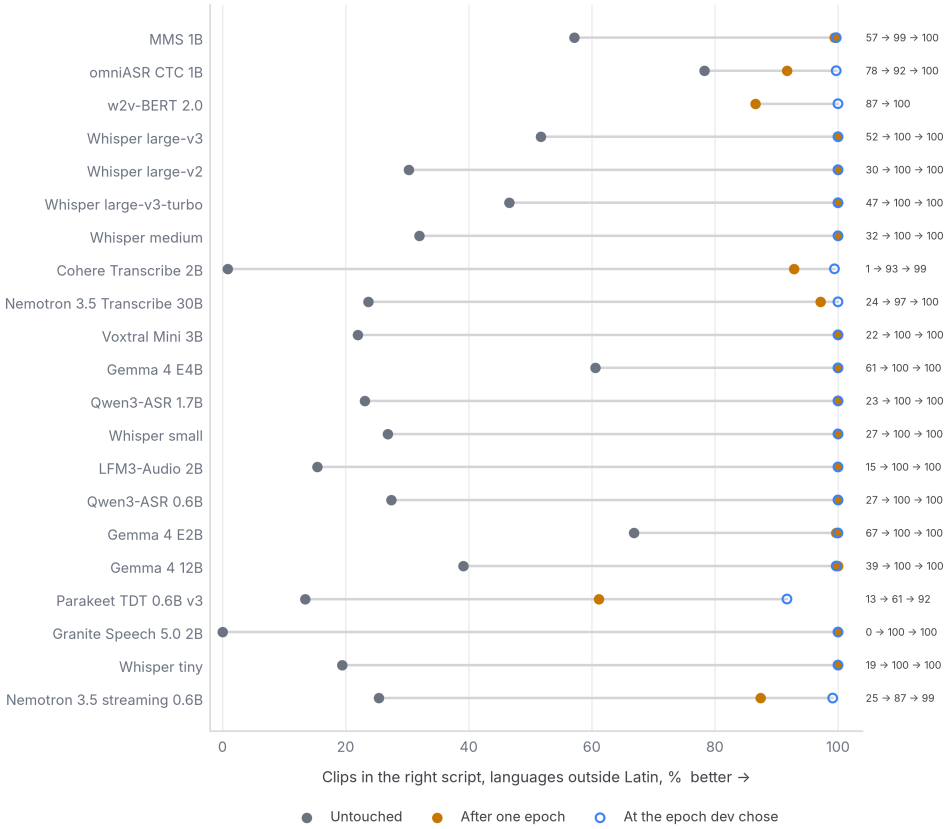

FIGURE 22 Whether the model writes in the language's script, by the Unicode script of most of a clip's letters. Only the languages outside Latin. w2v-BERT has no head as downloaded, so no untouched score; MMS untouched uses its own heads where it has them and English's elsewhere.

7.4 HOW A NEW ROW STARTS

Section 4 gave the transducers' new rows a silent start in place of a first start that was not silent. A start that changes nothing need not learn best, so both ran the full length, and the two models disagree (Figure 23): on Parakeet the silent start is ahead from the first epoch on, on Nemotron it is far behind, 66.3 after one epoch against 45.7.

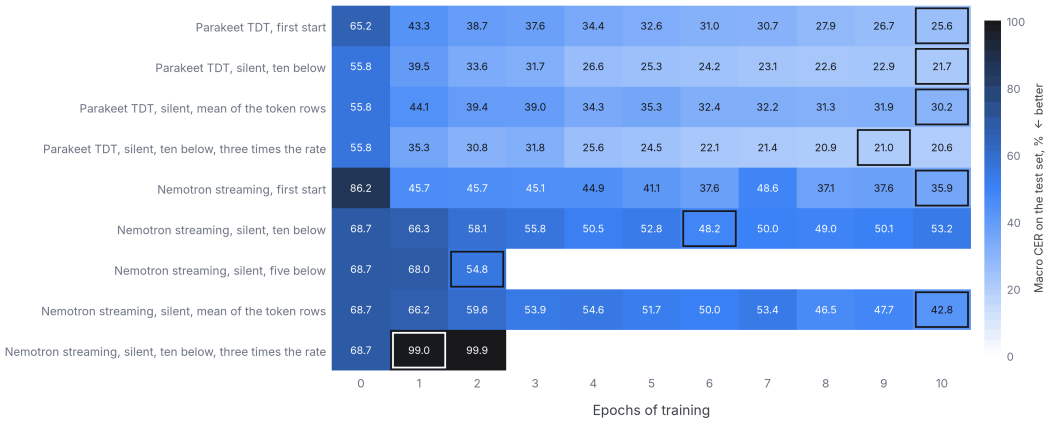

FIGURE 23 Ways to start the transducers' new rows, on the test set after every epoch. Epoch 0 is the extended model before any training: a silent start reads the untouched score there, the first start does not. A box marks the epoch dev chose; blank cells were not run.

The silent Nemotron writes too little, and not only where it needs new rows: after one epoch its output on the Latin-script languages is 50% as long as the reference, against 74% from the first start, and its new rows have hardly left the blank, at a median cosine of 0.997 with the blank's row after training. A start silent for another reason, each new row the mean of the pretrained token rows a point below it, leaves almost no clip empty and still learns slowly, as does five below the blank instead of ten; three times the rate makes the model fall silent, with 1,225 of the 1,250 test clips empty after 2 epochs. On Parakeet the mean start ends behind even the first, while the copy of the blank's row reaches 21.7, and at three times the rate 21.0, Parakeet's main run (Appendix C). Leaving the pretrained model intact at the start is not the same as starting well, and Nemotron's main run is the first start, because dev chose it.

7.5 LANGUAGES ON A MODEL'S LIST

A model that already lists a language starts ahead: as downloaded, Whisper large-v3 is at 32.4 on the 17 languages on its list and 59.6 on the rest. Part of that is the languages, so Figure 24 fits the CER of every fine-tuned model on every language as the model's level plus the language's difficulty, less a lead where the language is on the model's own list. Most models do not list a given language, so its difficulty does not rest on the ones that do, and the fit keeps to the 36 Latin-script languages, so that the script cannot decide it. A list is the card's for most models. For MMS and omniASR it is the model's own table of languages, and for w2v-BERT 2.0, which never saw a transcript, the languages of its unlabelled audio.

As downloaded a listed language is 15 to 29 points ahead. On Whisper that lead is gone after one epoch, and with it the 18% of Latin-script clips it wrote in another script as downloaded; omniASR's is gone after three epochs. The other lists keep some of theirs to the epoch dev chose: MMS 2.6 points (0.9 to 4.2, languages resampled), w2v-BERT 3.4 (1.5 to 5.3), and Parakeet and the two Nemotrons 5.1 (2.9 to 6.6) on the 4 they list, Estonian, Finnish, Latvian and Slovenian. Being on the list can even cost something: Finnish, the only language Qwen3-ASR lists, goes from 8.6 to 15.1 on the larger model when all 50 train together. Nor does its list save Whisper outside the Latin script: on Burmese, Tibetan and Lao, all three on its list, large-v3 ends 7.3 to 15.6 points behind the median of MMS 1B, omniASR CTC 1B and w2v-BERT 2.0, and twice the decoding budget wins back at most 4.0 of that (Section 14).

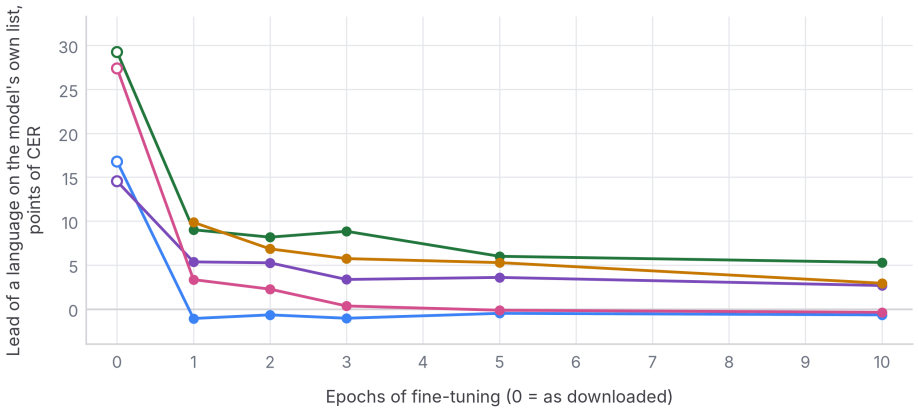

		Epoch						
	Of 36	0	1	2	3	5	10	Interval at 10
Whisper, 7 sizes	11	16.8	-1.1	-0.7	-1.1	-0.5	-0.7	-1.8 to 0.5
Parakeet and the two Nemotrons	4	29.3	9.0	8.2	8.8	6.0	5.3	3.0 to 7.2
MMS 1B	18	14.5	5.4	5.3	3.4	3.6	2.7	1.0 to 4.4
omniASR CTC, 3 sizes	25	27.4	3.3	2.3	0.3	-0.1	-0.4	-2.1 to 1.3
w2v-BERT 2.0	7	--	9.9	6.8	5.7	5.3	2.9	1.1 to 4.5

FIGURE 24 What a language on a model's own list is still worth, epoch by epoch, on the 36 Latin-script languages. The test CER of all 24 models fine-tuned with the main recipe is fitted as a model's level plus a language's difficulty, less a lead where the language is on the model's own list; the models that share a list share a lead. Epoch 0 (hollow) is the models as downloaded, and a run that stopped before an epoch stands at its last. Of 36: the languages on the list. The interval resamples the languages.

To see whether a relative on the list helps, Table 5 keeps to Latin-script languages and splits them by whether the model lists the language, another of its family, or neither. Before training a relative seems to help the models with short lists, Nemotron 3.5 Transcribe being at 50.7 with one and 61.2 without, but training leaves 1.7 points of that there and 1.4 on Qwen3-ASR 1.7B, and on Whisper the 4 languages with no relative on its list even end ahead. Family is a weak guide to what transfers.

TABLE 5 Mean CER of the Latin-script languages by what the model's card says about them. A relative is another language of the same top-level family.

MODEL	ON ITS LIST			A RELATIVE ON ITS LIST			NO RELATIVE		
	N	UNTOUCHED	TUNED	N	UNTOUCHED	TUNED	N	UNTOUCHED	TUNED
Whisper large-v3	11	18.1	8.5	21	56.5	14.6	4	55.2	10.2
Whisper large-v3-turbo	11	20.3	10.8	21	58.3	14.6	4	61.7	10.2
Whisper medium	11	35.0	10.5	21	66.9	15.6	4	74.0	10.9
Whisper small	11	35.0	13.4	21	70.2	17.5	4	68.5	11.4
Whisper tiny	11	41.3	22.9	21	65.8	26.5	4	68.1	17.1
Qwen3-ASR 1.7B	1	8.6	15.1	23	52.0	15.3	12	51.9	16.7
Qwen3-ASR 0.6B	1	13.0	18.4	23	51.5	16.8	12	59.2	18.2
Parakeet TDT 0.6B v3	4	5.3	5.4	10	36.0	16.2	22	53.2	18.3
Nemotron 3.5 streaming 0.6B	4	9.7	20.8	18	65.3	35.0	14	76.2	33.8
Nemotron 3.5 Transcribe 30B	4	3.7	4.0	20	50.7	15.0	12	61.2	16.7

7.6 LANGUAGES NO MODEL LISTS

For 12 of the 50, no model in the study starts with anything: Angami Naga, Ao Naga, Chokri Naga, Gronings, Gurani, Karbi, Ladino, Lule Sami, Moroccan Tamazight, Southern Kurdish,

Southern Rengma Naga and Swiss German are on no card and in neither MMS's nor omniASR's table, in any script. Nothing a model does there is mixed with what it knew of the language, so Figure 25 repeats the main result on them. As downloaded every model is at 46 to 71, the two systems built for the long tail included: omniASR CTC 1B, at 22.9 on the other 38 languages, is at 46.5 on these. Fine-tuned, the models gain 27.9 to 51.0 points and end in nearly the order they end in on the other 38 (Spearman's ρ 0.97), the best at 11.7 (omniASR CTC 7B).

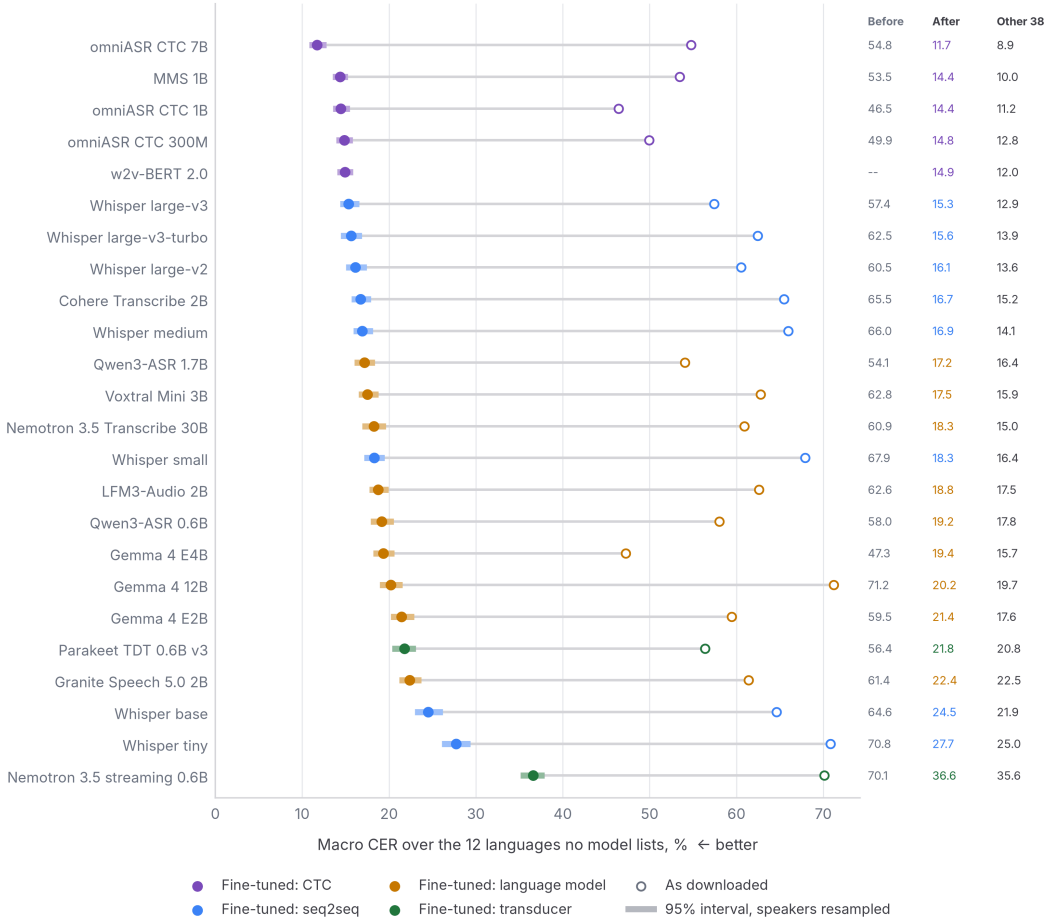

FIGURE 25 The 12 languages on no model's list, in any script: all 24 models fine-tuned with the main recipe, as downloaded and at the epoch dev chose, macro CER over the 12. Other 38: the fine-tuned model on the rest of the test set. MMS as downloaded has no head for these and uses English's; w2v-BERT has none at all.

So the order of Section 6.3 is not the lists' doing, though some of its margins are. Every CTC model is still ahead of every other model, but w2v-BERT 2.0 and Whisper large-v3 are 0.4 apart here (-0.3 to 1.3), and MMS's lead of 1.0 over omniASR CTC 1B on all 50 is 0.1 (-0.3 to 0.6). The twelve are no easier than the rest: 23 of the 24 models score worse on them, by a median of 2.2 points.

7.7 ONE LANGUAGE OR FIFTY

Most low-resource work trains one language at a time. Figure 26 trains Qwen3-ASR 0.6B on each language alone, with the same clips and recipe: alone is better by 2.8 points on average over the 50 languages, and only 8 do better with the other 49. Part of that is each language

choosing its own epoch on its own 25 dev clips, but at the last epoch of both, where nothing was chosen, alone is still ahead, 15.9 against 18.3, and the 7 languages with fewer than 500 clips also do better alone, so the scarce ones do not lean on the others. A model of 0.6B parameters shared by fifty languages may have less of itself for each. Fifty models are fifty times the storage, though, and the joint model is one.

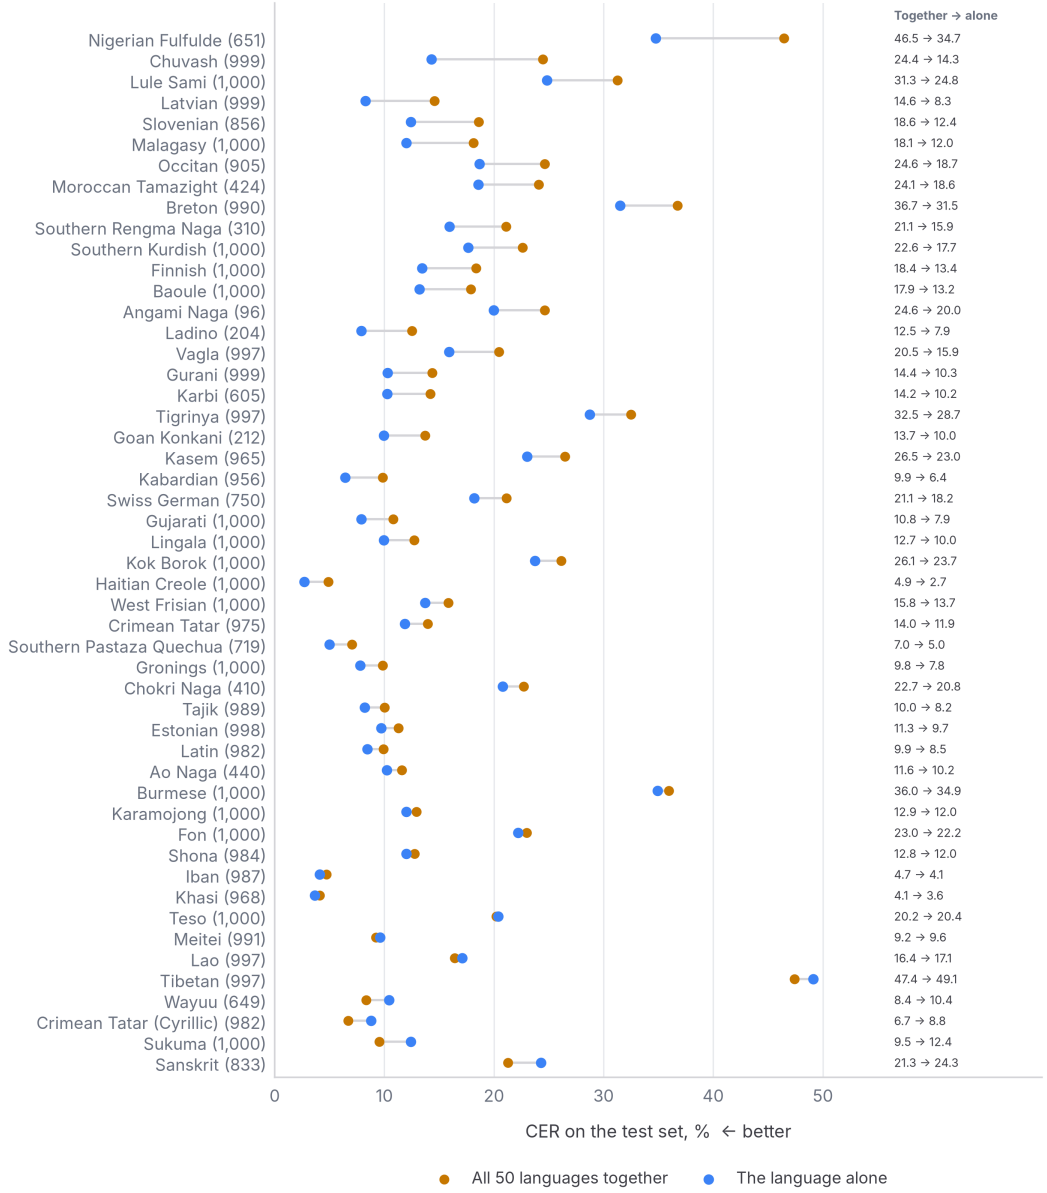

FIGURE 26 Each language trained alone against all 50 trained together, Qwen3-ASR 0.6B, ten epochs, each at the epoch its own dev chose. In brackets, the language's training clips. Sorted by the difference: at the bottom, the languages that did better with the other 49.

7.8 THE NEXT LANGUAGE

Does a model that has learned forty low-resource languages learn a forty-first faster than the model it started from? We drew 10 of the 50 at random, trained Qwen3-ASR 0.6B, Whisper small, Whisper large-v3-turbo and MMS on the other forty with the main recipe, and then trained each on the held-out ten with 25, 100 and 1,000 clips a language, from that first stage and from the pretrained model (Figure 27).

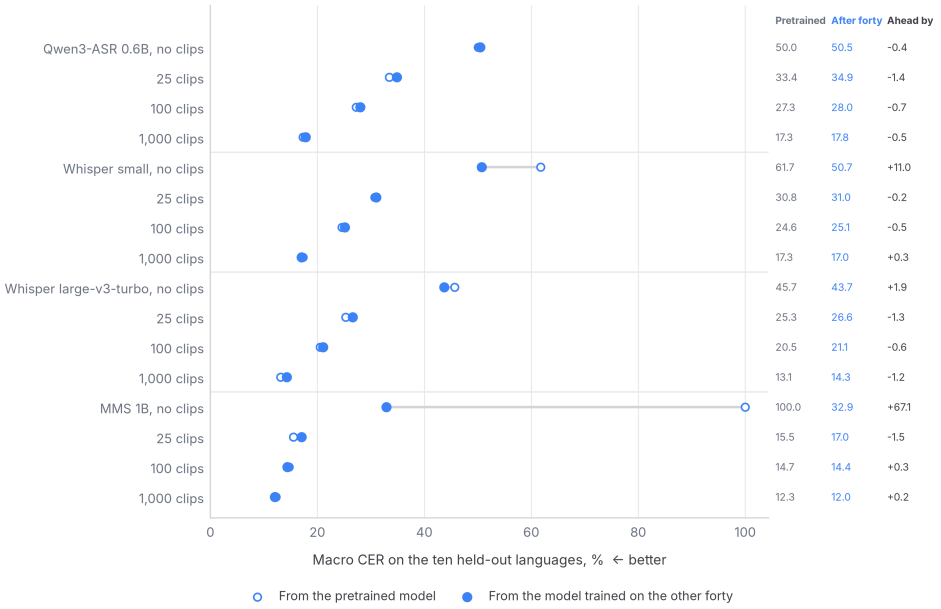

FIGURE 27 Ten languages held out of a first stage on the other forty: the second stage from the pretrained model (hollow) and from the first stage (filled), by how many clips of each held-out language it trains on. No clips is the model before the second stage. Test macro CER on the ten at the epoch dev chose; the last column is how far the first stage puts the second ahead. One training seed; the text adds a second.

Once there are clips to learn from, it does not. Before the second stage, the first puts Whisper small 11.0 points ahead on ten languages it has never heard and large-v3-turbo 1.9, and takes MMS, whose new head writes nothing until it is trained, from nothing to 32.9; Qwen3-ASR does not move. With clips of the ten, the head start is gone, and with very few it turns into a small cost. Over two training seeds on the same clips, the first stage leaves the model behind in 7 of the 8 pairs at 25 clips, by up to 1.7 points (4 with an interval that leaves out zero), and ahead in Whisper large-v3-turbo's second seed, by 0.2 (-0.6 to 1.0); at 100 and 1,000 clips the 16 differences go both ways, none larger than 1.9 points, 3 of them behind and 1 ahead with an interval that leaves out zero. Training the ten beside the other forty does not help either: Whisper small's main run, on all fifty at once, scores 17.9 on the ten, and the ten alone at 1,000 clips 17.3; large-v3-turbo's, 15.4 and 13.1.

8 Epochs

Whether training longer helps depends on the decoder, and for most of them the answer comes early. Figure 28 has dev after every epoch at a constant rate, and Table 6 has the same runs on the test set.

On seq2seq and language-model decoders, the first epoch has 90% of the whole gain in the median model and the third 99%, and dev is within a point of its best by epoch 3. After that the last column of Table 6 is flat for most of them, and Whisper large-v3, Whisper large-v3-turbo, Whisper medium, Whisper tiny and Cohere Transcribe 2B gain no more than 1.6 points. The CTC models and the transducers, which have new rows to learn, keep going: from epoch 3 to the last, w2v-BERT improves by 3.8 points, MMS by 3.5, omniASR CTC 1B by 3.2, Parakeet by 11.2 and Nemotron streaming by 9.2.

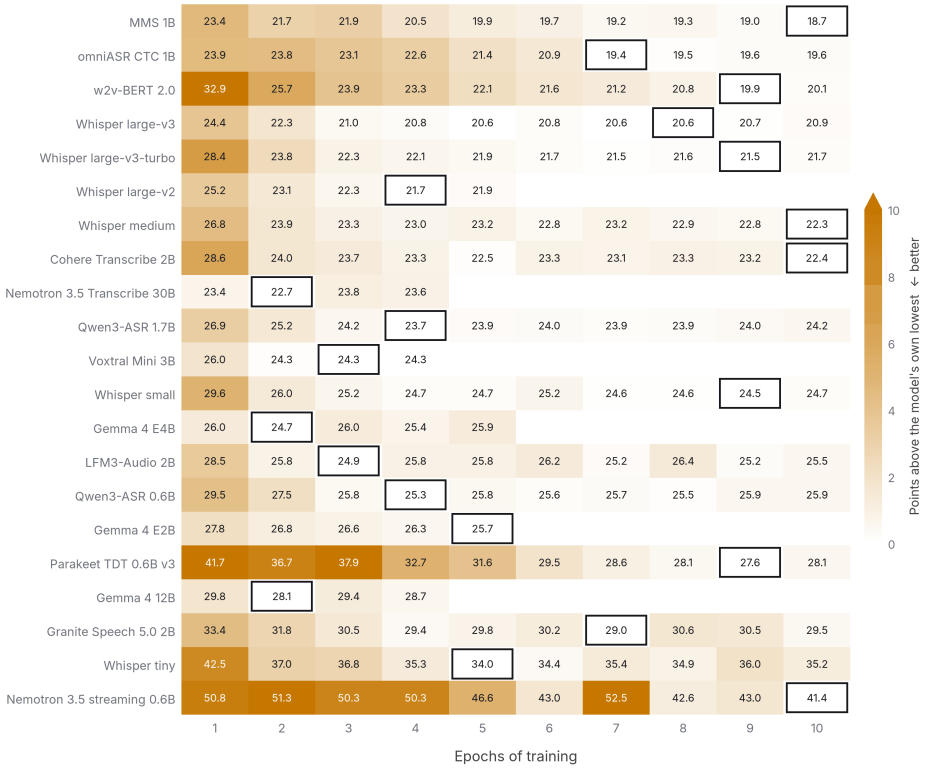

FIGURE 28 Dev after every epoch at a constant learning rate, one row a model. Each cell is the dev macro CER; its colour is how far it is above that model's own lowest. A box marks the epoch dev chose; blank cells were not run.

TABLE 6 Test macro CER after 1, 2, 3, 5 and 10 epochs of the same runs, shown for the record: no epoch was chosen on it. Best: the epoch dev chose. Within 1 pt: the first epoch within one point of that best on dev. The last column is how much the run's last saved epoch improves on epoch 3 on the test set, with an interval from resampling the same speakers on both sides.

MODEL	EPOCH 1	EPOCH 2	EPOCH 3	EPOCH 5	EPOCH 10	BEST (DEV)	WITHIN 1 PT AT	LAST GAINS ON EPOCH 3 (95% CI)
Whisper large-v3	17.1	14.8	13.8	13.6	13.4	8	3	+0.5 (+0.2 to +0.7)
Whisper large-v3-turbo	21.4	16.5	15.1	14.2	14.3	9	3	+0.8 (+0.4 to +1.2)
Whisper large-v2	18.5	15.4	14.7	14.3	–	4	3	+0.4 (-0.0 to +0.8)
Whisper medium	19.2	16.2	15.6	15.2	14.8	10	3	+0.8 (+0.5 to +1.2)
Whisper small	22.2	18.9	17.6	16.9	17.4	9	3	+0.2 (-0.3 to +0.6)
Whisper tiny	34.3	30.1	28.5	25.7	26.9	5	5	+1.6 (+0.8 to +2.3)
Qwen3-ASR 1.7B	19.4	18.0	17.3	16.6	17.3	4	3	+0.1 (-0.4 to +0.6)
Qwen3-ASR 0.6B	22.3	19.4	18.7	18.2	18.3	4	3	+0.4 (-0.1 to +0.9)
Nemotron 3.5 Transcribe 30B	16.7	15.8	16.3	–	–	2	1	+0.0 (-0.4 to +0.4)
Parakeet TDT 0.6B v3	35.3	30.8	31.8	24.5	20.6	9	8	+11.2 (+10.2 to +12.3)
Nemotron 3.5 streaming 0.6B	45.7	45.7	45.1	41.1	35.9	10	10	+9.2 (+8.6 to +9.9)
w2v-BERT 2.0	26.6	18.7	16.5	14.6	12.7	9	8	+3.8 (+3.4 to +4.1)
MMS 1B	16.4	14.4	14.5	12.3	11.0	10	6	+3.5 (+3.1 to +3.9)
omniASR CTC 1B	16.6	15.6	14.9	14.2	11.7	7	7	+3.2 (+3.0 to +3.4)
Voxtral Mini 3B	18.1	16.7	16.3	–	–	3	2	+0.2 (-0.4 to +0.7)
Granite Speech 5.0 2B	26.9	24.4	22.8	23.1	22.5	7	4	+0.3 (-0.2 to +1.0)
Cohere Transcribe 2B	22.0	17.1	16.3	15.5	15.6	10	4	+0.7 (+0.3 to +1.2)
LFM3-Audio 2B	21.2	18.2	17.8	18.1	18.2	3	2	-0.4 (-0.9 to +0.1)
Gemma 4 E2B	18.9	18.1	18.4	18.5	–	5	3	-0.1 (-0.8 to +0.5)
Gemma 4 E4B	17.4	16.6	17.4	18.3	–	2	2	-0.9 (-1.5 to -0.3)
Gemma 4 12B	22.6	19.8	21.4	–	–	2	2	+0.4 (-0.4 to +1.1)

To find where that stops we ran the longest curves again on another seed (Figure 29). Thirty epochs of Whisper small, Whisper large-v3-turbo and Qwen3-ASR 0.6B are no better than

the first ten (on turbo dev chooses epoch 9), though Whisper's training loss falls from 0.27 at epoch 3 to 0.01 at epoch 20: it is learning its training set by heart. In the CTC models' longer runs, w2v-BERT, whose head is new, goes on improving to epoch 30 of 30, 2.2 points below its epoch 10, while omniASR CTC 1B, whose head came with the model, has flattened out, dev choosing epoch 20 by only 0.3 points over epoch 10. MMS, whose head is also new, is still improving at epoch 20, the last it ran: dev chooses it, 0.8 points below epoch 10. Parakeet at 1e-4, copied at epoch 7 and carried on to twenty, ends 1.2 points below its own epoch 10.

Epochs of training	Qwen3-ASR 0.6B		Whisper small		Whisper large-v3-turbo		w2v-BERT 2.0		omniASR CTC 1B		Parakeet TDT 0.6B v3		MMS 1B	
	dev	test	dev	test	dev	test	dev	test	dev	test	dev	test	dev	test
1	31.1	23.4	29.4	21.9	29.1	21.5	37.2	31.2	24.9	17.1	43.9	39.5	23.2	15.9
2	27.6	20.2	26.5	18.9	23.5	16.3	26.7	19.7	23.8	16.3	39.0	33.6	22.0	14.7
3	25.6	18.4	25.1	17.4	22.4	15.1	24.9	17.7	22.8	15.5	38.2	31.7	20.9	13.6
4	25.4	18.5	24.5	17.2	21.7	14.4	23.4	16.2	22.4	14.6	32.7	26.6	20.6	13.1
5	26.0	18.2	24.9	17.0	21.9	14.8	21.8	14.5	22.0	13.9	31.2	25.3	20.0	12.4
6	25.2	18.1	25.0	17.0	21.6	14.3	21.5	14.2	20.0	12.3	30.2	24.2	19.1	11.9
7	25.1	18.2	24.6	17.2	21.8	14.1	21.2	13.6	19.5	12.0	29.3	23.1	19.1	11.8
8	25.8	18.5	24.9	17.0	21.6	14.4	20.6	13.5	19.4	11.9	28.2	22.5	19.7	12.1
9	25.8	17.9	25.1	17.2	21.2	13.8	20.4	12.8	19.7	11.8	28.2	22.7	18.8	11.3
10	25.8	18.6	24.9	17.1	21.9	14.4	20.1	12.8	19.0	11.6	27.6	21.9	18.5	11.3
11	25.2	18.3	24.6	17.2	22.1	14.3	19.9	12.5	19.2	11.8	27.7	21.7	18.4	11.1
12	25.6	18.4	24.7	17.3	21.2	14.2	19.4	11.8	19.2	11.7	27.6	21.7	18.3	10.9
13	25.2	17.7	24.3	16.9	21.7	14.2	19.7	11.9	19.0	11.5	27.0	21.2	18.0	10.7
14	25.8	18.2	24.5	16.9	21.7	14.1	19.5	11.7	18.8	11.6	27.5	21.1	18.5	11.0
15	25.7	18.0	24.7	17.1	22.1	14.3	19.0	11.5	18.8	11.4	27.1	20.7	18.0	10.8
16	25.5	18.2	25.1	17.0	21.3	14.1	19.2	11.7	18.9	11.3	26.9	20.8	18.0	10.7
17	25.4	17.7	24.8	16.9	21.6	13.8	18.8	11.3	18.8	11.3	27.1	20.9	17.8	10.6
18	25.8	18.3	24.8	17.1	21.5	13.8	19.2	11.2	18.8	11.4	26.4	20.5	17.8	10.3
19	25.9	17.7	24.8	17.3	21.3	13.6	18.7	11.2	18.9	11.6	27.2	20.9	17.9	10.5
20	25.5	17.8	25.4	17.5	22.0	14.5	18.8	11.0	18.7	11.6	26.4	20.8	17.7	10.3
21	26.9	19.5	24.9	16.7	21.9	13.9	18.4	10.9						
22	25.8	18.3	24.5	17.3	21.4	14.0	18.6	11.1						
23	25.6	18.0	24.6	16.9	21.2	13.8	18.3	11.1						
24	25.6	18.0	24.8	17.5	21.4	14.0	18.6	11.0						
25	26.2	18.3	25.6	17.5	21.3	14.1	18.2	11.0						
26	25.8	17.8	25.2	17.5	21.7	14.2	17.7	10.6						
27	25.6	18.0	24.7	17.2	21.6	14.1	18.0	10.7						
28	26.3	18.6	25.2	17.4	21.8	14.2	17.9	10.8						
29	25.7	18.2	24.9	17.2	21.8	14.5	18.0	10.7						
30	26.0	18.2	24.8	17.7	21.5	13.7	17.6	10.6						

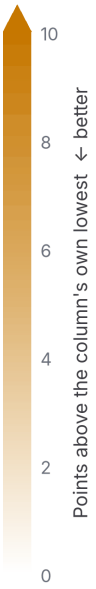

FIGURE 29 The longest runs at a constant rate, on another seed, except Parakeet's, which carries its run at 1e-4 on from epoch 7. Dev and test macro CER after every epoch; the test columns are for the record. The colour is how far a cell is above its column's lowest. A box marks the epoch dev chose; blank cells were not run.

Scarce languages might need more passes, since they see fewer clips in each. Figure 30 says not many more: on dev the 7 languages with under 500 clips come within a point of their best at epoch 6, the other two groups at 3 and 3.

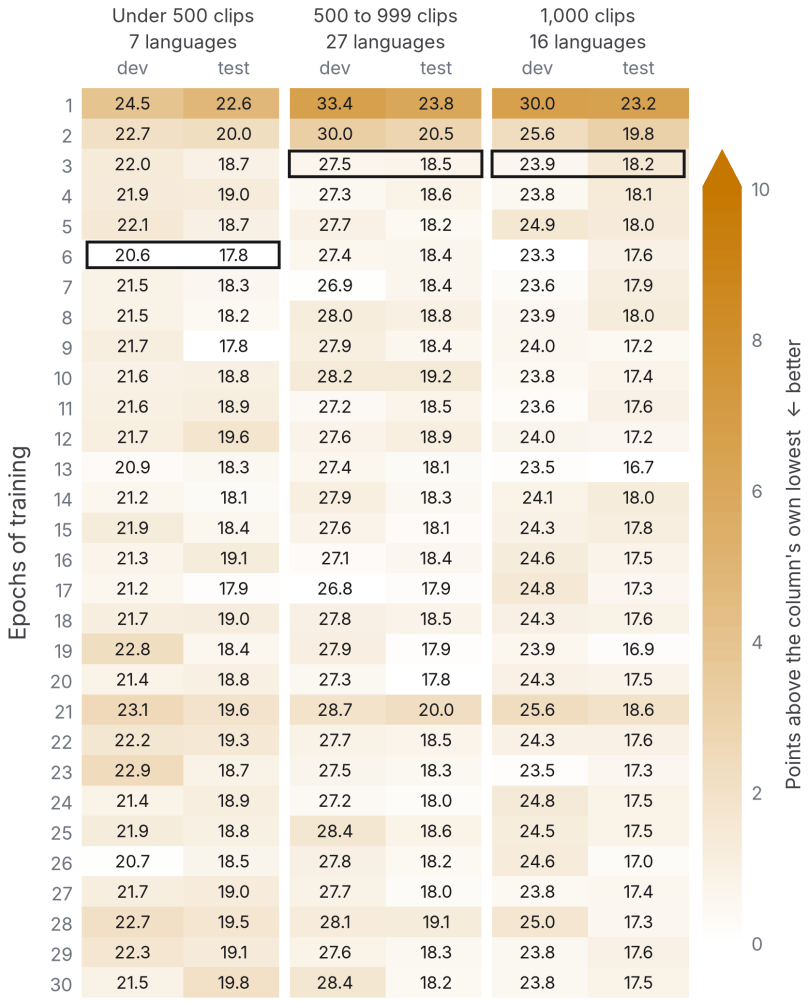

FIGURE 30 The long run of Qwen3-ASR 0.6B, by how many clips a language has to train on: the mean dev and test CER of each group's languages after every epoch. The colour is how far a cell is above its column's lowest. A box marks the first epoch within a point of the group's lowest on dev.

A curve through one long run is not a training planned for that length, so we also ran the usual recipe at each length, the rate falling to zero at the last epoch (Figure 31). From three epochs on it beats the constant rate by 0.6 to 1.1 points on Qwen3-ASR 0.6B and 0.2 to 0.9 on Whisper small, where two seeds of the constant recipe lie up to 0.5 and 0.4 points apart, and it hardly moves the answer about epochs: Qwen planned for three epochs scores 17.6 and planned for ten 17.5. The two models still learning at ten epochs lose from the falling rate instead: Parakeet at every length, by more than its seeds differ except at three, and w2v-BERT until ten, where the cost is gone. That ten-epoch w2v-BERT run is on seed 1, because on seed 0 its new head stalled, writing the same output for every clip for 3 epochs until we stopped it; the same command run again broke through within its first epoch, so the stall was chance, not the seed.

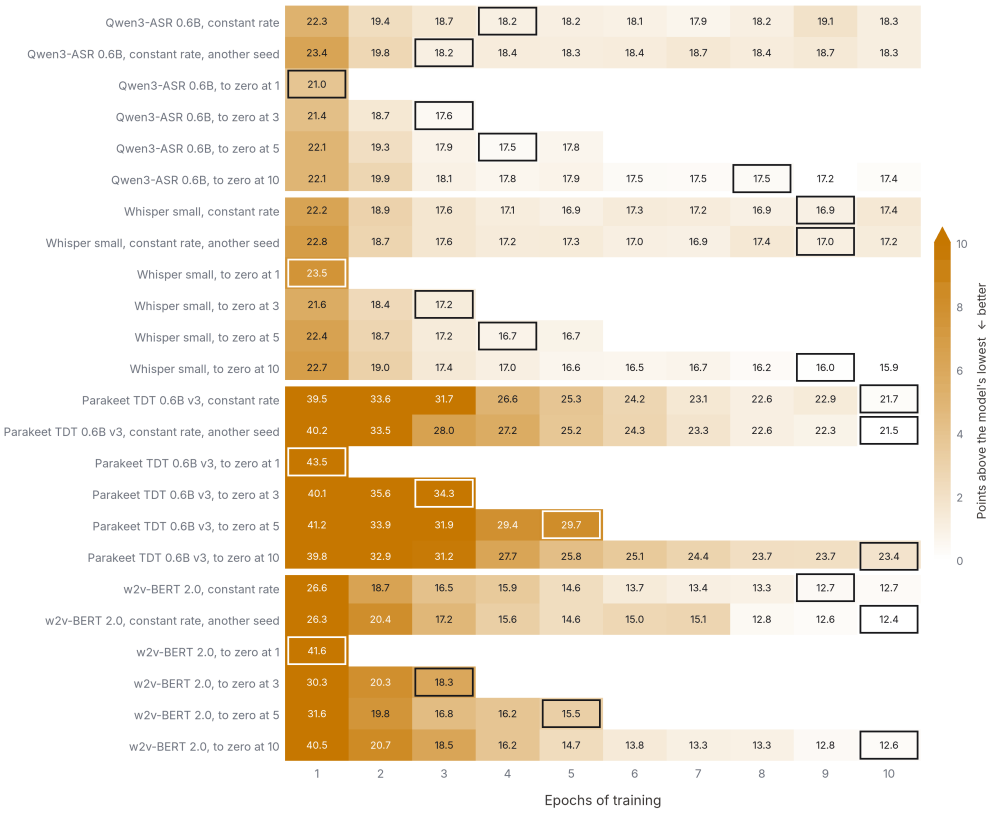

FIGURE 31 Each model's main run at a constant rate, the same recipe on another seed, and separate runs planned for 1, 3, 5 and 10 epochs whose rate falls linearly to zero at their last. Test macro CER after every epoch, for the record; a box marks the epoch dev chose. The colour is how far a cell is above the lowest of that model's rows; blank cells were not run. w2v-BERT's ten-epoch run is on seed 1, after a stall on seed 0 (see text).

9 Data size

A few clips already transfer (Table 7, Figure 32): twenty-five clips a language, about 2 hours over all 50, take Qwen3-ASR 0.6B from 60.9 to 39.3, Whisper small from 65.7 to 38.4 and Whisper large-v3-turbo from 54.6 to 32.3, and every doubling after that buys less. Parakeet TDT 0.6B v3 gets less from the first few clips, 45.6 after 25.

TABLE 7 The data-size runs. Clips presented is clips times epochs. Macro CER on the test set at the epoch dev chose. A language with fewer clips than a row's size uses all it has.

CLIPS A LANGUAGE, AT MOST	CLIPS	HOURS	EPOCHS	CLIPS PRESENTED	QWEN3-ASR 0.6B	WHISPER SMALL	WHISPER LARGE-V3-TURBO	PARAKEET TDT 0.6B V3
25	1,209	2	60	72,540	39.3	38.4	32.3	45.6
50	2,419	4	40	96,760	33.4	31.7	26.7	43.2
100	4,820	8	24	115,680	28.4	26.7	22.4	34.7
200	9,542	17	15	143,130	24.4	22.5	19.1	32.8
500	22,896	42	10	228,960	20.4	18.7	15.9	27.3
1,000	42,816	82	10	428,160	18.2	16.9	14.3	21.7
2,000	77,278	159	5	386,390	16.1	15.4	12.7	21.4
4,000	129,236	279	3	387,708	15.3	14.4	12.8	20.7

Data also beats epochs at the same cost. On Qwen3-ASR 0.6B a thousand clips a language for 3 epochs are 9.7 points better than a hundred for 24, the same number of clips presented; Whisper small gains 9.1 the same way, large-v3-turbo 7.3, and Parakeet, which keeps improving for longer, 2.9. At the top of Table 7, four thousand clips a language for three

epochs present fewer clips than a thousand for ten and still beat them, by 2.9, 2.5, 1.5 and 1.0 points in the same order, each interval above zero. New clips are worth more than further passes over the old ones, and on every model but Parakeet time past the third epoch would be better spent collecting clips.

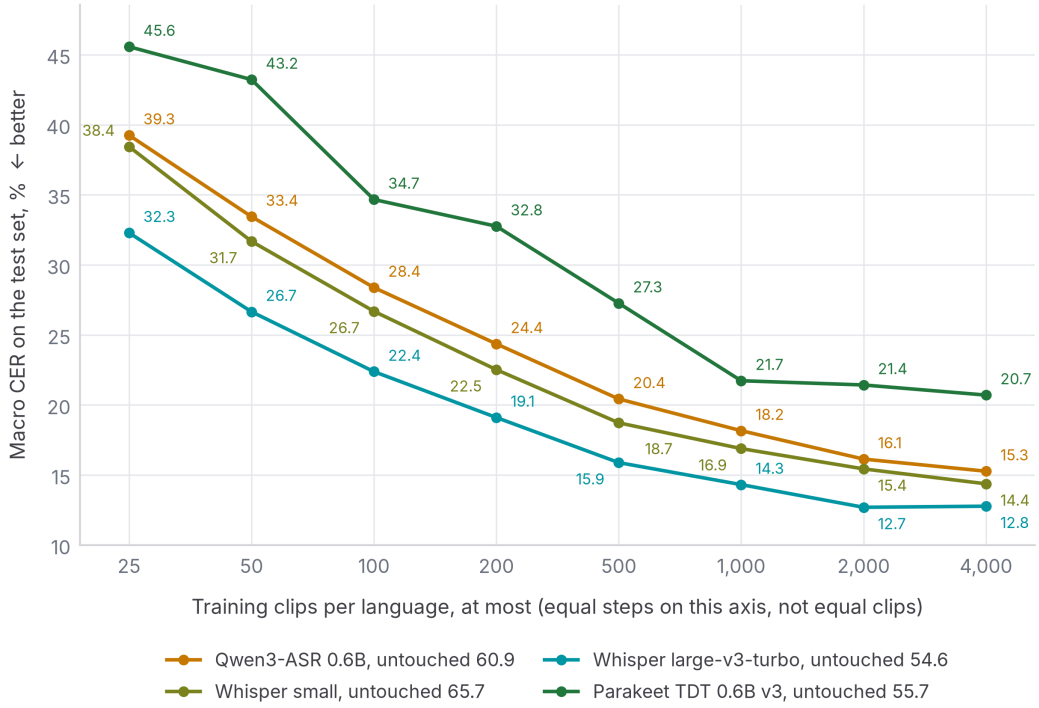

FIGURE 32 Macro CER against the number of clips per language, at the epoch dev chose; the legend gives each model untouched. Smaller sets train for more epochs, and the two largest for fewer (Table 7).

10 Common Voice

Common Voice 22 has at least 100 clips of 11 of our languages, and Figure 33 trains on them three ways: on this corpus, on Common Voice, and on both.

We took at most 4,000 a language, leaving out any sentence in our dev or test set, 30,535 clips and 41.9 hours against this corpus's 10,157 and 23.4 of the same languages, and trained Qwen3-ASR 0.6B with the usual recipe. On Common Voice alone, with three times the clips, it ends 3.5 points behind this corpus alone, partly a change of kind, since our test clips come from the same corpus as our training clips: Common Voice alone is better on 6 of the 8 languages where it has more clips than this corpus and far worse where it has fewer. Together they beat either, 4.1 points below this corpus alone (3.2 to 4.9), and Whisper small and Parakeet TDT 0.6B v3, trained the same three ways, agree. Where Common Voice has a language, it belongs in the training set.

The run on this corpus alone also checks Section 7.7 at a larger scale: trained on these 11 languages only, Qwen reaches 16.9 on them, 2.6 points better than the main run trained on all 50 (1.4 to 4.0; Figure 33, purple), and Whisper small does 1.5 better. On Parakeet, still learning at ten epochs (Section 8), there is no difference. Fewer languages sharing a model again do better on each.

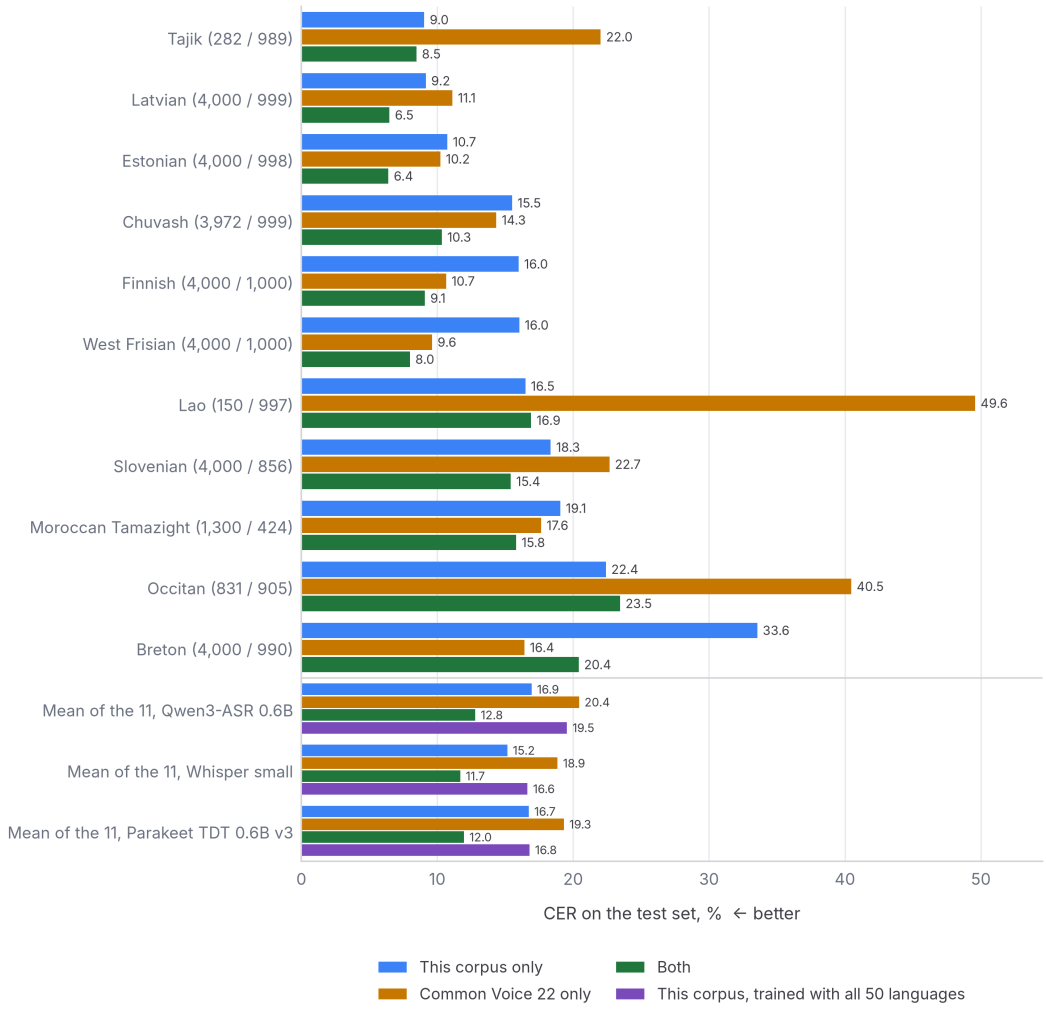

FIGURE 33 The languages Common Voice 22 has at least 100 clips of, trained on this study's corpus, on Common Voice, and on both. Test CER at the epoch dev chose: each language on Qwen3-ASR 0.6B, then the mean of the same three runs on each model, beside the model trained on all 50 languages with this corpus's clips and scored on the same 11 (purple). In brackets, the language's clips in Common Voice (at most 4,000 taken) and in this corpus (at most 1,000).

11 Another corpus

Every test clip so far comes from the corpus that also gave the training clips, so a model could learn the corpus rather than the language. FLEURS [Conneau et al., 2022] records 11 of our languages, but for 7 of them the corpus took clips from FLEURS itself: each of our clips that reads a FLEURS sentence is as long as a FLEURS recording of that sentence, to within 0.02 seconds, and 181 of 200 FLEURS test clips of Tajik read a sentence our training clips read. That leaves Finnish, Gujarati, Lingala and Shona, where FLEURS is another corpus: other texts, other speakers, other microphones. We score every model there on 200 FLEURS test clips a language, as downloaded and from the epoch dev chose in its main run, beside the same checkpoint on our own test clips of those languages (Figure 34).

Most of the gain is the language. On FLEURS the median model keeps 79% of what fine-tuning gained it on our clips, 33% at least (among the 16 that gained 5 points or more there), and FLEURS ranks the fine-tuned models as our test does (Spearman's ρ 0.94). The four

that end worse on FLEURS than they started, MMS 1B, omniASR CTC 1B, Gemma 4 E4B and Gemma 4 E2B, are the ones that knew these languages already, at 11.0 or better on our clips as downloaded: with little to gain there, fine-tuning on fifty languages trades some of it away.

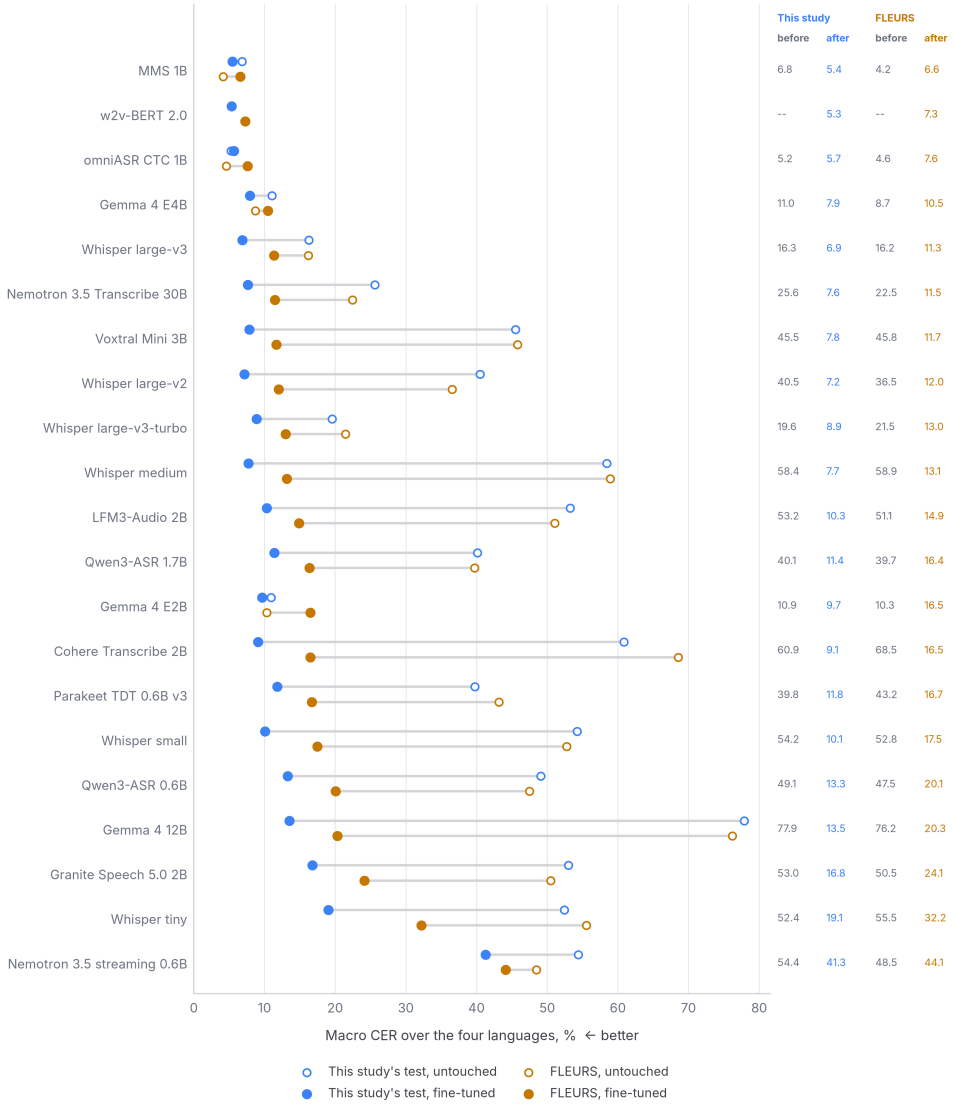

FIGURE 34 Finnish, Gujarati, Lingala and Shona, the languages FLEURS shares with this study whose clips here are not FLEURS recordings: as downloaded and fine-tuned, on this study's test clips and on FLEURS's. Each fine-tuned model is one checkpoint, scored on both; w2v-BERT has no head as downloaded.

11.1 THE LANGUAGES THEY KNEW

Learning fifty new languages costs a model some of the ones it knew (Figure 35). On English, Spanish, French and German from FLEURS, the median model's CER rises by 5.7 points, and by 46.4 for Nemotron 3.5 streaming 0.6B, which goes from 5.1 to 51.4. All 20 end worse, Nemotron 3.5 Transcribe 30B the least, by 0.4 (-0.1 to 1.0); its first run, which kept the experts as downloaded, even ended better, from 3.7 to 1.8. Adapters, which leave the pretrained weights as they were, lose less on 2 of the 3 models we tried them on (Section 12).

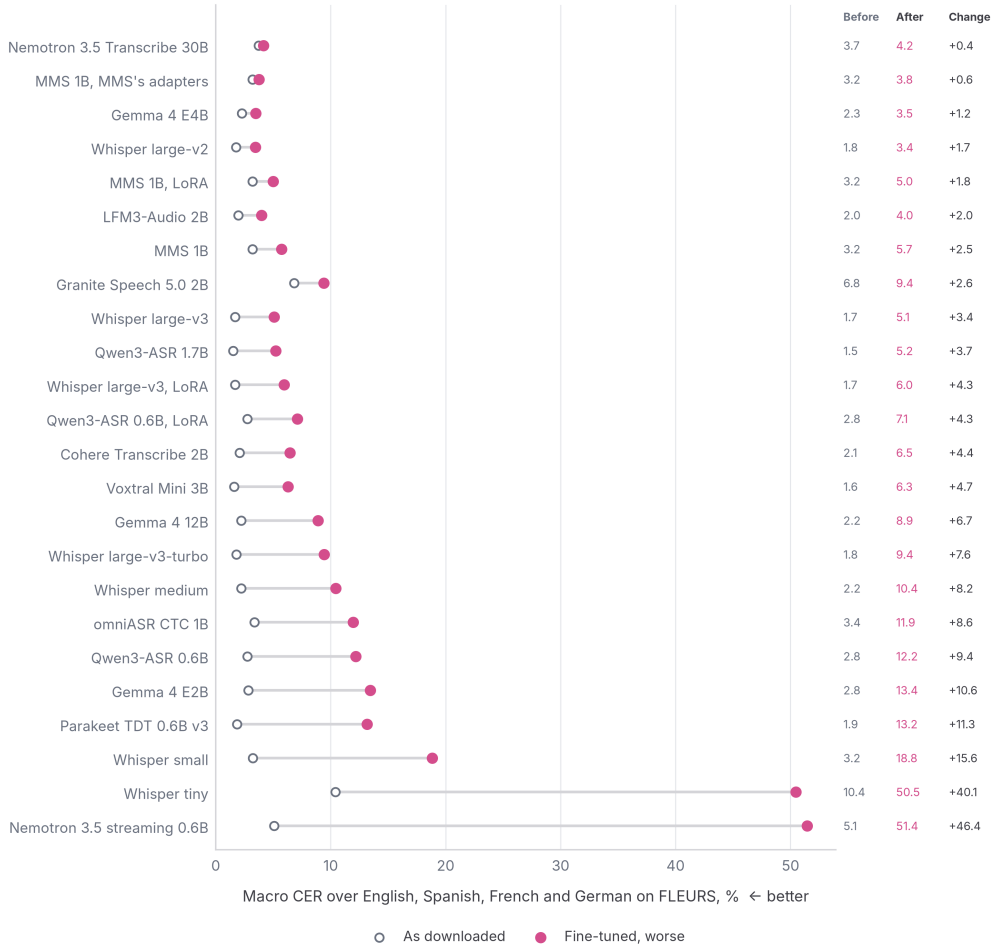

FIGURE 35 English, Spanish, French and German on FLEURS, as downloaded and after fine-tuning on the fifty languages, with the adapters of Section 12 where they ran.

11.2 SWITCHING LANGUAGES

Every clip so far is in one language, and speech often mixes two. CS-FLEURS [Yan et al., 2025] rewrites FLEURS sentences with English words and phrases mixed in and voices each language's stretch with MMS's TTS; five of its pairs are languages of ours with English. We score all 21 models on 200 clips of each, told the language that is not English, beside the same five languages spoken alone on FLEURS (Figure 36). The voices are synthetic and the switches are joins, so this asks whether a model follows a switch, not whether it hears one speaker make it, and the Latvian, Tajik and Burmese sentences come from FLEURS, which our corpus partly repackages (Section 11).

FIGURE 36 Five of the study's languages with English mixed in (CS-FLEURS) and the same languages alone (FLEURS), as downloaded and fine-tuned: macro CER over the five. English kept: of the English words in the Gujarati, Burmese and Tajik clips, where English has a script of its own, the share the output writes.

Fine-tuning helps with mixed speech too: 17 of the 20 models we can compare end lower, by a median of 22.7 points. It stays far harder than one language, though: the median fine-tuned model is at 36.2 on the mixed clips against 15.0 on the same languages alone. Much of the gap is the English. In Gujarati, Burmese and Tajik, where English is 22 to 27% of a reference's characters, the median model writes 12% of the English words as downloaded and 1% fine-tuned: told the other language, it writes the English in that language's script or drops it.

Malay with English is a switch the Malaysian arm trains on: in the Scicom telephony dataset of Section 13, 26% of the training words are English and 38% of the clips have two or more. On CS-FLEURS's Malay pair, against the same model as downloaded and told Malay, that training helps only 11 of the 19 models, by a median of 7.4 points. The public Malay benchmark has switching of its own: 171 of its 820 clips have two or more English words, 45

of them among its 66 telephony clips. Within each category those clips gain about as much from the training as the clips with none (median difference -0.7 points), so what the training adds there is the domain, not the switching.

12 Adapters and a language model

Two cheap additions to the main recipe: training a small part of the model instead of all of it, and decoding a CTC model with a word language model.

12.1 ADAPTERS

Tuning every weight means a full copy of the model for every set of new languages. Low-rank adapters of rank 32 on every attention and feed-forward projection [Hu et al., 2021] train instead, the rest frozen except what the new languages add (Whisper's new token rows, MMS's new CTC head), and for MMS also its own adapters, one small one in every layer, with a head [Pratap et al., 2023]. MMS keeps one such set for each language and swaps in the set of the language a user names. We train it both ways: one set for all 50 languages, as in every other arm, and one set for each language, trained on that language alone and swapped in for it at test, which needs the language named. Each arm has its own rate sweep and then trains ten epochs on the main data (Table 8). The sets for one language train longer, since a language alone gets a fiftieth of the updates: each sees 20,000 clips, so a language with fewer than 1,000 trains for more epochs, up to 210 for the scarcest.

TABLE 8 Tuning every weight against adapters on a frozen model, ten epochs on the same data: the parameters trained, the rate the model's own sweep chose, and test macro CER at the epoch dev chose, with a 95% interval from resampling speakers. MMS's adapters one set a language: one language's set, trained on 20,000 clips of that language and named at test.

MODEL	TRAINED	PARAMETERS TRAINED	RATE	TEST CER	95% CI
Whisper large-v3	everything	1,543.5M (100.0%)	1e-5	13.5	13.0–14.2
Whisper large-v3	LoRA, rank 32	57.7M (3.6%)	3e-4	14.6	14.0–15.4
Qwen3-ASR 0.6B	everything	782.4M (100.0%)	1e-5	18.2	17.5–18.9
Qwen3-ASR 0.6B	LoRA, rank 32	29.5M (3.6%)	1e-4	18.4	17.8–19.2
MMS 1B	everything	965.7M (100.0%)	1e-5	11.0	10.6–11.6
MMS 1B	LoRA, rank 32	36.4M (3.6%)	1e-4	11.7	11.3–12.4
MMS 1B	MMS's adapters, one set for all 50	3.2M (0.3%)	3e-4	14.0	13.6–14.6
MMS 1B	MMS's adapters, one set a language	2.2M (0.2%)	3e-4	16.8	16.4–17.3

LoRA comes close to tuning everything. On MMS, training 3.6% of the weights costs 0.7 points of CER (0.5 to 0.9), and MMS's adapters, one set for all 50 languages and 0.3% of the weights, cost 3.0. MMS's own way, one set for each language, does worse still, 2.8 points behind one set for all (2.4 to 3.1) although each is told its language, and 19.3 after the same ten epochs as the shared set. On the 26 languages MMS already had, the sets trained here reach 16.5, against 14.7 for its own; what MMS gains from a set of its own per language is the languages it knew, which keep theirs untouched. On Qwen3-ASR 0.6B LoRA is level with tuning everything, 0.2 points behind (-0.3 to 0.7), and on Whisper large-v3 it costs 1.1 (0.7 to 1.5). On two of the three, what they buy is the languages the model knew. On English, Spanish, French and German, Qwen3-ASR tuned in full loses 9.4 points of CER and with LoRA 4.3; MMS loses 2.5 tuned in full, 1.8 with LoRA and 0.6 with one set of its adapters for all 50. Whisper large-v3 loses more with LoRA, 4.3 against 3.4 tuned in full (Figure 35).

12.2 A LANGUAGE MODEL

A CTC model writes each frame without knowing which words exist. We searched the three CTC models' output with a word 4-gram model of each language [Heafield, 2011], built from the transcripts of the whole training pool, up to 4,000 clips a language where the models heard at most 1,000, less every sentence that dev or test also has (median 3,213 sentences and 30,102 words a language, the fewest 96). The beam is 64 wide [Kensho Technologies, 2021], and one weight for the language model and one bonus per word serve every language, chosen on dev (Table 9).

TABLE 9 The CTC models decoded greedily and with a word 4-gram, from the same saved output of the epoch dev chose. The interval of the WER gain resamples the same speakers on both sides; the weight and bonus were chosen on dev.

MODEL	CER, GREEDY	CER, LM	WER, GREEDY	WER, LM	WER GAIN, 95% CI	WEIGHT, BONUS
MMS 1B	11.1	**9.8**	40.1	**34.4**	5.7 (4.8 to 6.4)	0.25, 2
w2v-BERT 2.0	12.7	**11.3**	44.8	**38.8**	6.0 (5.0 to 6.7)	0.25, 2
omniASR CTC 1B	12.0	**10.8**	43.1	**38.3**	4.8 (4.1 to 5.4)	0.25, 1

It helps all three by about the same: 1.2 to 1.4 points of CER and 4.8 to 6.0 of WER, every interval clear of zero, and dev chose the same weight for each. With it MMS 1B reaches 9.8, the lowest macro CER on our test set, yet its WER is still 34.4: in the median language 20.7% of the test words never occur in the text the language model was built from, and up to 68.8% in one.

13 Malaysian

Malay is not a language with little to train on, but the question of what a model brings to it is the same, and one thing is different: the training data and the test data do not come from the same kind of speech. We train on the Scicom telephony dataset, 31,622 private clips and 60 hours of human-transcribed call-centre speech in Malay and English with some Mandarin, cut at the speakers' pauses, and test on the public split of the Revolab ASR Benchmark [Revolab Research, 2026]: 820 clips and 2.1 hours of Malaysian Malay and English in 12 categories, from telephone calls to singing, each counting once as each language did. Dev is 454 held-out clips of the telephony dataset and alone chooses the epoch; each model trains at its main run's rate, w2v-BERT excepted (below). As the benchmark's card asks, listed word variants count as correct, and a number may be spelled out or written in digits.

A model with Malay on its list can be run told the language or left to detect it; the two differ enough that Figure 37 and Table 10 have both, and a gain is counted from the better. Trained on telephone calls alone, Whisper large-v3 goes from 8.9 to 6.0, a gain of 2.9 points (1.6 to 4.4) and the lowest score of any model. Of the 24 models, 18 gain half a point or more and Gemma 4 E2B and Qwen3-ASR 0.6B end worse by as much; a model that starts well has little left for these calls to give. w2v-BERT's new CTC head stalls at its main rate, writing the same thing for every clip for 7 and 6 epochs on two seeds, while at a third of the rate it is out by epoch 3 and ends ninth best, at 9.3. MMS's new head beats the Malay head it came with, 12.2 against 15.3.

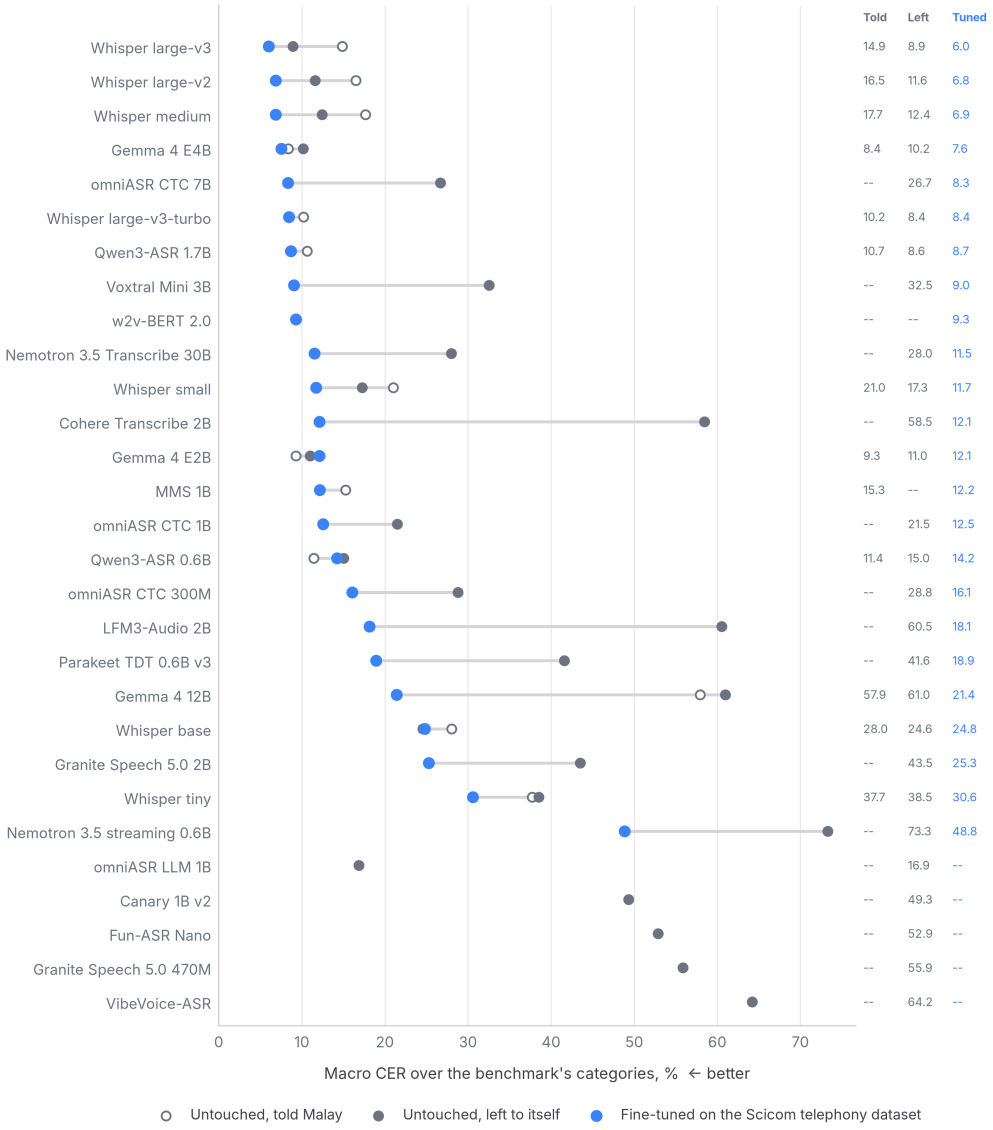

FIGURE 37 Every model on the public Malay benchmark: as downloaded, told Malay where its list has it and left to itself, and after training on the Scicom telephony dataset. Macro over the benchmark's categories. Dashes were not run.

Most of the gap between the two untouched columns is translation (Figure 38): told that a clip is Malay, an untouched Whisper translates the English in it. Of the 91 clips whose reference is mostly English by a count of common words, large-v3 told Malay answers 57 in Malay, and they hold 46% of its errors; left to detect the language it does so on 9, and fine-tuned, still told Malay, on none. Counted from the model told Malay, its gain would read 8.9 points instead of 2.9. Detection is not free either: Qwen3-ASR 0.6B is better told, 11.4 against 15.0.

TABLE 10 The Malaysian arm, best first. The test set is the Revolab benchmark, macro over its categories, capped per clip. MMS untouched is told Malay through its Malay adapter and head. The last columns, WER and the benchmark's telephony category, go from the better untouched score to the fine-tuned one.

MODEL	TOLD MALAY	LEFT TO ITSELF	FINE-TUNED	95% CI	WER	BEST EPOCH	TELEPHONY
Whisper large-v3	14.9	8.9	**6.0**	5.0–7.2	14.9 to 13.3	2	17.8 to 11.7
Whisper large-v2	16.5	11.6	**6.8**	6.2–7.6	18.7 to 16.2	4	26.6 to 12.8
Whisper medium	17.7	12.4	**6.9**	6.0–8.0	22.0 to 16.1	2	30.4 to 12.1
Gemma 4 E4B	8.4	10.2	**7.6**	6.8–8.4	17.1 to 16.2	2	20.9 to 13.8
omniASR CTC 7B	–	26.7	**8.3**	7.6–9.1	43.1 to 24.8	3	48.5 to 15.5
Whisper large-v3-turbo	10.2	8.4	**8.4**	7.5–9.4	14.9 to 19.4	3	17.3 to 14.1
Qwen3-ASR 1.7B	10.7	8.6	**8.7**	7.9–9.5	17.6 to 20.7	1	7.3 to 11.5
Voxtral Mini 3B	–	32.5	**9.0**	8.0–10.3	44.4 to 19.6	2	25.0 to 11.2
w2v-BERT 2.0	–	–	**9.3**	8.4–10.3	27.9	12	–
Nemotron 3.5 Transcribe 30B	–	28.0	**11.5**	10.6–12.5	48.7 to 26.7	2	17.1 to 12.8
Whisper small	21.0	17.3	**11.7**	10.6–13.1	30.8 to 27.1	3	28.0 to 15.4
Cohere Transcribe 2B	–	58.5	**12.1**	11.1–13.1	82.2 to 29.6	3	23.1 to 12.7
Gemma 4 E2B	9.3	11.0	**12.1**	11.1–13.3	20.0 to 24.3	2	21.4 to 16.5
MMS 1B	15.3	–	**12.2**	11.3–13.1	40.2 to 32.3	10	36.7 to 20.1
omniASR CTC 1B	–	21.5	**12.5**	11.5–13.6	39.9 to 34.2	8	36.5 to 19.9
Qwen3-ASR 0.6B	11.4	15.0	**14.2**	13.2–15.4	25.2 to 31.4	4	13.1 to 17.0
omniASR CTC 300M	–	28.8	**16.1**	15.0–17.2	53.5 to 43.6	10	46.6 to 24.2
LFM3-Audio 2B	–	60.5	**18.1**	16.9–19.4	82.0 to 37.9	3	24.6 to 16.6
Parakeet TDT 0.6B v3	–	41.6	**18.9**	17.9–20.1	77.8 to 41.2	6	15.1 to 16.3
Gemma 4 12B	57.9	61.0	**21.4**	19.9–23.0	73.6 to 39.5	2	40.1 to 19.3
Whisper base	28.0	24.6	**24.8**	23.5–26.1	47.8 to 48.7	4	33.4 to 23.5
Granite Speech 5.0 2B	–	43.5	**25.3**	24.0–26.7	79.6 to 49.2	4	21.5 to 15.4
Whisper tiny	37.7	38.5	**30.6**	29.2–32.0	65.8 to 56.8	4	69.5 to 26.1
Nemotron 3.5 streaming 0.6B	–	73.3	**48.8**	47.5–50.1	88.4 to 77.3	10	33.5 to 42.9
omniASR LLM 1B	–	16.9	–	–	–	–	–
Canary 1B v2	–	49.3	–	–	–	–	–
Fun-ASR Nano	–	52.9	–	–	–	–	–
Granite Speech 5.0 470M	–	55.9	–	–	–	–	–
VibeVoice-ASR	–	64.2	–	–	–	–	–

Beyond translation, the calls help large-v3 most on short inputs, parliament, telephony and street interviews, and cost it a little on singing, news and FLEURS (Figure 38; every model is in Figure 39). They do not even win the benchmark's own calls: the best fine-tuned telephony score is 11.2 (Voxtral Mini 3B), and Qwen3-ASR 1.7B scores 7.3 there untouched. These are not the same calls, and a model tuned on one call centre writes what that call centre's transcribers write.

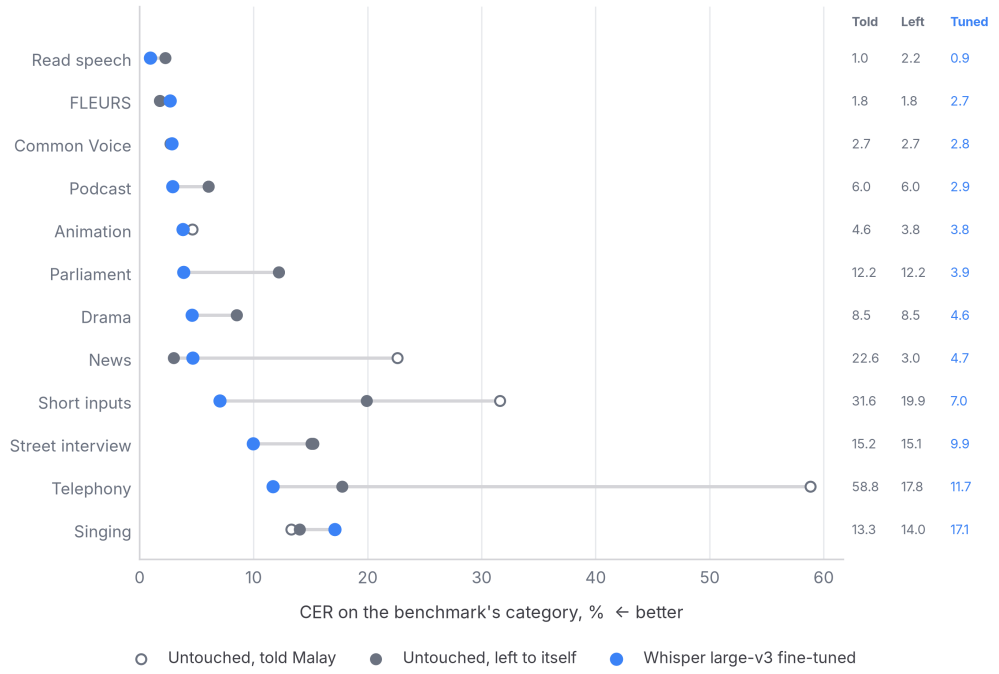

FIGURE 38 The best fine-tuned model on each of the benchmark's categories: as downloaded both ways, and after training on telephone speech alone.

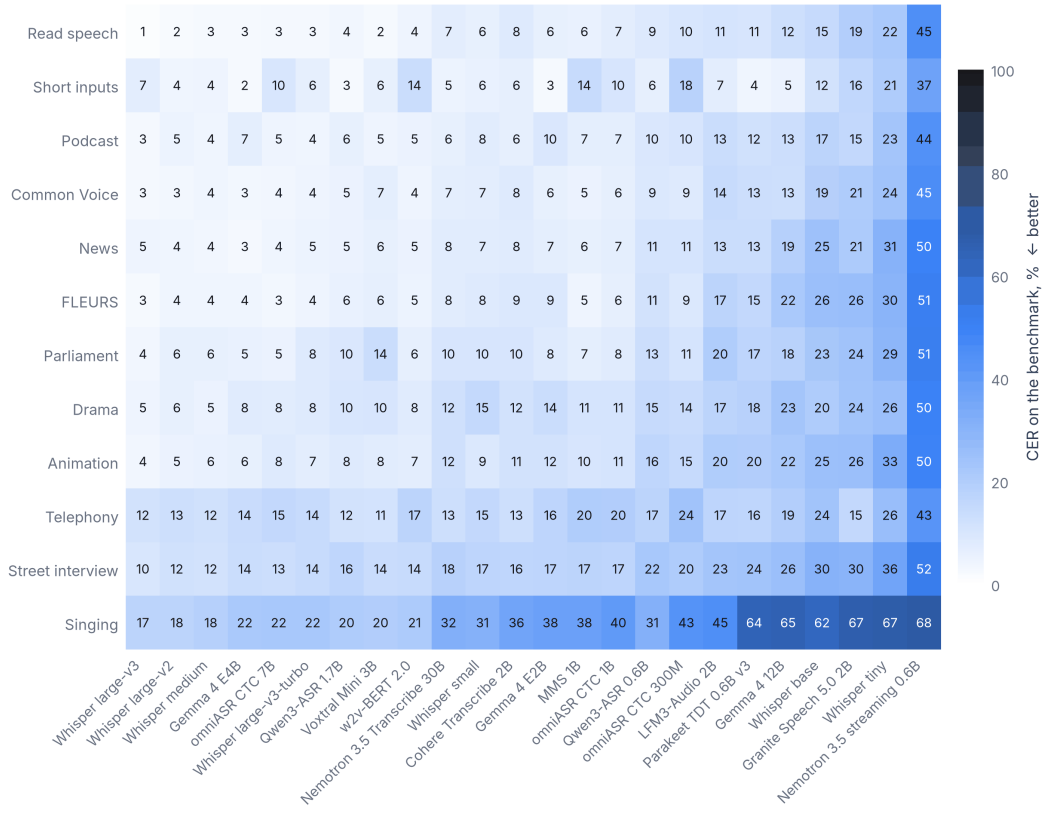

FIGURE 39 Every fine-tuned model on every category of the benchmark.

More epochs can hurt here, and dev does not warn of it (Figure 40): 16 of the 24 models

reach their lowest benchmark score within three epochs, and Whisper tiny climbs from 28.4 at epoch 2 to 36.6 by epoch 10 while its telephony dev improves from 39.7 to 30.9. It is becoming a better recogniser of these calls and a worse one of everything else.

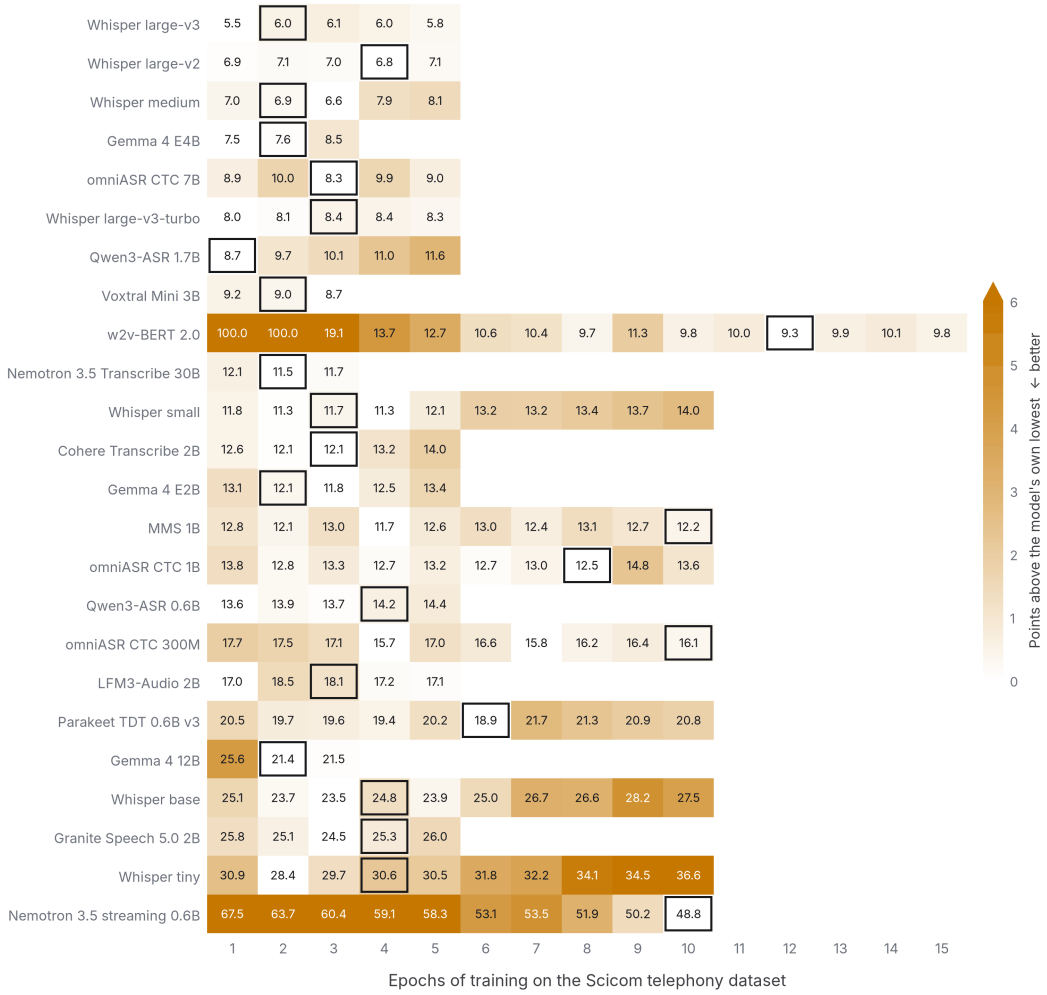

FIGURE 40 The benchmark after every epoch of training on the Scicom telephony dataset, one row a model. The colour is how far a cell is above that model's own lowest. A box marks the epoch telephony dev chose. Blank cells were not run.

Nor does more telephony help the benchmark (Figure 41): Whisper small ends at 11.0 after a thousand clips, 2.0 hours, and at 11.7 after all 60.1, while telephony dev improves the whole way. Whatever the benchmark gains from telephone calls, a small sample of them already has it.

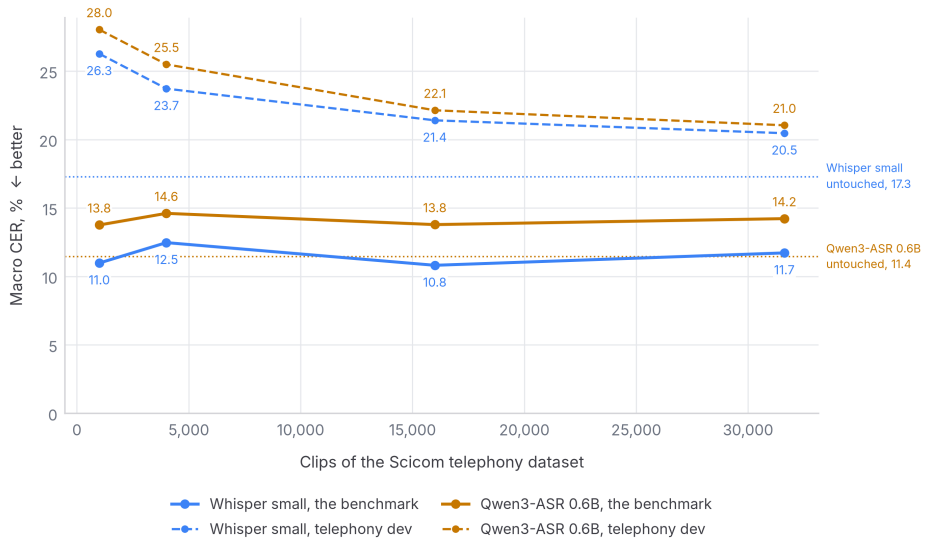

FIGURE 41 More of the Scicom telephony dataset, at the epoch telephony dev chose. Solid: the public benchmark. Dashed: telephony dev. Dotted: the benchmark, the model's better untouched score.

14 Limitations

The test set is 25 clips a language: a language's own CER is a rough number, and the macro mean is what we trust. Its intervals resample speakers and take the 50 languages as given; with the languages drawn again, 27 of the main result's 210 pairs of models could change places (Section 6.4). What a model had heard before is known only from its card or its table, and a language on no list may still have been in a model's training audio. Its labels come from a language detector applied to transcripts, verified by confidence, so a language here means what GlotLID says [Kargaran et al., 2023]; our Teso and Karamojong share a corpus labelled Turkana. For 41 languages every test clip comes from one corpus, and in all 50 a test corpus also supplies training clips, from other speakers where it can: the score says how well a model learns a language as that corpus records it, not how it will do on other speech in that language. For 7 languages that corpus is partly or wholly FLEURS, and FLEURS can test other speech in only 4 of the languages it shares with us (Section 11). Crimean Tatar, Goan Konkani, Gurani, Meitei and Wayuu have no other speakers, so their test clips share speakers with training, and every model scores them at 54 to 85% of its CER on the rest (Figure 42). Of the test clips, 2 pair a long text with very short audio. Some corpora are readings of the Bible, as Goan Konkani's, Malagasy's and Wayuu's are, and such readings repeat their vocabulary.

The test set's rows were verified one by one, by language and by length; training and dev rows carry the parent corpus's labels, gated on length alone. Dev is harder than test for every model, by 5.5 to 8.3 points (Figure 43). Of dev's clips, 72% come from speakers with at most five clips in the corpus, against 11% of test's, and in the median model 5.0% fail outright, with three quarters of their characters wrong or worse, against 1.0% of test's. So a dev score and a test score here are not comparable, and the test score flatters how these models would do on a language's rarer speakers.

Decoding stops at 40 tokens plus 30 for each second of audio, and Whisper's byte-level vocabulary needs more than that on 17 of the 25 Tibetan test clips and 12 of the Lao ones (Figure 44); decoding the chosen epoch again at 60 moves the macro CER by at most 0.3 points on the 7 models we tried it on.

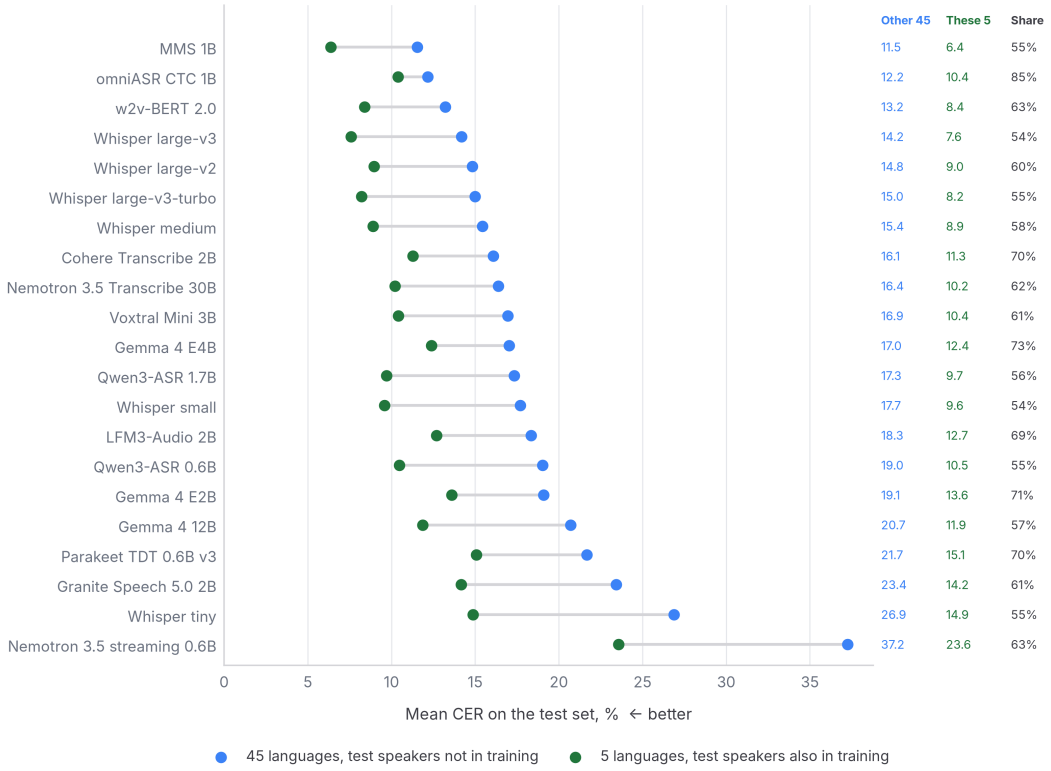

FIGURE 42 Test CER on the 5 languages whose test speakers also appear in training, since they have no others, against the rest, each model at the epoch dev chose. Share: the first as a percentage of the second.

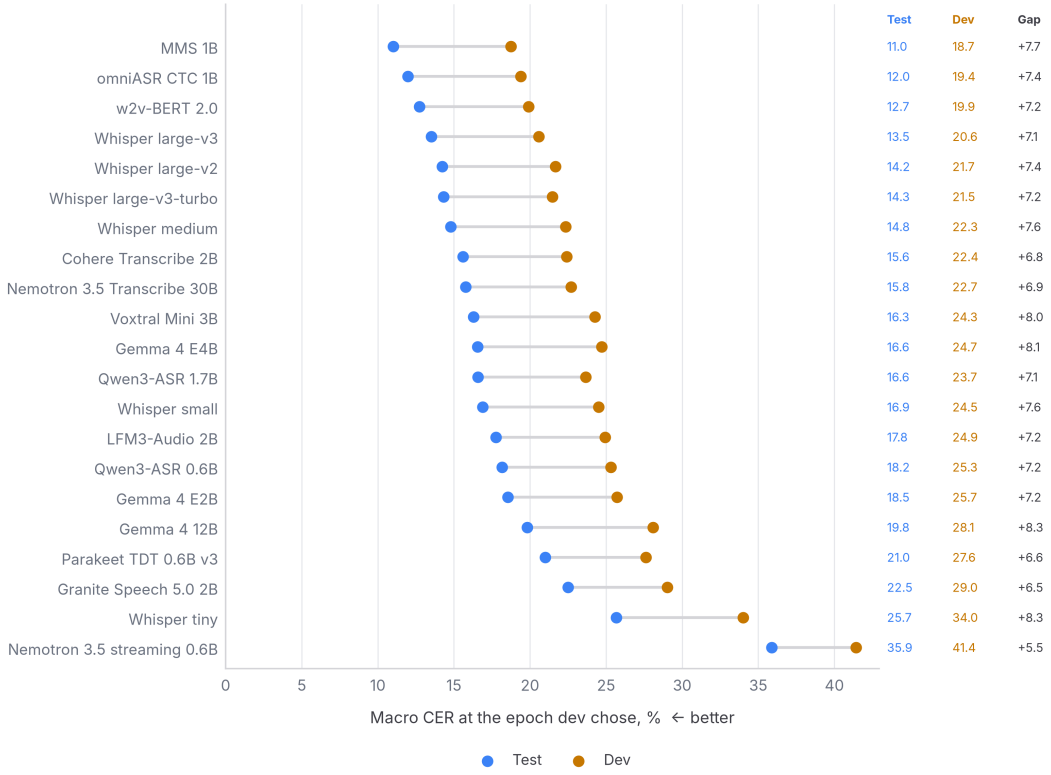

FIGURE 43 Dev and test macro CER of each model at the epoch dev chose.

We tuned the learning rate per model but no other setting, so a model with a better recipe may rank higher. Granite Speech, MMS and Voxtral were swept late, and the sweep found the rate MMS and Voxtral first trained at too fast (Appendix C). Run again at 1e-5, MMS went from 16.2 to 11.0, the lowest of any model, and Voxtral at 3e-6 from 17.2 to 16.3. Outside Section 12 every score is greedy decoding with no language model. In the Malaysian arm w2v-BERT's new head stalls at its main rate, and we did not find out why; its row there is at a third of that rate, the only row of that arm not at the model's main rate.

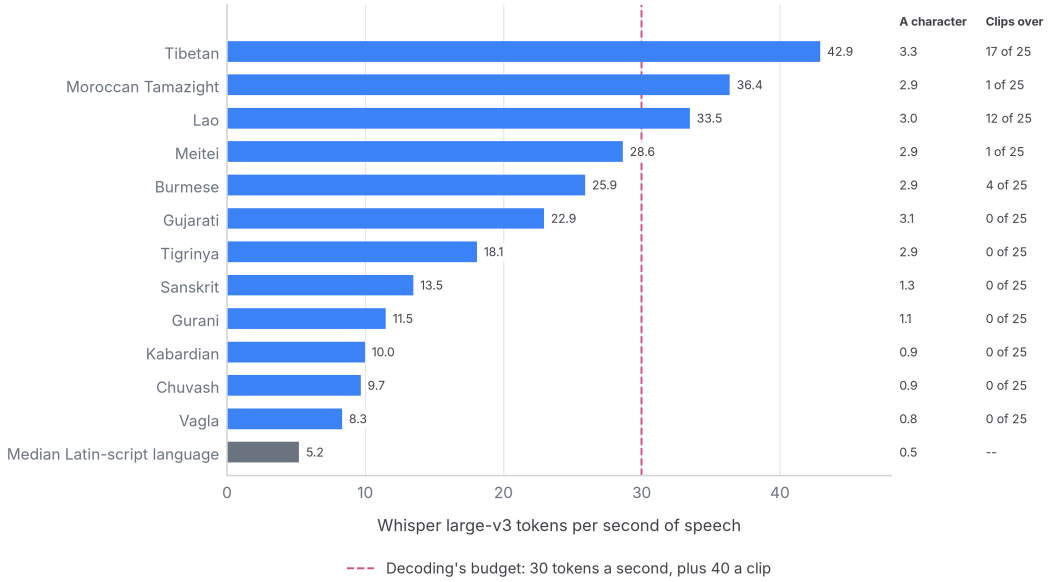

FIGURE 44 Tokens Whisper large-v3 needs per second of speech for each language's test transcripts, as it trains on them, for the twelve languages that need the most and the median Latin-script language. The columns give tokens per character and the test clips whose transcript needs more tokens than decoding allows.

15 Conclusion

A multilingual recogniser fine-tuned on an average of 98 minutes of a language it has never heard learns to write it. On 50 low-resource languages every model gained, by 16.6 to 53.1 points of macro CER, and the same architectures from random weights barely got started. On the 12 that no model lists, every model starts at 46 or worse and the best ends at 11.7. What transfers is hearing speech in general: an encoder that never heard these languages already gives a Whisper or Qwen decoder most of what it needs, though not a transducer's, and the script, the thing an untouched model most often gets wrong, is the first thing training puts right. Most of it is the language, not the corpus: on FLEURS, in the languages it shares with ours without sharing recordings, the median model keeps 79% of its gain.

With a little of a new language, start from a model with a strong encoder and an output that can spell anything, as large as you can train; a CTC model or a Whisper did best here, and omniASR CTC 7B best of all, at 9.5. On a seq2seq or language-model decoder, stop after three to five epochs; CTC models and transducers want ten or more. New clips are worth more than further epochs at the same number presented, so spend the time collecting them; where Common Voice has the language, adding it took 4.1 points off. If one model for each language is no burden, train the language alone. On a transducer, check what the new rows write before trusting the loss. On a mixture of experts, train the audio encoder with the rest;

the routed experts, most of the weights, were worth under a point, less than letting the rate fall to zero (Section 6.8). Before crediting the training, score the untouched model both told the language and left to detect it: told that a clip was Malay, an untouched Whisper translated the English that the speakers mixed in, and lost more to that than training gained. A CTC model should be decoded with a word language model; a few thousand sentences a language took 4.8 to 6.0 points off its WER. Expect to lose some of the languages the model knew, 5.7 points of CER in the median model, and check them before shipping it. Expect a switch into English to go badly: trained on one language a clip, the median model wrote 1% of the English words in speech that mixed them into a language with a script of its own. Do not expect a model that has learned forty other low-resource languages to learn the next one faster. Where a full copy of the model for every set of languages is a burden, LoRA on a few percent of the weights does nearly as well, and on 2 of our 3 models it also keeps more of the languages the model knew. Read the output as a draft: even with a language model, the lowest WER here is 34.4.

We tried one recipe per model, a language model and adapters on three models each, and one seed for most runs, and FLEURS could check our corpus in only 4 languages. The languages we could test are the ones a public corpus happened to record, as it recorded them. What the paper does show is that the first hour or two of a new language is worth far more on top of a multilingual model than on its own, and that it is enough to start.

16 Acknowledgement

This work was carried out at Scicom (MSC) Berhad. The author thanks Scicom for the compute.

References

Abdulhady Abas Abdullah, Shima Tabibian, Hadi Veisi, Aso Mahmudi, and Tarik Rashid. End-to-End Transformer-based Automatic Speech Recognition for Northern Kurdish: A Pioneering Approach. arXiv:2410.16330, 2024. URL https://arxiv.org/abs/2410.16330.

Oliver Adams, Matthew Wiesner, Shinji Watanabe, and David Yarowsky. Massively Multilingual Adversarial Speech Recognition. arXiv:1904.02210, 2019. URL https://arxiv.org/abs/1904.02210.

Benjamin Akera, Evelyn Nafula, Patrick Walukagga, Gilbert Yiga, John Quinn, and Ernest Mwebaze. How much speech data is necessary for ASR in African languages? An evaluation of data scaling in Kinyarwanda and Kikuyu. arXiv:2510.07221, 2025. URL https://arxiv.org/abs/2510.07221.

Rosana Ardila, Megan Branson, Kelly Davis, Michael Henretty, Michael Kohler, Josh Meyer, Reuben Morais, Lindsay Saunders, et al. Common Voice: A Massively-Multilingual Speech Corpus. arXiv:1912.06670, 2019. URL https://arxiv.org/abs/1912.06670.

Arun Babu, Changhan Wang, Andros Tjandra, Kushal Lakhotia, Qiantong Xu, Naman Goyal, Kritika Singh, Patrick von Platen, et al. XLS-R: Self-supervised Cross-lingual Speech Representation Learning at Scale. arXiv:2111.09296, 2021. URL https://arxiv.org/abs/2111.09296.

Martijn Bartelds, Nay San, Bradley McDonnell, Dan Jurafsky, and Martijn Wieling. Making More of Little Data: Improving Low-Resource Automatic Speech Recognition Using Data Augmentation. arXiv:2305.10951, 2023. URL https://arxiv.org/abs/2305.10951.

Zhehuai Chen, He Huang, Andrei Andrusenko, Oleksii Hrinchuk, Krishna C. Puvvada, Jason Li, Subhankar Ghosh, Jagadeesh Balam, et al. SALM: Speech-augmented Language Model with In-context Learning for Speech Recognition and Translation. arXiv:2310.09424, 2023. URL https://arxiv.org/abs/2310.09424.

Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, et al. FLEURS: Few-shot Learning Evaluation of Universal Representations of Speech. arXiv:2205.12446, 2022. URL https://arxiv.org/abs/2205.12446.

Siddharth Dalmia, Ramon Sanabria, Florian Metze, and Alan W. Black. Sequence-based Multi-lingual Low Resource Speech Recognition. arXiv:1802.07420, 2018. URL https://arxiv.org/abs/1802.07420.

Gemma Team. Gemma 4 Technical Report. arXiv:2607.02770, as cited on the model card, 2026. URL https://huggingface.co/google/gemma-4-E4B-it.

Dawit Ketema Gete, Bedru Yimam Ahmed, Tadesse Destaw Belay, Yohannes Ayana Ejigu, Sukairaj Hafiz Imam, Alemu Belay Tessema, Mohammed Oumer Adem, Tadesse Amare Belay, et al. Whispering in Amharic: Fine-tuning Whisper for Low-resource Language. arXiv:2503.18485, 2025. URL https://arxiv.org/abs/2503.18485.

Kenneth Heafield. KenLM: Faster and Smaller Language Model Queries. Proceedings of the Sixth Workshop on Statistical Machine Translation, 2011. URL https://aclanthology.org/W11-2123/.

Jui-Yang Hsu, Yuan-Jui Chen, and Hung yi Lee. Meta Learning for End-to-End Low-Resource Speech Recognition. arXiv:1910.12094, 2019. URL https://arxiv.org/abs/1910.12094.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. arXiv:2106.09685, 2021. URL https://arxiv.org/abs/2106.09685.

IBM. Granite Speech 5.0. Hugging Face model card, 2026. URL https://huggingface.co/ibm-granite/granite-speech-5.0-2b.

Amir Hossein Kargaran, Ayyoob Imani, François Yvon, and Hinrich Schütze. GlotLID: Language Identification for Low-Resource Languages. arXiv:2310.16248, 2023. URL https://arxiv.org/abs/2310.16248.

Kensho Technologies. pyctcdecode. GitHub repository, 2021. URL https://github.com/kensho-technologies/pyctcdecode.

James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, et al. Overcoming catastrophic forgetting in neural networks. arXiv:1612.00796, 2016. URL https://arxiv.org/abs/1612.00796.

Julius Kunze, Louis Kirsch, Ilia Kurenkov, Andreas Krug, Jens Johannsmeier, and Sebastian Stober. Transfer Learning for Speech Recognition on a Budget. arXiv:1706.00290, 2017. URL https://arxiv.org/abs/1706.00290.

Liquid AI. Introducing LFM3: Every Modality for Specialized Intelligence. Liquid AI Blog, as cited on the model card of LFM3-Audio-2B, 2026. URL https://huggingface.co/LiquidAI/LFM3-Audio-2B.

Alexander H. Liu, Andy Ehrenberg, Andy Lo, Clément Denoix, Corentin Barreau, Guillaume Lample, Jean-Malo Delignon, Khyathi Raghavi Chandu, et al. Voxtral. arXiv:2507.13264, 2025. URL https://arxiv.org/abs/2507.13264.

Julian Mack, Ekagra Ranjan, Walter Beller-Morales, Bharat Venkitesh, and Pierre Richemond. Cohere Transcribe (cohere-transcribe-03-2026). Hugging Face model card, doi:10.57967/hf/8653, 2026. URL https://huggingface.co/CohereLabs/cohere-transcribe-03-2026.

Malaysia-AI. Low-Language-TTS. Hugging Face dataset card, 2026a. URL https://huggingface.co/datasets/malaysia-ai/Low-Language-TTS.

Malaysia-AI. Multilingual-TTS. Hugging Face dataset card, 2026b. URL https://huggingface.co/datasets/malaysia-ai/Multilingual-TTS.

Mozilla. Common Voice 22.0, as mirrored in fsicoli/common_voice_22_0. Hugging Face dataset, 2025. URL https://huggingface.co/datasets/fsicoli/common_voice_22_0.

Alvin Nahabwe, Sulaiman Kagumire, Denis Musinguzi, Bruno Beijuka, Jonah Mubuuke Kyagaba, Peter Nabende, Andrew Katumba, and Joyce Nakatumba-Nabende. Benchmarking Automatic Speech Recognition Models for African Languages. arXiv:2512.10968, 2025. URL https://arxiv.org/abs/2512.10968.

NVIDIA. Canary 1B v2. Hugging Face model card, 2025a. URL https://huggingface.co/nvidia/canary-1b-v2.

NVIDIA. Parakeet TDT 0.6B v3. Hugging Face model card, 2025b. URL https://huggingface.co/nvidia/parakeet-tdt-0.6b-v3.

NVIDIA. Fine-tuning Nemotron 3.5 ASR. Hugging Face blog, 2026a. URL https://huggingface.co/blog/nvidia/fine-tuning-nemotron-35-asr.

NVIDIA. Nemotron 3.5 ASR Streaming 0.6B. Hugging Face model card, 2026b. URL https://huggingface.co/nvidia/nemotron-3.5-asr-streaming-0.6b.

NVIDIA. Nemotron 3.5 Transcribe 30B A4B. Hugging Face model card, 2026c. URL https://huggingface.co/nvidia/NVIDIA-Nemotron-3.5-Transcribe-30B-A4B.

Omnilingual ASR team, Gil Keren, Artyom Kozhevnikov, Yen Meng, Christophe Ropers, Matthew Setzler, Skyler Wang, Ife Adebara, et al. Omnilingual ASR: Open-Source Multilingual Speech Recognition for 1600+ Languages. arXiv:2511.09690, 2025. URL https://arxiv.org/abs/2511.09690.

Leena G Pillai, Kavya Manohar, Basil K Raju, and Elizabeth Sherly. Multistage Fine-tuning Strategies for Automatic Speech Recognition in Low-resource Languages. arXiv:2411.04573, 2024. URL https://arxiv.org/abs/2411.04573.

Vineel Pratap, Andros Tjandra, Bowen Shi, Paden Tomasello, Arun Babu, Sayani Kundu, Ali Elkahky, Zhaoheng Ni, et al. Scaling Speech Technology to 1,000+ Languages. arXiv:2305.13516, 2023. URL https://arxiv.org/abs/2305.13516.

Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust Speech Recognition via Large-Scale Weak Supervision. arXiv:2212.04356, 2022. URL https://arxiv.org/abs/2212.04356.

Dima Rekesh, Nithin Rao Koluguri, Samuel Kriman, Somshubra Majumdar, Vahid Noroozi, He Huang, Oleksii Hrinchuk, Krishna Puvvada, et al. Fast Conformer with Linearly Scalable Attention for Efficient Speech Recognition. arXiv:2305.05084, 2023. URL https://arxiv.org/abs/2305.05084.

Revolab Research. Revolab ASR Benchmark. Hugging Face dataset card, 2026. URL https://huggingface.co/datasets/Revolab/ASR-Benchmark-Public.

George Saon, Avihu Dekel, Alexander Brooks, Tohru Nagano, Abraham Daniels, Aharon Satt, Ashish Mittal, Brian Kingsbury, et al. Granite-speech: open-source speech-aware LLMs with strong English ASR capabilities. arXiv:2505.08699, 2025. URL https://arxiv.org/abs/2505.08699.

Scicom AI Enterprise Organization. Fused linear cross-entropy in fp32 (small-ablation, cross-entropy). GitHub repository, 2026. URL https://github.com/Scicom-AI-Enterprise-Organization/small-ablation/tree/main/cross-entropy.

Xian Shi, Xiong Wang, Zhifang Guo, Yongqi Wang, Pei Zhang, Xinyu Zhang, Zishan Guo, Hongkun Hao, et al. Qwen3-ASR Technical Report. arXiv:2601.21337, 2026. URL https://arxiv.org/abs/2601.21337.

Steven Vander Eeckt and Hugo Van hamme. Using Adapters to Overcome Catastrophic Forgetting in End-to-End Automatic Speech Recognition. arXiv:2203.16082, 2022. URL https://arxiv.org/abs/2203.16082.

Erik Wijmans, Brody Huval, Alexander Hertzberg, Vladlen Koltun, and Philipp Krähenbühl. Cut Your Losses in Large-Vocabulary Language Models. arXiv:2411.09009, 2024. URL https://arxiv.org/abs/2411.09009.

Hainan Xu, Fei Jia, Somshubra Majumdar, He Huang, Shinji Watanabe, and Boris Ginsburg. Efficient Sequence Transduction by Jointly Predicting Tokens and Durations. arXiv:2304.06795, 2023. URL https://arxiv.org/abs/2304.06795.

Brian Yan, Injy Hamed, Shuichiro Shimizu, Vasista Lodagala, William Chen, Olga Iakovenko, Bashar Talafha, Amir Hussein, et al. CS-FLEURS: A Massively Multilingual and Code-Switched Speech Dataset. arXiv:2509.14161, 2025. URL https://arxiv.org/abs/2509.14161.

Hao Yang, Jinming Zhao, Gholamreza Haffari, and Ehsan Shareghi. Investigating Pre-trained Audio

Encoders in the Low-Resource Condition. arXiv:2305.17733, 2023. URL https://arxiv.org/abs/2305.17733.

A Untouched scores

Table 11 has every model as it was downloaded, and Figure 45 what the cap takes away: the part of each model's error that comes from clips writing more than their own length, mostly decoders stuck in a loop.

TABLE 11 Every model as downloaded, on the test set. CER is capped per clip; the uncapped column counts every character of a looping decoder. Loops: clips that end by repeating a piece of themselves four times or more.

MODEL	DECODER	LANGUAGES	CER	95% CI	CER, UNCAPPED	WER	LOOPS
MMS 1B, a head per language	CTC	26	14.7	14.0–15.7	14.9	42.7	0
omniASR LLM 1B	language model	50	27.8	26.9–28.8	28.3	56.4	2
omniASR CTC 1B	CTC	50	28.5	27.8–29.6	28.8	63.1	0
omniASR CTC 7B	CTC	50	29.7	28.9–30.7	29.9	59.2	0
MMS 1B, English's head where it has none	CTC	50	30.6	30.1–31.2	31.1	62.9	2
omniASR CTC 300M	CTC	50	31.5	30.8–32.4	31.7	69.5	1
Gemma 4 E4B	language model	50	39.8	39.1–40.7	41.7	77.0	5
Gemma 4 E2B	language model	50	43.5	42.7–44.4	43.7	80.2	1
Whisper large-v3	seq2seq	50	50.3	49.5–51.4	59.5	83.7	101
Whisper large-v3-turbo	seq2seq	50	54.6	53.6–55.8	89.7	84.8	154
Parakeet TDT 0.6B v3	transducer	50	55.7	55.0–56.7	56.4	88.0	0
Whisper large-v2	seq2seq	50	57.8	56.7–58.9	72.3	86.8	94
Nemotron 3.5 Transcribe 30B	language model	50	58.1	57.2–59.3	59.4	86.3	3
Qwen3-ASR 1.7B	language model	50	59.3	58.5–60.2	59.8	92.9	0
Qwen3-ASR 0.6B	language model	50	60.9	60.1–61.9	61.5	94.2	1
Voxtral Mini 3B	language model	50	61.8	60.8–62.9	67.2	91.2	16
Granite Speech 5.0 2B	language model	50	62.1	61.4–63.0	68.3	95.9	7
Canary 1B v2	seq2seq	50	62.4	61.5–63.4	84.3	90.4	52
Whisper base	seq2seq	50	62.7	61.8–63.8	94.3	94.5	89
Whisper medium	seq2seq	50	63.7	62.7–64.8	80.0	89.3	170
LFM3-Audio 2B	language model	50	64.3	63.5–65.2	78.1	89.5	15
Whisper small	seq2seq	50	65.7	64.7–66.7	82.7	91.5	303
Whisper tiny	seq2seq	50	66.6	65.6–67.6	102.0	96.1	130
Cohere Transcribe 2B	seq2seq	50	68.0	67.0–69.1	93.0	95.8	40
Nemotron 3.5 streaming 0.6B	transducer	50	68.7	68.0–69.5	69.7	92.2	8
Granite Speech 5.0 470M	CTC	50	69.7	68.8–70.5	69.8	95.8	0
Gemma 4 12B	language model	50	72.9	71.9–73.9	97.7	96.4	67
Fun-ASR Nano	language model	50	73.5	72.7–74.6	81.4	97.7	26
VibeVoice-ASR	language model	50	84.9	84.1–85.9	135.2	98.3	146

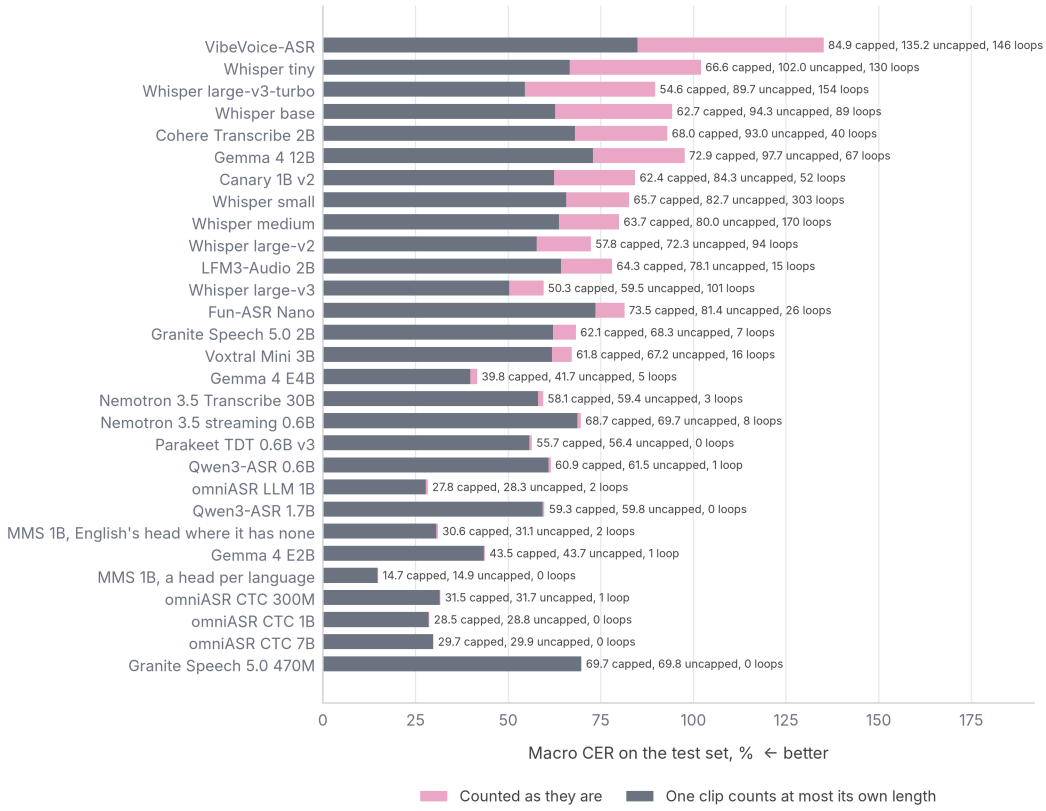

FIGURE 45 Every model as downloaded, with and without the cap on a clip's errors. The pink part is what clips add beyond their own length, which is mostly decoders stuck in a loop.

B The 50 languages

Table 12 lists the study's languages, ordered by how little the parent corpus holds.

TABLE 12 The study's languages. Corpus rows are from the Low-Language-TTS card. Training clips are those left after every rule of Section 3, at most 4,000 a language. Disjoint: no test speaker in training.

CODE	LANGUAGE	FAMILY	CORPUS ROWS	TRAIN CLIPS	TRAIN H	TRAIN SPEAKERS	TEST SPEAKERS	DISJOINT
gom_Latn	Goan Konkani	Indo-European	262	212	0.6	1	1	no
njm_Latn	Angami Naga	Sino-Tibetan	443	96	0.1	30	12	yes
guc_Latn	Wayuu	Arawakan	736	649	0.7	48	10	no
nre_Latn	Southern Rengma Naga	Sino-Tibetan	992	310	0.4	70	13	yes
mni_Mtei	Meitei	Sino-Tibetan	1,098	991	1.8	4	4	no
nri_Latn	Chokri Naga	Sino-Tibetan	1,246	410	0.5	79	16	yes
njo_Latn	Ao Naga	Sino-Tibetan	1,282	440	0.7	80	12	yes
zgh_Tfng	Moroccan Tamazight	Afro-Asiatic	1,283	424	0.3	95	13	yes
bci_Latn	Baoule	Atlantic-Congo	1,317	1,031	1.8	53	11	yes
qup_Latn	Southern Pastaza Quechua	Quechuan	1,529	719	0.9	118	19	yes
lad_Latn	Ladino	Indo-European	1,635	204	0.2	73	10	yes
mjw_Latn	Karbi	Sino-Tibetan	1,644	605	0.7	208	15	yes
gos_Latn	Gronings	Indo-European	1,895	1,690	3.1	1,646	19	yes
iba_Latn	Iban	Austronesian	2,942	2,006	5.6	36	8	yes
teo_Latn	Teso	Nilotic	3,195	1,186	1.0	90	9	yes
ekk_Latn	Estonian	Uralic	3,774	1,790	5.2	4	3	yes
lvs_Latn	Latvian	Indo-European	3,847	1,266	3.9	5	4	yes
tgk_Cyrl	Tajik	Indo-European	4,349	3,165	11.2	94	10	yes
hac_Arab	Gurani	Indo-European	4,448	2,272	1.8	1	2	no
oci_Latn	Occitan	Indo-European	4,800	2,696	11.0	15	7	yes
kdj_Latn	Karamojong	Nilotic	5,230	1,737	1.6	94	15	yes
trp_Latn	Kok Borok	Sino-Tibetan	5,384	2,435	3.2	701	19	yes
bre_Latn	Breton	Indo-European	5,790	3,967	2.0	2,822	20	yes
crh_Latn	Crimean Tatar	Turkic	6,784	3,237	4.6	311	15	yes
slv_Latn	Slovenian	Indo-European	9,769	1,384	4.2	6	7	yes
fon_Latn	Fon	Atlantic-Congo	11,201	3,997	11.7	923	7	yes
san_Deva	Sanskrit	Indo-European	12,197	3,664	10.3	80	10	yes
suk_Latn	Sukuma	Atlantic-Congo	12,330	3,453	9.9	36	9	yes
hat_Latn	Haitian Creole	Indo-European	12,988	3,994	10.8	3,994	25	yes
lin_Latn	Lingala	Atlantic-Congo	13,887	4,000	16.6	8	11	yes
crh_Cyrl	Crimean Tatar	Turkic	14,219	3,930	4.8	1	1	no
kha_Latn	Khasi	Austroasiatic	14,234	1,865	4.9	21	8	yes
smj_Latn	Lule Sami	Uralic	14,715	3,999	4.6	503	11	yes
sna_Latn	Shona	Atlantic-Congo	15,476	3,946	19.9	27	15	yes
kbd_Cyrl	Kabardian	Abkhaz-Adyge	17,393	3,348	5.3	124	10	yes
sdh_Arab	Southern Kurdish	Indo-European	17,998	3,999	6.6	71	9	yes
fin_Latn	Finnish	Uralic	20,414	4,000	6.8	45	1	yes
plt_Latn	Malagasy	Austronesian	21,705	3,997	11.7	97	12	yes
lat_Latn	Latin	Indo-European	22,862	3,916	11.1	190	1	yes
mya_Mymr	Burmese	Sino-Tibetan	24,226	3,998	6.2	275	10	yes
chv_Cyrl	Chuvash	Turkic	29,796	3,998	4.5	508	14	yes
guj_Gujr	Gujarati	Indo-European	32,373	4,000	6.9	220	11	yes
fry_Latn	West Frisian	Indo-European	35,956	4,000	5.2	2,526	8	yes
lao_Laoo	Lao	Tai-Kadai	38,586	3,986	13.2	124	9	yes
fuv_Latn	Nigerian Fulfulde	Atlantic-Congo	39,284	3,417	11.2	823	12	yes
vag_Latn	Vagla	Atlantic-Congo	39,303	3,989	2.3	2,521	12	yes
bod_Tibt	Tibetan	Sino-Tibetan	48,956	3,994	2.5	3,994	12	yes
gsw_Latn	Swiss German	Indo-European	51,710	3,000	5.7	3	1	yes
tir_Ethi	Tigrinya	Afro-Asiatic	53,835	3,992	16.7	583	12	yes
xsm_Latn	Kasem	Atlantic-Congo	58,196	3,832	2.4	2,883	13	yes

C Learning rates

Table 13 has the sweep that chose each model's rate.

The rows a transducer adds have a rate of their own, since they start with nothing to keep. Table 14 shows it hardly matters: dev moves by a median of 0.9 points across three rates for the rows at a given rate for the rest of the model, against more than ten points between the

rest of the model's rates in Table 13. The main runs give the rows ten times the rate of the rest of the model, and at least 1e-3.

TABLE 13 The learning-rate sweep: dev macro CER on 200 clips a language after three epochs. Bold is the rate the main run used. A sweep this short favours a fast rate, so where the lowest score was not clearly lowest, or the rate was the fastest tried, a full-length run at a lower rate decided on dev; Qwen3-ASR 0.6B, which every ablation starts from, kept the faster of its two tied rates, and Nemotron streaming's lower-rate run slowed only its new rows (Table 14). Where the lowest score sat at an end of the grid, the sweep tried one rate past that end. Each did worse, except on Parakeet, where 3e-4 beat 1e-4 and 1e-3 did not, so 3e-4 also ran the full length, and dev chose it. Over the full length, 3e-4 silences Nemotron streaming (Section 7.4). Gemma 4 E4B and 12B also ran the next rate down for the full length, and dev chose 3e-6 for E4B (24.7 against 26.5 at 1e-6) and 3e-6 for the 12B (28.1 against 28.4 at 1e-6). w2v-BERT at 1e-4 stayed at 100.0, writing the same output for every clip, when run a second time too. omniASR CTC 7B takes the same 120 seconds of audio per update as the rest, in pieces of 40 seconds in the sweep and 20 in its full run, because larger pieces did not fit. Dashes are rates not run.

MODEL	1E-6	3E-6	1E-5	3E-5	1E-4	3E-4	1E-3
Cohere Transcribe 2B	–	49.9	42.0	**37.0**	38.3	–	–
Gemma 4 12B	36.3	**35.2**	41.4	100.0	–	–	–
Gemma 4 E2B	–	32.2	**31.5**	35.1	45.9	–	–
Gemma 4 E4B	30.5	**30.2**	31.2	36.4	–	–	–
Granite Speech 5.0 2B	–	43.3	**40.4**	42.6	–	–	–
LFM3-Audio 2B	–	33.0	**32.0**	33.0	41.5	–	–
MMS 1B	–	24.9	**23.5**	23.8	29.0	–	–
Nemotron 3.5 streaming 0.6B	–	80.3	83.1	73.1	**69.0**	–	–
omniASR CTC 1B	–	–	30.7	**24.4**	29.2	–	–
omniASR CTC 300M	–	–	36.5	**28.1**	31.9	49.6	–
omniASR CTC 7B	–	29.0	**20.8**	24.5	34.3	–	–
Parakeet TDT 0.6B v3	–	58.9	53.2	50.5	47.0	**43.6**	47.8
Qwen3-ASR 0.6B	39.2	34.8	**34.8**	36.7	43.2	–	–
Qwen3-ASR 1.7B	36.2	30.3	**31.4**	32.2	–	–	–
Voxtral Mini 3B	30.8	**29.5**	32.9	40.1	–	–	–
w2v-BERT 2.0	–	–	60.3	**40.8**	100.0	–	–
Whisper base	–	–	–	43.4	**39.3**	49.1	–
Whisper large-v3-turbo	–	34.9	**32.1**	35.8	–	–	–
Whisper large-v3	–	30.8	**29.3**	29.6	–	–	–
Whisper medium	–	–	**33.0**	32.7	37.8	–	–
Whisper small	–	–	40.3	**35.0**	37.8	–	–
Whisper tiny	–	–	53.1	48.3	**47.1**	54.6	–

TABLE 14 The transducers' new rows, on the sweep of Table 13: dev macro CER, by the rate of the pretrained weights (Rate) and of the new rows.

MODEL	RATE	ROWS 1E-4	ROWS 3E-4	ROWS 1E-3
Nemotron 3.5 streaming 0.6B	1e-5	83.2	83.3	83.1
Nemotron 3.5 streaming 0.6B	3e-5	76.7	75.8	73.1
Nemotron 3.5 streaming 0.6B	1e-4	71.3	72.1	69.0
Parakeet TDT 0.6B v3	3e-5	50.3	51.2	50.5
Parakeet TDT 0.6B v3	1e-4	46.9	47.6	47.0